

TxProbe: Discovering Bitcoin's Network Topology Using Orphan Transactions

Sergi Delgado-Segura, Surya Bakshi, Cristina Pérez-Solà, James Litton, Andrew Pachulski, Andrew Miller and Bobby Bhattacharjee

MOTIVATION



What is the big deal with knowing the network topology?

- Is the network really decentralized?
- Are there supernodes in the network?
- Are there weak spots that can be easily isolated?

Currently, we do not know

THE TOPOLOGY SHOULD LOOK RANDOM



How Bitcoin (Core client) nodes choose their peers?

- Pseudorandomly from the **addrman**
- **8 outbound** connections by default
 - No pair of nodes in the same **/16** (IPv4)
- **117 inbound** connection by default (no IP restriction here)

Bitcoin forks based on the Core client follow the same approach

REACHABLE AND NON-REACHABLE NODES

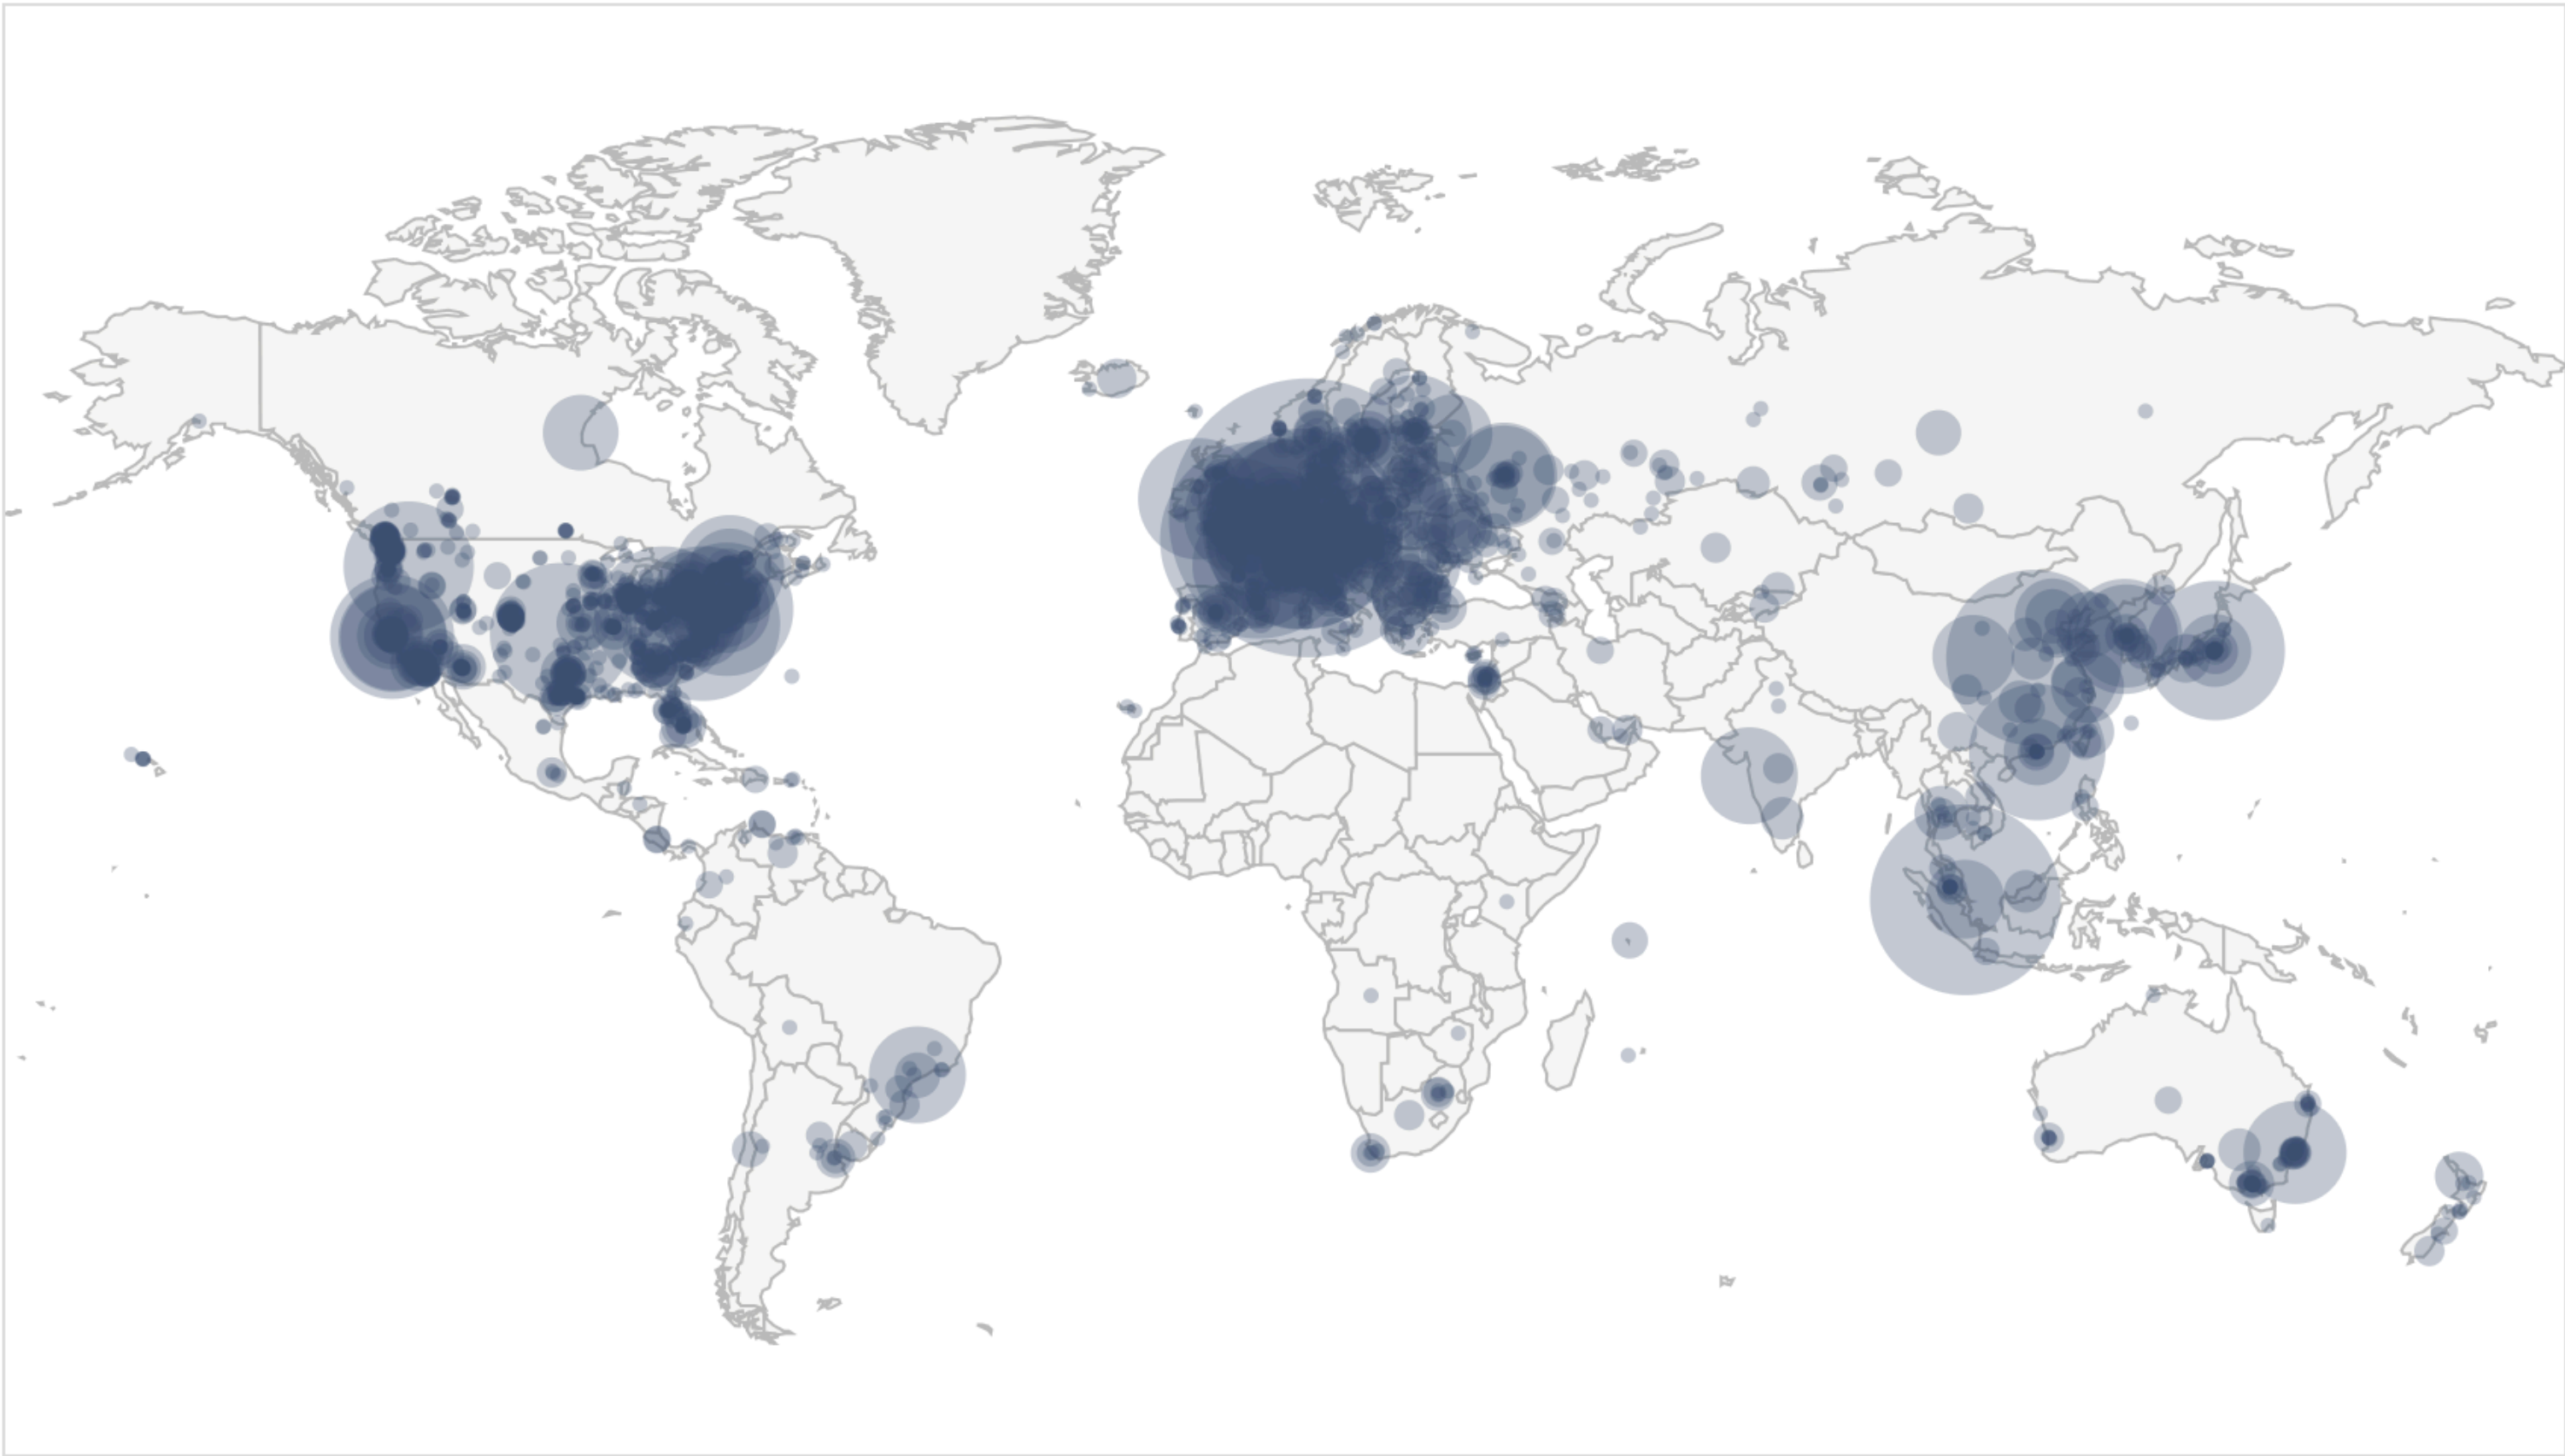
GLOBAL BITCOIN NODES DISTRIBUTION
Reachable nodes as of Thu Feb 07 2019
10:26:44 GMT+0000 (Greenwich Mean Time).

10365 NODES
[24-hour charts »](#)

Top 10 countries with their respective number of reachable nodes are as follow.

| RANK | COUNTRY | NODES |
|------|--------------------|---------------|
| 1 | United States | 2570 (24.79%) |
| 2 | Germany | 1968 (18.99%) |
| 3 | France | 689 (6.65%) |
| 4 | Netherlands | 514 (4.96%) |
| 5 | China | 411 (3.97%) |
| 6 | Canada | 384 (3.70%) |
| 7 | United Kingdom | 355 (3.42%) |
| 8 | Singapore | 321 (3.10%) |
| 9 | Russian Federation | 277 (2.67%) |
| 10 | Japan | 228 (2.20%) |

[More \(100\) »](#)



Map shows concentration of reachable Bitcoin nodes found in countries around the world.

[LIVE MAP](#)

Source: <https://bitnodes.earn.com/>

REACHABLE AND NON-REACHABLE NODES



**GLOBAL BITCOIN NODES
DISTRIBUTION**
Reachable nodes as of Thu Feb 07 2019
10:26:44 GMT+0000 (Greenwich Mean Time).

10365 NODES

241 countries

Difference between reachable and non-reachable?

| | | |
|--------------|-------|-------------|
| 10 | Japan | 228 (2.20%) |
| More (100) » | | |



Map shows concentration of reachable Bitcoin nodes found in countries around the world.

LIVE MAP

Source: <https://bitnodes.earn.com/>

REACHABLE AND NON-REACHABLE NODES

GLOBAL BITCOIN NODES DISTRIBUTION

Reachable nodes as of Thu Feb 07 2019
10:26:44 GMT+0000 (Greenwich Mean Time).

10365 NODES

Difference between reachable and non-reachable?

Reachable nodes accept incoming connections, non-reachable do not

| | | |
|----|-------|-------------|
| 10 | Japan | 228 (2.20%) |
|----|-------|-------------|

[More \(100\) »](#)

Map shows concentration of reachable Bitcoin nodes found in countries around the world.

LIVE MAP

Source: <https://bitnodes.earn.com/>

REACHABLE AND NON-REACHABLE NODES

GLOBAL BITCOIN NODES DISTRIBUTION

Reachable nodes as of Thu Feb 07 2019
10:26:44 GMT+0000 (Greenwich Mean Time).

10365 NODES

Difference between reachable and non-reachable?

Reachable nodes accept incoming connections, non-reachable do not

How many non-reachable nodes are out there?

10 Japan 228 (2.20%)

[More \(100\) »](#)

Map shows concentration of reachable Bitcoin nodes found in countries around the world.

LIVE MAP

Source: <https://bitnodes.earn.com/>

REACHABLE AND NON-REACHABLE NODES

GLOBAL BITCOIN NODES DISTRIBUTION

Reachable nodes as of Thu Feb 07 2019
10:26:44 GMT+0000 (Greenwich Mean Time).

10365 NODES

Difference between reachable and non-reachable?

Reachable nodes accept incoming connections, non-reachable do not

How many non-reachable nodes are out there?

There is not good approximation (AFAIK)

10 Japan 228 (2.20%)

[More \(100\) »](#)

Map shows concentration of reachable Bitcoin nodes found in countries around the world.

LIVE MAP

Source: <https://bitnodes.earn.com/>

REACHABLE AND NON-REACHABLE NODES

GLOBAL BITCOIN NODES DISTRIBUTION

Reachable nodes as of Thu Feb 07 2019
10:26:44 GMT+0000 (Greenwich Mean Time).

10365 NODES

Difference between reachable and non-reachable?

Reachable nodes accept incoming connections, non-reachable do not

How many non-reachable nodes are out there?

There is not good approximation (AFAIK)

Probing techniques generally focus on the reachable network

10 Japan 228 (2.20%)

[More \(100\) »](#)

Map shows concentration of reachable Bitcoin nodes found in countries around the world.

LIVE MAP

Source: <https://bitnodes.earn.com/>

TRANSACTION PROPAGATION IN BITCOIN



Our inferring technique is based on transaction propagation



We take advantage of how some kind of transactions (**orphans** and **double-spending**) are handled by nodes

How does it roughly work?

TRANSACTION PROPAGATION IN BITCOIN



Transactions are shared between peers in a push manner




TRANSACTION PROPAGATION IN BITCOIN



Transactions are shared between peers in a push manner

When a peer receives / generates a new transaction he announce it to his neighbors (**announce**)



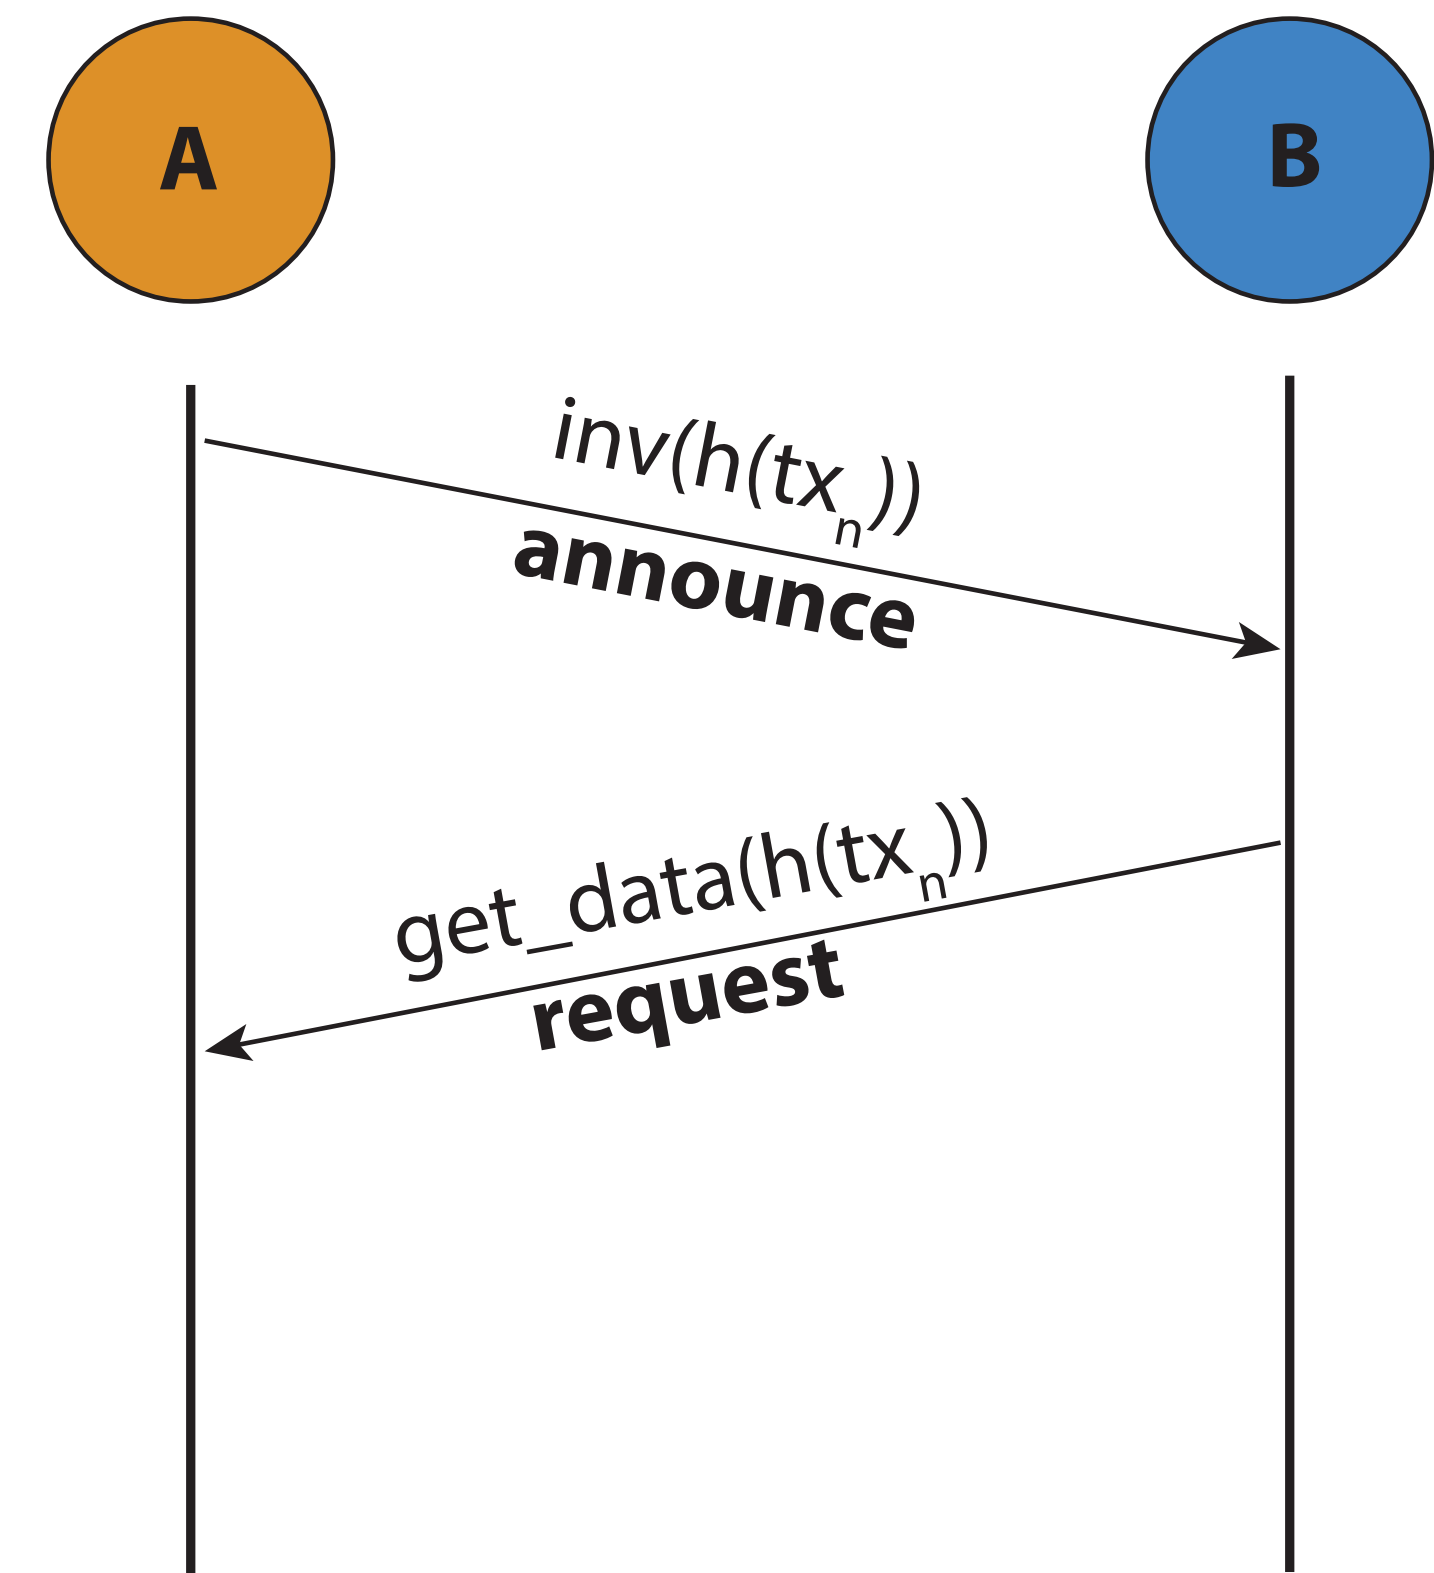
TRANSACTION PROPAGATION IN BITCOIN



Transactions are shared between peers in a push manner

When a peer receives / generates a new transaction he announce it to his neighbors (**announce**)

Upon receiving an announce of an item, a node that does not know about it will request the item back to the announcer (**request**)

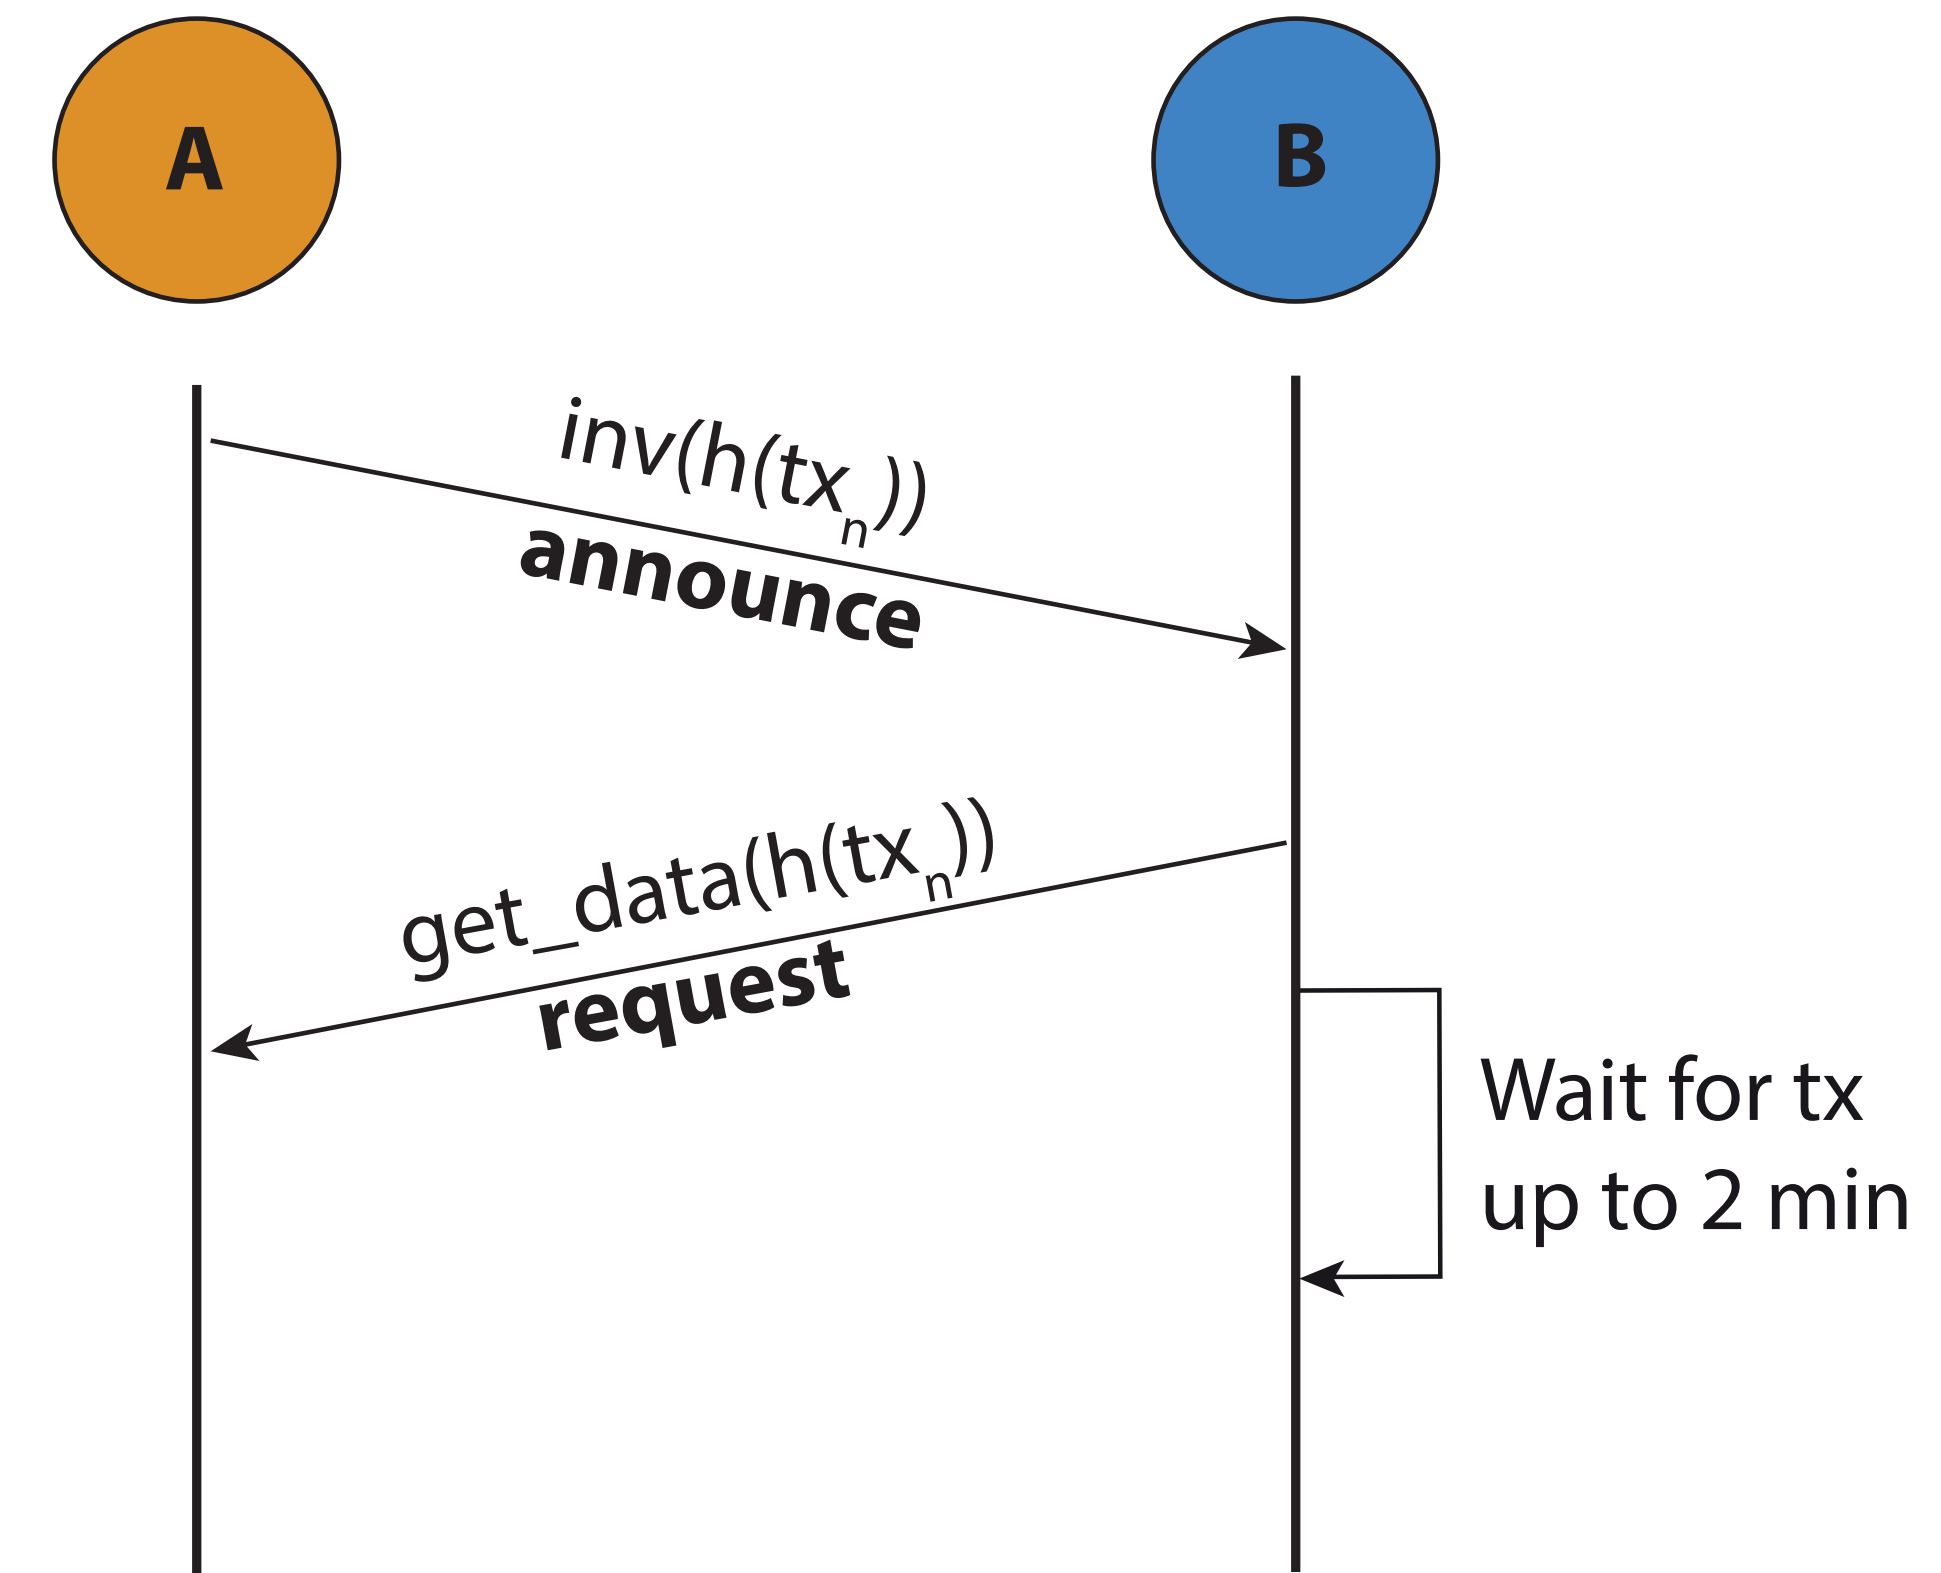


TRANSACTION PROPAGATION IN BITCOIN

Transactions are shared between peers in a push manner

When a peer receives / generates a new transaction he announce it to his neighbors (**announce**)

Upon receiving an announce of an item, a node that does not know about it will request the item back to the announcer (**request**)



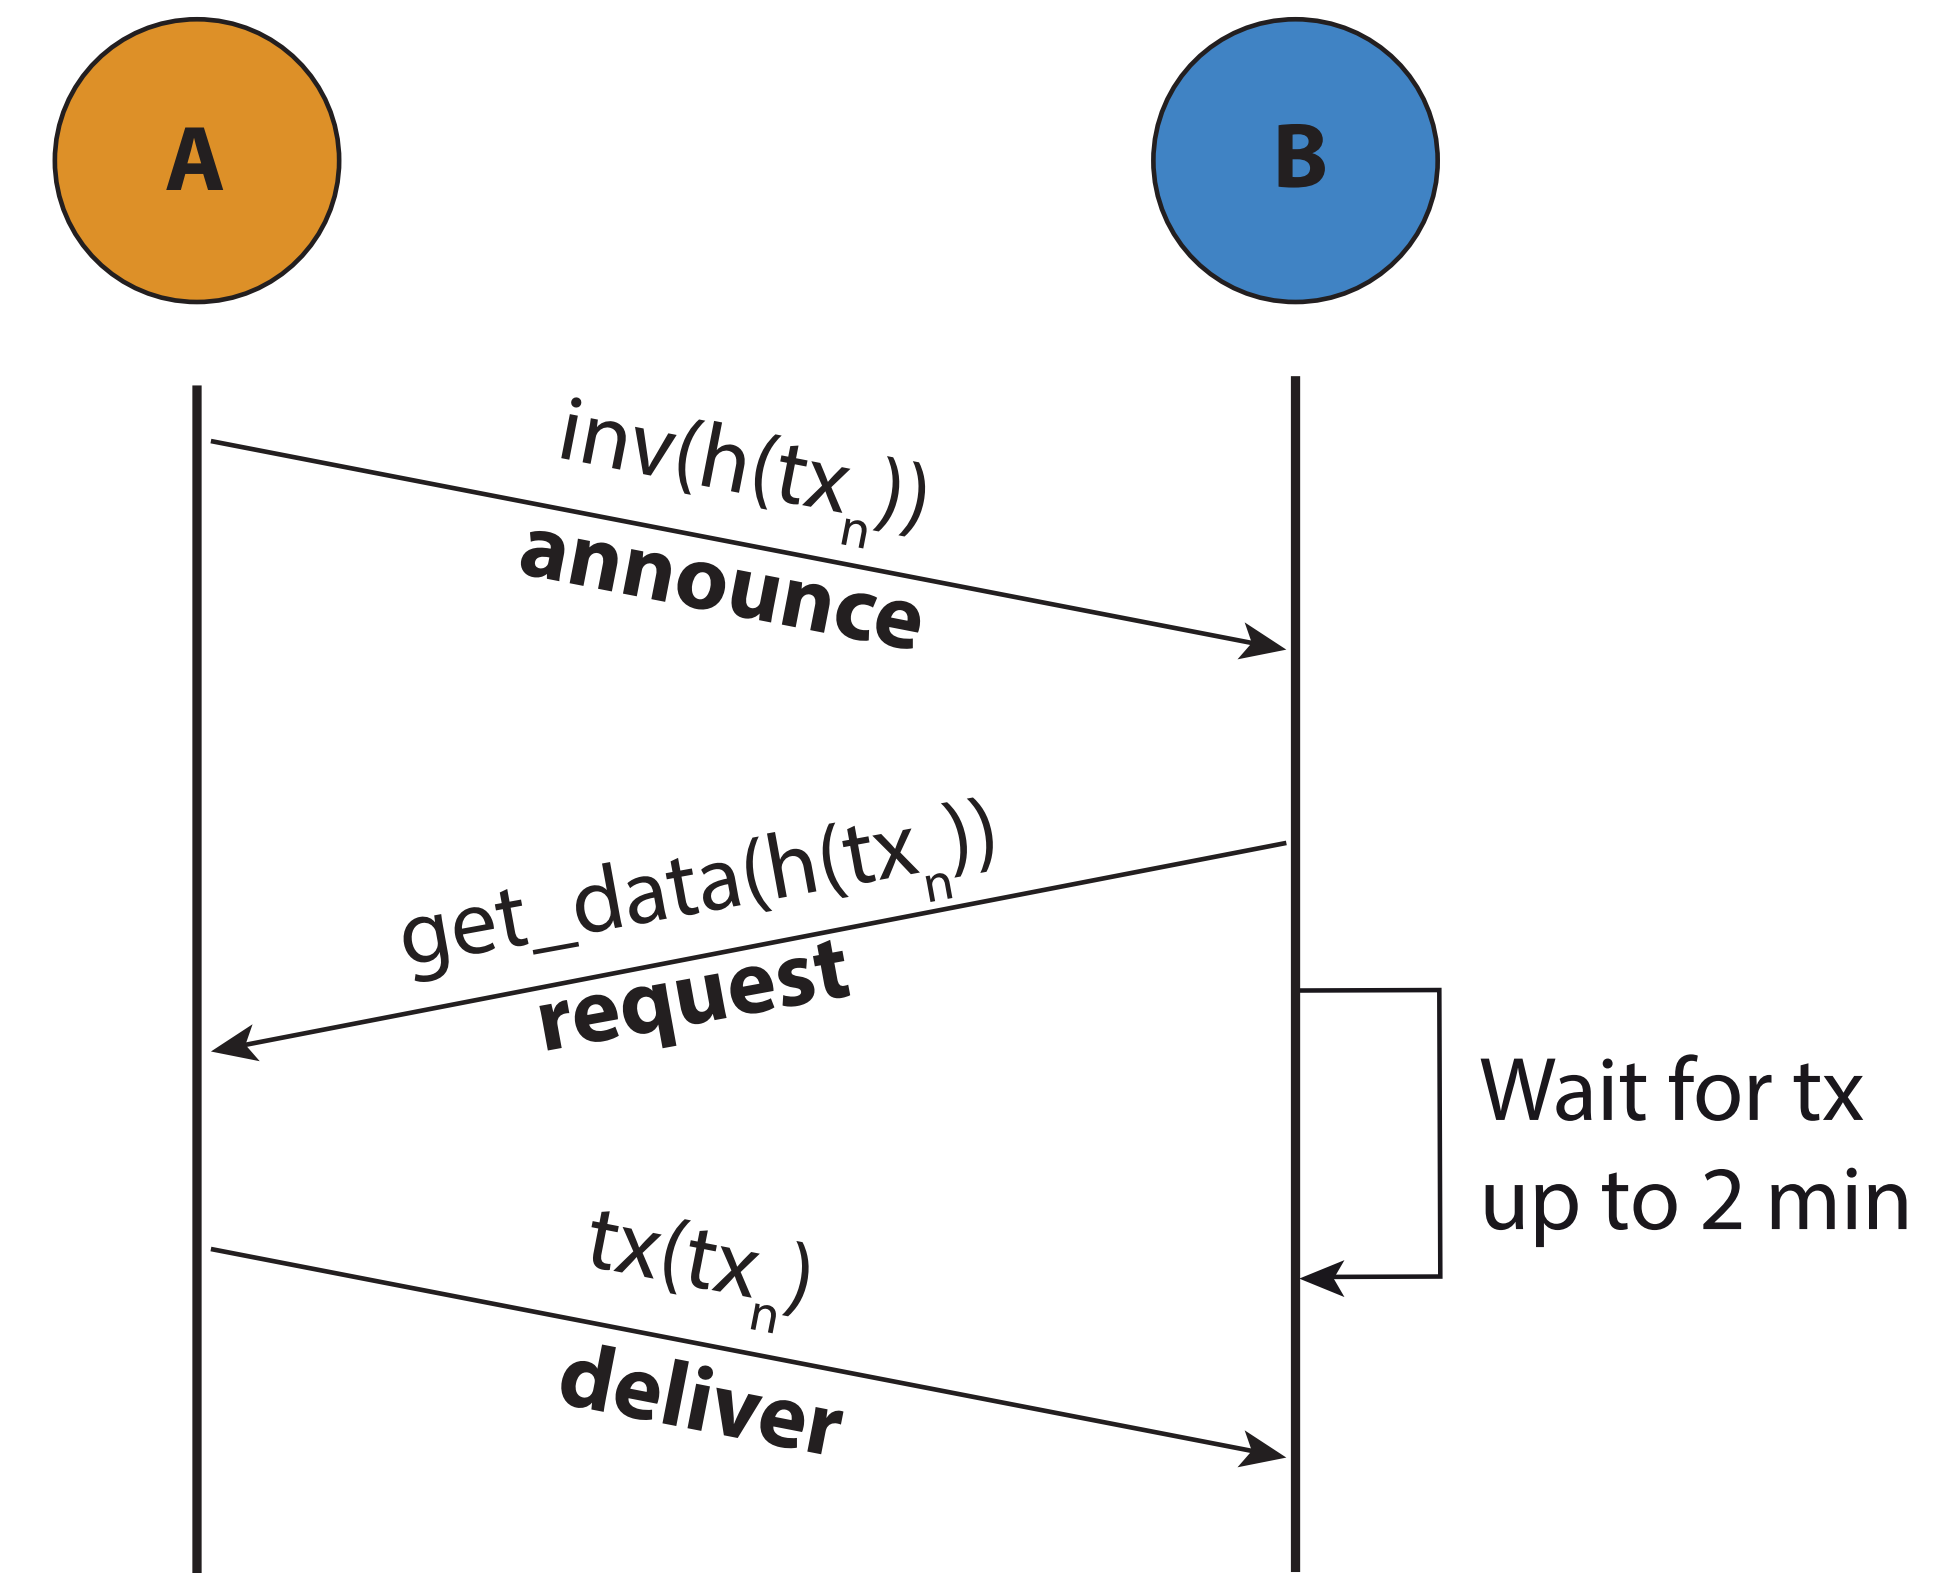
TRANSACTION PROPAGATION IN BITCOIN

Transactions are shared between peers in a push manner

When a peer receives / generates a new transaction he announce it to his neighbors (**announce**)

Upon receiving an announce of an item, a node that does not know about it will request the item back to the announcer (**request**)



Upon receiving a request of a known item, a node will reply back with it (**deliver**)



TRANSACTION PROPAGATION IN BITCOIN



Transactions are shared between peers in a push manner



When
he a

Valid transaction are stored in **mempool**

Upon
does
the



Upon receiving a request for a transaction, a node
will reply back with it (**deliver**)

Wait for tx
up to 2 min

TRANSACTION PROPAGATION IN BITCOIN



Transactions are shared between peers in a push manner



When
he a

Valid transaction are stored in **mempool**

Transaction in **mempool** are eventually **propagated throughout the node neighborhood**



Upon
does
the

Upon receiving a request for a transaction, a node
will reply back with it (**deliver**)

Wait for tx
up to 2 min

TRANSACTION PROPAGATION IN BITCOIN

Transactions are shared between peers in a push manner



When
he a

Valid transaction are stored in **mempool**

Transaction in **mempool** are eventually **propagated throughout the node neighborhood**

Upon
does
the

A node will wait **up to 2 minutes** for a response to a **getaddr message** before asking any other peer

Upon receiving a response or a timeout, a node will reply back with it (**deliver**)

Wait for tx
up to 2 min

ORPHAN TRANSACTIONS



ORPHAN TRANSACTIONS



A transaction is flagged as orphan when a nodes receiving it does not know about some of it's parents (**some of the UTXOs are unknown**)

ORPHAN TRANSACTIONS



A transaction is flagged as orphan when a nodes receiving it does not know about some of it's parents (**some of the UTXOs are unknown**)

Orphan transactions can not be validated

ORPHAN TRANSACTIONS



A transaction is flagged as orphan when a nodes receiving it does not know about some of it's parents (**some of the UTXOs are unknown**)

Orphan transactions can not be validated


They are stored in a separated data structure known as **MapOrphanTransactions**

ORPHAN TRANSACTIONS

A transaction is flagged as orphan when a nodes receiving it does not know about some of it's parents (**some of the UTXOs are unknown**)

Orphan transactions can not be validated

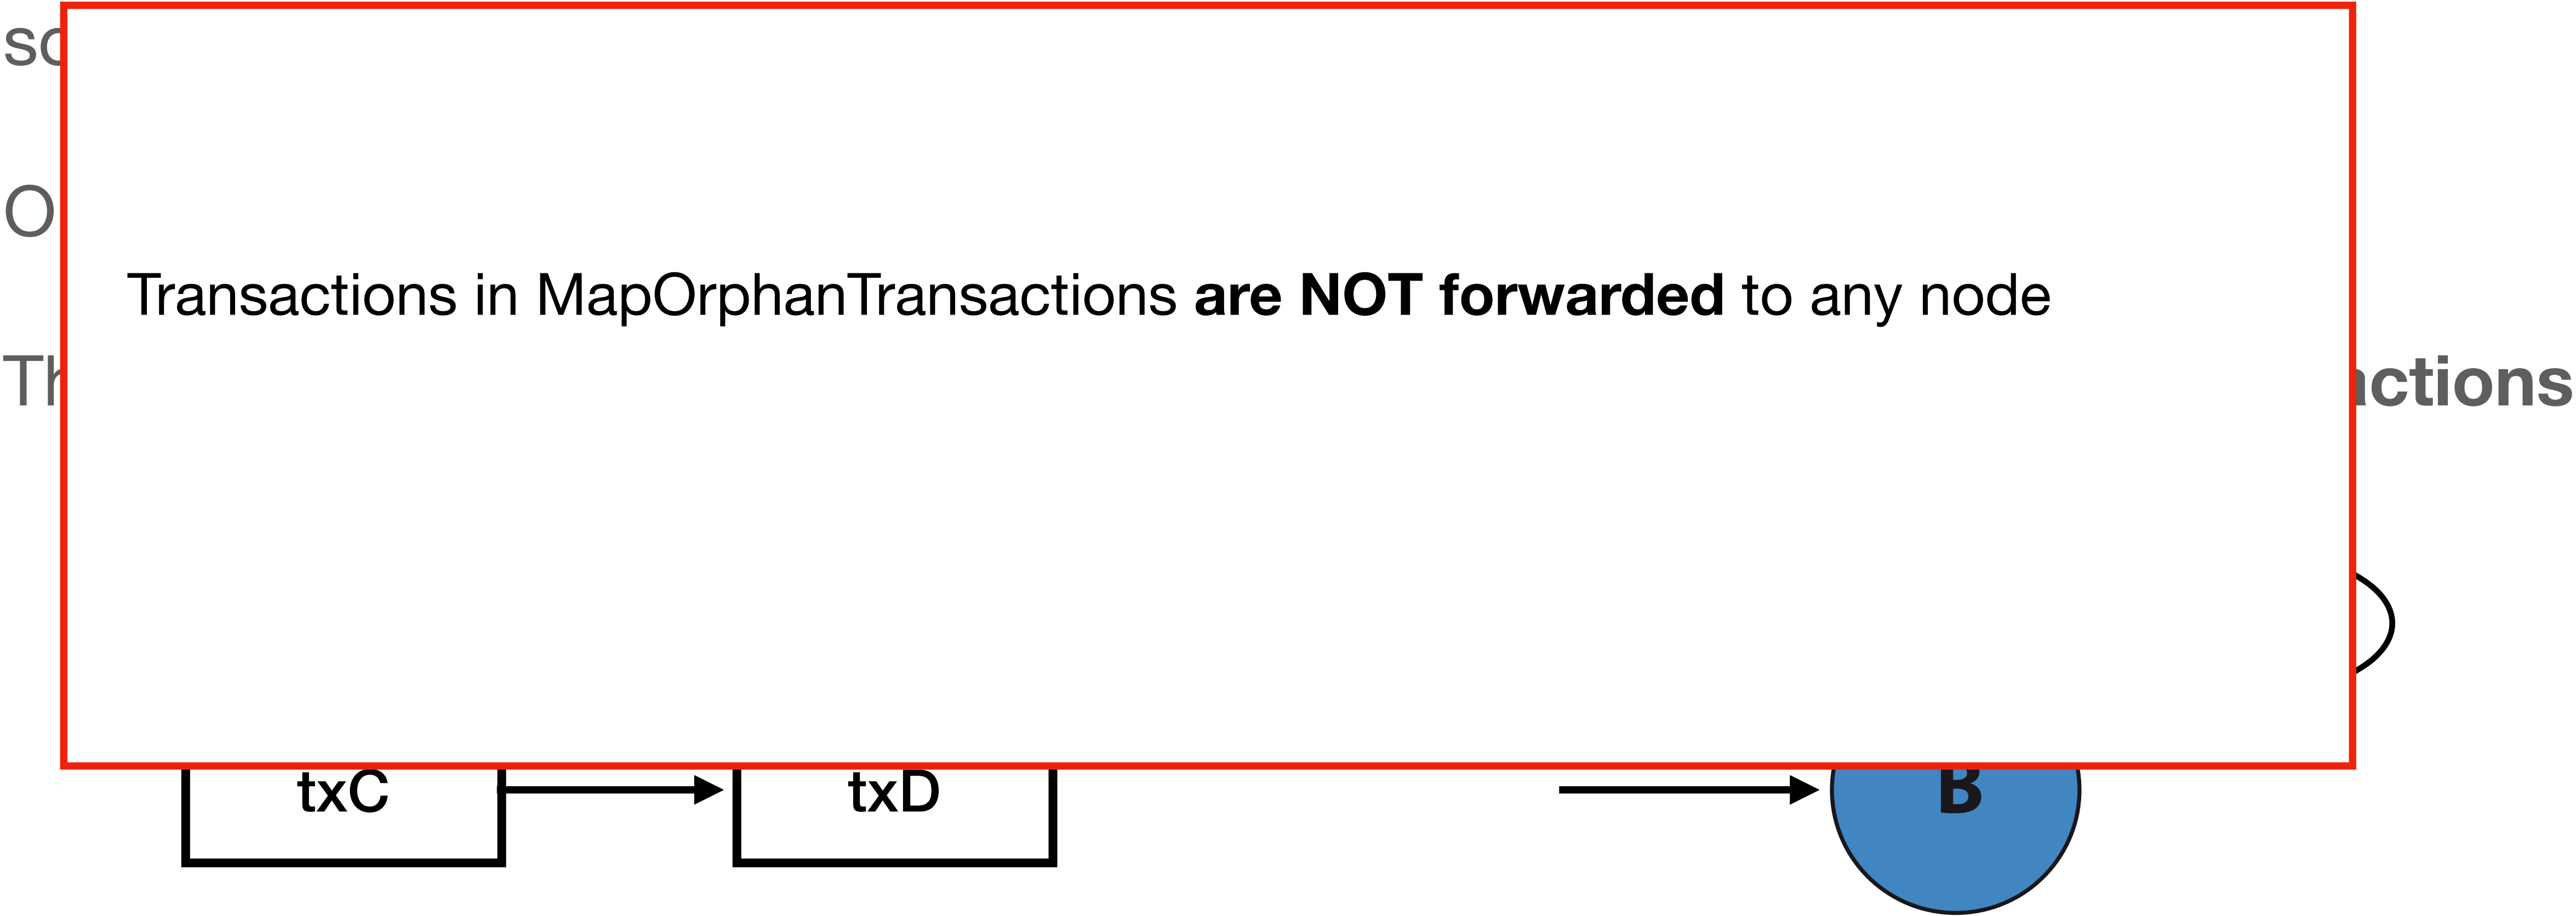
They are stored in a separated data structure known as **MapOrphanTransactions**



ORPHAN TRANSACTIONS



A transaction is flagged as orphan when a nodes receiving it does not know about




txD is orphan from
B's point of view

ORPHAN TRANSACTIONS

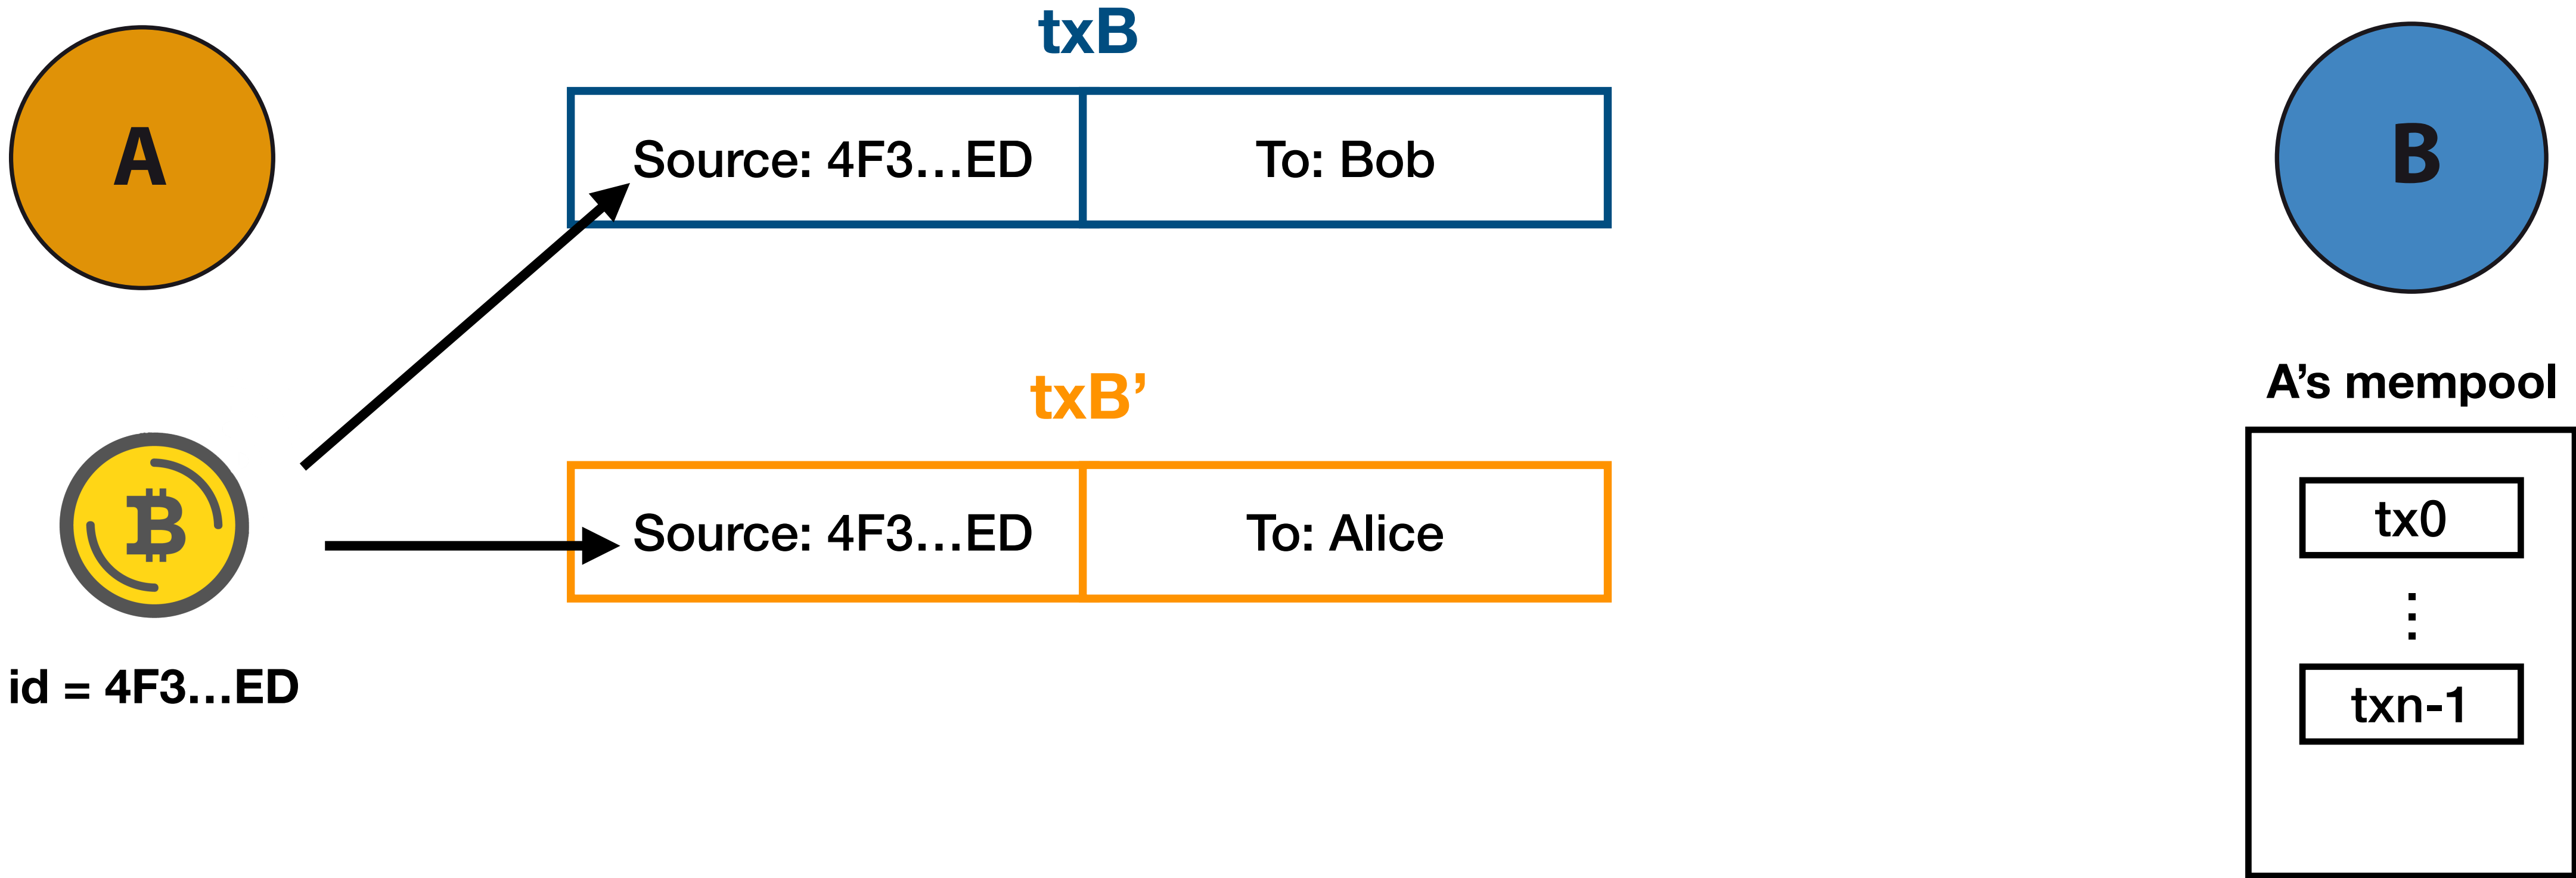
A transaction is flagged as orphan when a node receiving it does not know about

Transactions in MapOrphanTransactions **are NOT forwarded** to any node
If the same transactions is offered again to the node (**inv message**), it will not ask back for it (**getaddr**)

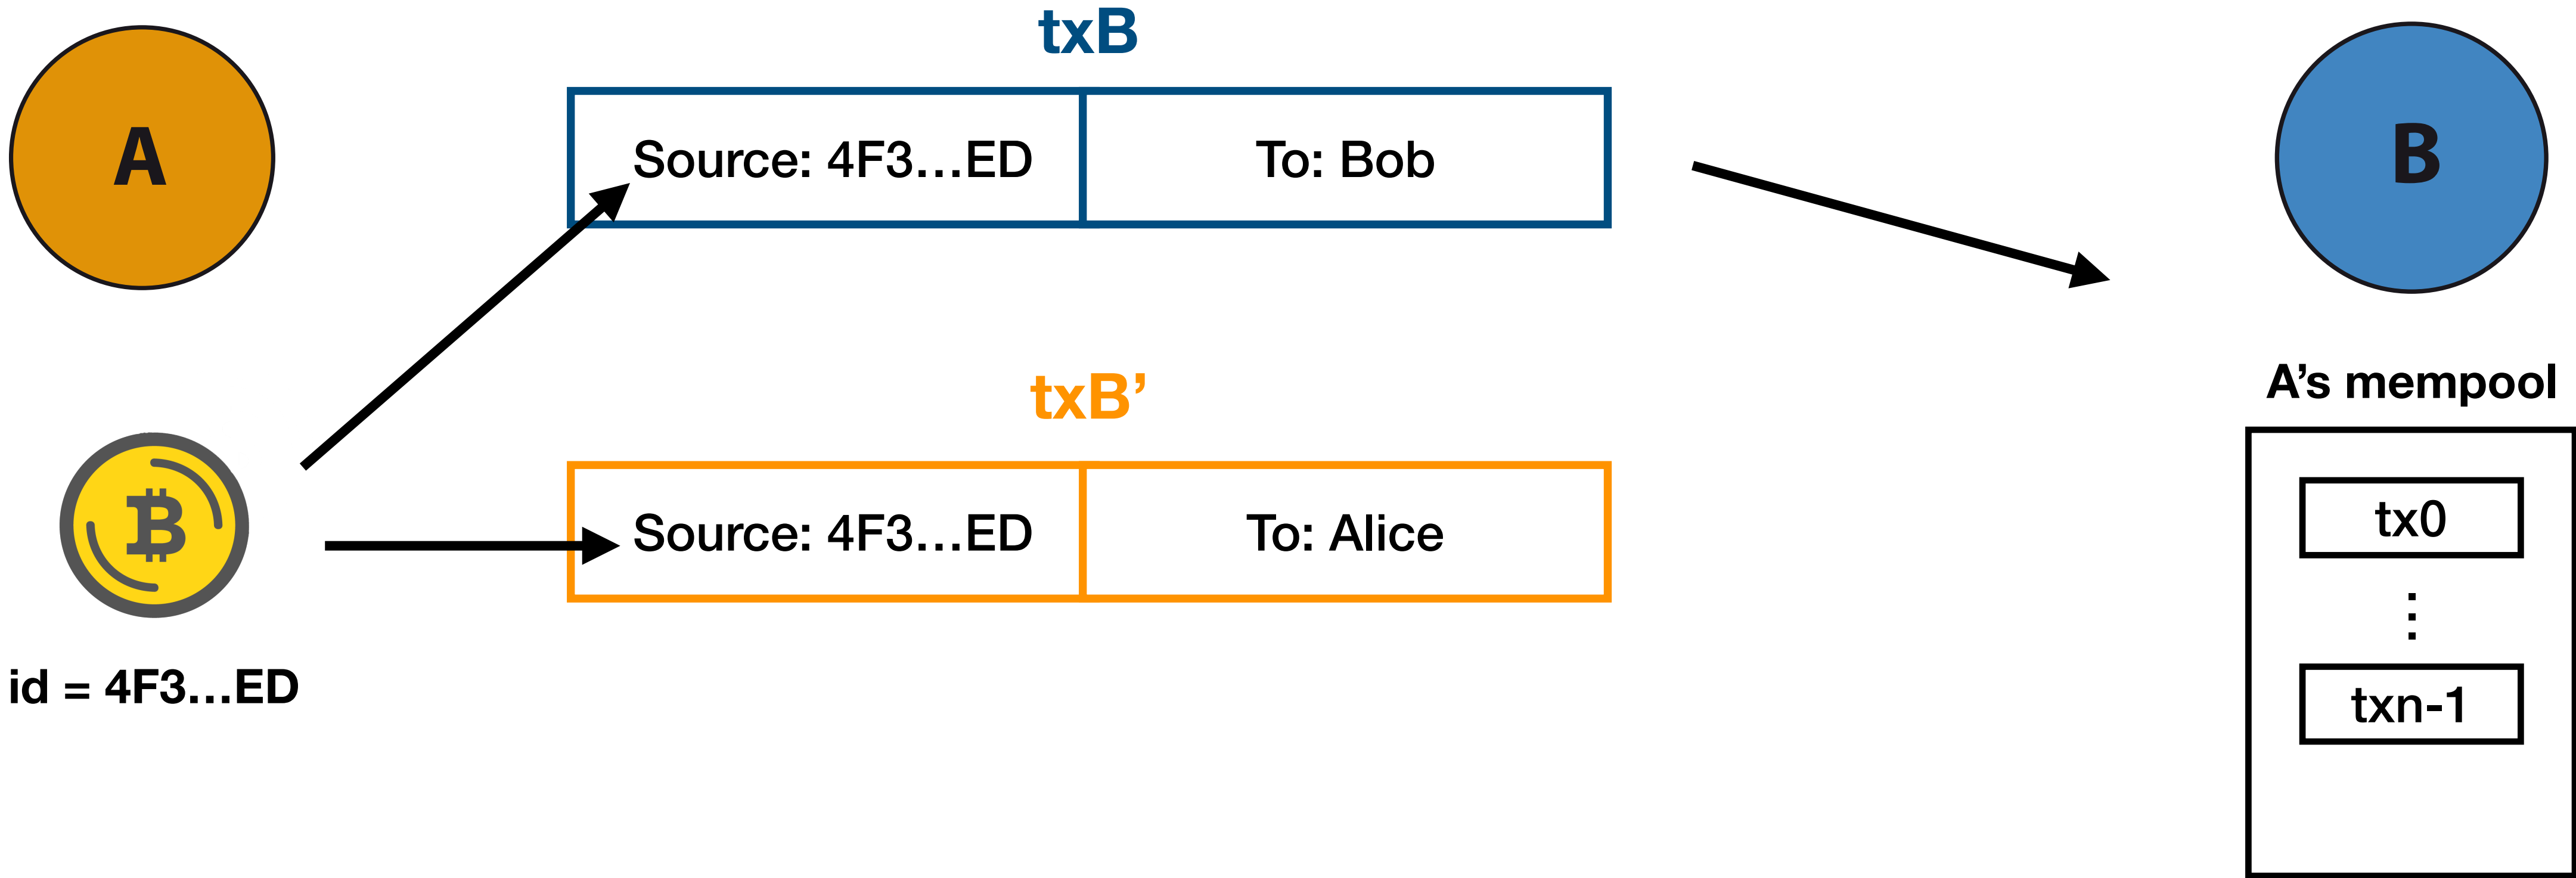


**txD is orphan from
B's point of view**

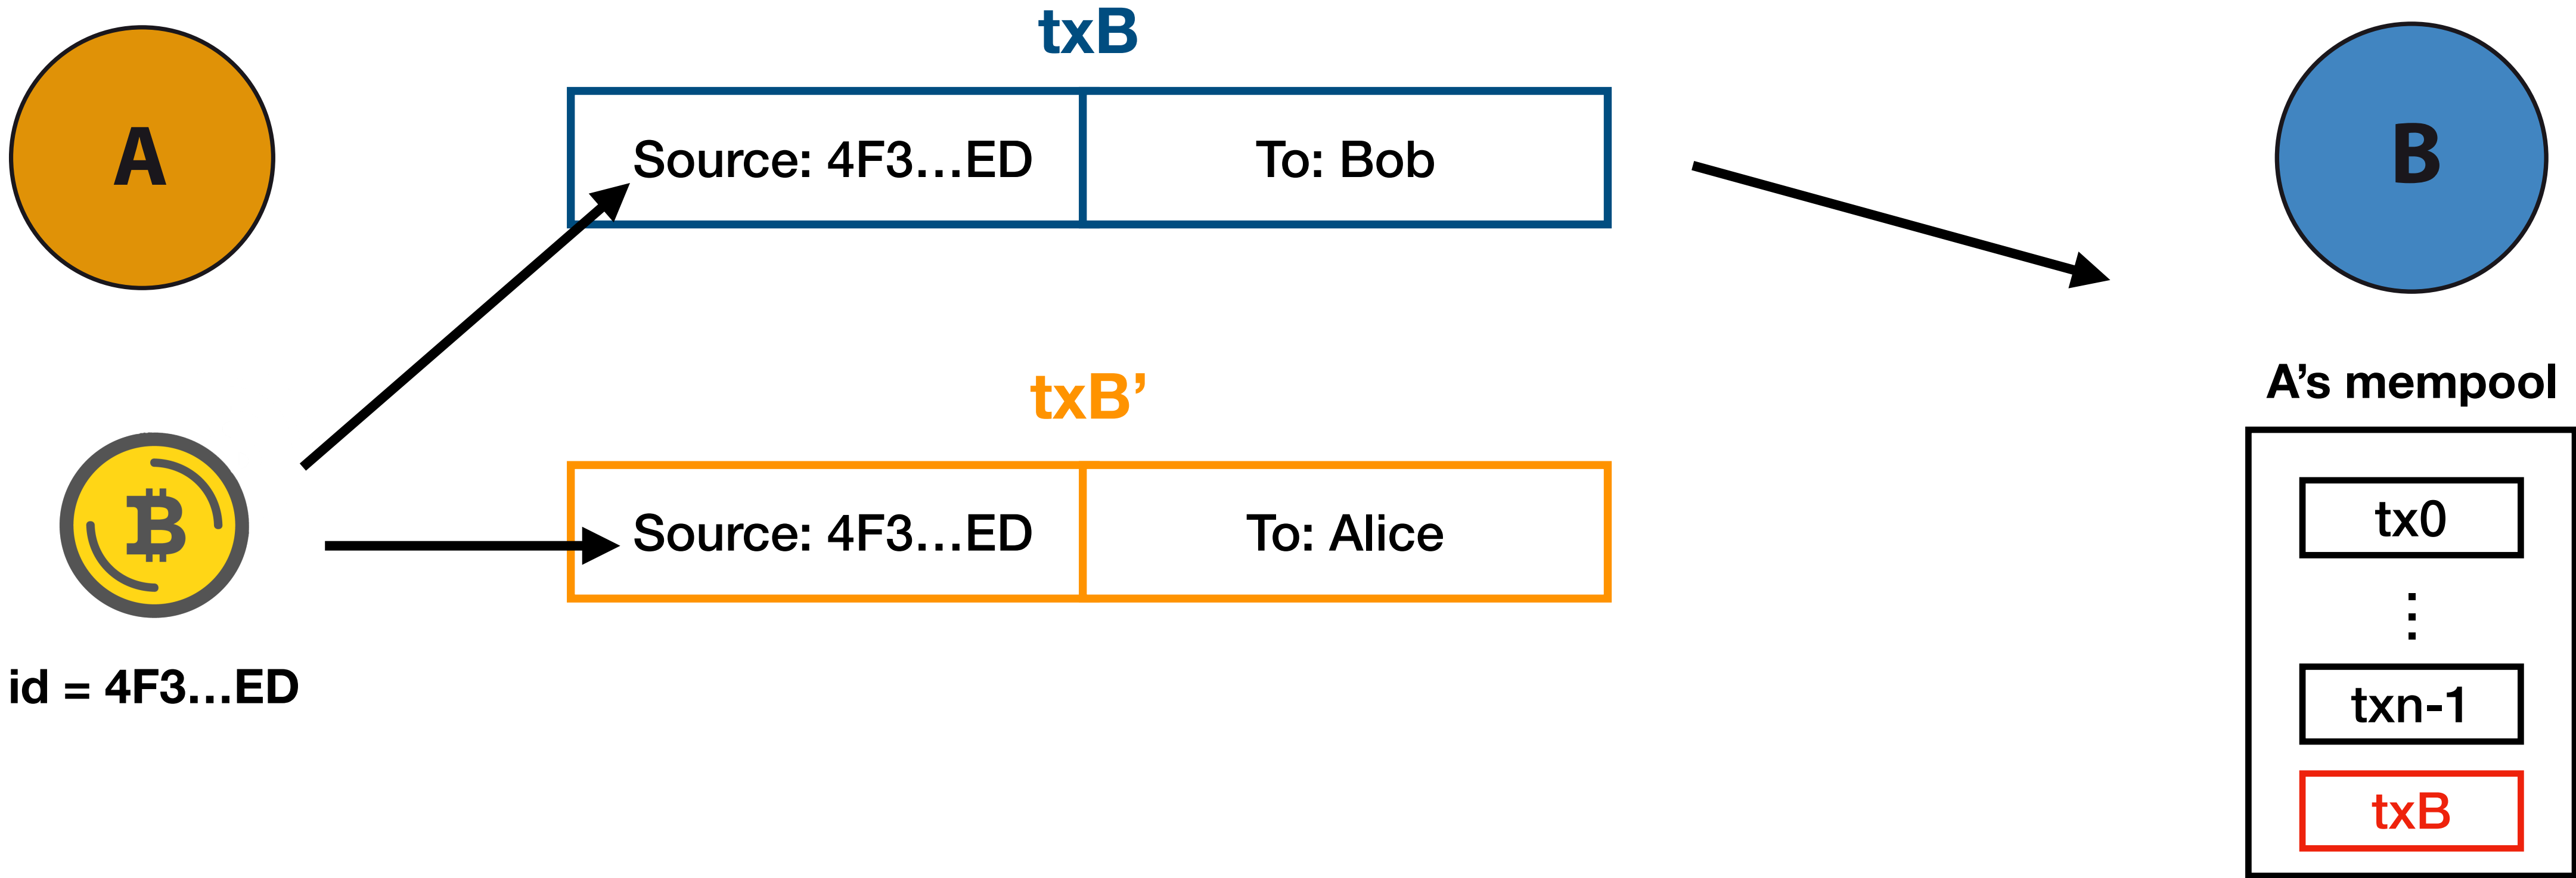
DOUBLE-SPENDING TRANSACTIONS



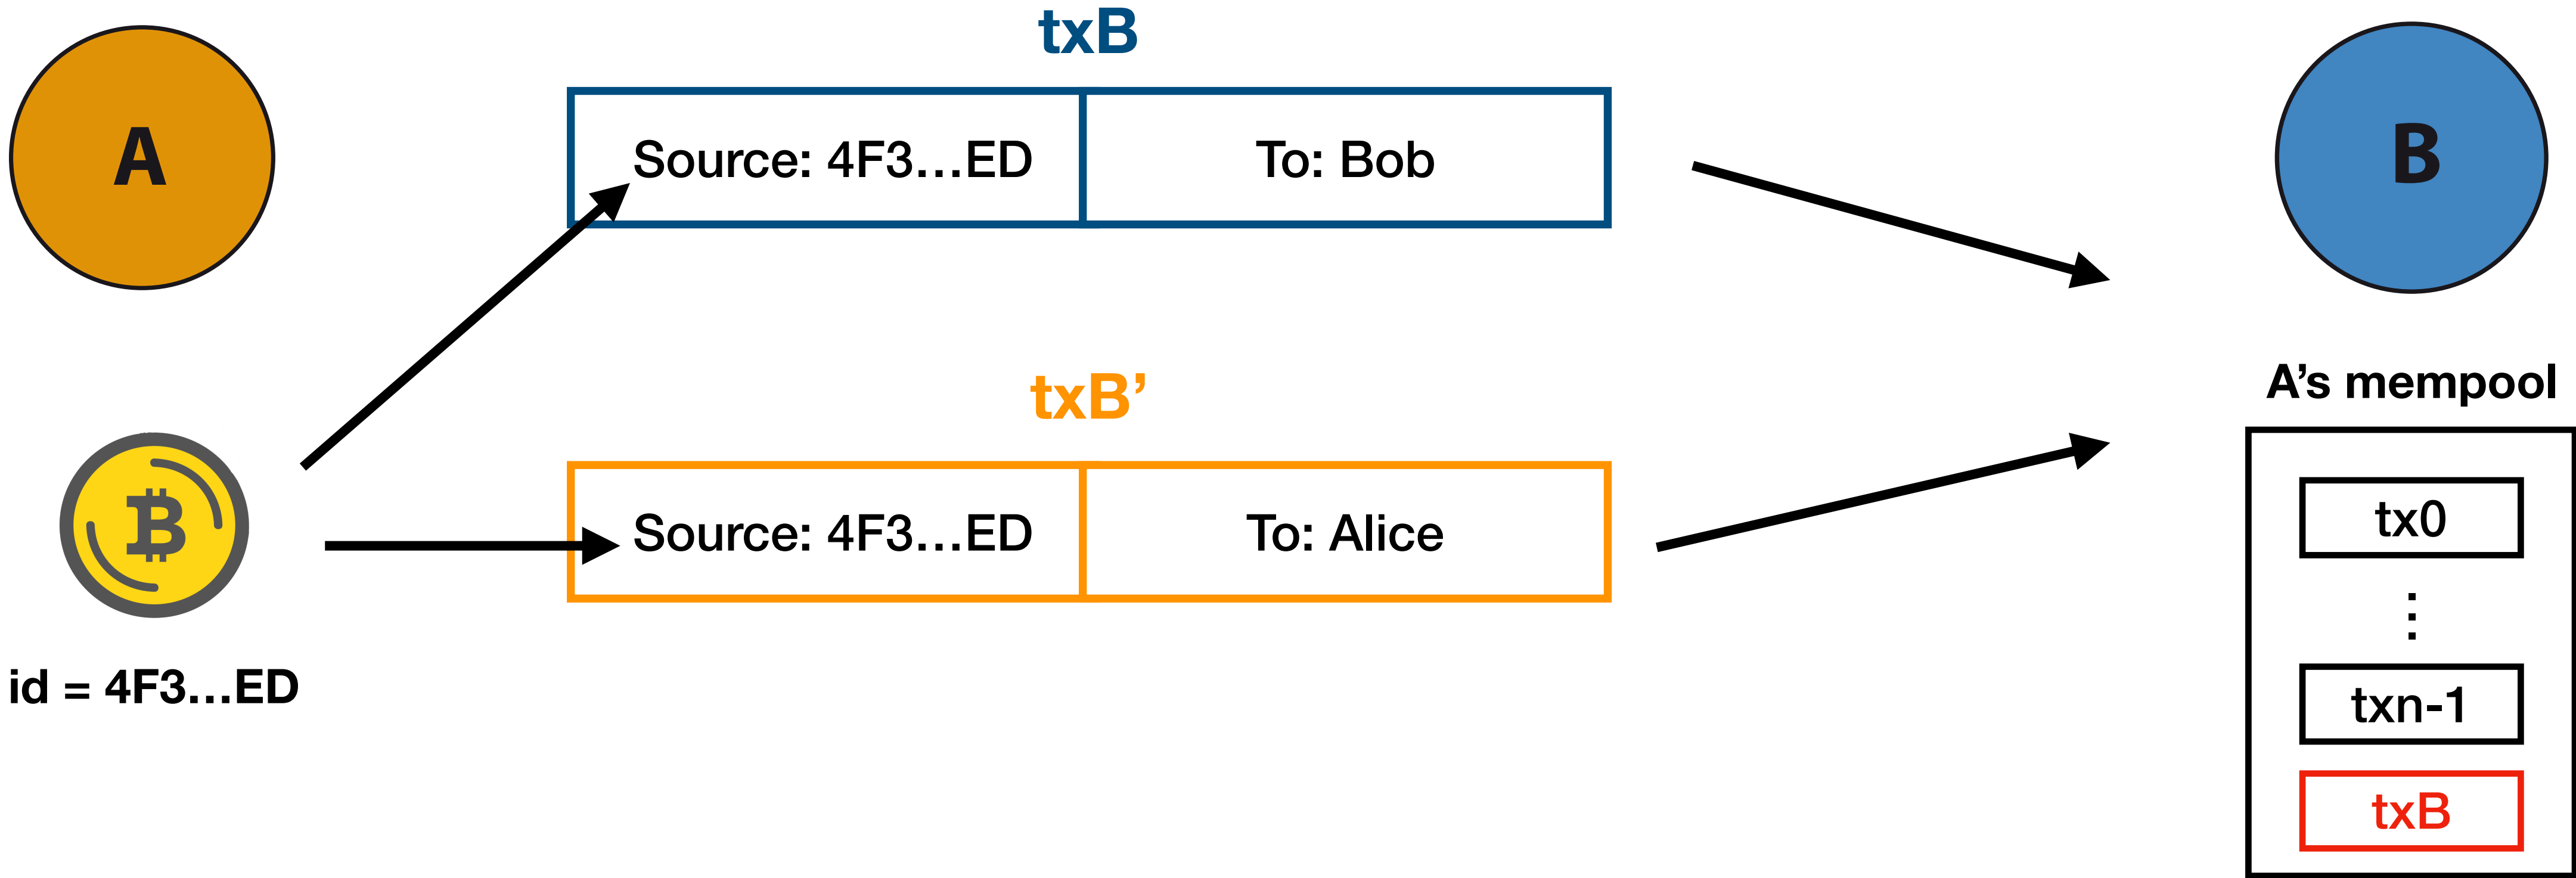
DOUBLE-SPENDING TRANSACTIONS



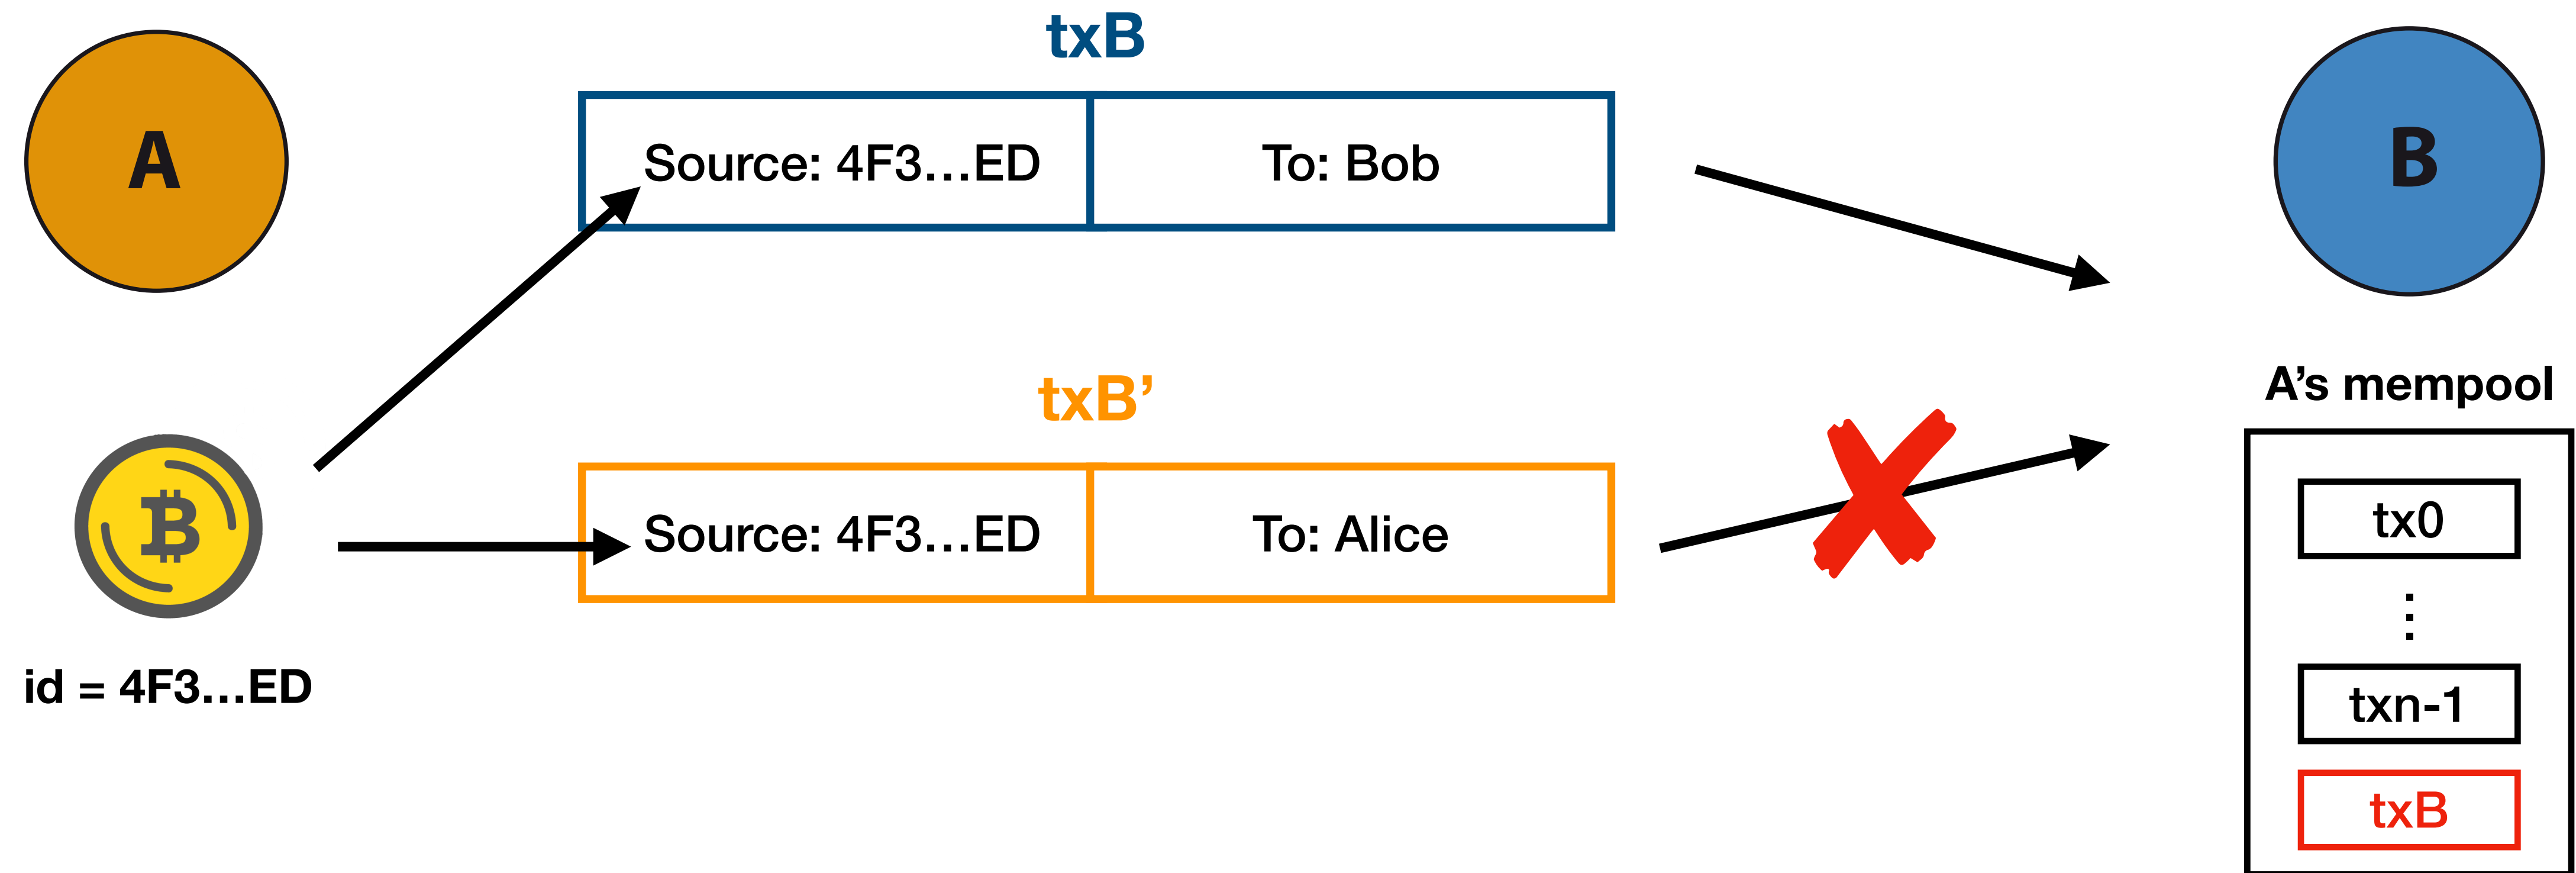
DOUBLE-SPENDING TRANSACTIONS



DOUBLE-SPENDING TRANSACTIONS



DOUBLE-SPENDING TRANSACTIONS



DOUBLE-SPENDING TRANSACTIONS



Given a pair of double-spending transactions, a node will **pick the first one** it learns about

id = 4F3...ED

:

txn-1

txB

A BASIC TOPOLOGY INFERRING TECHNIQUE



Two nodes


Three transactions

Observation tool

A BASIC TOPOLOGY INFERRING TECHNIQUE



Two nodes




Observation tool

Three transactions

A BASIC TOPOLOGY INFERRING TECHNIQUE




Two nodes



Observation tool


Three transactions



A BASIC TOPOLOGY INFERRING TECHNIQUE



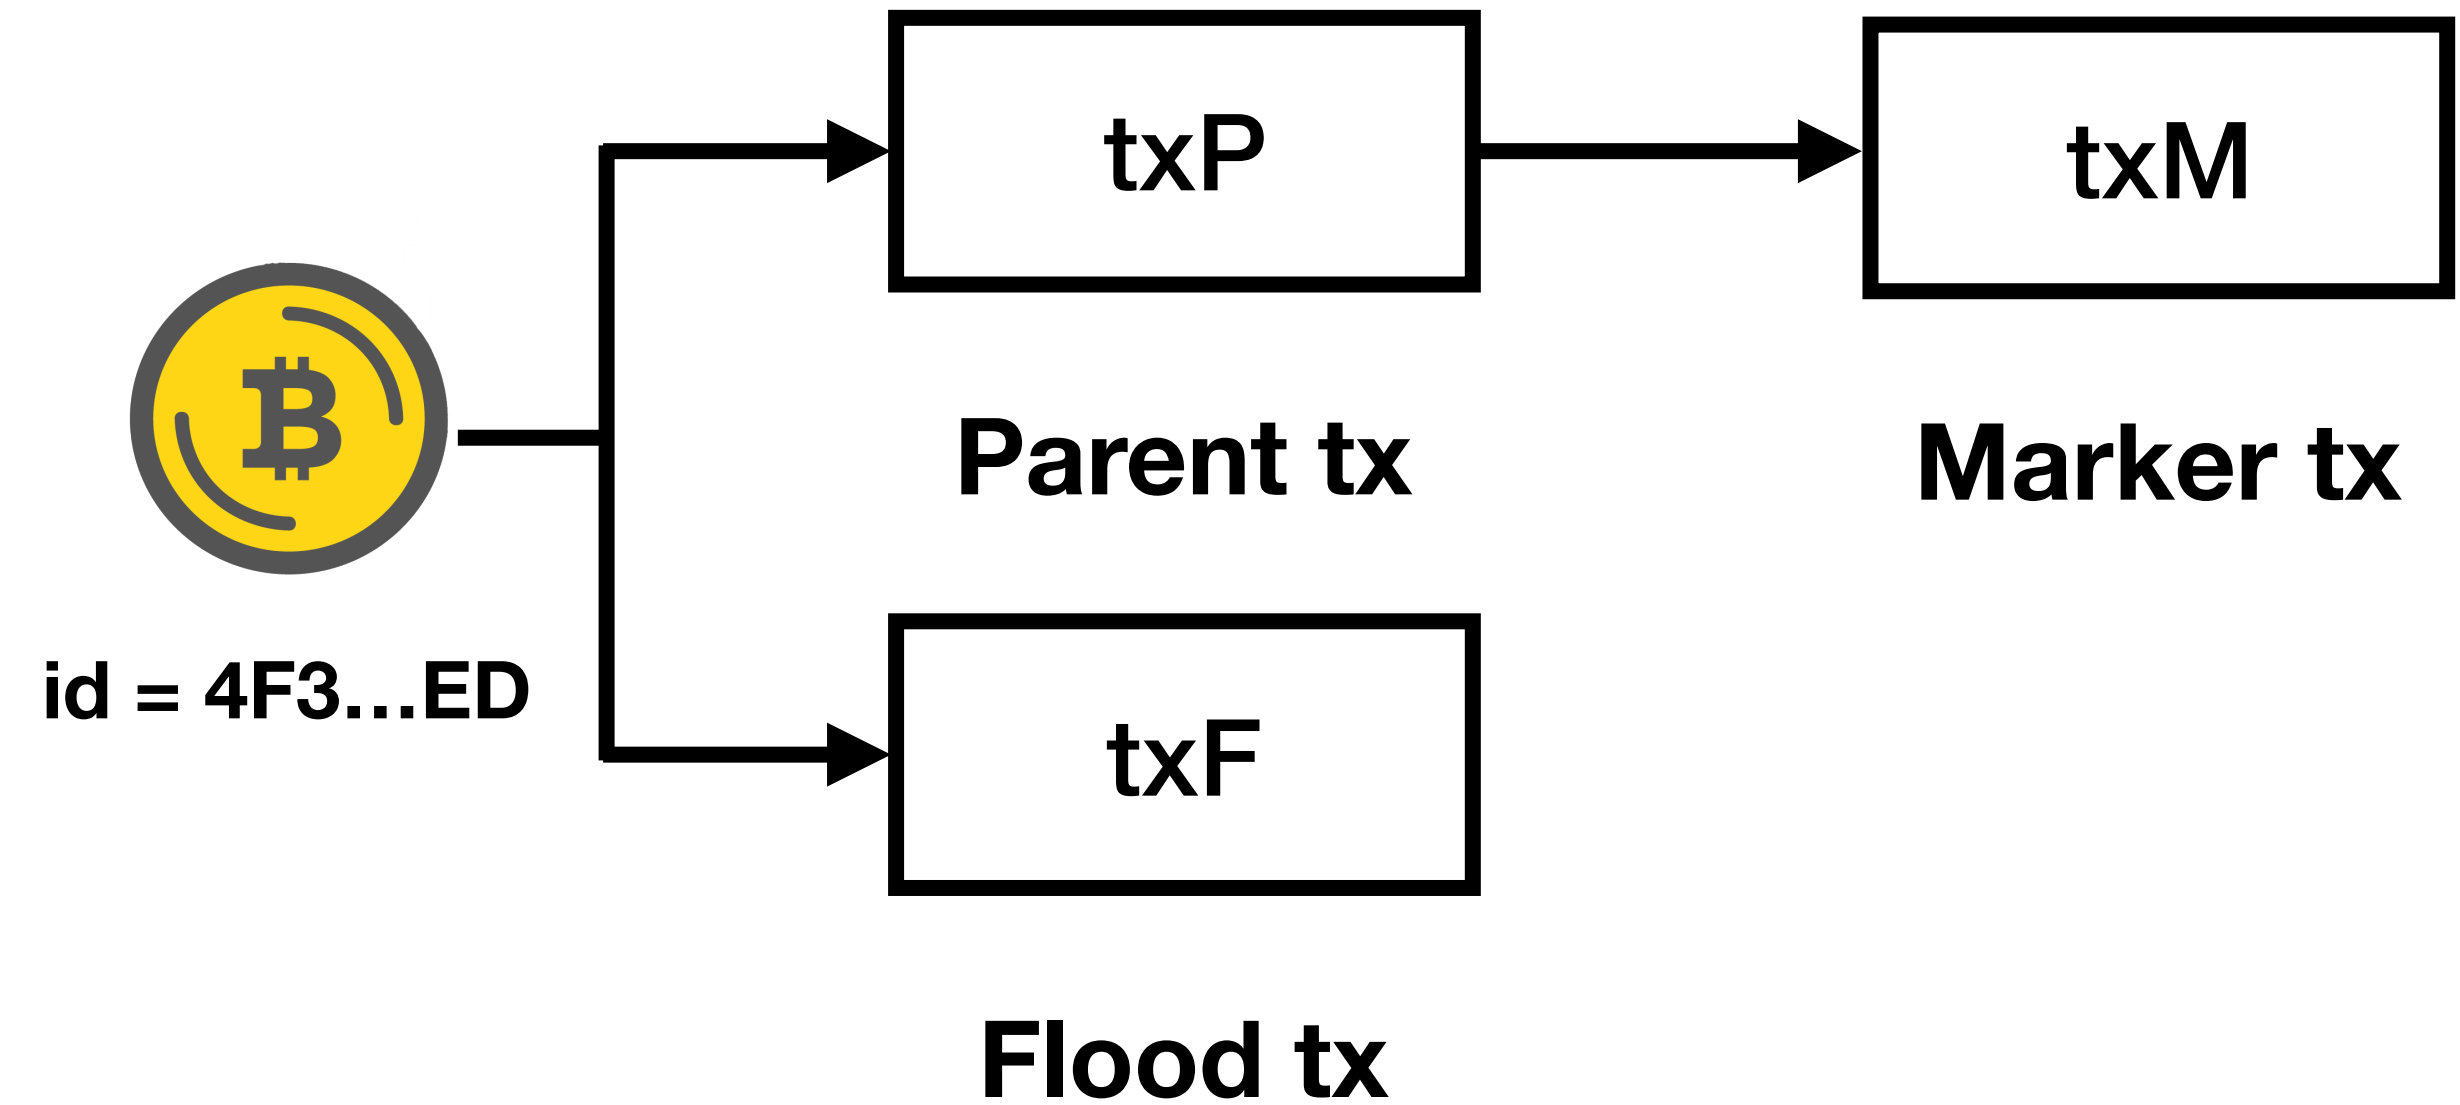
Two nodes



Observation tool




Three transactions



A BASIC TOPOLOGY INFERRING TECHNIQUE

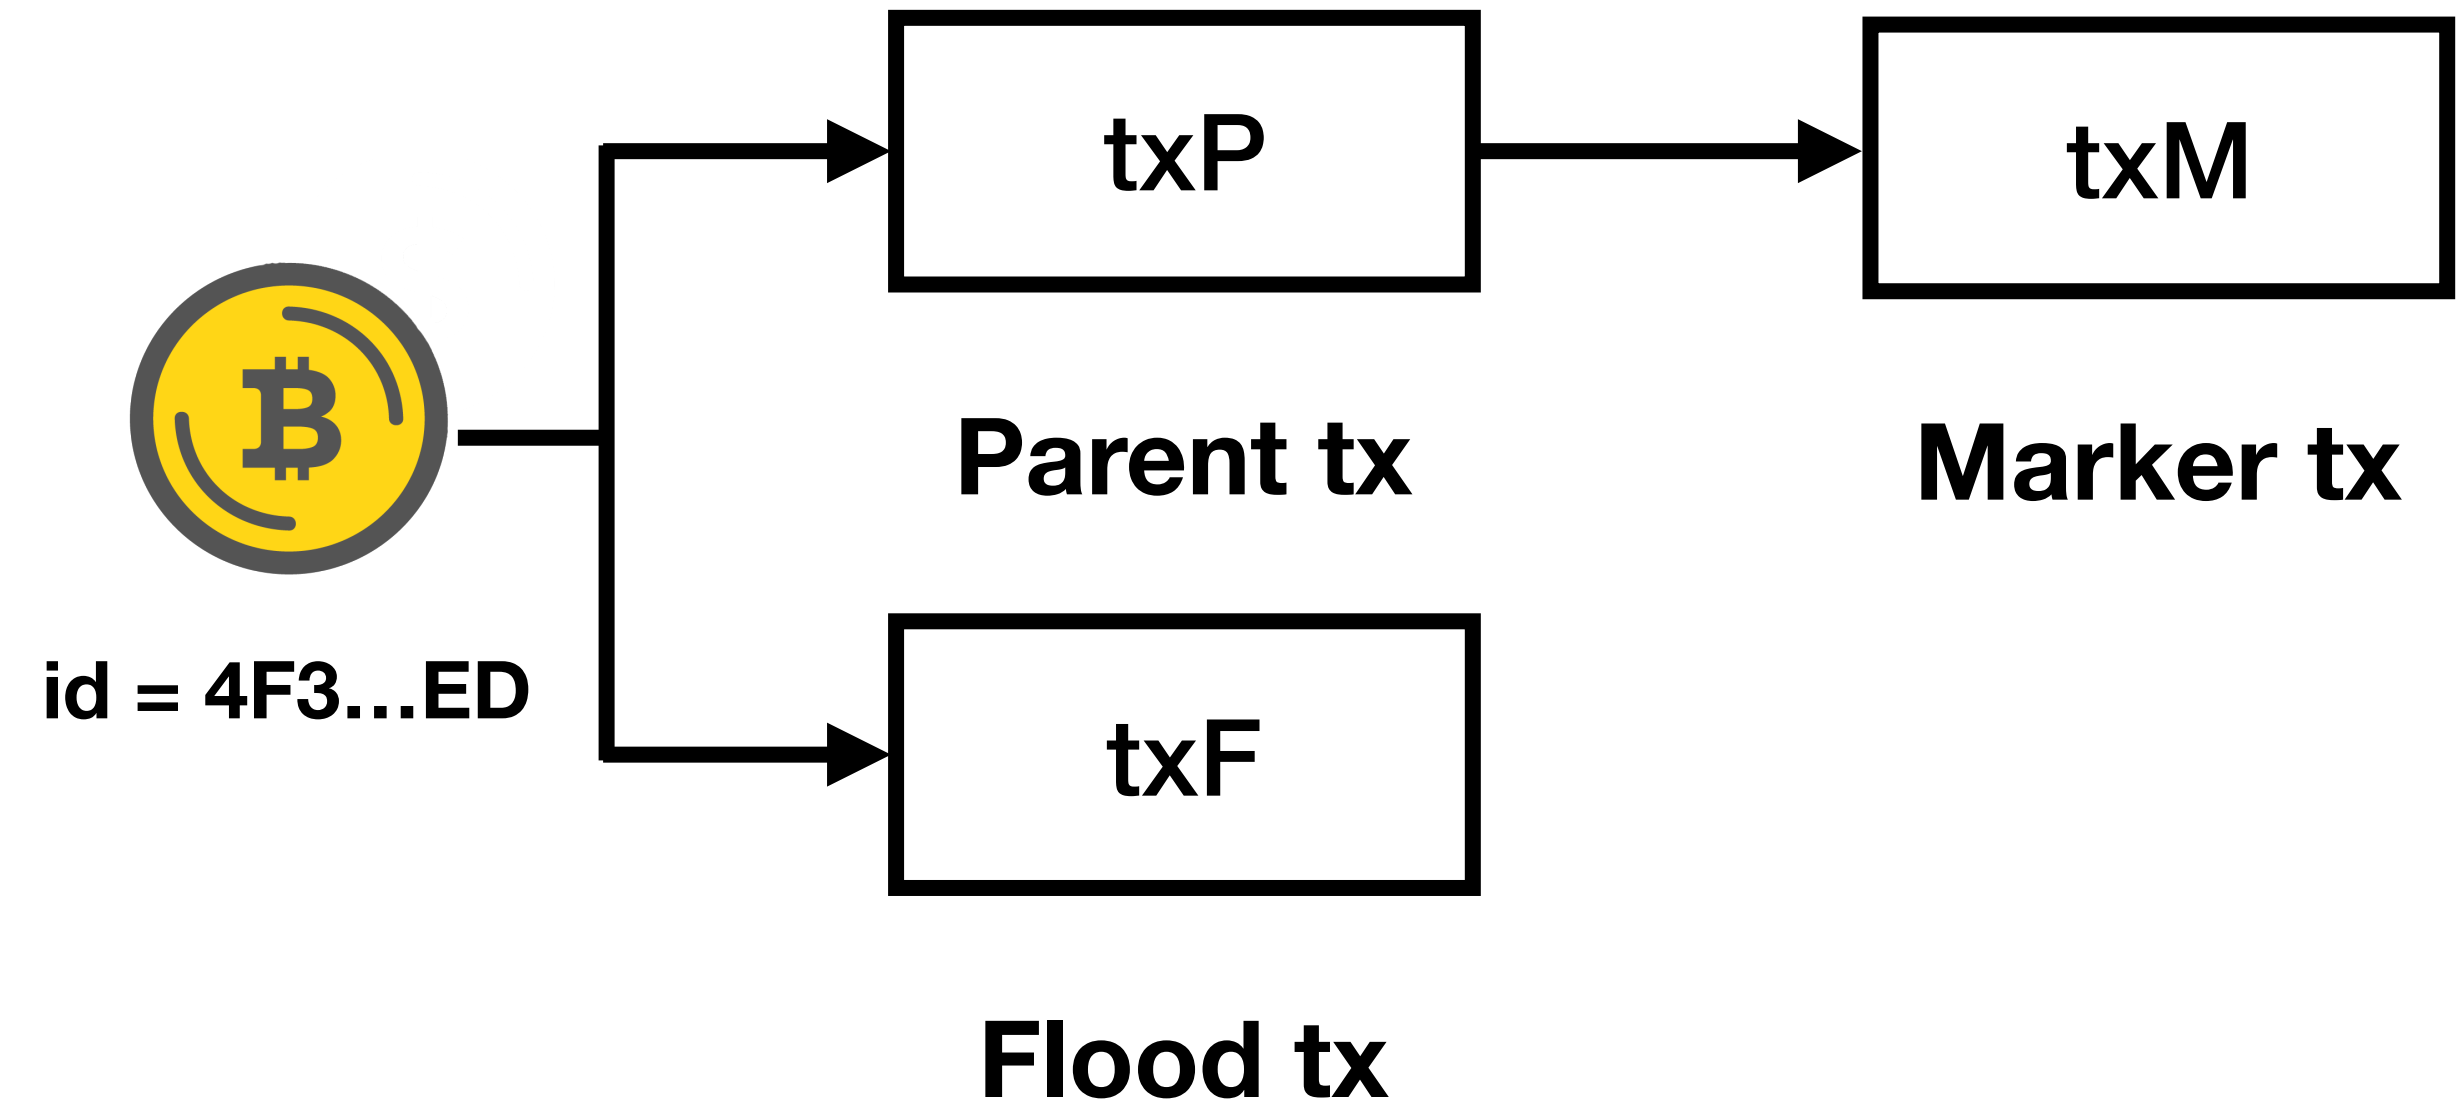
Two nodes



Observation tool





Three transactions






Andrew Miller, James Litton, Andrew Pachulski, Neal Gupta, Dave Levin, Neil Spring,
Bobby Bhattacharjee
Discovering Bitcoin's Public Topology and Influential Nodes


NEGATIVE INFERRING TECHNIQUE




NEGATIVE INFERRING TECHNIQUE




NEGATIVE INFERRING TECHNIQUE




NEGATIVE INFERRING TECHNIQUE




NEGATIVE INFERRING TECHNIQUE




NEGATIVE INFERRING TECHNIQUE




NEGATIVE INFERRING TECHNIQUE




NEGATIVE INFERRING TECHNIQUE




NEGATIVE INFERRING TECHNIQUE




NEGATIVE INFERRING TECHNIQUE




NEGATIVE INFERRING TECHNIQUE



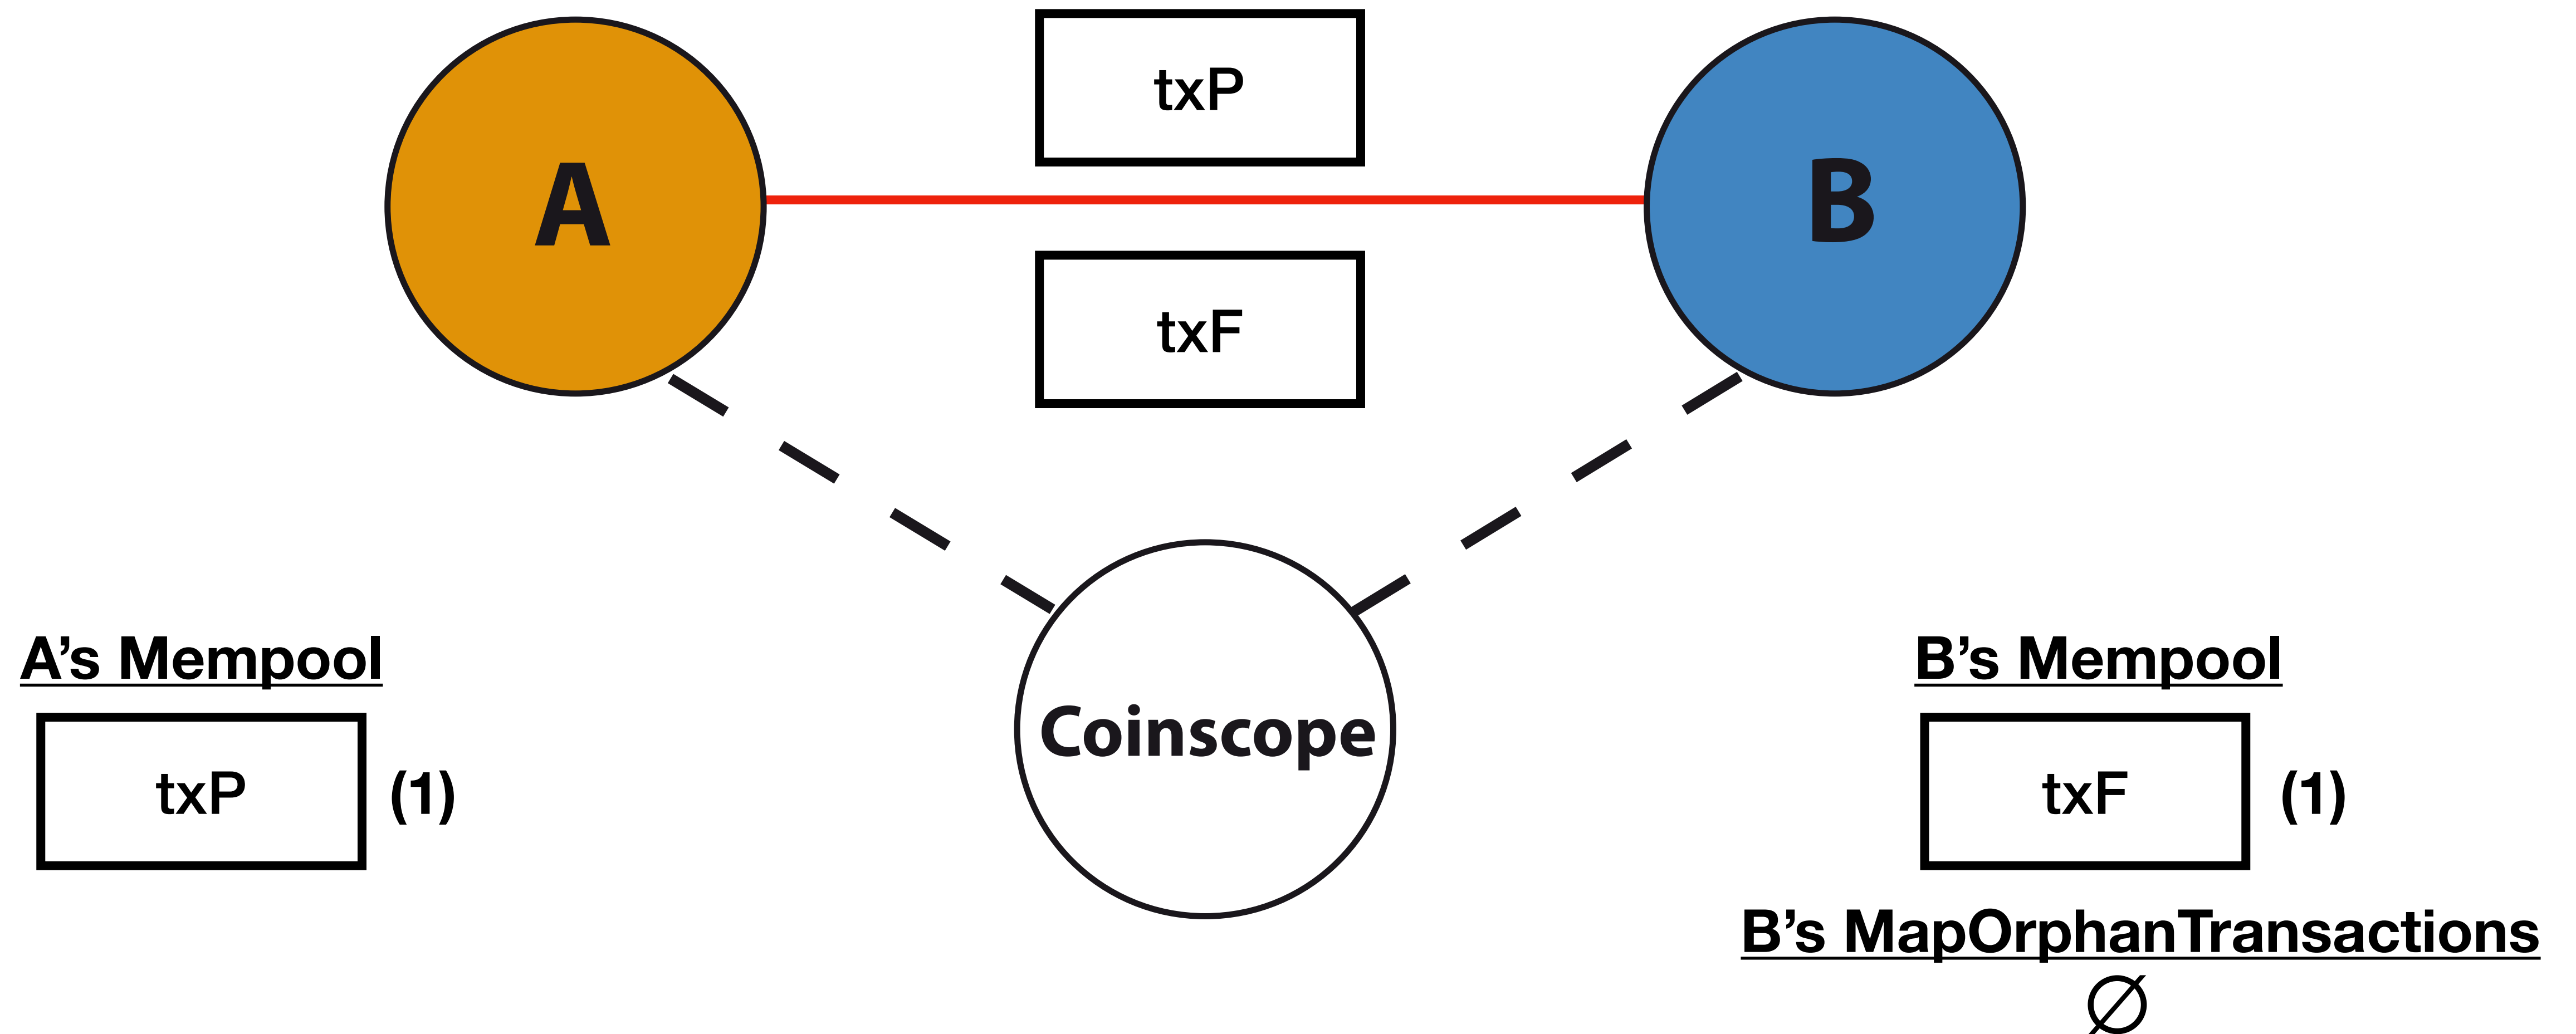
POSITIVE INFERRING TECHNIQUE




POSITIVE INFERRING TECHNIQUE




POSITIVE INFERRING TECHNIQUE




POSITIVE INFERRING TECHNIQUE




POSITIVE INFERRING TECHNIQUE




POSITIVE INFERRING TECHNIQUE



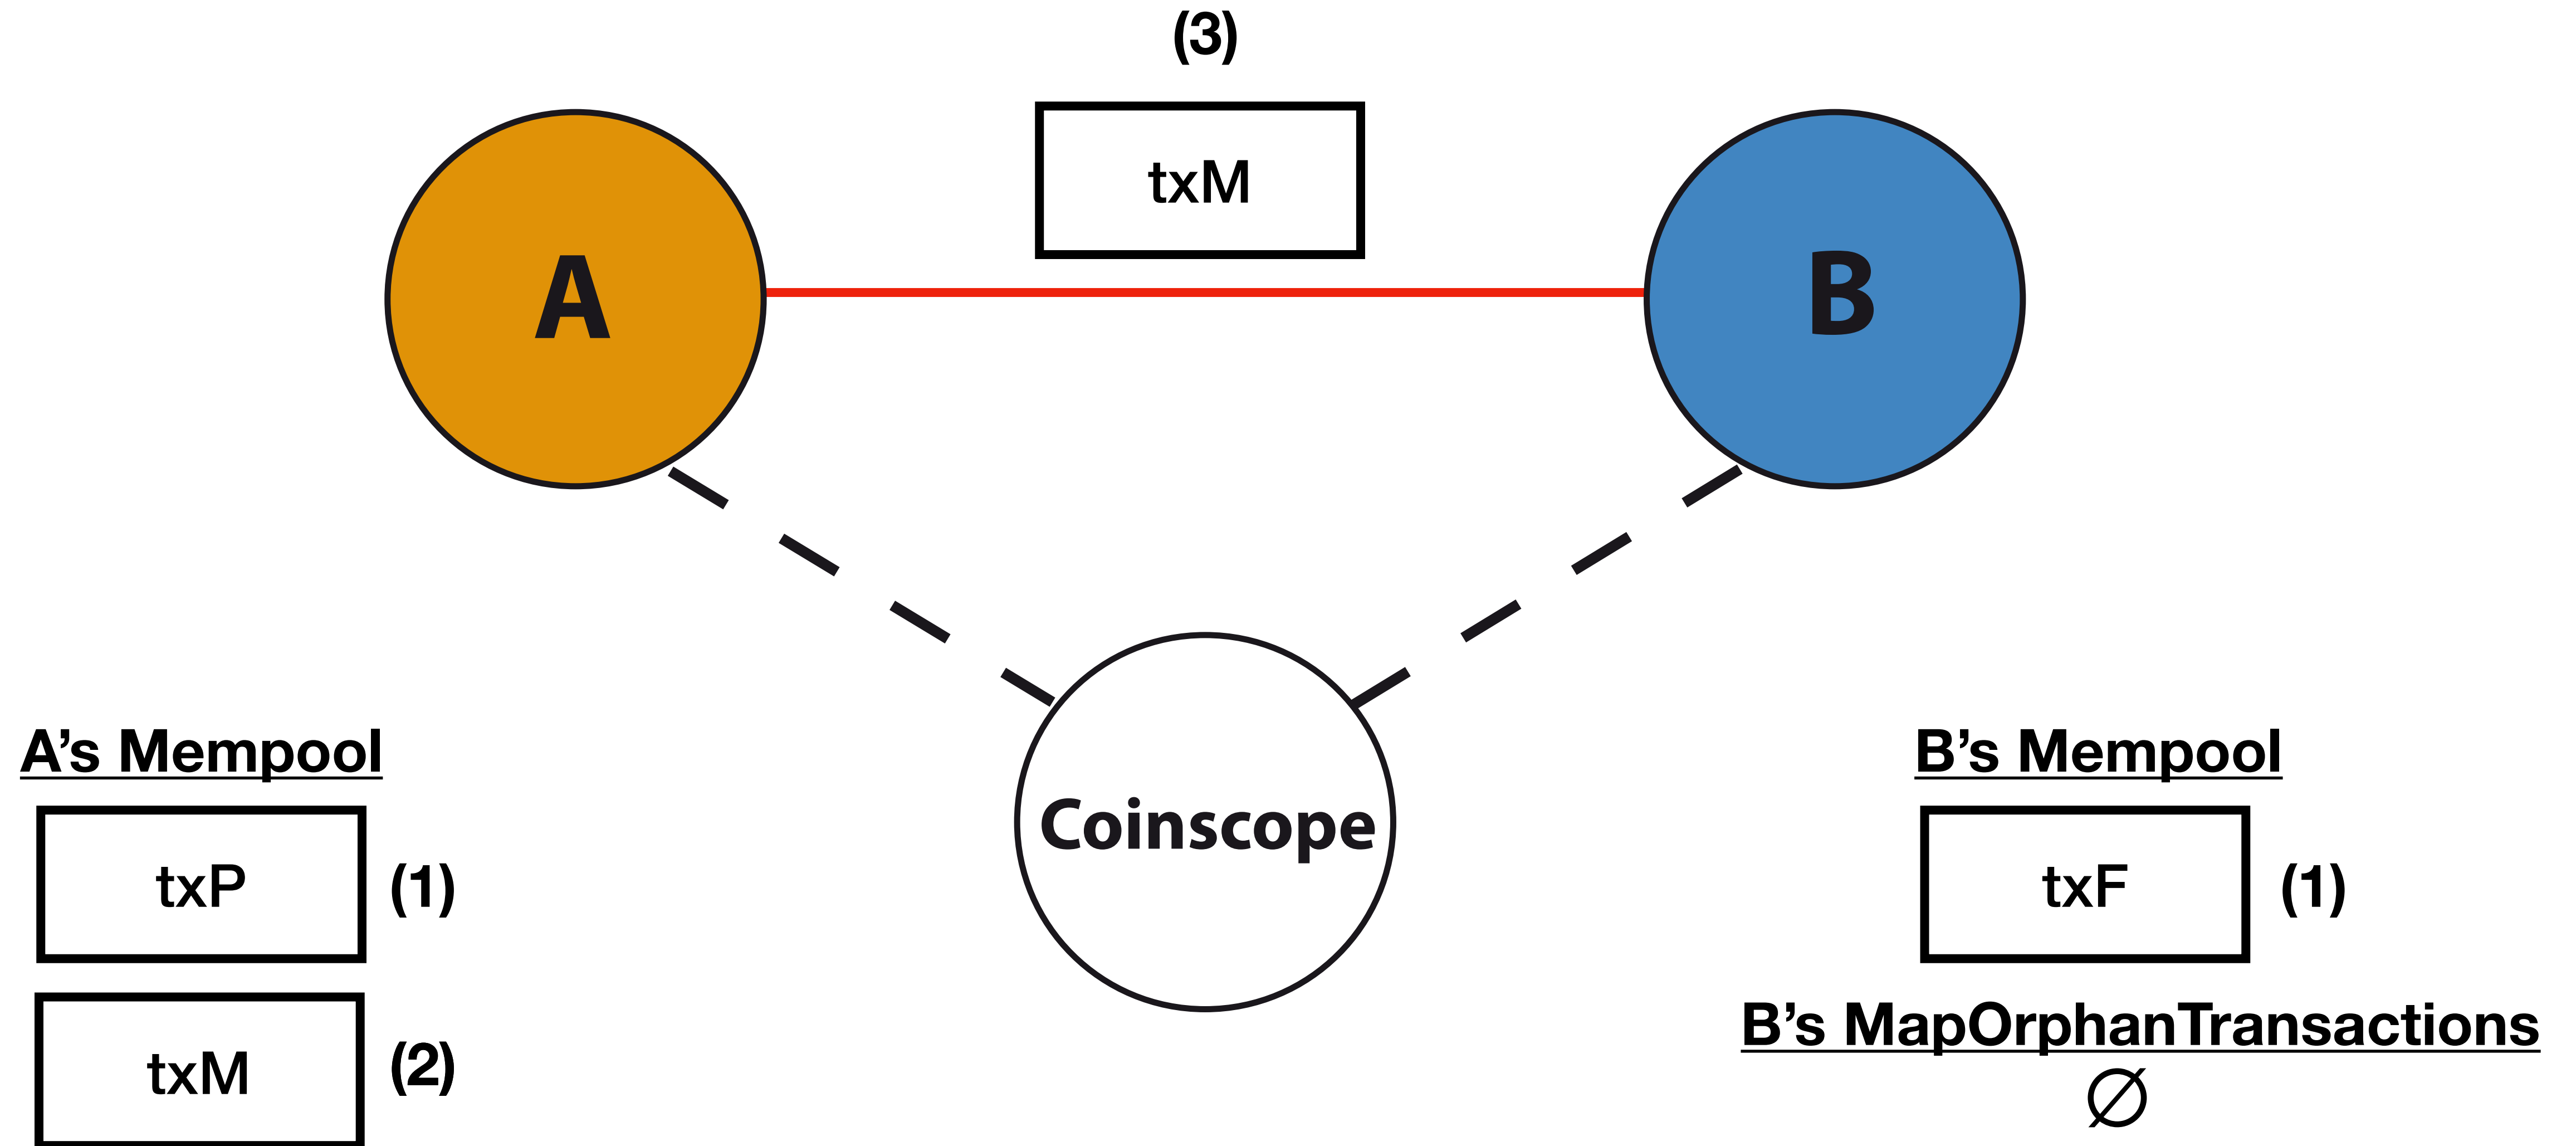
POSITIVE INFERRING TECHNIQUE




POSITIVE INFERRING TECHNIQUE



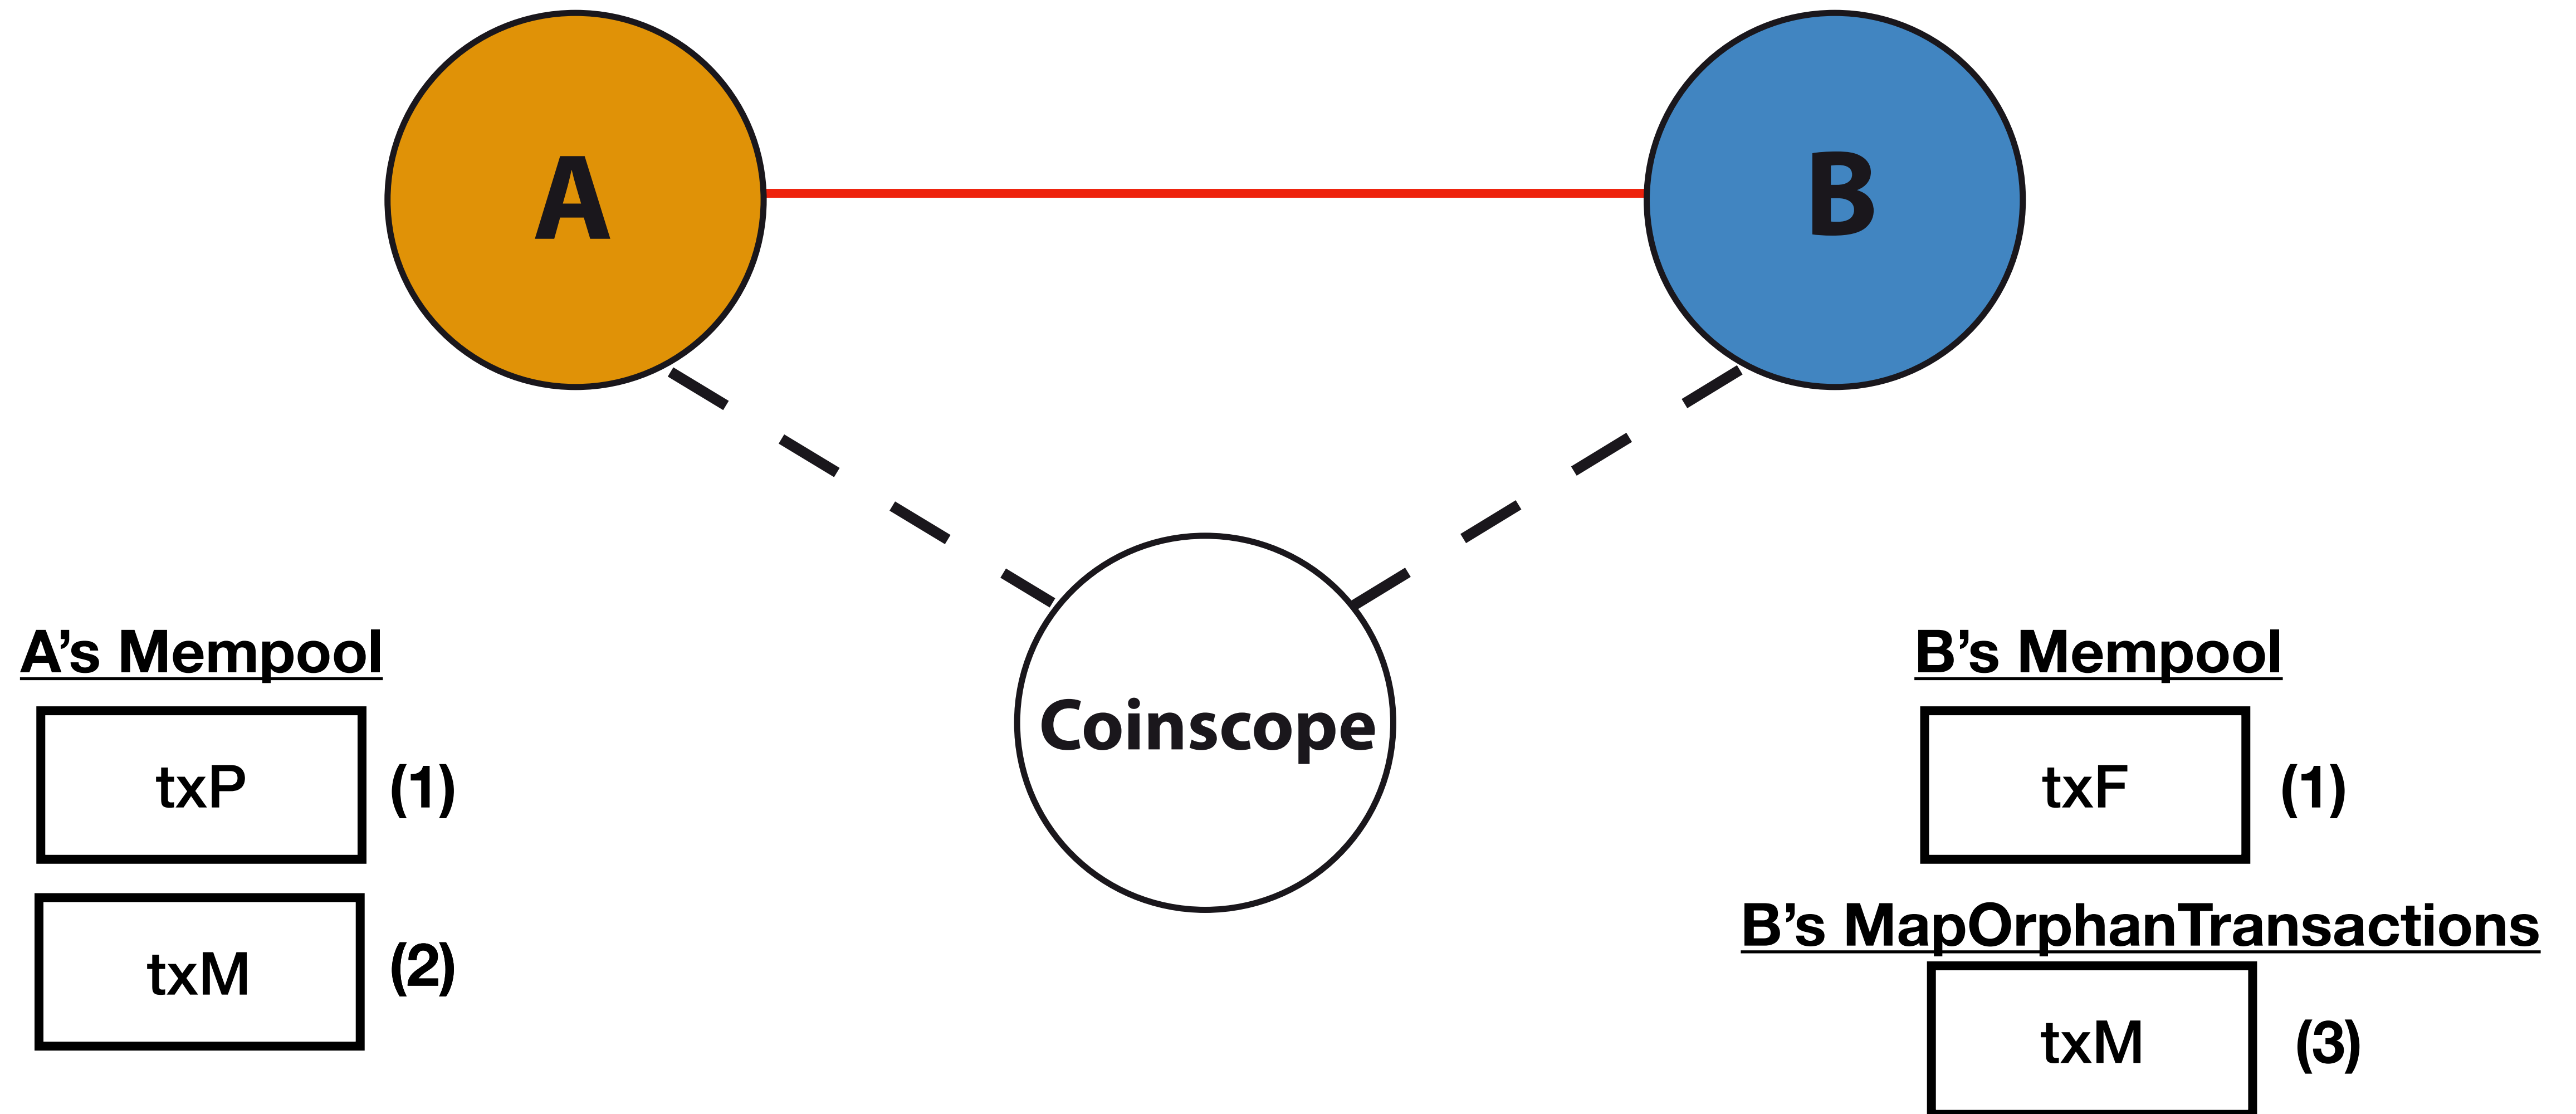
POSITIVE INFERRING TECHNIQUE



POSITIVE INFERRING TECHNIQUE




POSITIVE INFERRING TECHNIQUE




A BASIC TOPOLOGY INFERRING TECHNIQUE



Positive edge inferring




Negative edge inferring



A BASIC TOPOLOGY INFERRING TECHNIQUE




Position




A's Mempool

| | |
|-----|-----|
| txP | (1) |
| txM | (2) |



$inv(h(txM))$

inferring



B's Mempool

| | |
|-----|-----|
| txF | (1) |
|-----|-----|


B's MapOrphanTransactions

| |
|-------------|
| \emptyset |
|-------------|

A BASIC TOPOLOGY INFERRING TECHNIQUE





Position




A's Mempool

| | |
|-----|-----|
| txP | (1) |
| txM | (2) |



inferring



B's Mempool

| | |
|-----|-----|
| txF | (1) |
|-----|-----|



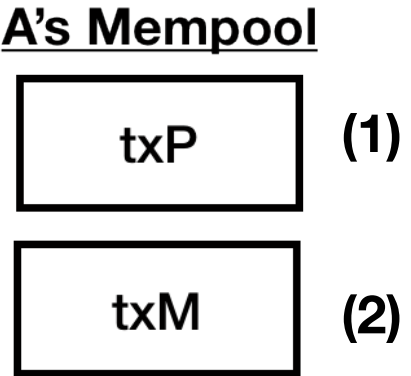
B's MapOrphanTransactions

∅



A BASIC TOPOLOGY INFERRING TECHNIQUE



Position




inferring







A BASIC TOPOLOGY INFERRING TECHNIQUE



Position




A's Mempool



edge does not exist

inferring



B's Mempool







B's MapOrphanTransactions

\emptyset



A BASIC TOPOLOGY INFERRING TECHNIQUE



Position






inferring



A BASIC TOPOLOGY INFERRING TECHNIQUE





Position



$inv(h(txM))$

\emptyset



inferring





A BASIC TOPOLOGY INFERRING TECHNIQUE



Position




inferring



ITS NOT THAT EASY




ITS NOT THAT EASY



ITS NOT THAT EASY




Long story short, if you add an additional node to the equation it will fail



ITS NOT THAT EASY




Long story short, if you add an additional node to the equation it will fail



ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

\emptyset

C's Mempool

\emptyset

B's Mempool

\emptyset


B's MapOrphanTransactions

\emptyset

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

\emptyset

C's Mempool

\emptyset

B's Mempool

\emptyset


B's MapOrphanTransactions

\emptyset

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

\emptyset

C's Mempool

\emptyset

B's Mempool

\emptyset


B's MapOrphanTransactions

\emptyset

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

txP (1)

C's Mempool

∅

B's Mempool

txF (1)


B's MapOrphanTransactions

∅

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

txP (1)

C's Mempool

∅

B's Mempool

txF (1)


B's MapOrphanTransactions

∅

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

txP (1)

C's Mempool

∅

B's Mempool

txF (1)


B's MapOrphanTransactions

∅

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

txP (1)

C's Mempool

txP (2)

B's Mempool

txF (1)

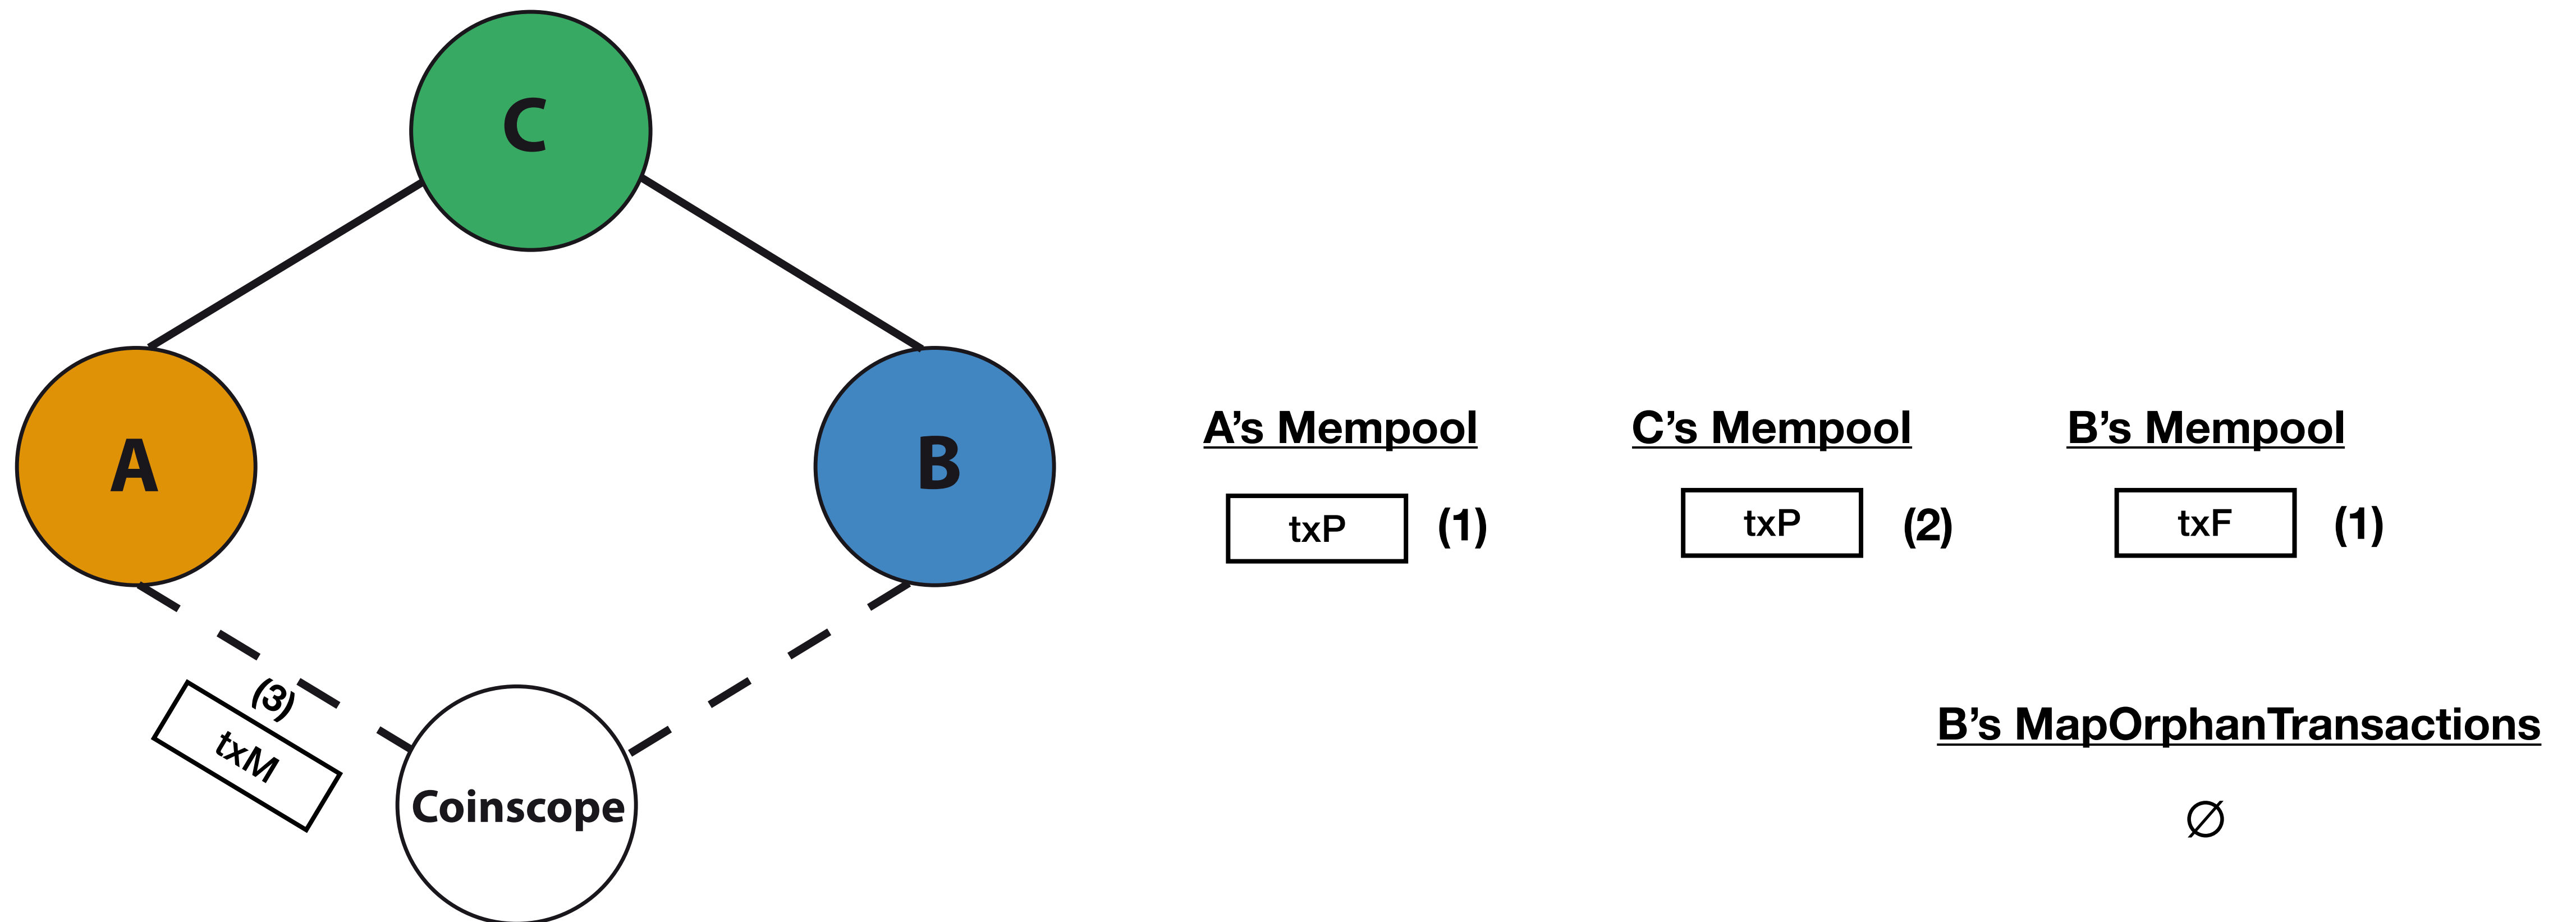
B's MapOrphanTransactions

∅

ITS NOT THAT EASY




Long story short, if you add an additional node to the equation it will fail



ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

txP (1)

C's Mempool

txP (2)

B's Mempool

txF (1)


B's MapOrphanTransactions

∅

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

txP (1)

txM (3)

C's Mempool

txP (2)

B's Mempool

txF (1)


B's MapOrphanTransactions

∅

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

| | |
|-----|-----|
| txP | (1) |
| txM | (3) |

C's Mempool

| | |
|-----|-----|
| txP | (2) |
|-----|-----|

B's Mempool

| | |
|-----|-----|
| txF | (1) |
|-----|-----|


B's MapOrphanTransactions

∅

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

txP (1)

txM (3)

C's Mempool

txP (2)

B's Mempool

txF (1)


B's MapOrphanTransactions

∅

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

txP (1)

txM (3)

C's Mempool

txP (2)

txM (4)

B's Mempool

txF (1)


B's MapOrphanTransactions

∅

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

- txP (1)
- txM (3)

C's Mempool

- txP (2)
- txM (4)

B's Mempool

- txF (1)


B's MapOrphanTransactions

∅

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

- txP (1)
- txM (3)

C's Mempool

- txP (2)
- txM (4)

B's Mempool

- txF (1)


B's MapOrphanTransactions

∅

ITS NOT THAT EASY



Long story short, if you add an additional node to the equation it will fail



A's Mempool

| | |
|-----|-----|
| txP | (1) |
| txM | (3) |

C's Mempool

| | |
|-----|-----|
| txP | (2) |
| txM | (4) |

B's Mempool


| | |
|-----|-----|
| txF | (1) |
|-----|-----|

B's MapOrphanTransactions


| | |
|-----|-----|
| txM | (5) |
|-----|-----|

MAKE THIS WORK IN A REAL NETWORK

Isolation



Synchronicity



Scalability


$$\approx 3n \text{ txs}$$




$$\approx 2\sqrt{n} \text{ txs}$$

$$n = \#nodes$$

ACHIEVING ISOLATION AND SYNCHRONICITY

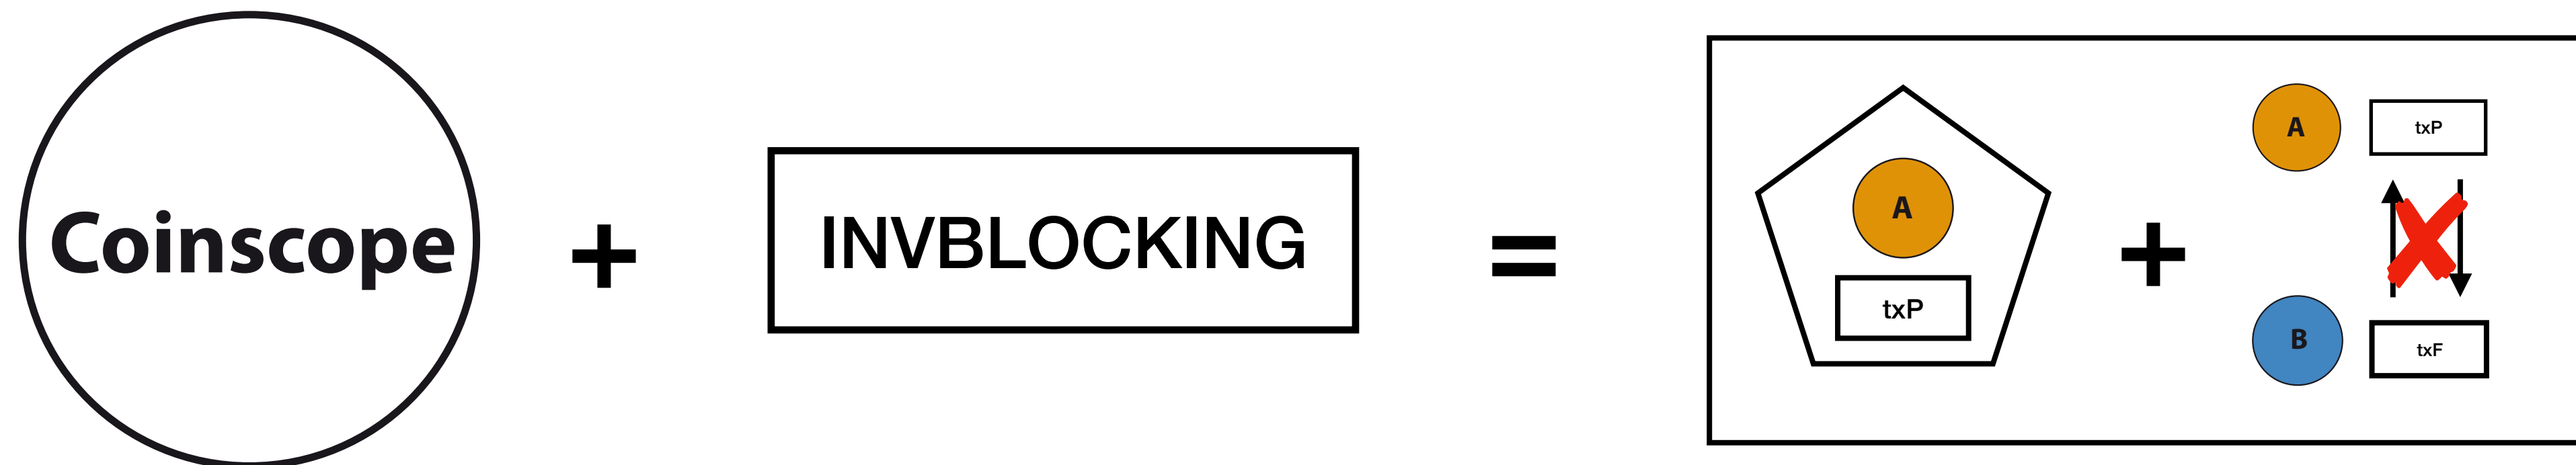


ACHIEVING ISOLATION AND SYNCHRONICITY



INVBLOCKING


ACHIEVING ISOLATION AND SYNCHRONICITY



INVBLOCKING

- Send **INV** messages with **txP** and **txF** to the whole network


ACHIEVING ISOLATION AND SYNCHRONICITY



INVBLOCKING

- Send **INV** messages with **txP** and **txF** to the whole network
- Nodes will **ask us** about **txP** and **txF**



ACHIEVING ISOLATION AND SYNCHRONICITY





INVBLOCKING

- Send **INV** messages with **txP** and **txF** to the whole network
- Nodes will **ask us** about **txP** and **txF**
- We **withhold** the information effectively **blocking the propagation of txP and txF**

ACHIEVING SCALABILITY



source set



sink set

$$\approx 3n \text{ txs}$$



$$\approx 2\sqrt{n} \text{ txs}$$

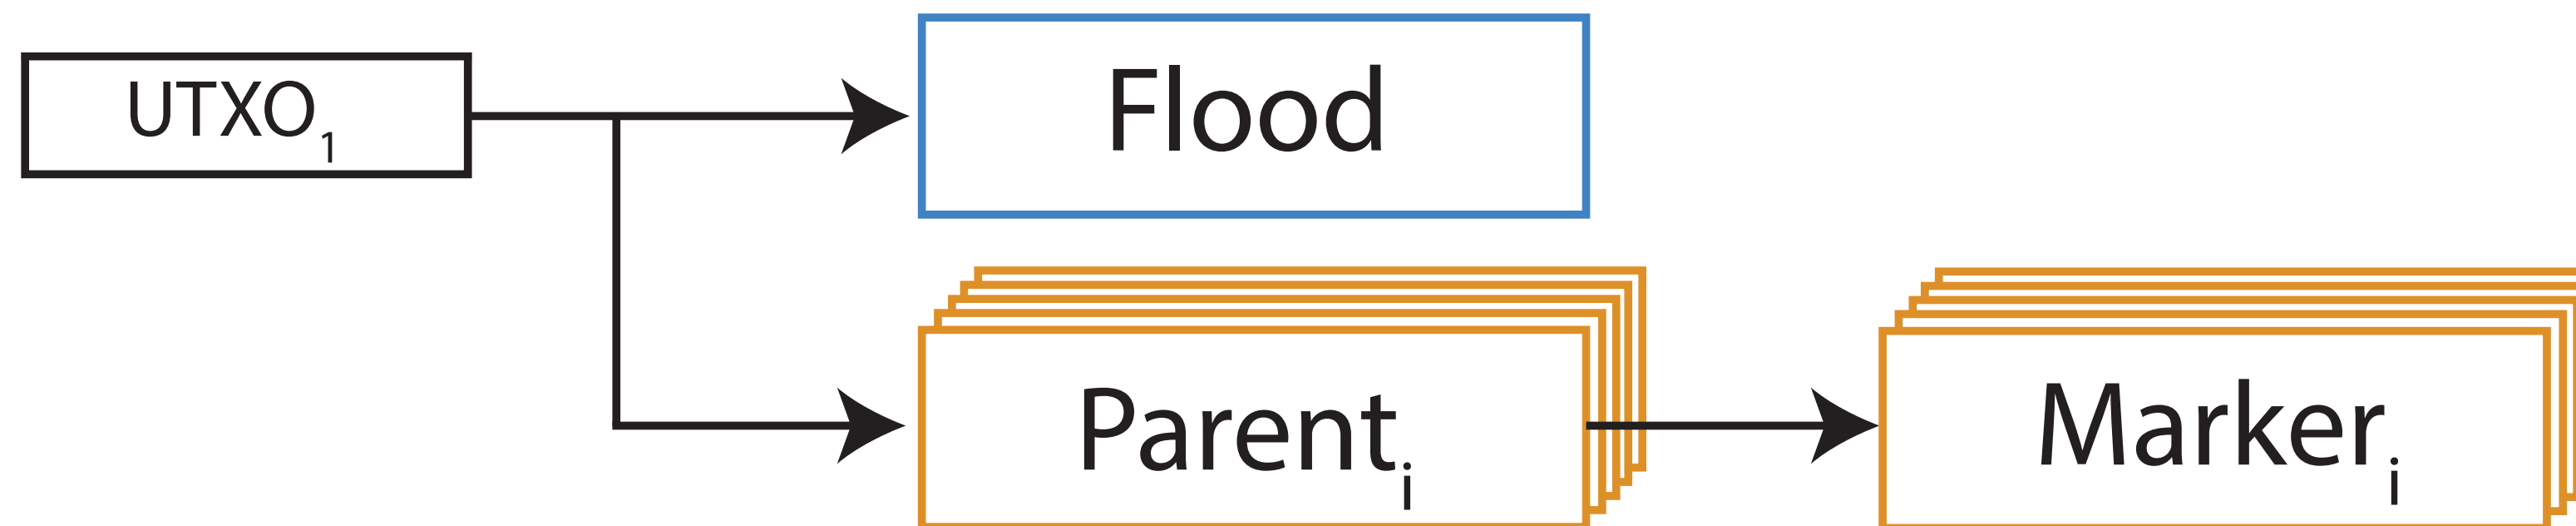
ACHIEVING SCALABILITY



source set

$\approx 3n \text{ txs}$

$$n = \text{source set} + \text{sink set}$$



edges are inferred by checking the **propagation of markers** from the **source set** to the **sink set**

ACHIEVING SCALABILITY (CONT)



To satisfy the **scalability property**, we are inferring several edges at the same time

We need to make sure that the MapOrphanTransactions pool **have enough room** to store all our orphans

Otherwise **false negatives** could occur

ACHIEVING SCALABILITY (CONT)



To satisfy the **scalability property**, we are inferring several edges at the same time

We need to make sure that the MapOrphanTransactions pool **have enough room** to store all our orphans

Otherwise **false negatives** could occur



Refer to the paper for details about the orphan pool handling

TXPROBE - PROTOCOL OVERVIEW



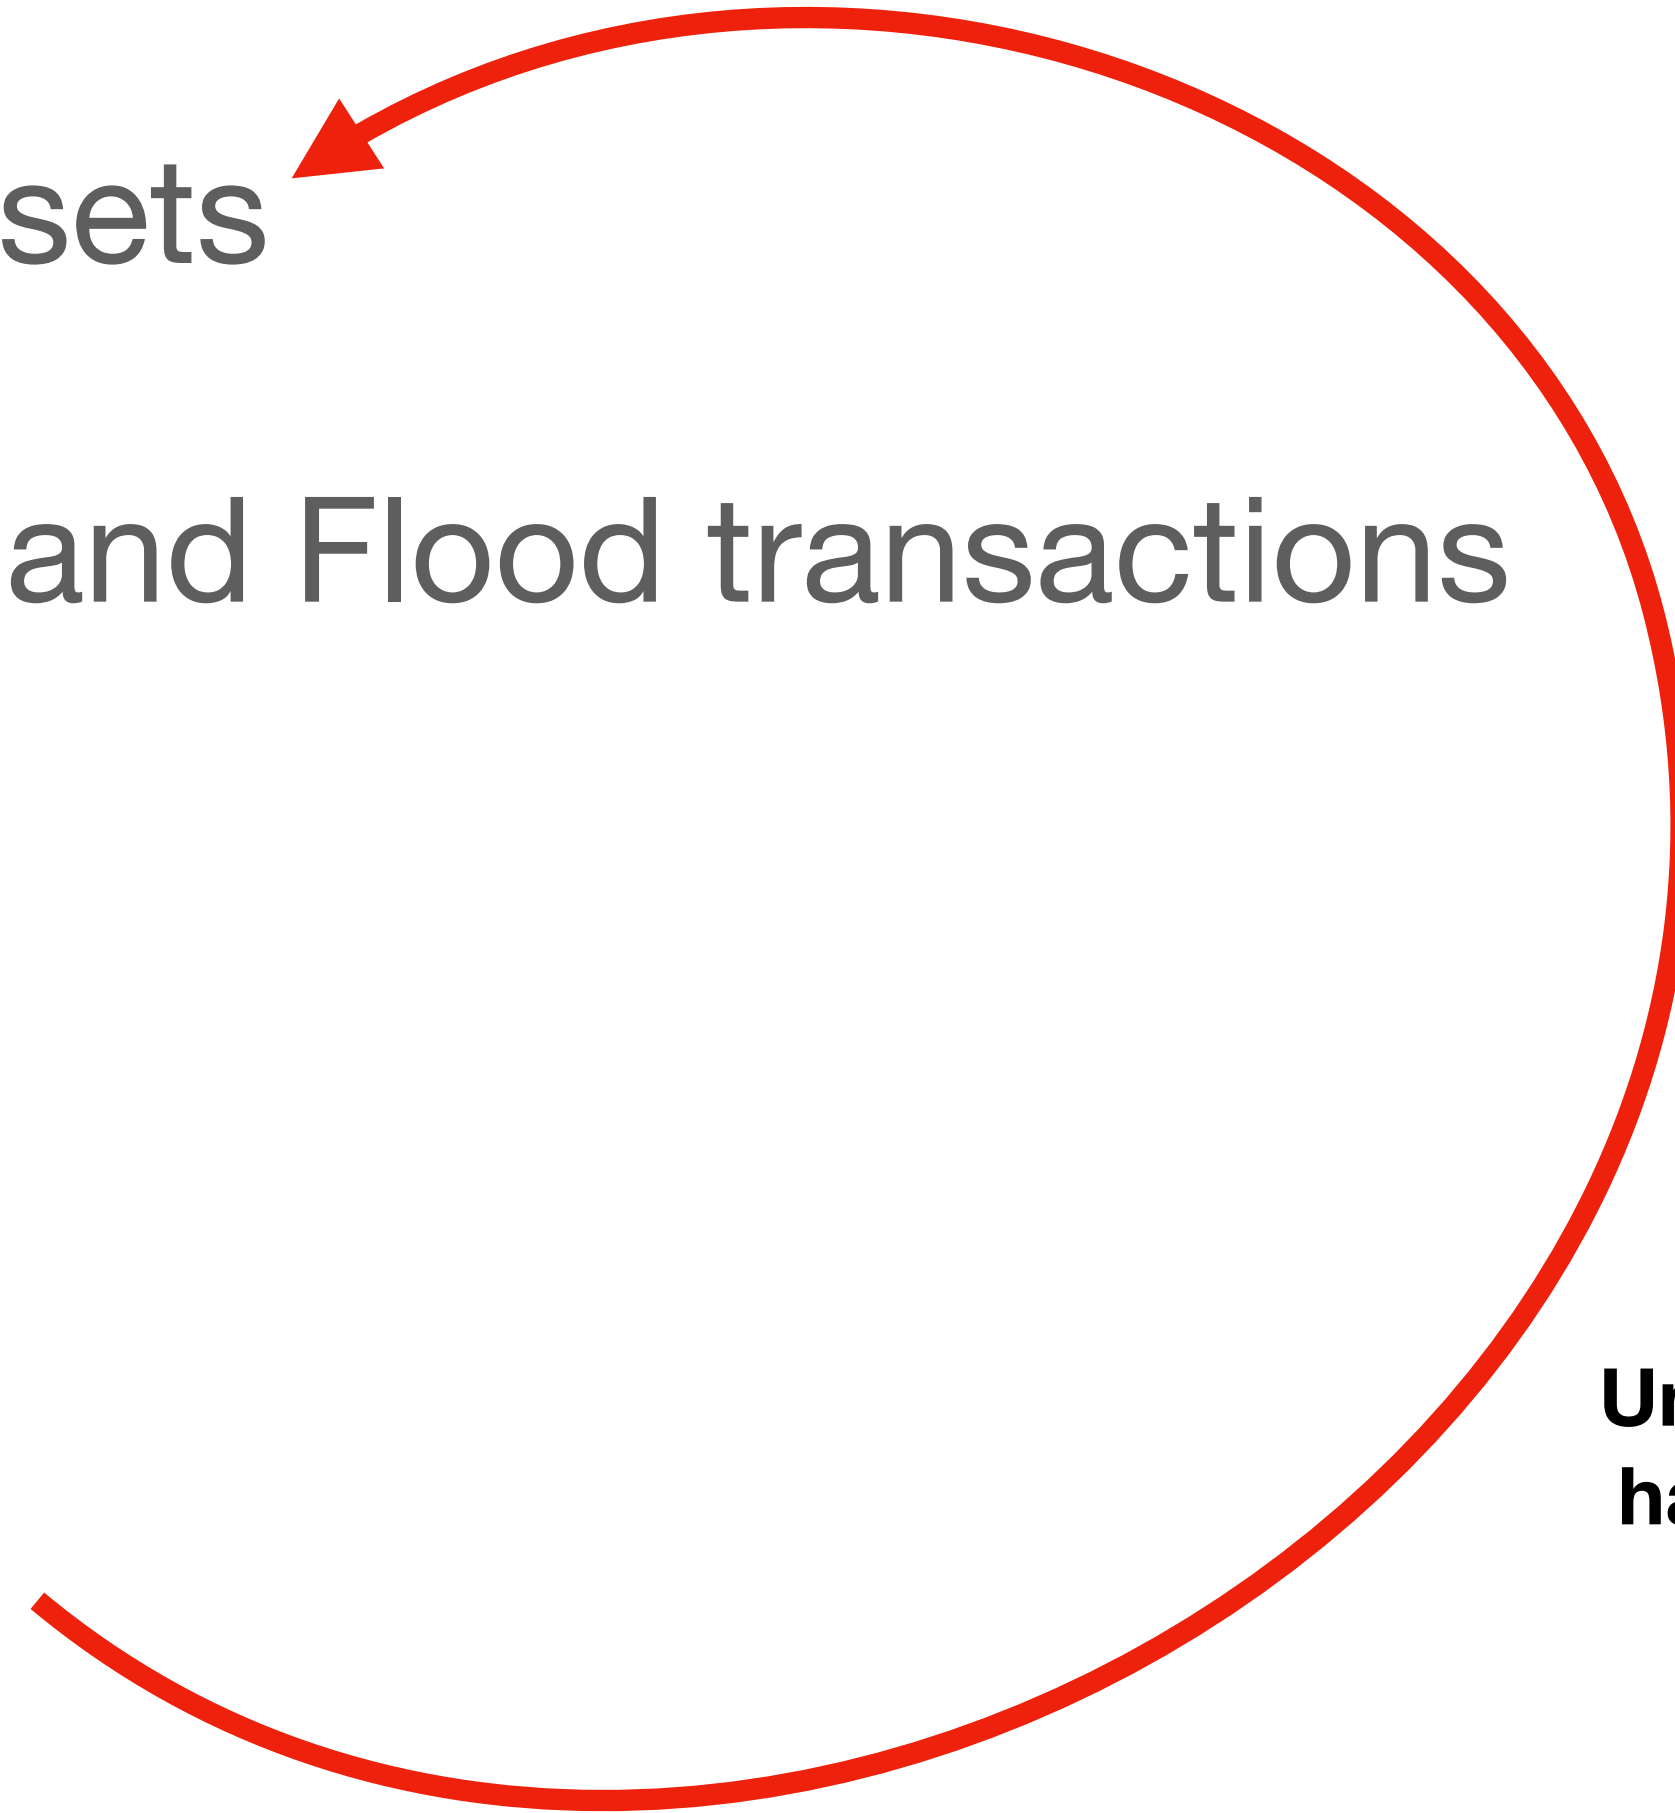
- **Choose** source and sink sets
- **Create** Parents, Markers and Flood transactions
- **INVBLOCK** the network
- **Send** transactions
- **Request** markers back

TXPROBE - PROTOCOL OVERVIEW



- **Choose** source and sink sets
- **Create** Parents, Markers and Flood transactions
- **INVBLOCK** the network
- **Send** transactions
- **Request** markers back

**Until every pair of nodes
have been in a different
set at least once**



TXPROBE - DATA VALIDATION



- We run 5 Bitcoin Core nodes as **ground truth**
- Nodes are included as part of the source set
- We define our precision / recall by checking how well can we infer the ground truth nodes connections

TXPROBE - COSTS



For a network like **Bitcoin mainnet**:

nodes \approx 10000

time \approx 8.25 hours

cost = 573210-764280 satoshi (5 sat/byte) \approx **\$(20-30)**

**We can actually do better, add
some notes about this**

WHY TESTNET AND NO MAINNET?



- TxProbe is rather invasive: it empties the **MapOrphanTransactions pool** of all nodes in the network every round
- We could not measure the implication that such behavior may have had on the **propagation of regular transactions**
- The technique **could** have also **be seen as an attack** to the network