

Bitcoin network layer and information propagation

Sergi Delgado Segura



WHAT ARE WE GOING TO COVER?

Differences between client/server and peer-to-peer paradigms

How a new node joins the network

- How it learns about the network
- How others learn about it

Actors and their role in the network

The gossip protocol

Data propagation

- Transactions and blocks
- 0-conf and double-spending

Node misbehavior

Network based attacks

Network topology

BEFORE WE START

We will use Bitcoin as an example when explaining how certain parts of the network work. However, the same mechanisms apply to most of the existing cryptocurrencies with slight modifications (some times even without any).

Also keep in mind that for most things within cryptocurrencies there is no formal specification but the live code. Therefore some details may change in the near future.



Introduction

CLIENT-SERVER PARADIGM (1/2)

Classic paradigm where actors are split into **clients** and **servers**

Servers:

- serve specific **resources** upon request
- can also provide different types of **services**

CLIENT-SERVER PARADIGM (2/2)

Clients:

- resource/service requesters
- do not share resources or provide any service

Clients initiate the communication and need to know the server endpoint

Classical examples: WWW, DNS, Email, etc

PEER-TO-PEER (P2P) PARADIGM

All actors (**peers**) are equal and have both client and server capabilities

Services / resources can be shared between several peers or found in a single location

Each peer can choose what to serve/request

Quite usual paradigm for distributed file sharing (e.g: BitTorrent)

Usual problems: Bootstrapping and file searching

P2P BOOTSTRAPPING

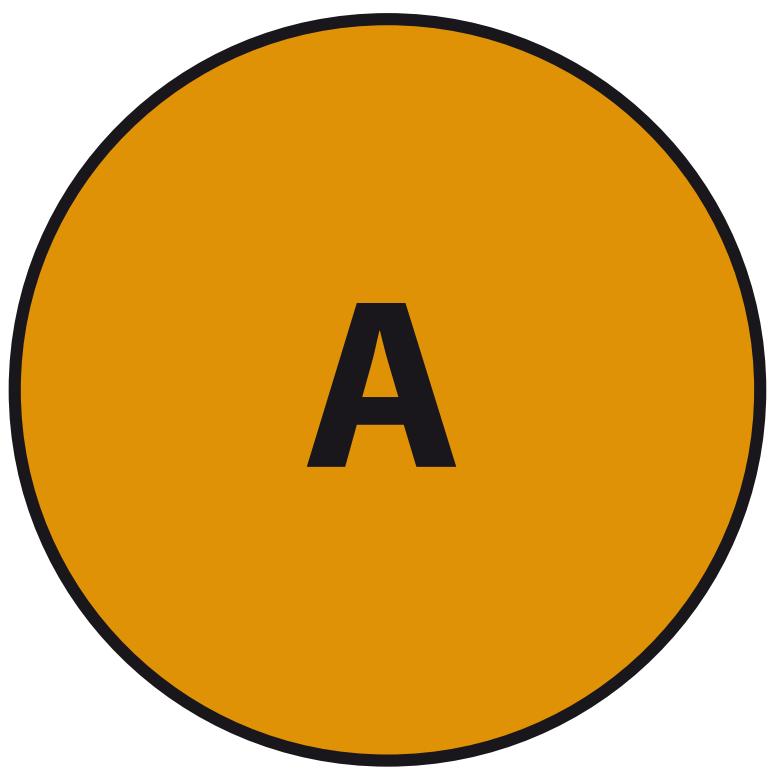
How do you find peers when you run a new node in the network?

P2P BOOTSTRAPPING

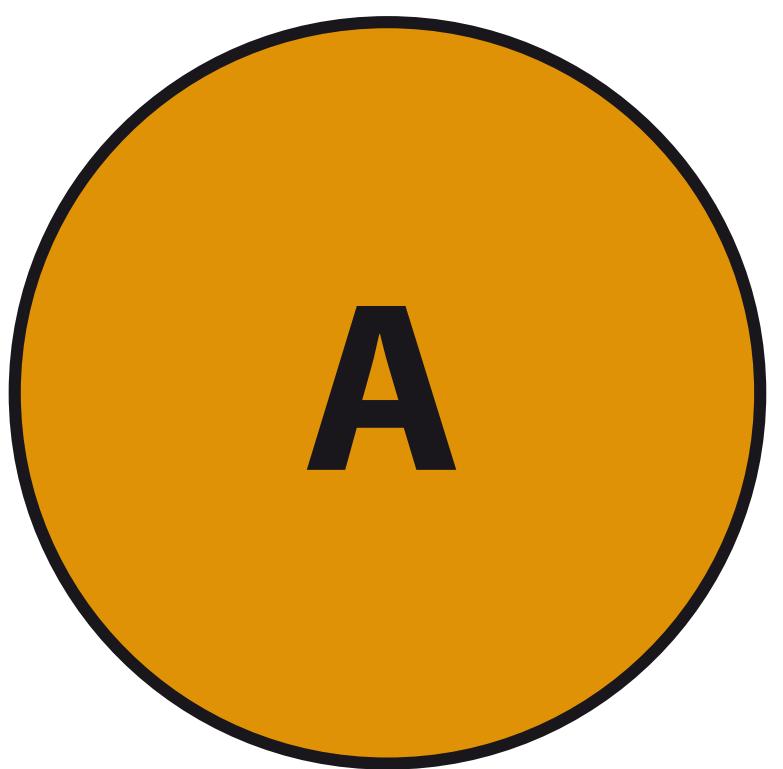
How do you find peers when you run a new node in the network?

How do peers announce their presence in the network?

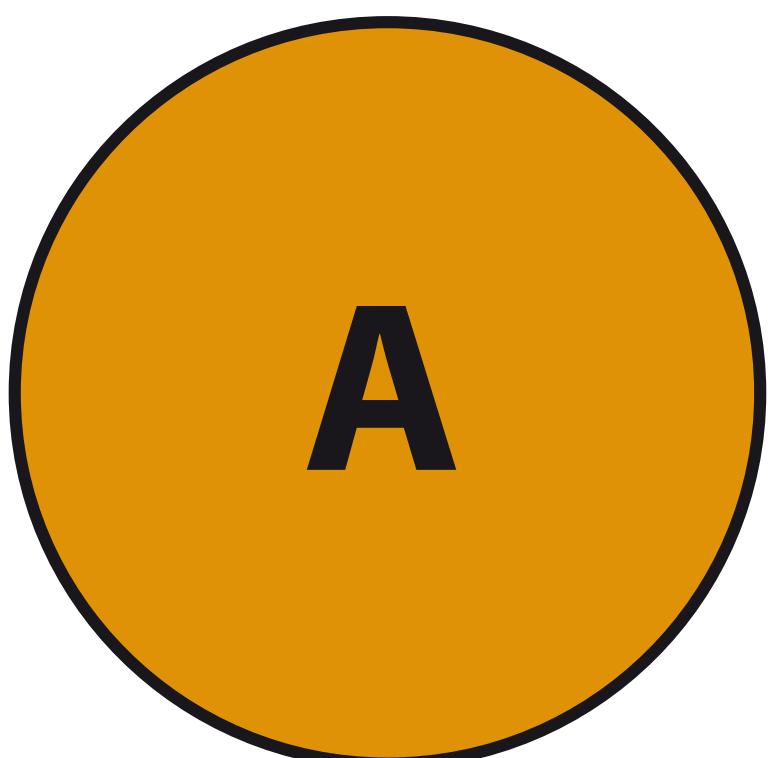
PEER DISCOVERY?



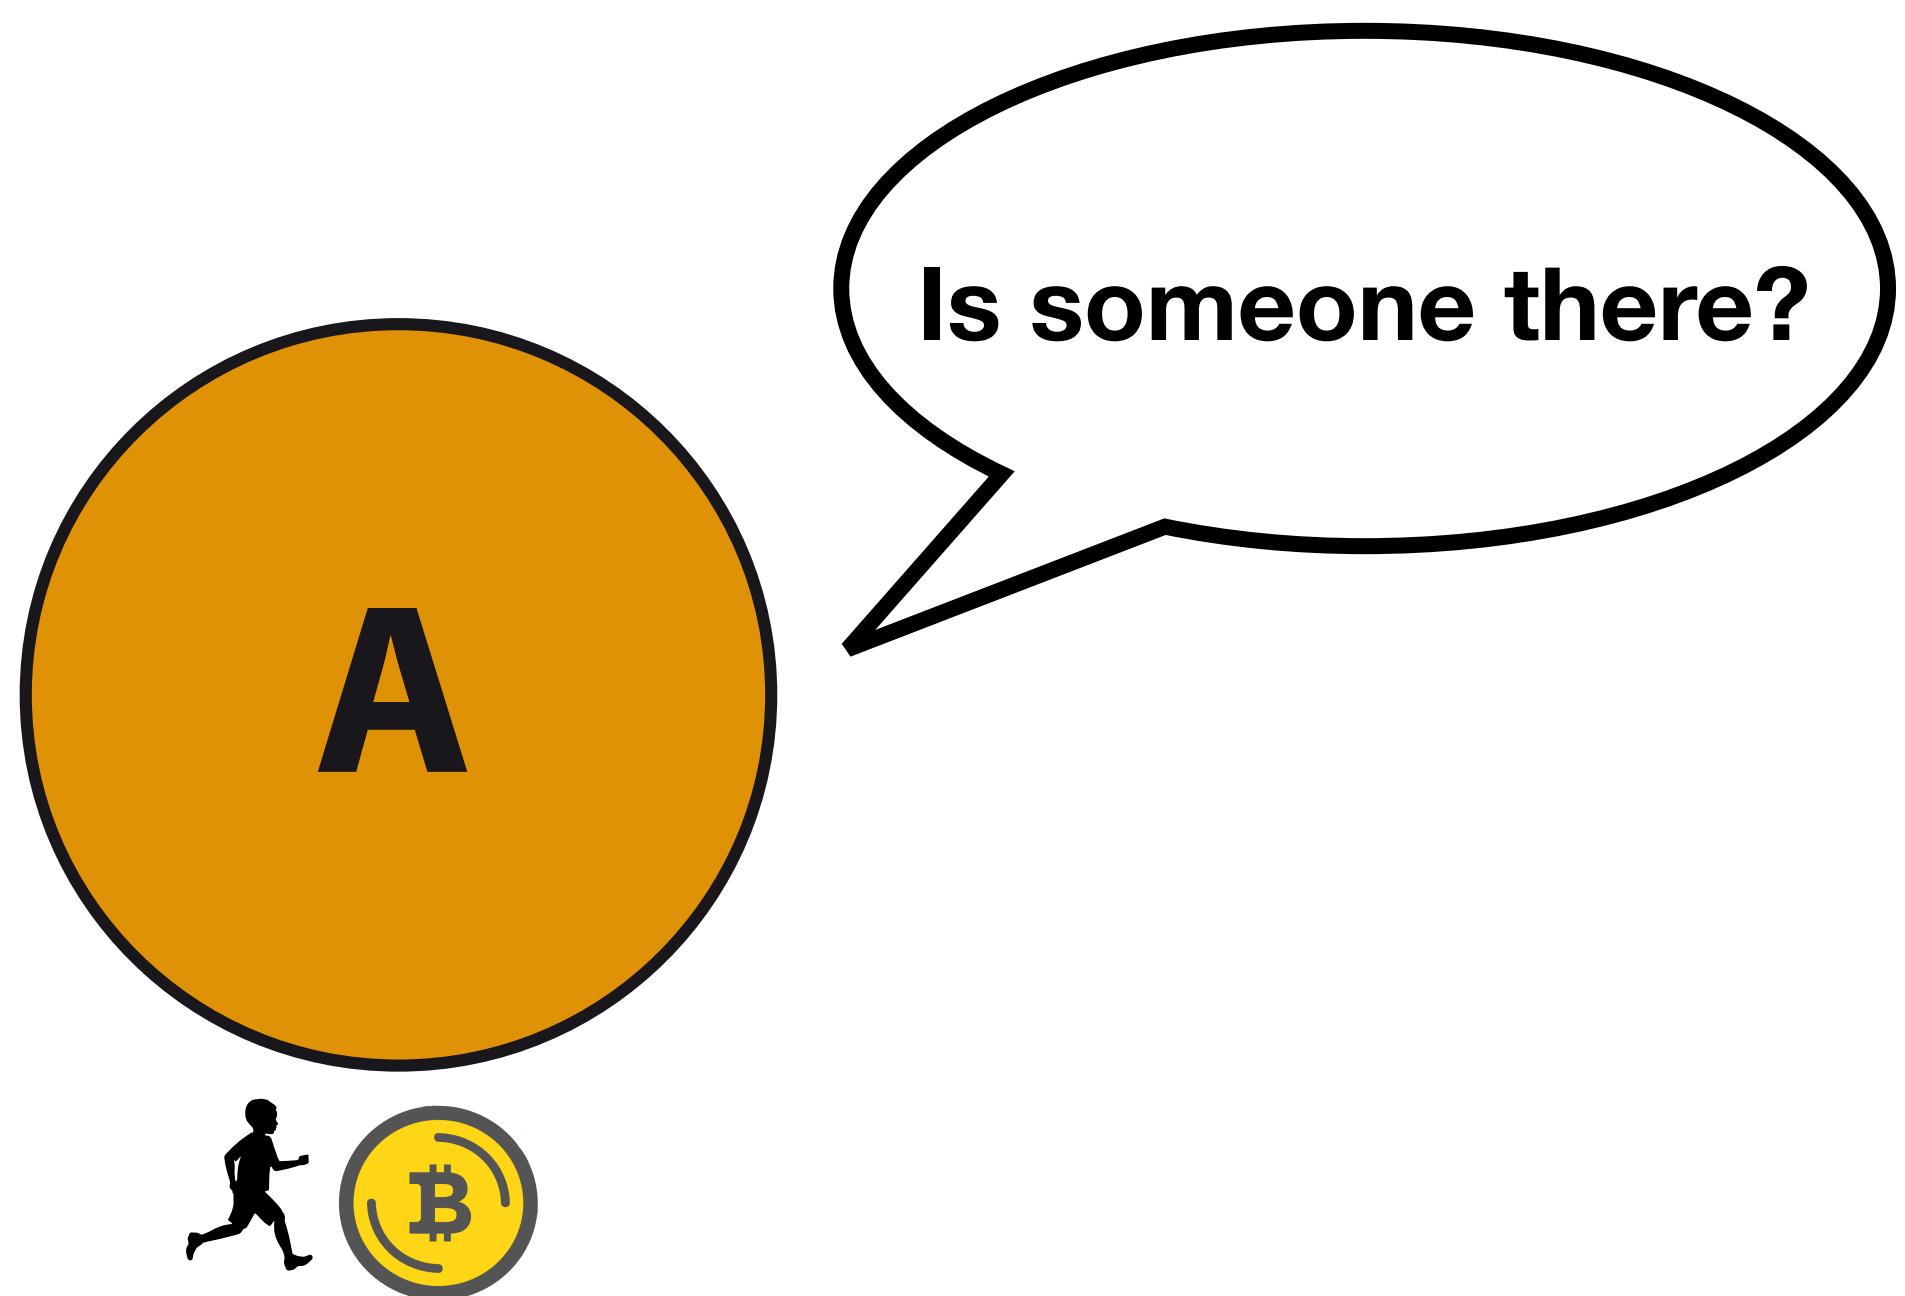
PEER DISCOVERY?



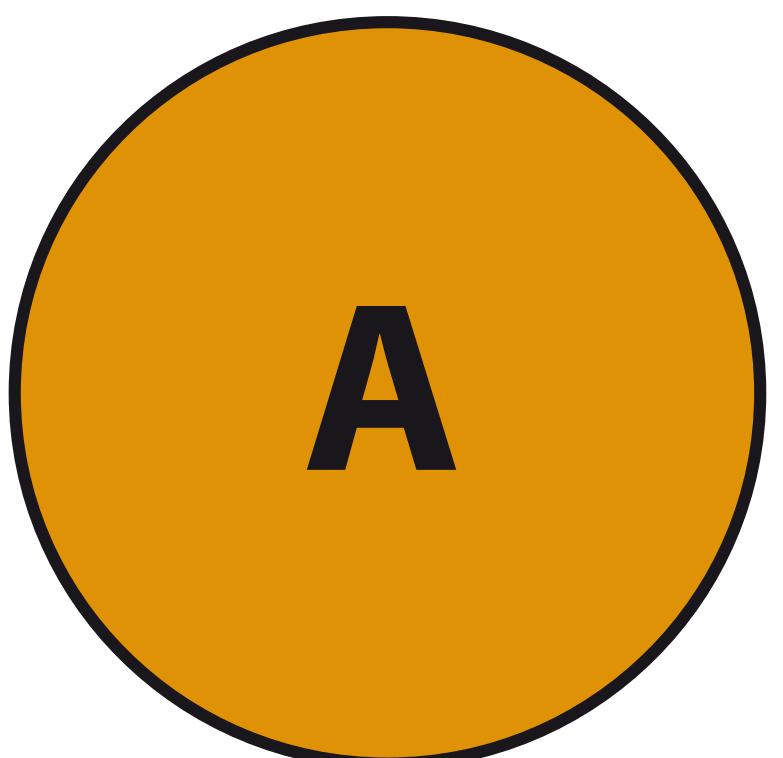
PEER DISCOVERY?



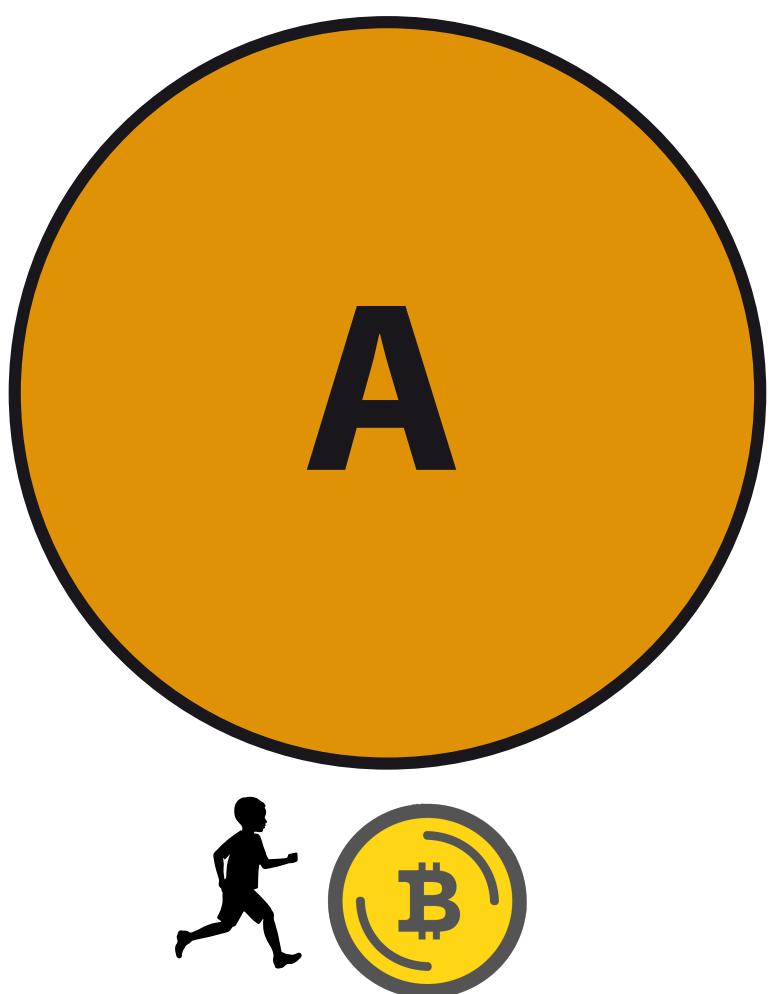
PEER DISCOVERY?



PEER DISCOVERY?



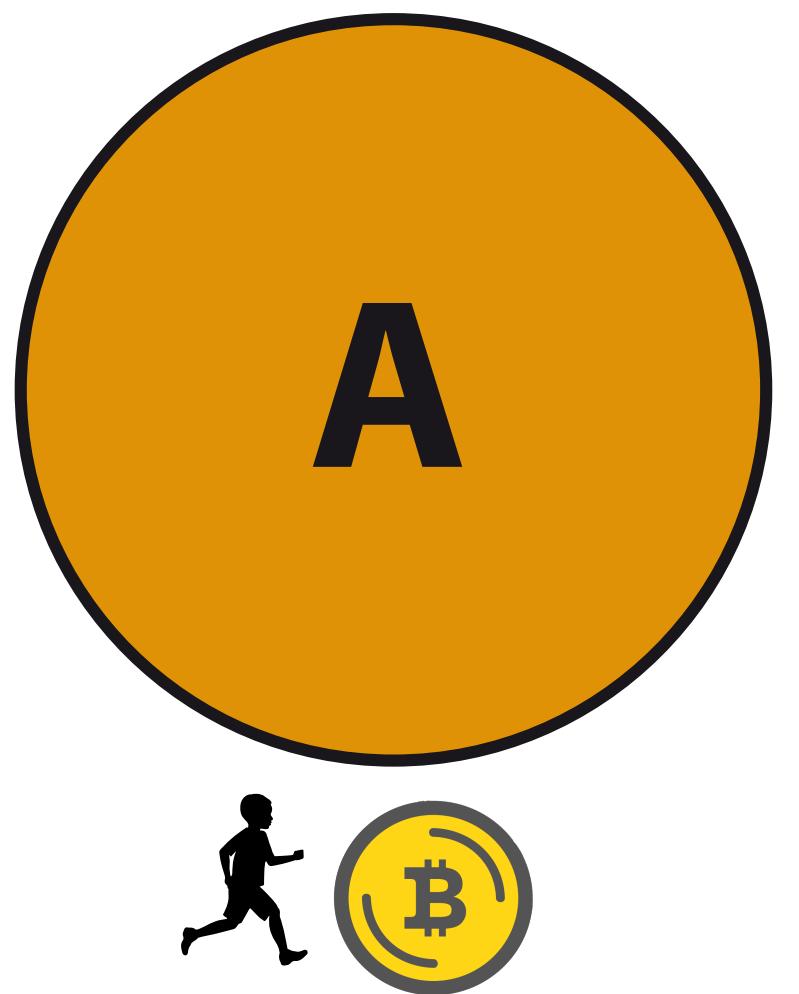
PEER DISCOVERY?



****tumbleweed****

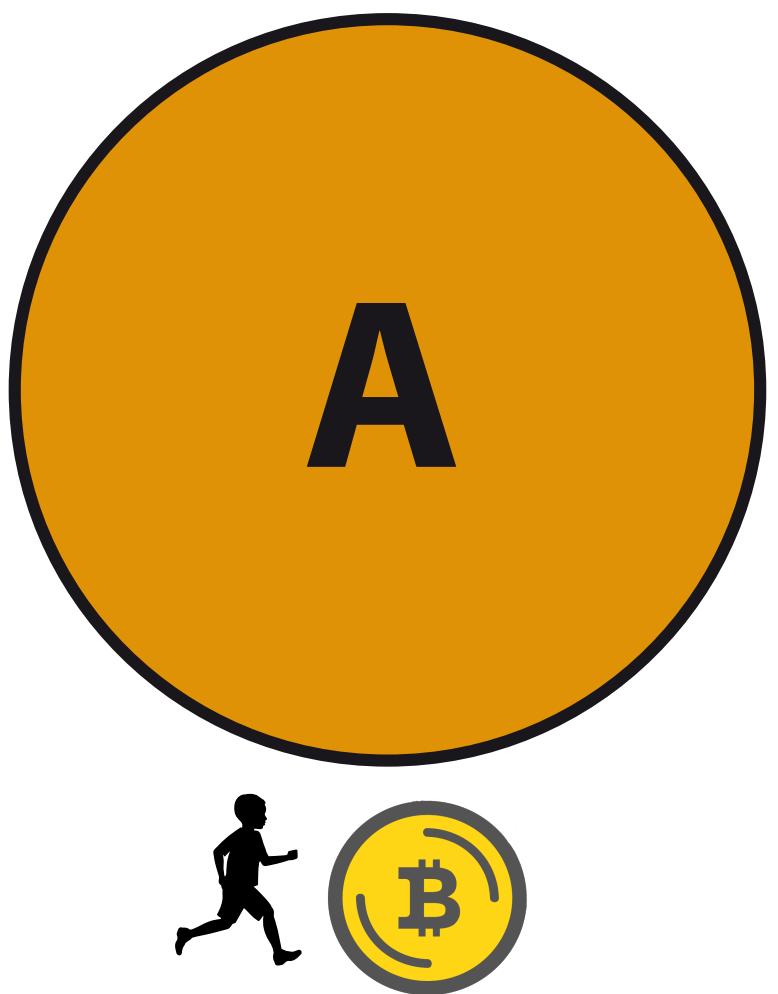


PEER DISCOVERY?



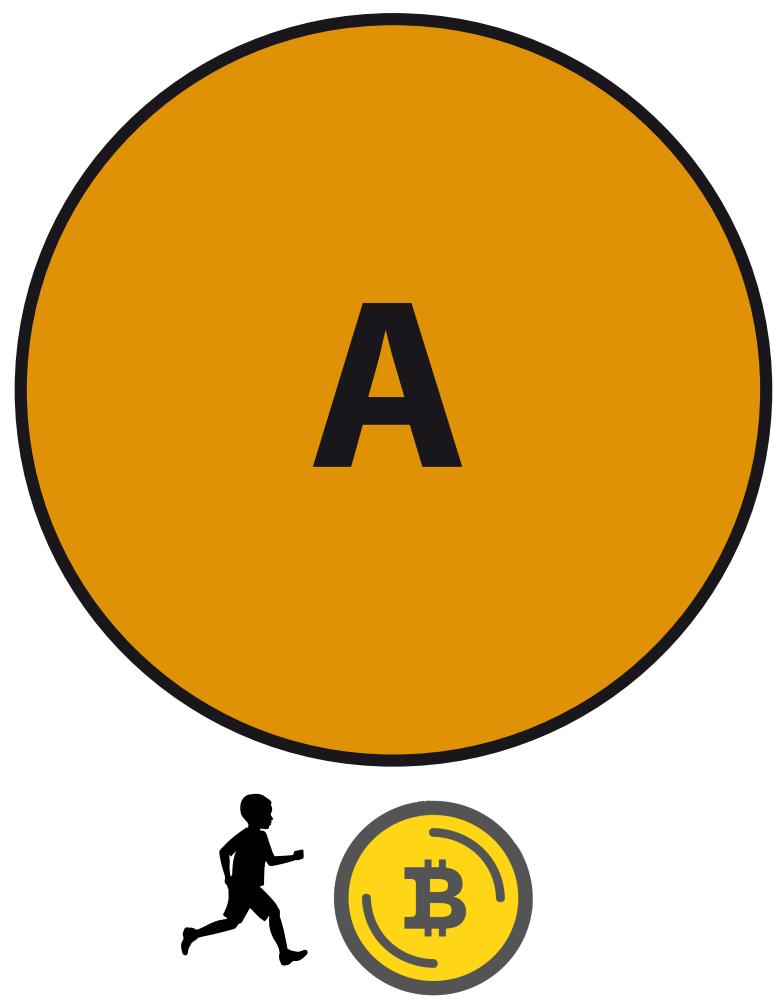
****tumbleweed****

PEER DISCOVERY?



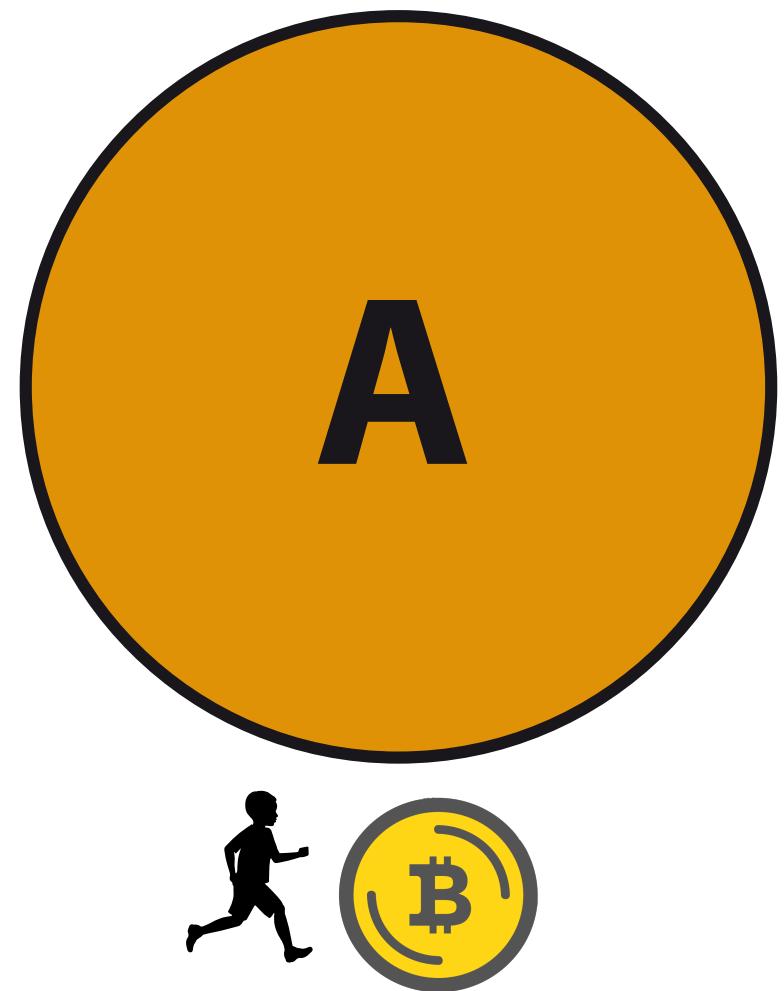
****tumbleweed****

PEER DISCOVERY?



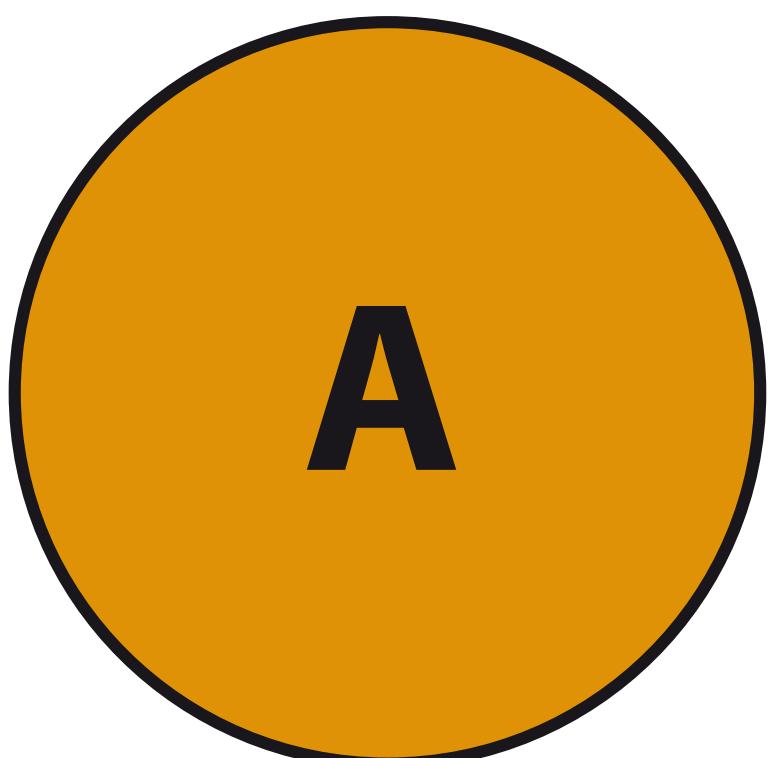
****tumbleweed****

PEER DISCOVERY?

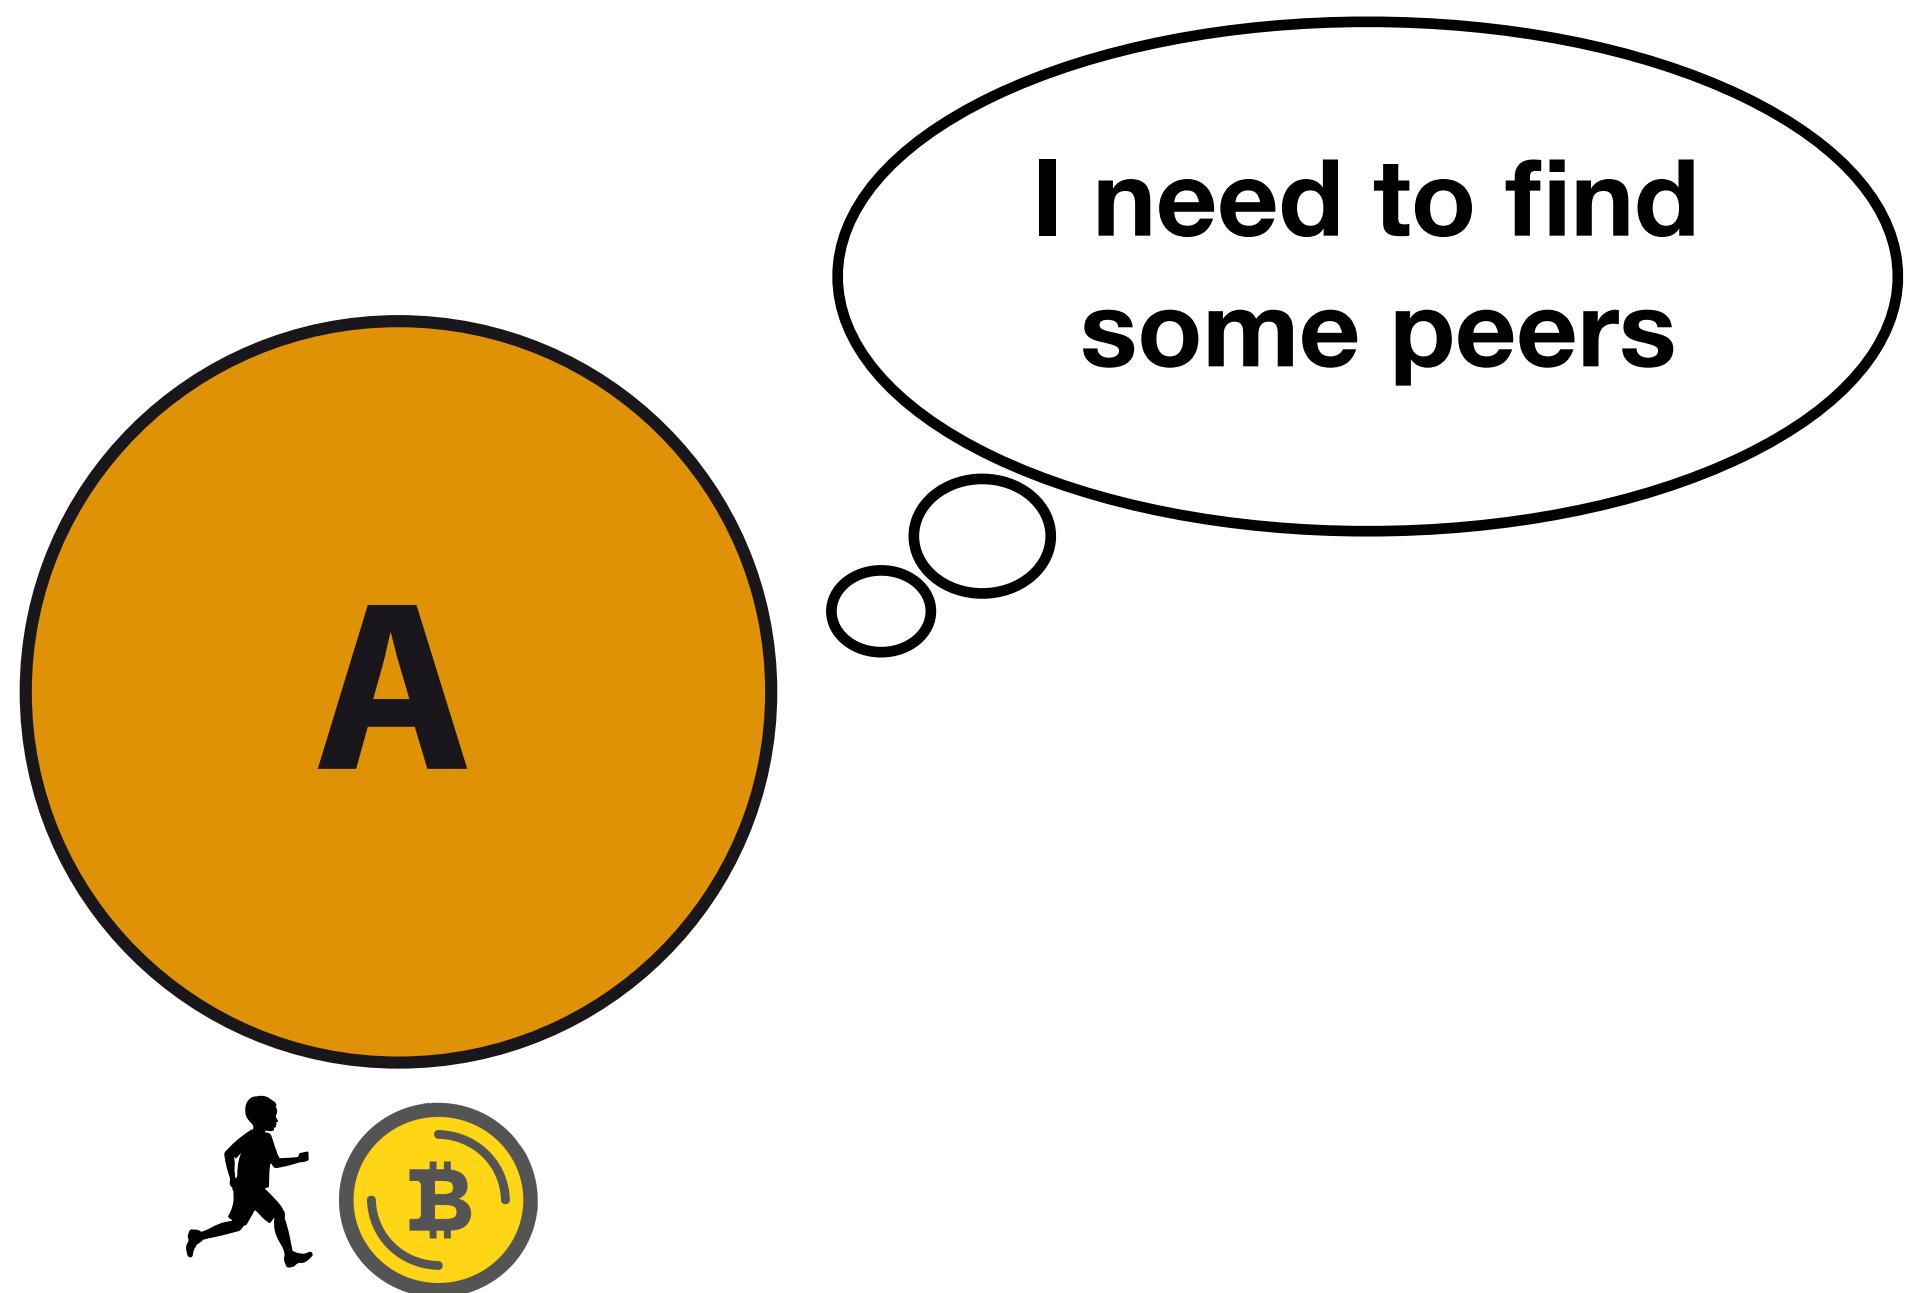


****tumbleweed****

PEER DISCOVERY?



PEER DISCOVERY?



P2P BOOTSTRAPPING

How do you find peers when you run a new node in the network?

How do peers announce their presence in the network?

P2P BOOTSTRAPPING

How do you find peers when you run a new node in the network?

How do peers announce their presence in the network?

Hardcoded trusted addresses / IRC bootstrapping / Trusted DNS seeds / etc

P2P FILE SHARING (1/2)

How to identify what other nodes are sharing (who knows what)?

How are files served?

P2P FILE SHARING (1/2)

How to identify what other nodes are sharing (who knows what)?

How are files served?

Announce / Request

P2P FILE SHARING (2/2)

Request paradigm: Files are requested by peers, so the network needs a lookup protocol to identify who knows what (e.g: DHT, trackers, etc)

Announce paradigm: Files are announced to peers, which will decide whether they would like a copy or not. No lookup protocol is required (e.g: gossip protocols)

P2P FILE SHARING (2/2)

Request paradigm: Files are requested by peers, so the network needs a lookup protocol to identify who knows what (e.g: DHT, trackers, etc)

Announce paradigm: Files are announced to peers, which will decide whether they would like a copy or not. No lookup protocol is required (e.g: gossip protocols)

What paradigm do cryptocurrency networks follow?

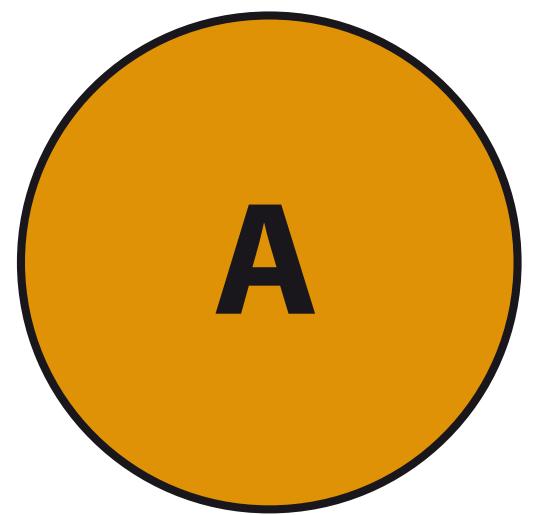
P2P FILE SHARING (2/2)

Request paradigm: Files are requested by peers, so the network needs a lookup protocol to identify who knows what (e.g: DHT, trackers, etc)

Announce paradigm: Files are announced to peers, which will decide whether they would like a copy or not. No lookup protocol is required (e.g: gossip protocols)

What paradigm do cryptocurrency networks follow? Announce

REQUEST PARADIGM (BitTorrent)

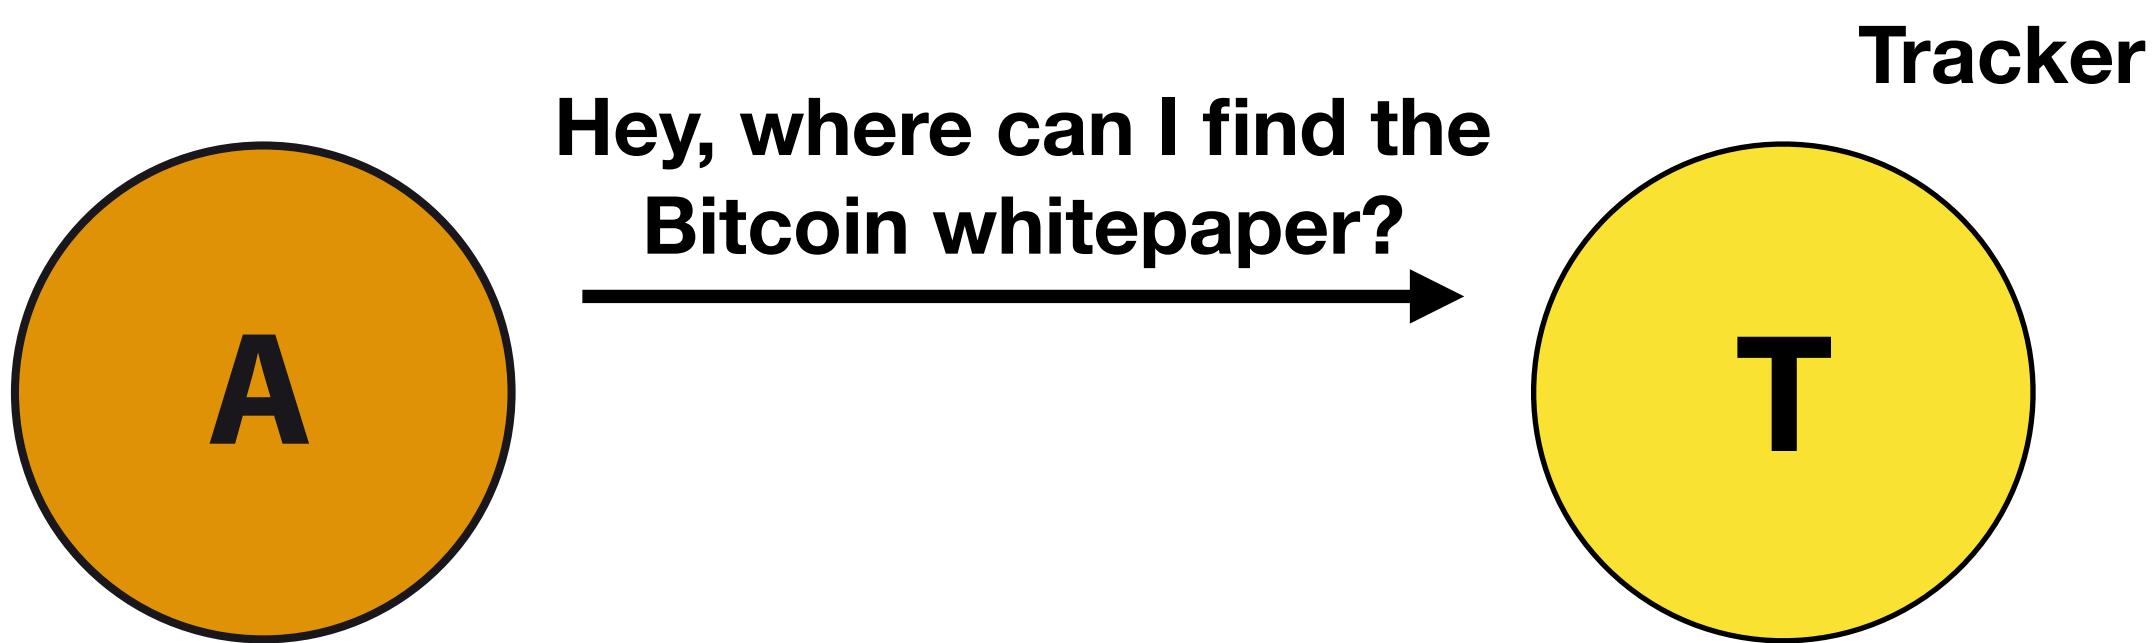


REQUEST PARADIGM (BitTorrent)



REQUEST PARADIGM (BitTorrent)

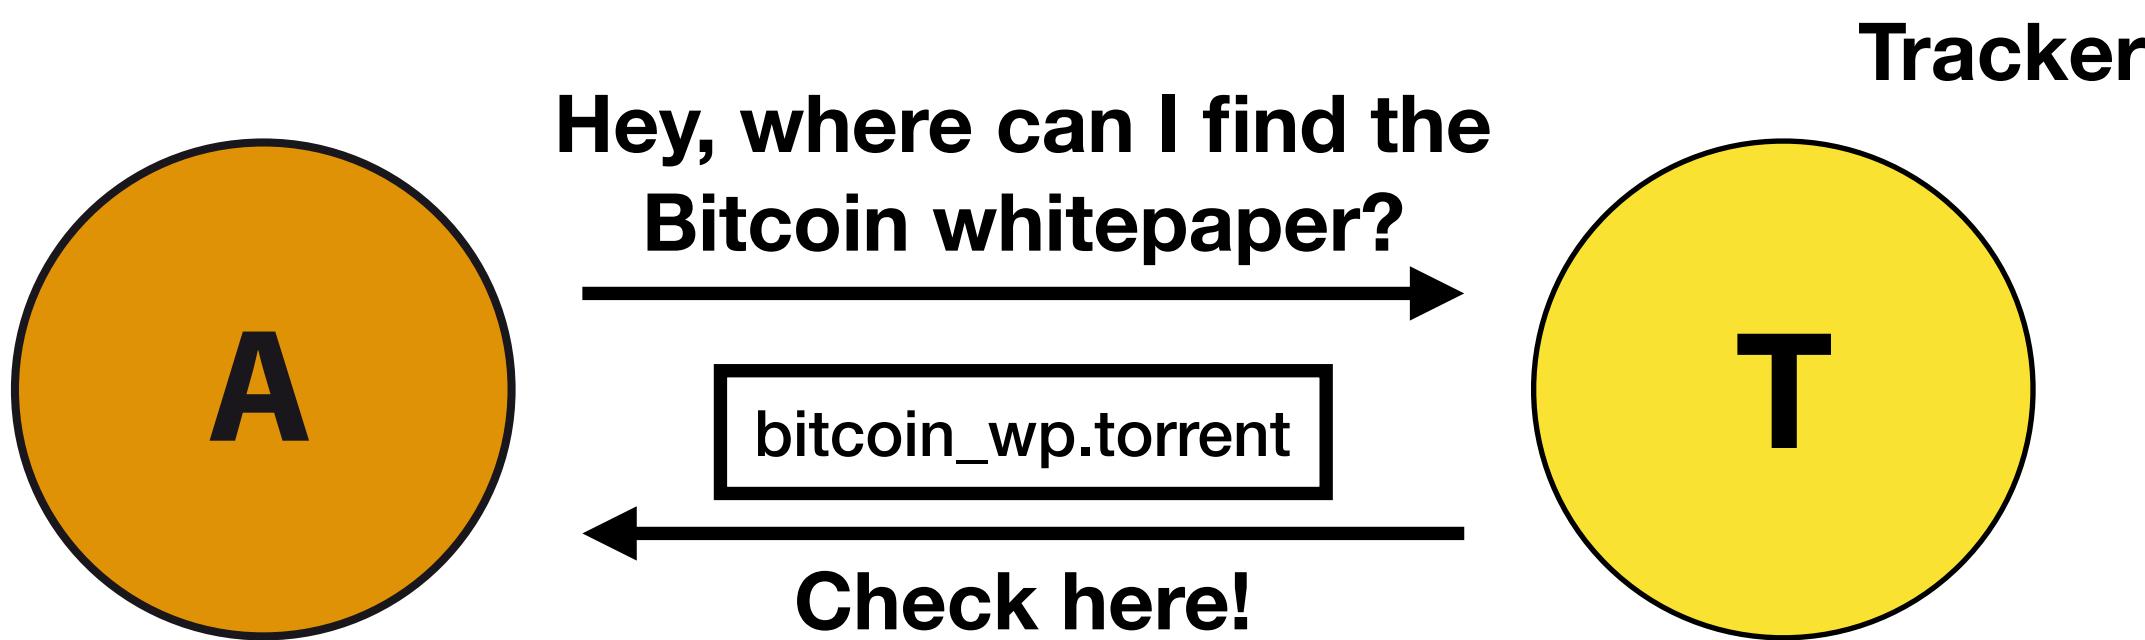
- Get file information from a tracker



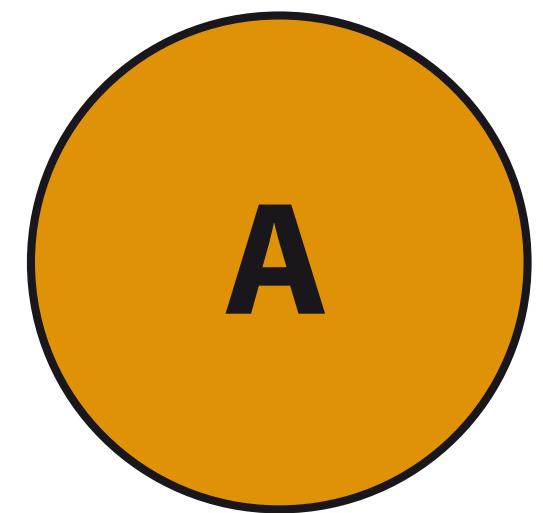
REQUEST PARADIGM (BitTorrent)



- Get file information from a tracker



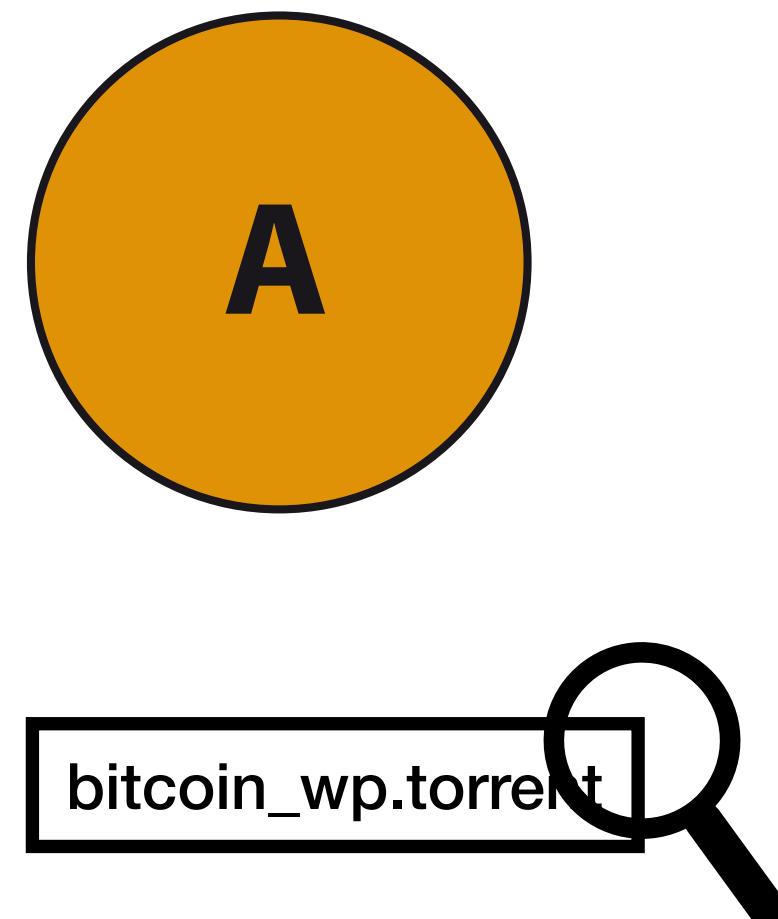
REQUEST PARADIGM (BitTorrent)



bitcoin_wp.torrent

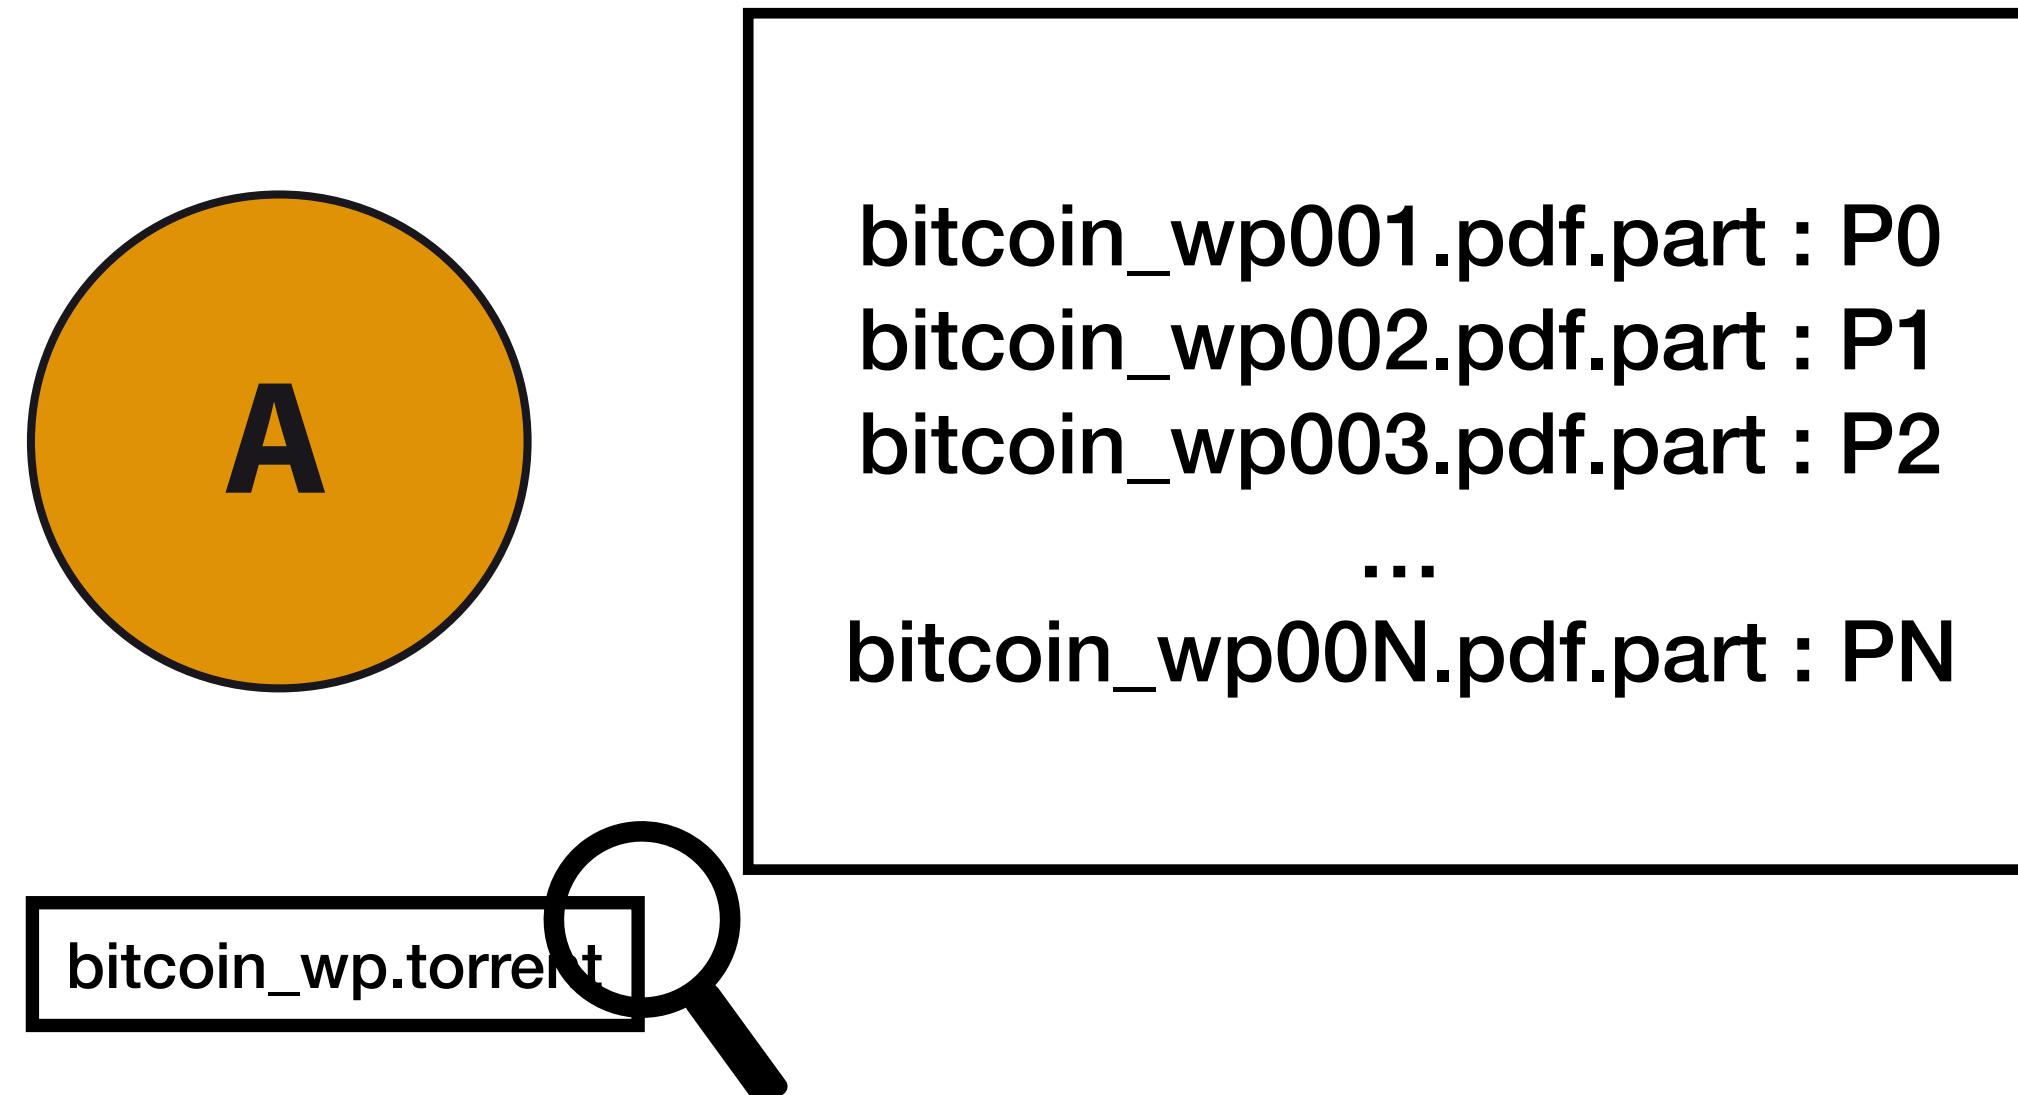
- Get file information from a tracker

REQUEST PARADIGM (BitTorrent)



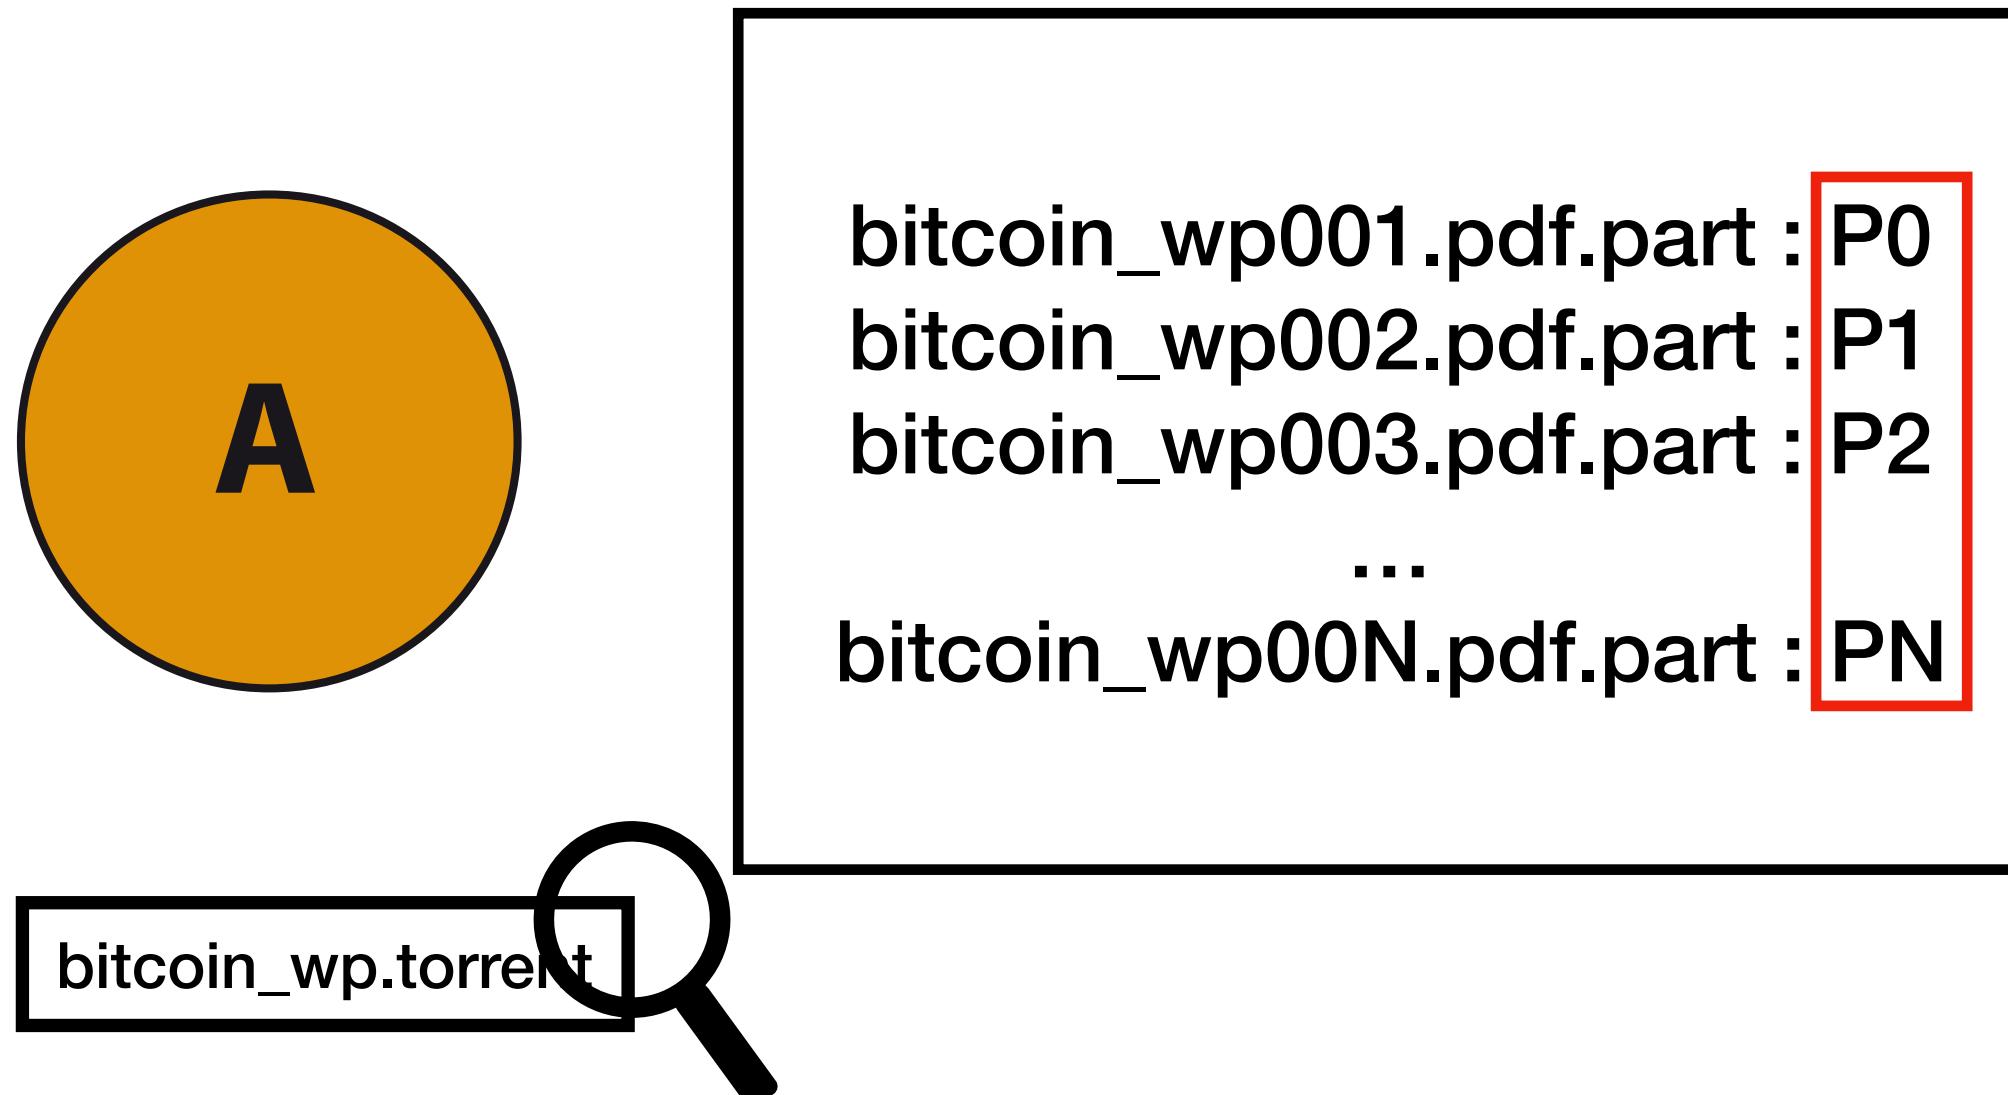
- Get file information from a tracker
- Check the .torrent file

REQUEST PARADIGM (BitTorrent)



- Get file information from a tracker
- Check the .torrent file

REQUEST PARADIGM (BitTorrent)



- Get file information from a tracker
- Check the .torrent file
- Connect to peers and retrieve the file parts

ANNOUNCE VS REQUEST

Why would a request paradigm (like the one we just saw) not work for cryptocurrency networks?

ANNOUNCE VS REQUEST

Why would a request paradigm (like the one we just saw) not work for cryptocurrency networks?

New items (transactions and blocks) can be created by others, so we can't know about them if they are not offered

ANNOUNCE VS REQUEST

Why would a request paradigm (like the one we just saw) not work for cryptocurrency networks?

New items (transactions and blocks) can be created by others, so we can't know about them if they are not offered

What information should a node know about the system?

ANNOUNCE VS REQUEST

Why would a request paradigm (like the one we just saw) not work for cryptocurrency networks?

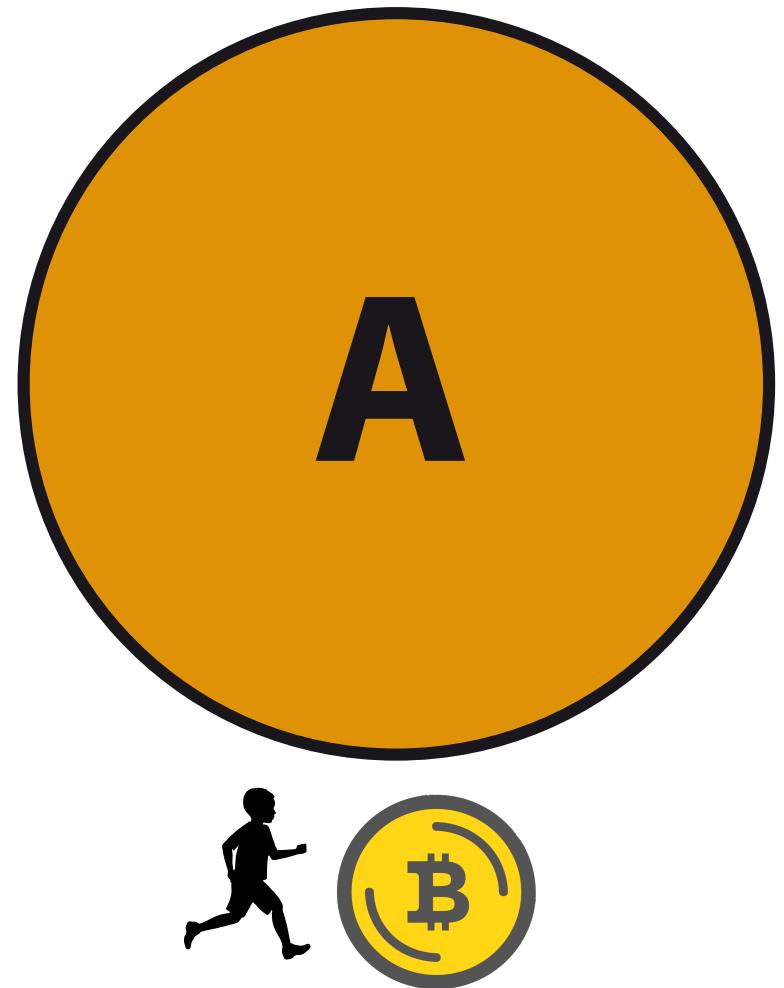
New items (transactions and blocks) can be created by others, so we can't know about them if they are not offered

What information should a node know about the system?

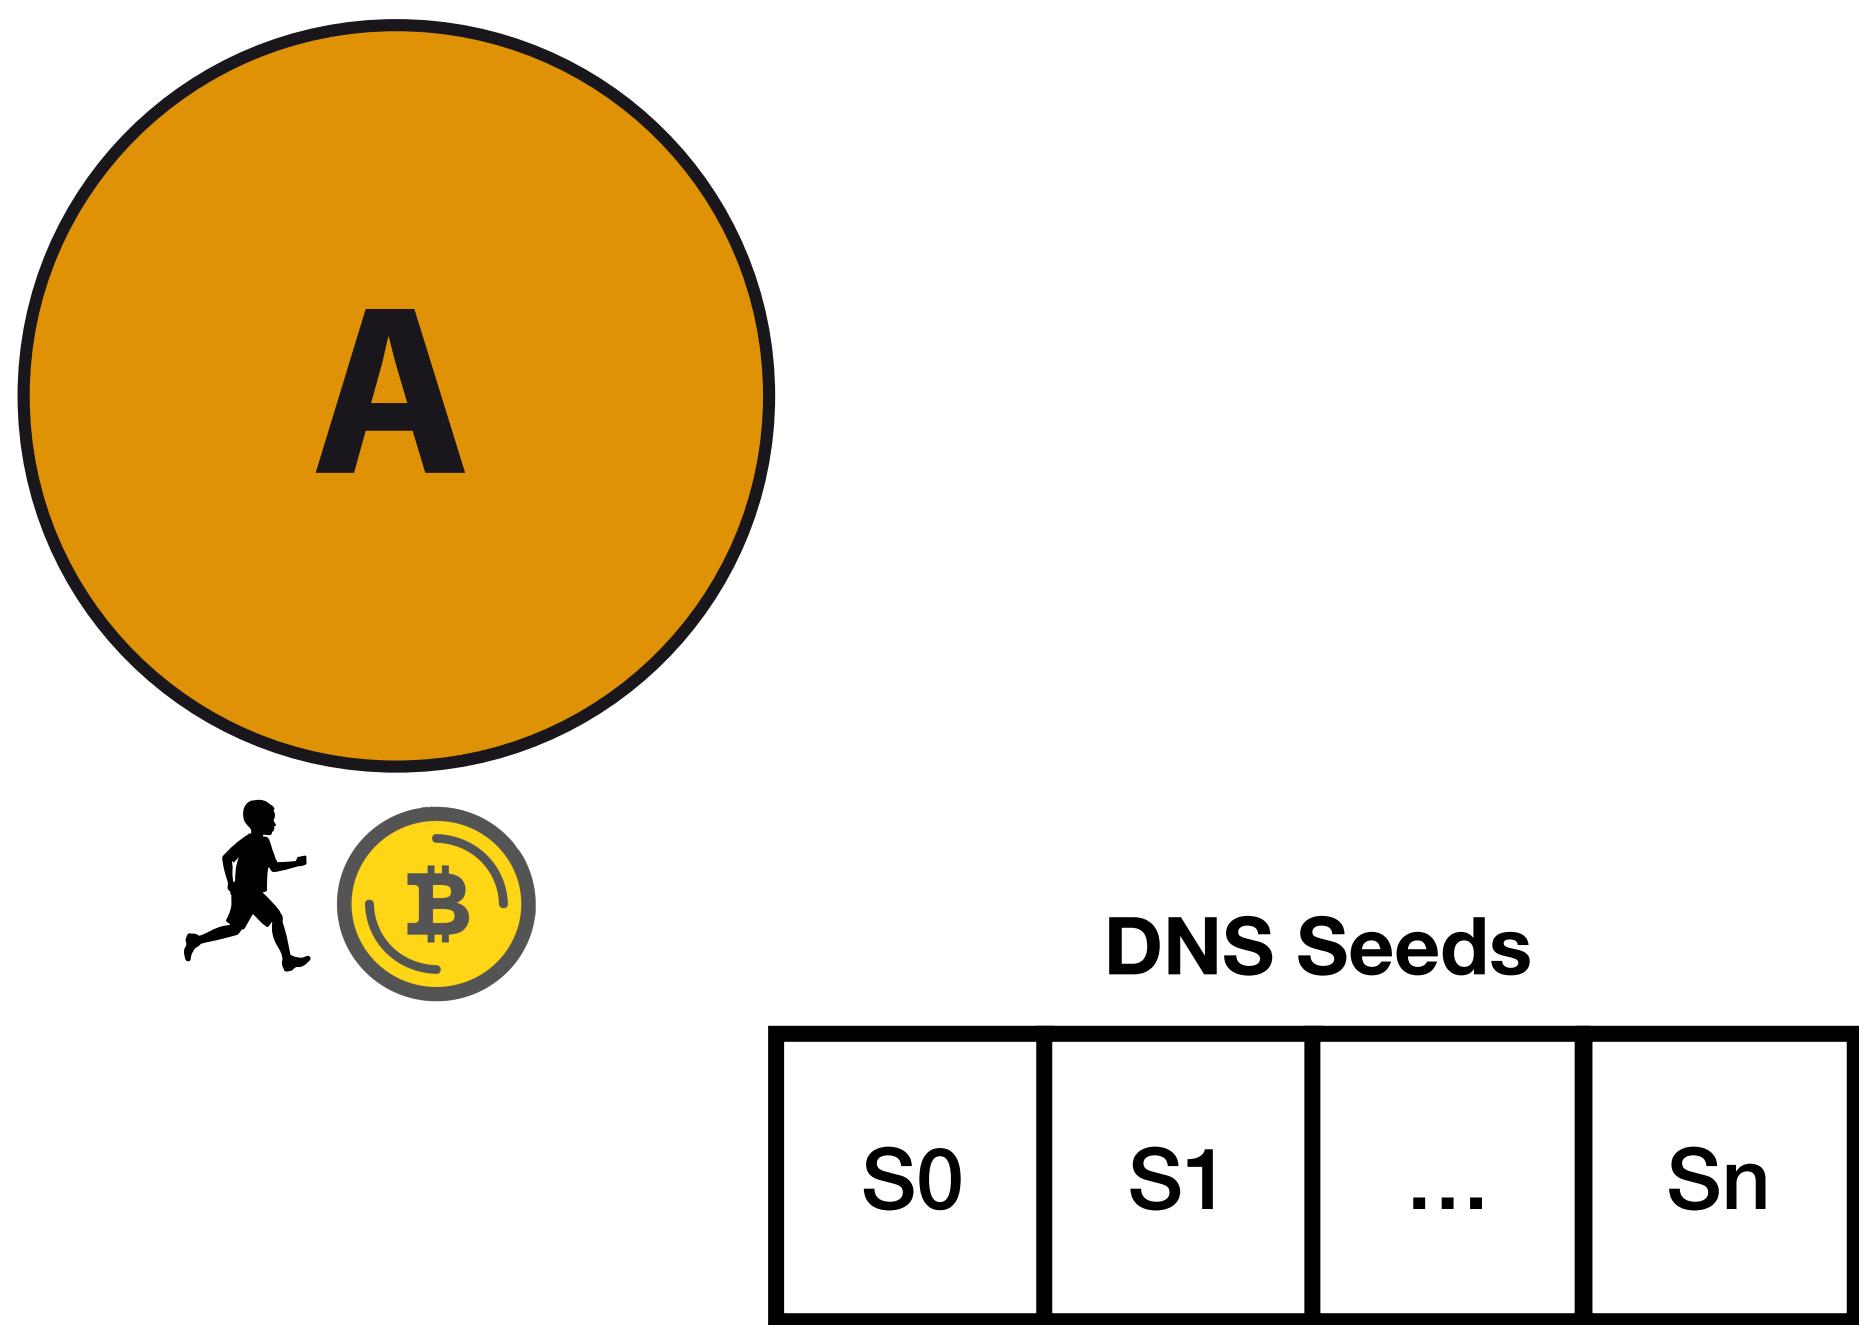
A (full) node needs all the information in order to validate new items

Node bootstrapping and peer discovery

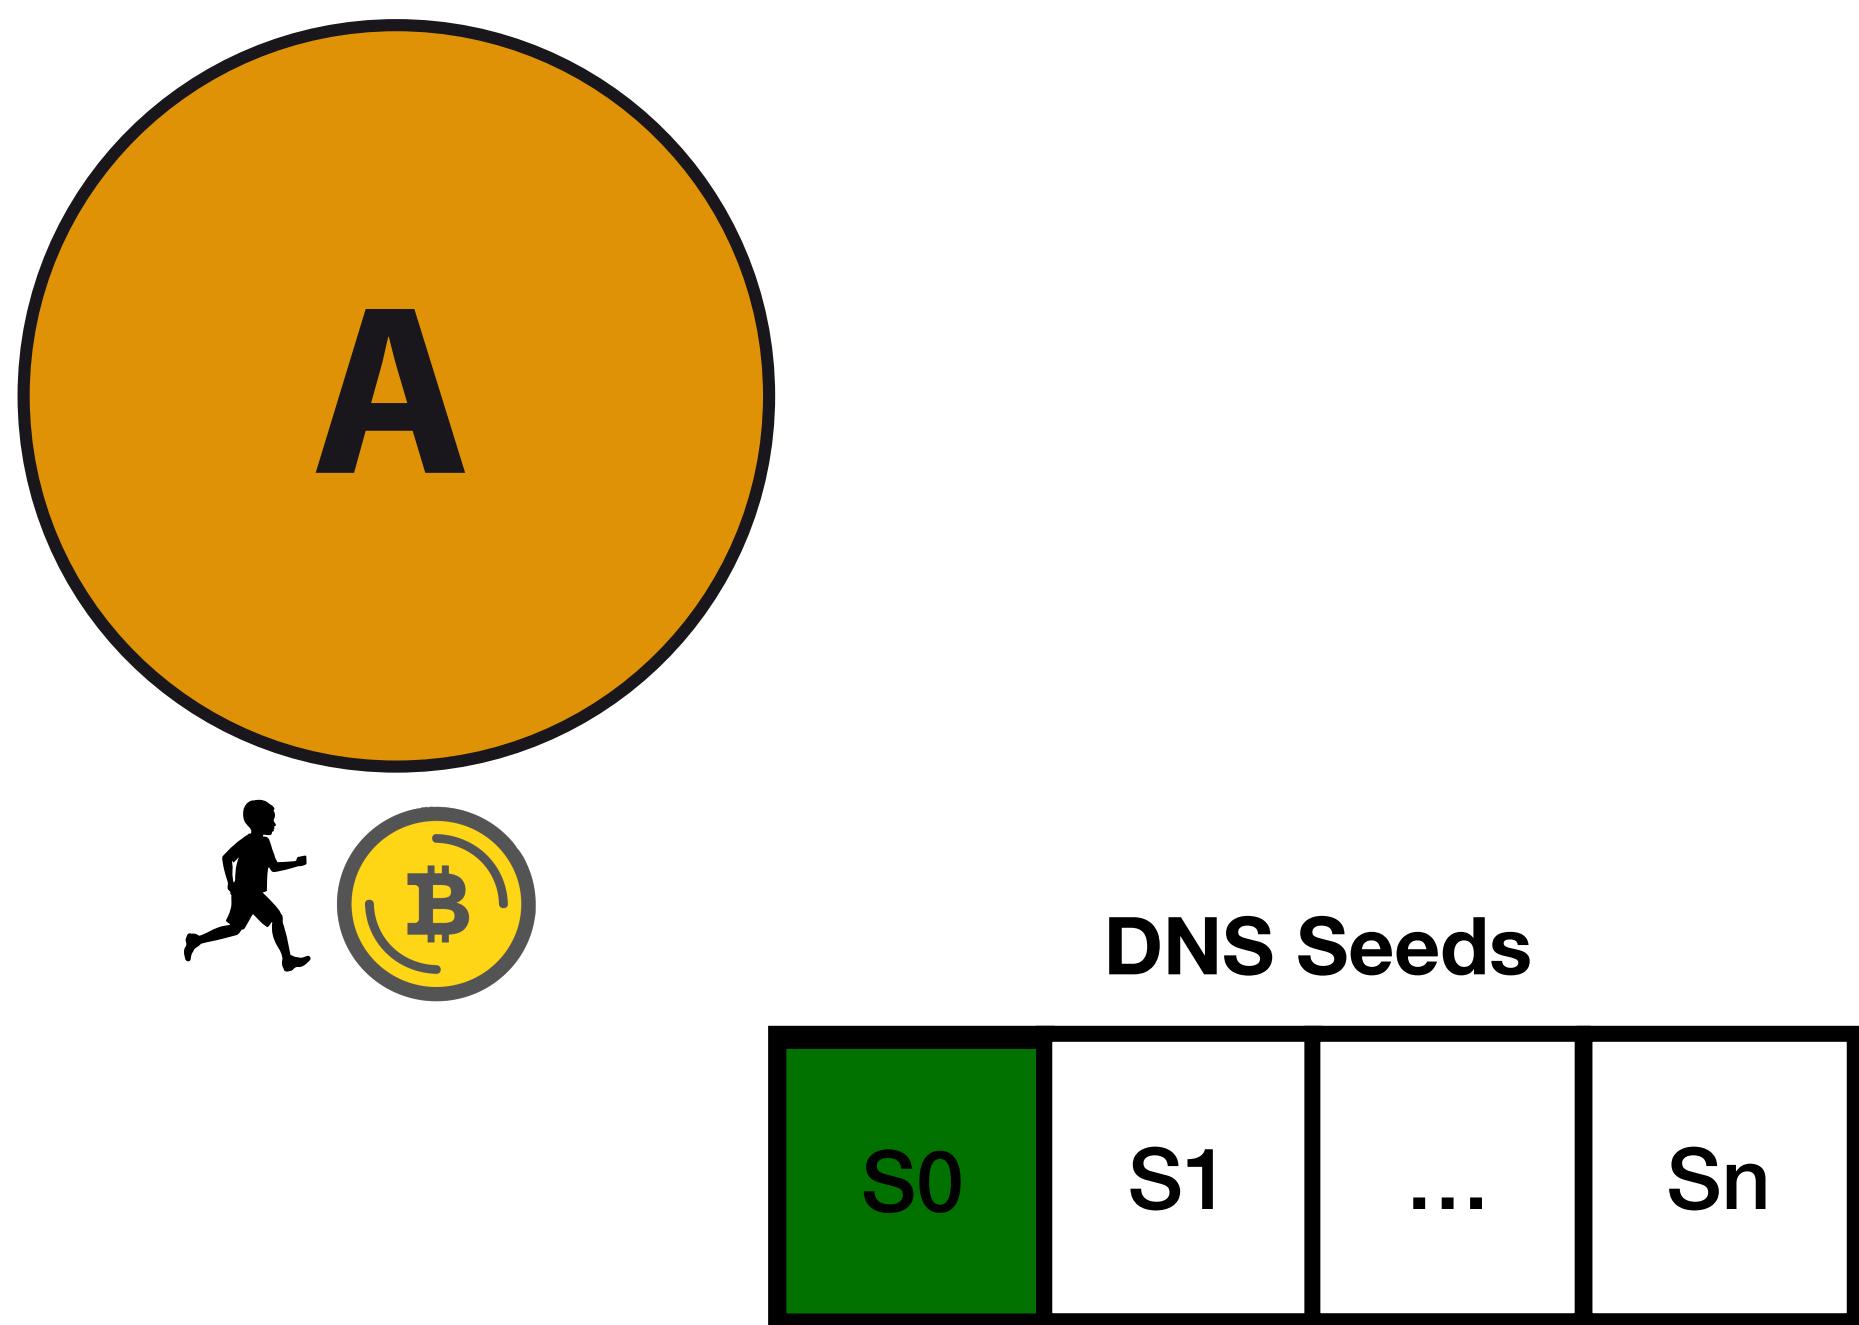
BITCOIN PEER DISCOVERY



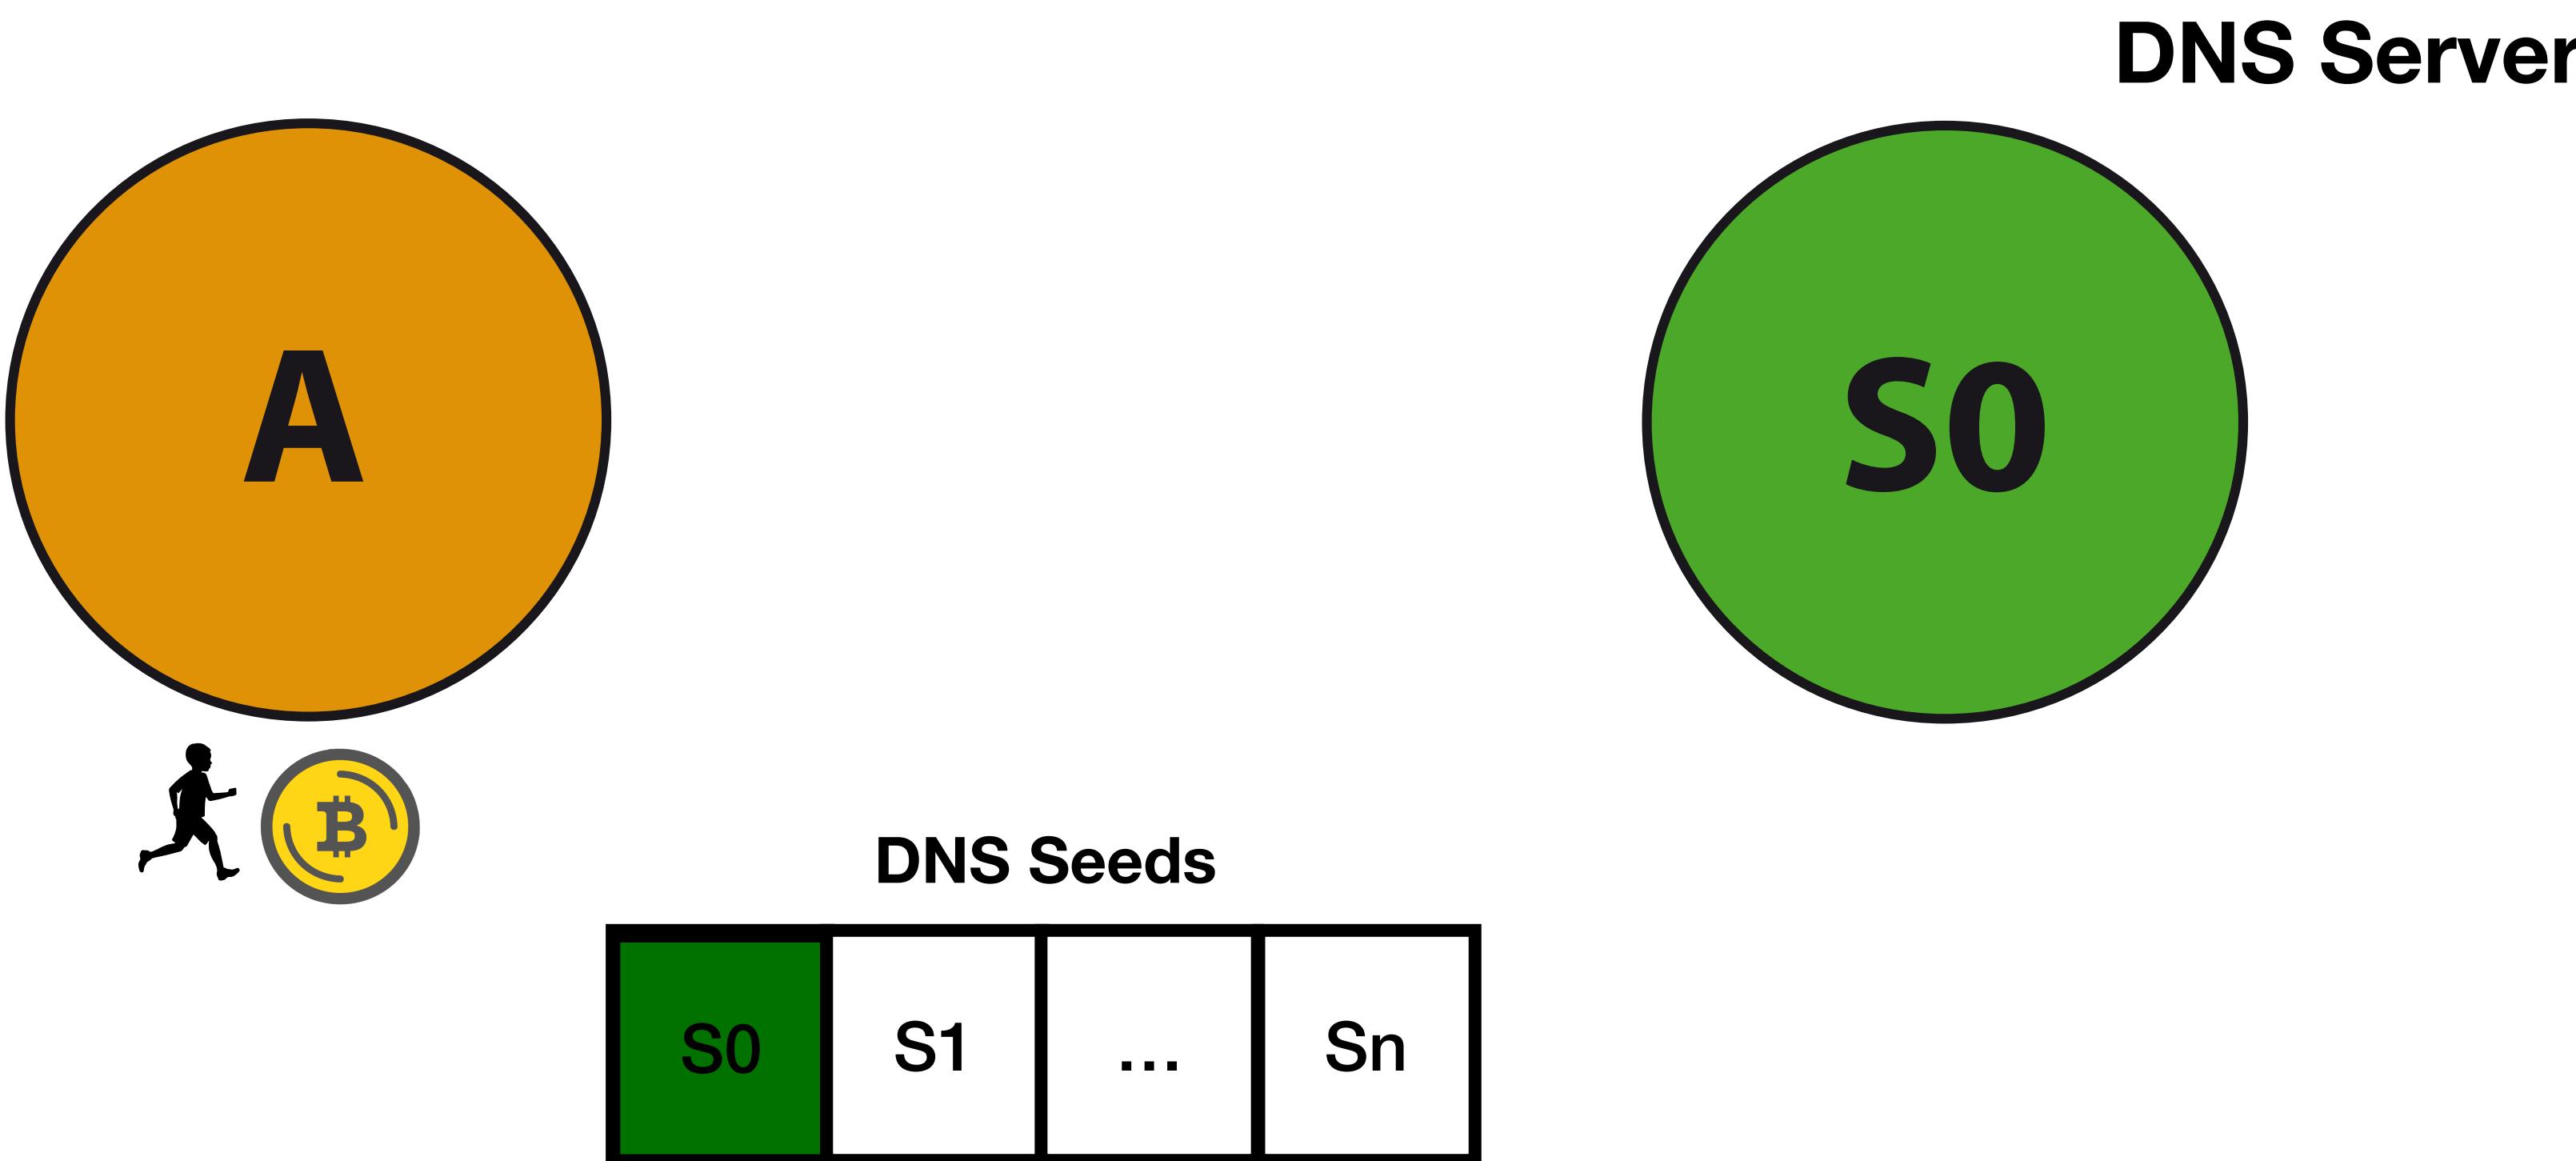
BITCOIN PEER DISCOVERY



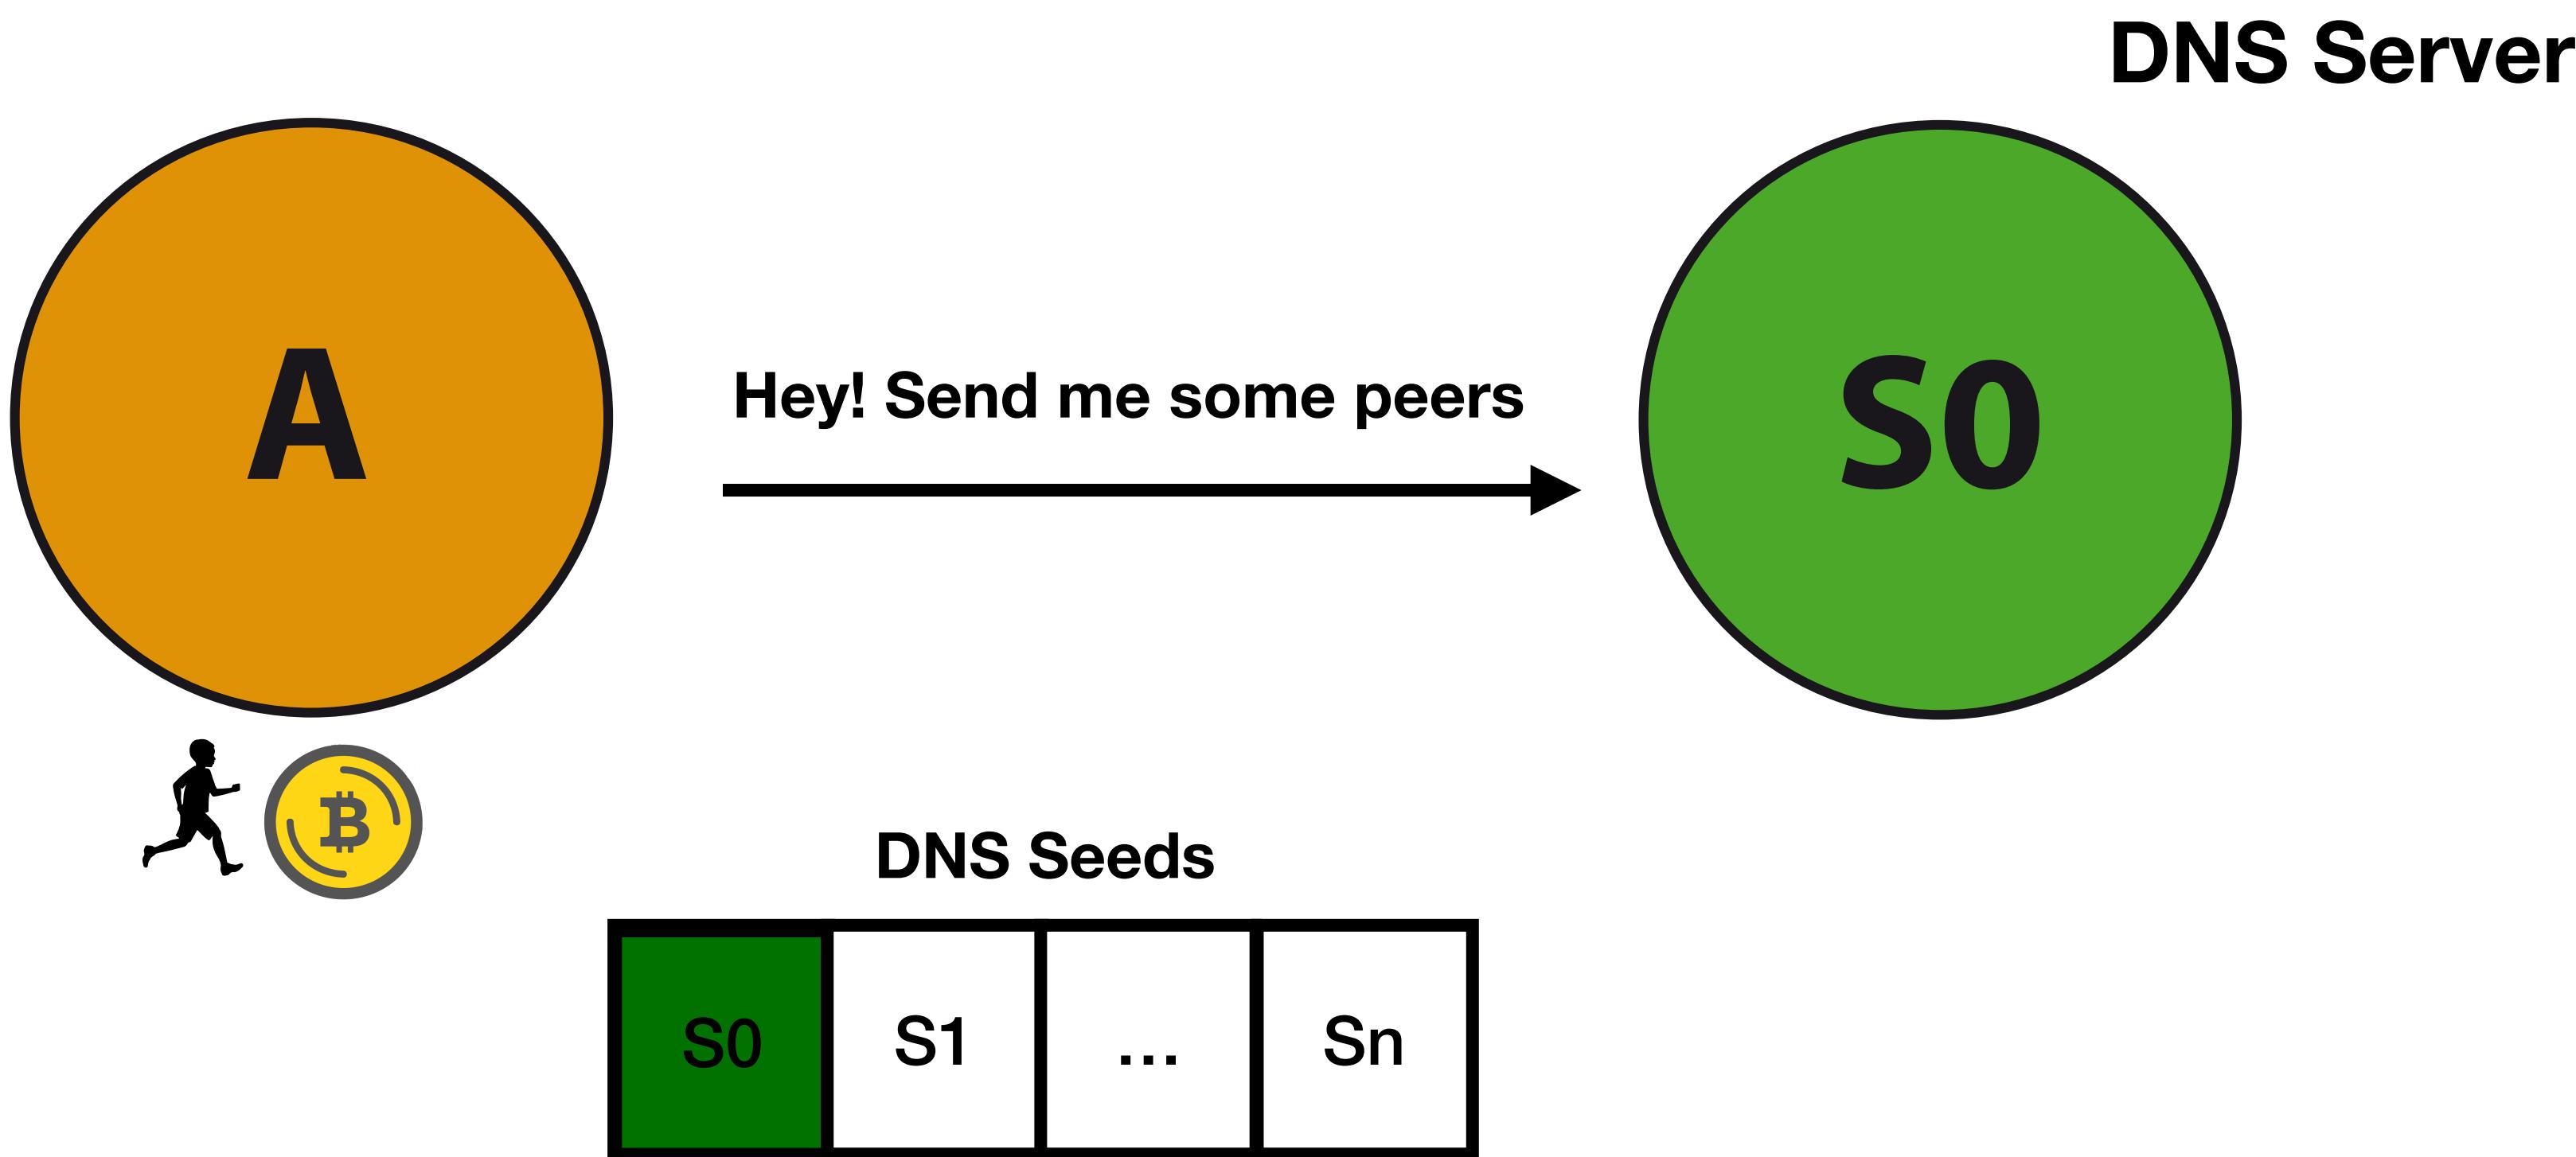
BITCOIN PEER DISCOVERY



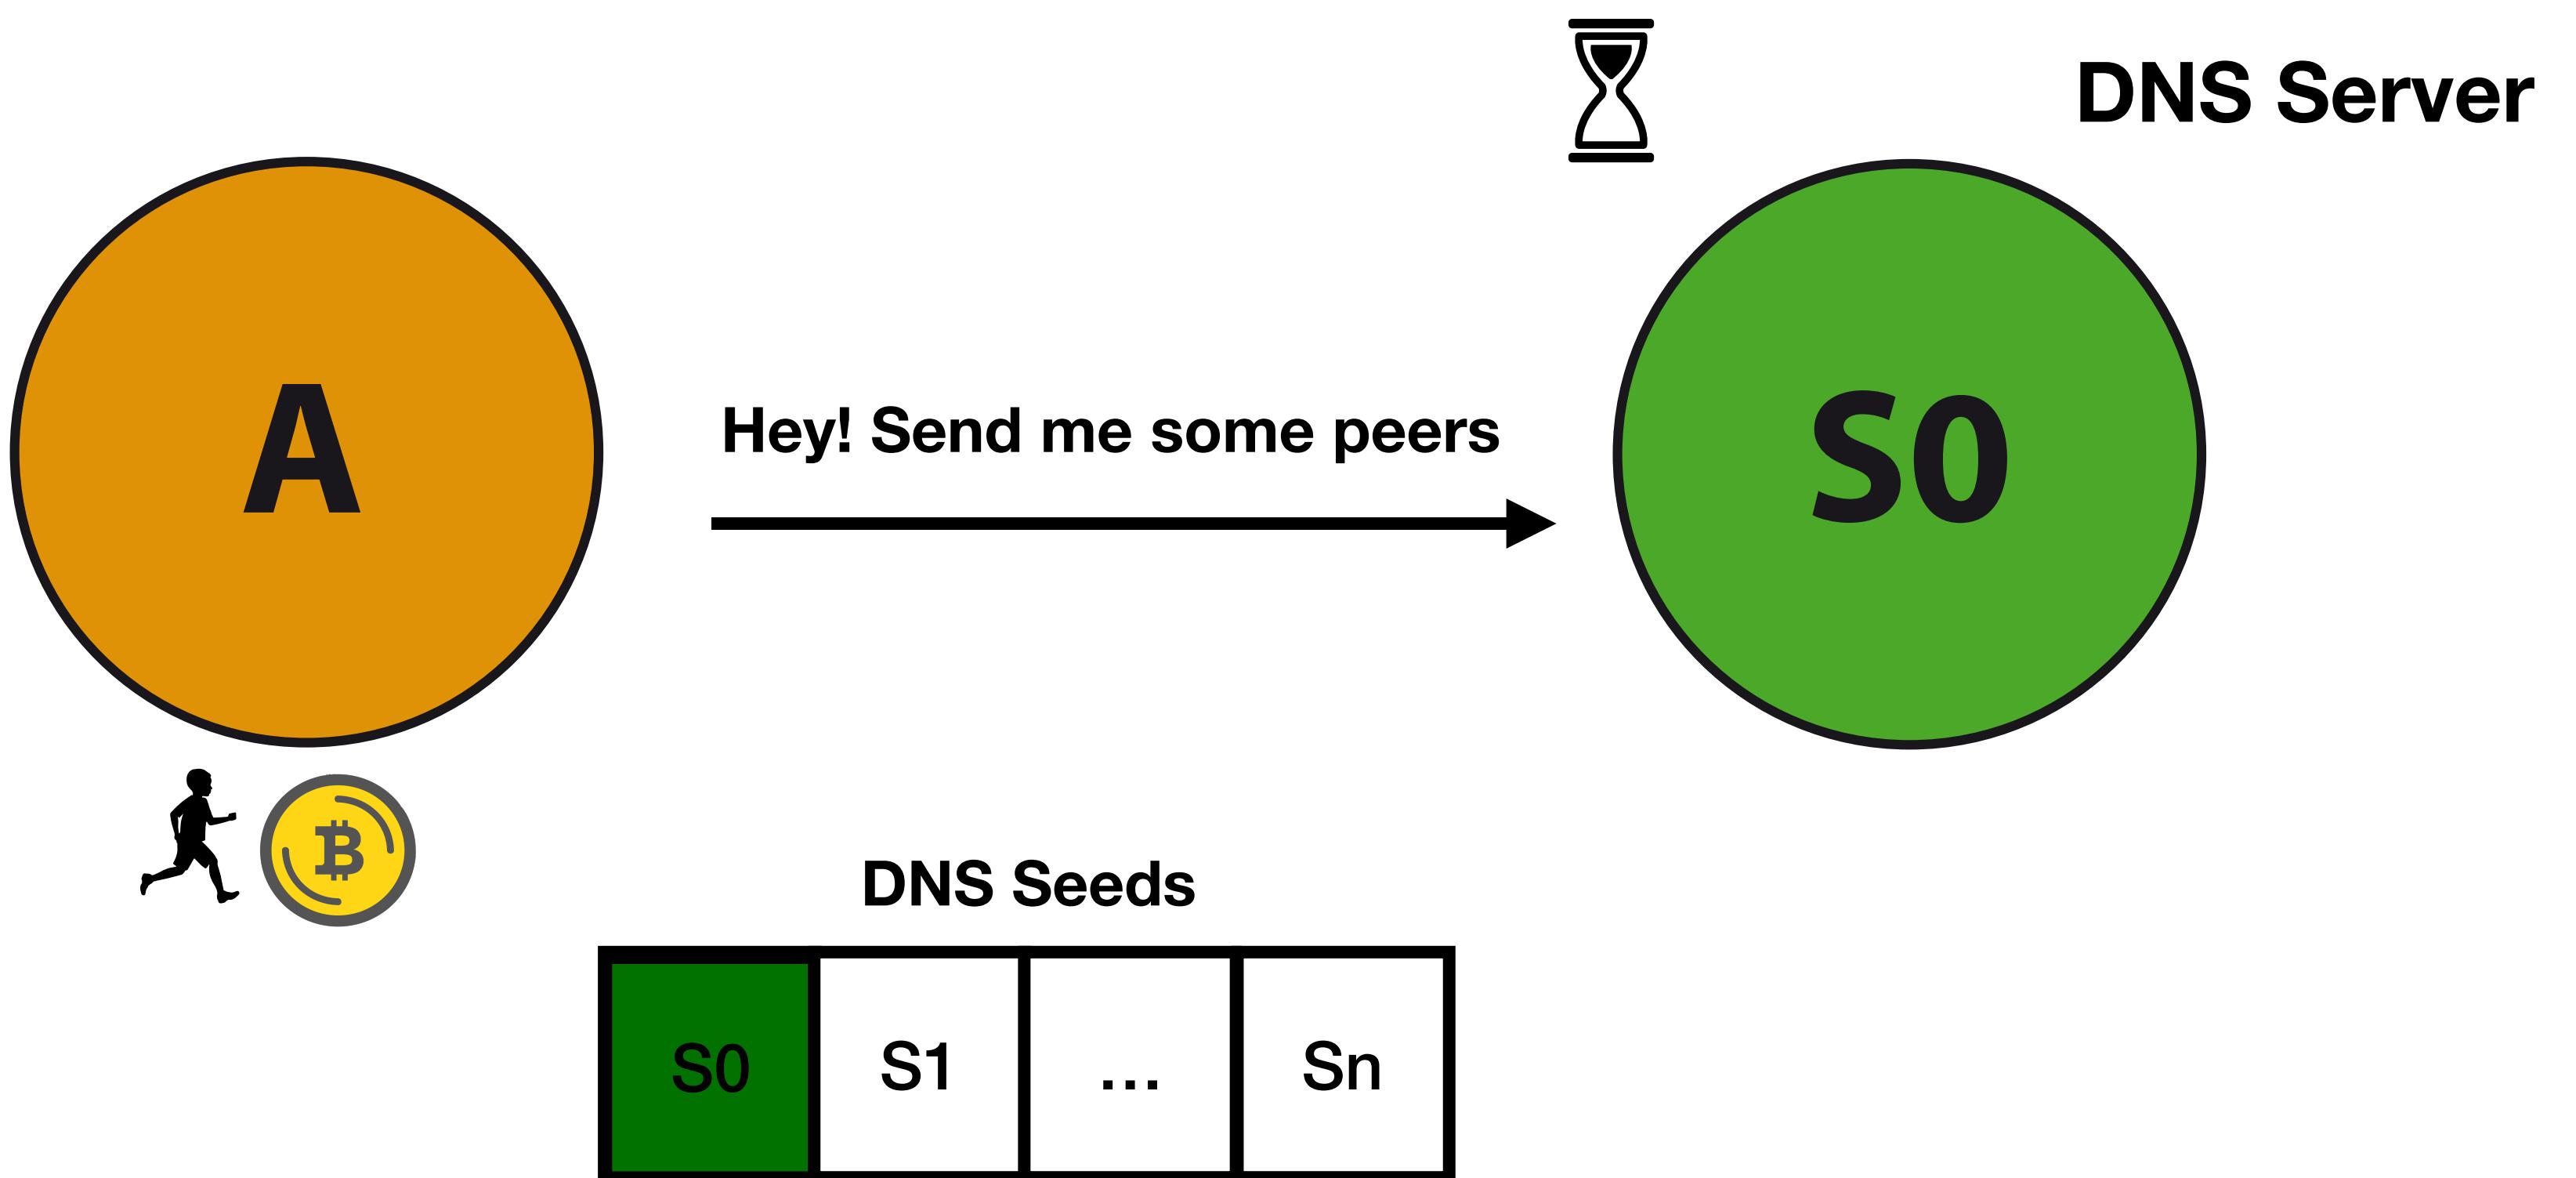
BITCOIN PEER DISCOVERY



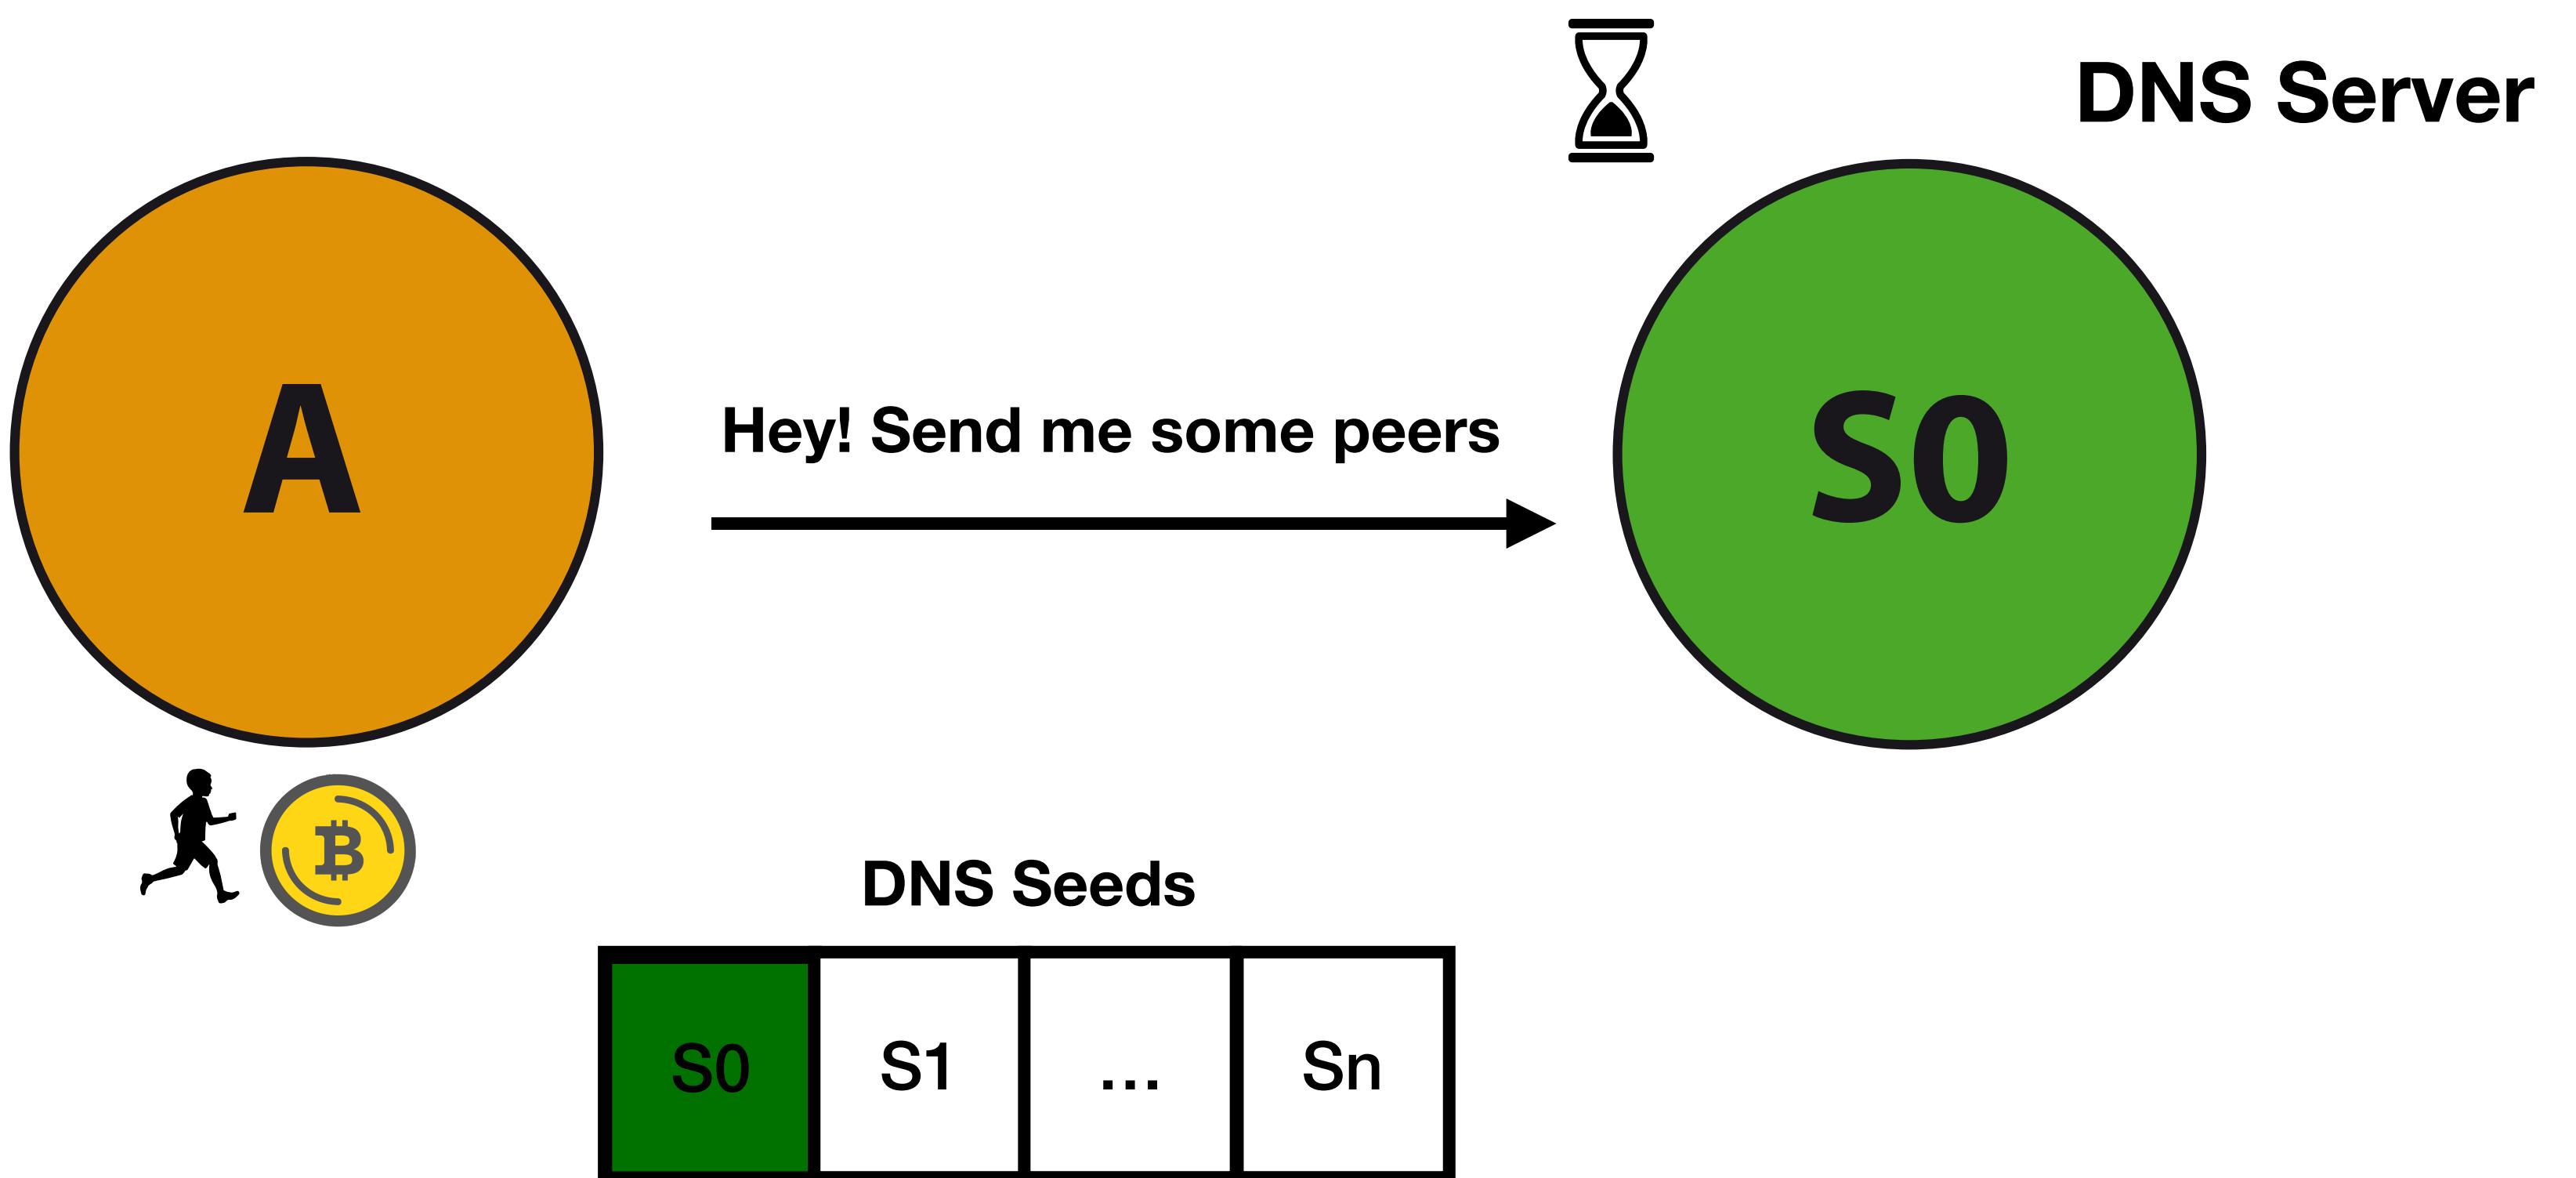
BITCOIN PEER DISCOVERY



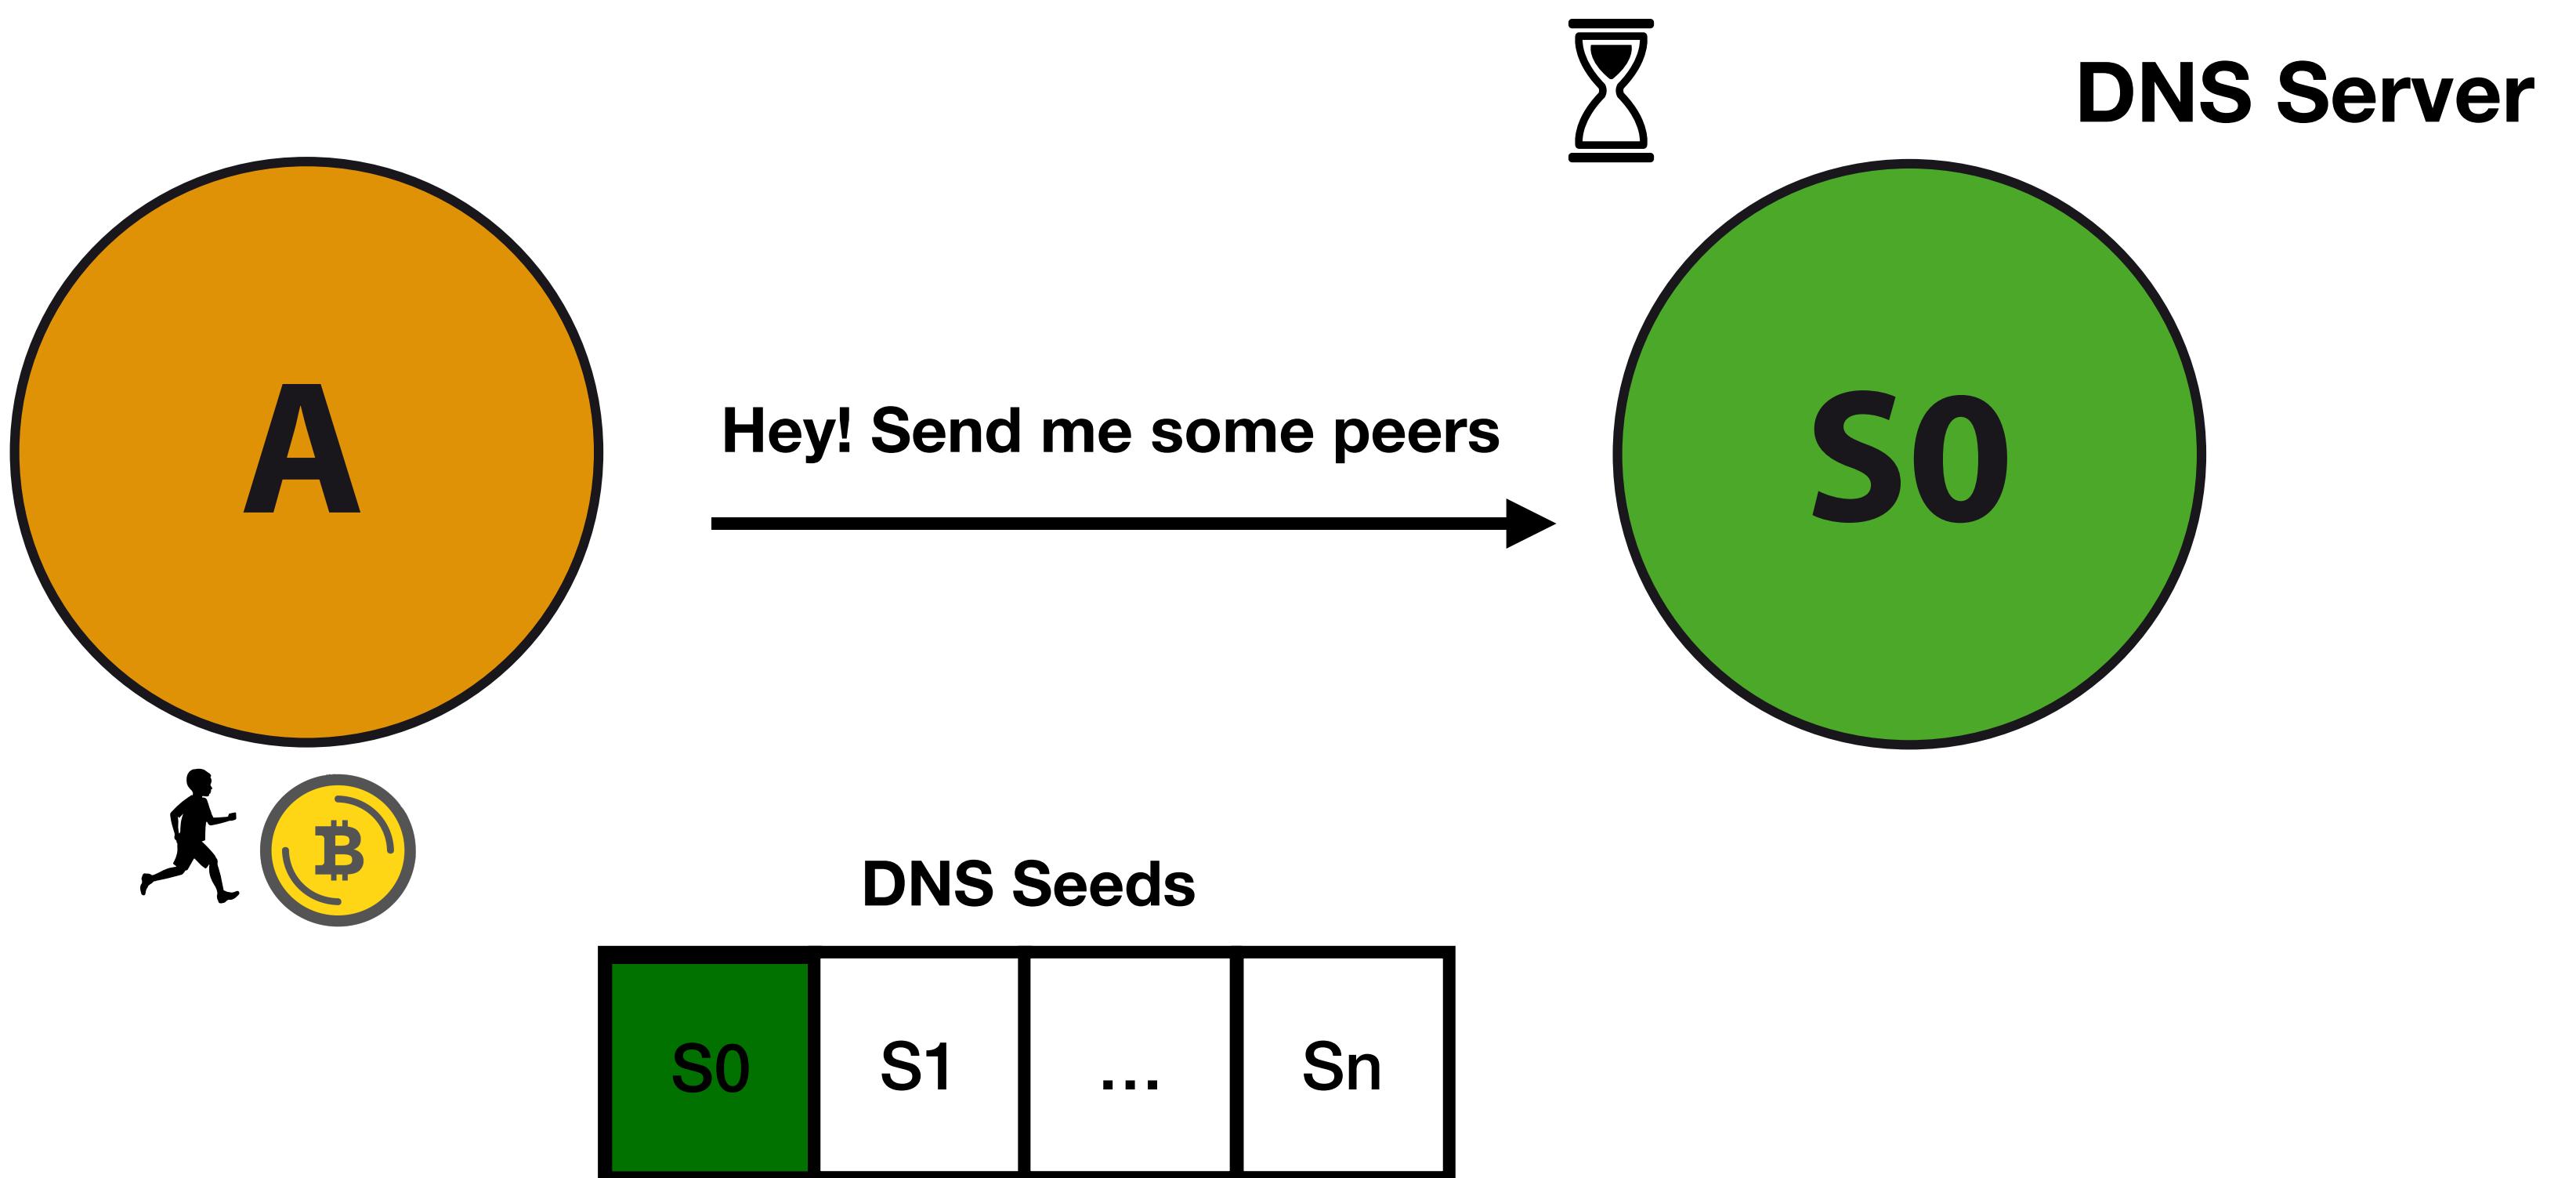
BITCOIN PEER DISCOVERY



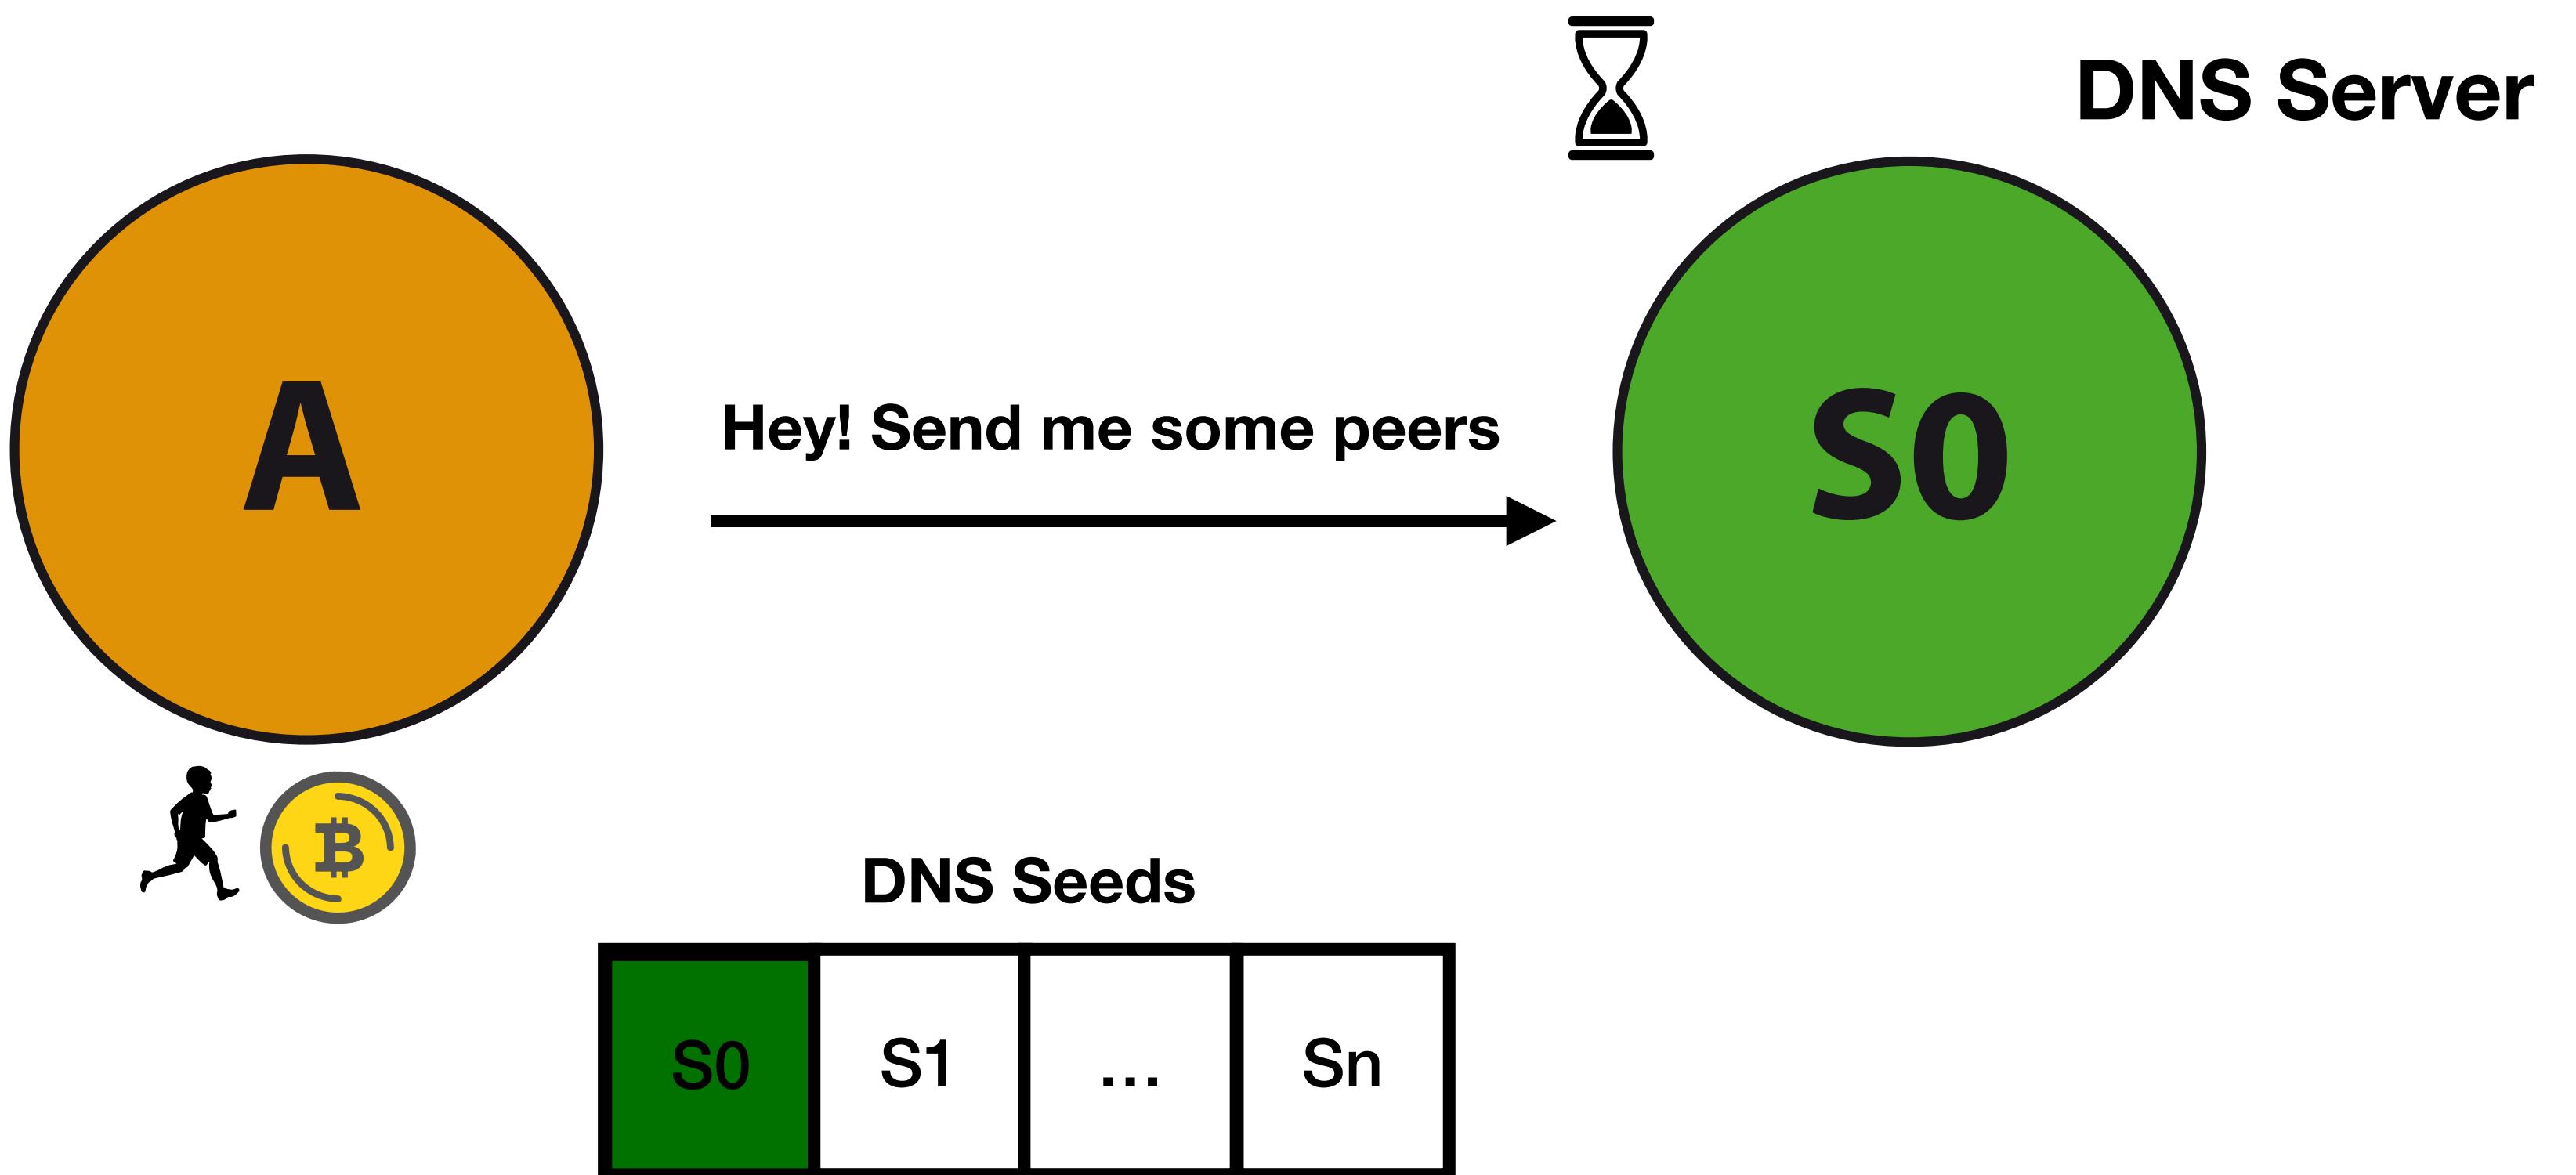
BITCOIN PEER DISCOVERY



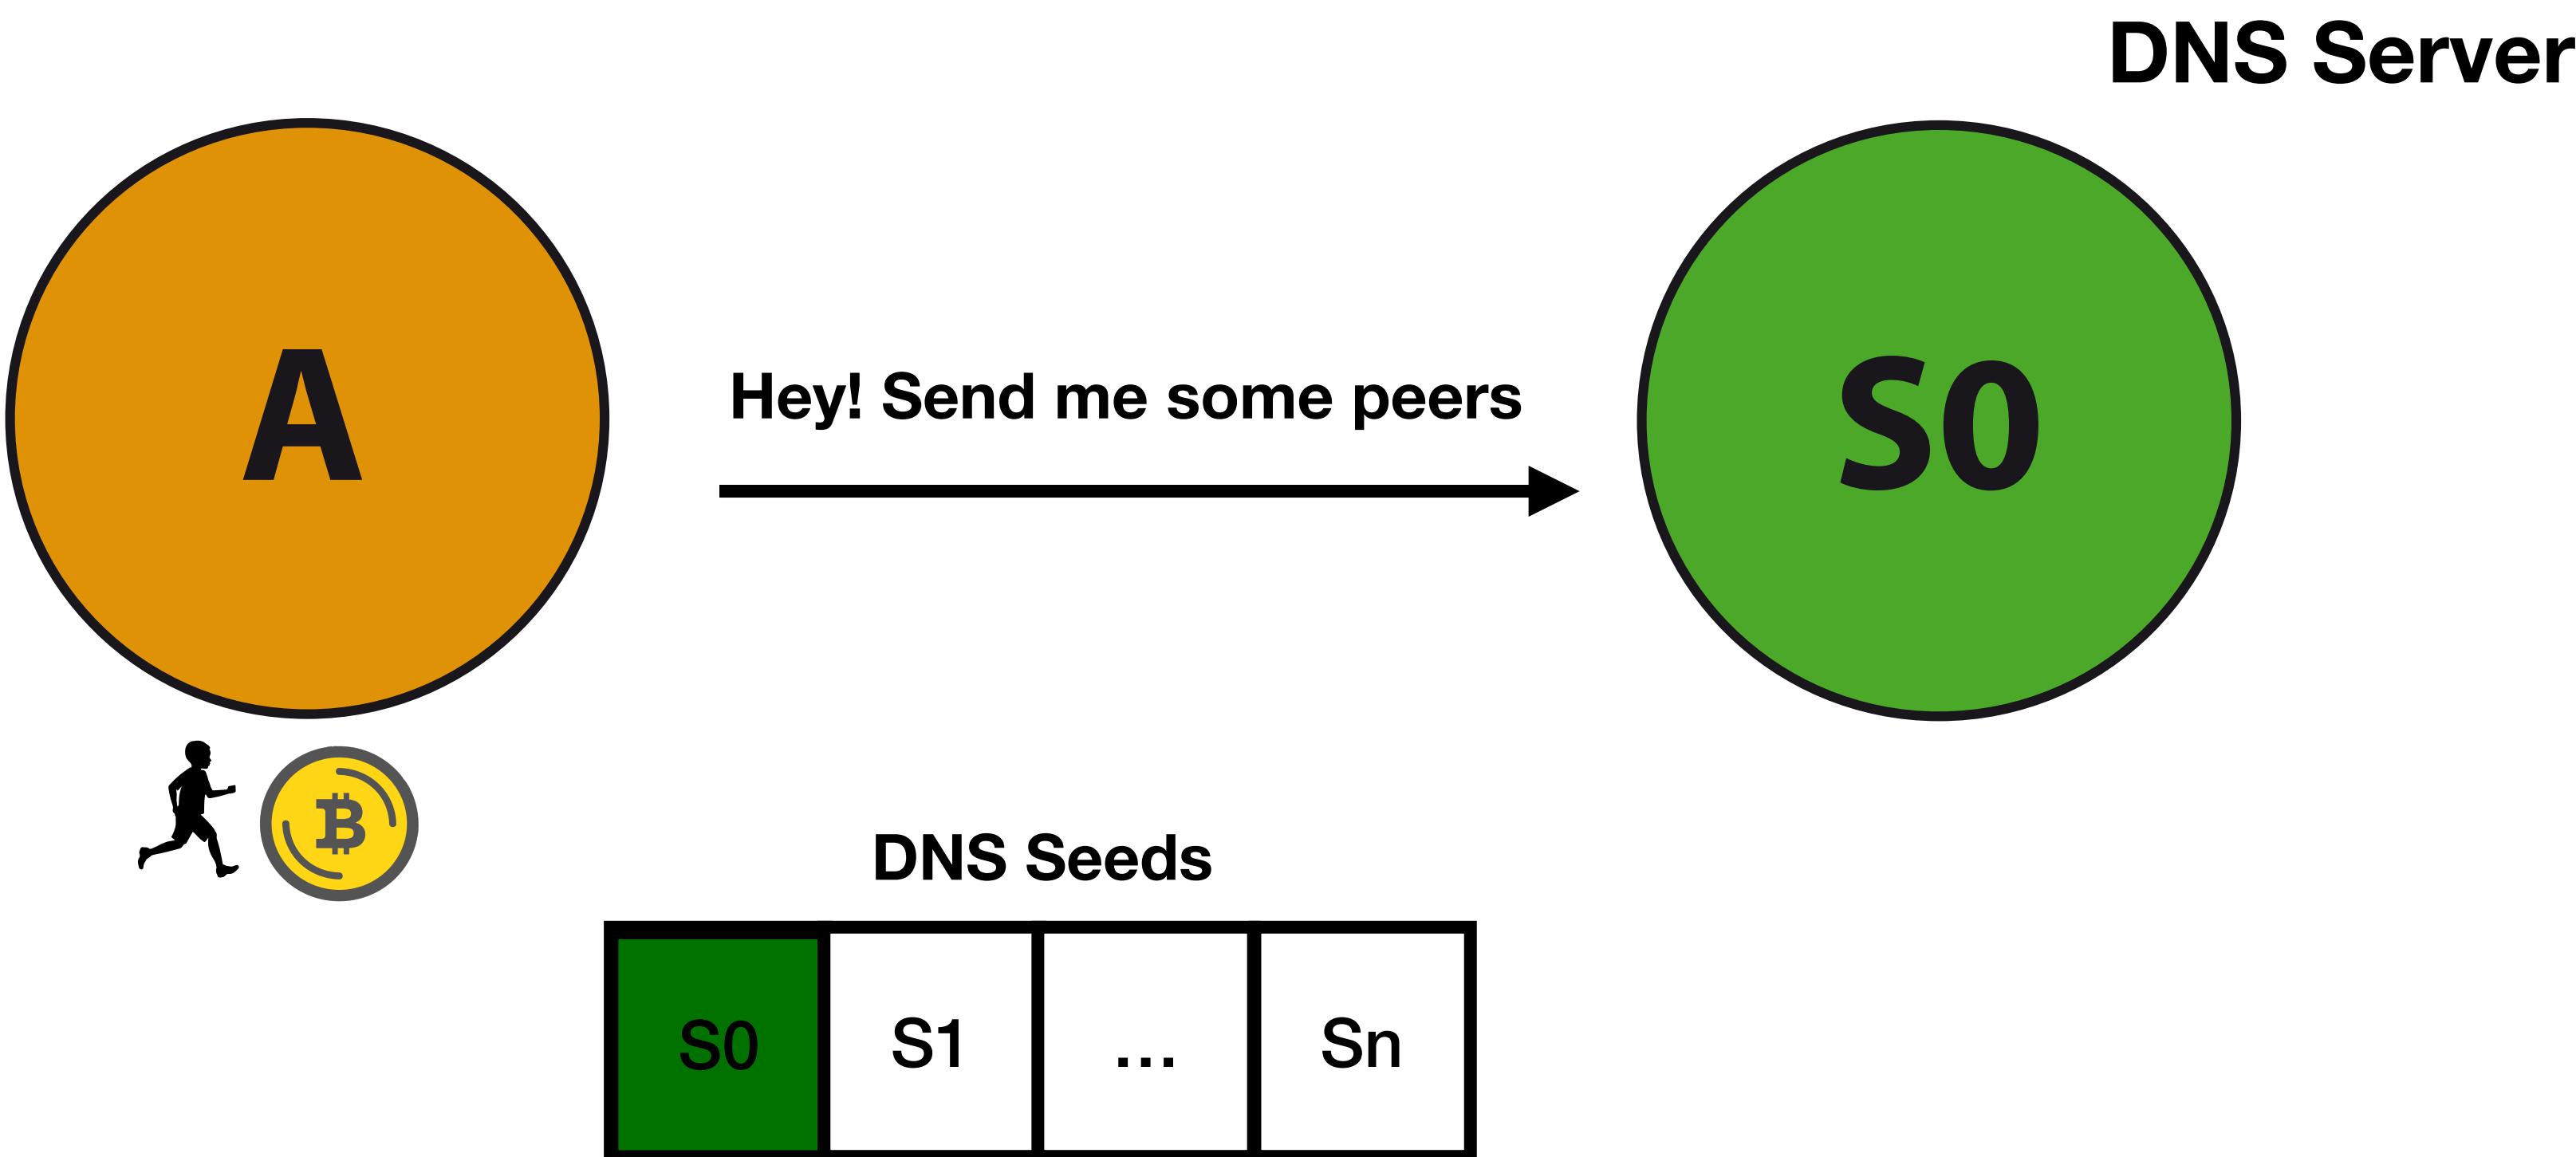
BITCOIN PEER DISCOVERY



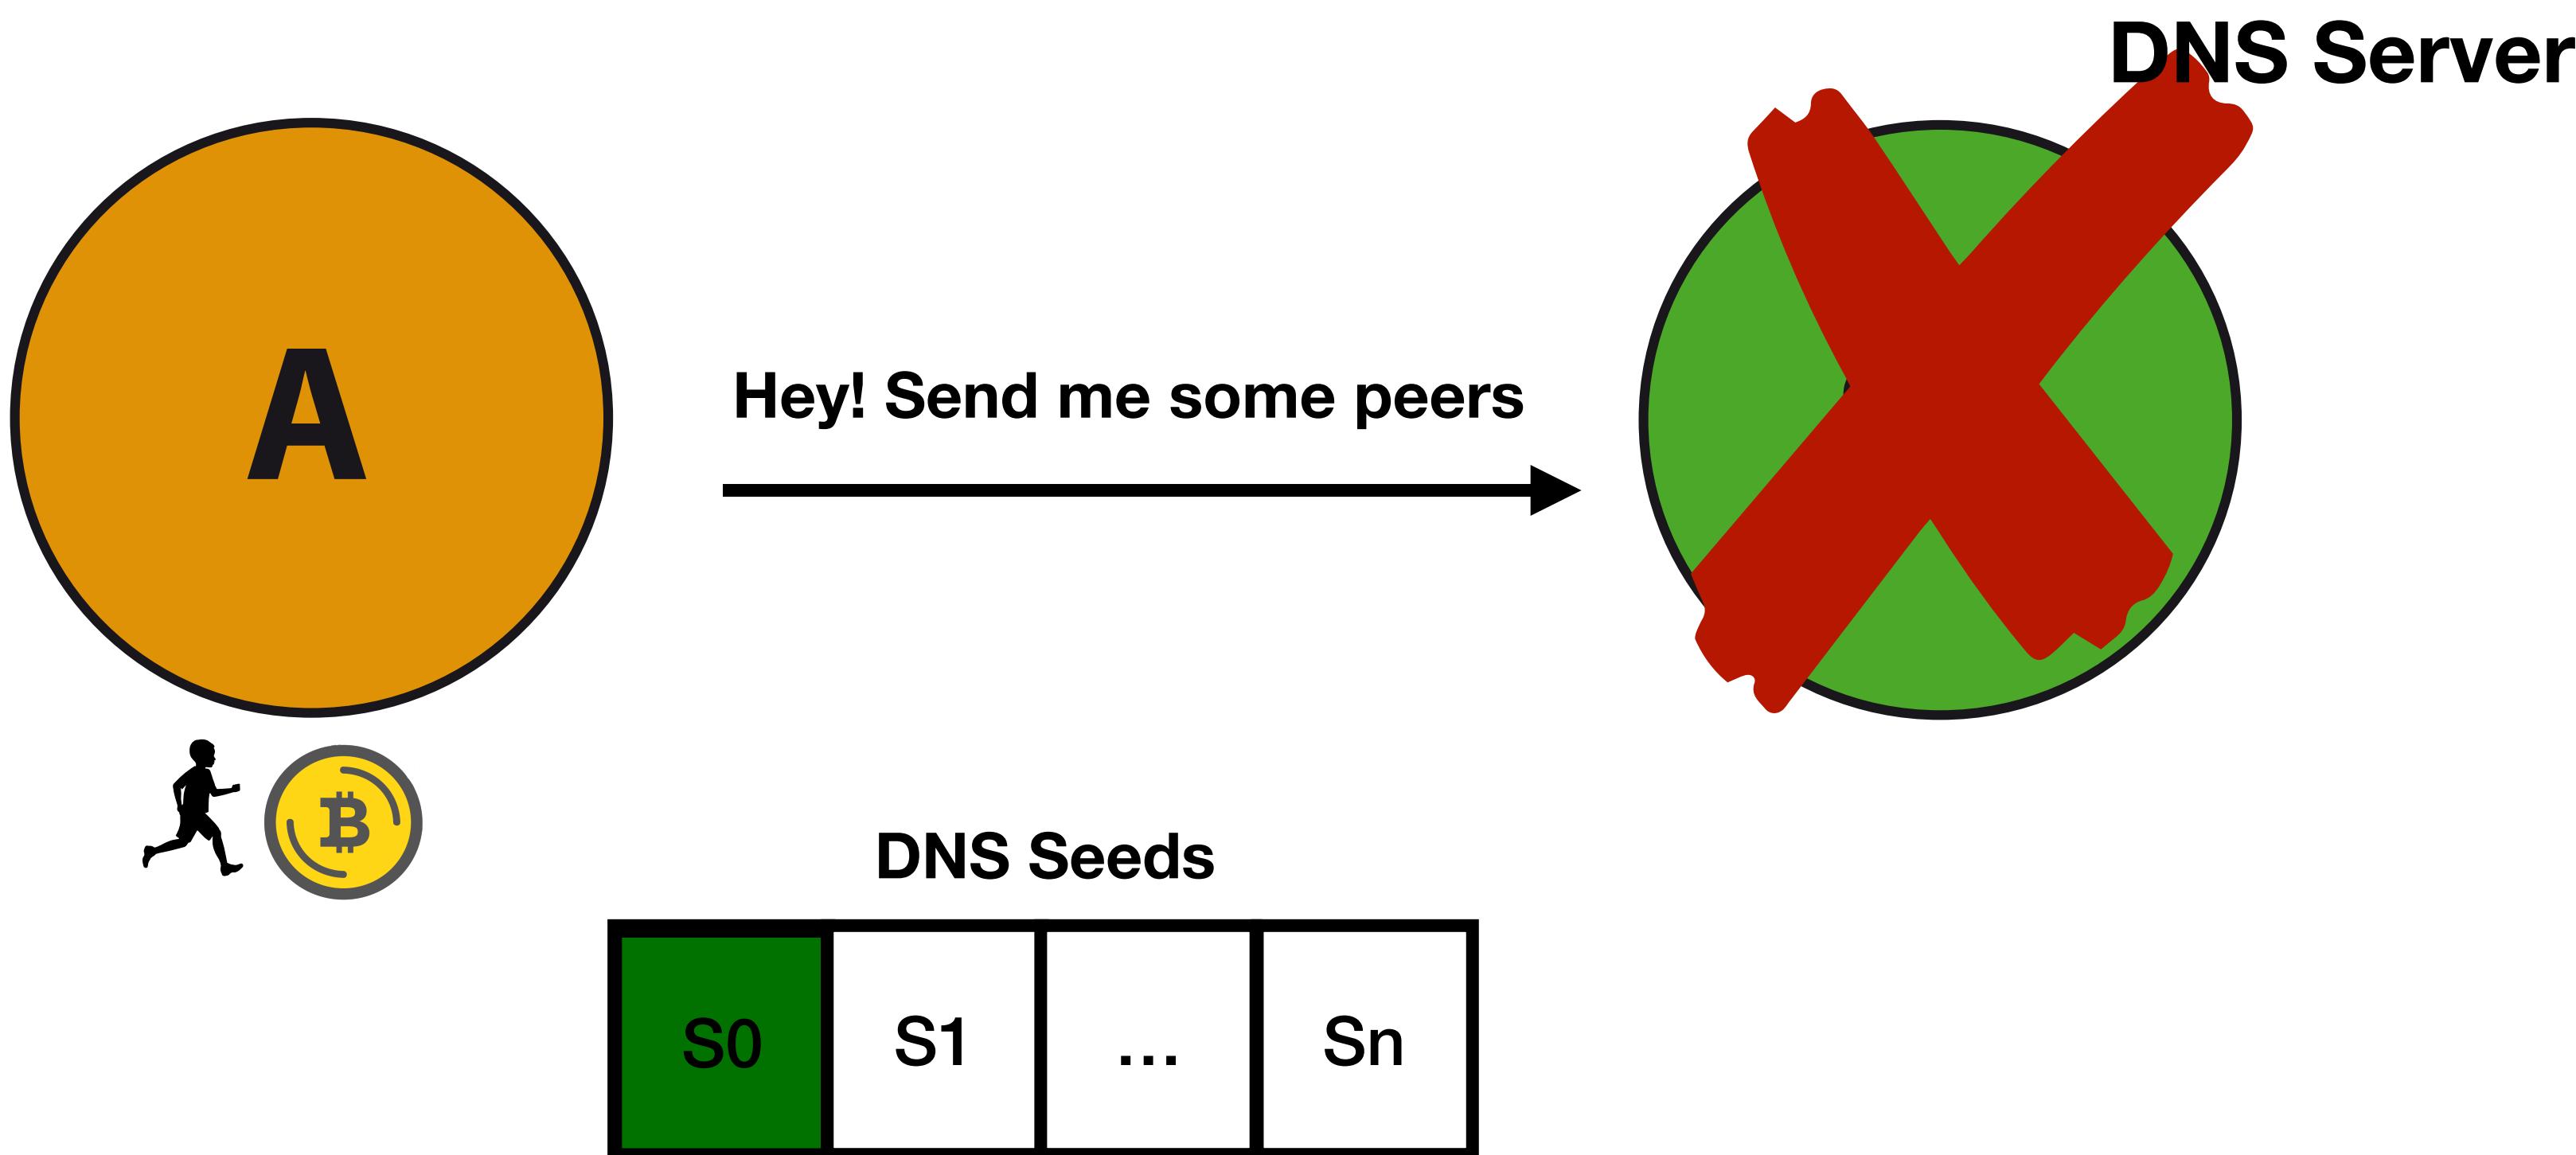
BITCOIN PEER DISCOVERY



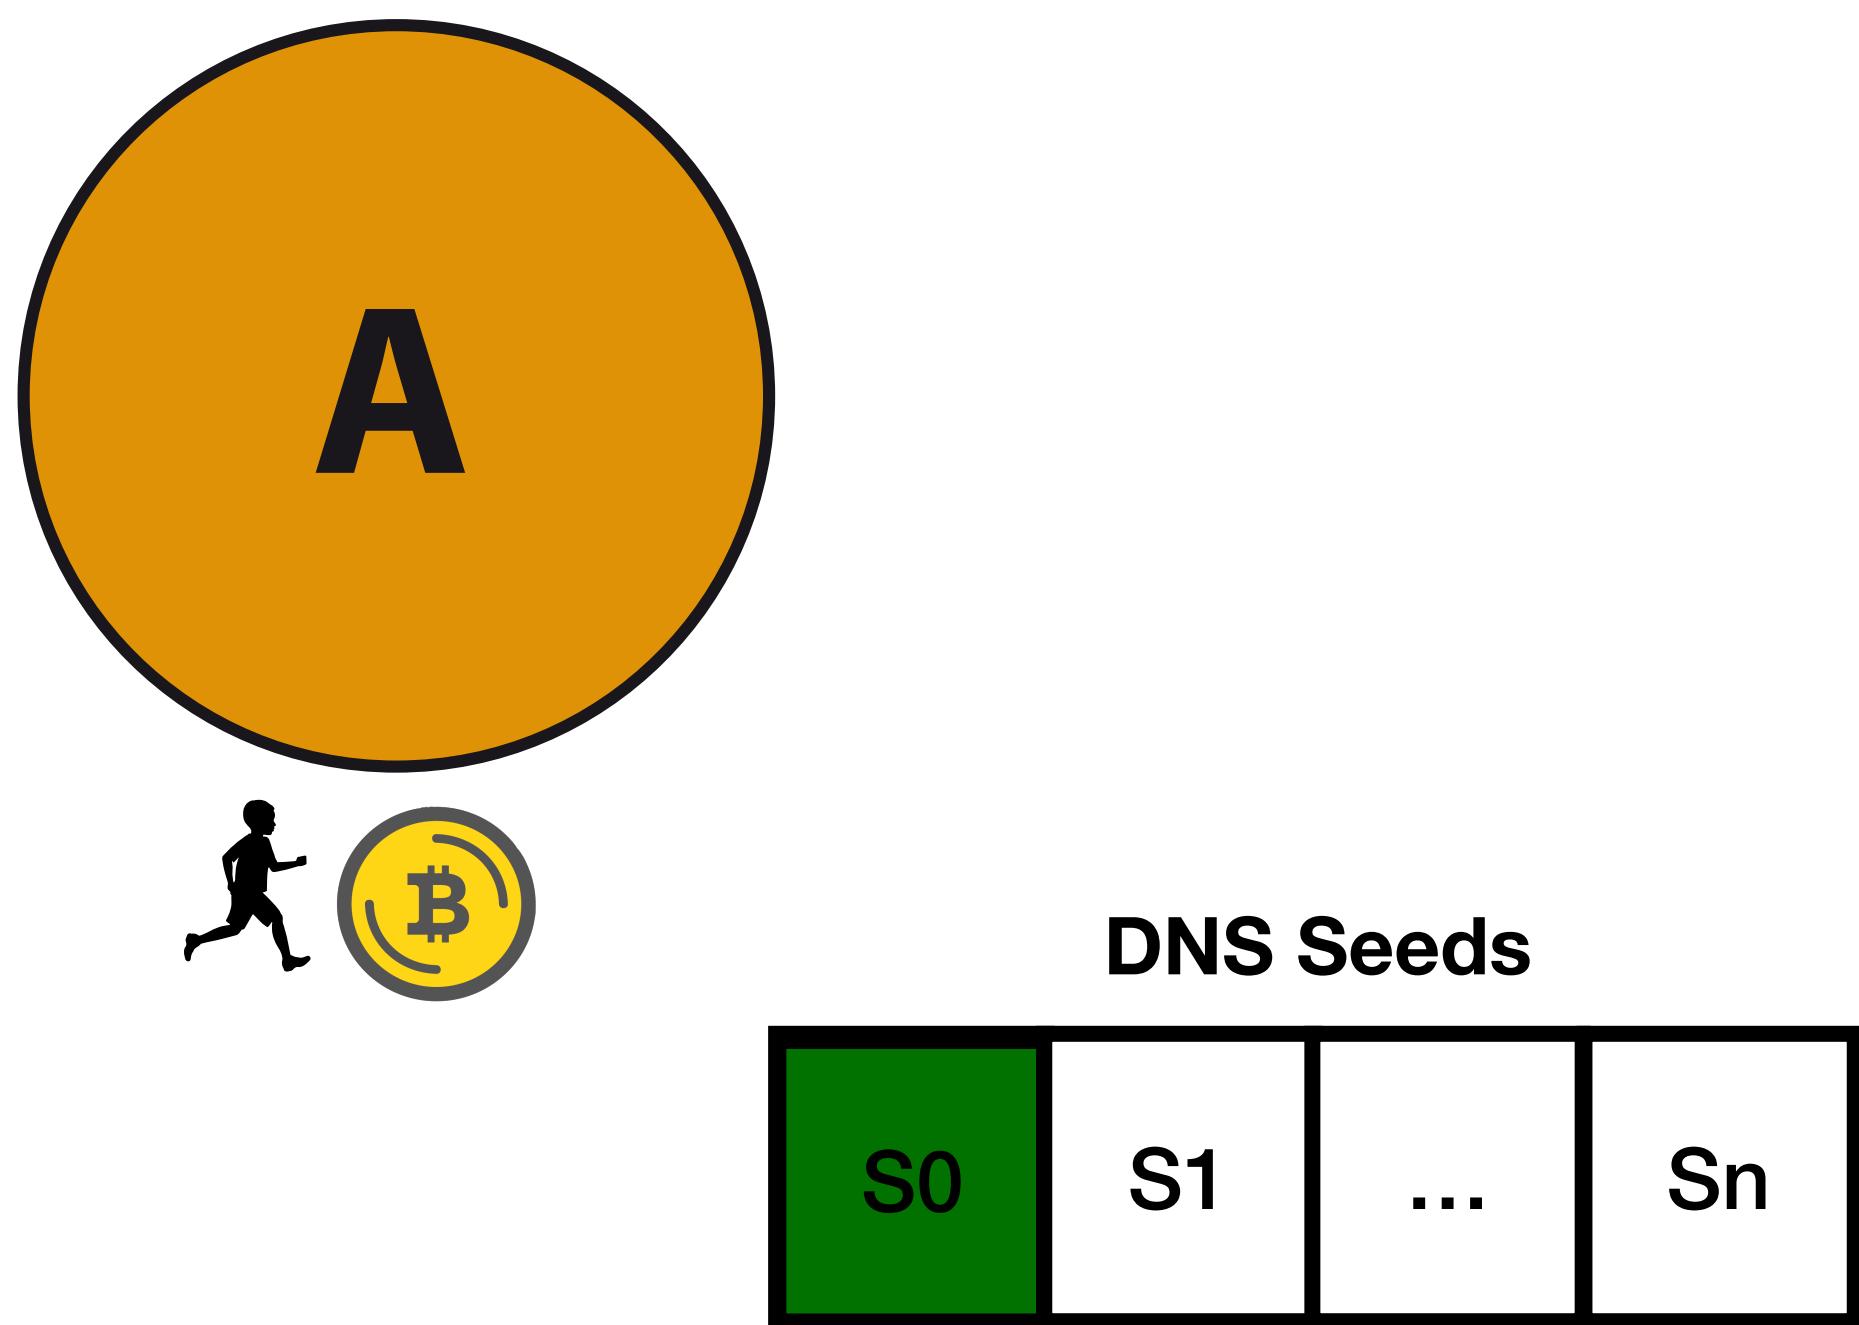
BITCOIN PEER DISCOVERY



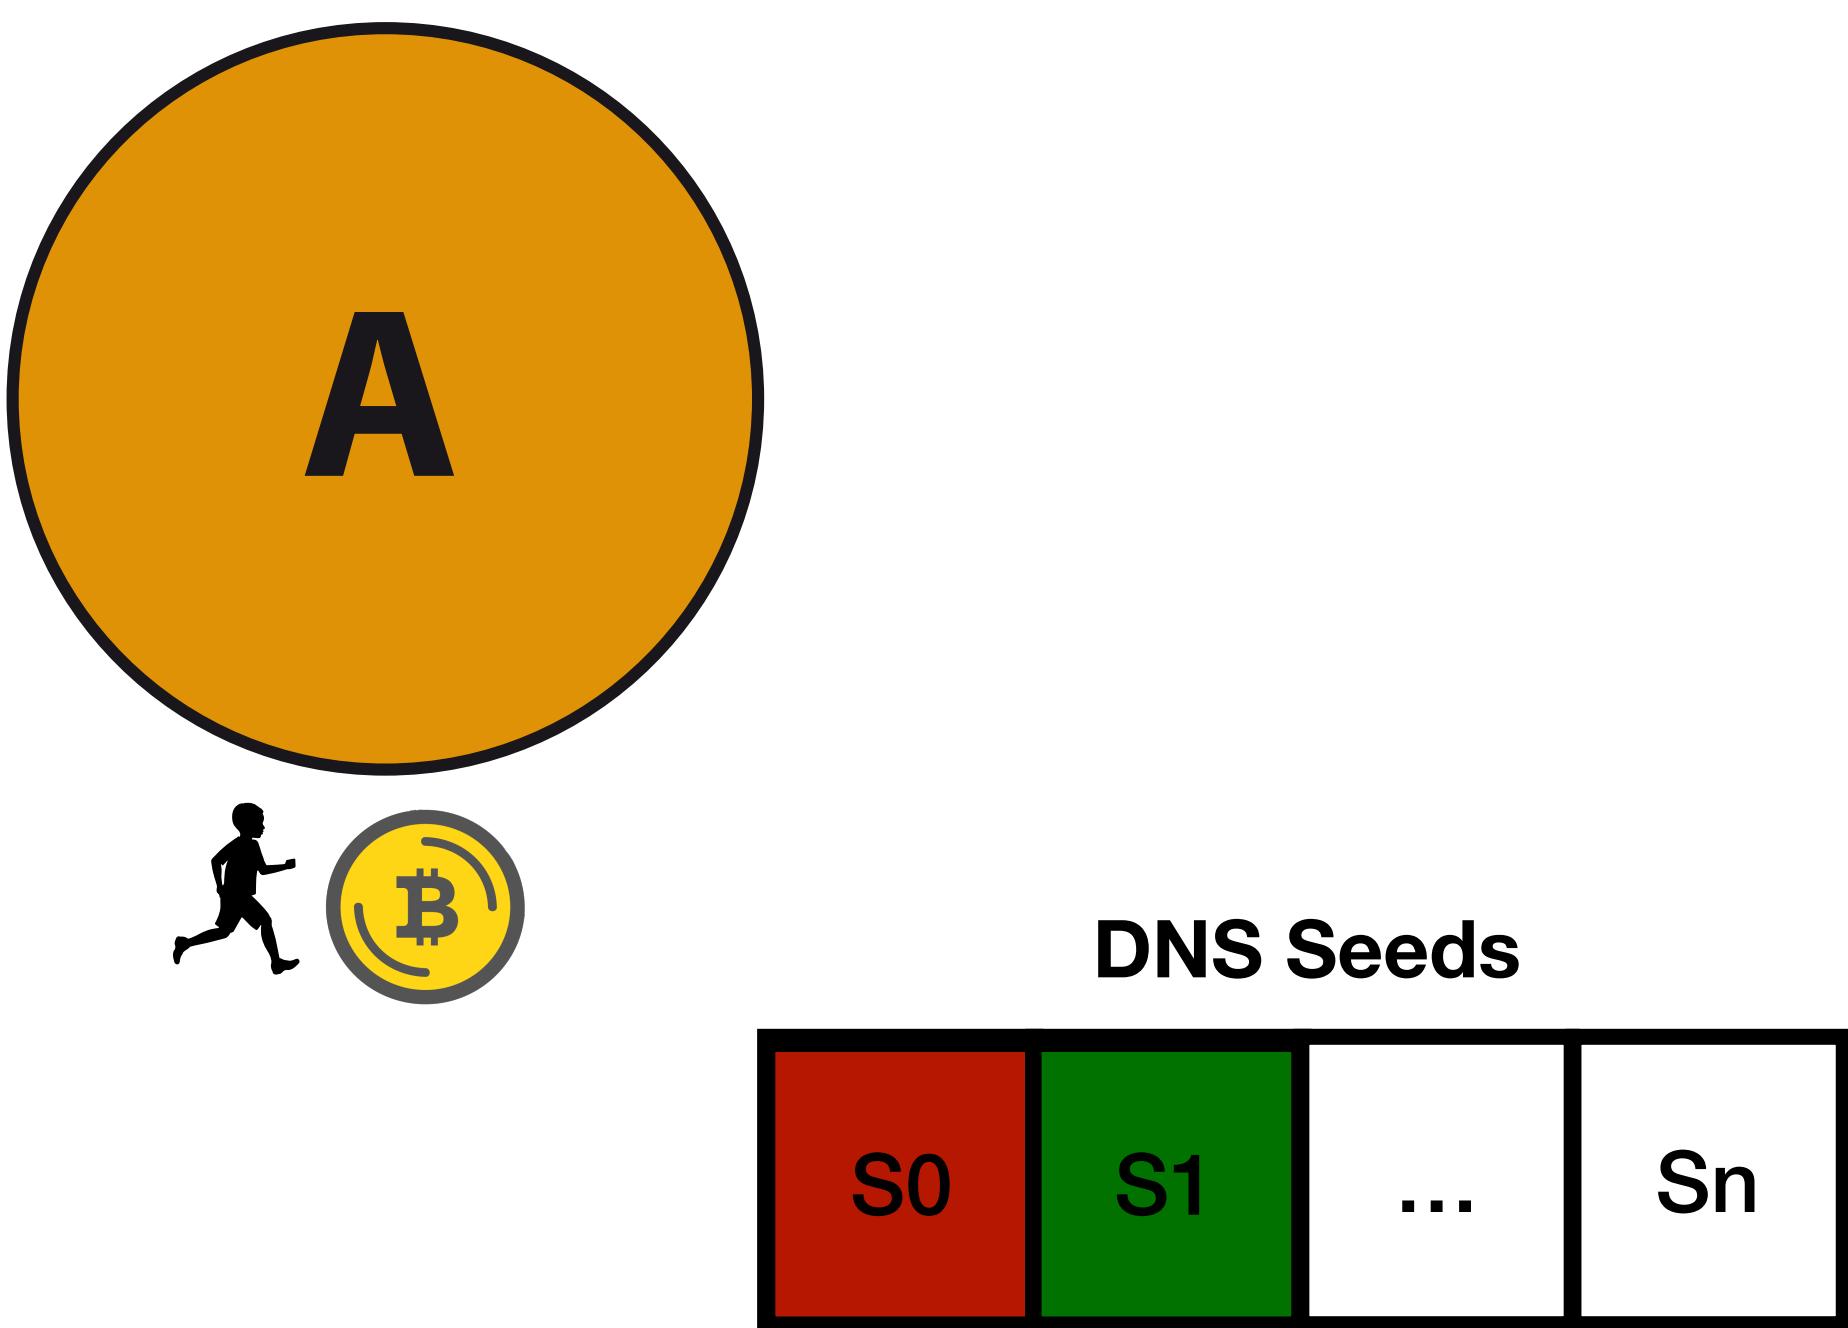
BITCOIN PEER DISCOVERY



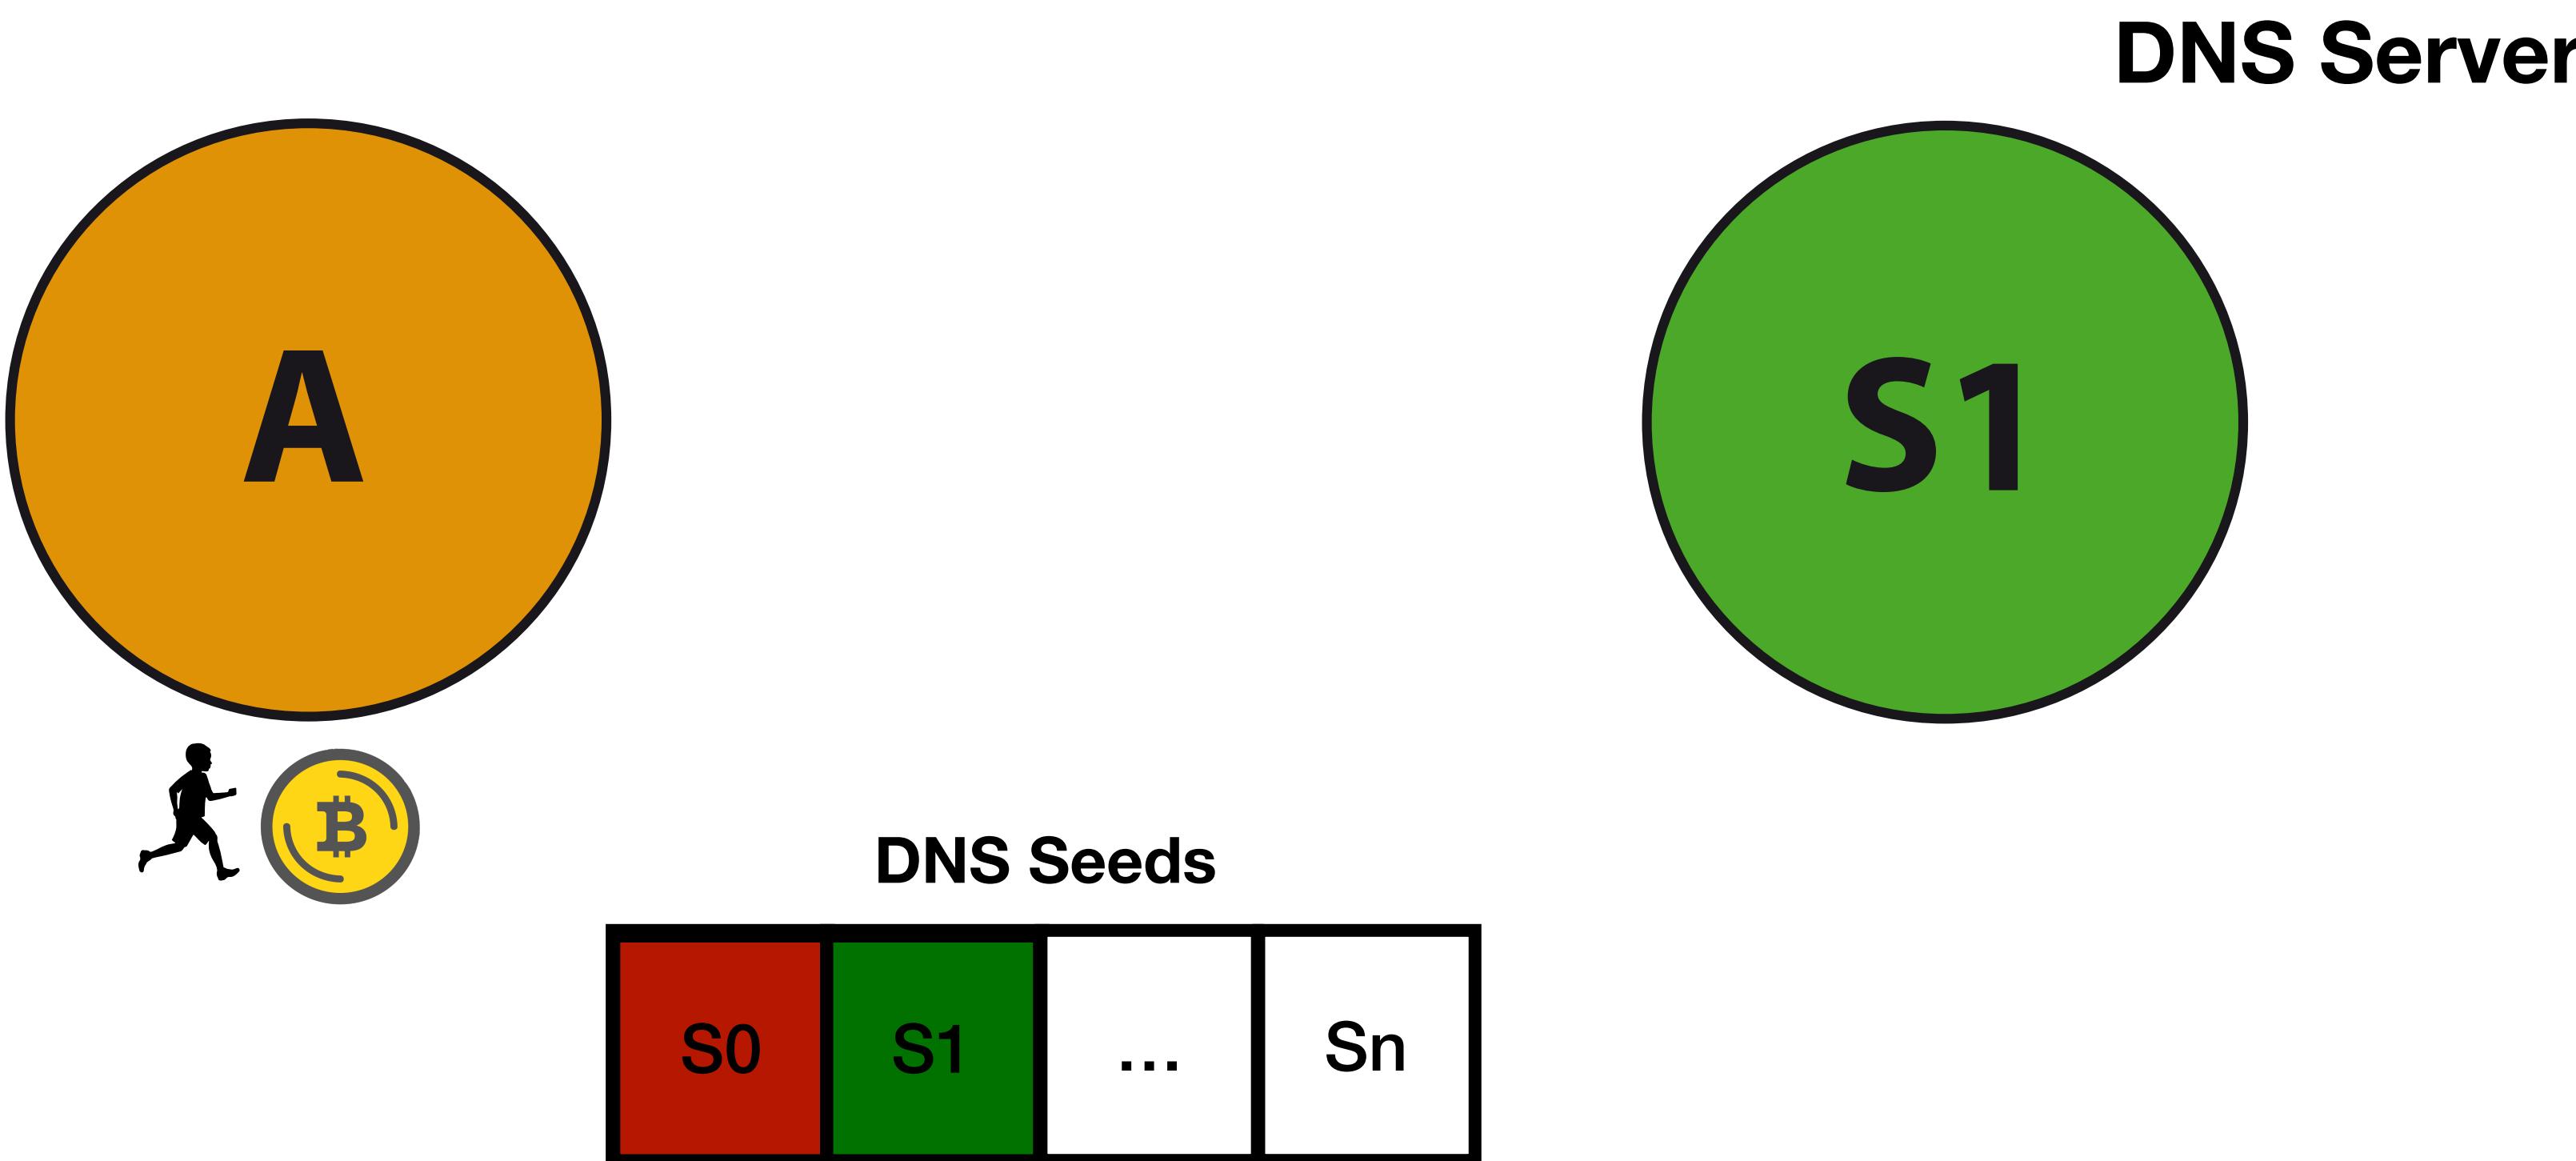
BITCOIN PEER DISCOVERY



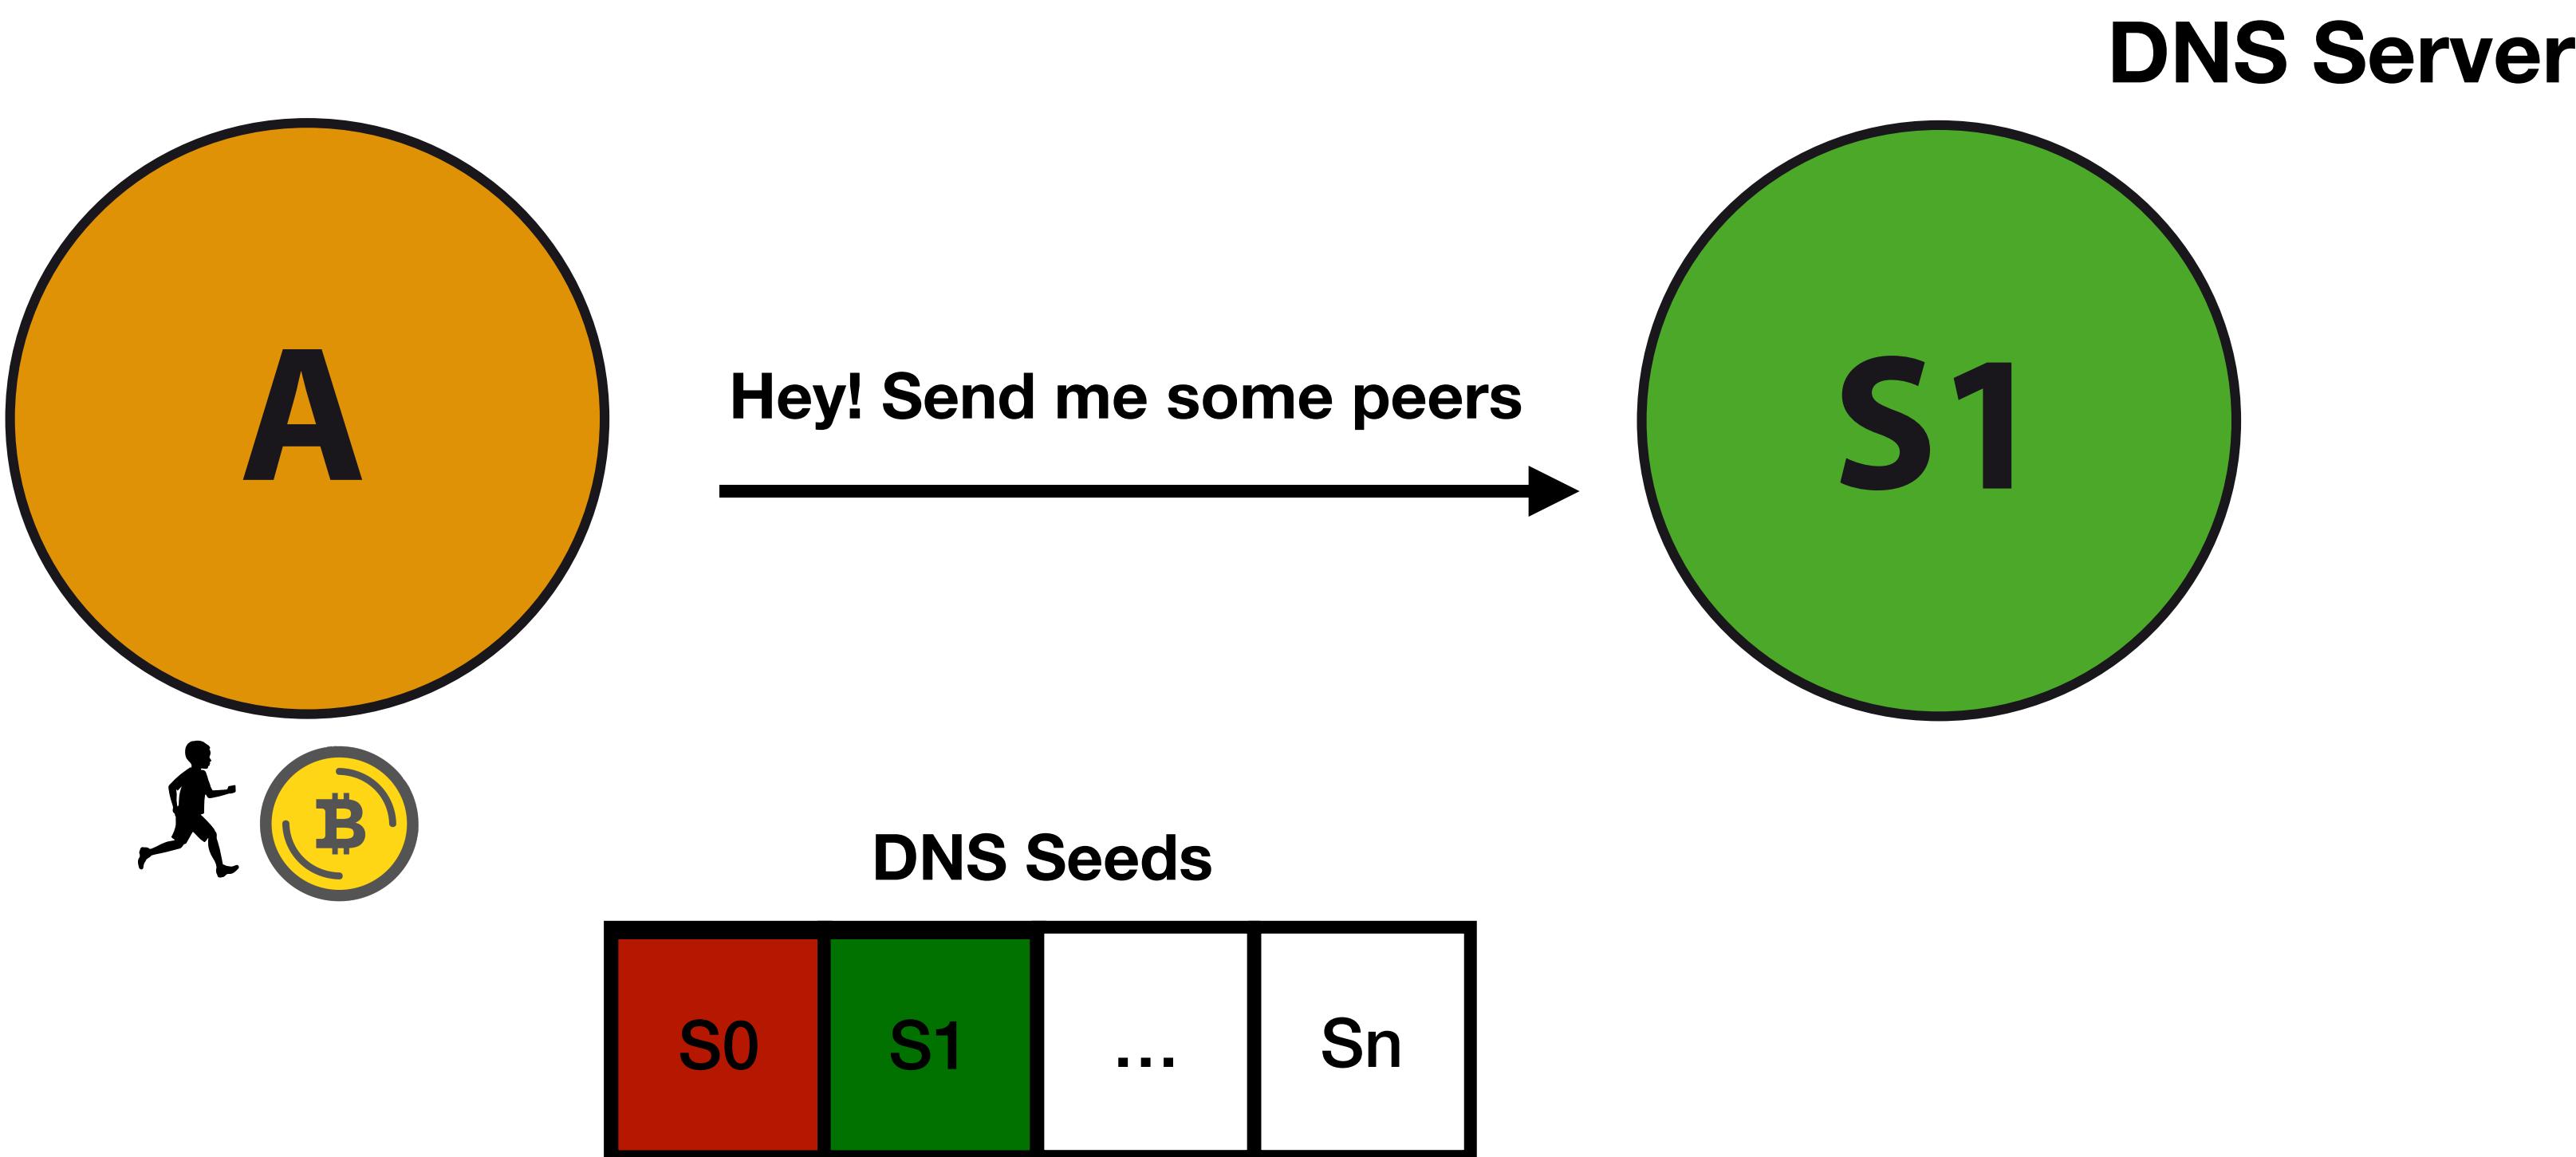
BITCOIN PEER DISCOVERY



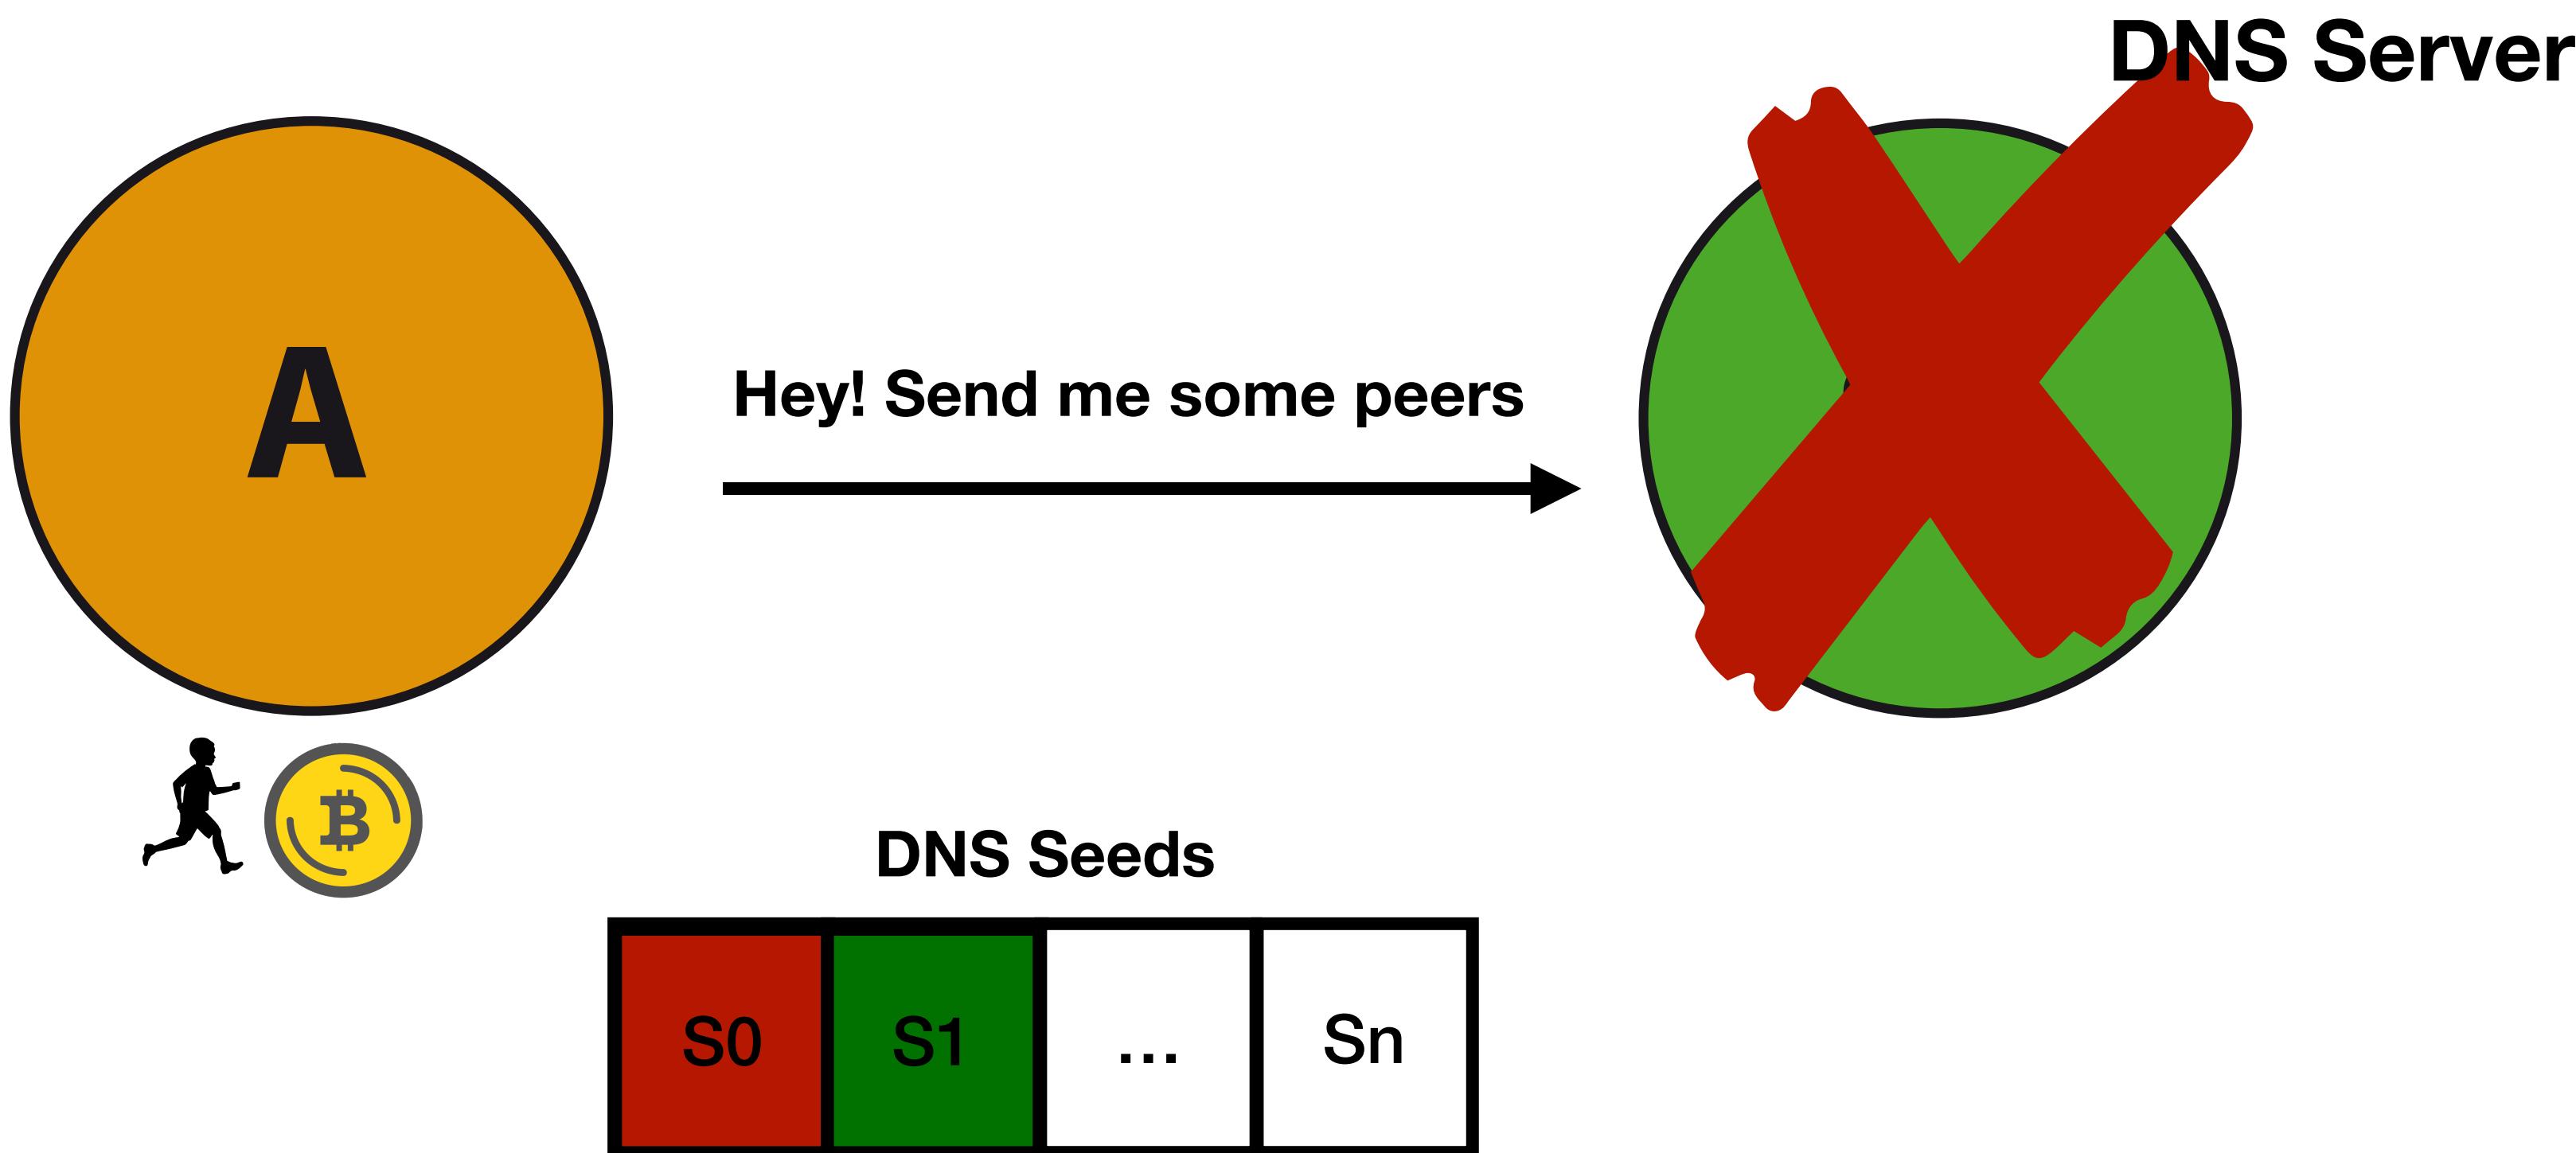
BITCOIN PEER DISCOVERY



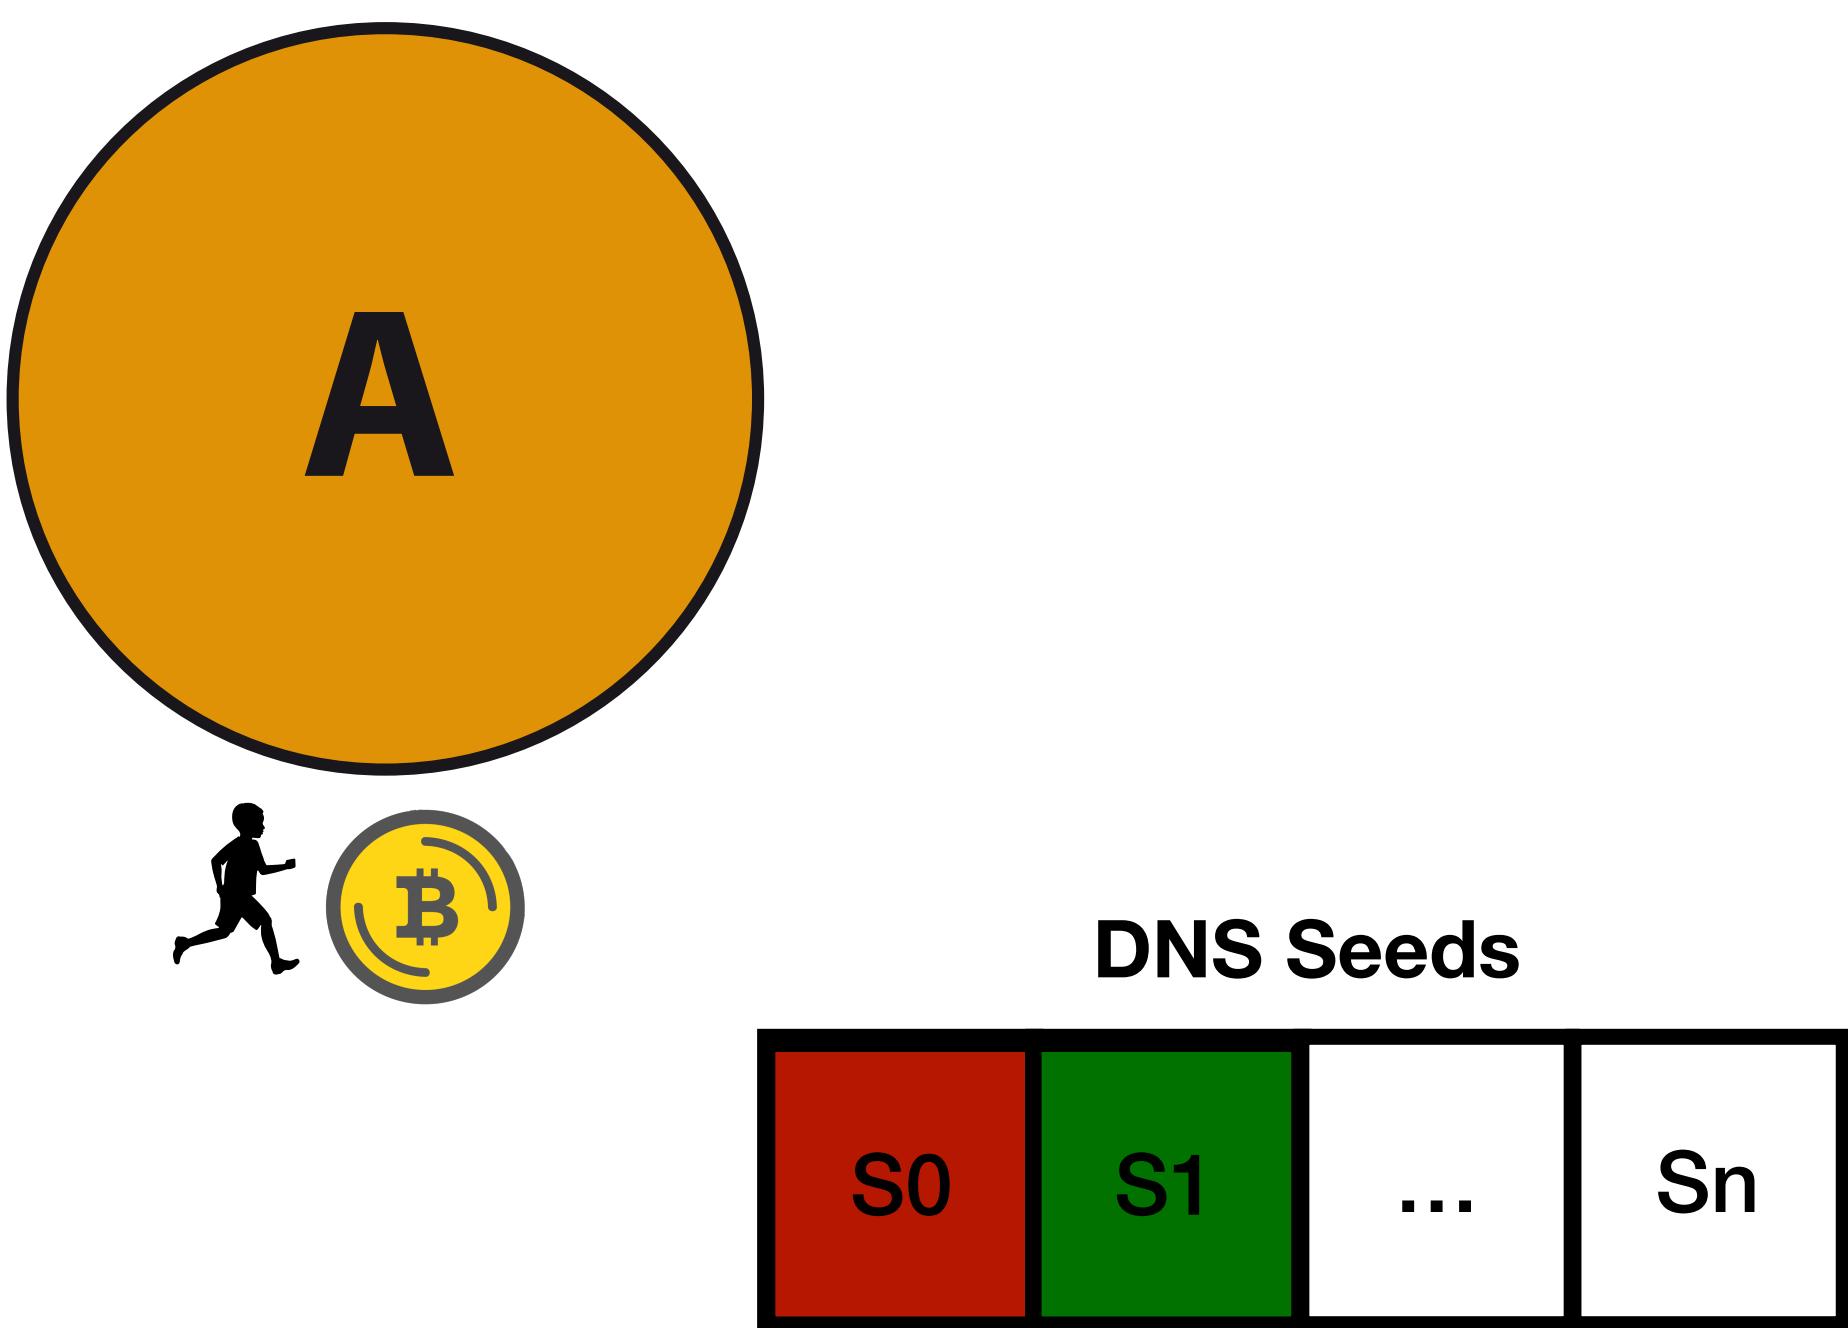
BITCOIN PEER DISCOVERY



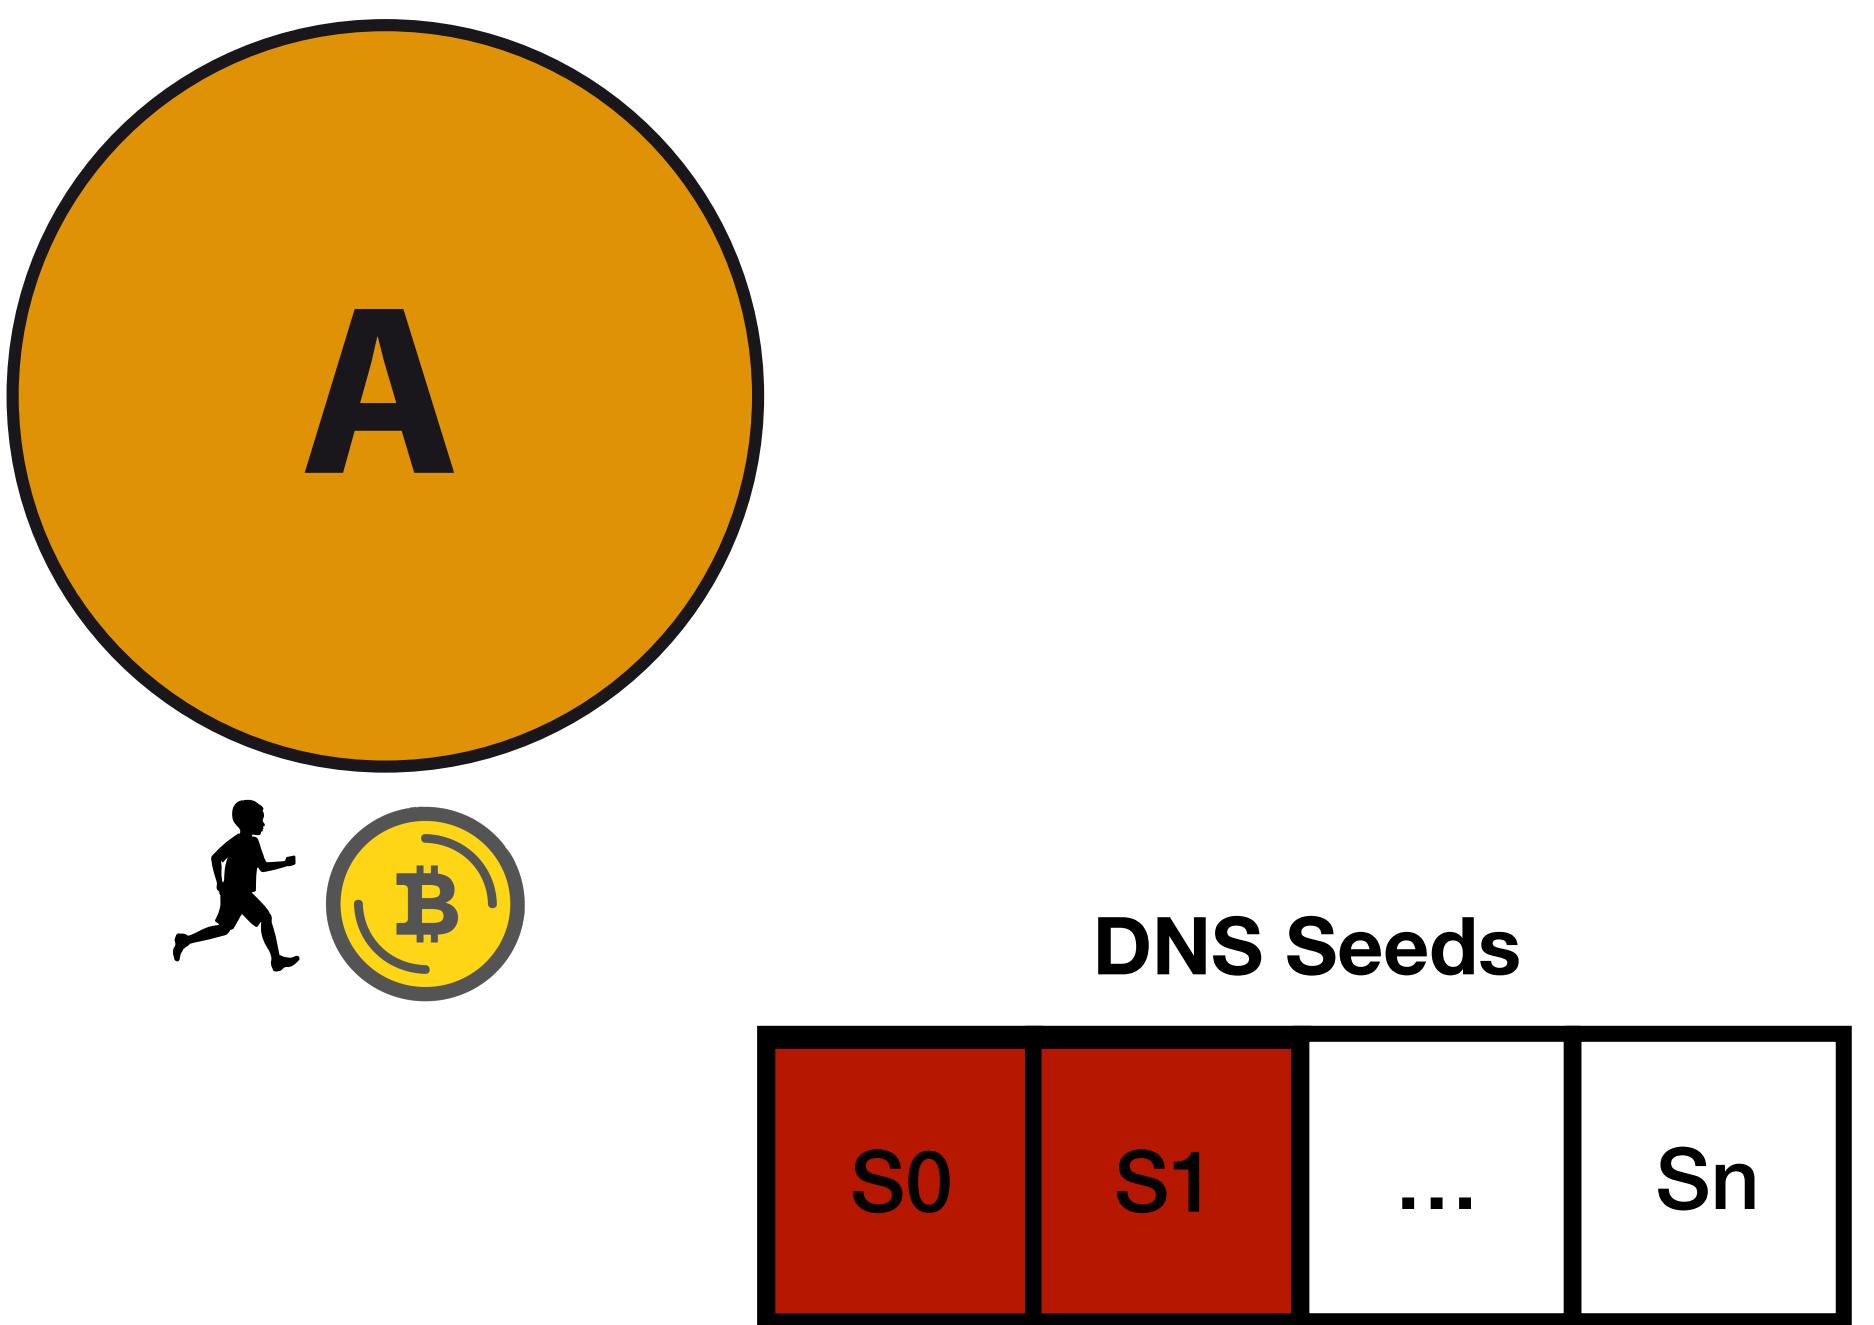
BITCOIN PEER DISCOVERY



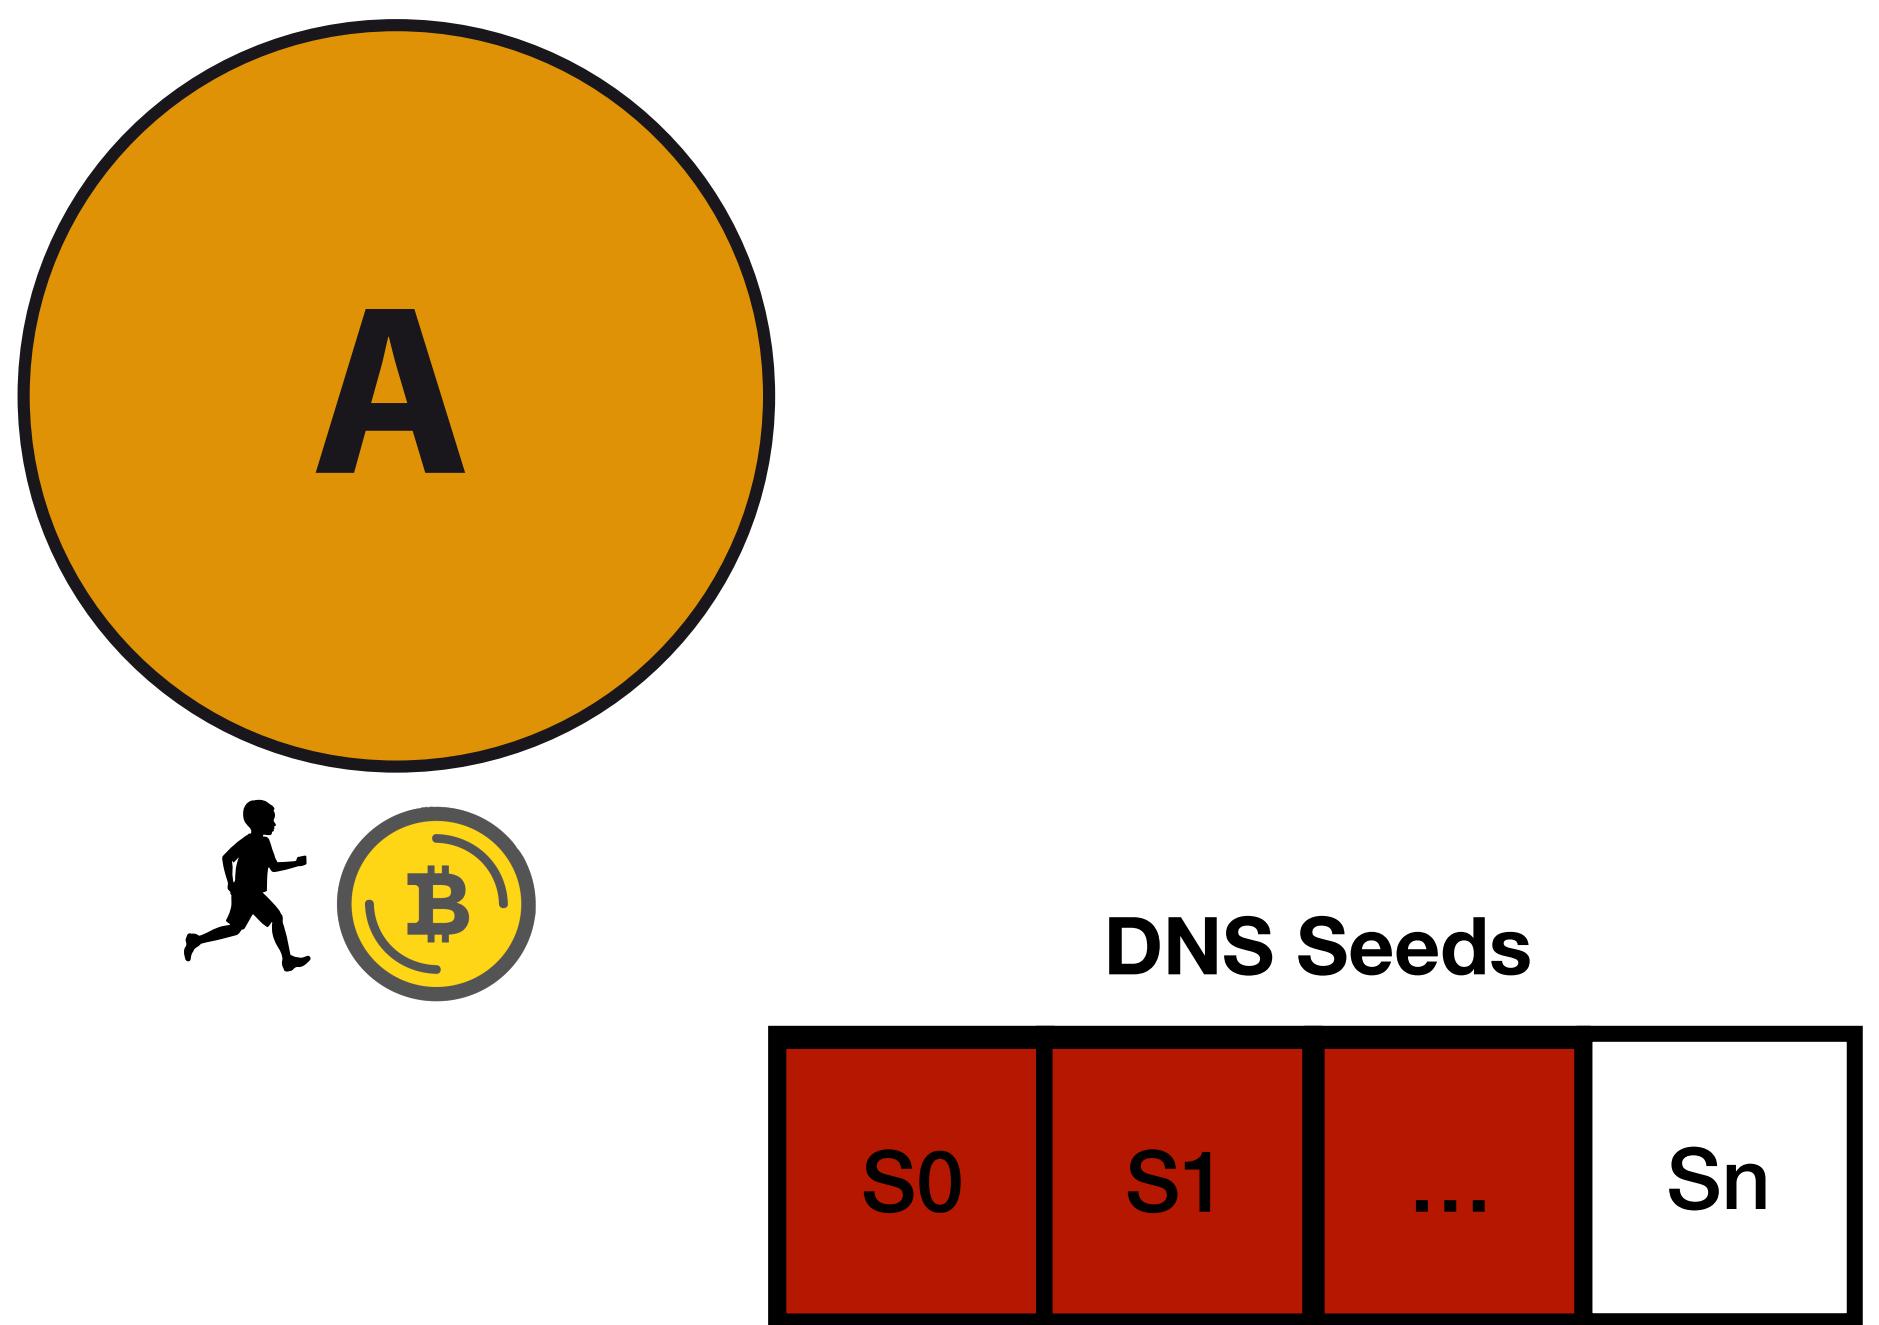
BITCOIN PEER DISCOVERY



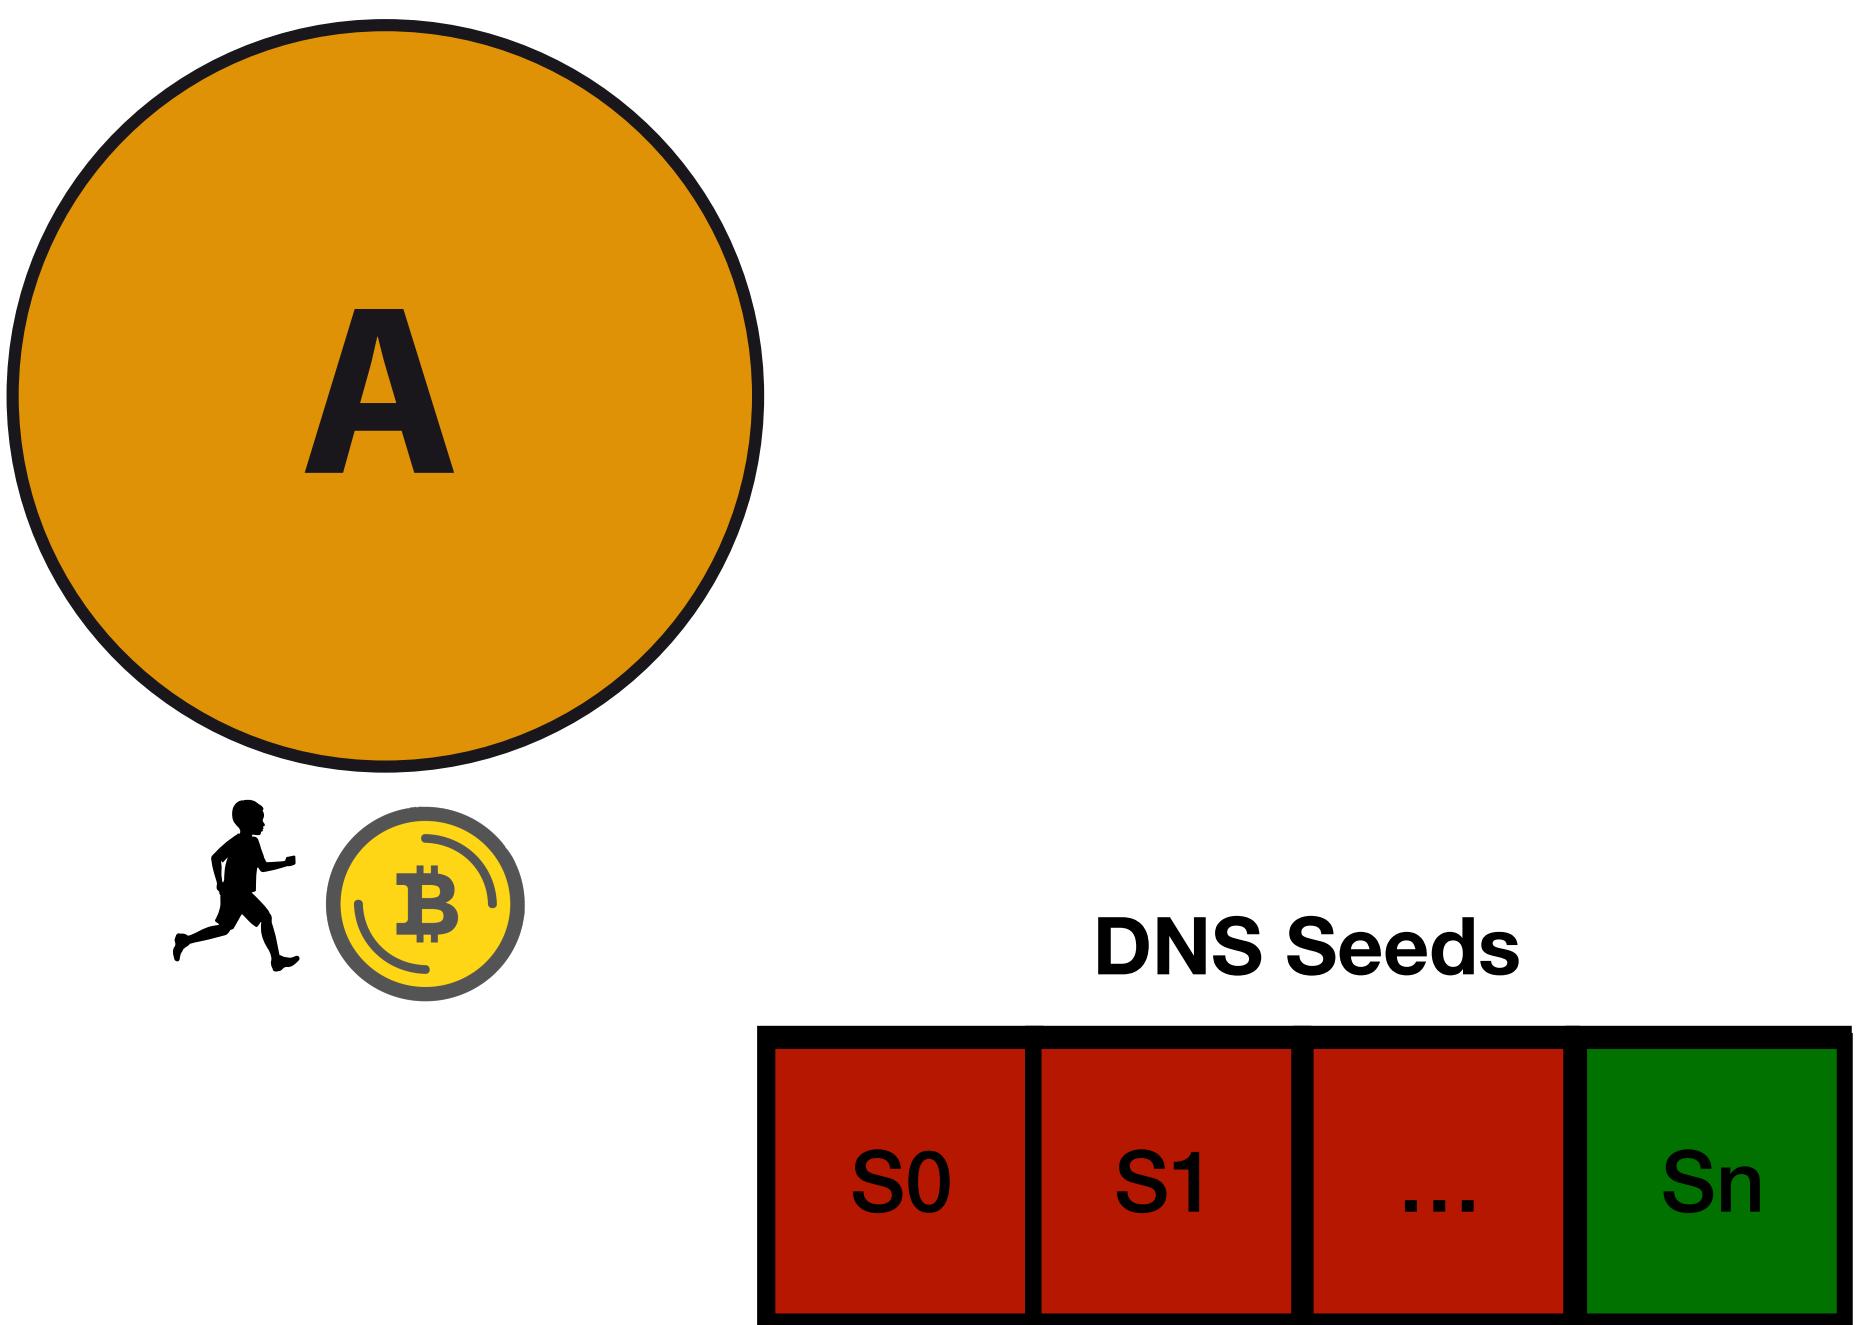
BITCOIN PEER DISCOVERY



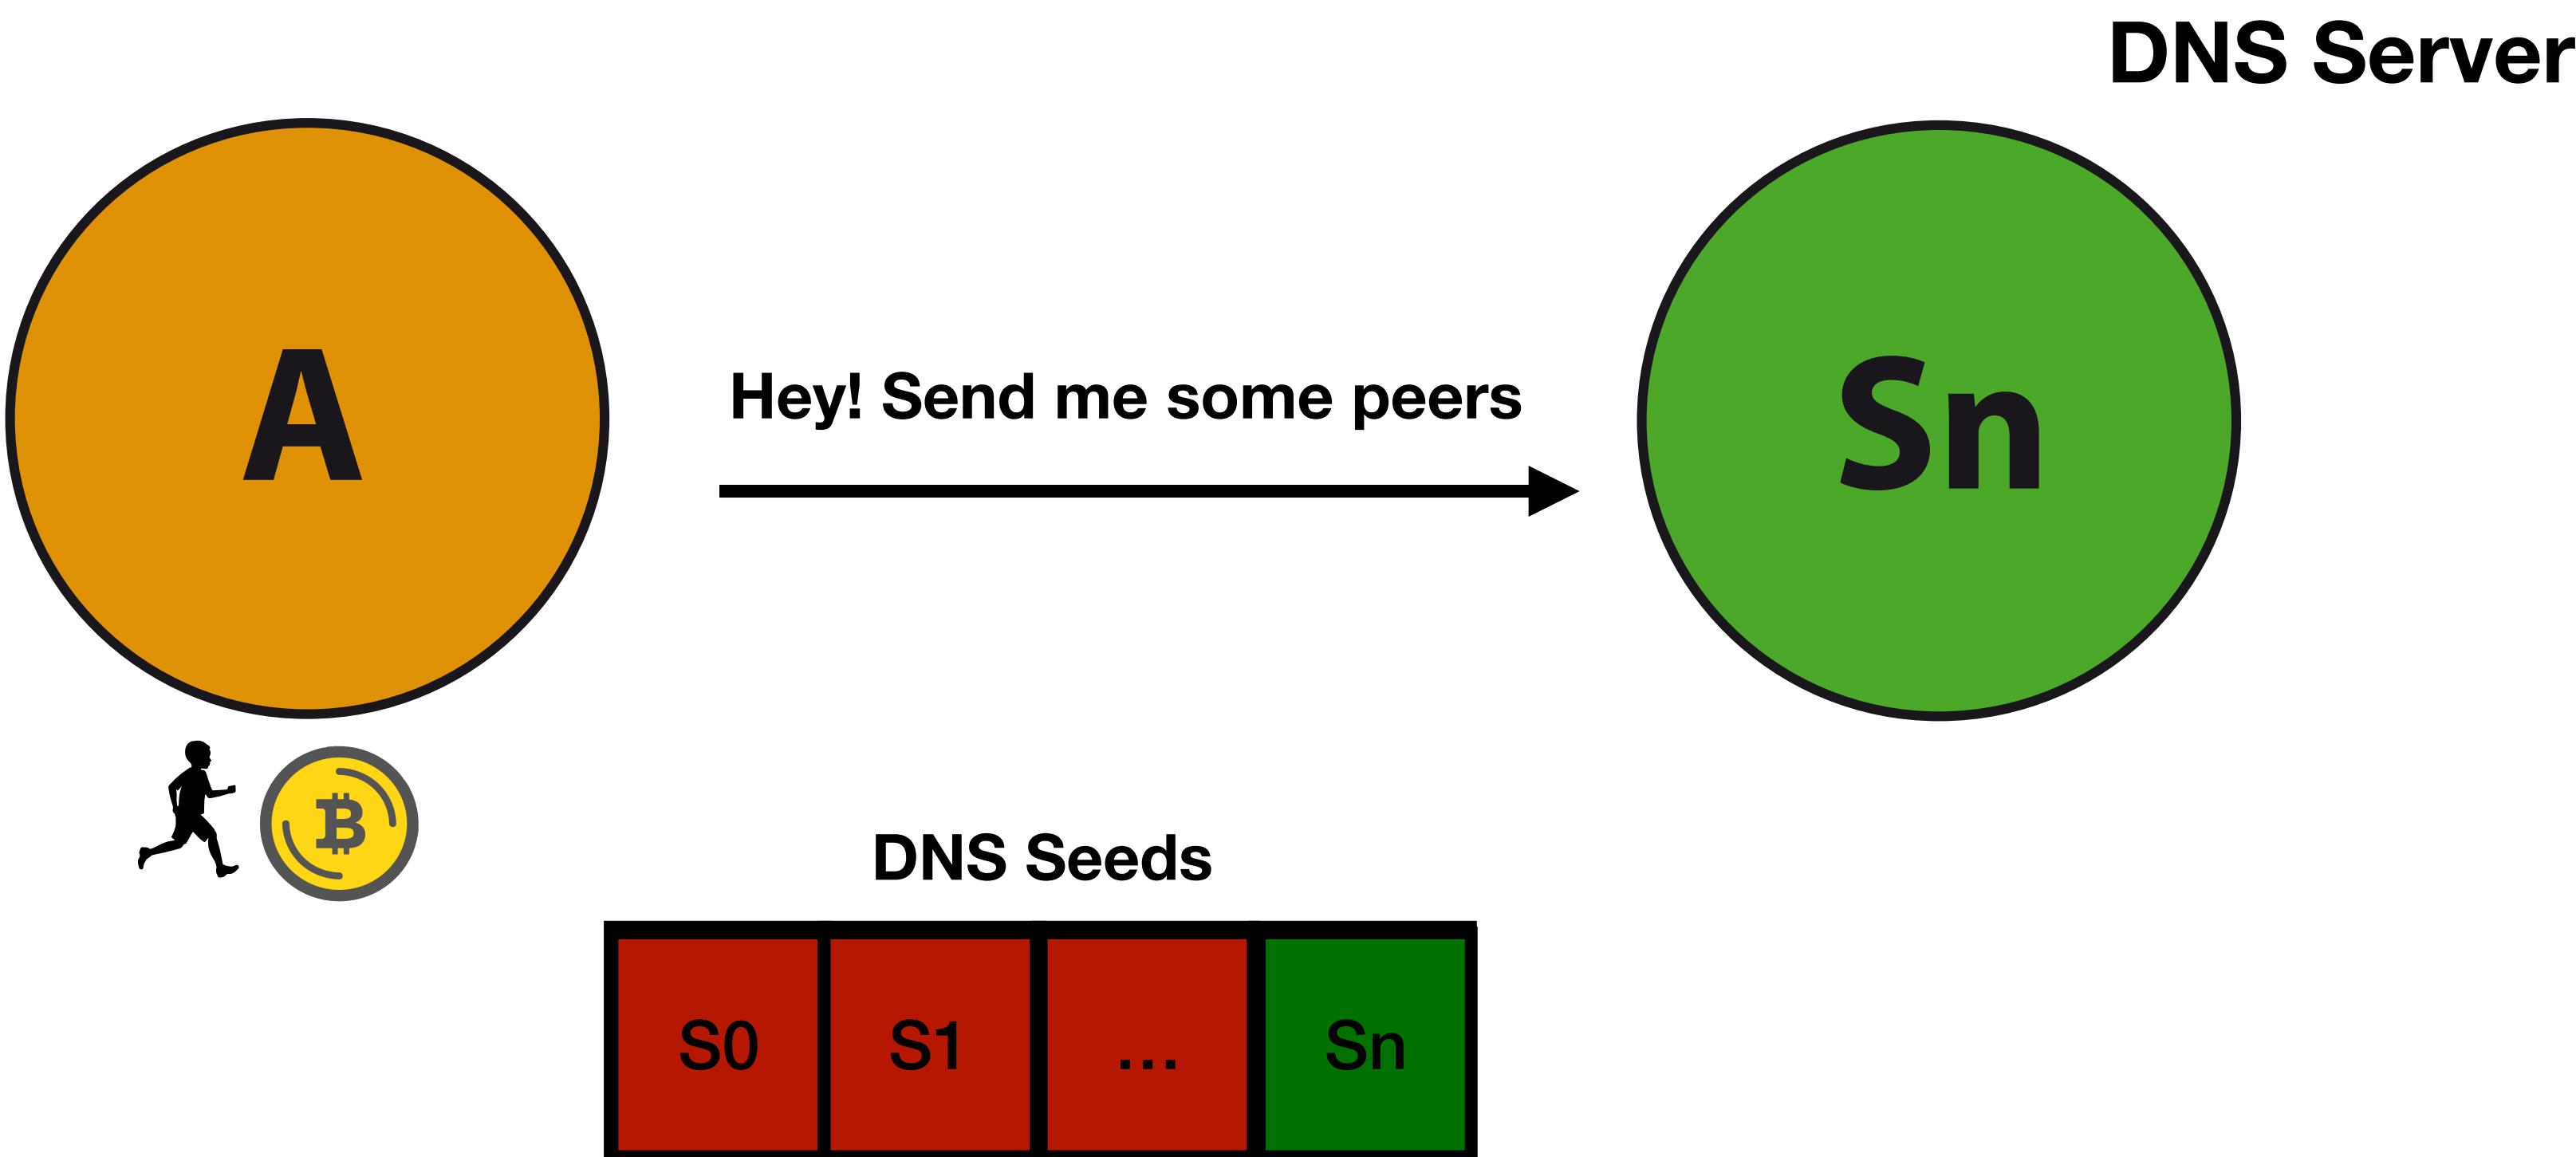
BITCOIN PEER DISCOVERY



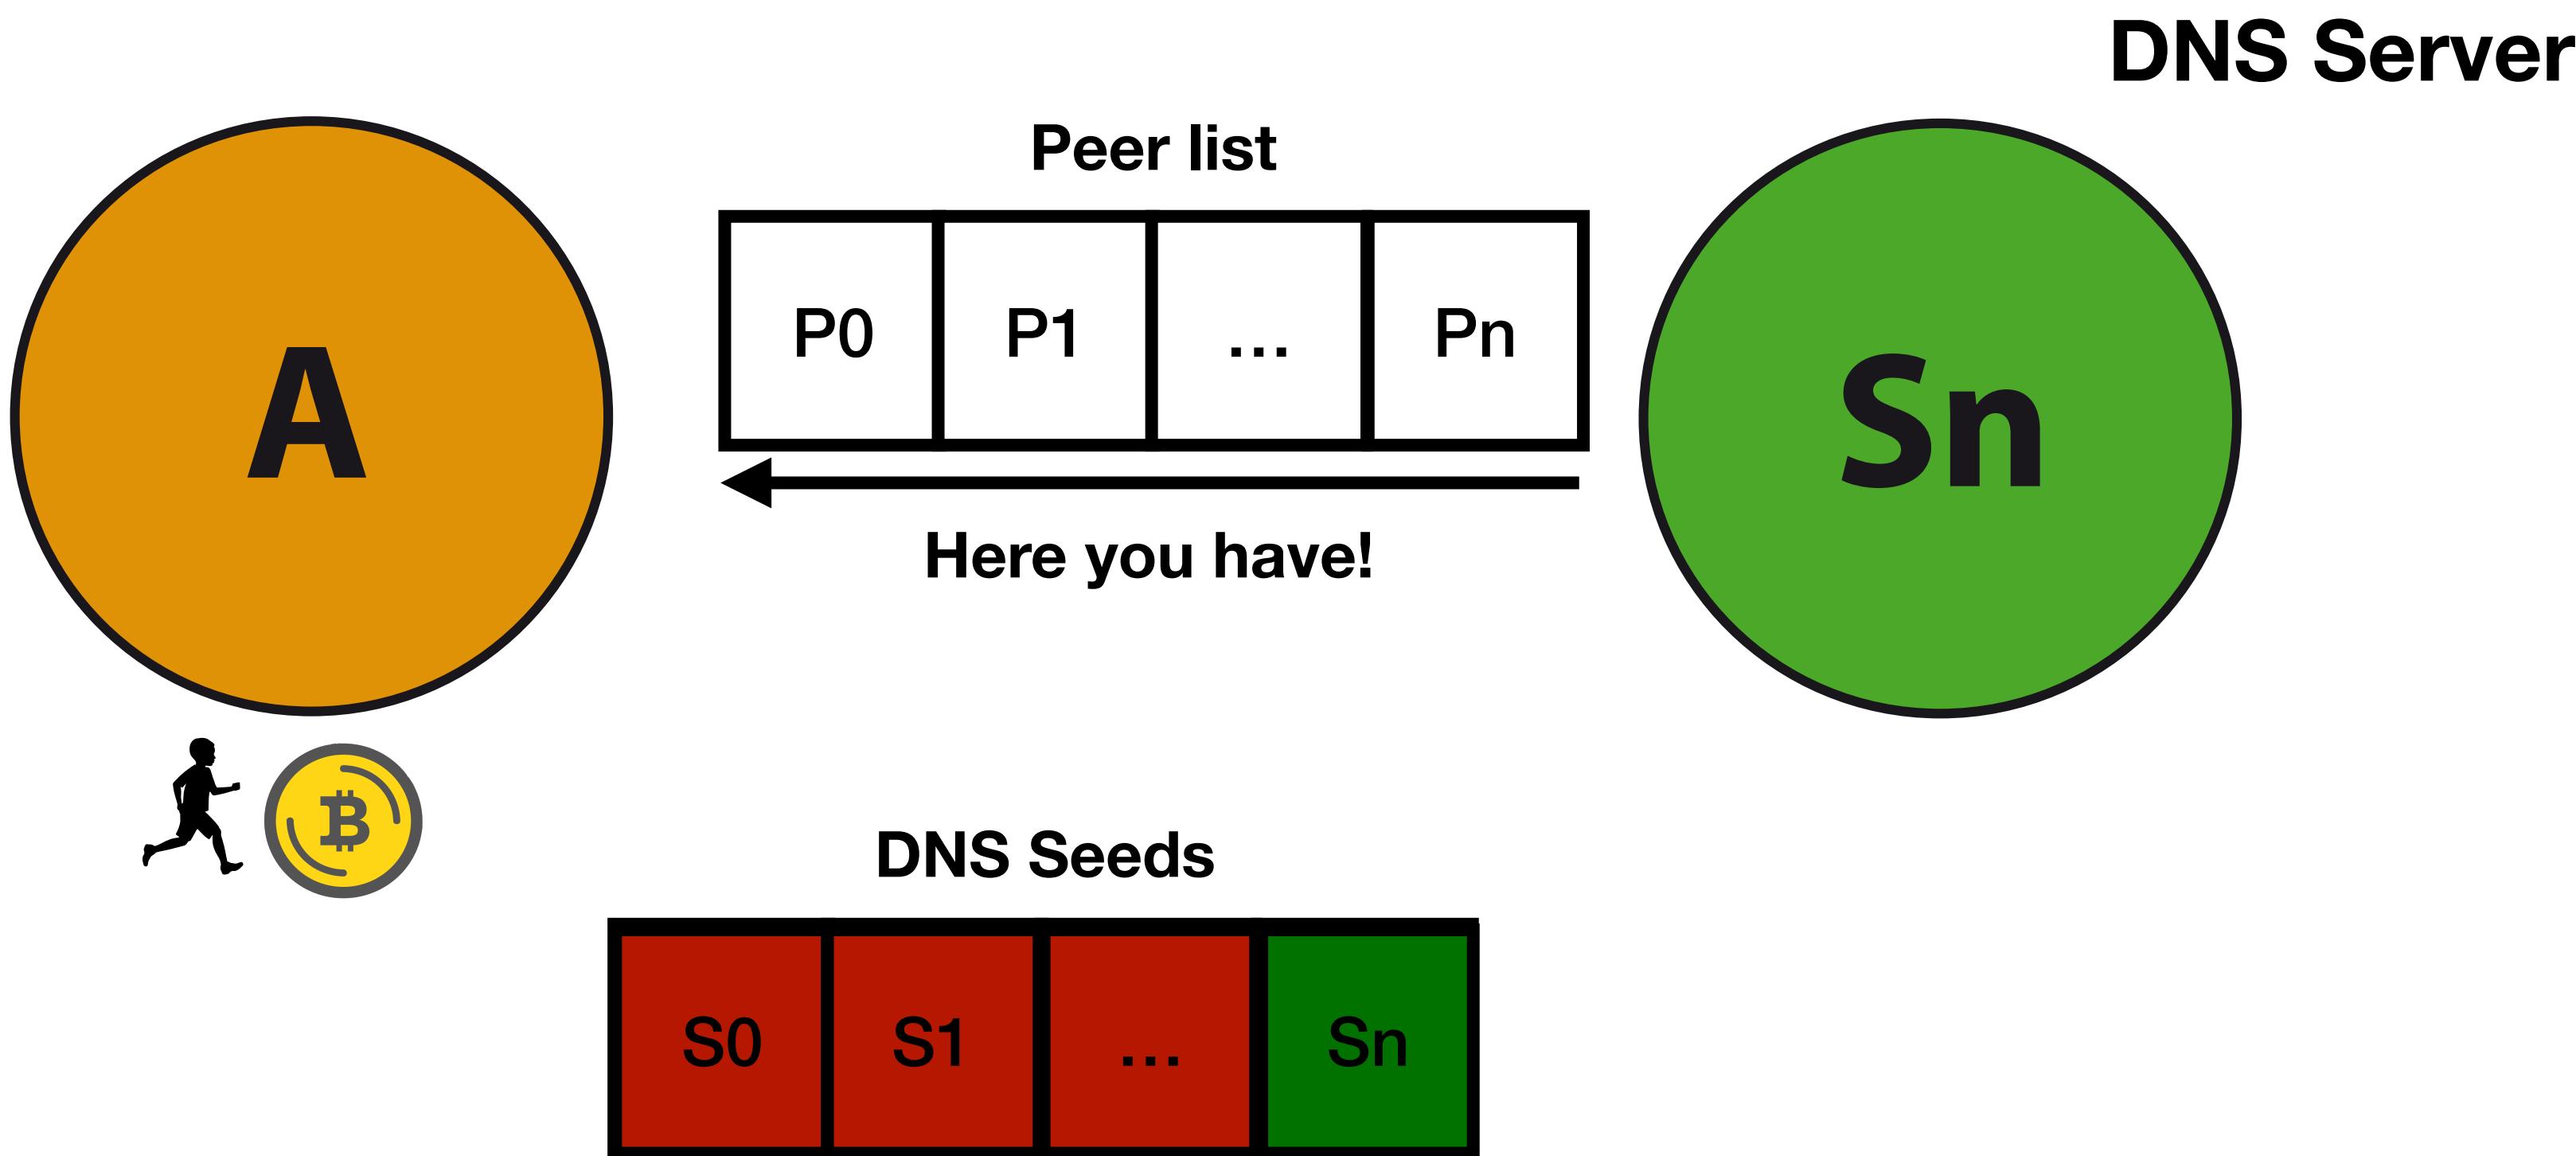
BITCOIN PEER DISCOVERY



BITCOIN PEER DISCOVERY

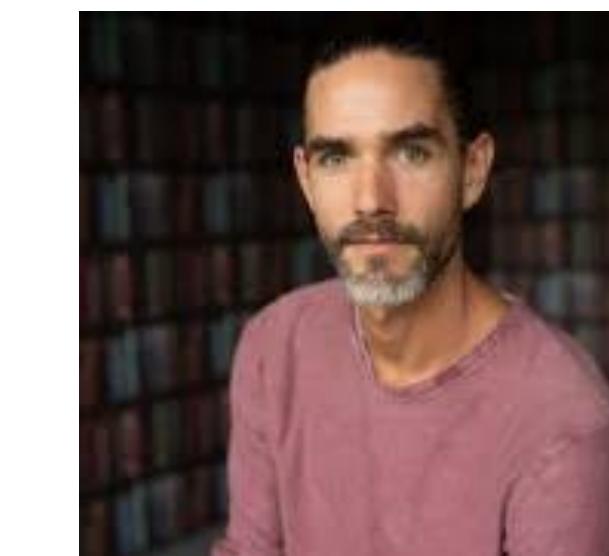


BITCOIN PEER DISCOVERY



BITCOIN DNS SERVER HOSTS

```
vSeeds.emplace_back("seed.bitcoin.sipa.be"); // Pieter Wuille  
vSeeds.emplace_back("dnsseed.bluematt.me"); // Matt Corallo  
vSeeds.emplace_back("dnsseed.bitcoin.dashjr.org"); // Luke Dashjr  
vSeeds.emplace_back("seed.bitcoinstats.com"); // Christian Decker  
vSeeds.emplace_back("seed.bitcoin.jonasschnelli.ch"); // Jonas Schnelli  
vSeeds.emplace_back("seed.btc.petertodd.org"); // Peter Todd  
vSeeds.emplace_back("seed.bitcoin.sprovoost.nl"); // Sjors Provoost
```

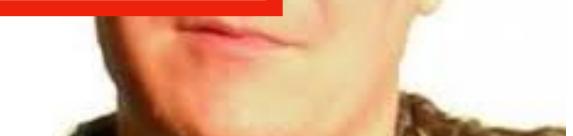


BITCOIN DNS SERVER HOSTS



If DNS seeds do not work,
a node will try to connect
to a hardcoded list of
nodes (fixed seed)

```
vSeeds.emplace_back("seed.bitcoin.sipa.be"); // Pieter Wuille
vSeeds.static SeedSpec6 pnSeed6_main[] = {
    {{0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0xff,0xff,0x02,0x84,0x64,0x2f}, 8333},
    {{0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0xff,0xff,0x05,0x01,0x61,0x04}, 8333},
    {{0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0xff,0xff,0x05,0x27,0xae,0x74}, 8333},
    {{0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0xff,0xff,0x05,0x2d,0x4f,0x0e}, 8333},
    {{0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0xff,0xff,0x05,0x35,0x10,0x85}, 8333},
    {{0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0x00,0xff,0xff,0x05,0x65,0x8b,0xa6}, 8333},
    ...
}
```



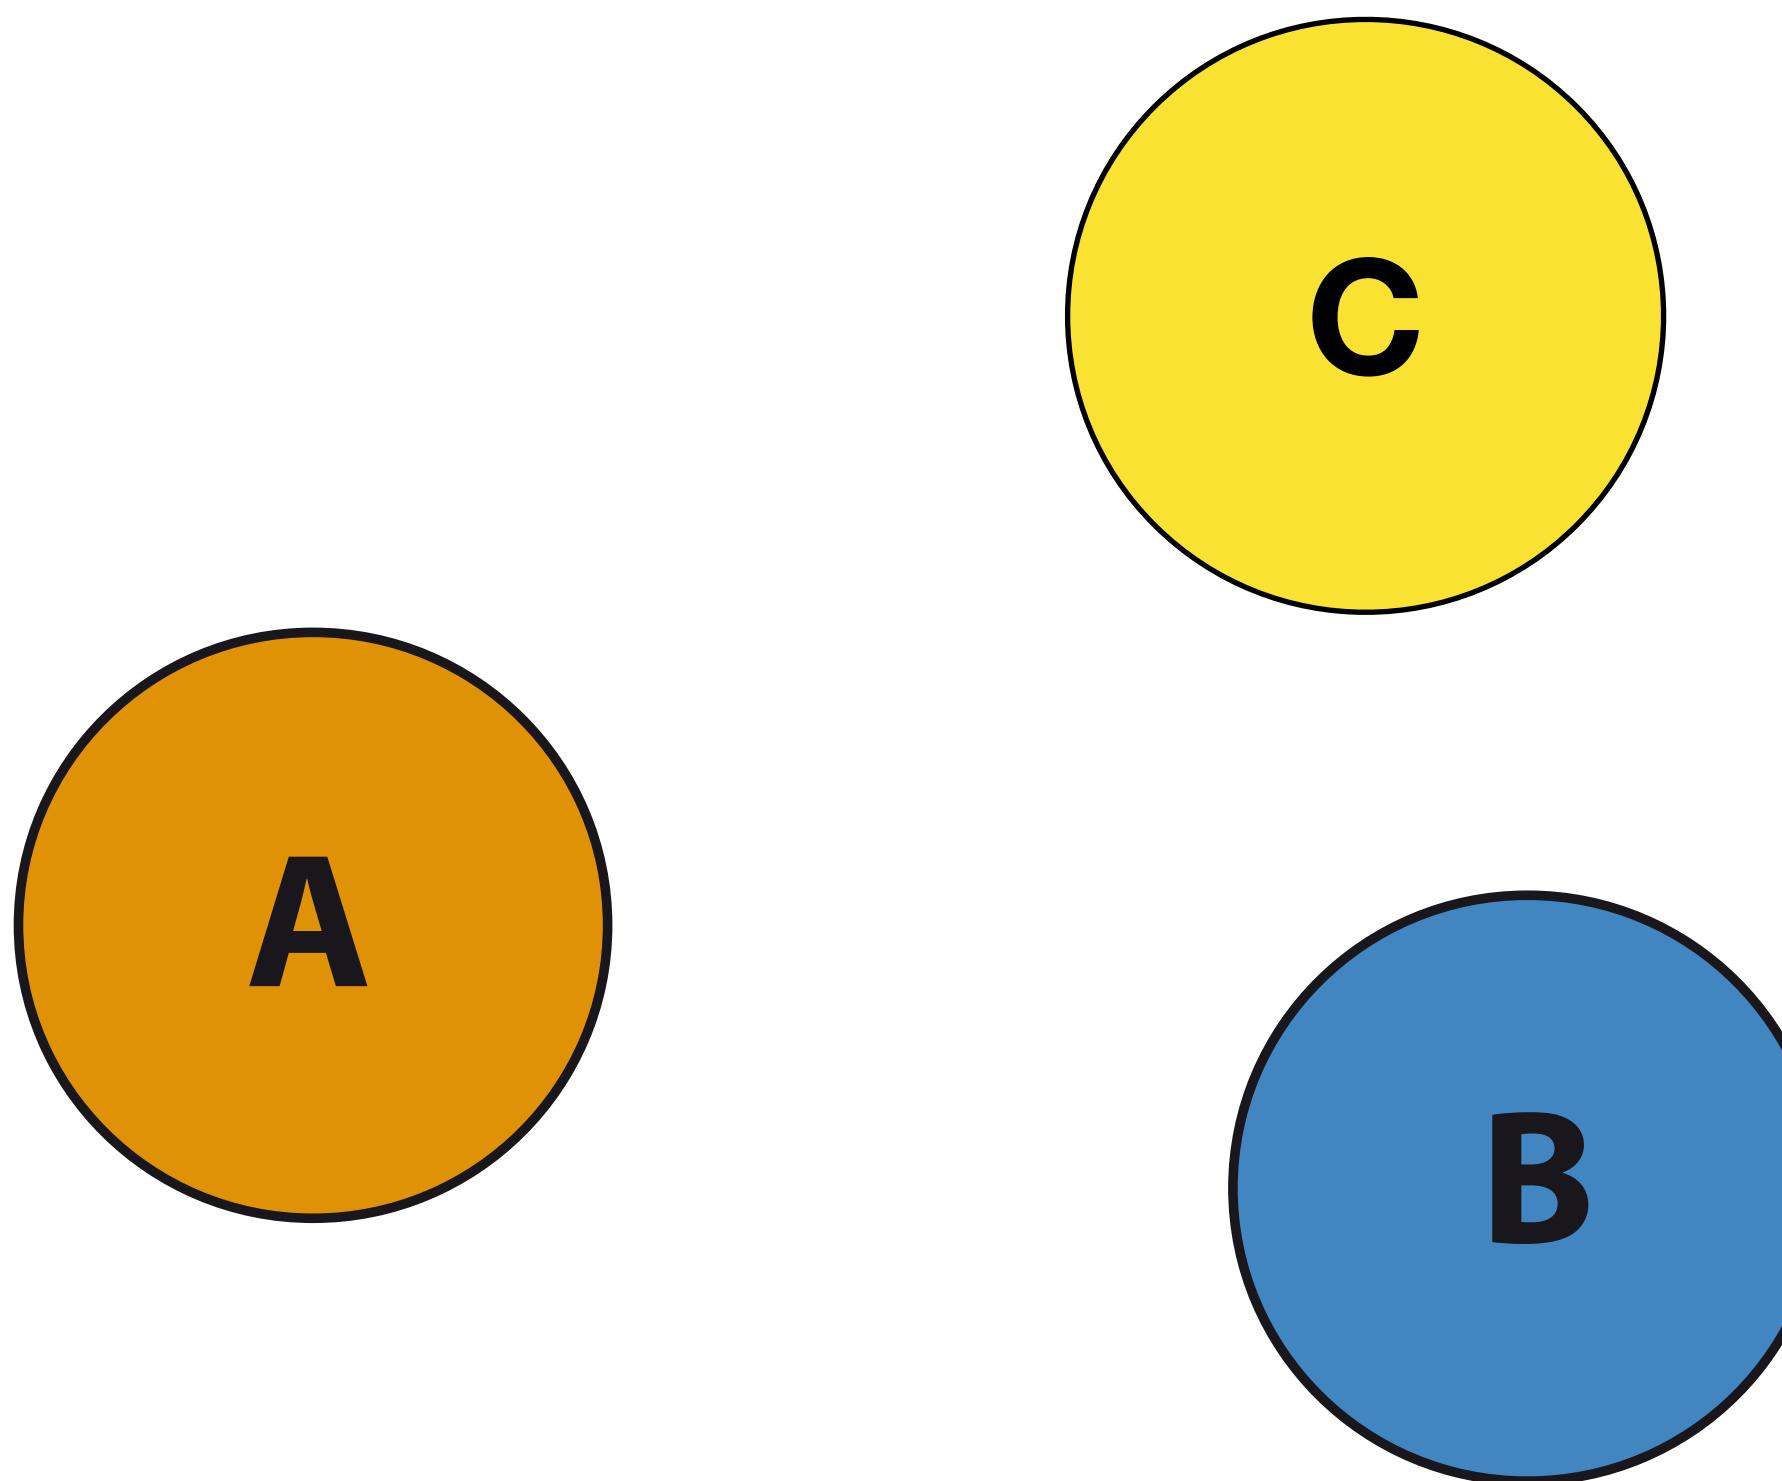
BITCOIN P2P BOOTSTRAPPING (RECAP)

A node bootstraps with no known peers

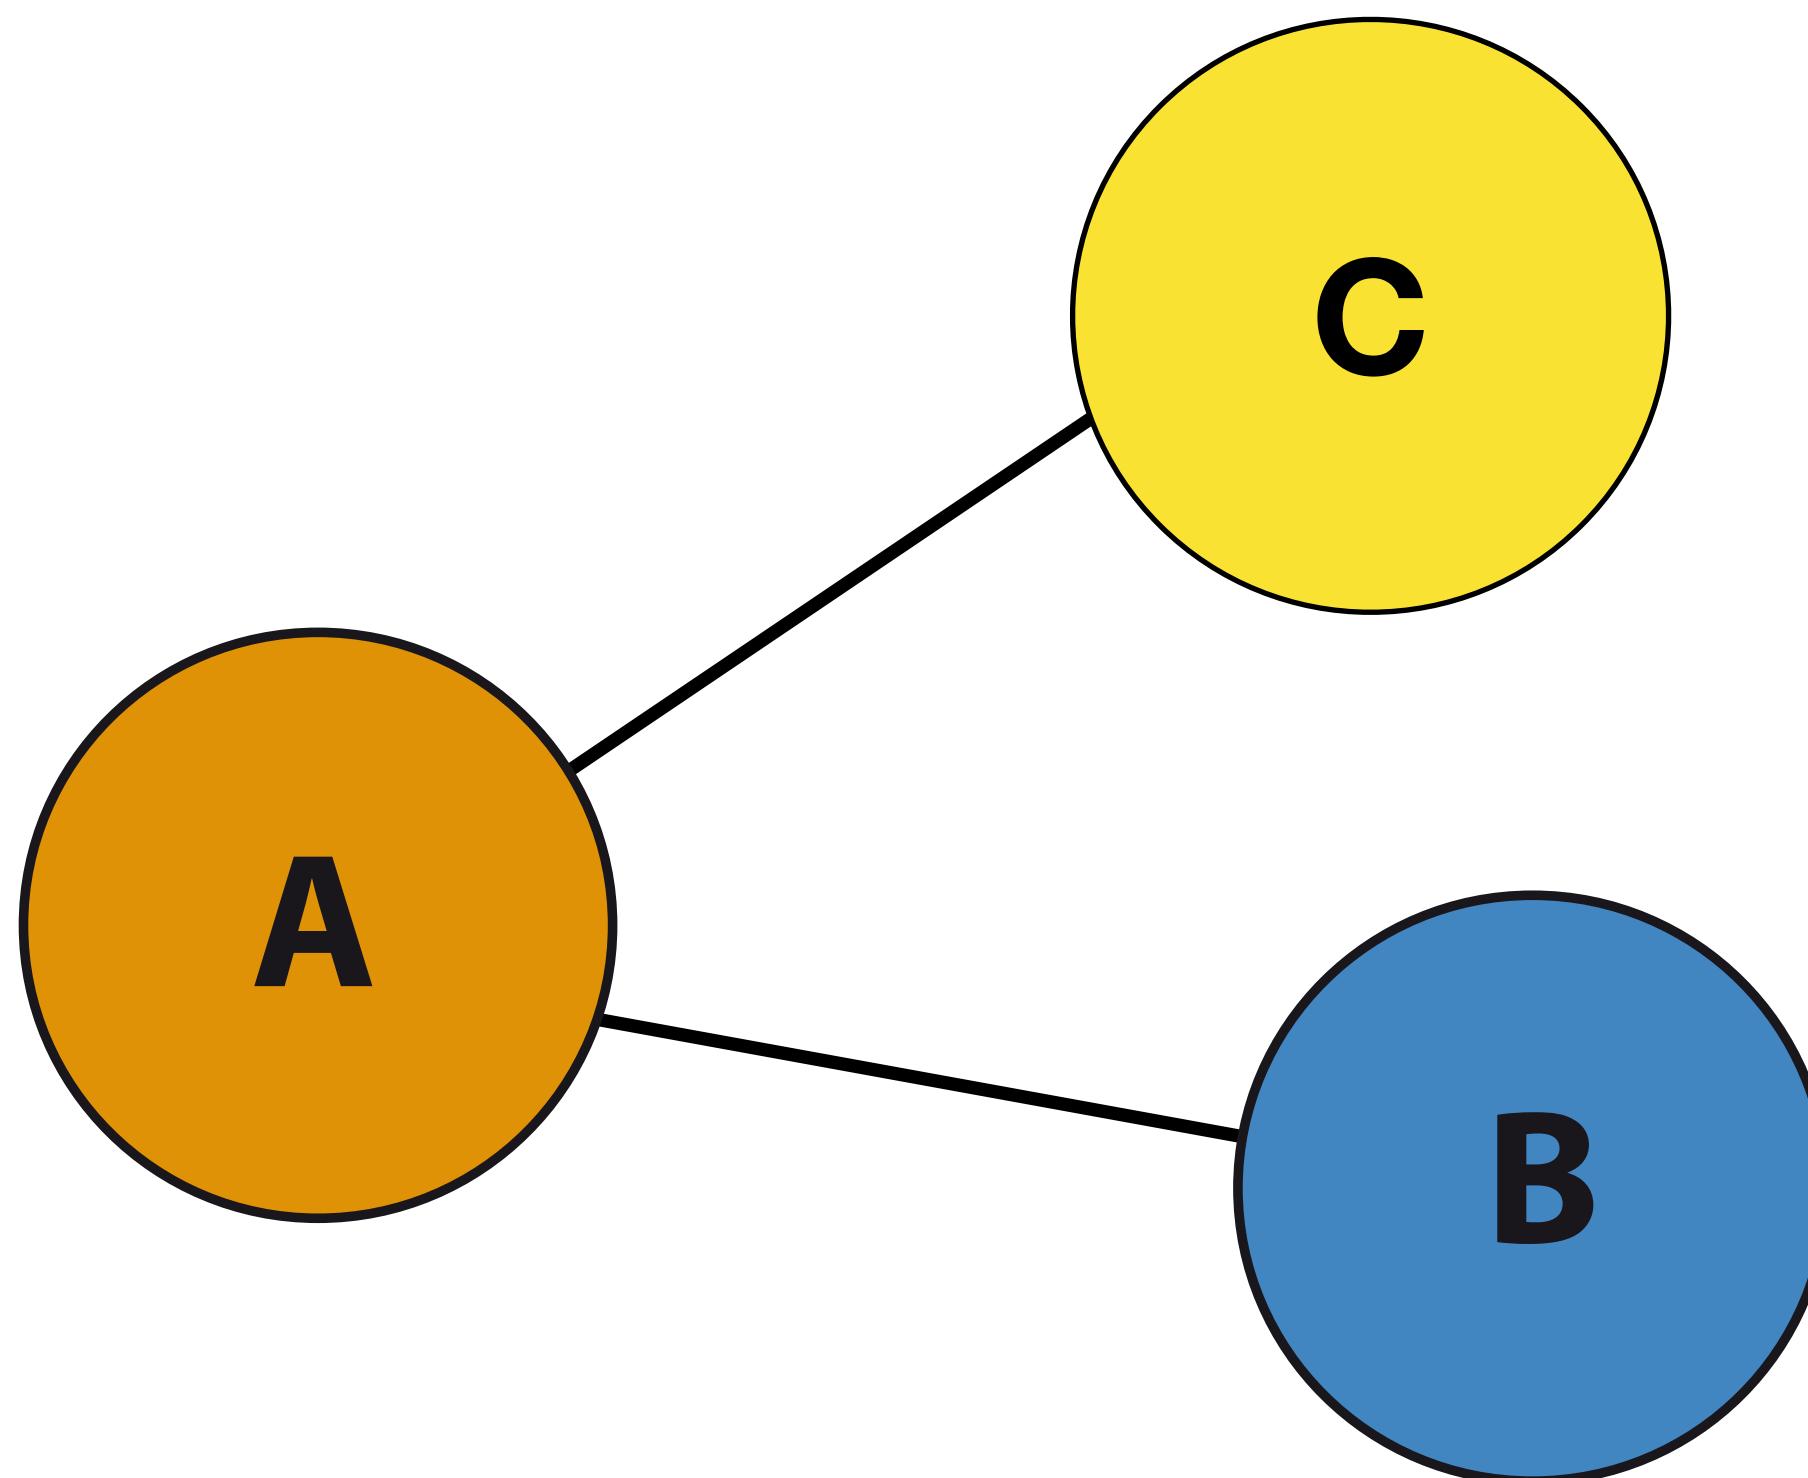
First it tries to query a list of well known DNS seeds

As a **last resource** it uses a hardcoded seed

POPULATING THE PEERS DATABASE

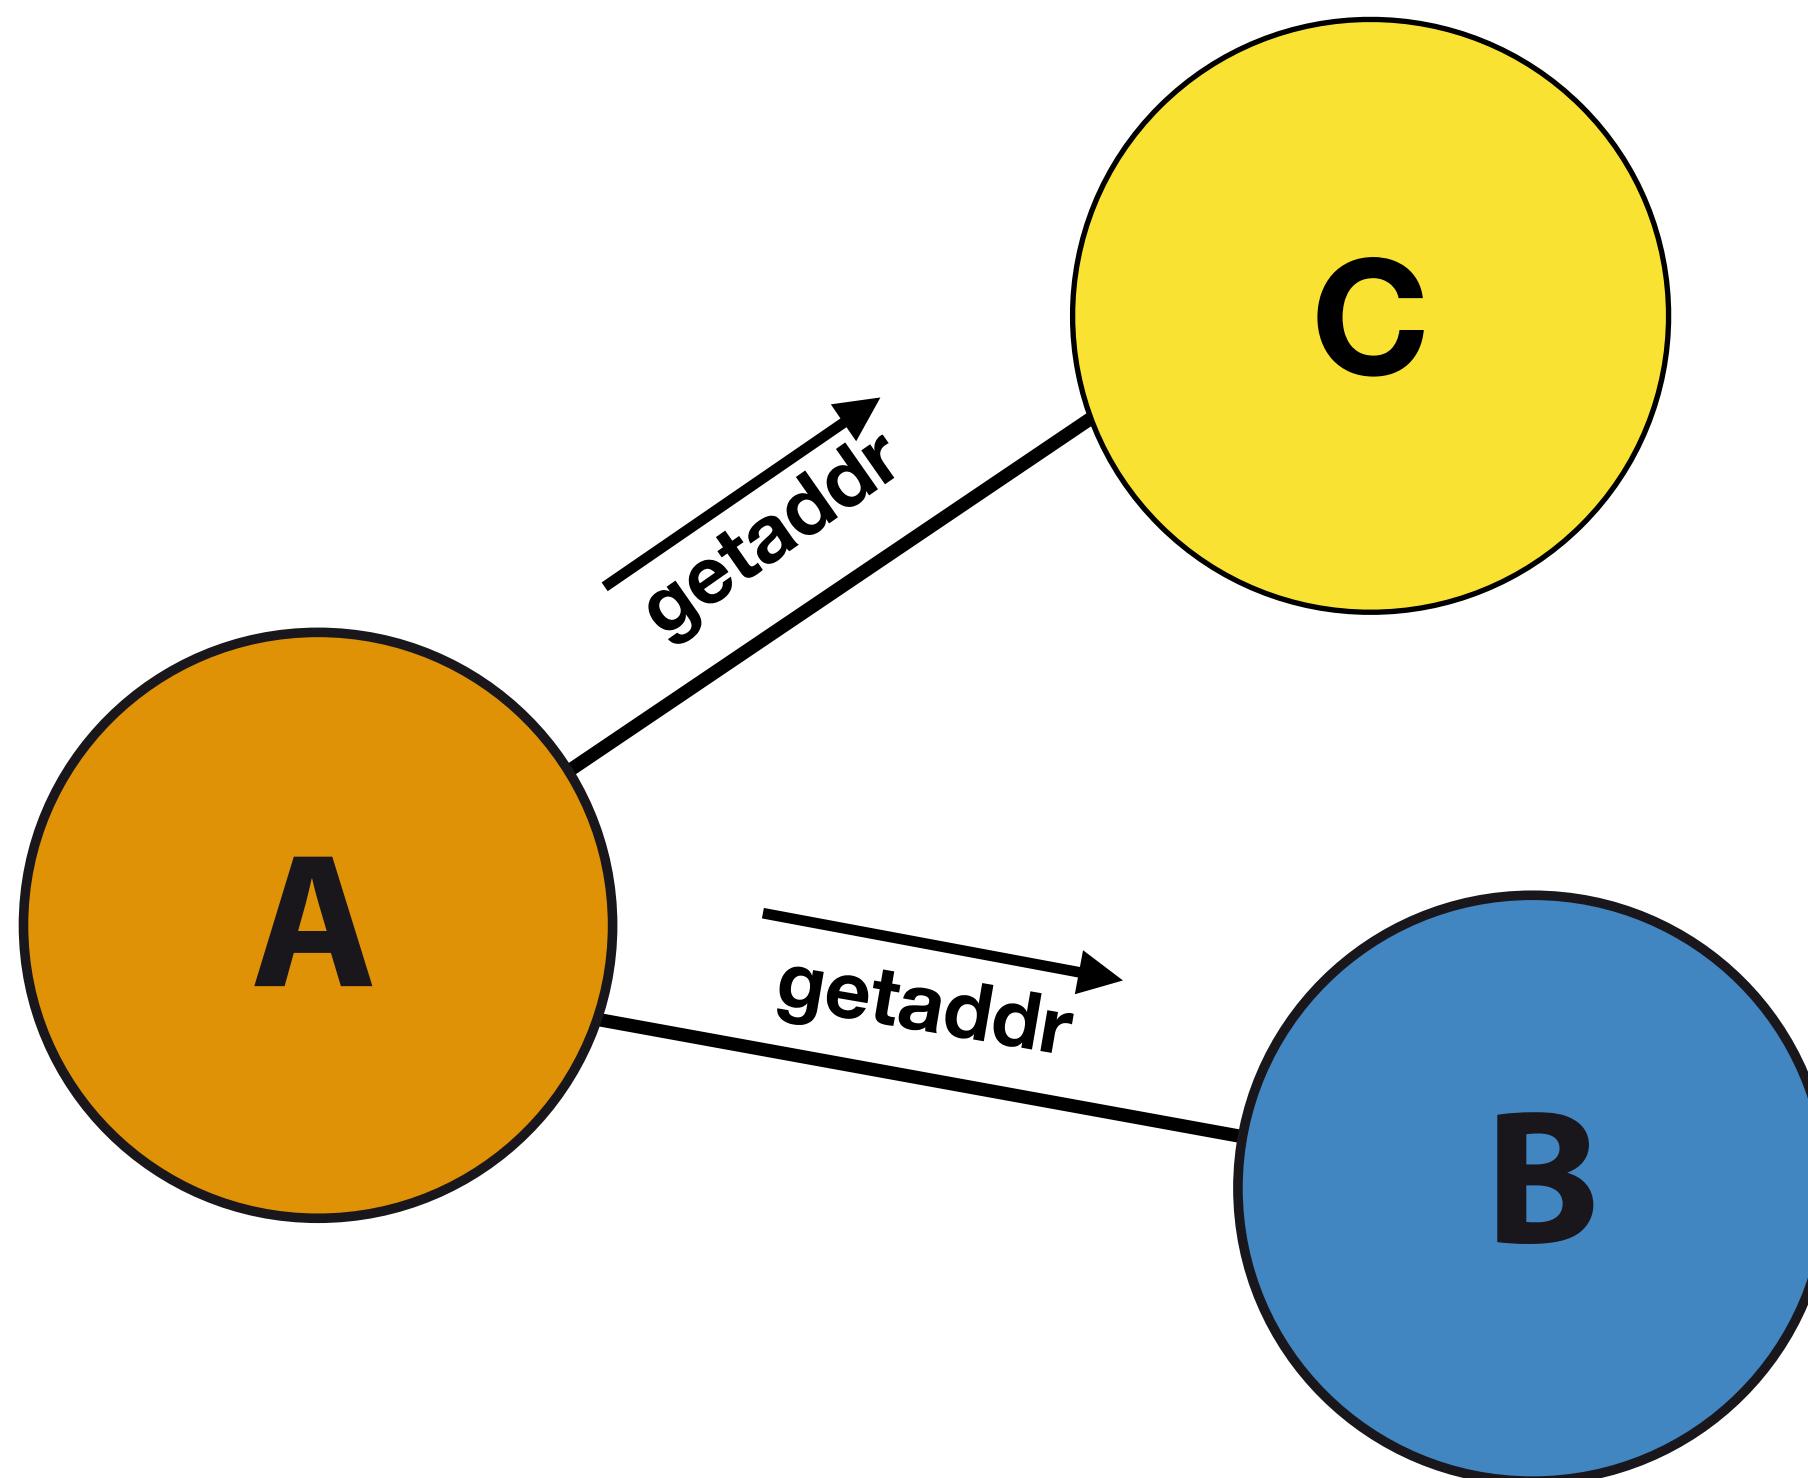


POPULATING THE PEERS DATABASE



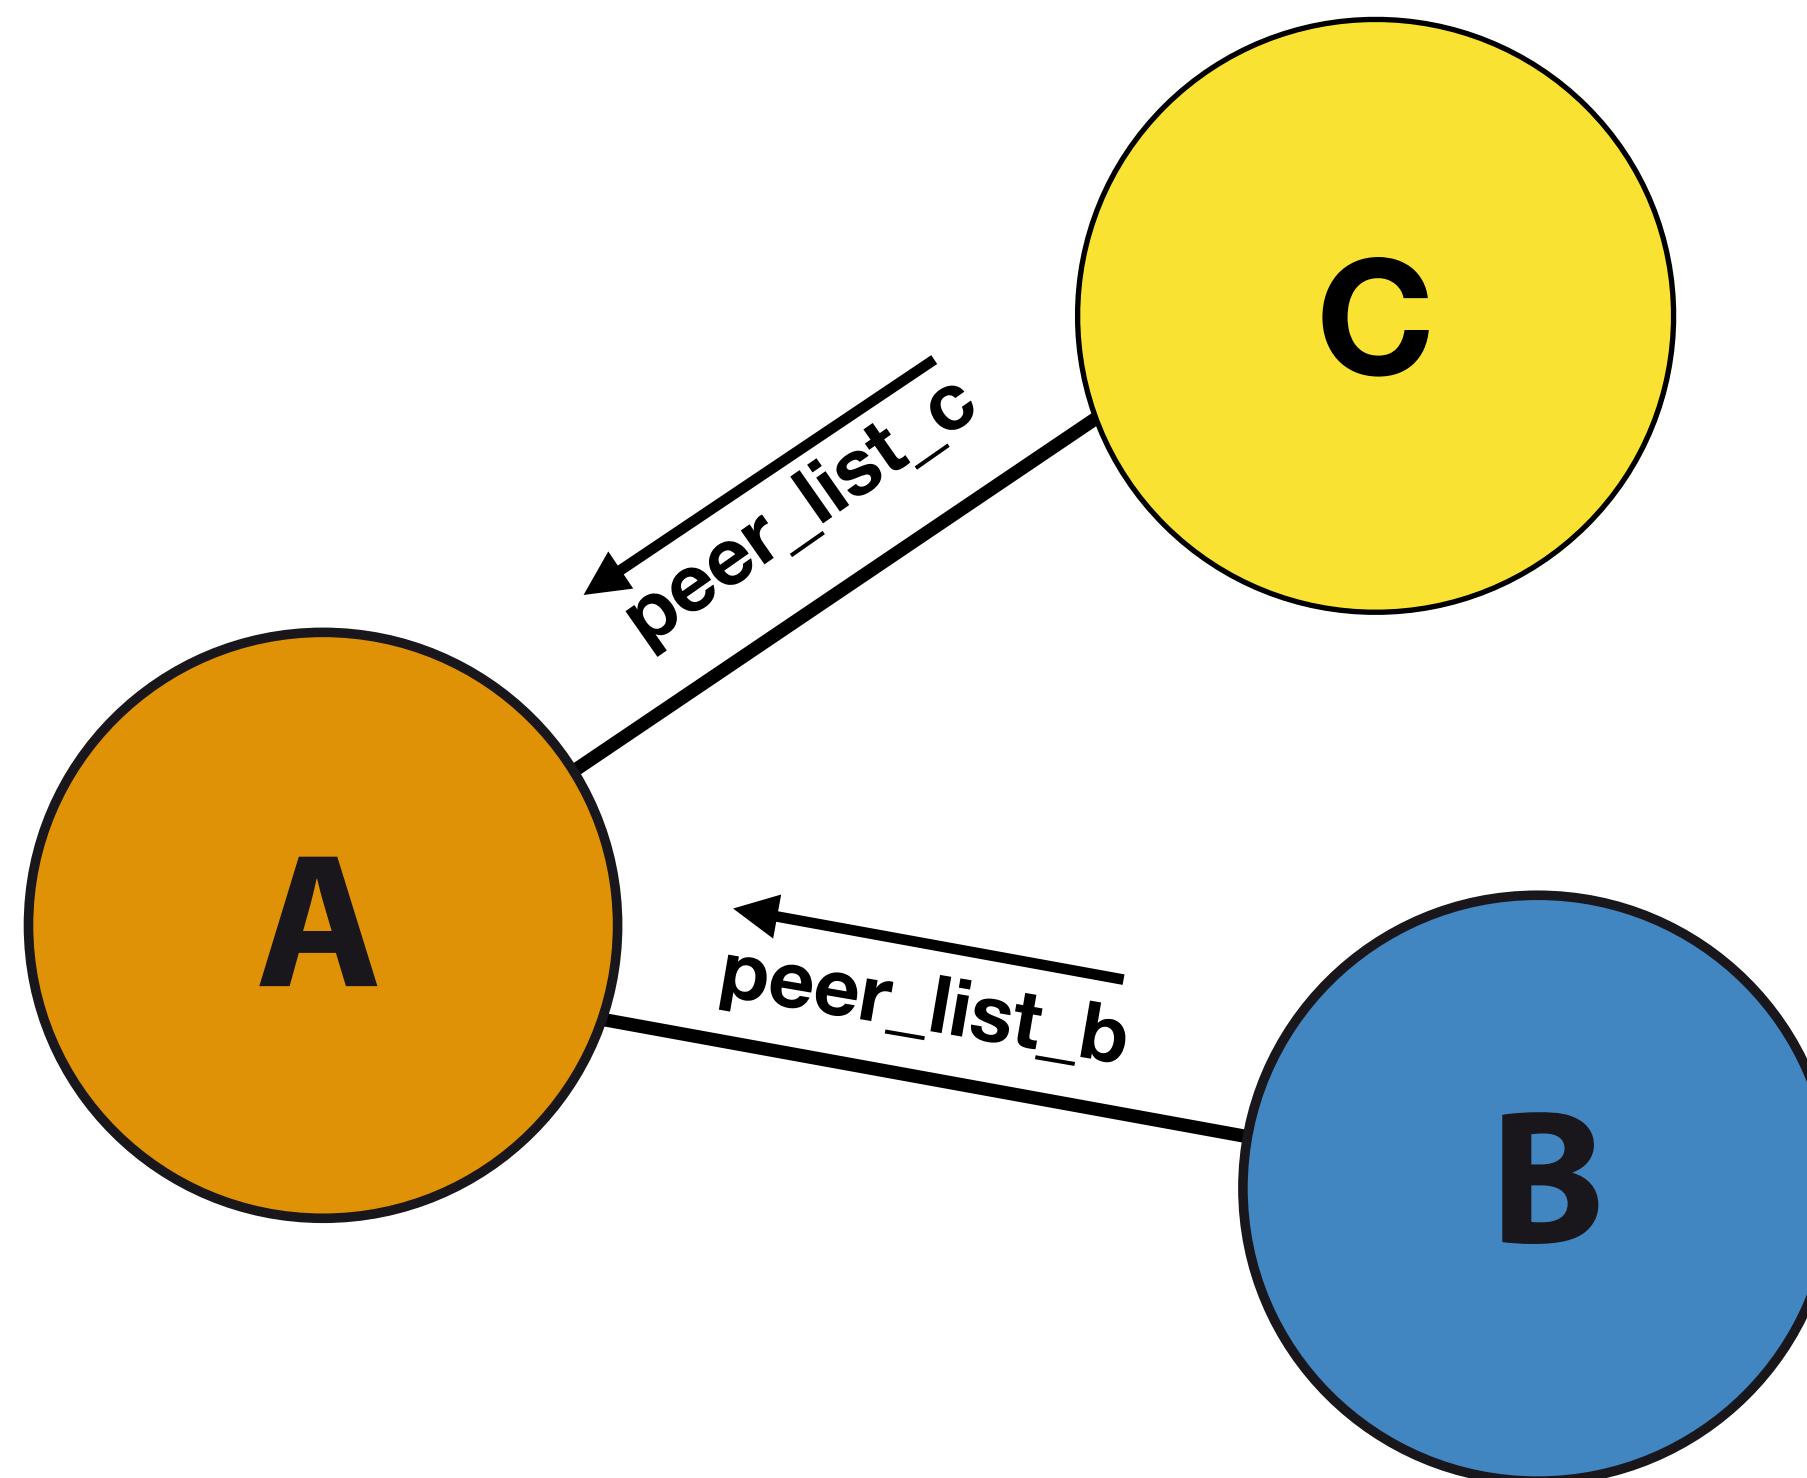
- A connects a subset of peers from the ones learned from the DNS seeds

POPULATING THE PEERS DATABASE



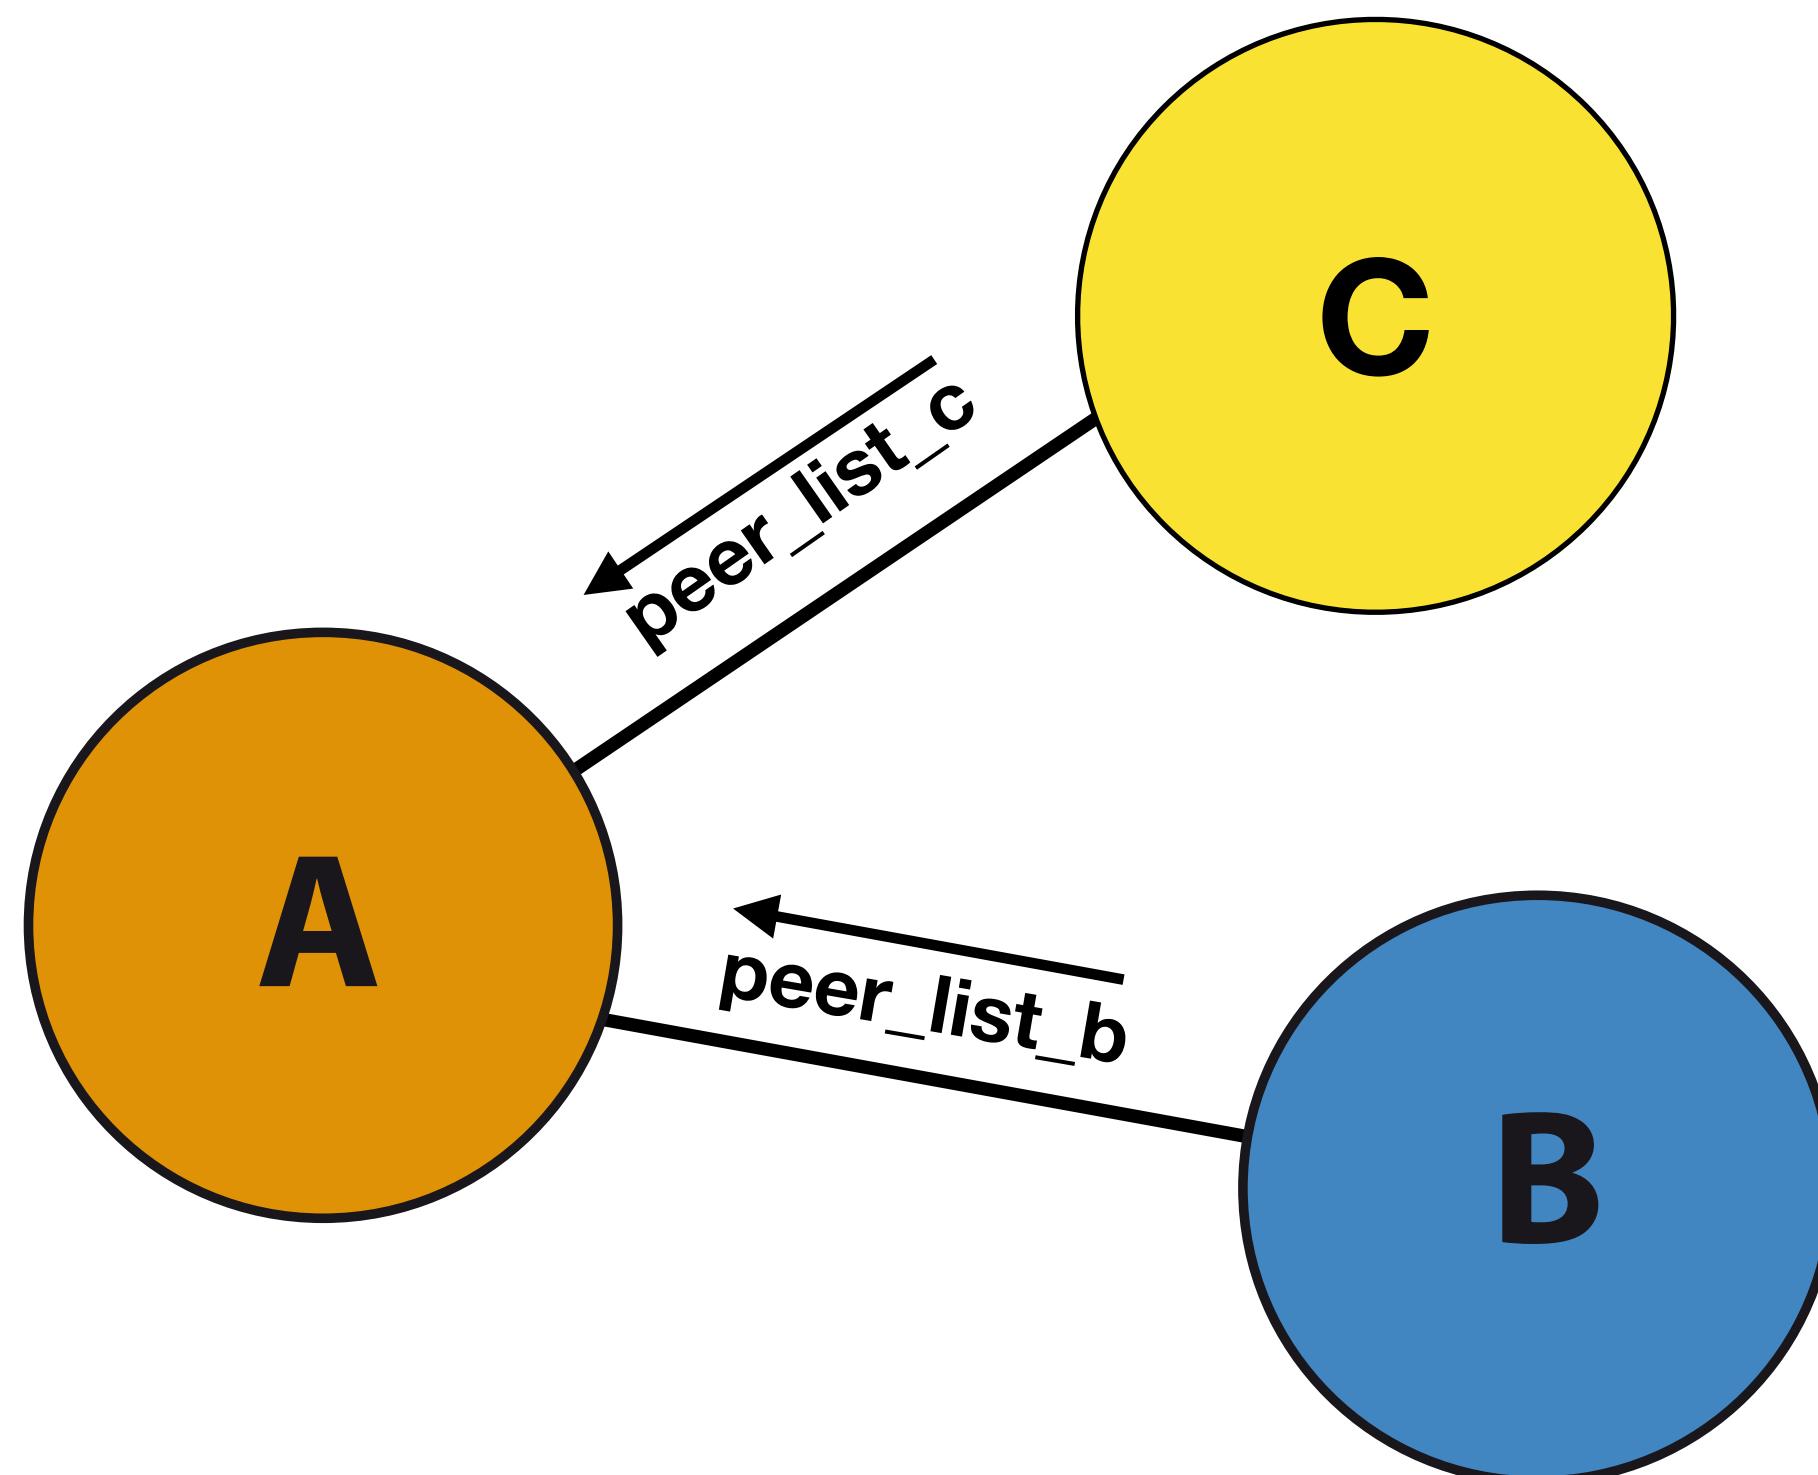
- A connects a subset of peers from the ones learned from the DNS seeds
- A requests more peers to his neighbors (**getaddr**)

POPULATING THE PEERS DATABASE



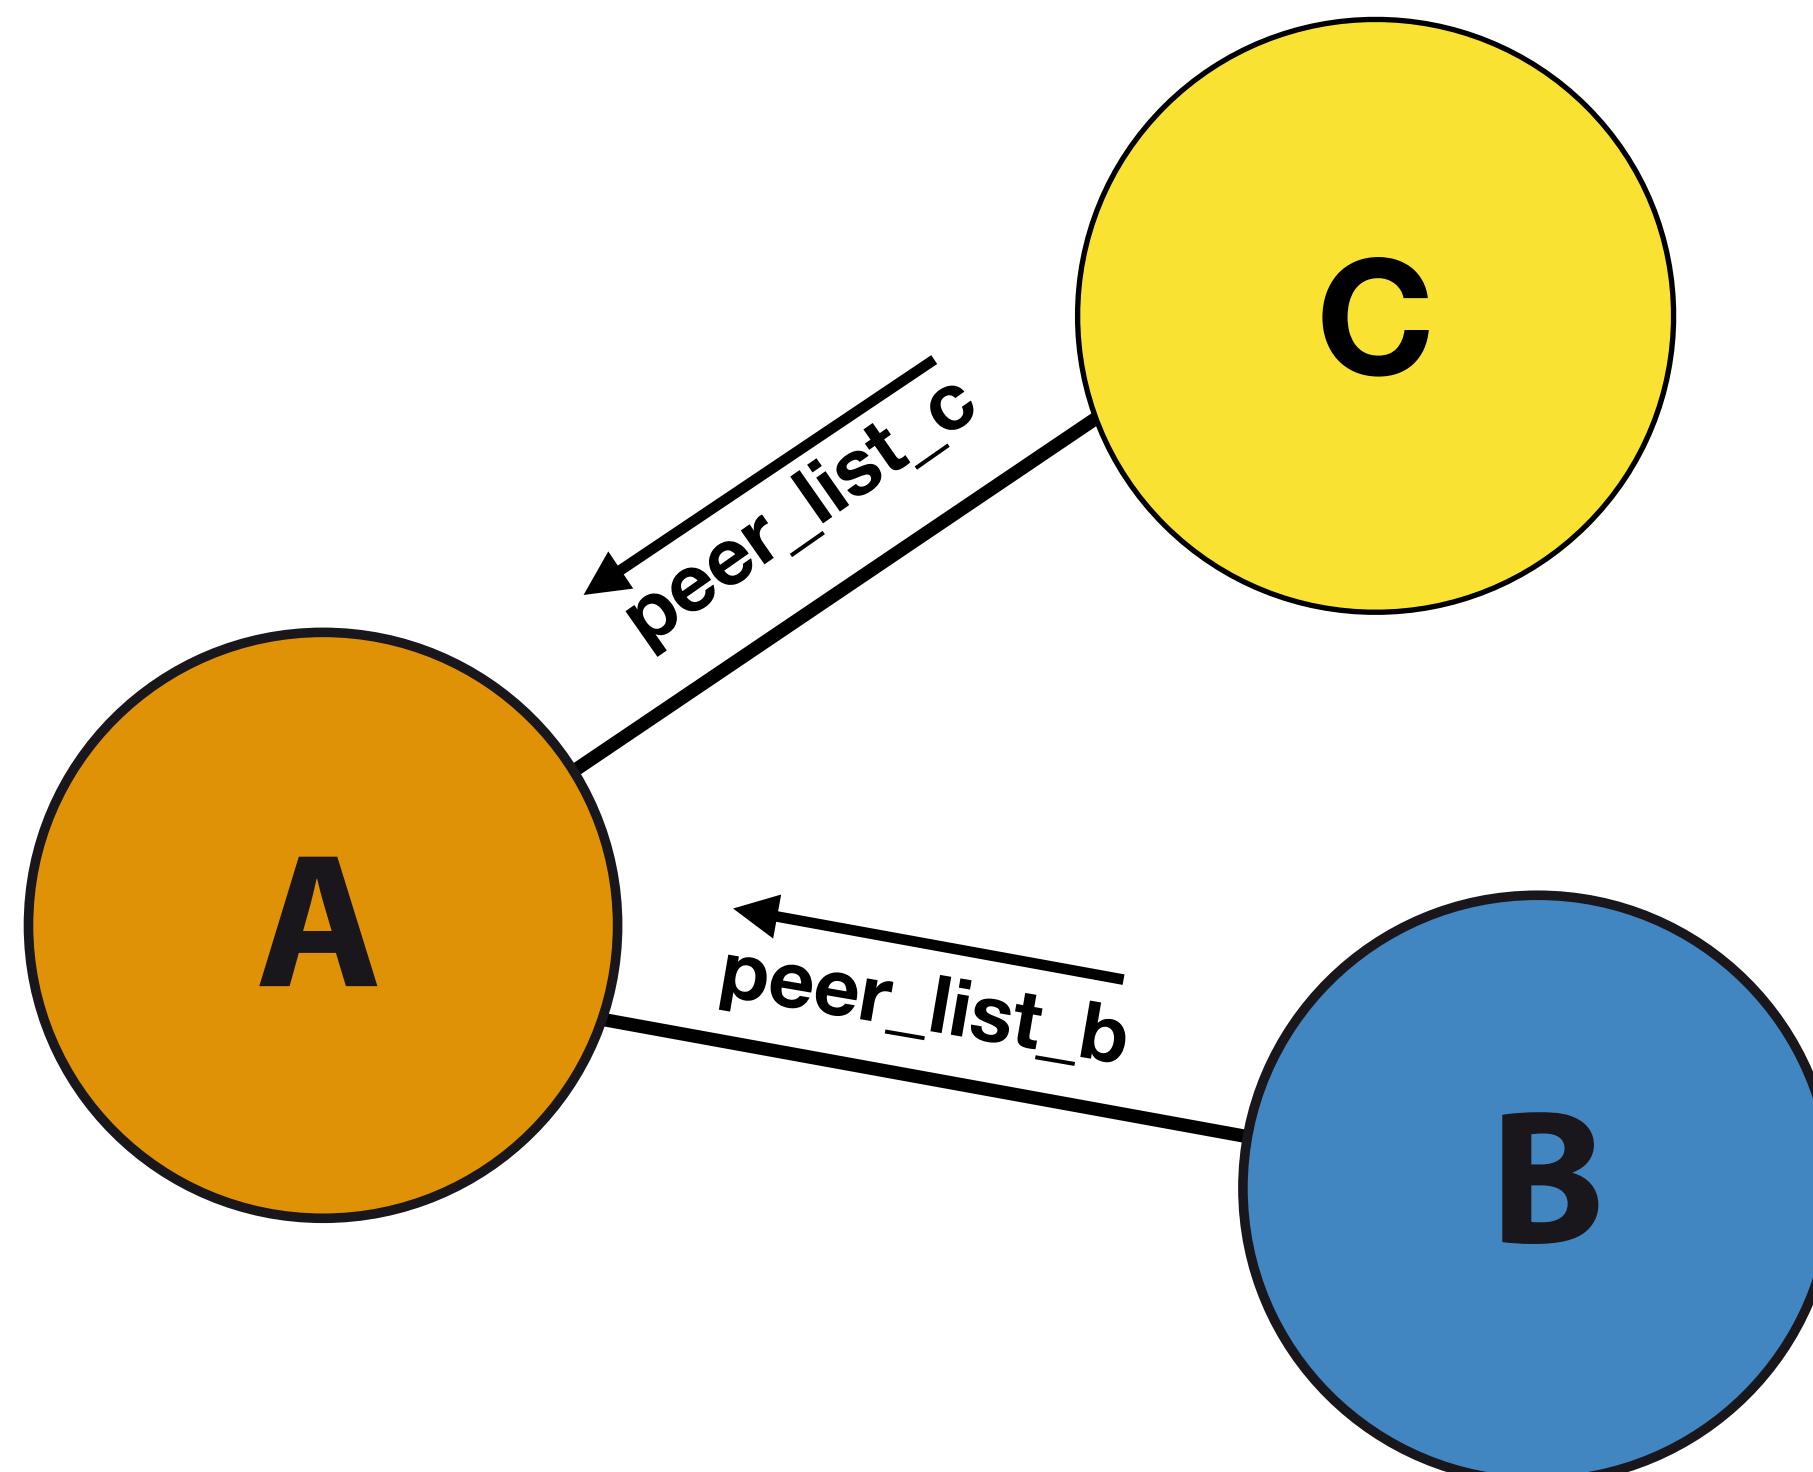
- A connects a subset of peers from the ones learned from the DNS seeds
- A requests more peers to his neighbors (**getaddr**)
- Peers reply with some addresses they know about (**addr, up to 1000 addresses**)

POPULATING THE PEERS DATABASE



- A connects a subset of peers from the ones learned from the DNS seeds
- A requests more peers to his neighbors (**getaddr**)
- Peers reply with some addresses they know about (**addr, up to 1000 addresses**)
- A adds the new addresses to its peers database (or updates the existing ones)

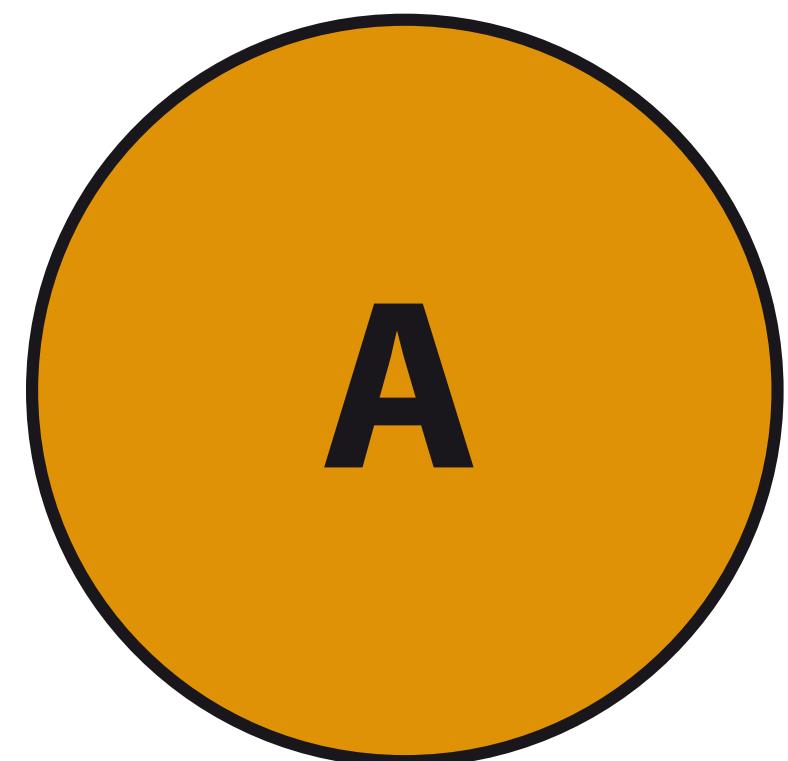
POPULATING THE PEERS DATABASE



- A connects a subset of peers from the ones learned from the DNS seeds
- A requests more peers to his neighbors (**getaddr**)
- Peers reply with some addresses they know about (**addr, up to 1000 addresses**)
- A adds the new addresses to its peers database (or updates the existing ones)
- The database is known as **the addrman**

A's addrman = A's addrman U peer_list_c U peer_list_b

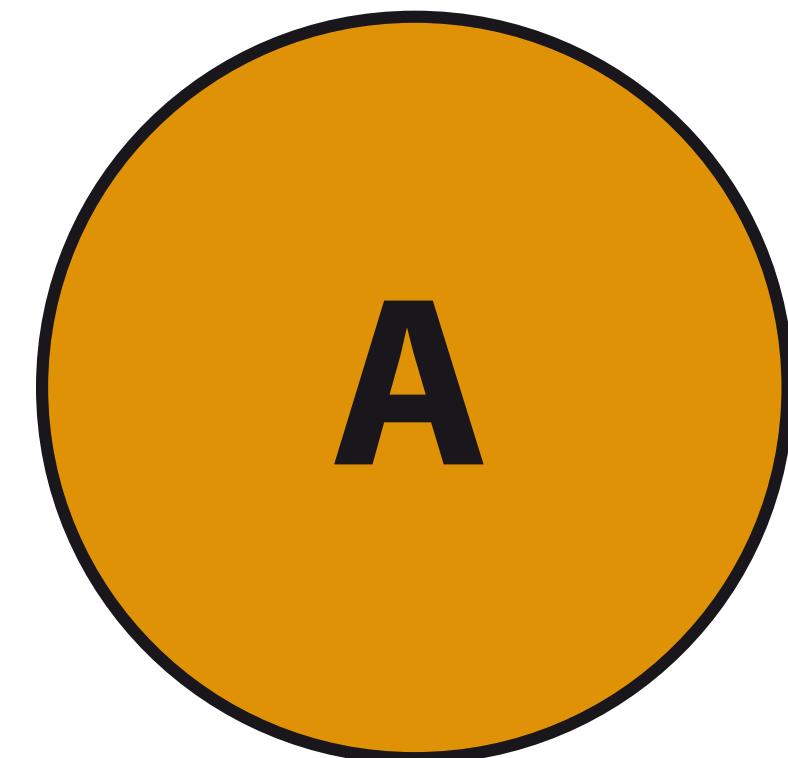
INCOMING/OUTGOING CONNECTIONS



Peer database (addrman)

P0	P1	...	Pn
----	----	-----	----

INCOMING/OUTGOING CONNECTIONS

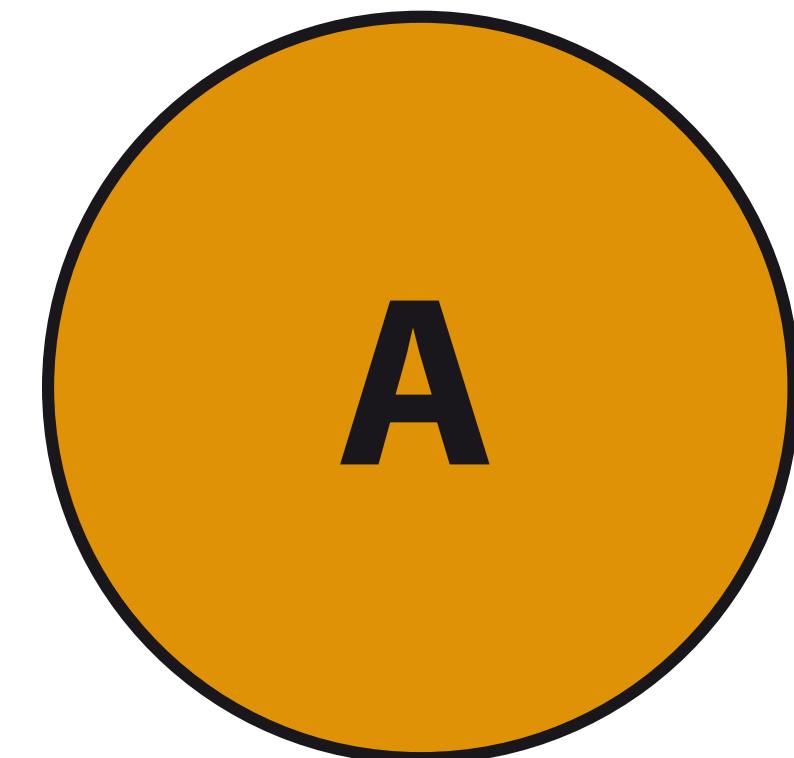


- During bootstrap, a node will start some **outgoing connections** with peers it has learnt about (**8 by default**) and tries to maintain them

Peer database (addrman)

P0	P1	...	Pn
----	----	-----	----

INCOMING/OUTGOING CONNECTIONS



- During bootstrap, a node will start some **outgoing connections** with peers it has learnt about (**8 by default**) and tries to maintain them
- A node will also accept **some incoming** connections (**117 by default**)

Peer database (addrman)

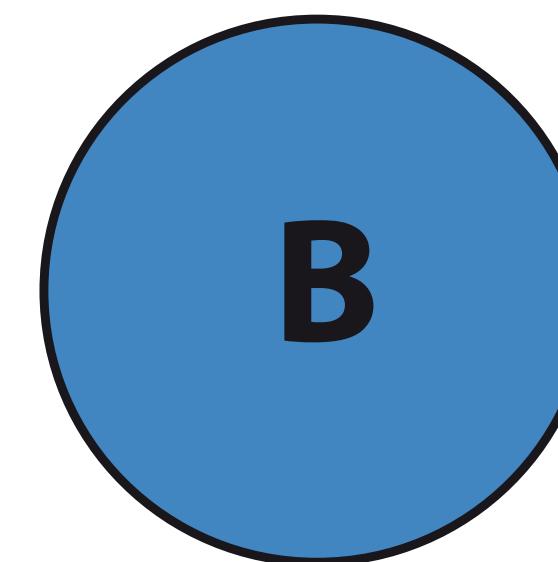
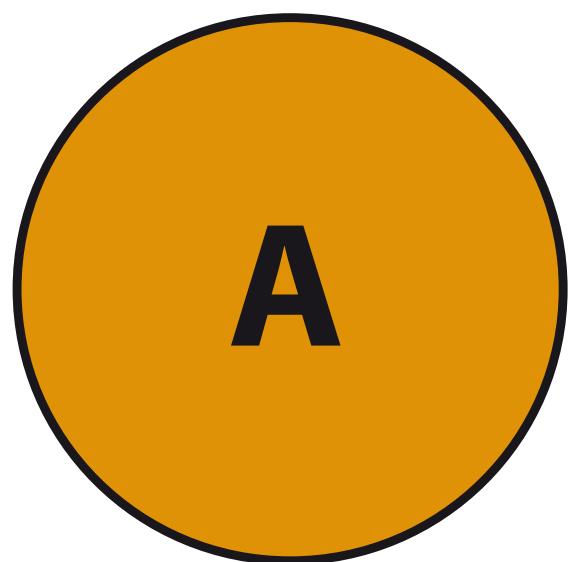
P0	P1	...	Pn
----	----	-----	----

ADDRESS PROPAGATION (1/2)

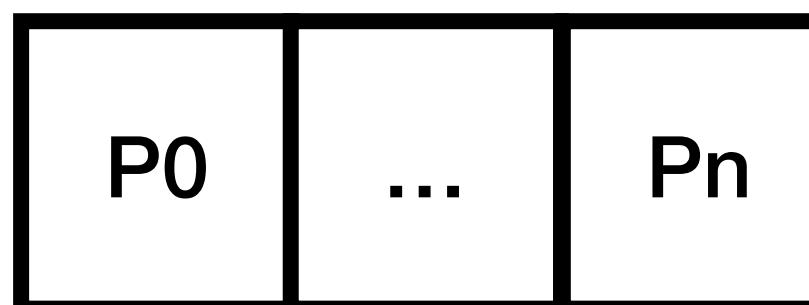
How does a node announce his presence to the rest of the network?

ADDRESS PROPAGATION (1/2)

How does a node announce his presence to the rest of the network?

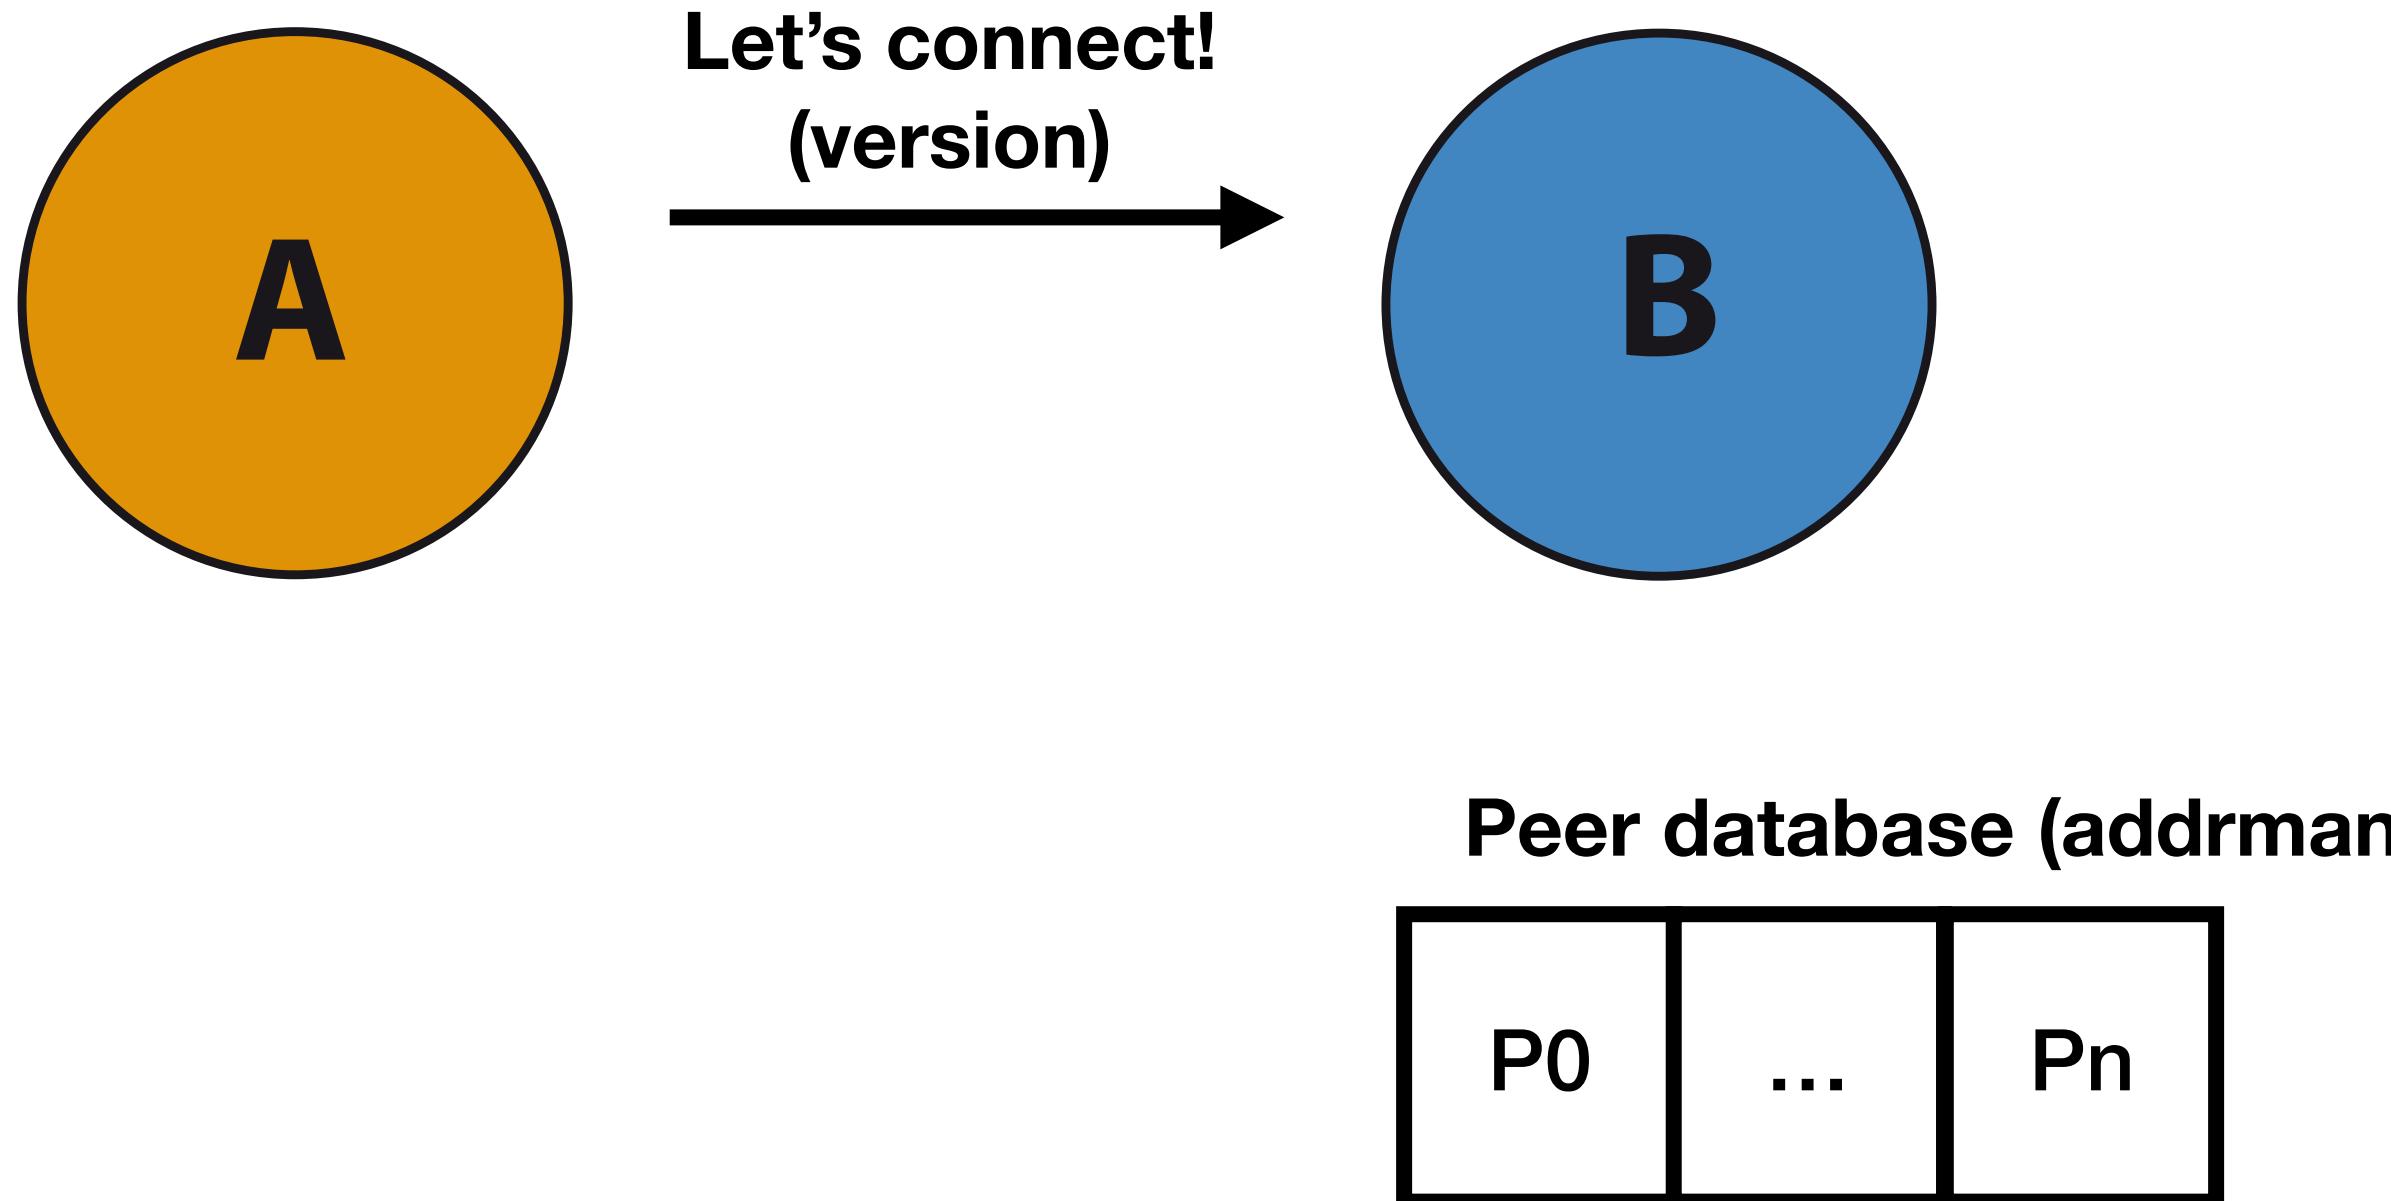


Peer database (addrman)



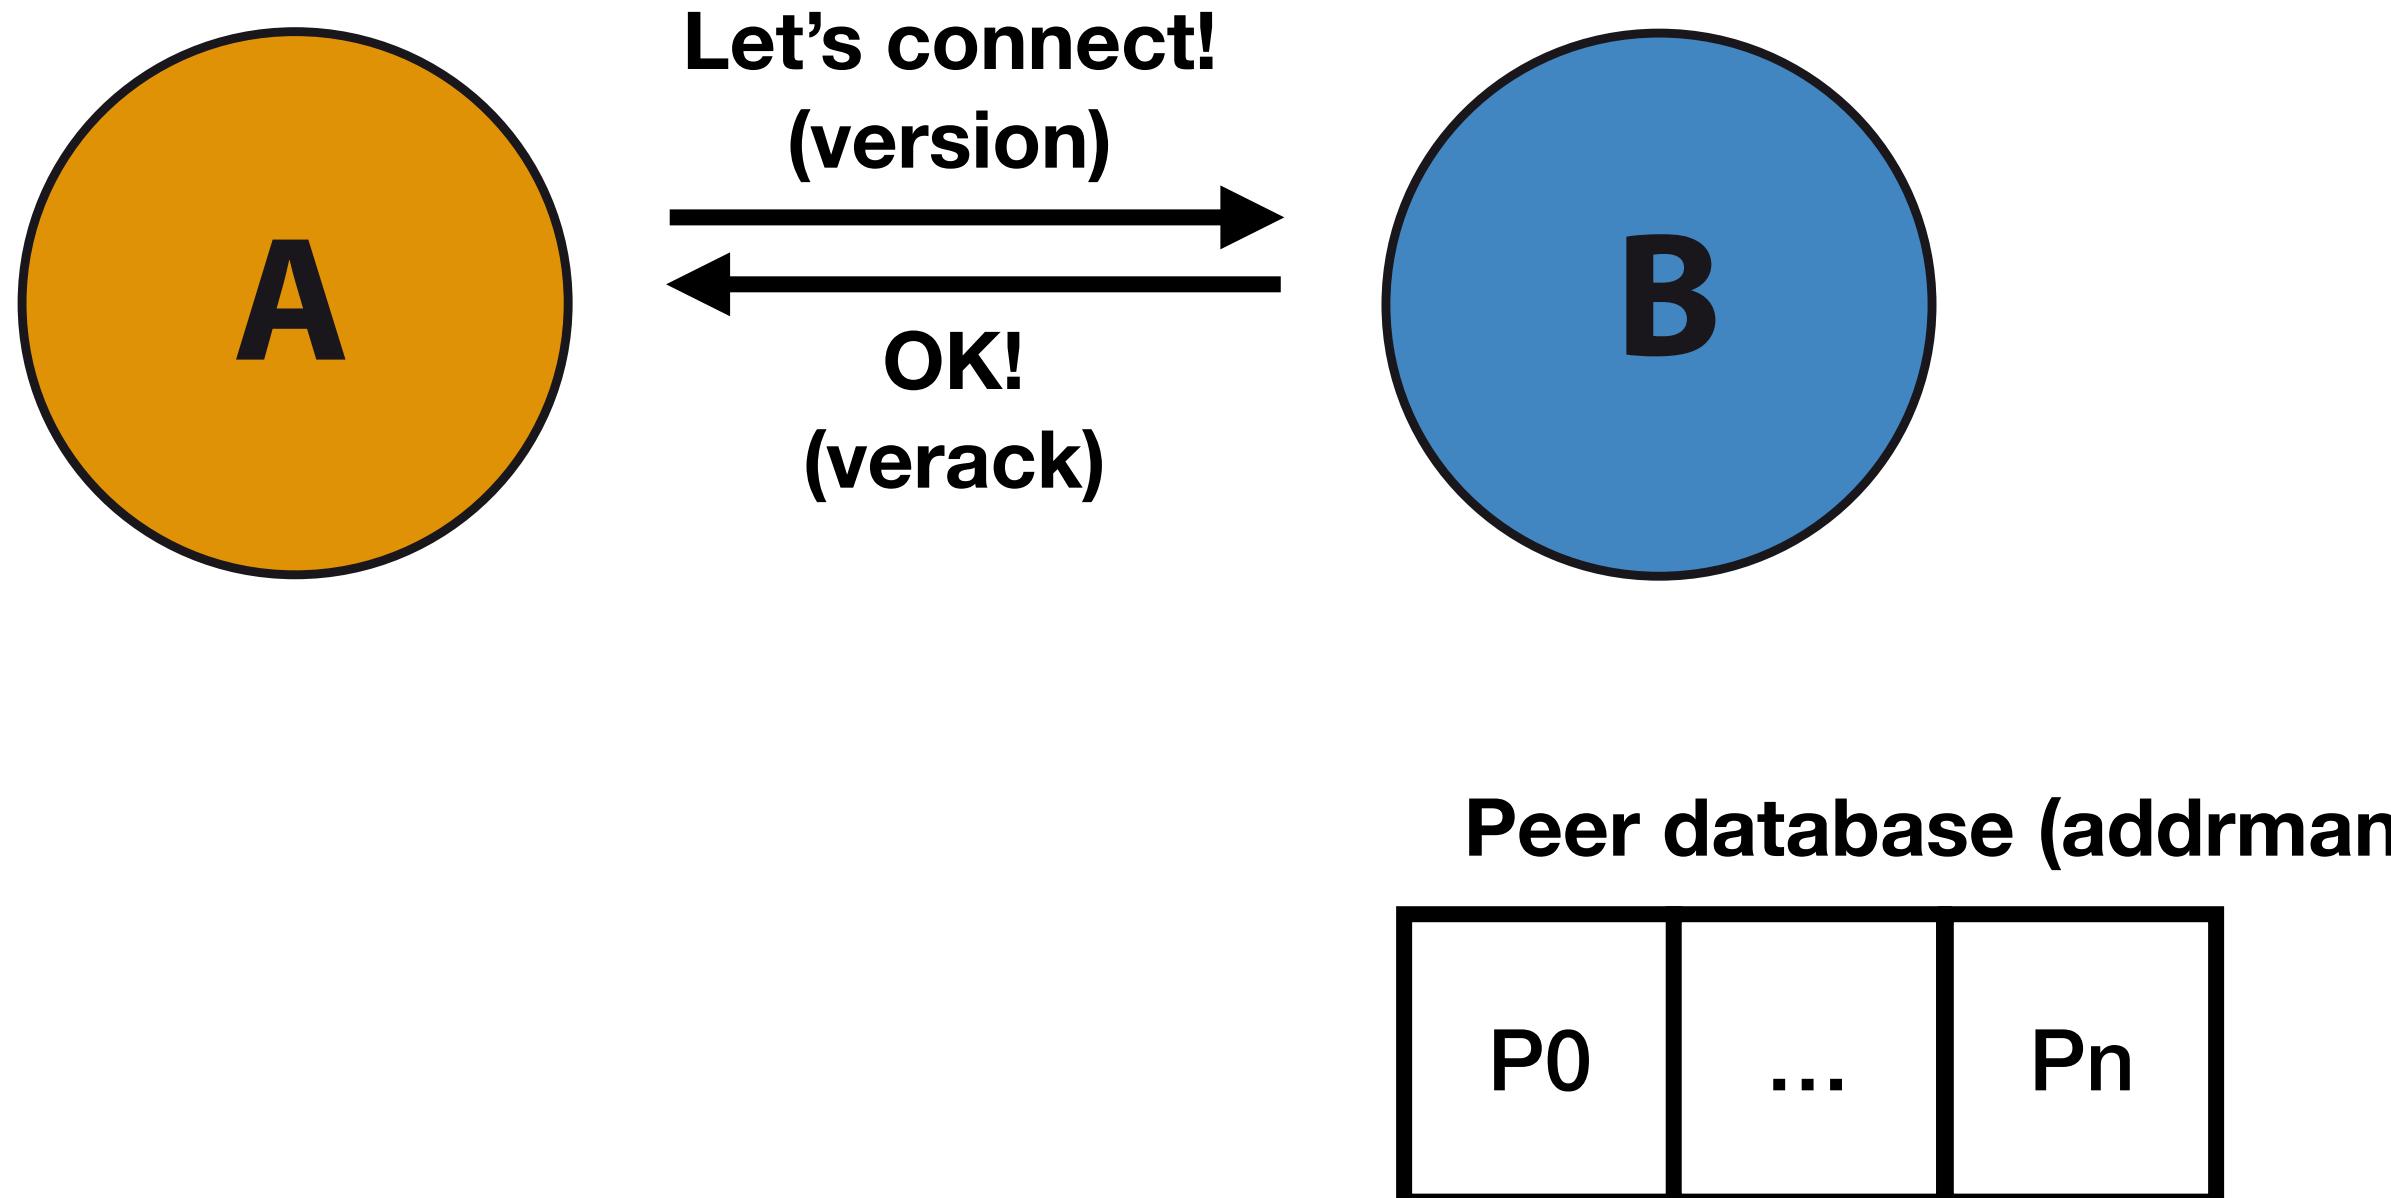
ADDRESS PROPAGATION (1/2)

How does a node announce his presence to the rest of the network?



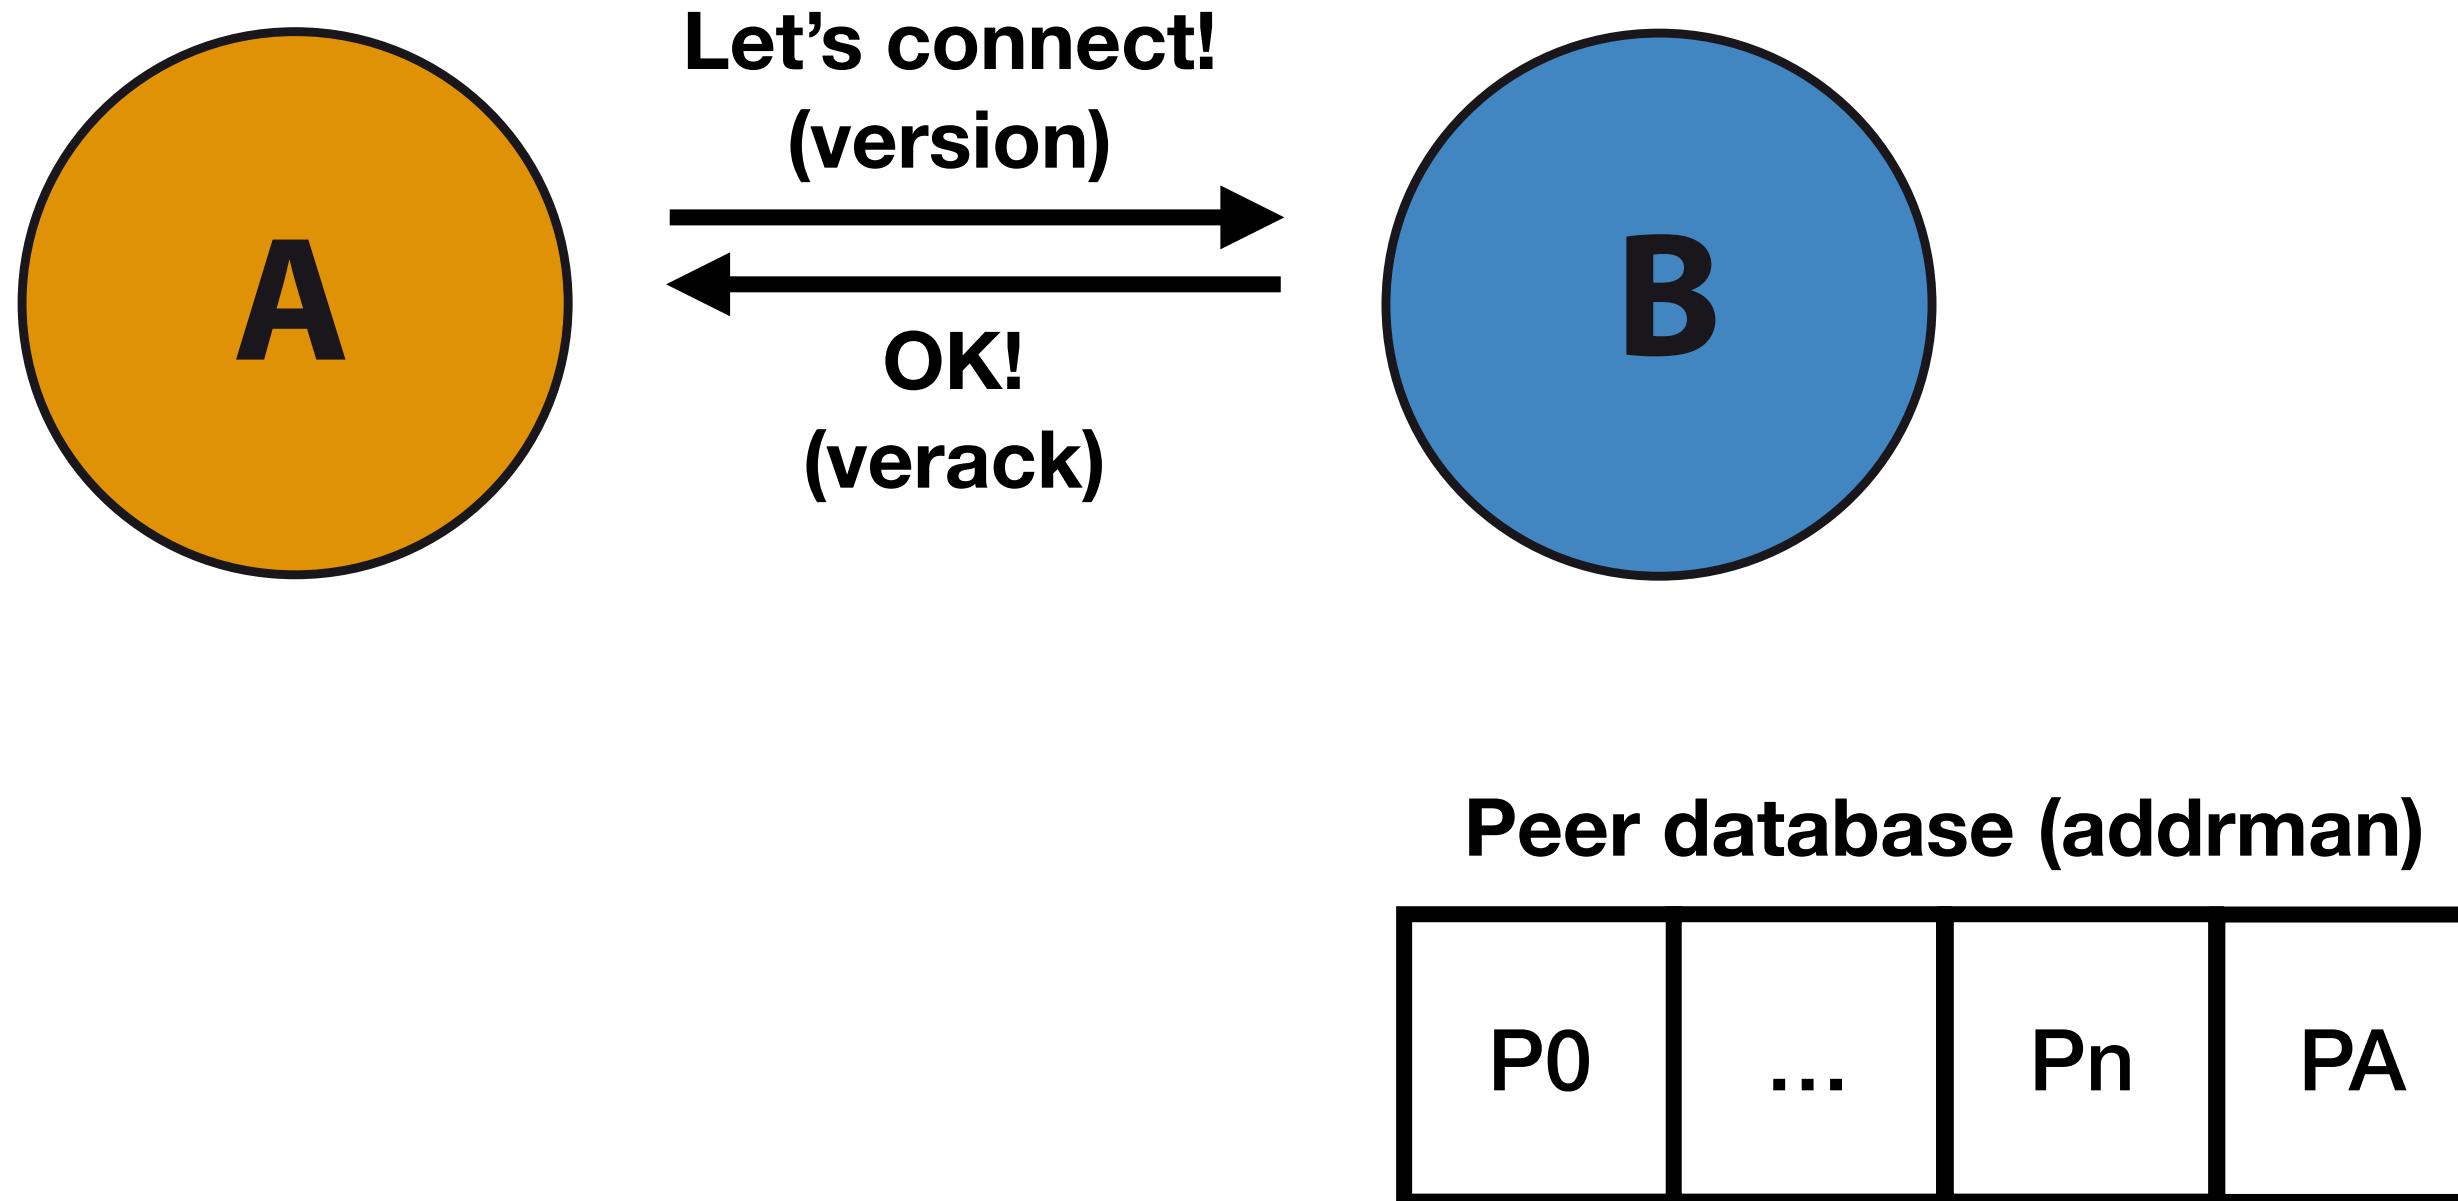
ADDRESS PROPAGATION (1/2)

How does a node announce his presence to the rest of the network?



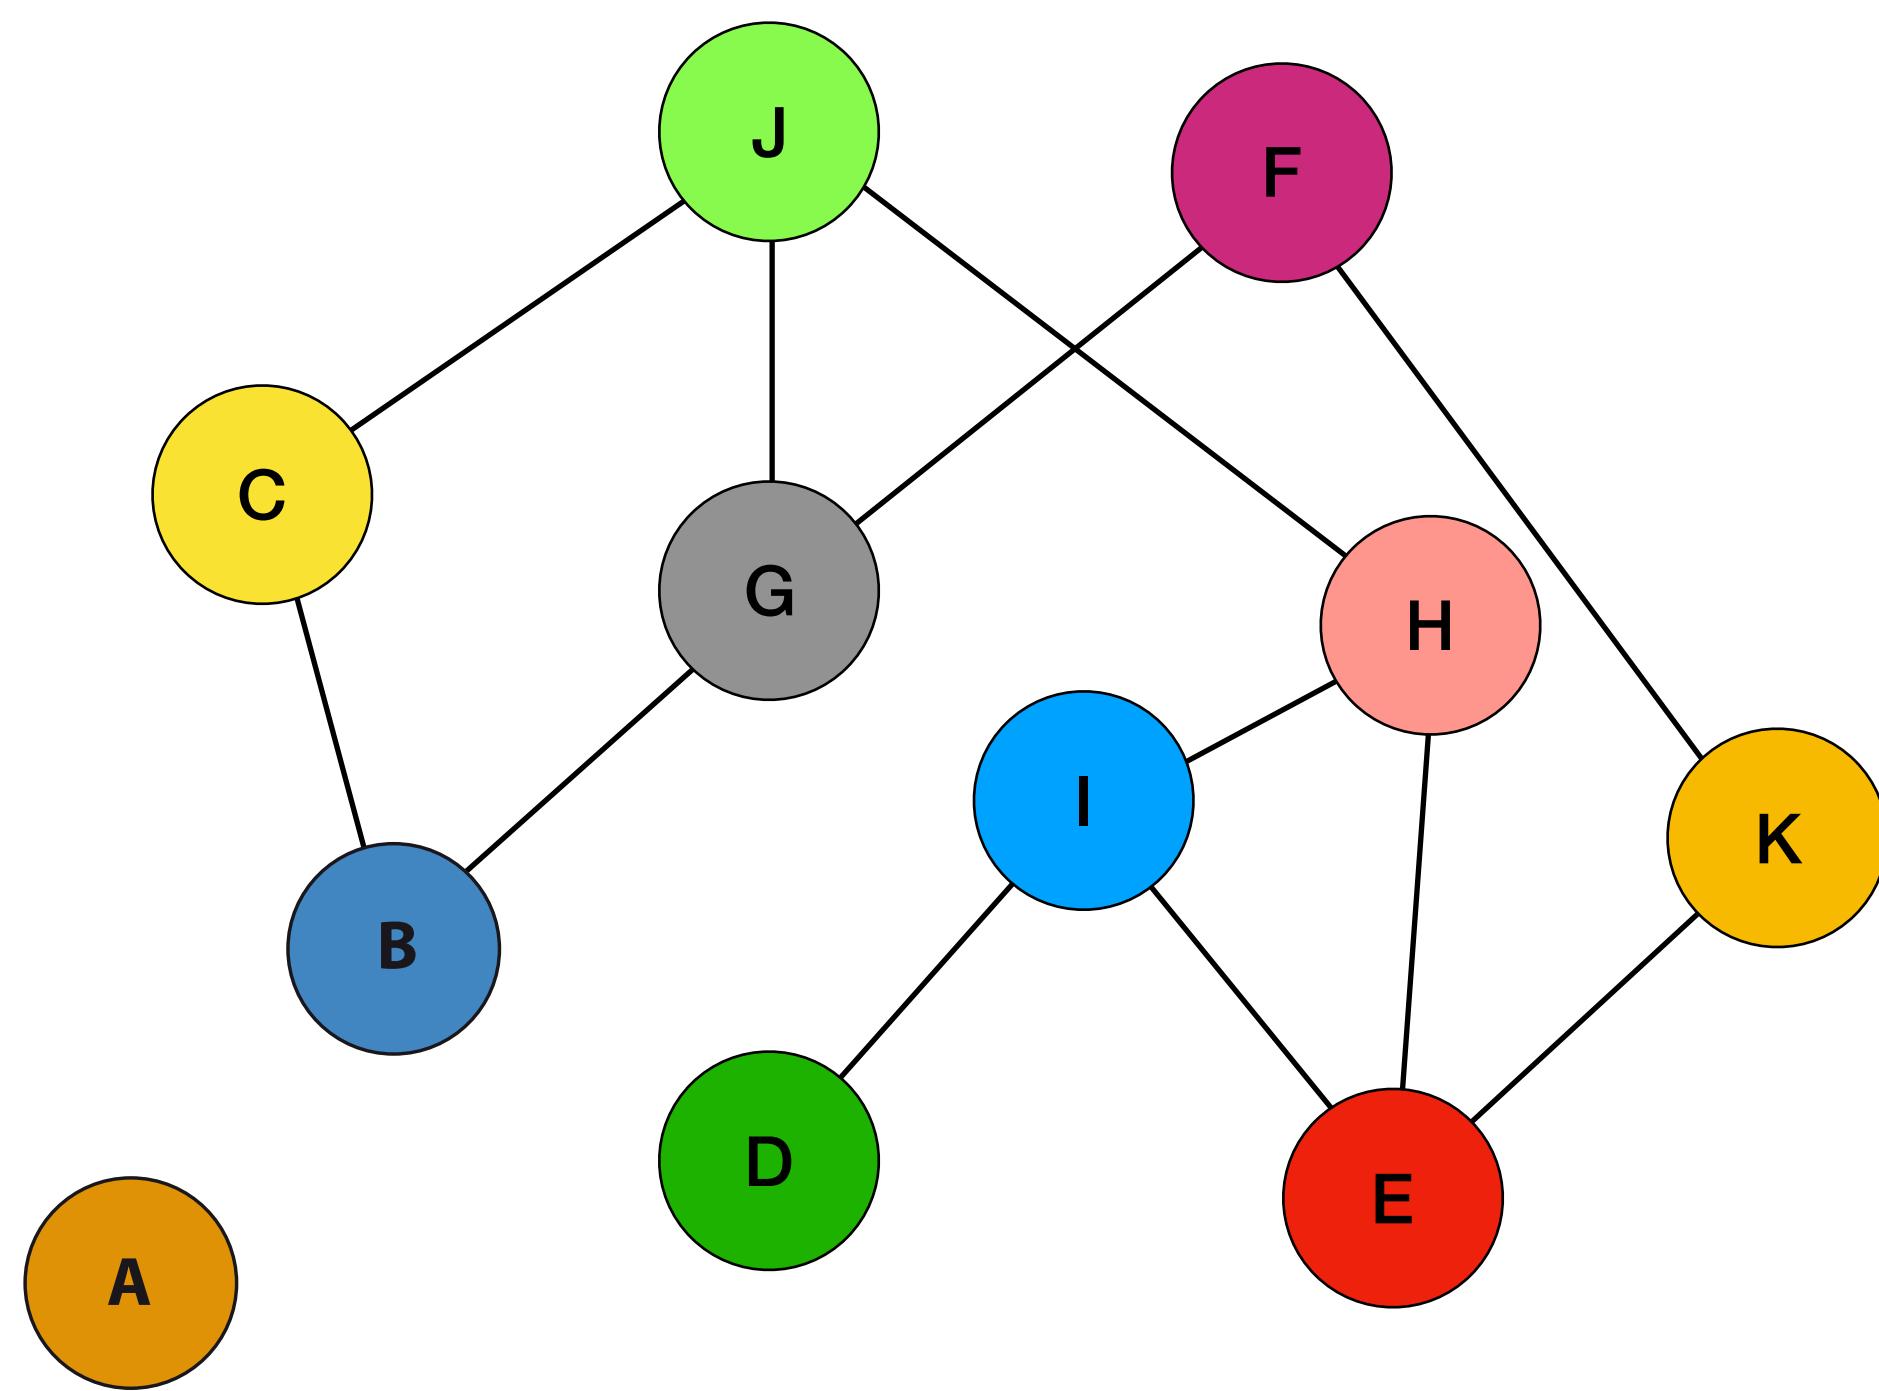
ADDRESS PROPAGATION (1/2)

How does a node announce his presence to the rest of the network?



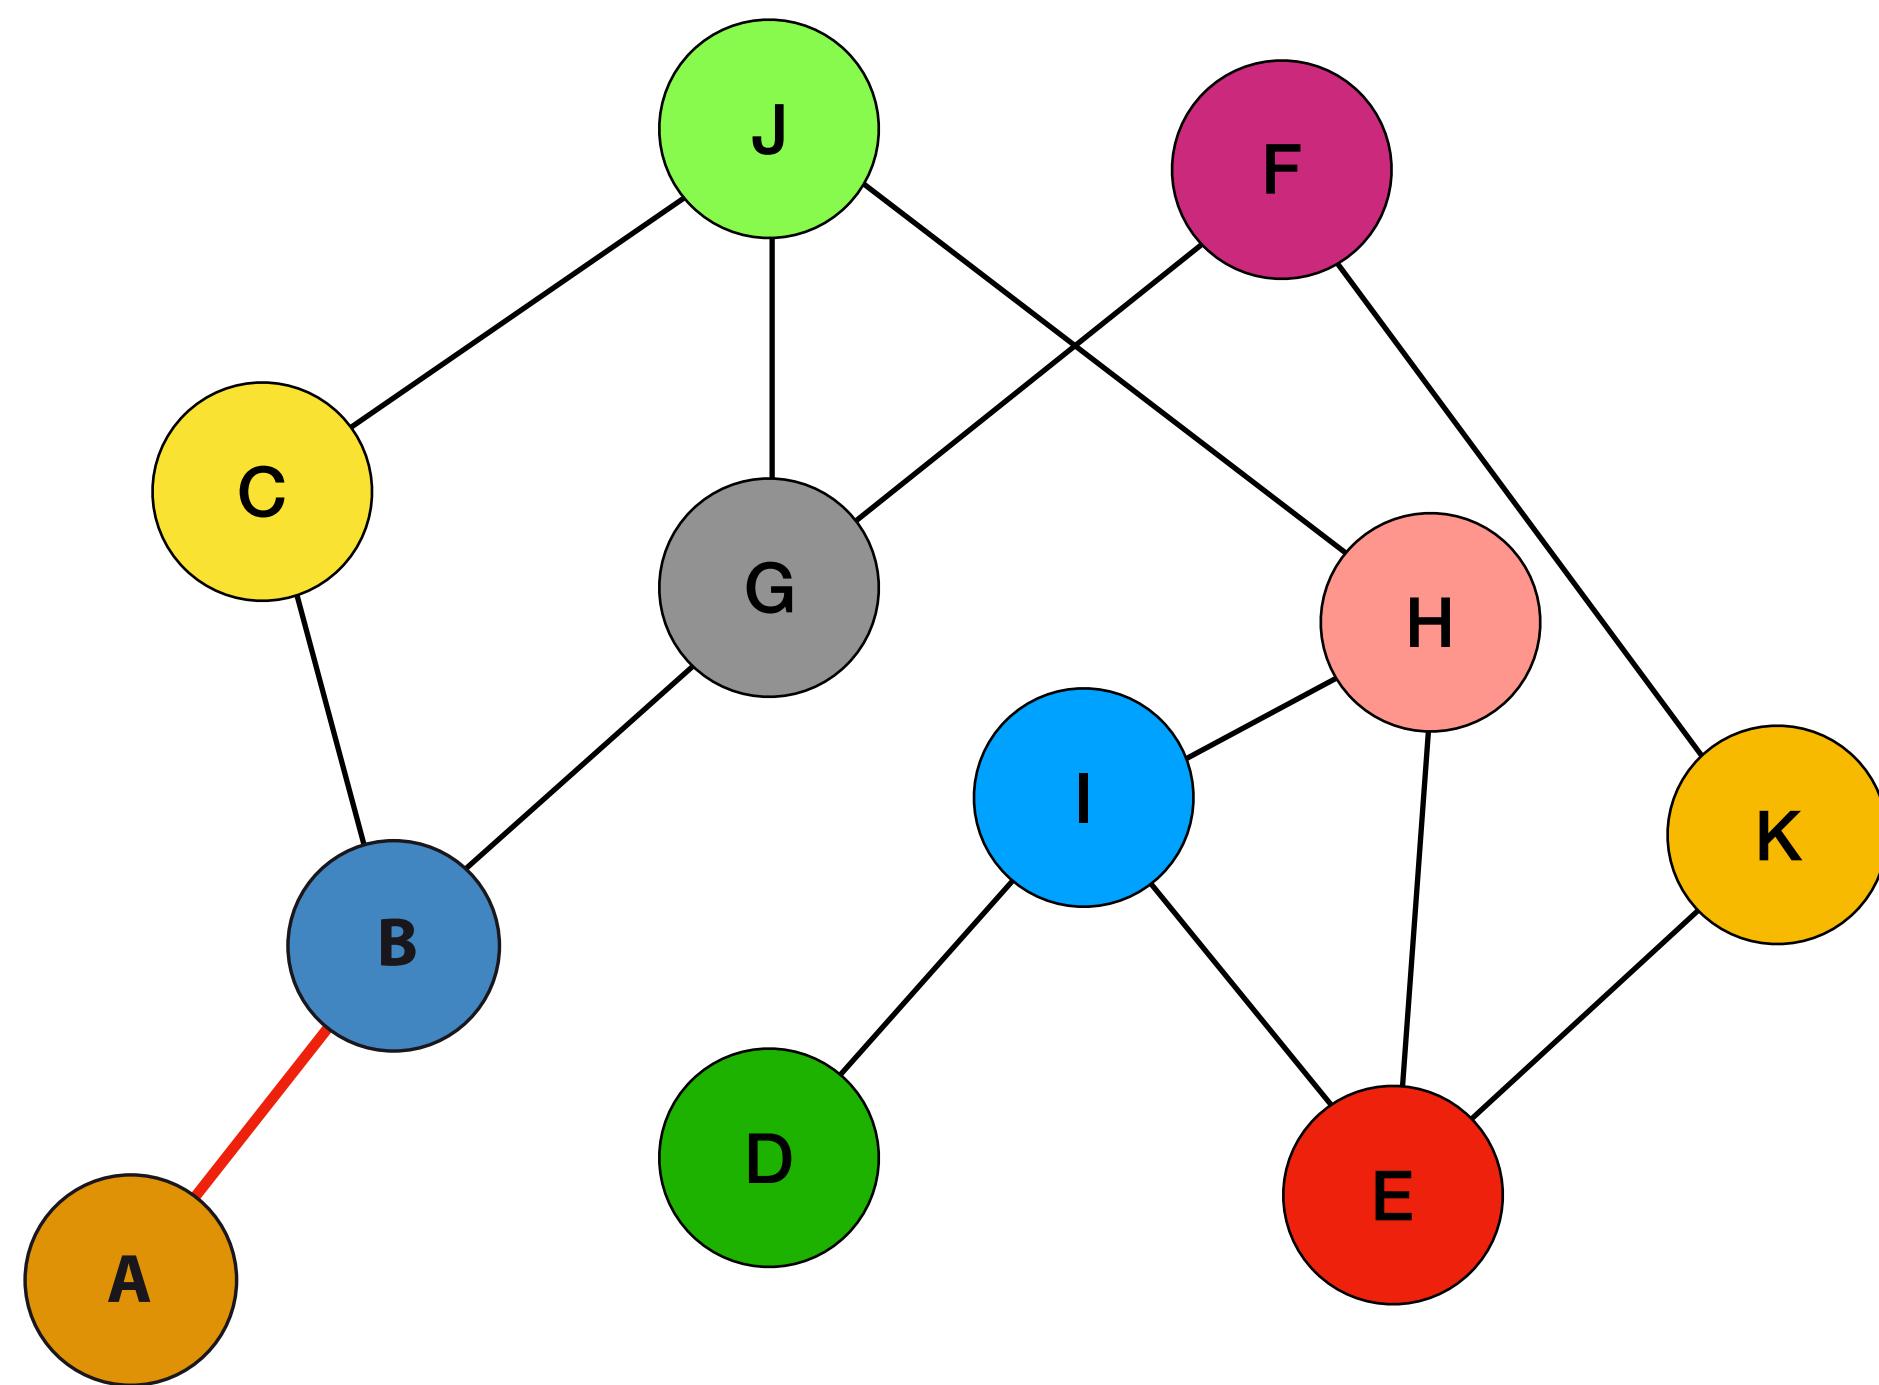
ADDRESS PROPAGATION (2/2)

How does a node announce his presence to the rest of the network?



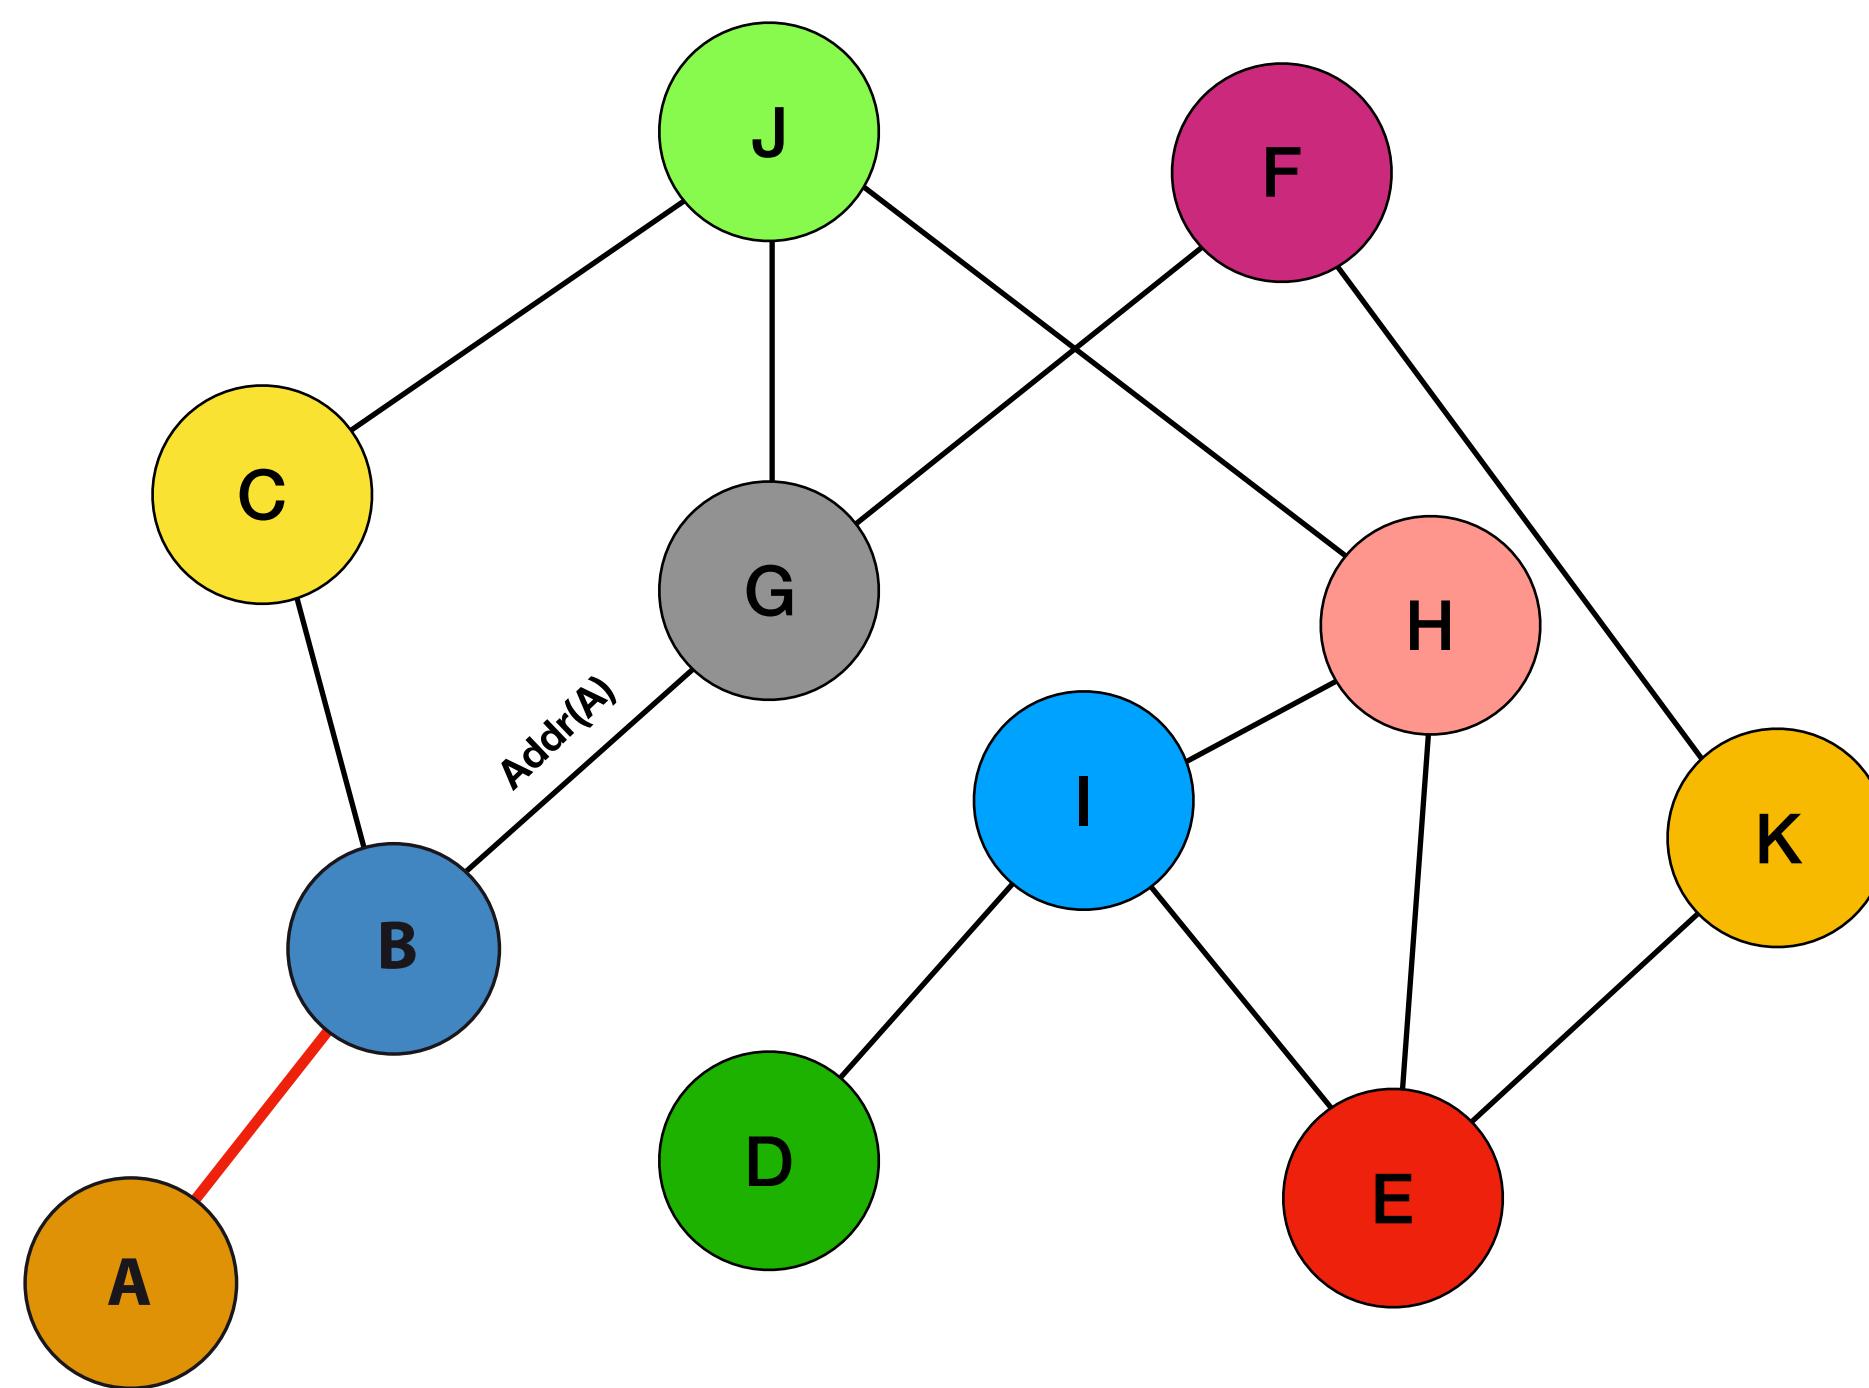
ADDRESS PROPAGATION (2/2)

How does a node announce his presence to the rest of the network?



ADDRESS PROPAGATION (2/2)

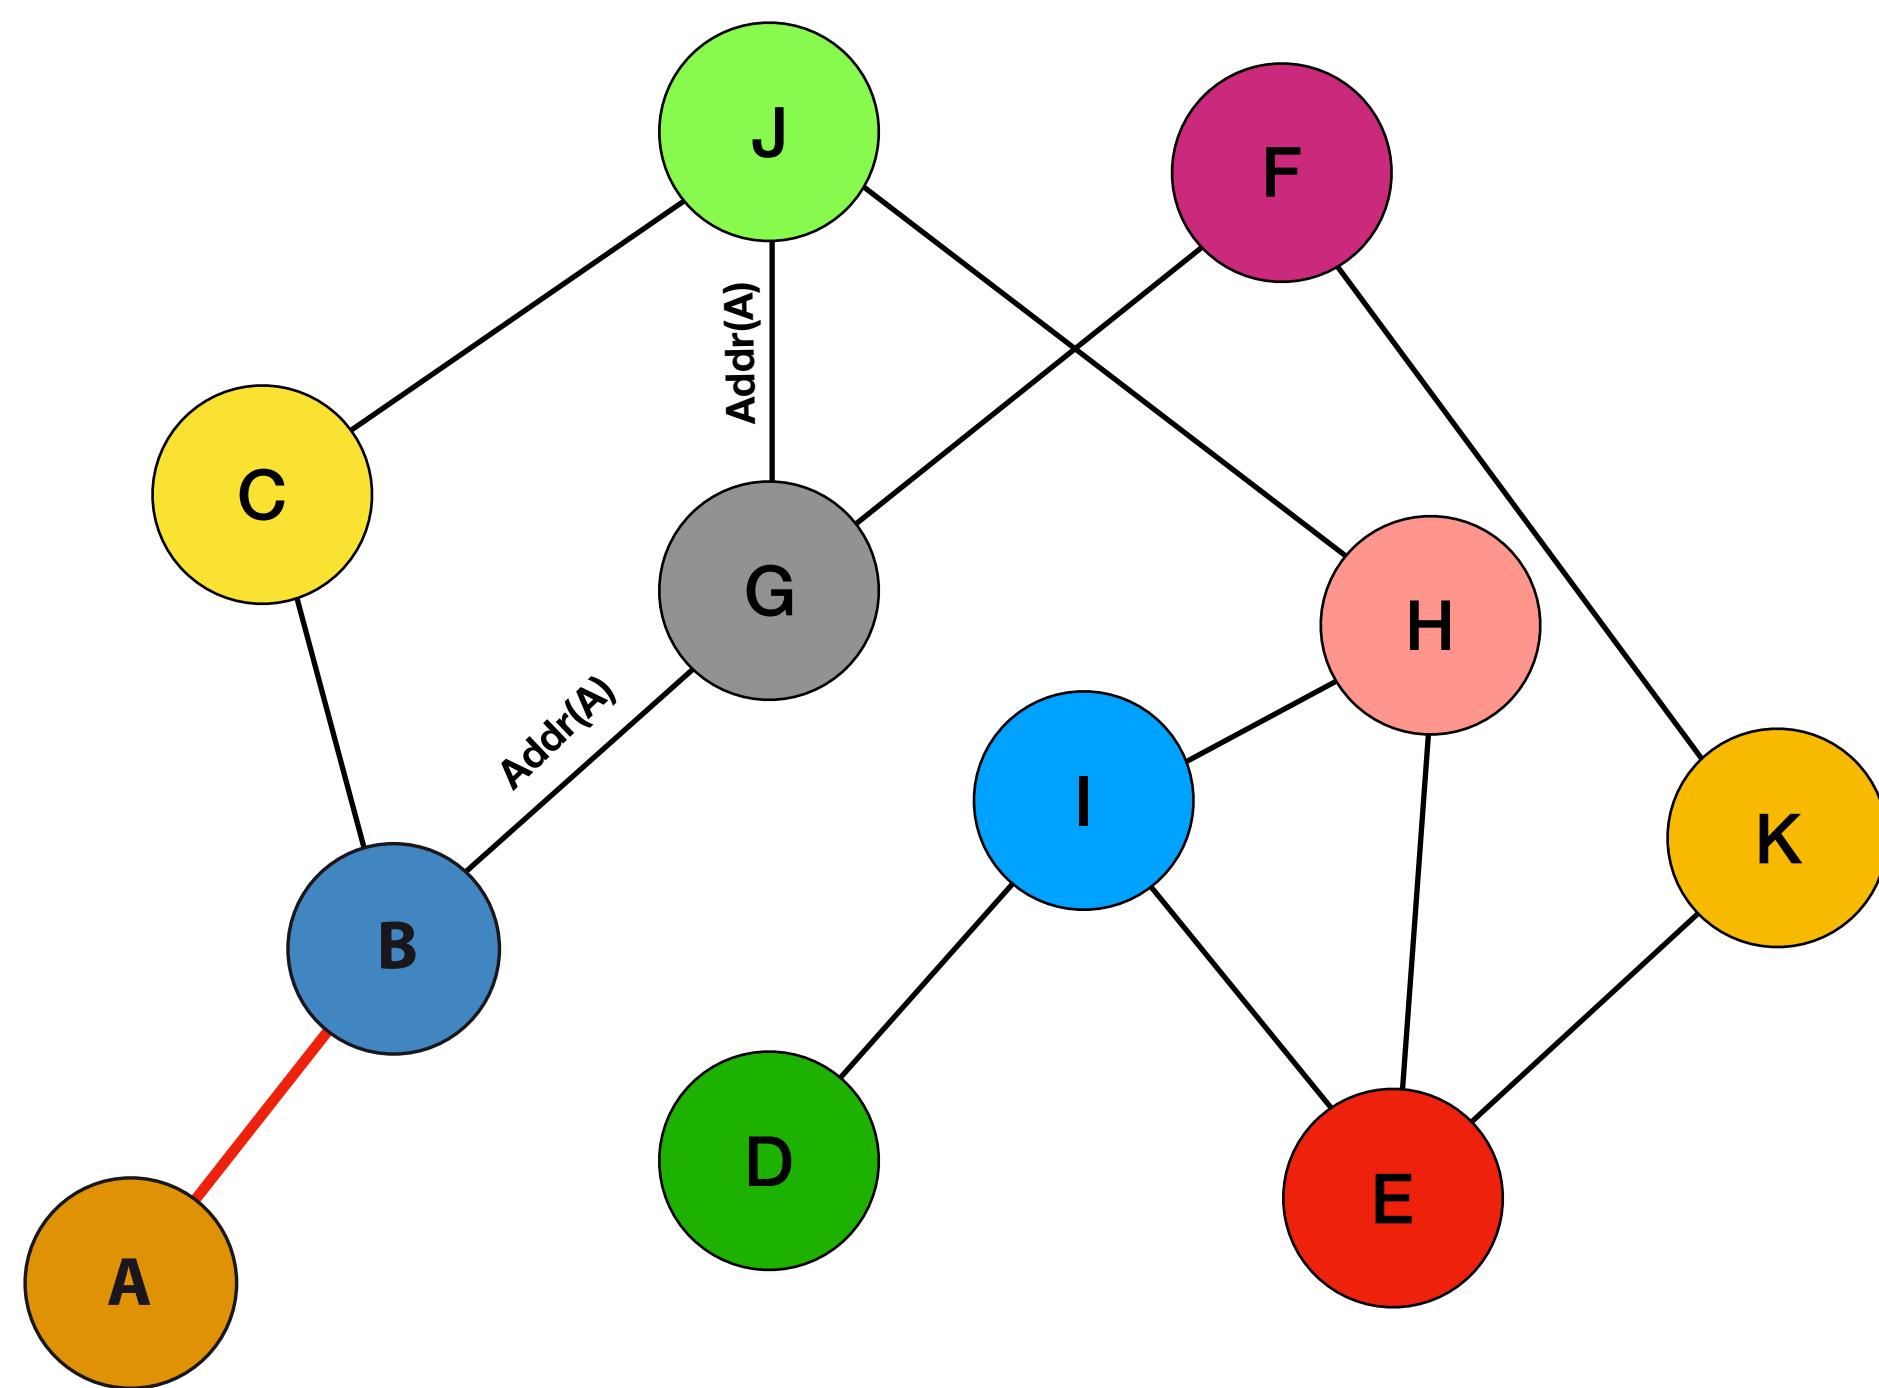
How does a node announce his presence to the rest of the network?



- B picks a random subset of its neighbors and relays A's address

ADDRESS PROPAGATION (2/2)

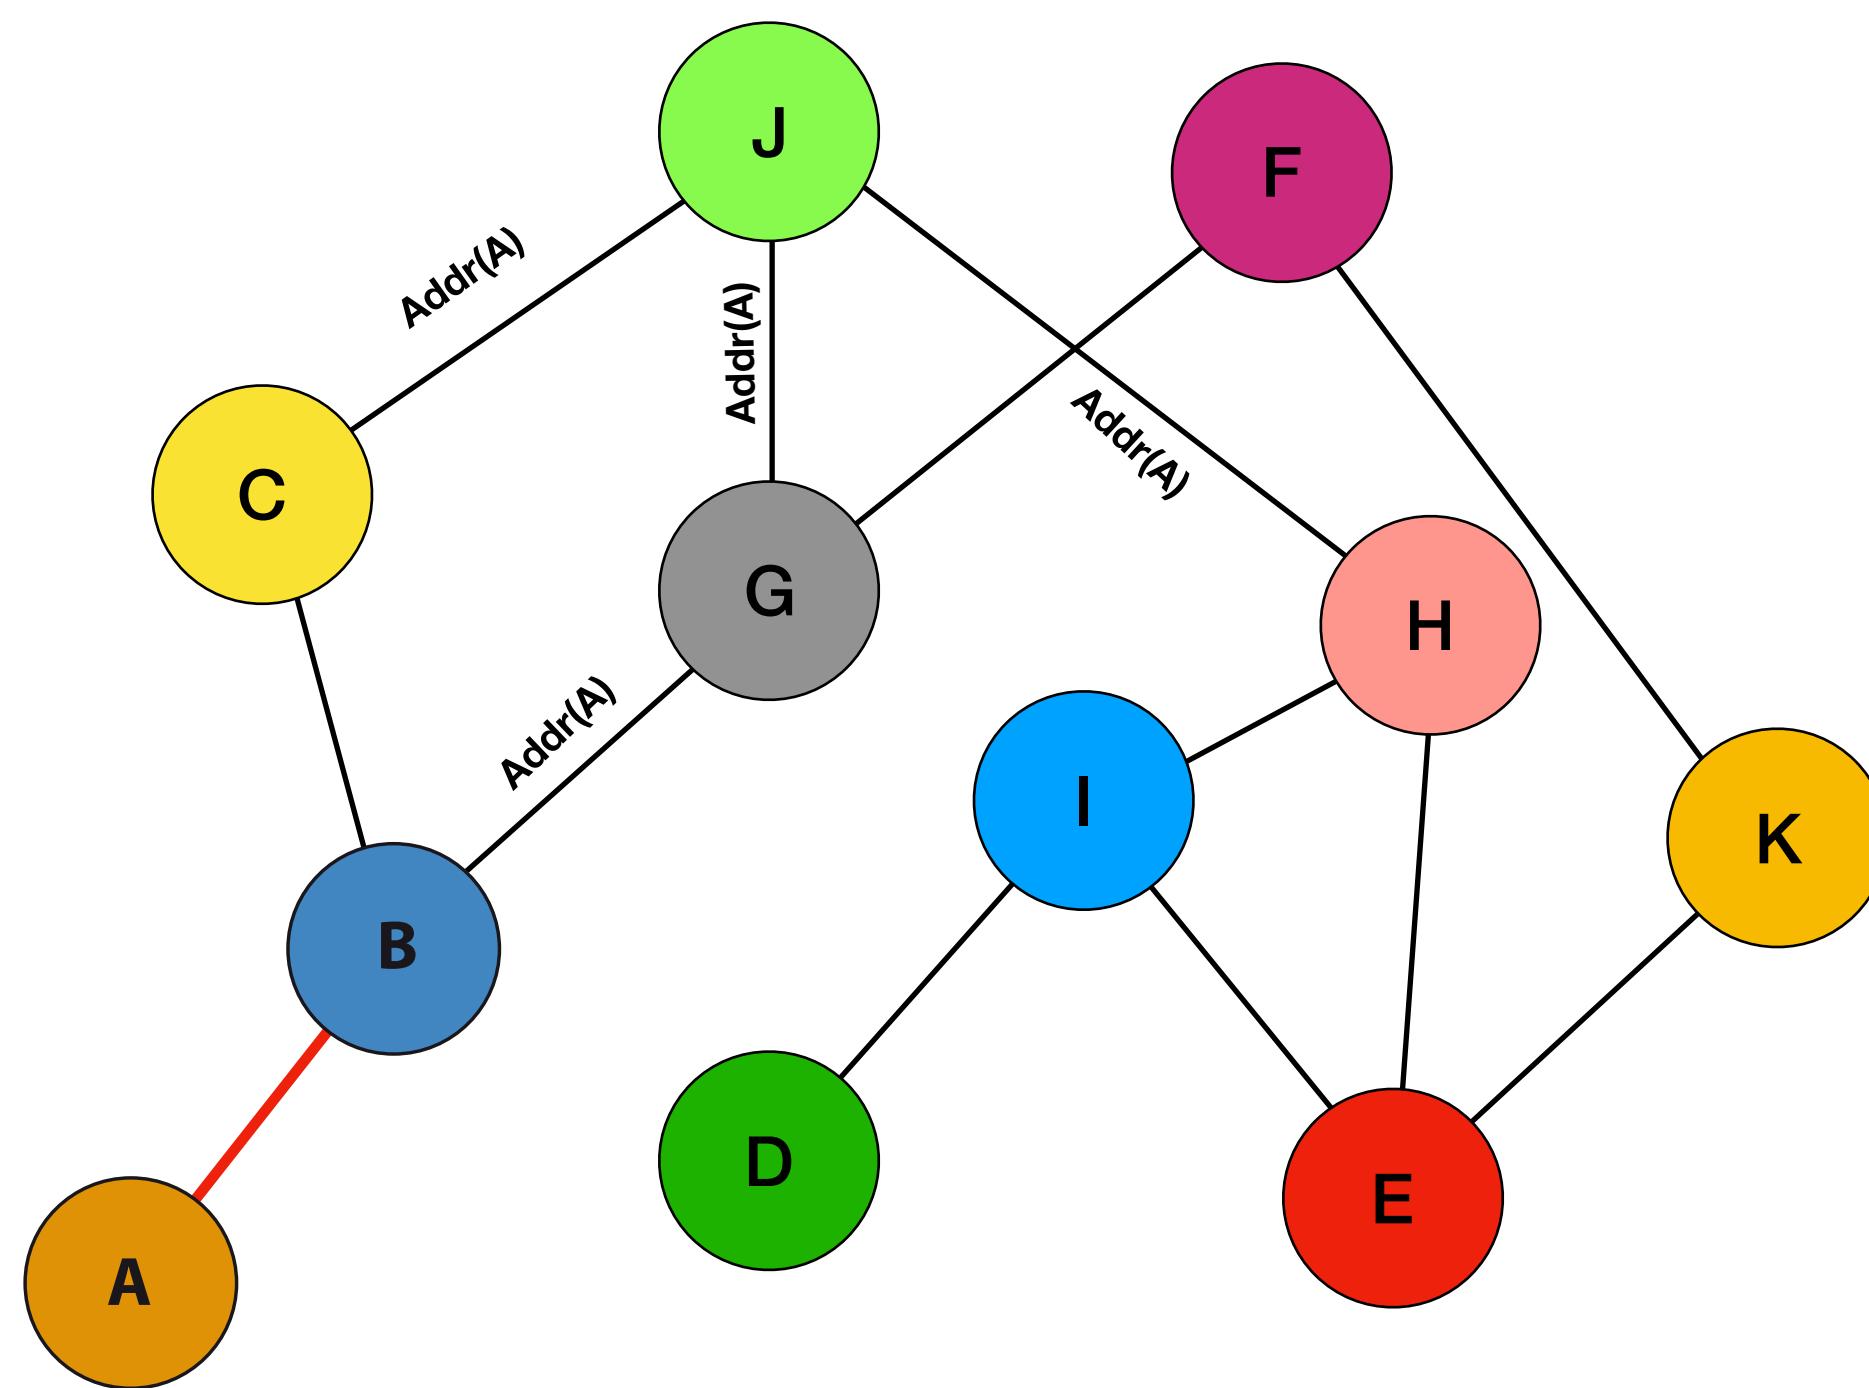
How does a node announce his presence to the rest of the network?



- B picks a random subset of its neighbors and relays A's address
- The nodes picked by B pick a random subset of their neighbors and relay A's address

ADDRESS PROPAGATION (2/2)

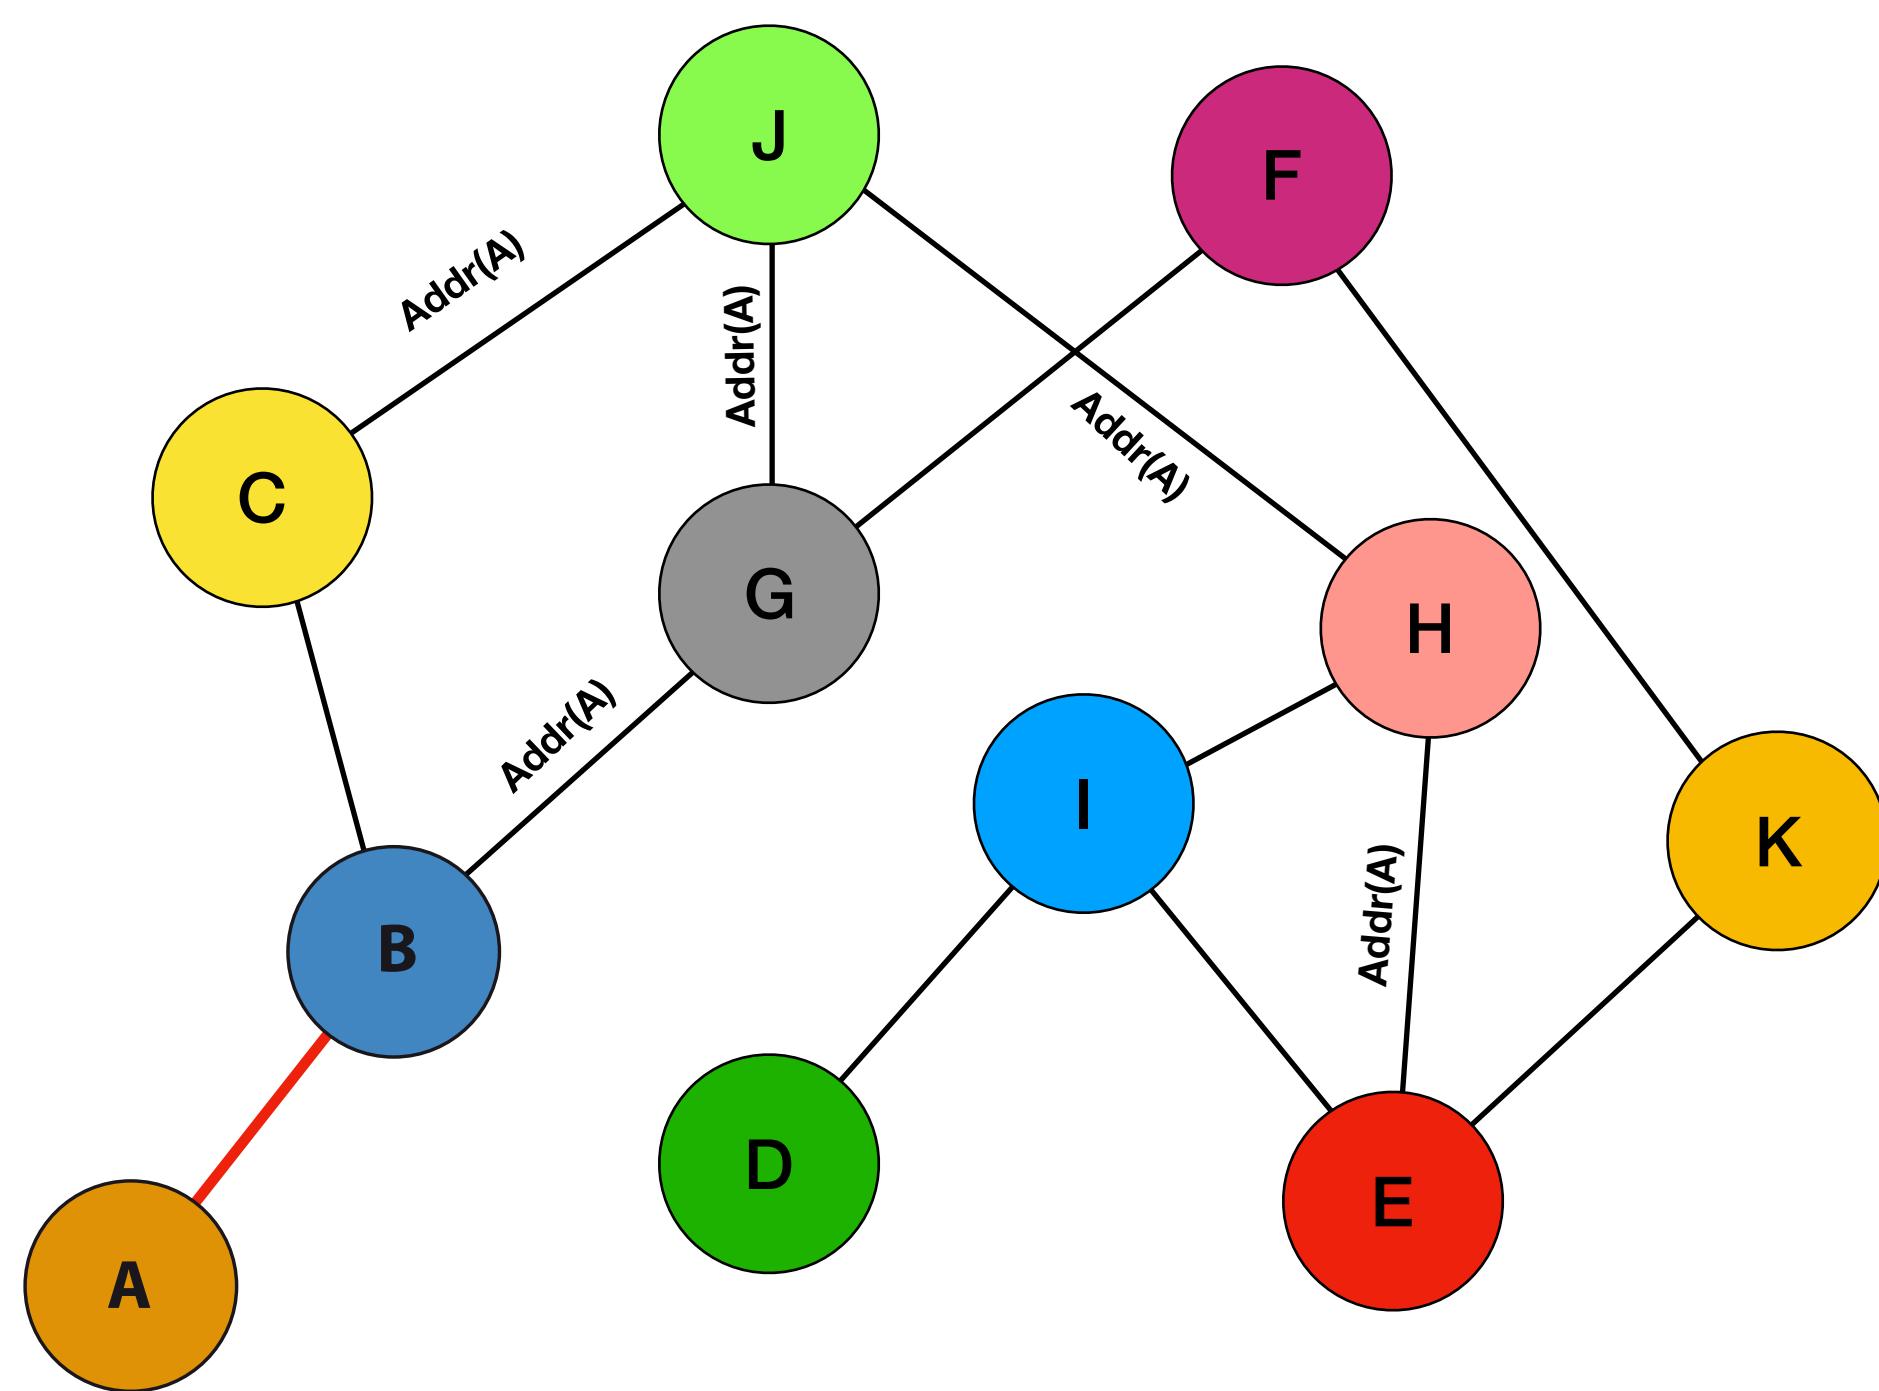
How does a node announce his presence to the rest of the network?



- B picks a random subset of its neighbors and relays A's address
- The nodes picked by B pick a random subset of their neighbors and relay A's address
- And so on and so forth...

ADDRESS PROPAGATION (2/2)

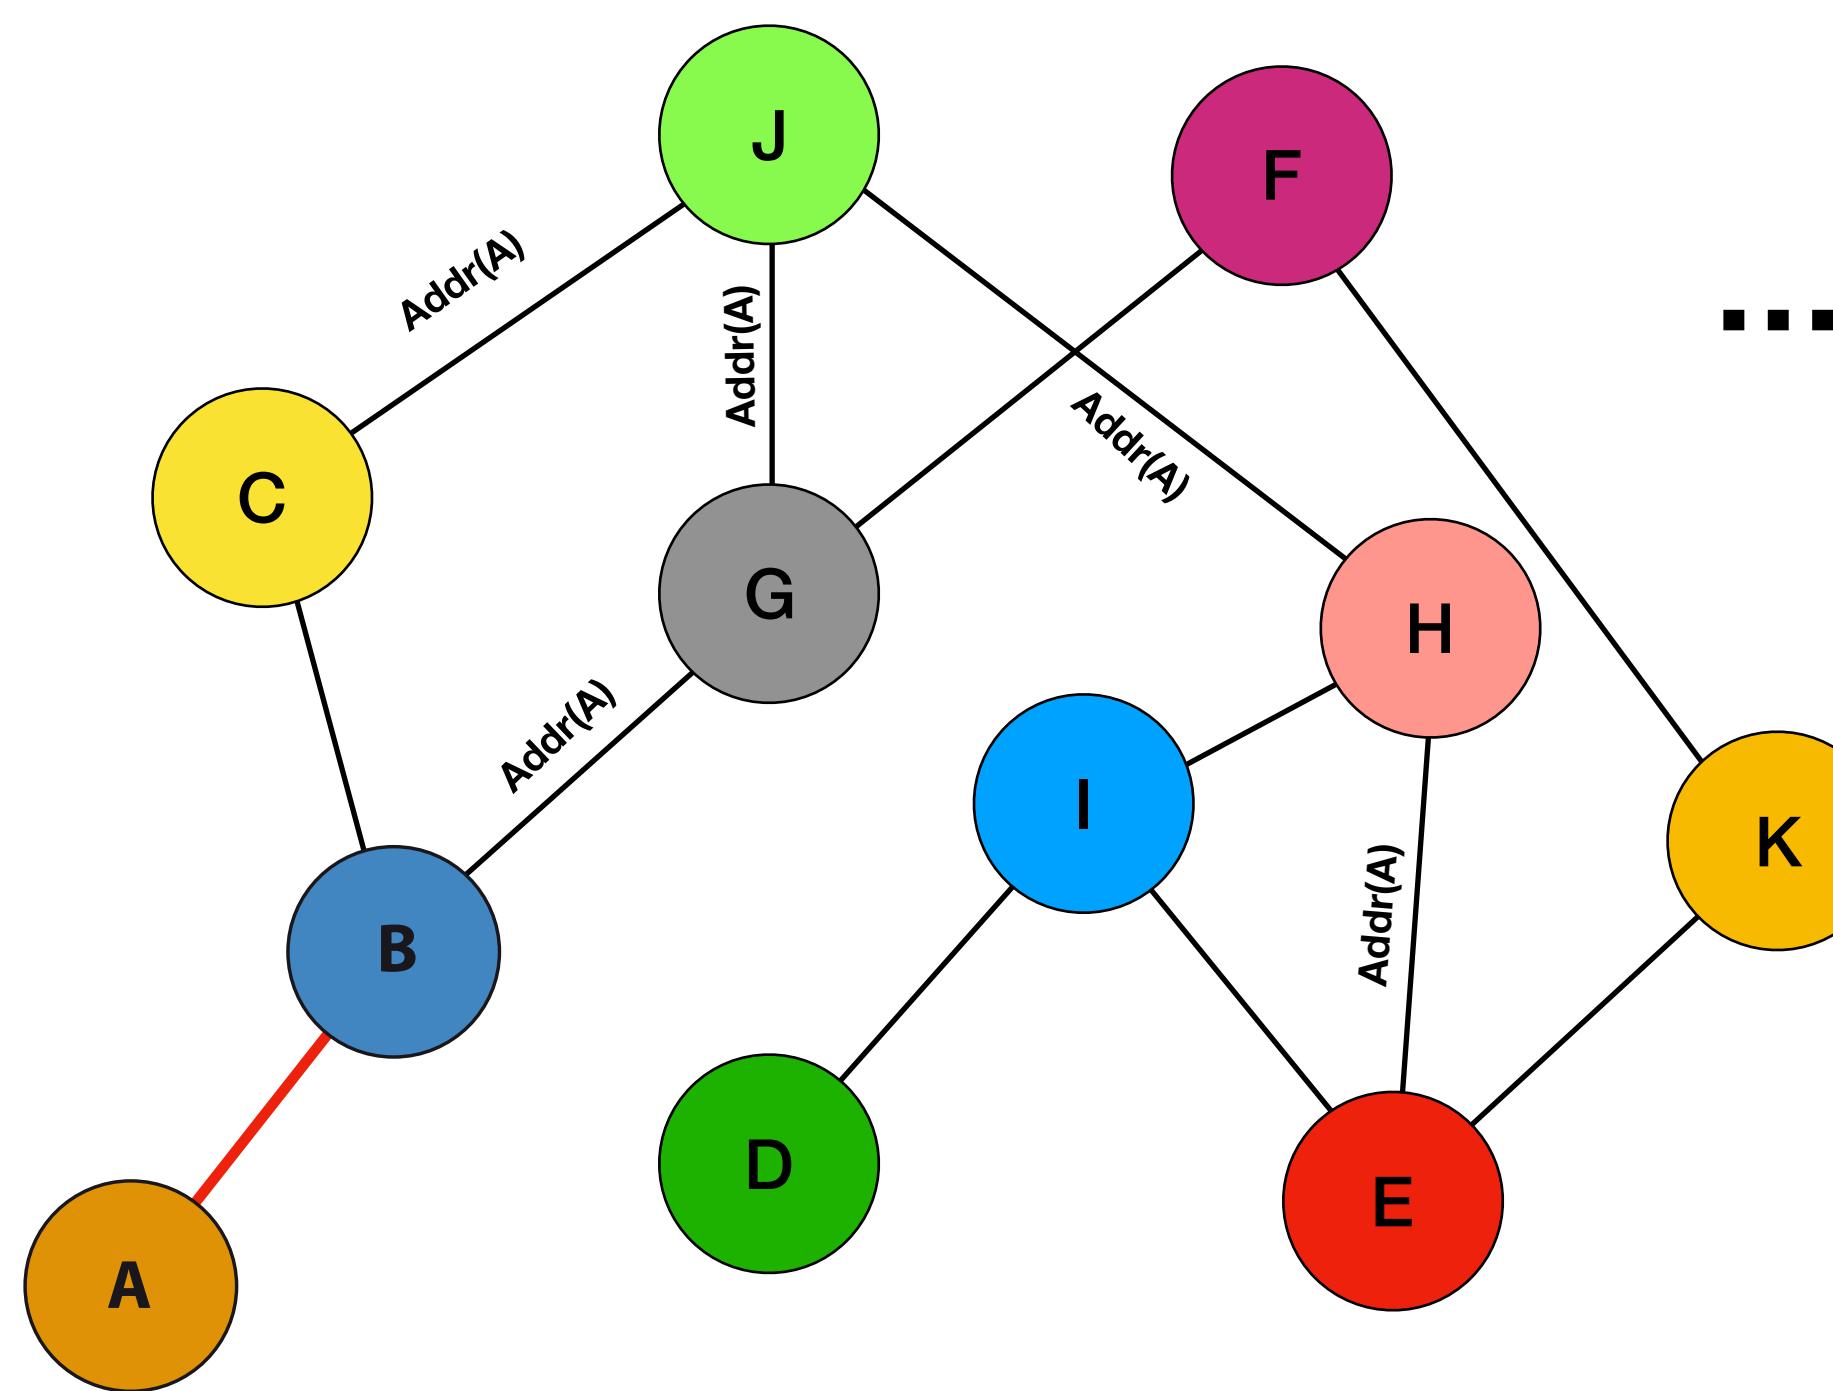
How does a node announce his presence to the rest of the network?



- B picks a random subset of its neighbors and relays A's address
- The nodes picked by B pick a random subset of their neighbors and relay A's address
- And so on and so forth...

ADDRESS PROPAGATION (2/2)

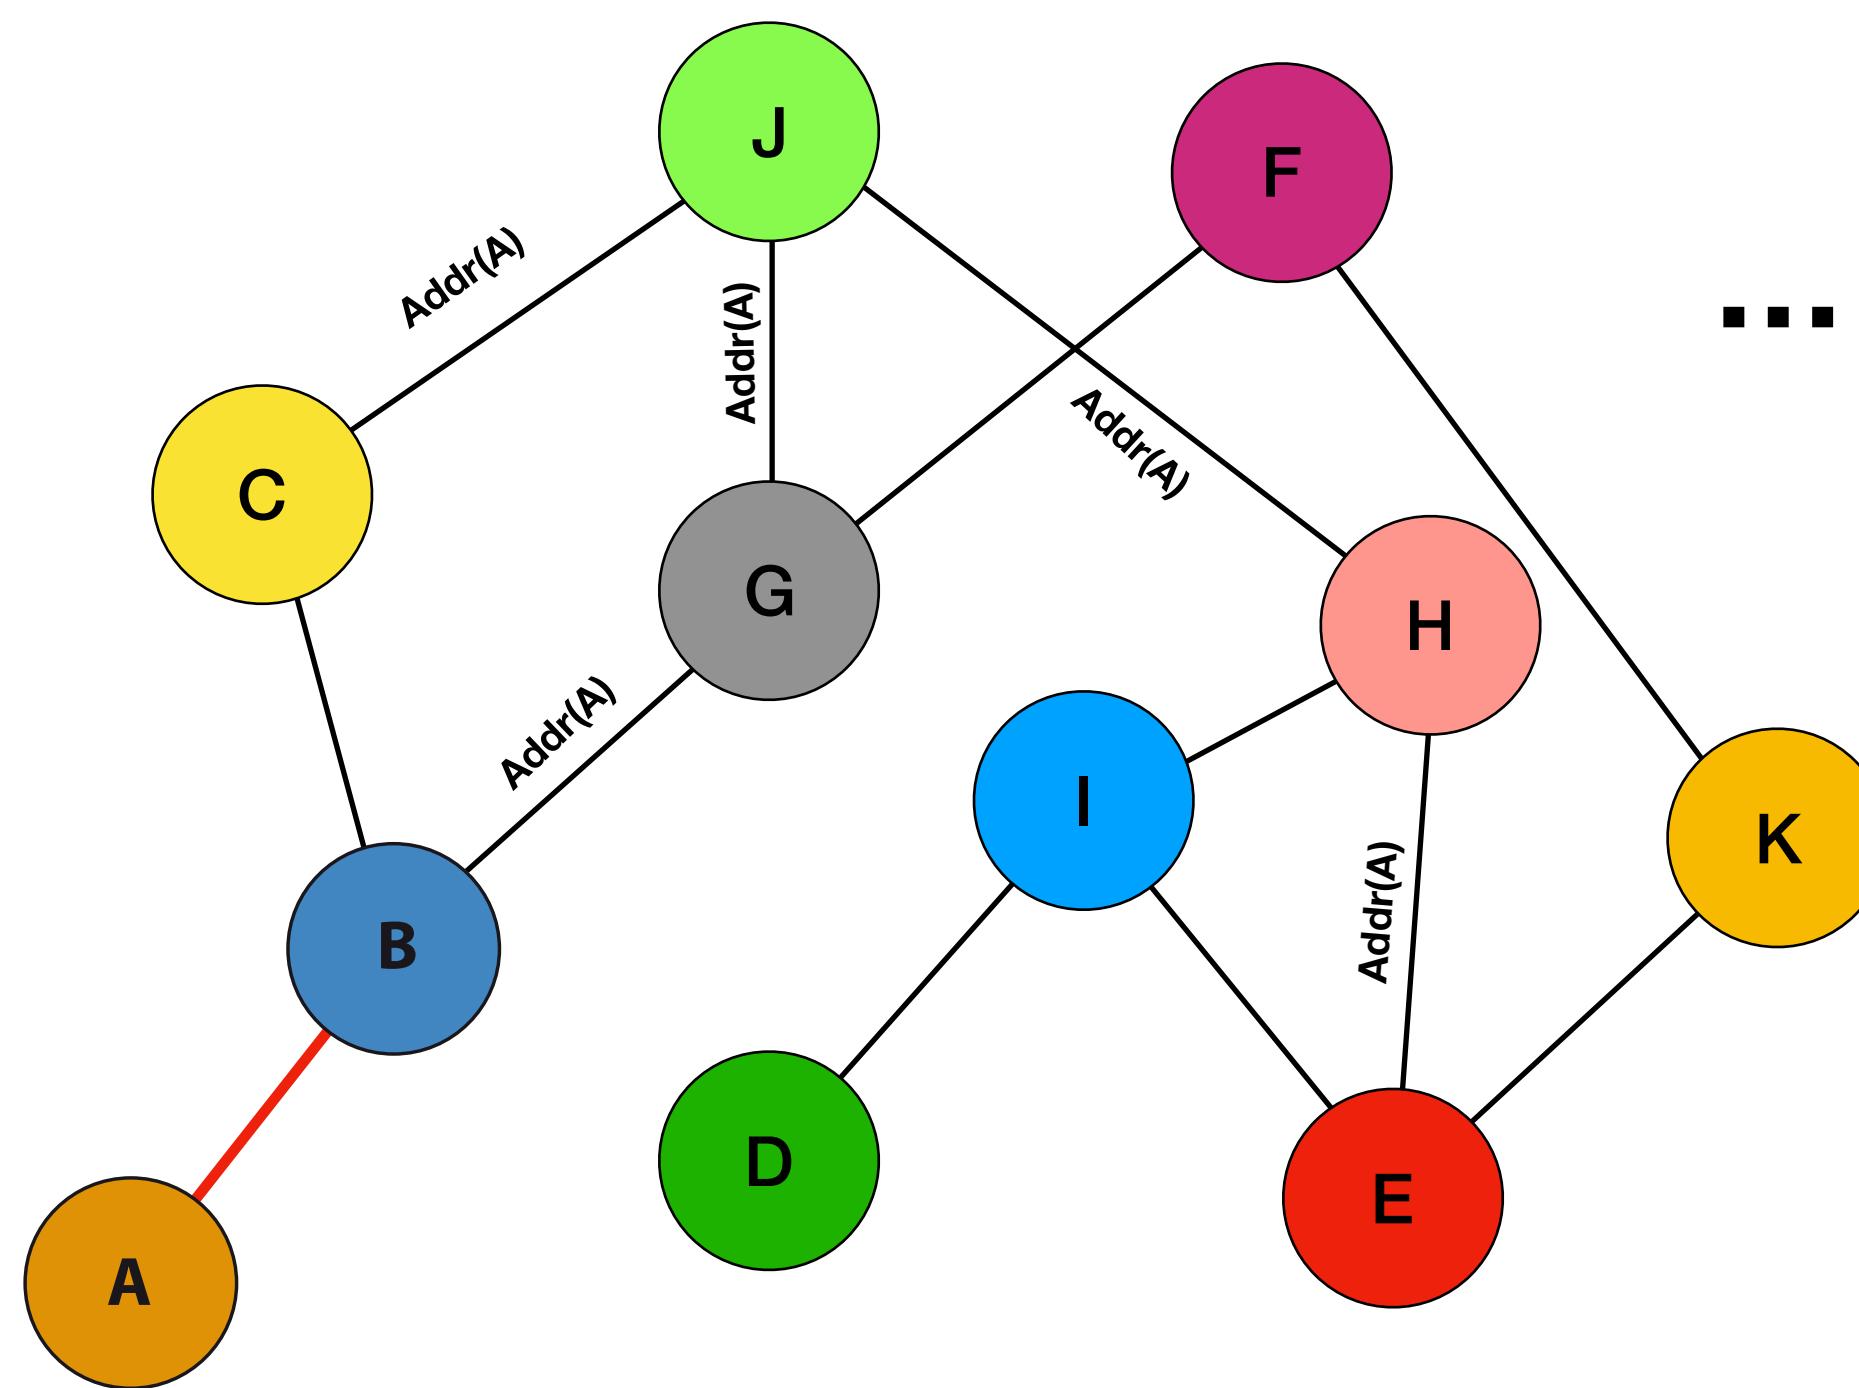
How does a node announce his presence to the rest of the network?



- B picks a random subset of its neighbors and relays A's address
- The nodes picked by B pick a random subset of their neighbors and relay A's address
- And so on and so forth...

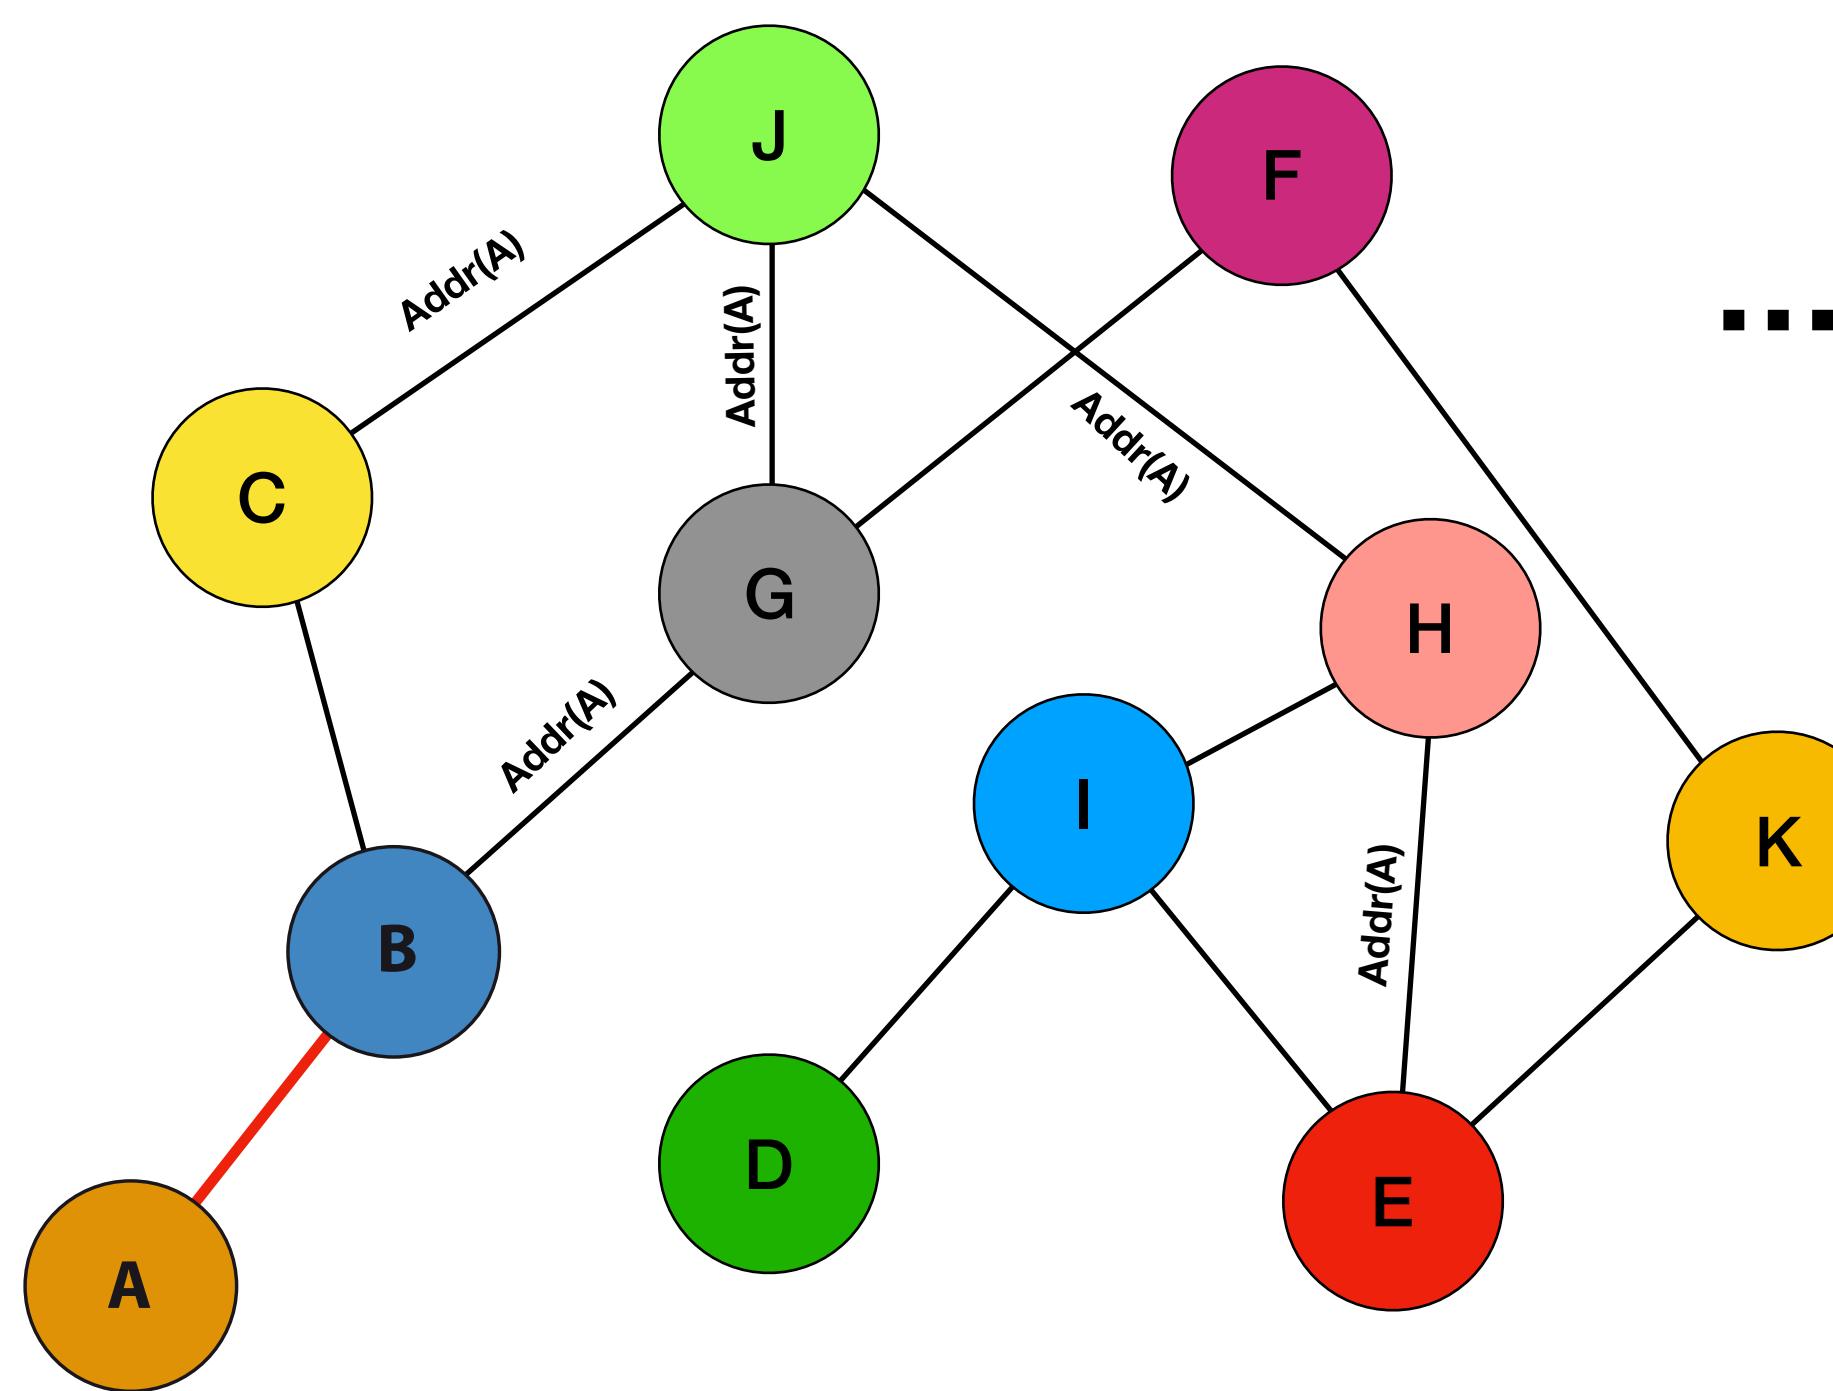
ADDRESS PROPAGATION (2/2)

How does a node announce his presence to the rest of the network?



ADDRESS PROPAGATION (2/2)

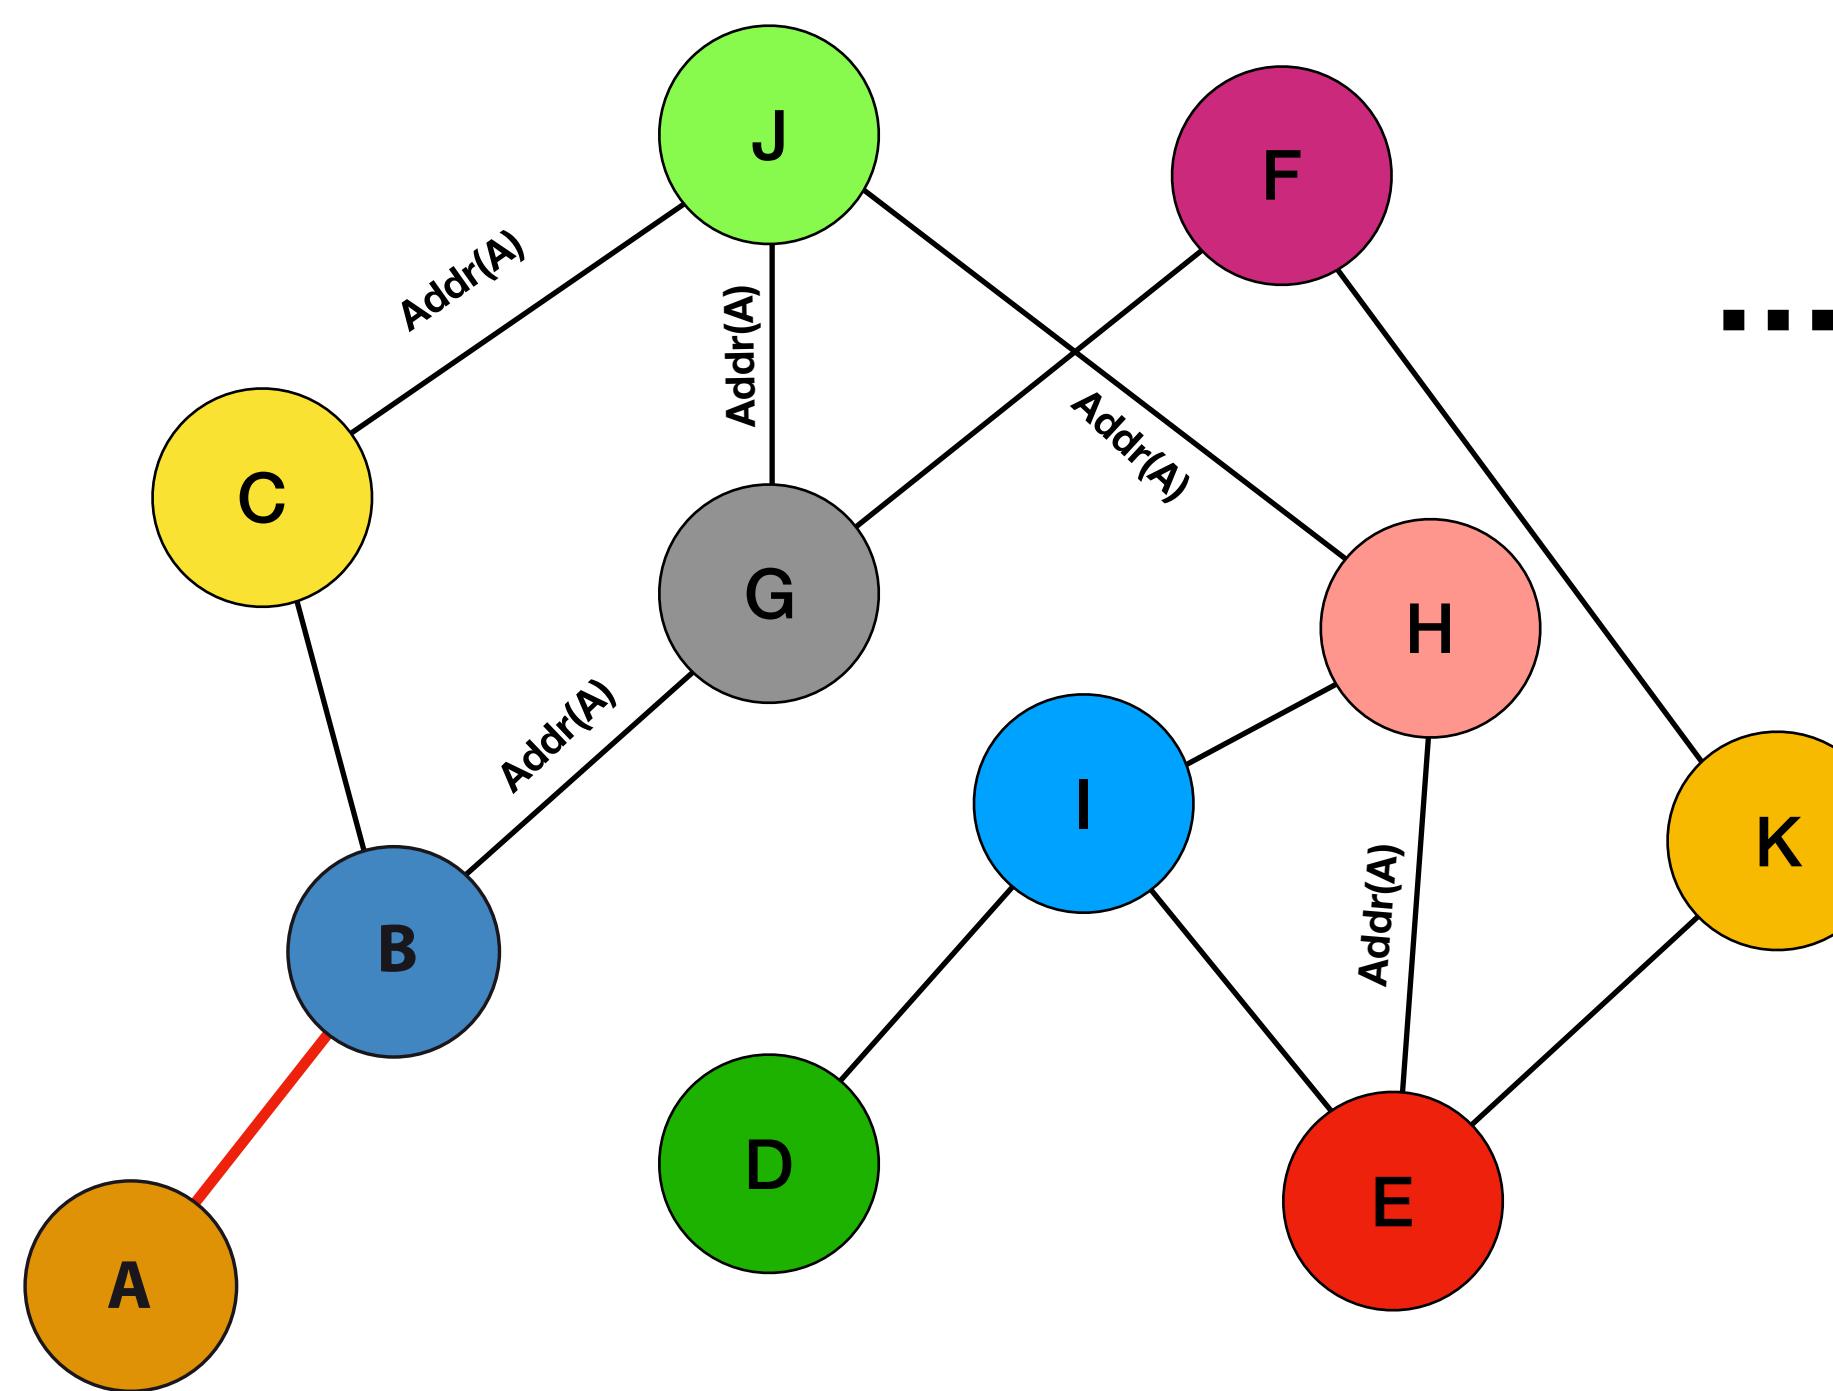
How does a node announce his presence to the rest of the network?



- The address will eventually be spread throughout the network

ADDRESS PROPAGATION (2/2)

How does a node announce his presence to the rest of the network?



- The address will eventually be spread throughout the network
- Nodes learning about the new peer will add it to their peers database

CONNECTIONS (RECAP)

A node learns about the peers in the network by asking other peers (after an initial bootstrap)

A node maintains a database of all the peers he has heard of and keeps populating it / updating it

A node initiates (and maintain) some outgoing connects and also accept some incoming ones

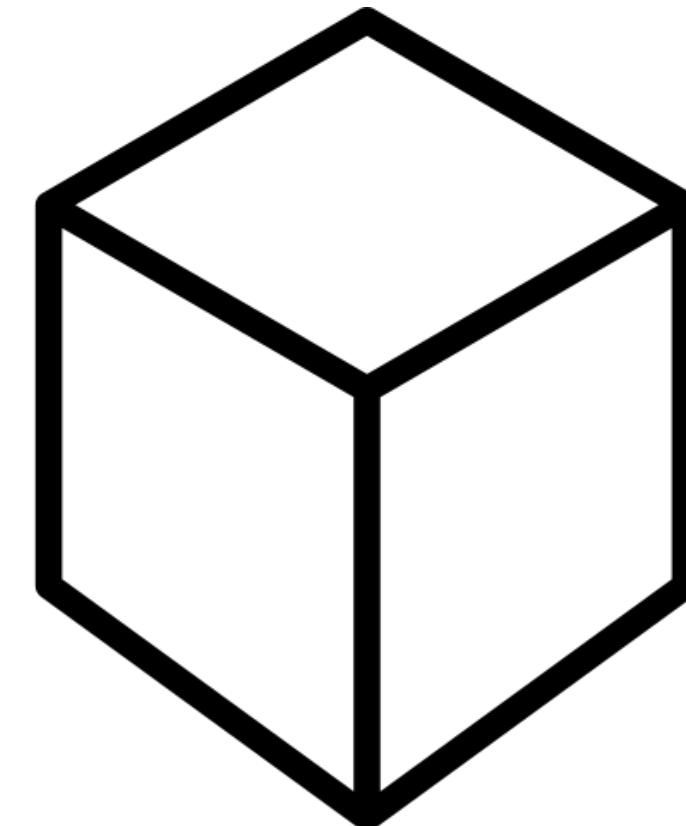
The address of a new node is propagated thought the network so all peers can know about it

Actors and purpose (what, who, why, and how)

THE DATA (WHAT?)

There are two main items that peers share in a cryptocurrency P2P network:
transactions and **blocks**

From: Ford	To: Arthur	42
------------	------------	----



THE ACTORS (WHO?) (1/2)

There are two main roles followed by nodes: **peers** and **miners**

THE ACTORS (WHO?) (1/2)

There are two main roles followed by nodes: **peers** and **miners**

(Normal) **Peers:**

THE ACTORS (WHO?) (1/2)

There are two main roles followed by nodes: **peers** and **miners**

(Normal) **Peers:**

- Can **create transactions** that spend some of their bitcoins

From: Alice	To: Bob	5
-------------	---------	---

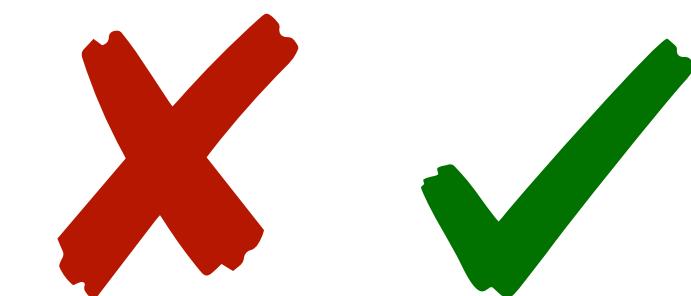
THE ACTORS (WHO?) (1/2)

There are two main roles followed by nodes: **peers** and **miners**

(Normal) Peers:

- Can **create transactions** that spend some of their bitcoins
- Do **verify** the correctness of received **transactions** and **blocks** (from other peers)

From: Alice	To: Bob	5
-------------	---------	---



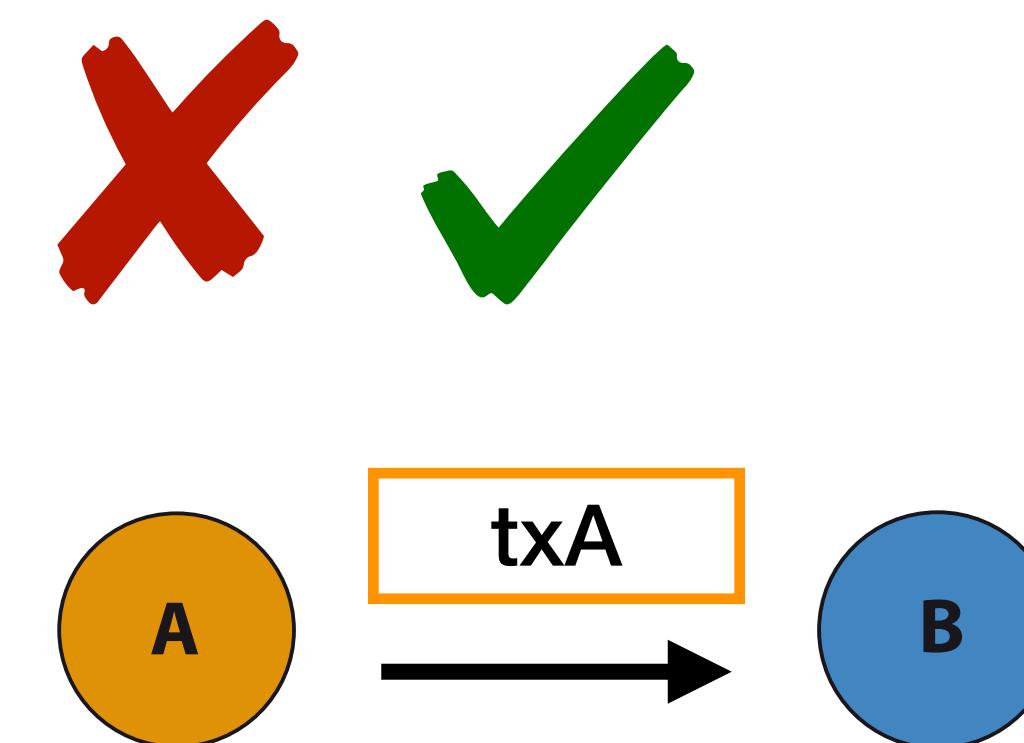
THE ACTORS (WHO?) (1/2)

There are two main roles followed by nodes: **peers** and **miners**

(Normal) Peers:

- Can **create transactions** that spend some of their bitcoins
- Do **verify** the correctness of received **transactions** and **blocks** (from other peers)
- Do **relay** valid **transactions** and **blocks** (created by them or obtained from other peers)

From: Alice	To: Bob	5
-------------	---------	---



THE ACTORS (WHO?) (2/2)

There are two main roles followed by nodes: **peers** and **miners**

THE ACTORS (WHO?) (2/2)

There are two main roles followed by nodes: **peers** and **miners**

Miners:



THE ACTORS (WHO?) (2/2)

There are two main roles followed by nodes: **peers** and **miners**

Miners:

- Can everything a **peer** could do*



THE ACTORS (WHO?) (2/2)

There are two main roles followed by nodes: **peers** and **miners**

Miners:

- Can everything a **peer** could do*
- Can **generate blocks** through a process known as mining



THE ACTORS (WHO?) (2/2)

There are two main roles followed by nodes: **peers** and **miners**

Miners:

- Can everything a **peer** could do*
- Can **generate blocks** through a process known as mining



* There are specific purpose miners (ASICs) that only perform mining

THE PURPOSE (WHY?)

Peers relay transactions in order to reach miners, which will include such transactions in future blocks

Miners generate blocks to obtain their reward (and also the transactions fees)

Blocks are relayed to ultimately achieve a consistent view of the blockchain

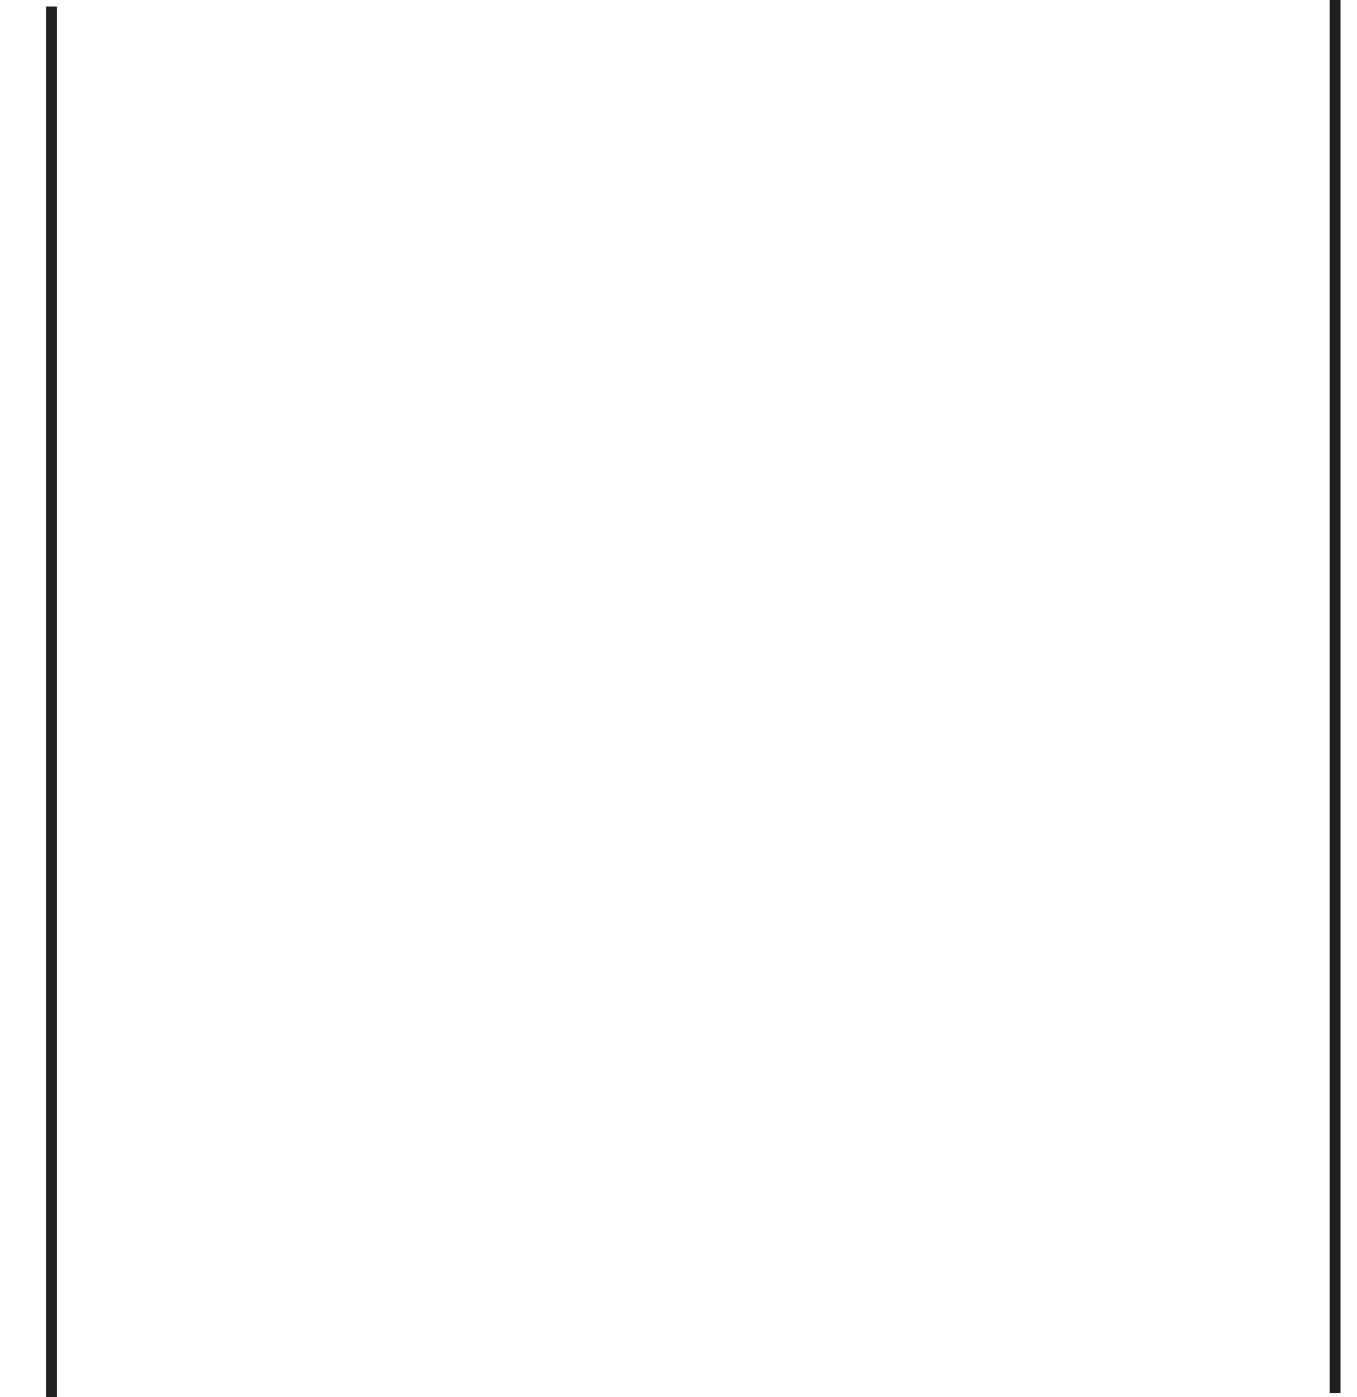
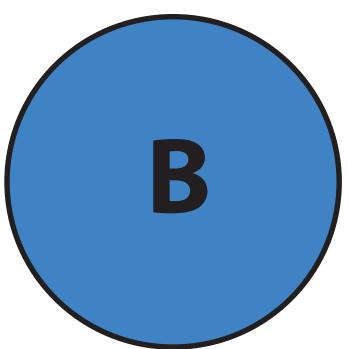
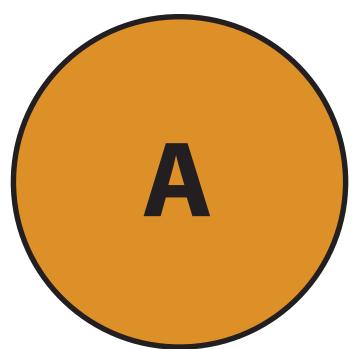
Peers validate transactions and blocks (and relay only the valid ones) in order to avoid cheating (e.g: double-spending, coin forgery, etc)

THE GOSSIP PROTOCOL (HOW?)



Items (transactions and blocks) are shared between peers in a push manner

Announce paradigm



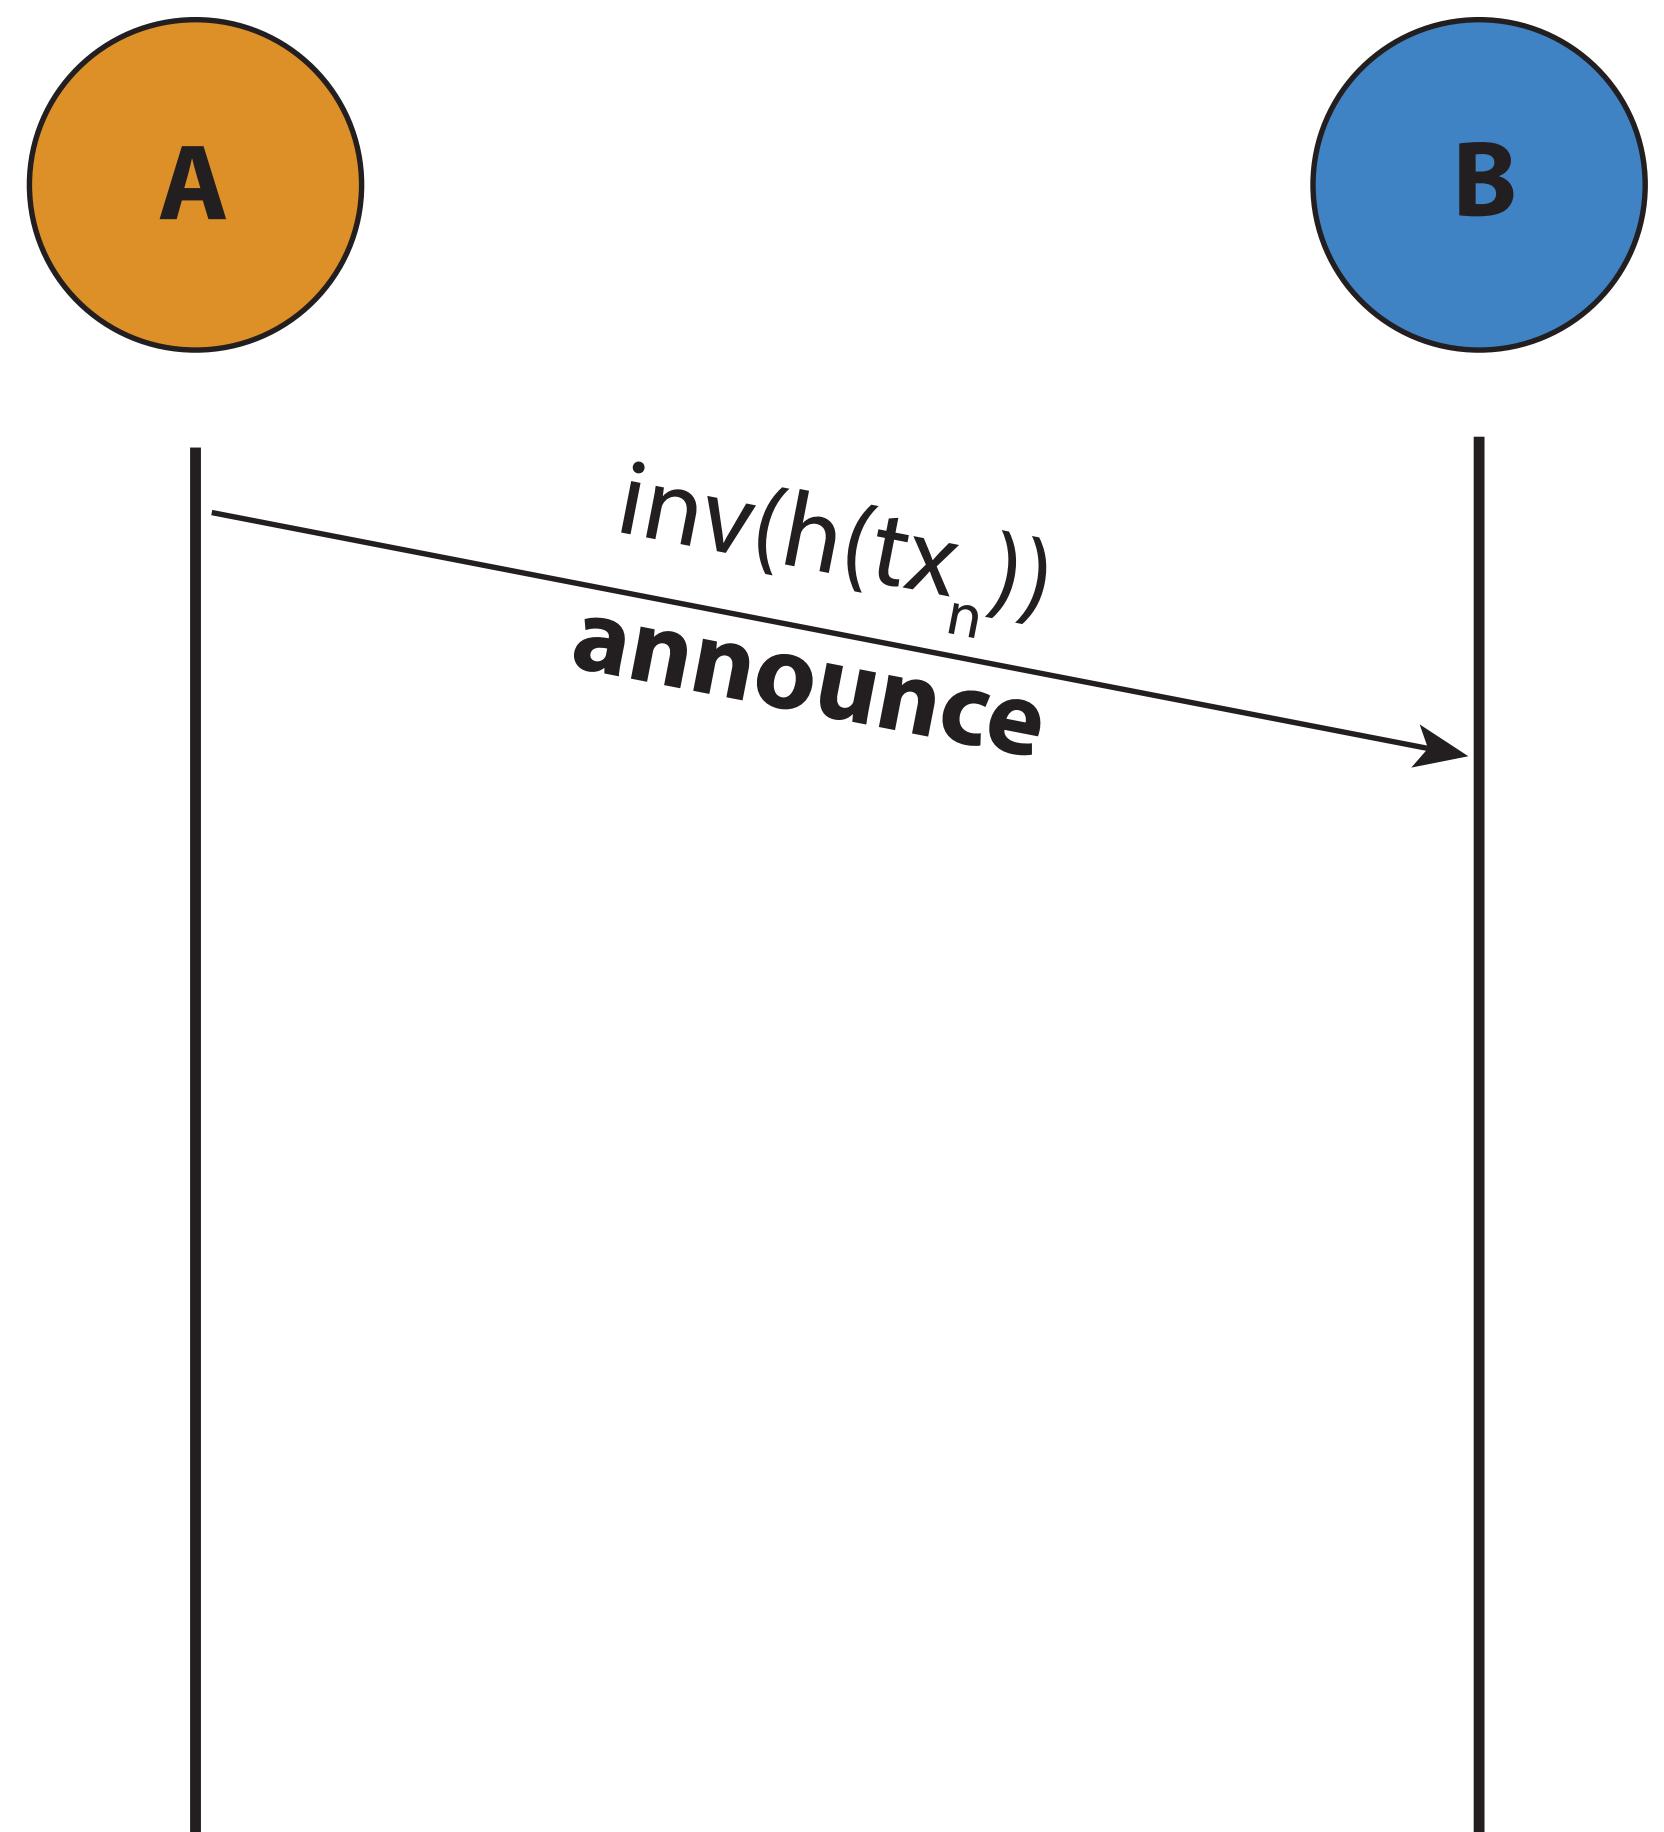
THE GOSSIP PROTOCOL (HOW?)



Items (transactions and blocks) are shared between peers in a push manner

When a peer receives / generates a new item he announce it to his neighbors (**announce**)

Announce paradigm



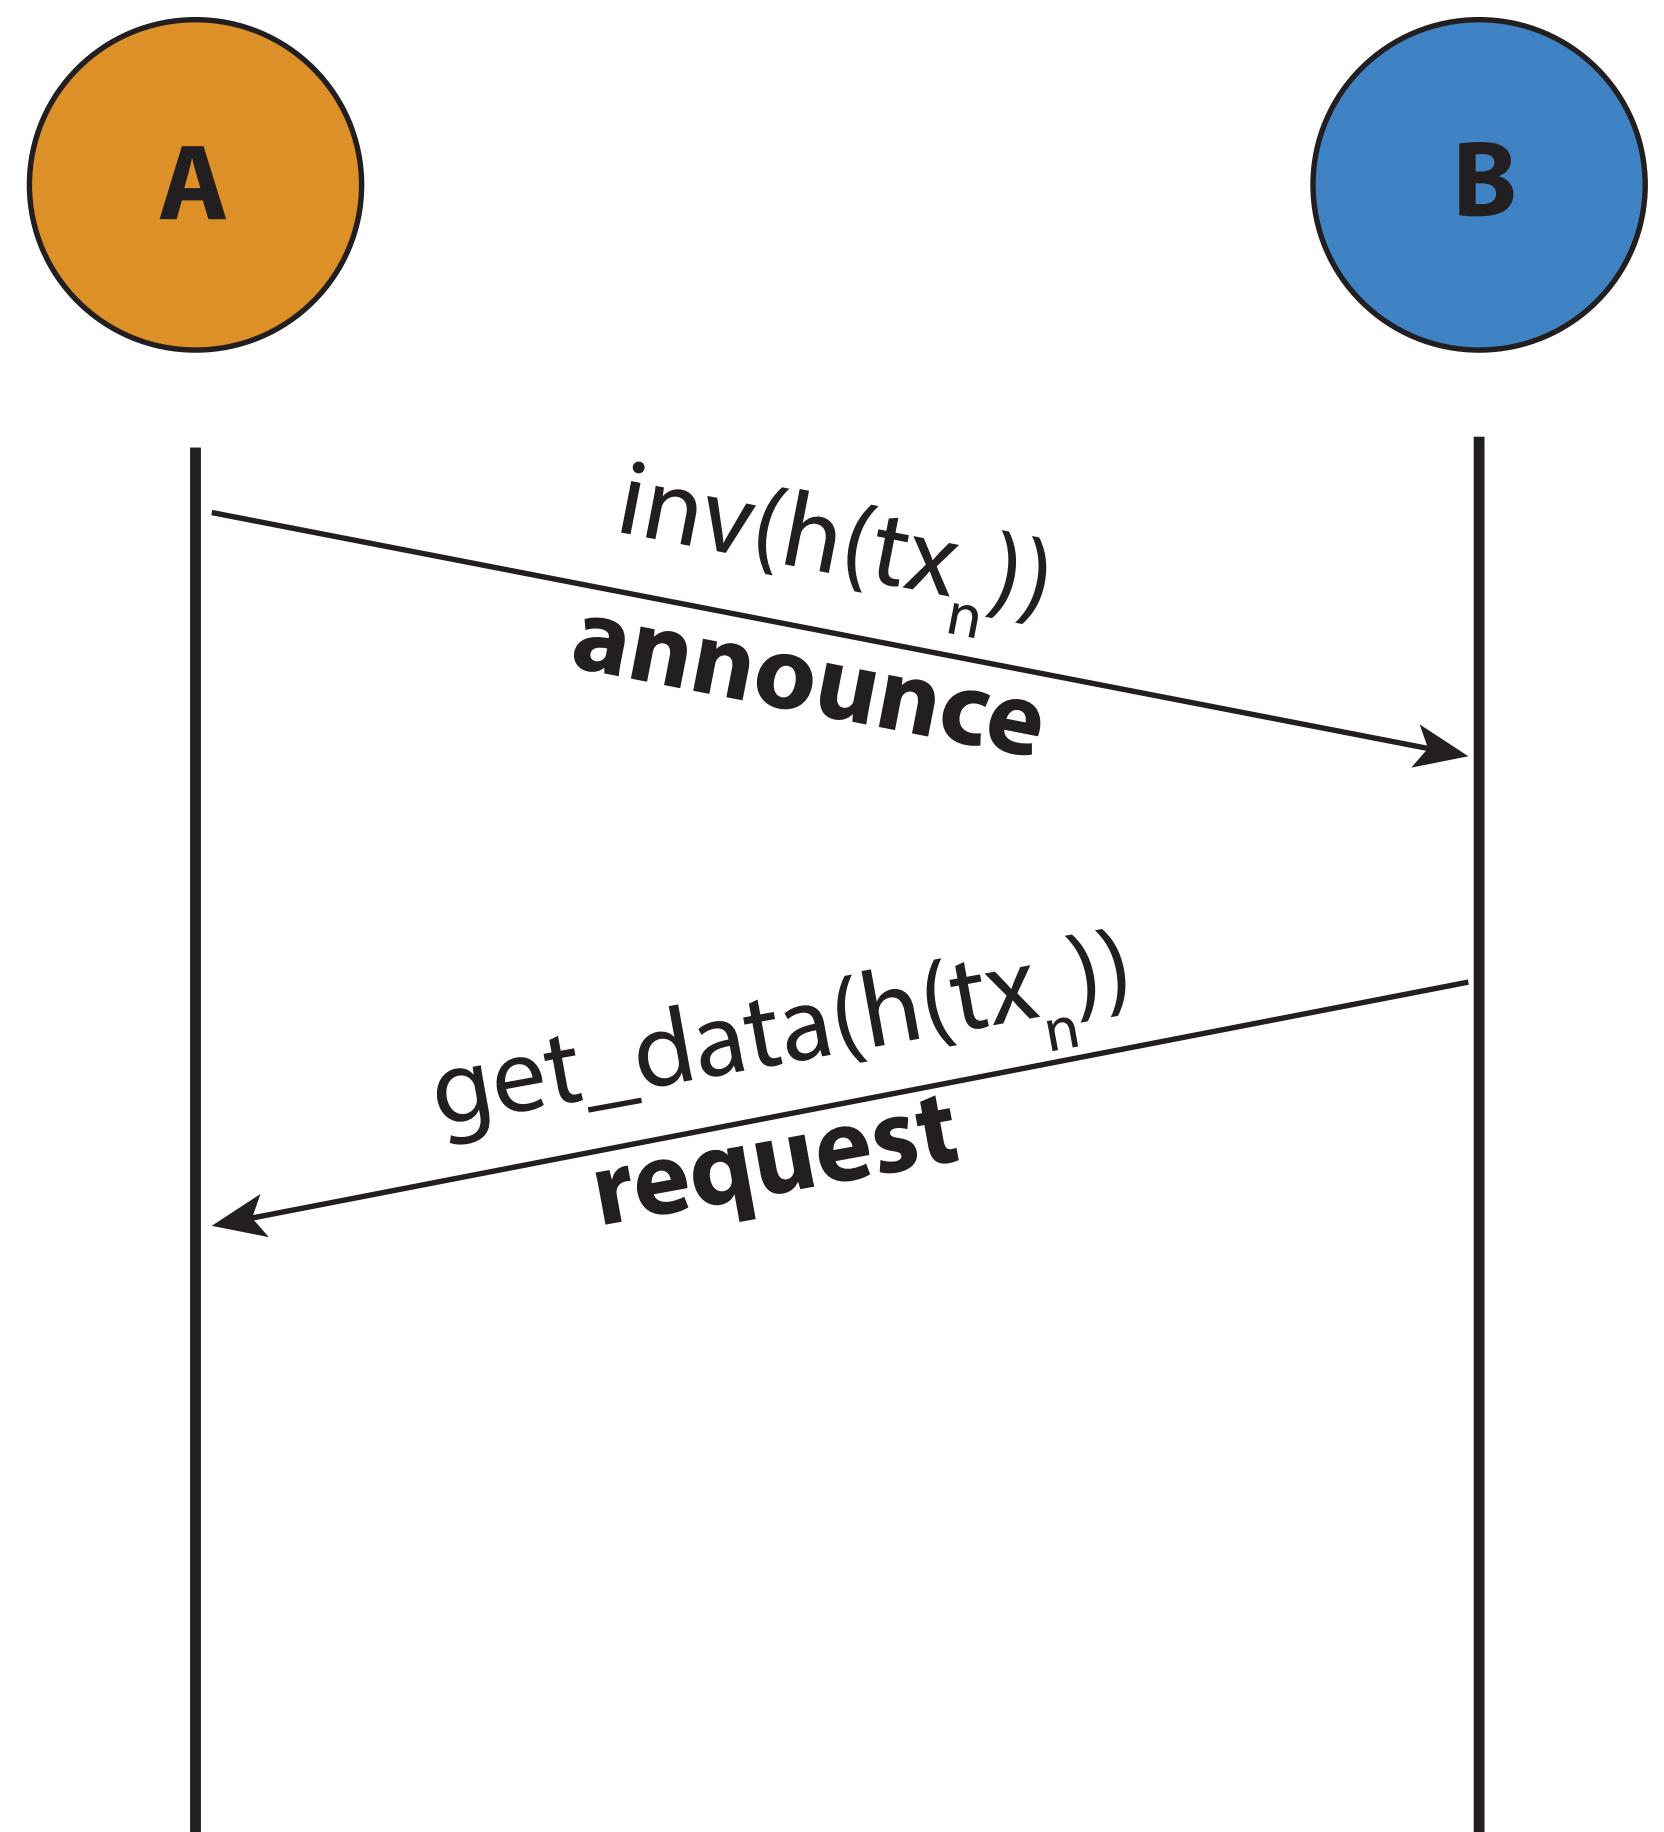
THE GOSSIP PROTOCOL (HOW?)

Items (transactions and blocks) are shared between peers in a push manner

When a peer receives / generates a new item he announce it to his neighbors (**announce**)

Upon receiving an announce of an item, a node that does not know about it will request the item back to the announcer (**request**)

Announce paradigm



THE GOSSIP PROTOCOL (HOW?)

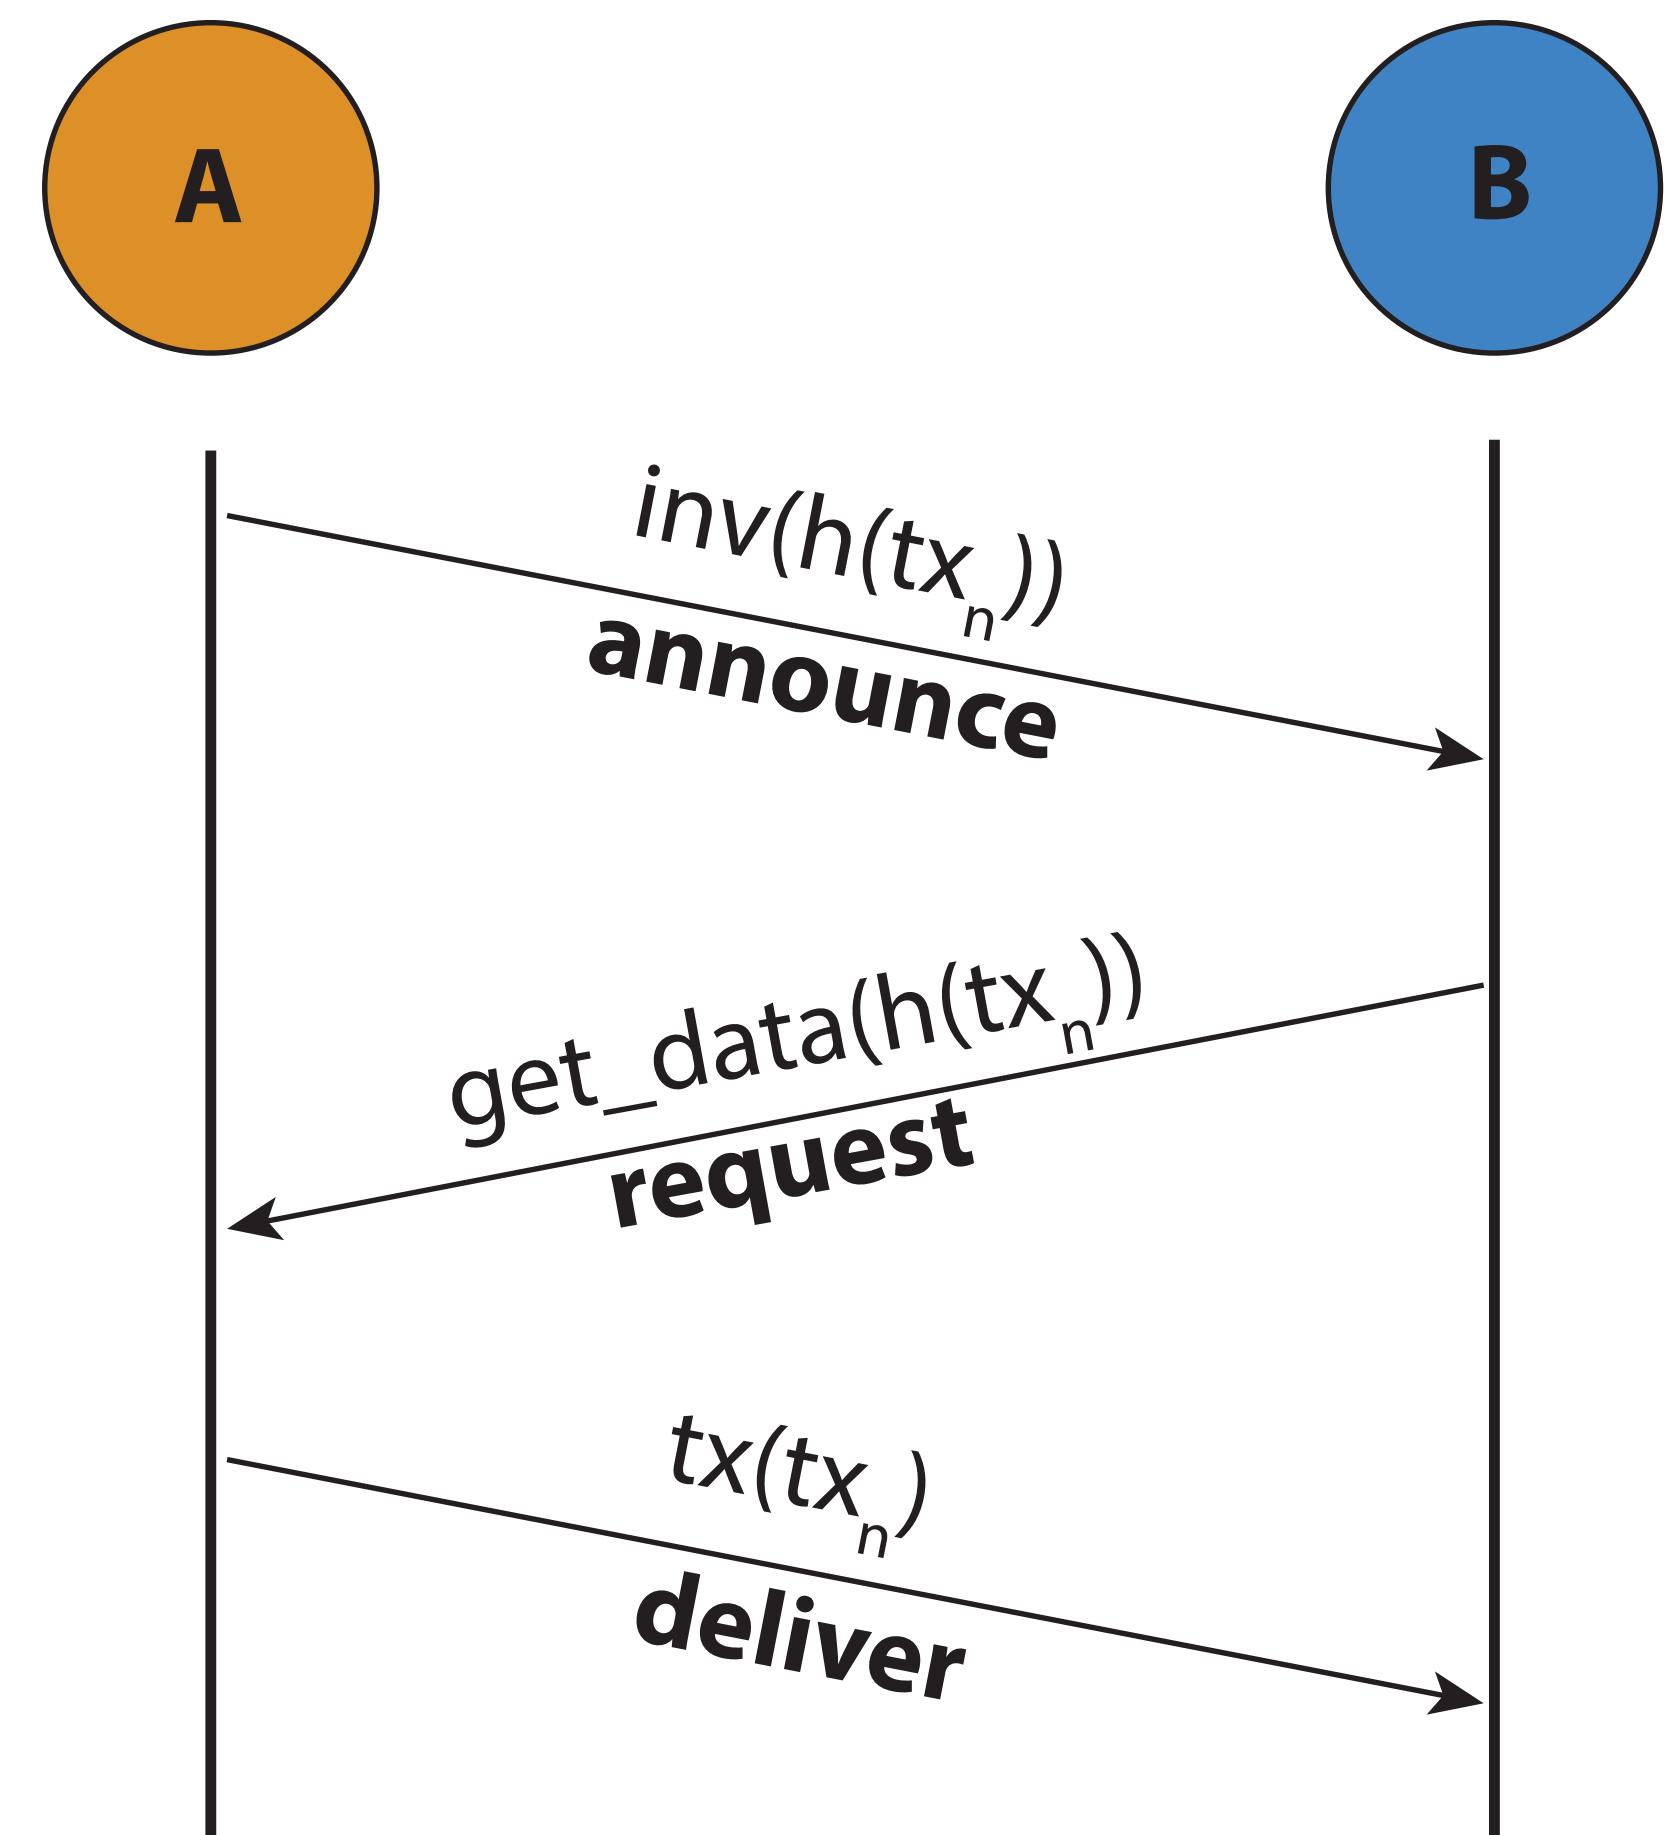
Items (transactions and blocks) are shared between peers in a push manner

When a peer receives / generates a new item he announce it to his neighbors (**announce**)

Upon receiving an announce of an item, a node that does not know about it will request the item back to the announcer (**request**)

Upon receiving a request of a known item, a node will reply back with it (**deliver**)

Announce paradigm

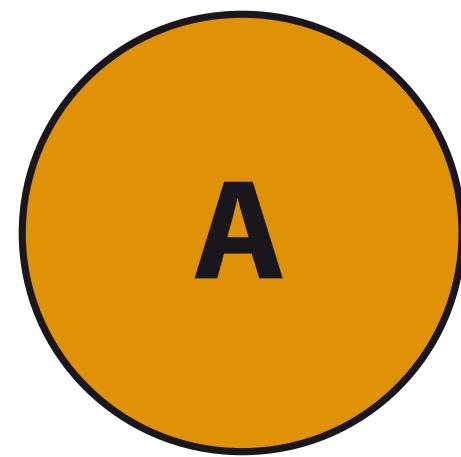




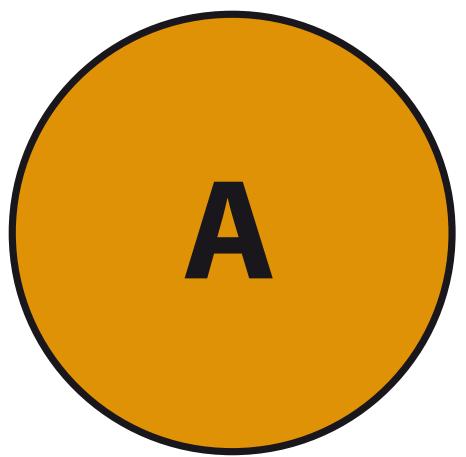
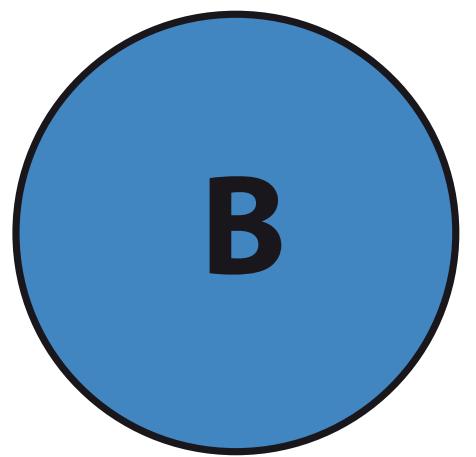
**BREAK
TIME**

Information propagation

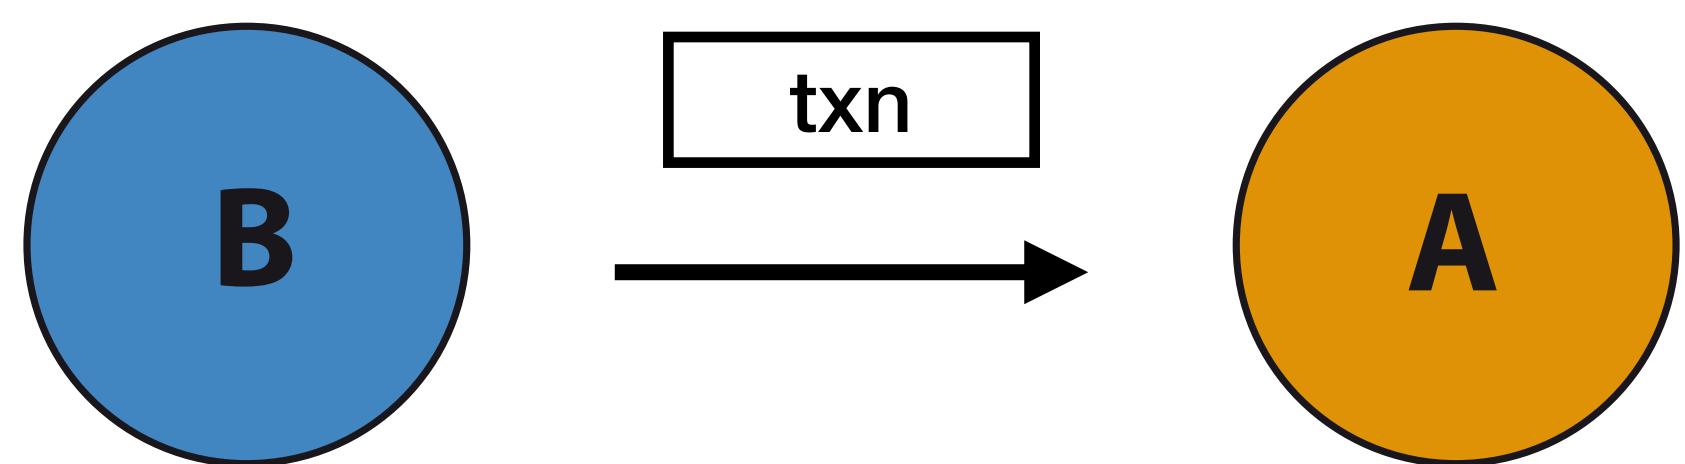
INFORMATION PROPAGATION (1/3)



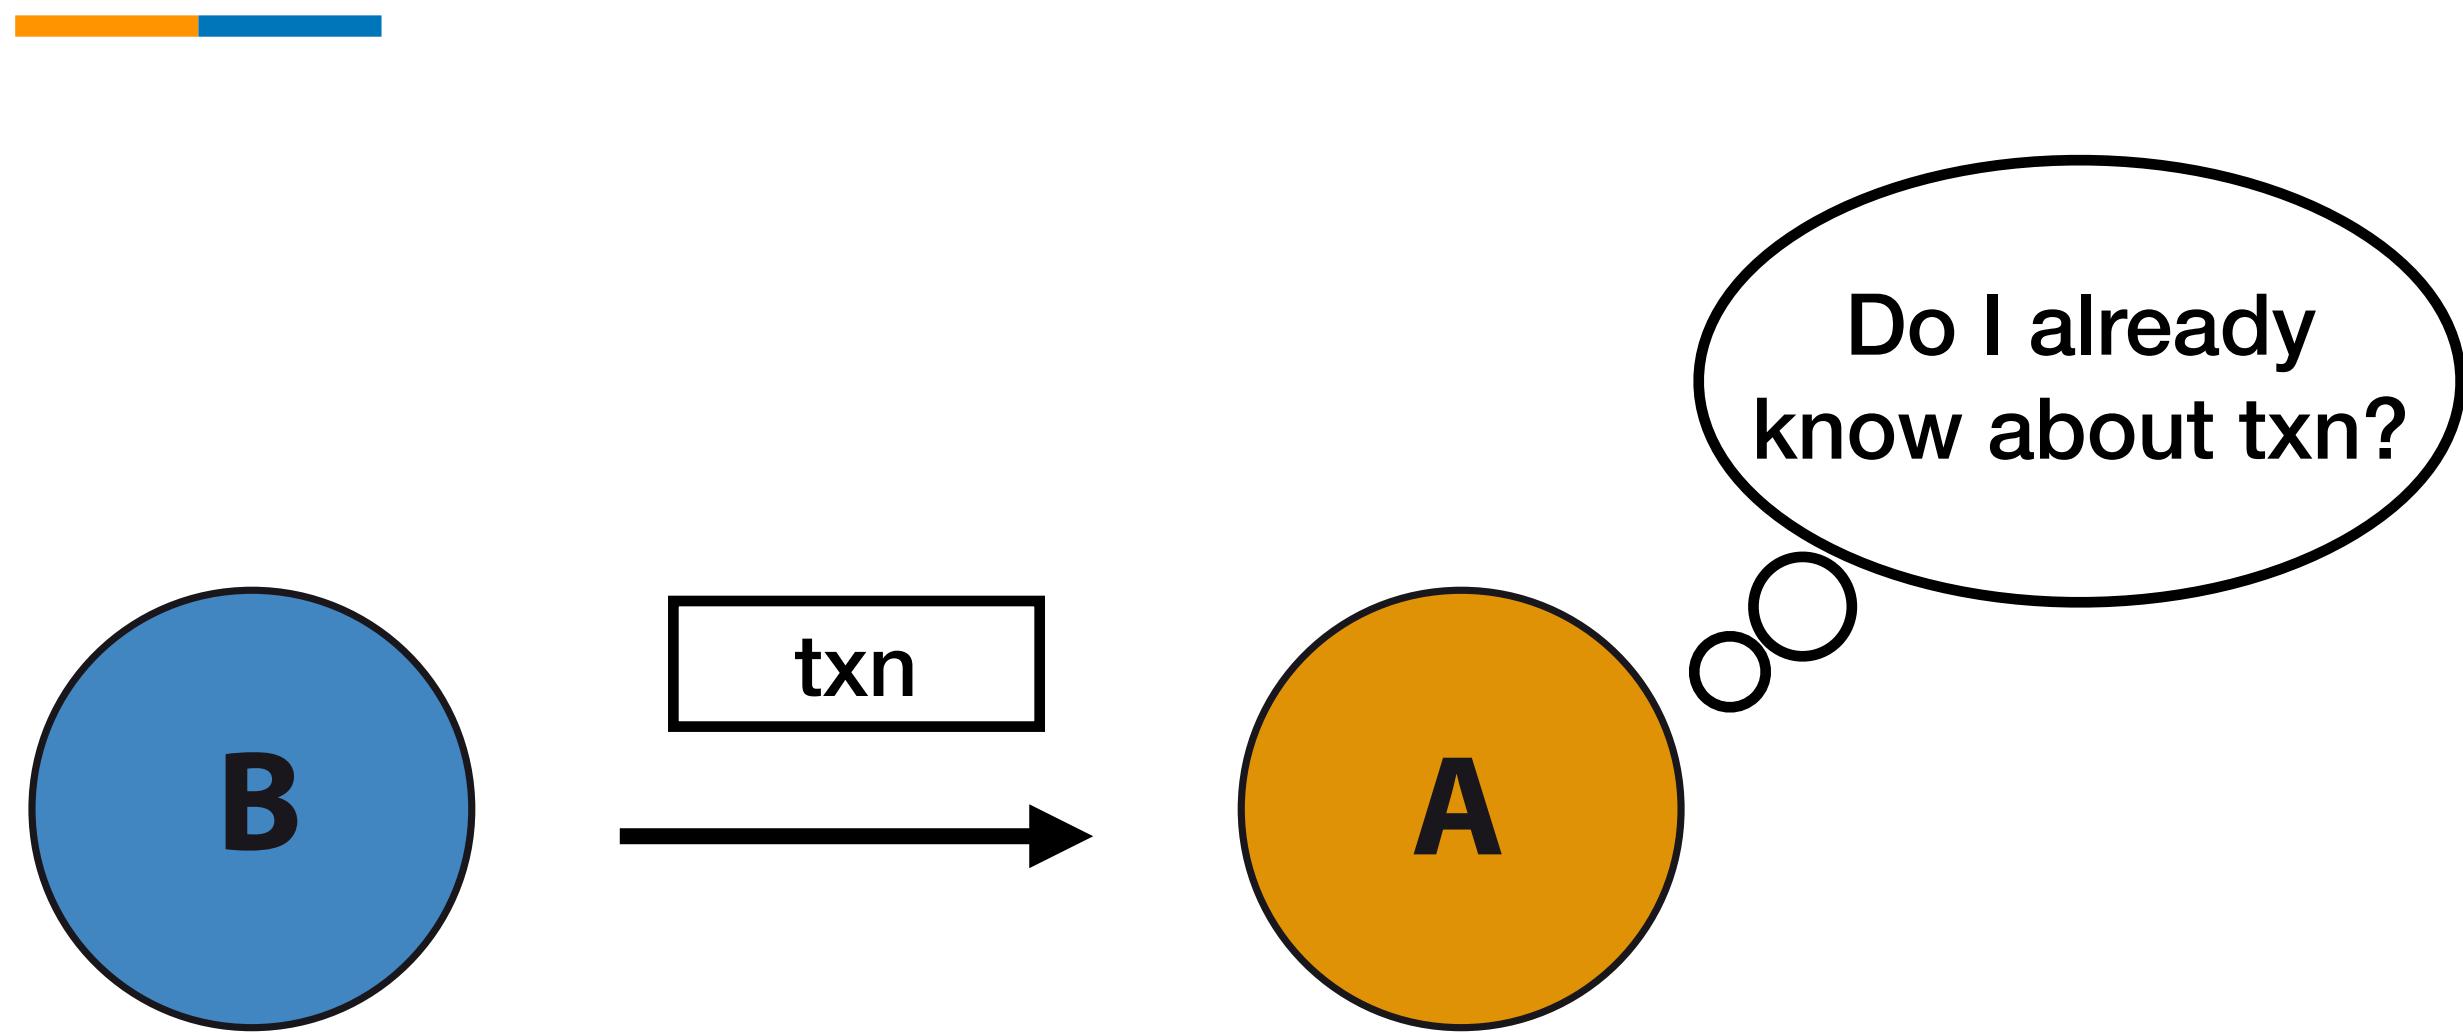
INFORMATION PROPAGATION (1/3)



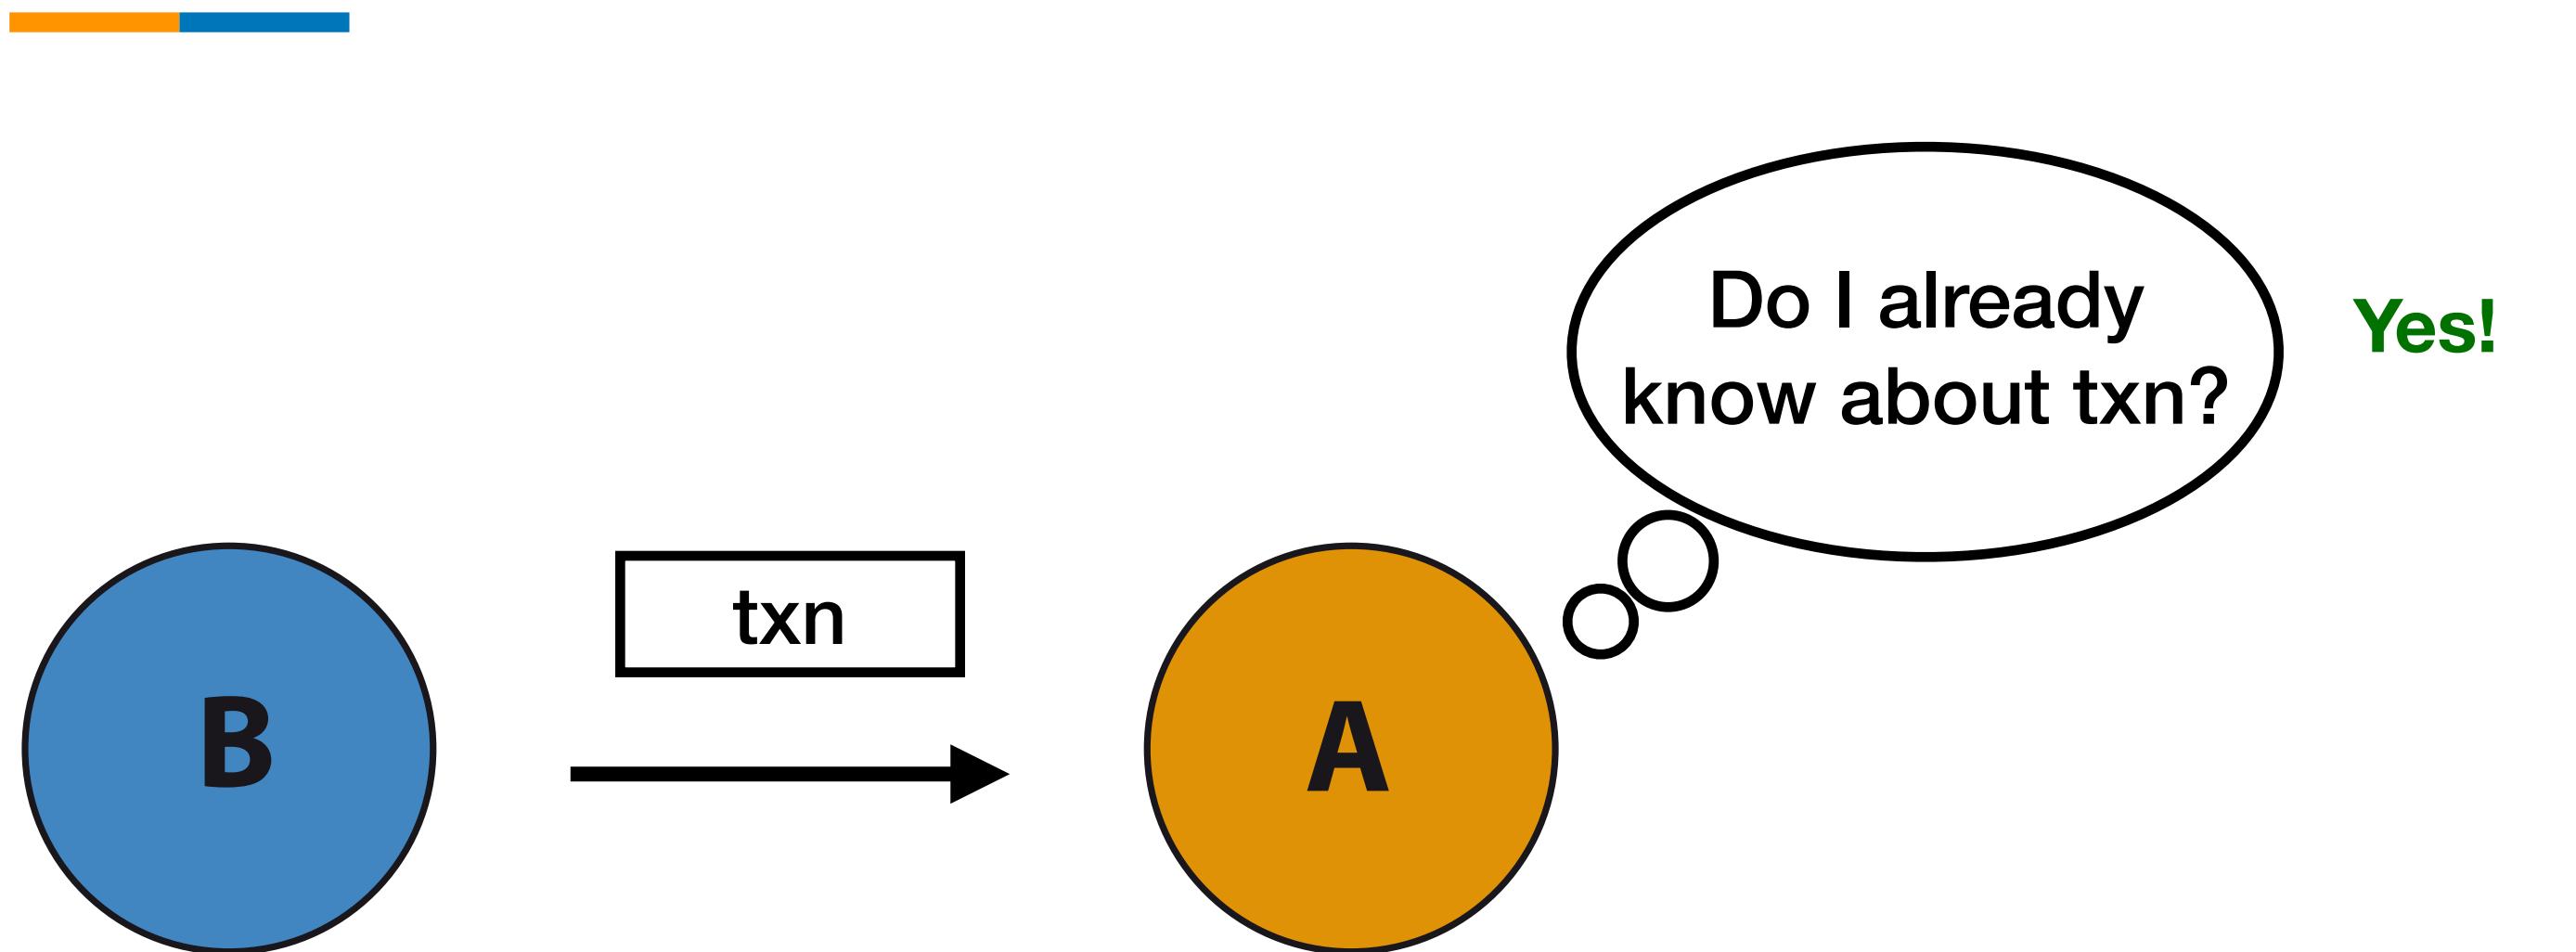
INFORMATION PROPAGATION (1/3)



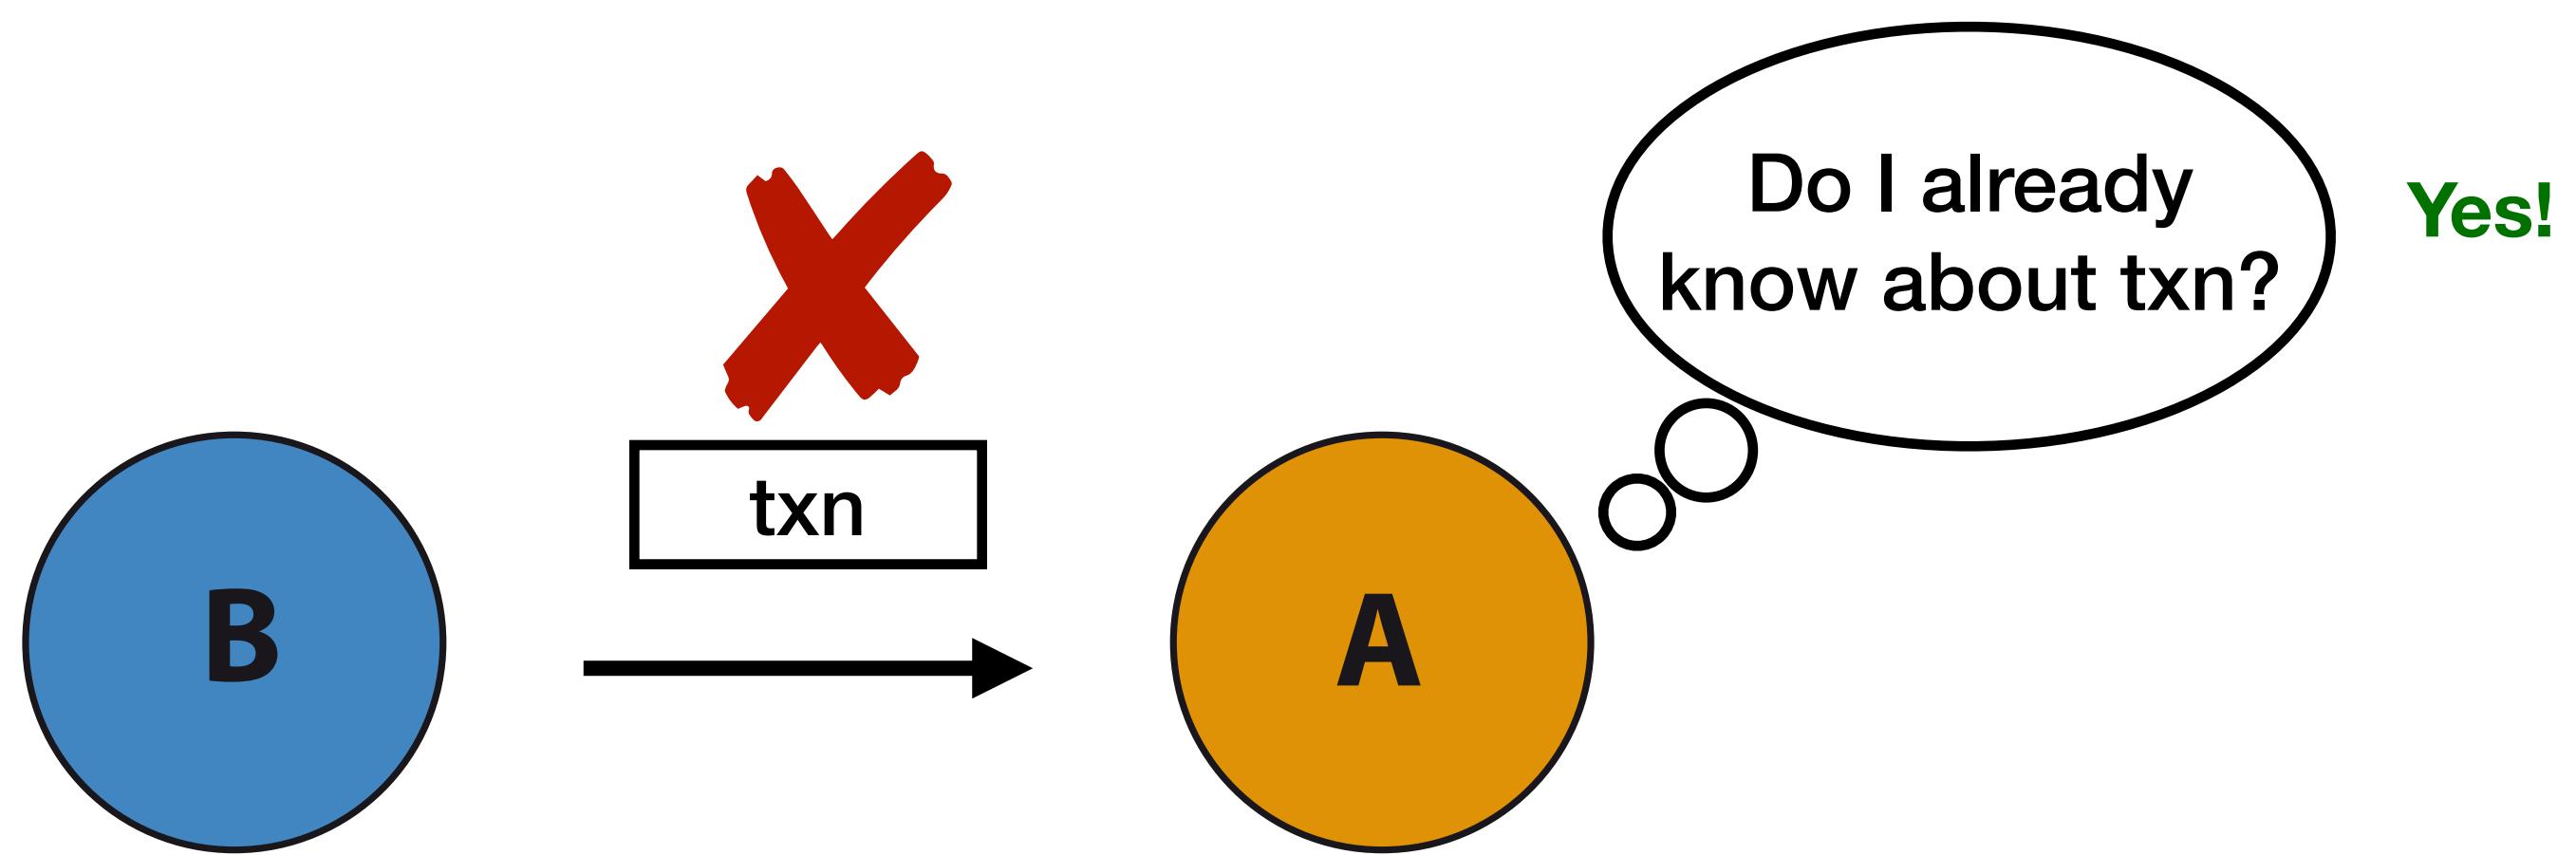
INFORMATION PROPAGATION (1/3)



INFORMATION PROPAGATION (1/3)

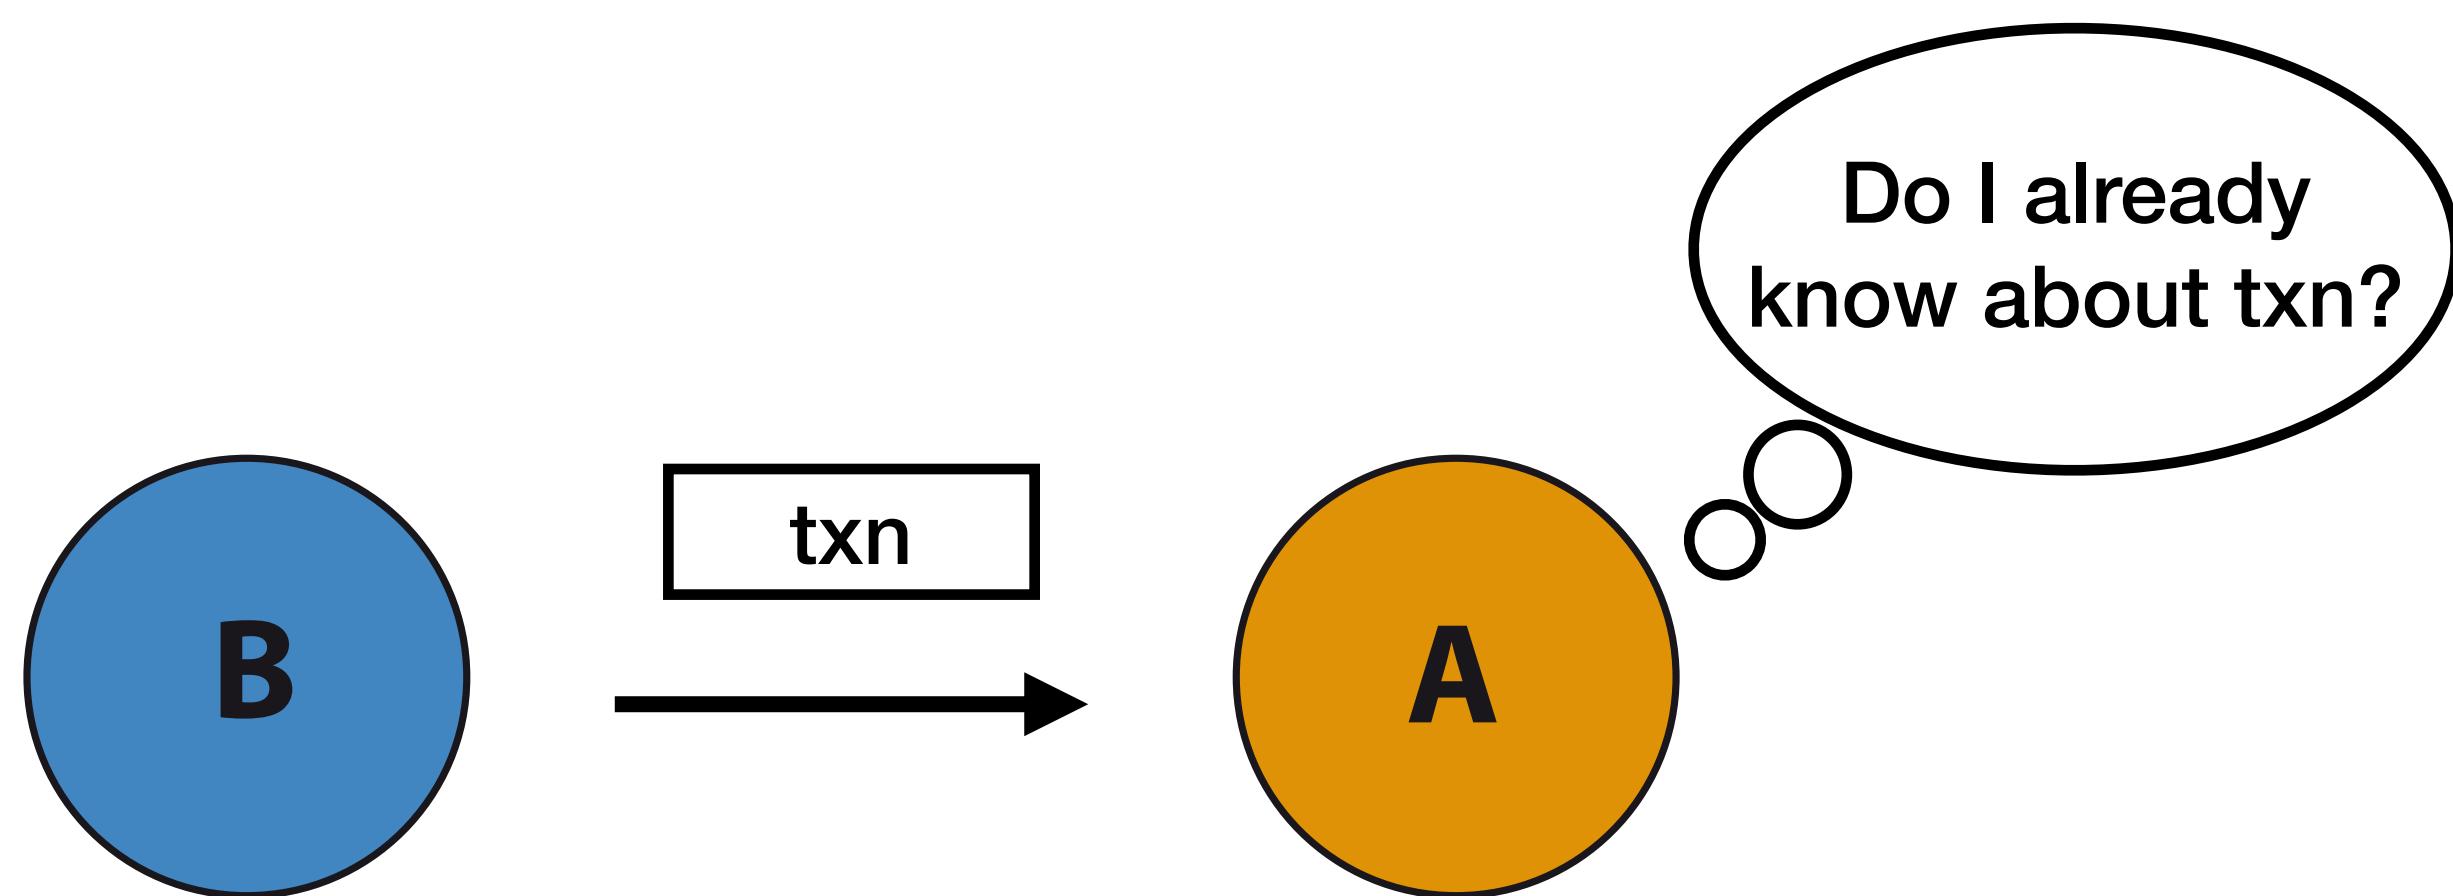


INFORMATION PROPAGATION (1/3)



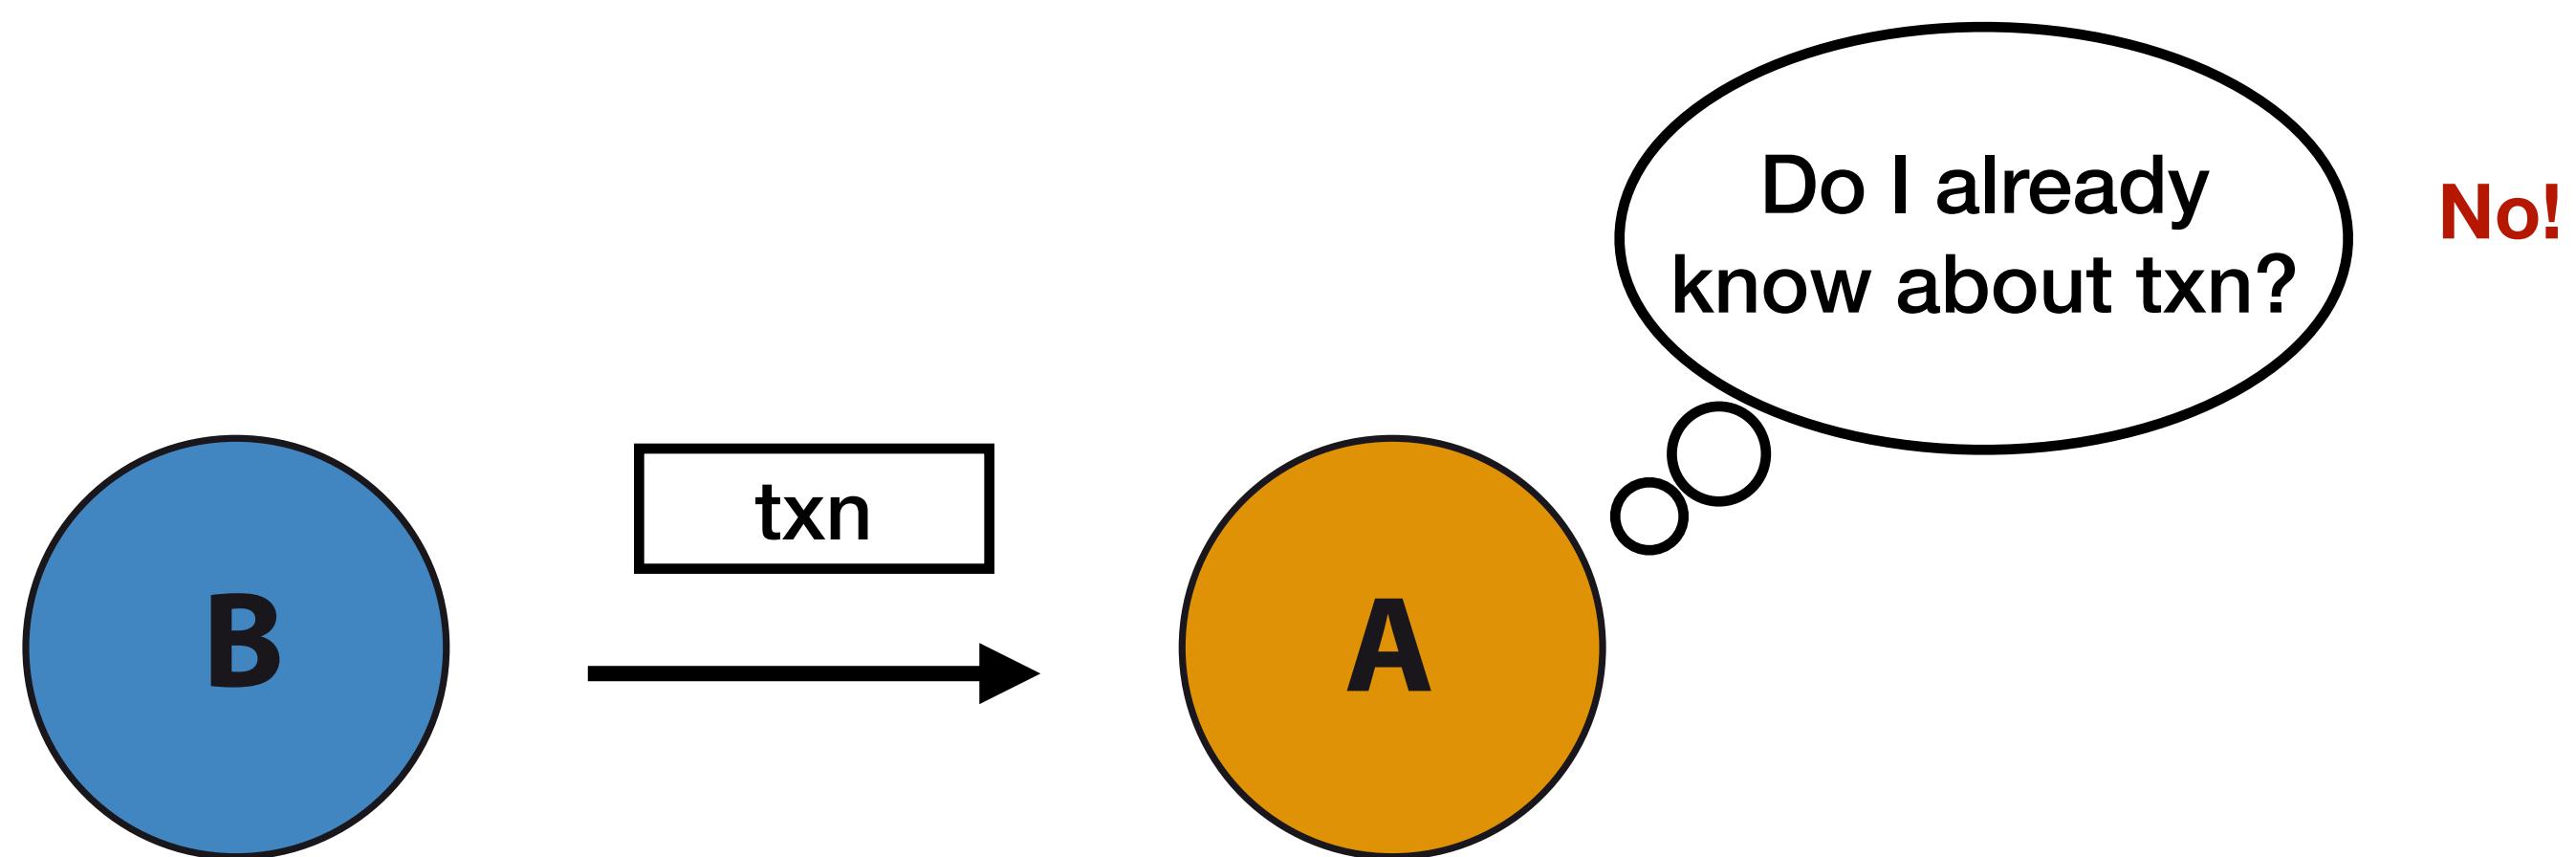
- Known transaction will be rejected

INFORMATION PROPAGATION (1/3)



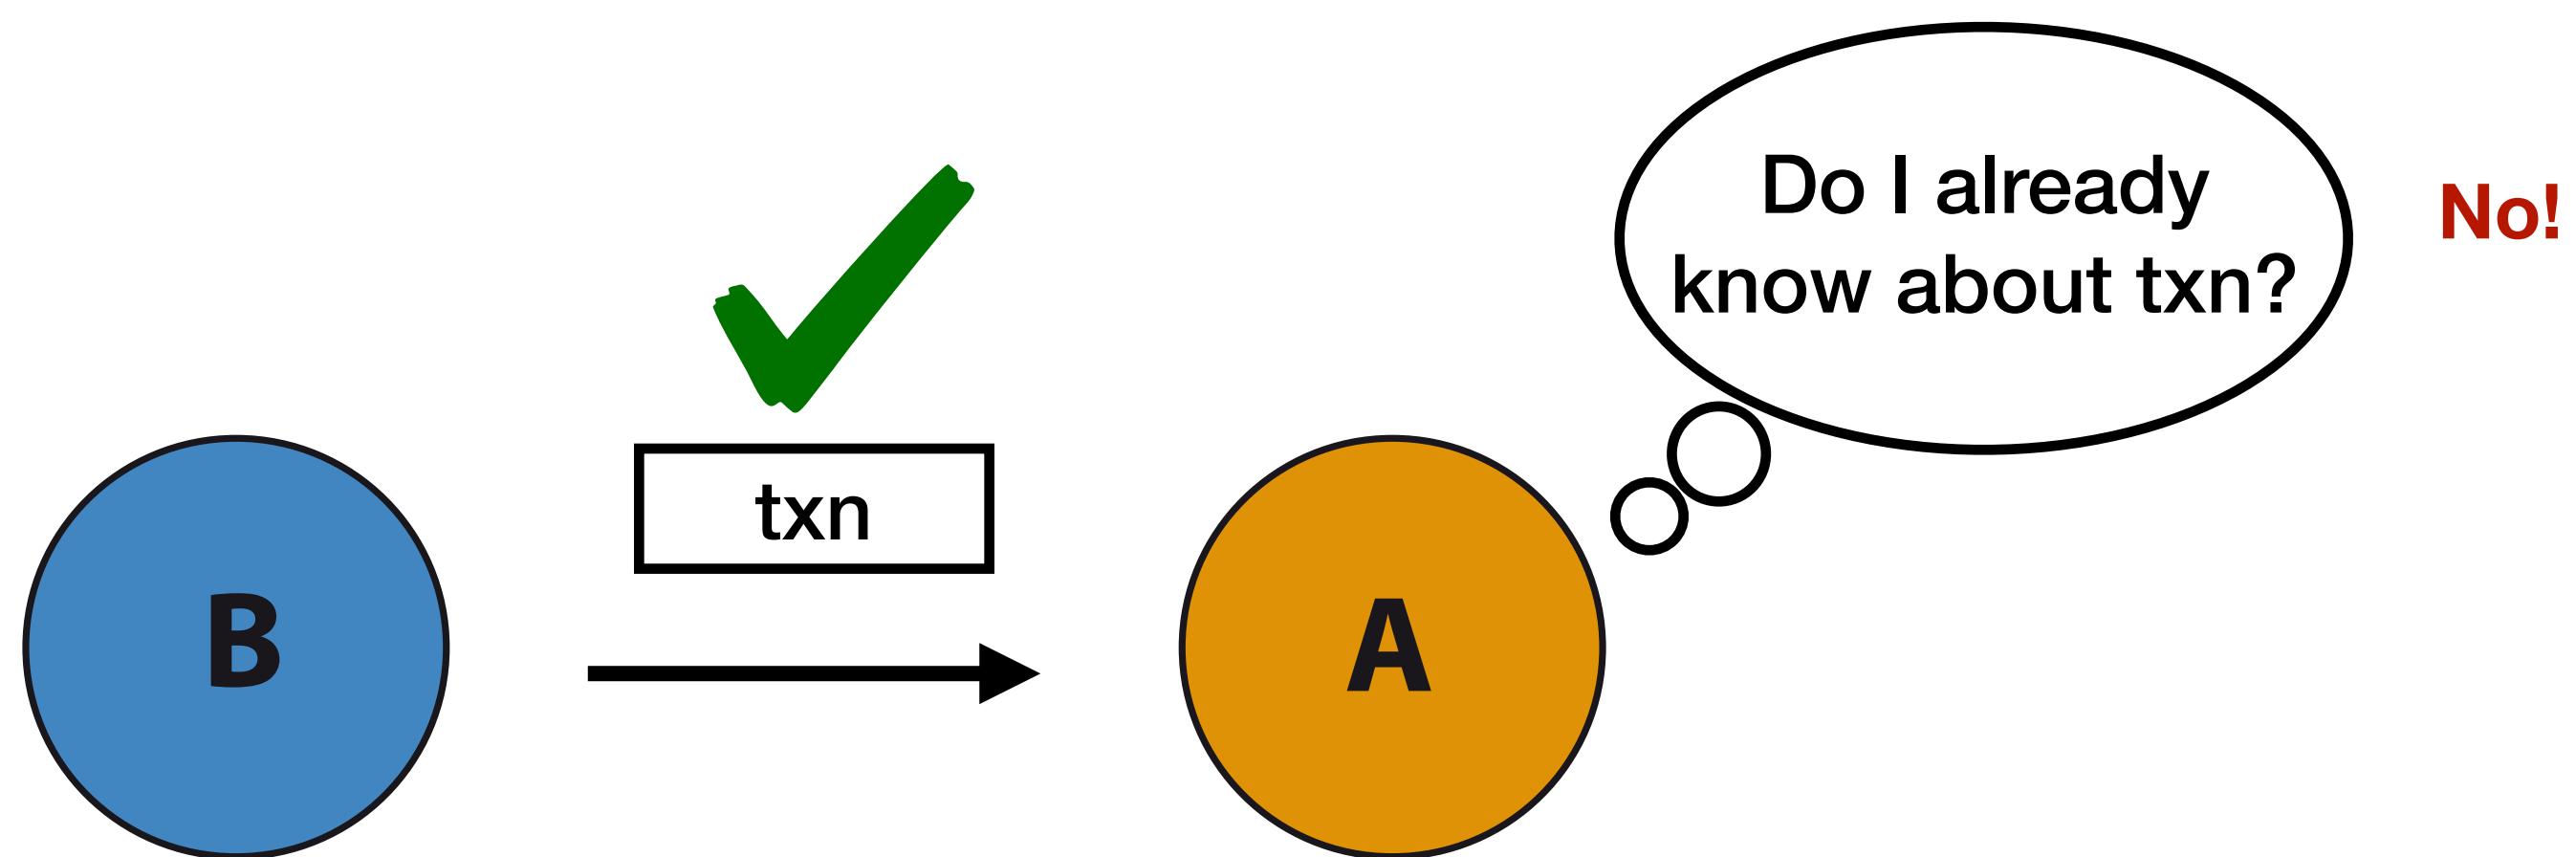
- Known transaction will be rejected

INFORMATION PROPAGATION (1/3)



- Known transaction will be rejected

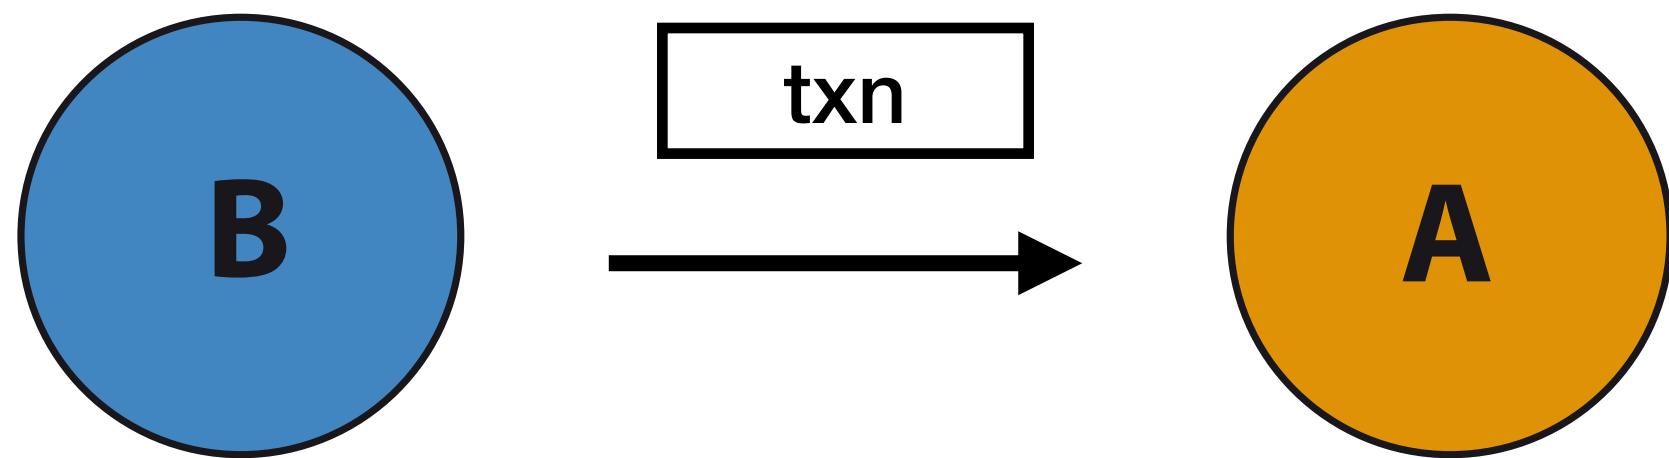
INFORMATION PROPAGATION (1/3)



- Known transaction will be rejected

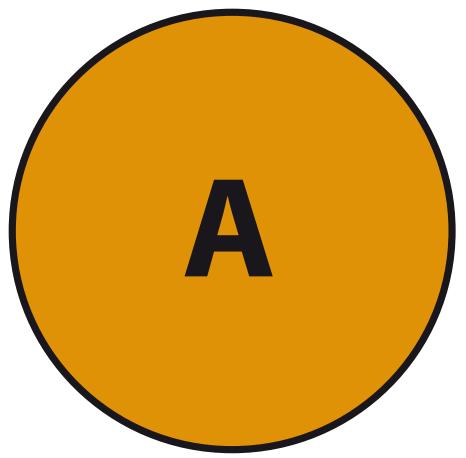
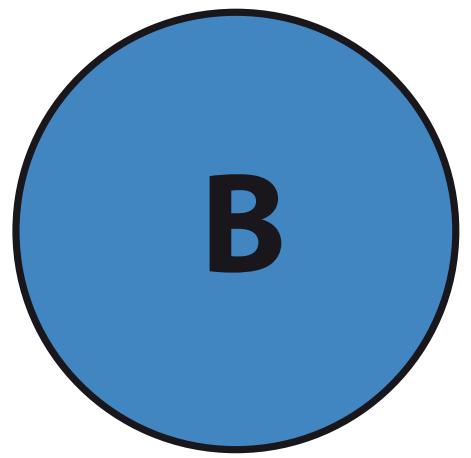
INFORMATION PROPAGATION (1/3)

- Known transaction will be rejected



INFORMATION PROPAGATION (1/3)

- Known transaction will be rejected

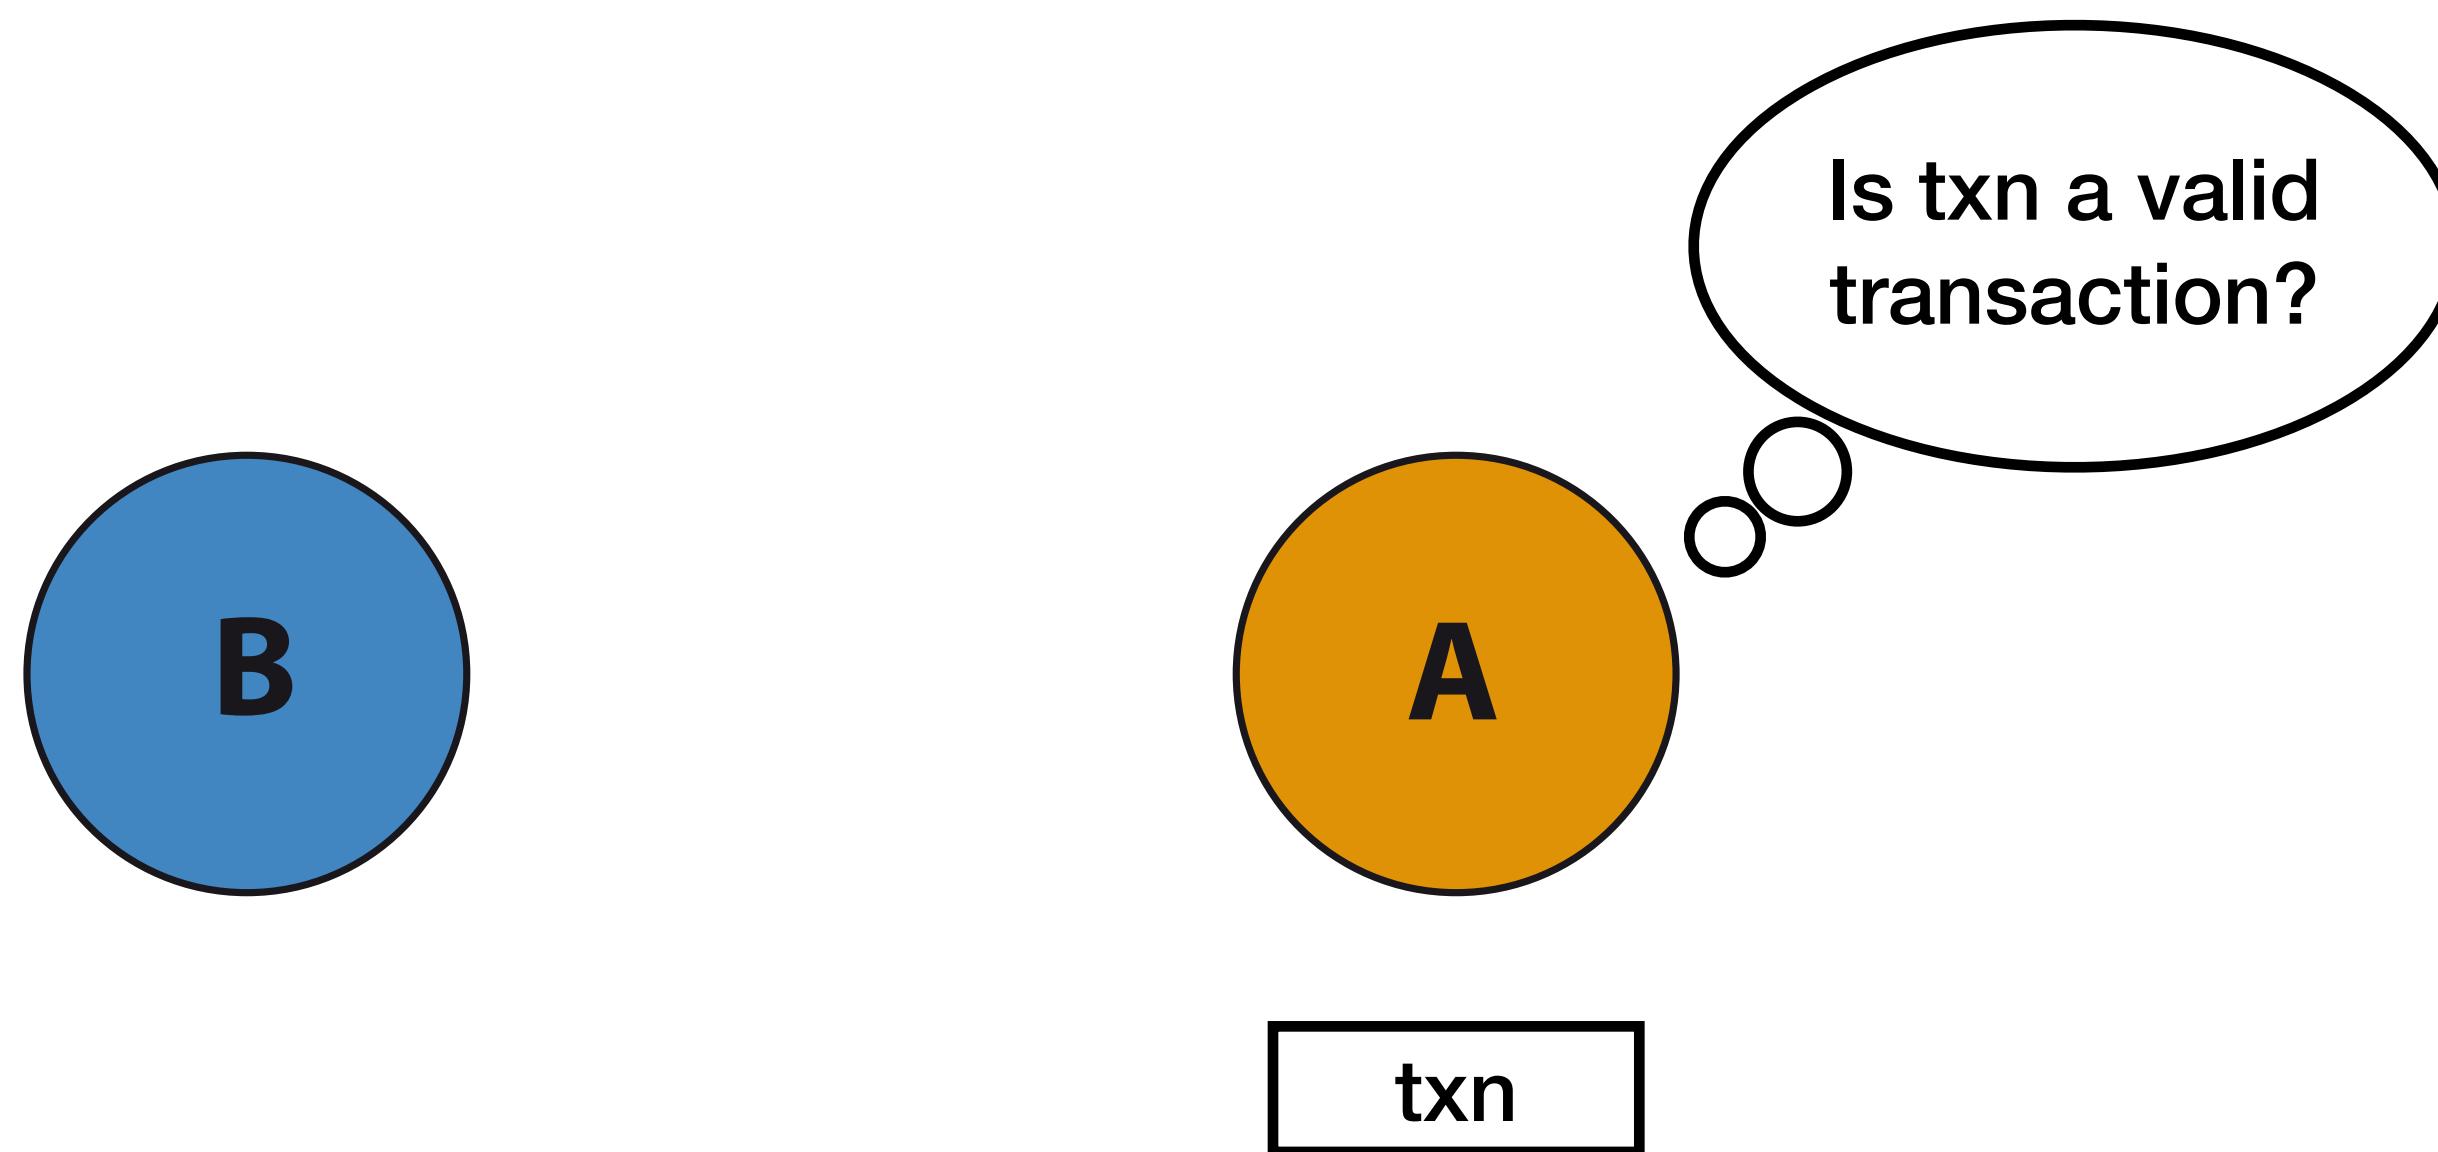


INFORMATION PROPAGATION (1/3)

- Known transaction will be rejected

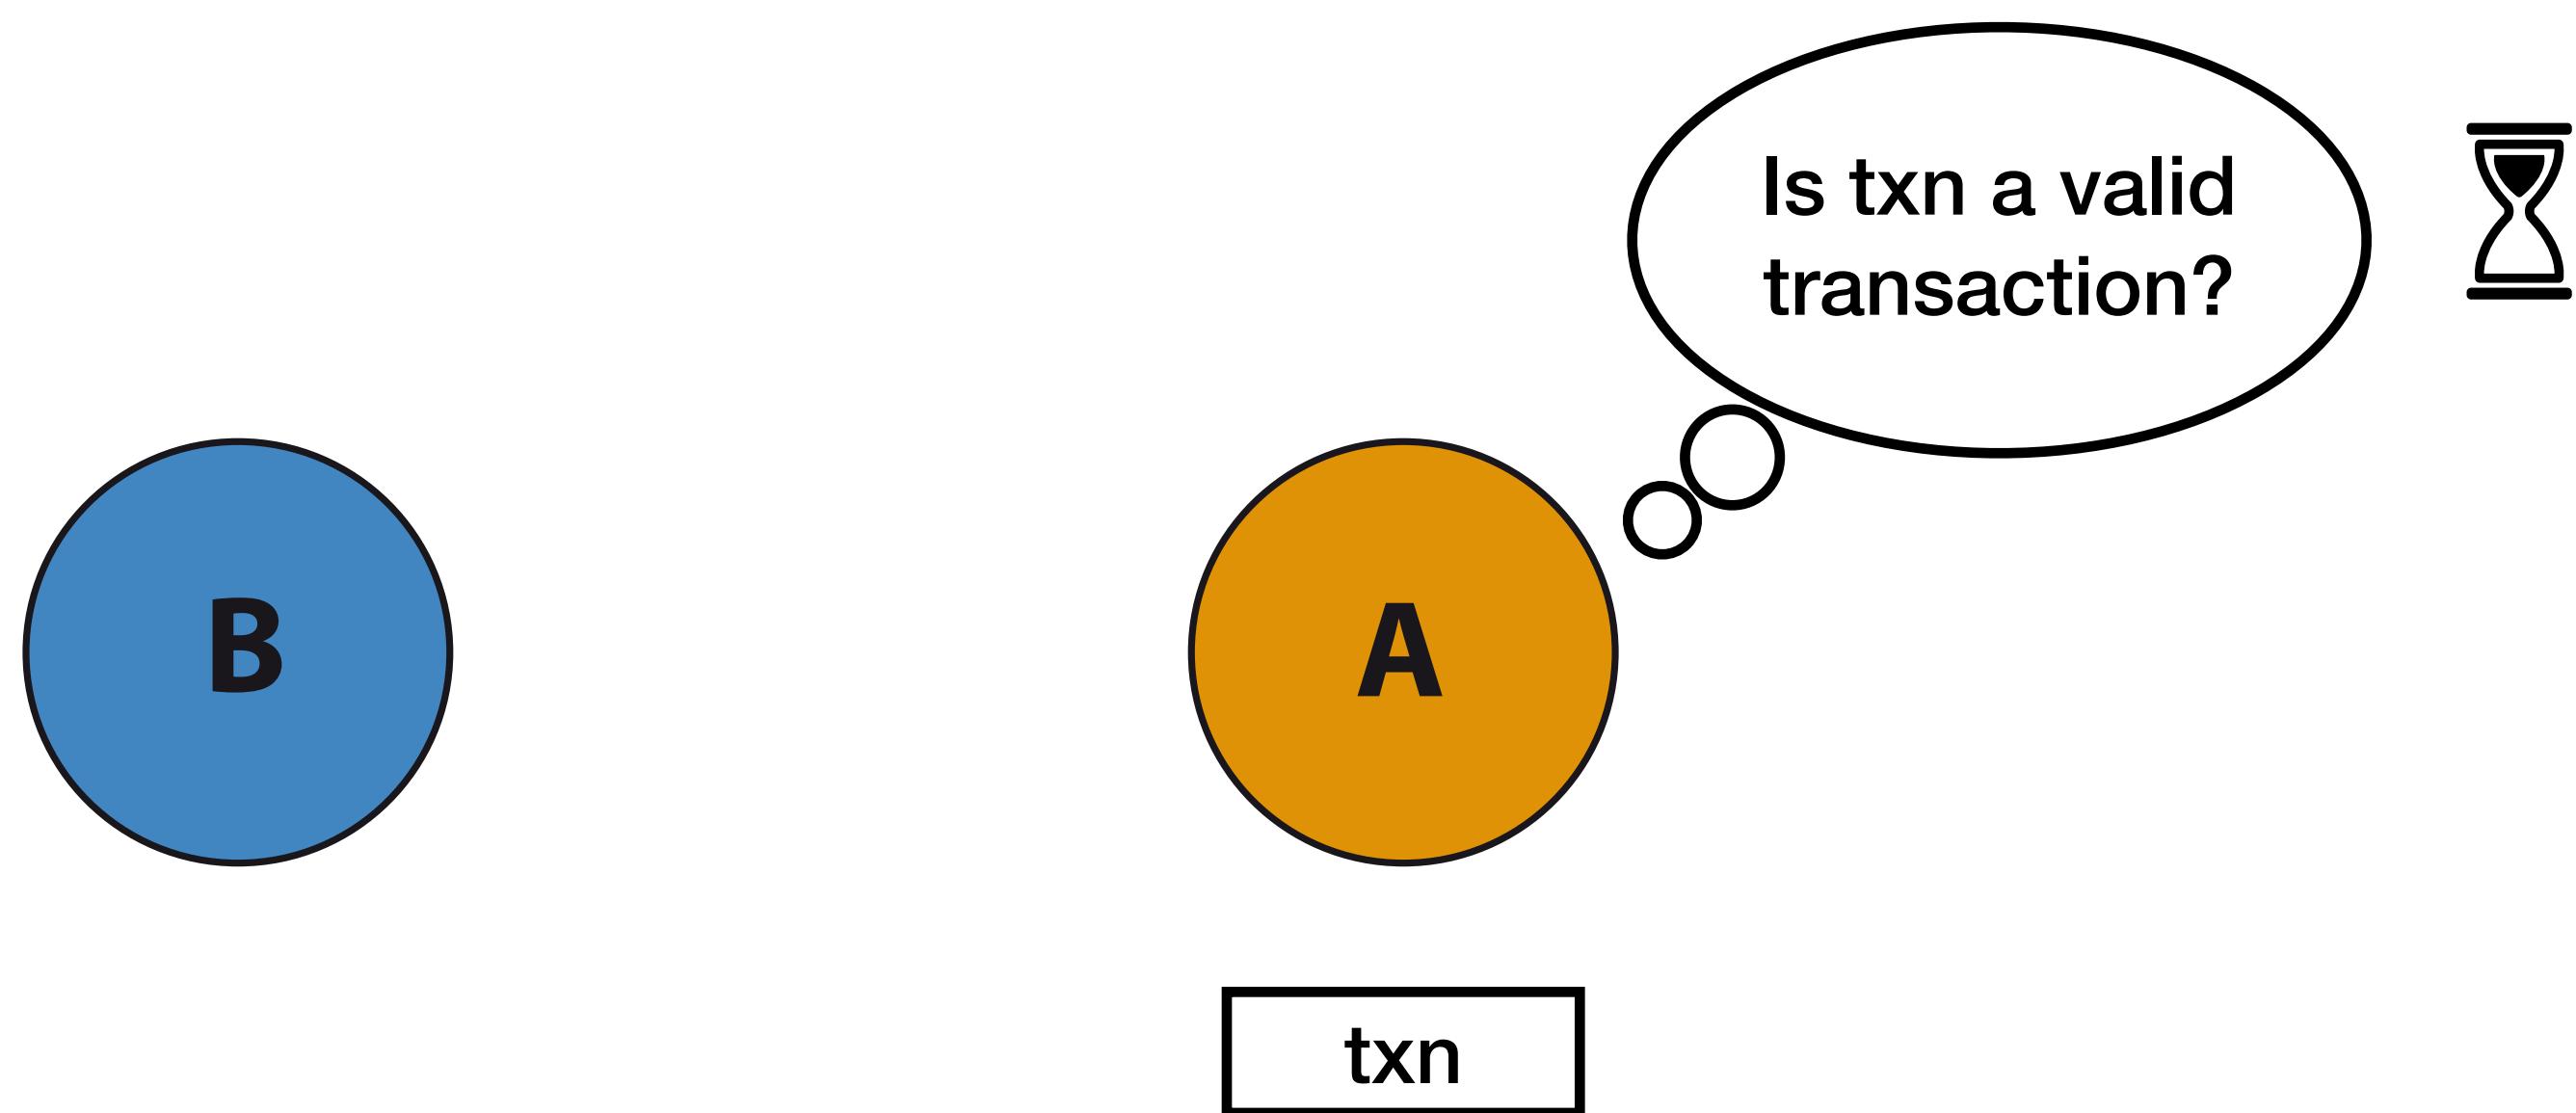


INFORMATION PROPAGATION (1/3)



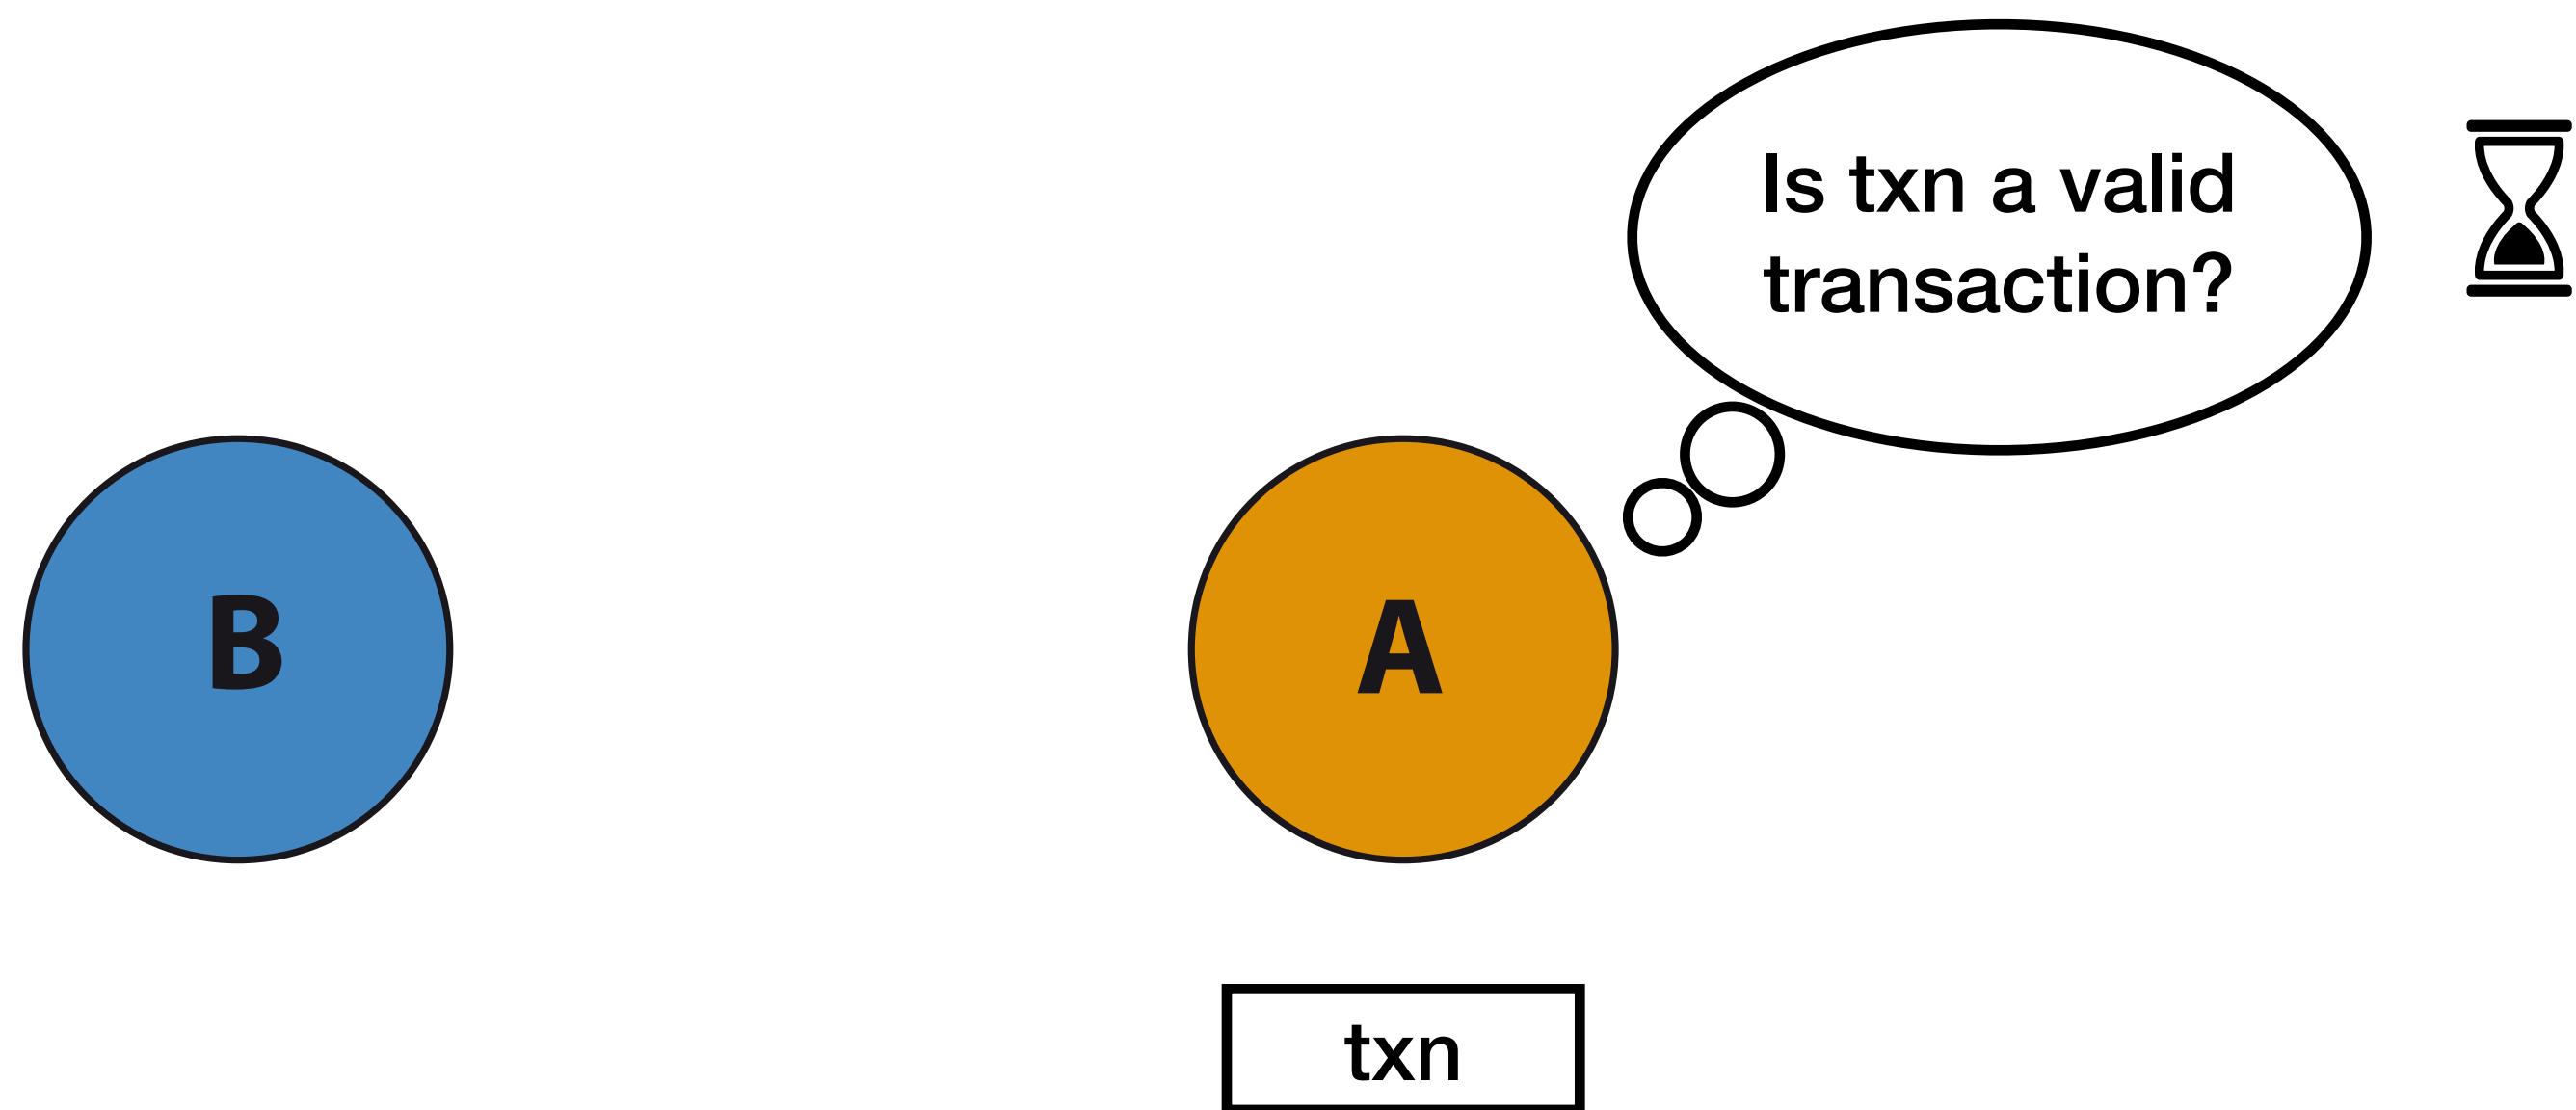
- Known transaction will be rejected

INFORMATION PROPAGATION (1/3)



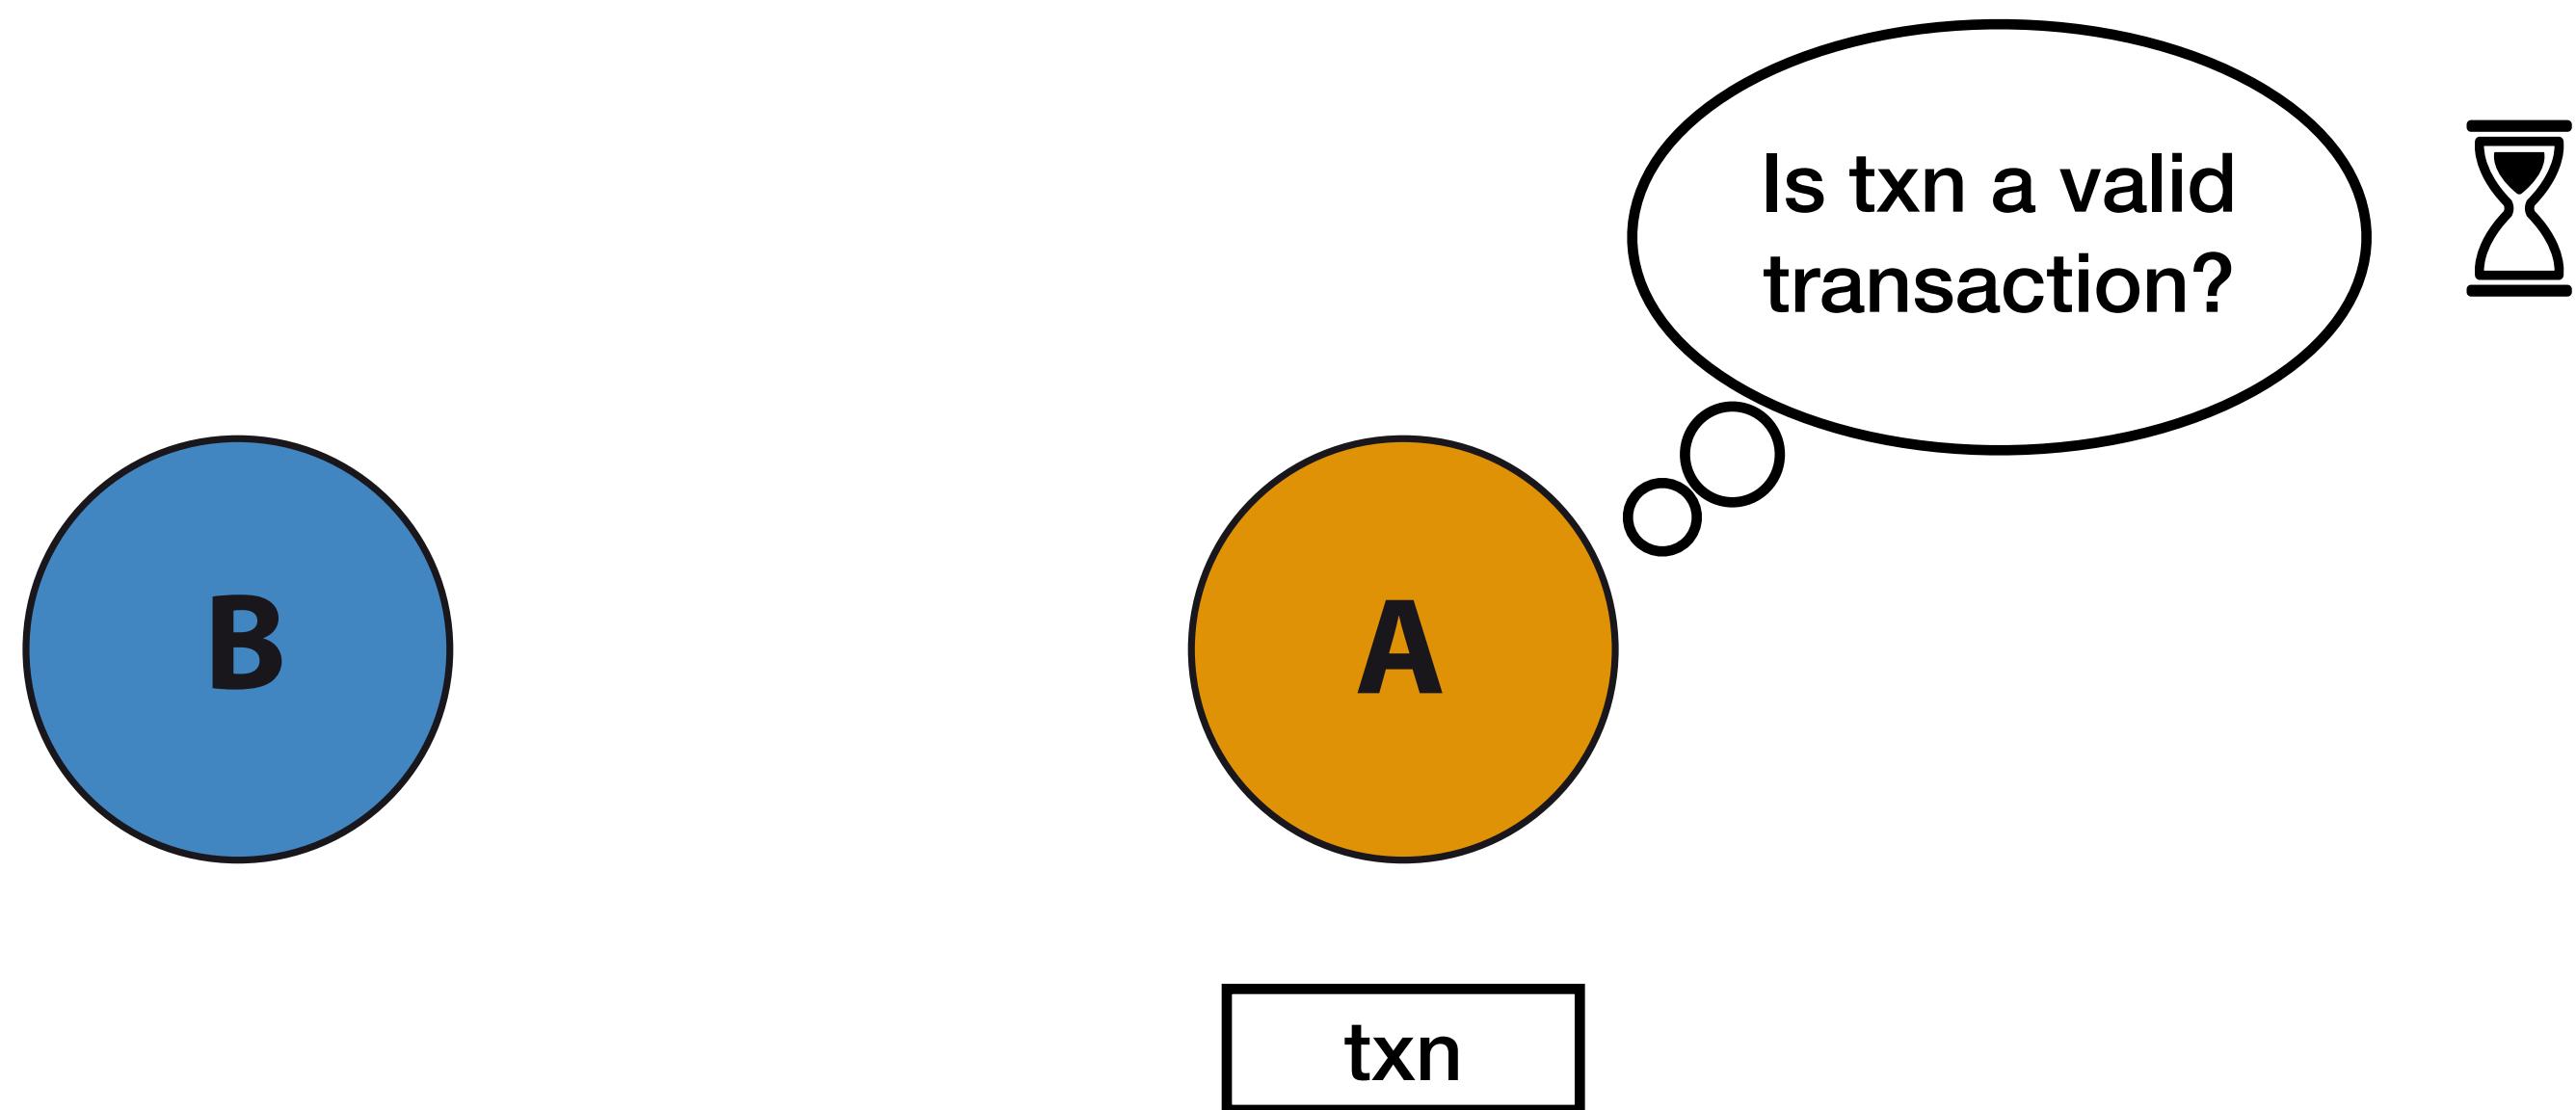
- Known transaction will be rejected

INFORMATION PROPAGATION (1/3)



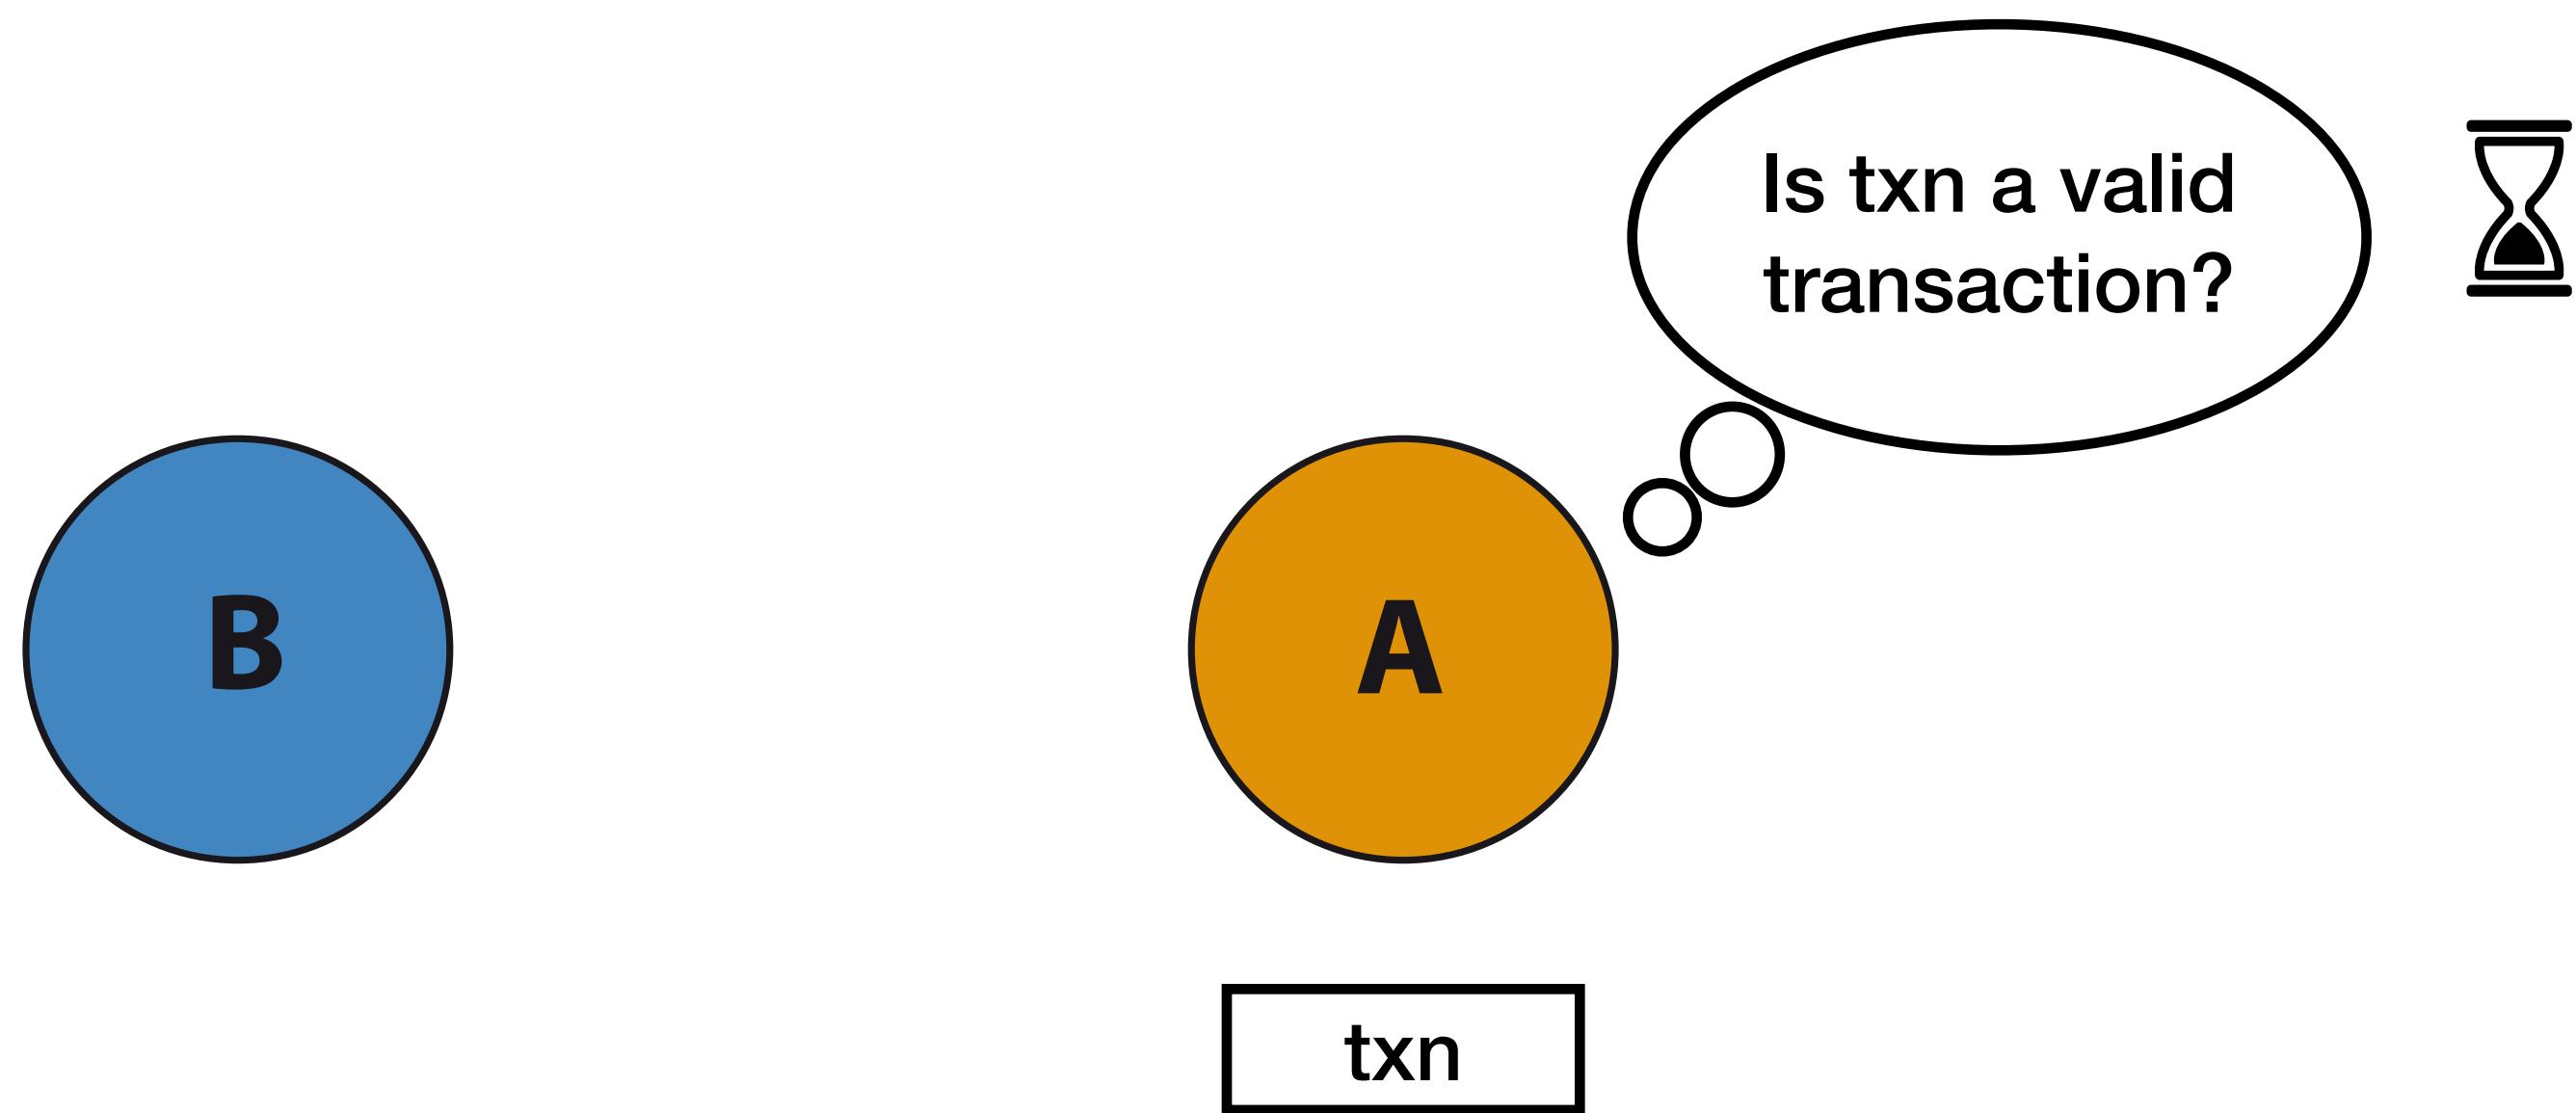
- Known transaction will be rejected

INFORMATION PROPAGATION (1/3)



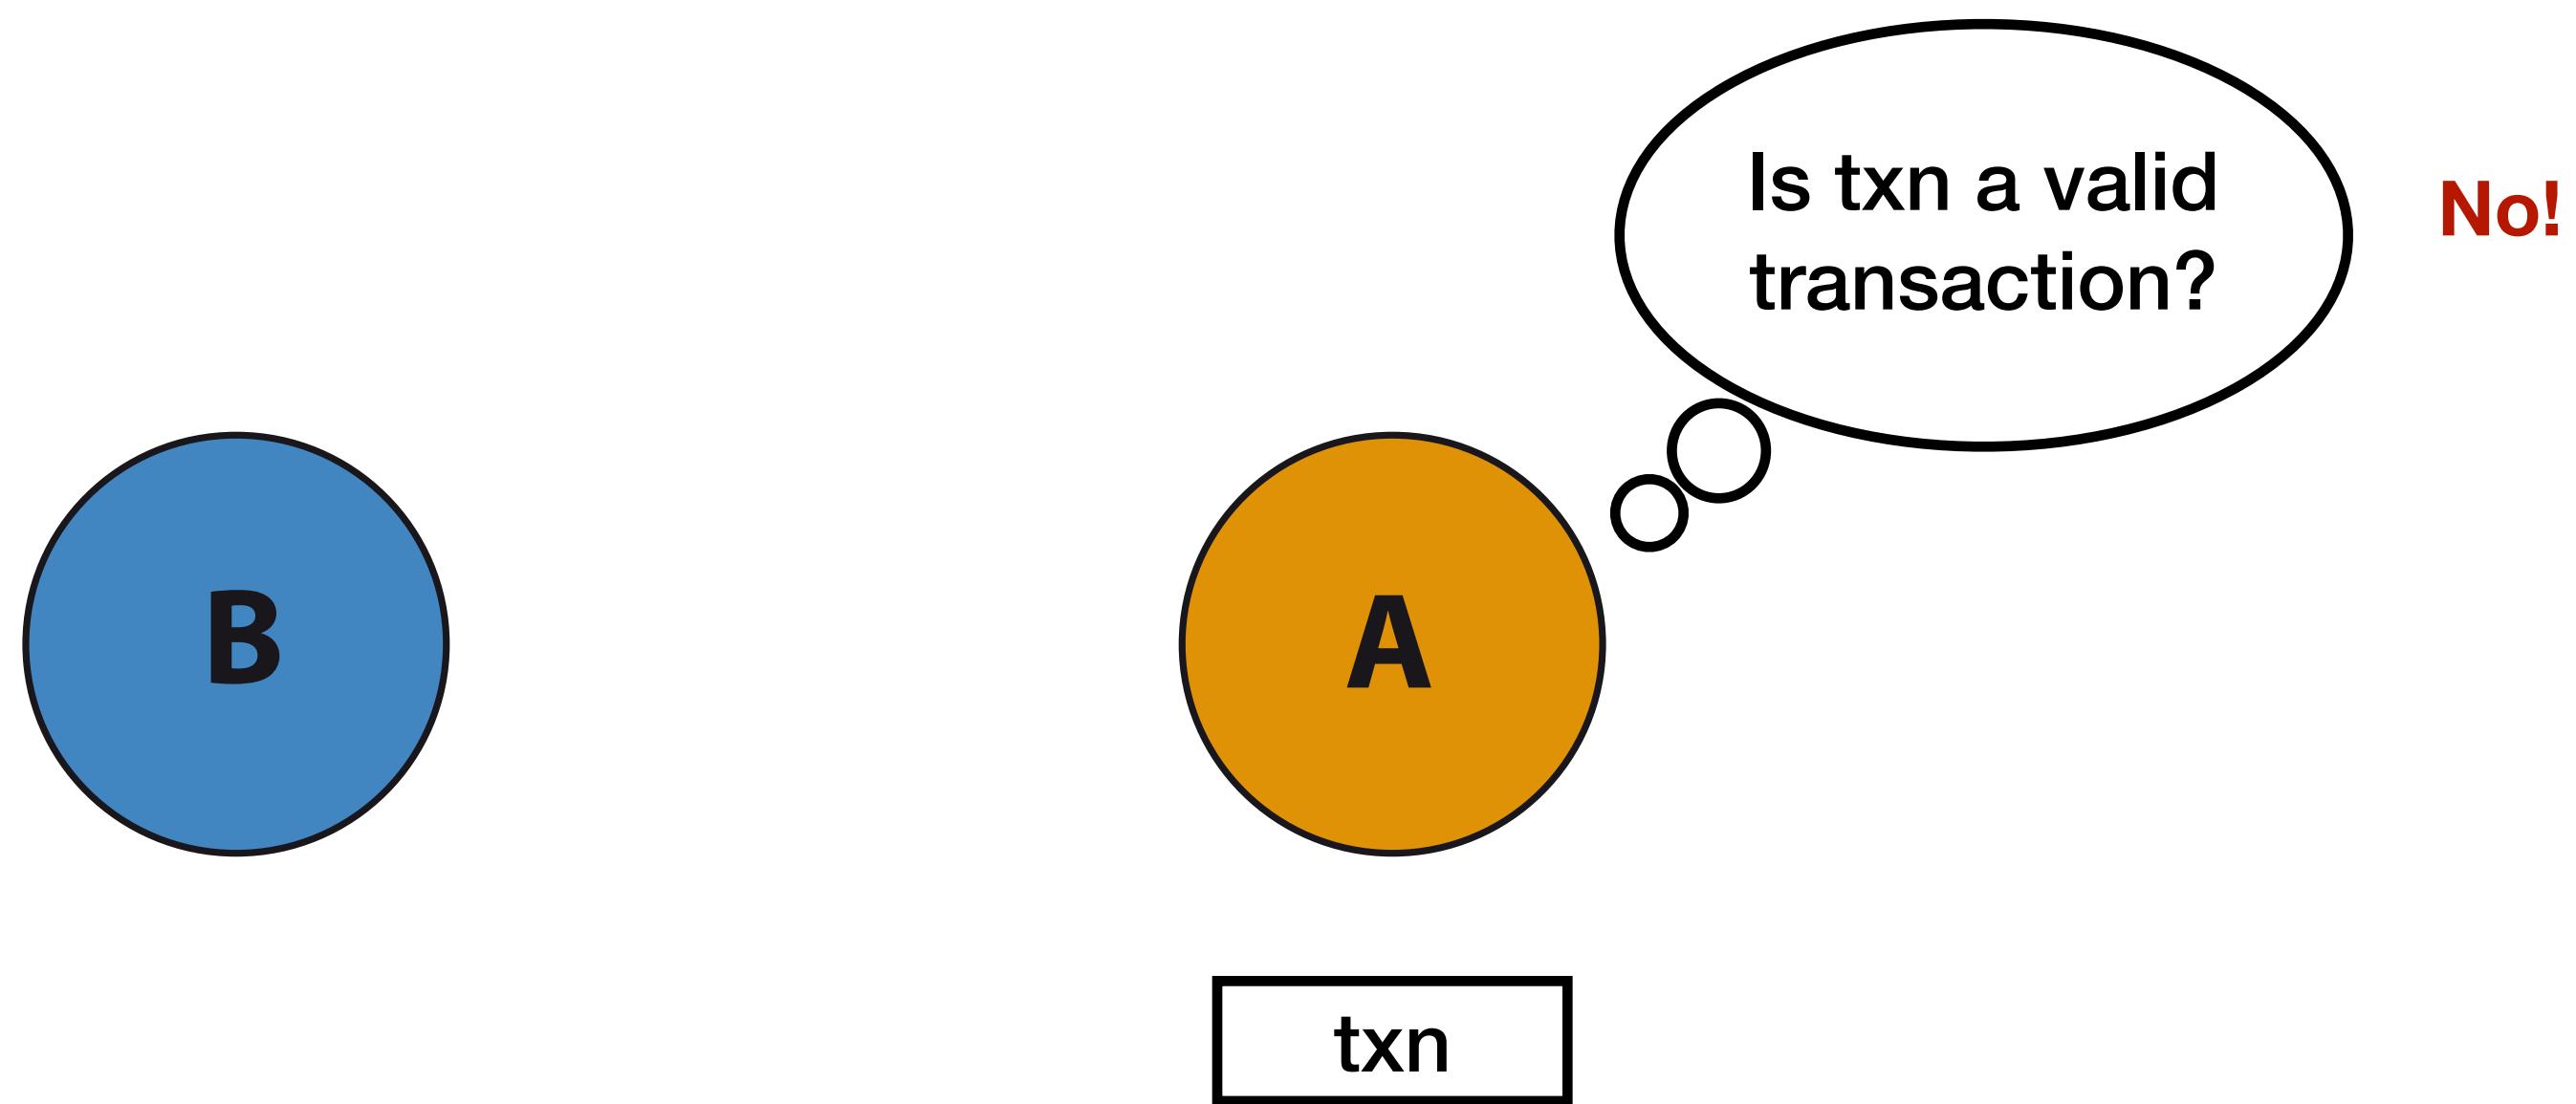
- Known transaction will be rejected

INFORMATION PROPAGATION (1/3)



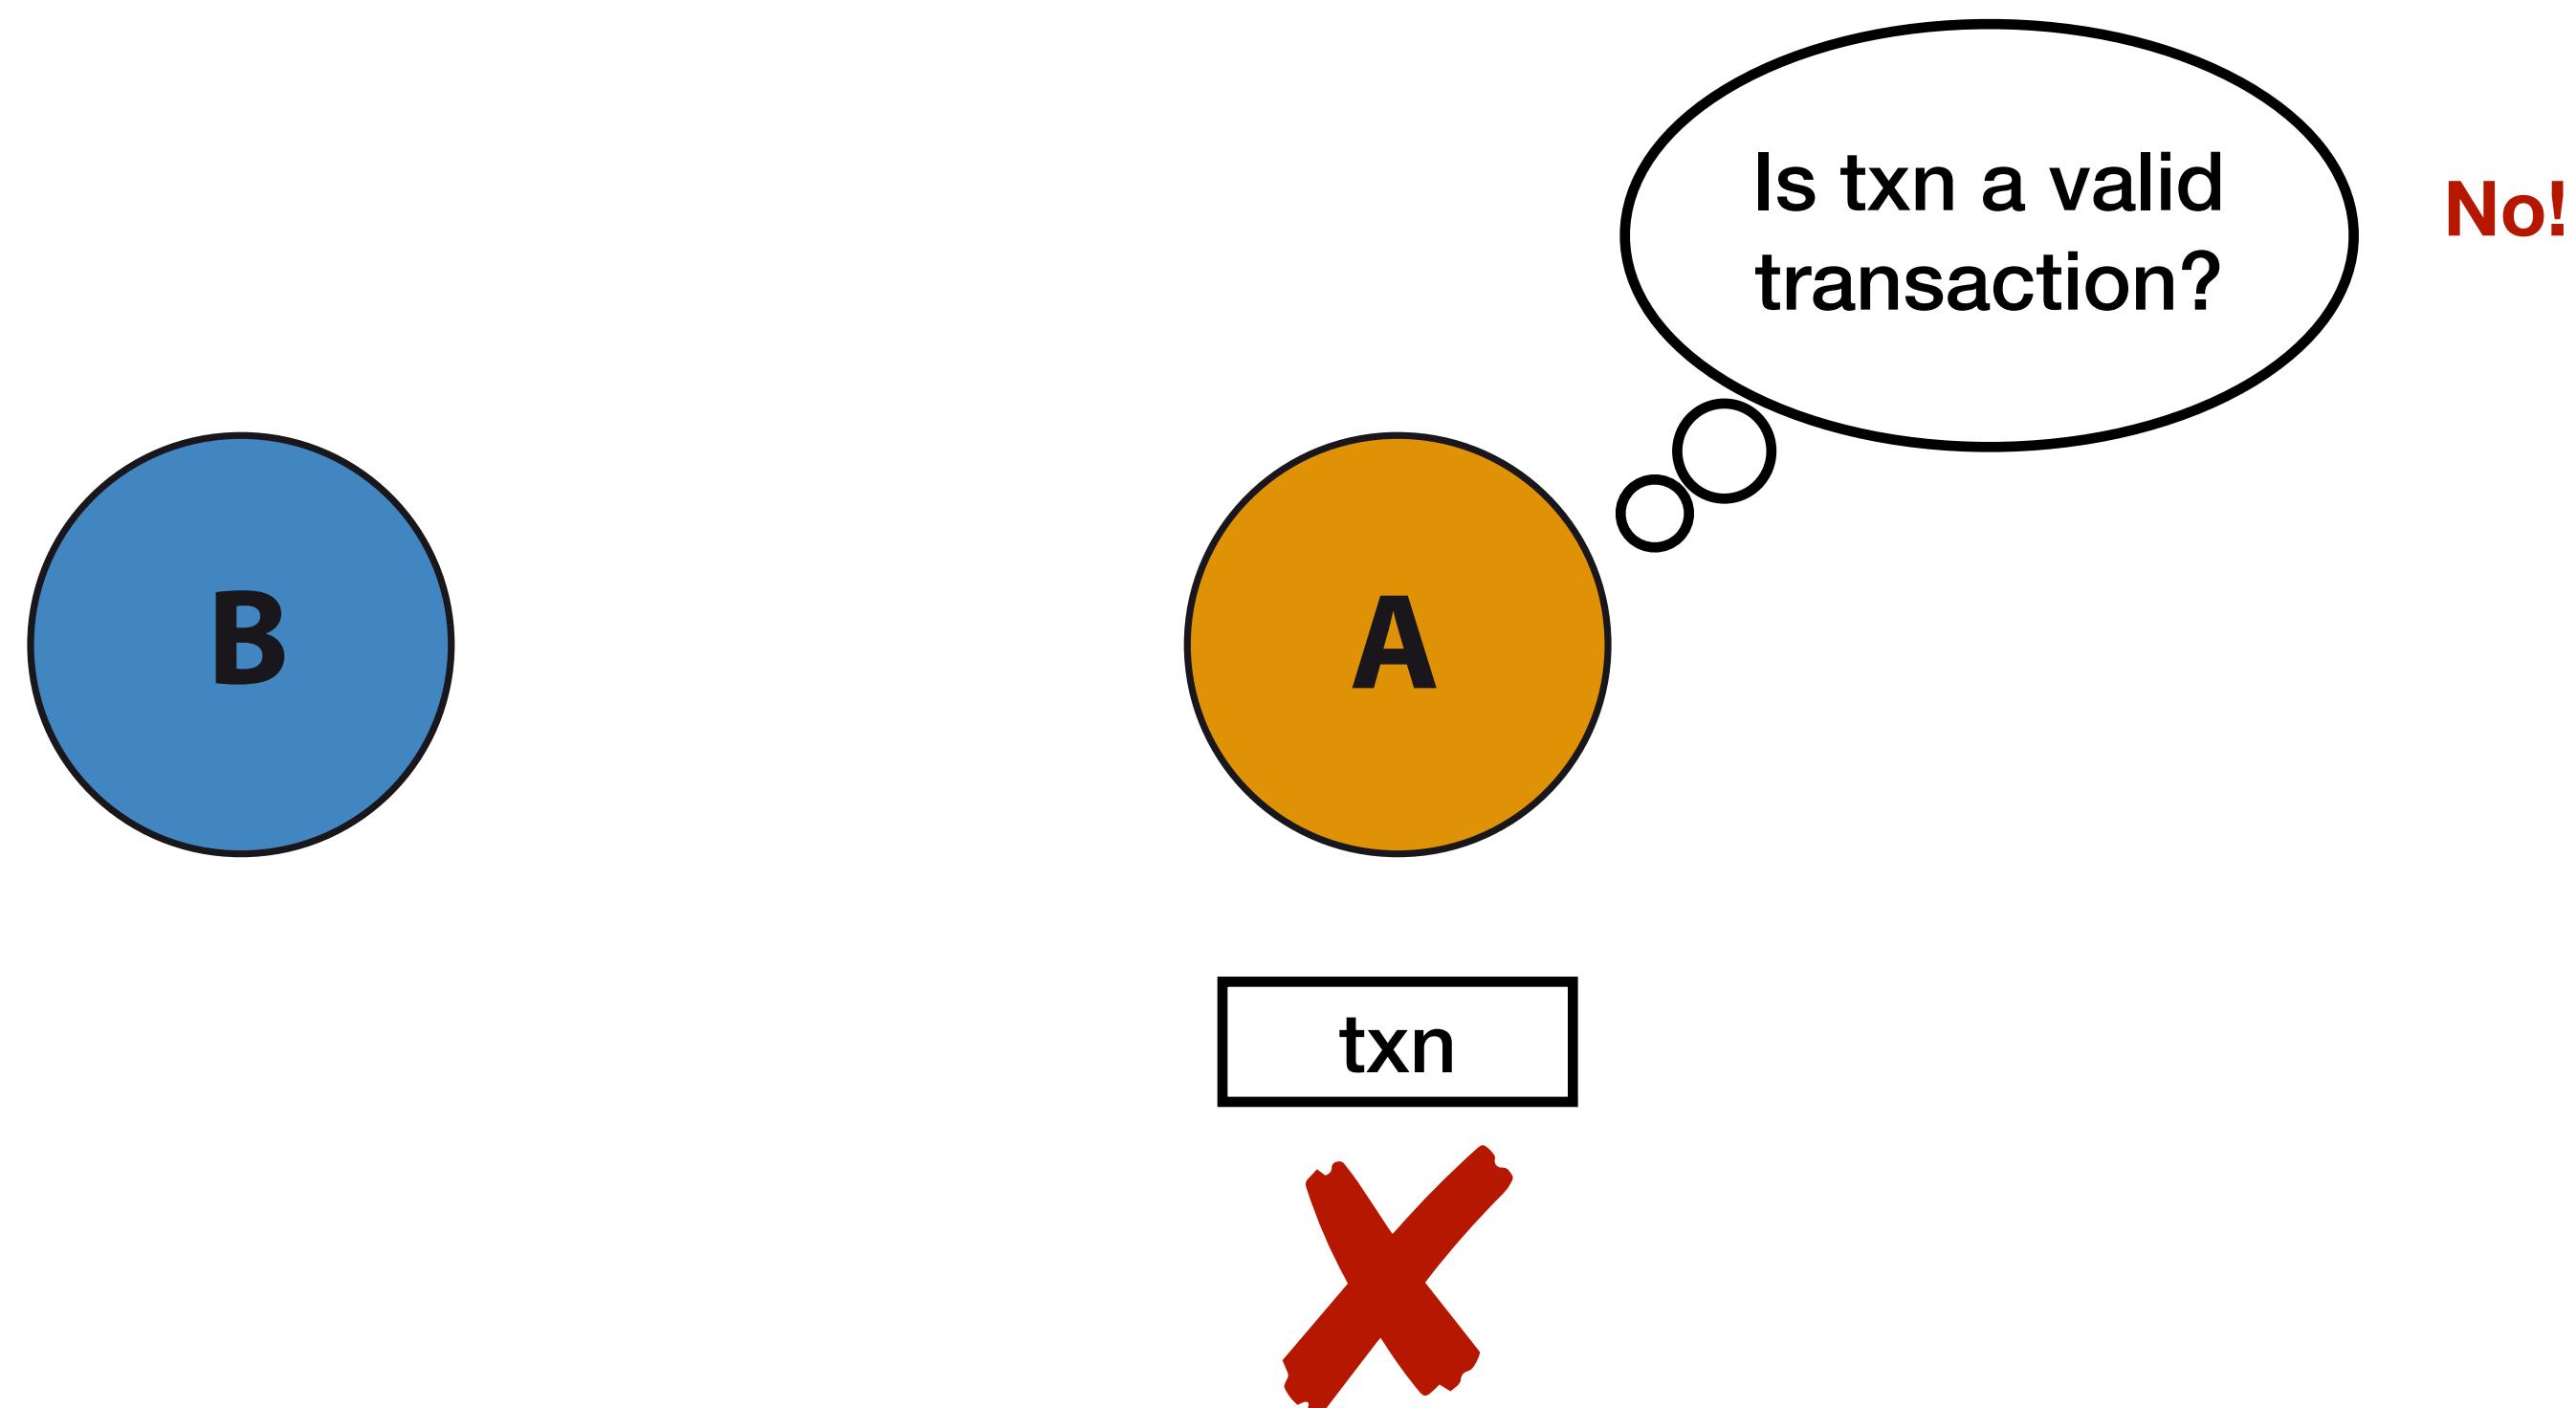
- Known transaction will be rejected

INFORMATION PROPAGATION (1/3)



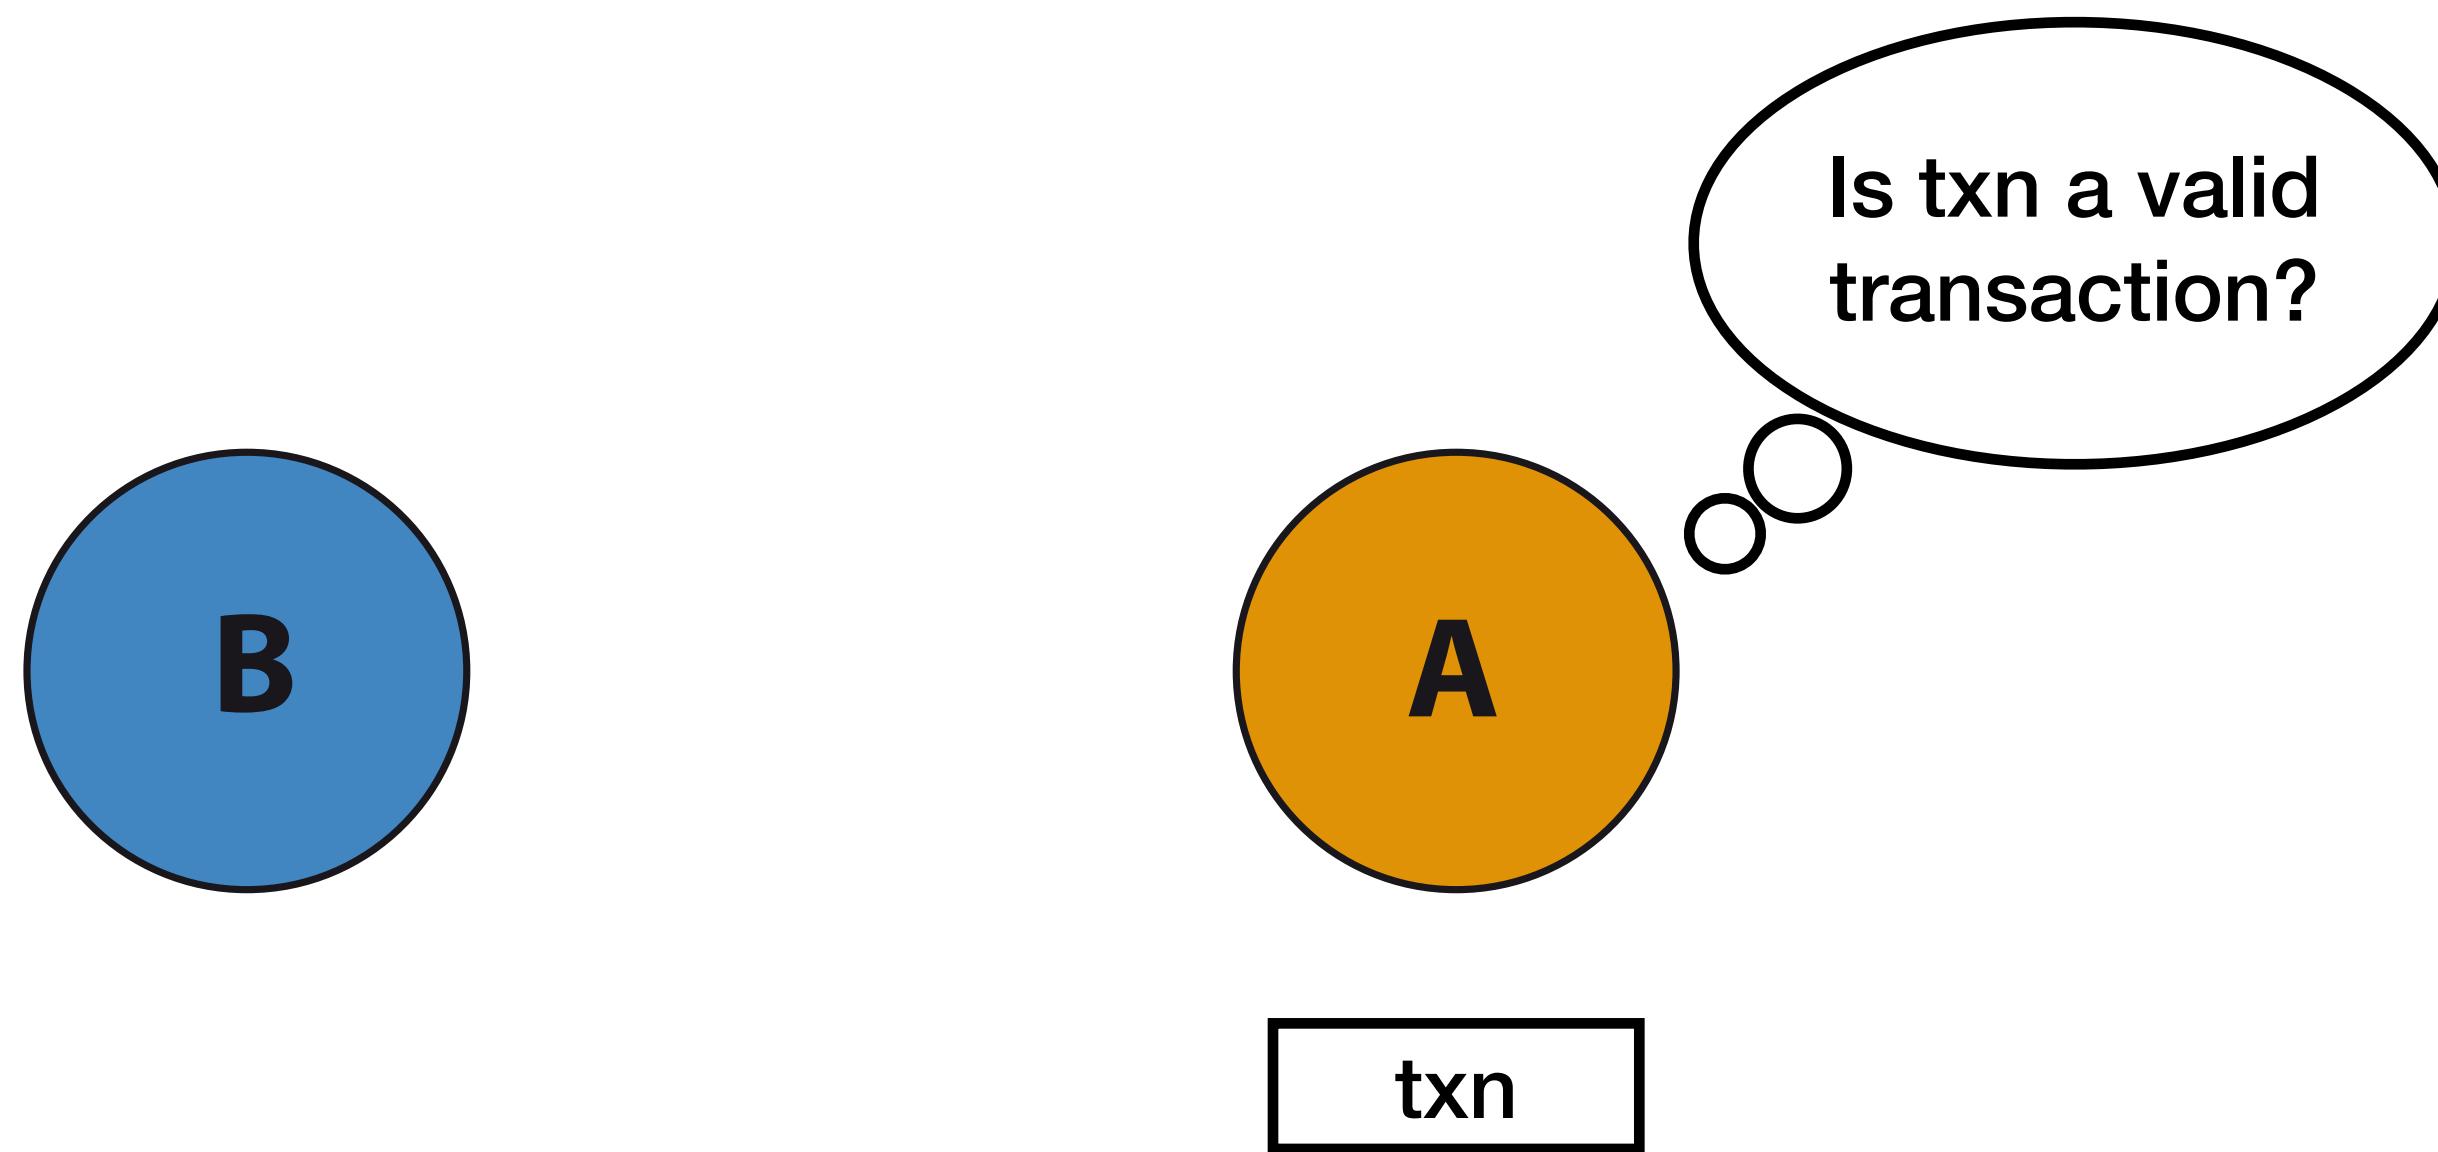
- Known transaction will be rejected

INFORMATION PROPAGATION (1/3)



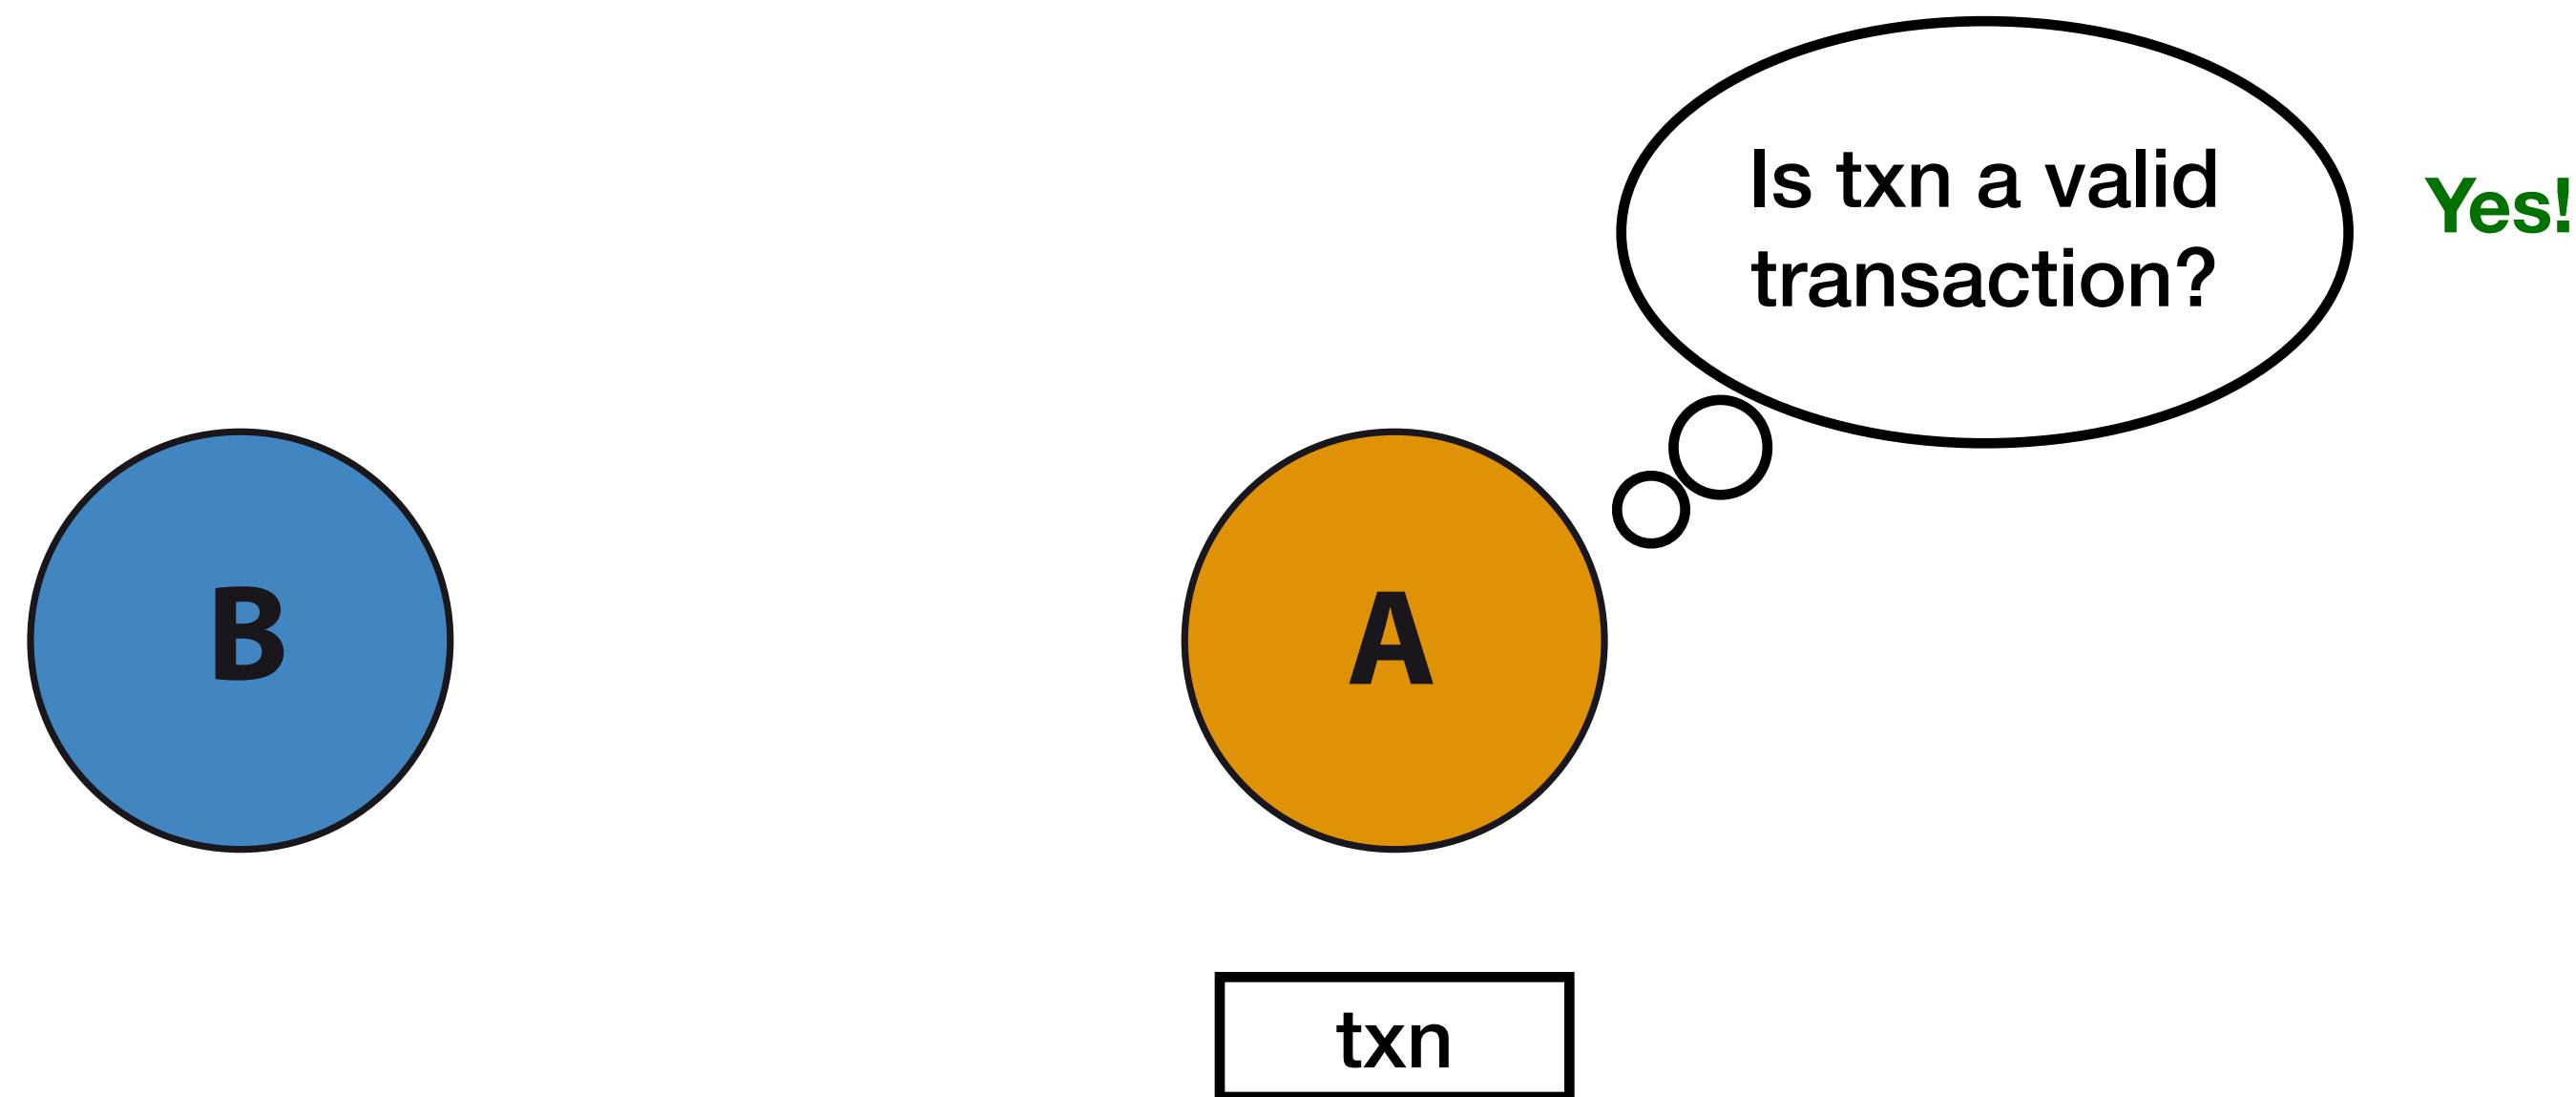
- Known transaction will be rejected
- Invalid transaction will also be rejected

INFORMATION PROPAGATION (1/3)



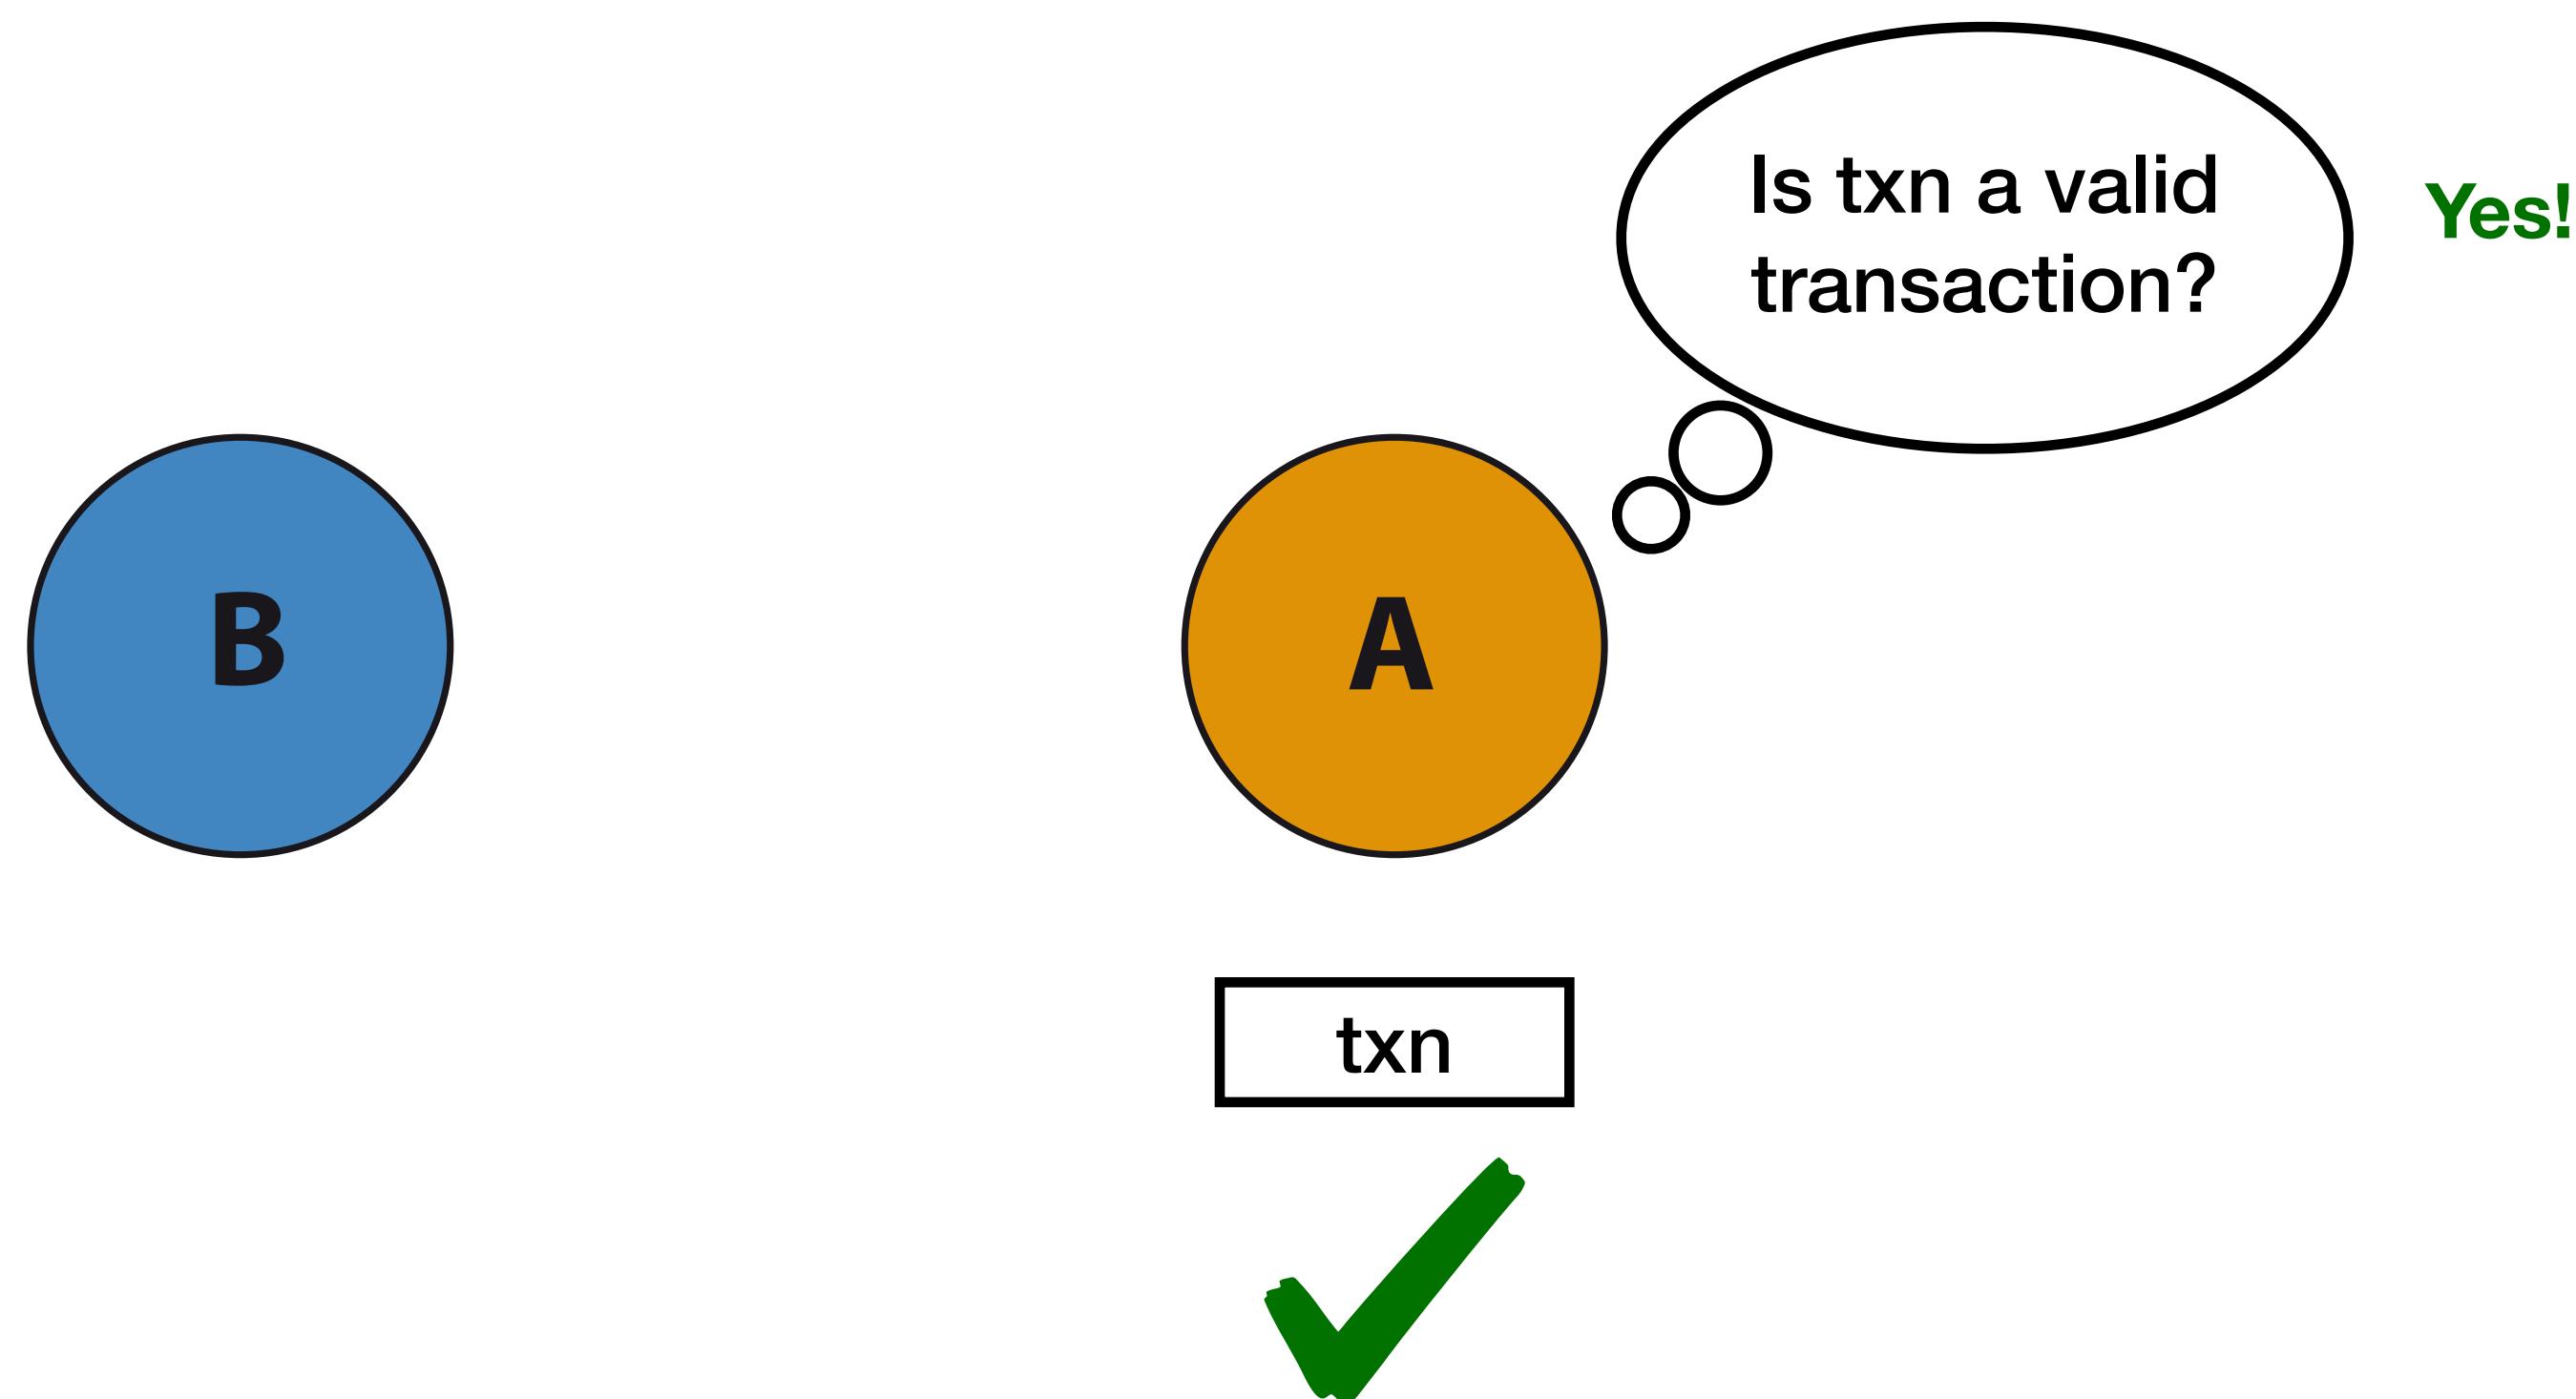
- Known transaction will be rejected
- Invalid transaction will also be rejected

INFORMATION PROPAGATION (1/3)



- Known transaction will be rejected
- Invalid transaction will also be rejected

INFORMATION PROPAGATION (1/3)



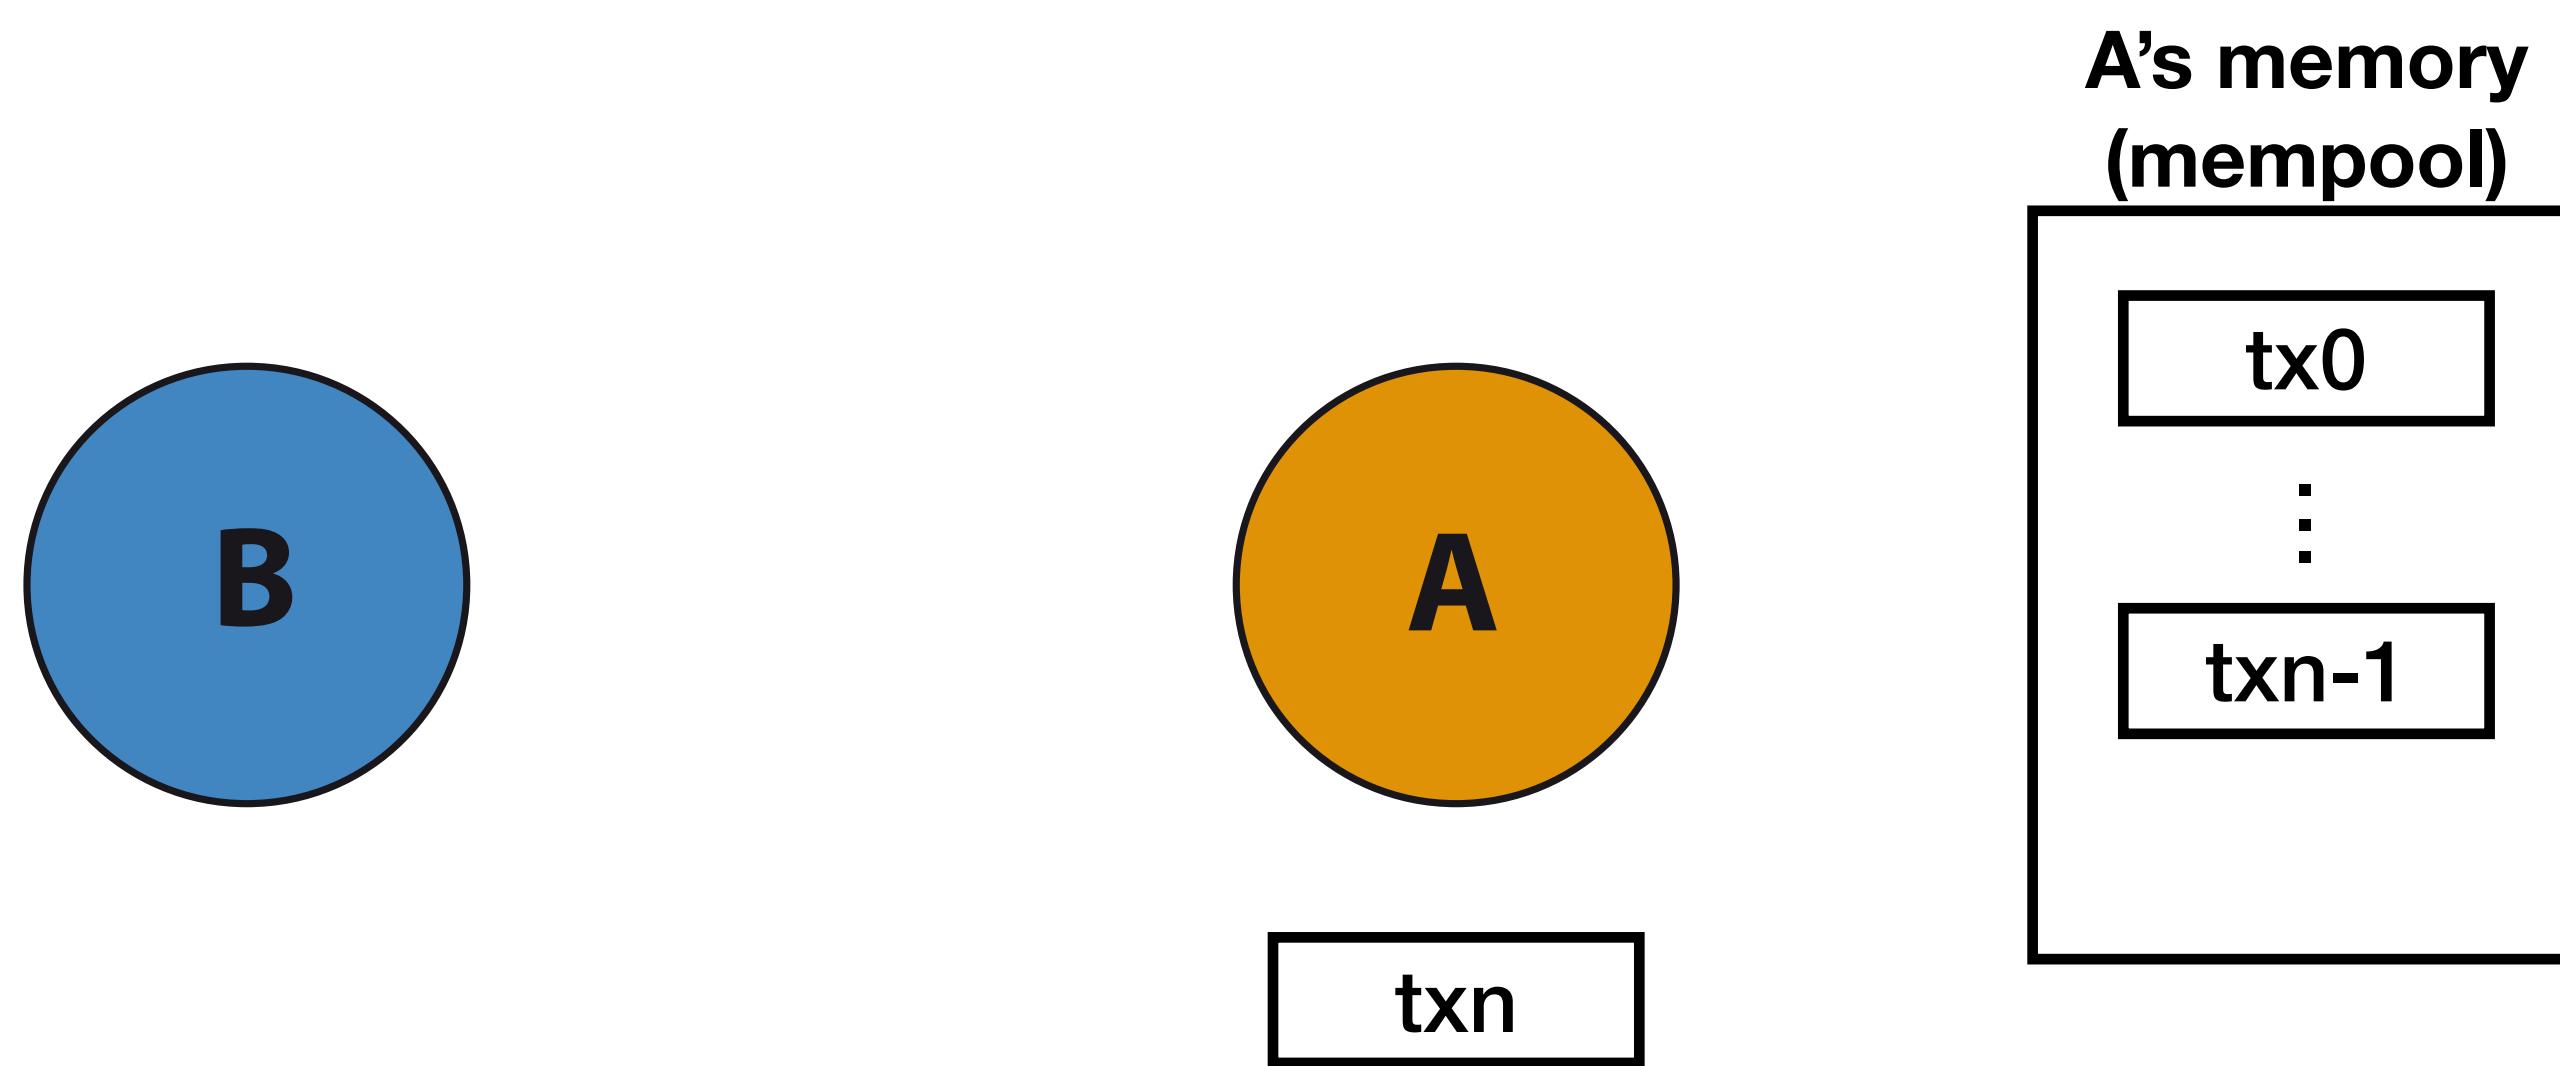
- Known transaction will be rejected
- Invalid transaction will also be rejected

INFORMATION PROPAGATION (1/3)



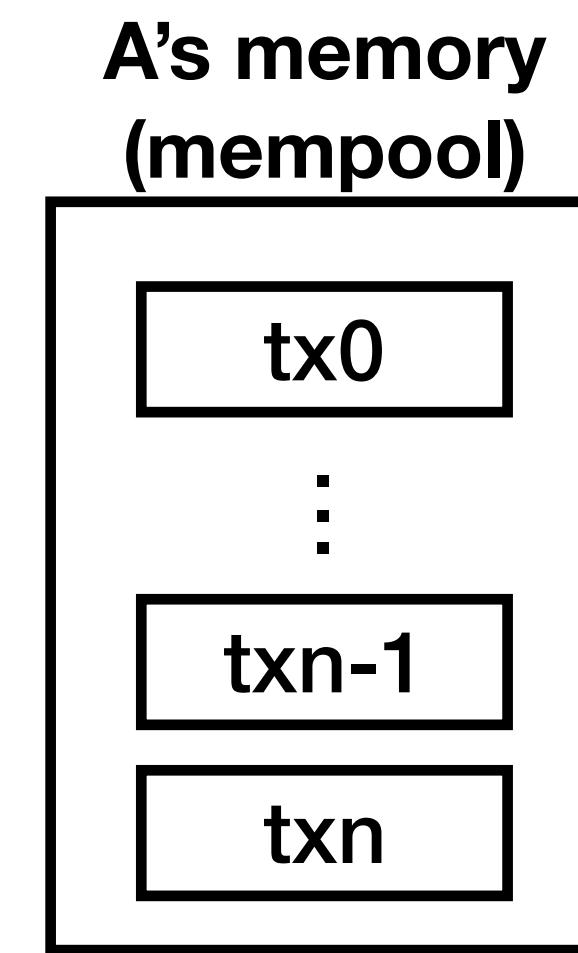
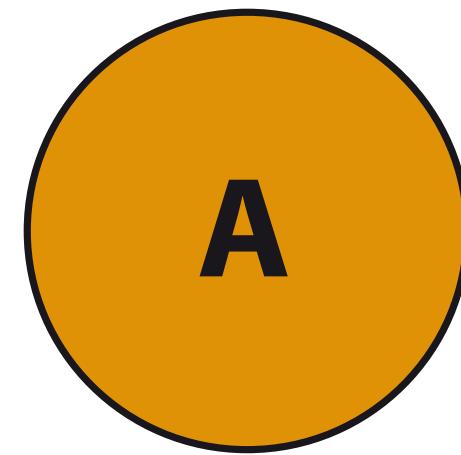
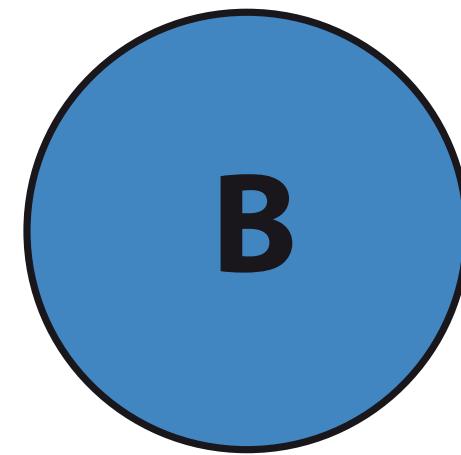
- Known transaction will be rejected
- Invalid transaction will also be rejected

INFORMATION PROPAGATION (1/3)



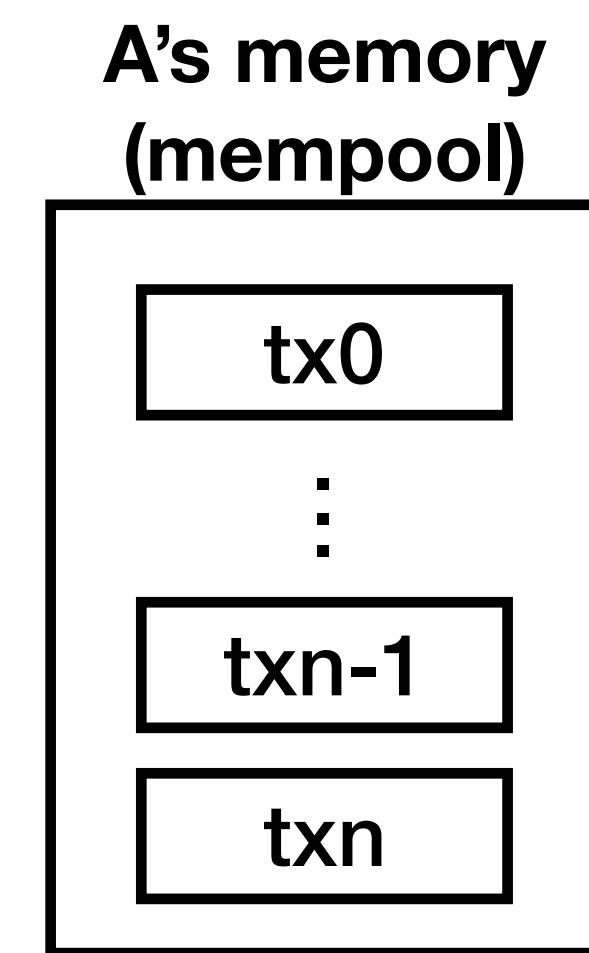
- Known transaction will be rejected
- Invalid transaction will also be rejected

INFORMATION PROPAGATION (1/3)



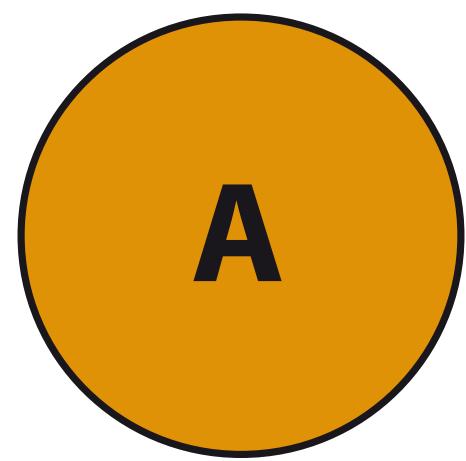
- Known transaction will be rejected
- Invalid transaction will also be rejected

INFORMATION PROPAGATION (1/3)

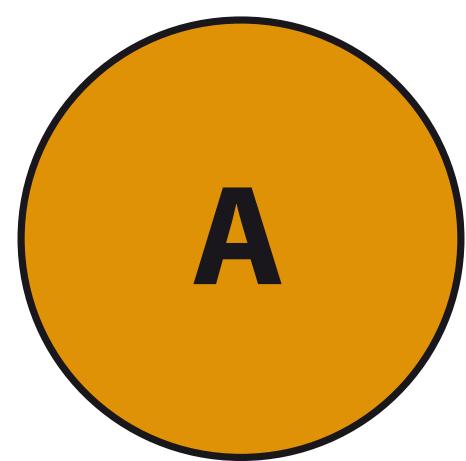
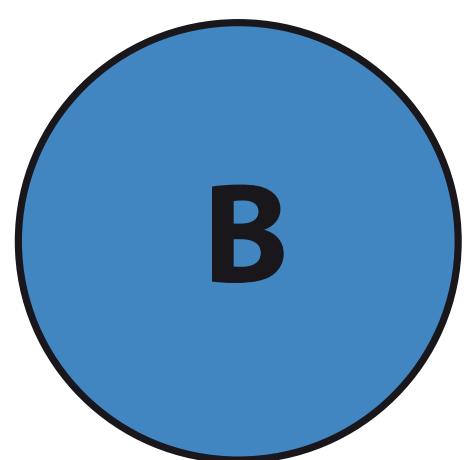


- Known transaction will be rejected
- Invalid transaction will also be rejected
- Valid (new) transactions will be kept in memory (mempool)

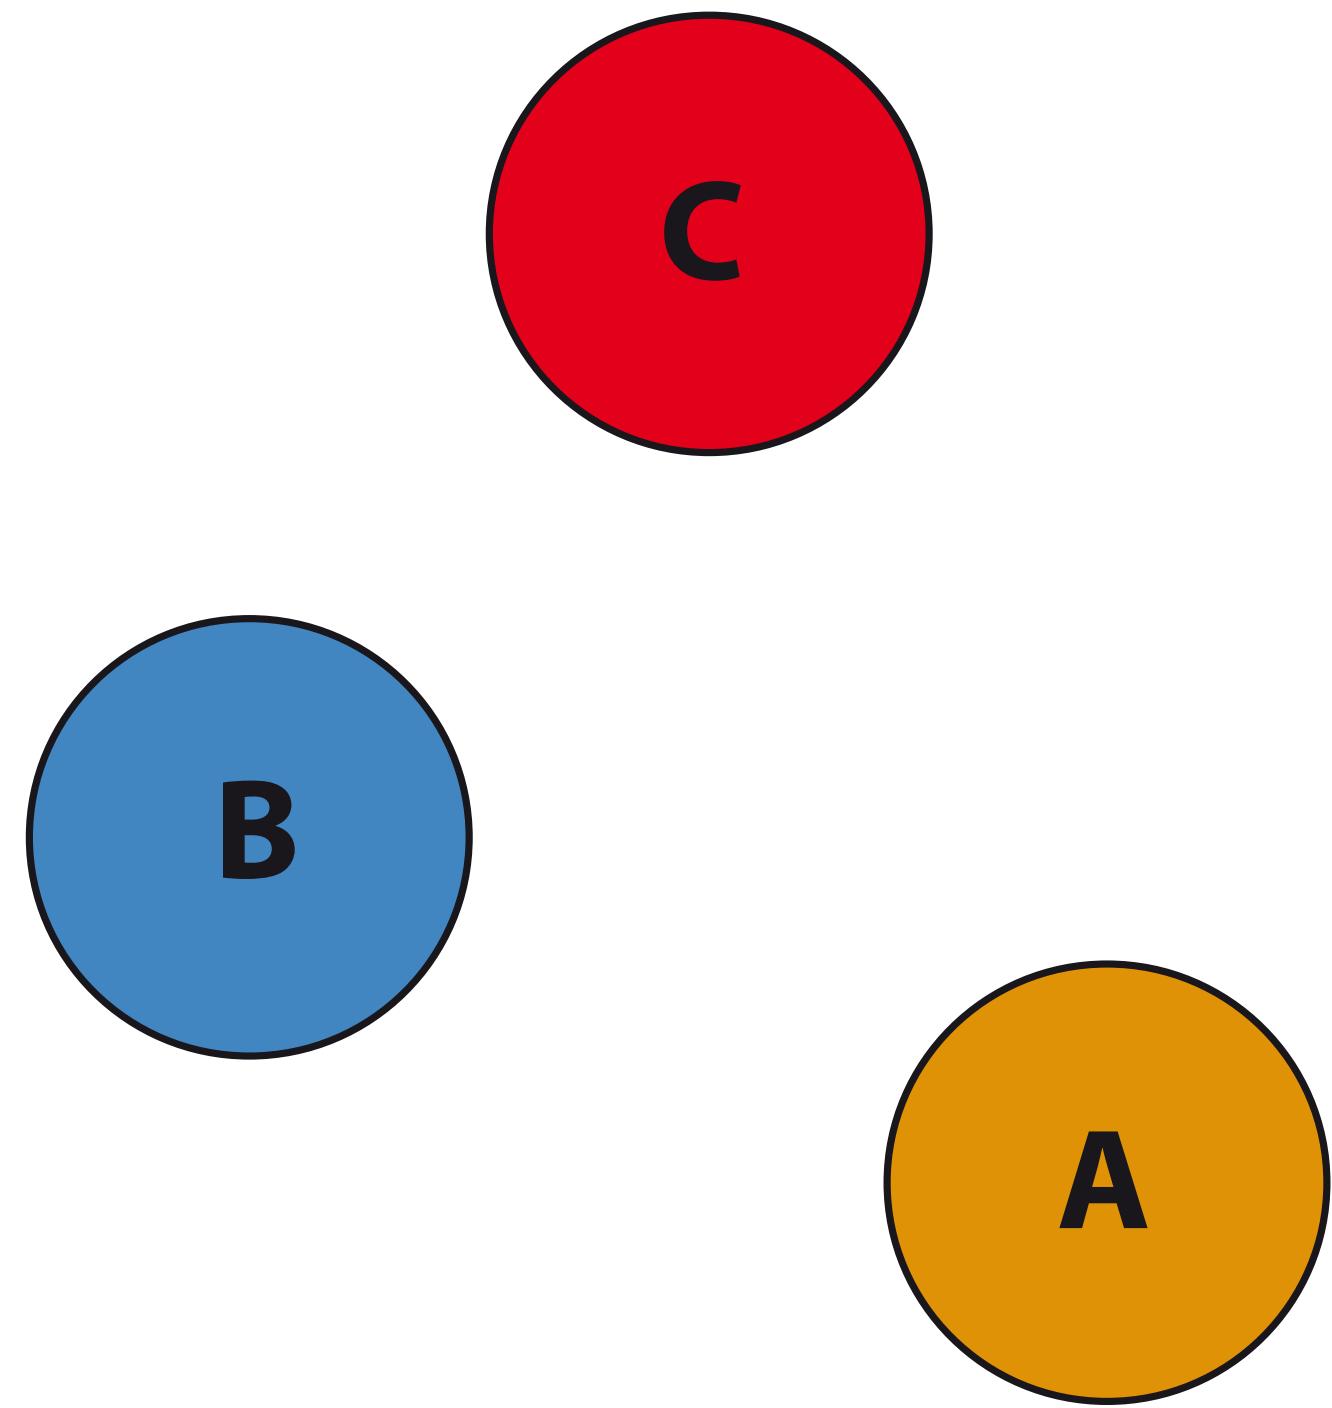
INFORMATION PROPAGATION (2/3)



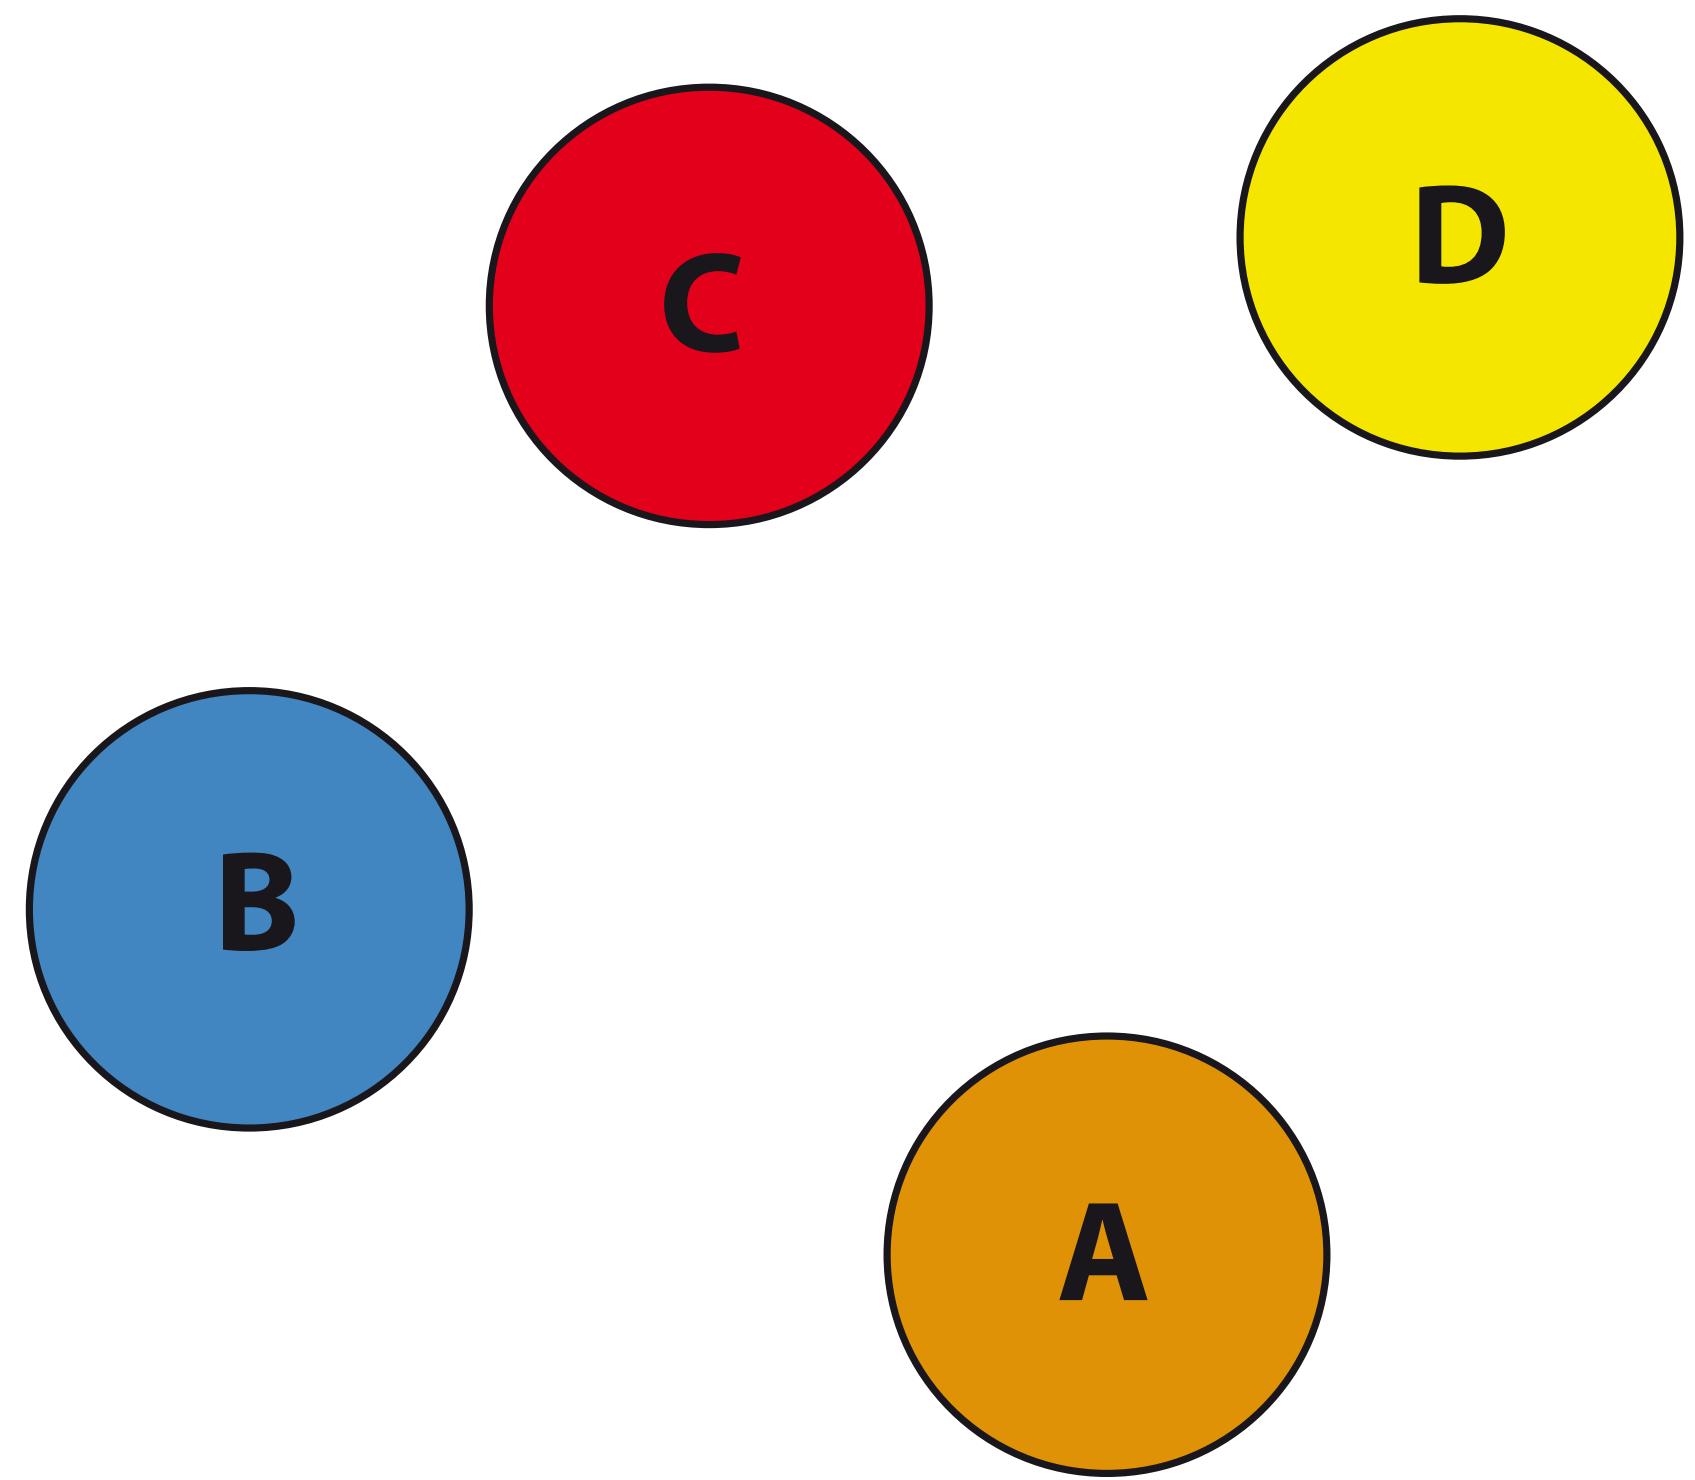
INFORMATION PROPAGATION (2/3)



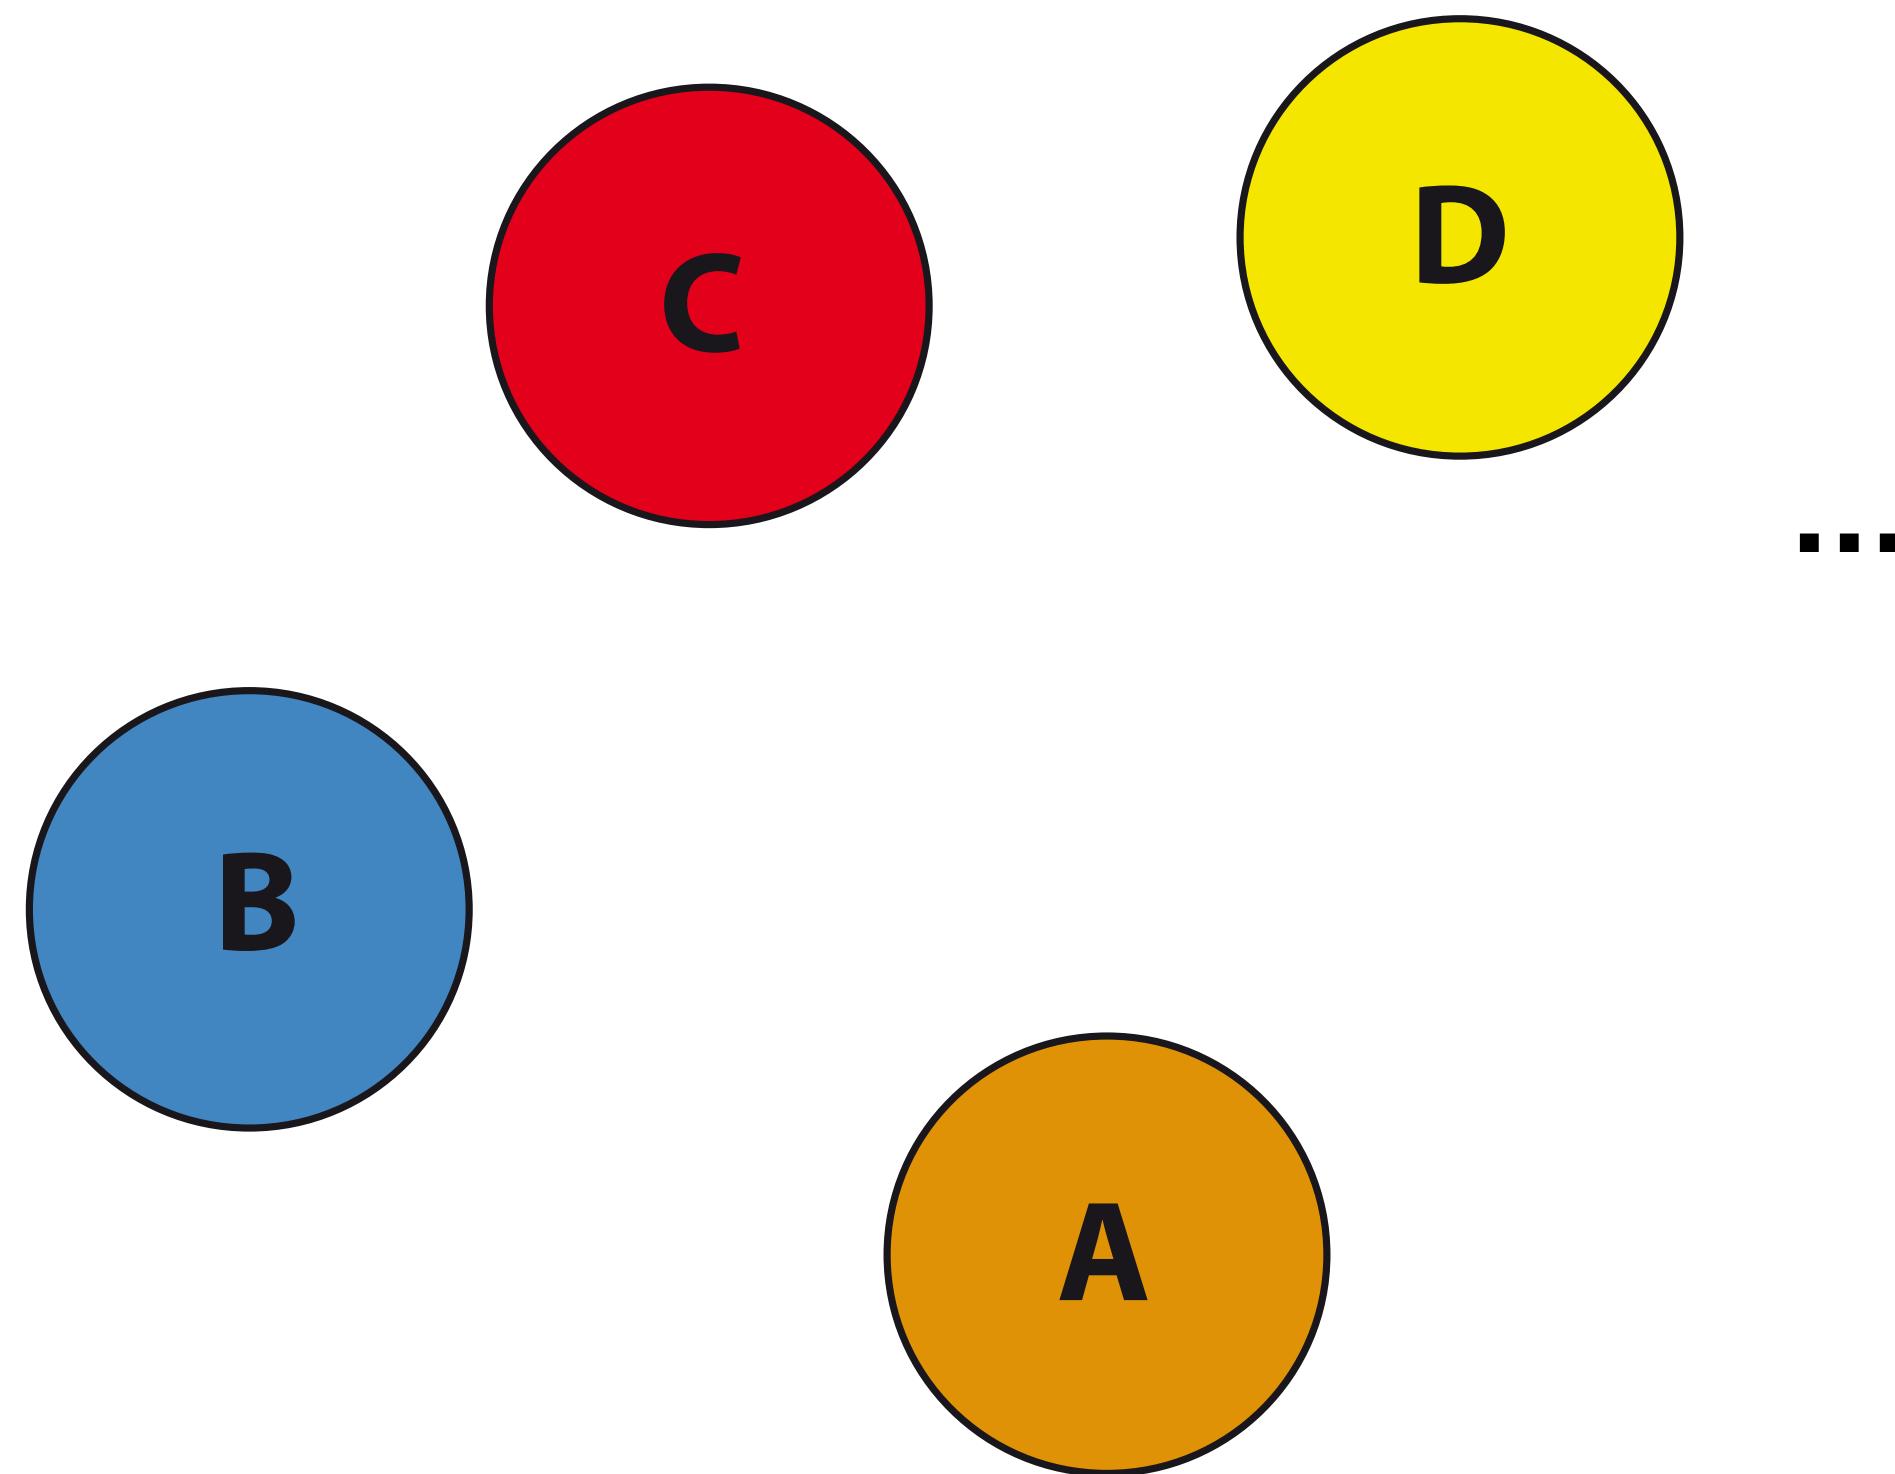
INFORMATION PROPAGATION (2/3)



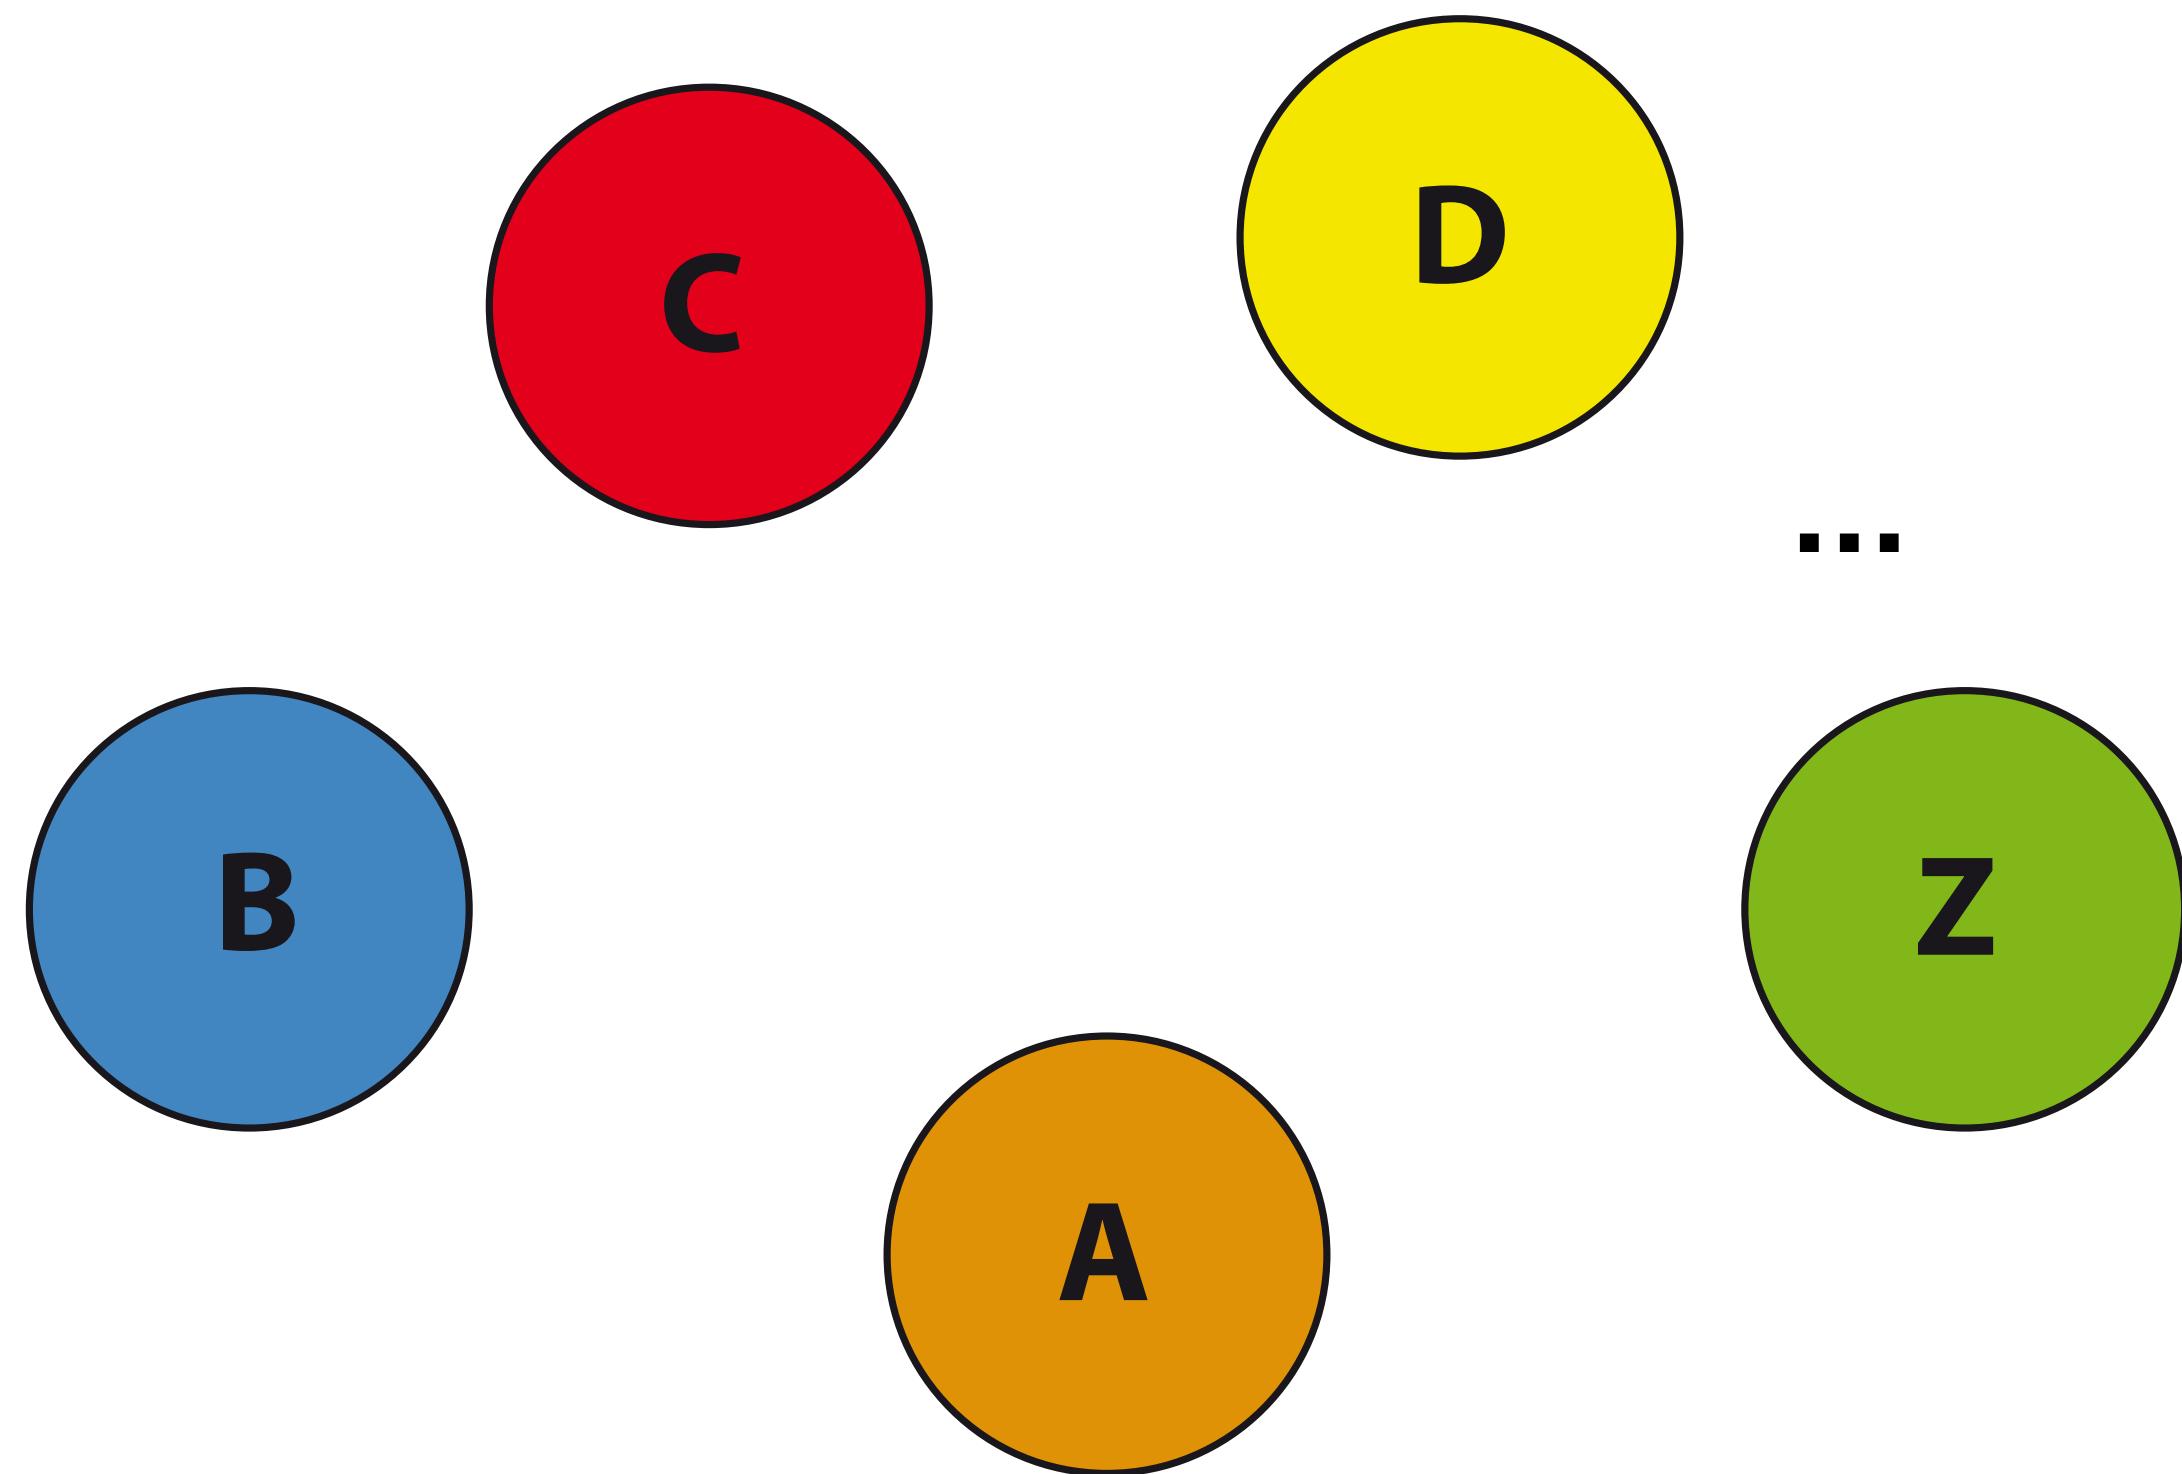
INFORMATION PROPAGATION (2/3)



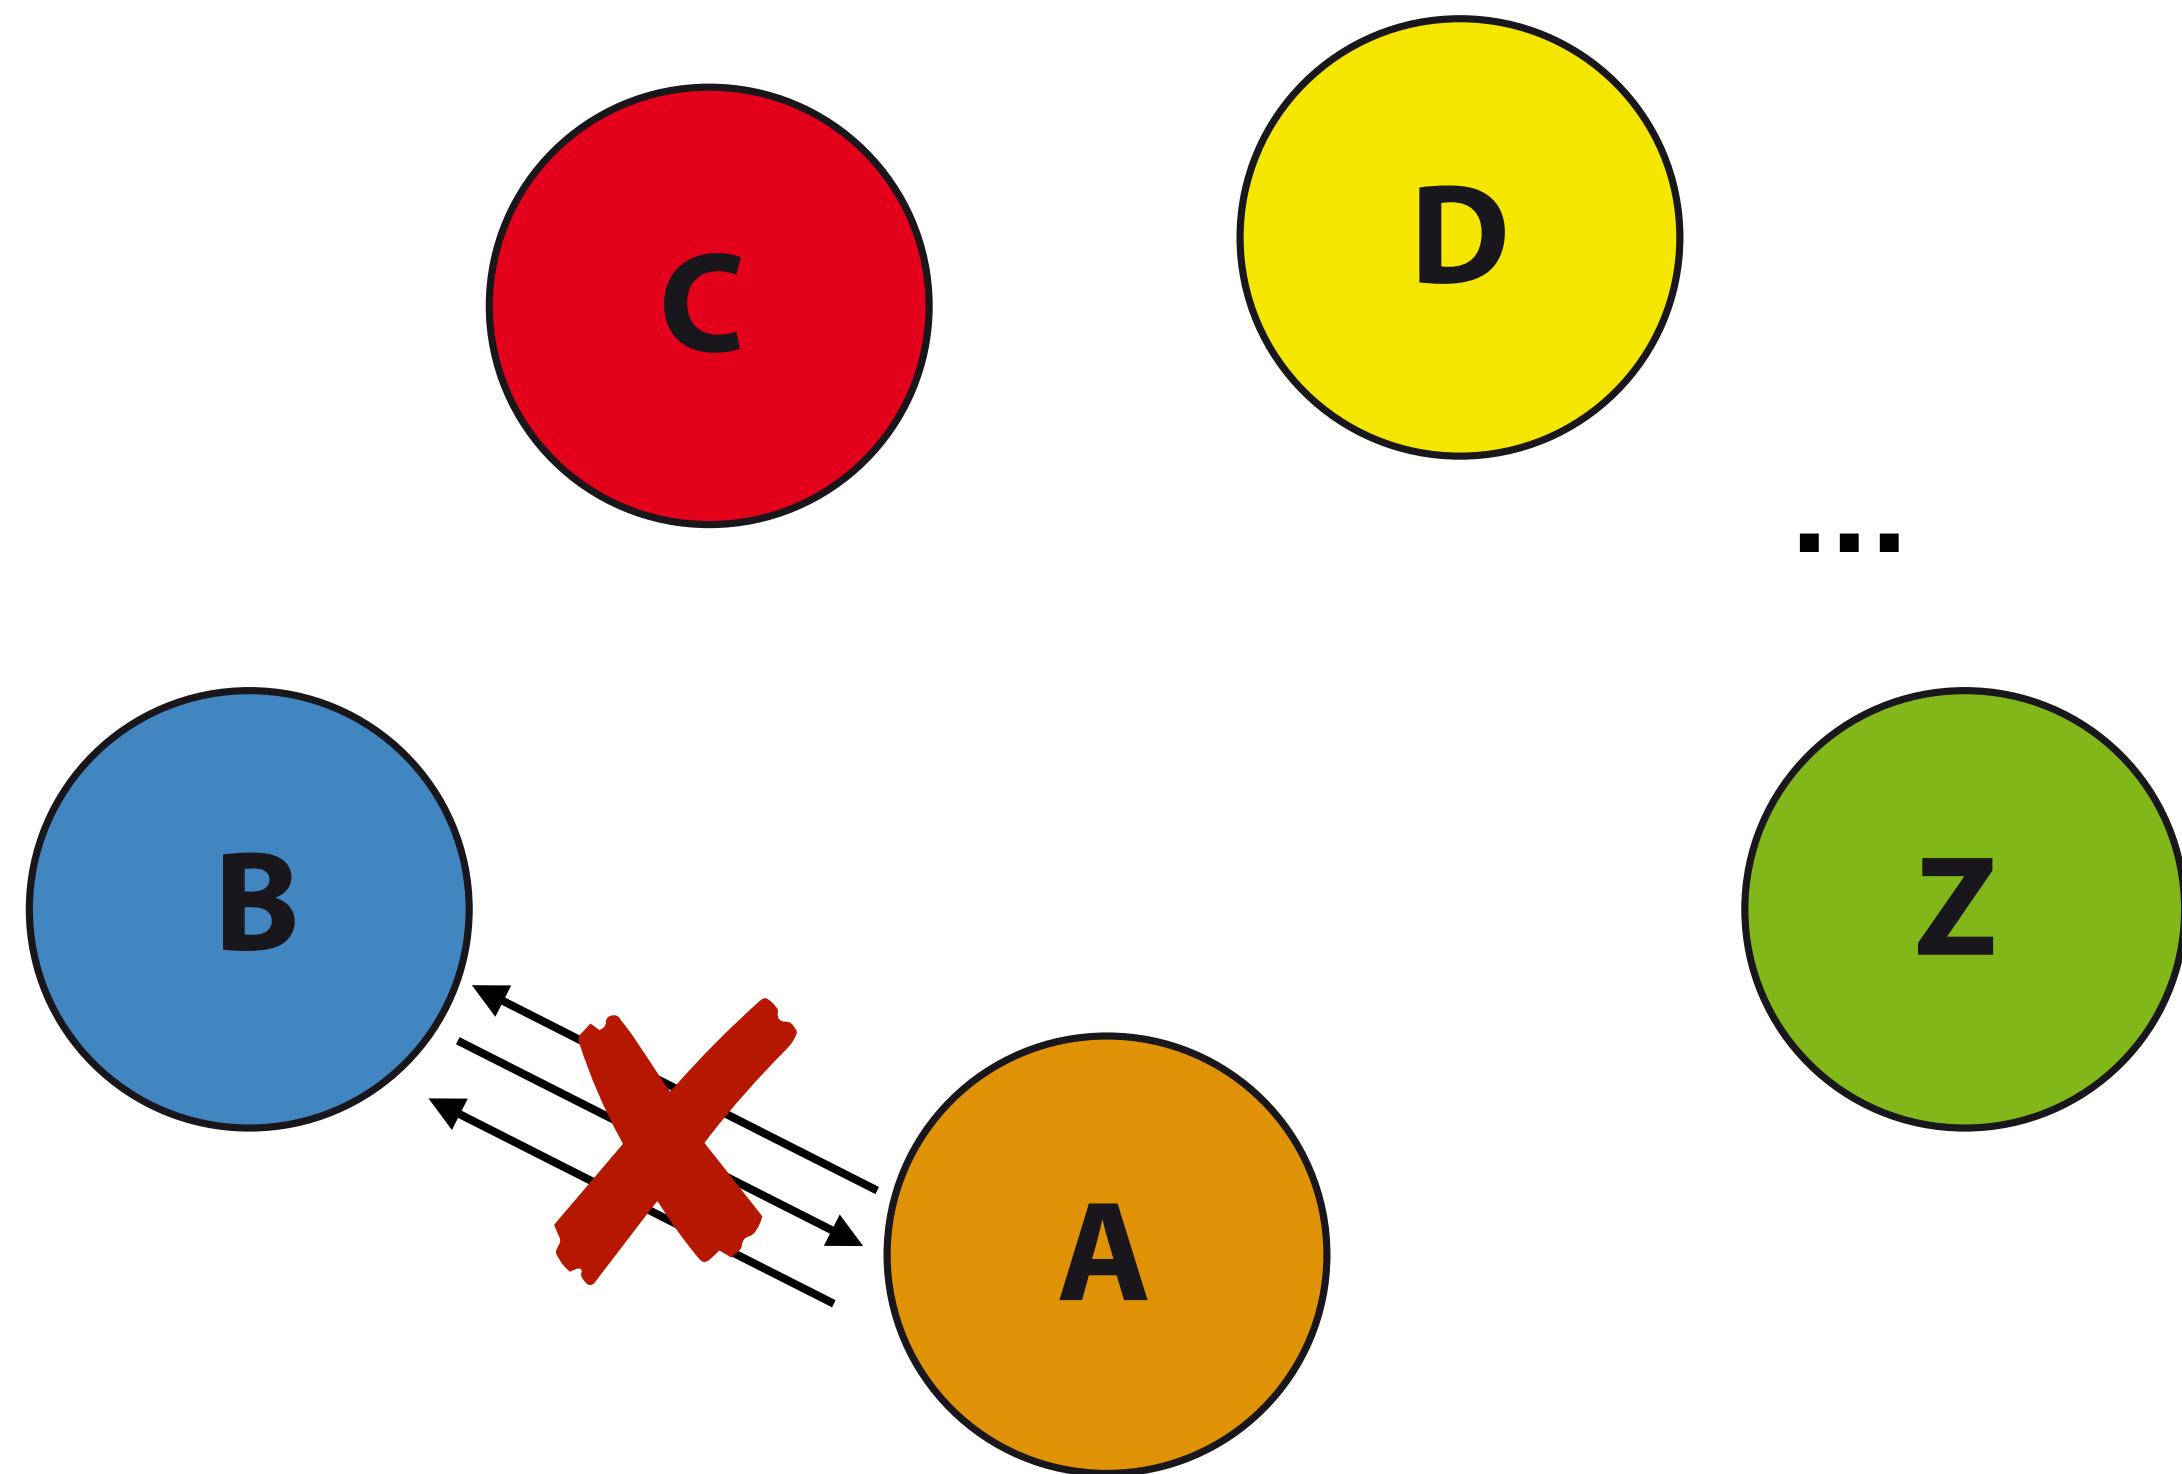
INFORMATION PROPAGATION (2/3)



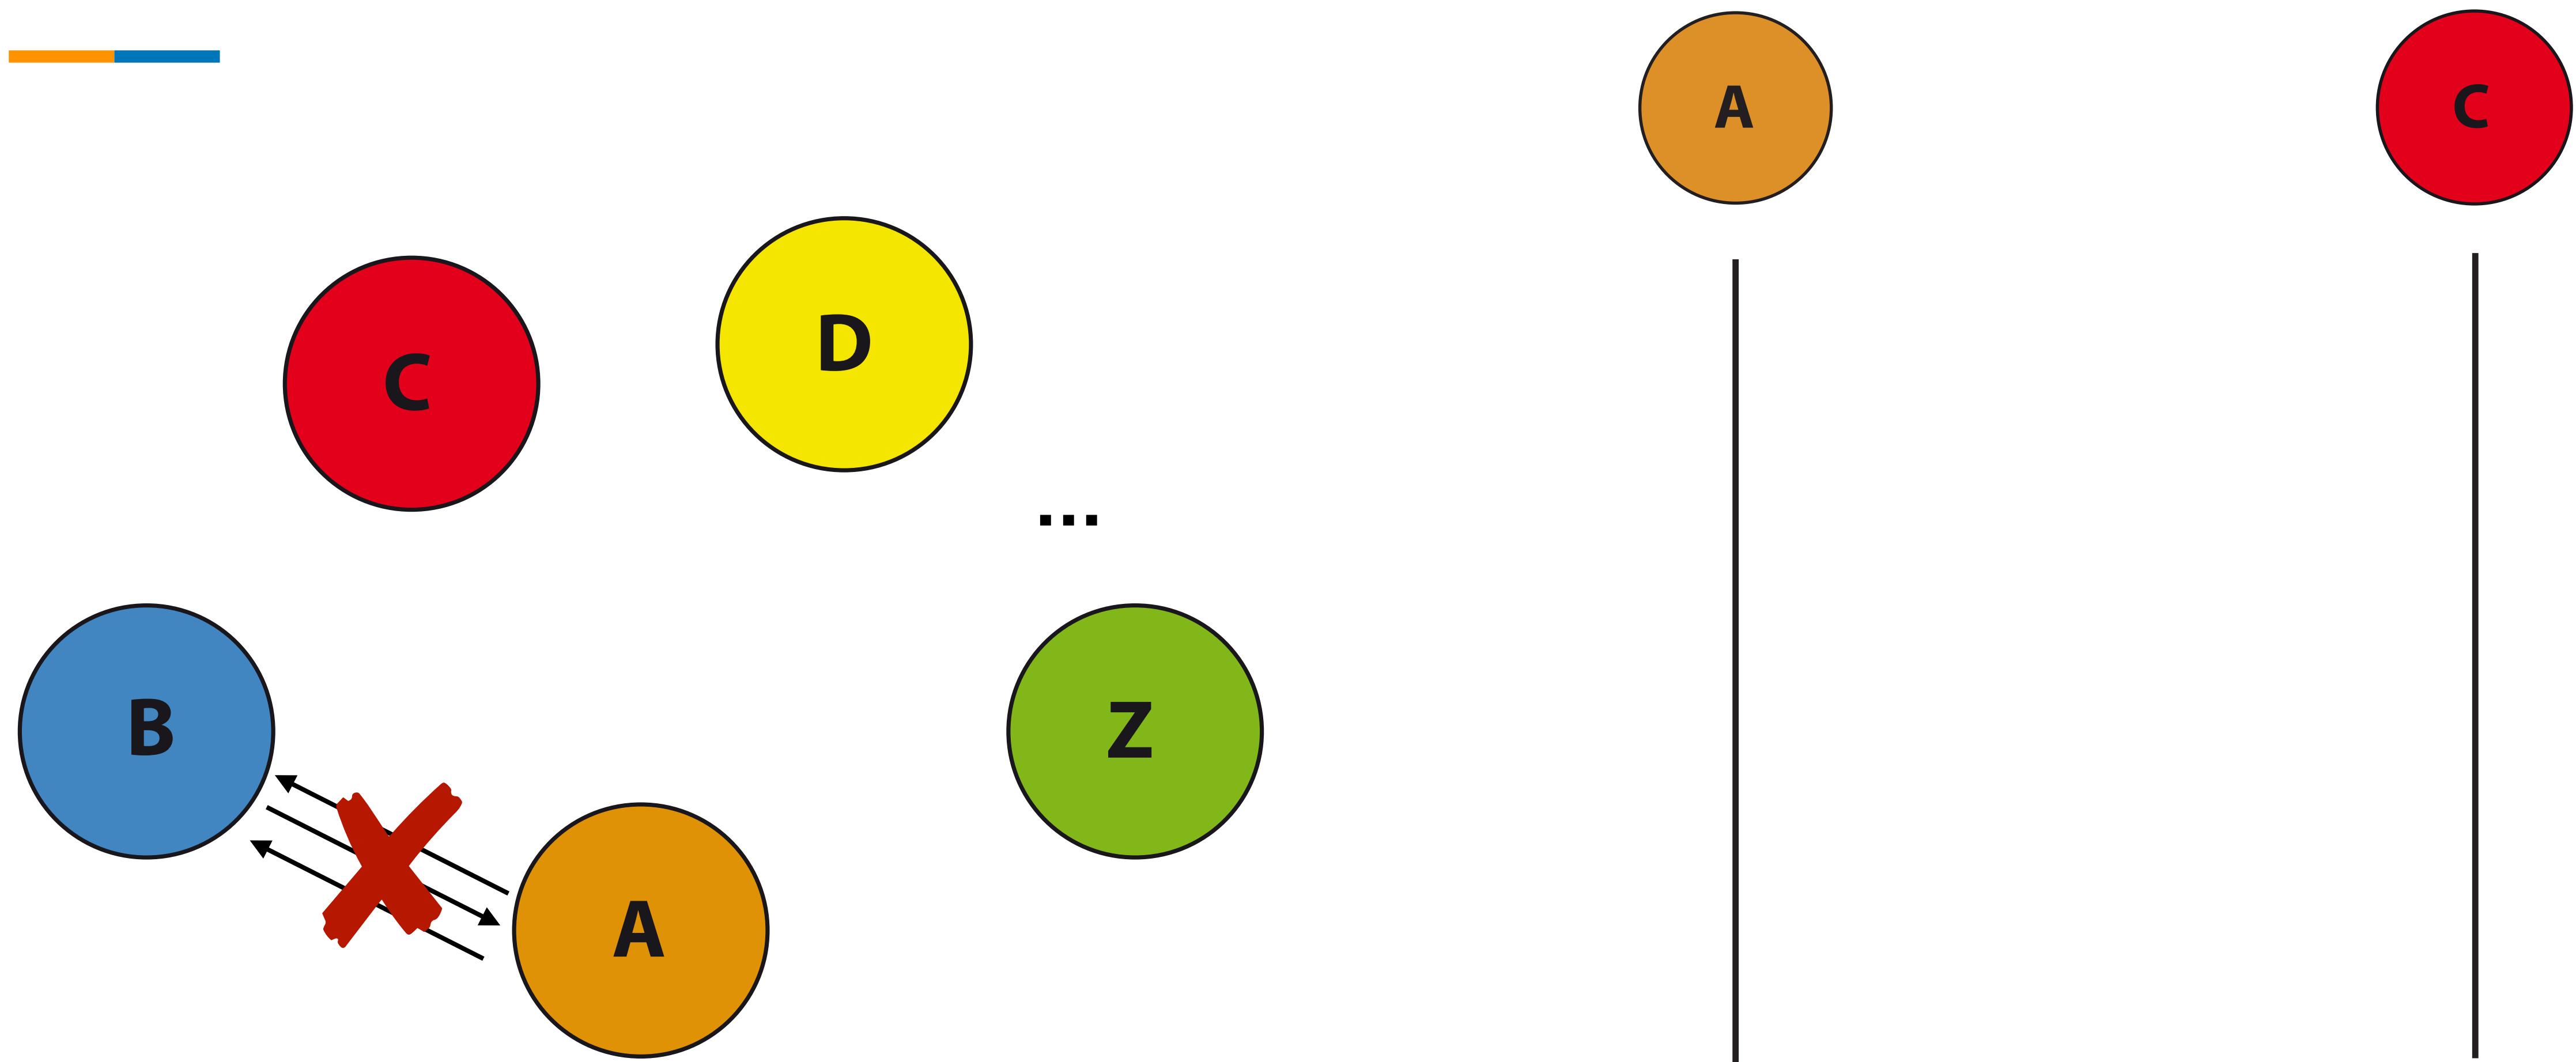
INFORMATION PROPAGATION (2/3)



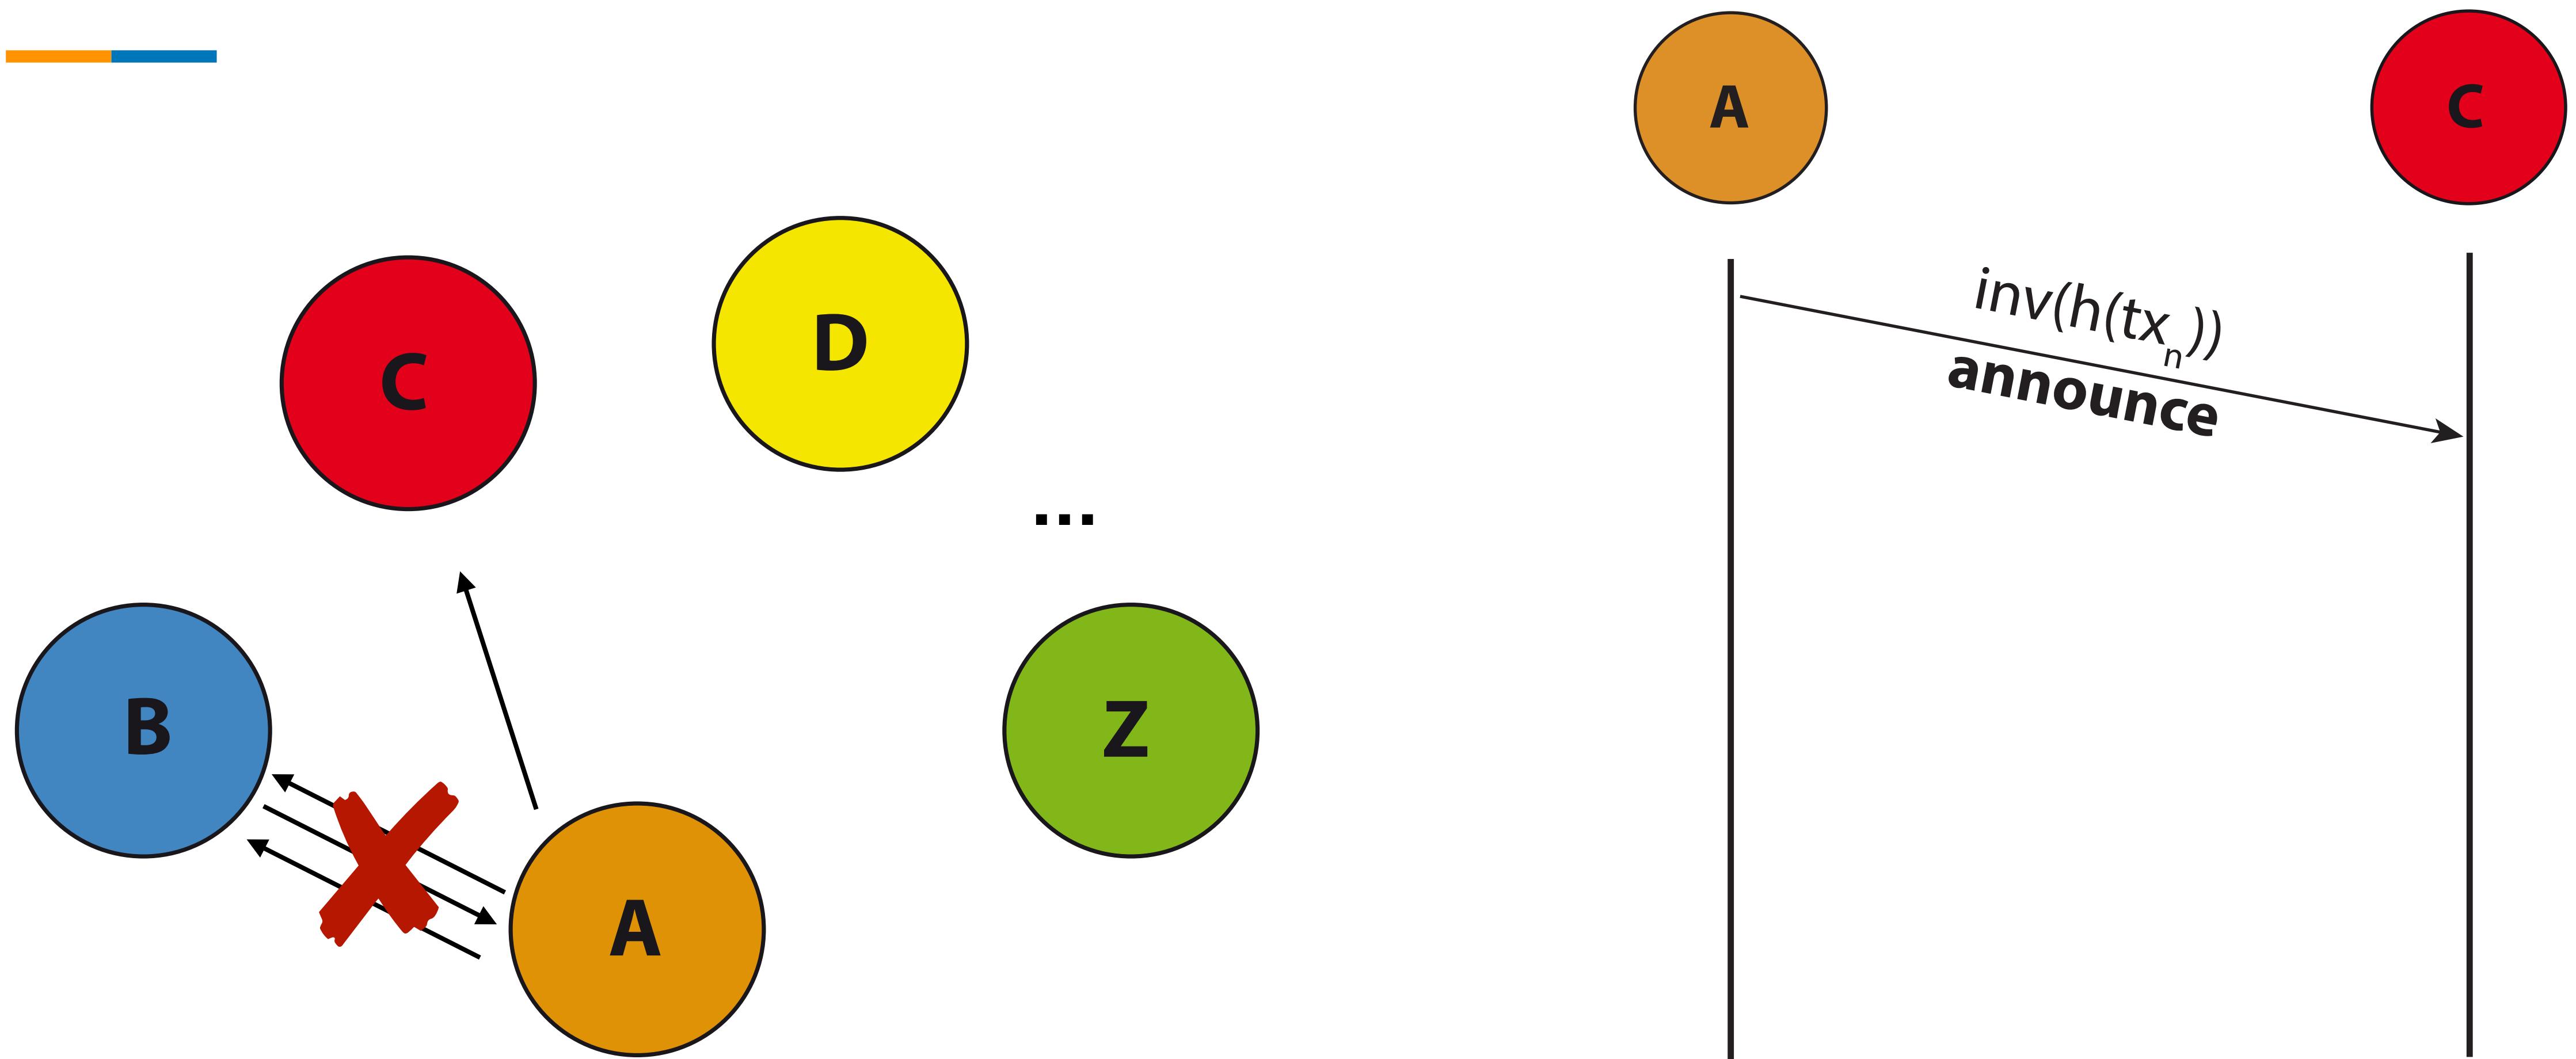
INFORMATION PROPAGATION (2/3)



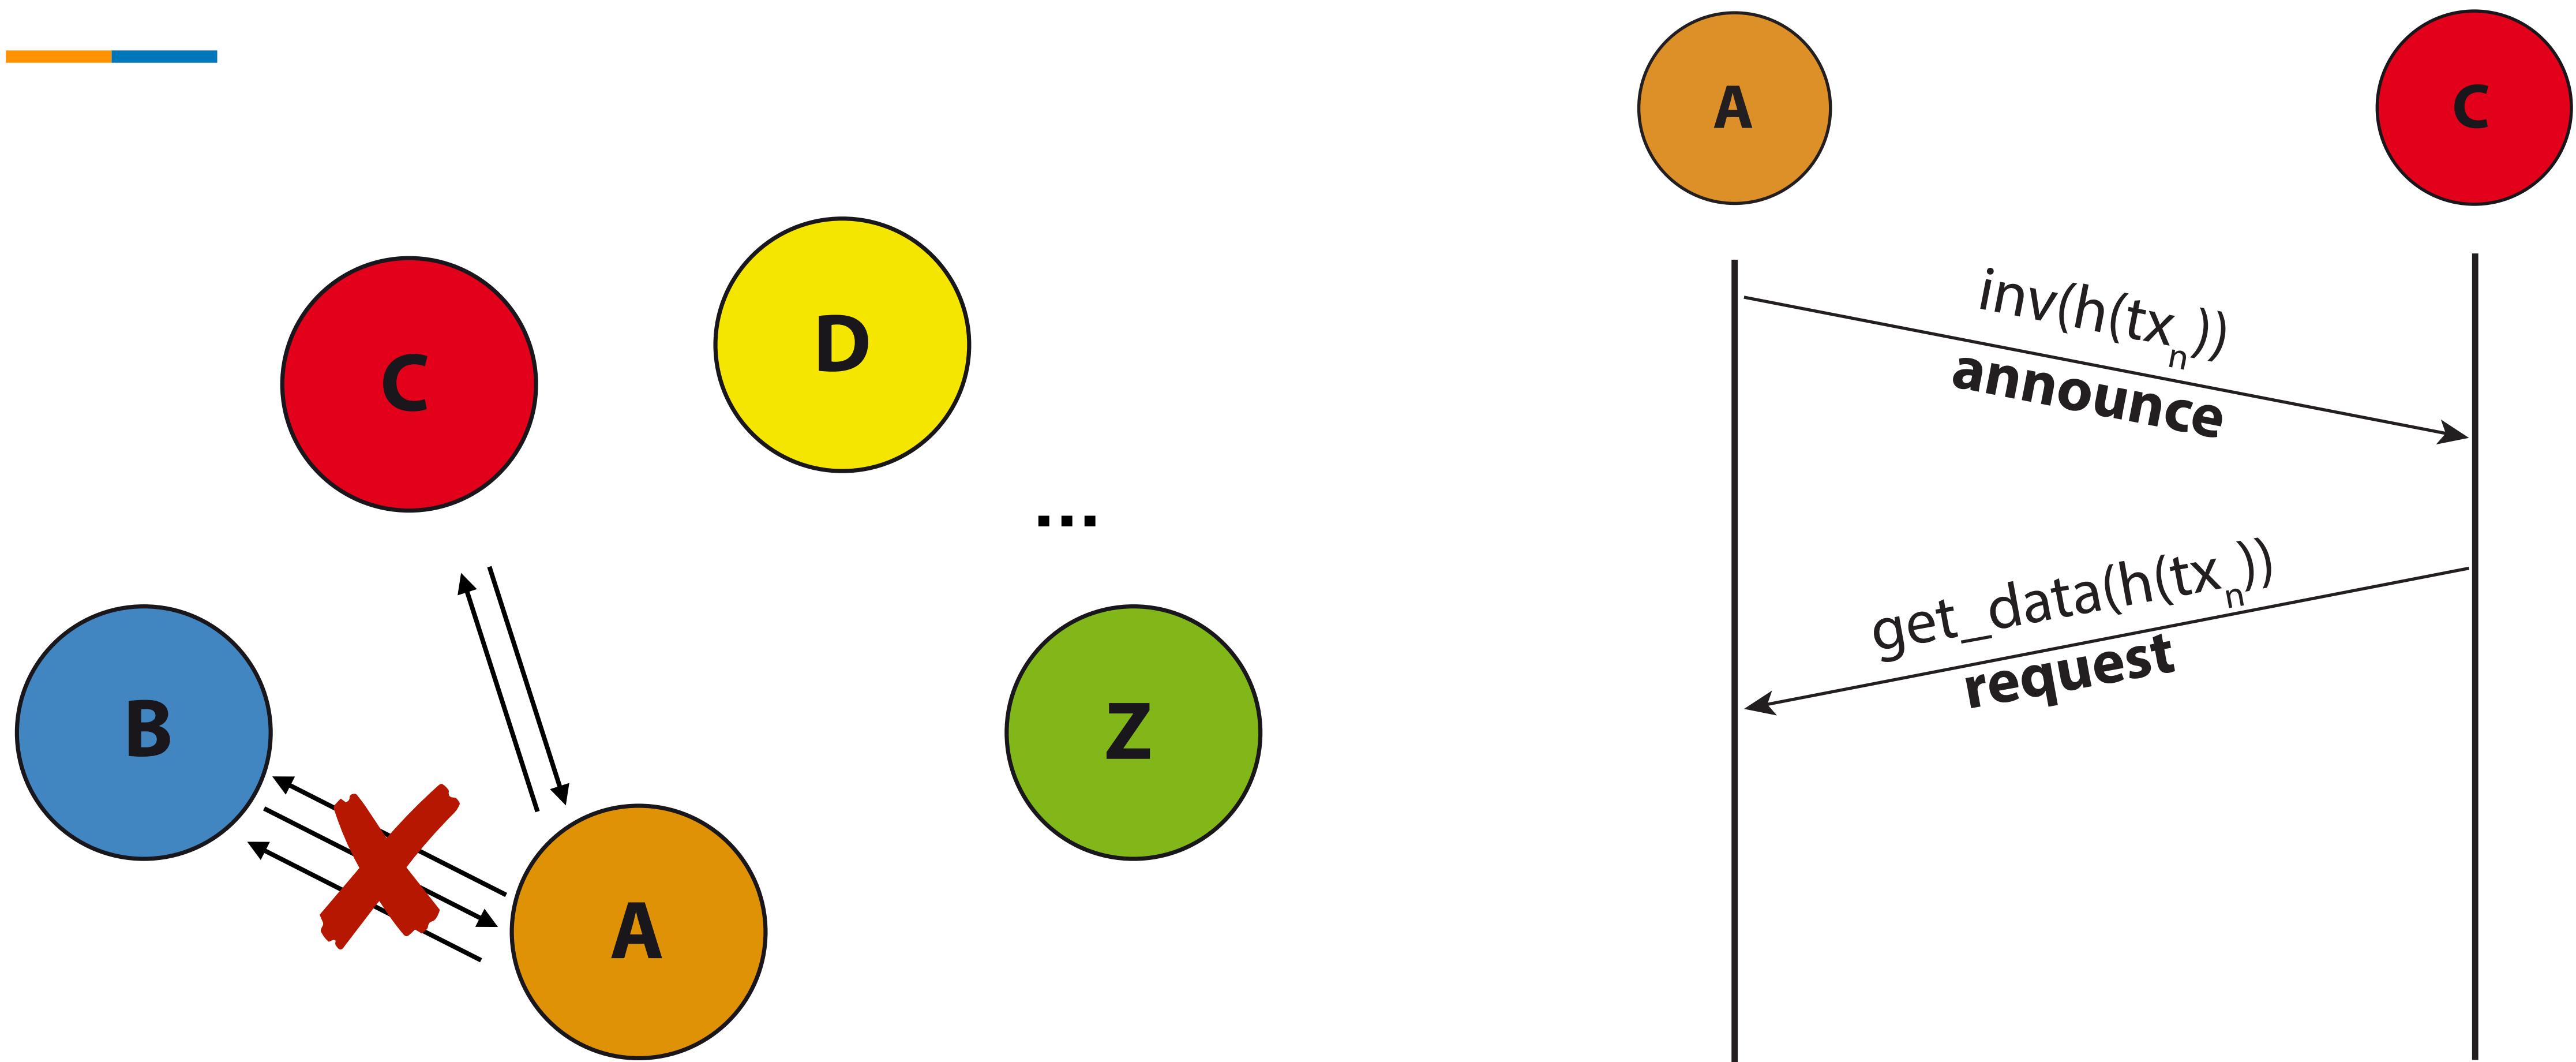
INFORMATION PROPAGATION (2/3)



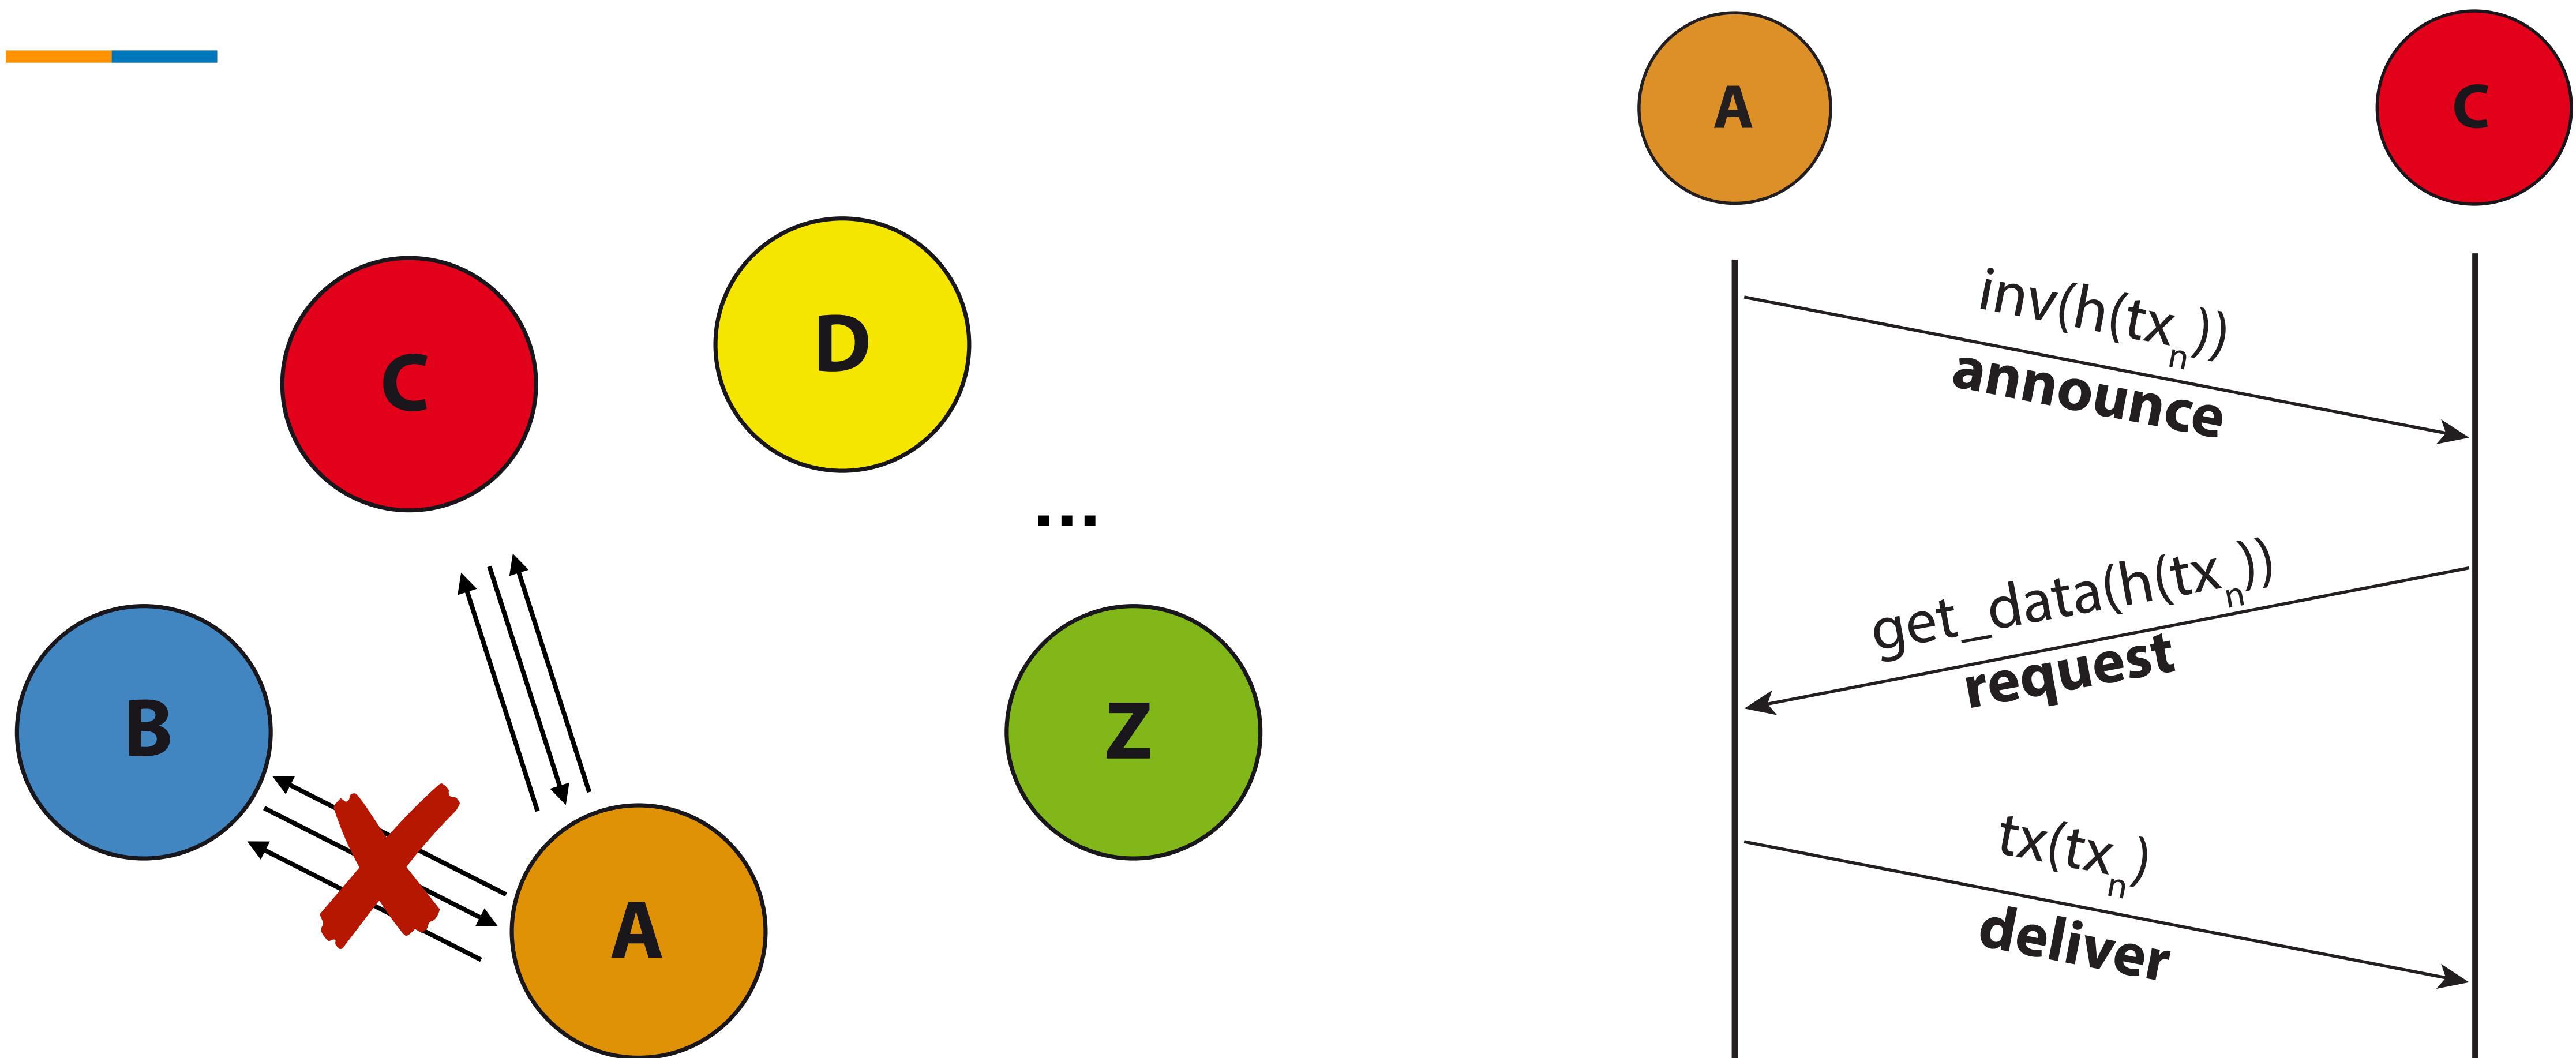
INFORMATION PROPAGATION (2/3)



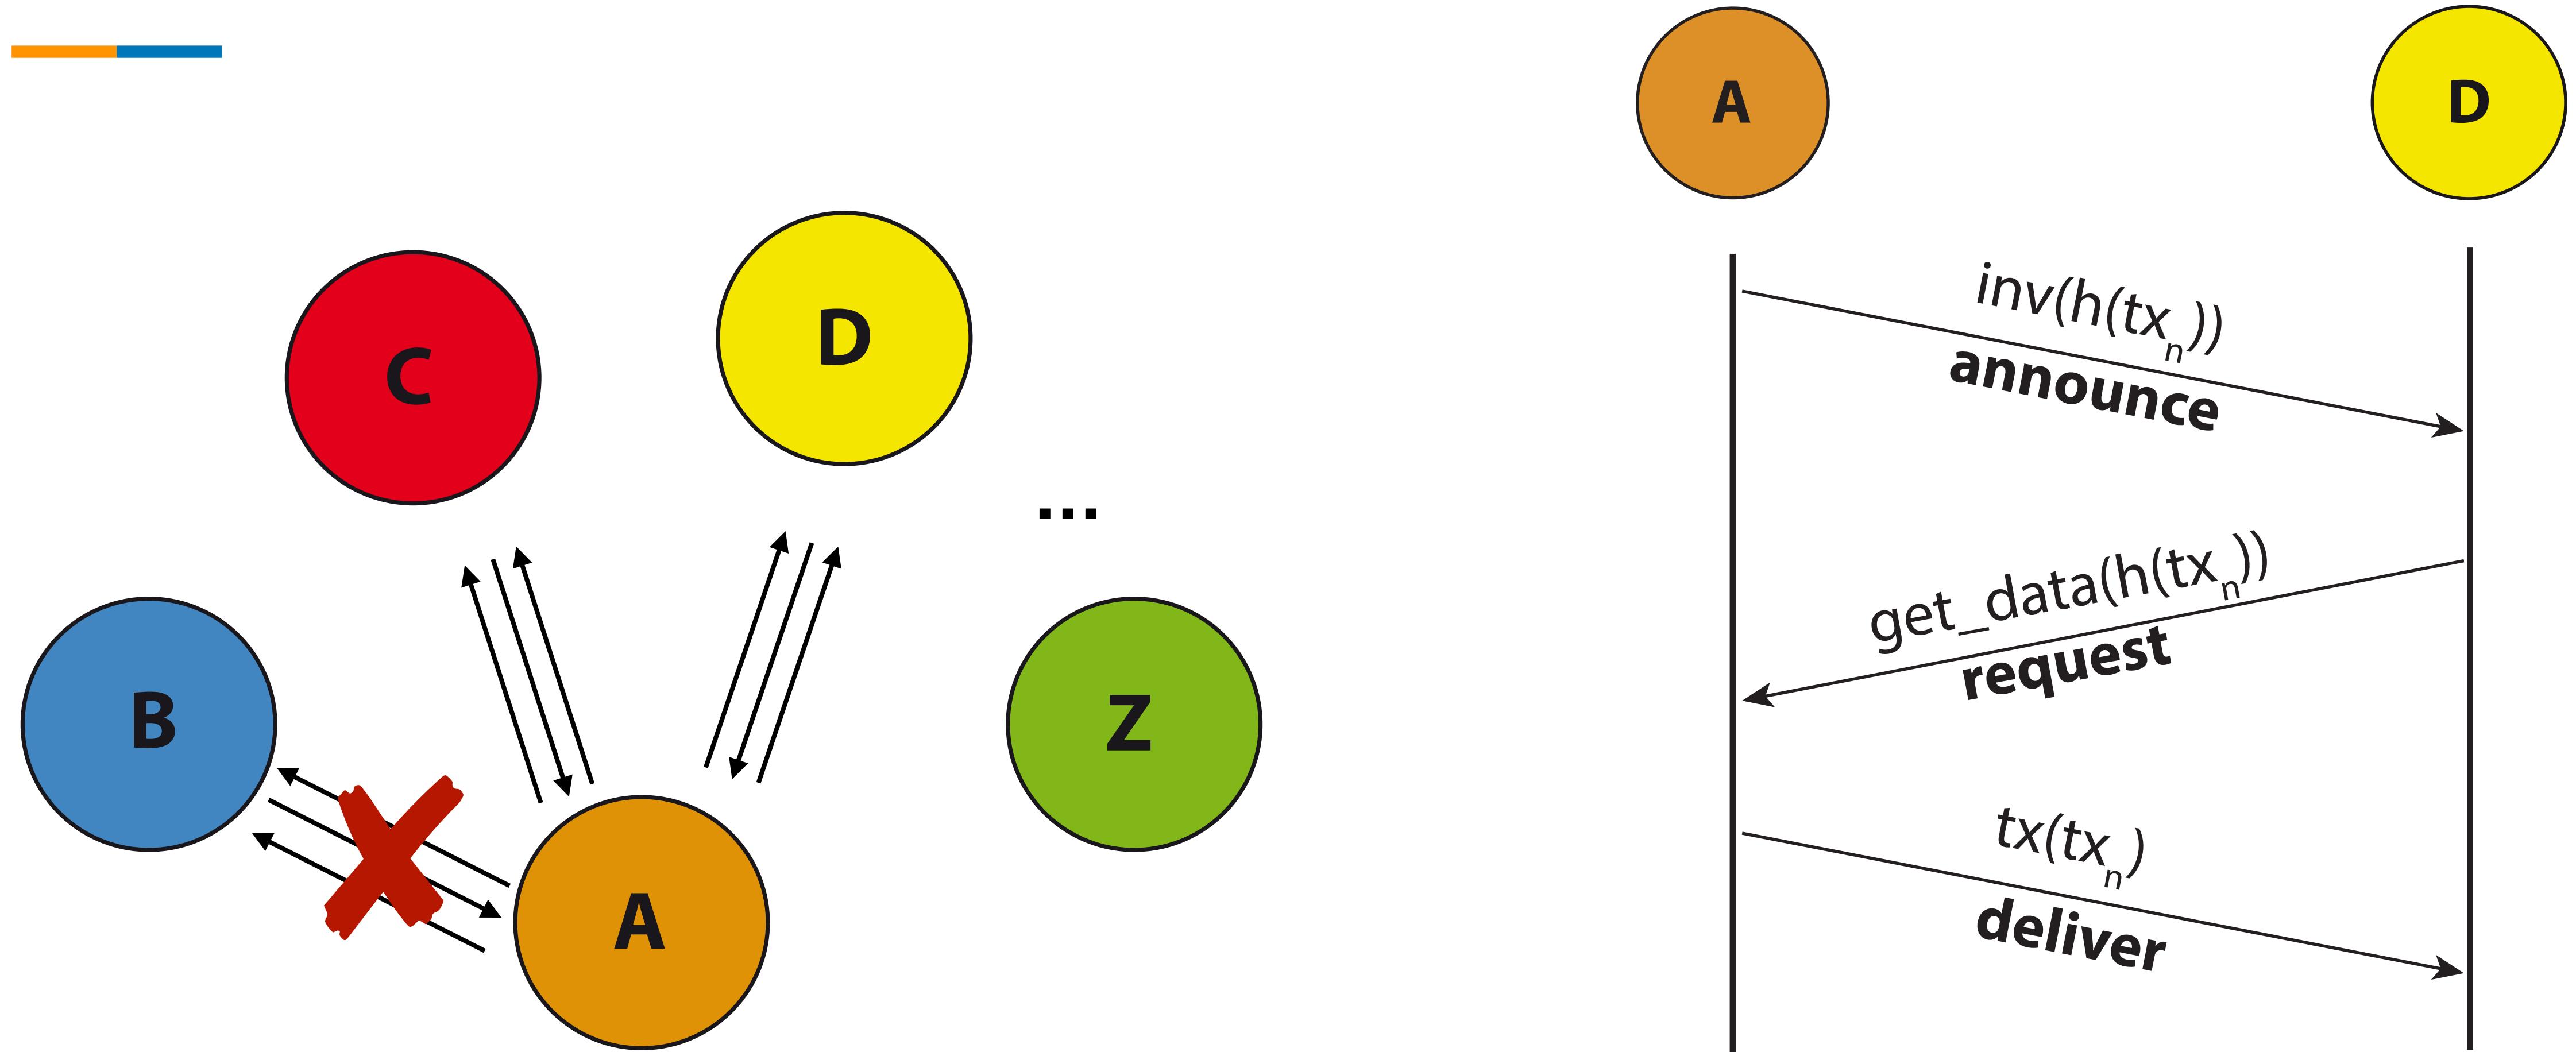
INFORMATION PROPAGATION (2/3)



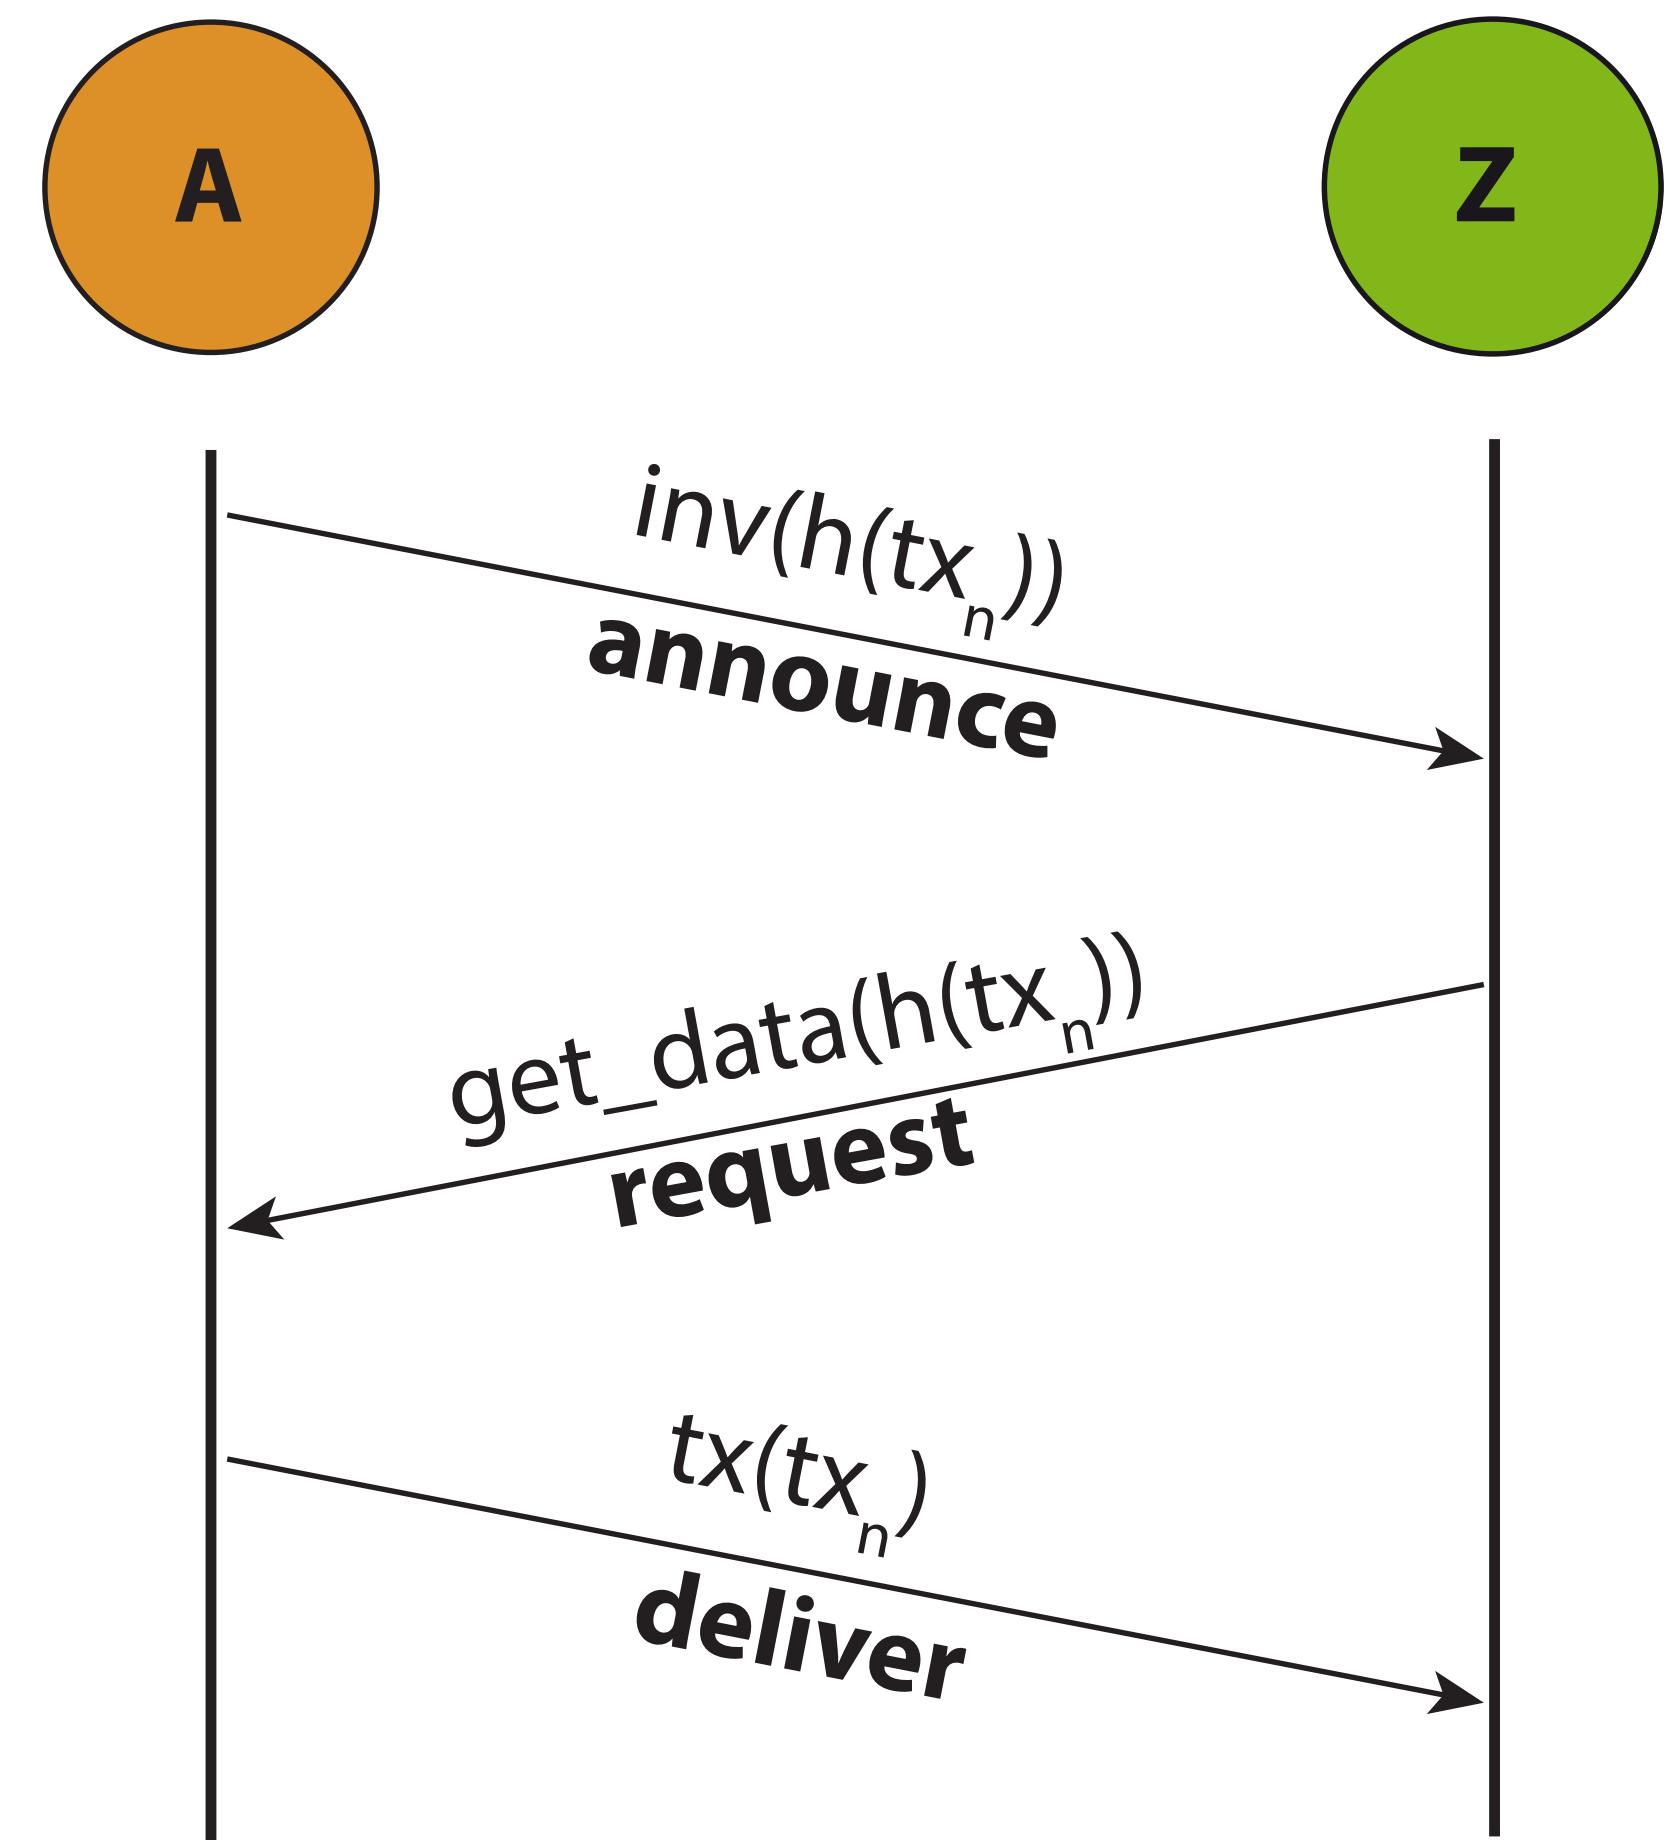
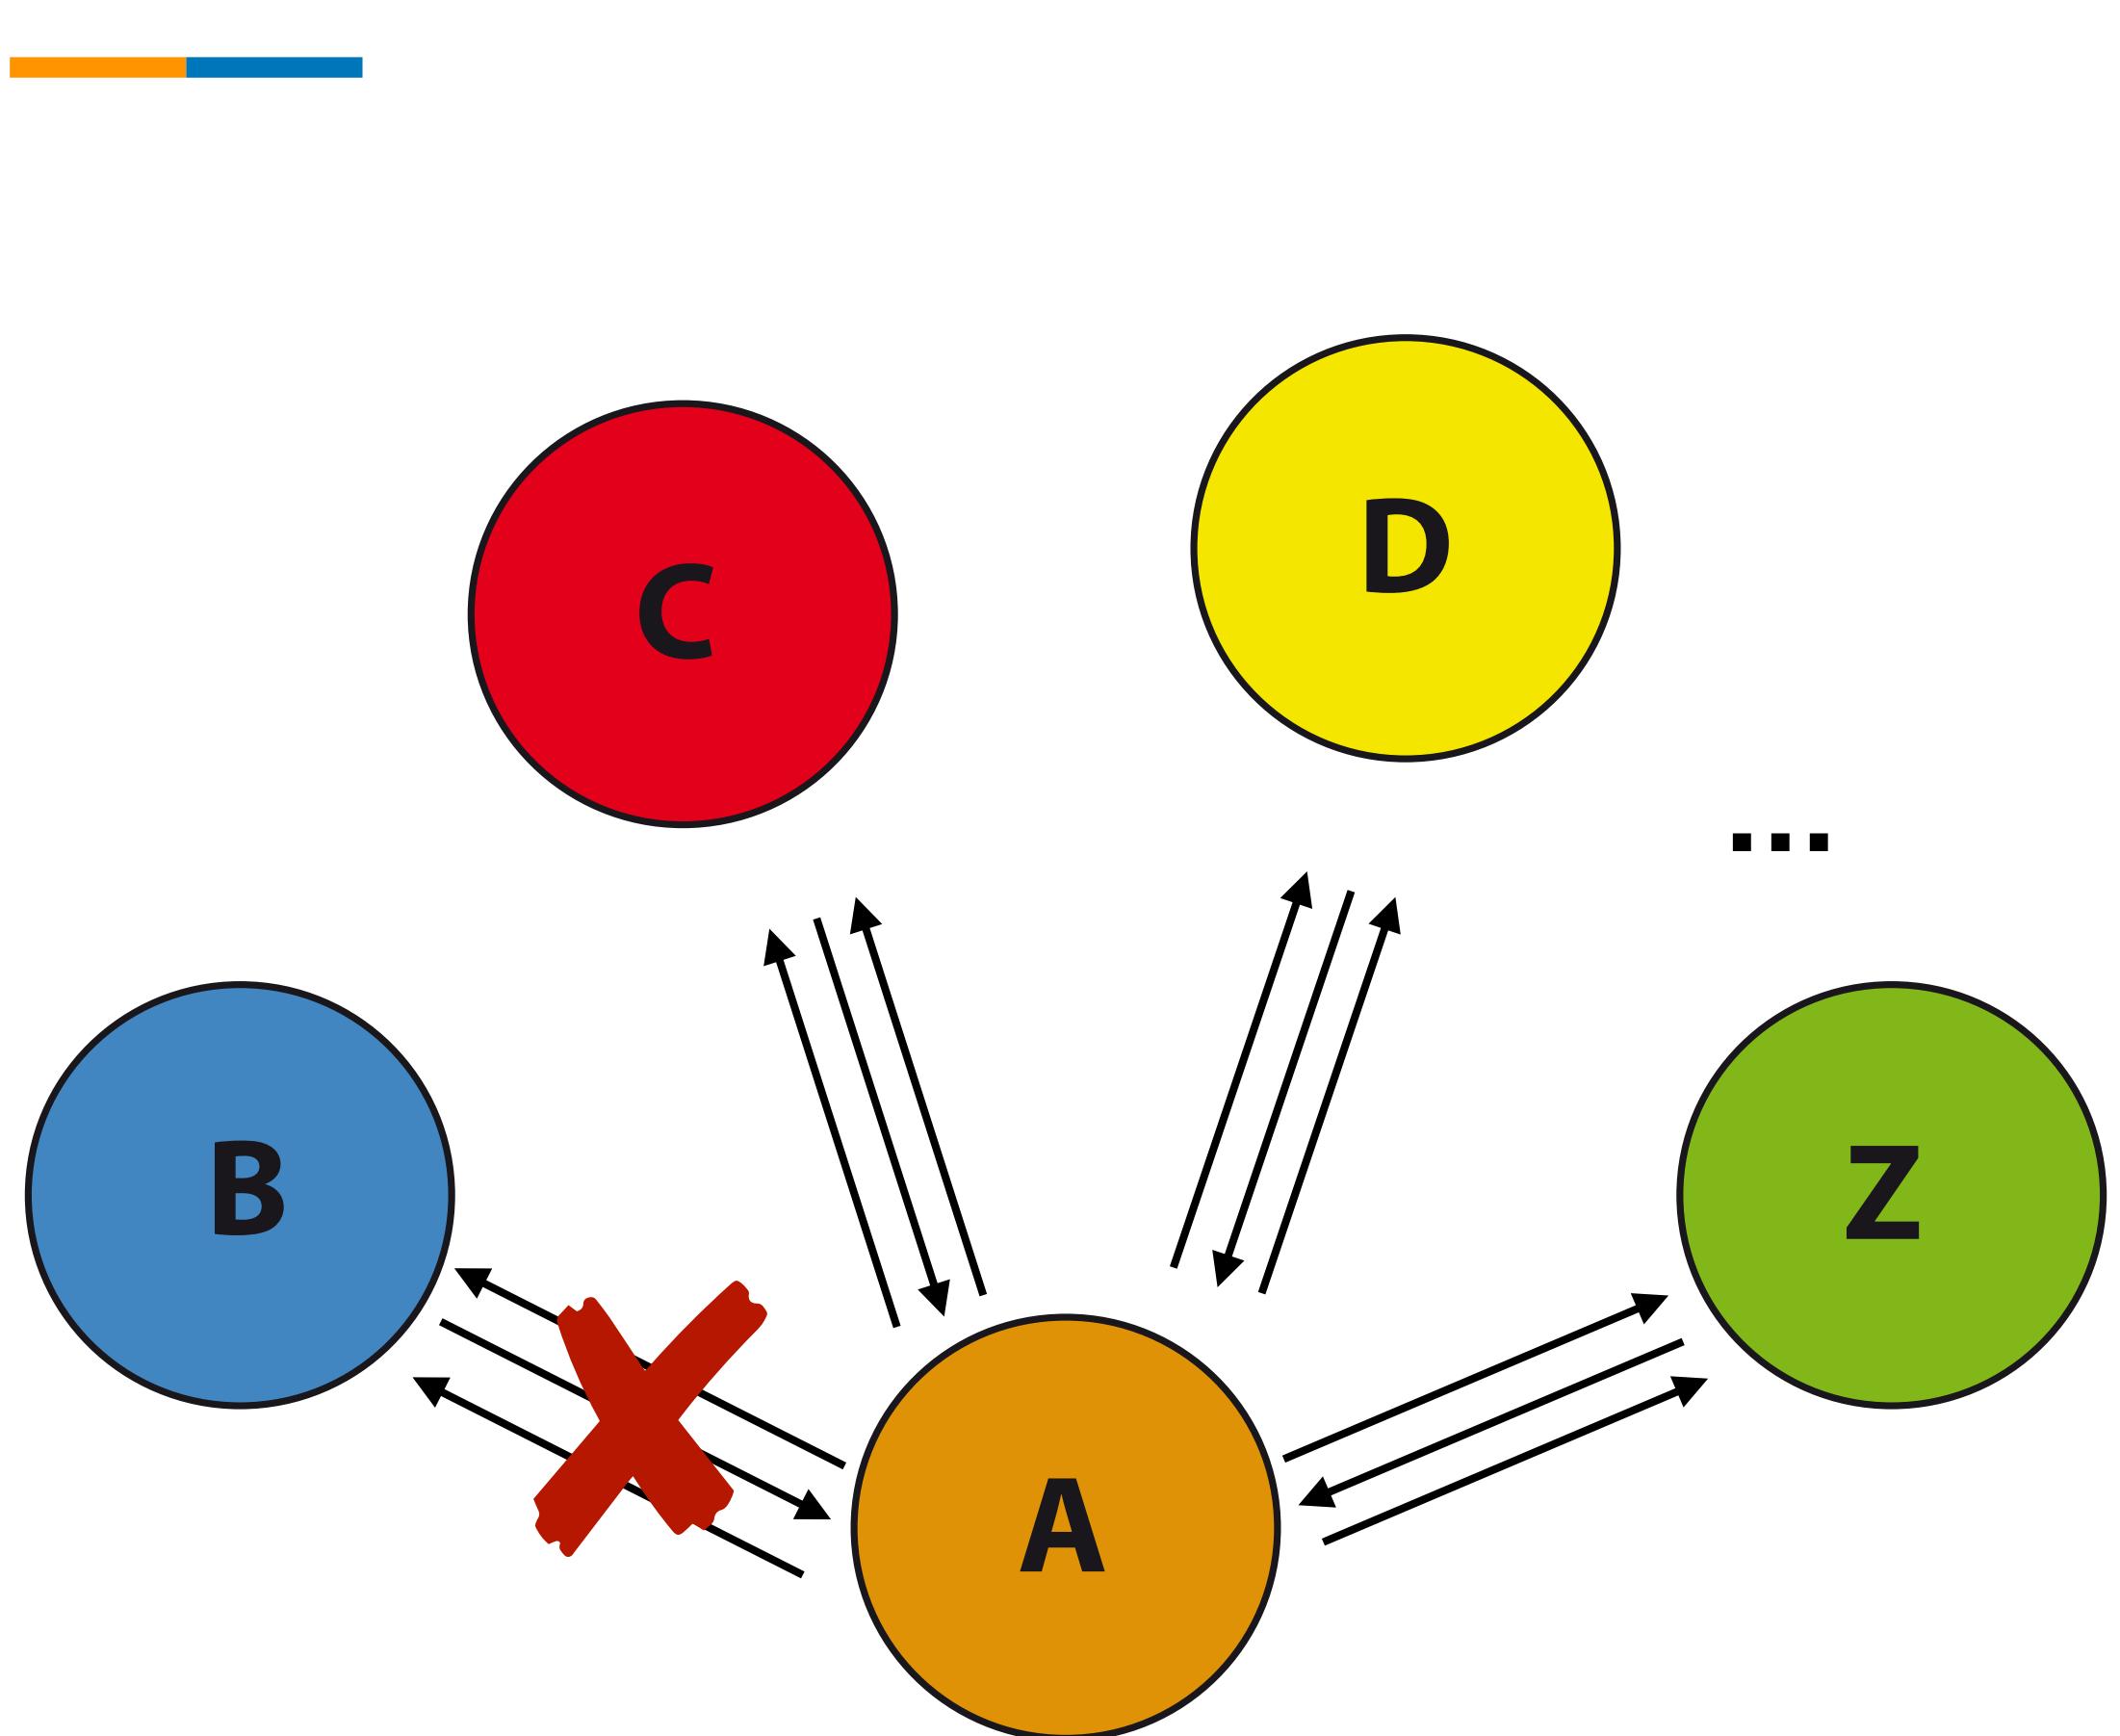
INFORMATION PROPAGATION (2/3)



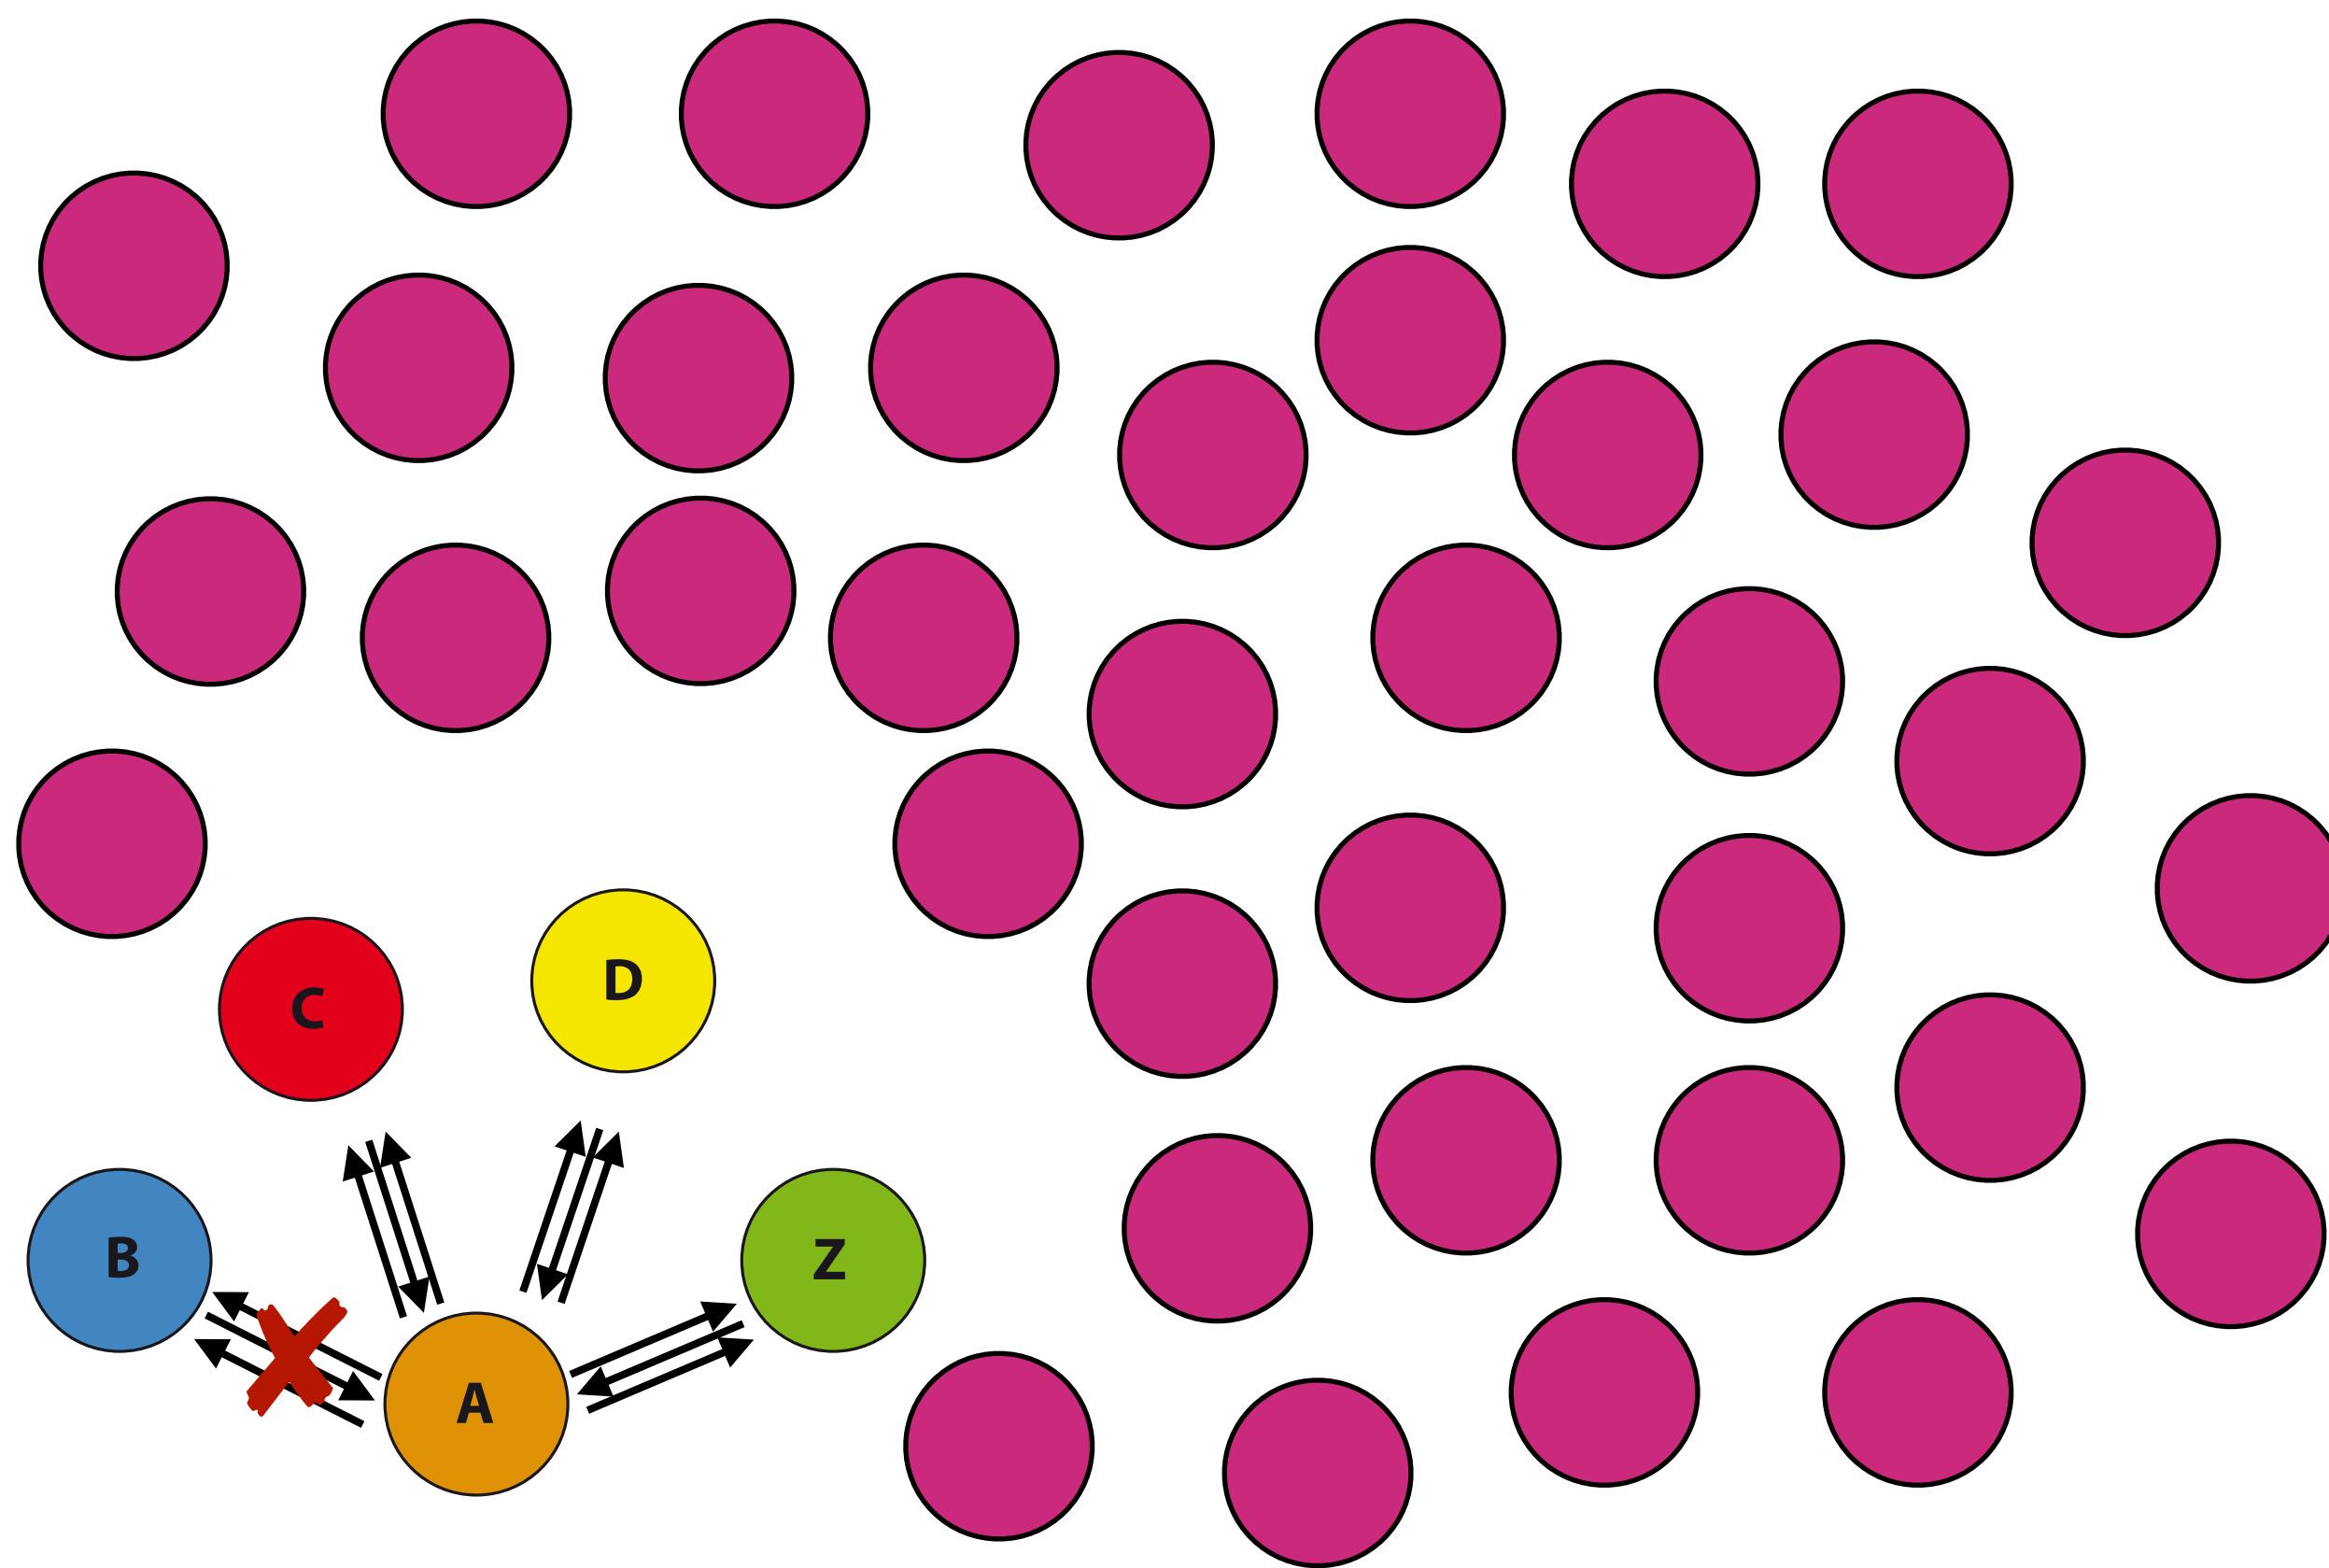
INFORMATION PROPAGATION (2/3)



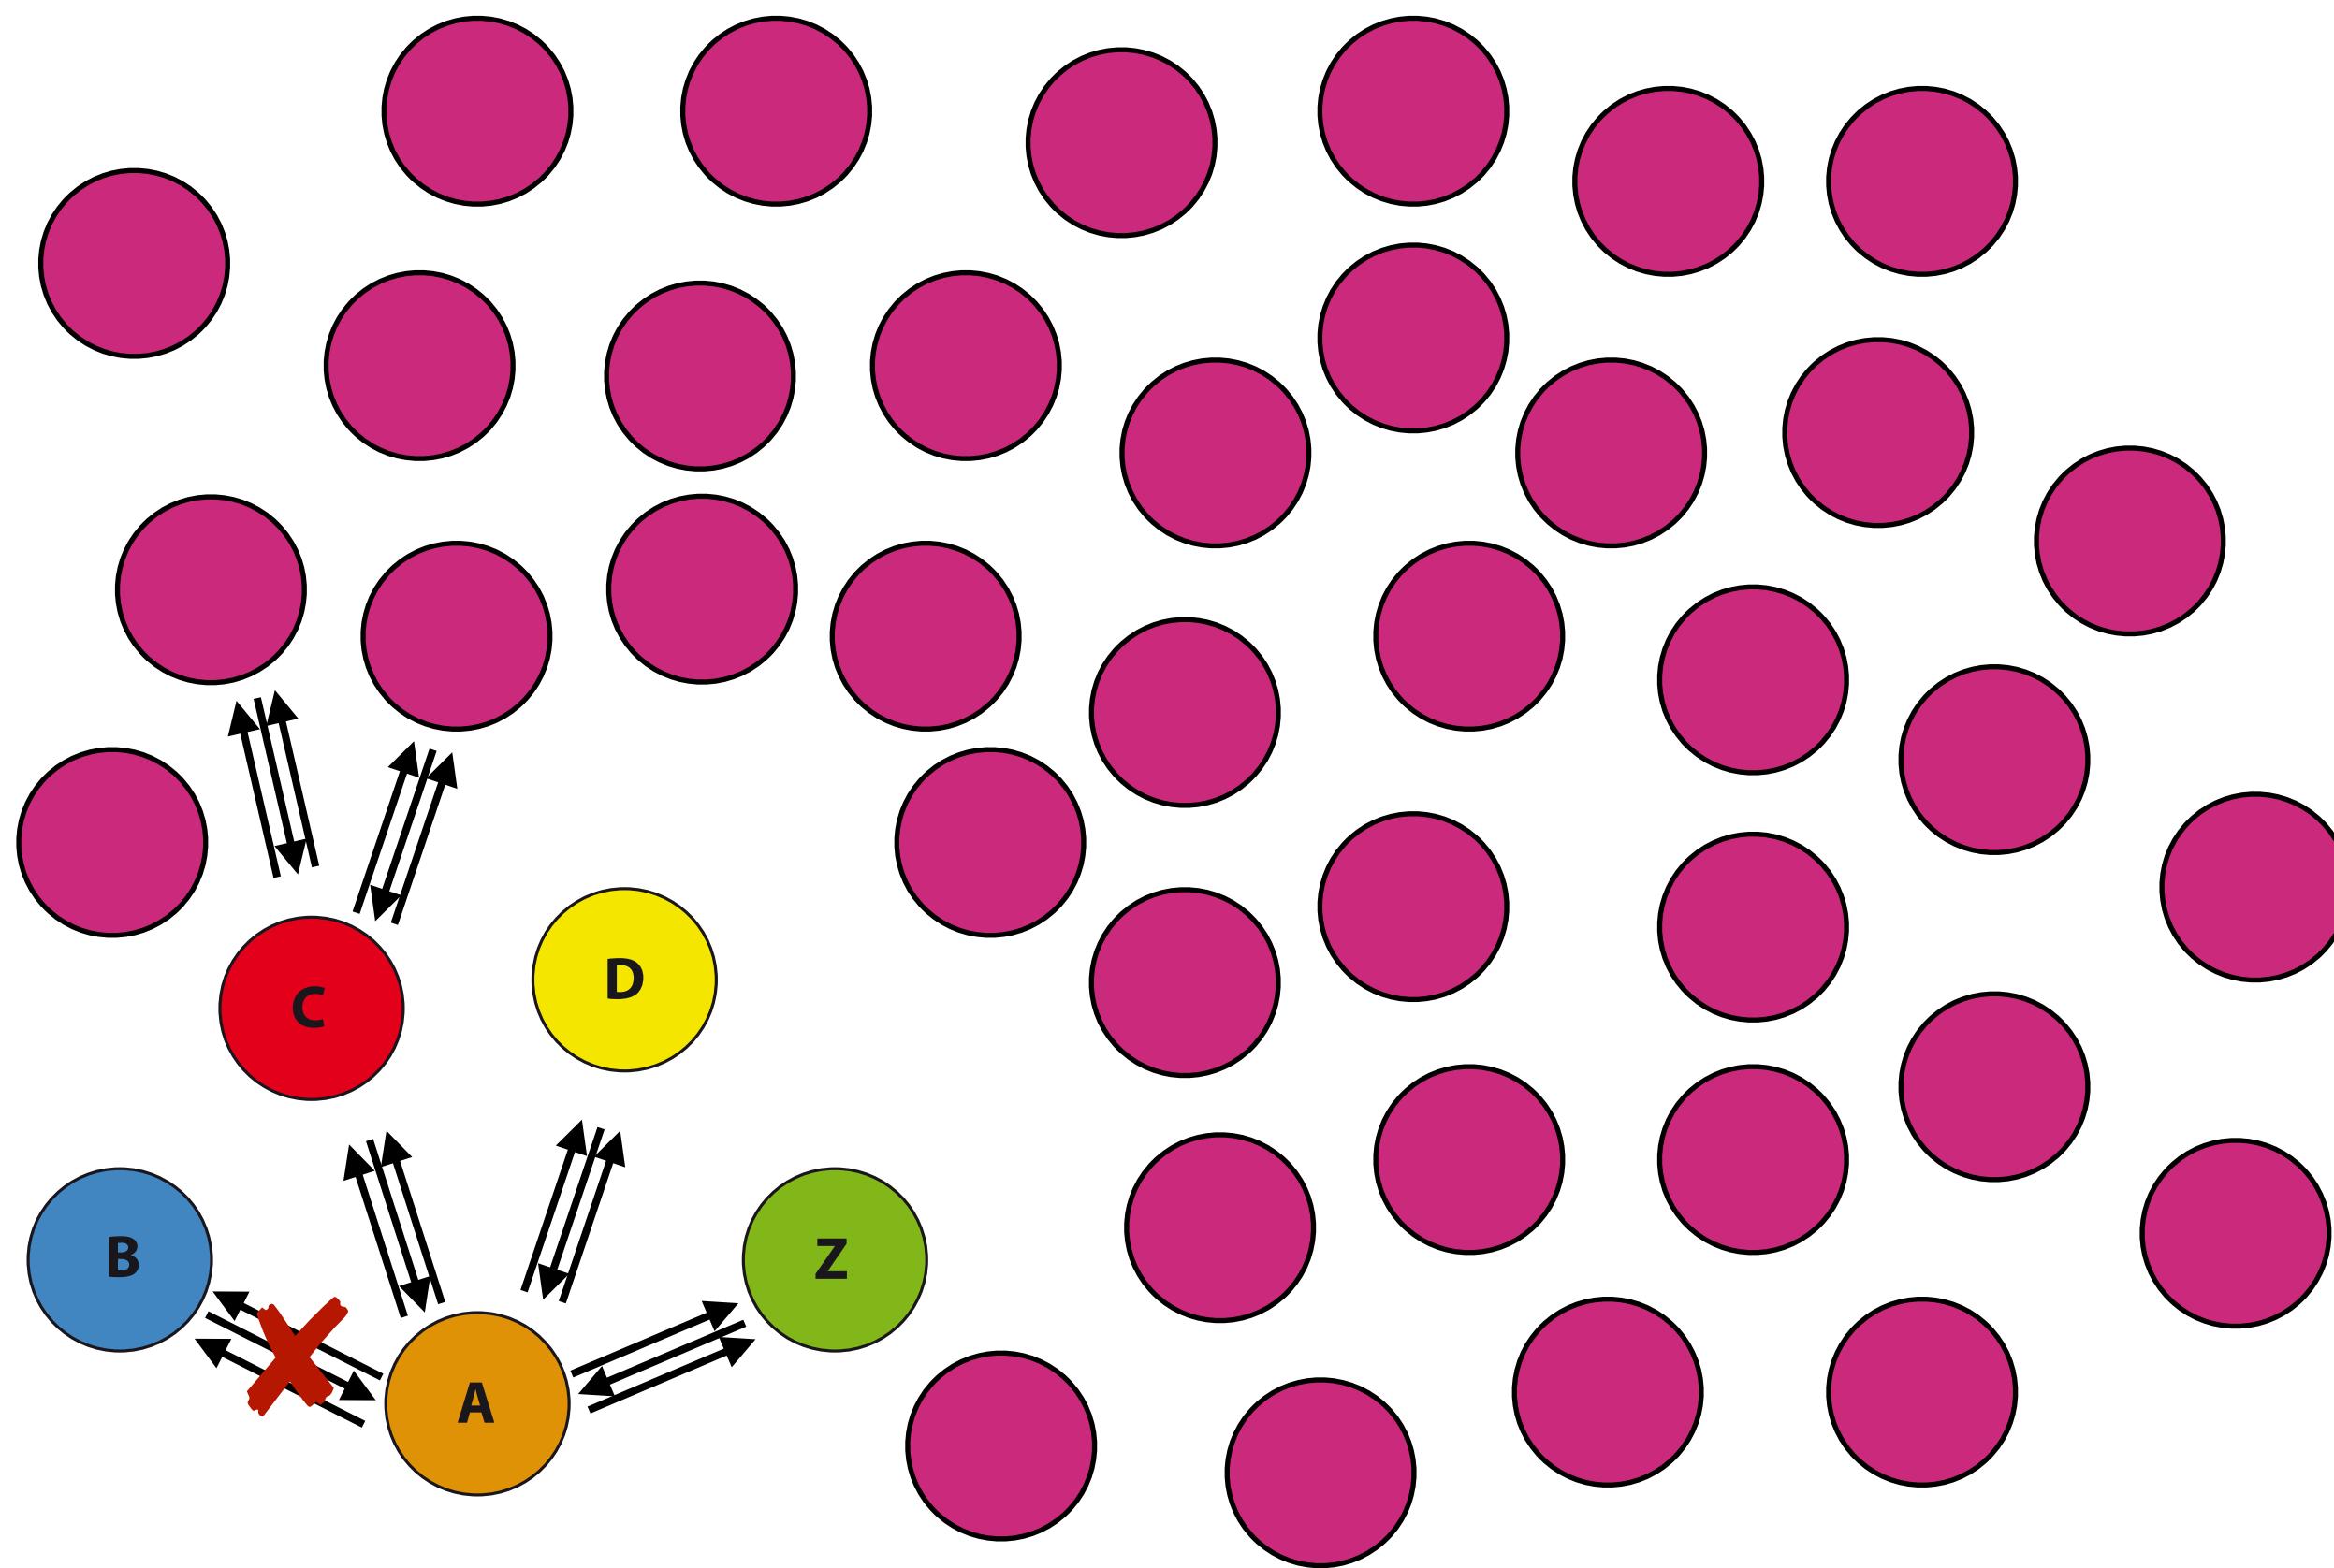
INFORMATION PROPAGATION (2/3)



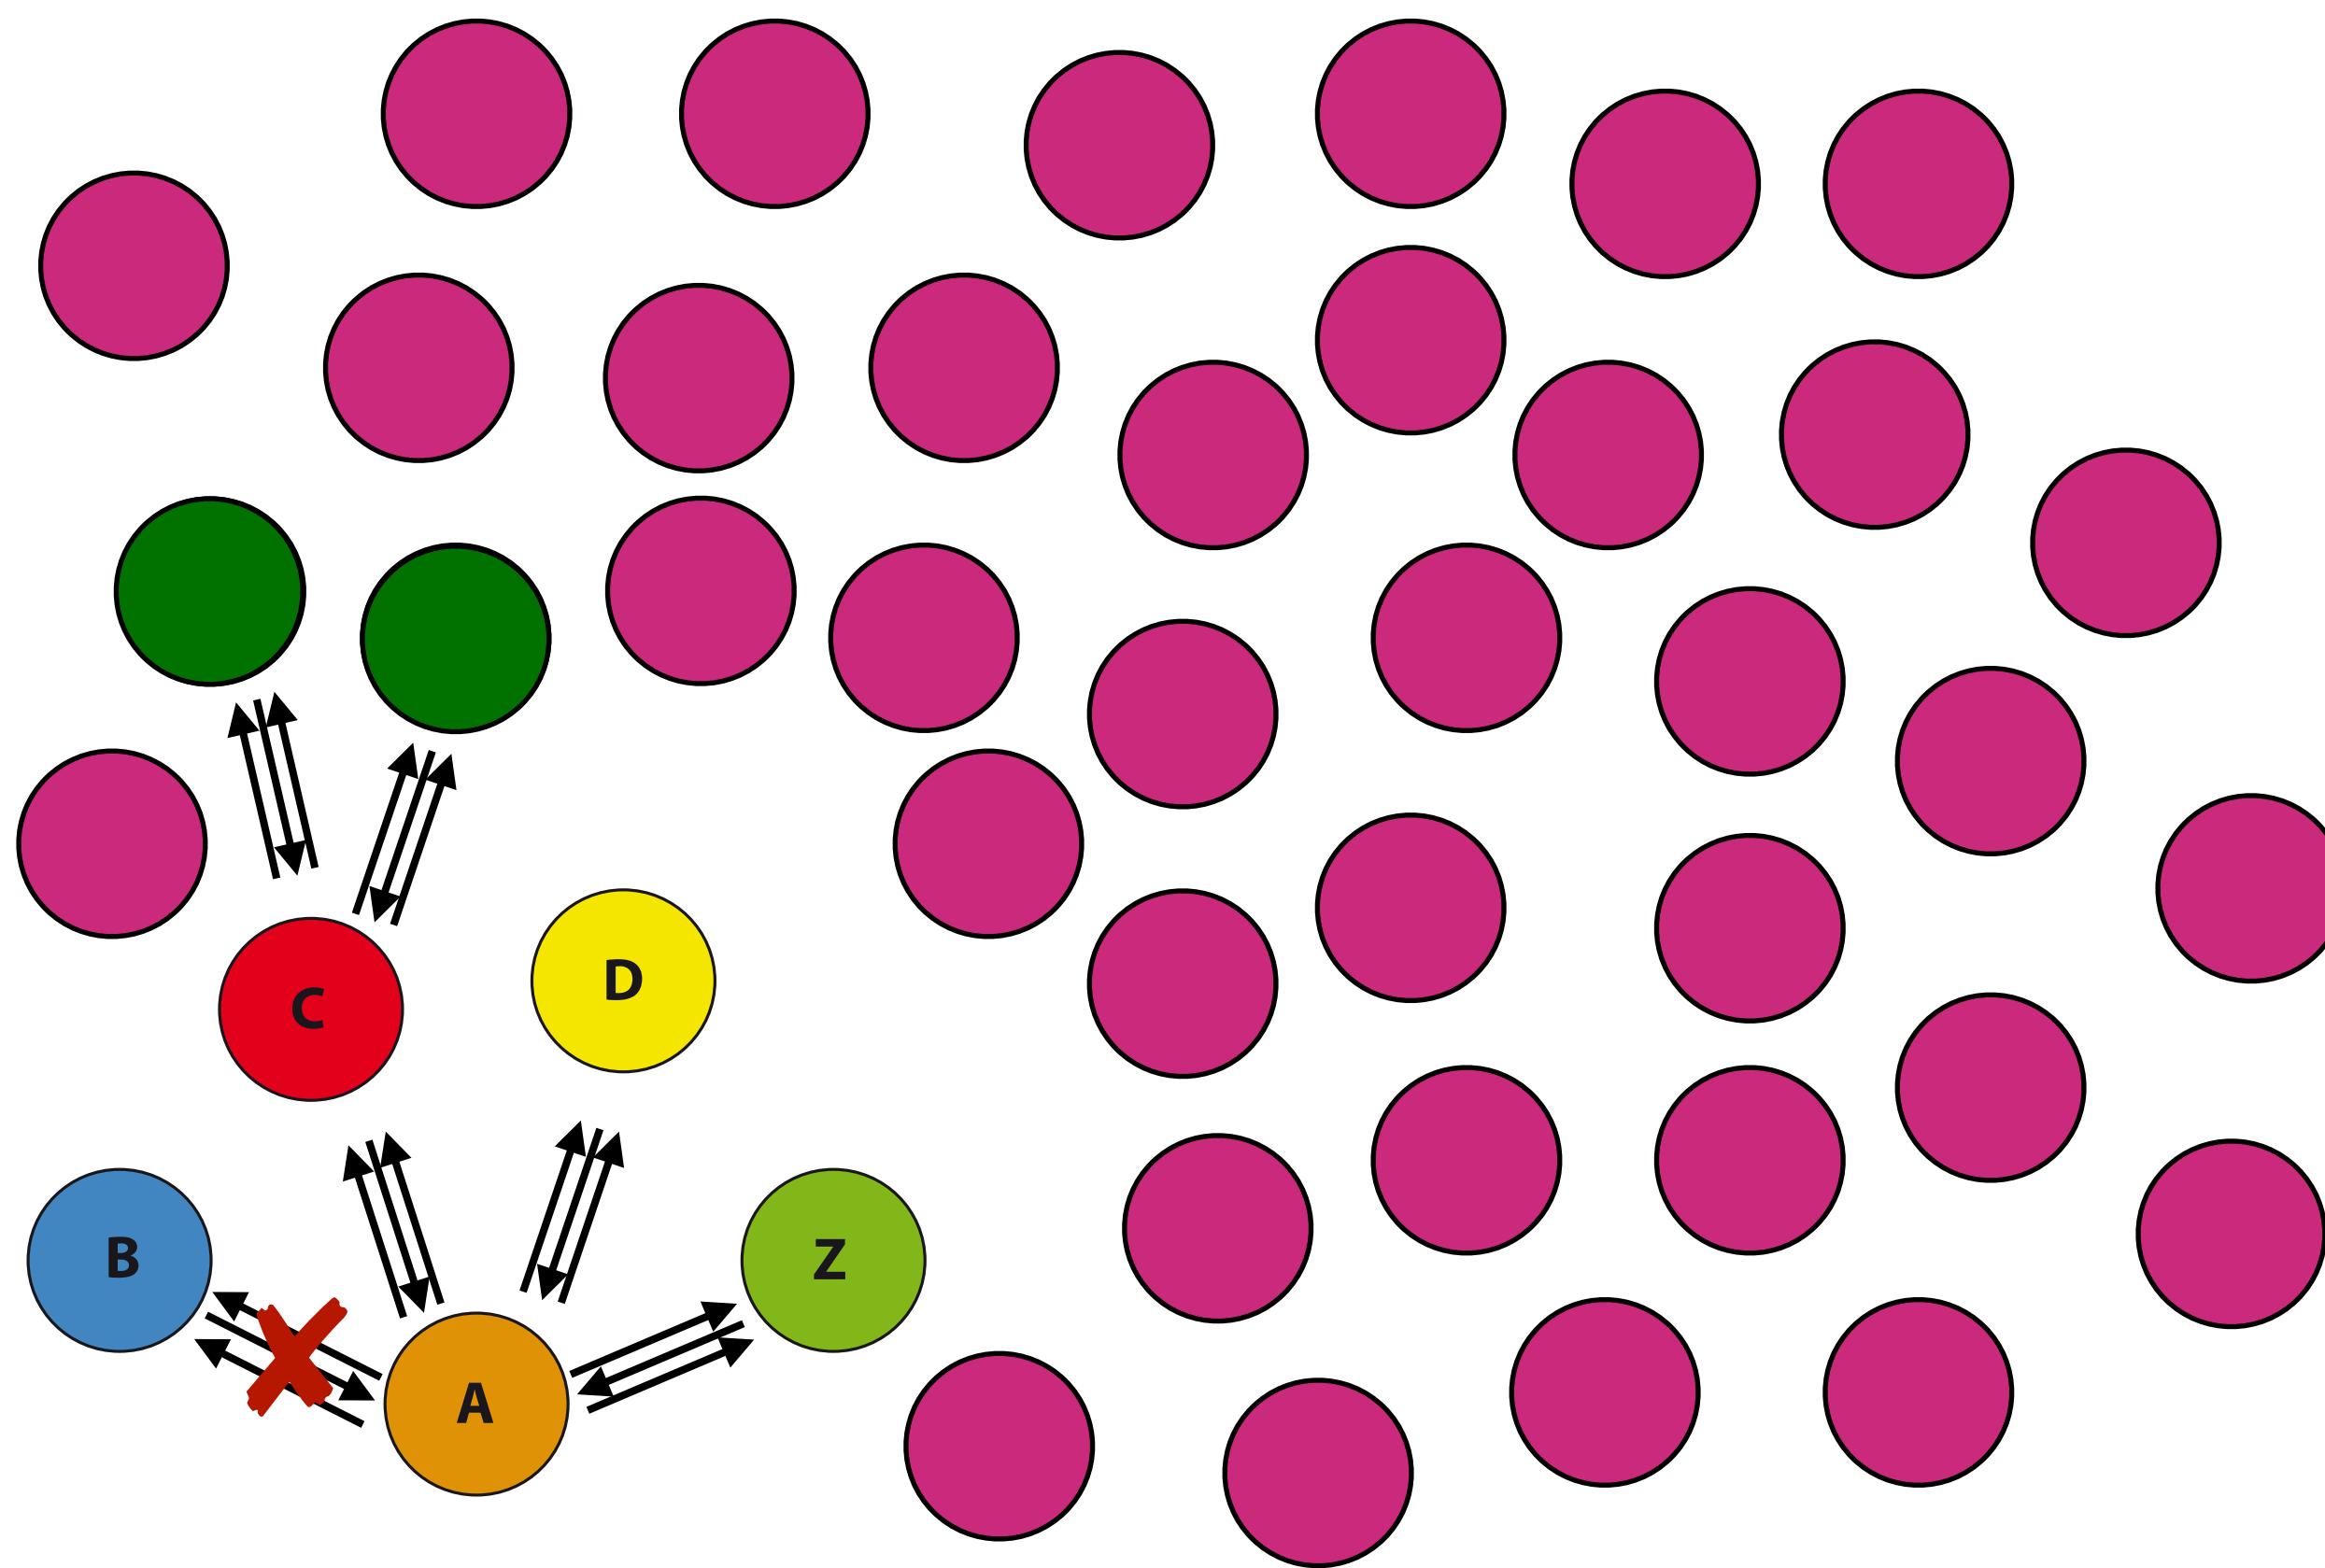
INFORMATION PROPAGATION (3/3)



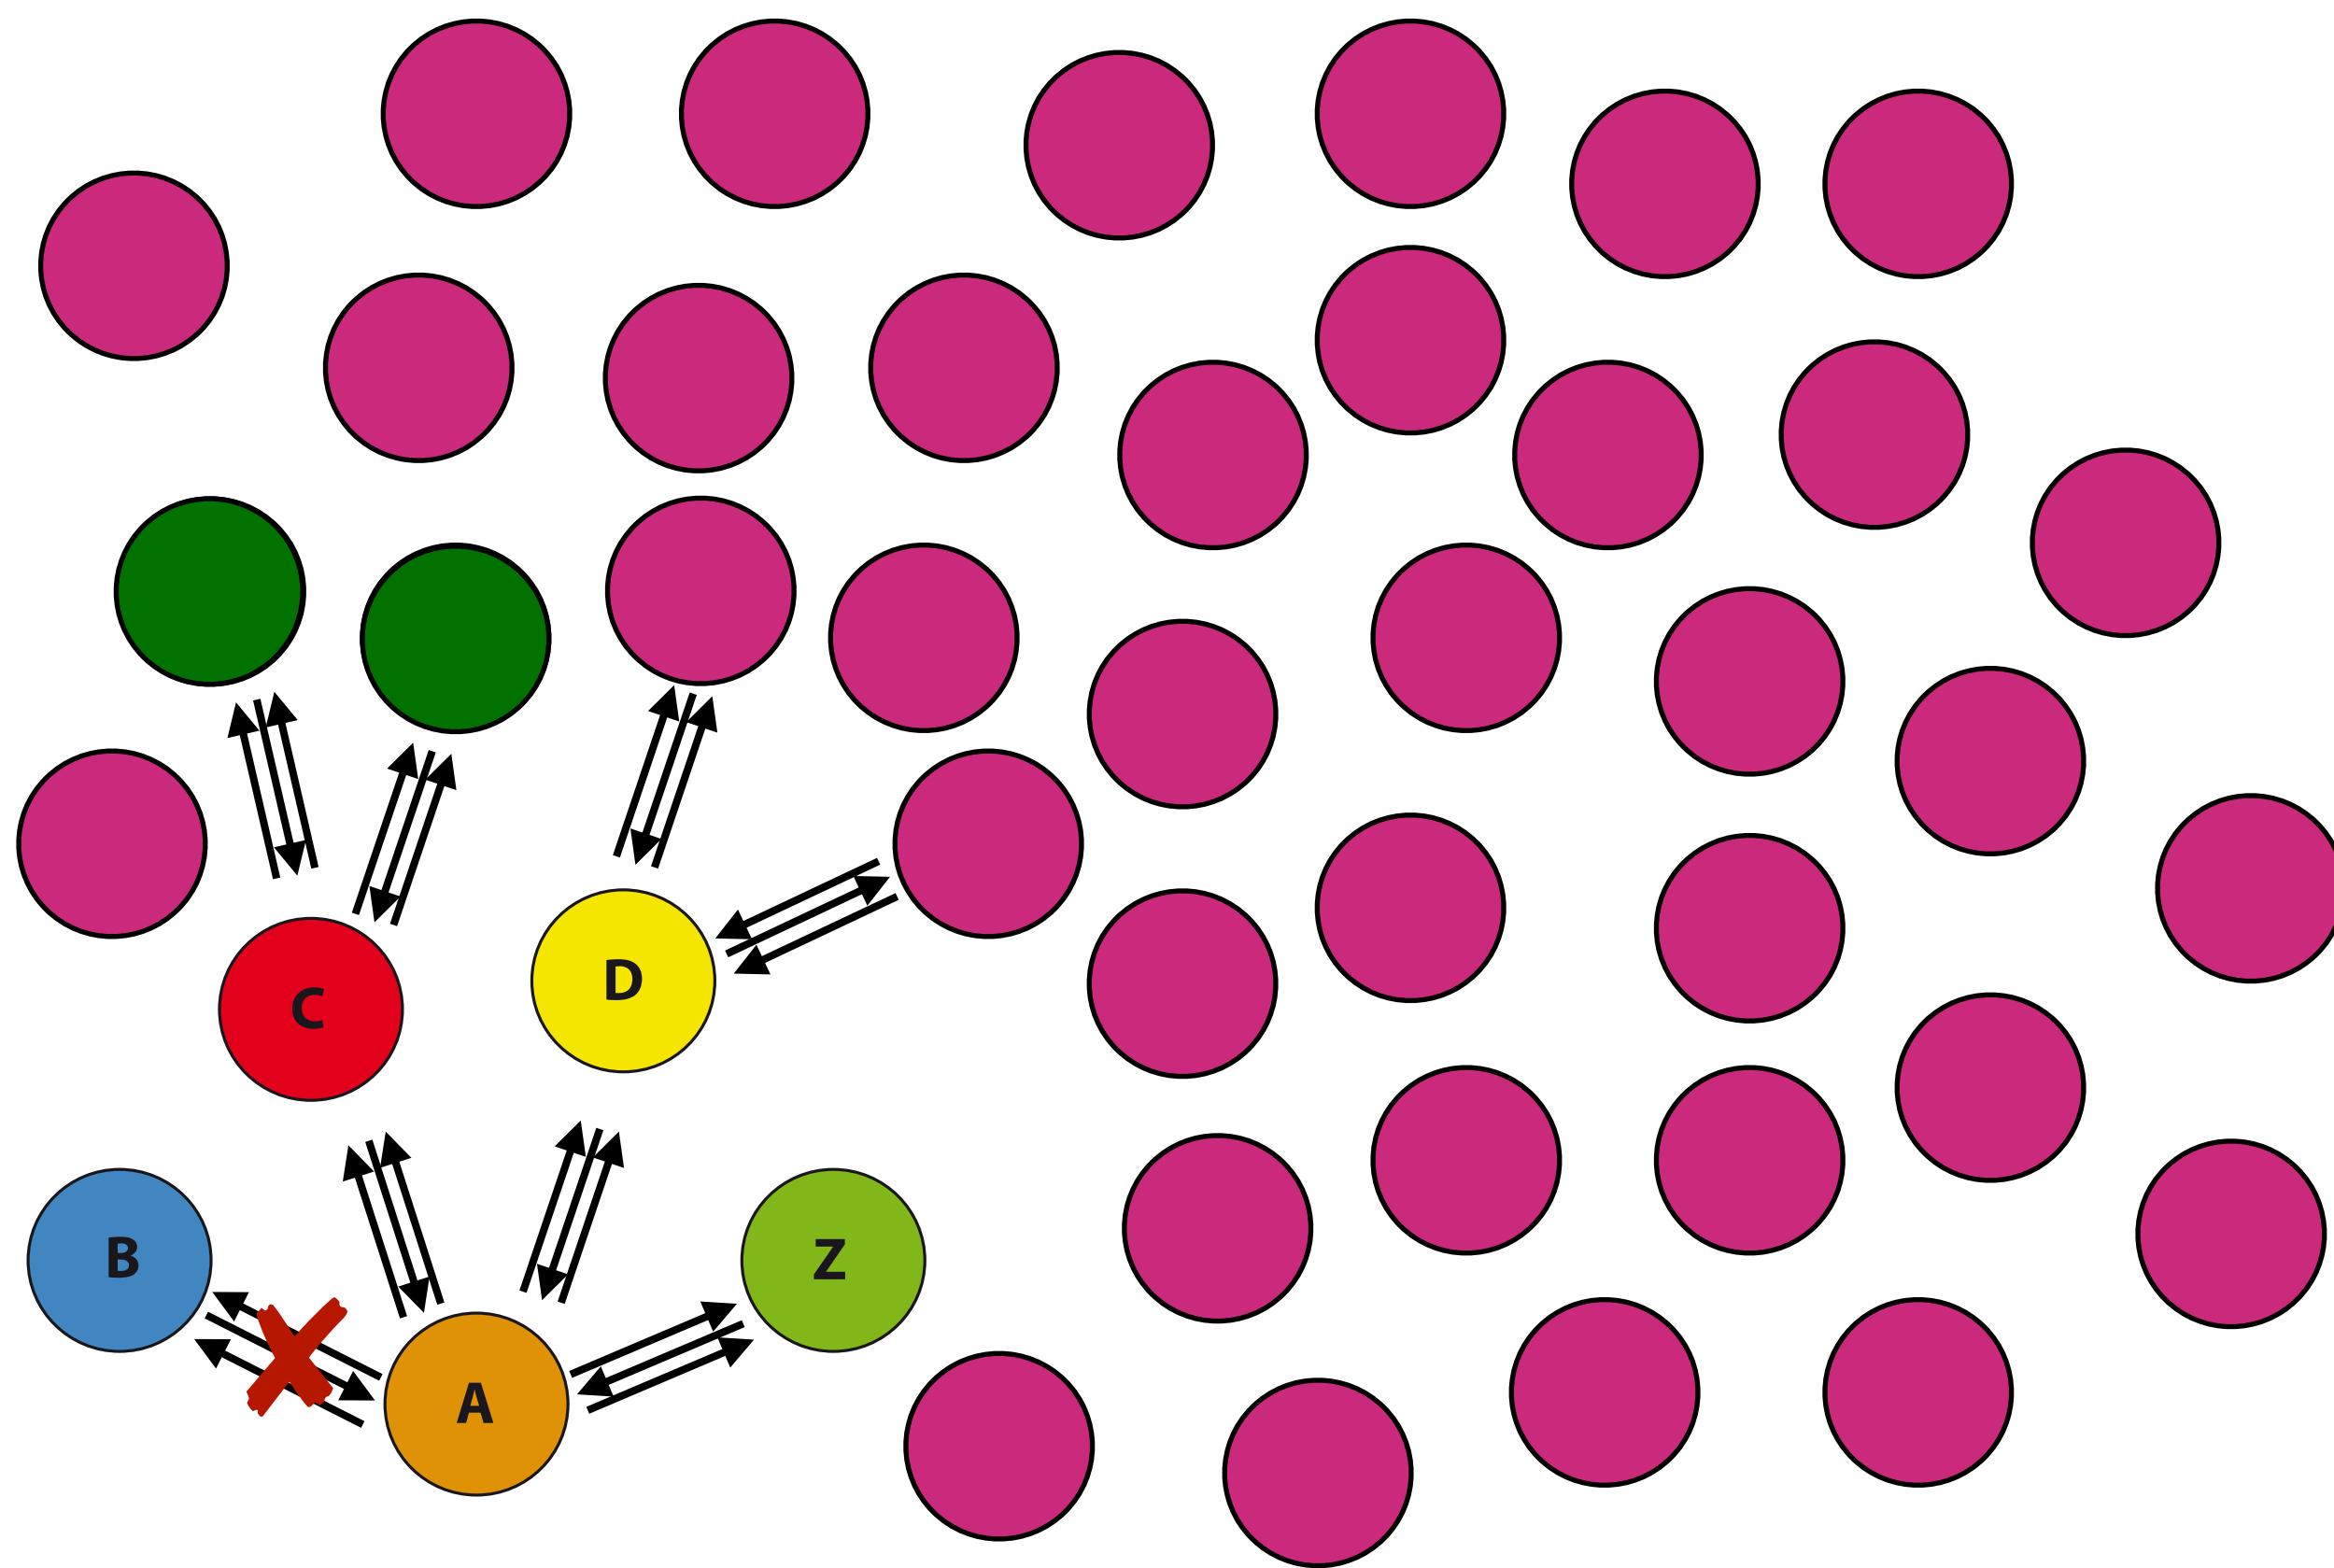
INFORMATION PROPAGATION (3/3)



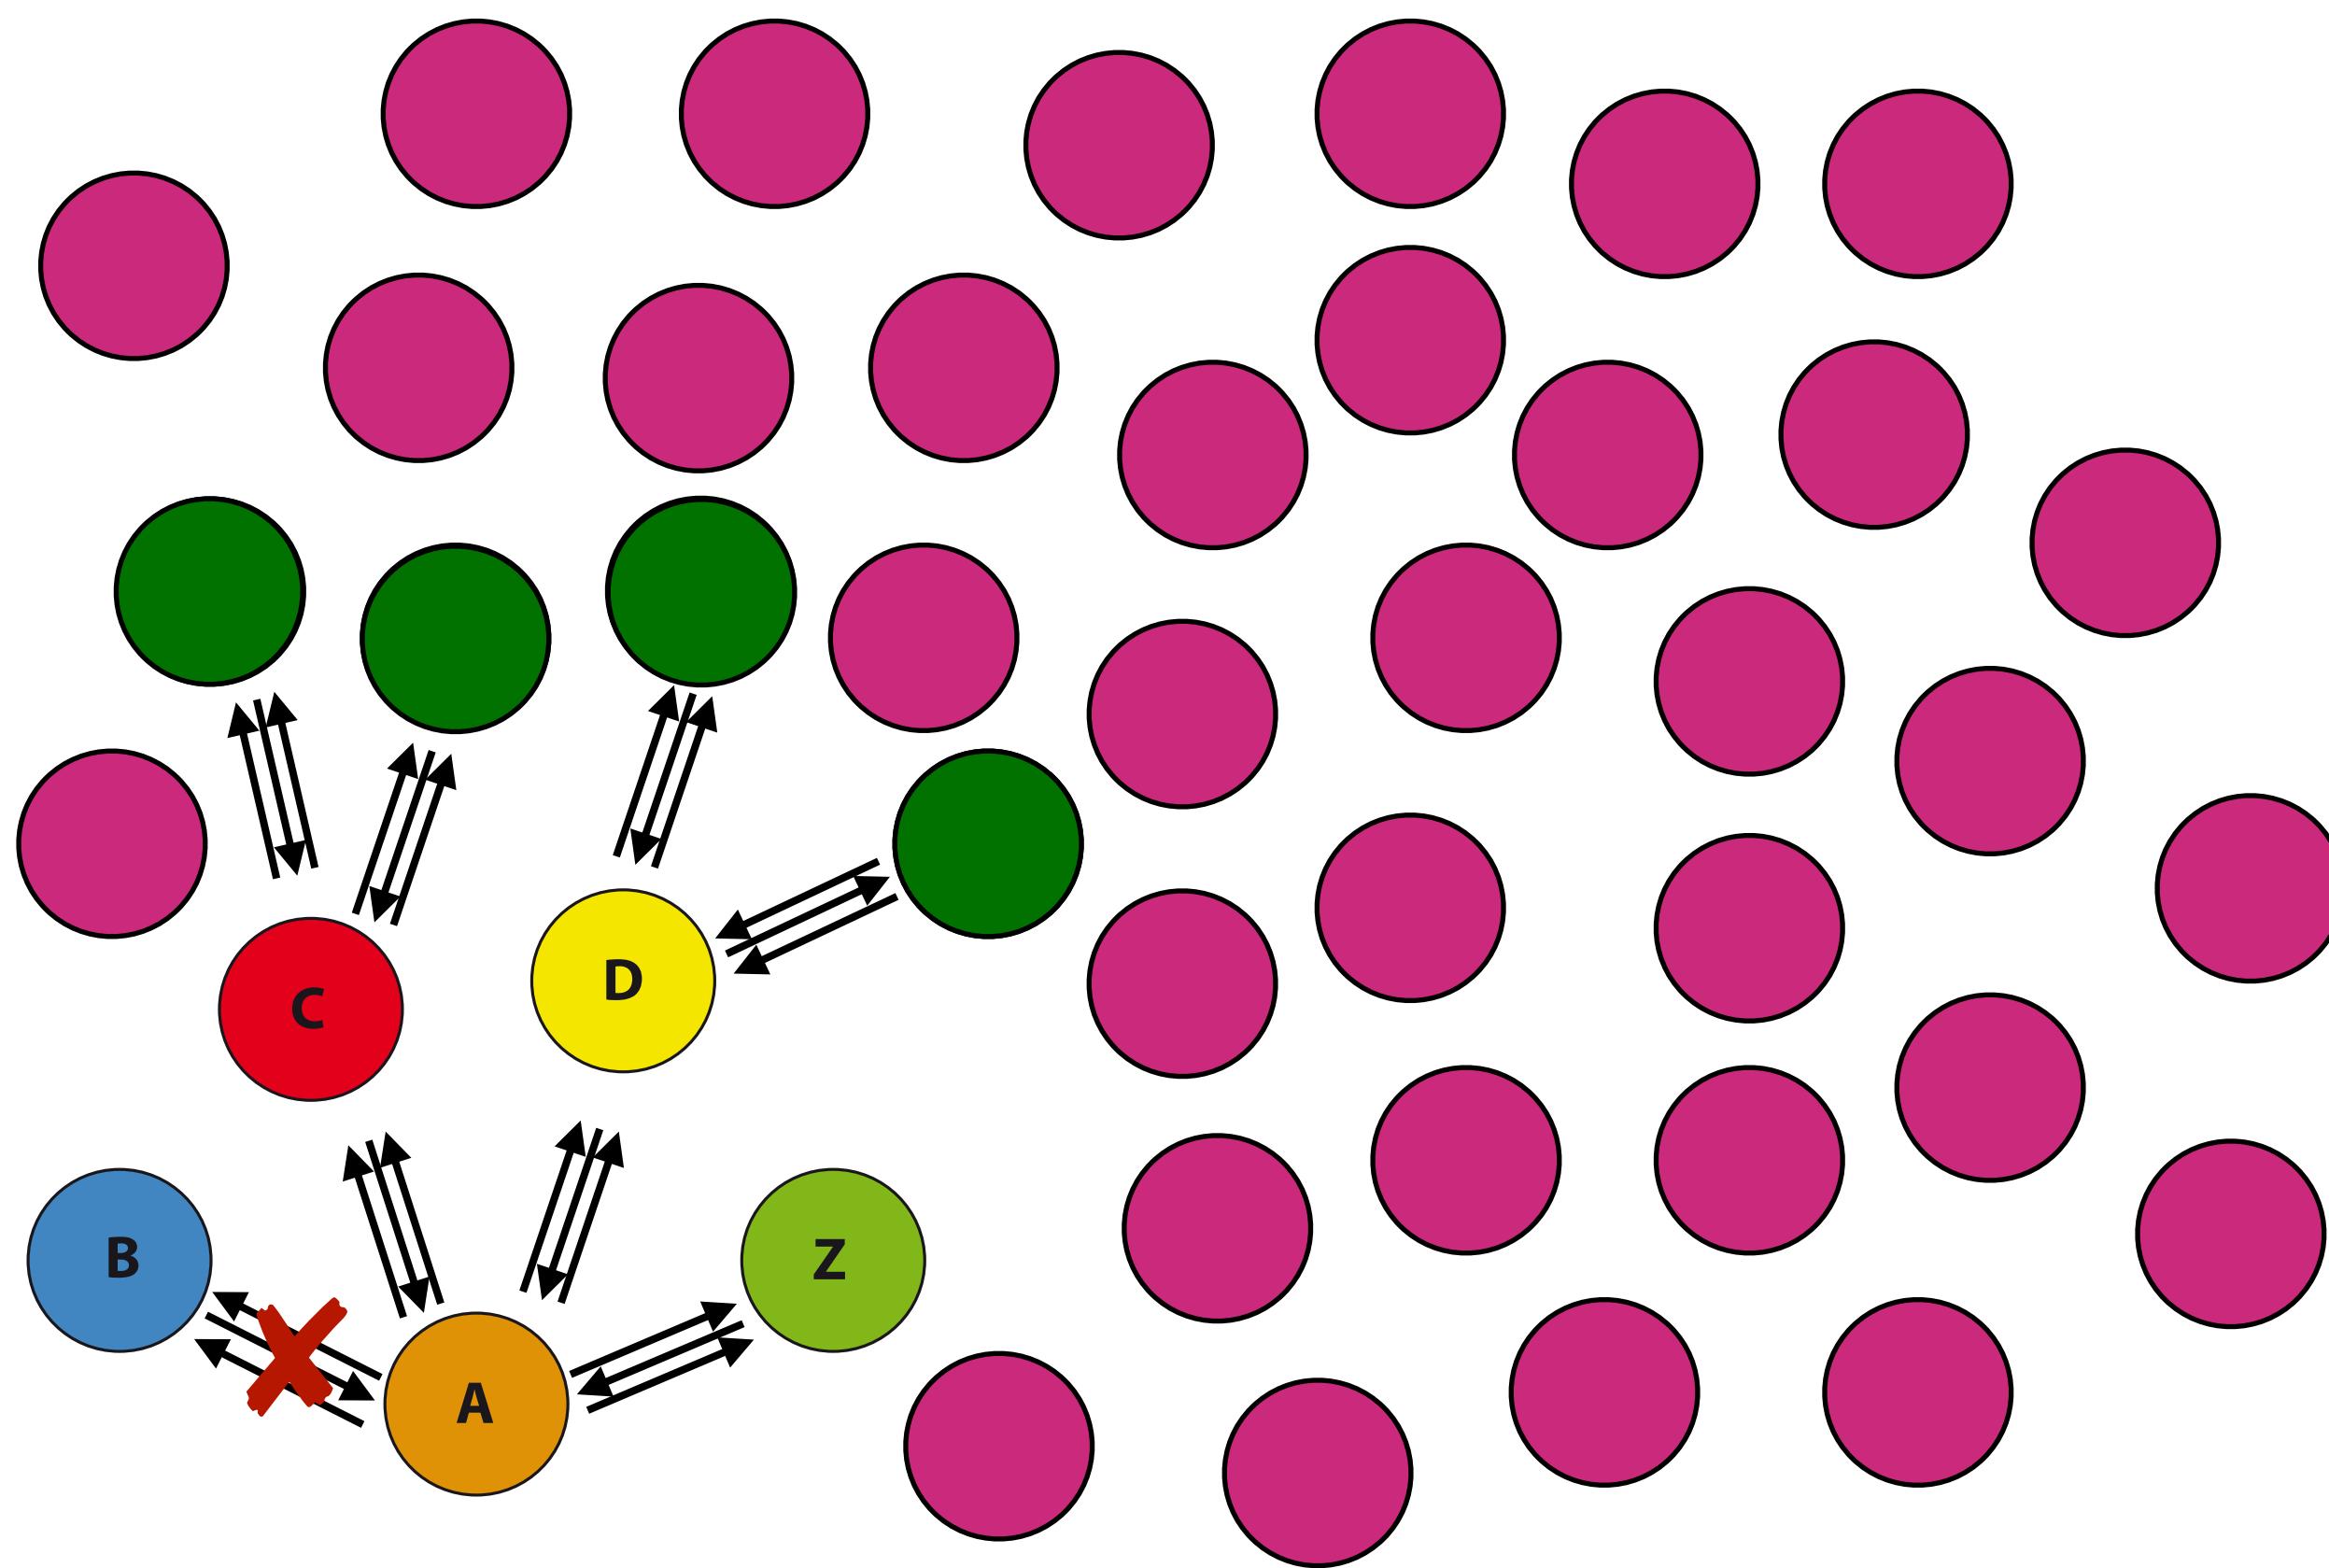
INFORMATION PROPAGATION (3/3)



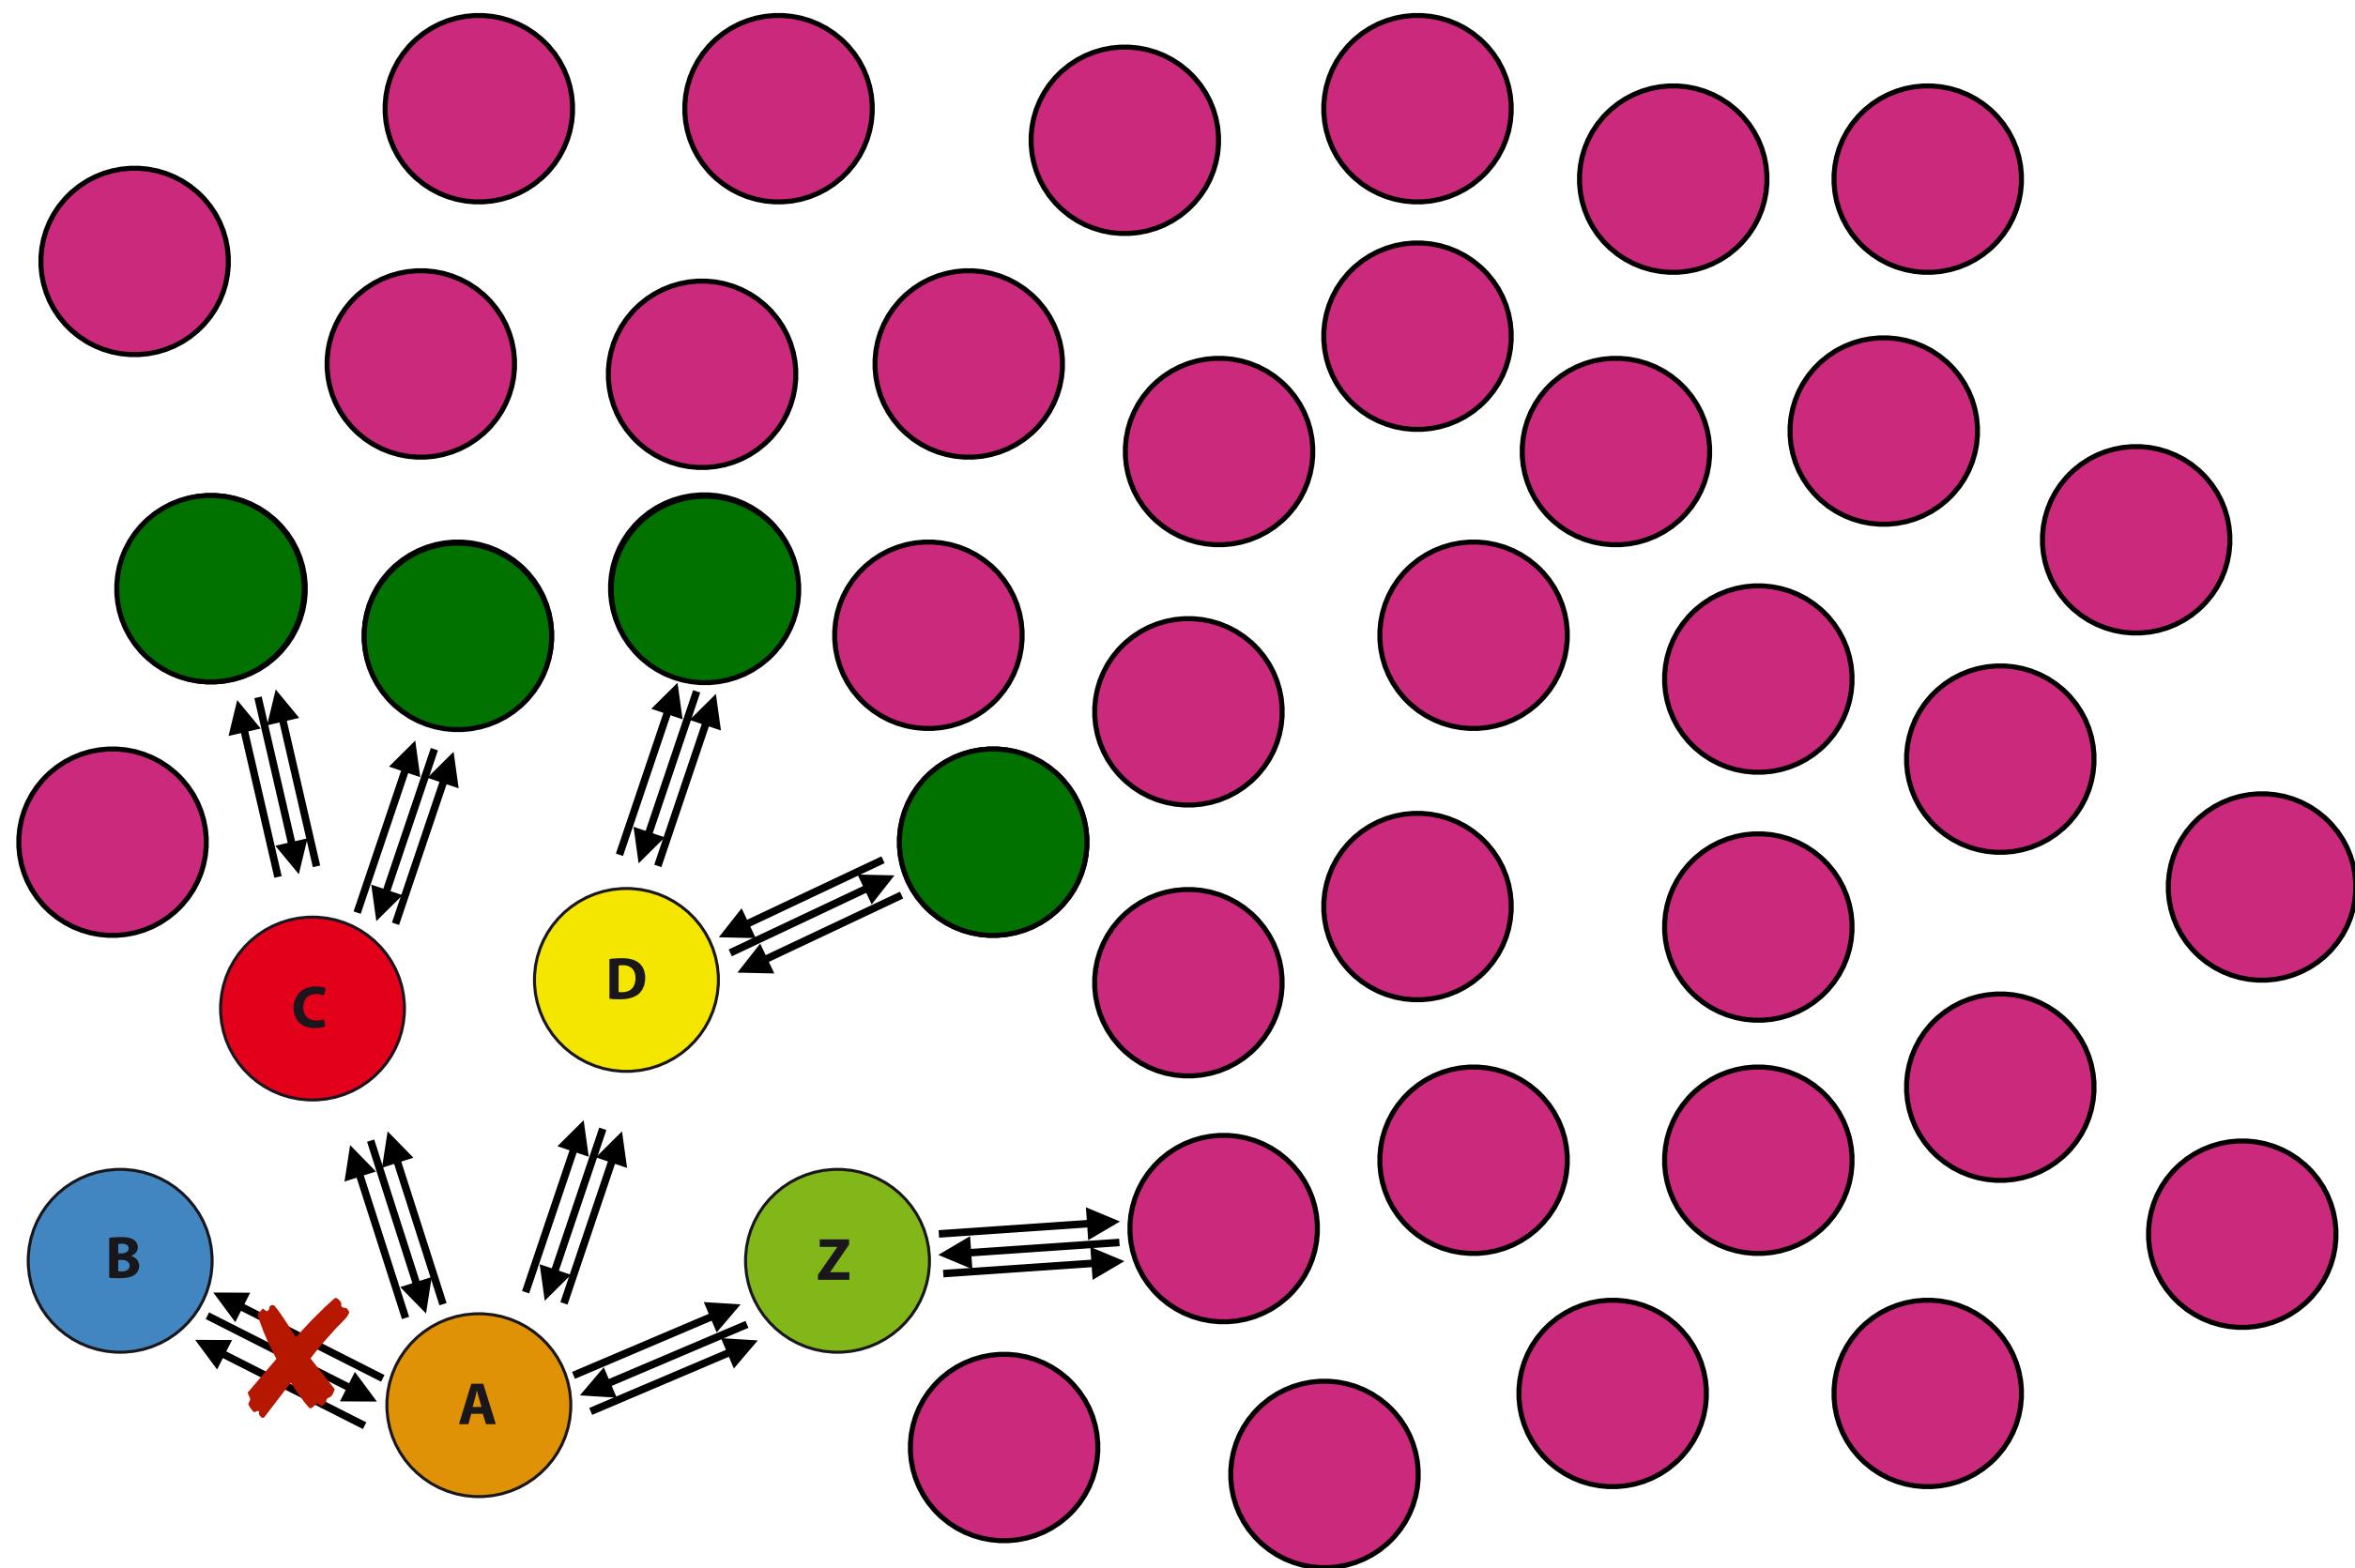
INFORMATION PROPAGATION (3/3)



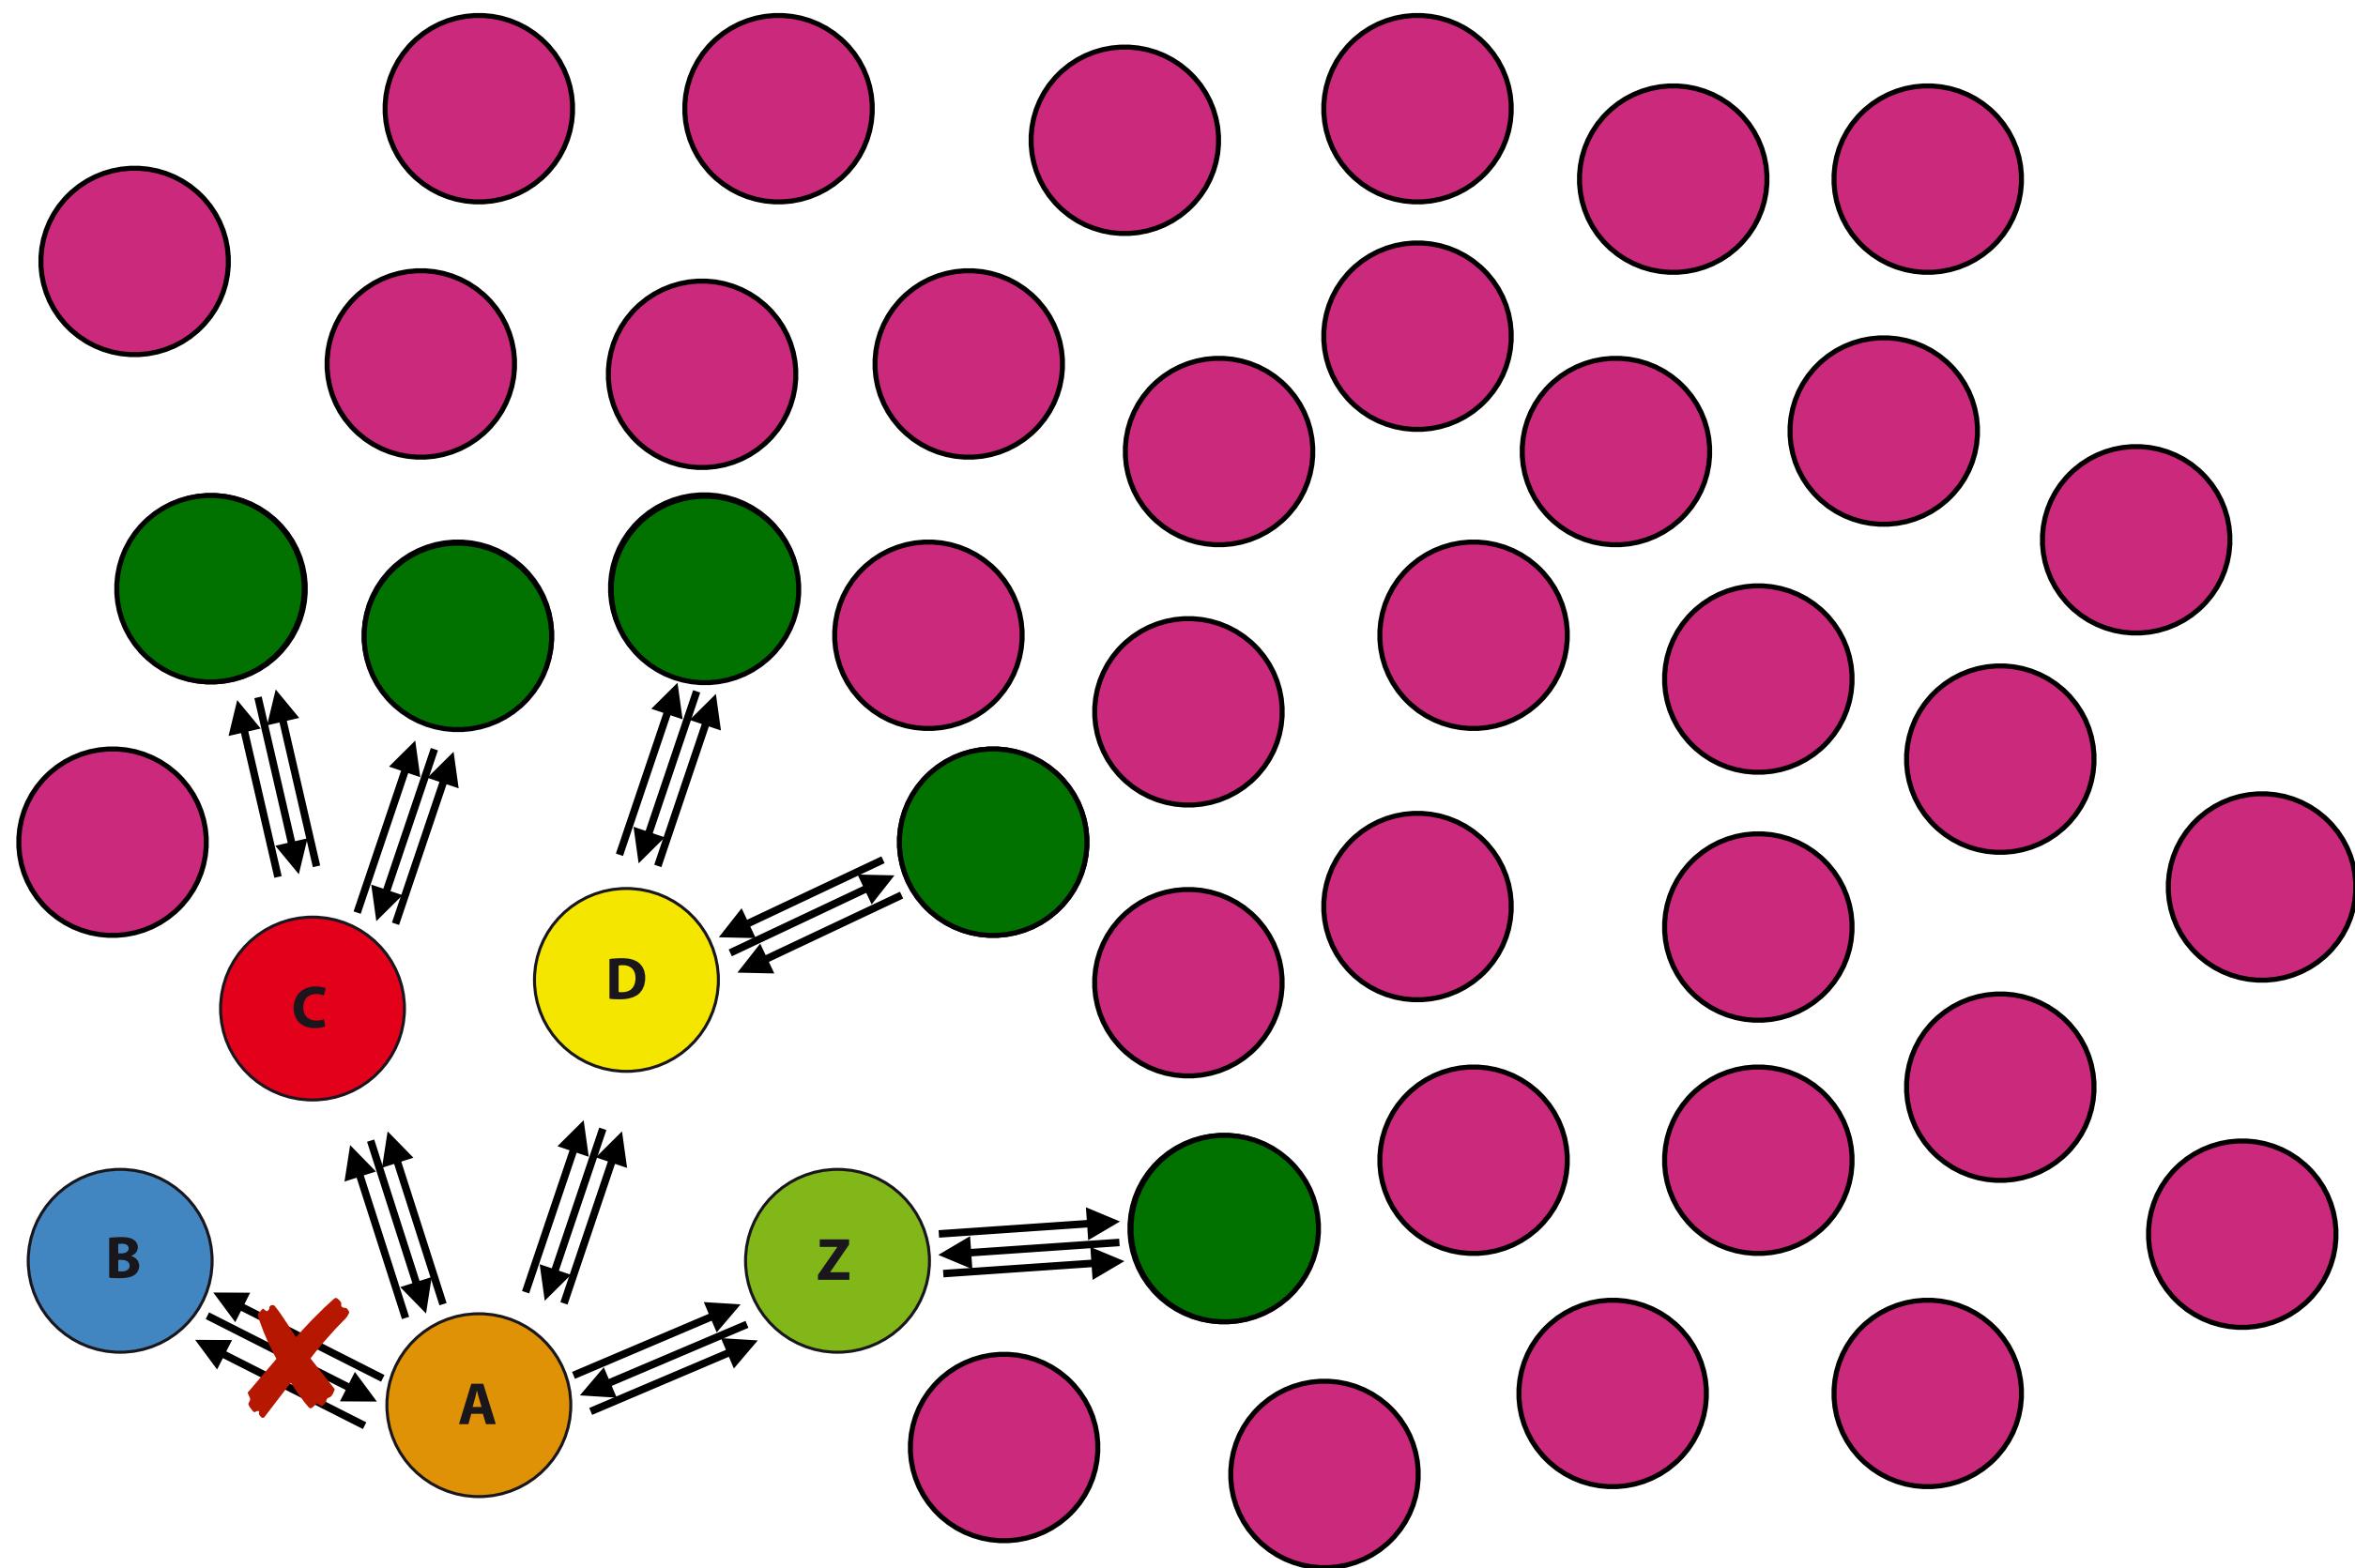
INFORMATION PROPAGATION (3/3)



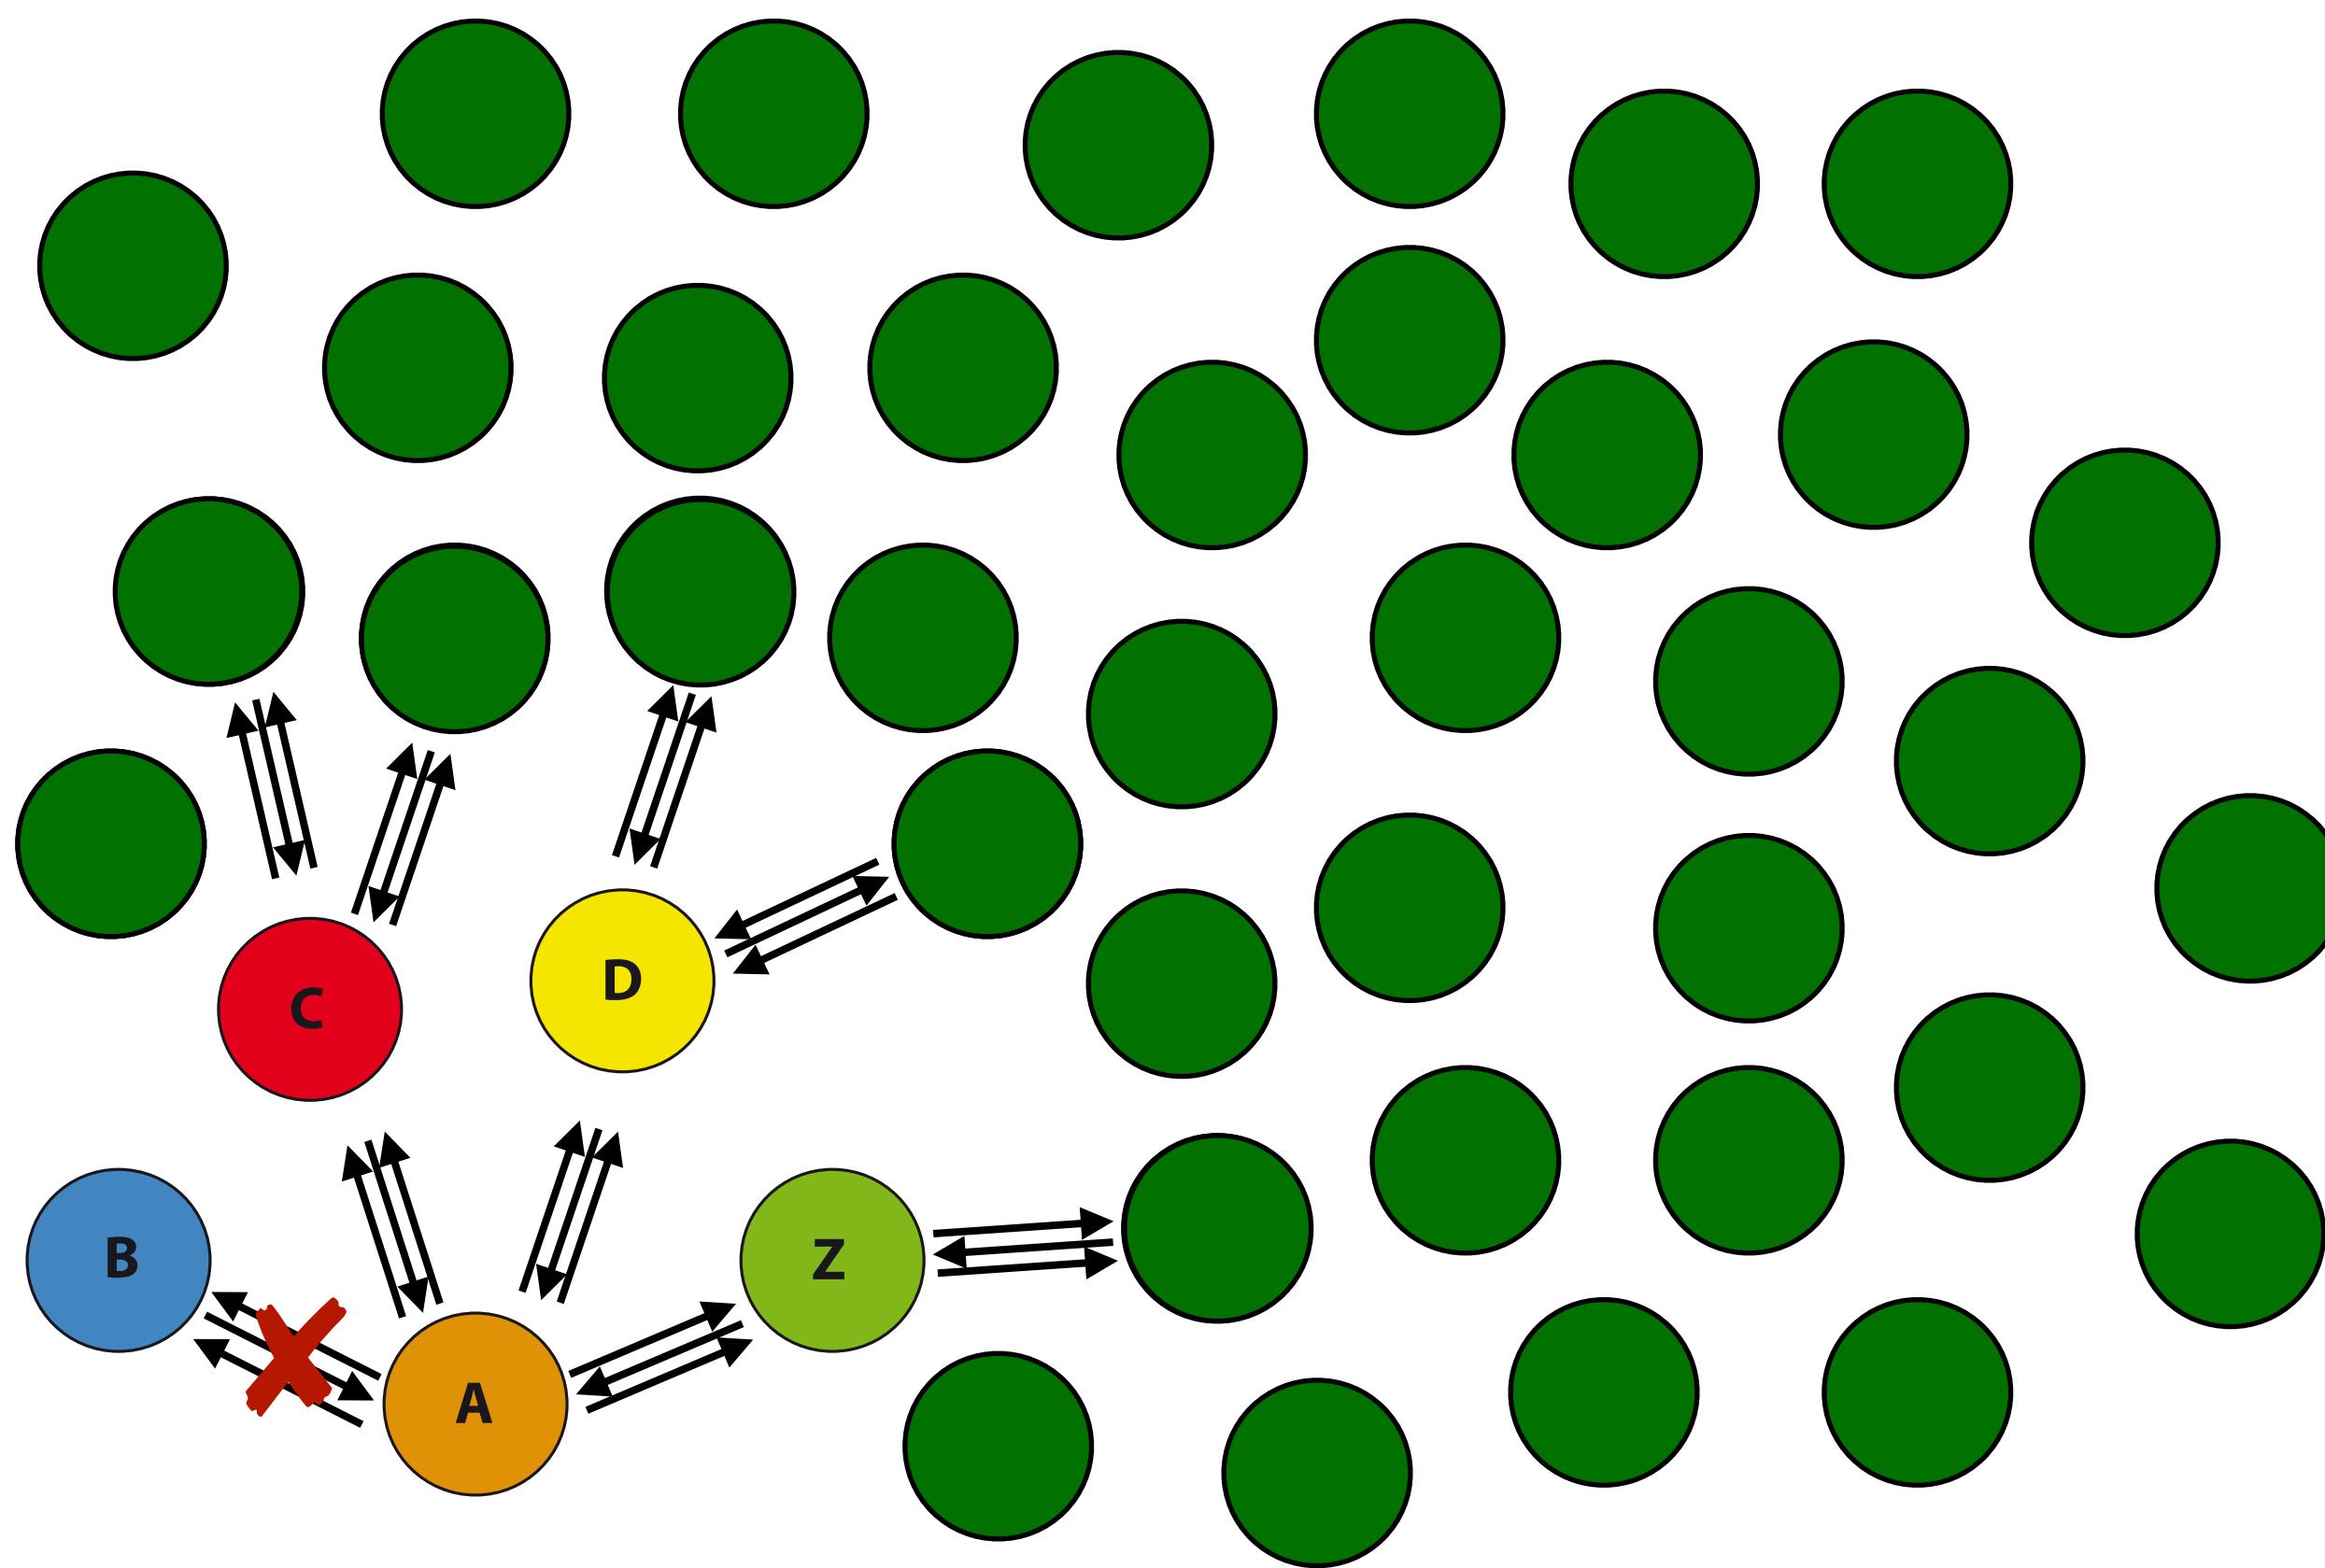
INFORMATION PROPAGATION (3/3)



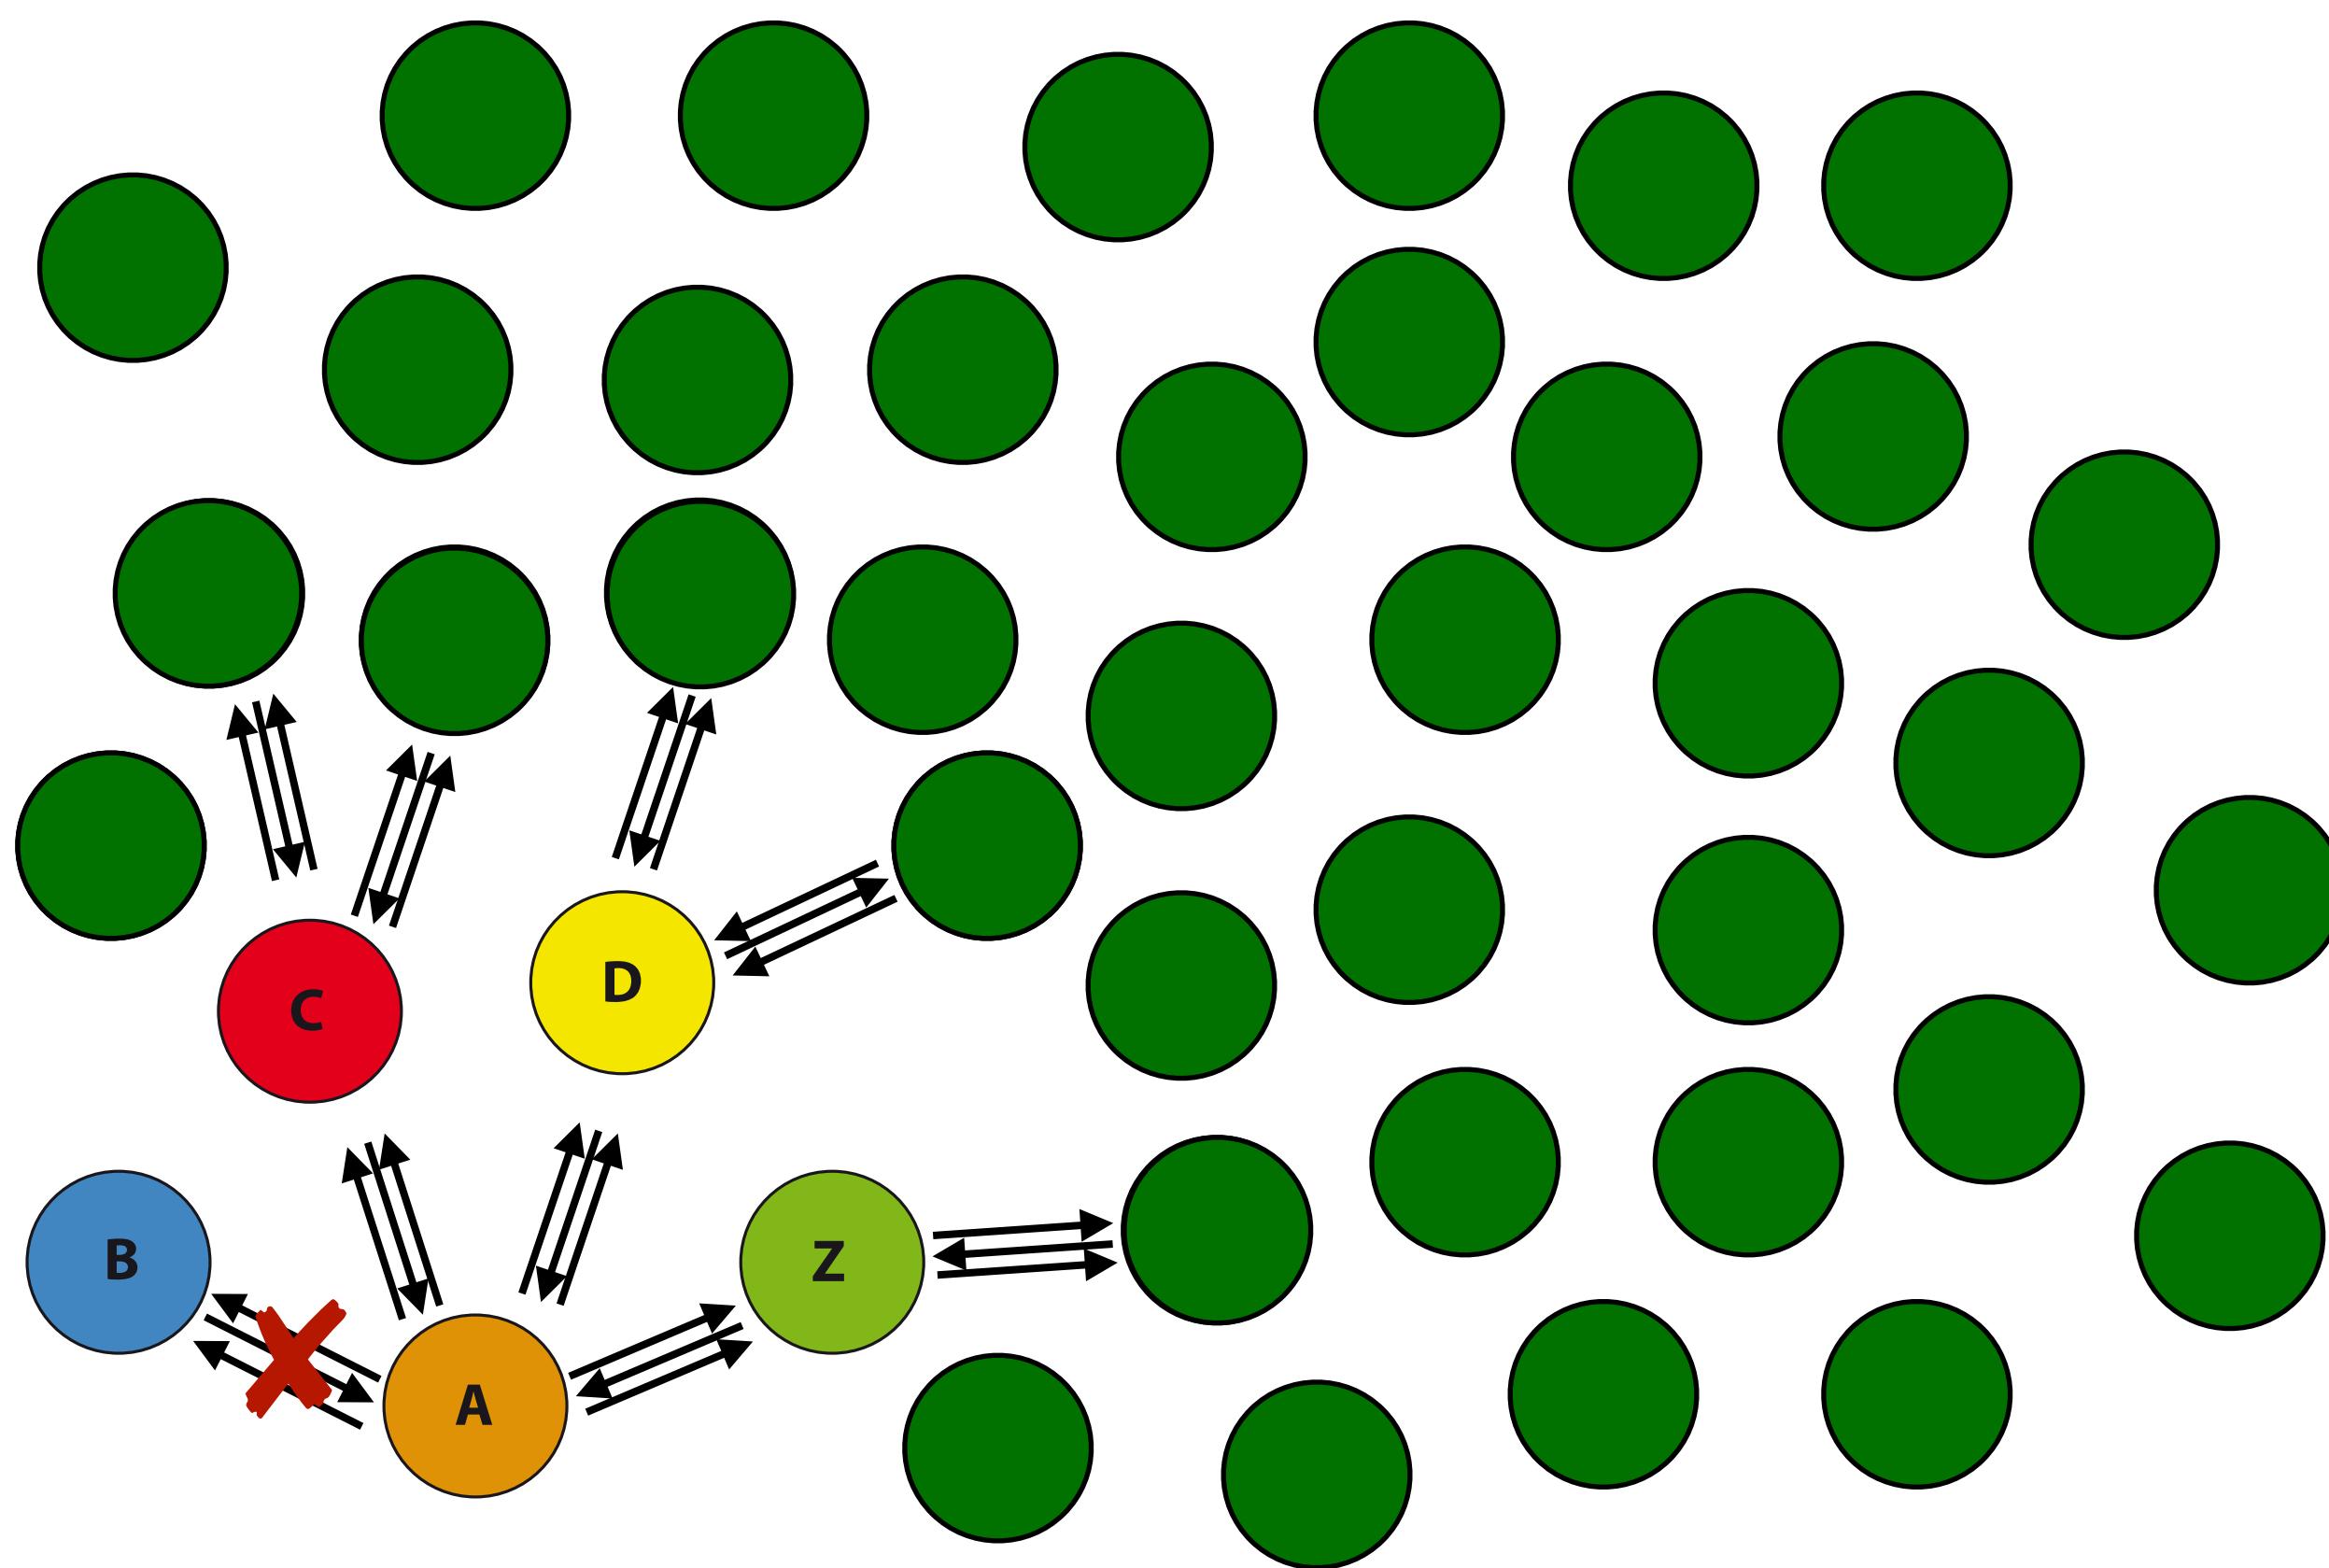
INFORMATION PROPAGATION (3/3)



INFORMATION PROPAGATION (3/3)

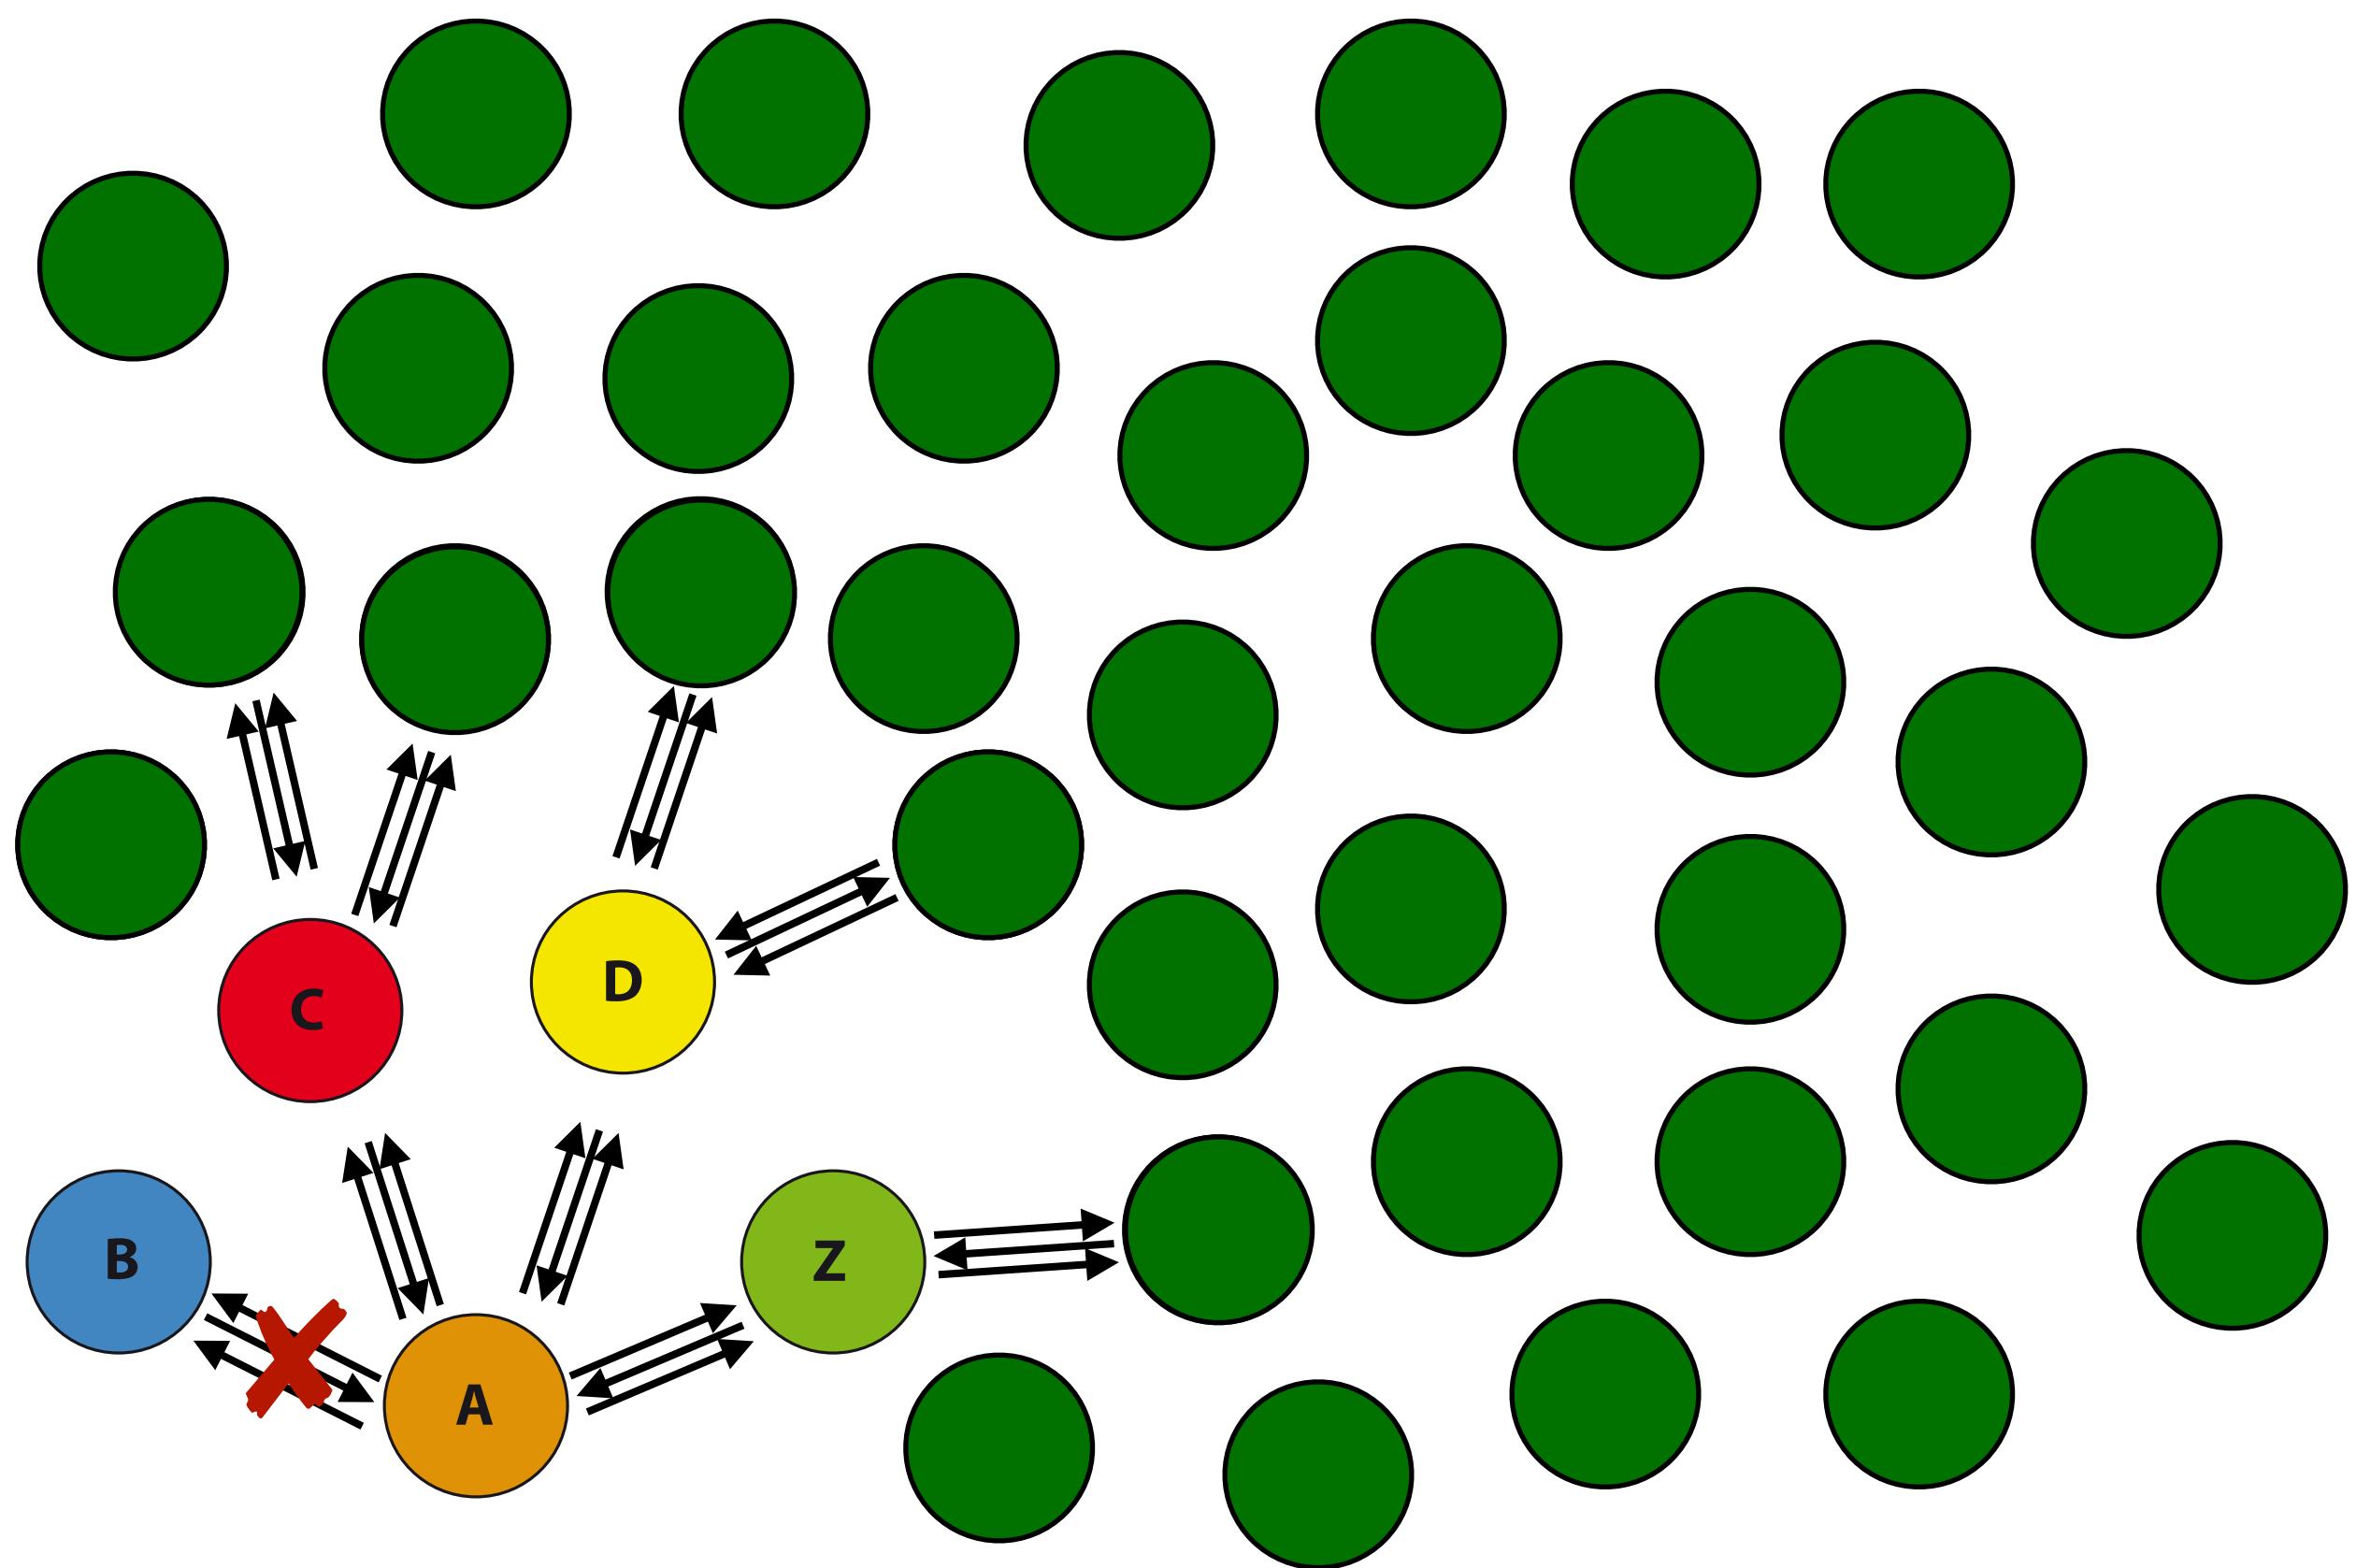


INFORMATION PROPAGATION (3/3)



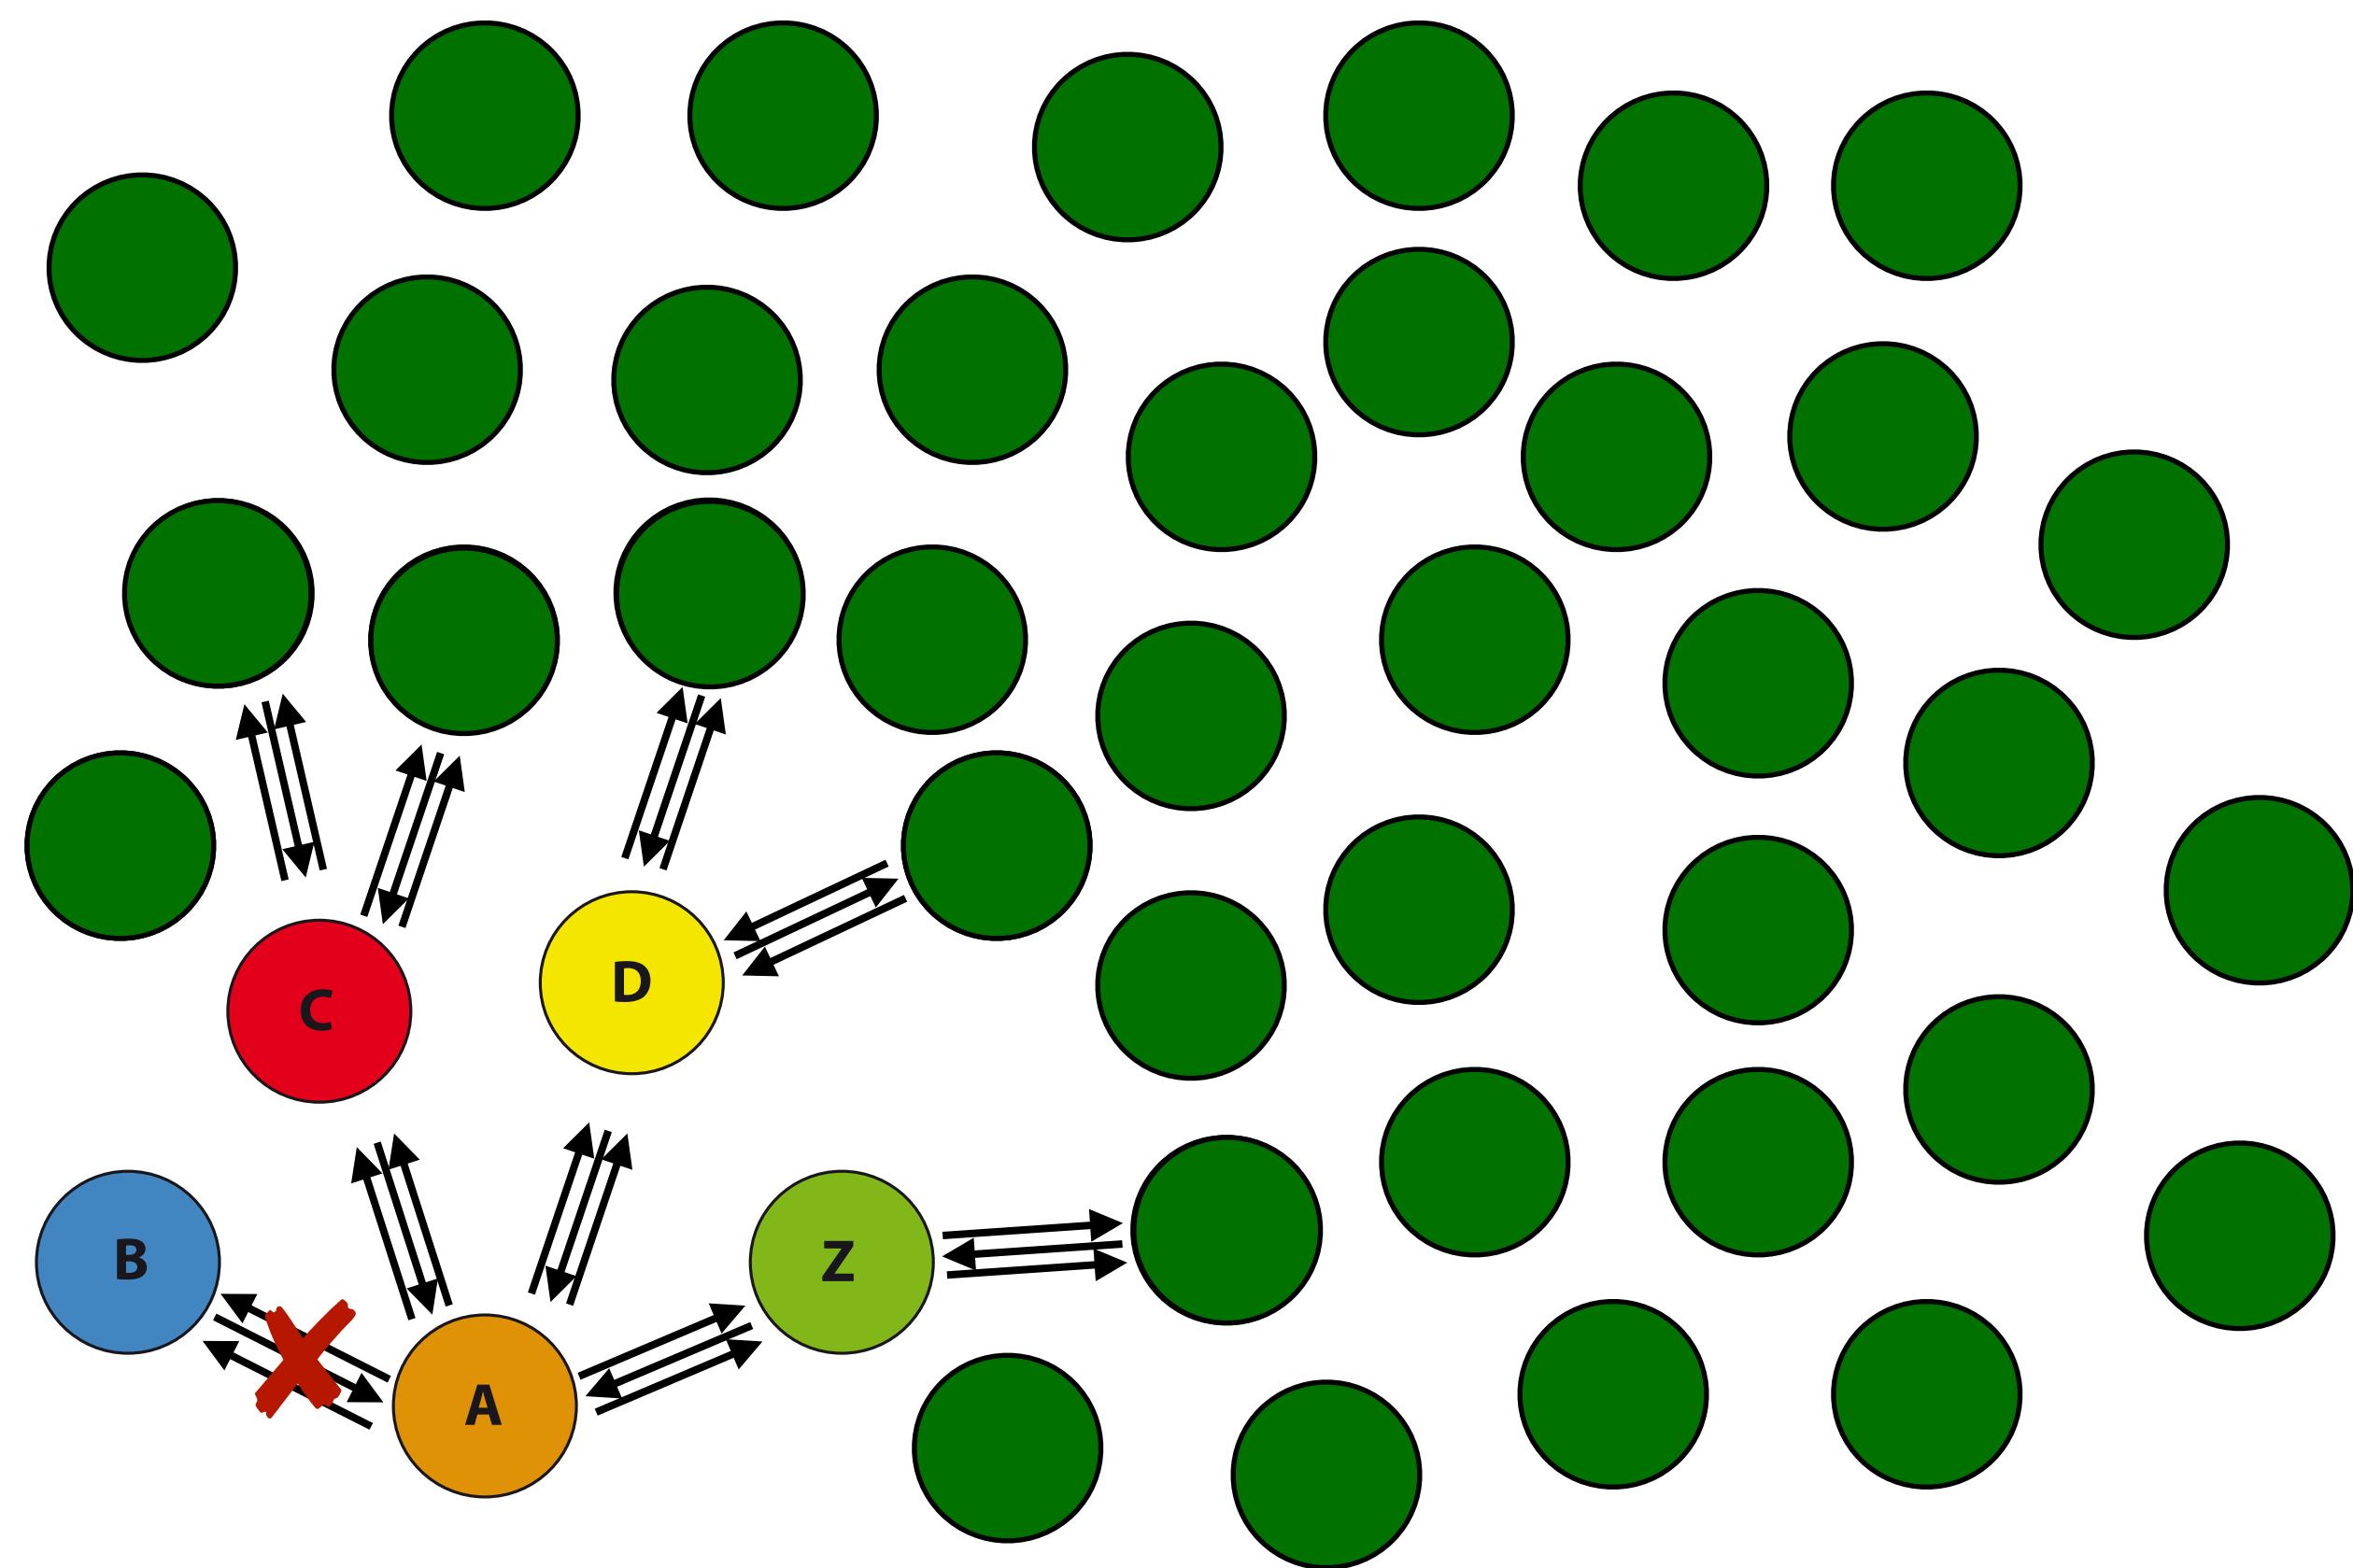
- And so on and so forth until all the nodes are reached

INFORMATION PROPAGATION (3/3)



- And so on and so forth until all the nodes are reached
- Recall that a node will reject a transaction if it has already learnt about it from any of its neighbors

INFORMATION PROPAGATION (3/3)



- And so on and so forth until all the nodes are reached
- Recall that a node will reject a transaction if it has already learnt about it from any of its neighbors
- The same procedure applies for blocks

IMPLICATIONS

The bigger the network the more it takes for an item to propagate (**this can be counterintuitive**)

Long propagation times (**for blocks**) imply bigger likelihood of forking the blockchain

IMPLICATIONS

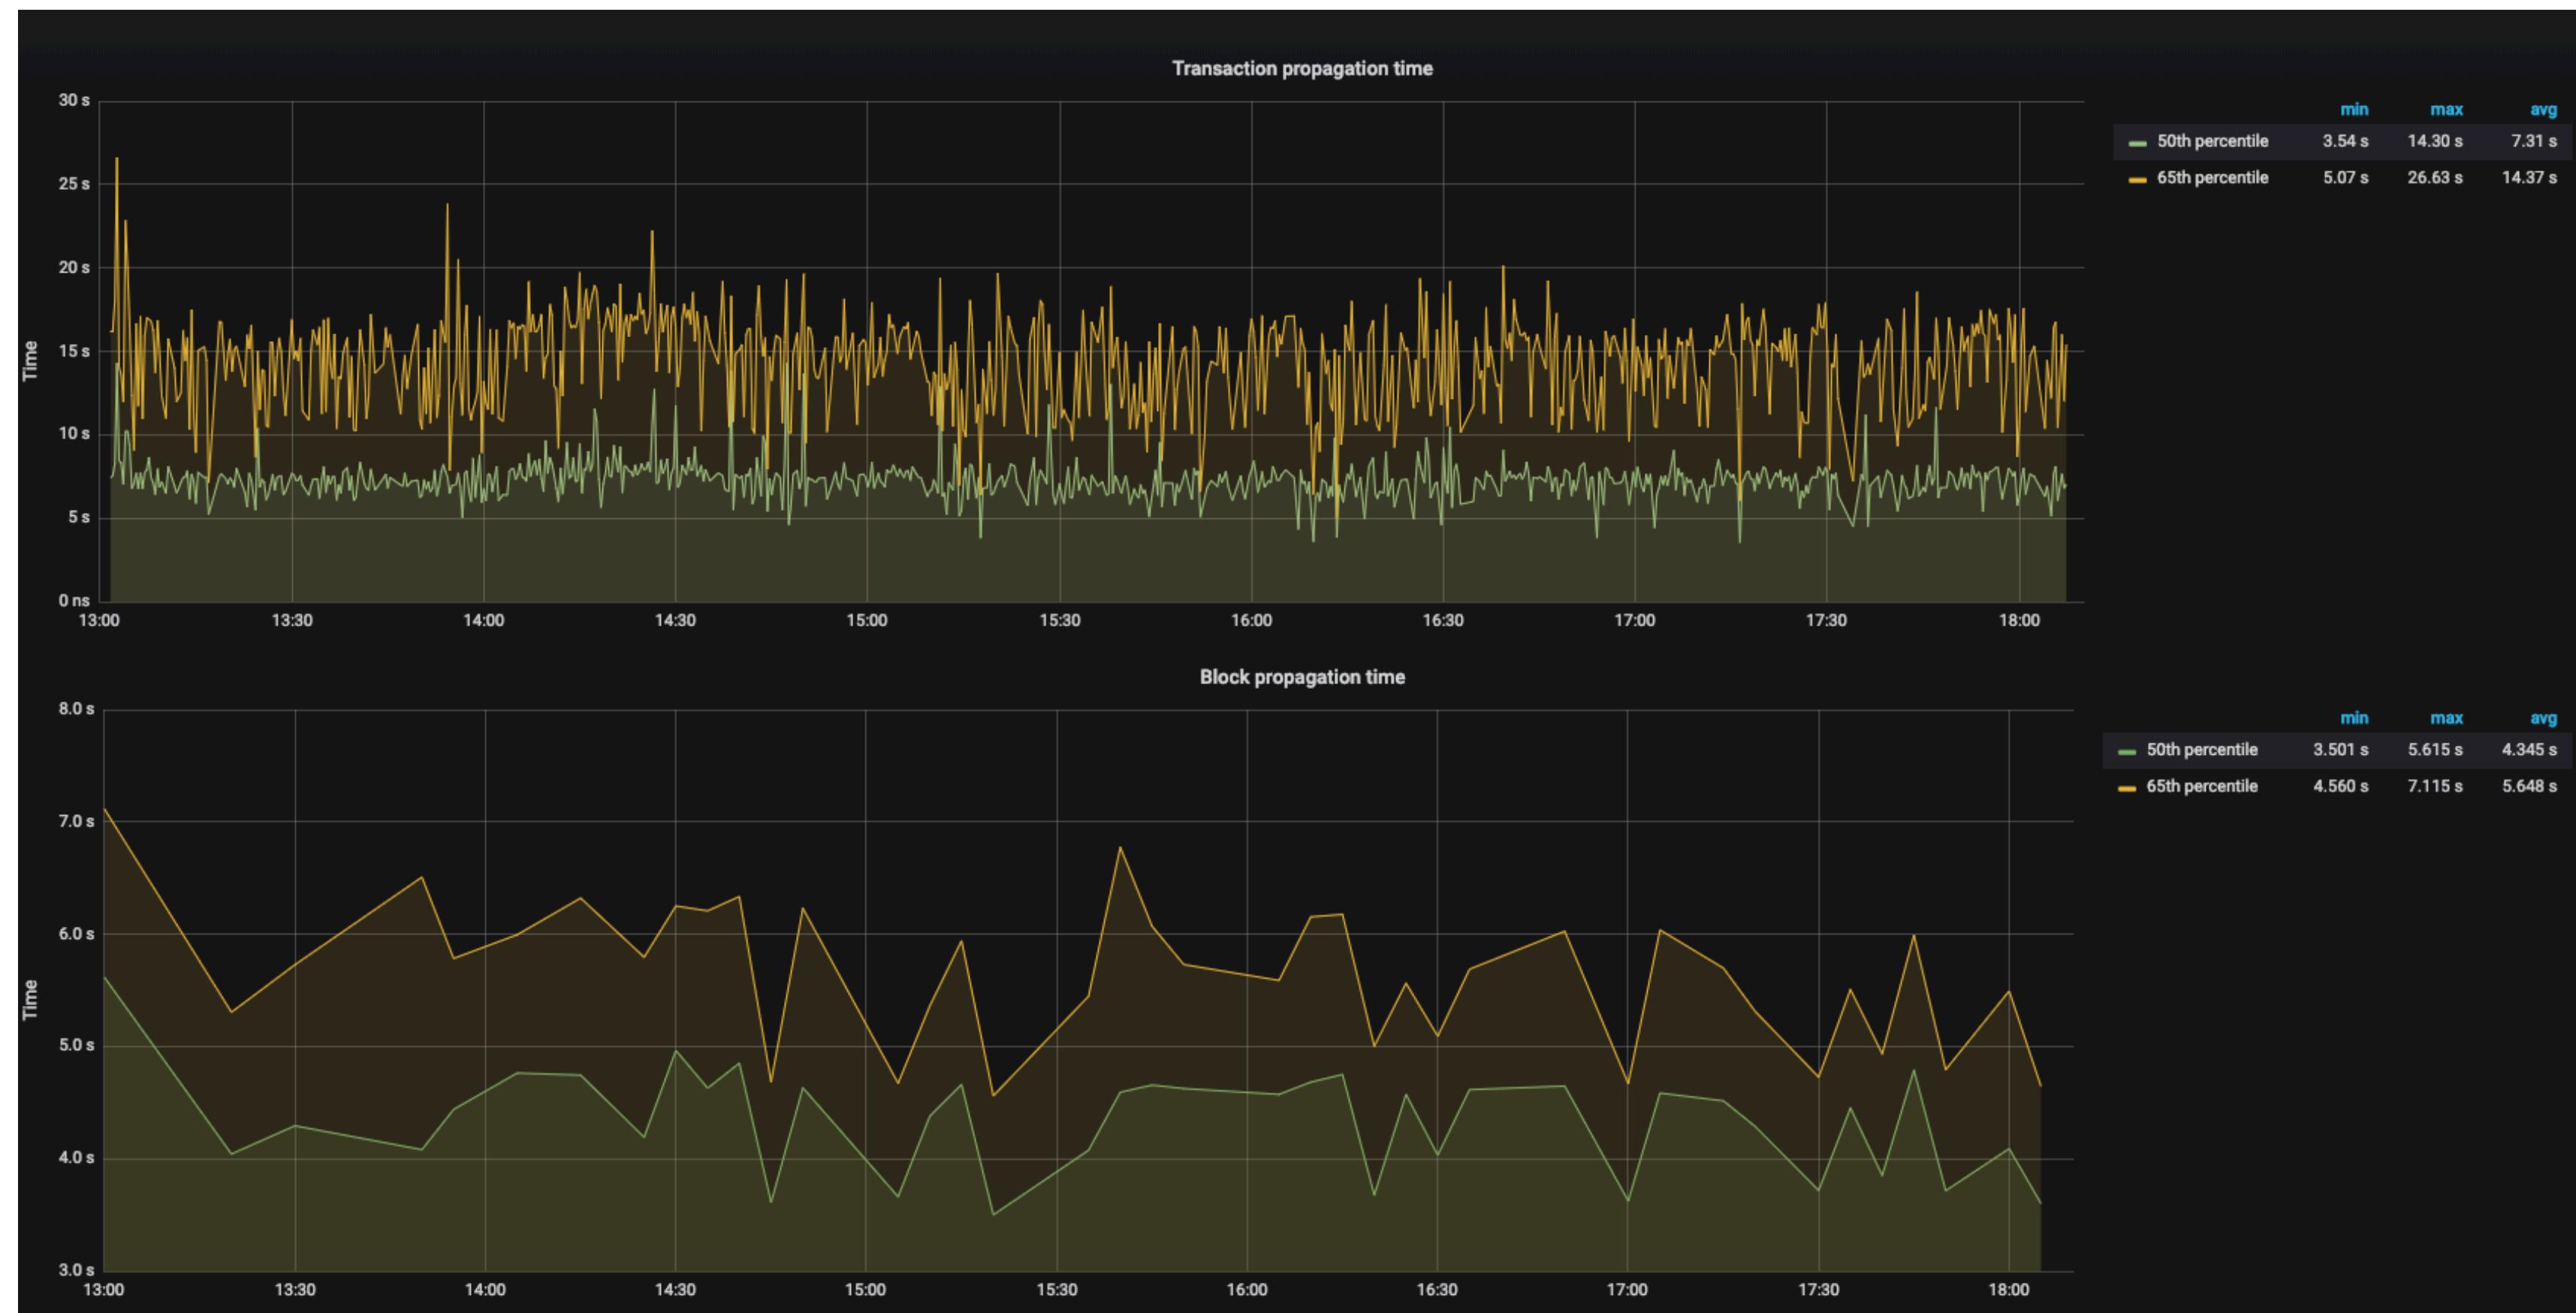
The bigger the network the more it takes for an item to propagate (**this can be counterintuitive**)

Long propagation times (**for blocks**) imply bigger likelihood of forking the blockchain



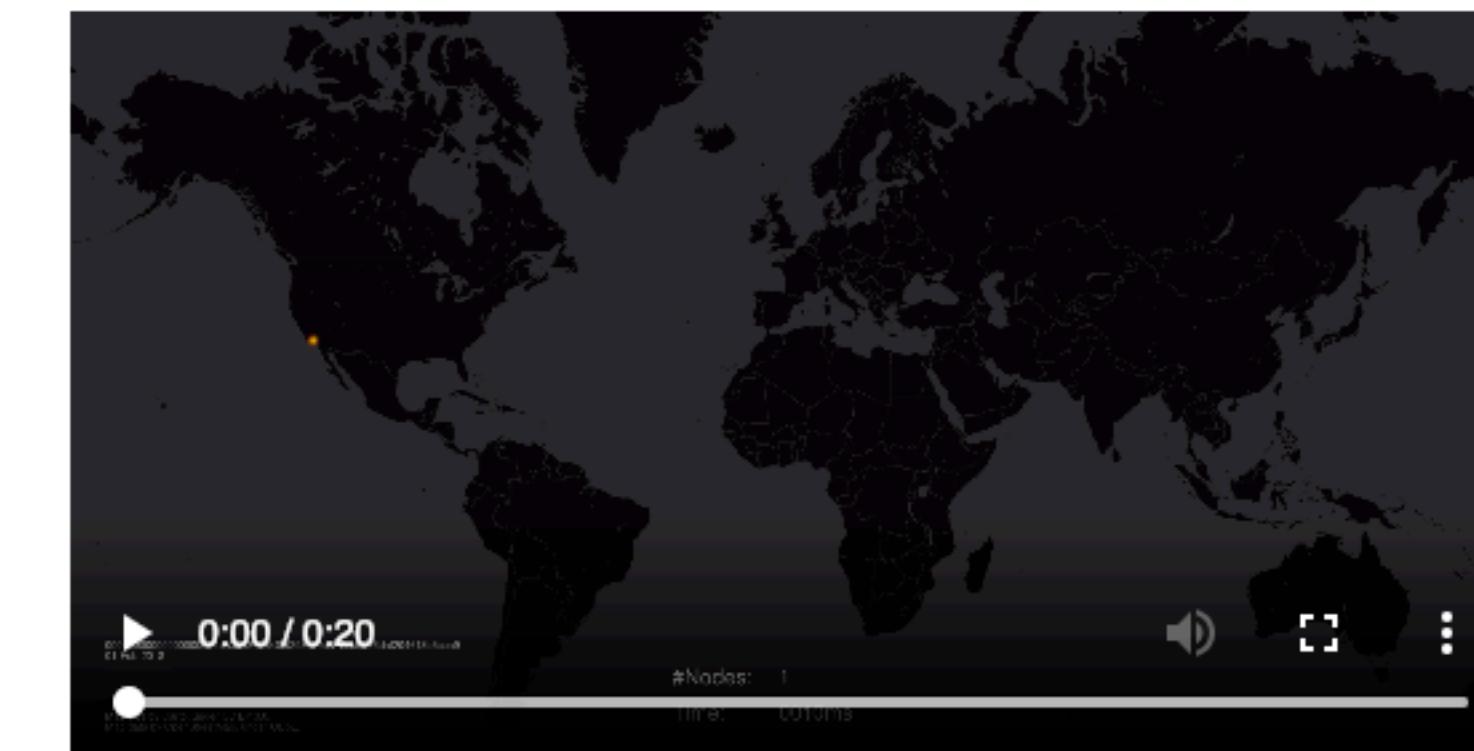
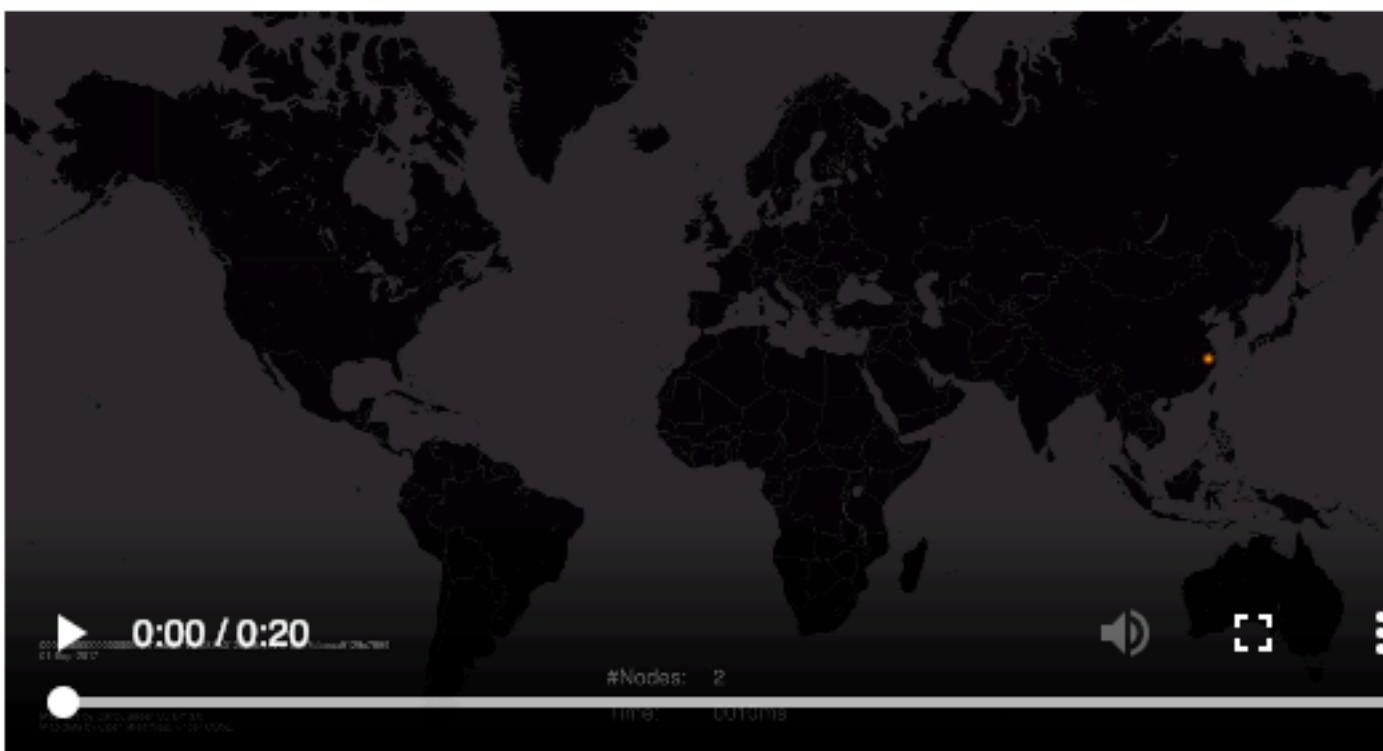
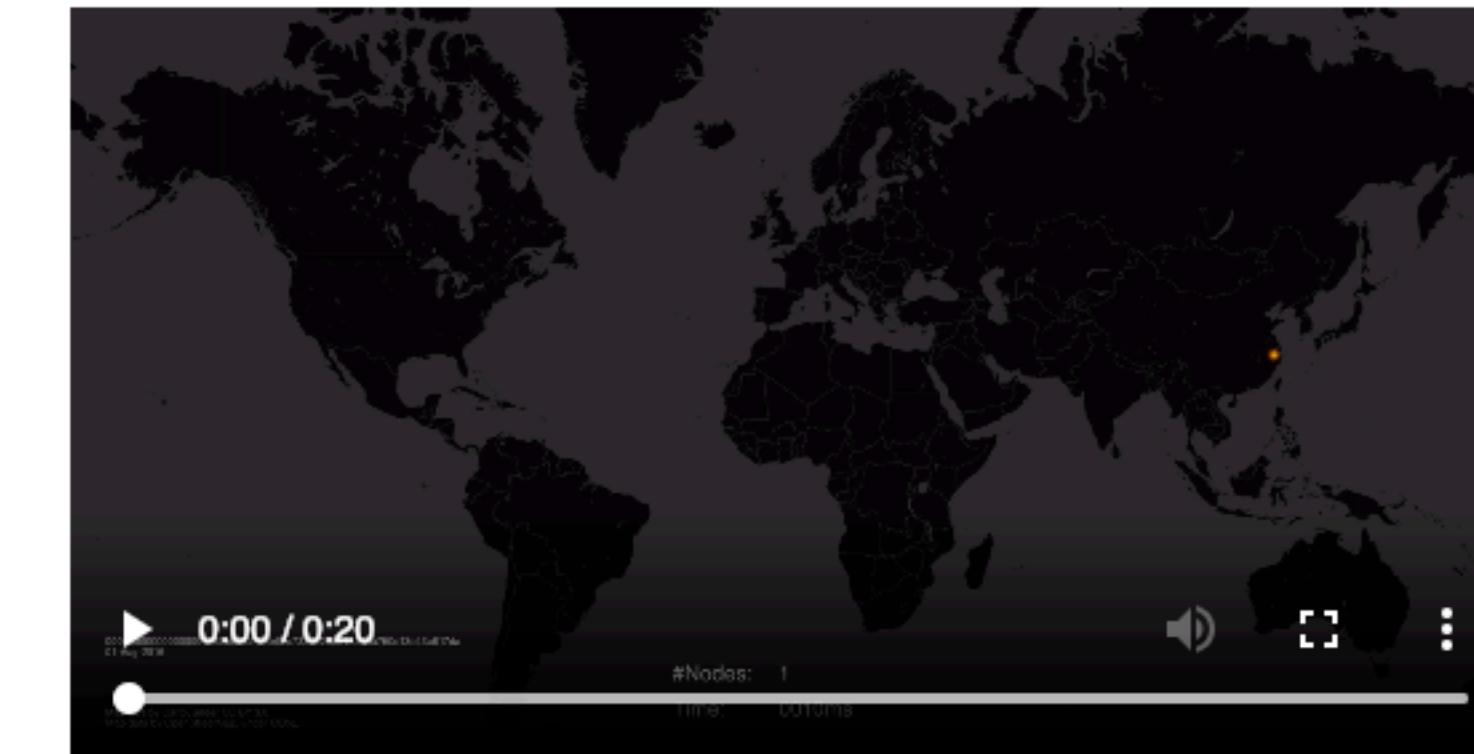
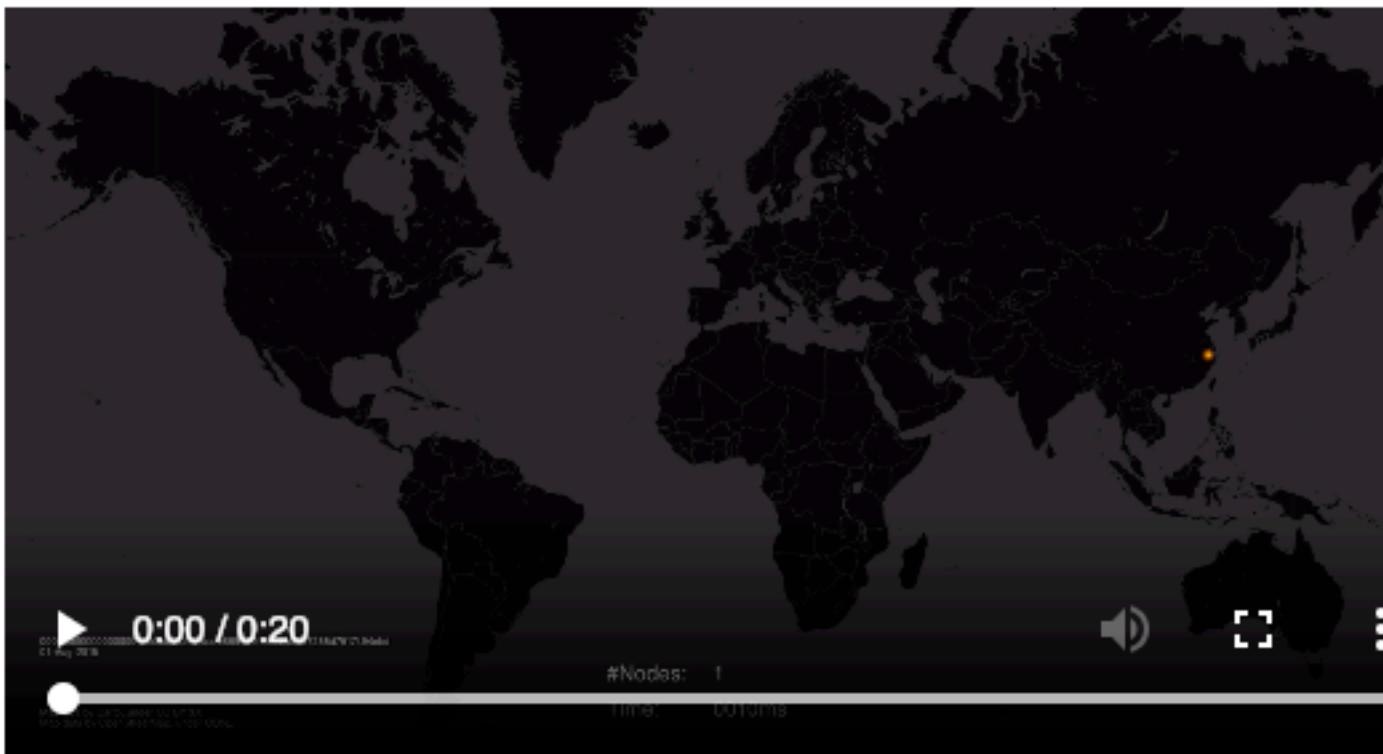
Christian Decker and Roger Wattenhofer
Information propagation in the Bitcoin network
<https://ieeexplore.ieee.org/document/6688704>

DATA PROPAGATION TIMES (TESTNET)



source: charts.satoshi.uab.cat

MORE ABOUT PROPAGATION TIMES



source: <https://dsn.tm.kit.edu/bitcoin/videos.html>

PROPAGATION DELAYS (1/2)

How can blocks propagate faster than transactions if the former are bigger than the later?

PROPAGATION DELAYS (1/2)

How can blocks propagate faster than transactions if the former are bigger than the later?

- Transactions are accumulated in buffers and forwarded in batches to break the link between first relayer and origin of a transaction

PROPAGATION DELAYS (1/2)

How can blocks propagate faster than transactions if the former are bigger than the later?

- Transactions are accumulated in buffers and forwarded in batches to break the link between first relayer and origin of a transaction
- The propagation of blocks is not delayed, in order to reach full network coverage as soon as possible

PROPAGATION DELAYS (2/2)

But blocks are way bigger than transactions, how can they be propagated so fast!?

PROPAGATION DELAYS (2/2)

But blocks are way bigger than transactions, how can they be propagated so fast!?

- Fast relay networks on top of Bitcoin exists (Falcon, FIBRE, etc) to enhance the propagation time of blocks

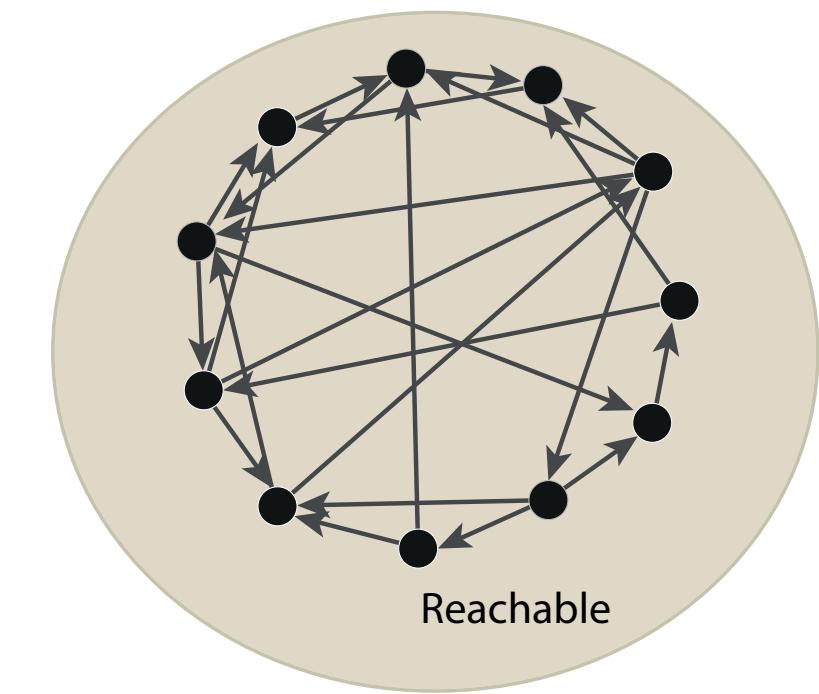
PROPAGATION DELAYS (2/2)

But blocks are way bigger than transactions, how can they be propagated so fast!?

- Fast relay networks on top of Bitcoin exists (Falcon, FIBRE, etc) to enhance the propagation time of blocks
- Miners use such networks to ensure minimal propagation times as well as ensure being mining on top of the most recent block

NETWORK TAXONOMY

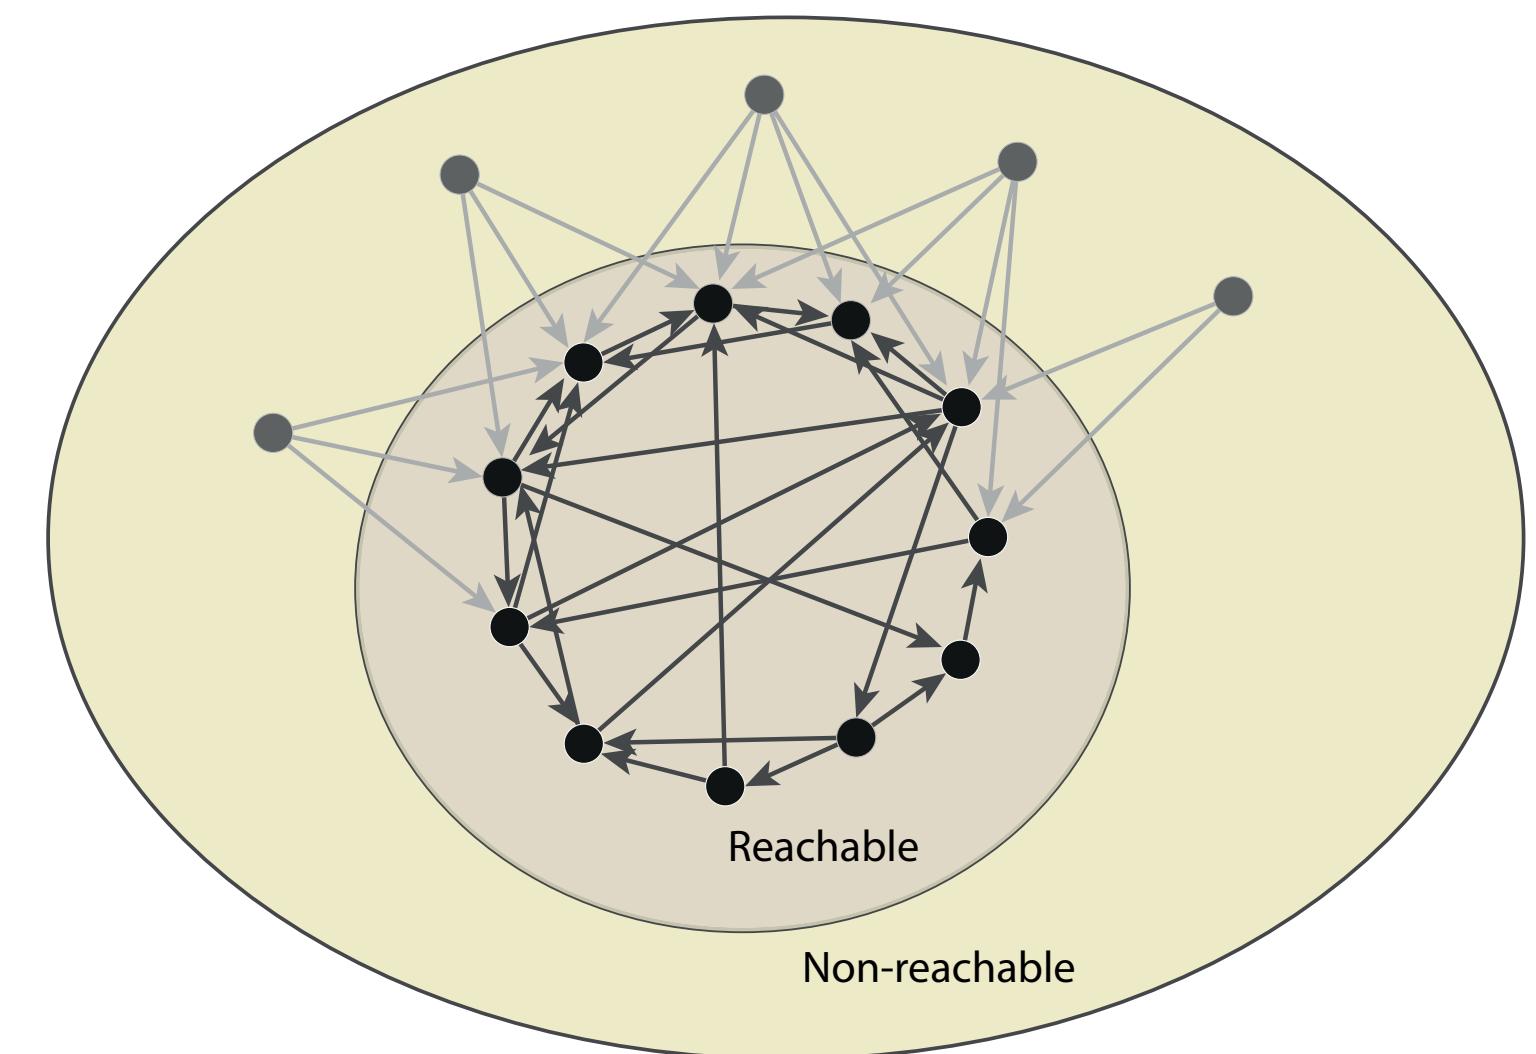
Reachable network: all nodes accept incoming / outgoing connections



NETWORK TAXONOMY

Reachable network: all nodes accept incoming / outgoing connections

Non-reachable: nodes do not accept incoming connections / cannot be reached (NAT/firewalls/...)

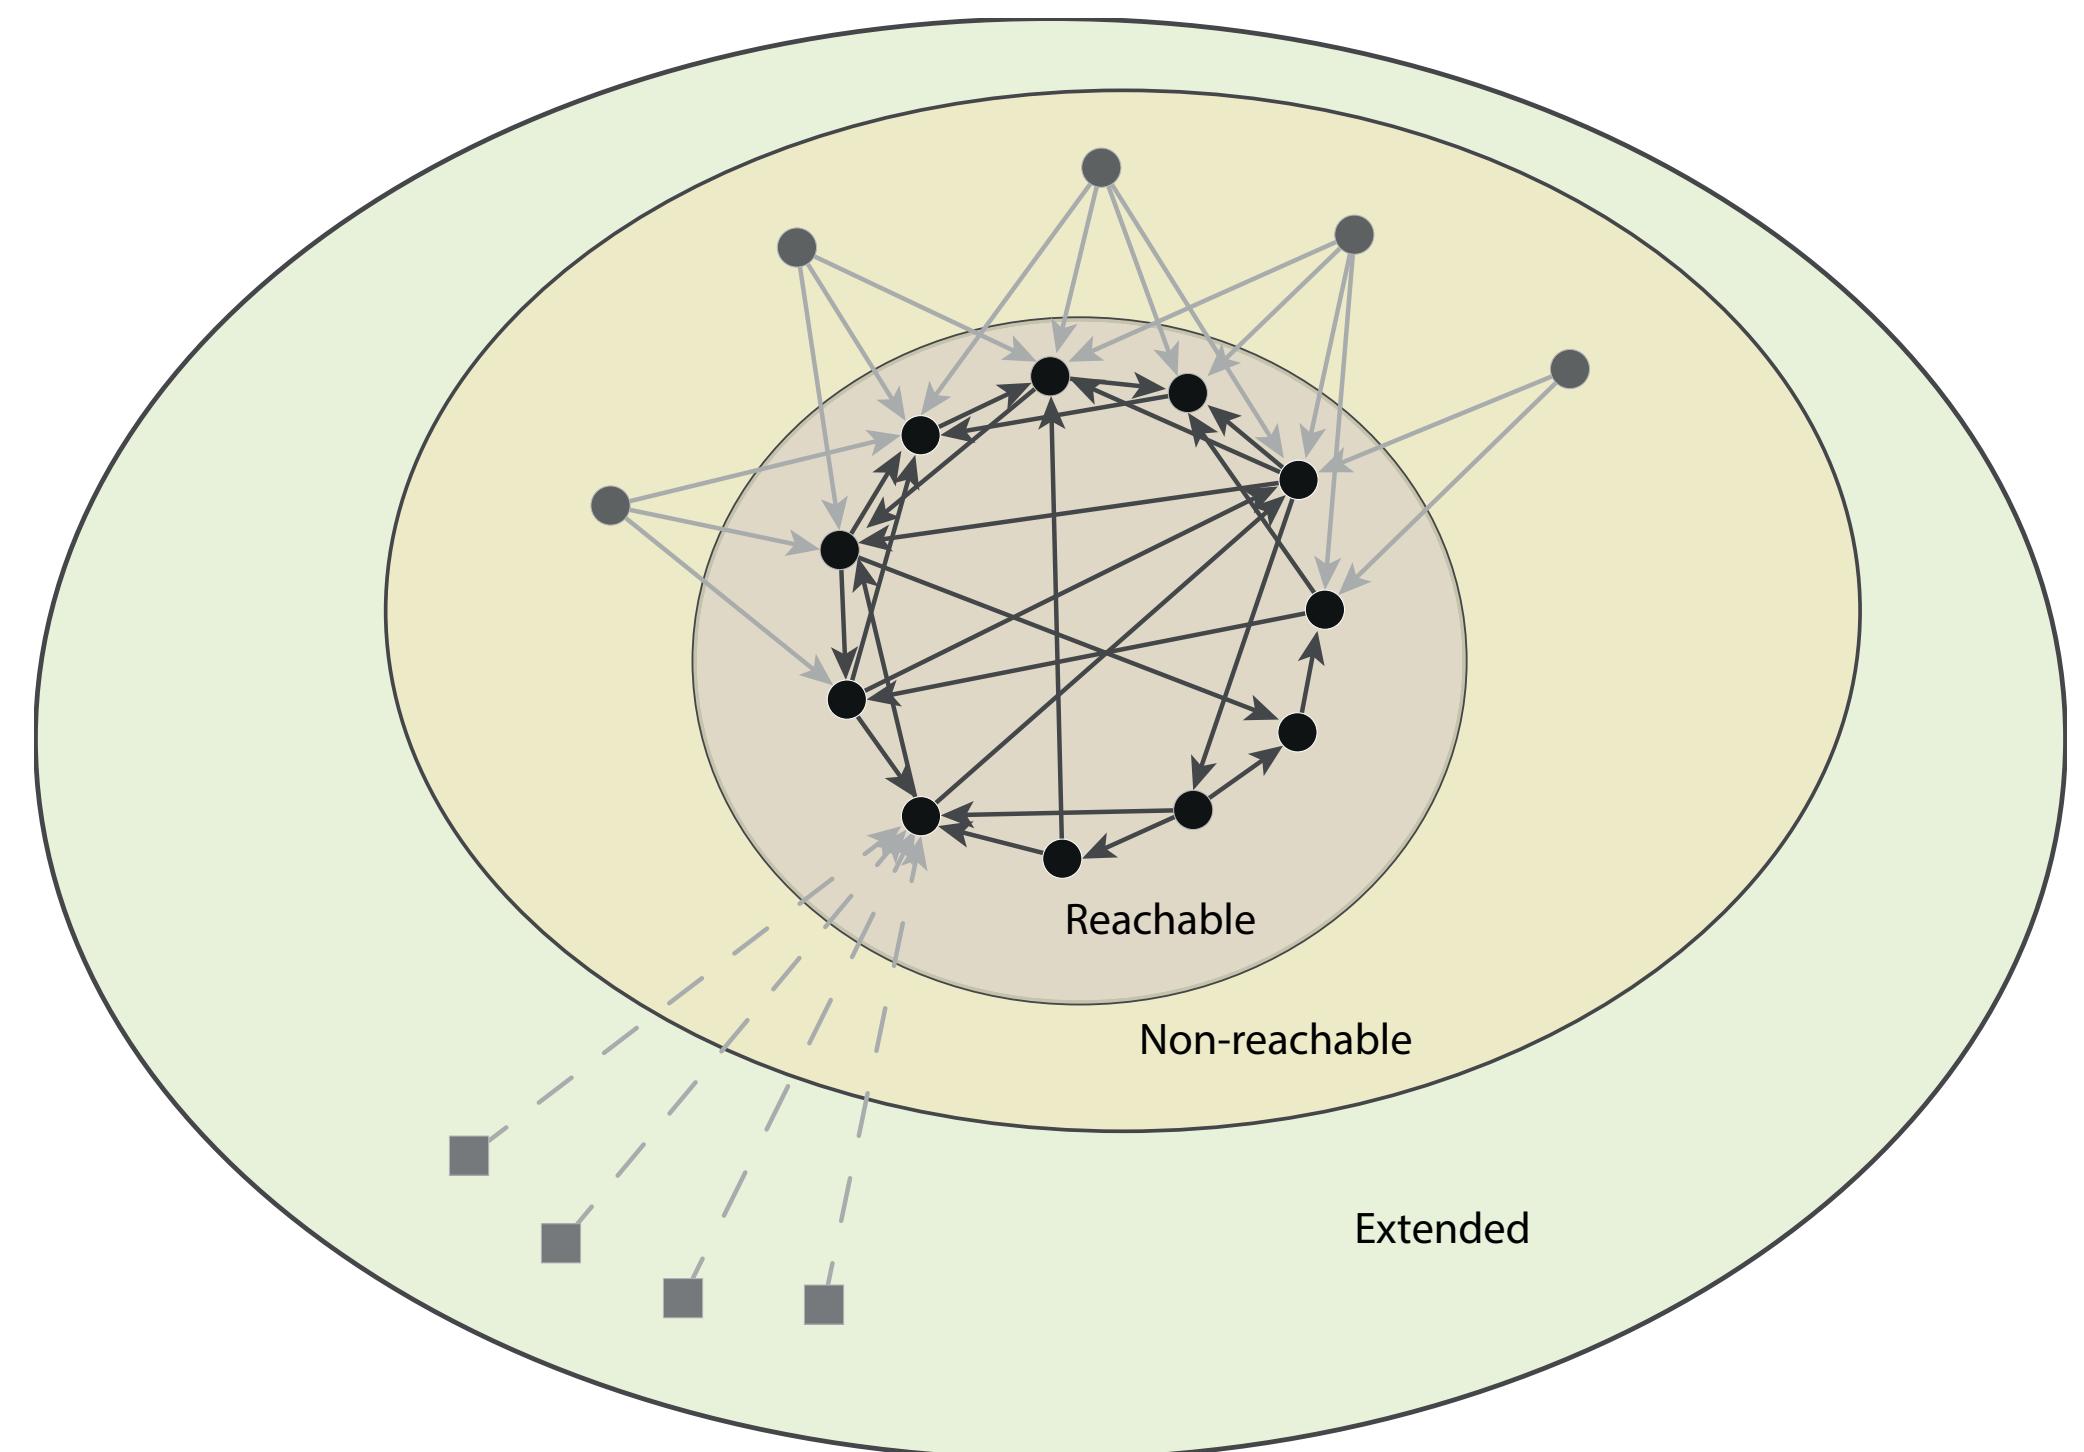


NETWORK TAXONOMY

Reachable network: all nodes accept incoming / outgoing connections

Non-reachable: nodes do not accept incoming connections / cannot be reached (NAT/firewalls/...)

Extended network: nodes use different protocol to communicate (not always P2P)





Nodes misbehavior

Nodes misbehavior

Every node maintains a **banscore** with each of its neighbors

If a node finds that one of its peers is misbehaving, the former will increase the banscore of the latter

If the banscore of a neighbor reaches (or surpasses) its maximum (**100 by default**), the node will ban that neighbor for a certain time (**24h by default**)

The banscore increase depends on how the neighbor is misbehaving

Banscore

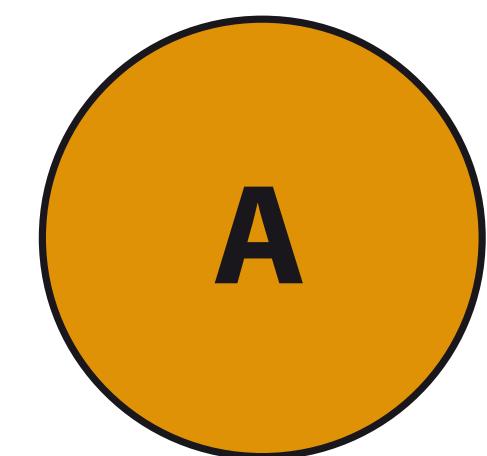
Examples of banscore increase:

- Not sending a version message as the first message in a handshake **(1)**
- Sending more than 1000 addresses in a single address message **(1)**
- Sending more than 50000 ids in a single inventory message **(20)**
- Sending a transaction with a script too big **(100)**

[src/net_processing.cpp](#) for more

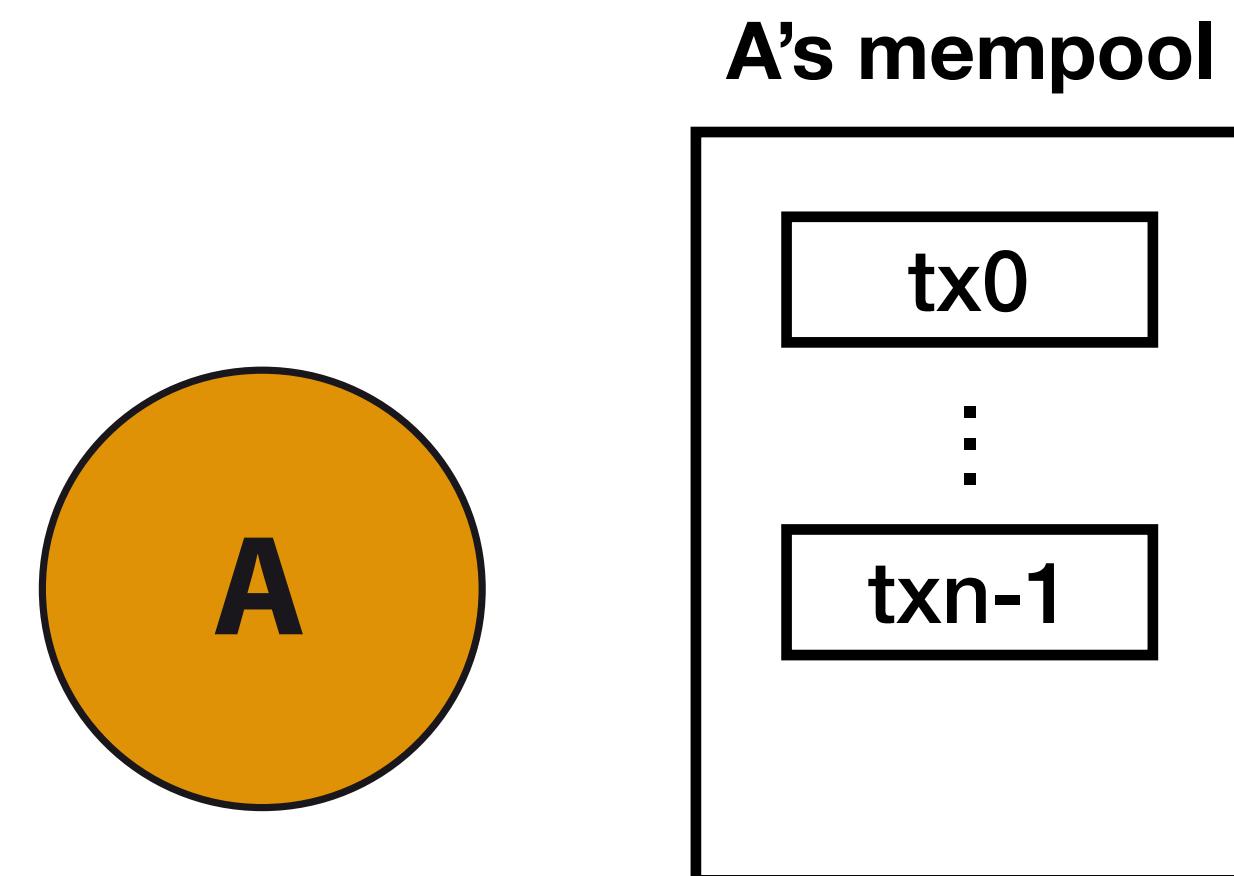
0-conf transactions and double-spending

UNCONFIRMED TRANSACTIONS



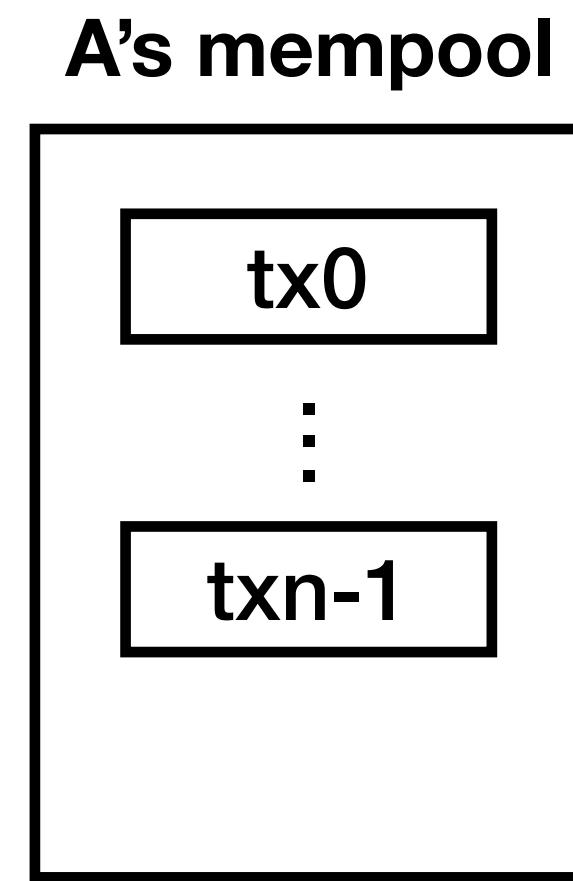
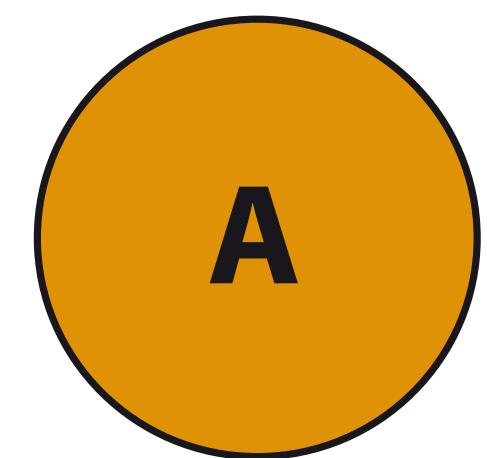
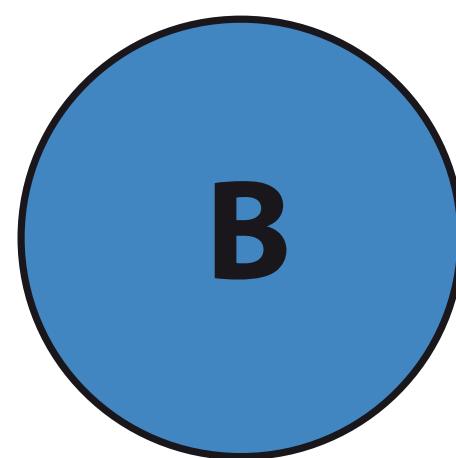
- **0-conf transactions / unconfirmed** transactions are those that are not part of the blockchain (**they are stored in the mempool**)
- 0-conf transactions are not covered by the double-spending protection offered by the blockchain (they are not part of it)
- Different nodes can have conflicting version of the “**same transaction**”

UNCONFIRMED TRANSACTIONS



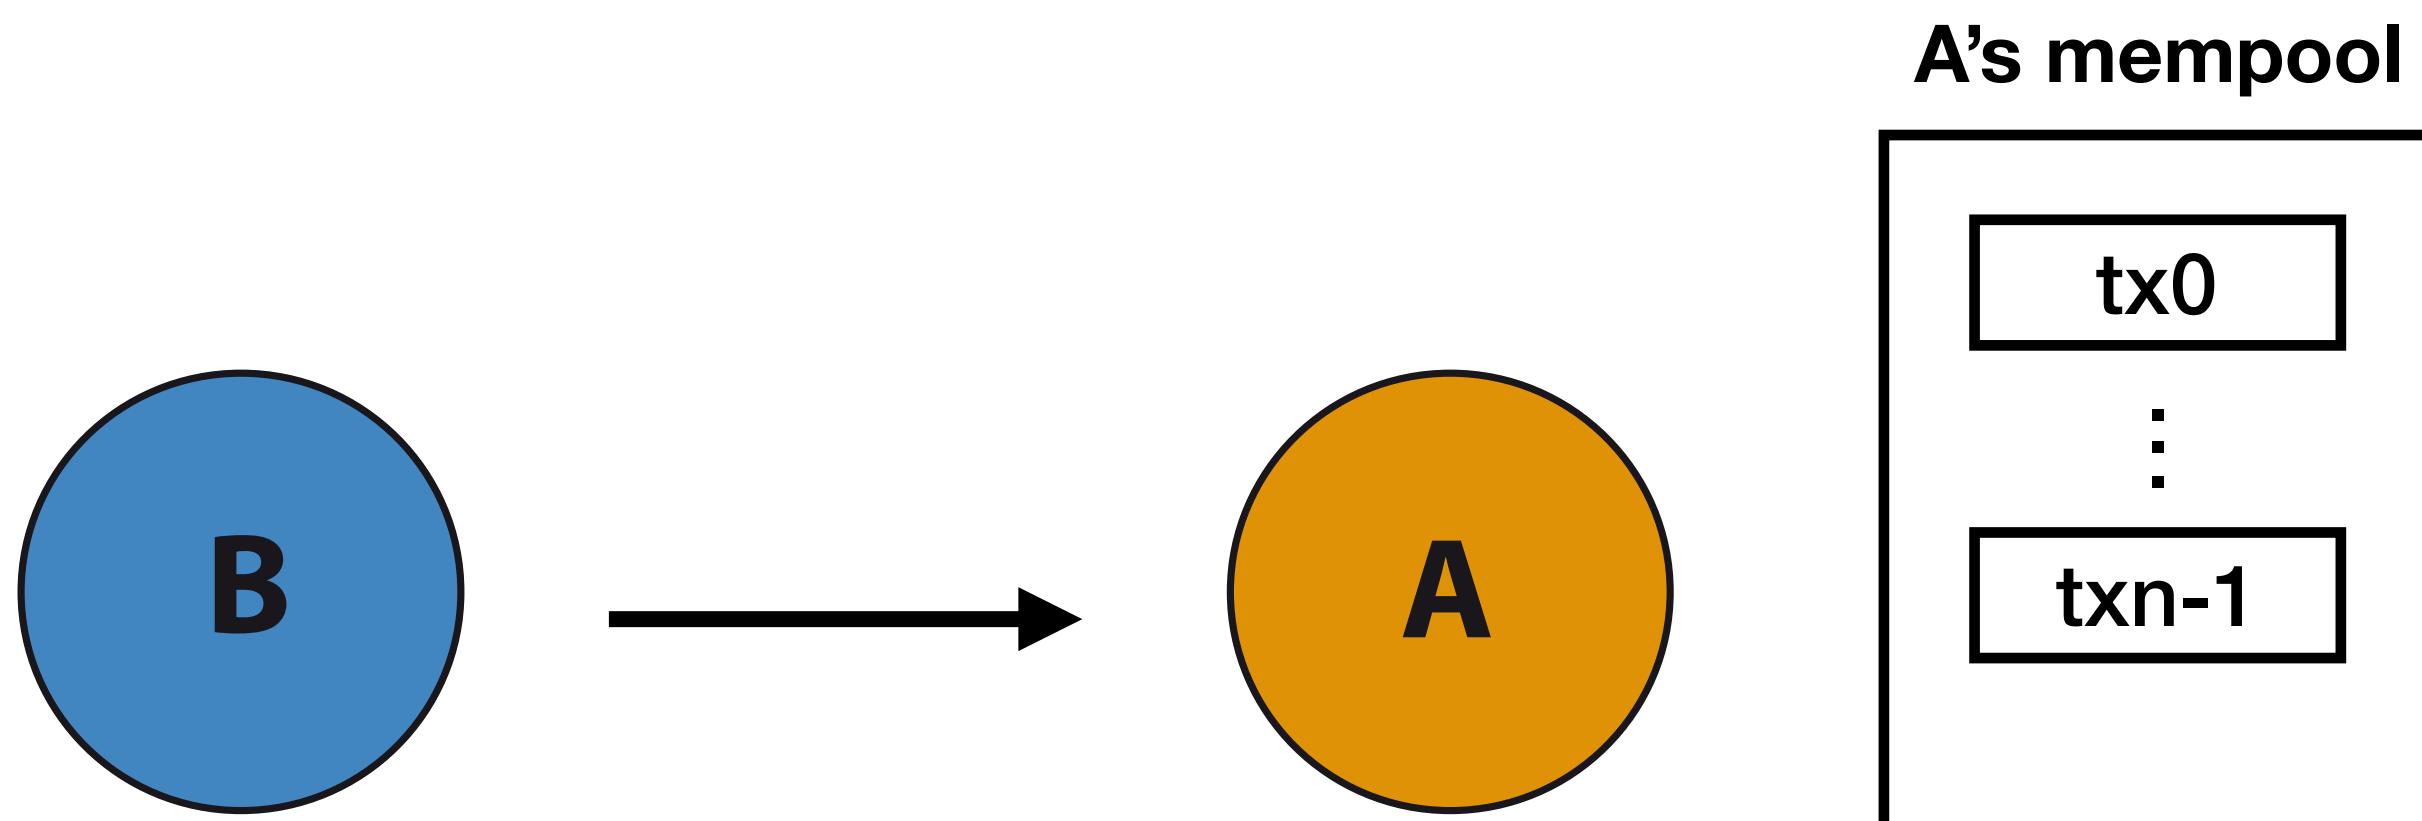
- **0-conf transactions / unconfirmed** transactions are those that are not part of the blockchain (**they are stored in the mempool**)
- 0-conf transactions are not covered by the double-spending protection offered by the blockchain (they are not part of it)
- Different nodes can have conflicting version of the “**same transaction**”

UNCONFIRMED TRANSACTIONS



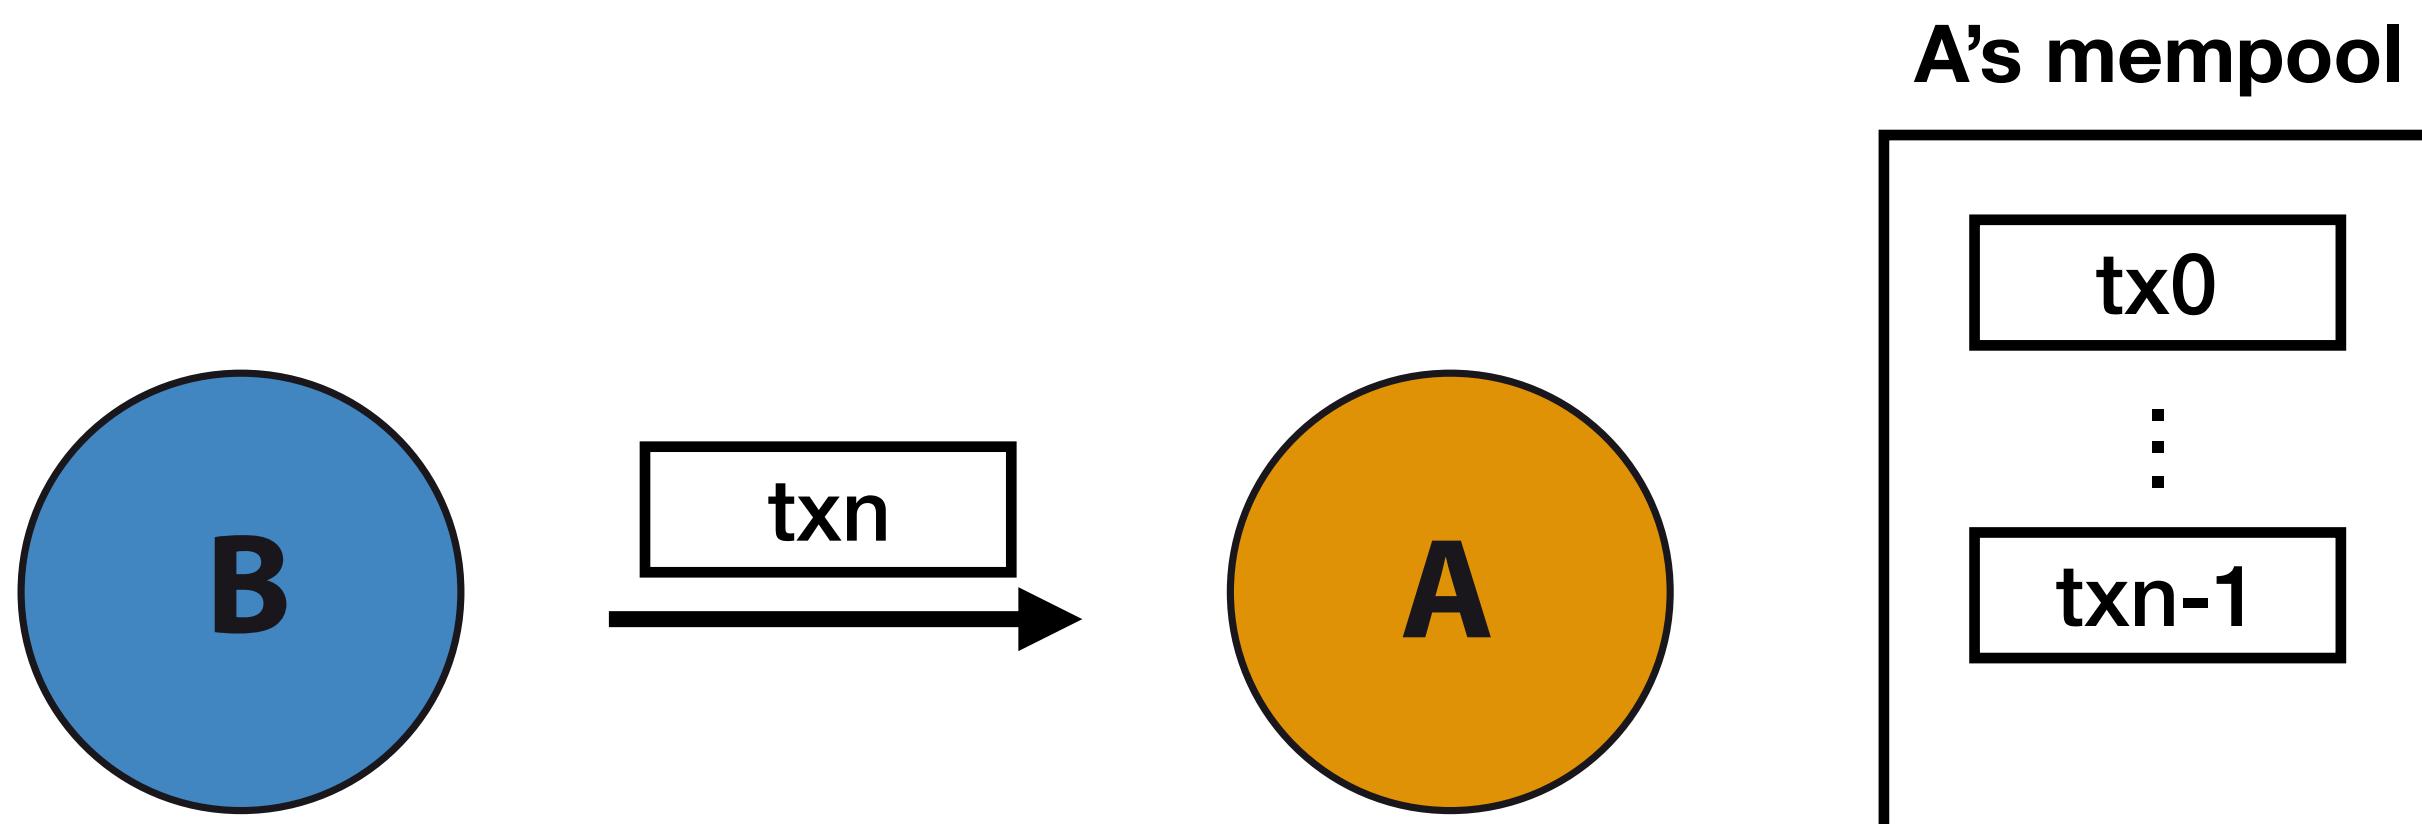
- **0-conf transactions / unconfirmed** transactions are those that are not part of the blockchain (**they are stored in the mempool**)
- 0-conf transactions are not covered by the double-spending protection offered by the blockchain (they are not part of it)
- Different nodes can have conflicting version of the “**same transaction**”

UNCONFIRMED TRANSACTIONS



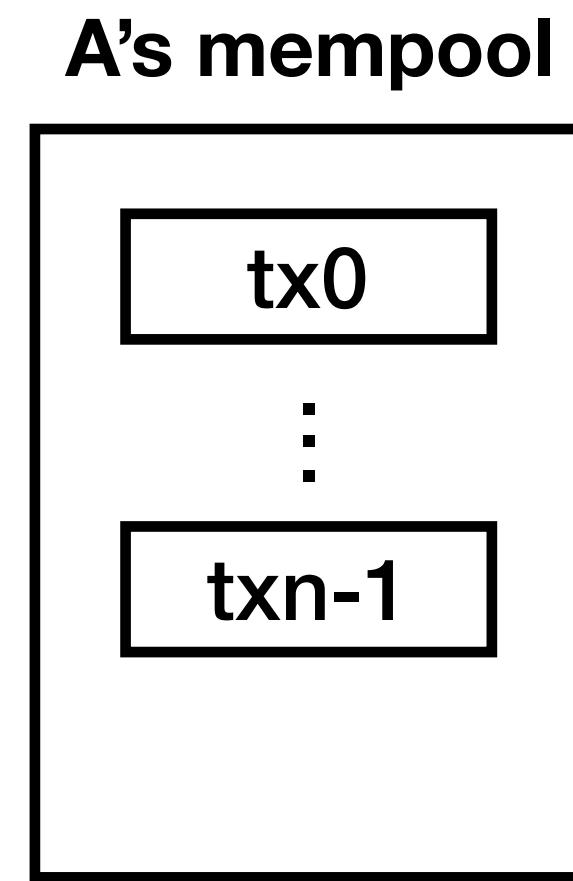
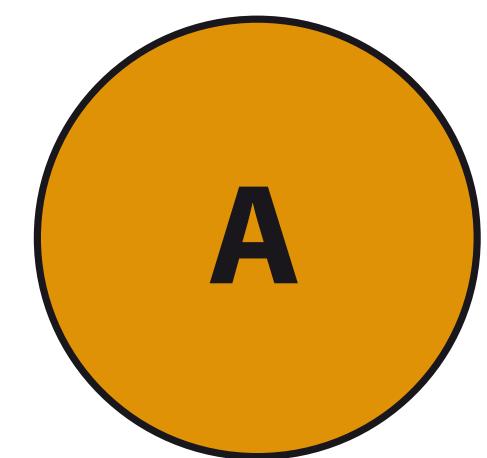
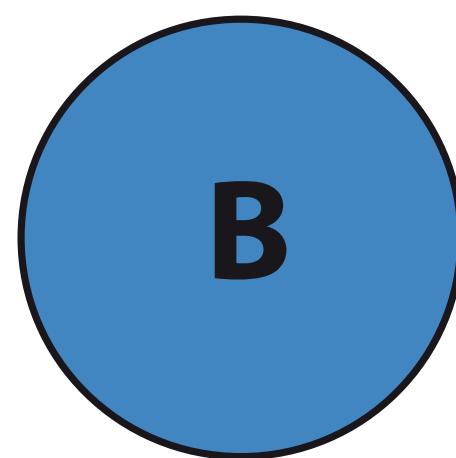
- **0-conf transactions / unconfirmed** transactions are those that are not part of the blockchain (**they are stored in the mempool**)
- 0-conf transactions are not covered by the double-spending protection offered by the blockchain (they are not part of it)
- Different nodes can have conflicting version of the “**same transaction**”

UNCONFIRMED TRANSACTIONS



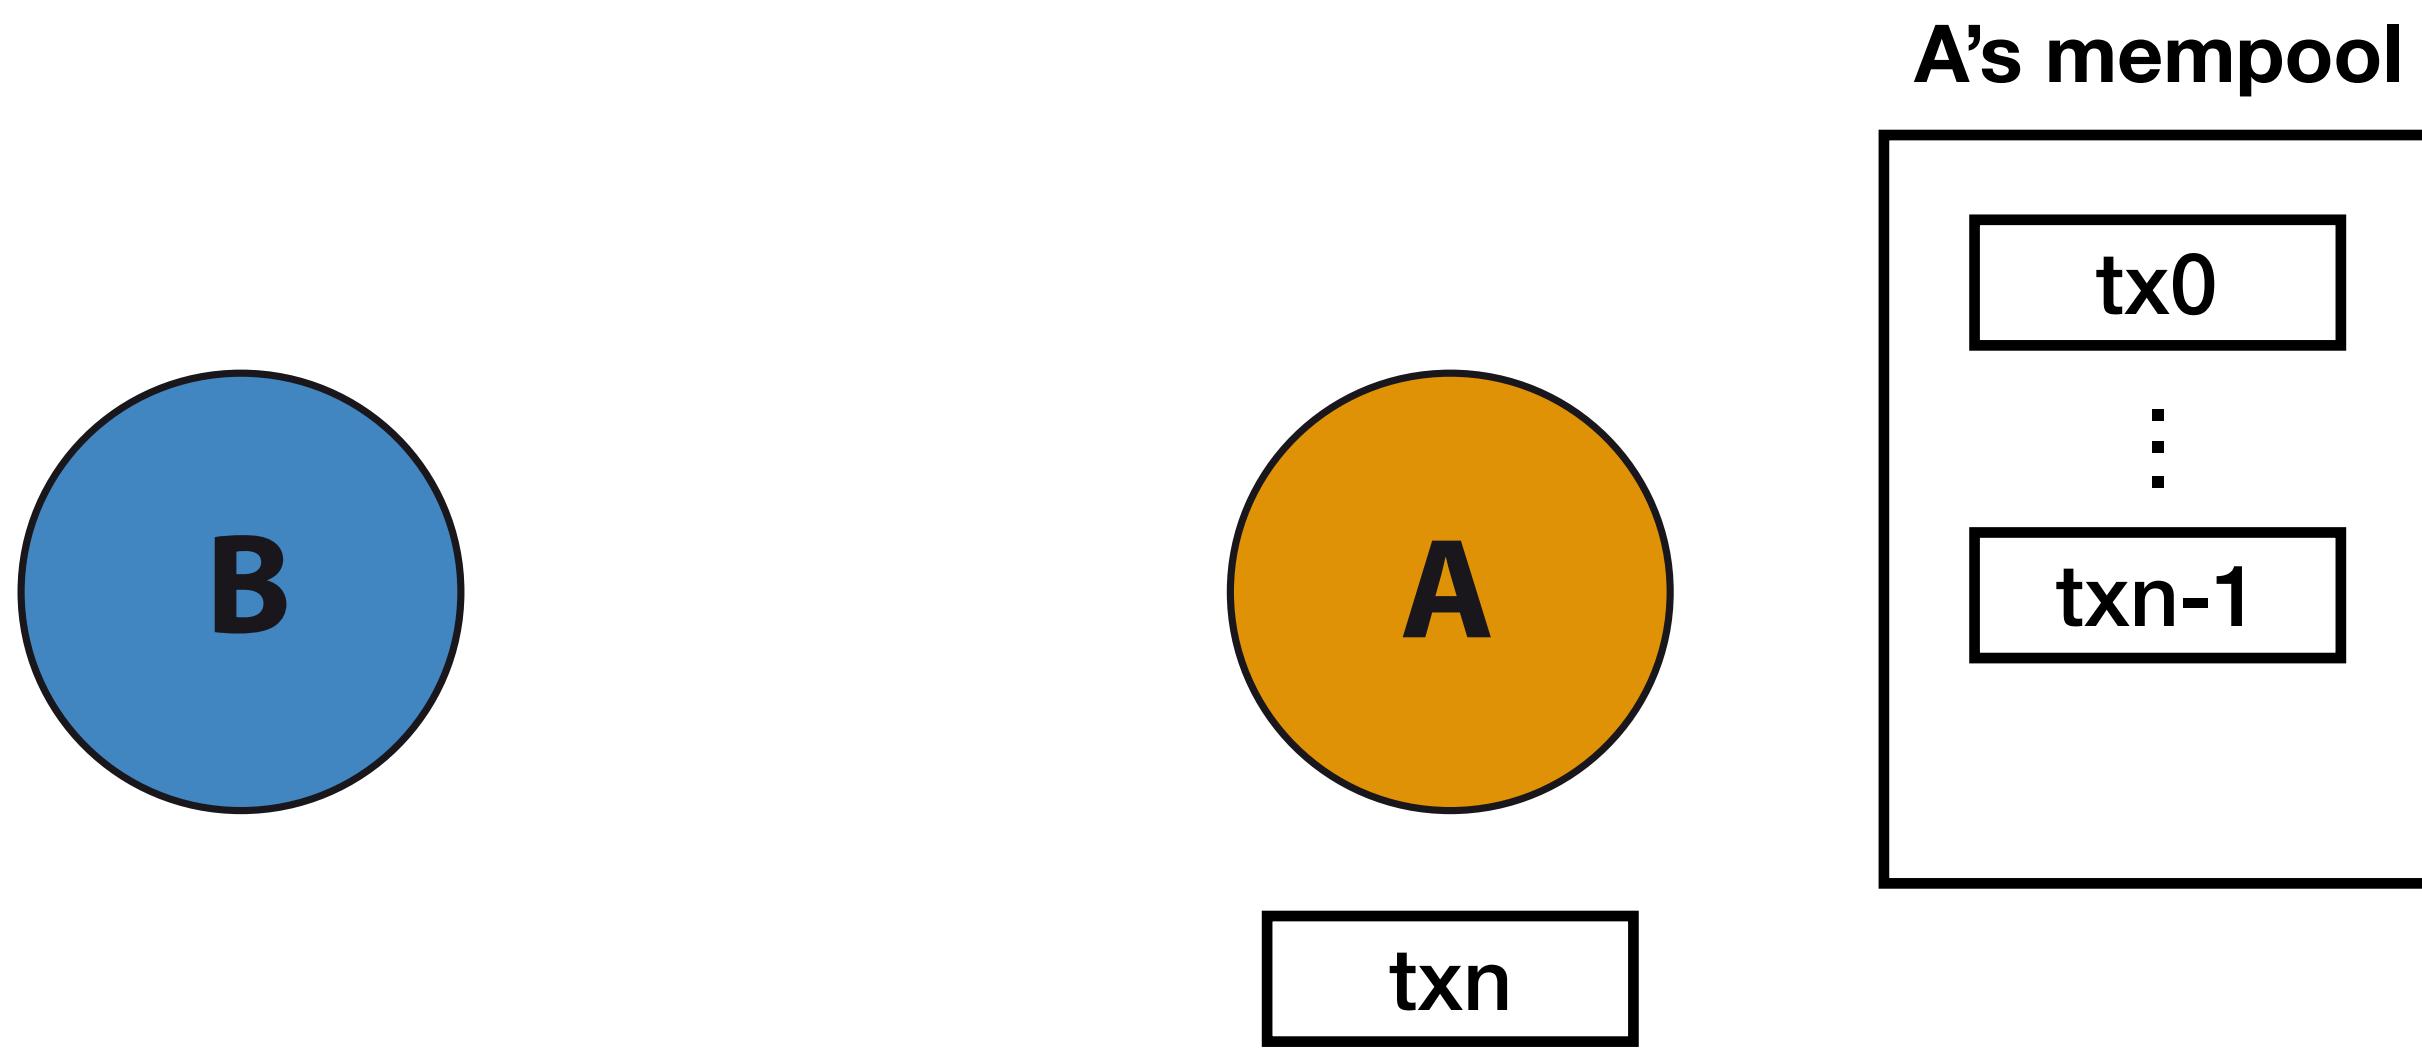
- **0-conf transactions / unconfirmed** transactions are those that are not part of the blockchain (**they are stored in the mempool**)
- 0-conf transactions are not covered by the double-spending protection offered by the blockchain (they are not part of it)
- Different nodes can have conflicting version of the “**same transaction**”

UNCONFIRMED TRANSACTIONS



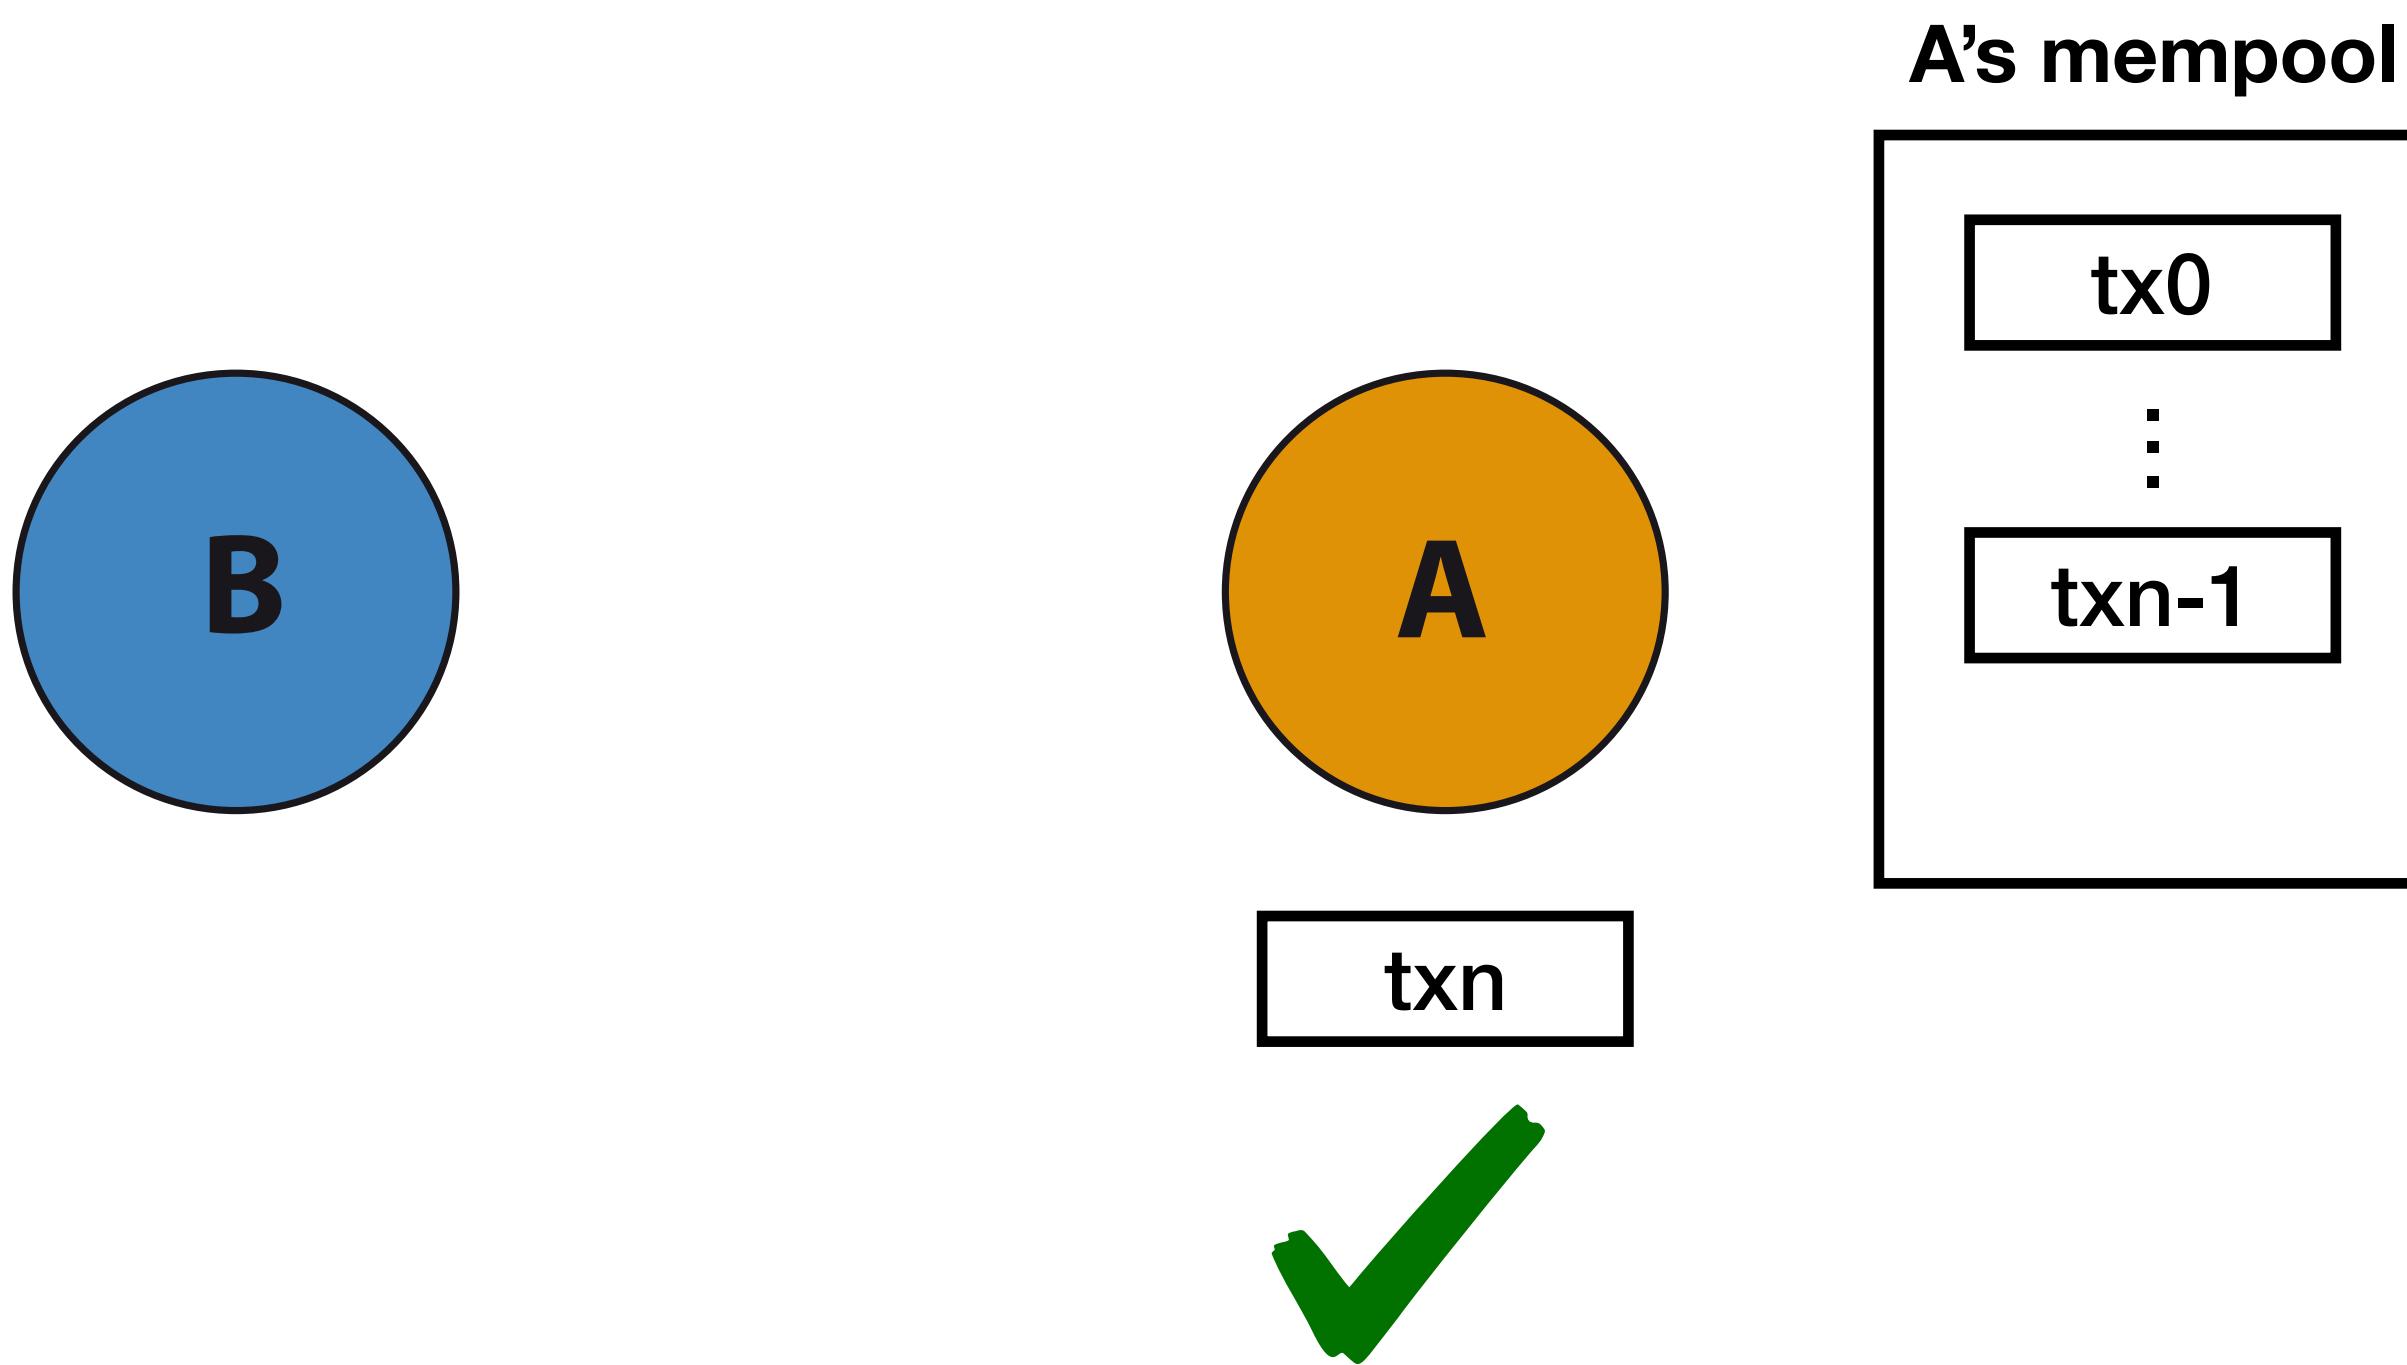
- **0-conf transactions / unconfirmed** transactions are those that are not part of the blockchain (**they are stored in the mempool**)
- 0-conf transactions are not covered by the double-spending protection offered by the blockchain (they are not part of it)
- Different nodes can have conflicting version of the “**same transaction**”

UNCONFIRMED TRANSACTIONS



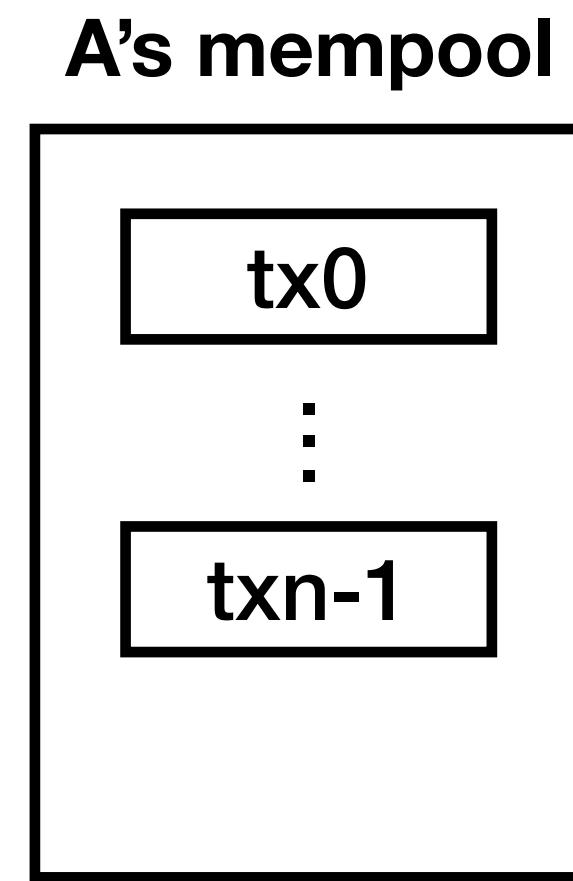
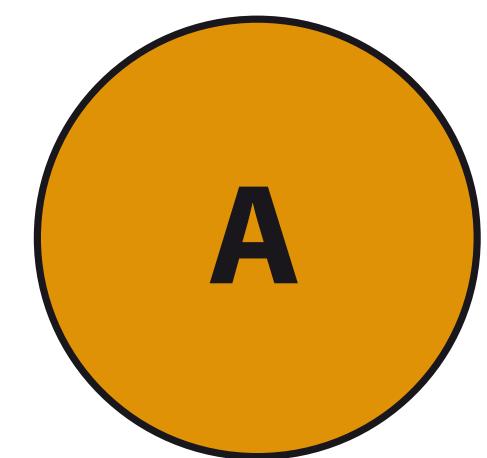
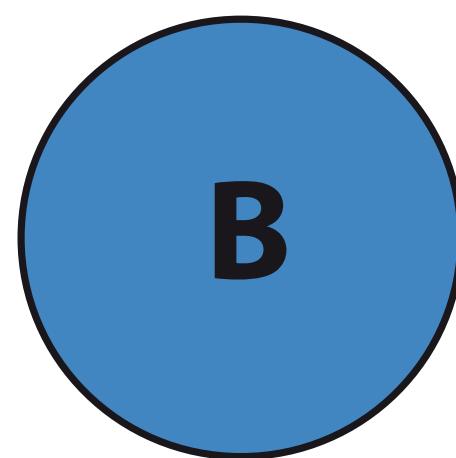
- **0-conf transactions / unconfirmed** transactions are those that are not part of the blockchain (**they are stored in the mempool**)
- 0-conf transactions are not covered by the double-spending protection offered by the blockchain (they are not part of it)
- Different nodes can have conflicting version of the “**same transaction**”

UNCONFIRMED TRANSACTIONS



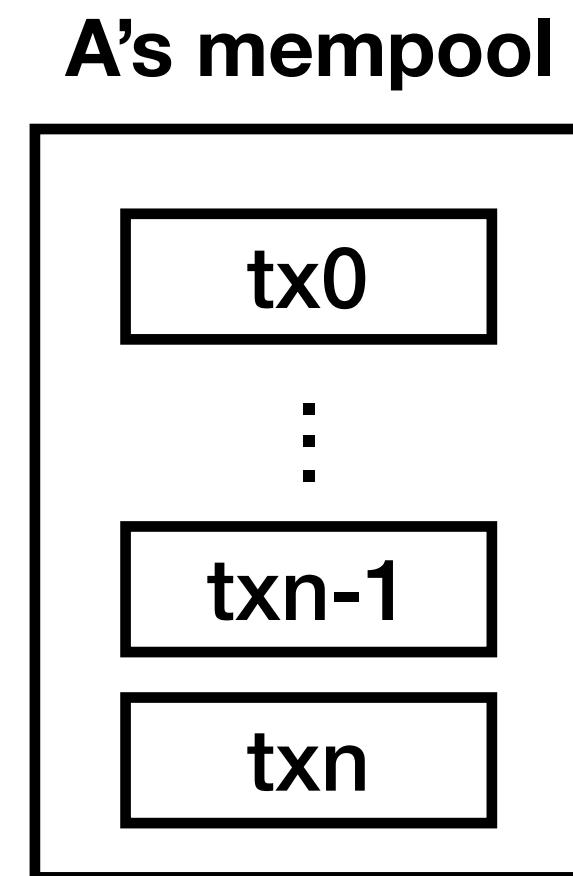
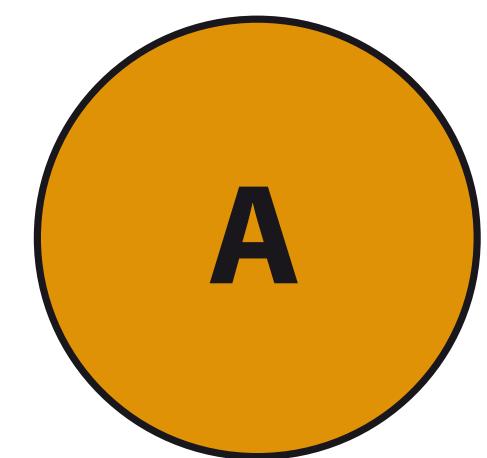
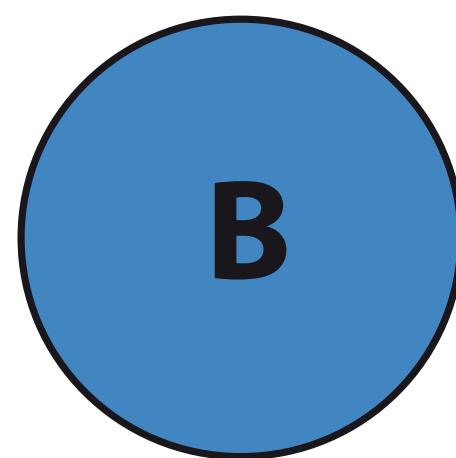
- **0-conf transactions / unconfirmed** transactions are those that are not part of the blockchain (**they are stored in the mempool**)
- 0-conf transactions are not covered by the double-spending protection offered by the blockchain (they are not part of it)
- Different nodes can have conflicting version of the “**same transaction**”

UNCONFIRMED TRANSACTIONS



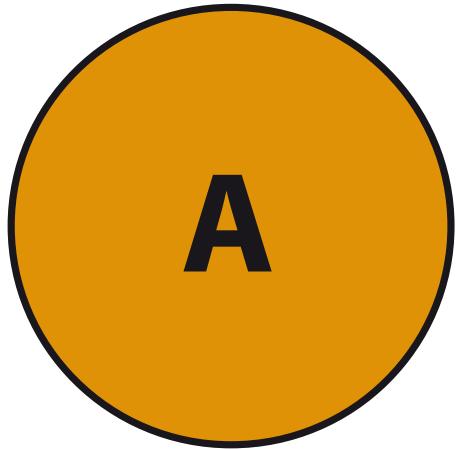
- **0-conf transactions / unconfirmed** transactions are those that are not part of the blockchain (**they are stored in the mempool**)
- 0-conf transactions are not covered by the double-spending protection offered by the blockchain (they are not part of it)
- Different nodes can have conflicting version of the “**same transaction**”

UNCONFIRMED TRANSACTIONS



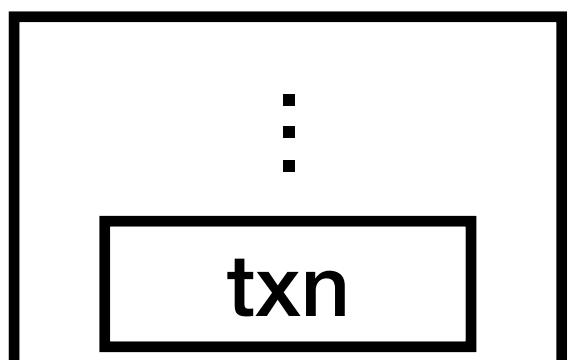
- **0-conf** transactions / **unconfirmed** transactions are those that are not part of the blockchain (**they are stored in the mempool**)
- 0-conf transactions are not covered by the double-spending protection offered by the blockchain (they are not part of it)
- Different nodes can have conflicting version of the “**same transaction**”

CONFIRMED TRANSACTIONS

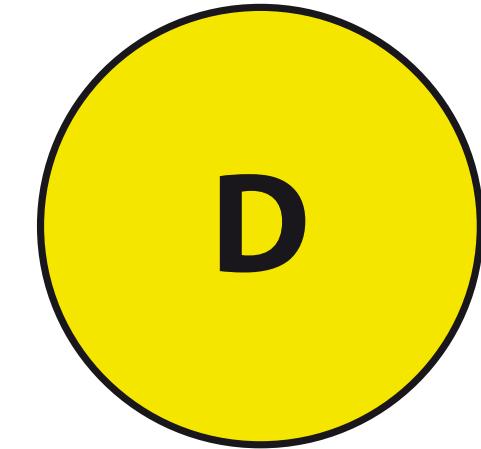
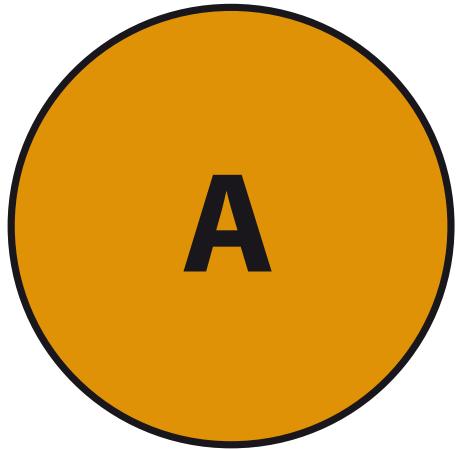


The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

A's mempool

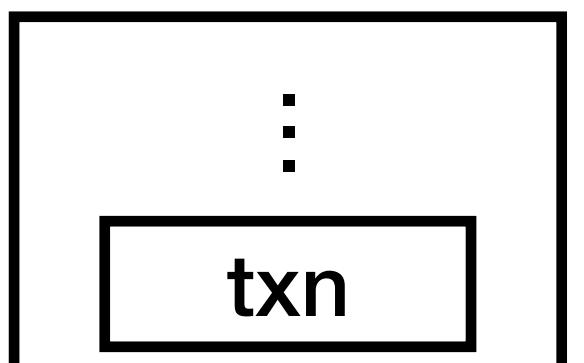


CONFIRMED TRANSACTIONS



The de facto confirmation time is **6 blocks** (5 on top of the one including a certain transaction)

A's mempool

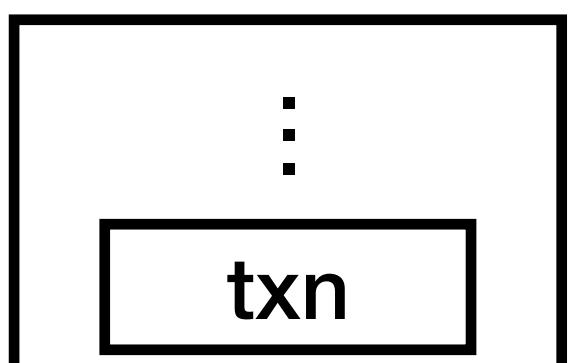


CONFIRMED TRANSACTIONS

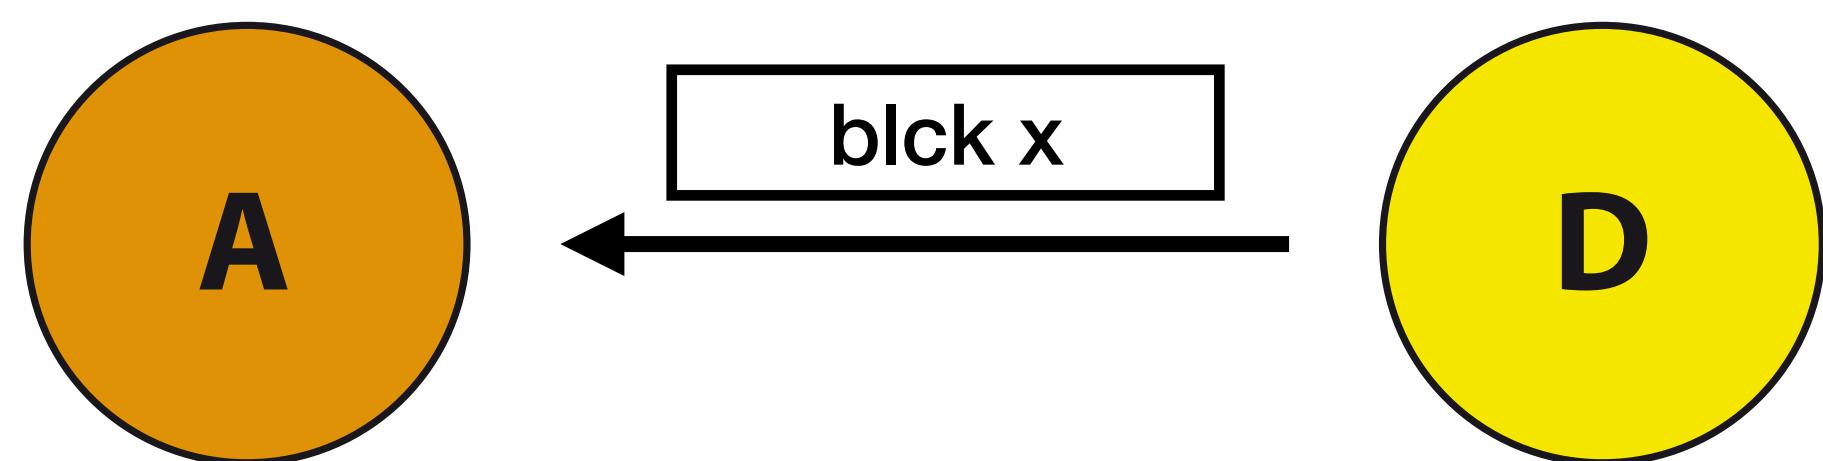


The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

A's mempool

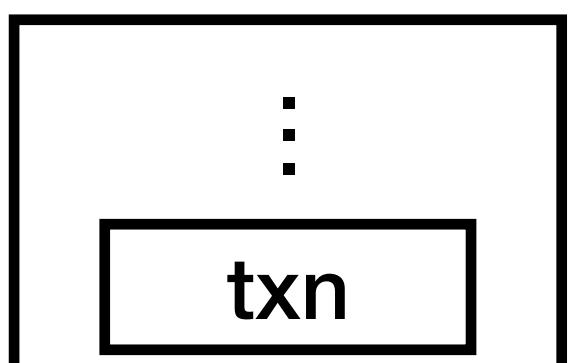


CONFIRMED TRANSACTIONS

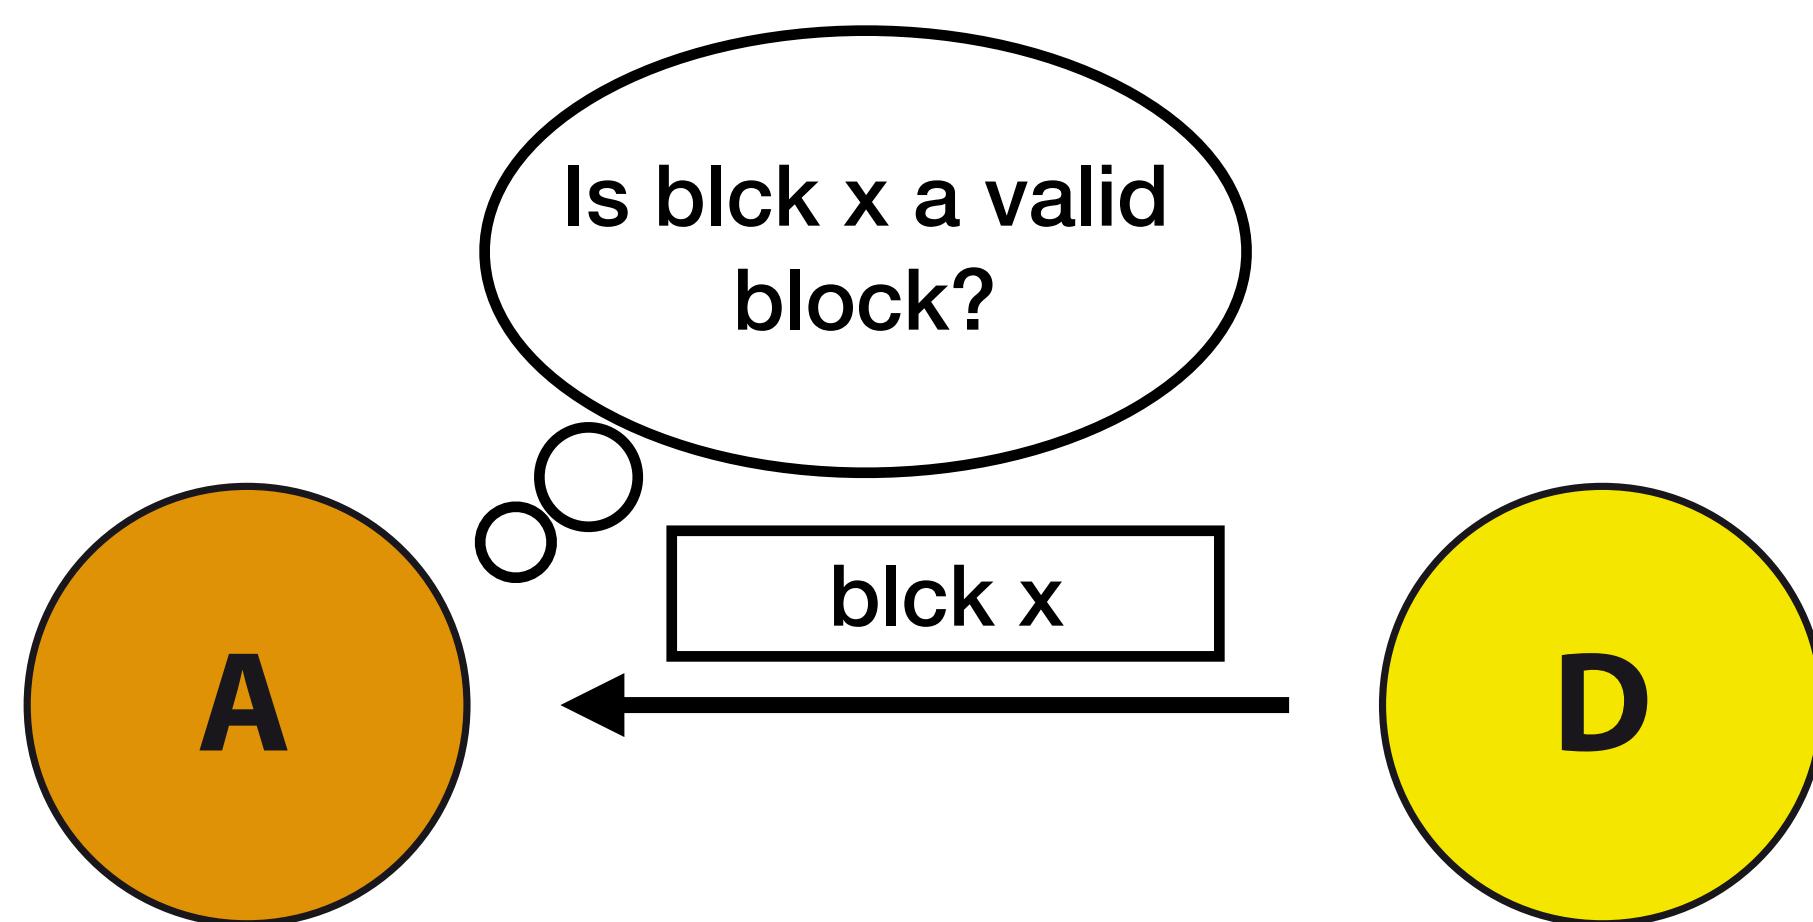


The de facto confirmation time is **6 blocks** (5 on top of the one including a certain transaction)

A's mempool

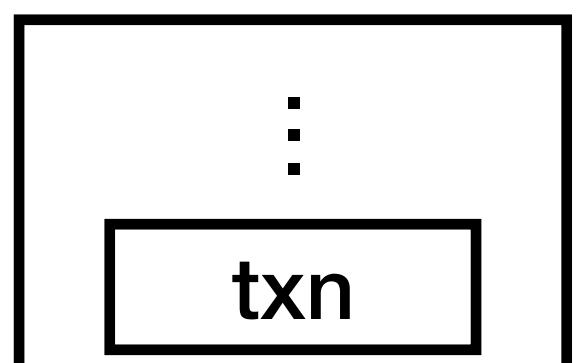


CONFIRMED TRANSACTIONS

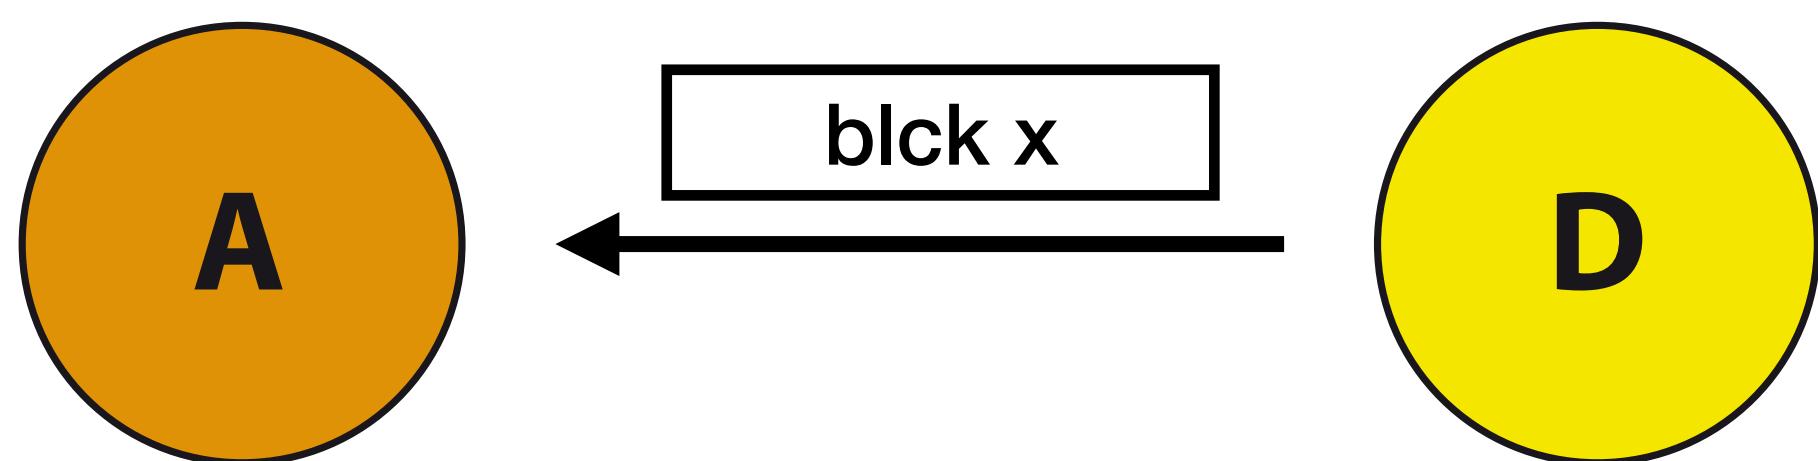


The de facto confirmation time is **6 blocks** (5 on top of the one including a certain transaction)

A's mempool

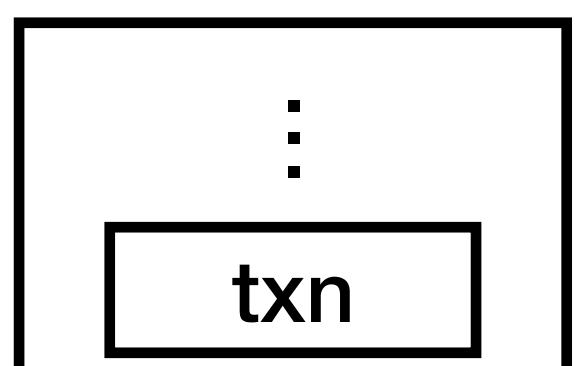


CONFIRMED TRANSACTIONS

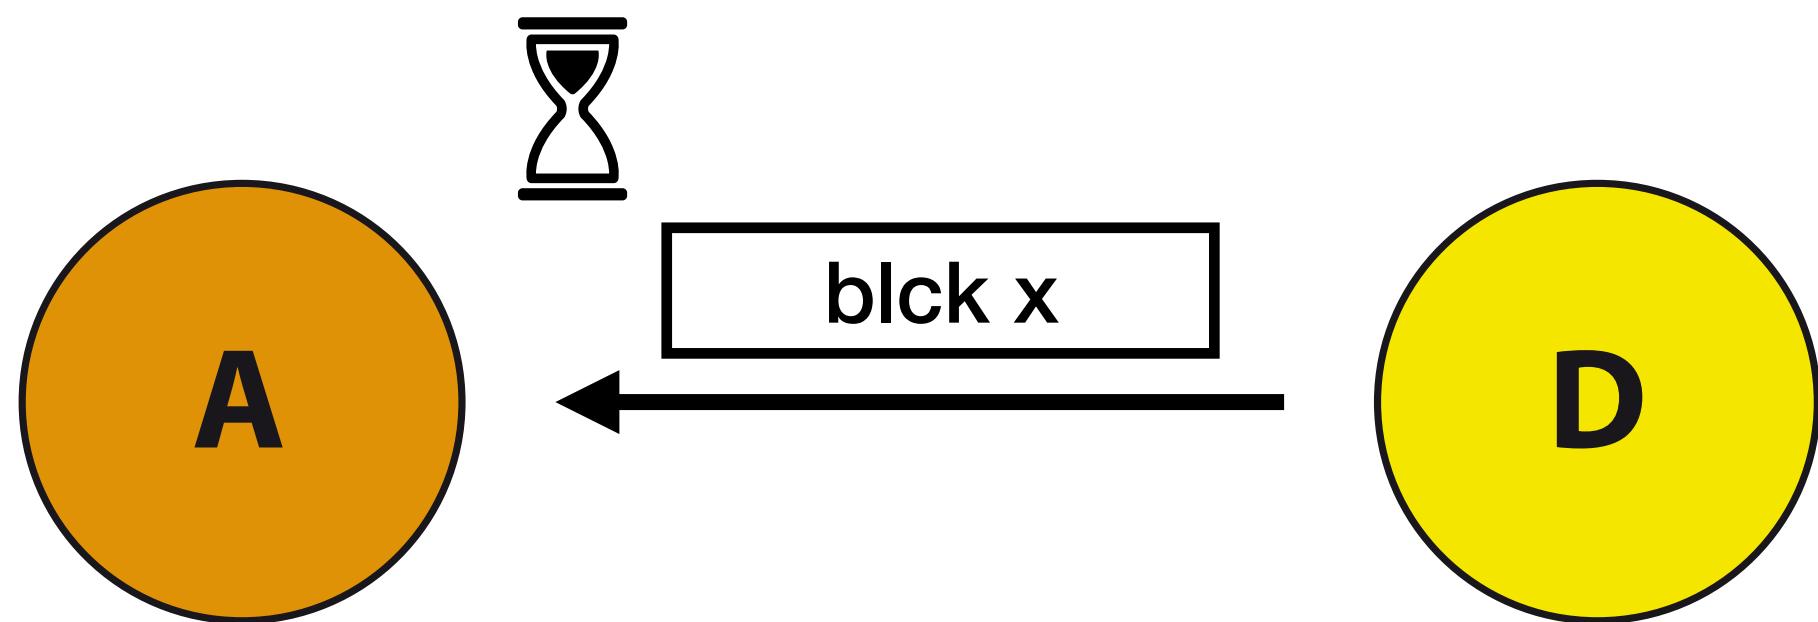


The de facto confirmation time is **6 blocks** (5 on top of the one including a certain transaction)

A's mempool

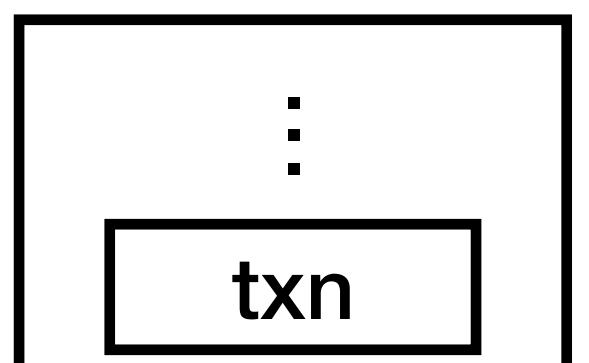


CONFIRMED TRANSACTIONS

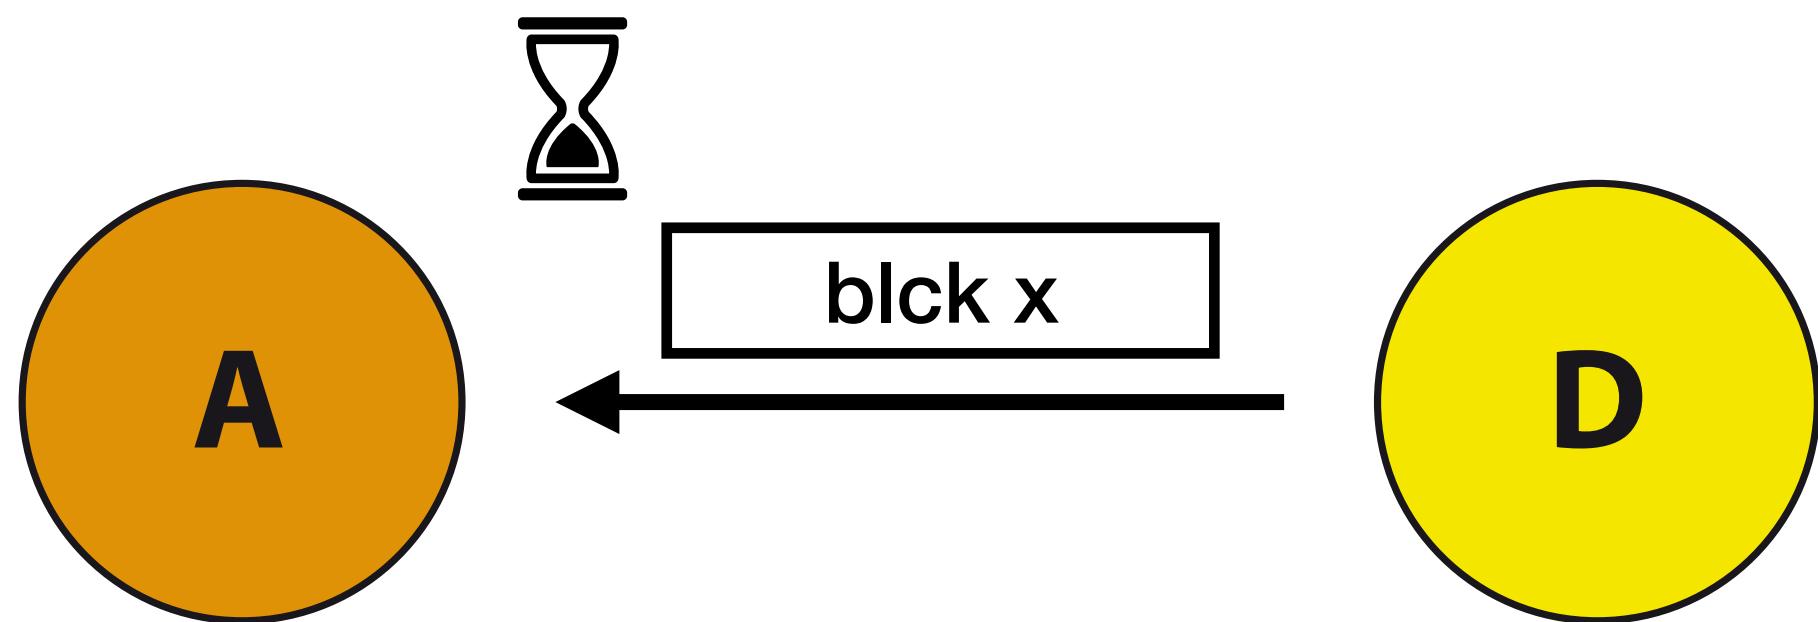


The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

A's mempool

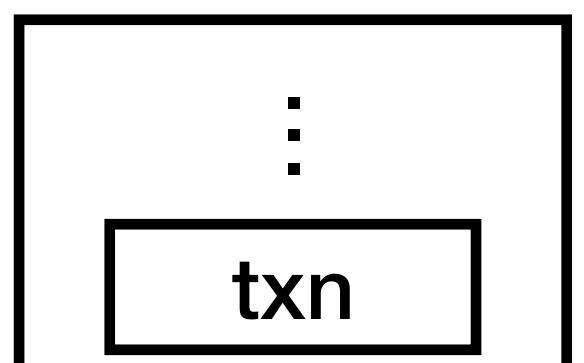


CONFIRMED TRANSACTIONS

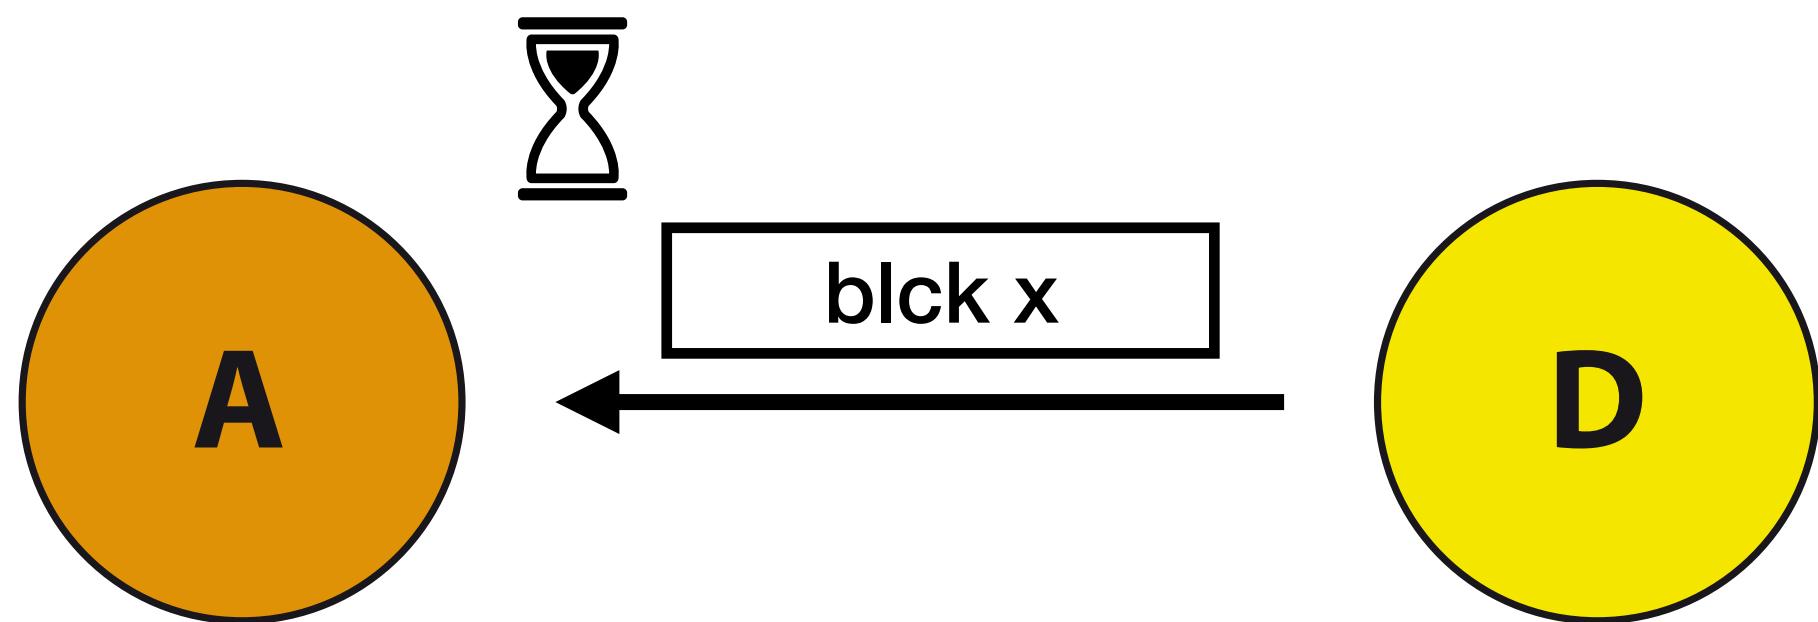


The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

A's mempool

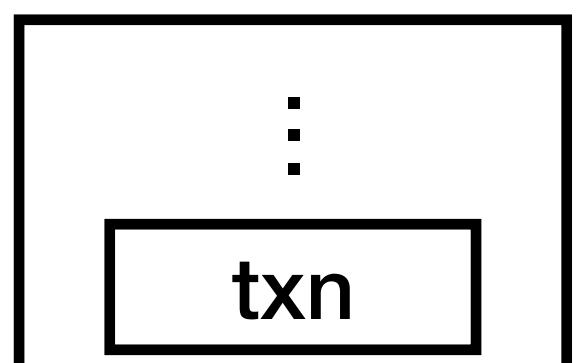


CONFIRMED TRANSACTIONS

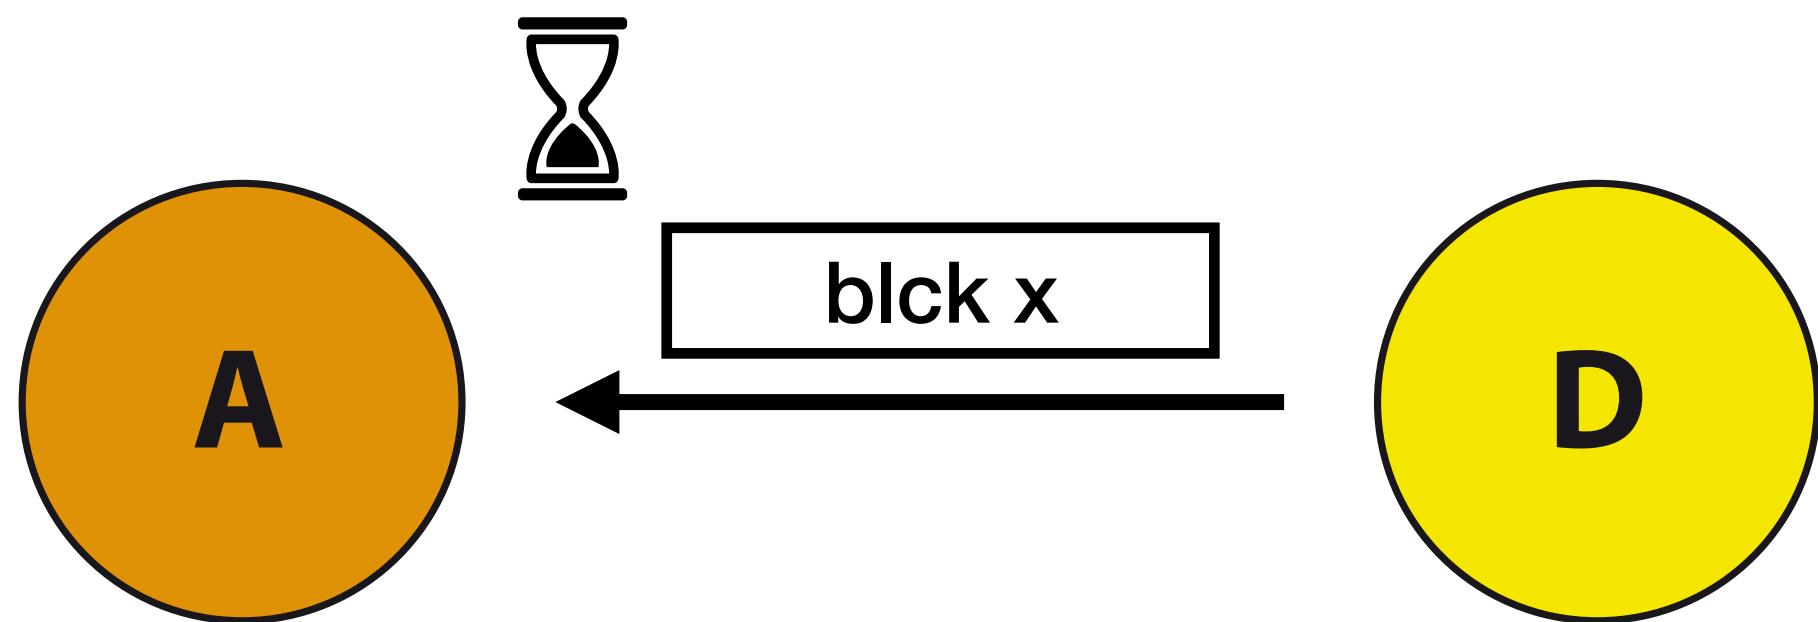


The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

A's mempool

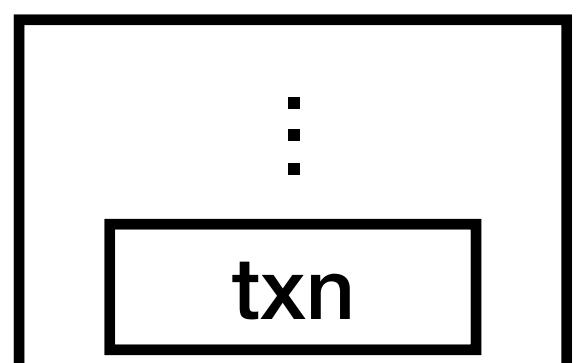


CONFIRMED TRANSACTIONS

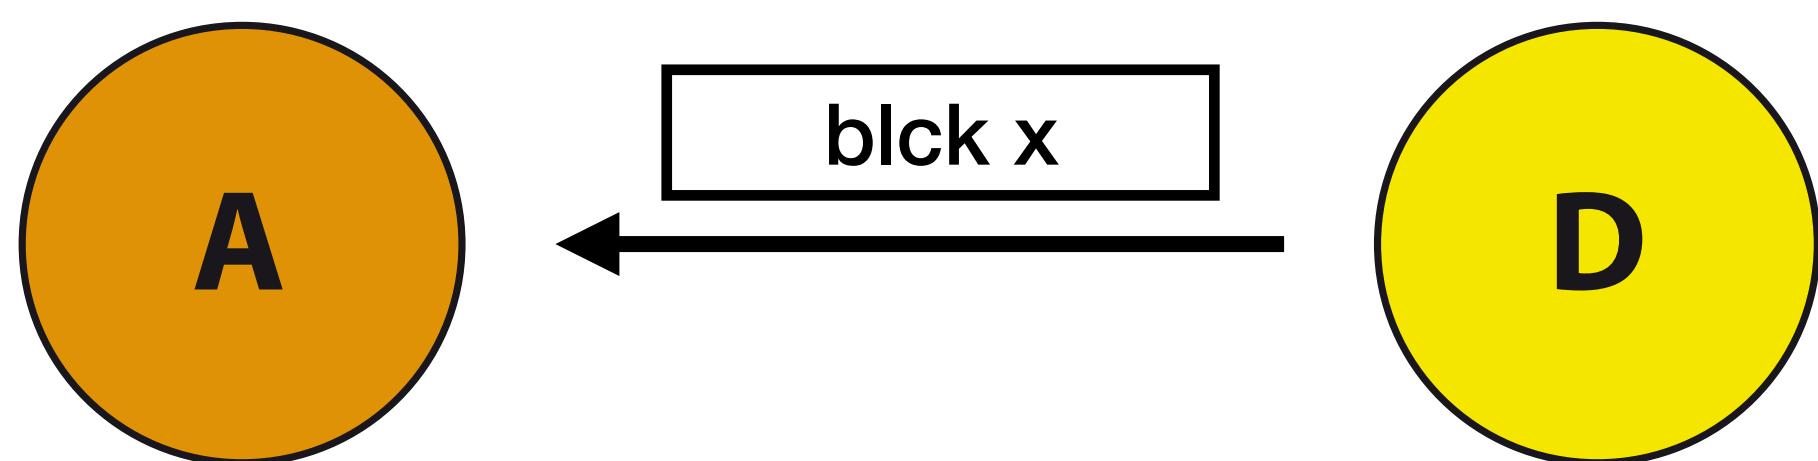


The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

A's mempool

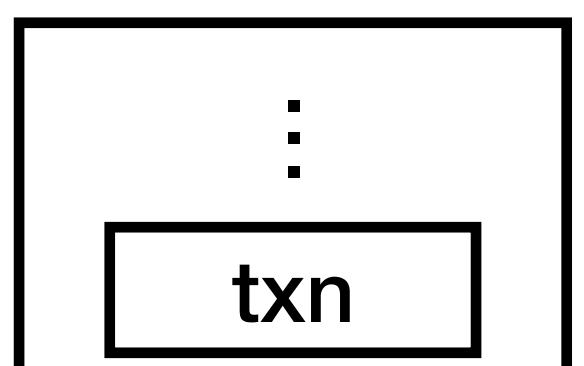


CONFIRMED TRANSACTIONS

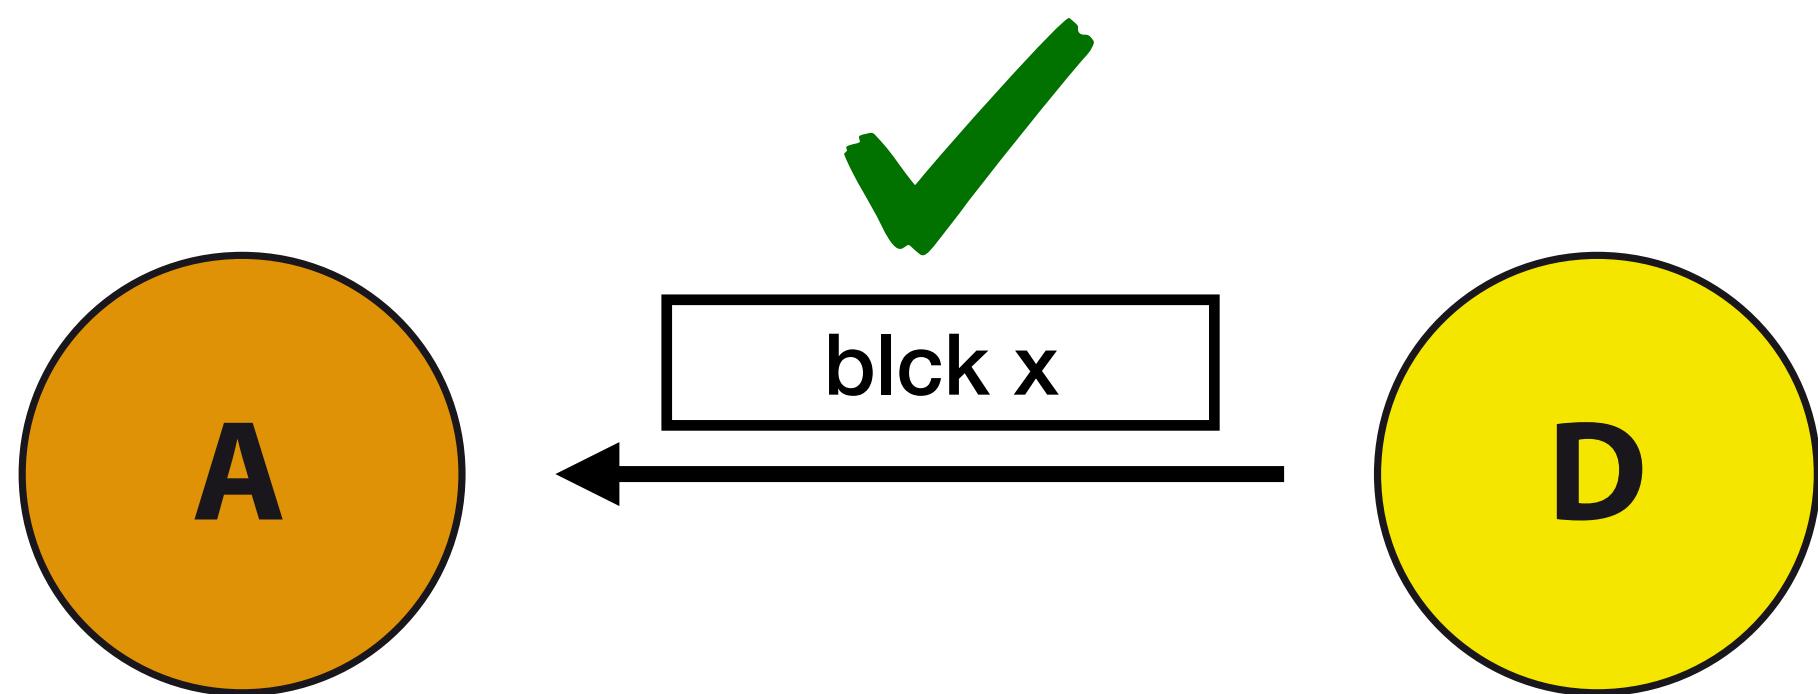


The de facto confirmation time is **6 blocks** (5 on top of the one including a certain transaction)

A's mempool

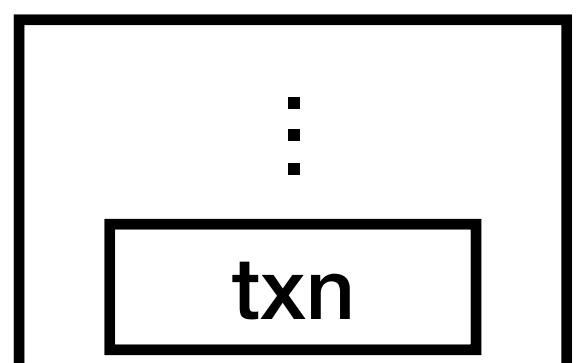


CONFIRMED TRANSACTIONS

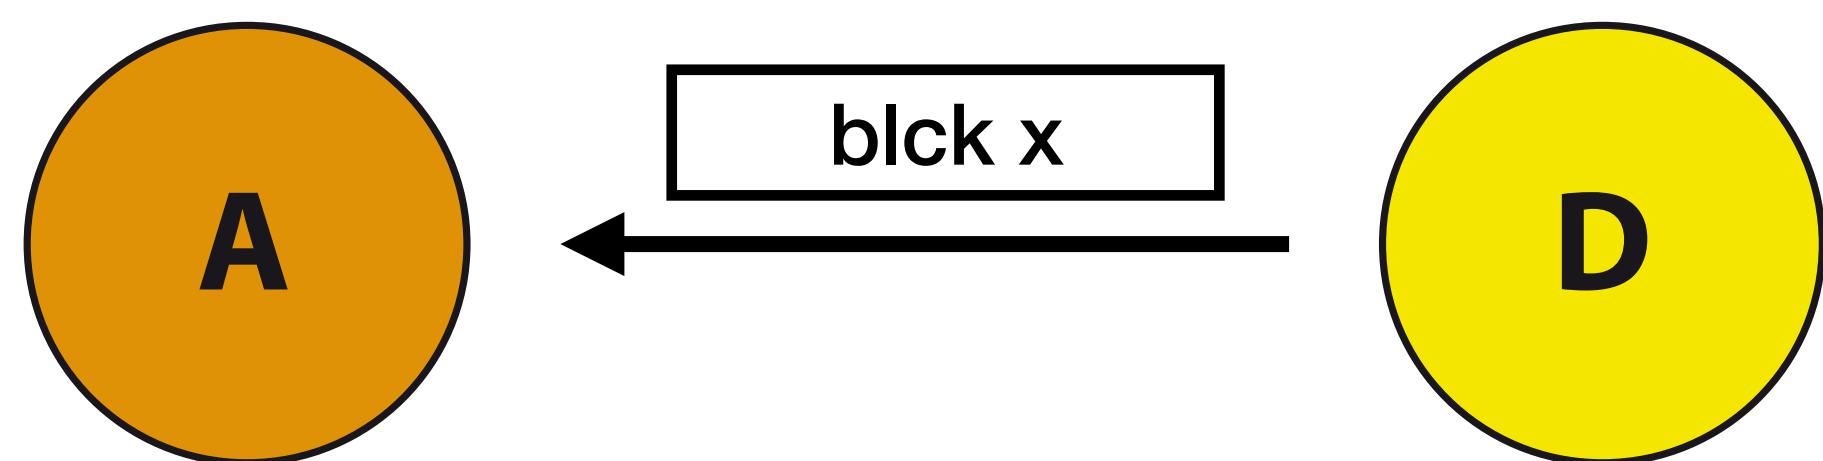


The de facto confirmation time is **6 blocks** (5 on top of the one including a certain transaction)

A's mempool

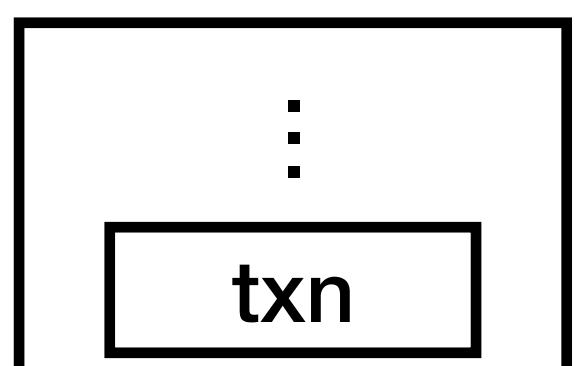


CONFIRMED TRANSACTIONS

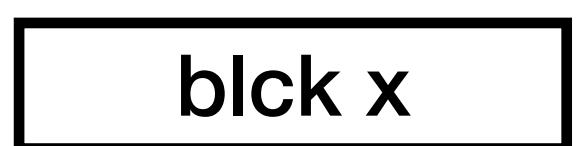
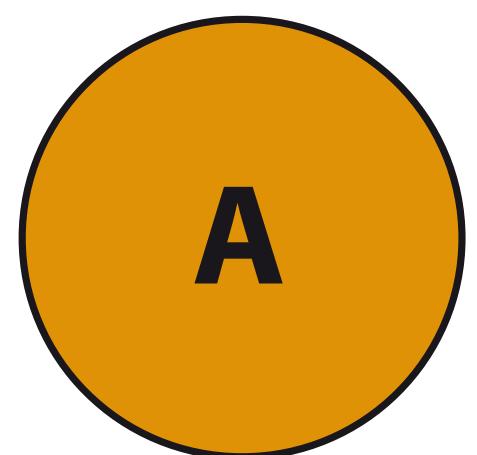


The de facto confirmation time is **6 blocks** (5 on top of the one including a certain transaction)

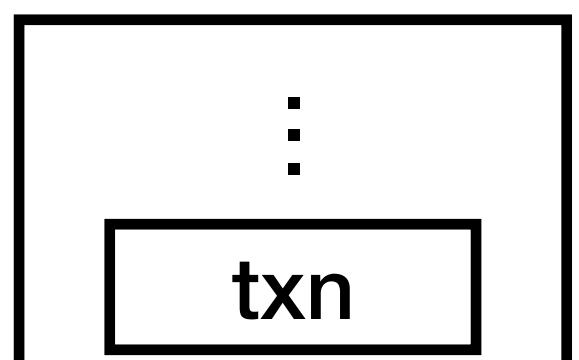
A's mempool



CONFIRMED TRANSACTIONS

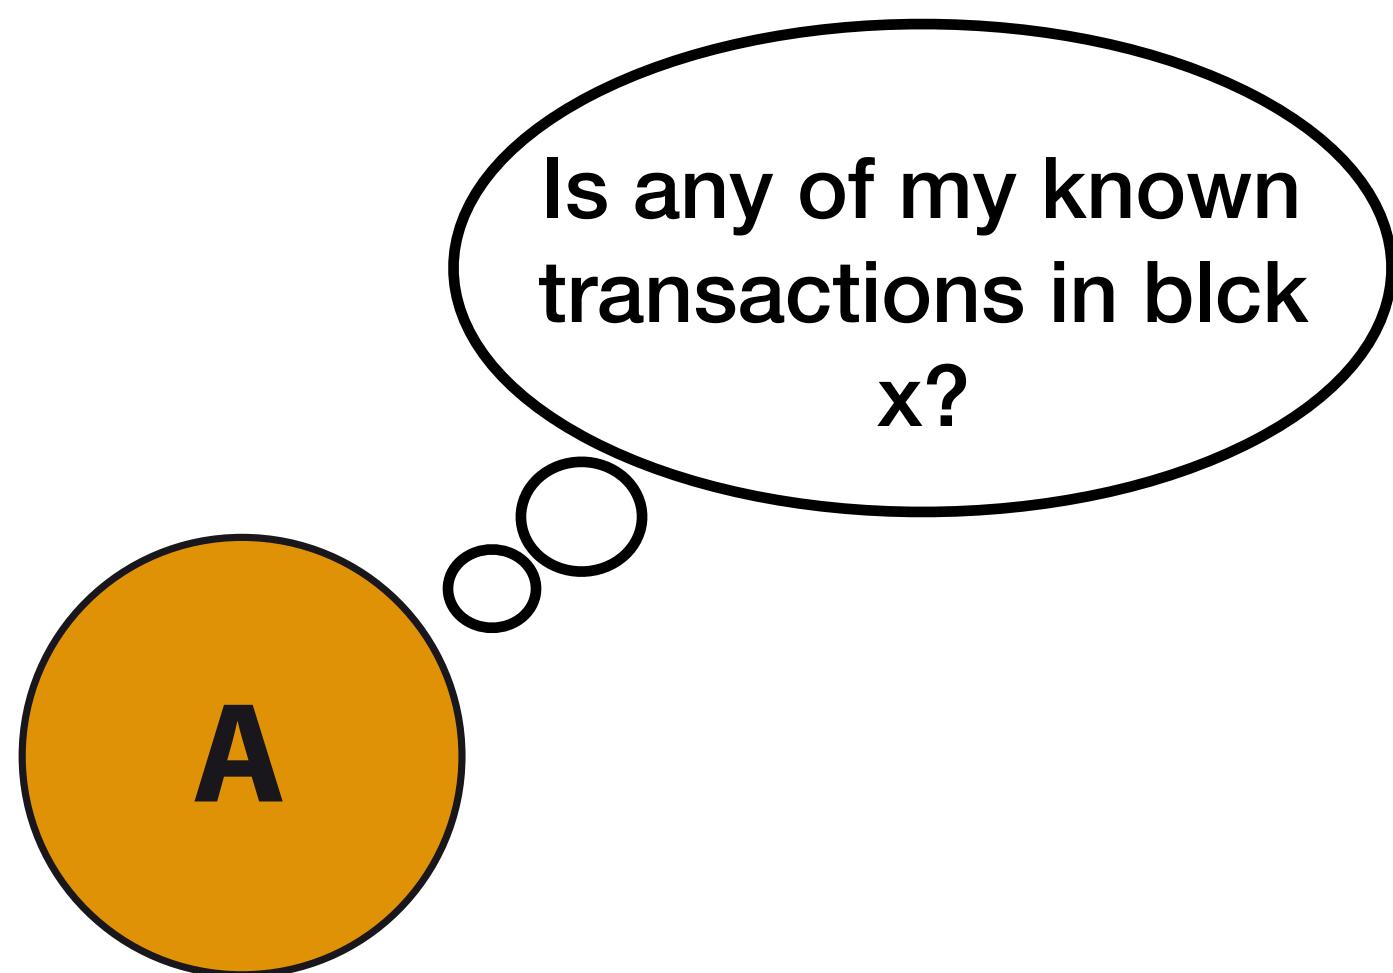


A's mempool

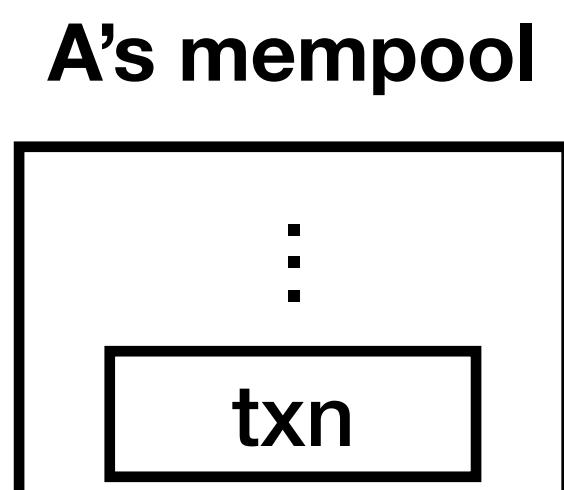


The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

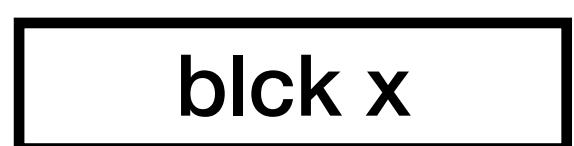
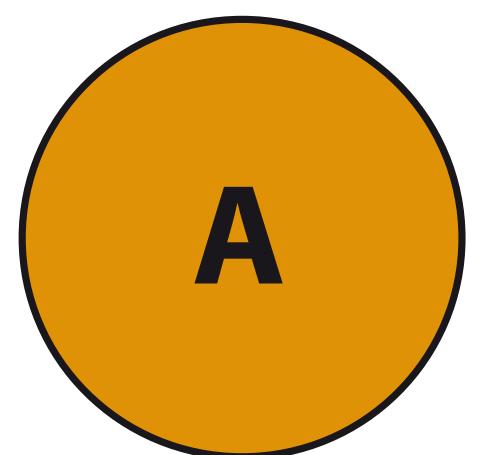
CONFIRMED TRANSACTIONS



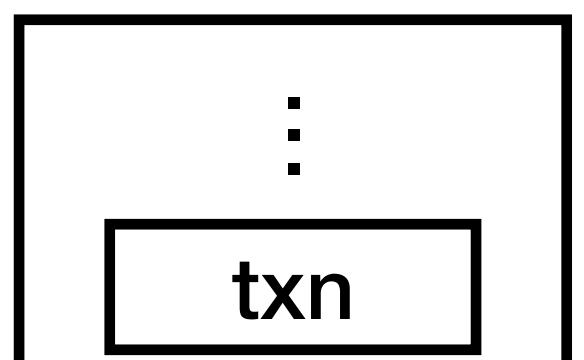
The de facto confirmation time is **6 blocks** (5 on top of the one including a certain transaction)



CONFIRMED TRANSACTIONS

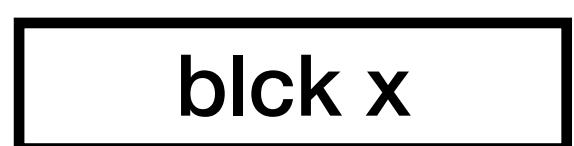
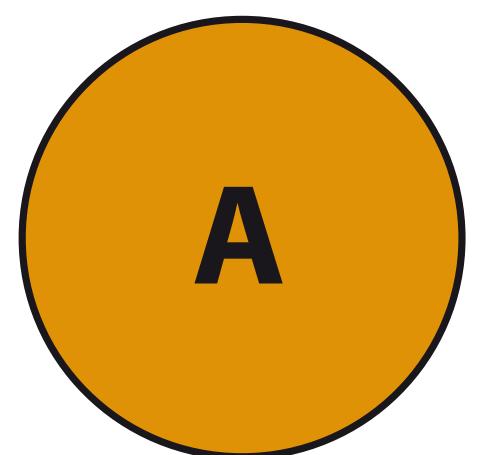


A's mempool

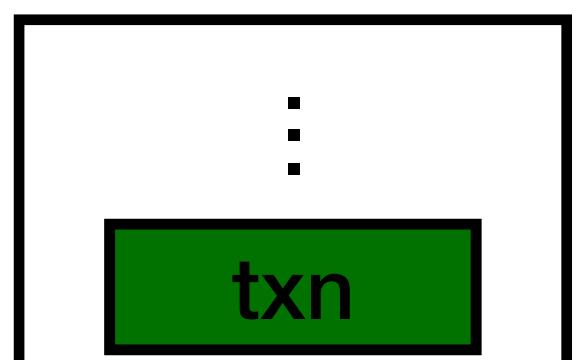


The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

CONFIRMED TRANSACTIONS

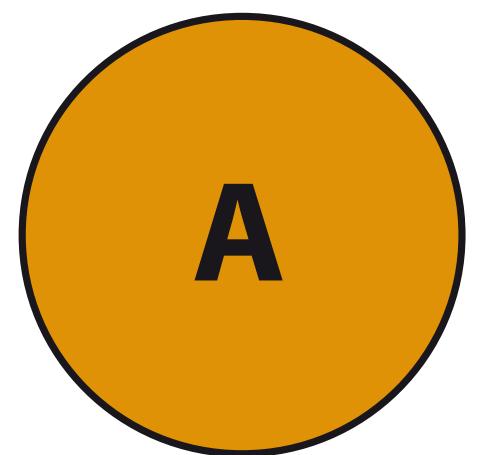


A's mempool



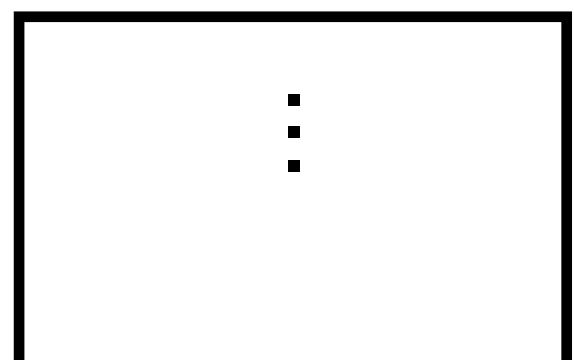
The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

CONFIRMED TRANSACTIONS



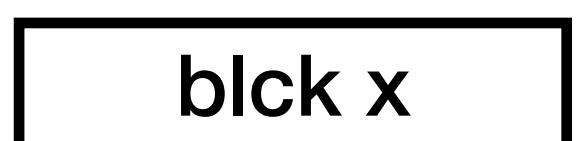
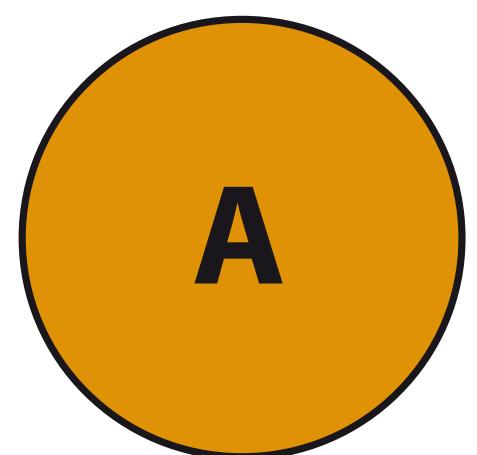
blk x

A's mempool

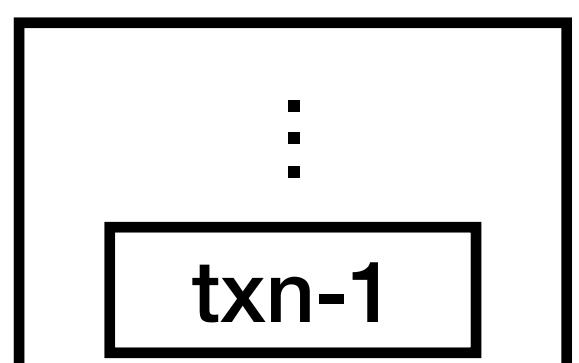


The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

CONFIRMED TRANSACTIONS

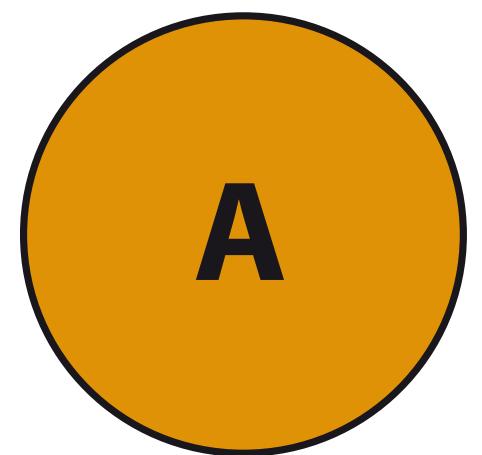


A's mempool

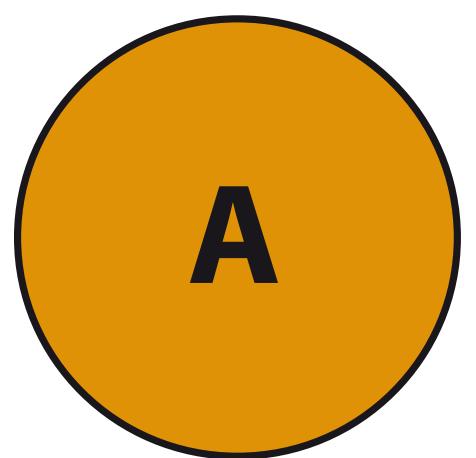


The de facto confirmation time is
6 blocks (5 on top of the one
including a certain transaction)

DOUBLE-SPENDING TRANSACTIONS (1/2)

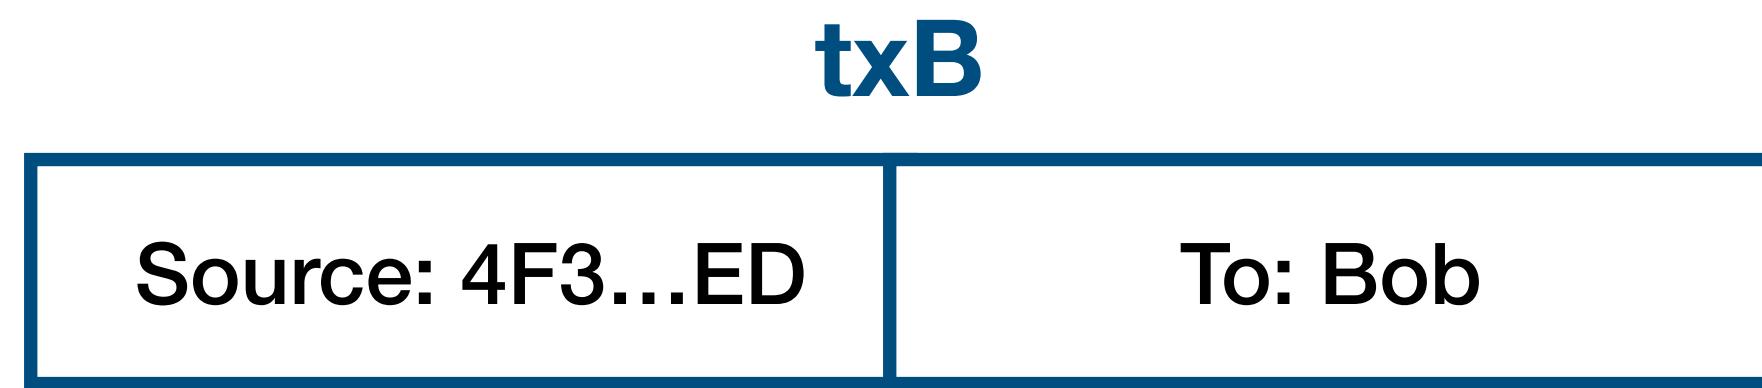
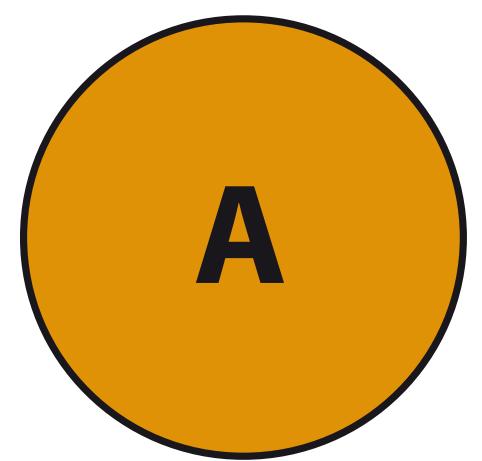


DOUBLE-SPENDING TRANSACTIONS (1/2)



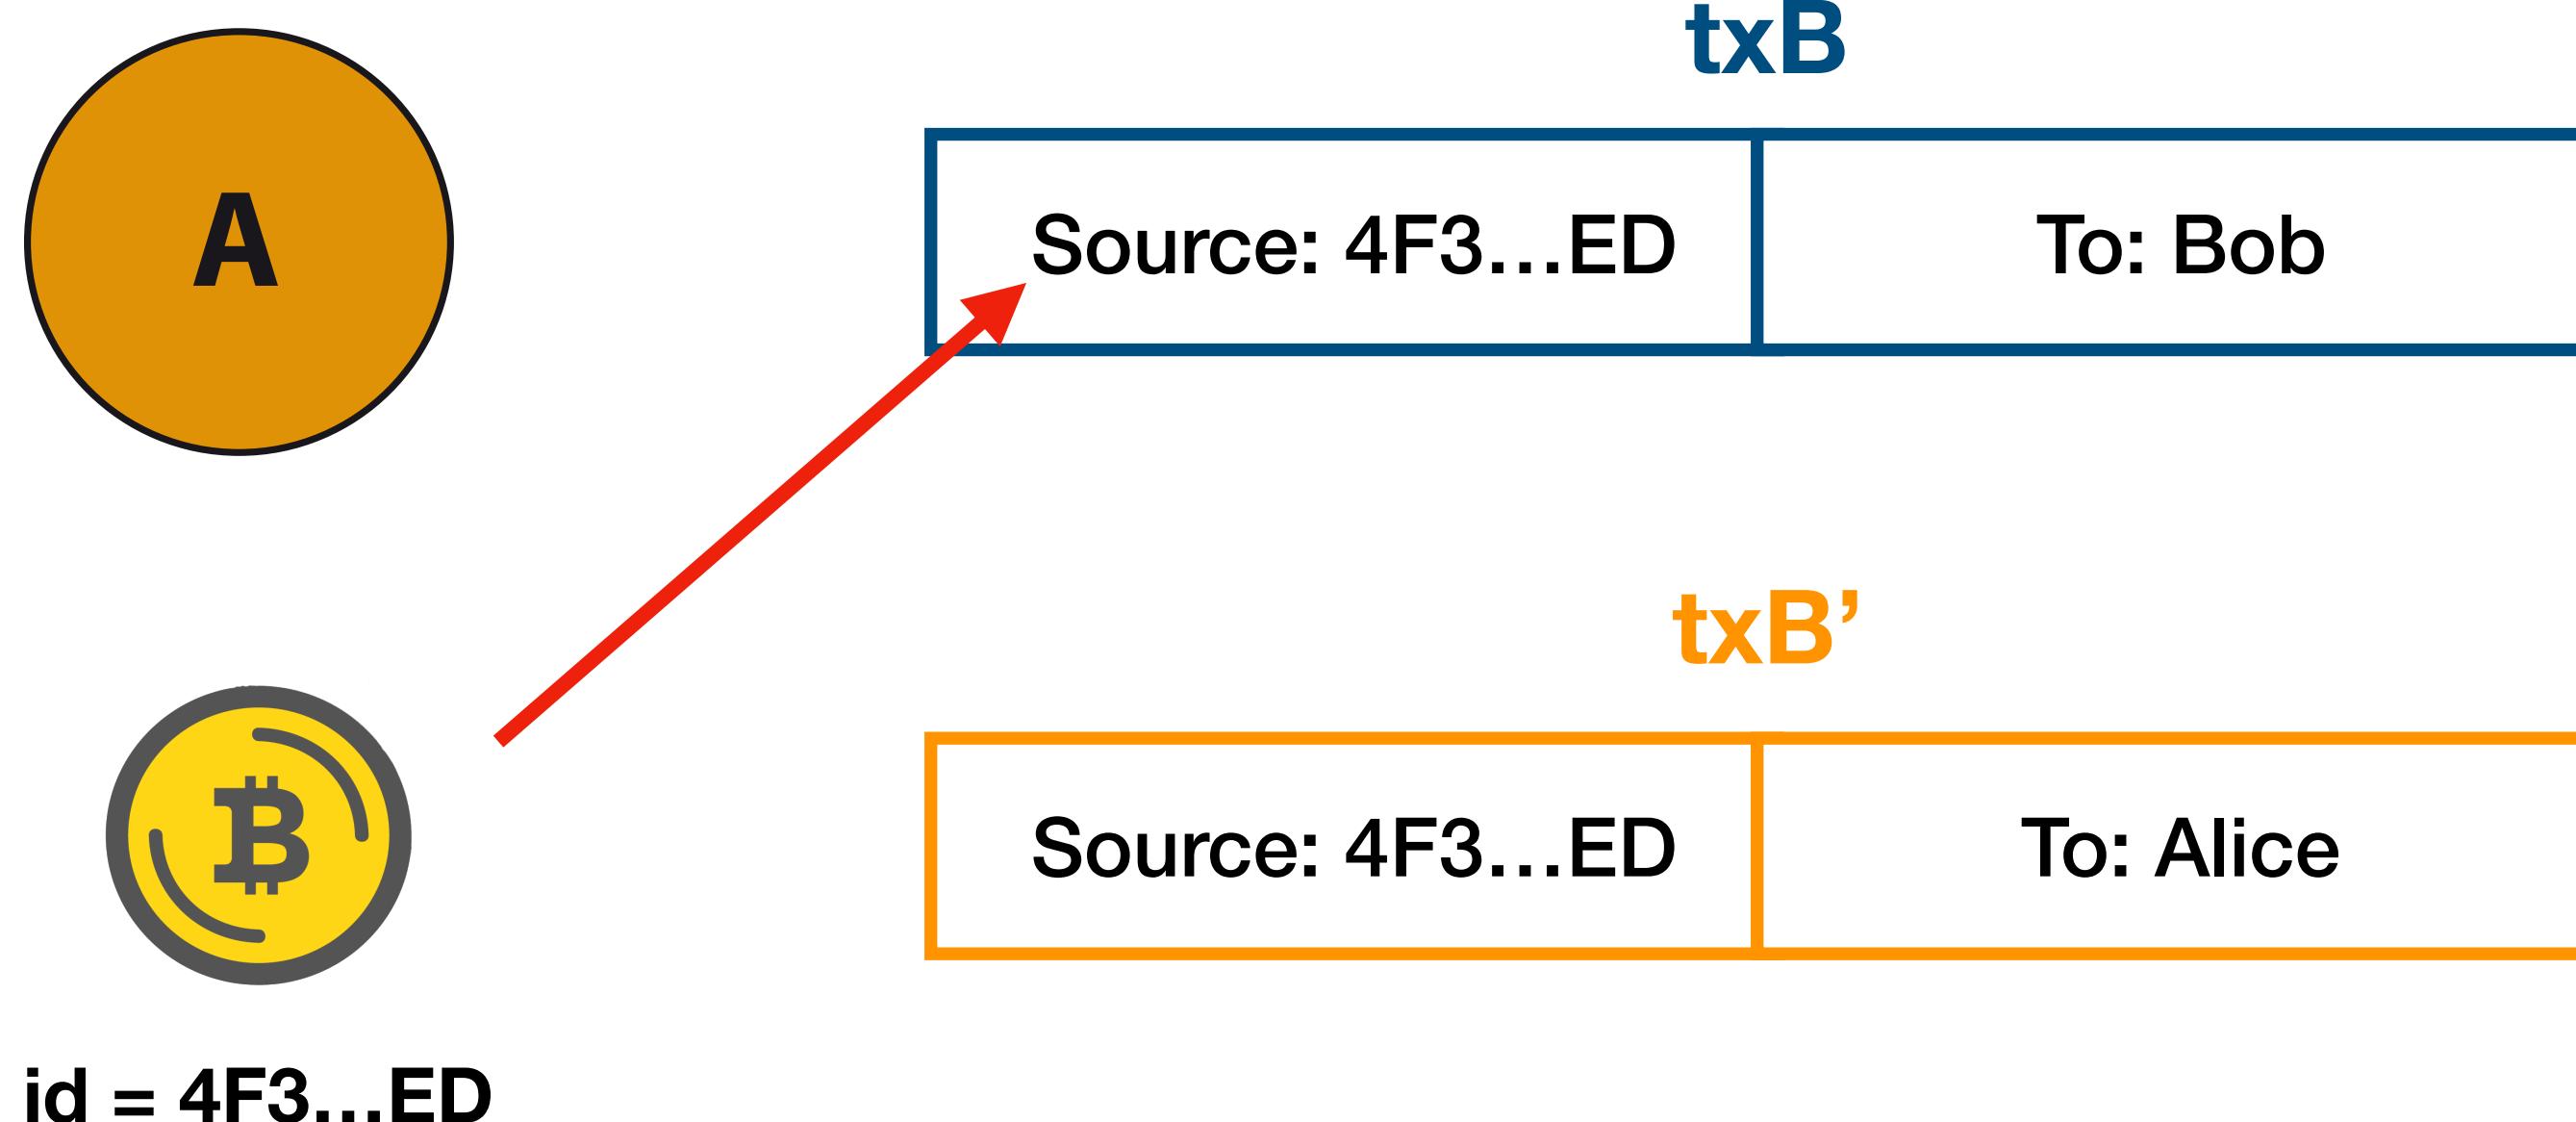
id = 4F3...ED

DOUBLE-SPENDING TRANSACTIONS (1/2)

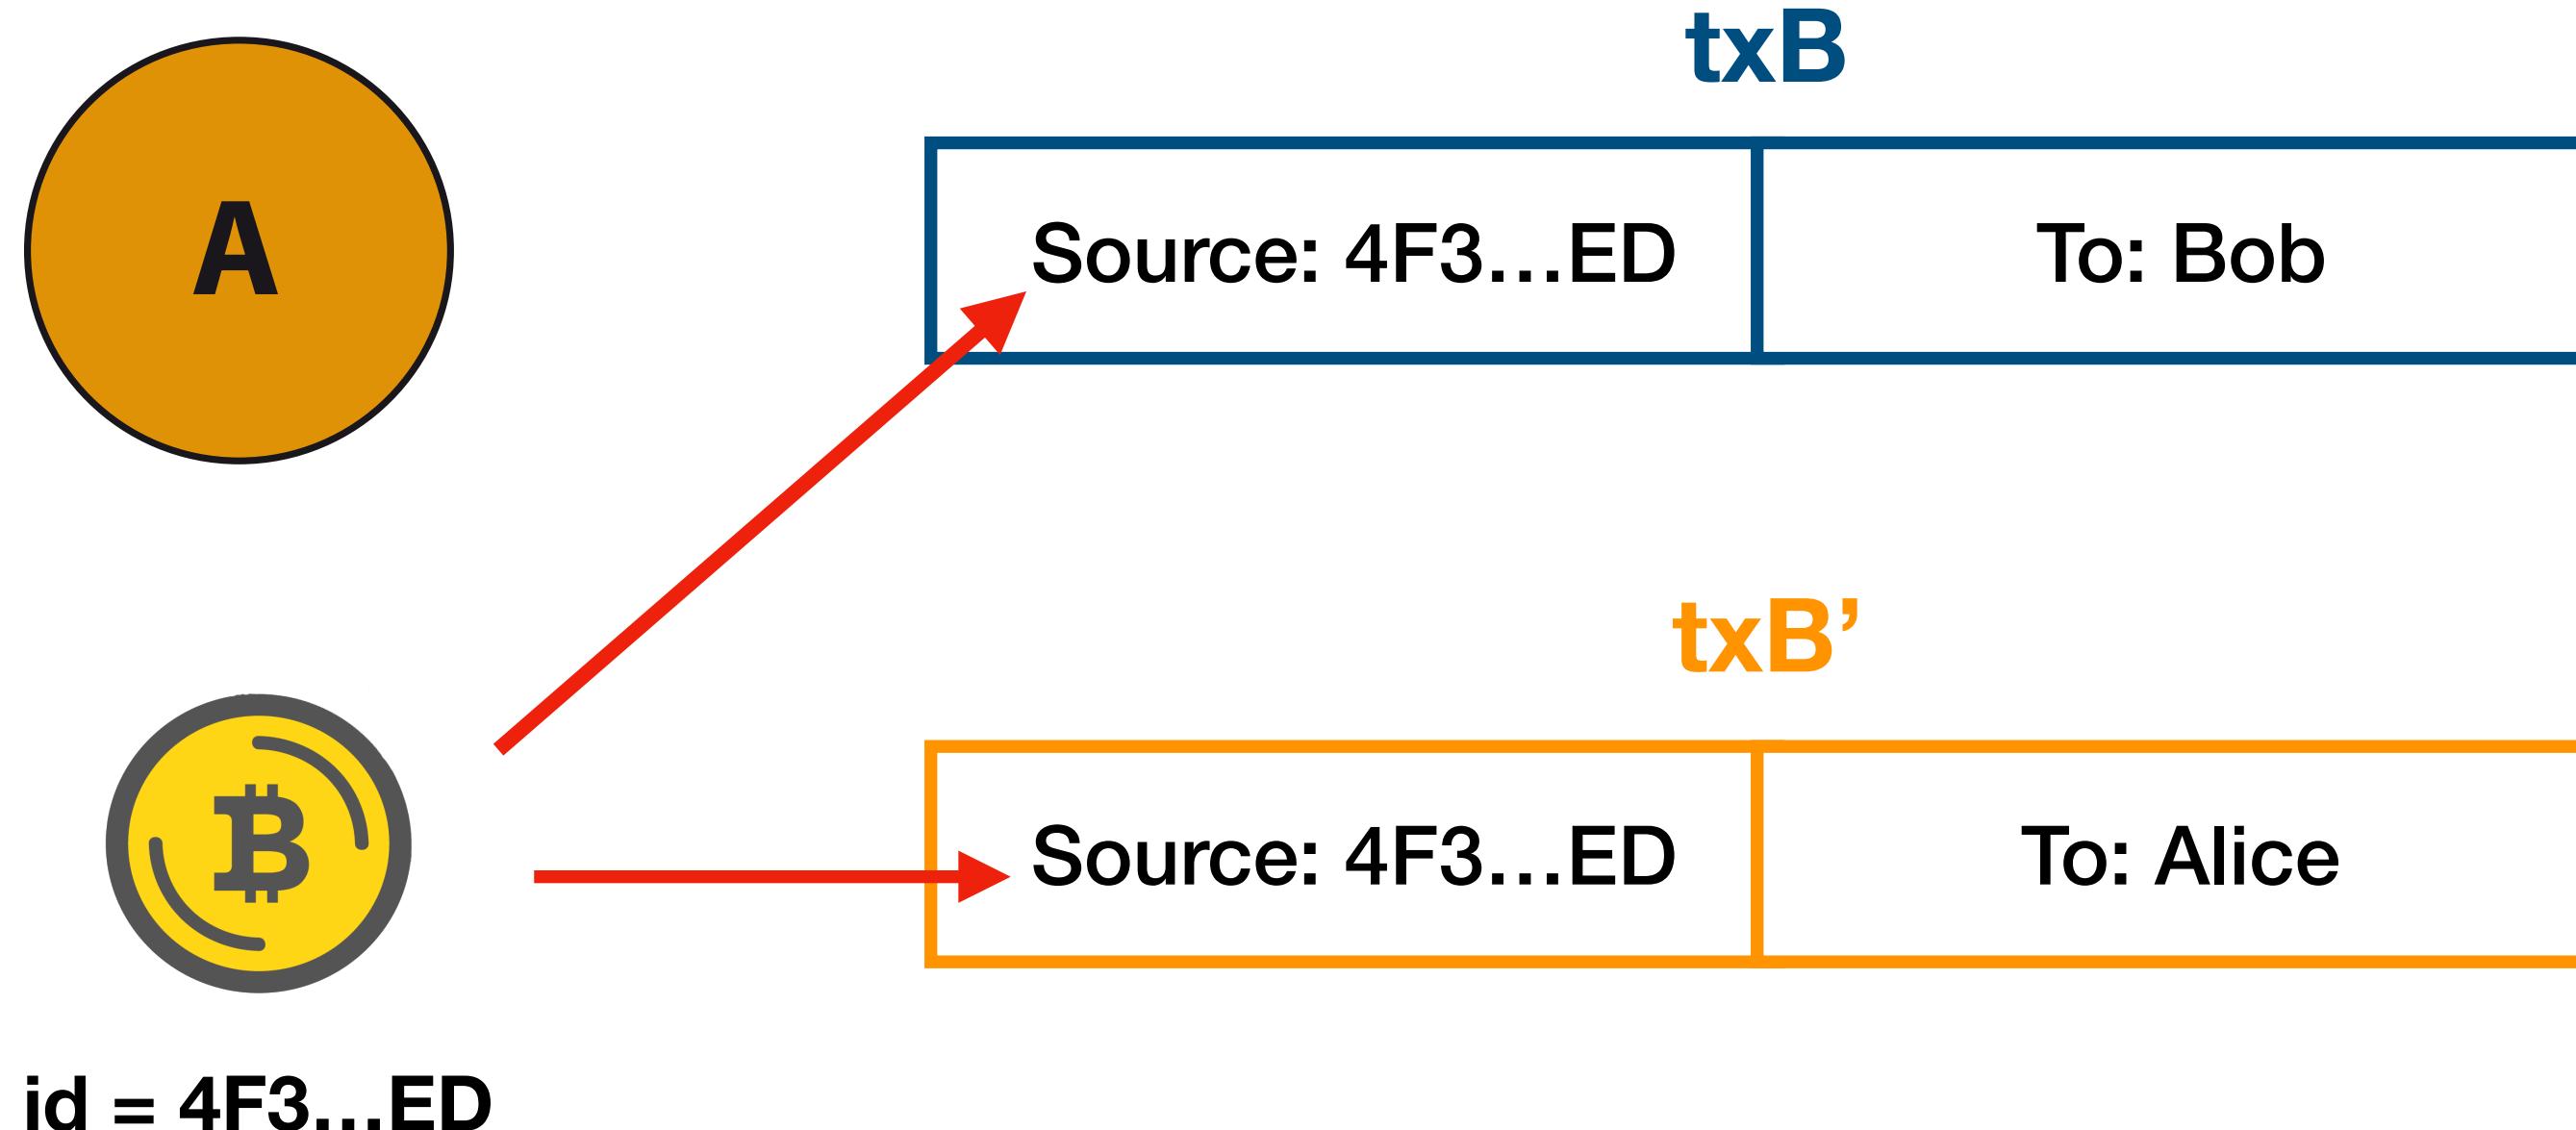


id = 4F3...ED

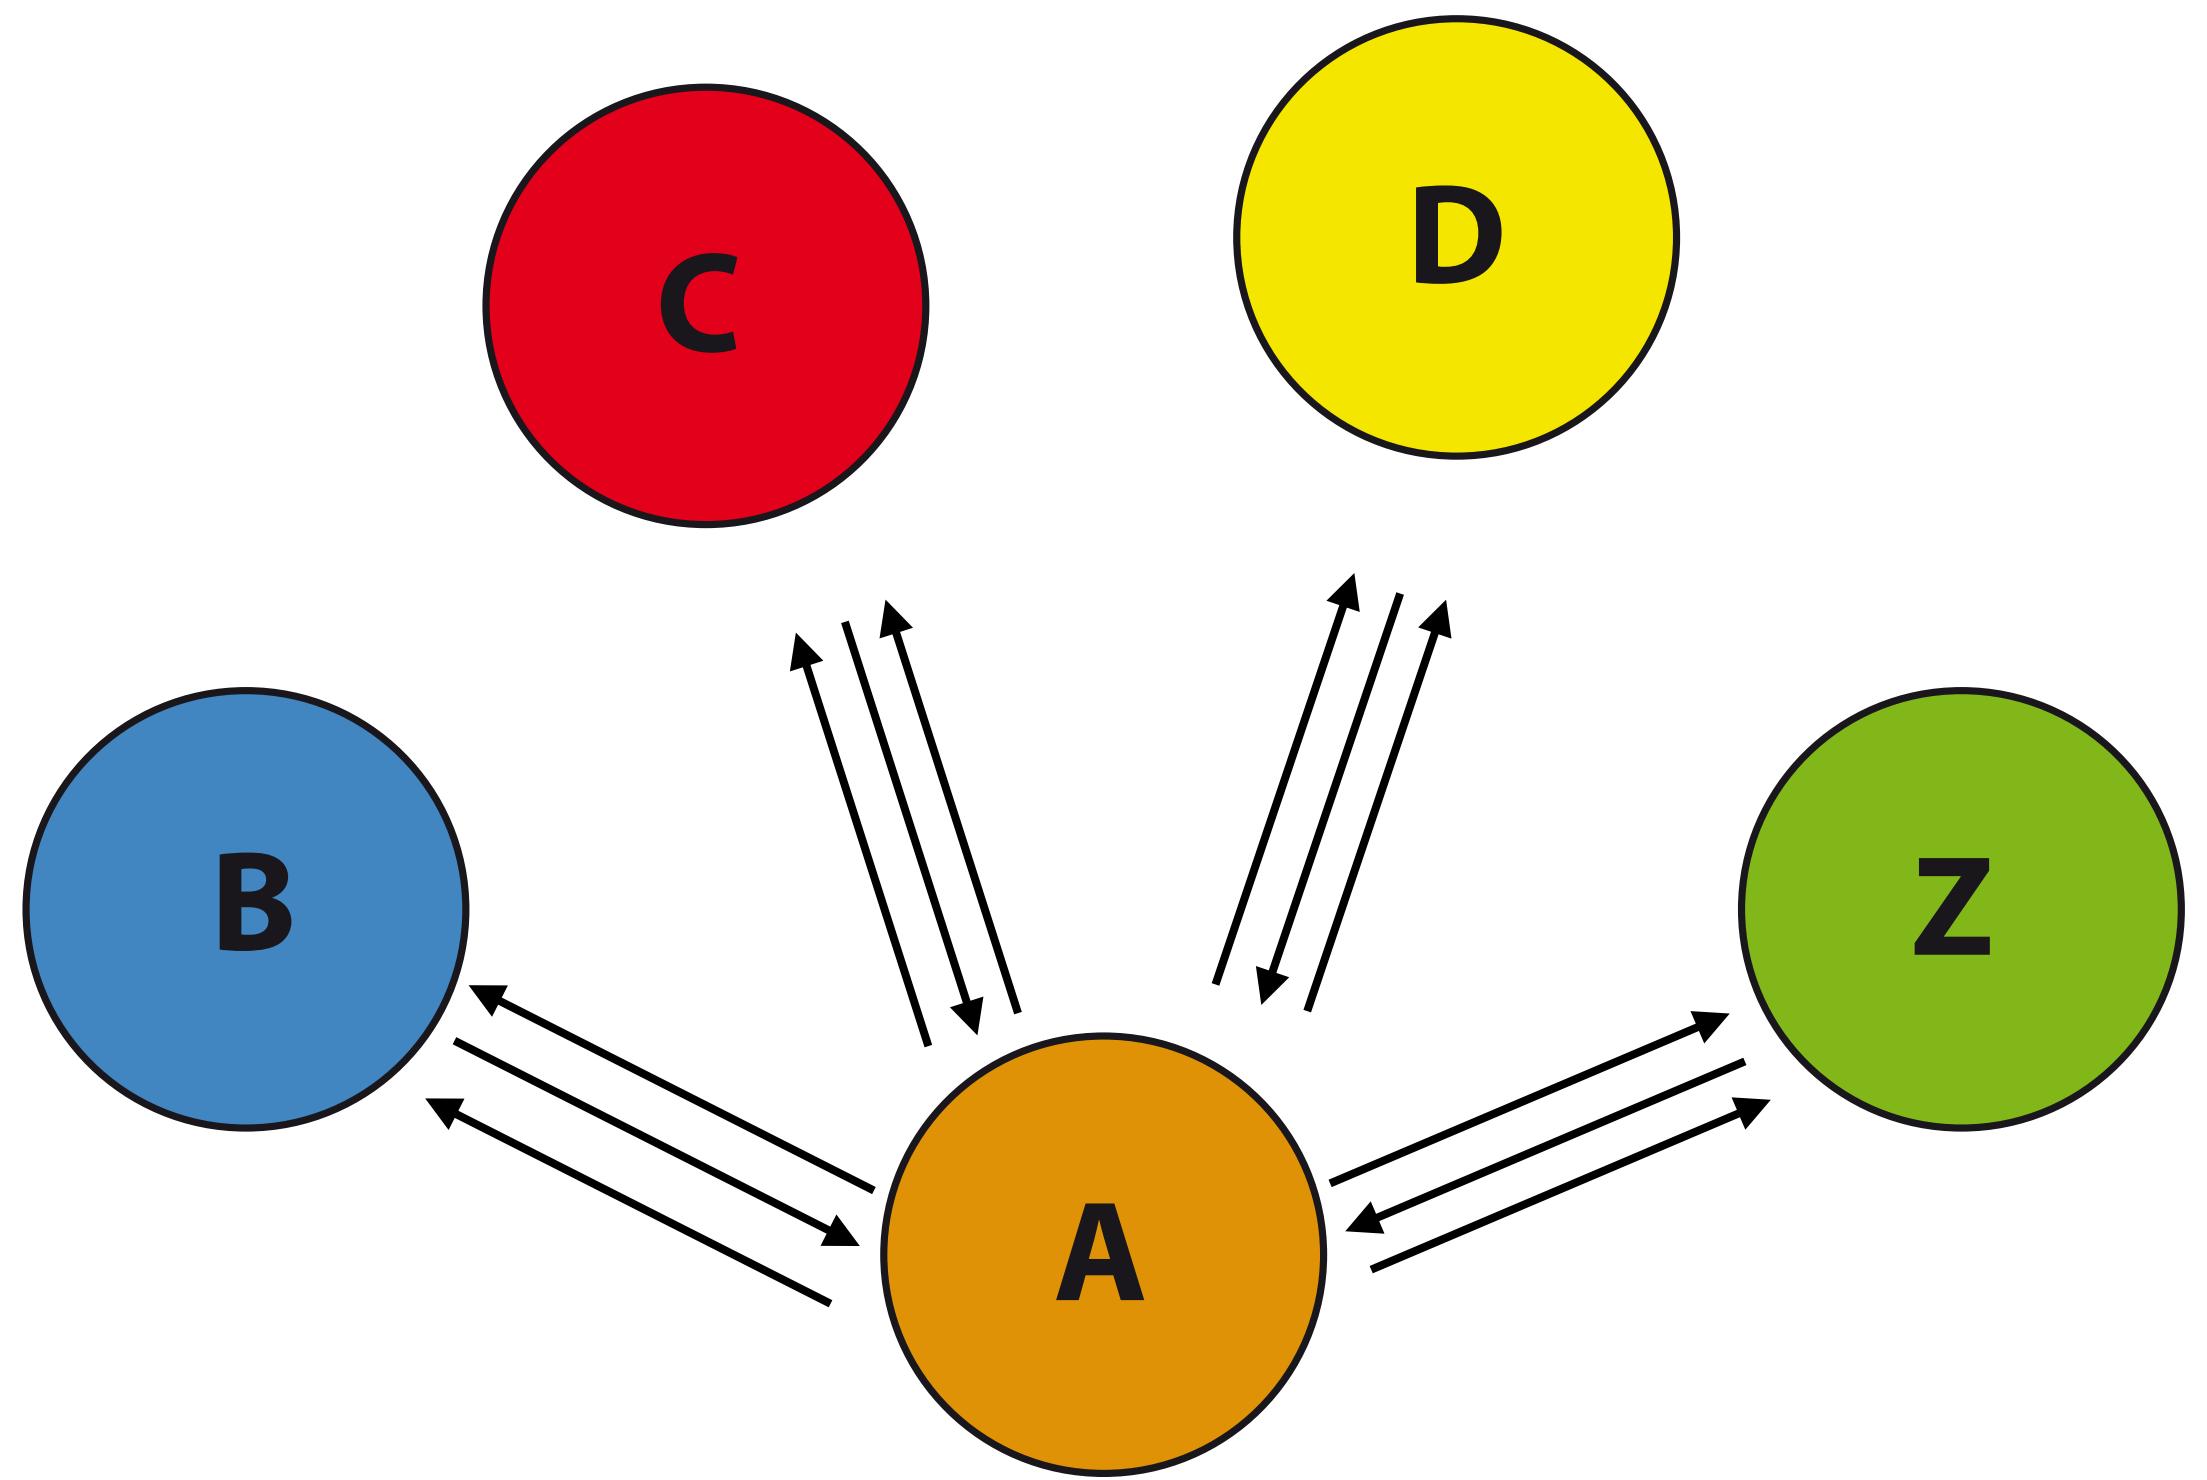
DOUBLE-SPENDING TRANSACTIONS (1/2)



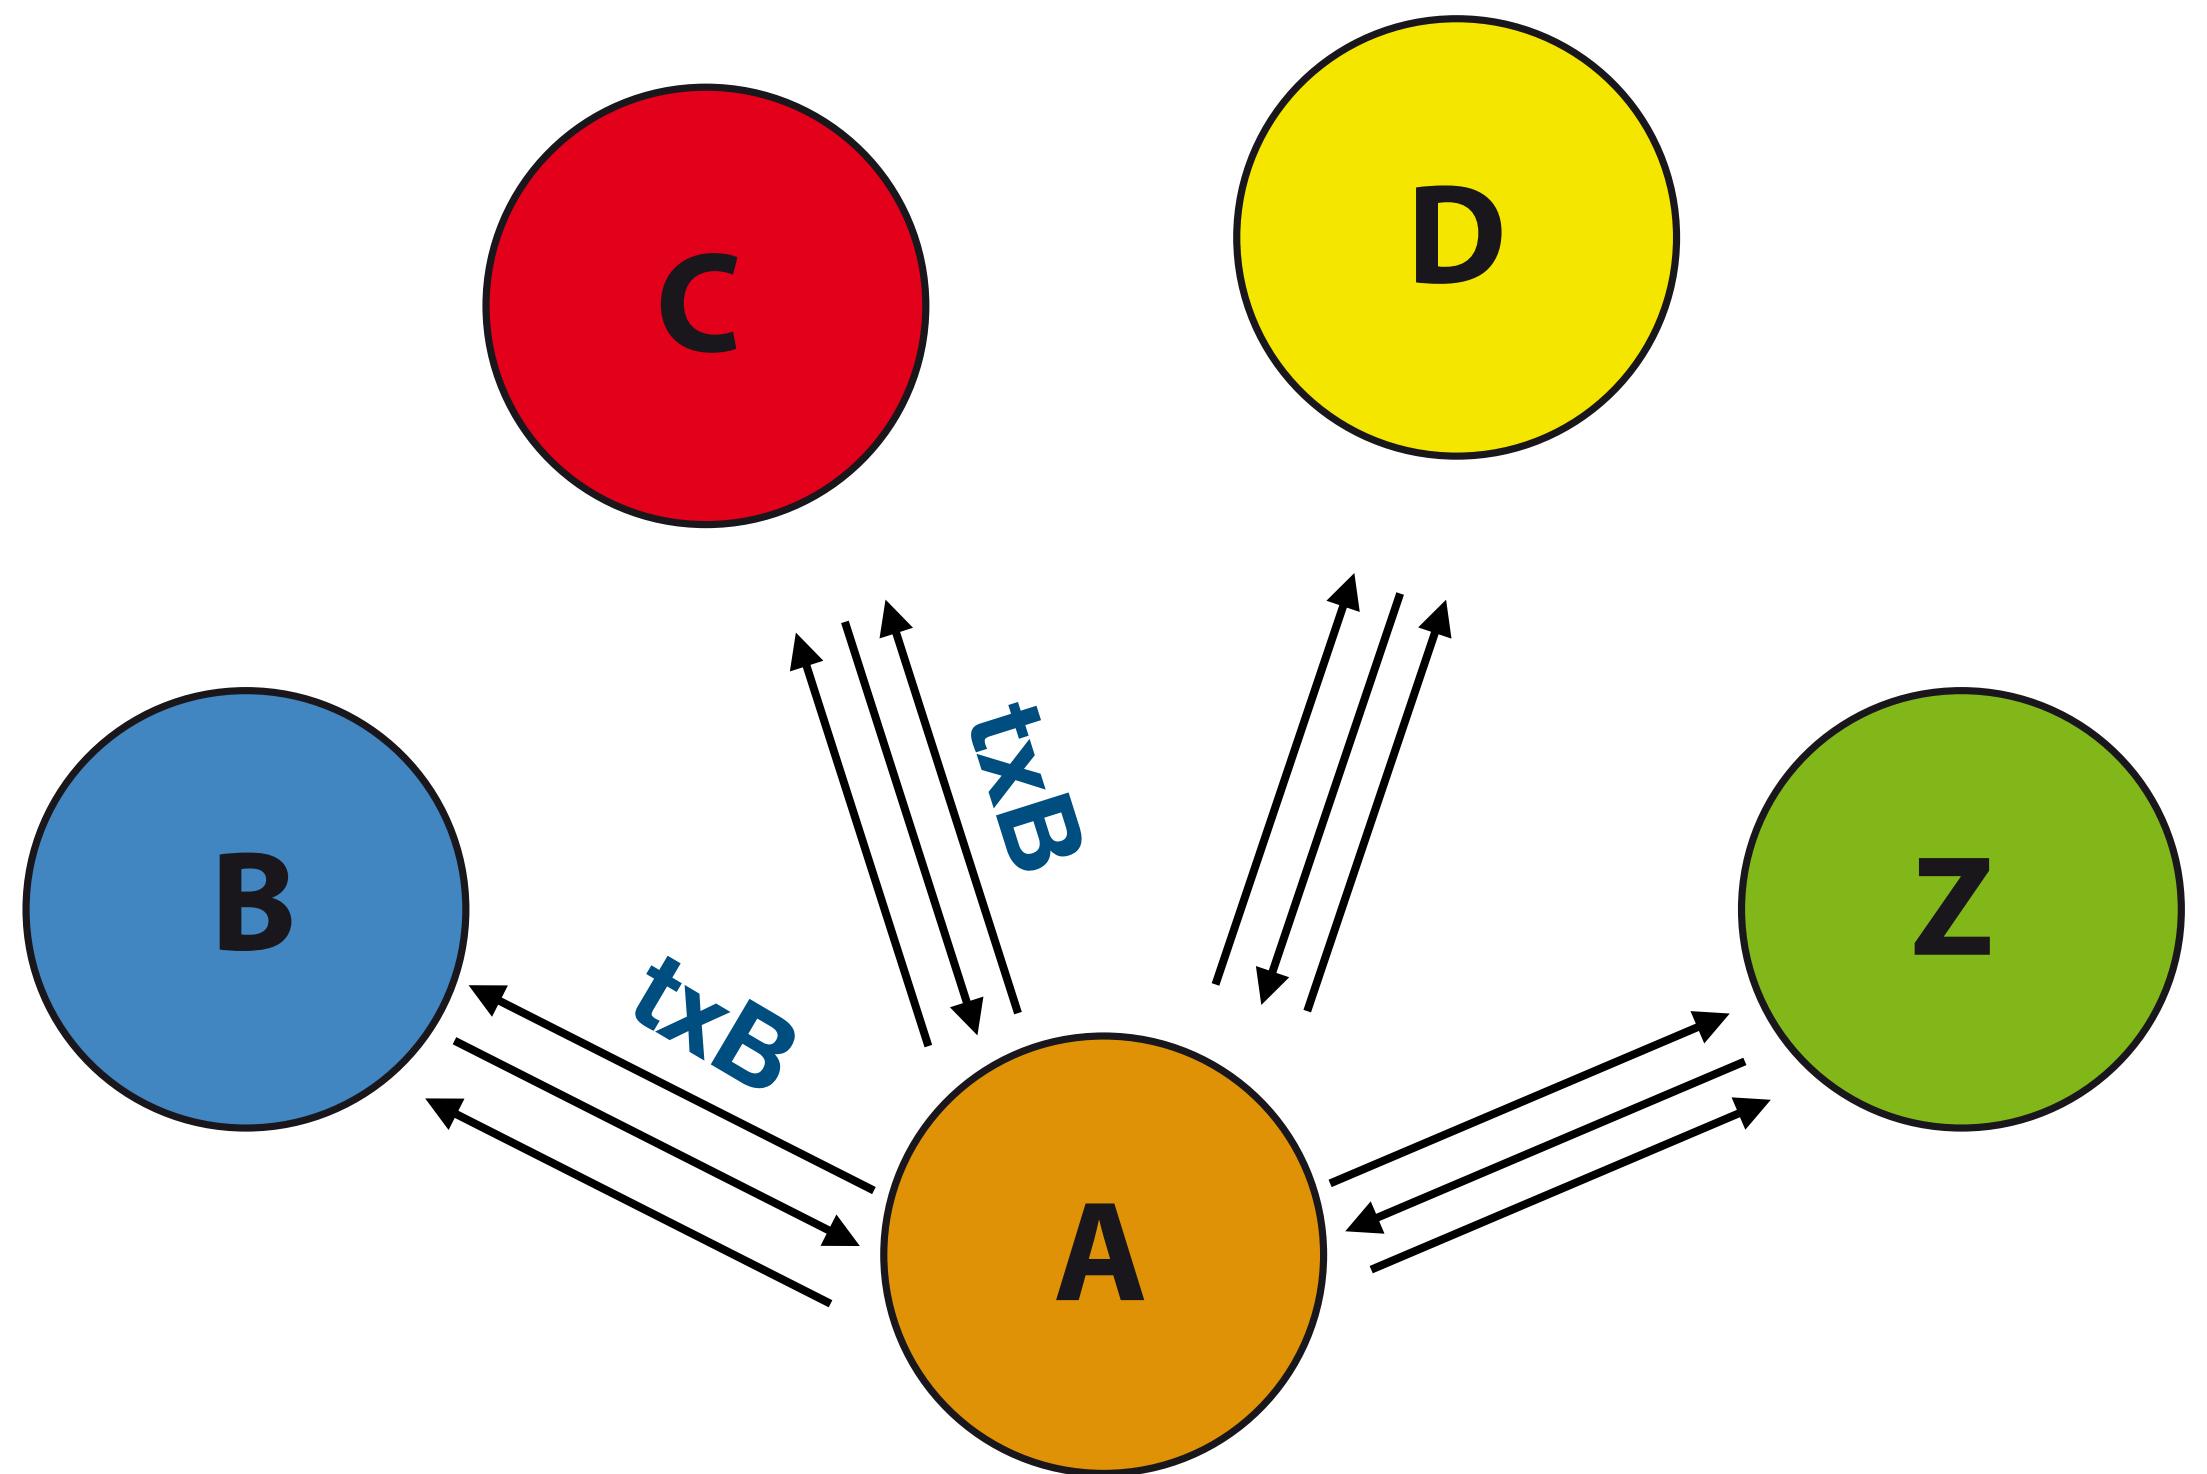
DOUBLE-SPENDING TRANSACTIONS (1/2)



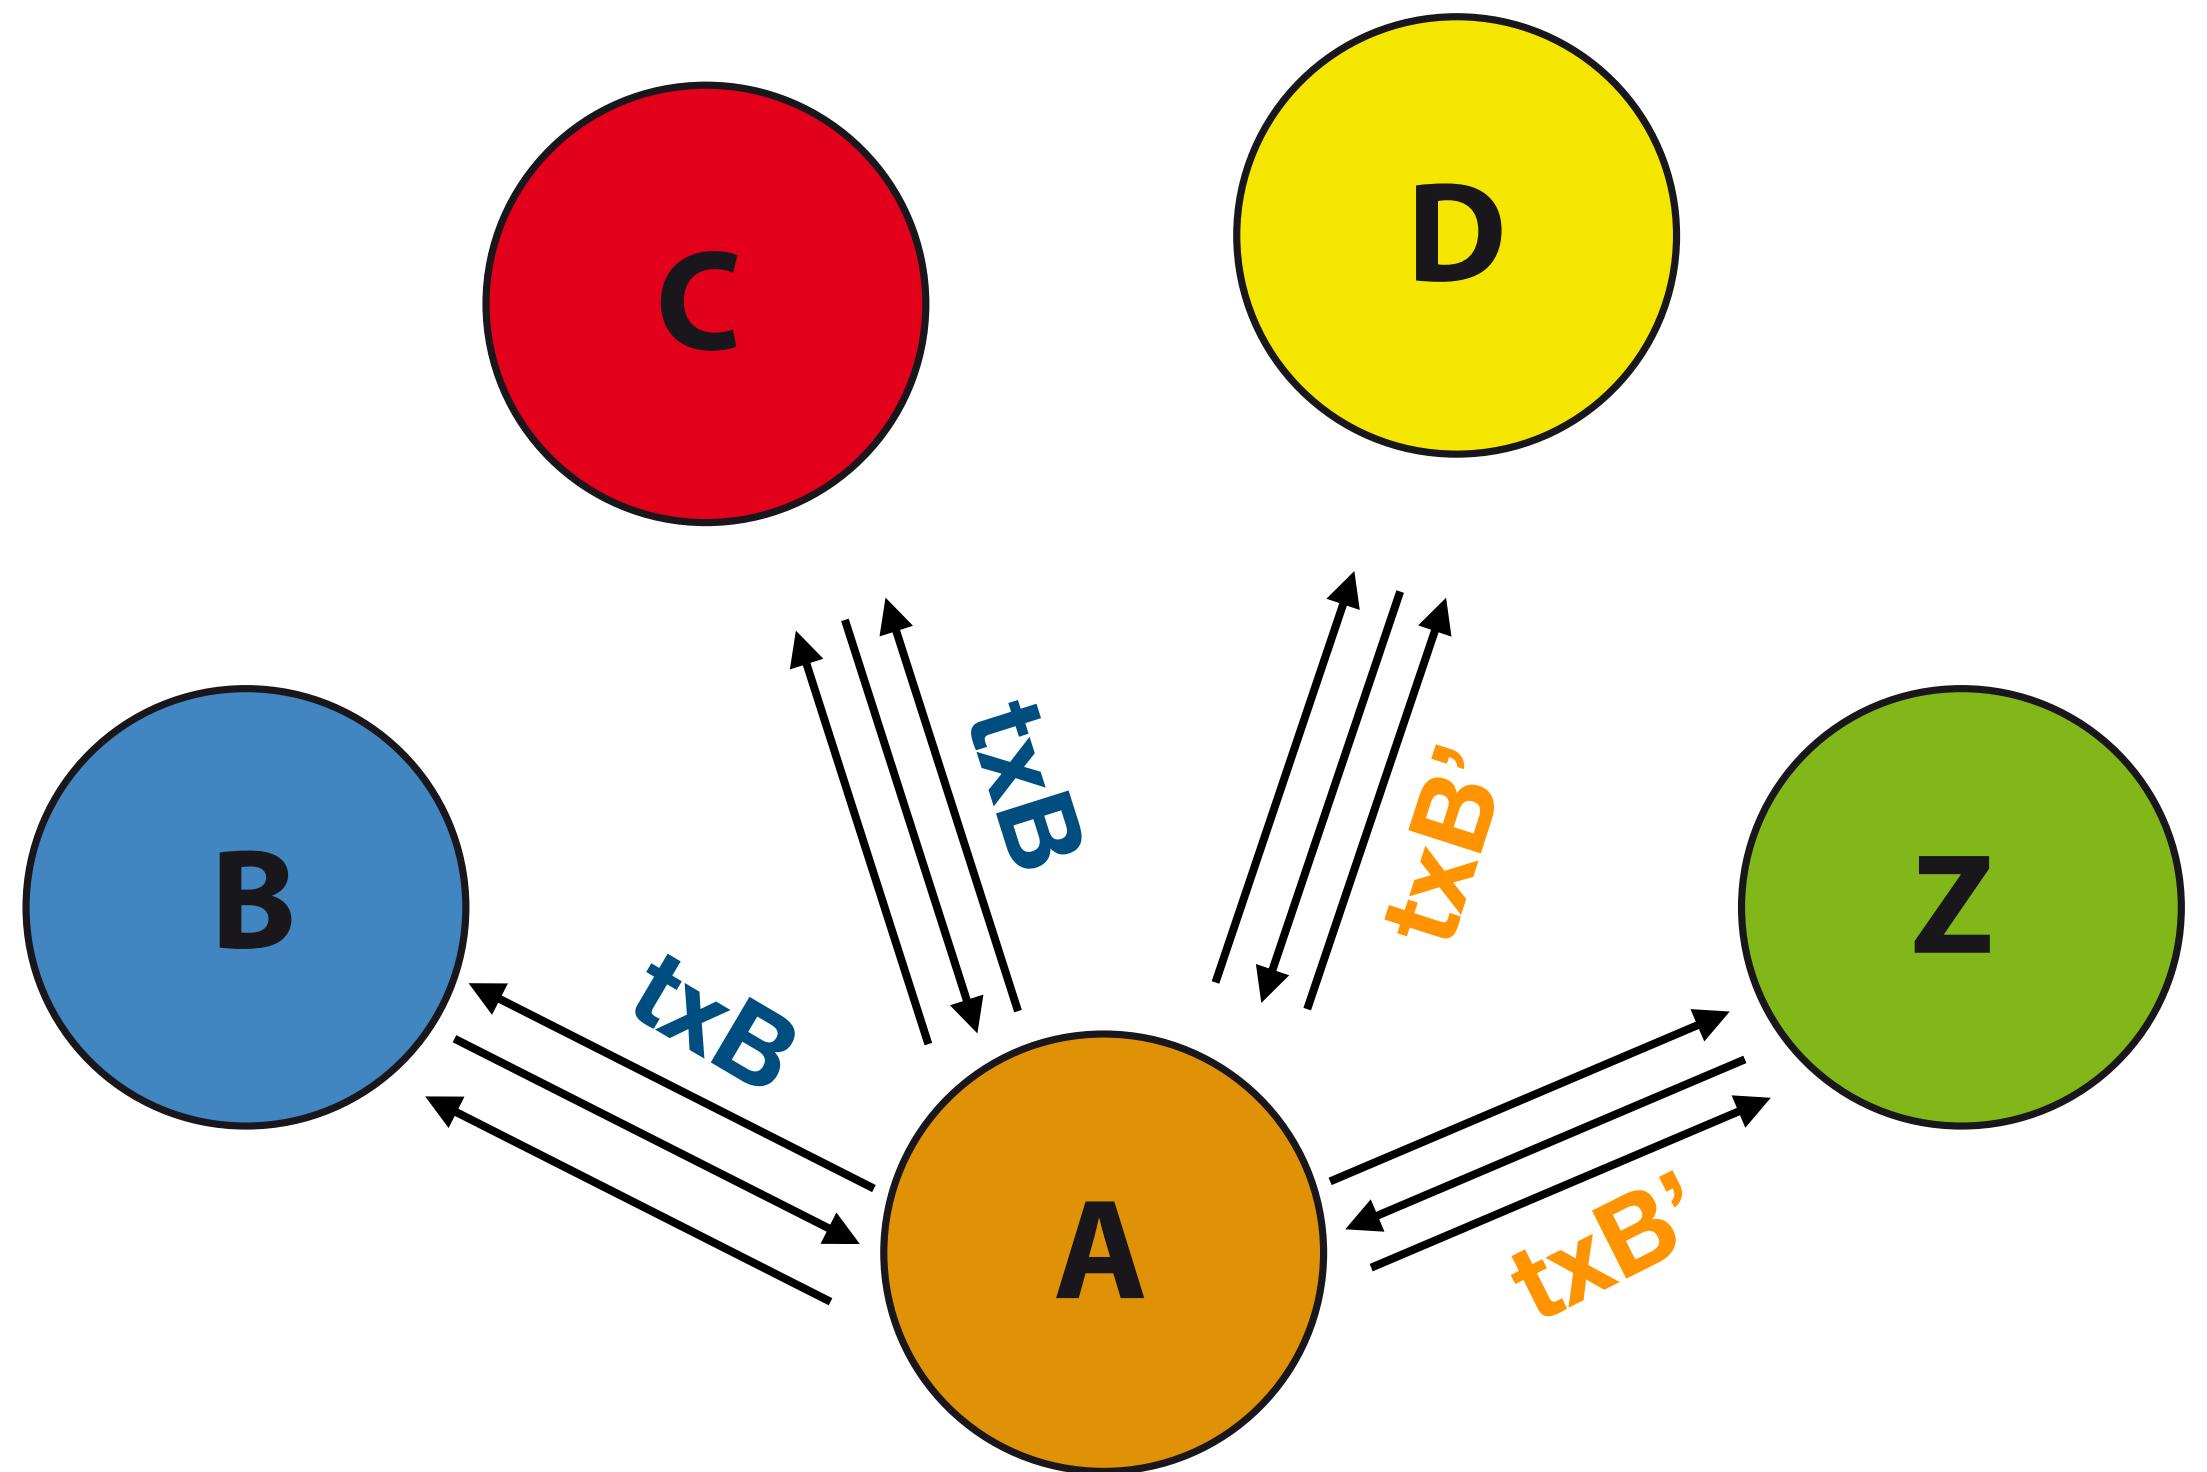
DOUBLE-SPENDING TRANSACTIONS (2/2)



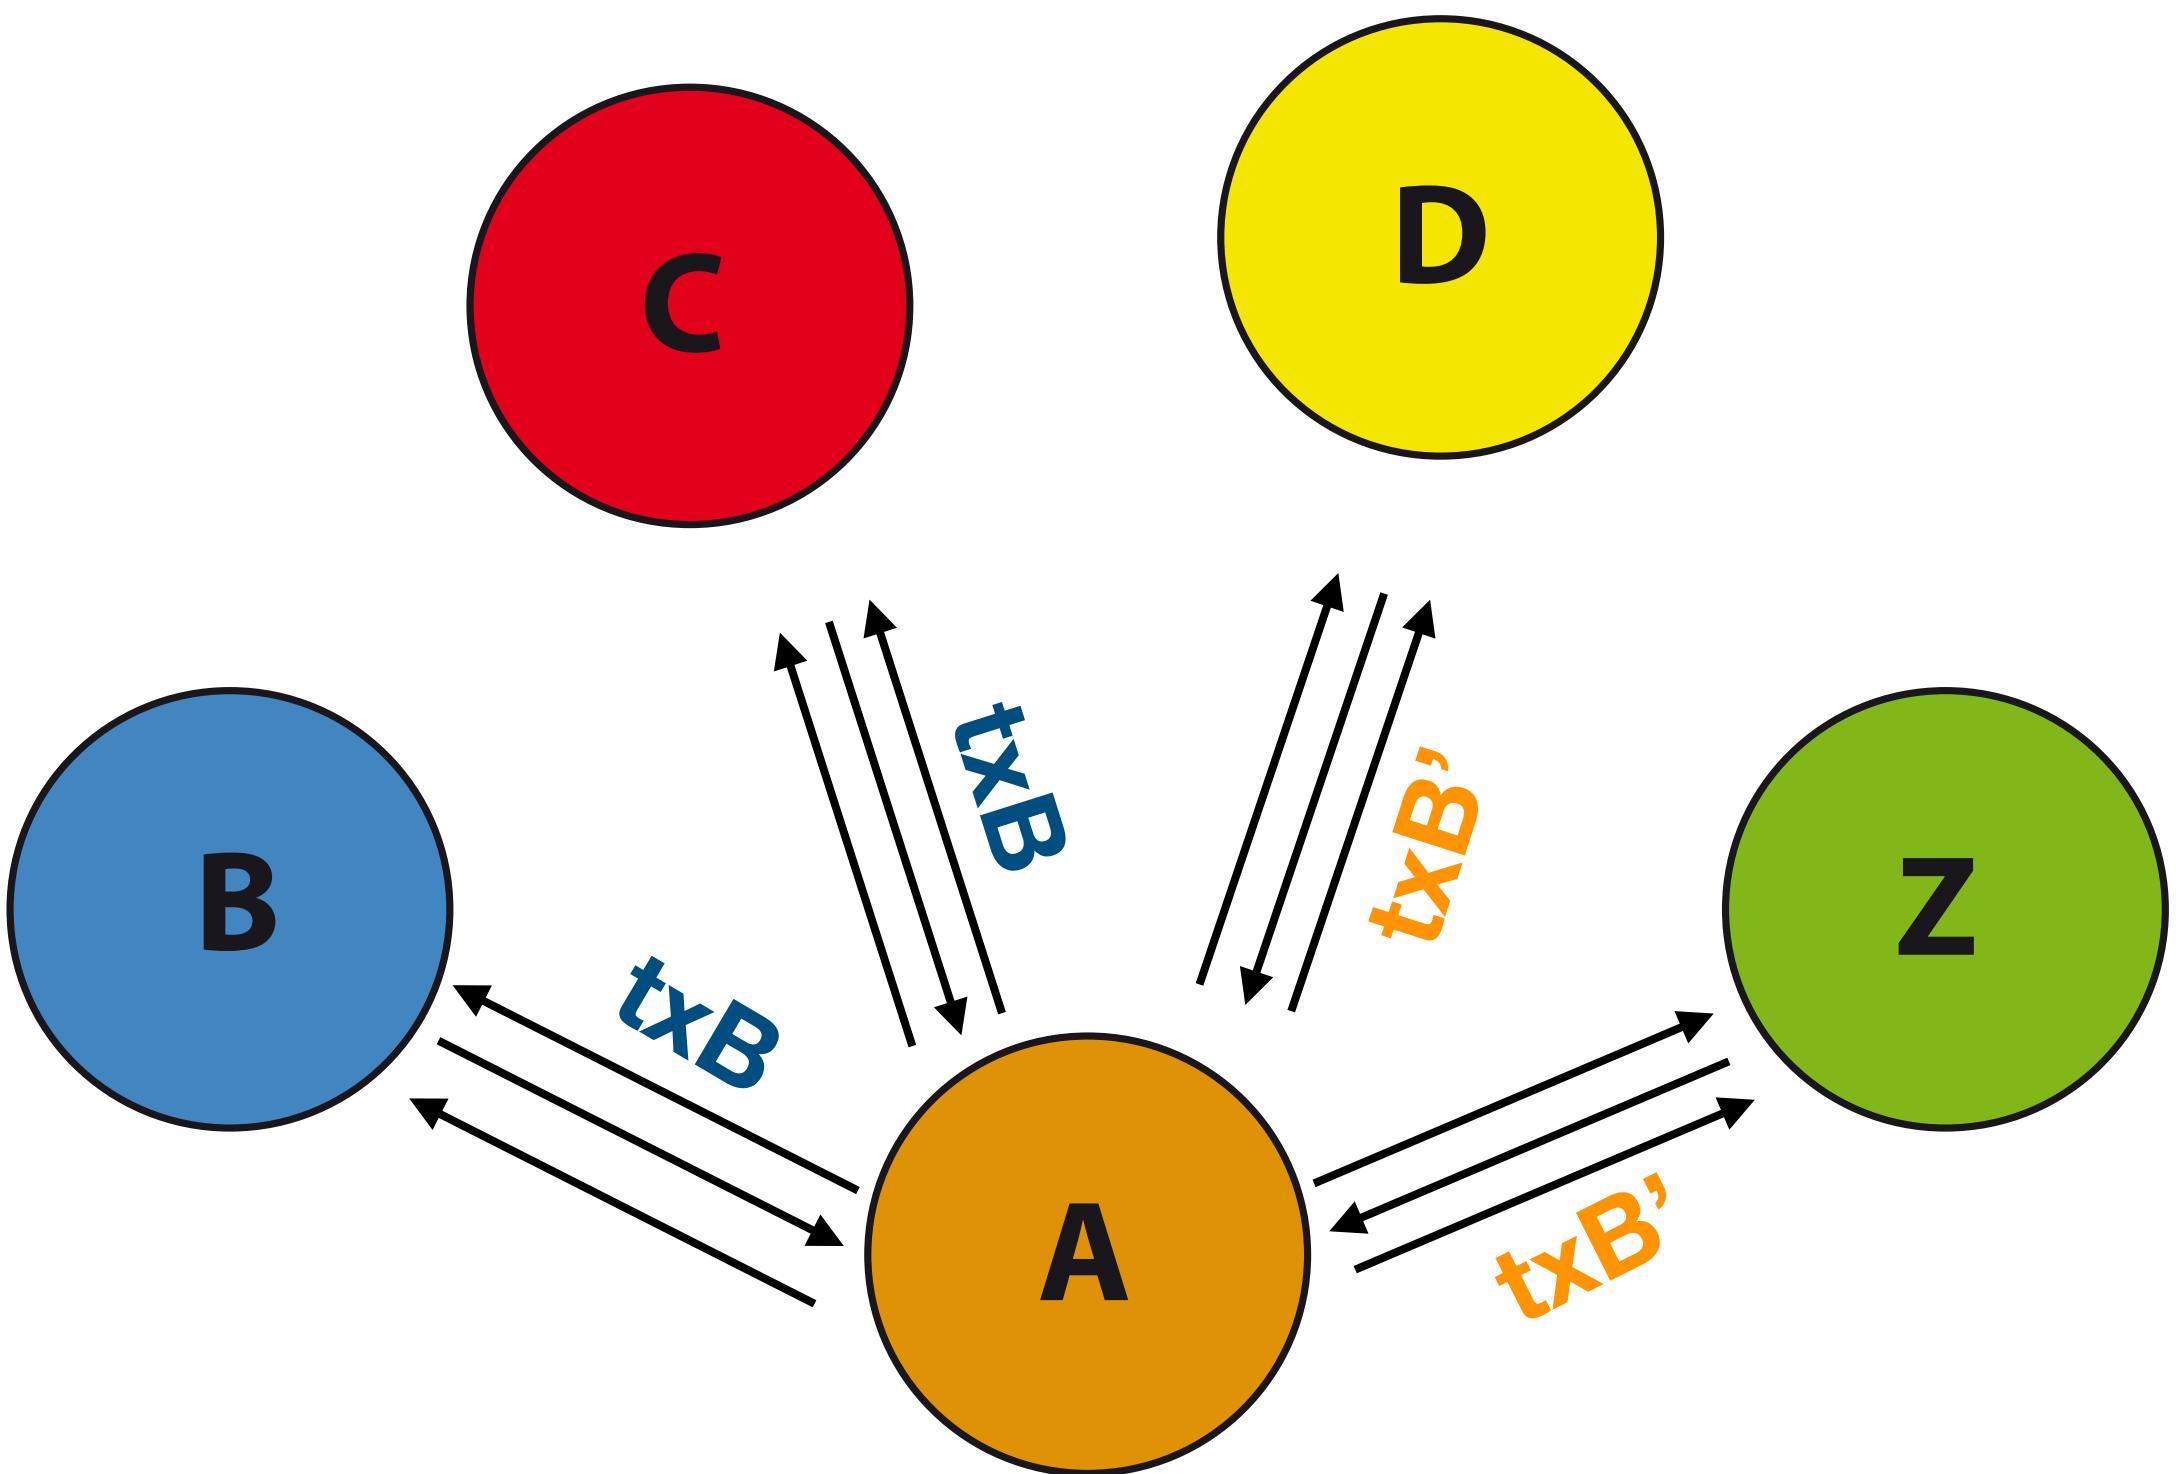
DOUBLE-SPENDING TRANSACTIONS (2/2)



DOUBLE-SPENDING TRANSACTIONS (2/2)

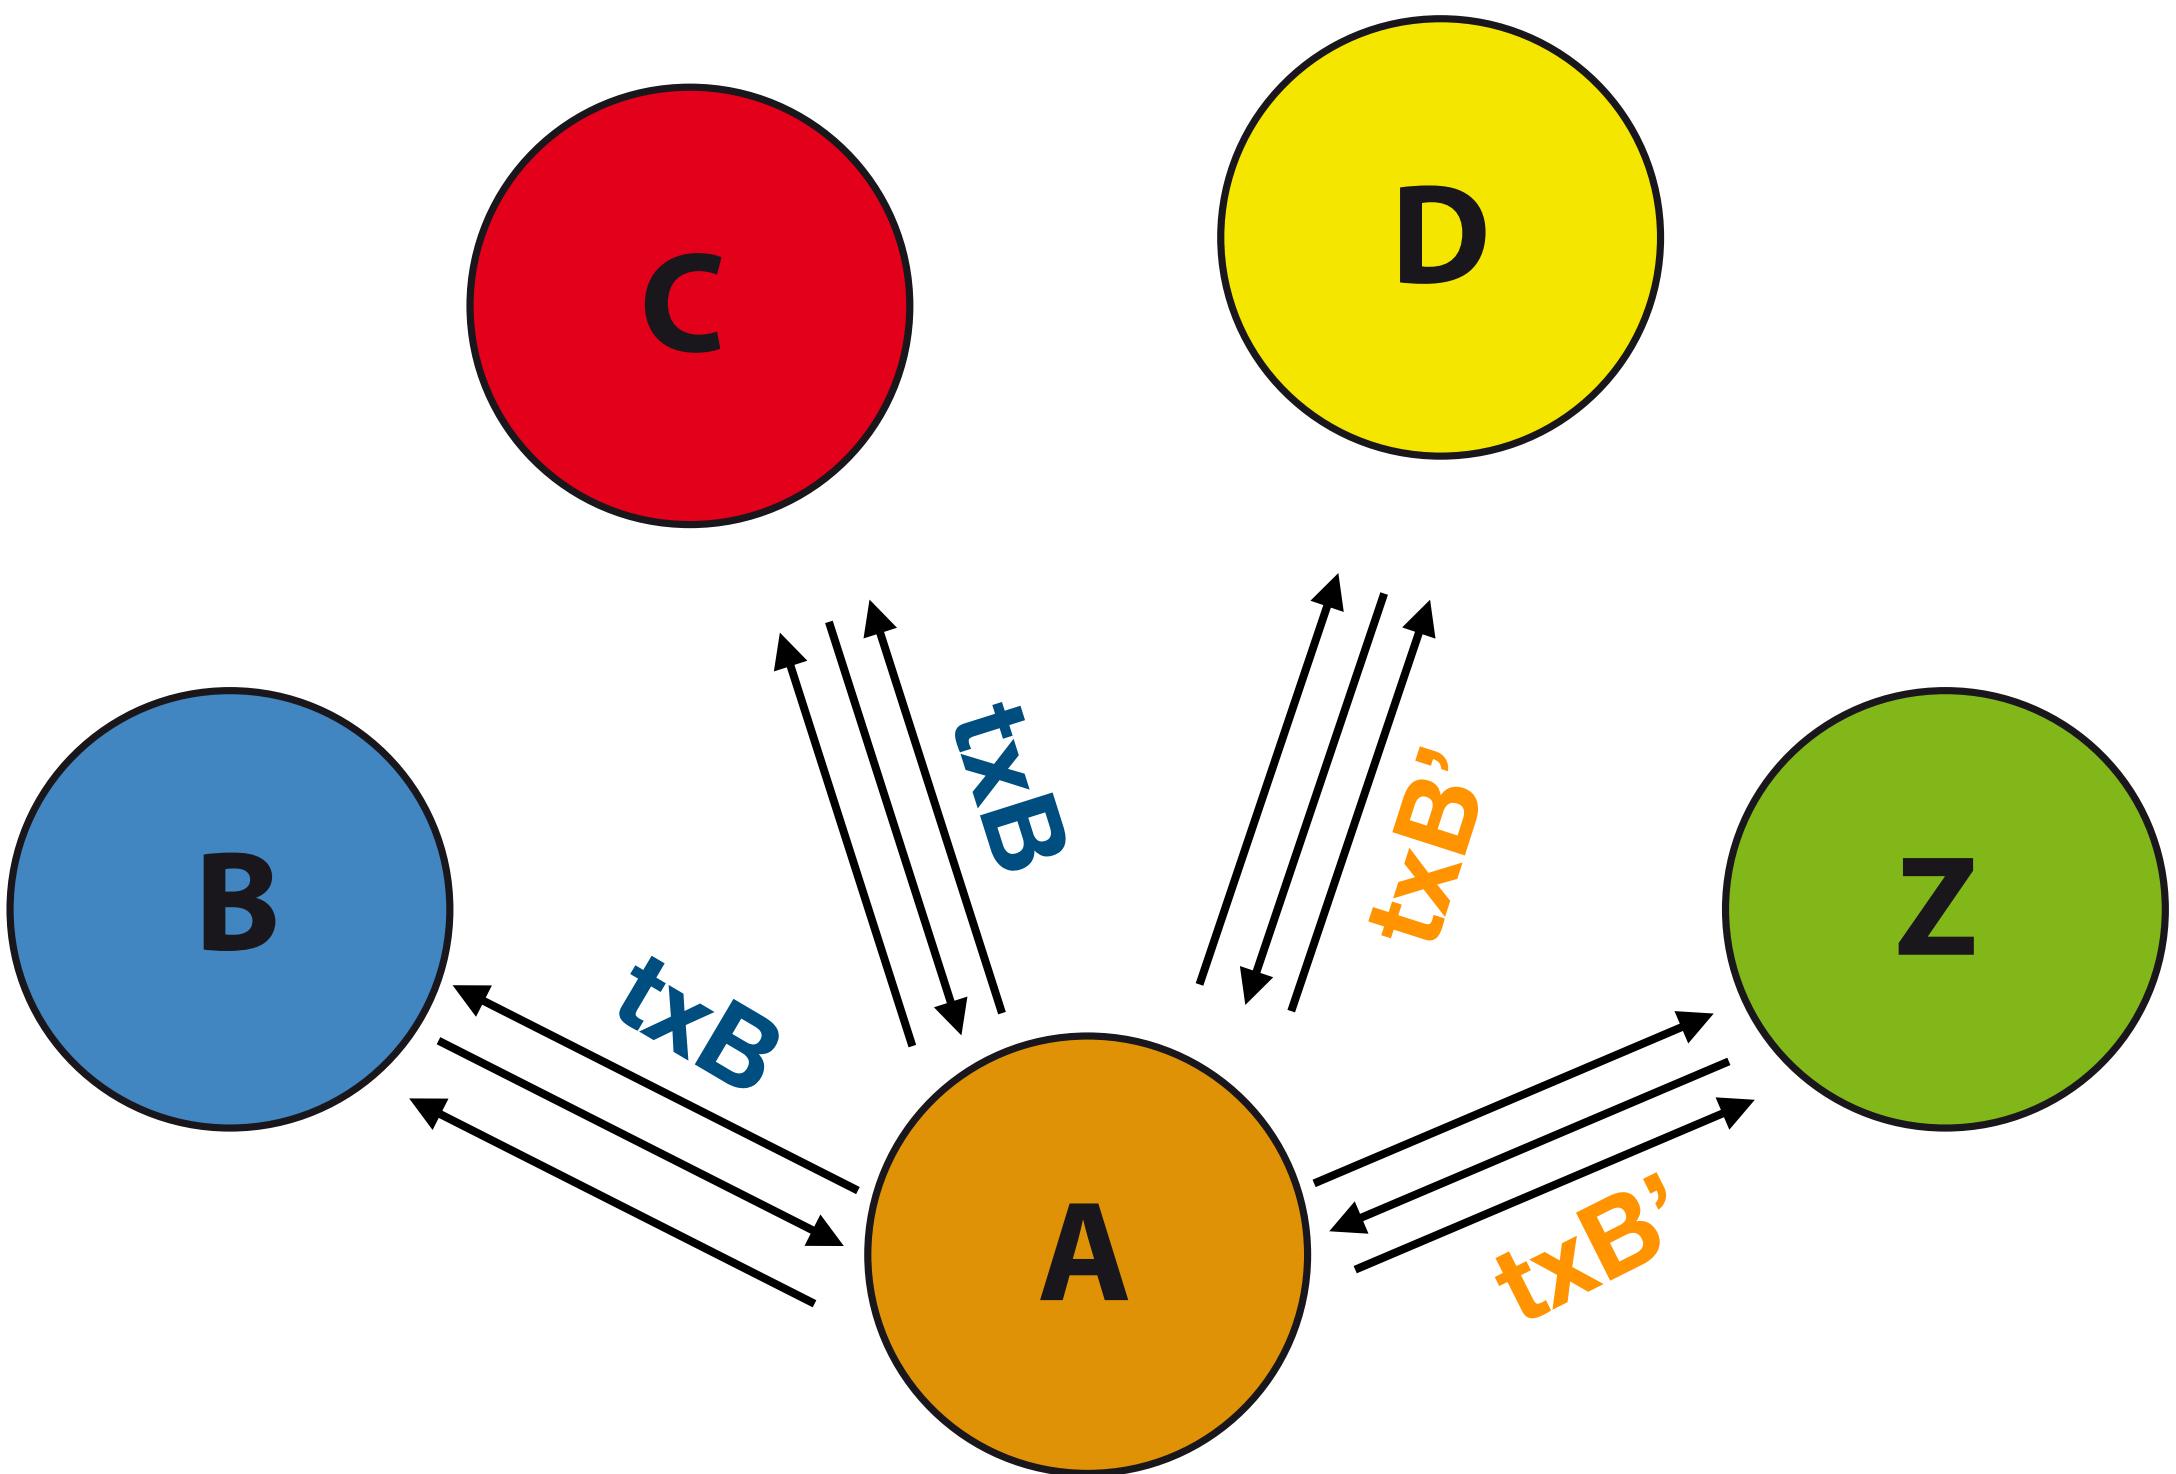


DOUBLE-SPENDING TRANSACTIONS (2/2)



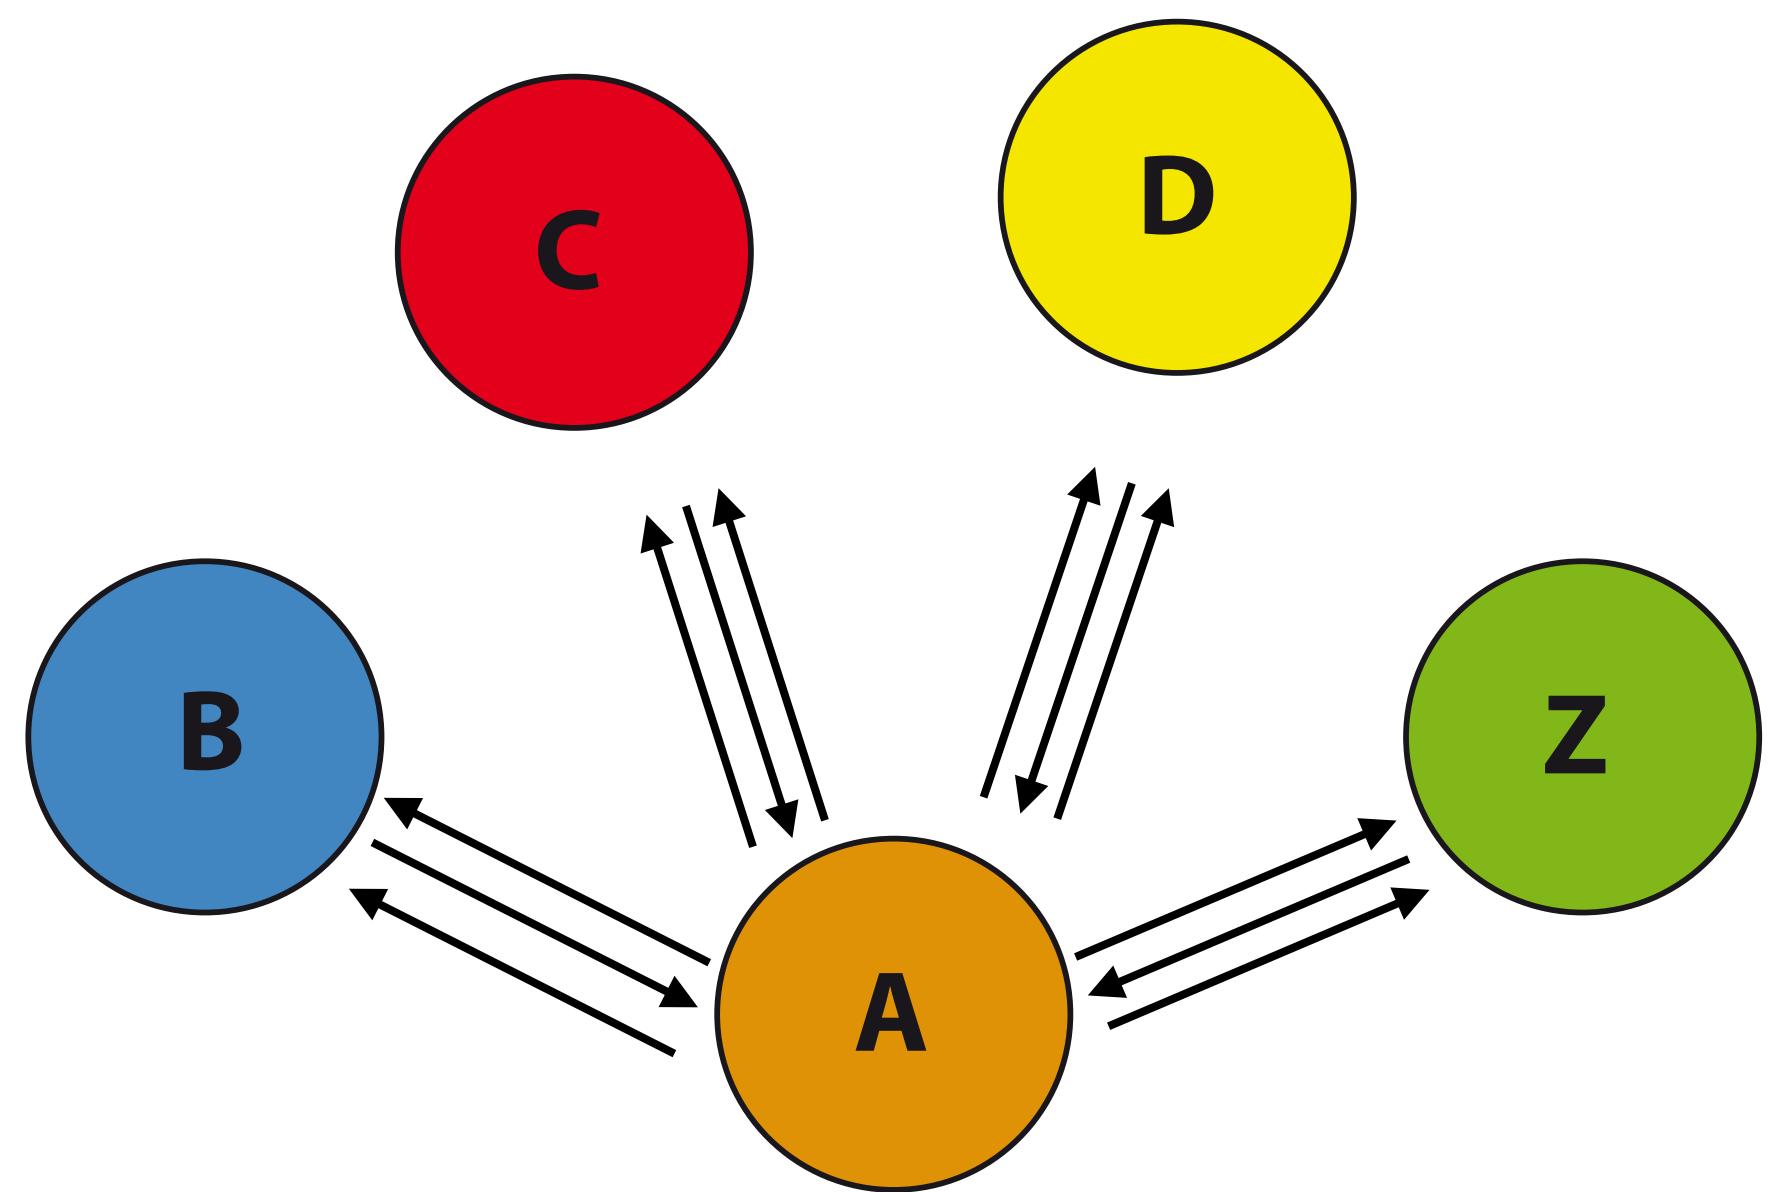
- 0-conf transactions should not be trusted

DOUBLE-SPENDING TRANSACTIONS (2/2)

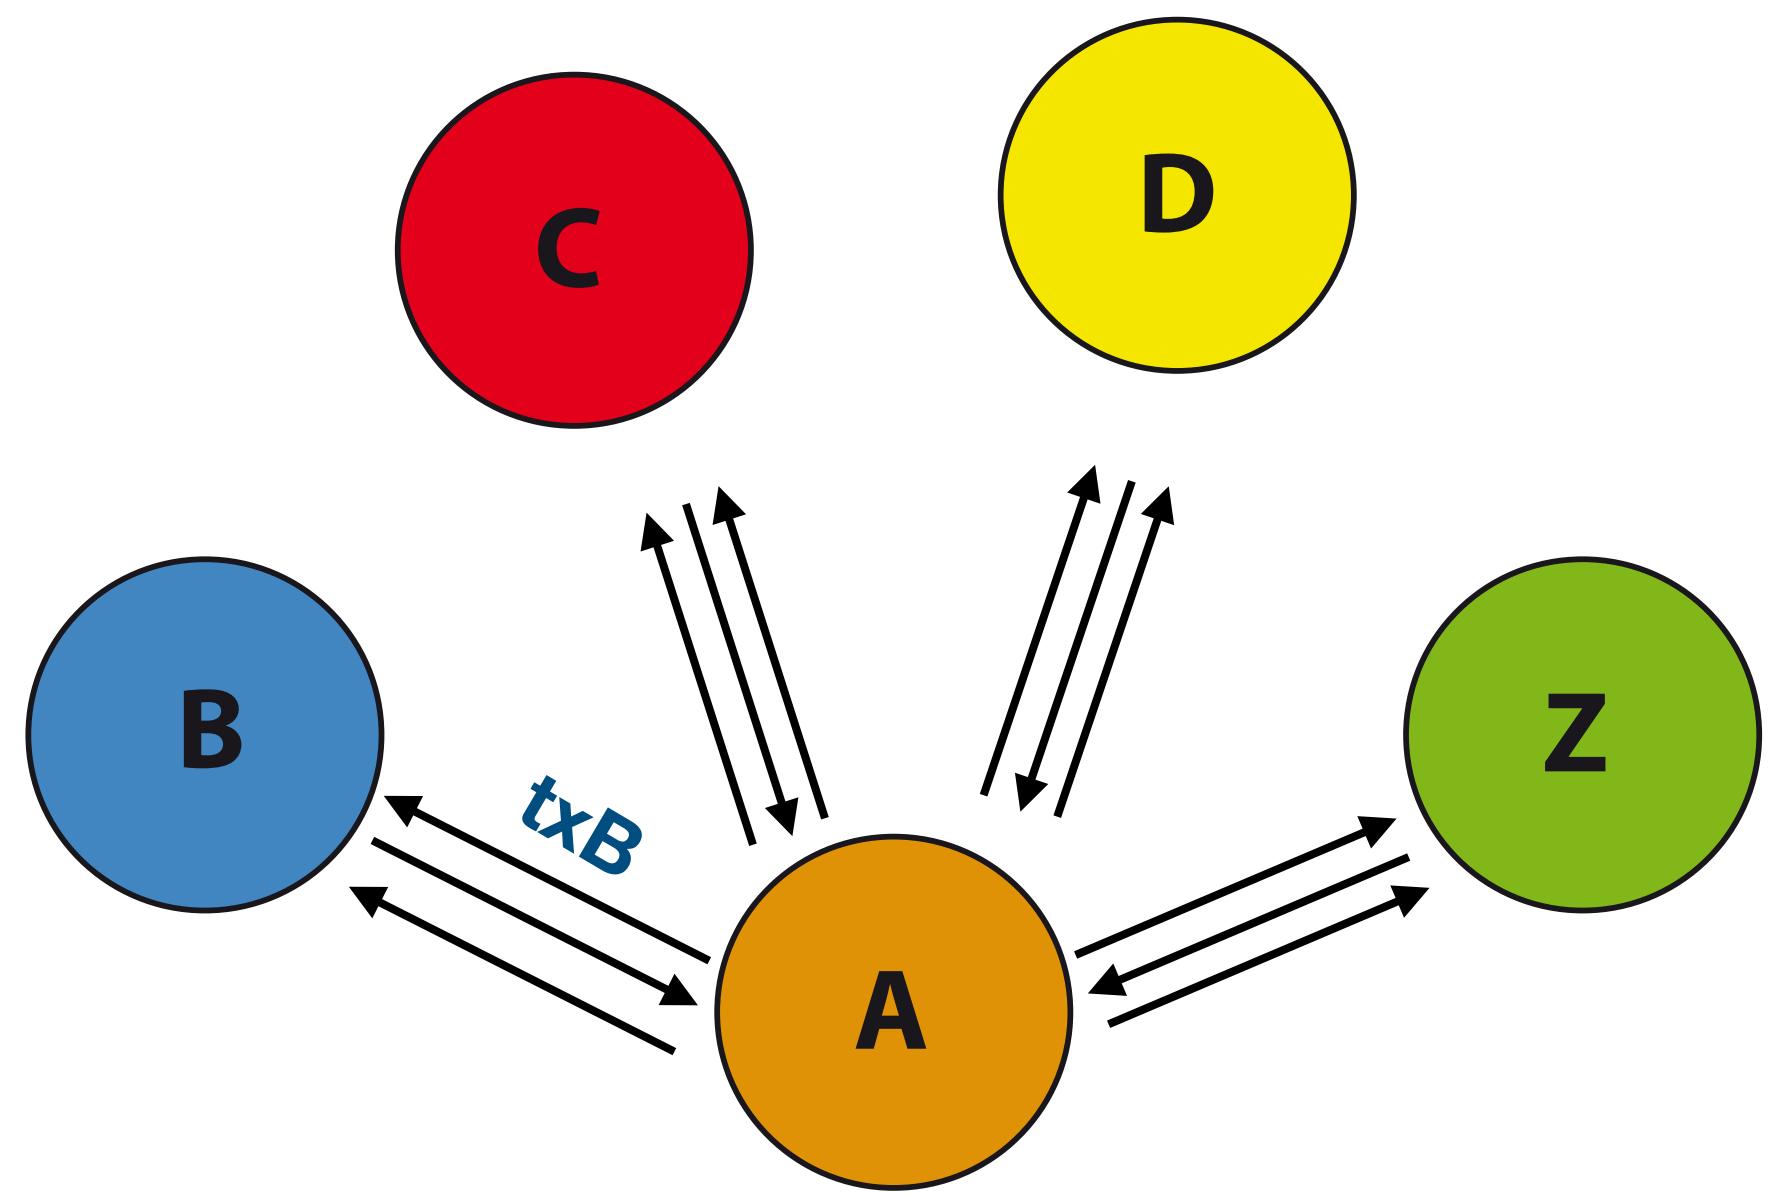


- 0-conf transactions should not be trusted
- If B accepts txB before it appears in a block he **can be deceived** by A

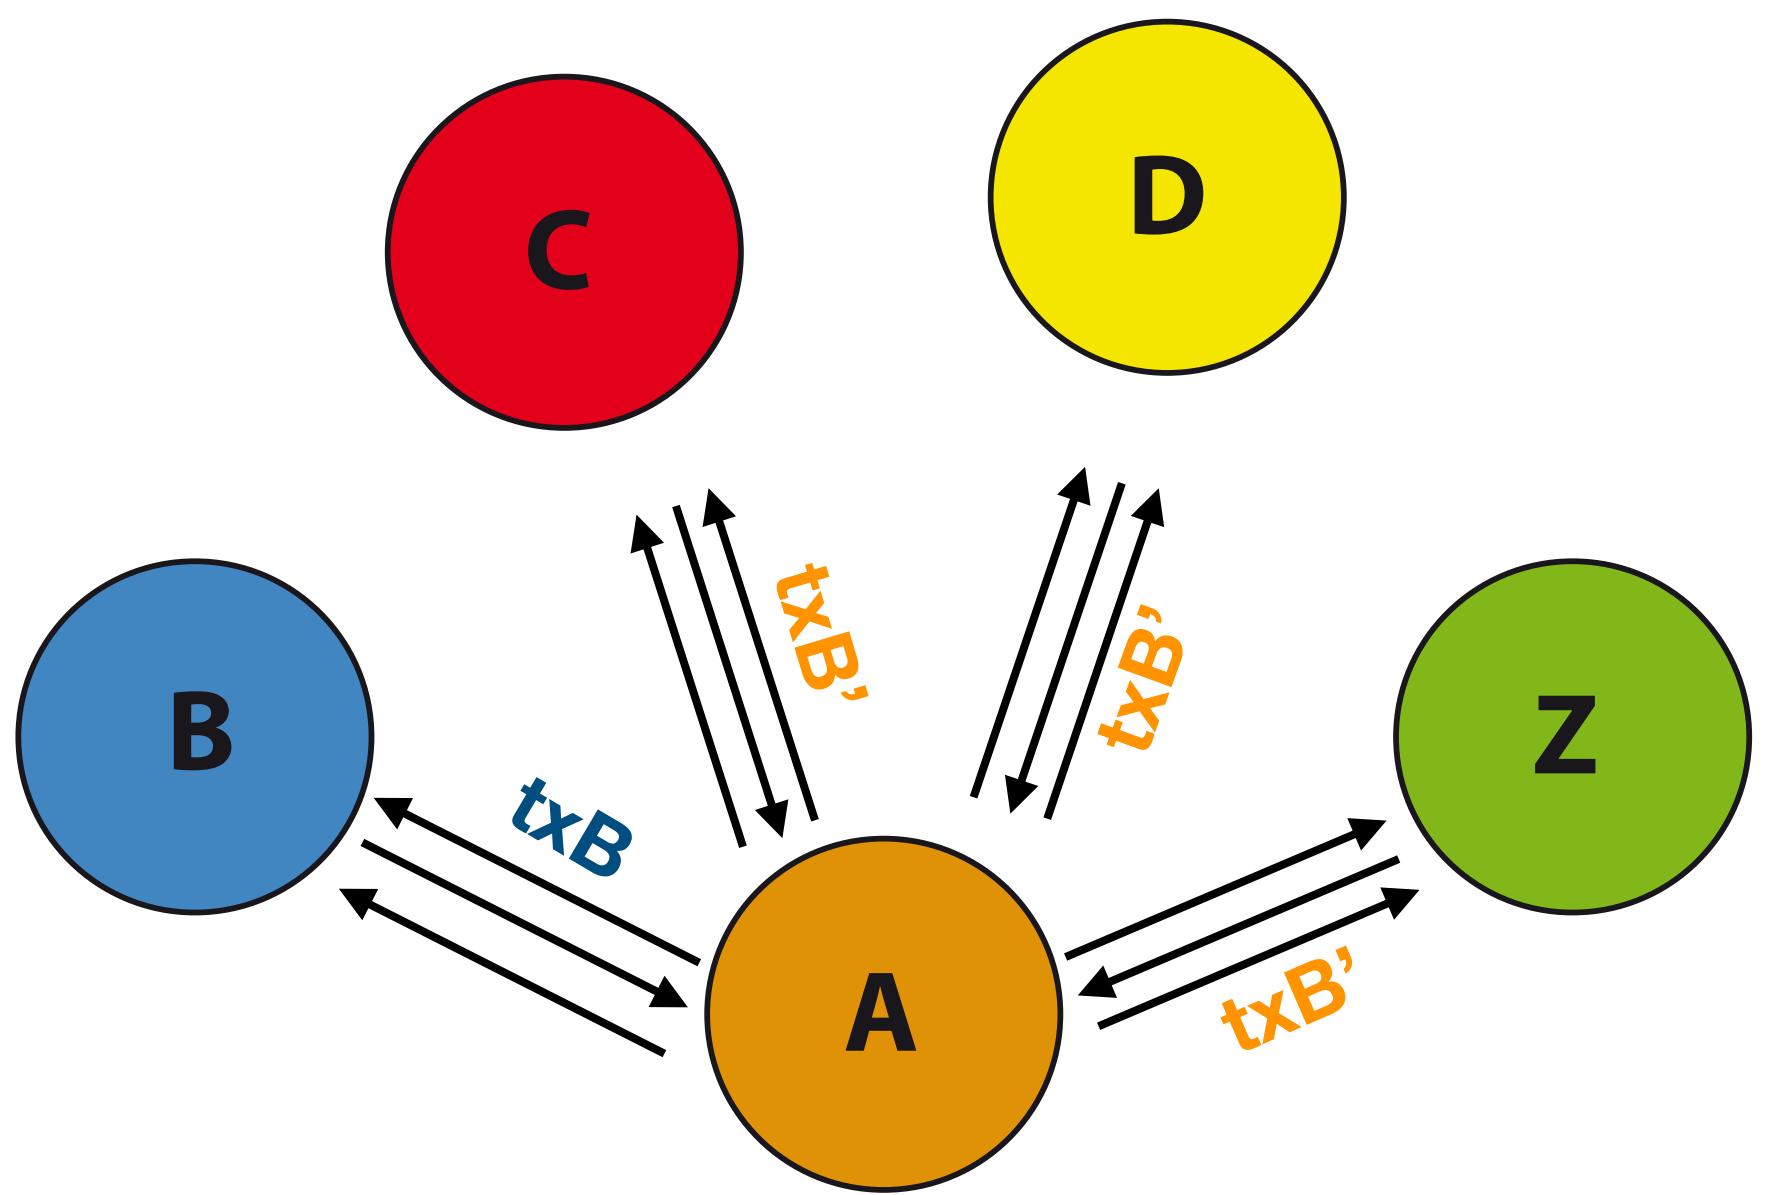
WHEN THINGS GO SOUTH



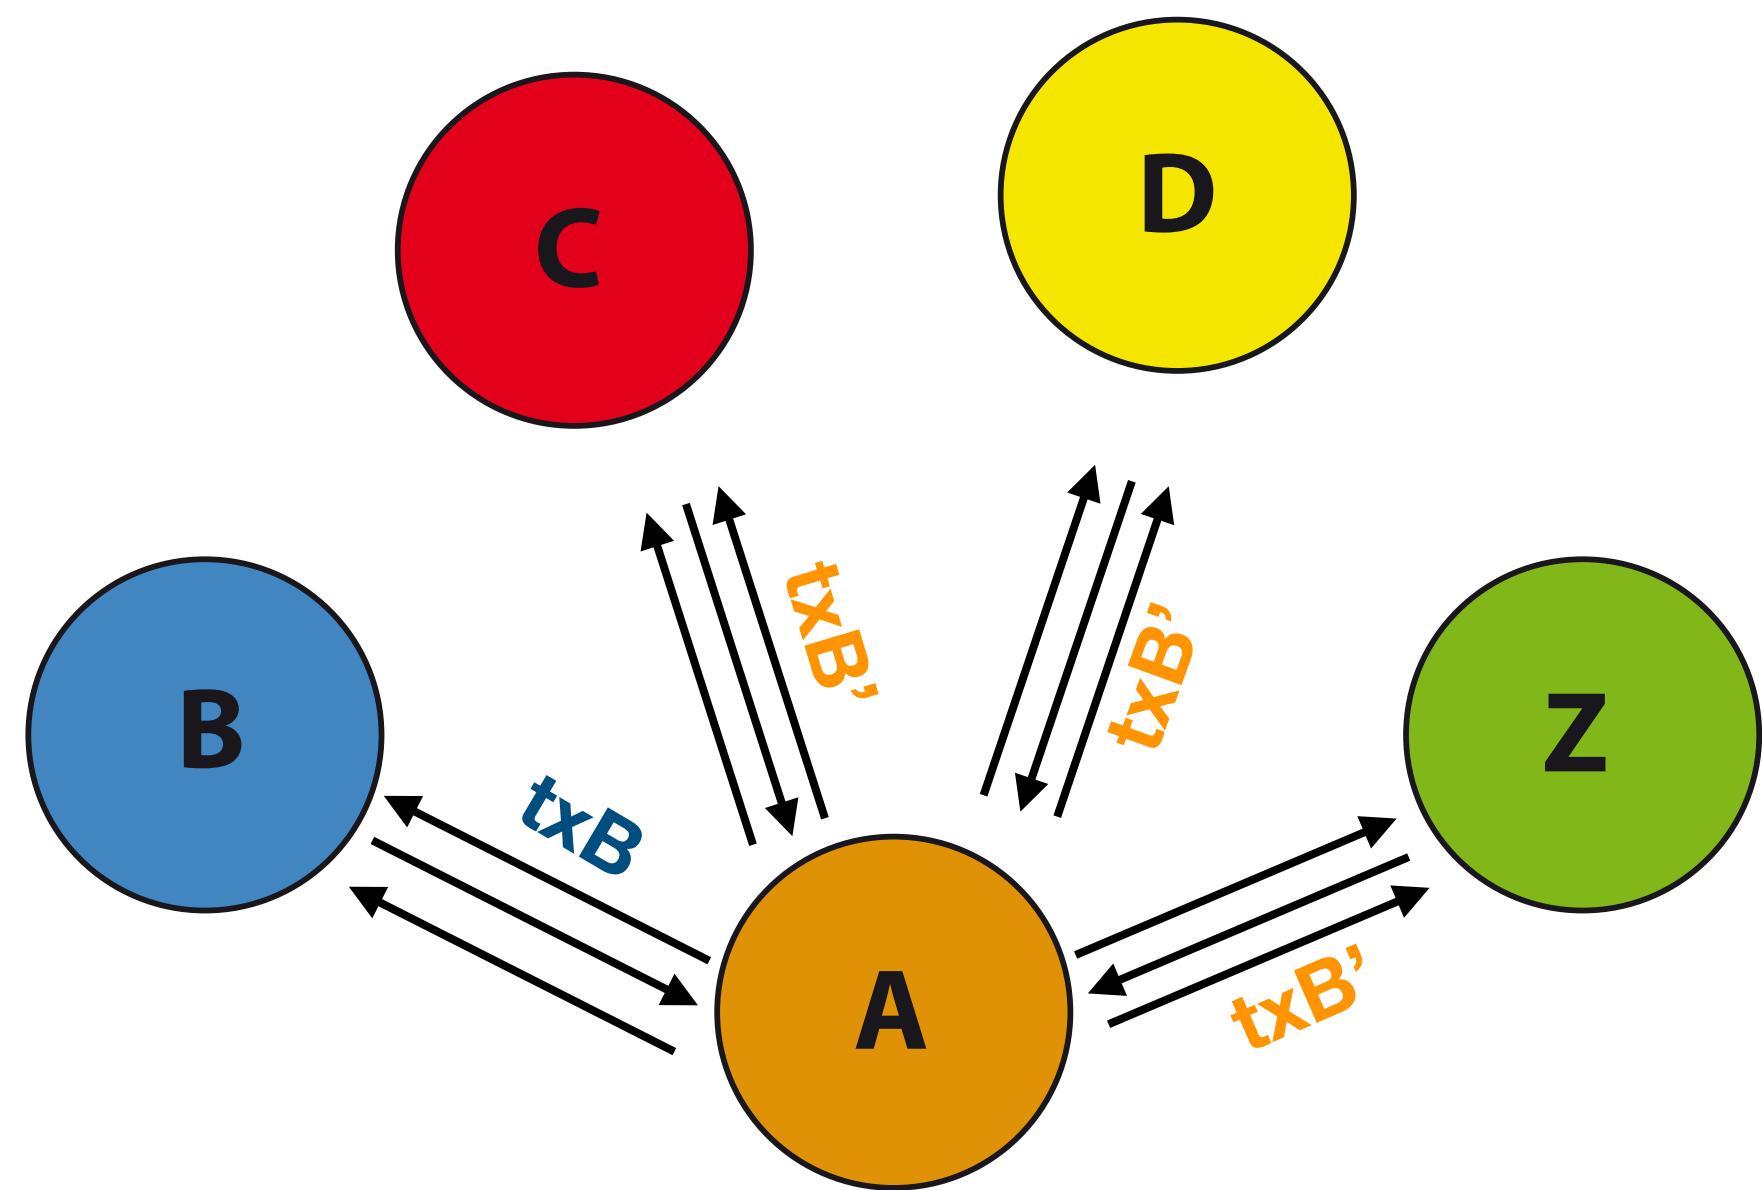
WHEN THINGS GO SOUTH



WHEN THINGS GO SOUTH

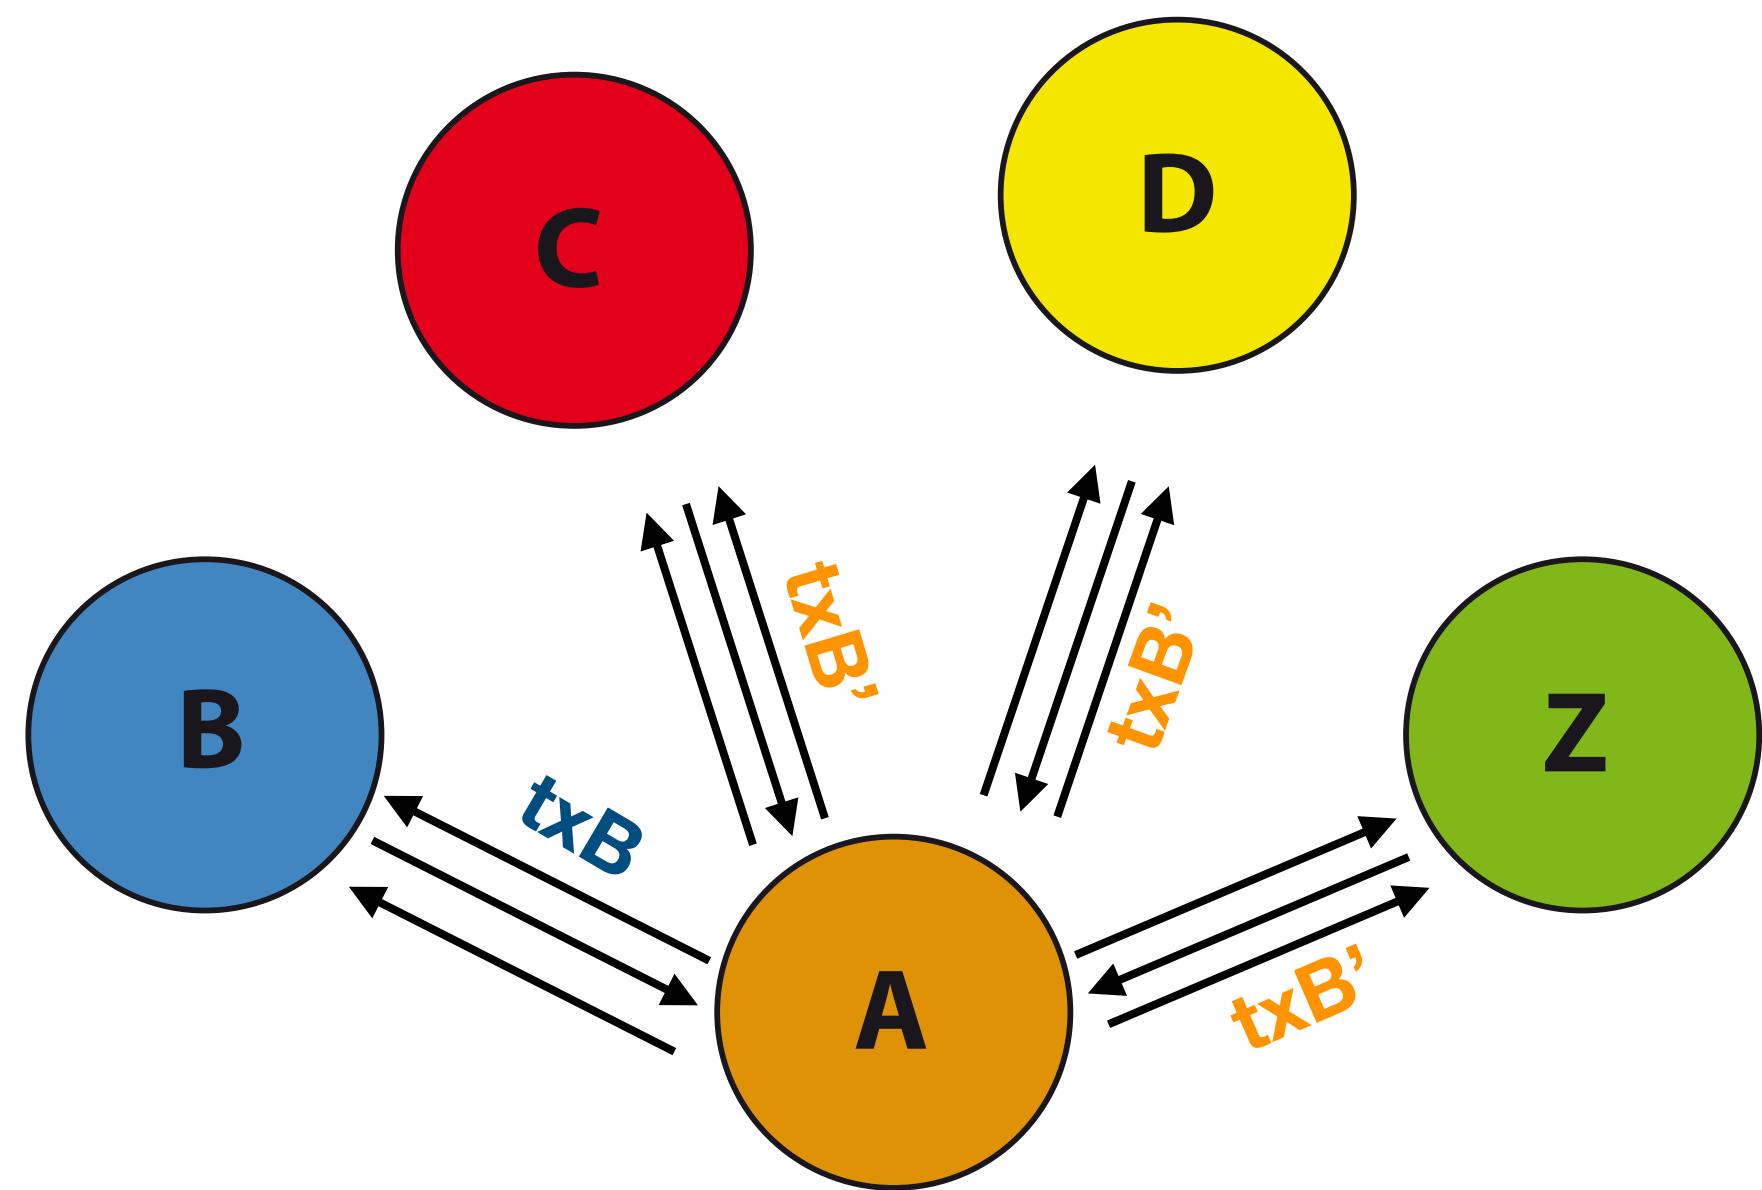


WHEN THINGS GO SOUTH



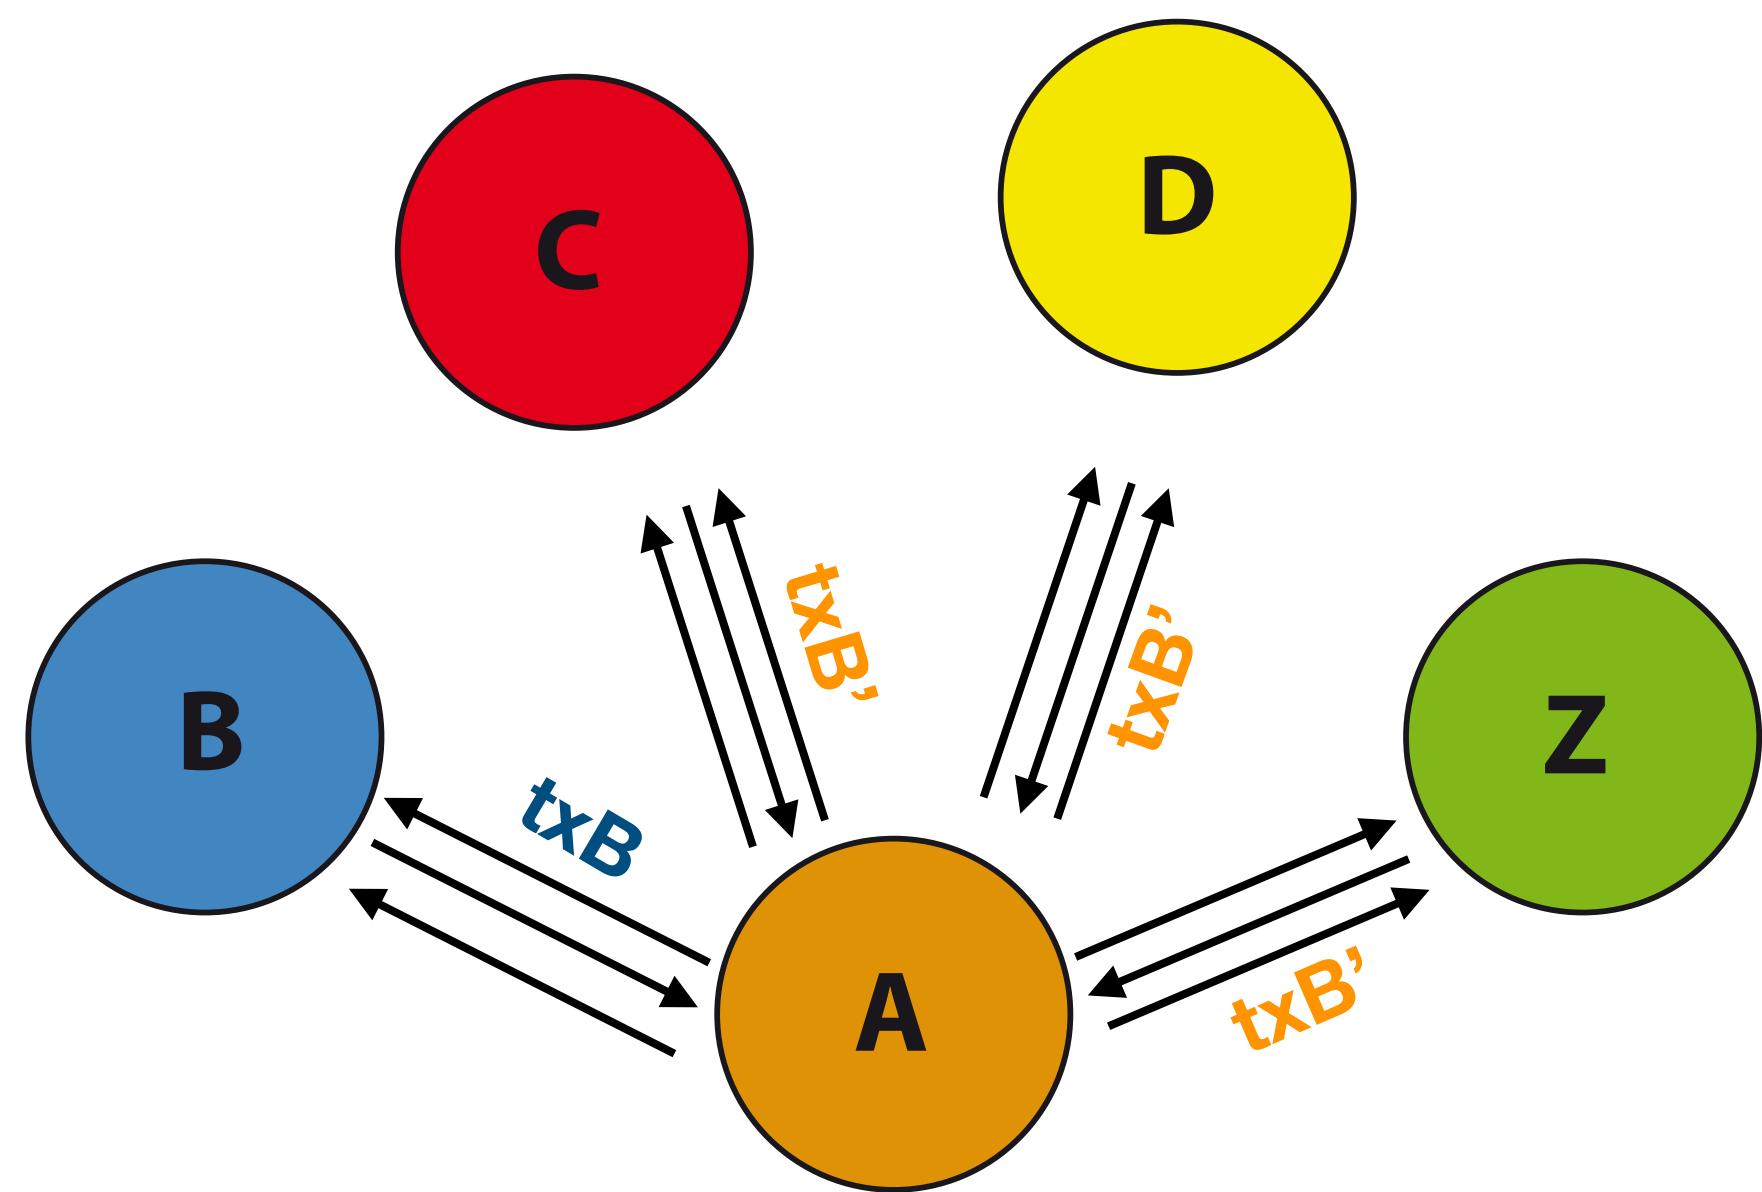
- If A manages to control the **network view** of B, A can easily deceive B

WHEN THINGS GO SOUTH



- If A manages to control the **network view** of B, A can easily deceive B
- When a node controls the view of another subset of nodes, the latter is said to be **eclipsed**

WHEN THINGS GO SOUTH

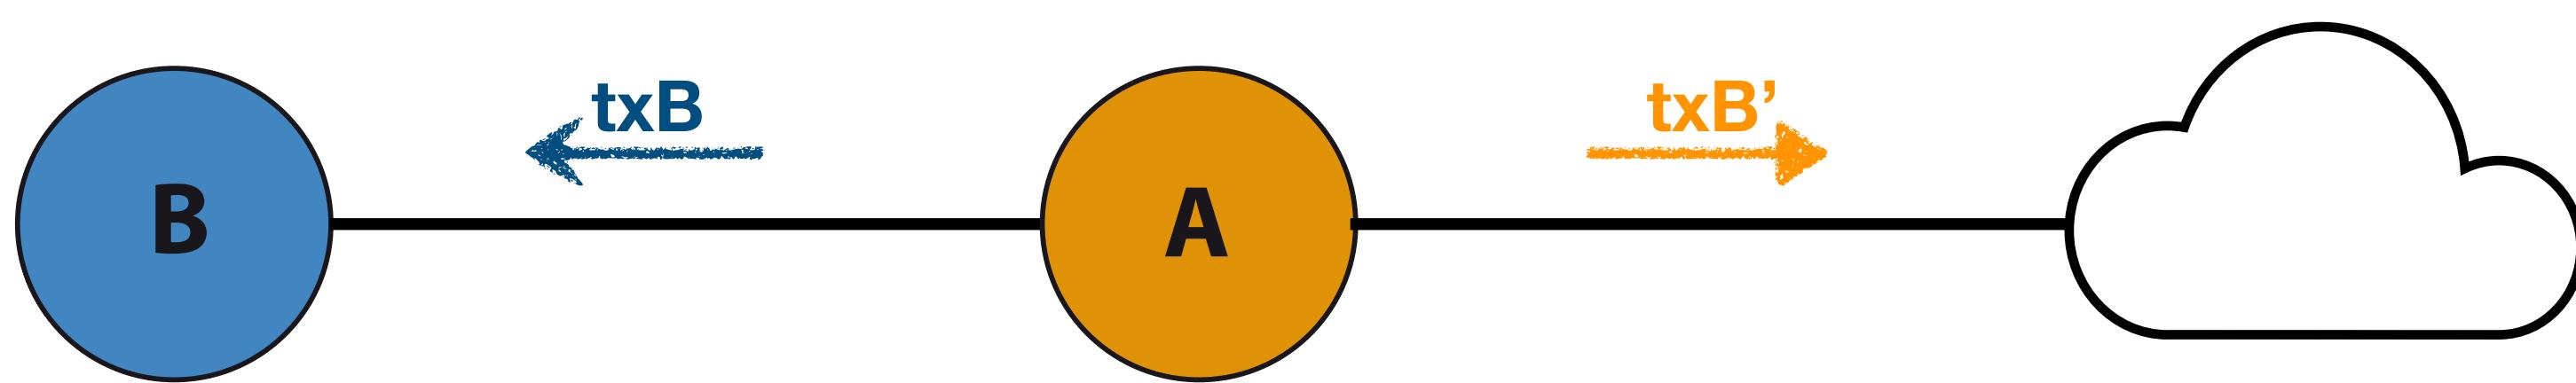


- If A manages to control the **network view** of B, A can easily deceive B
- When a node controls the view of another subset of nodes, the latter is said to be **eclipsed**

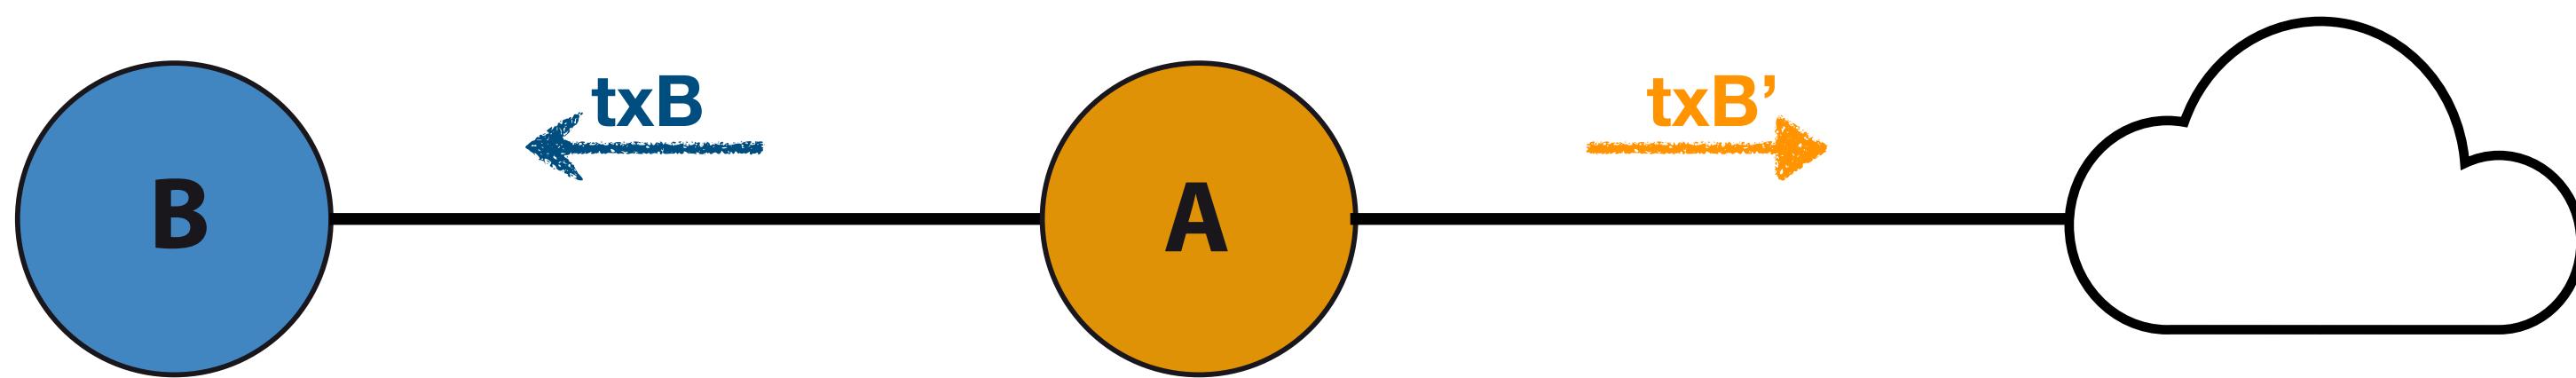


Ethan Heilman, Alison Kendler, Aviv Zohar and Sharon Goldberg
Eclipse Attacks on Bitcoin's Peer-to-Peer Network
<https://www.usenix.org/node/190891>

ECLIPSE ATTACKS (1/2)

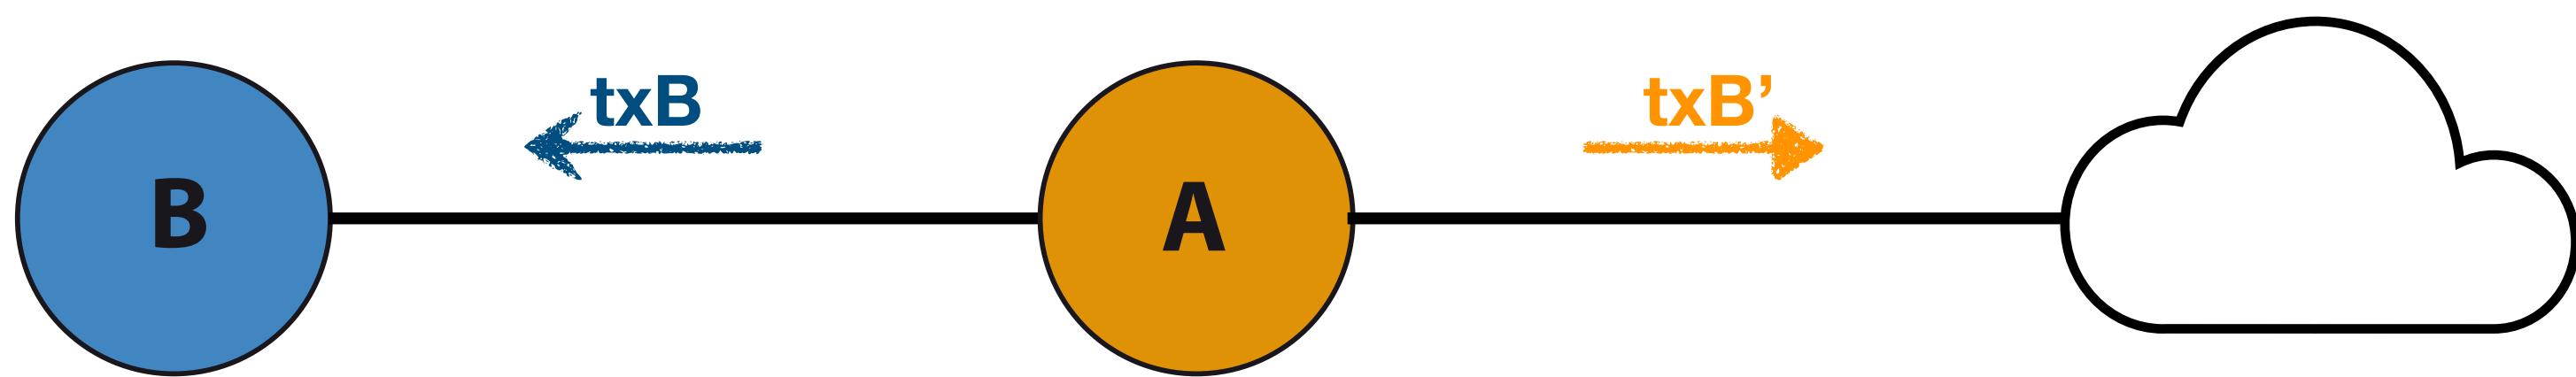


ECLIPSE ATTACKS (1/2)



B will be deceived provided:

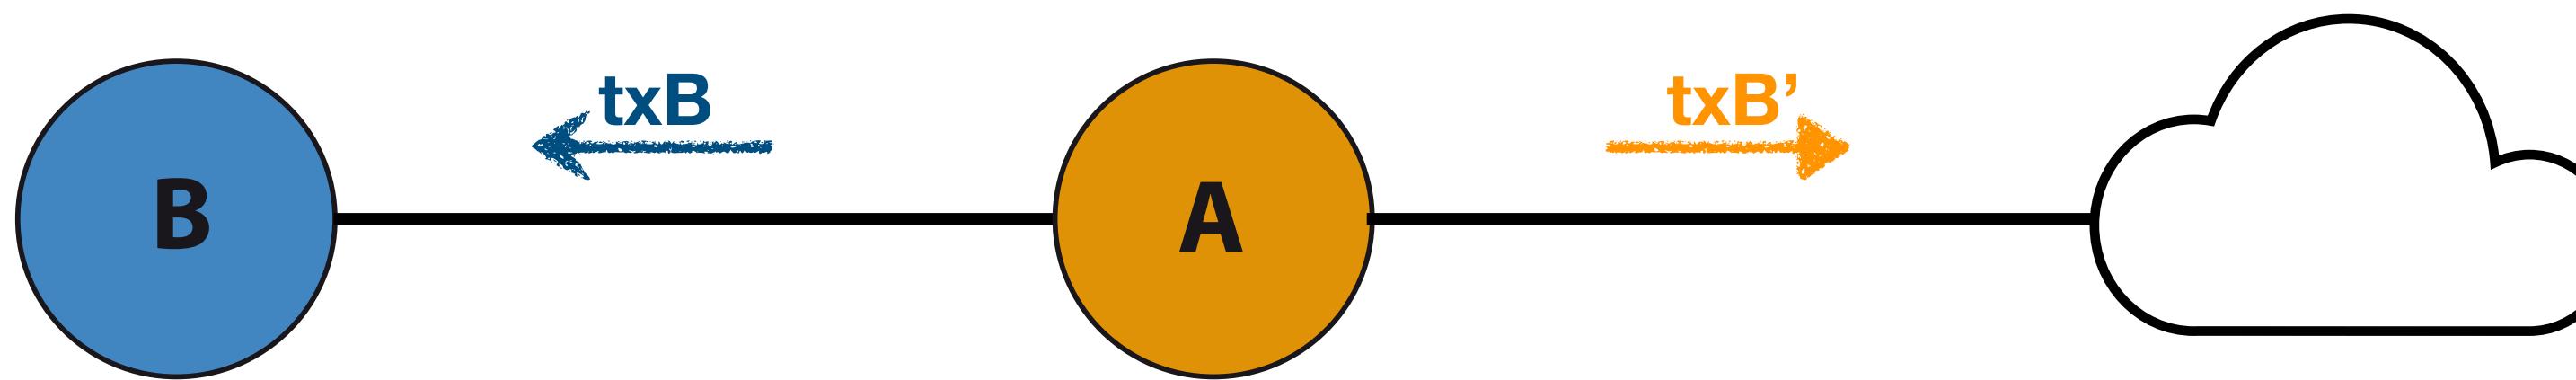
ECLIPSE ATTACKS (1/2)



B will be deceived provided:

- B accepts 0-conf transactions

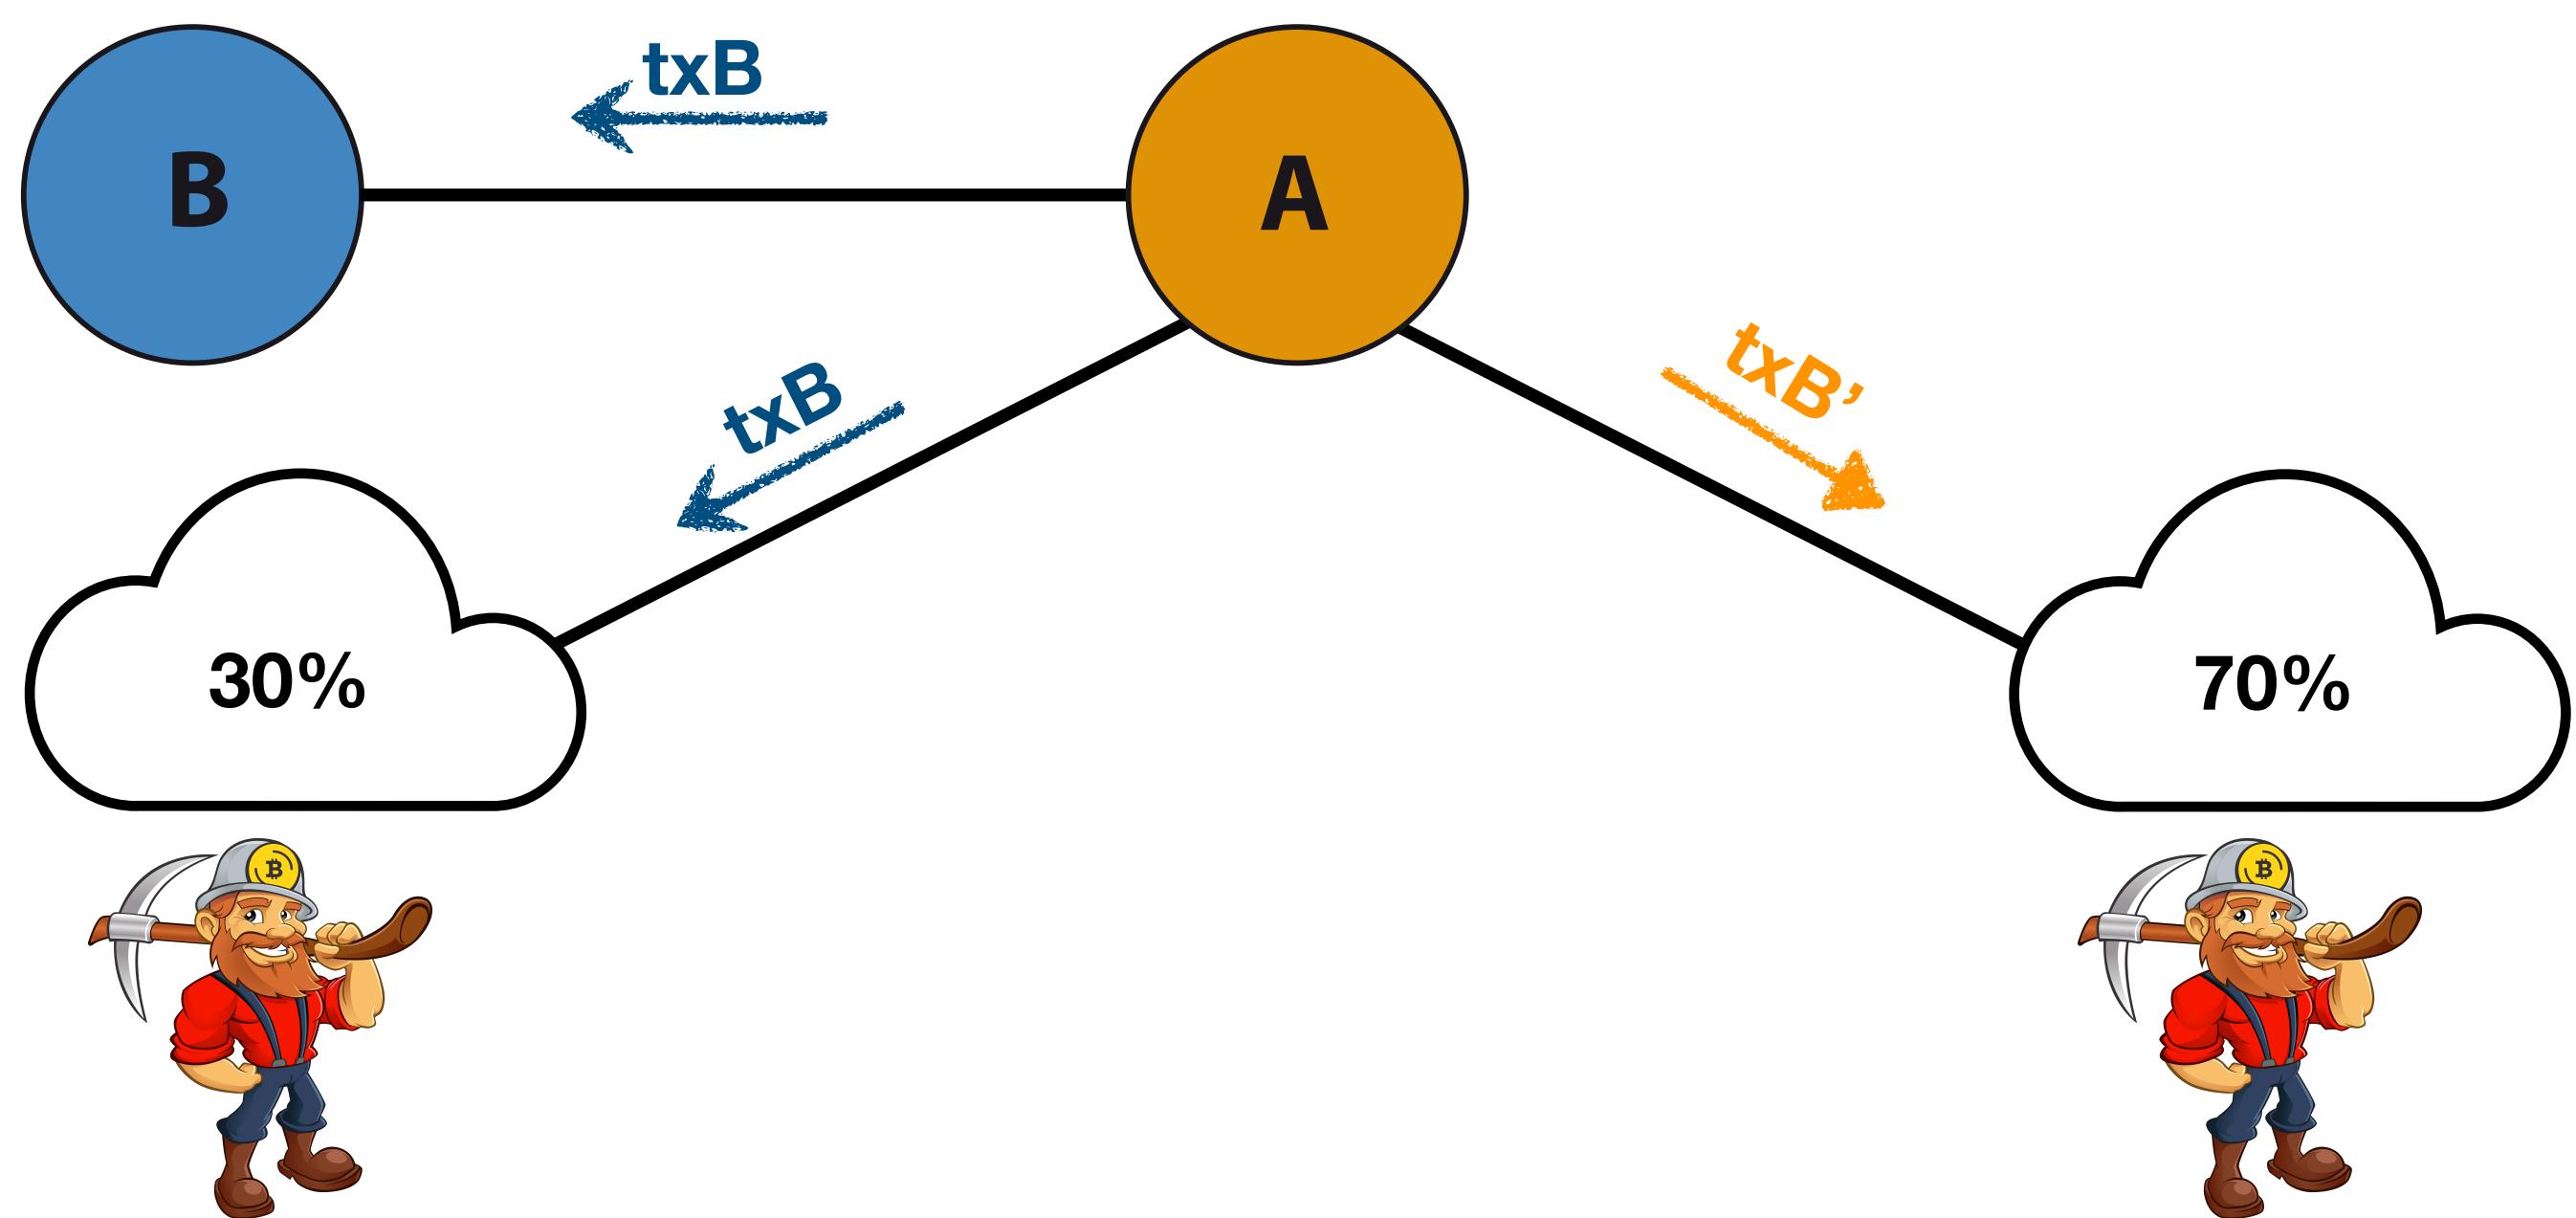
ECLIPSE ATTACKS (1/2)



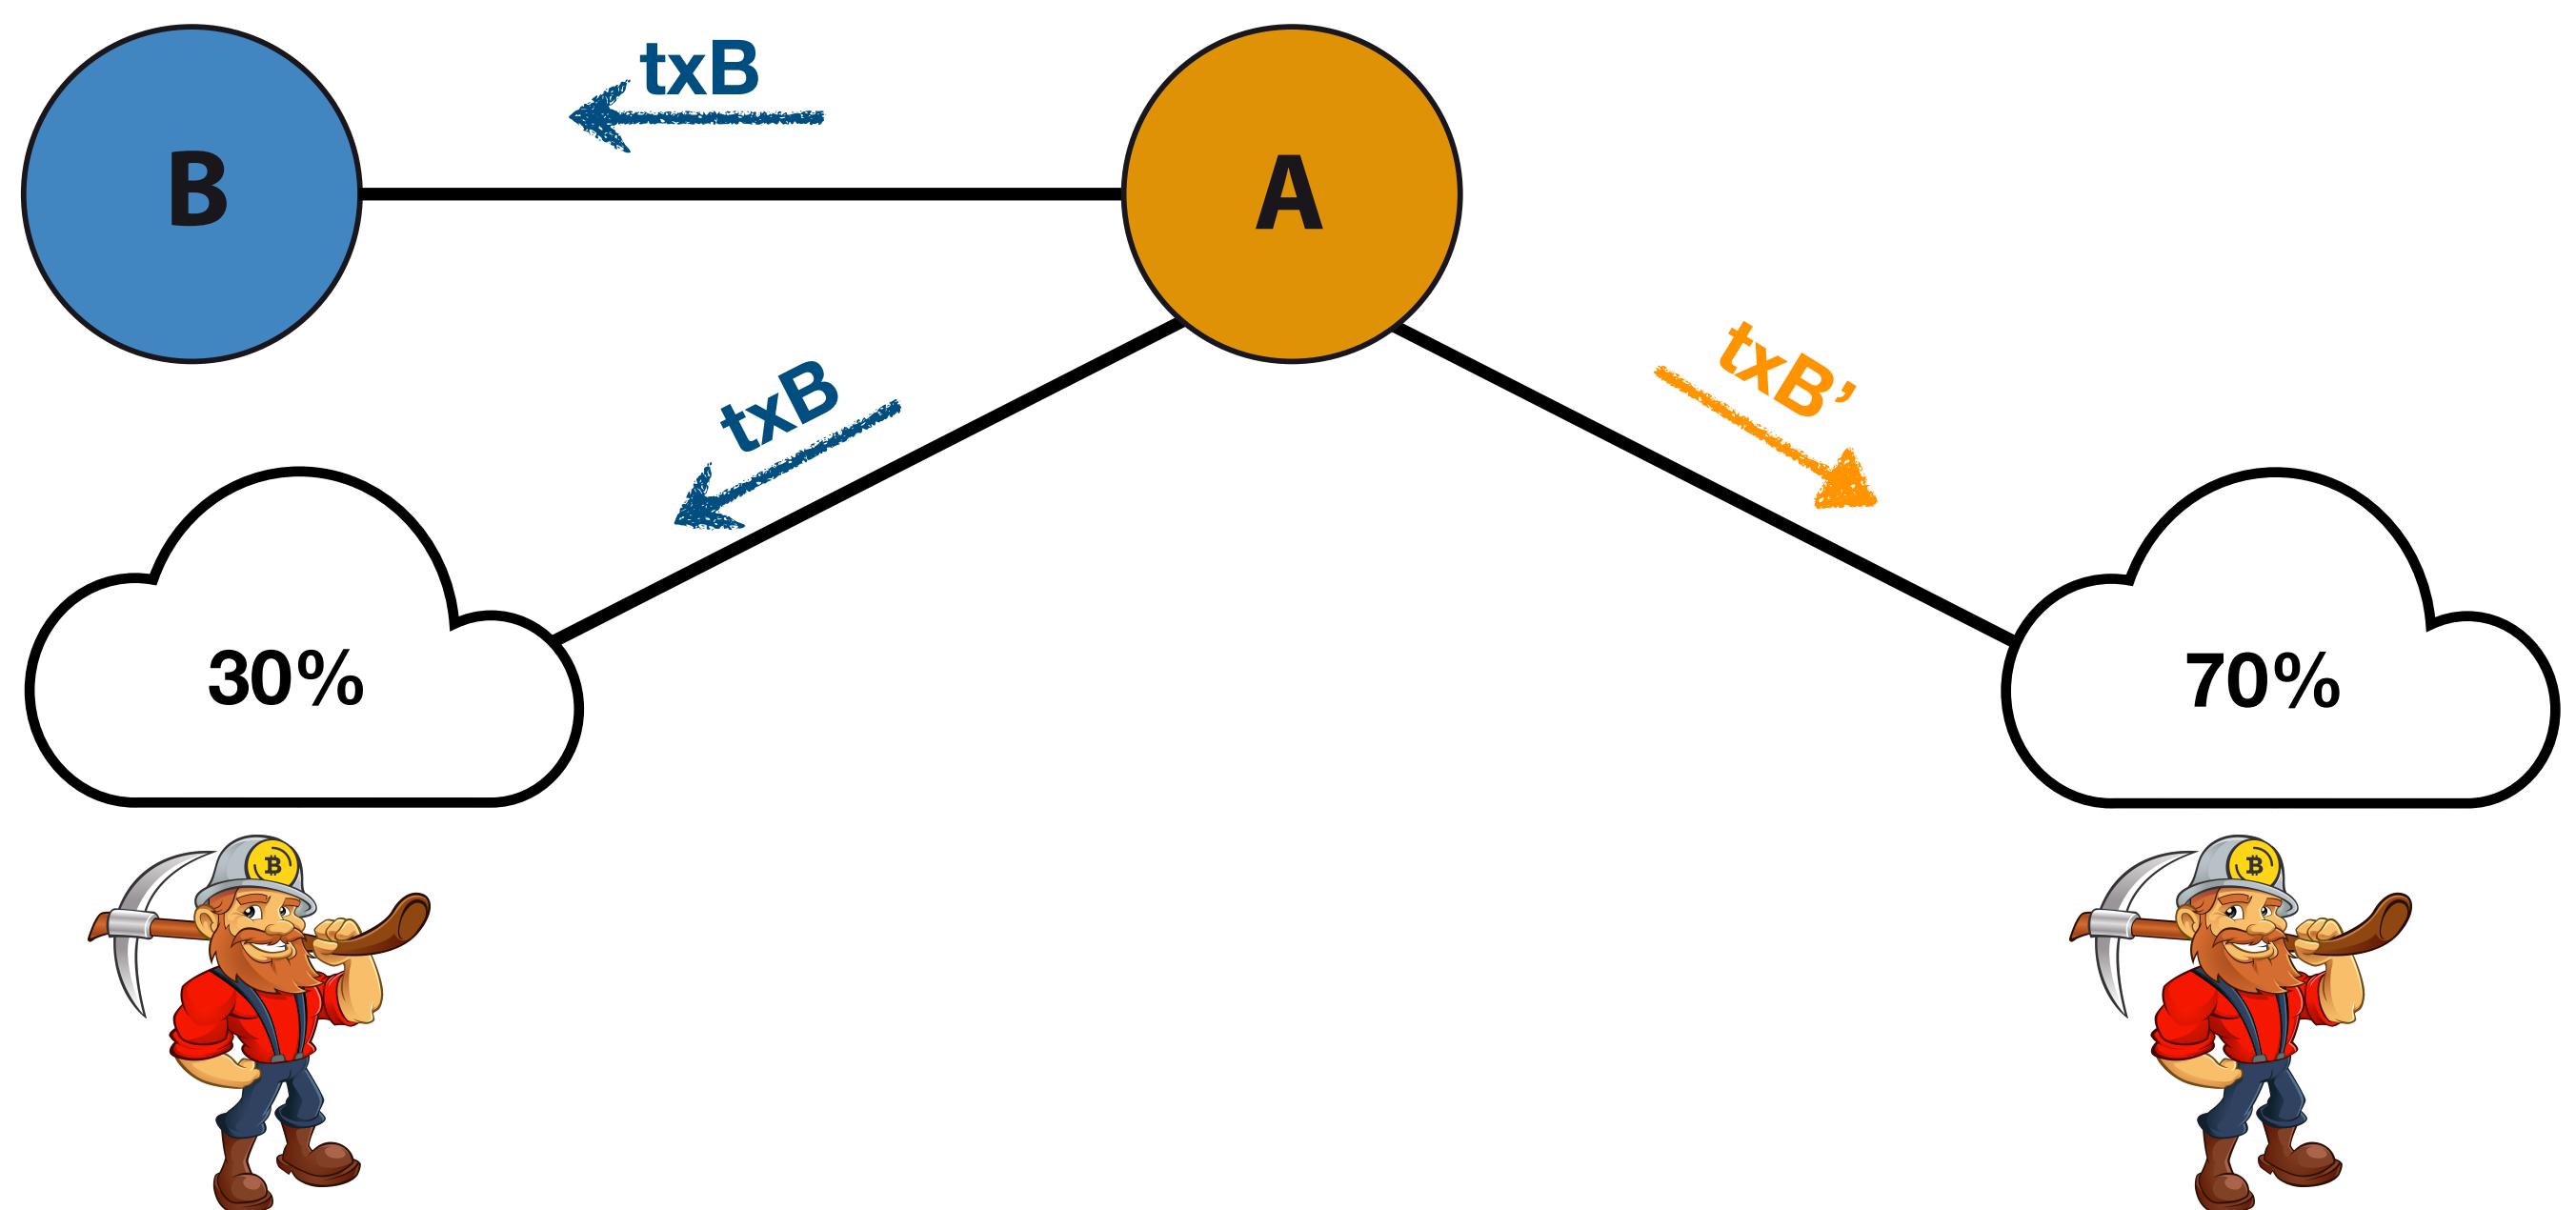
B will be deceived provided:

- B accepts 0-conf transactions
- A has enough hash power to generate blocks in a reasonable time

ECLIPSE ATTACKS (2/2)

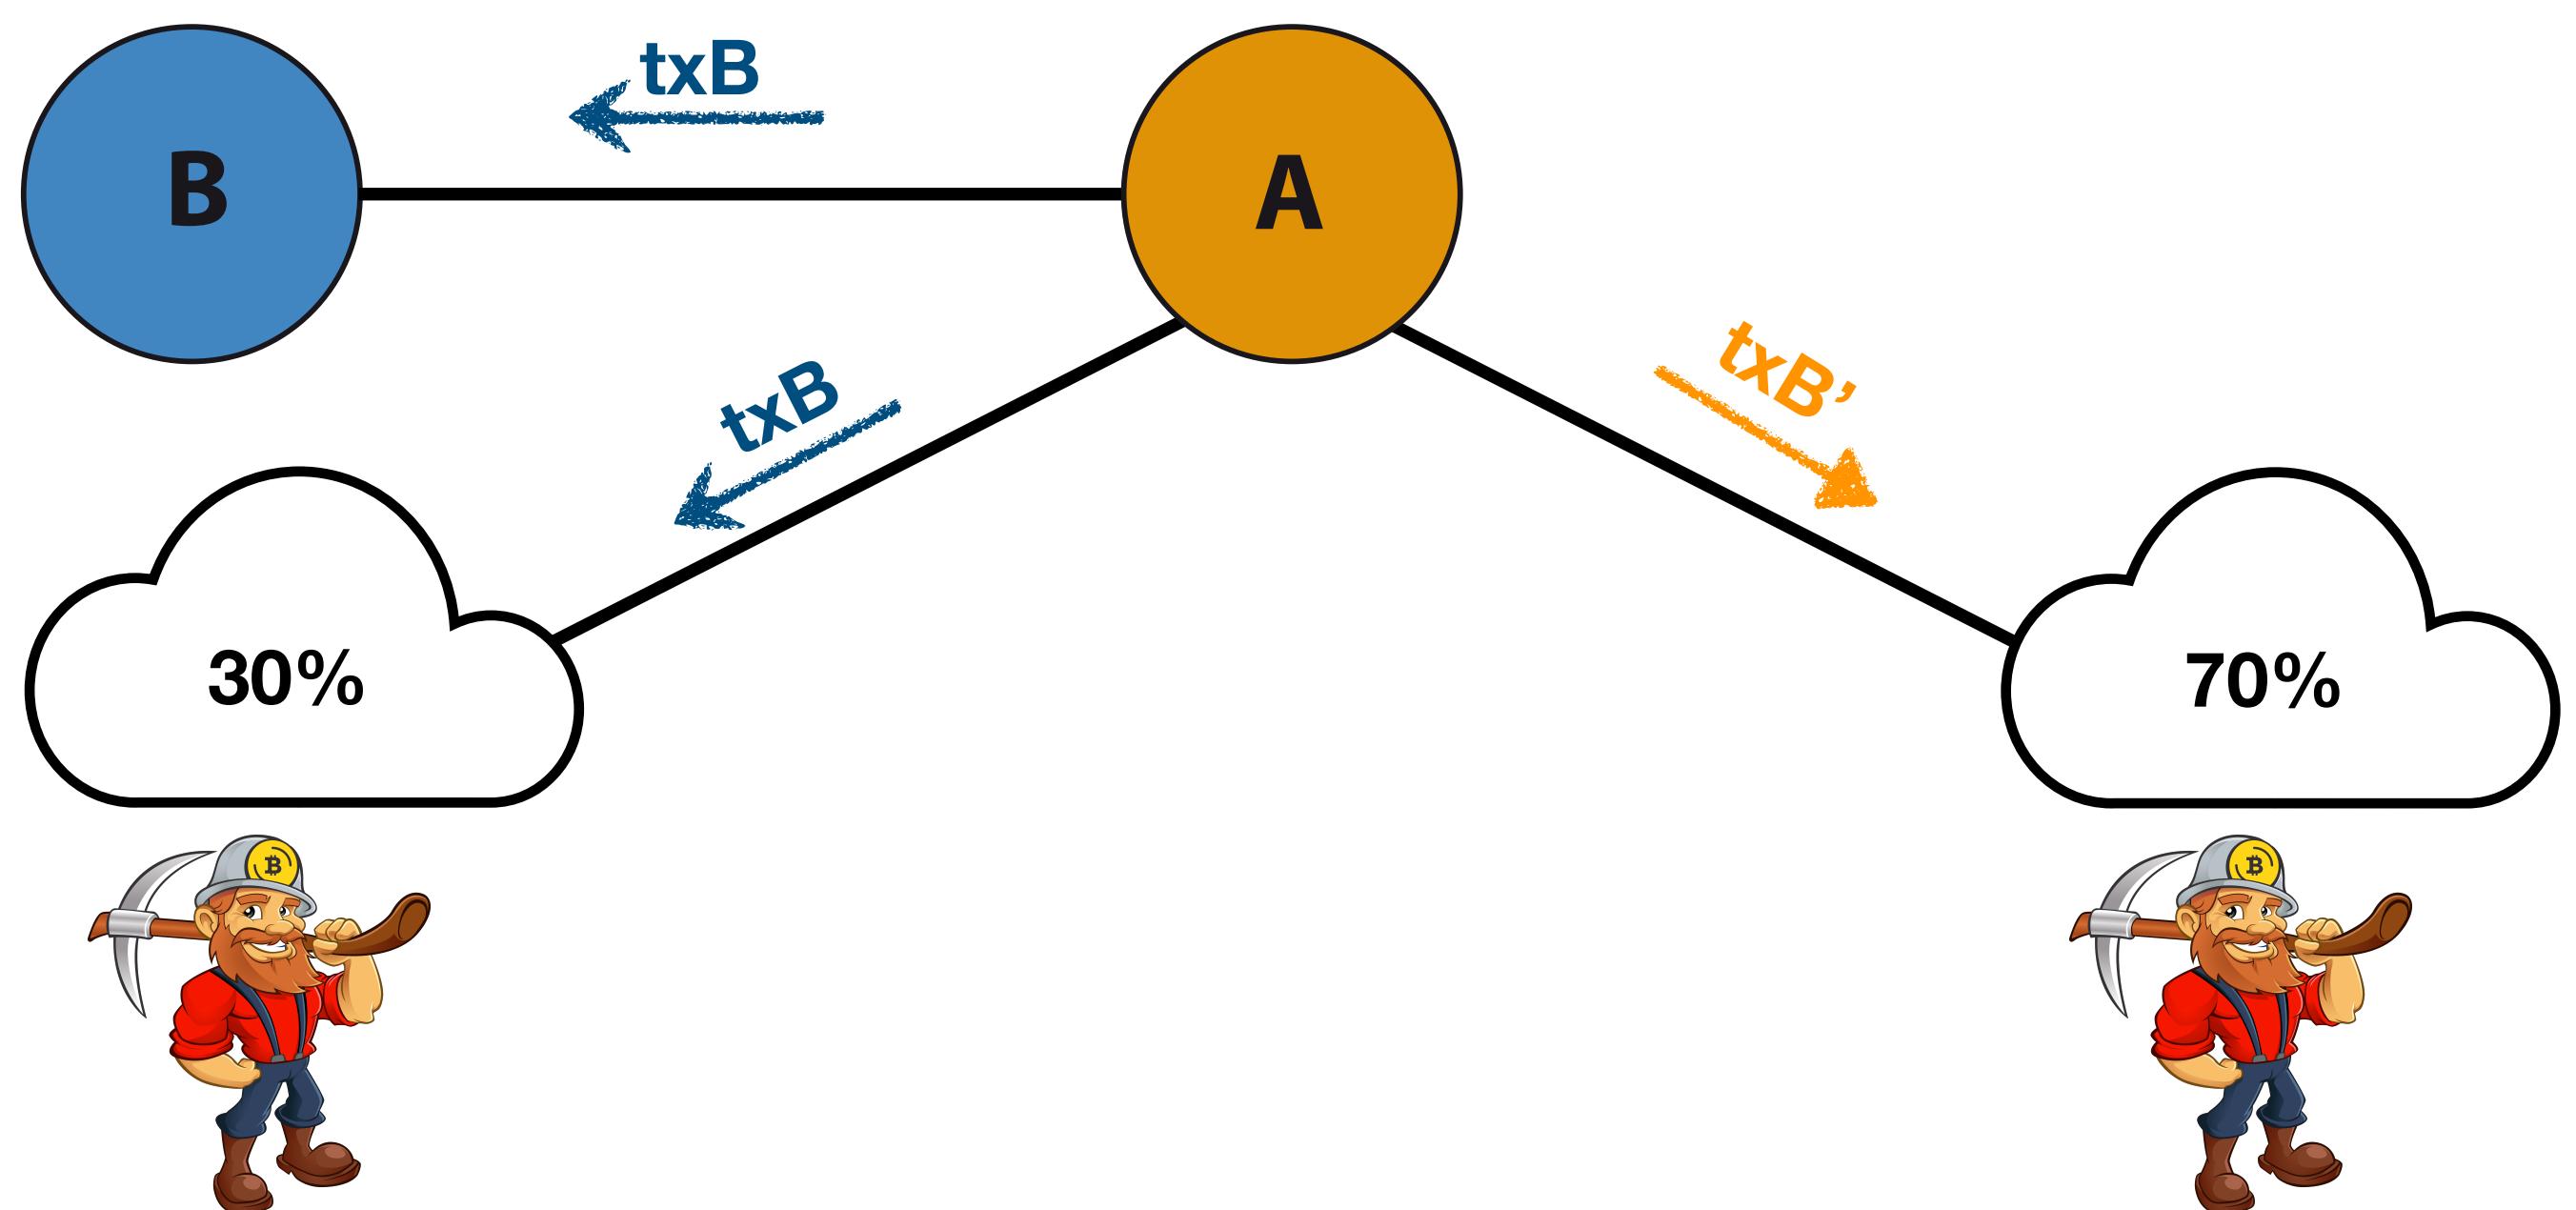


ECLIPSE ATTACKS (2/2)



- A does not even need to hold any mining power

ECLIPSE ATTACKS (2/2)



- A does not even need to hold any mining power
- With the right information it can participate the network in the most beneficial way for her

Network topology

UNKNOWN TOPOLOGY BY DESIGN

Peers are chosen pseudorandomly from the peer database of a node in order to become neighbors

Peers can be requested from other peers, but no information about whether the responder is (or has been) a neighbor of any of the provided peers is given

The network topology should mimic a random network

INFERRING THE TOPOLOGY

Does the network really look random?

INFERRING THE TOPOLOGY

Does the network really look random?

How can we known if we don't know what the topology looks like?

INFERRING THE TOPOLOGY

Does the network really look random?

How can we known if we don't know what the topology looks like?

Can we do anything to infer the topology?

INFERRING THE TOPOLOGY

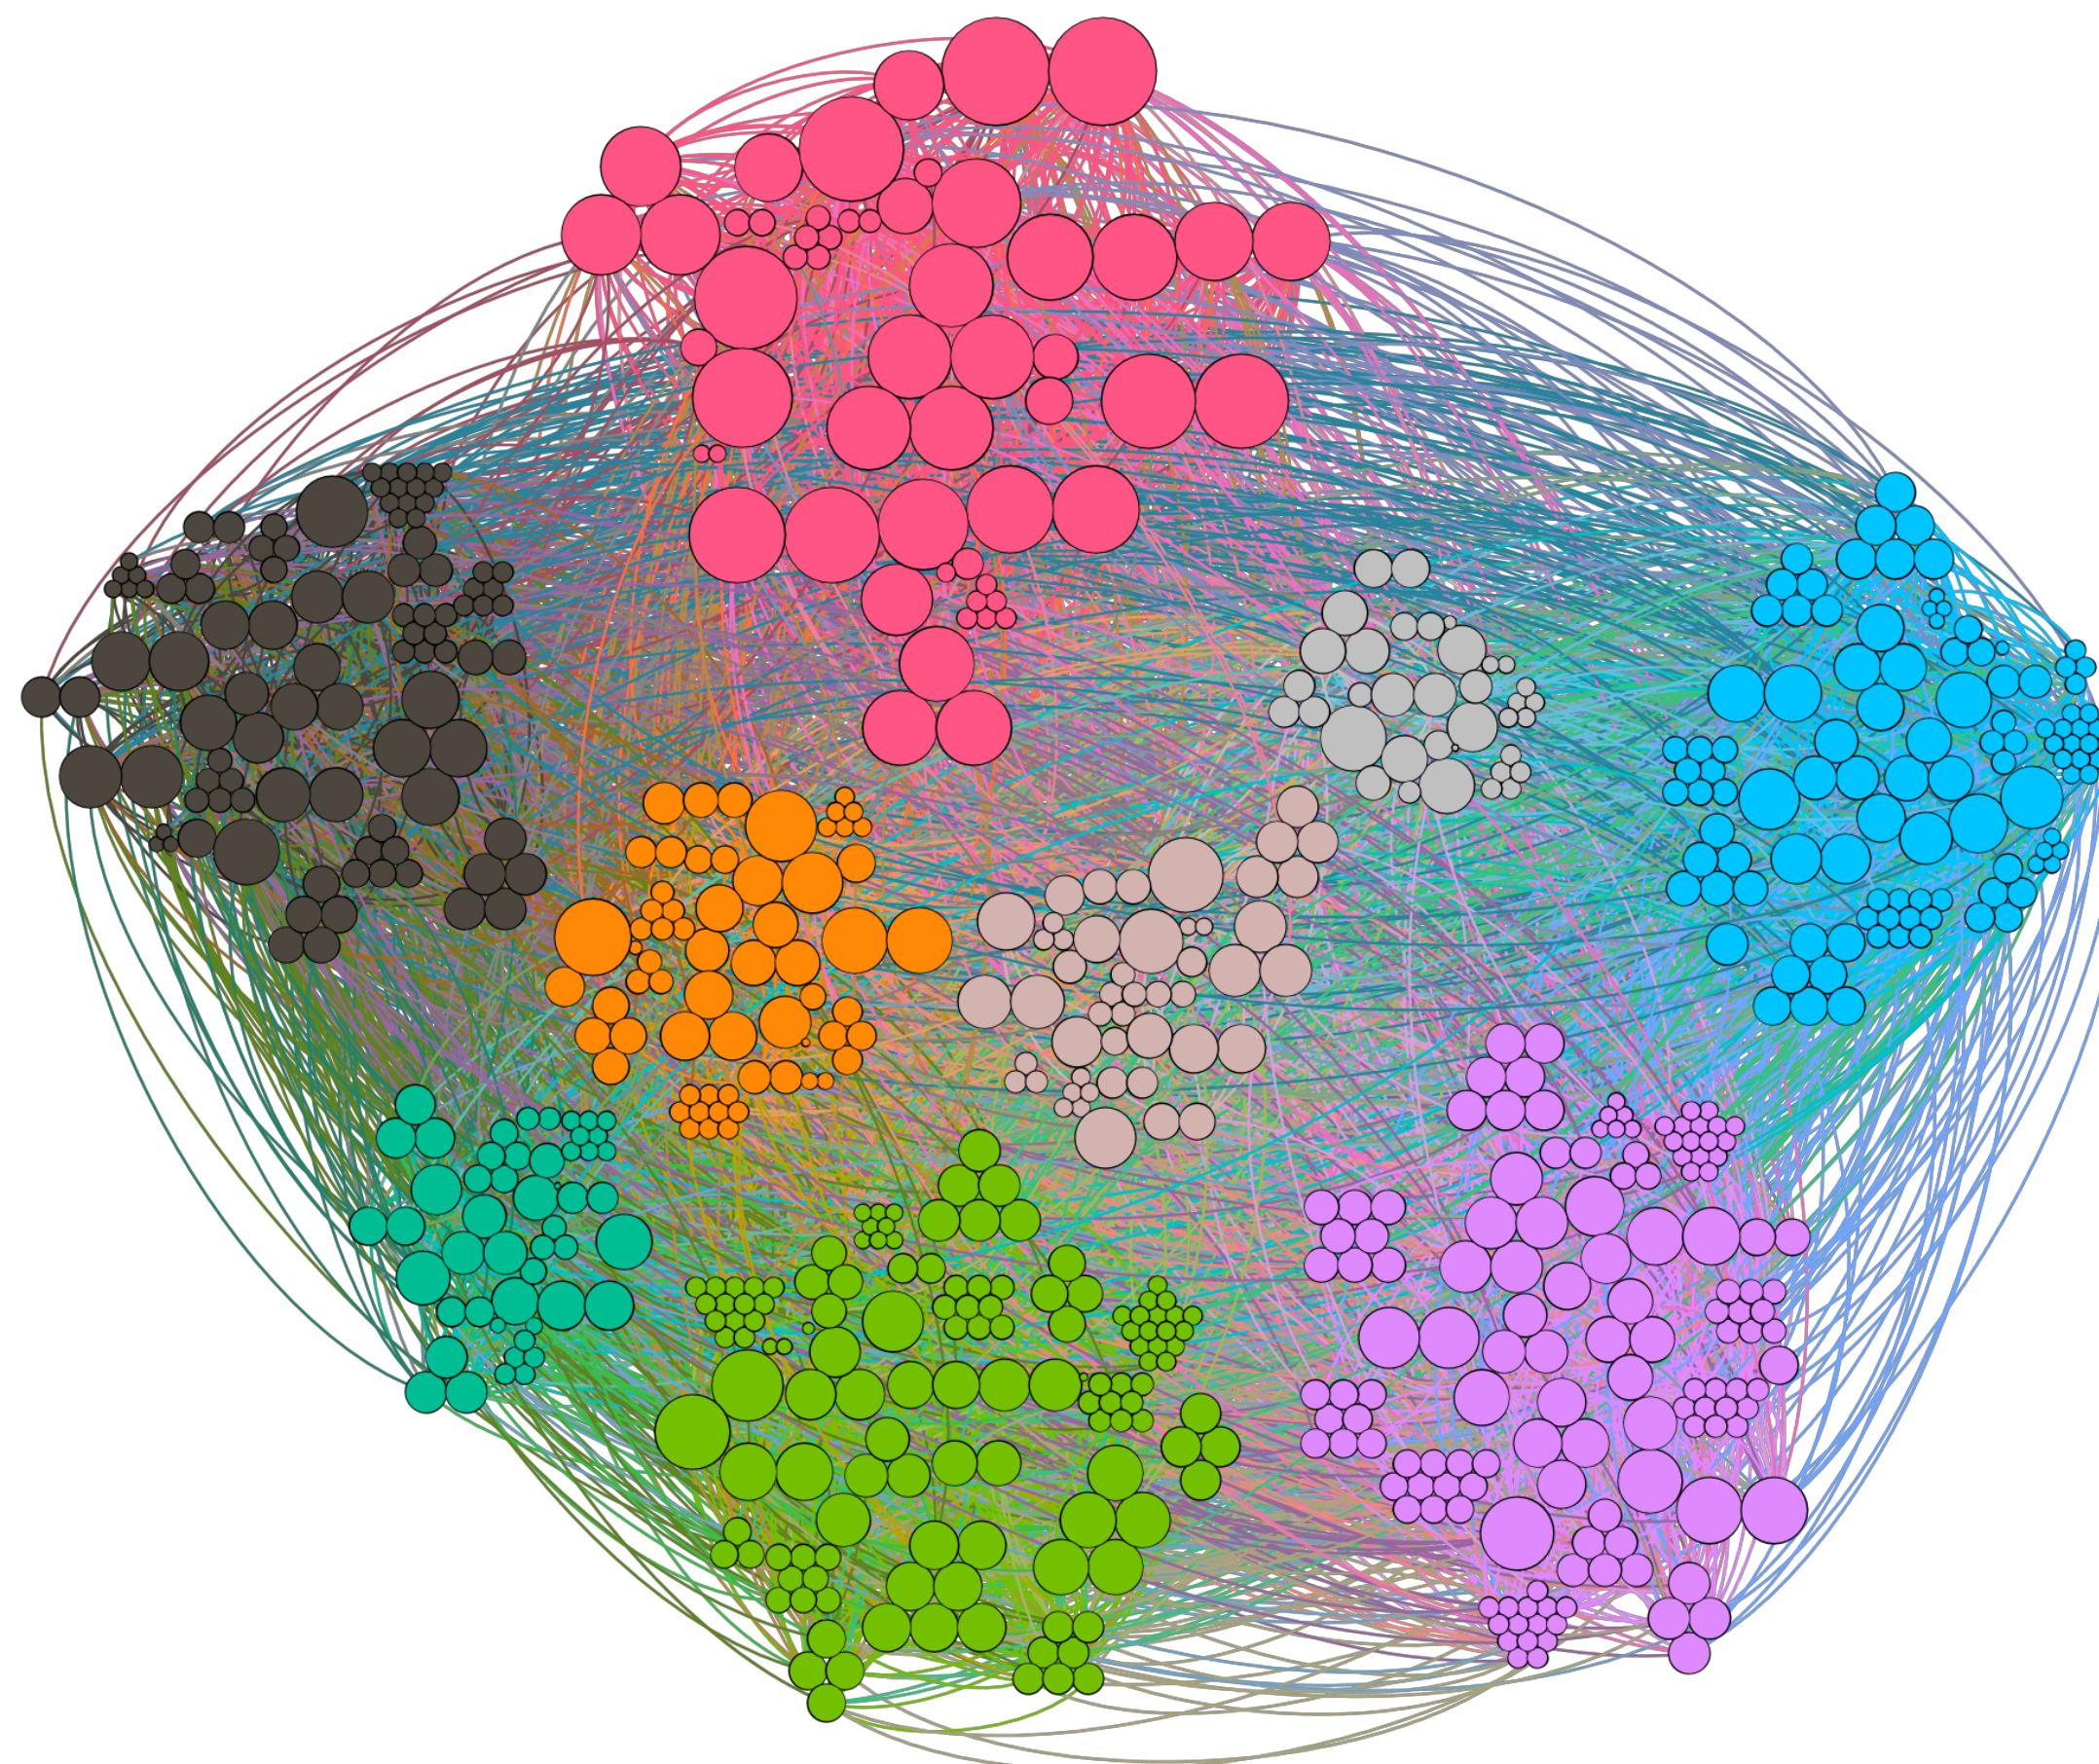
Does the network really look random?

How can we known if we don't know what the topology looks like?

Can we do anything to infer the topology?

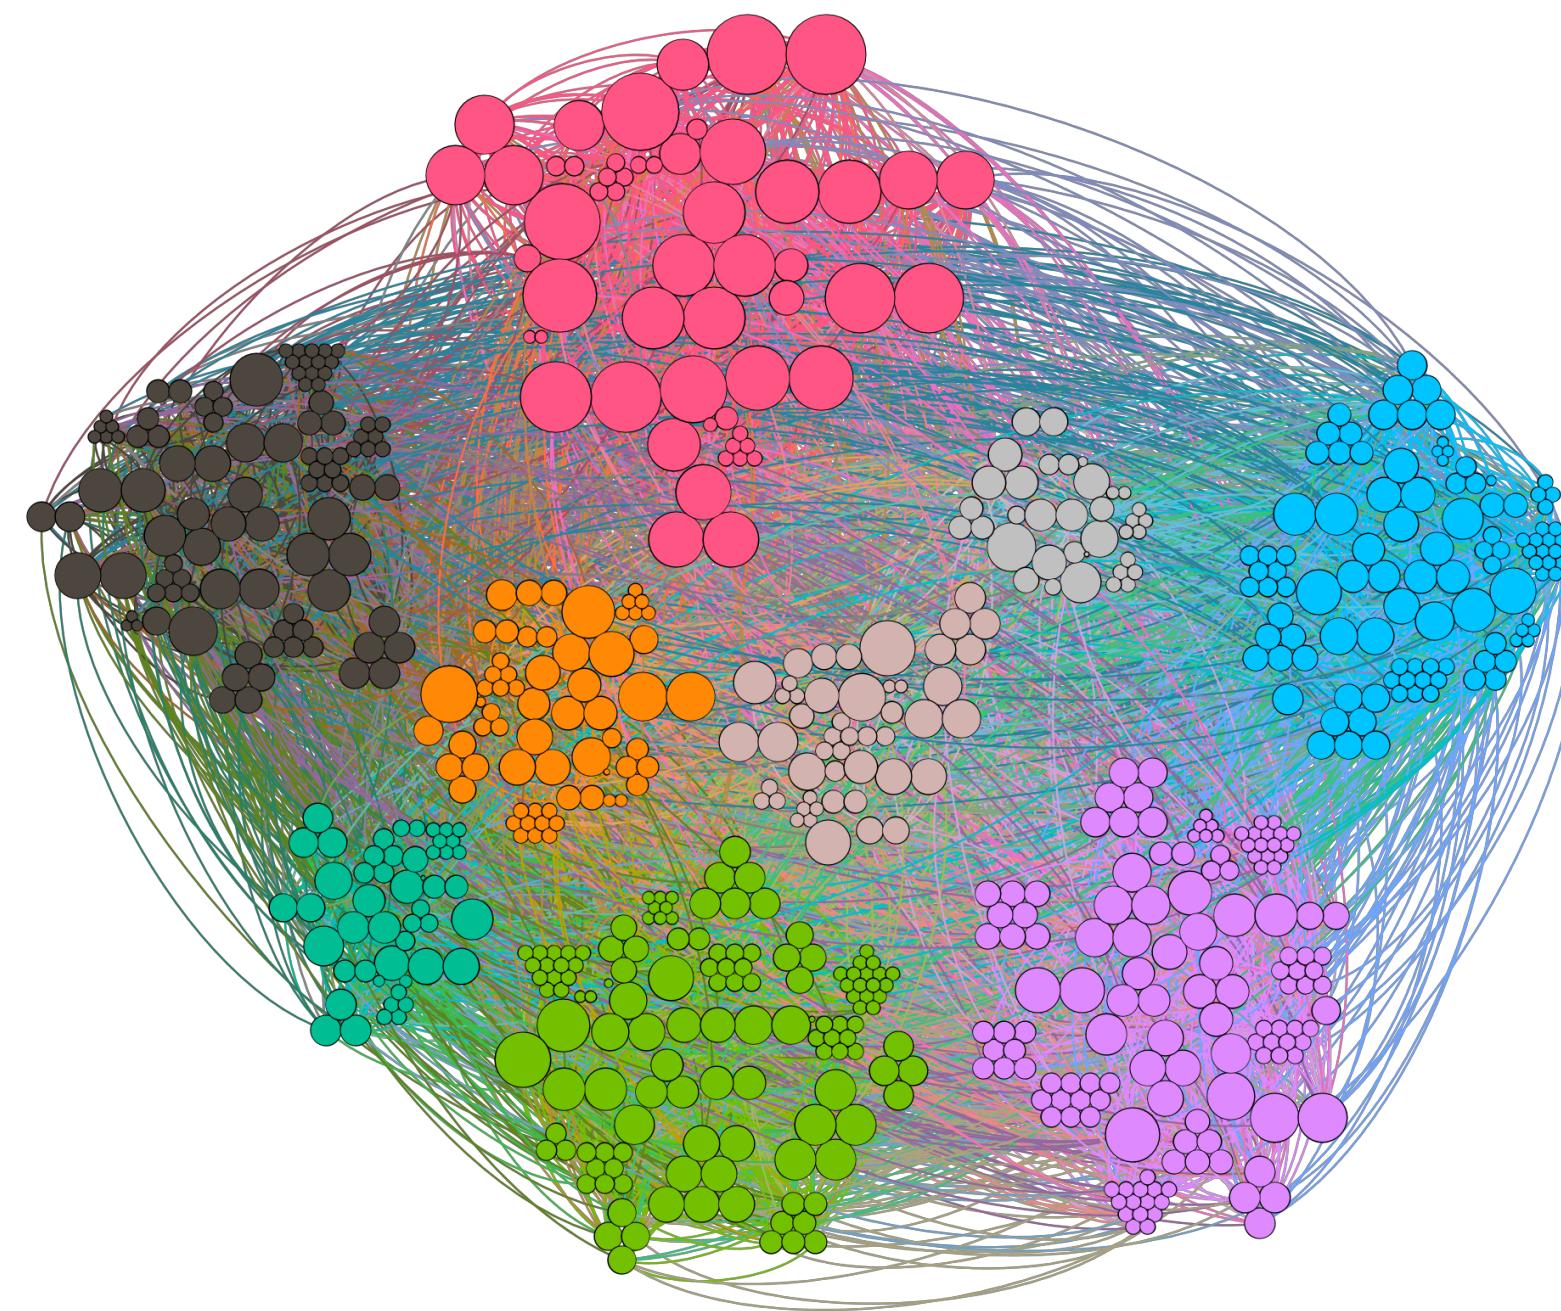


TESTNET TOPOLOGY



- Several communities can be easily identified
- The network looks far from a random graph of similar characteristics
- The topology can be used to identify undesired centralization
- But also to target some potential victims (e.g: Eclipse attacks)

TESTNET TOPOLOGY



- Several communities can be easily identified
- The network looks far from a random graph of similar characteristics
- The topology can be used to identify undesired centralization
- But also to target some potential victims (e.g: Eclipse attacks)

Sergi Delgado-Segura, Surya Bakshi, Cristina Pérez-Solà, James Litton, Andrew Pachulski, Andrew Miller, Bobby Bhattacharjee

TxProbe: Discovering Bitcoin's Network Topology Using Orphan Transactions

<https://fc19.ifca.ai/preproceedings/58-preproceedings.pdf>

