

TxProbe: Discovering Bitcoin's Network Topology Using Orphan Transactions

Sergi Delgado-Segura, Surya Bakshi, Cristina Pérez-Solà, James Litton, Andrew Pachulski, Andrew Miller
and Bobby Bhattacharjee



WHAT DO WE KNOW ABOUT THE TOPOLOGY?



WHAT DO WE KNOW ABOUT THE TOPOLOGY?

Number of nodes and location of them

WHAT DO WE KNOW ABOUT THE TOPOLOGY?

Number of nodes and location of them

GLOBAL BITCOIN NODES

DISTRIBUTION

Reachable nodes as of Thu Feb 07 2019
10:26:44 GMT+0000 (Greenwich Mean Time).

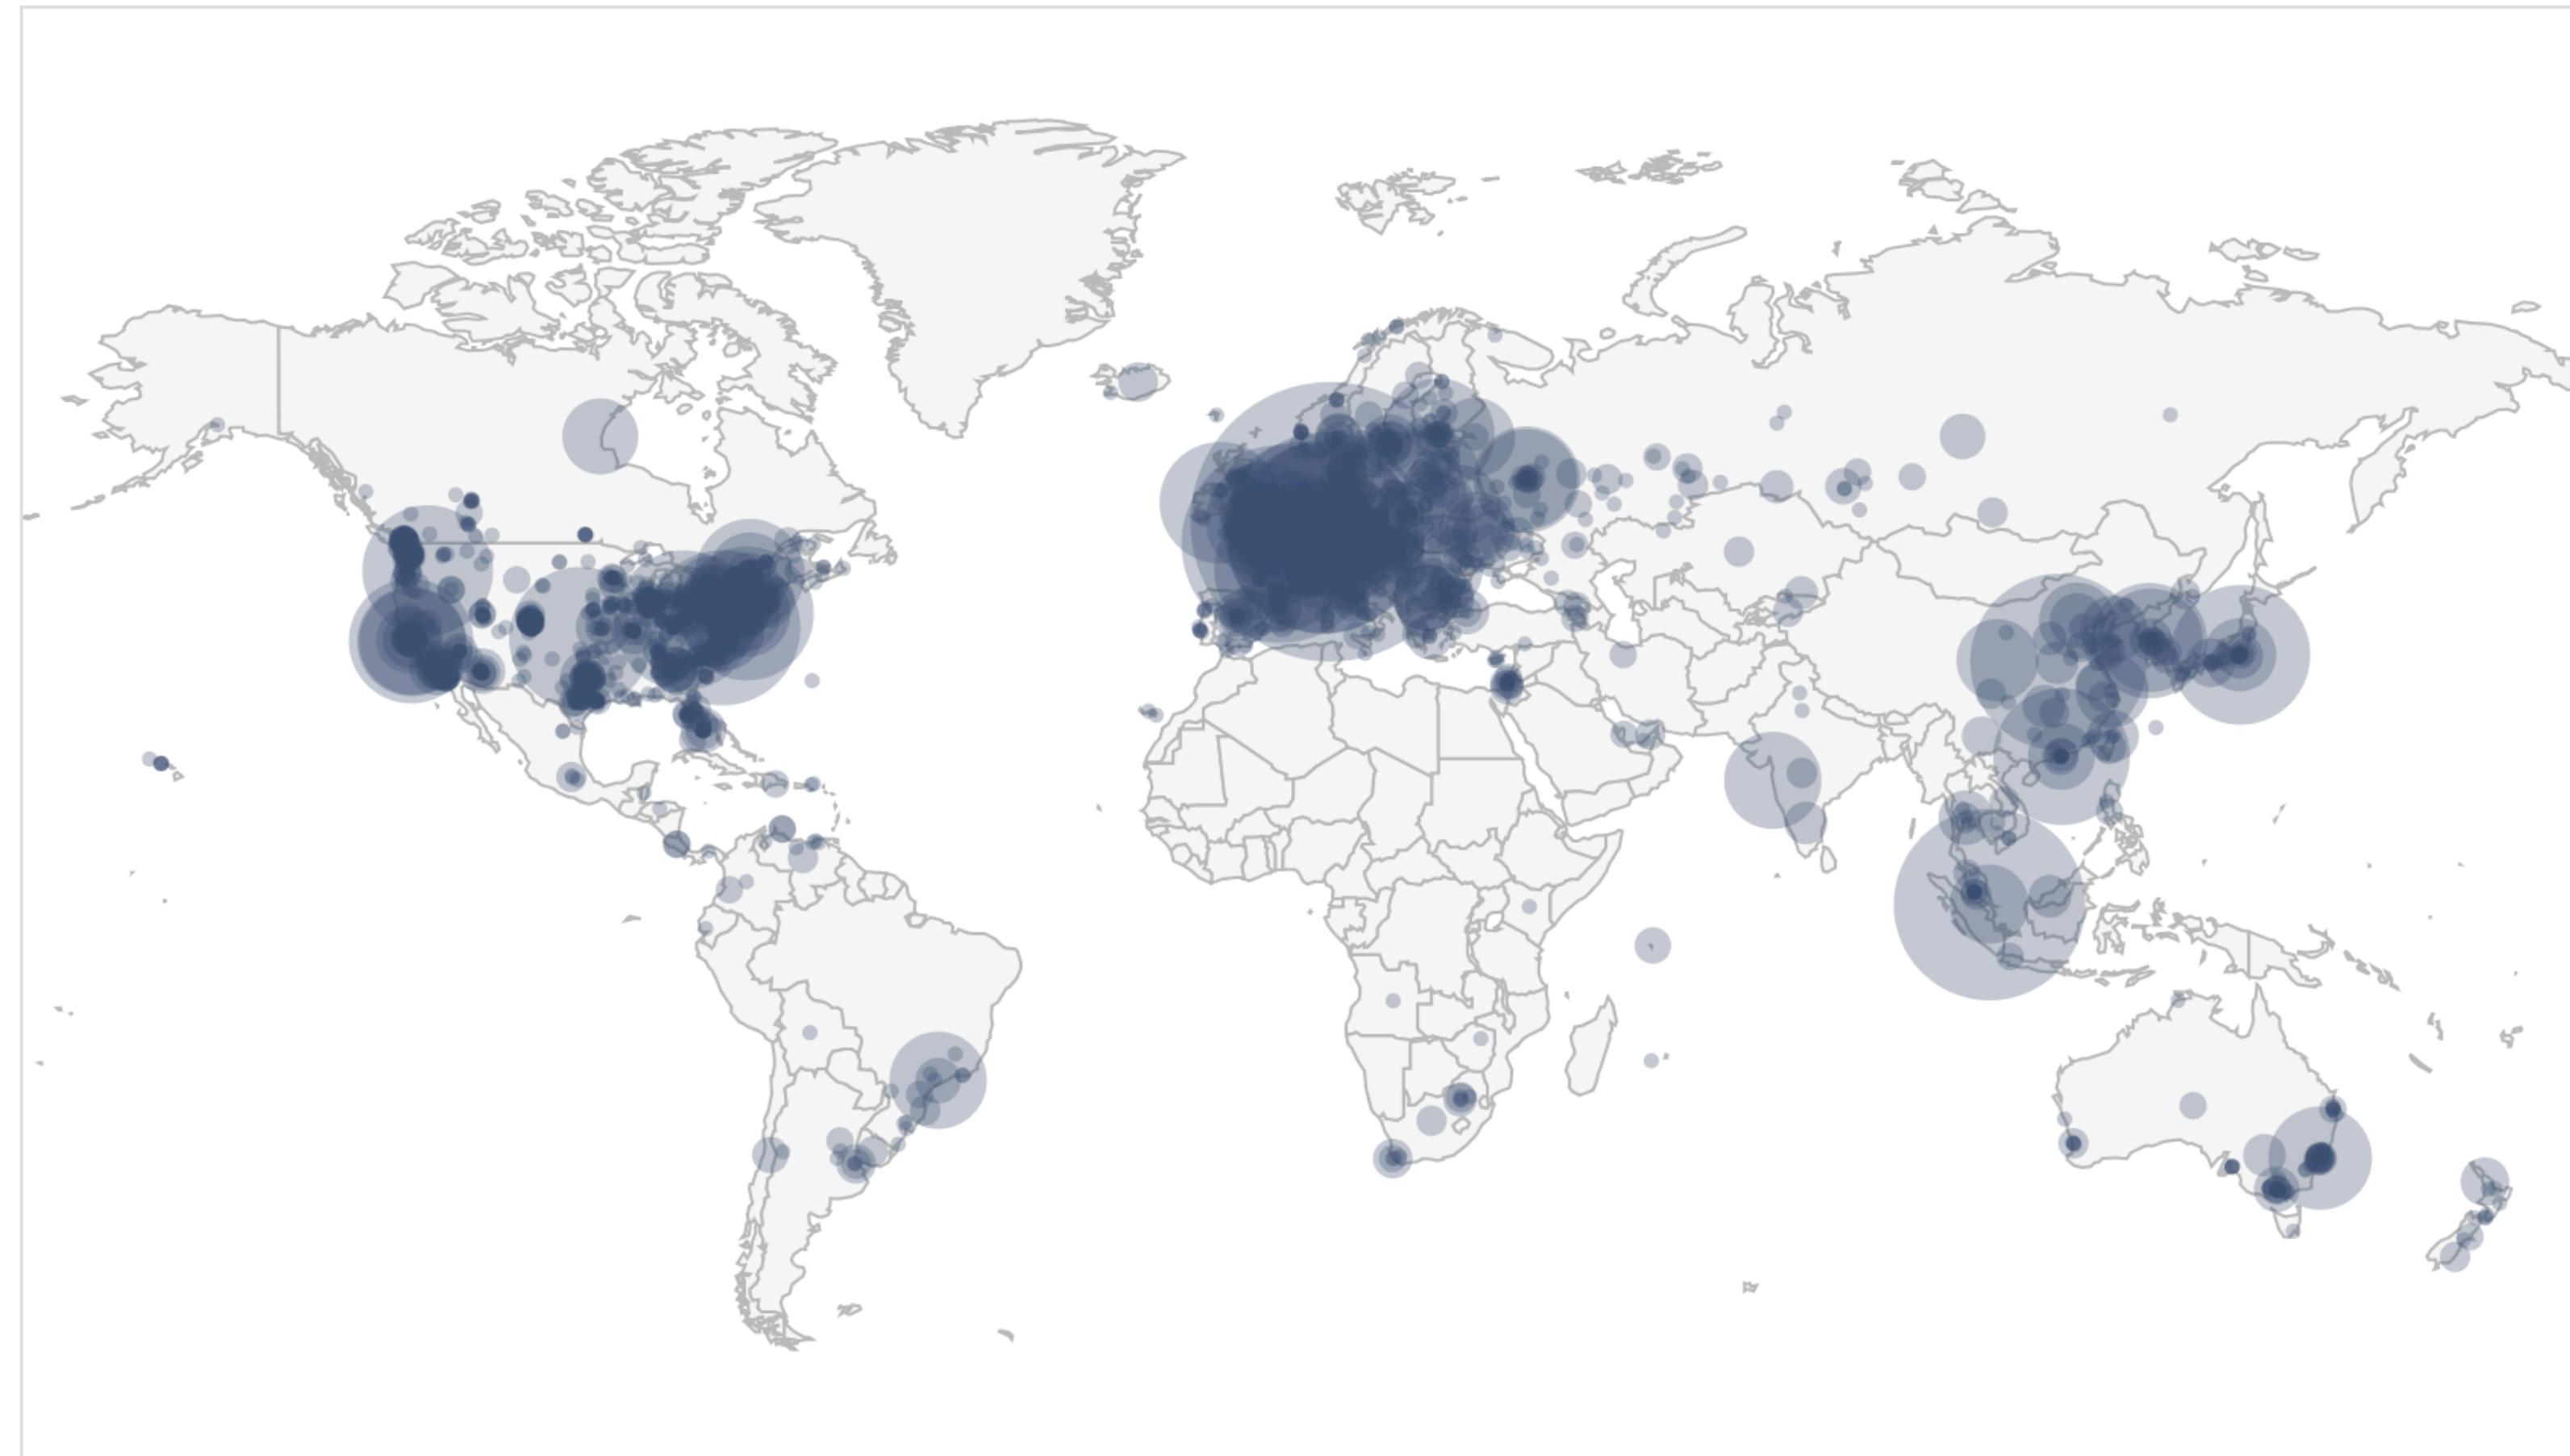
10365 NODES

[24-hour charts »](#)

Top 10 countries with their respective number of reachable nodes are as follow.

RANK	COUNTRY	NODES
1	United States	2570 (24.79%)
2	Germany	1968 (18.99%)
3	France	689 (6.65%)
4	Netherlands	514 (4.96%)
5	China	411 (3.97%)
6	Canada	384 (3.70%)
7	United Kingdom	355 (3.42%)
8	Singapore	321 (3.10%)
9	Russian Federation	277 (2.67%)
10	Japan	228 (2.20%)

[More \(100\) »](#)



WHAT DO WE KNOW ABOUT THE TOPOLOGY?

Number of nodes and location of them

GLOBAL BITCOIN NODES

DISTRIBUTION

Reachable nodes as of Thu Feb 07 2019
10:26:44 GMT+0000 (Greenwich Mean Time).

10365 NODES

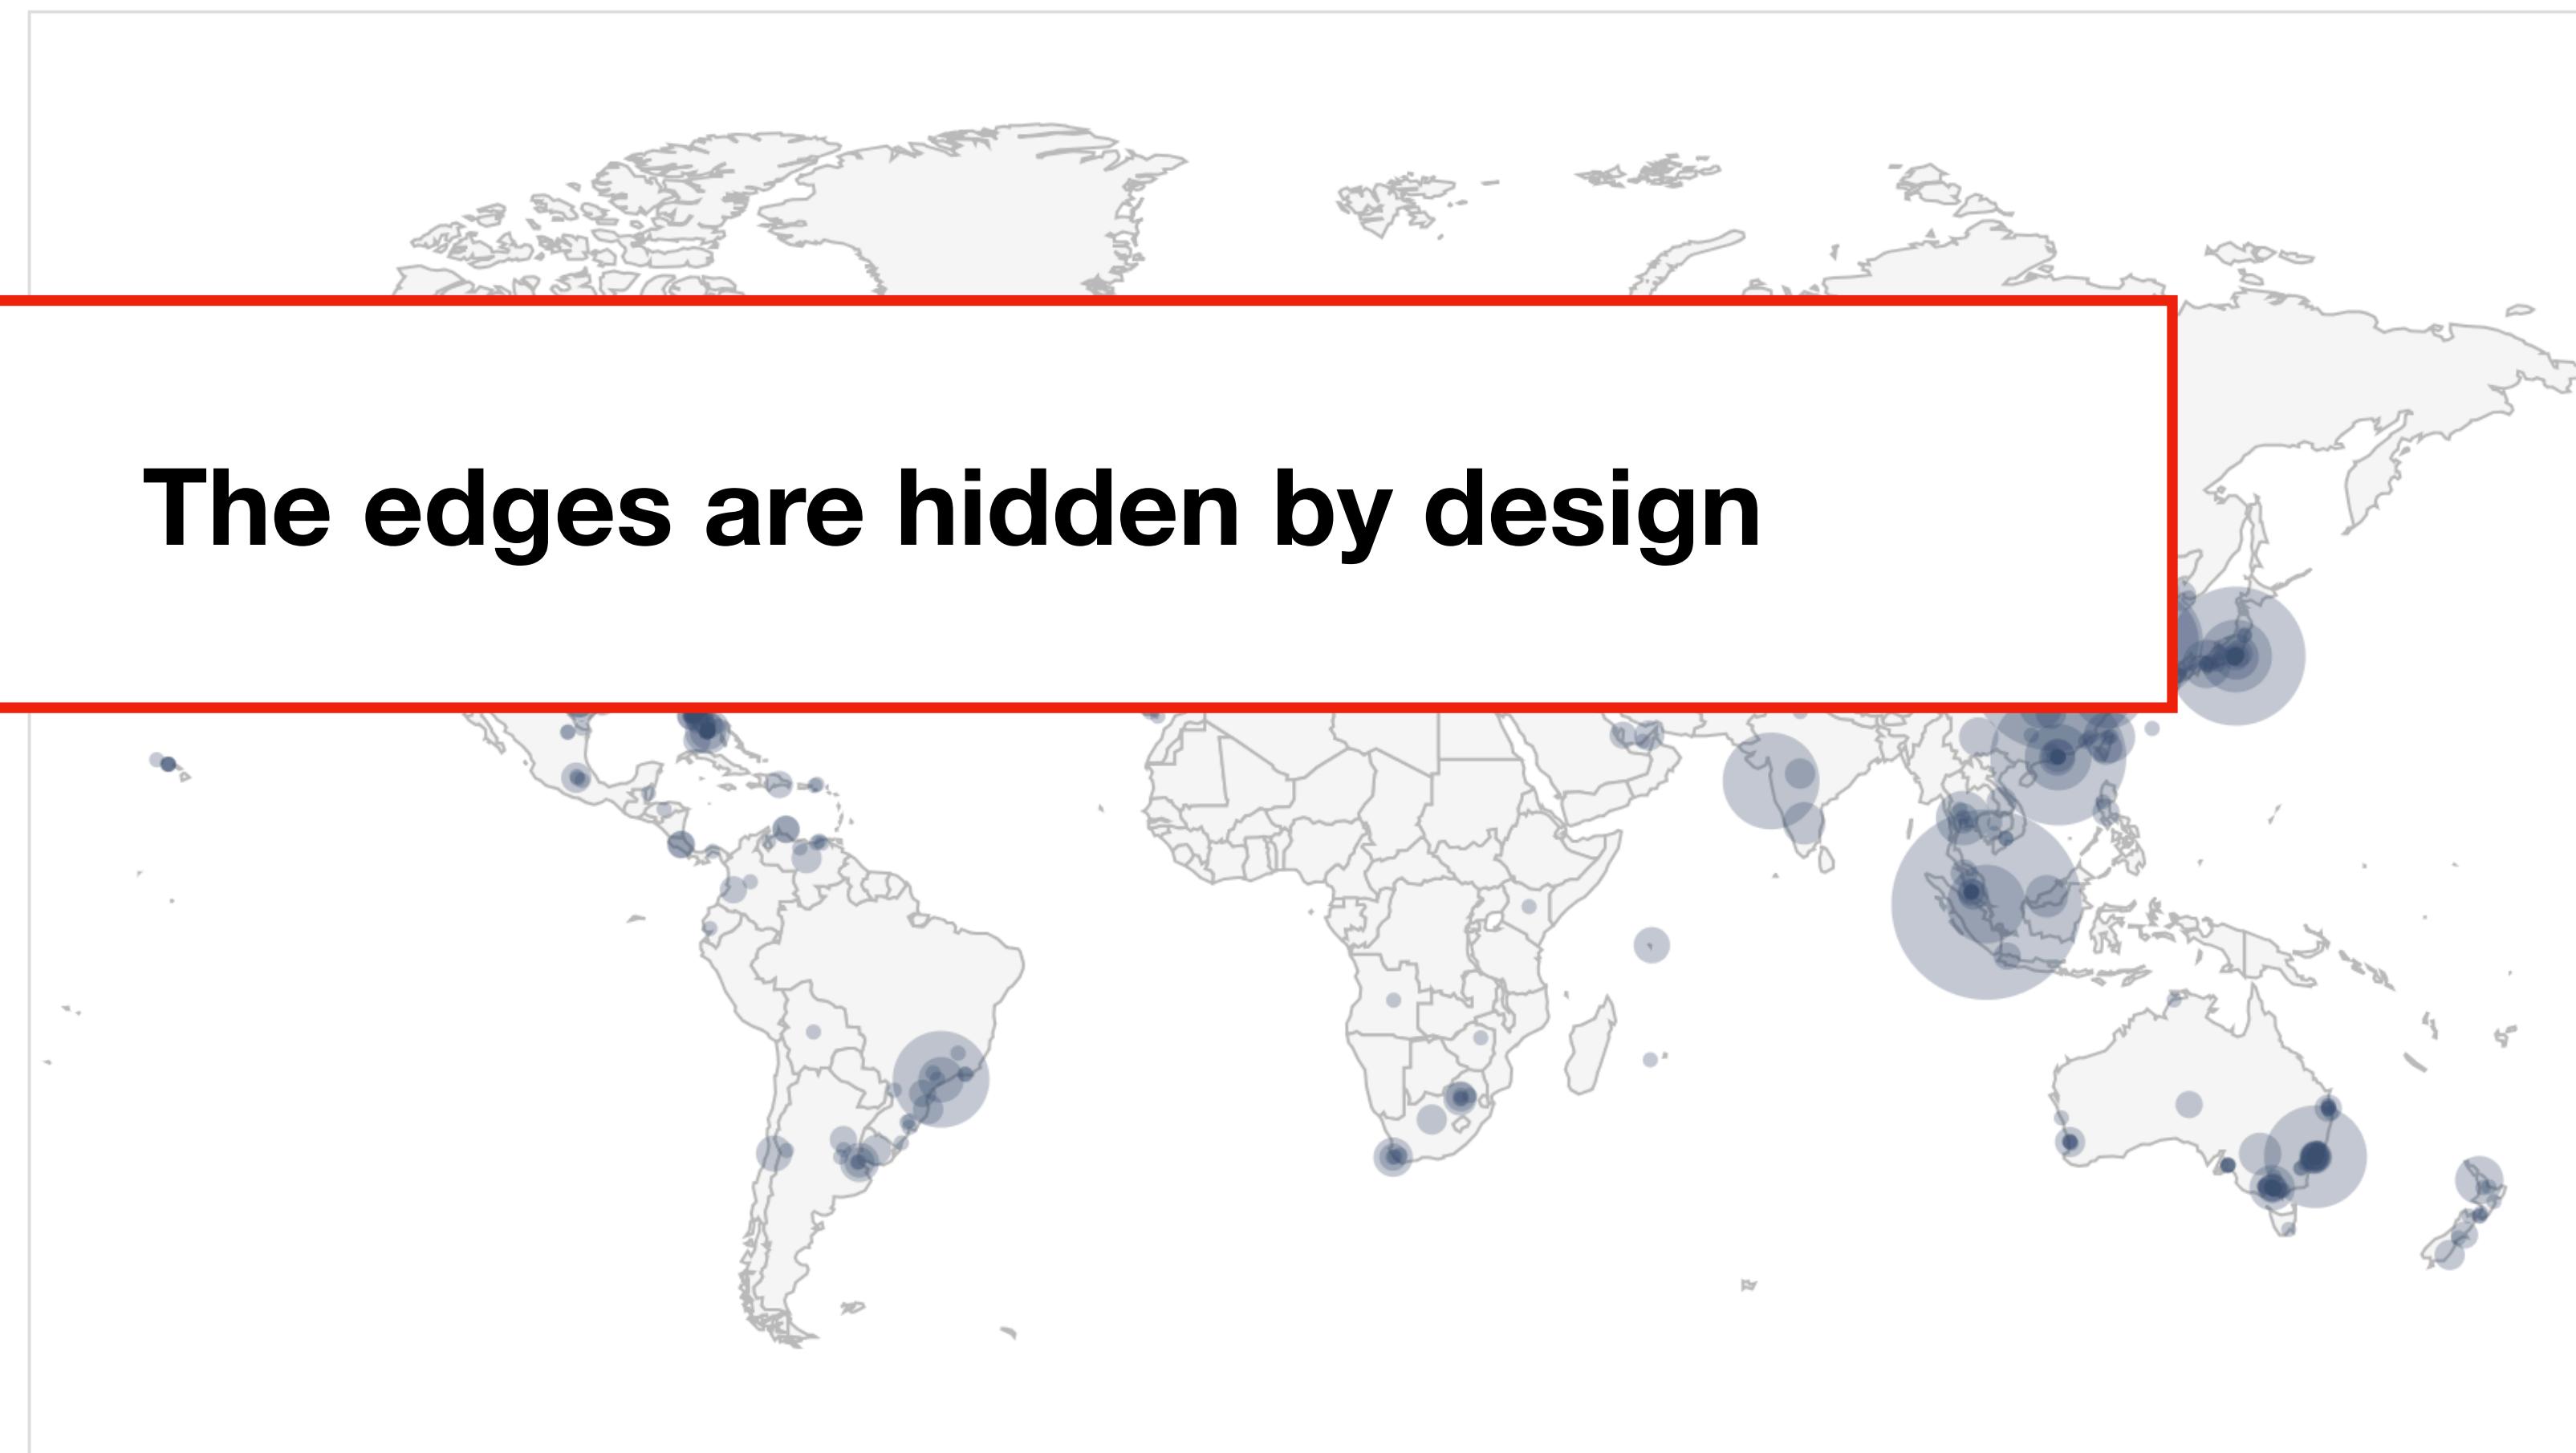
[24-hour charts »](#)

Top 10 countries with their respective
reachable nodes are as follow.

RANK	COUNTRY	NODES
1	United States	2570 (25.00%)
2	Germany	1968 (18.87%)
3	France	689 (6.65%)
4	Netherlands	514 (4.96%)
5	China	411 (3.97%)
6	Canada	384 (3.70%)
7	United Kingdom	355 (3.42%)
8	Singapore	321 (3.10%)
9	Russian Federation	277 (2.67%)
10	Japan	228 (2.20%)

[More \(100\) »](#)

The edges are hidden by design



WHY HAVE A **HIDDEN** TOPOLOGY?

An open topology **could ease** different types of attacks:

- Transaction deanonymization
- Network based attacks (e.g: Eclipse attacks)

The **current approach** of the Bitcoin Core is to **keep it hidden**

WHY HAVE AN **OPEN** TOPOLOGY?

We know nothings about how the network really is:

- Is the network decentralised?
- Are there supernodes controlling the network traffic?
 - Information withholding
 - Censorship
- Are there weak spots in the network that can be easily isolated?

Security by obscurity does not seem the proper way to go

THE TOPOLOGY SHOULD LOOK RANDOM

How Bitcoin (Core client) nodes choose their peers?

- Pseudorandomly from the **addrman**
- **8 outbound** connections by default

No pair of nodes in the same **/16** (IPv4)

- **117 inbound** connection by default (no IP restriction here)

Bitcoin forks based on the Core client follow the same approach

BACKGROUND

Our inferring technique is based on **transaction propagation**

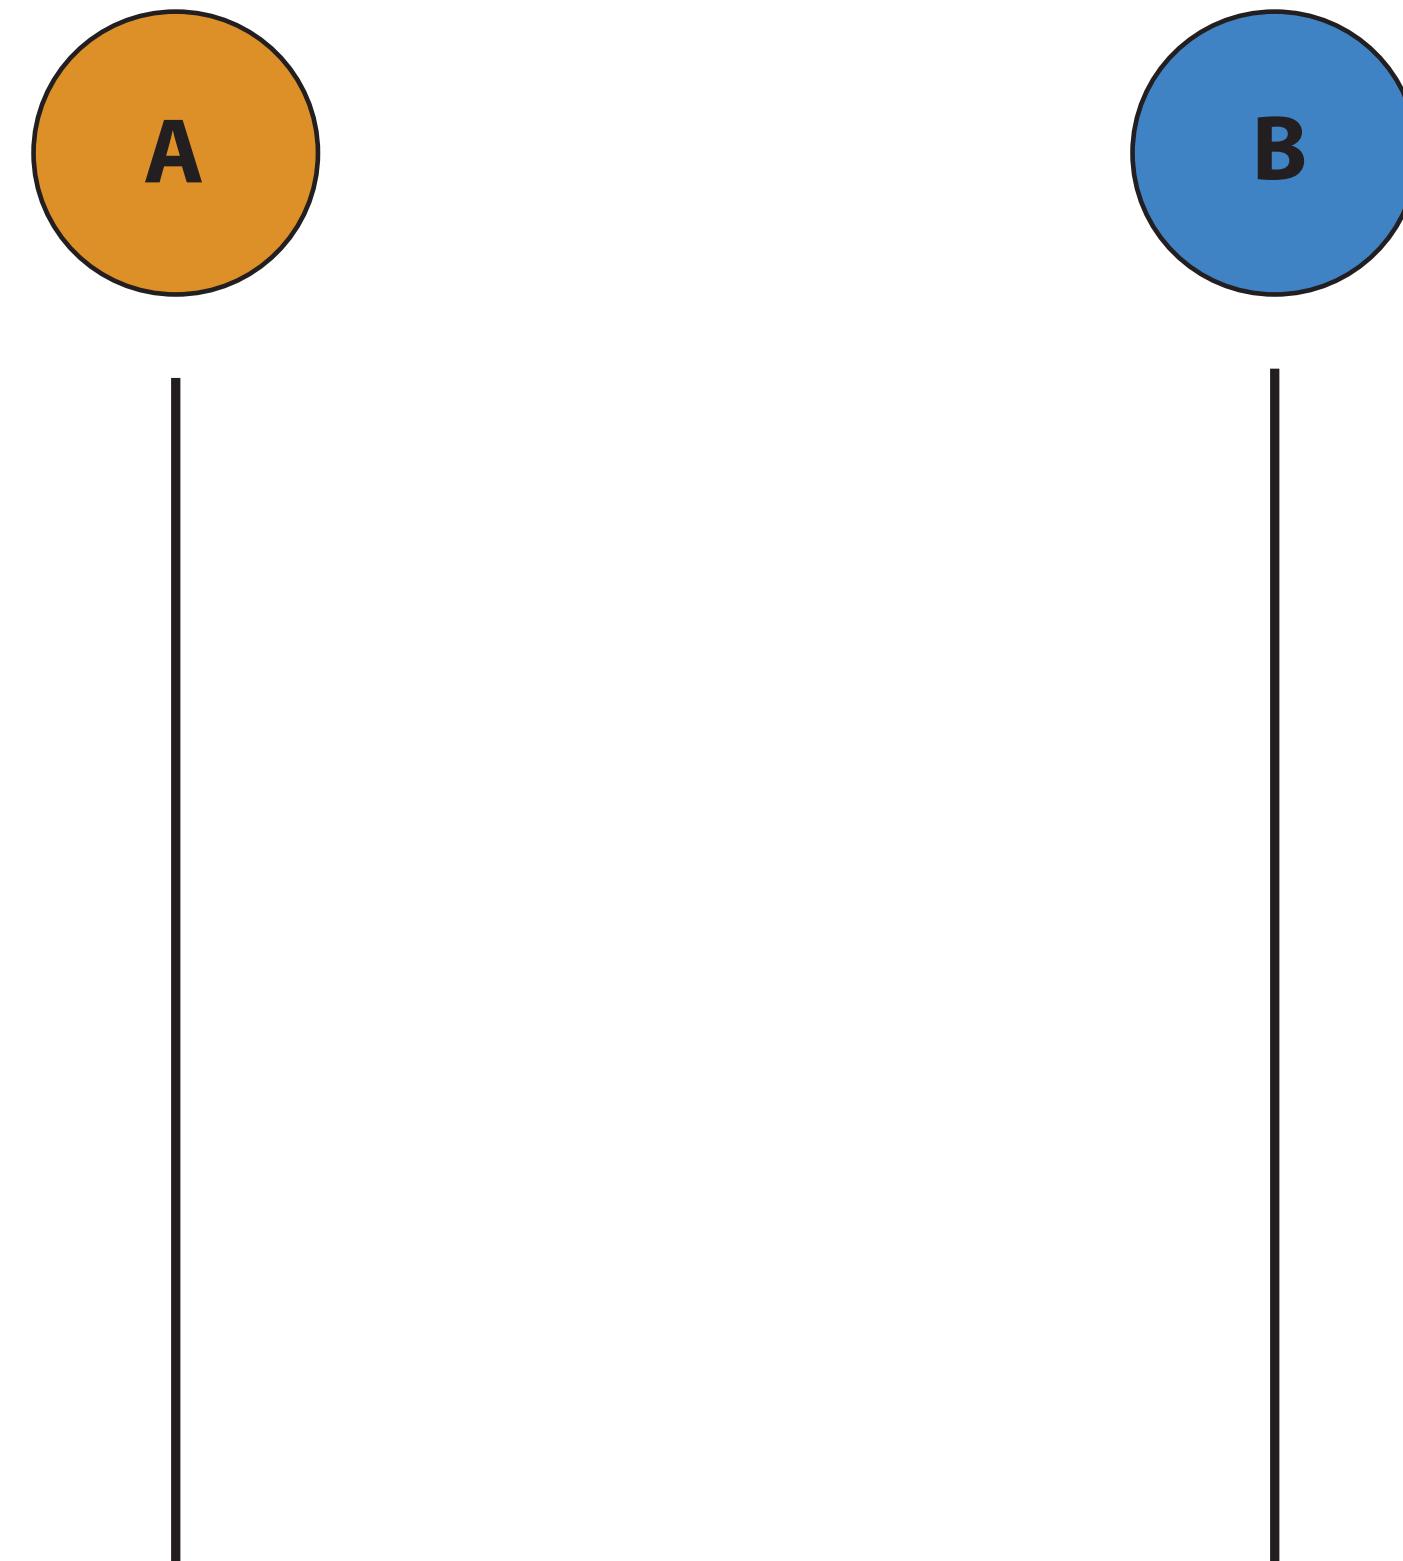
We take advantage of how transactions are handled by nodes:

- **orphans transactions**
- **double-spending transactions**

TRANSACTION PROPAGATION IN BITCOIN

Valid transaction are stored in **mempool**

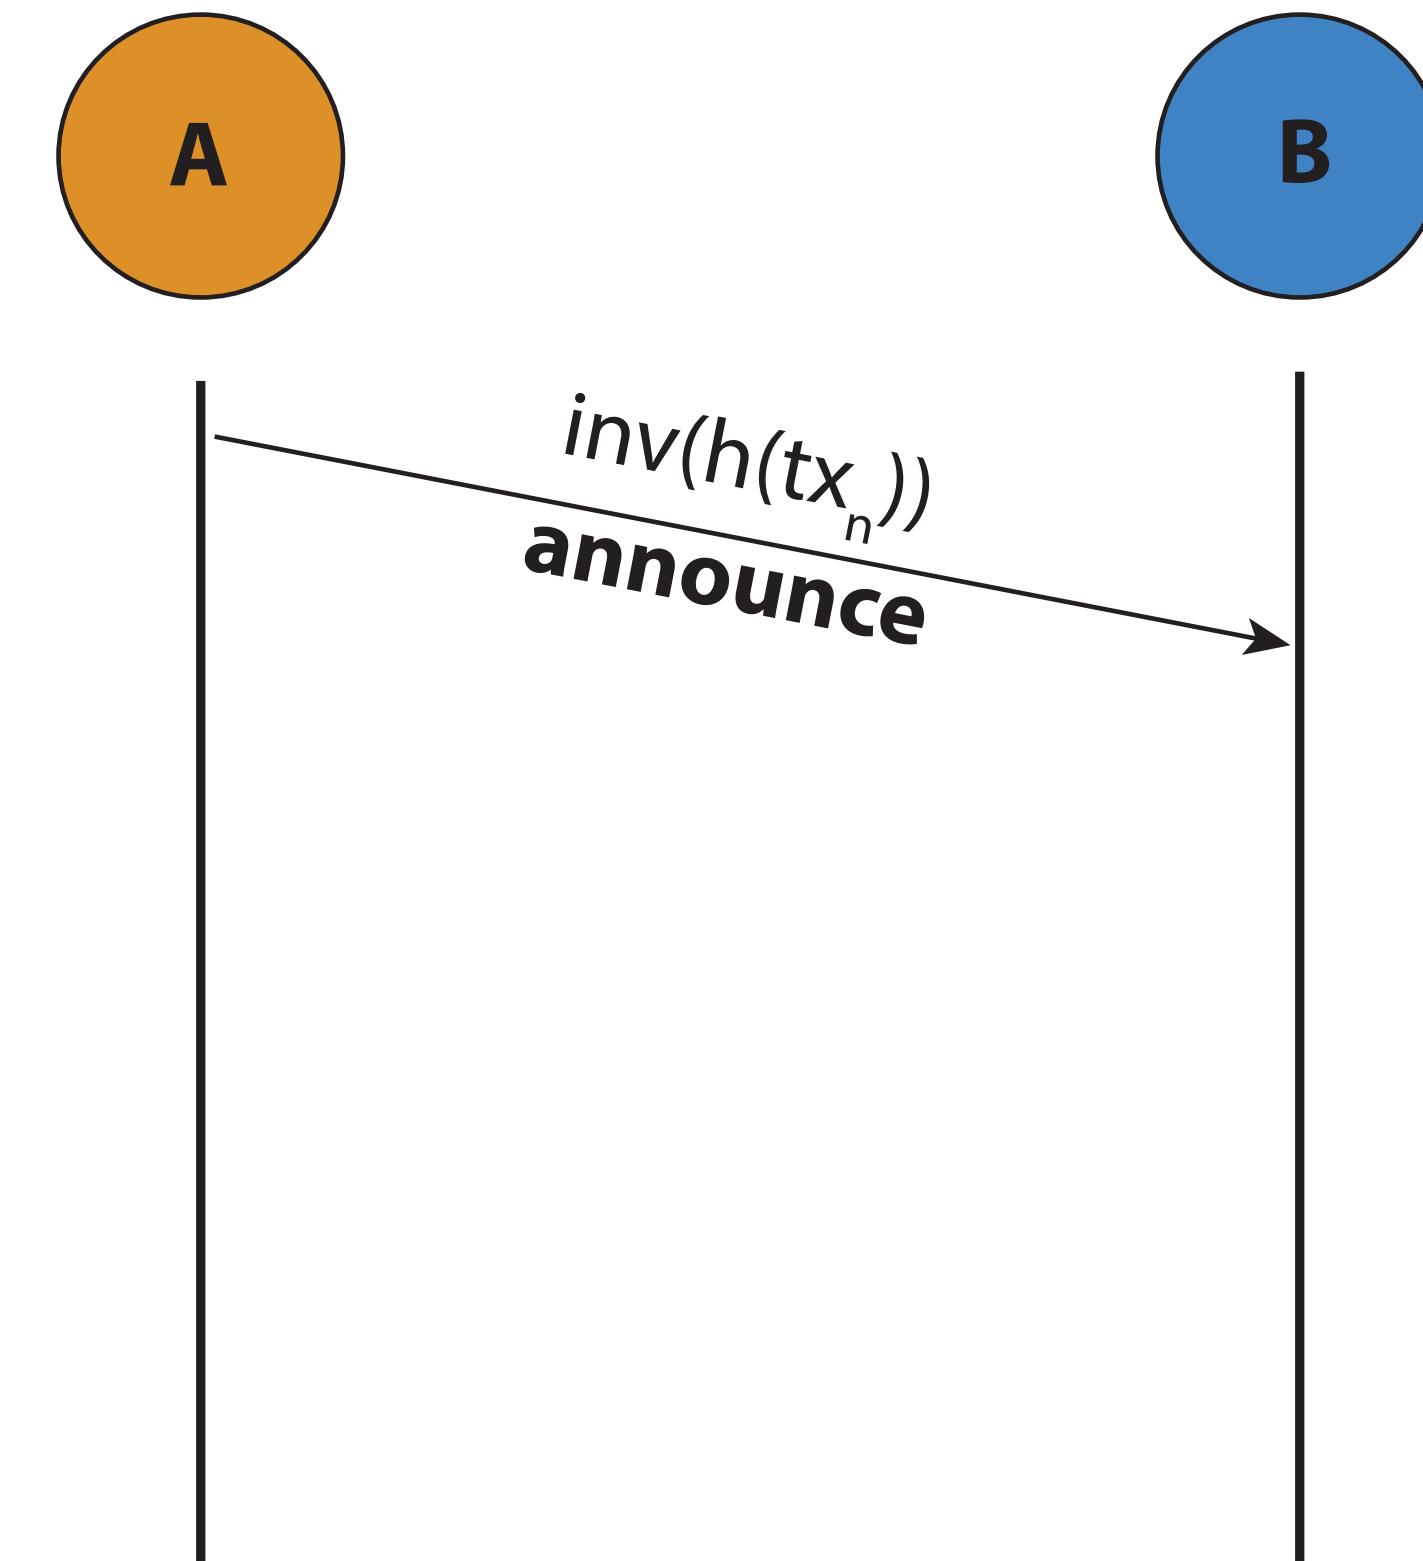
Transaction in mempool are eventually propagated throughout the node neighborhood



TRANSACTION PROPAGATION IN BITCOIN

Valid transaction are stored in **mempool**

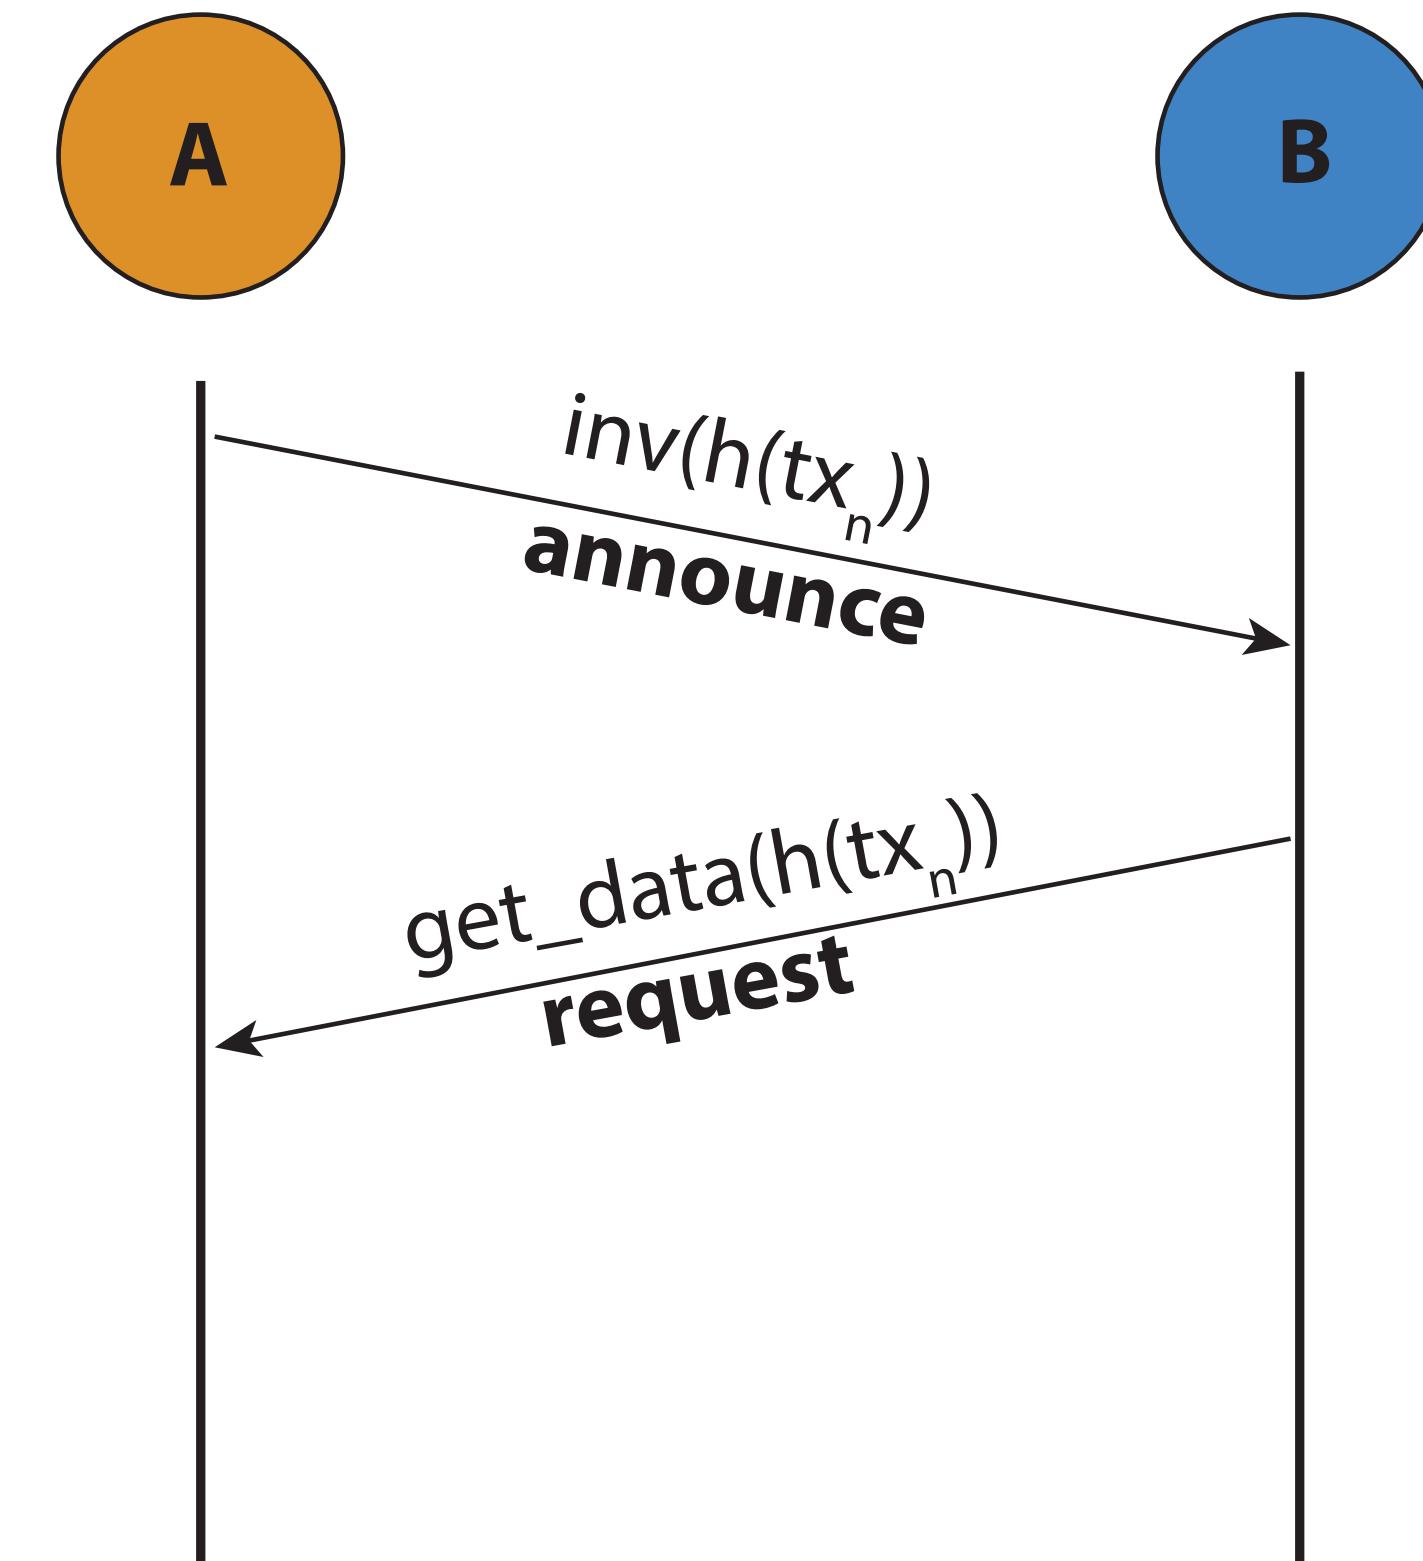
Transaction in mempool are eventually propagated throughout the node neighborhood



TRANSACTION PROPAGATION IN BITCOIN

Valid transaction are stored in **mempool**

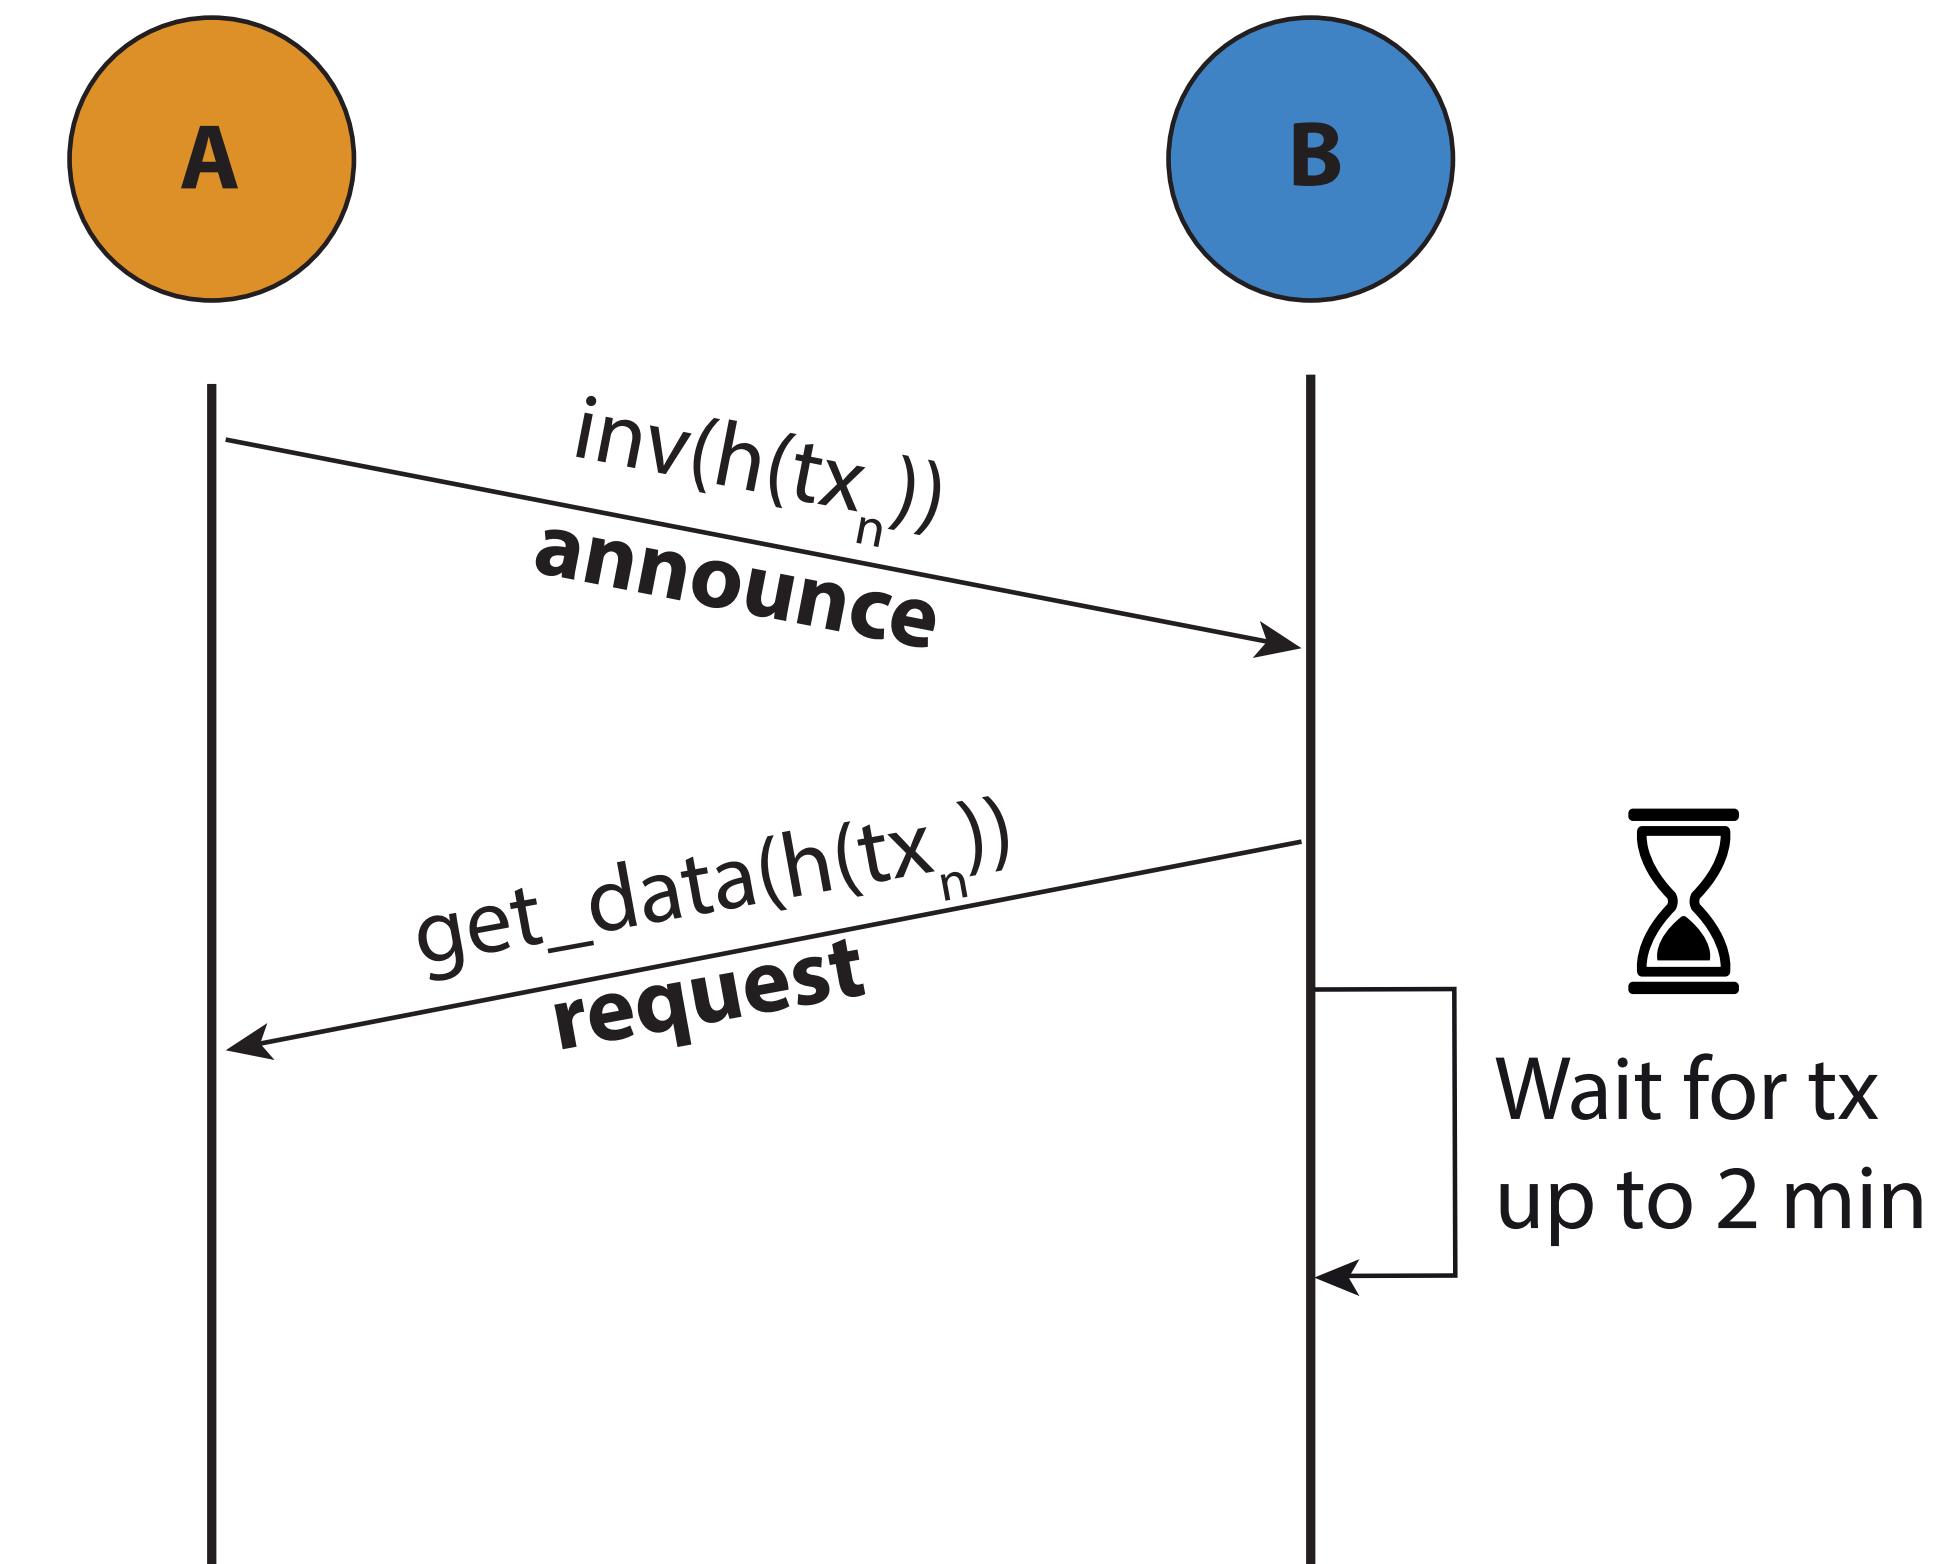
Transaction in mempool are eventually propagated throughout the node neighborhood



TRANSACTION PROPAGATION IN BITCOIN

Valid transaction are stored in **mempool**

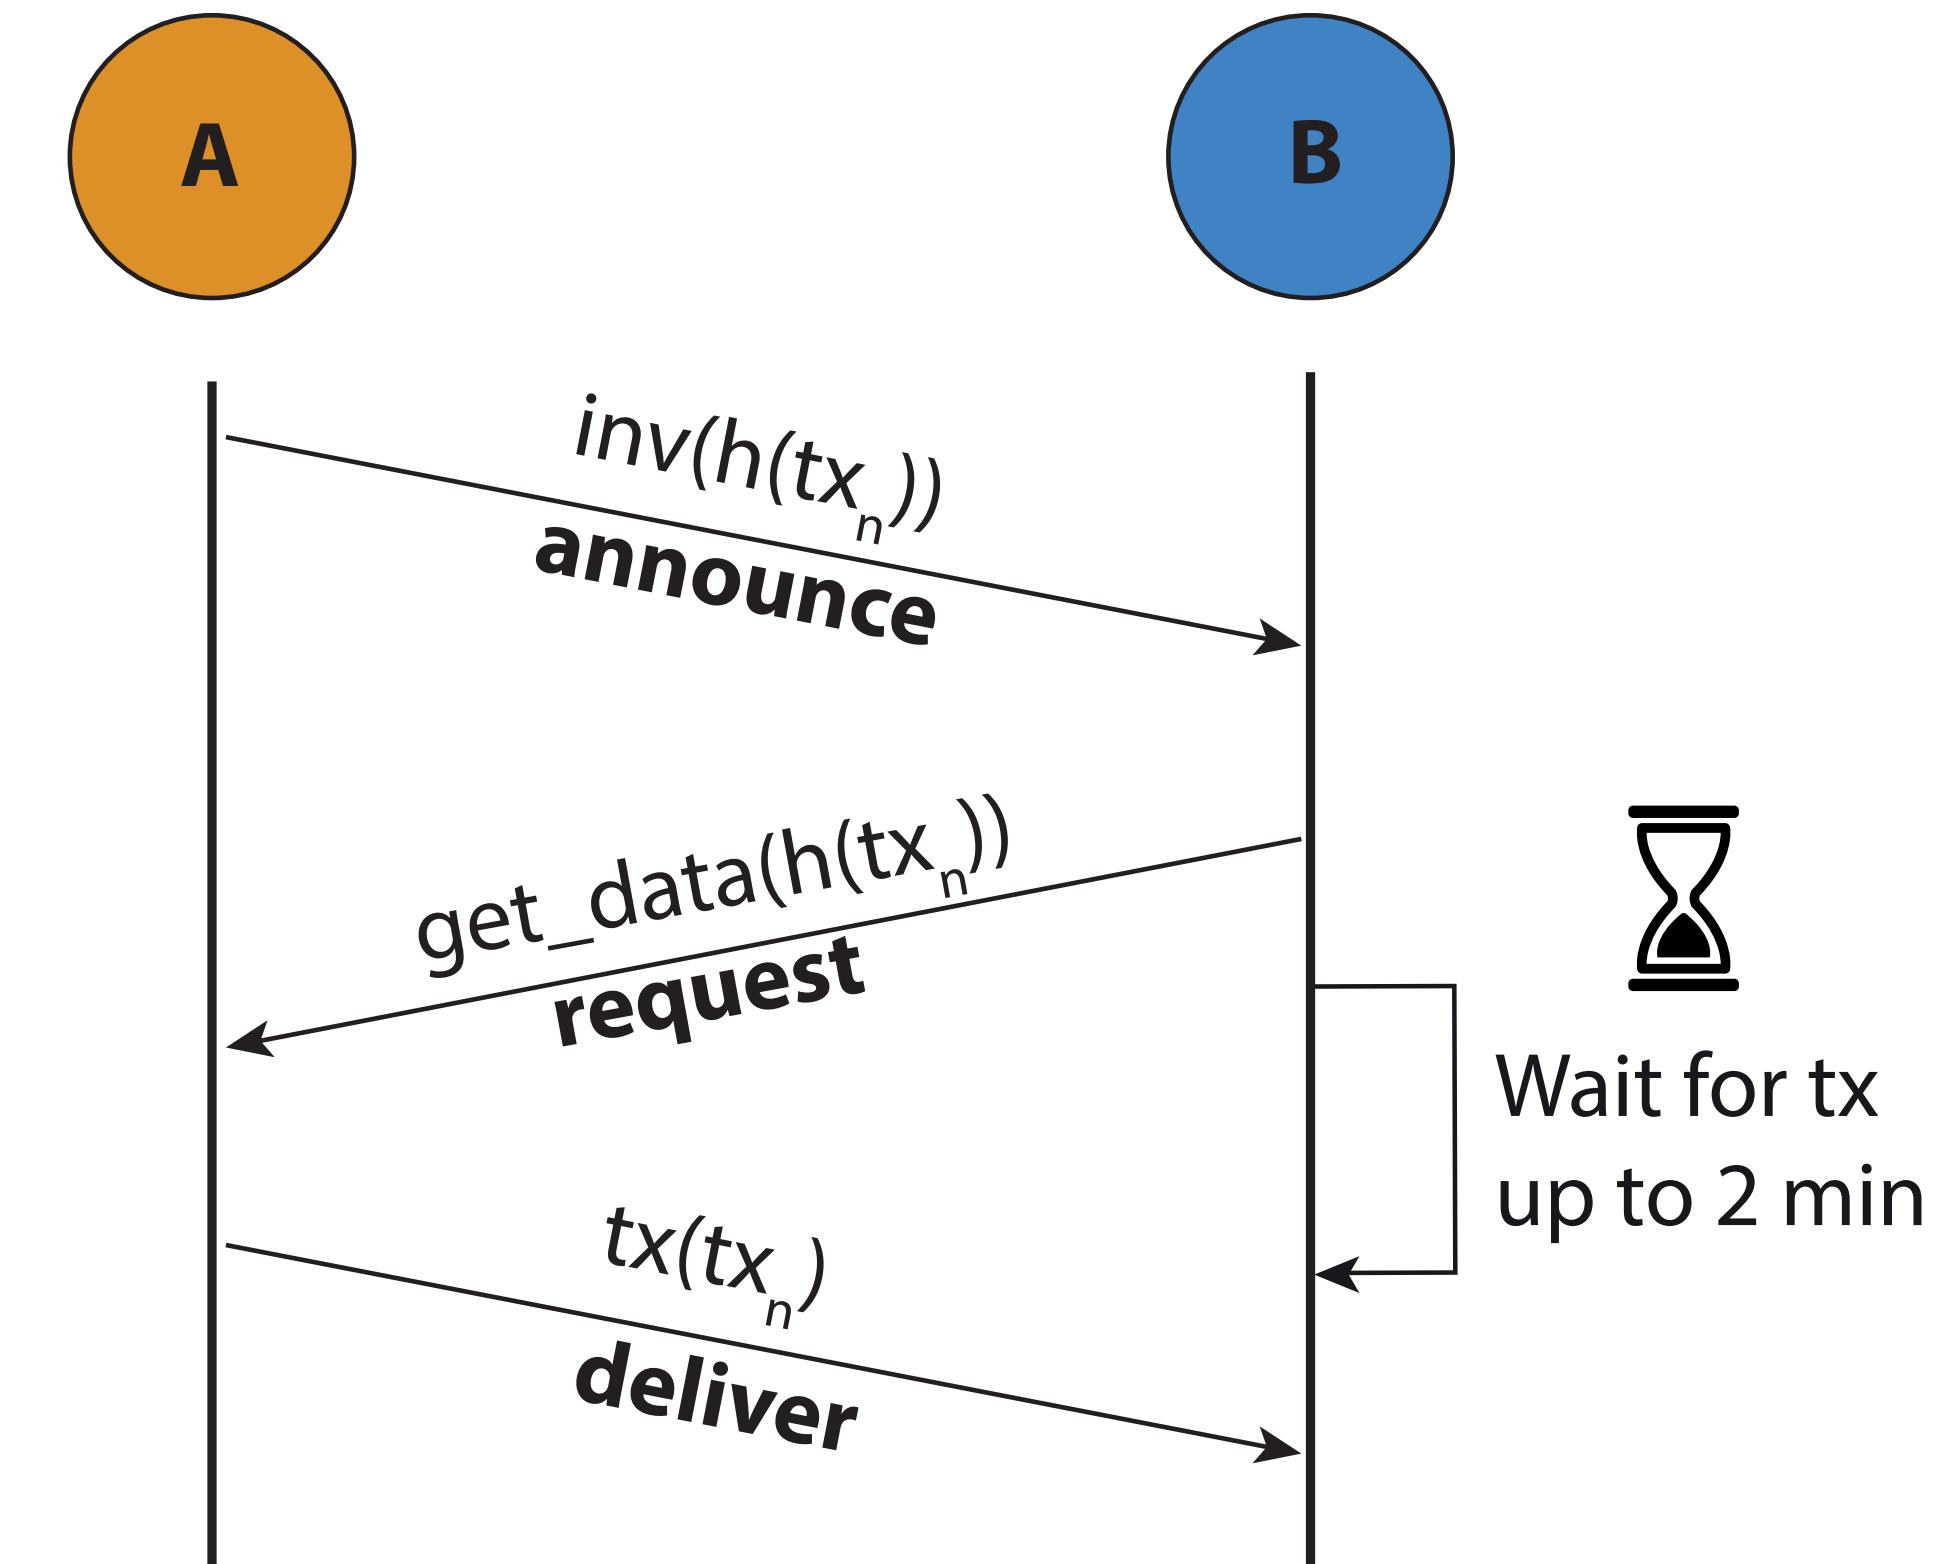
Transaction in mempool are eventually propagated throughout the node neighborhood



TRANSACTION PROPAGATION IN BITCOIN

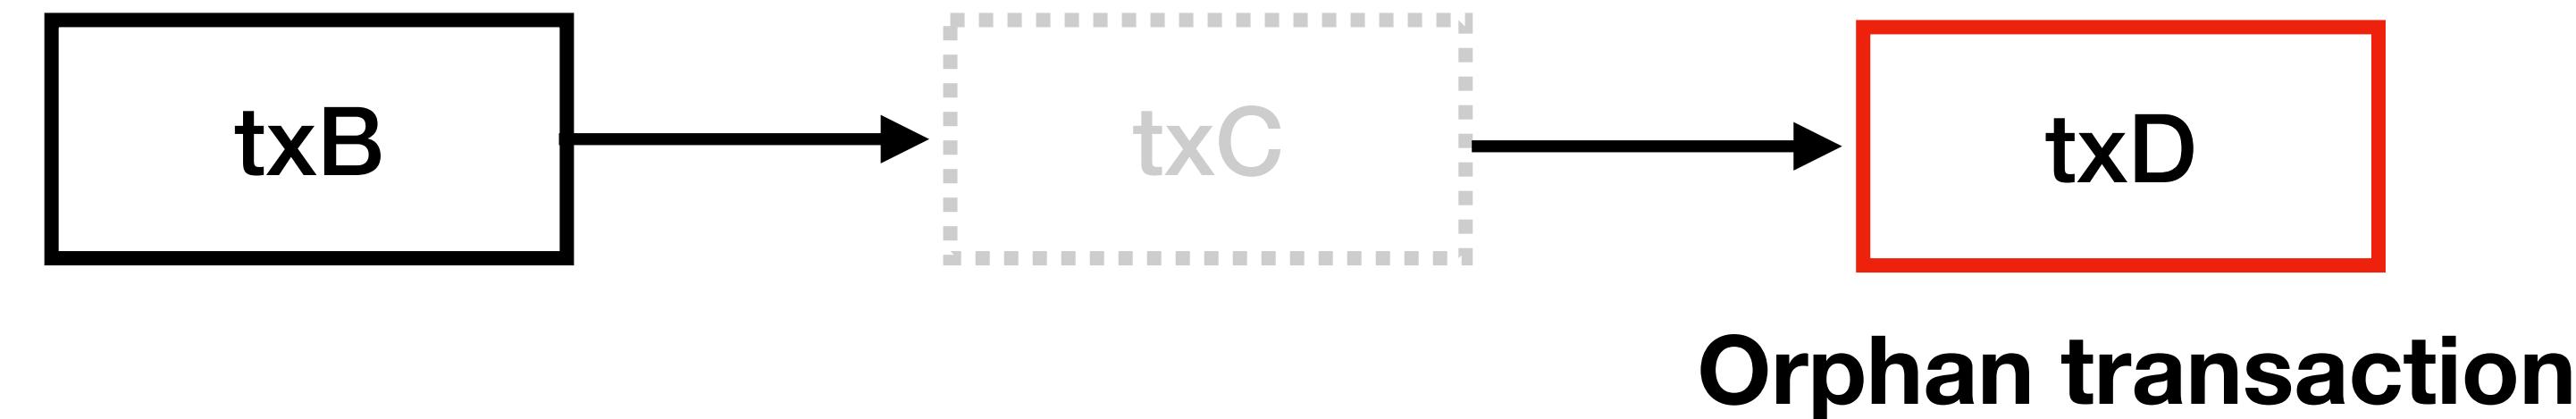
Valid transaction are stored in **mempool**

Transaction in mempool are eventually propagated throughout the node neighborhood



ORPHAN TRANSACTIONS

A transaction is orphan if **some of the referenced UTXOs are unknown**

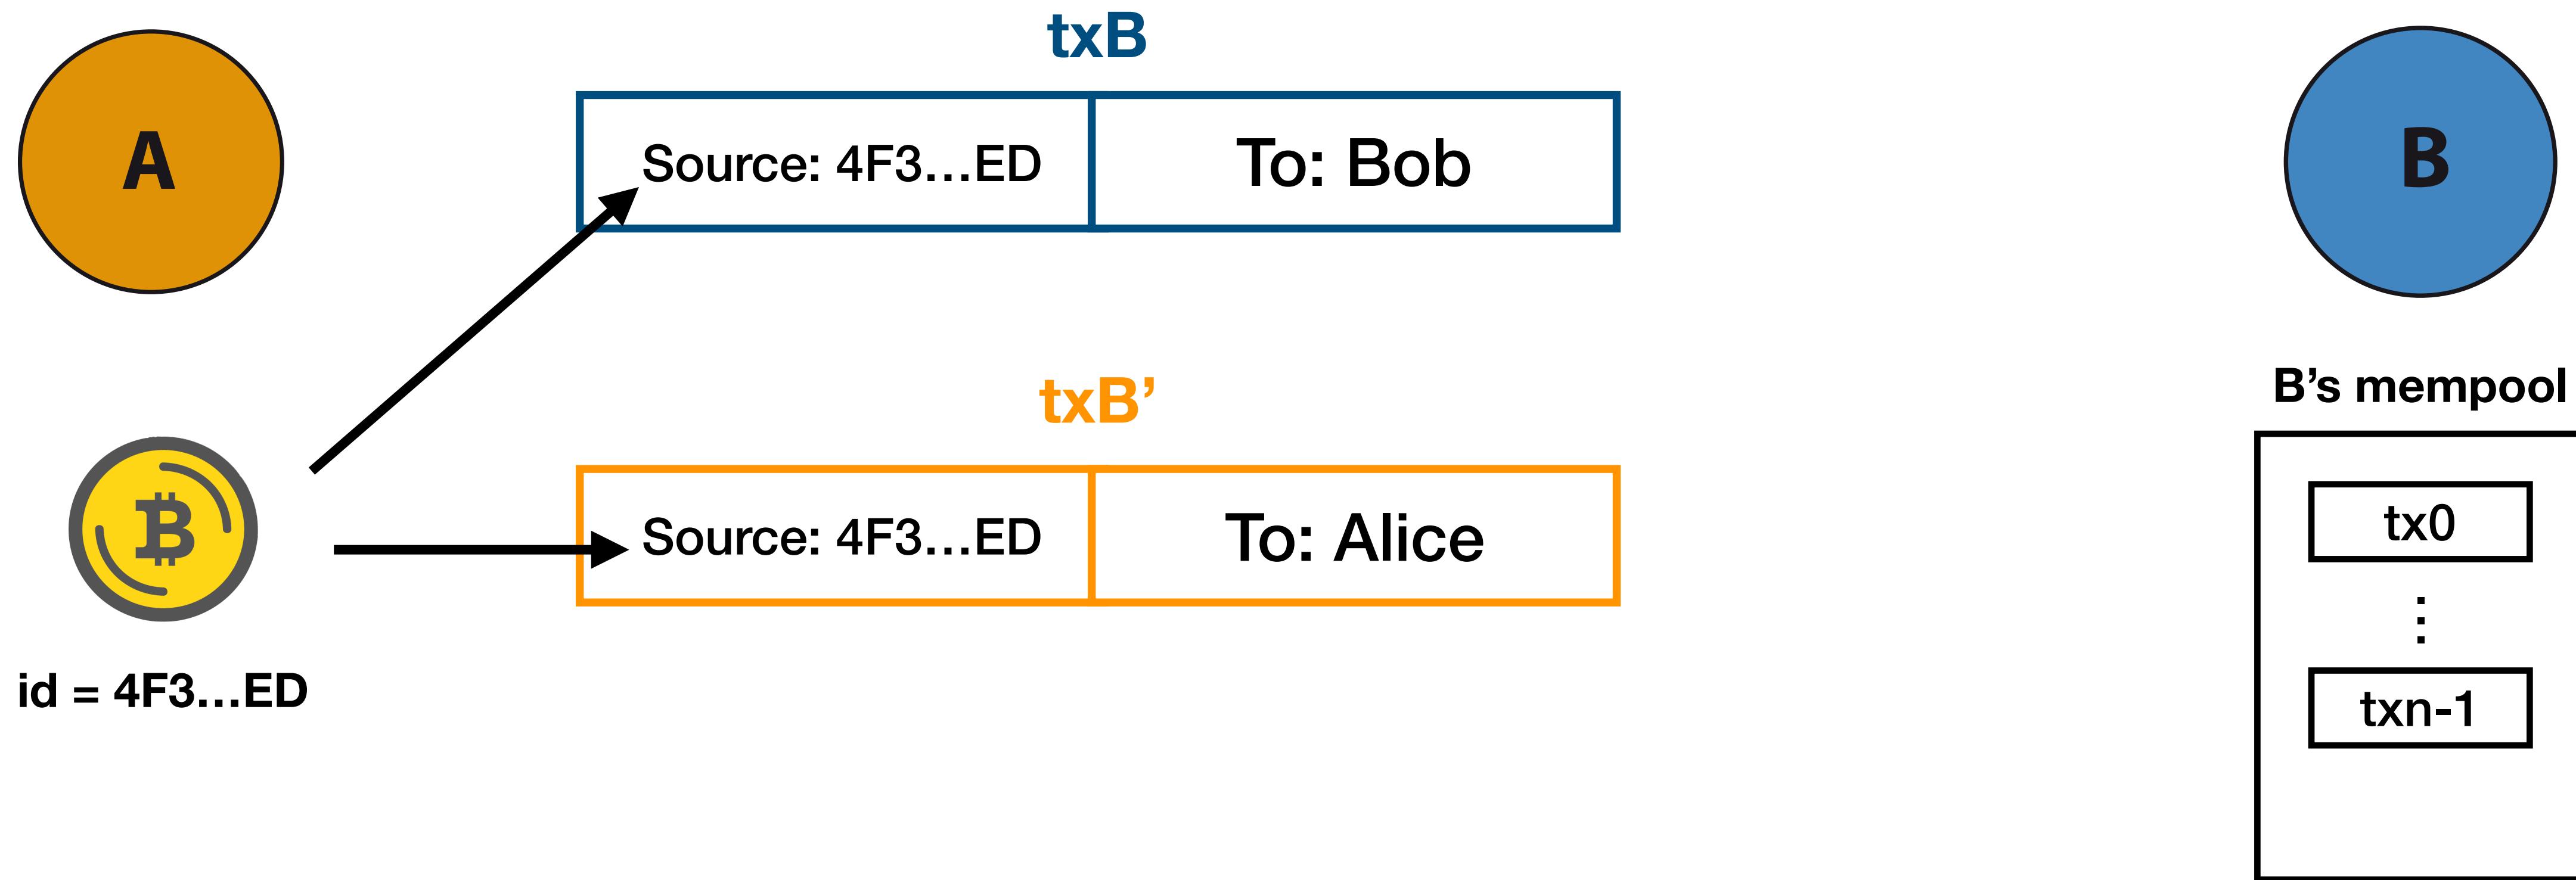


They can not be validated, so they are stored in a separated data structure known as **MapOrphanTransactions (or OrphanPool for short)**

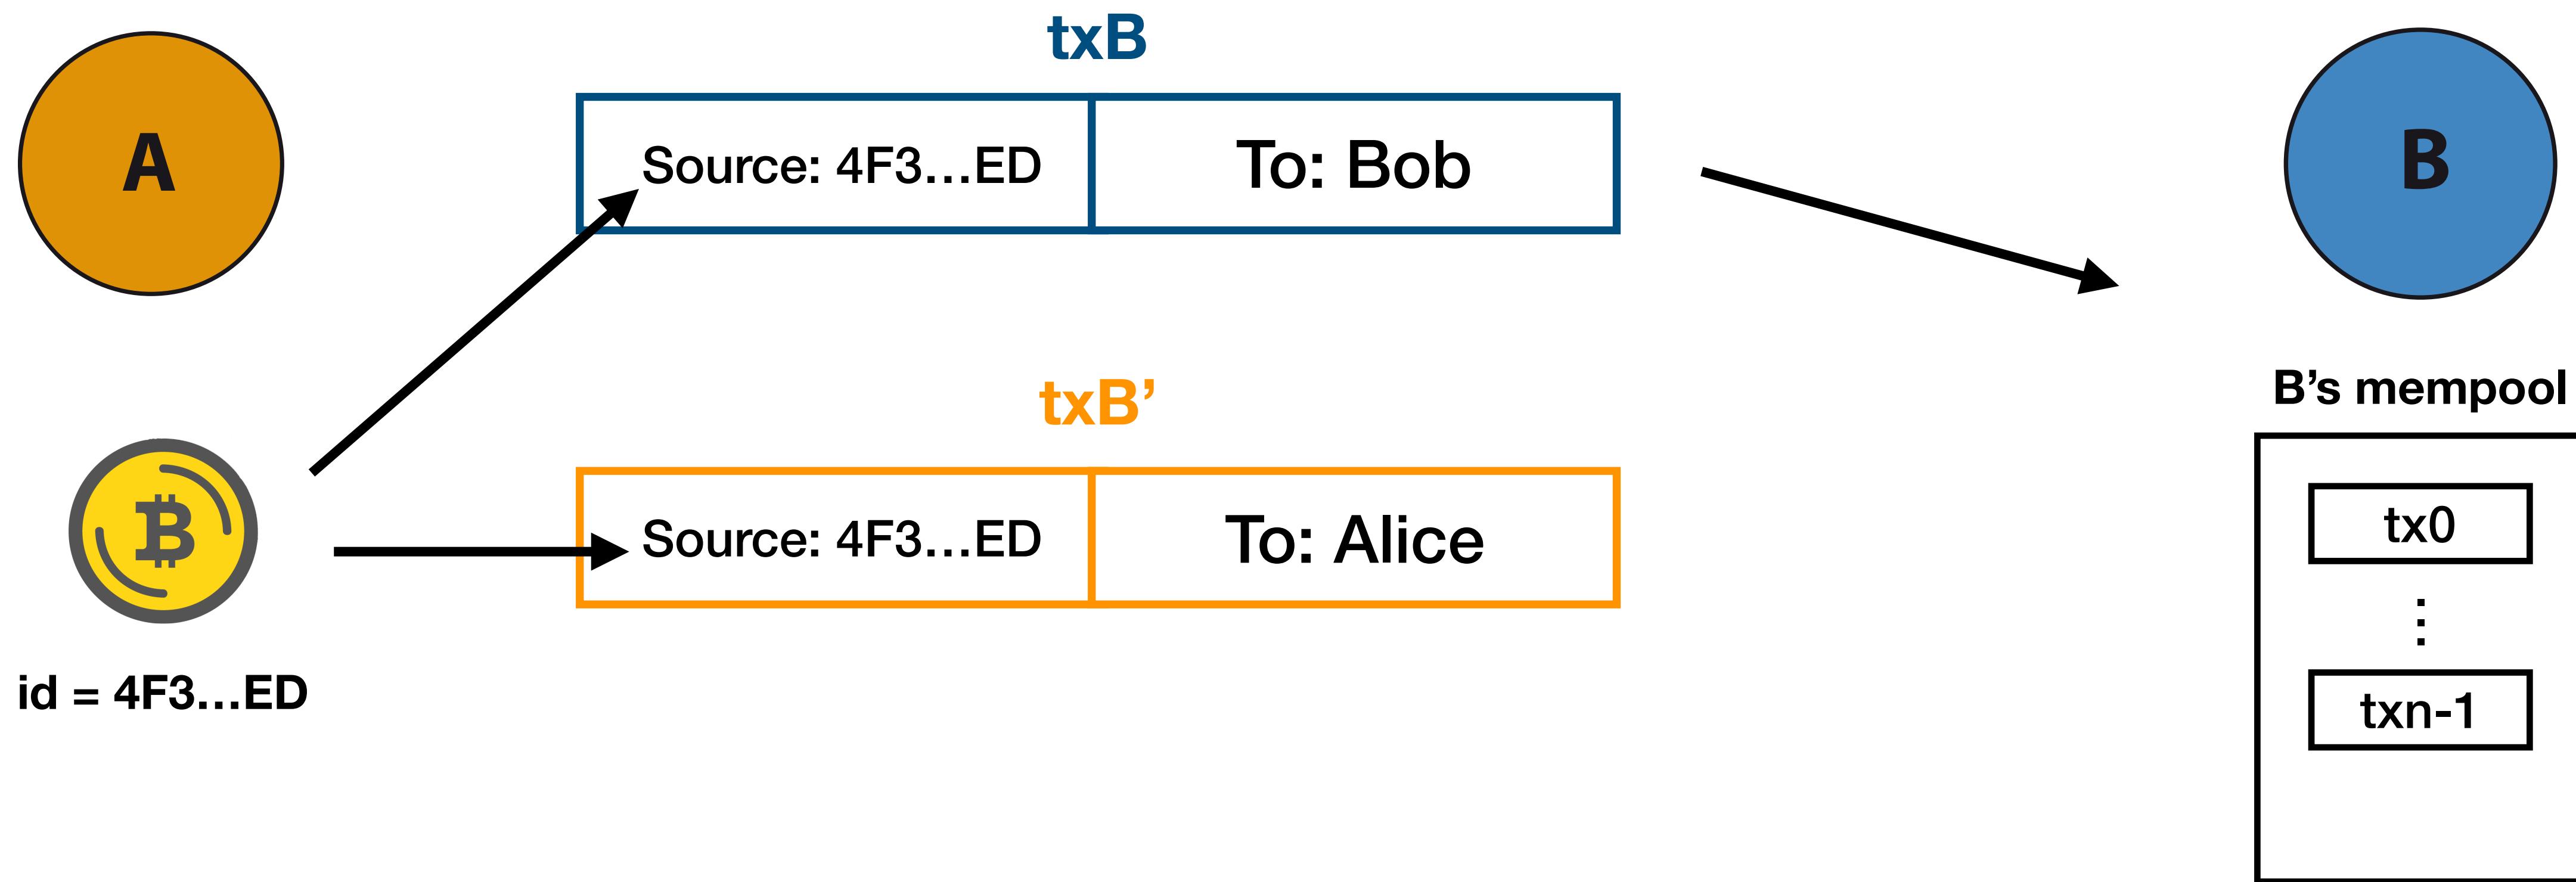
Transactions in MapOrphanTransactions are **NOT forwarded to any node**

If the same transaction is offered again to the node (**inv message**), it will not be requested back (**getaddr**)

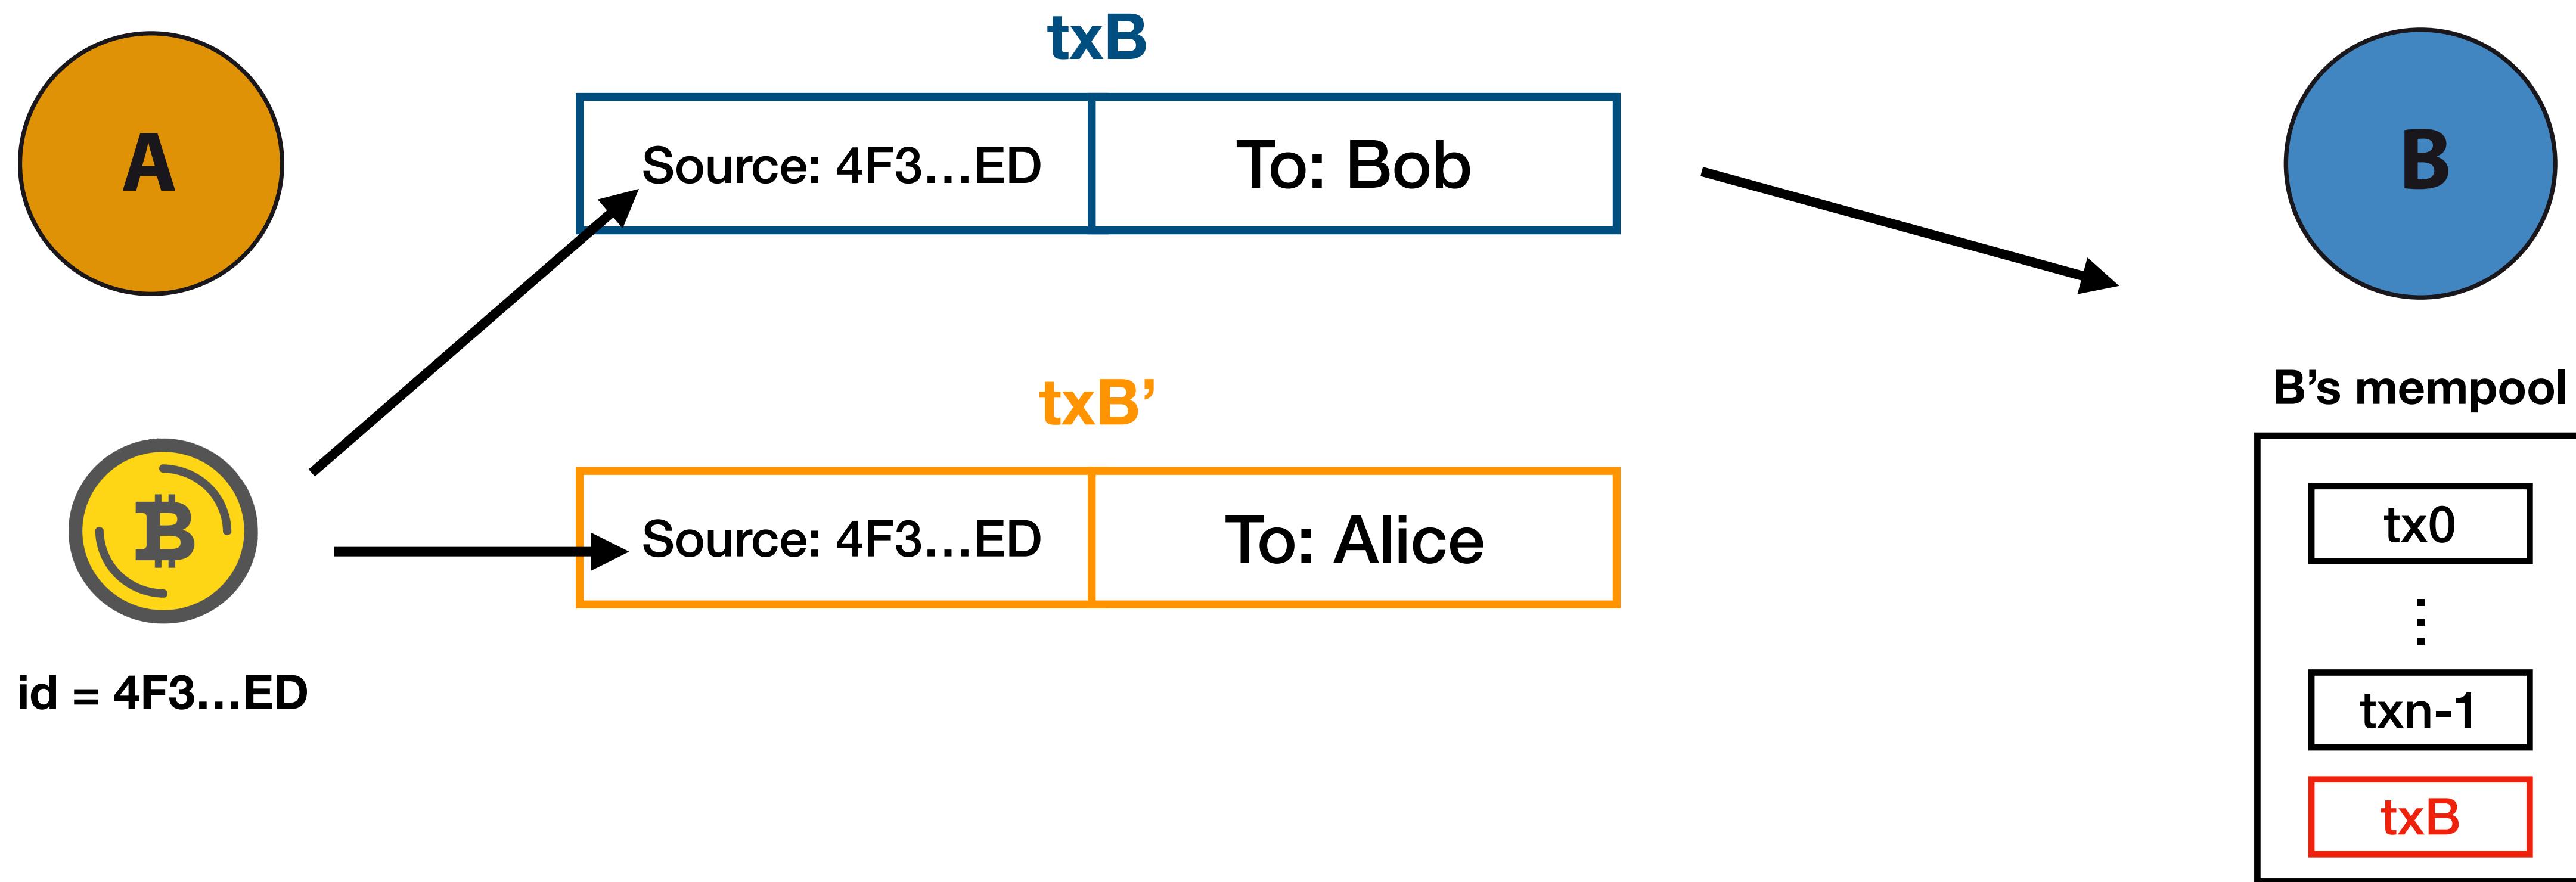
DOUBLE-SPENDING TRANSACTIONS



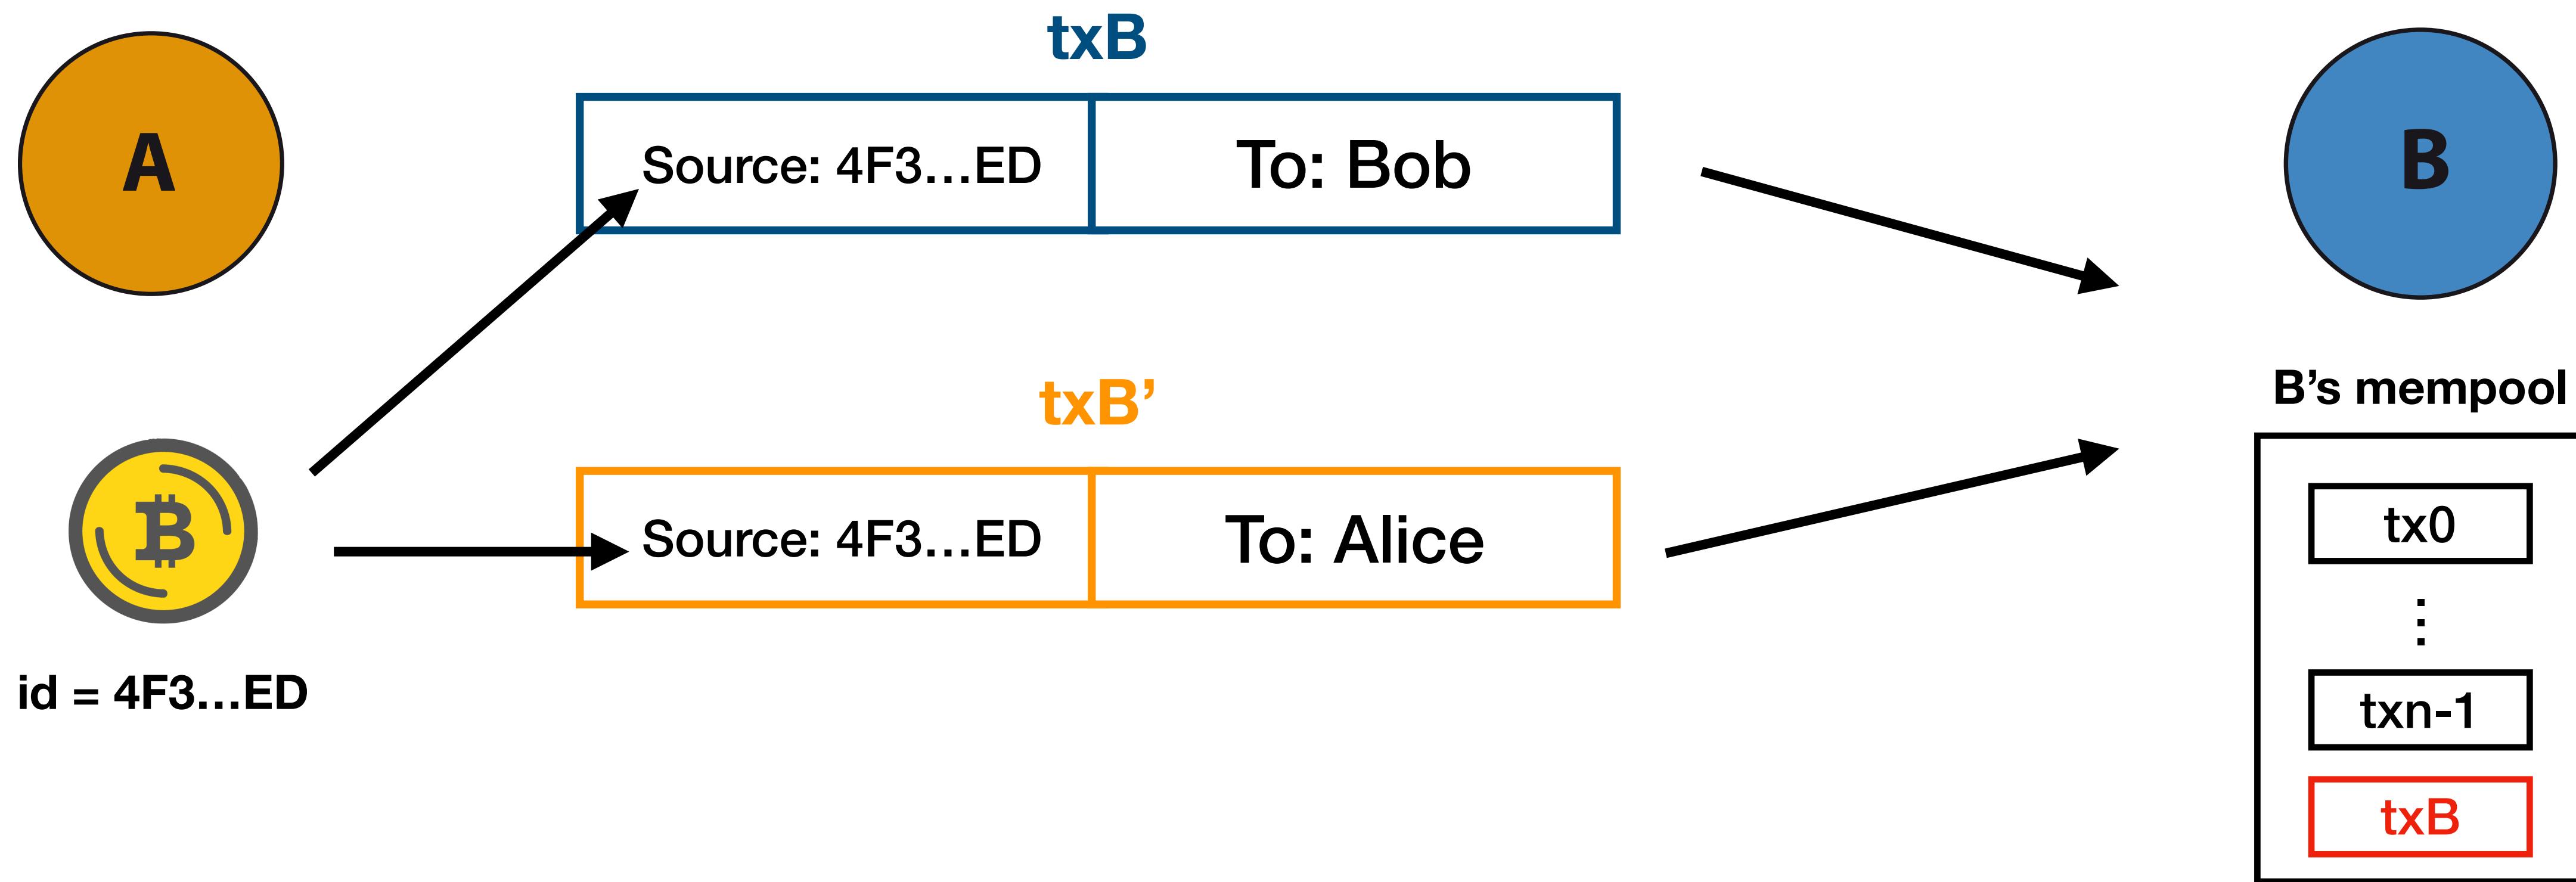
DOUBLE-SPENDING TRANSACTIONS



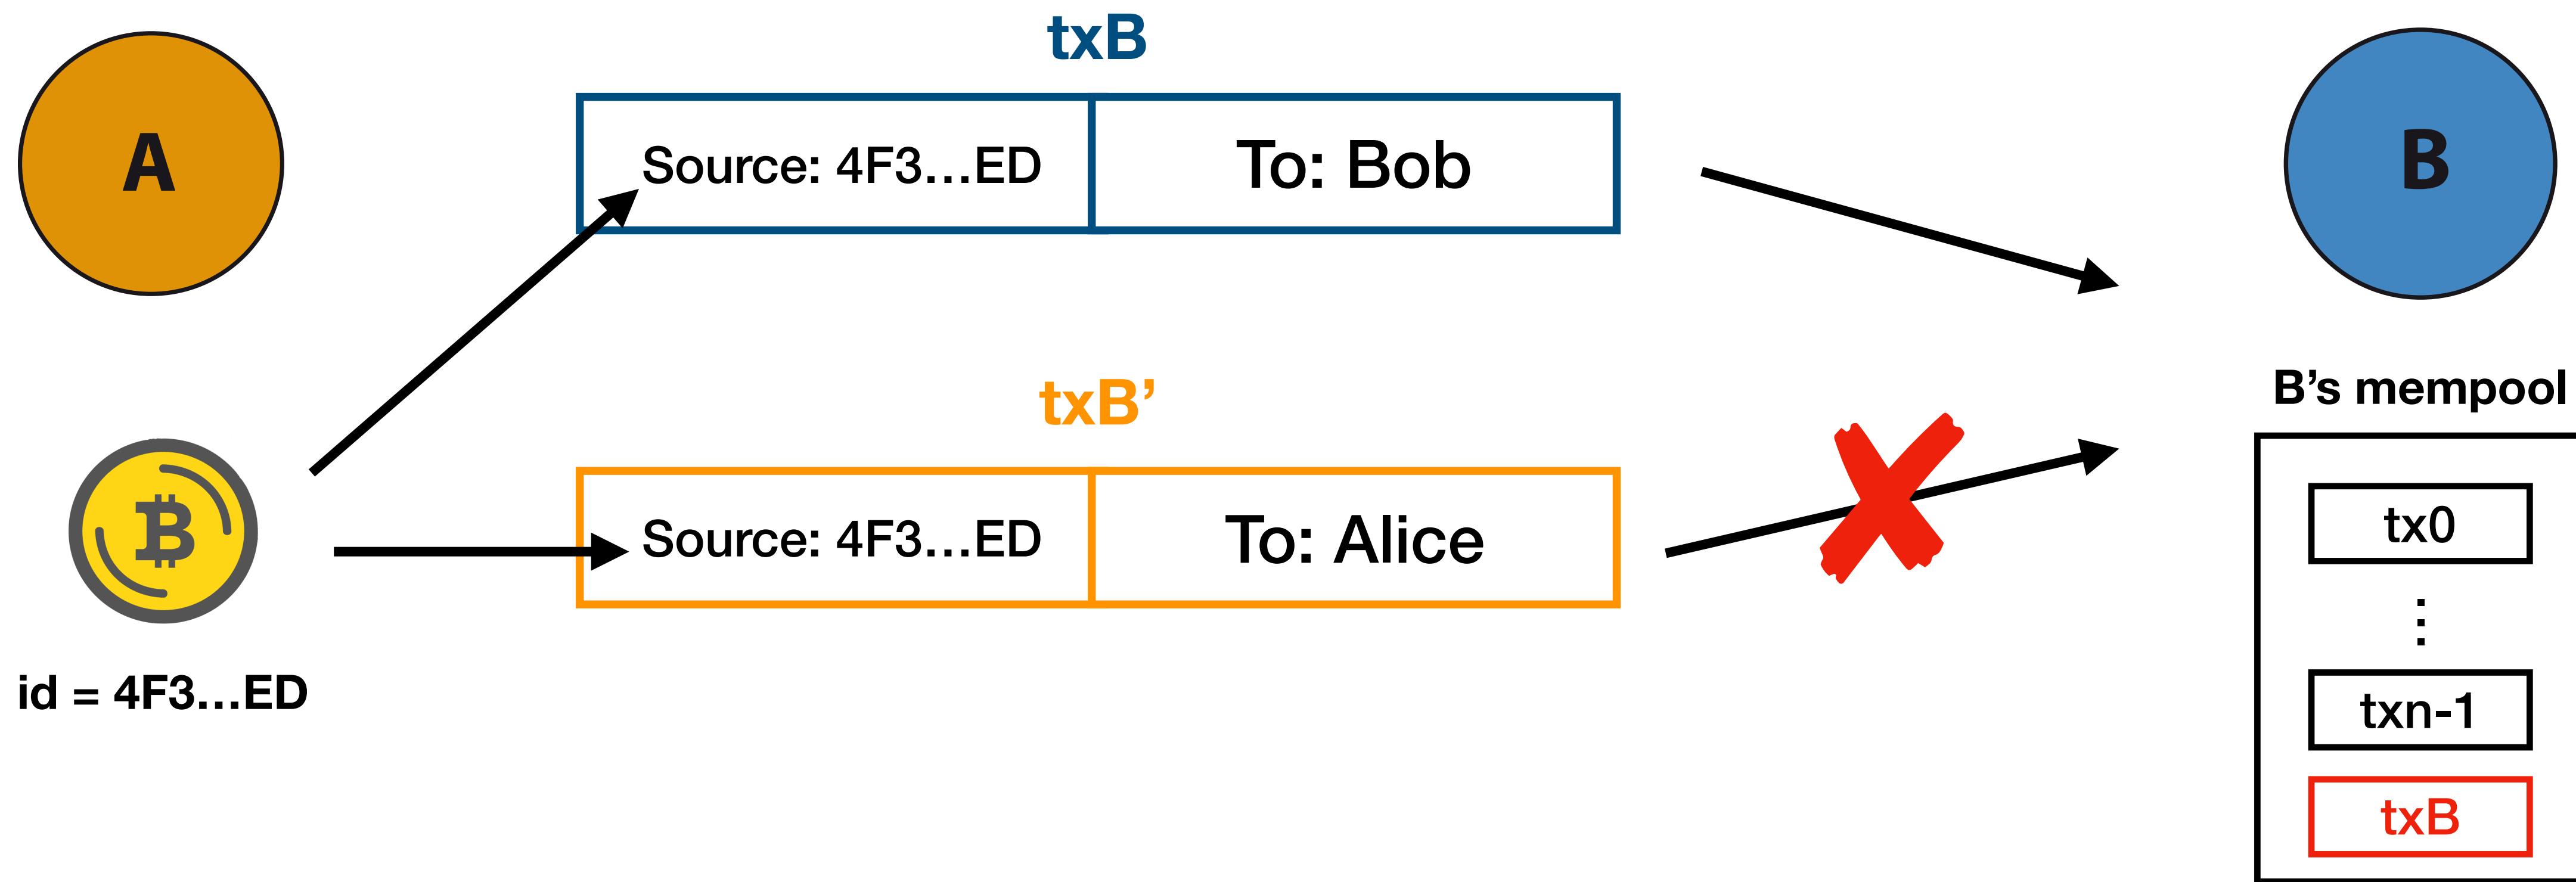
DOUBLE-SPENDING TRANSACTIONS



DOUBLE-SPENDING TRANSACTIONS



DOUBLE-SPENDING TRANSACTIONS



A BASIC TOPOLOGY INFERRING TECHNIQUE

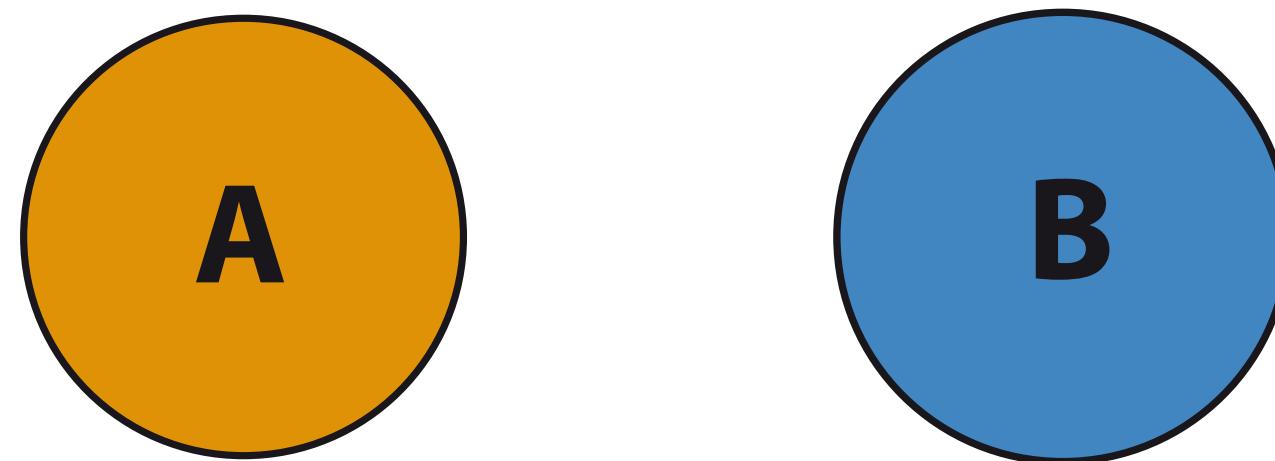
Two nodes

Three transactions

**Observation tool
(like coinscope)**

A BASIC TOPOLOGY INFERRING TECHNIQUE

Two nodes

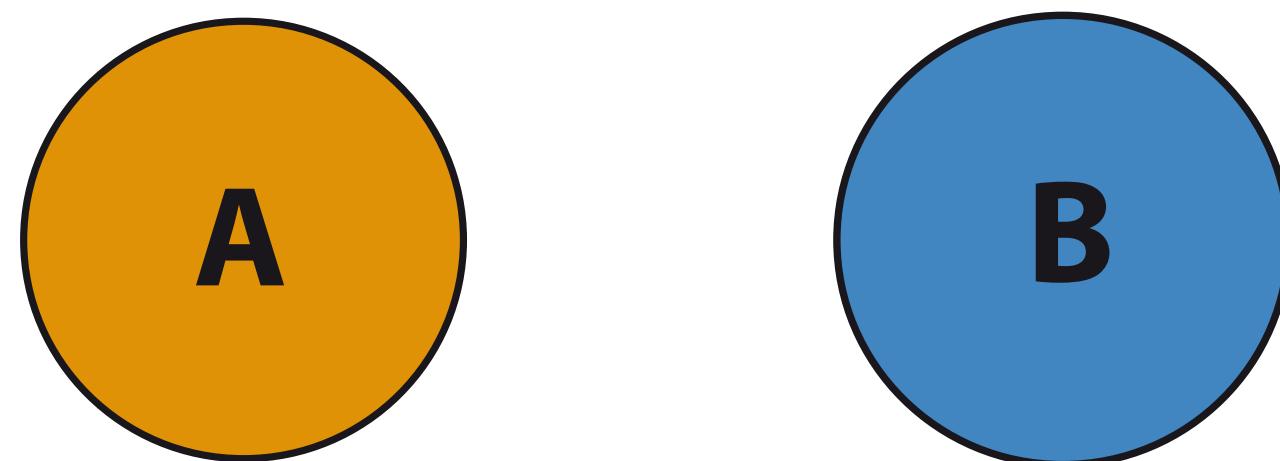


Three transactions

**Observation tool
(like coinscope)**

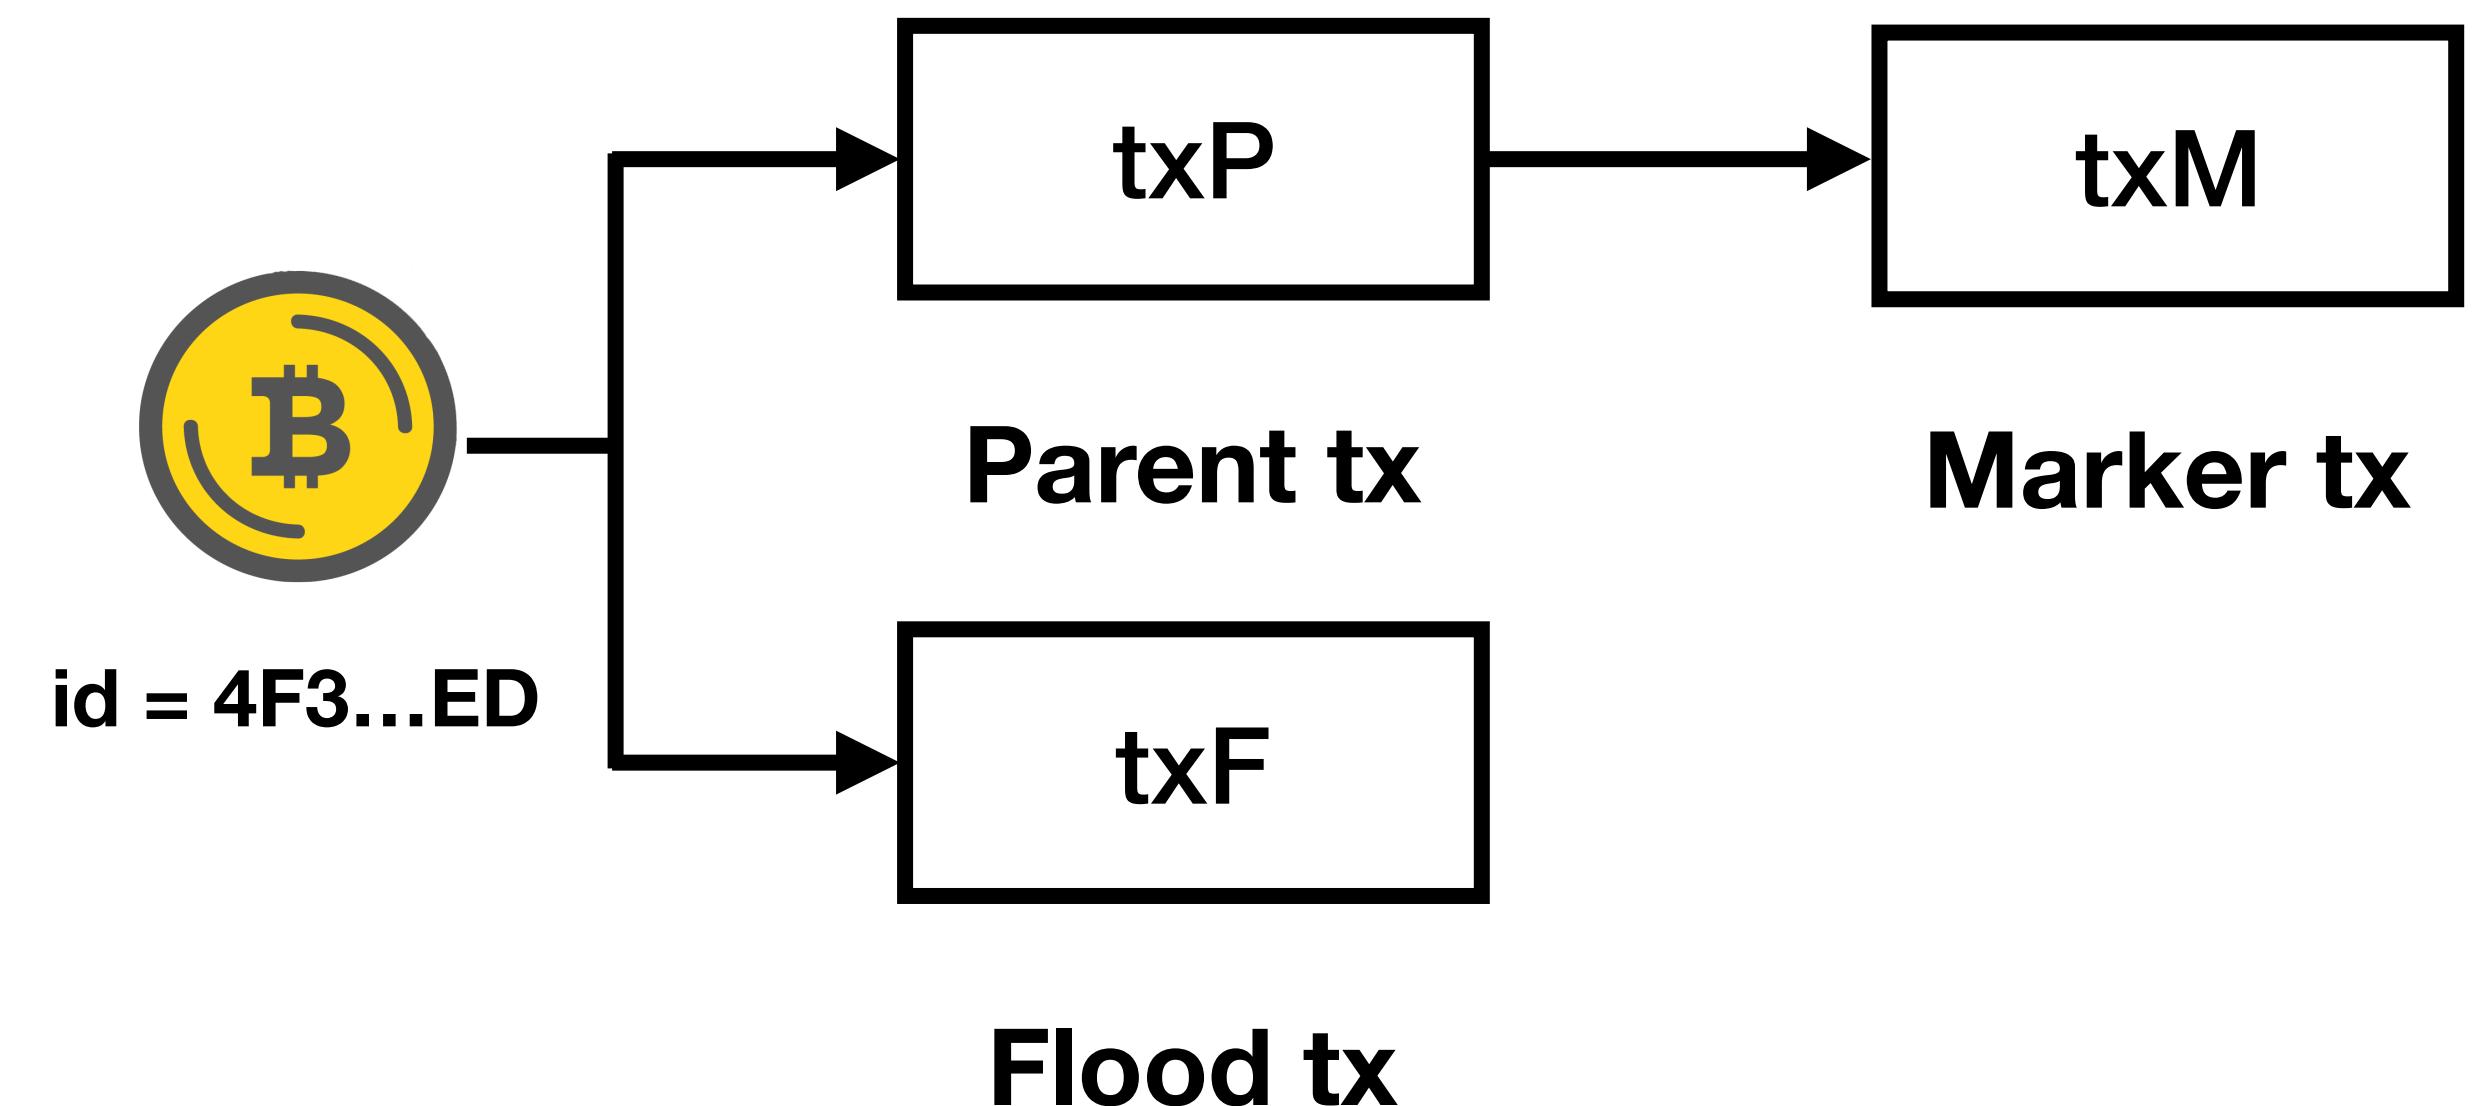
A BASIC TOPOLOGY INFERRING TECHNIQUE

Two nodes



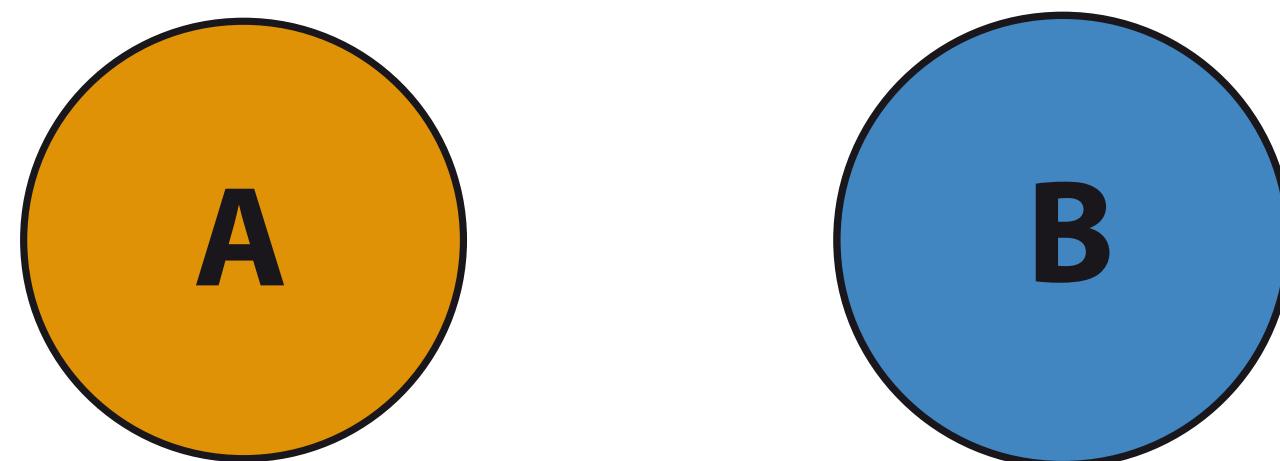
**Observation tool
(like coinscope)**

Three transactions

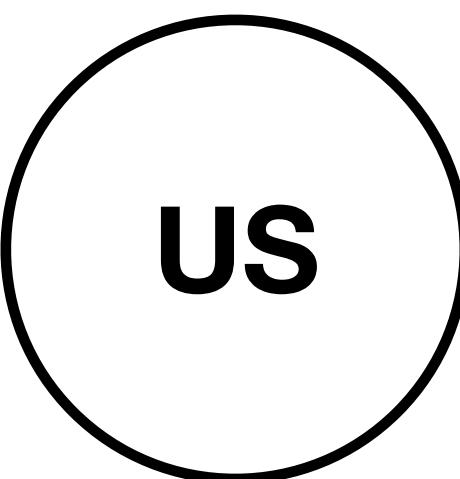


A BASIC TOPOLOGY INFERRING TECHNIQUE

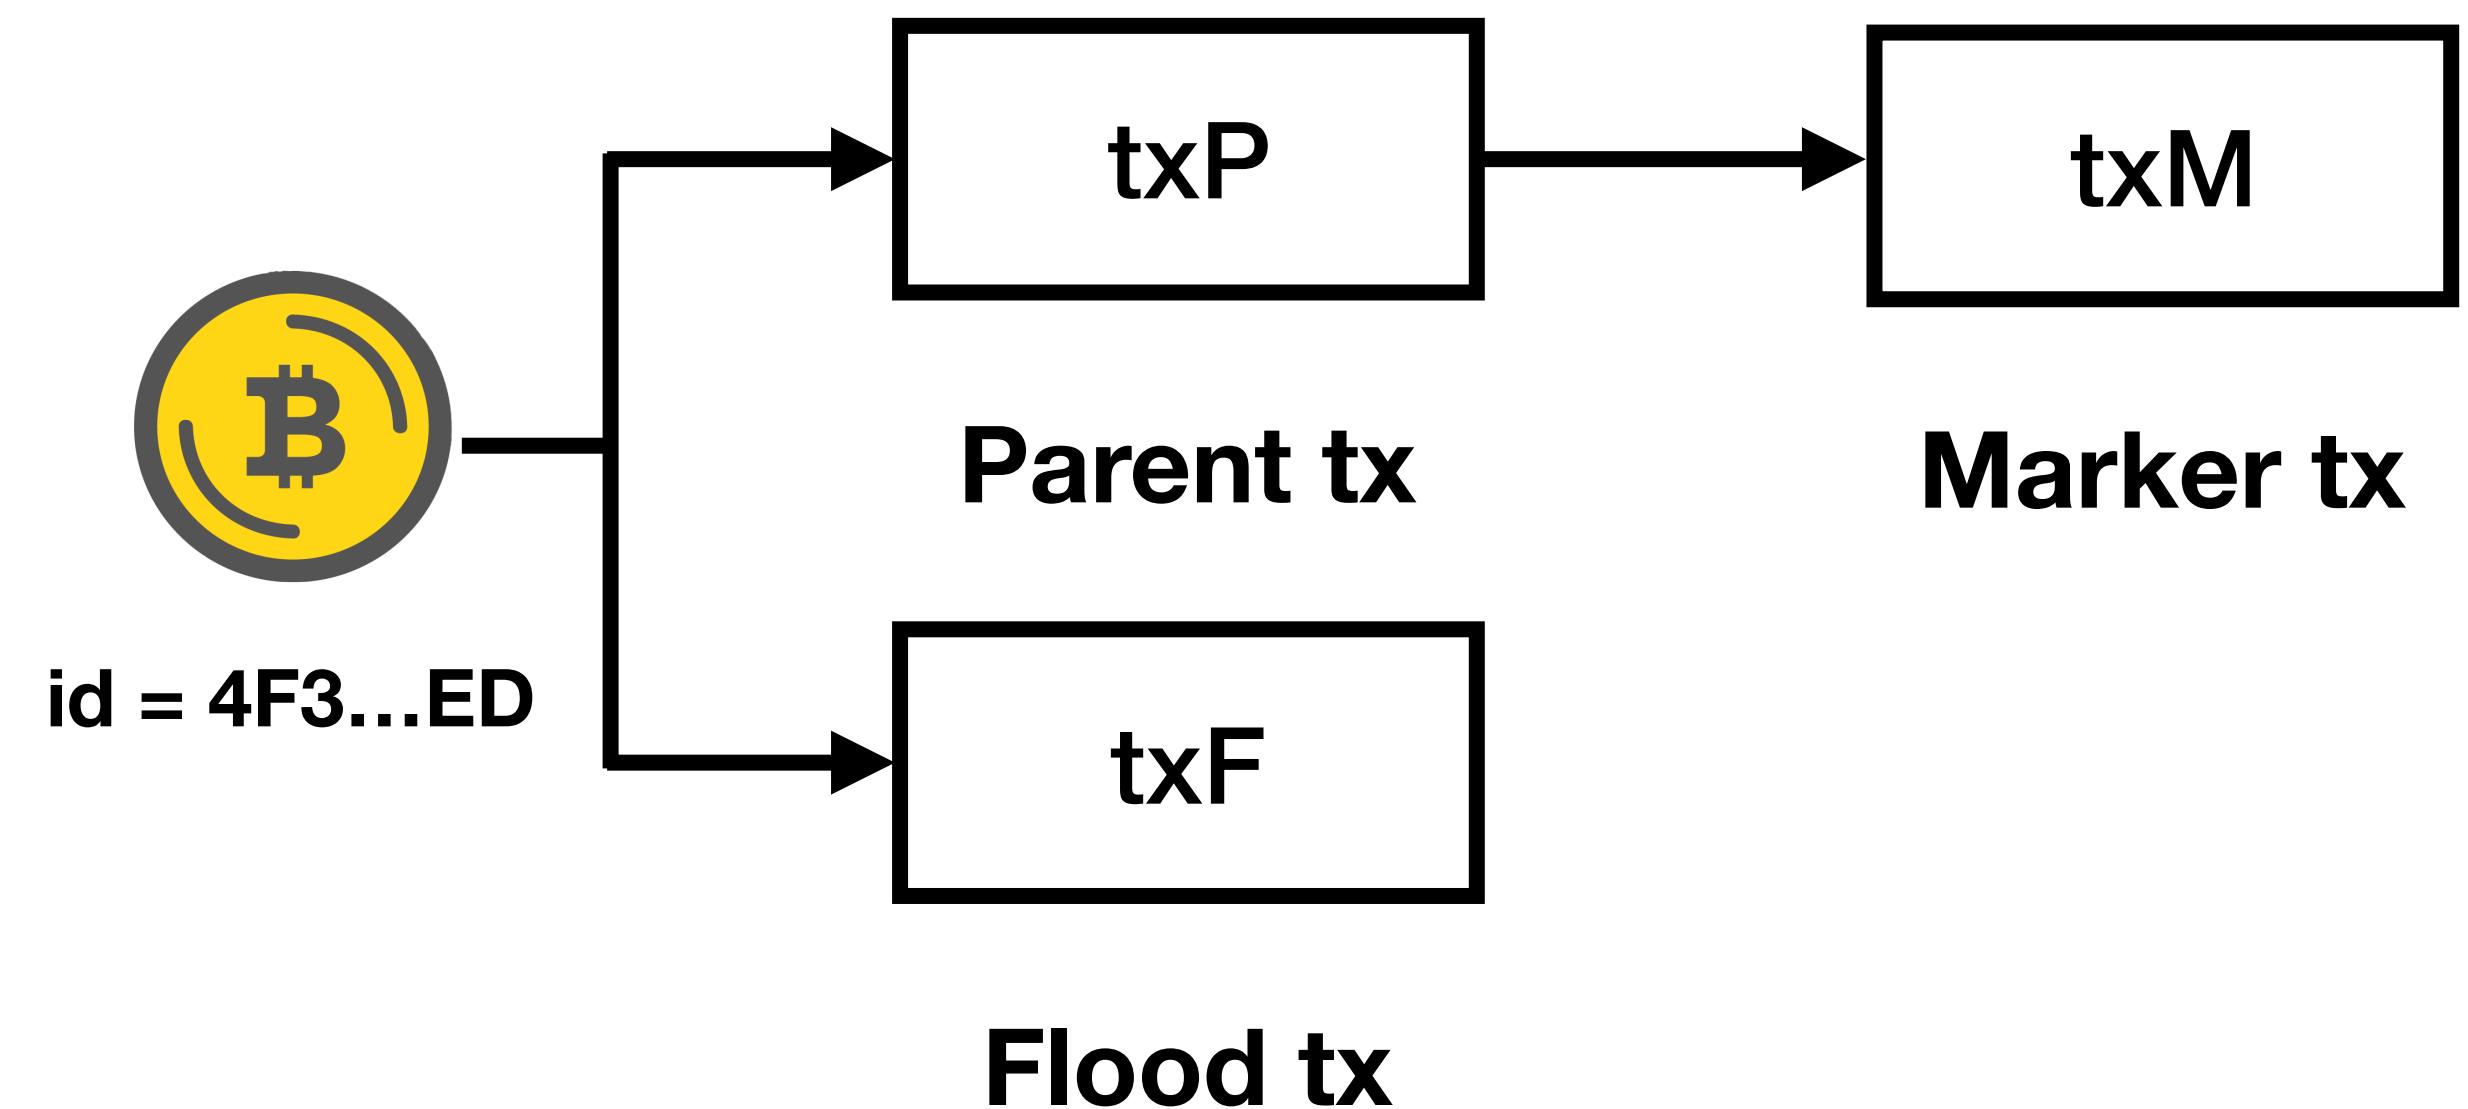
Two nodes



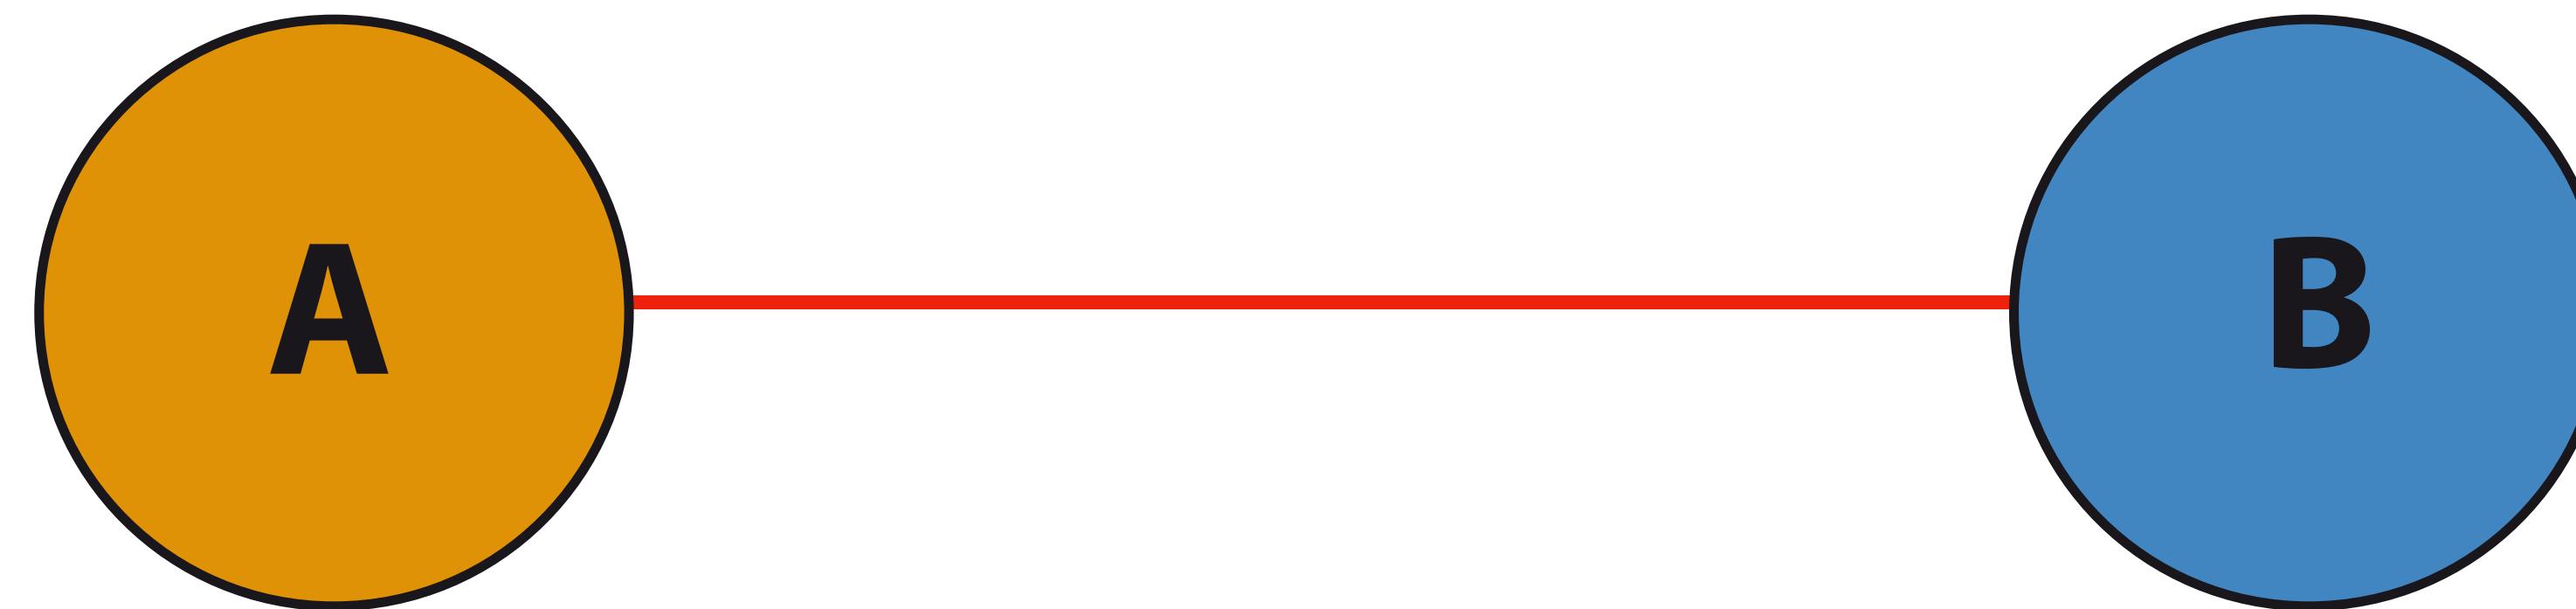
**Observation tool
(like coinscope)**



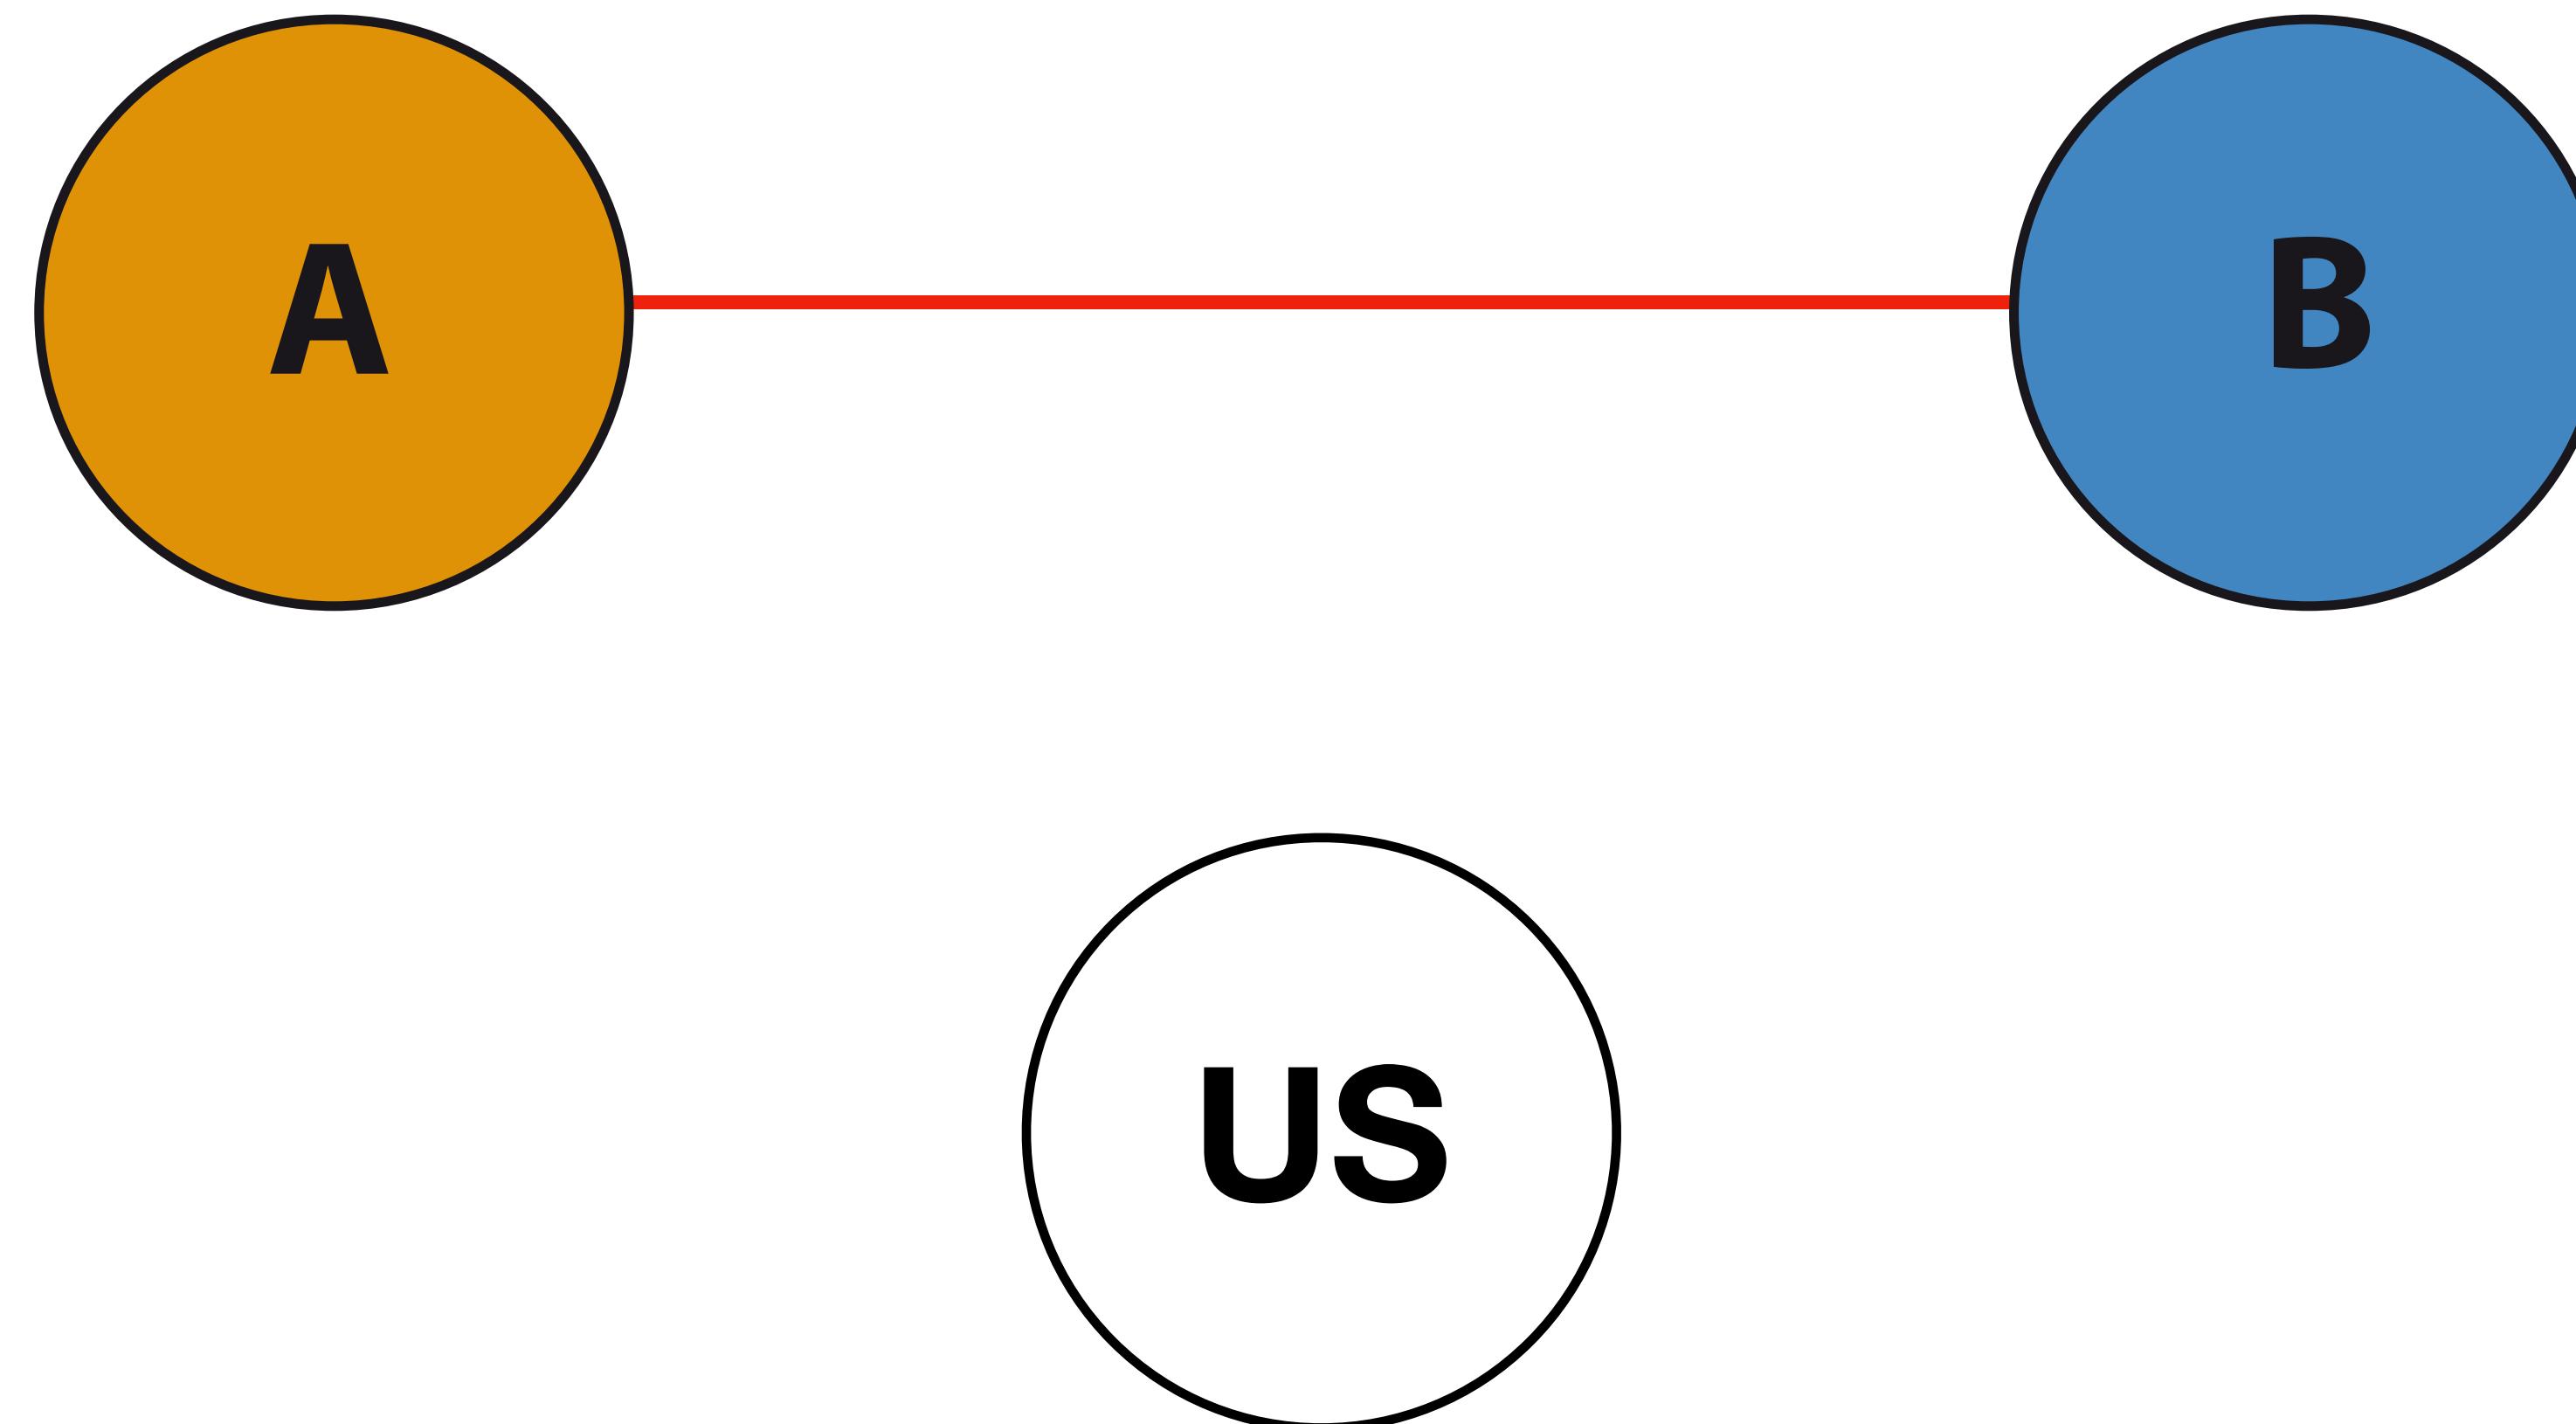
Three transactions



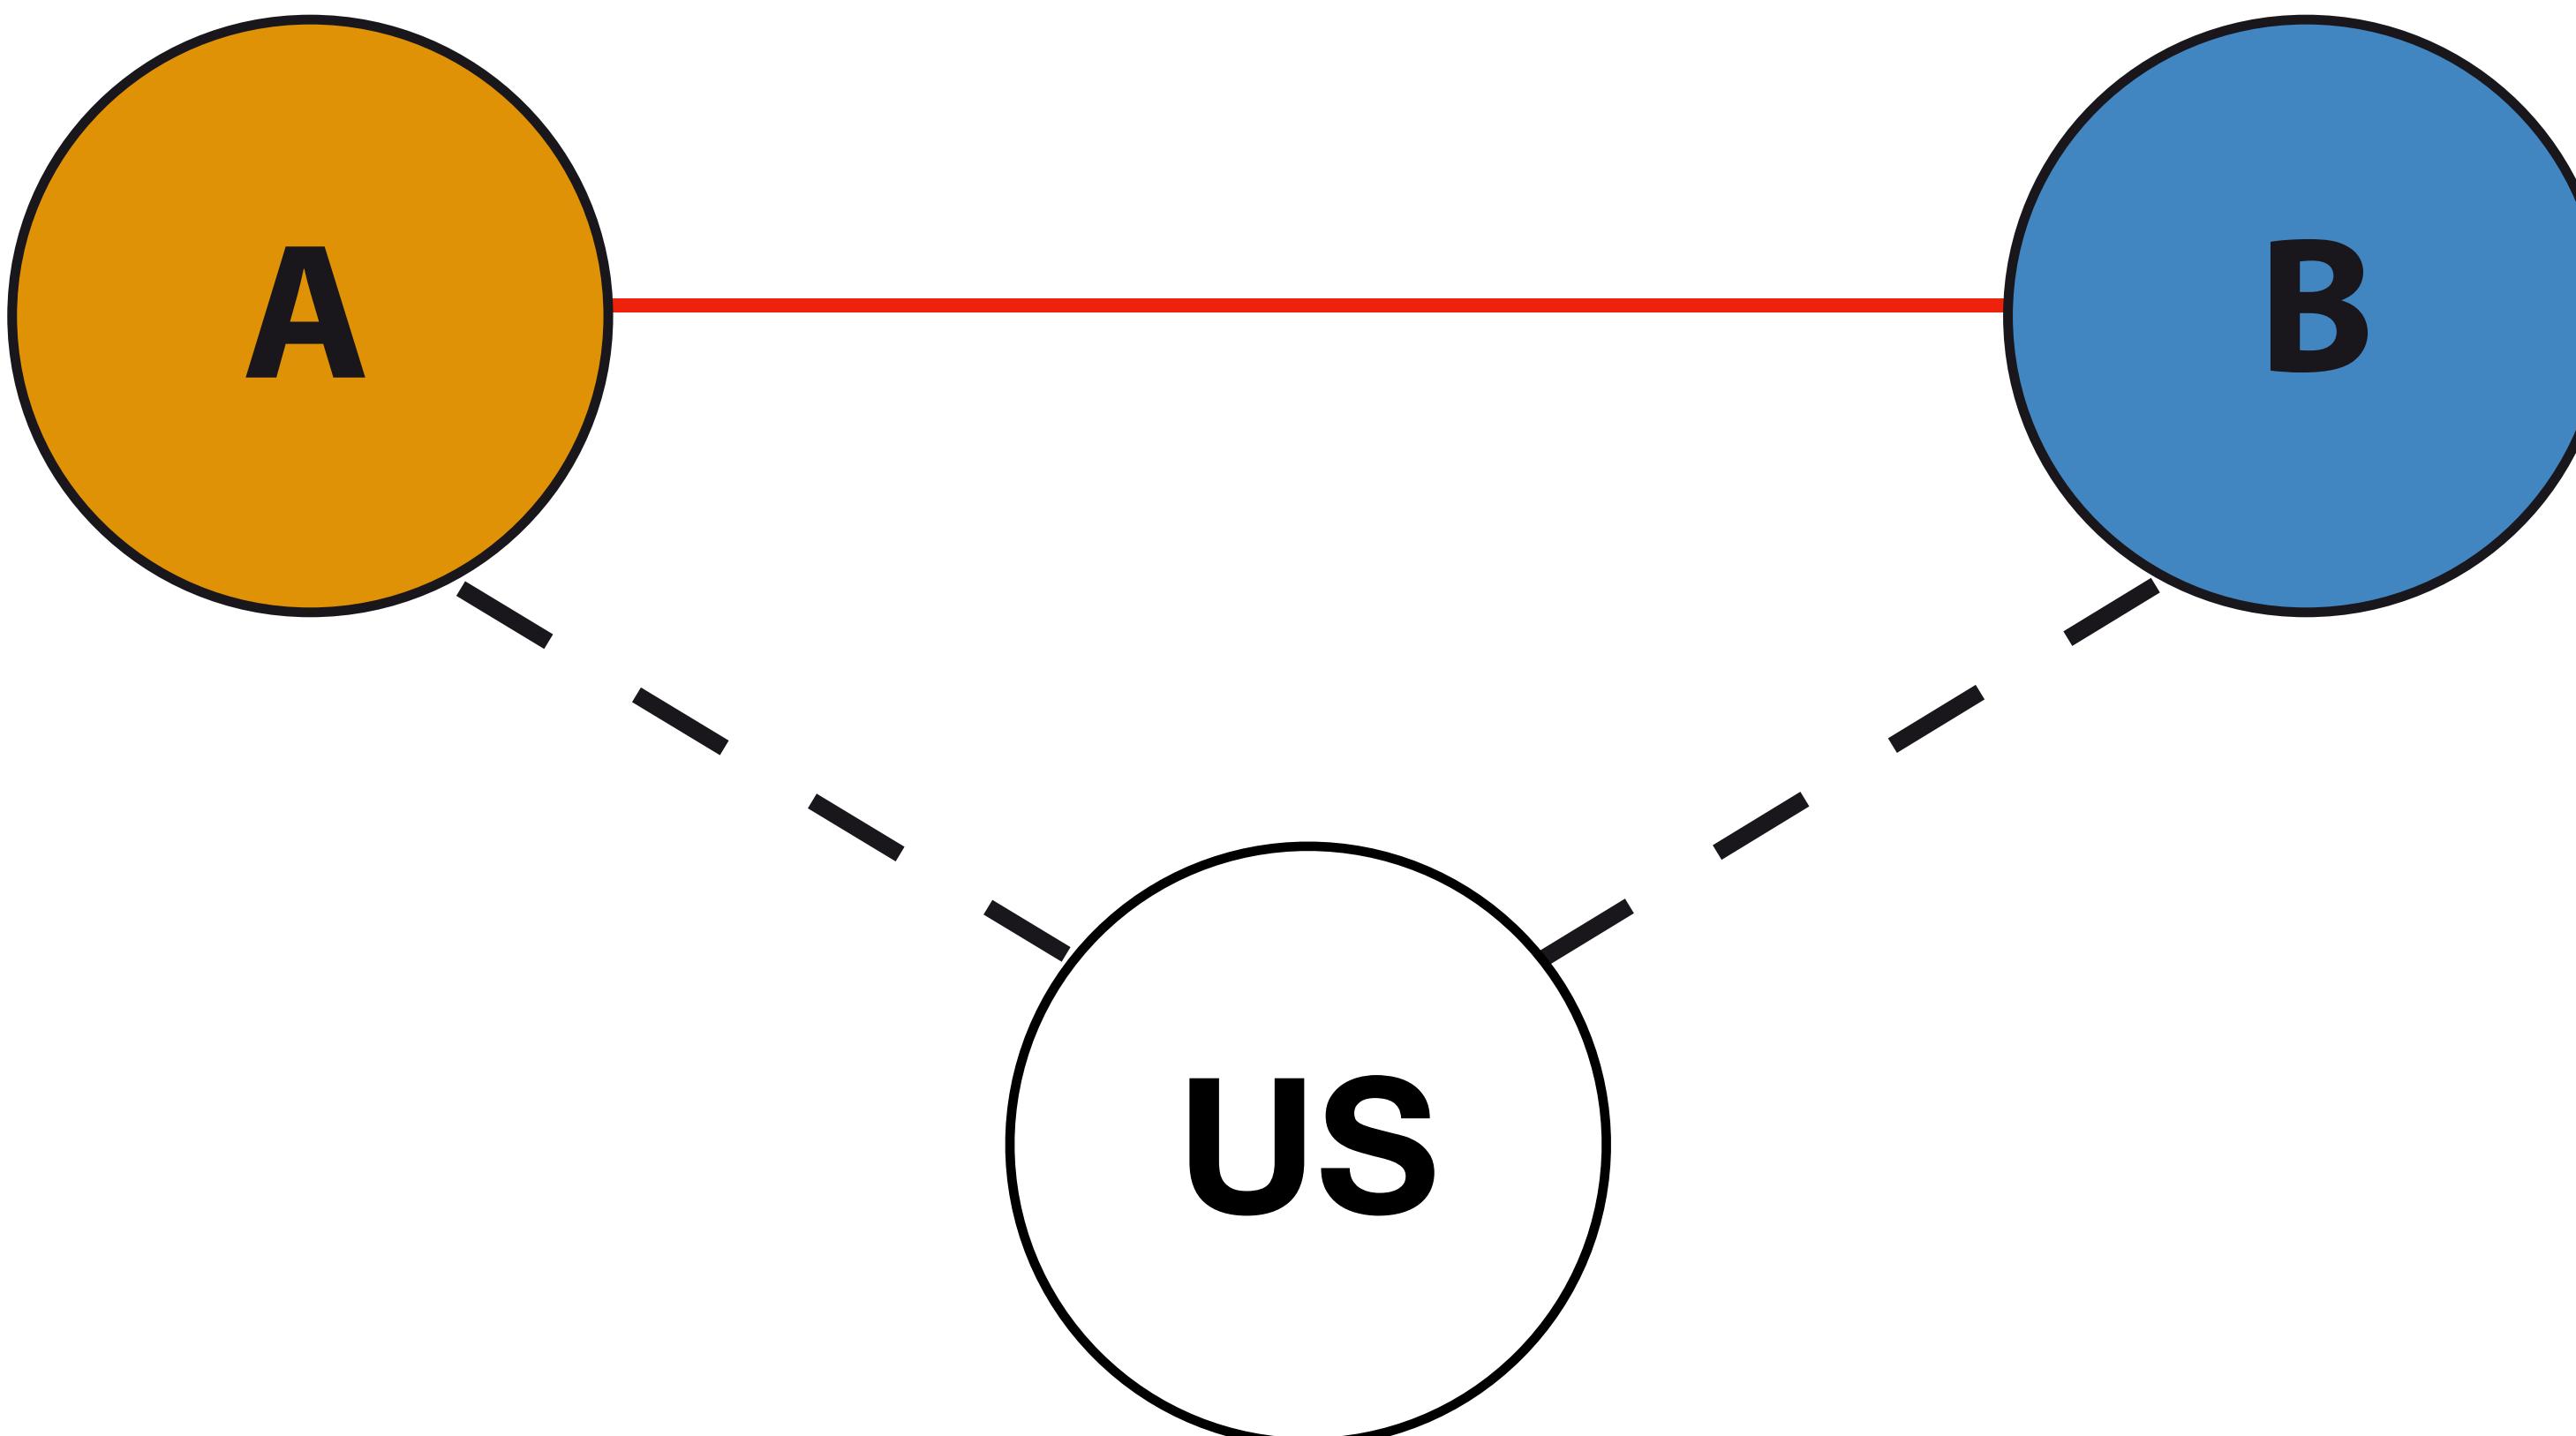
POSITIVE INFERRING TECHNIQUE



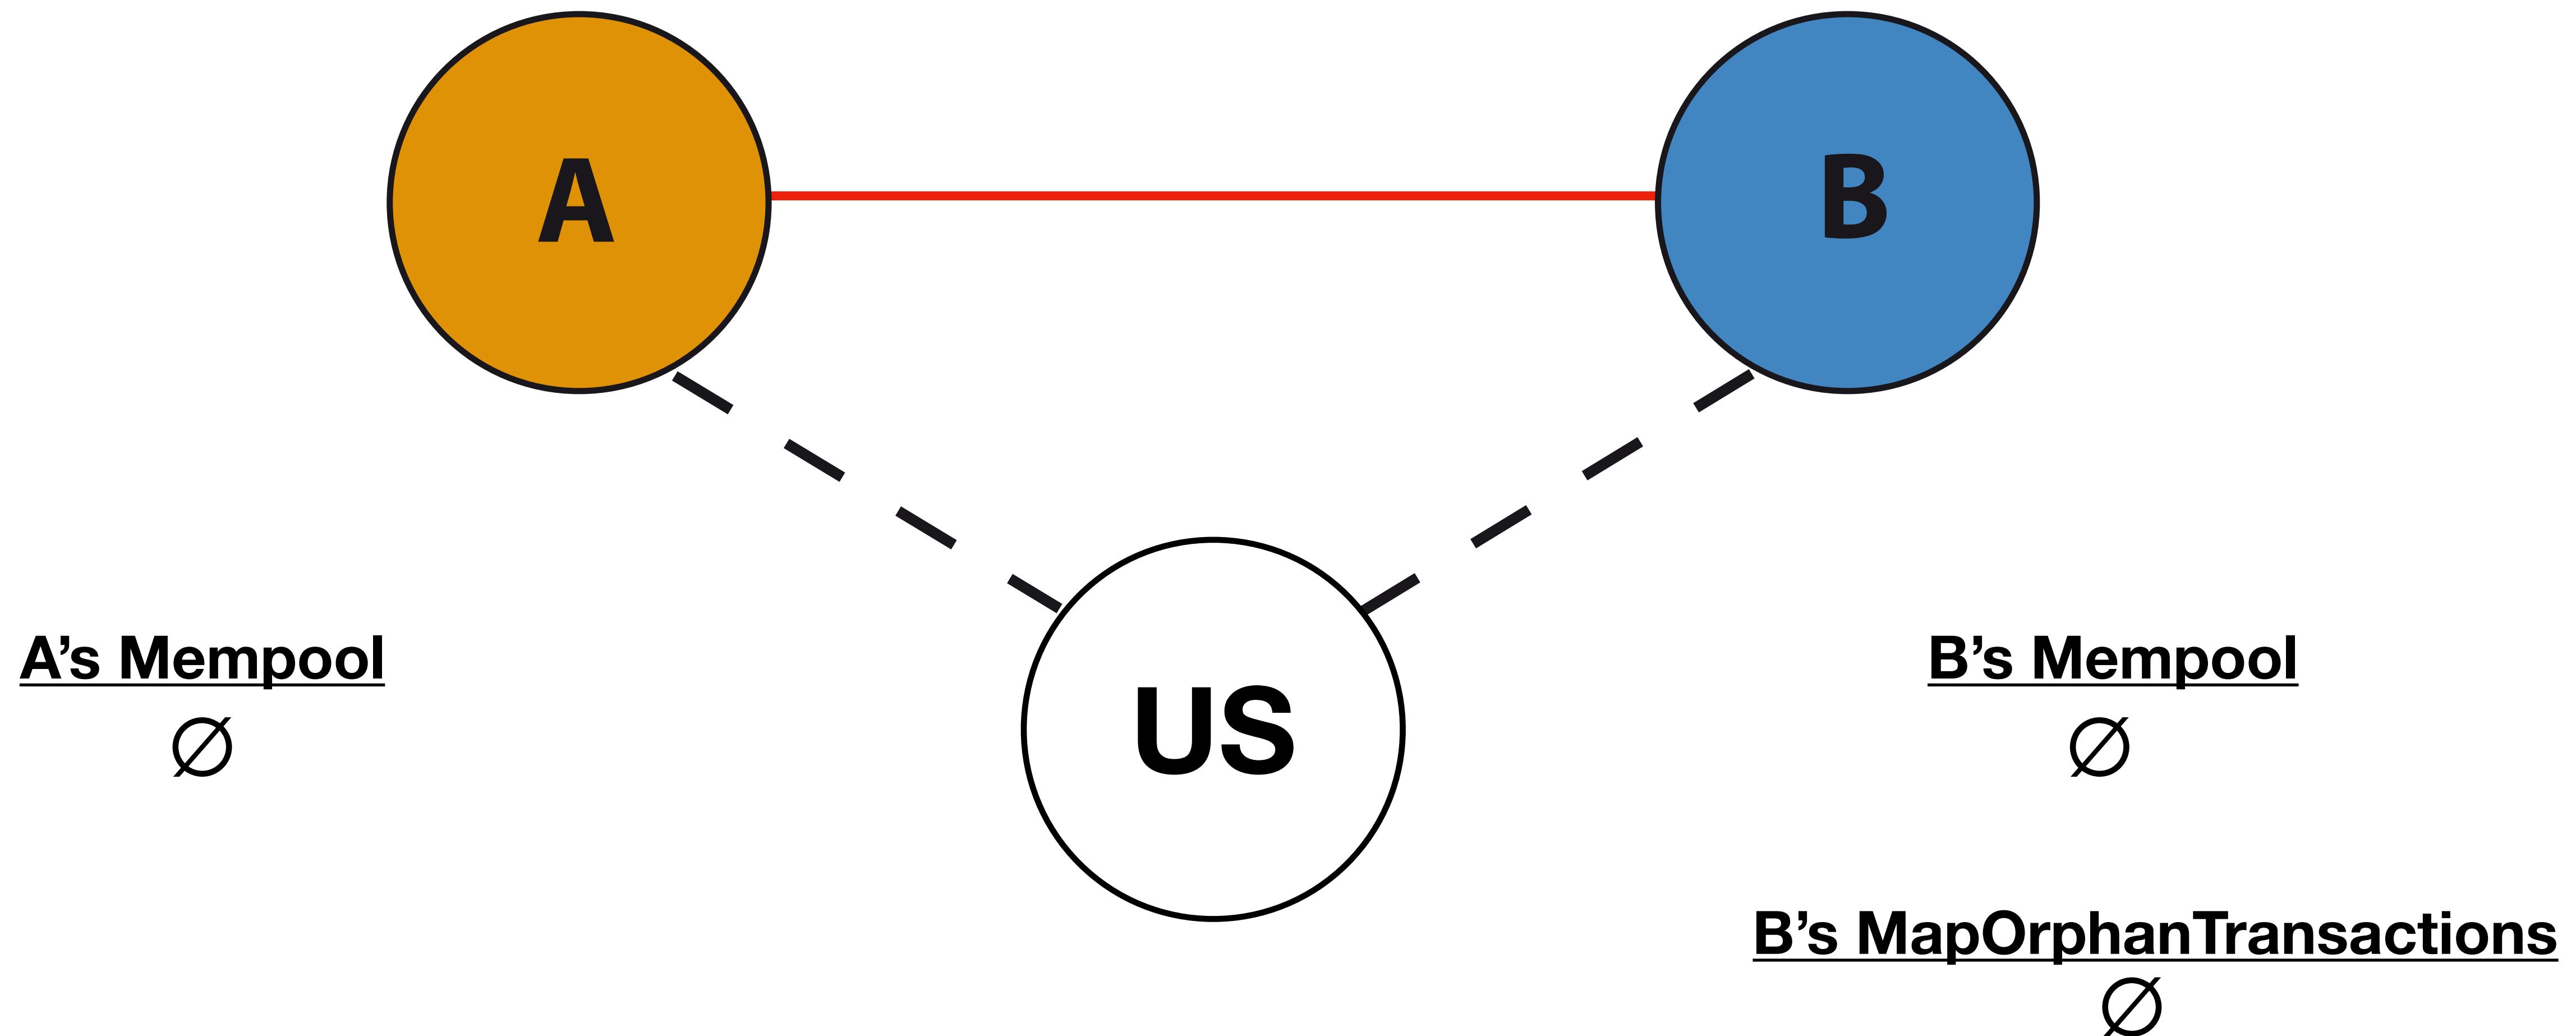
POSITIVE INFERRING TECHNIQUE



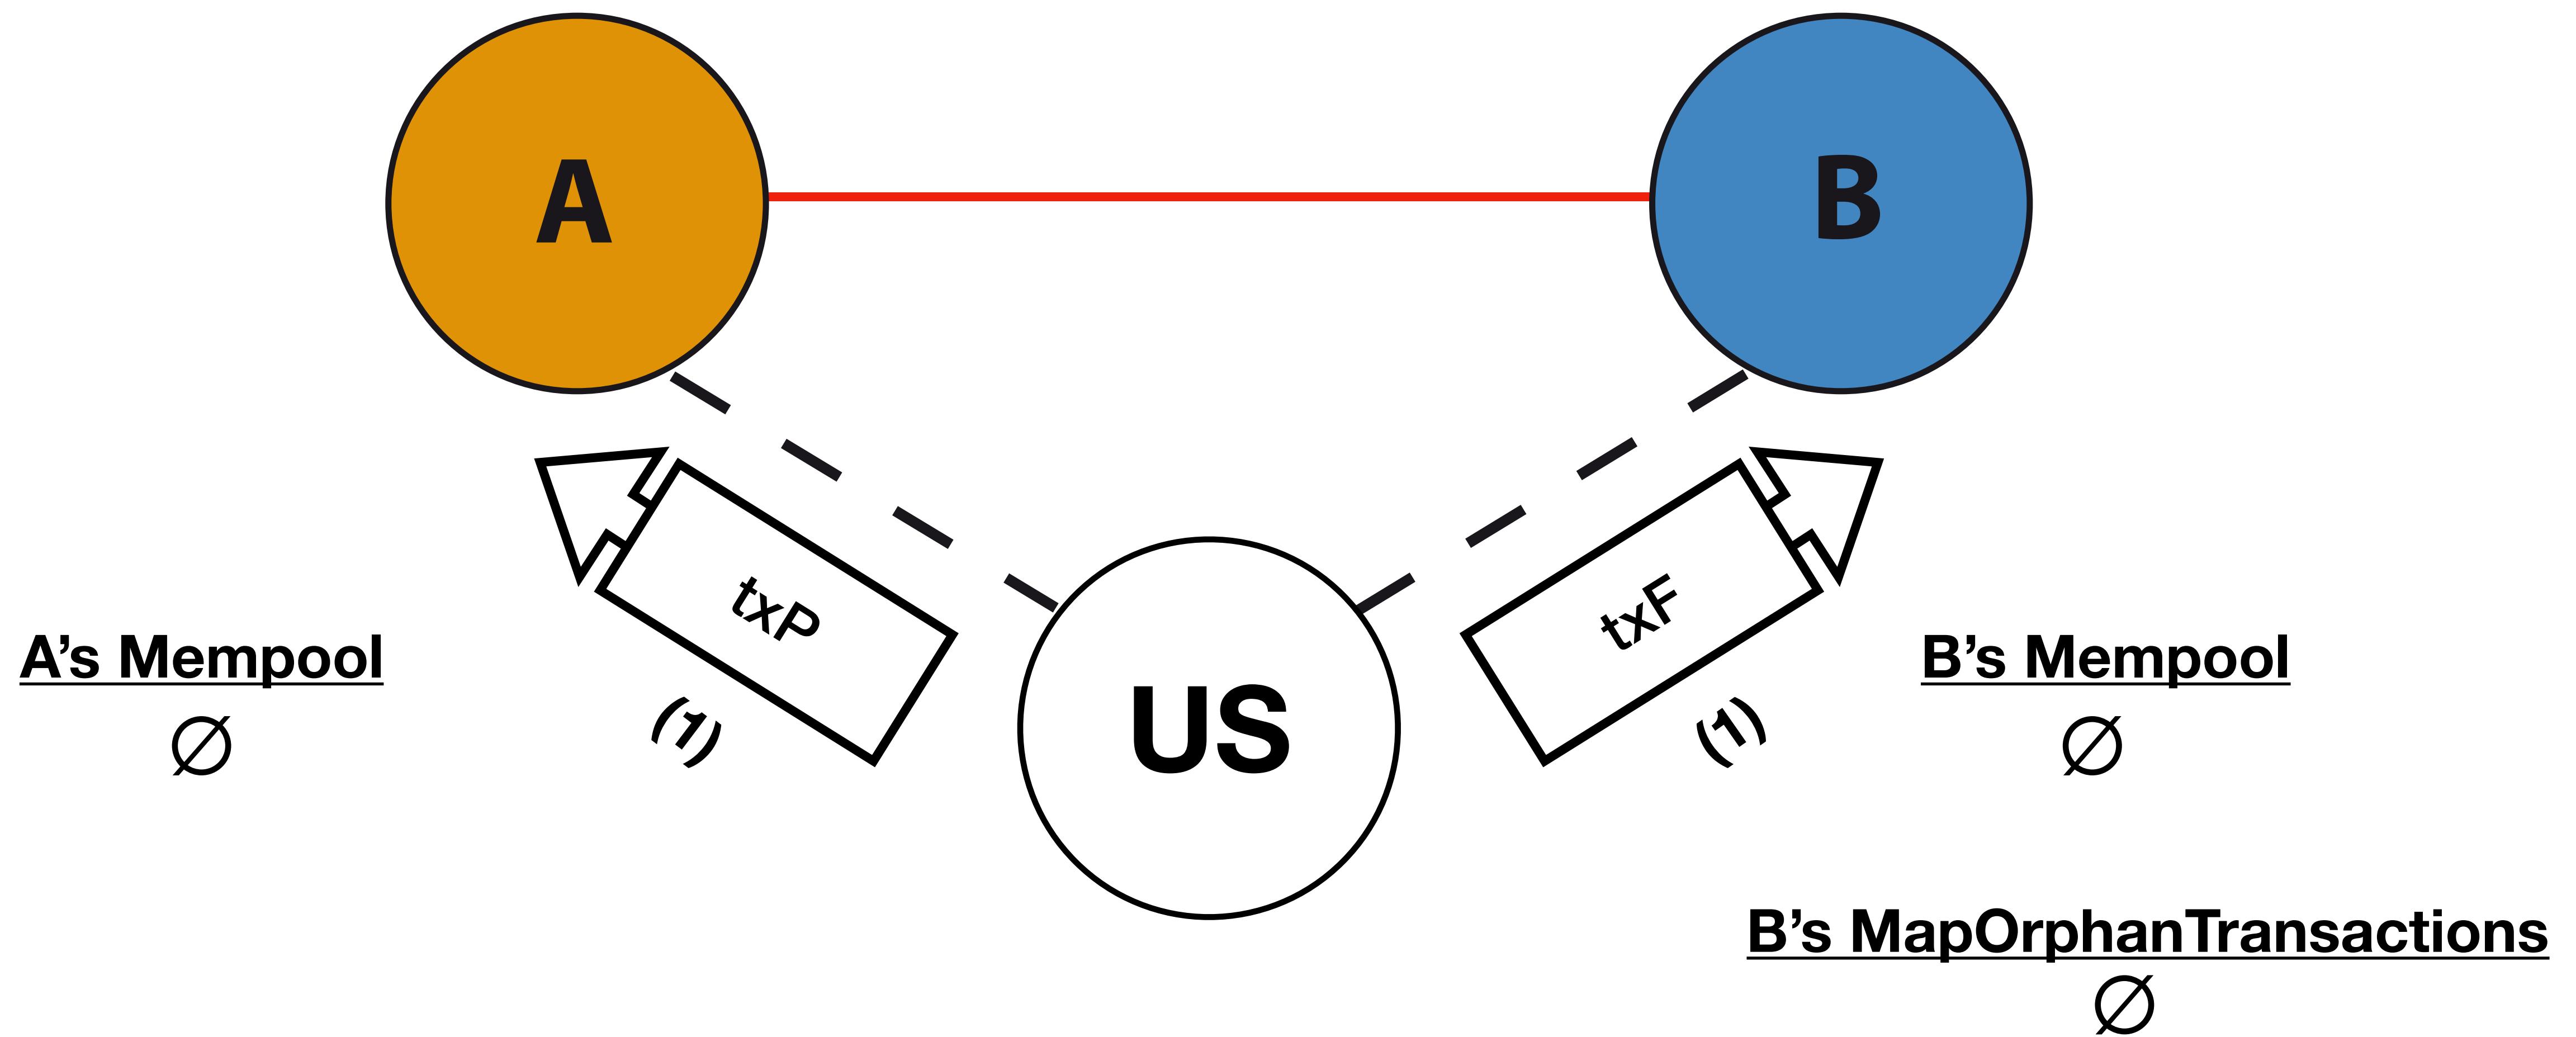
POSITIVE INFERRING TECHNIQUE



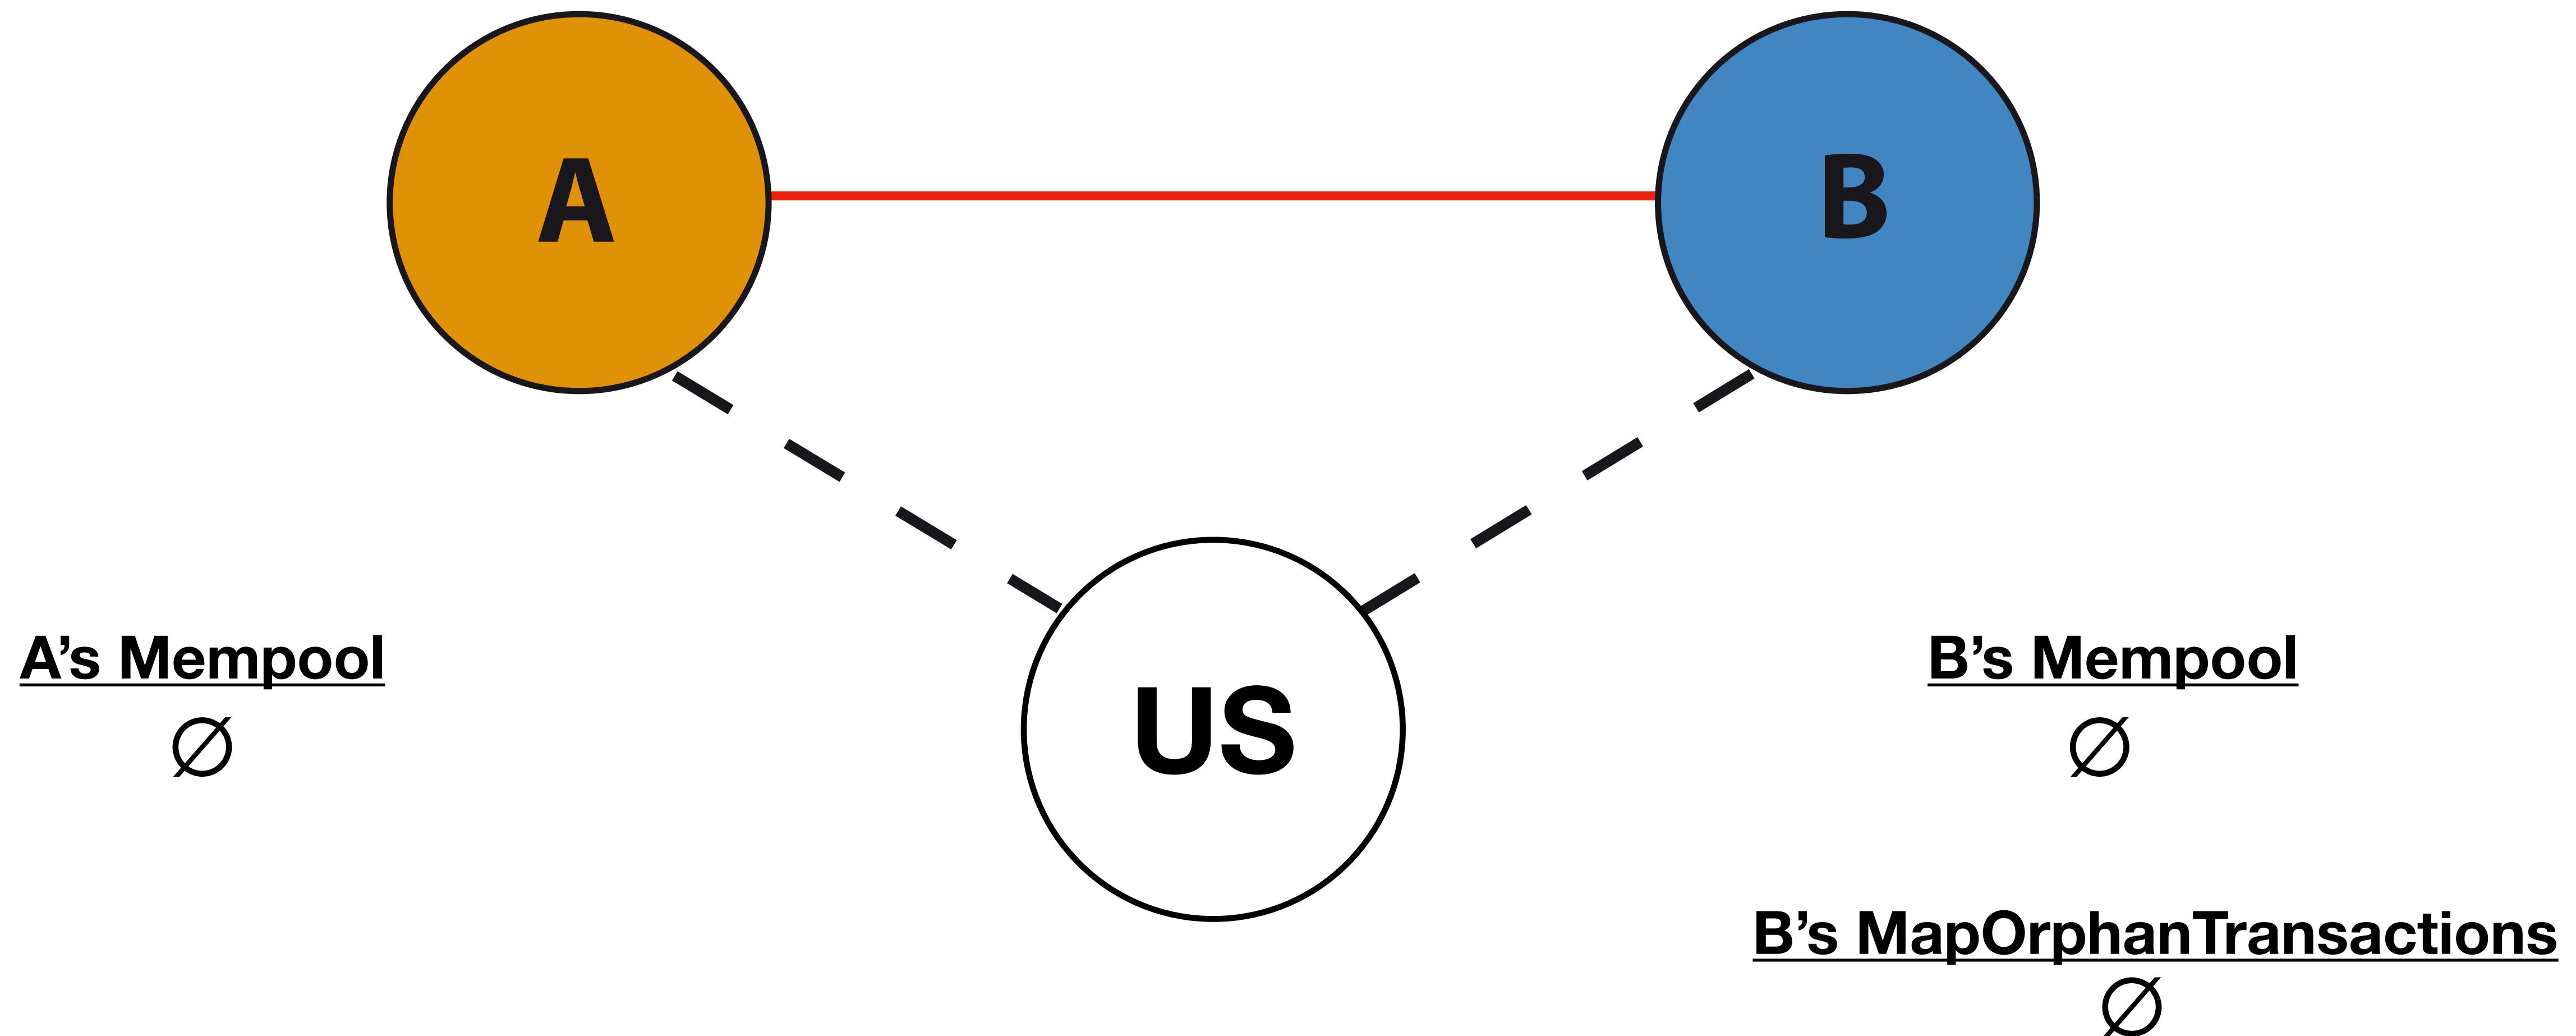
POSITIVE INFERRING TECHNIQUE



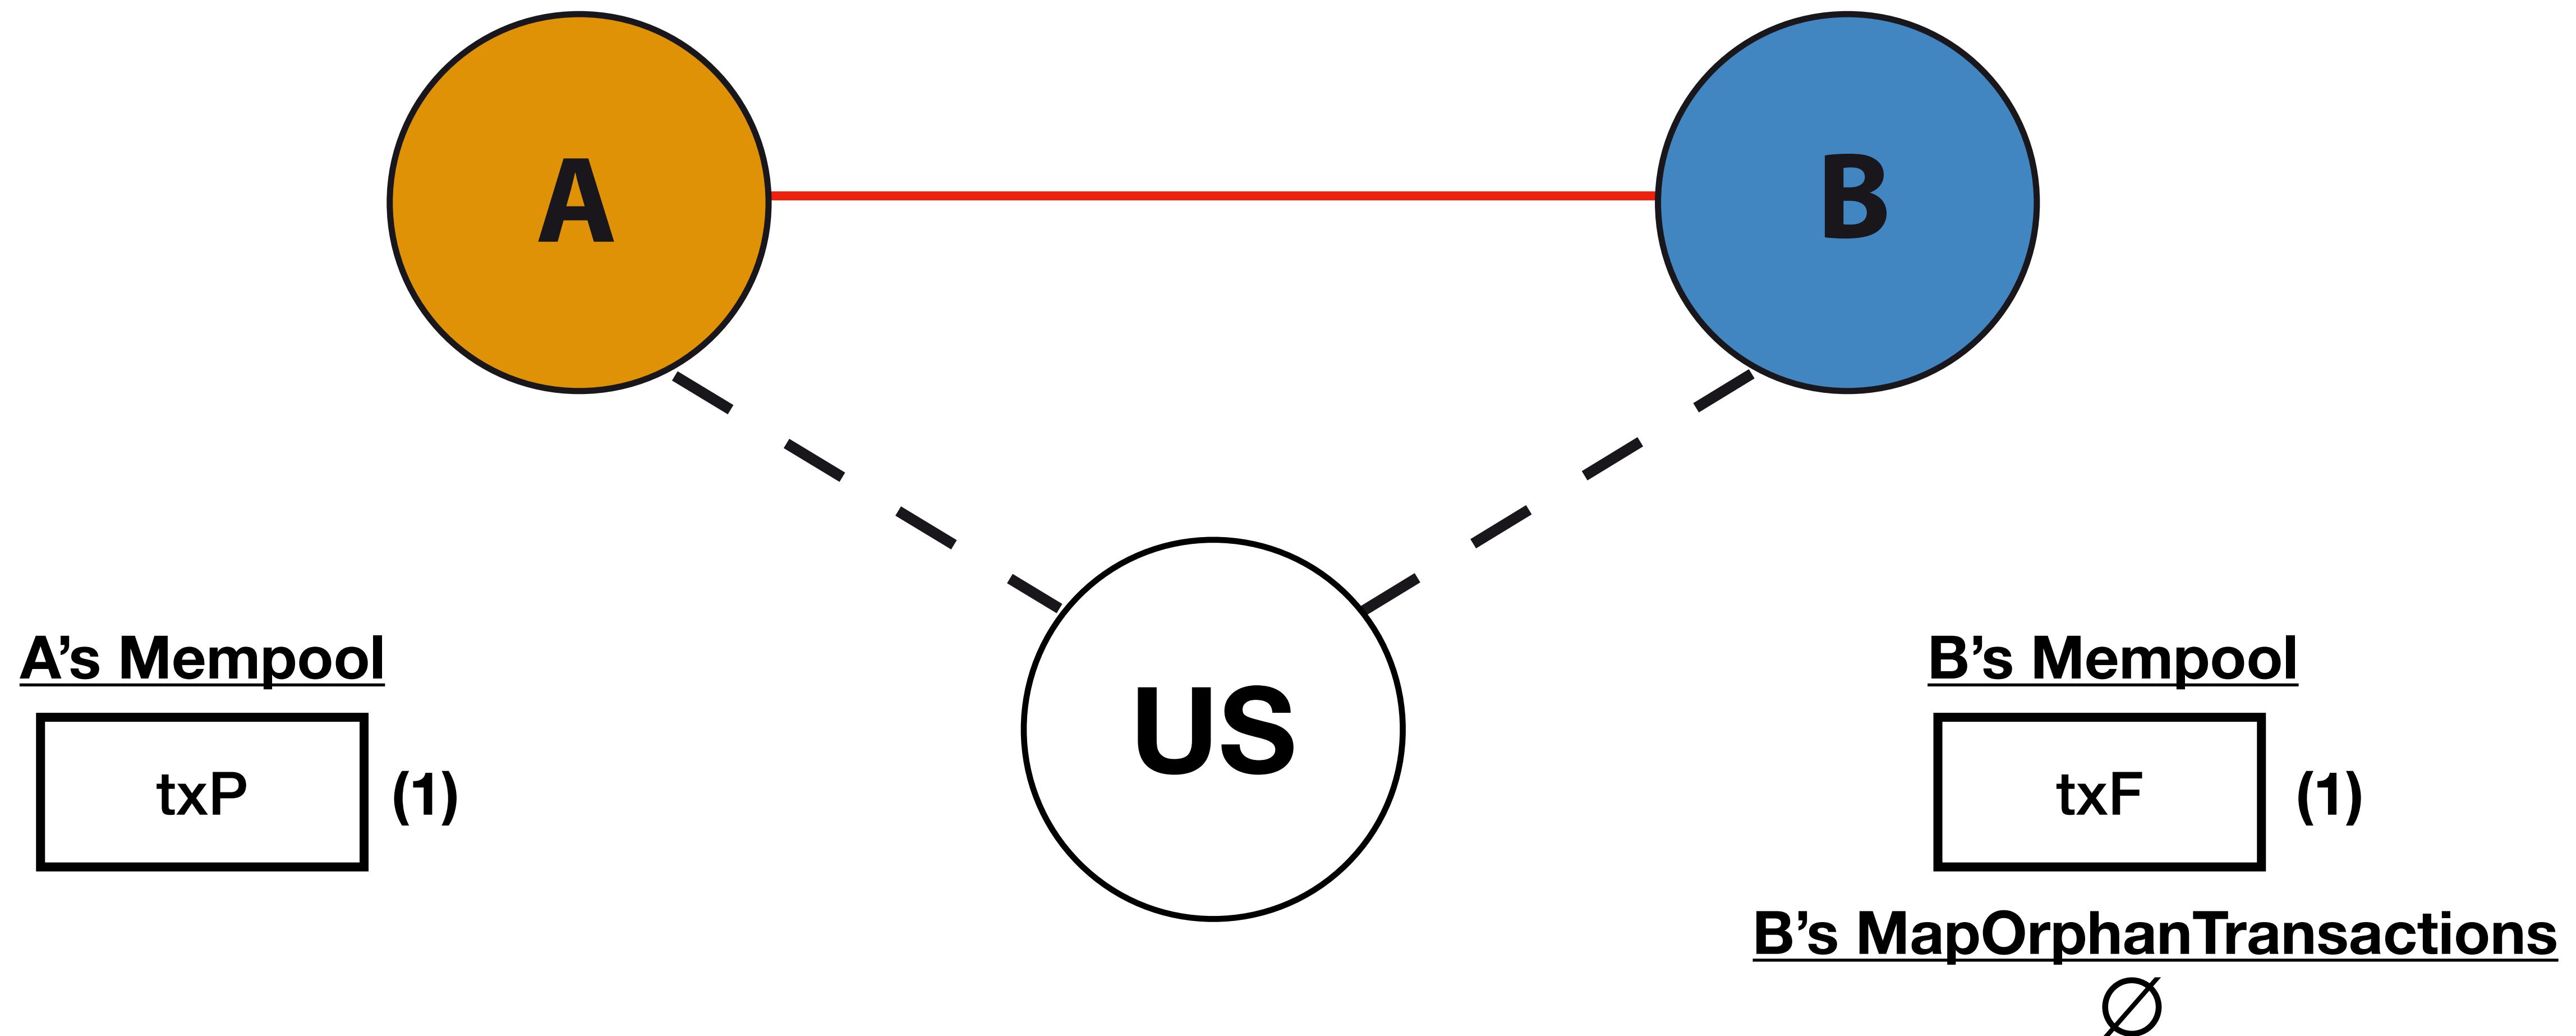
POSITIVE INFERRING TECHNIQUE



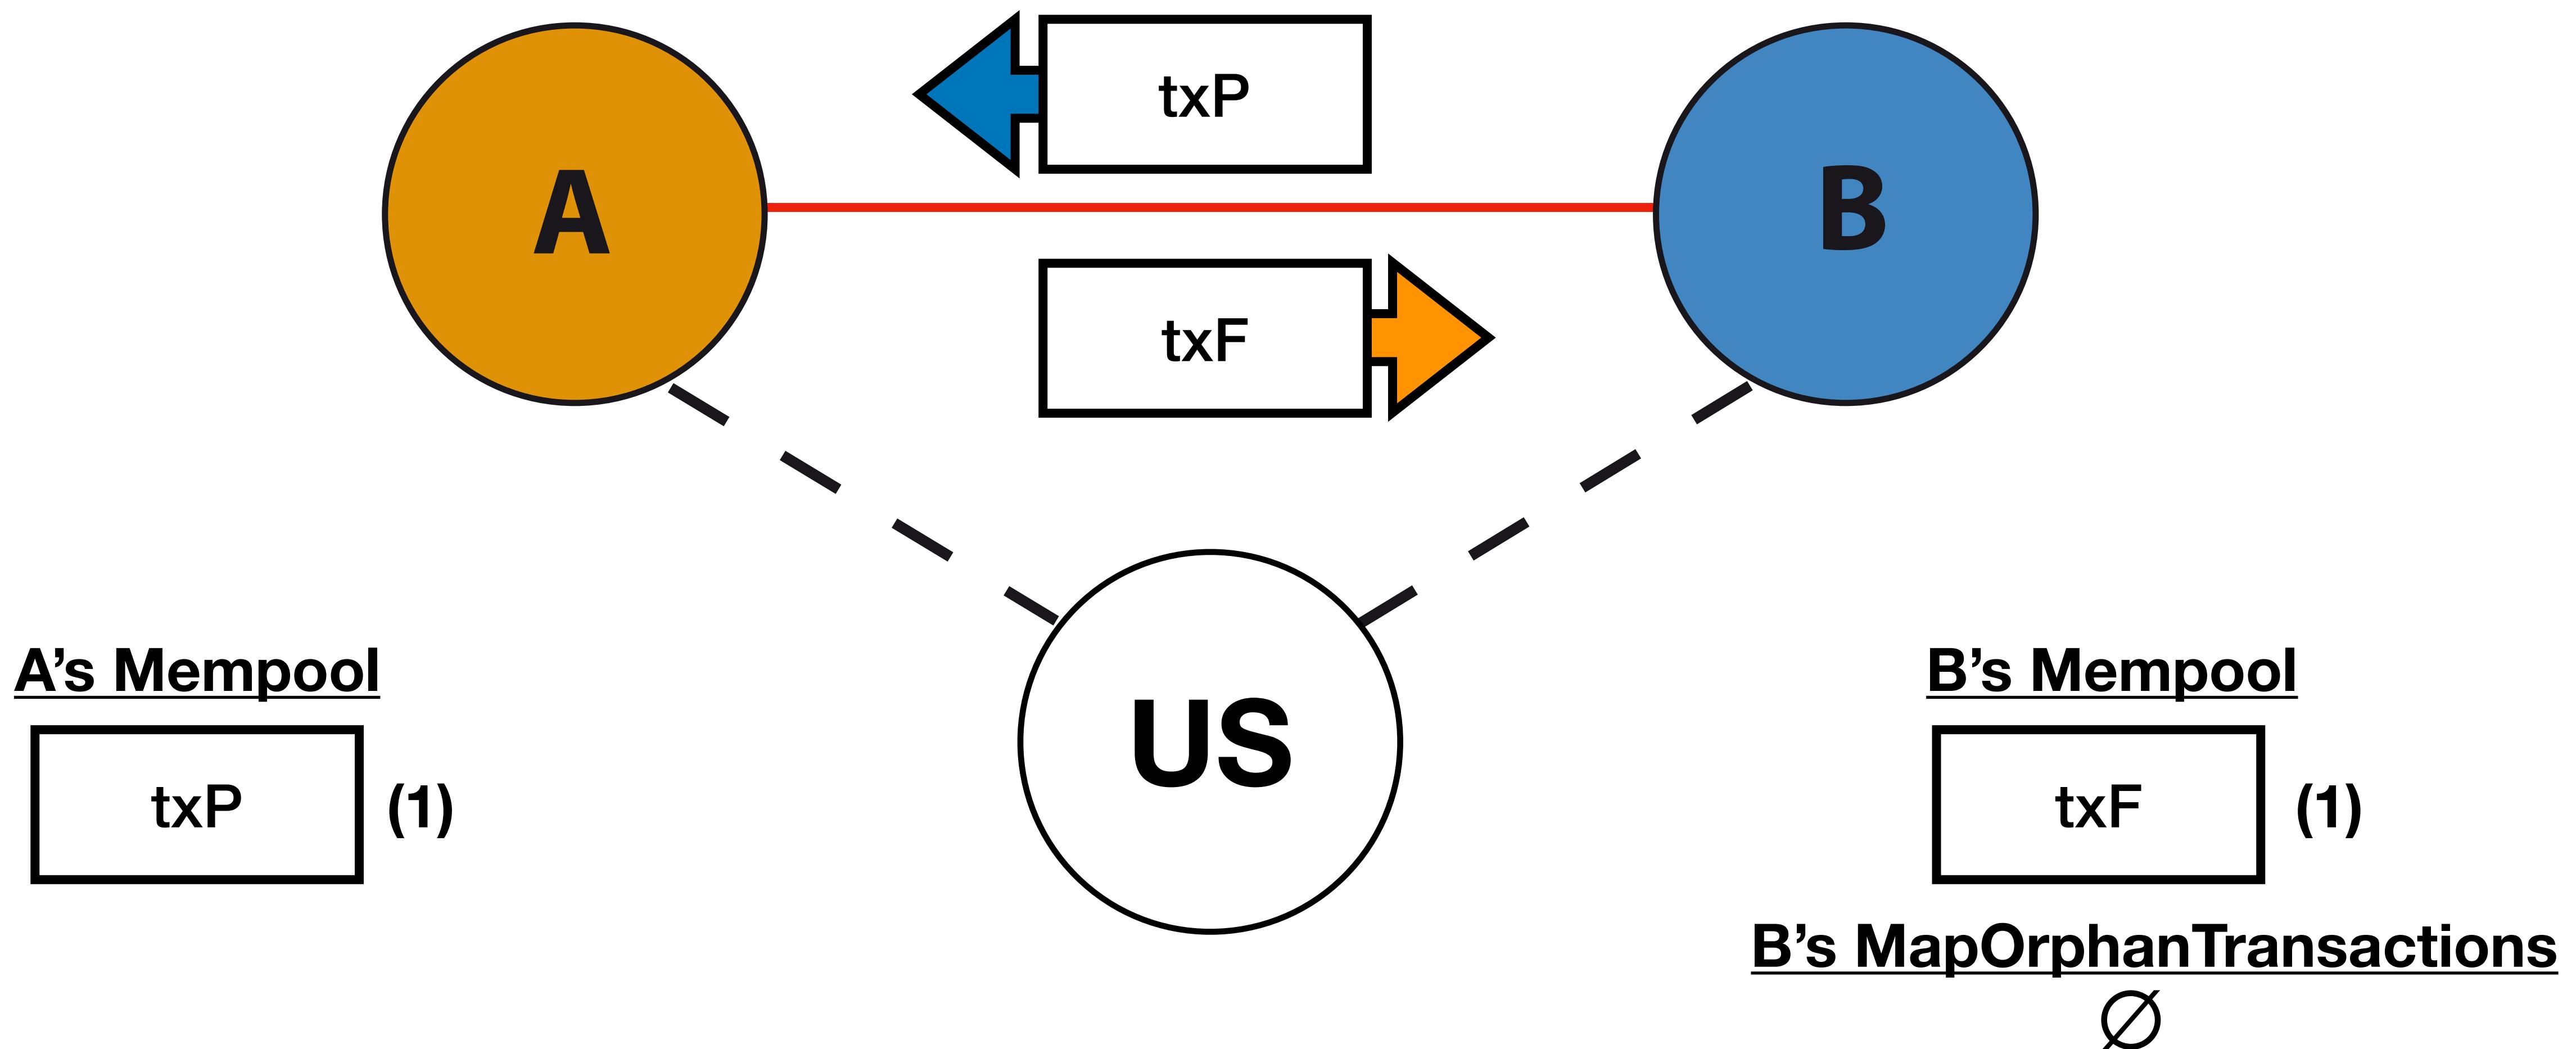
POSITIVE INFERRING TECHNIQUE



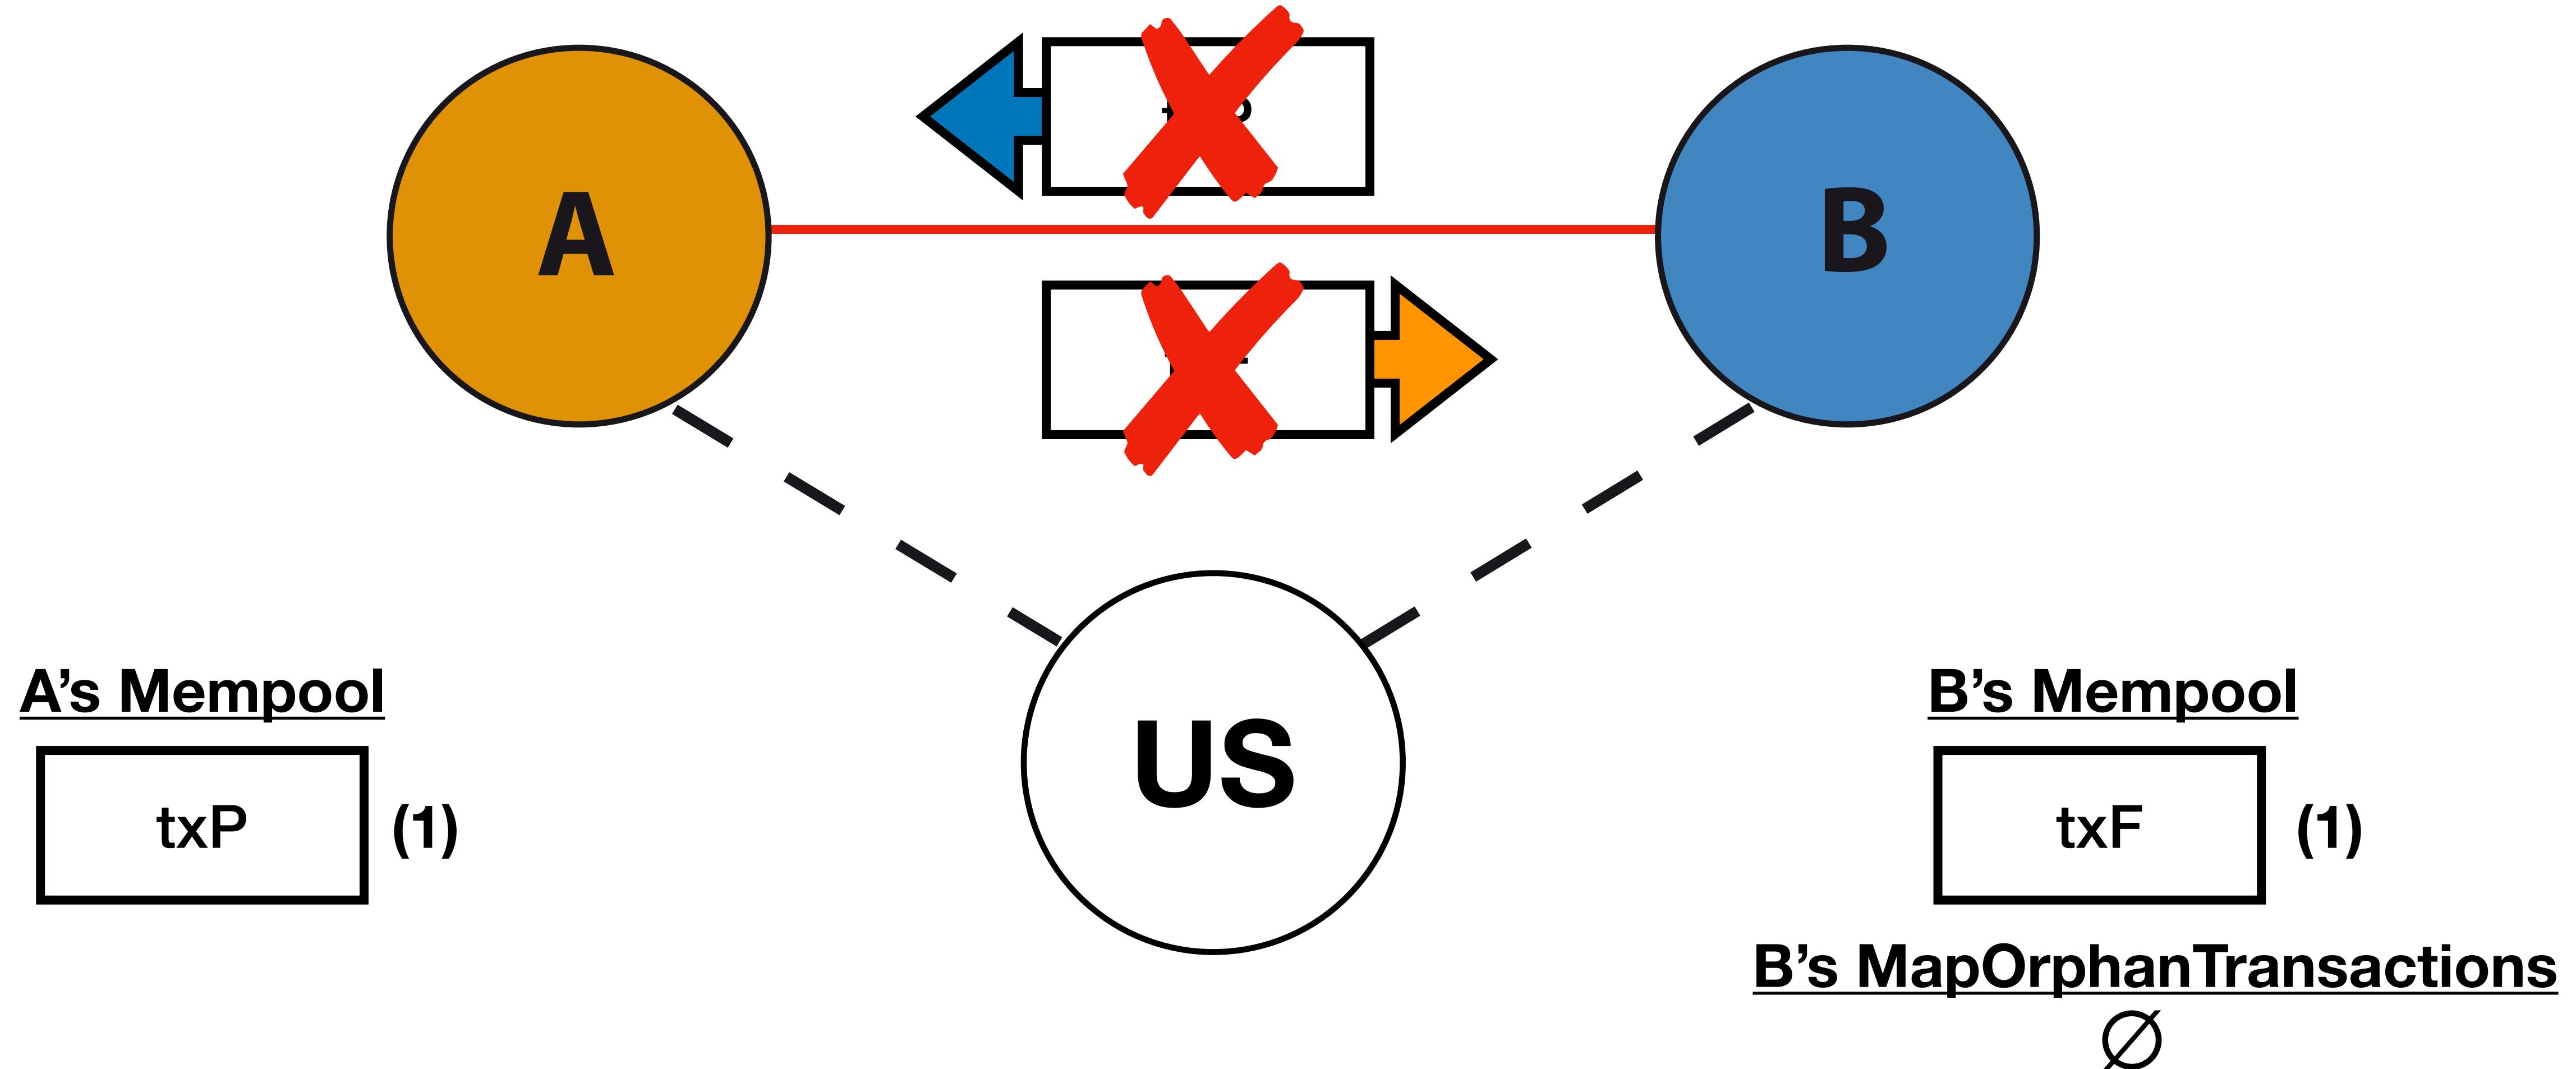
POSITIVE INFERRING TECHNIQUE



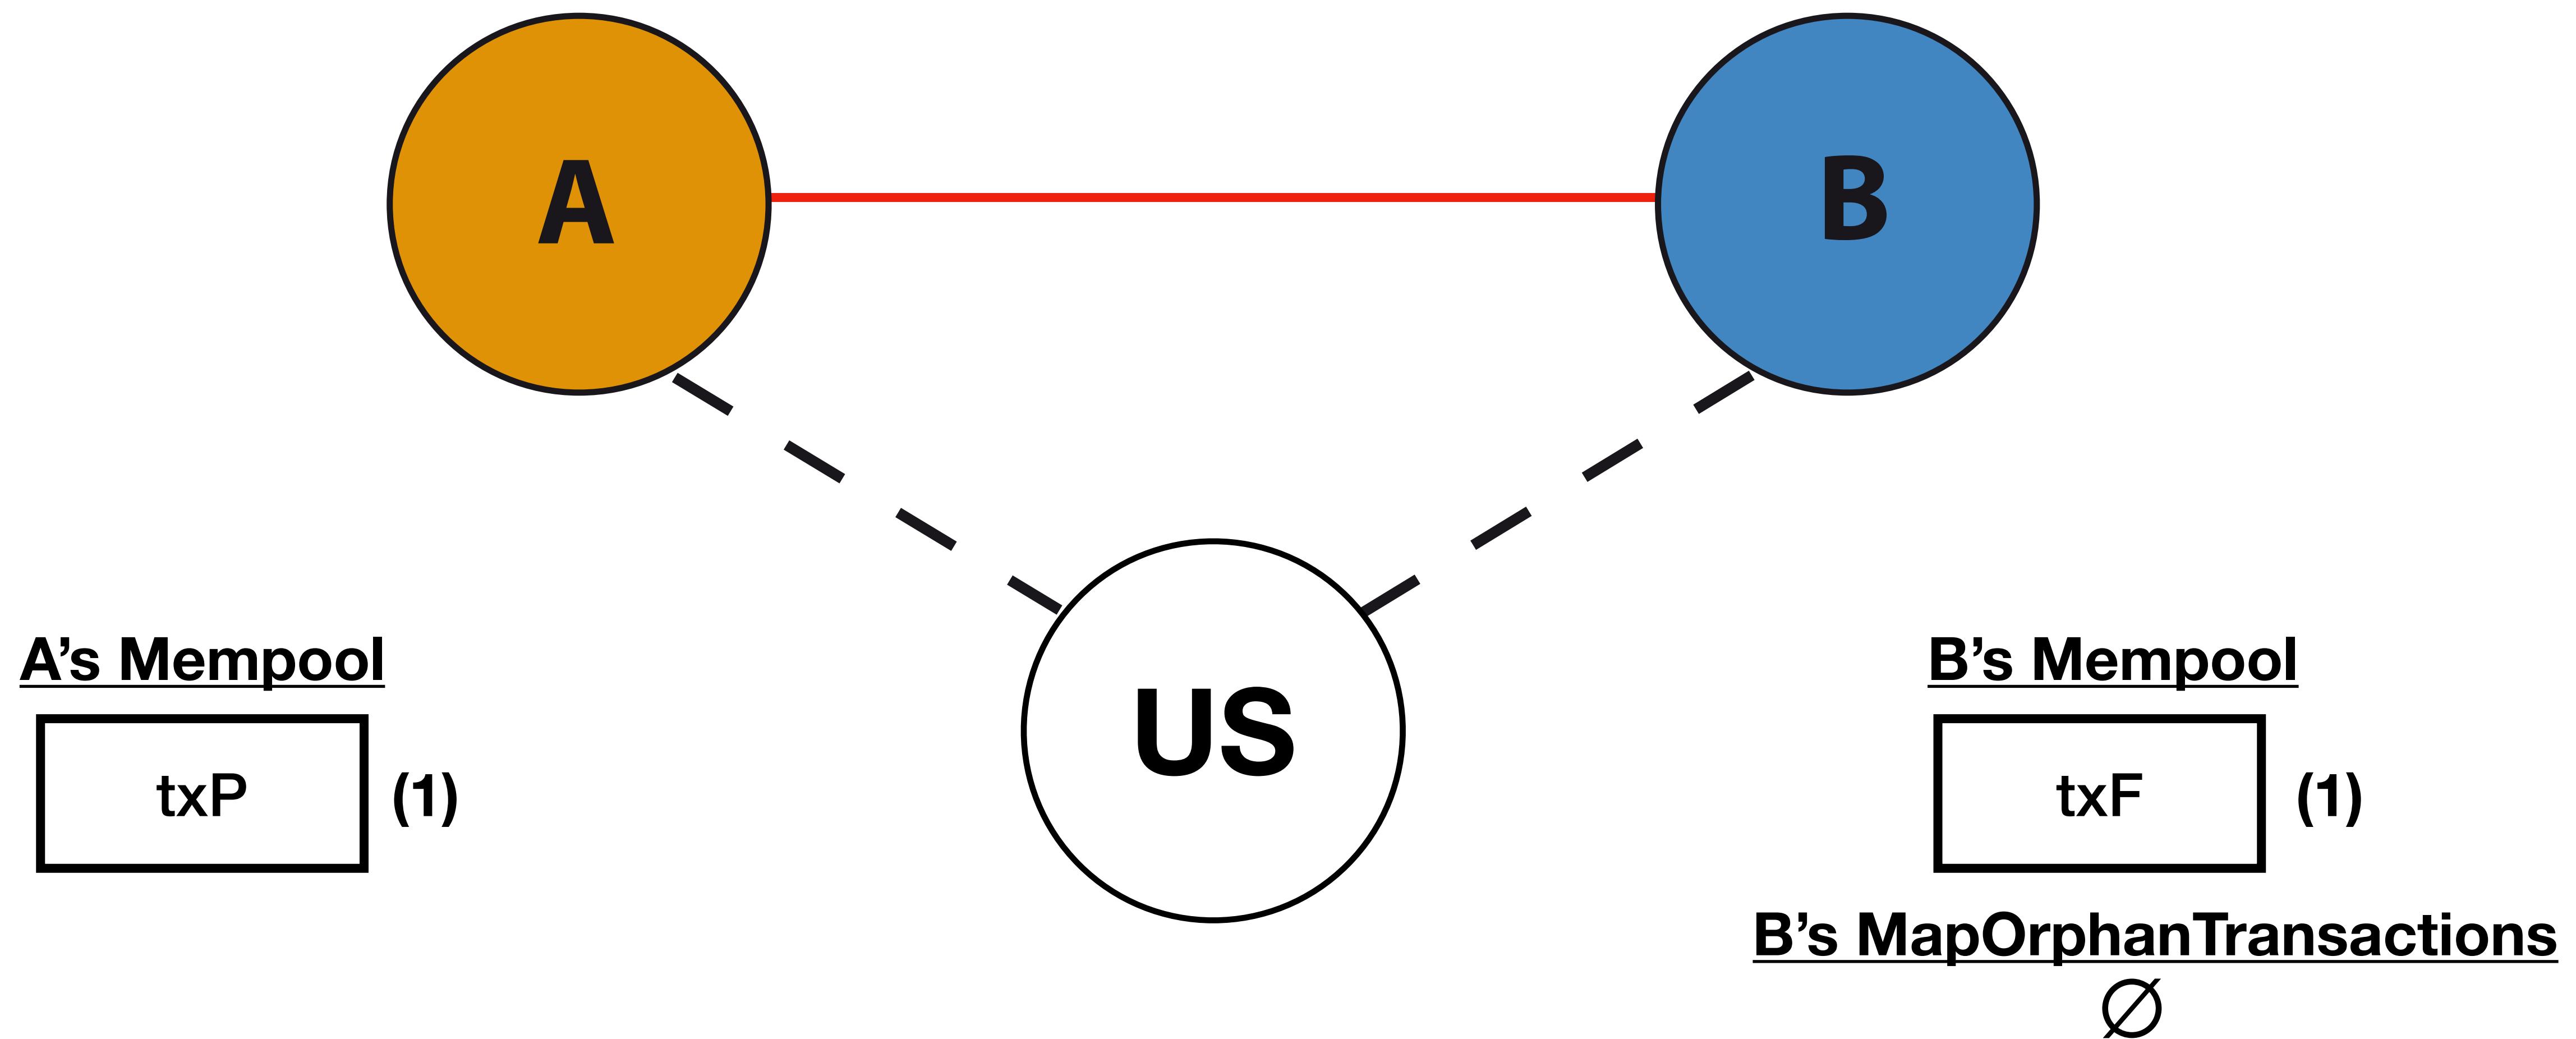
POSITIVE INFERRING TECHNIQUE



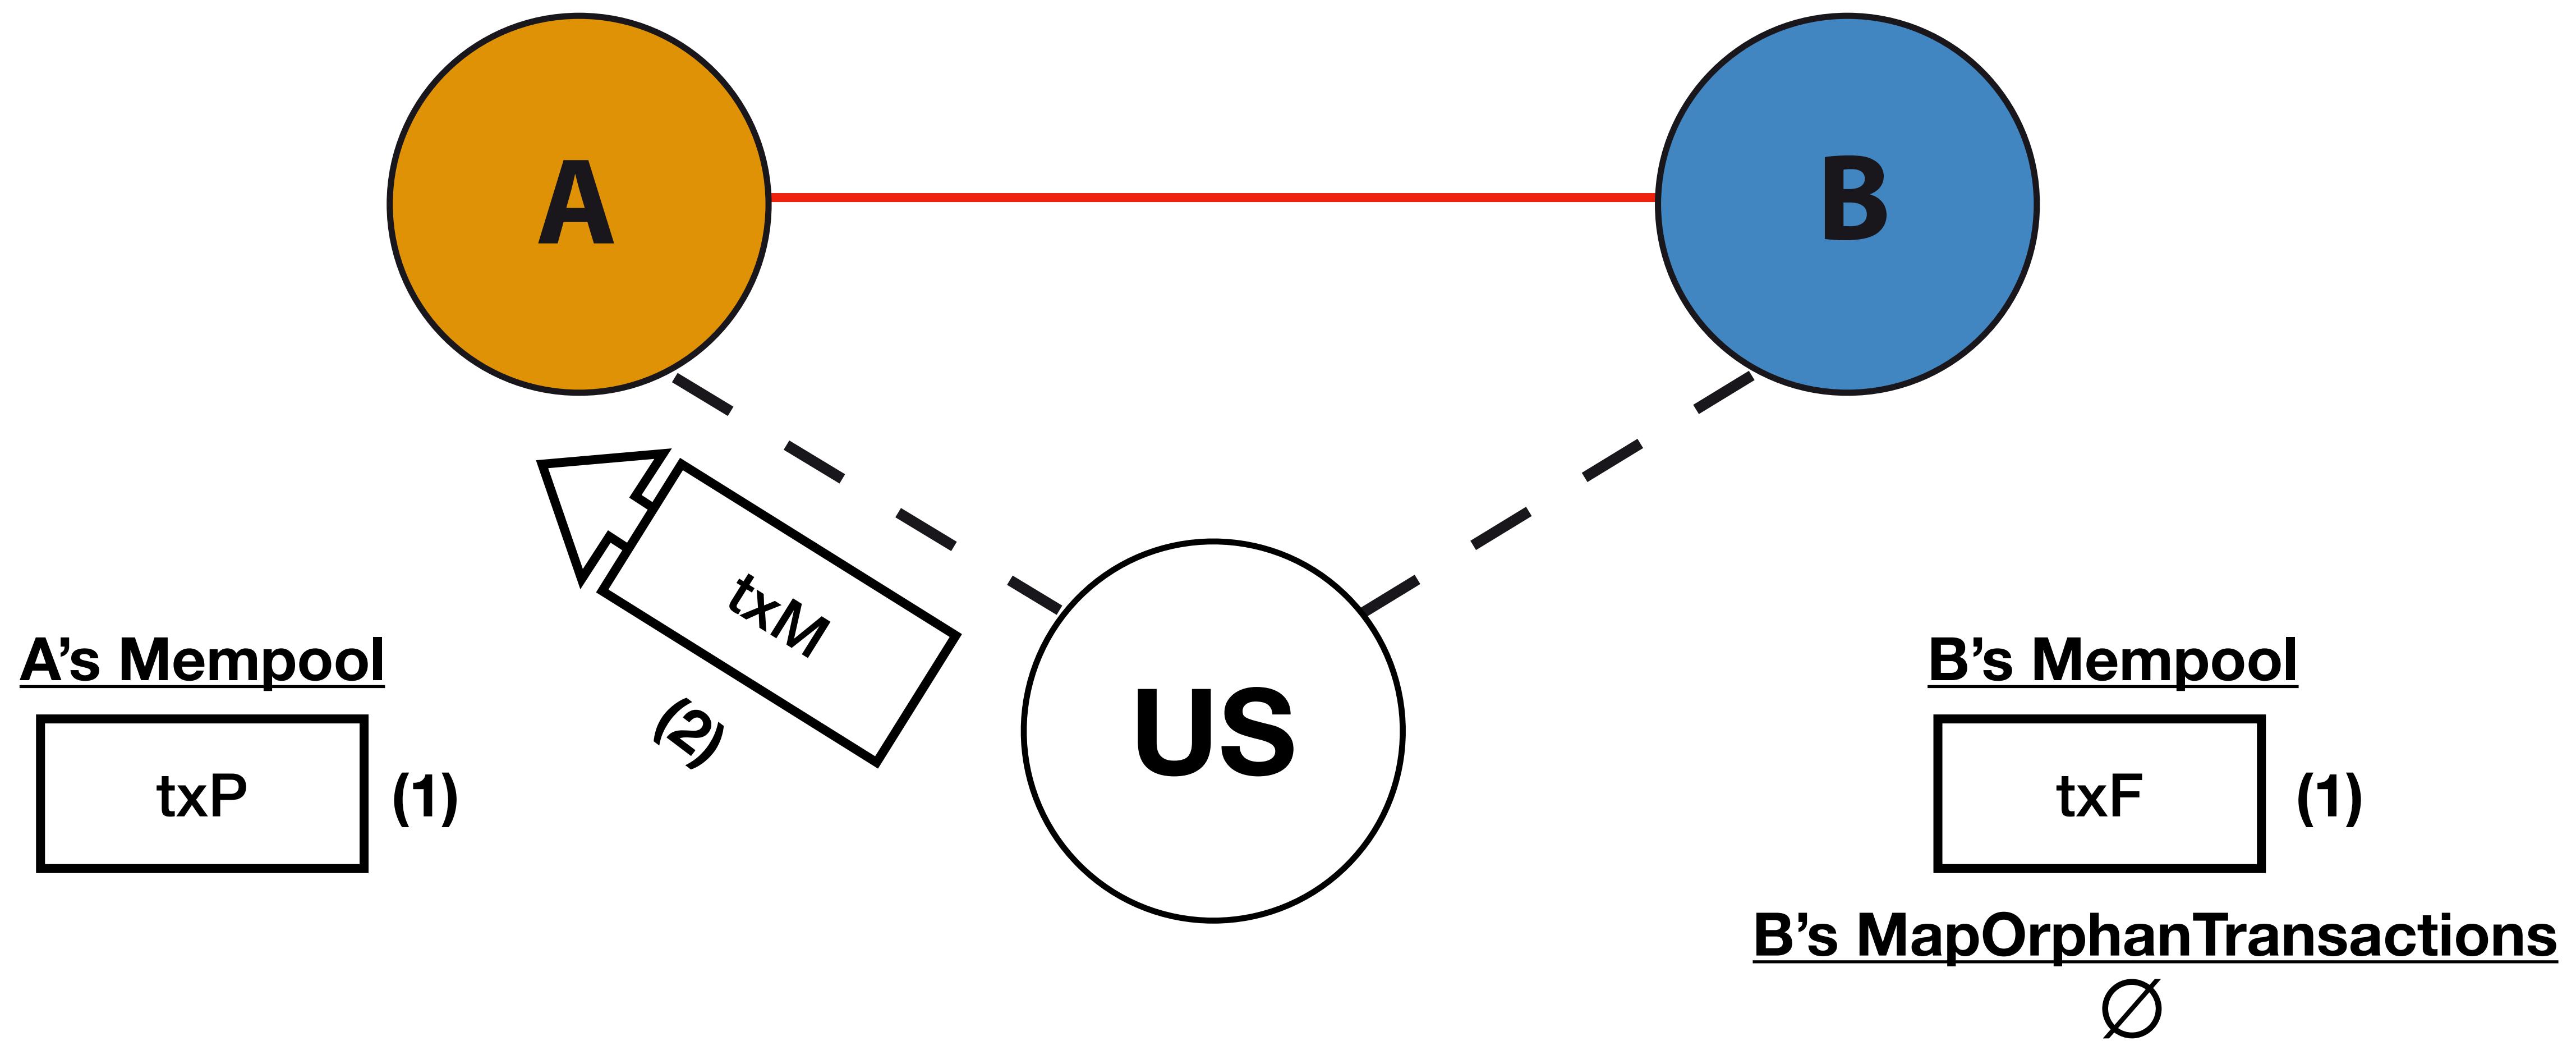
POSITIVE INFERRING TECHNIQUE



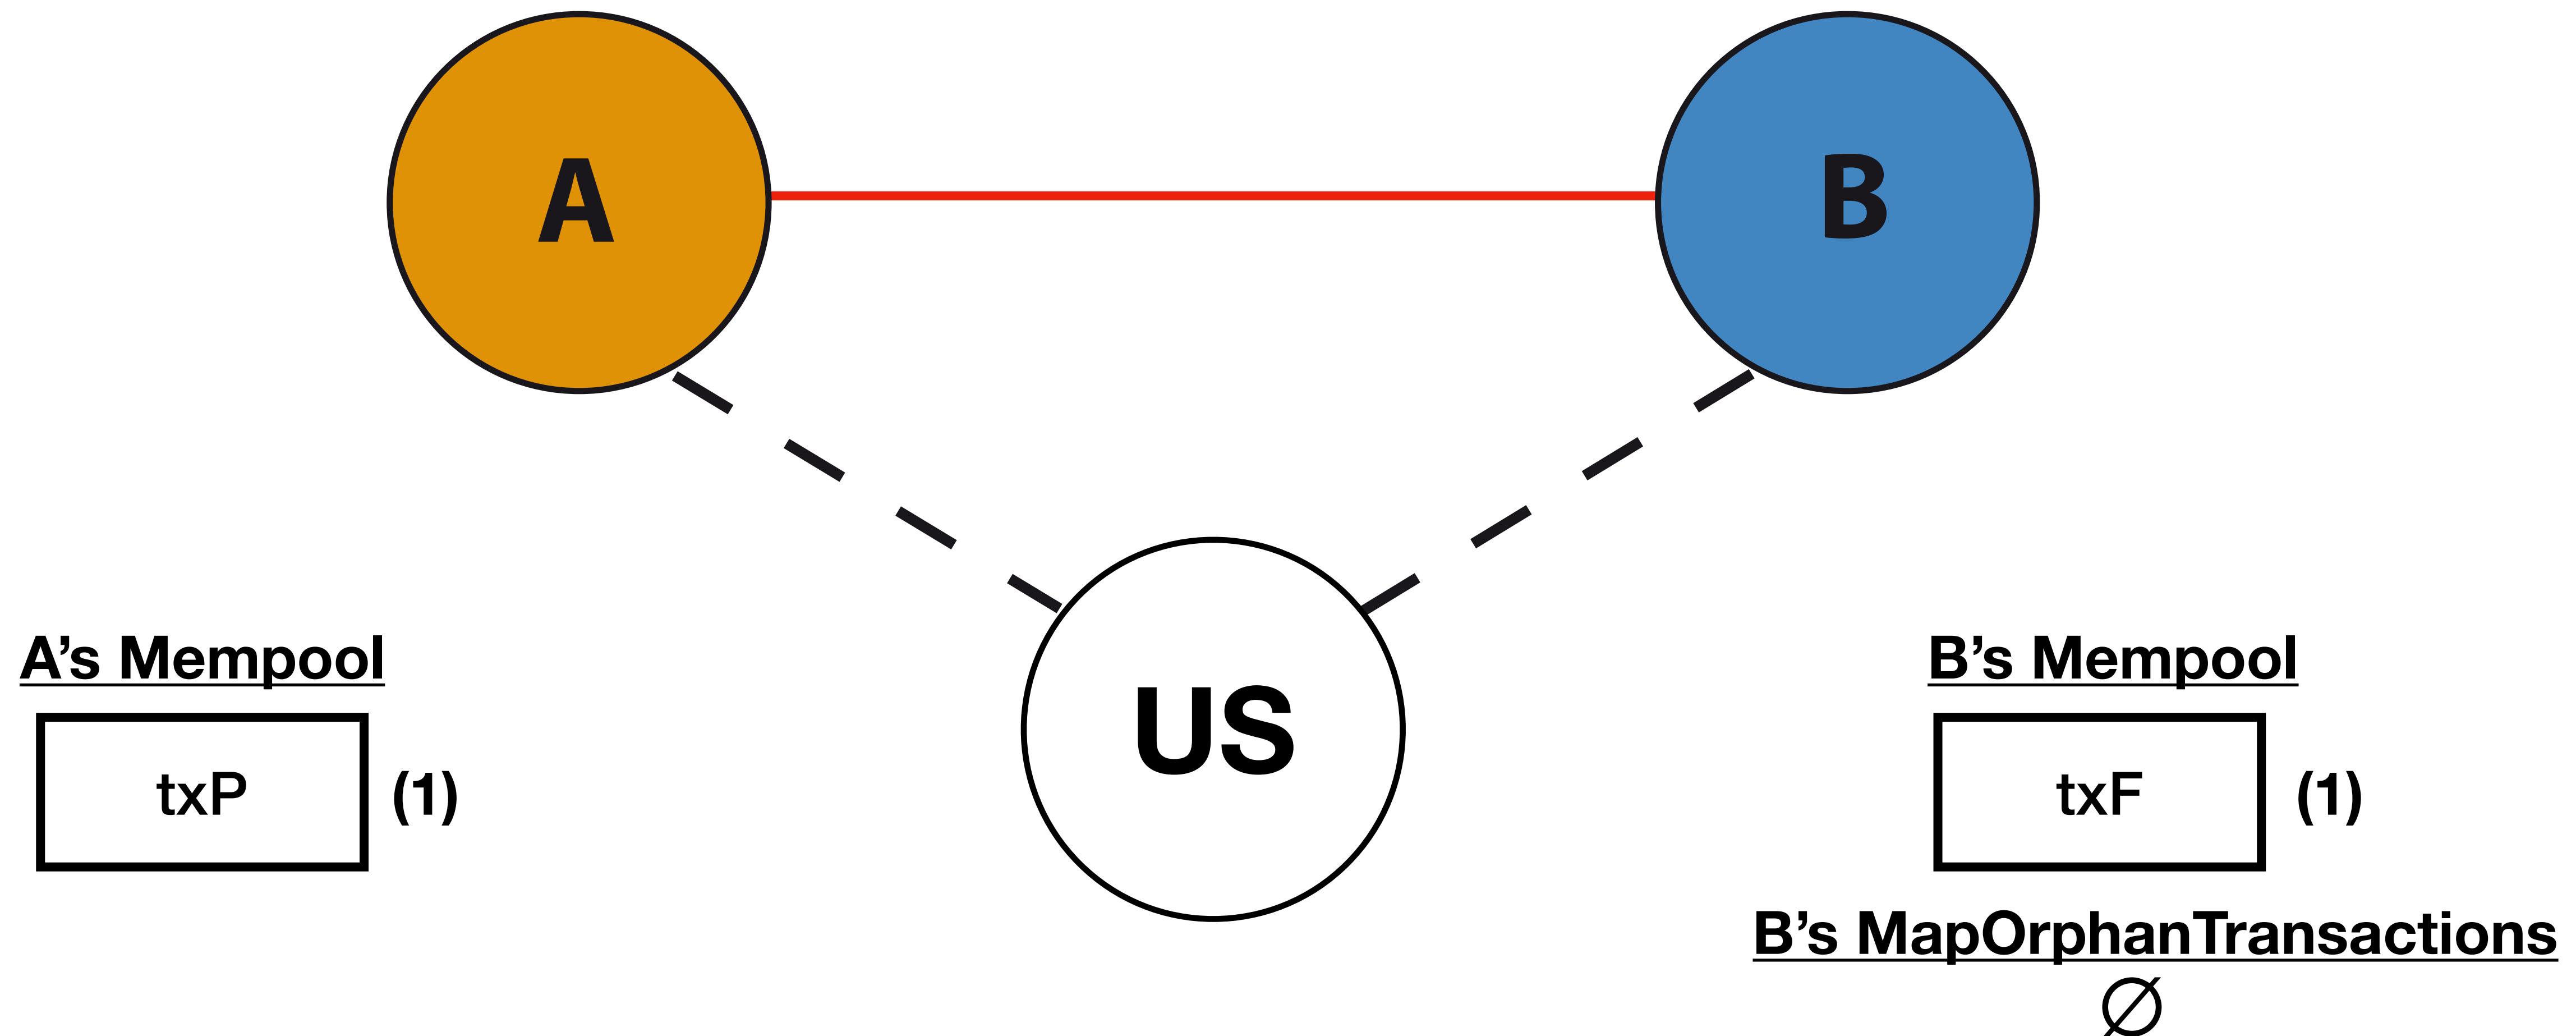
POSITIVE INFERRING TECHNIQUE



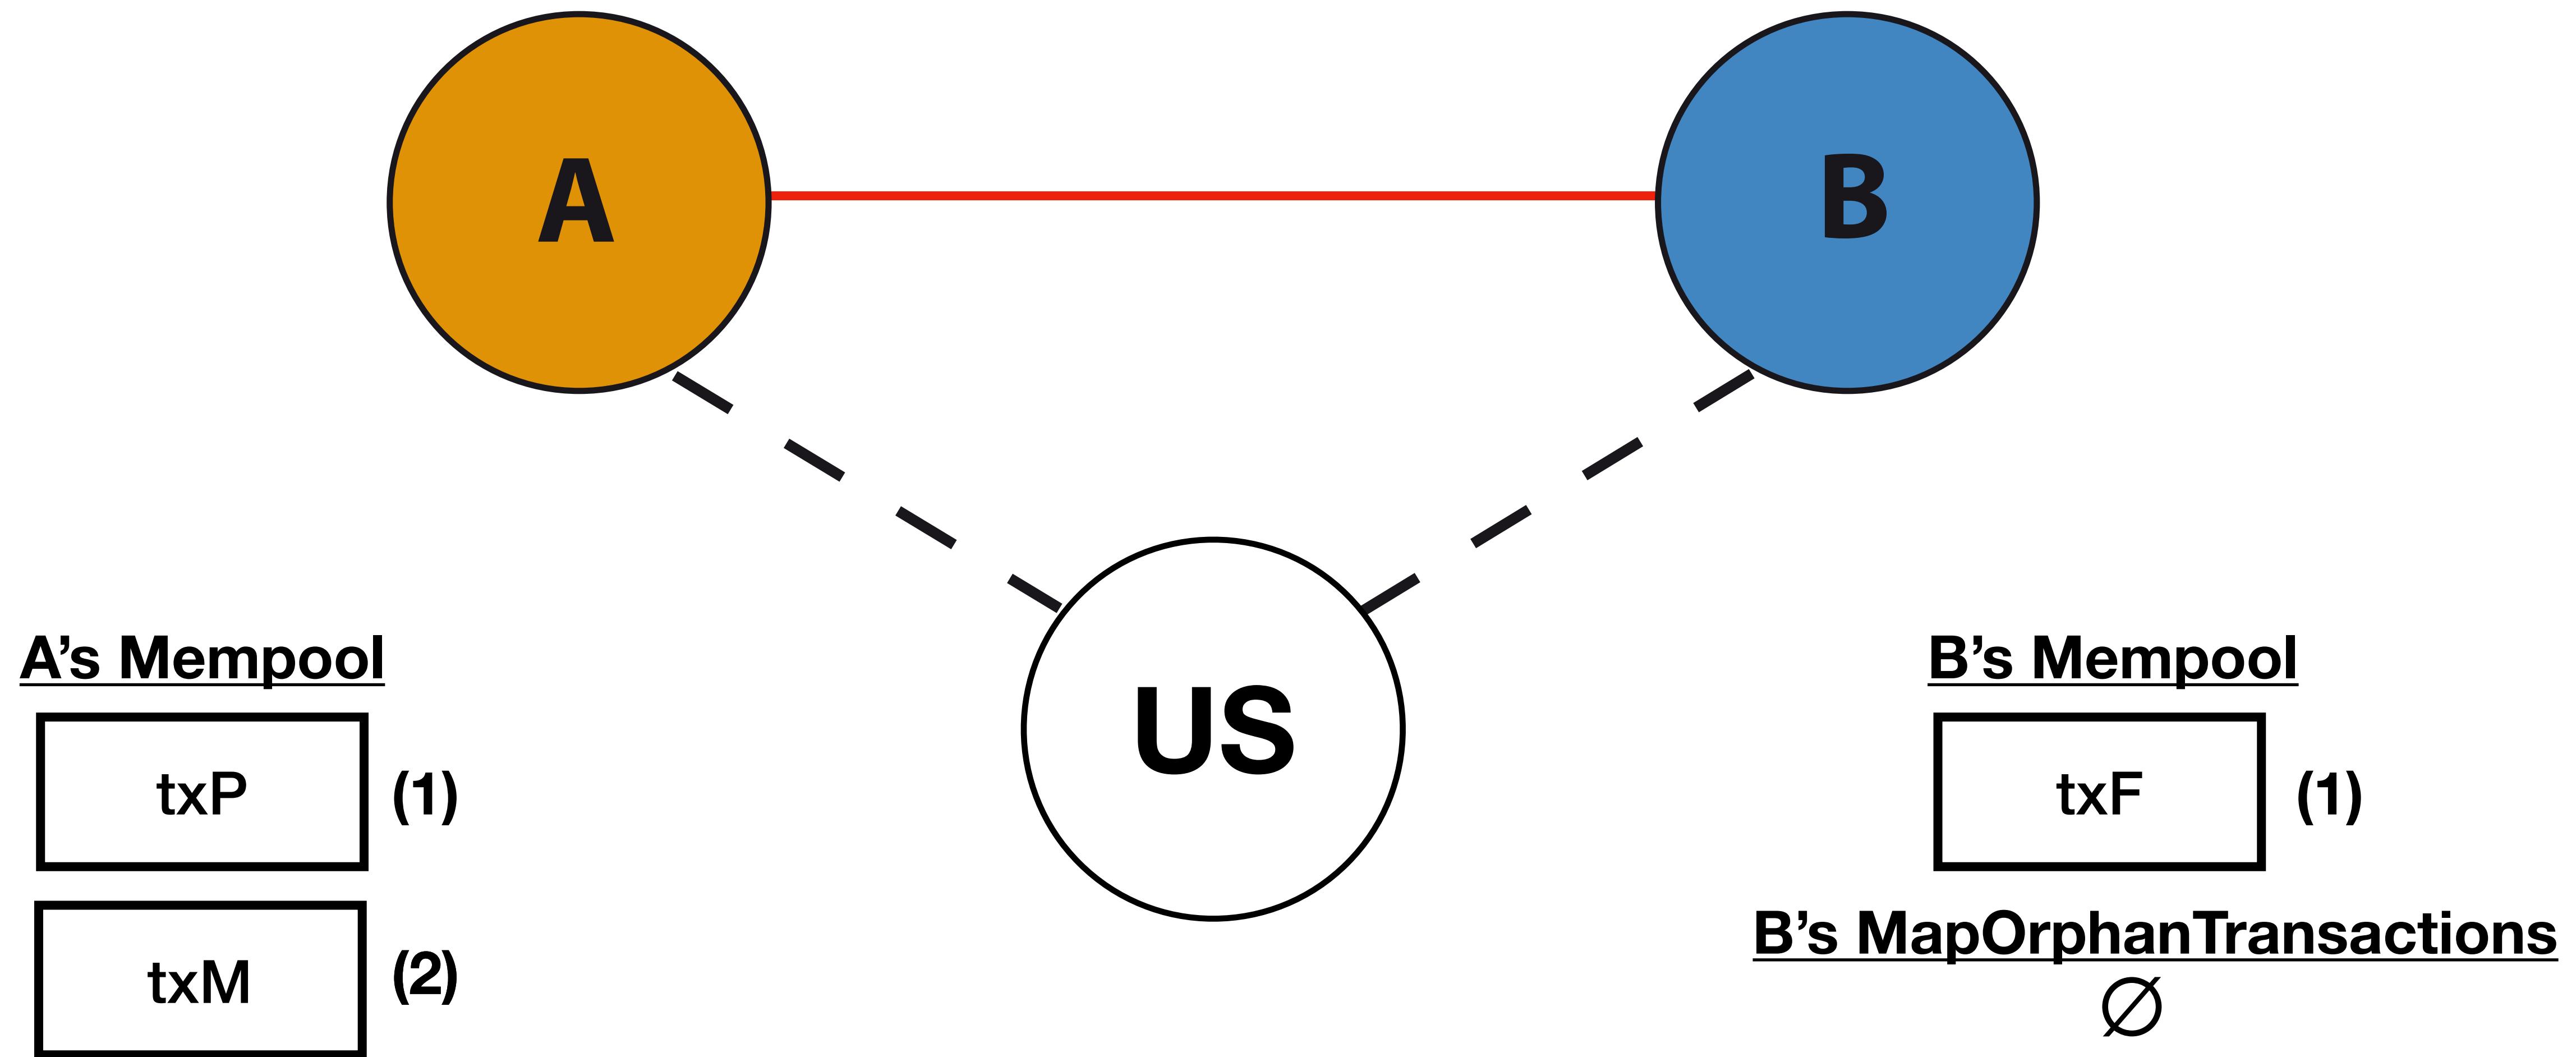
POSITIVE INFERRING TECHNIQUE



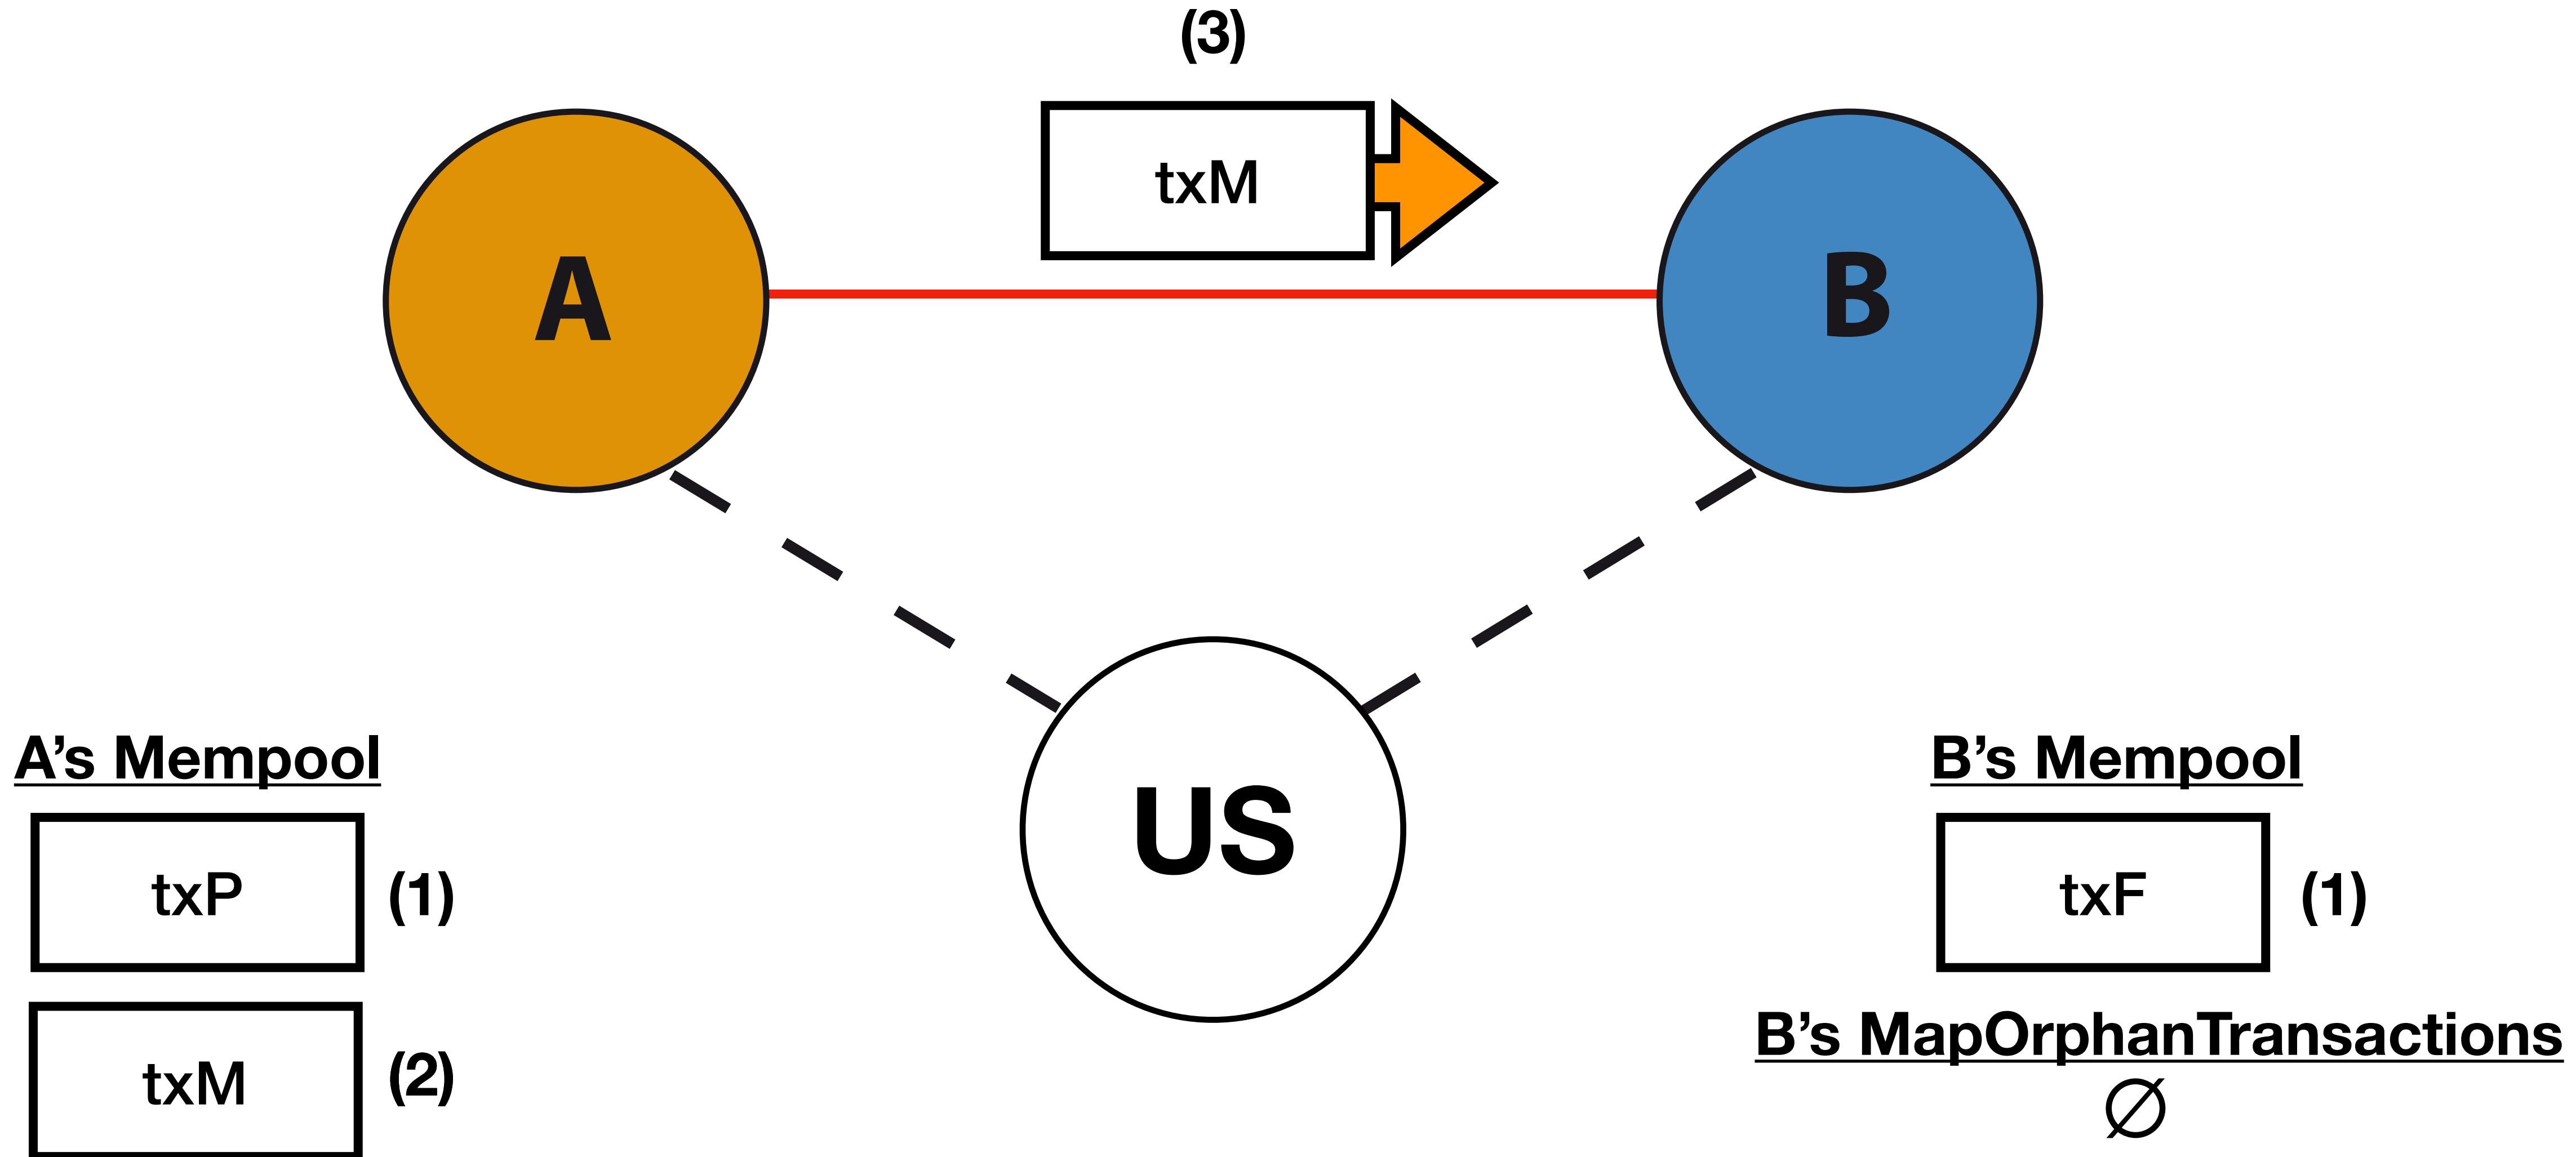
POSITIVE INFERRING TECHNIQUE



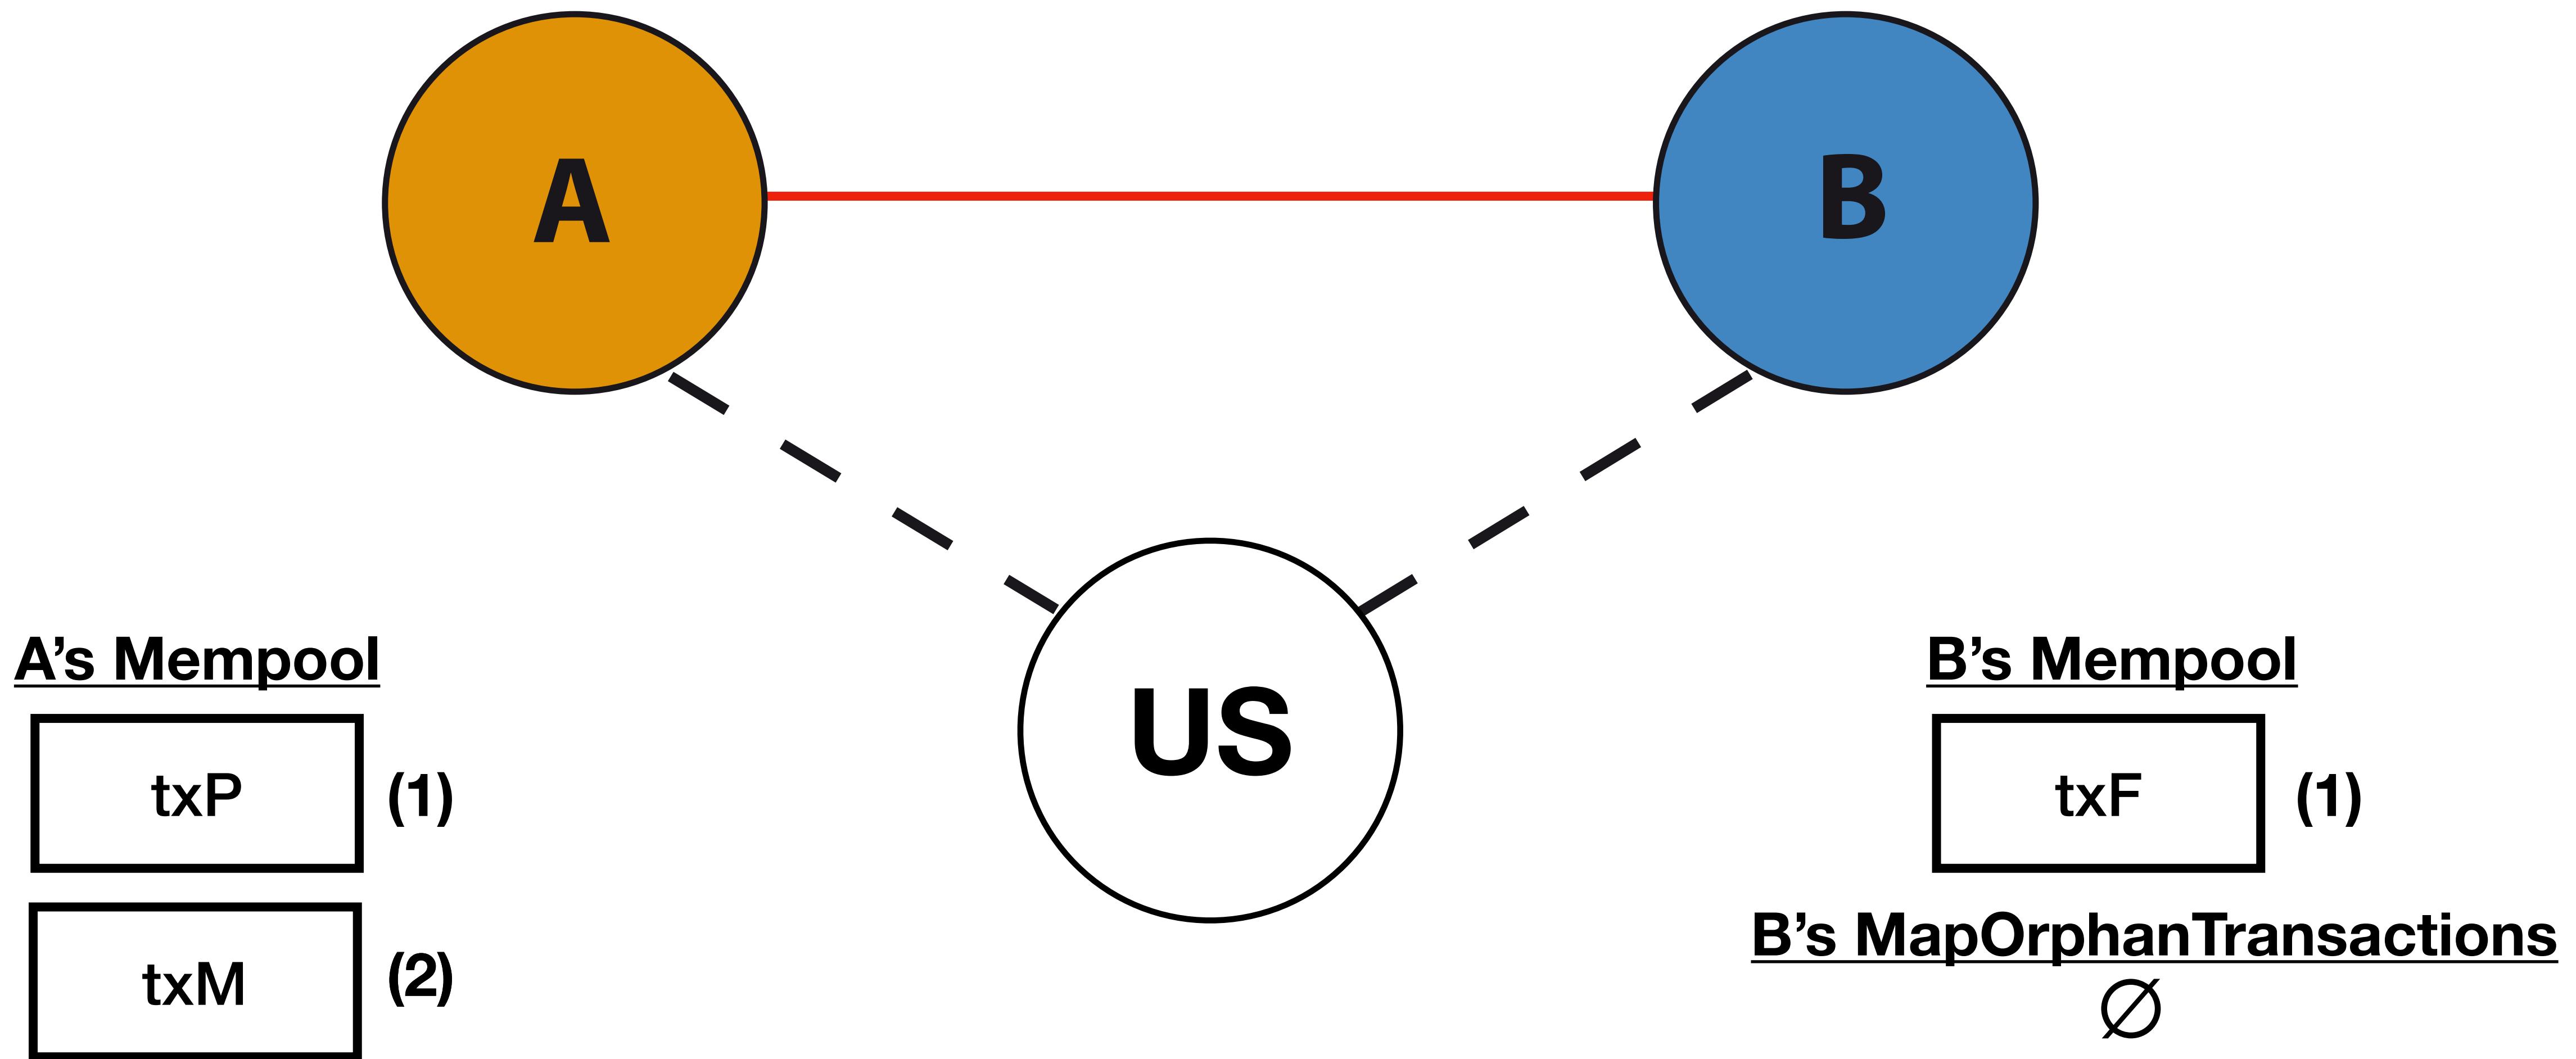
POSITIVE INFERRING TECHNIQUE



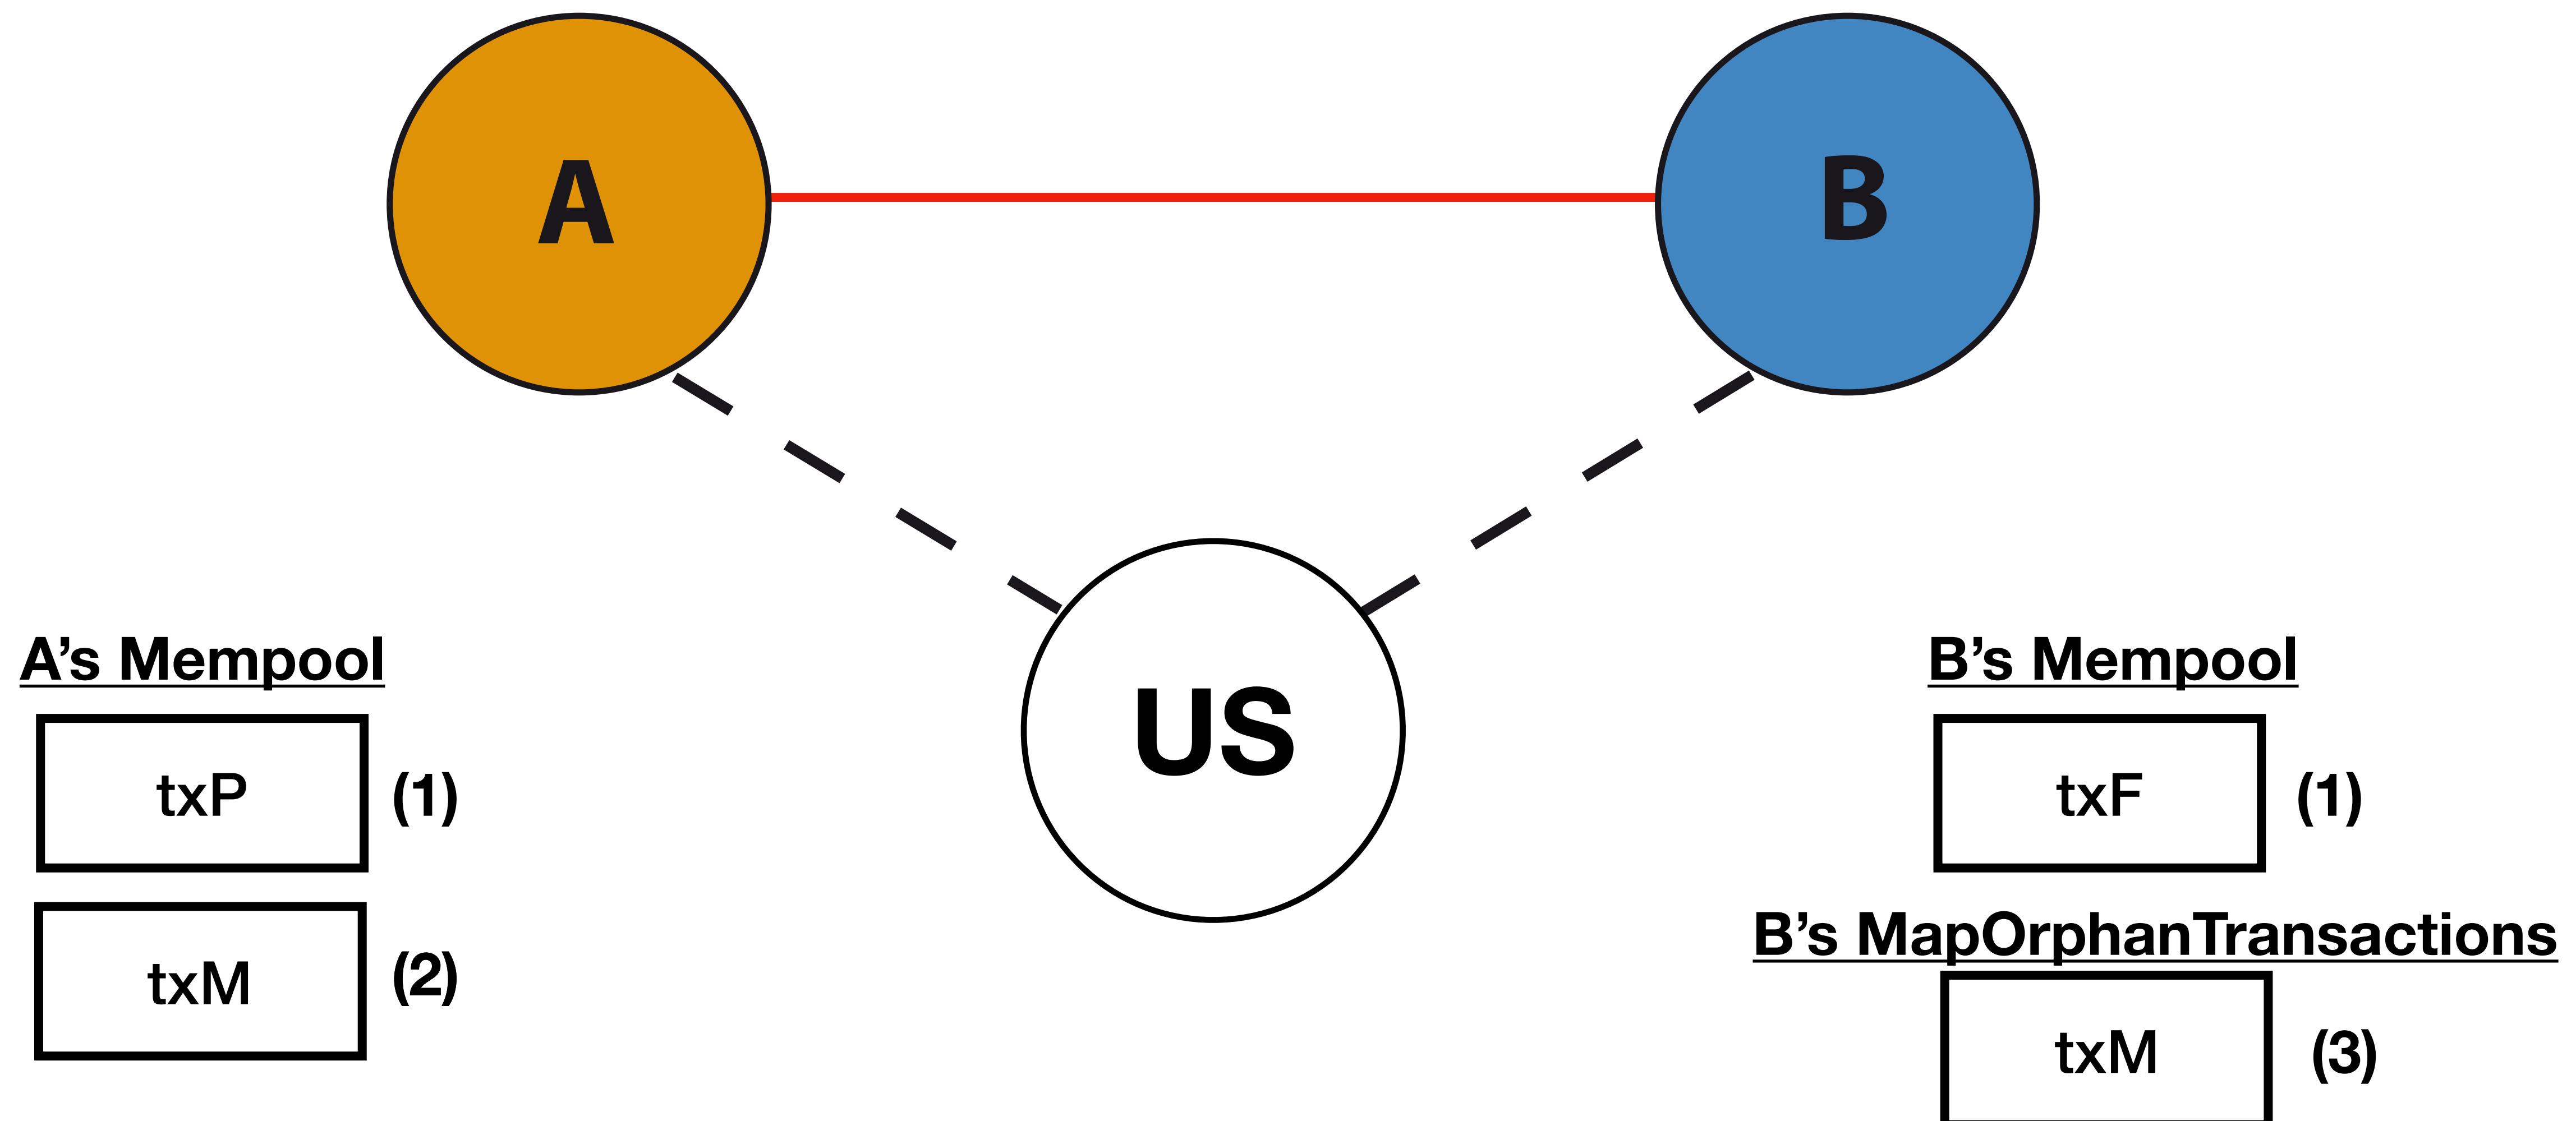
POSITIVE INFERRING TECHNIQUE



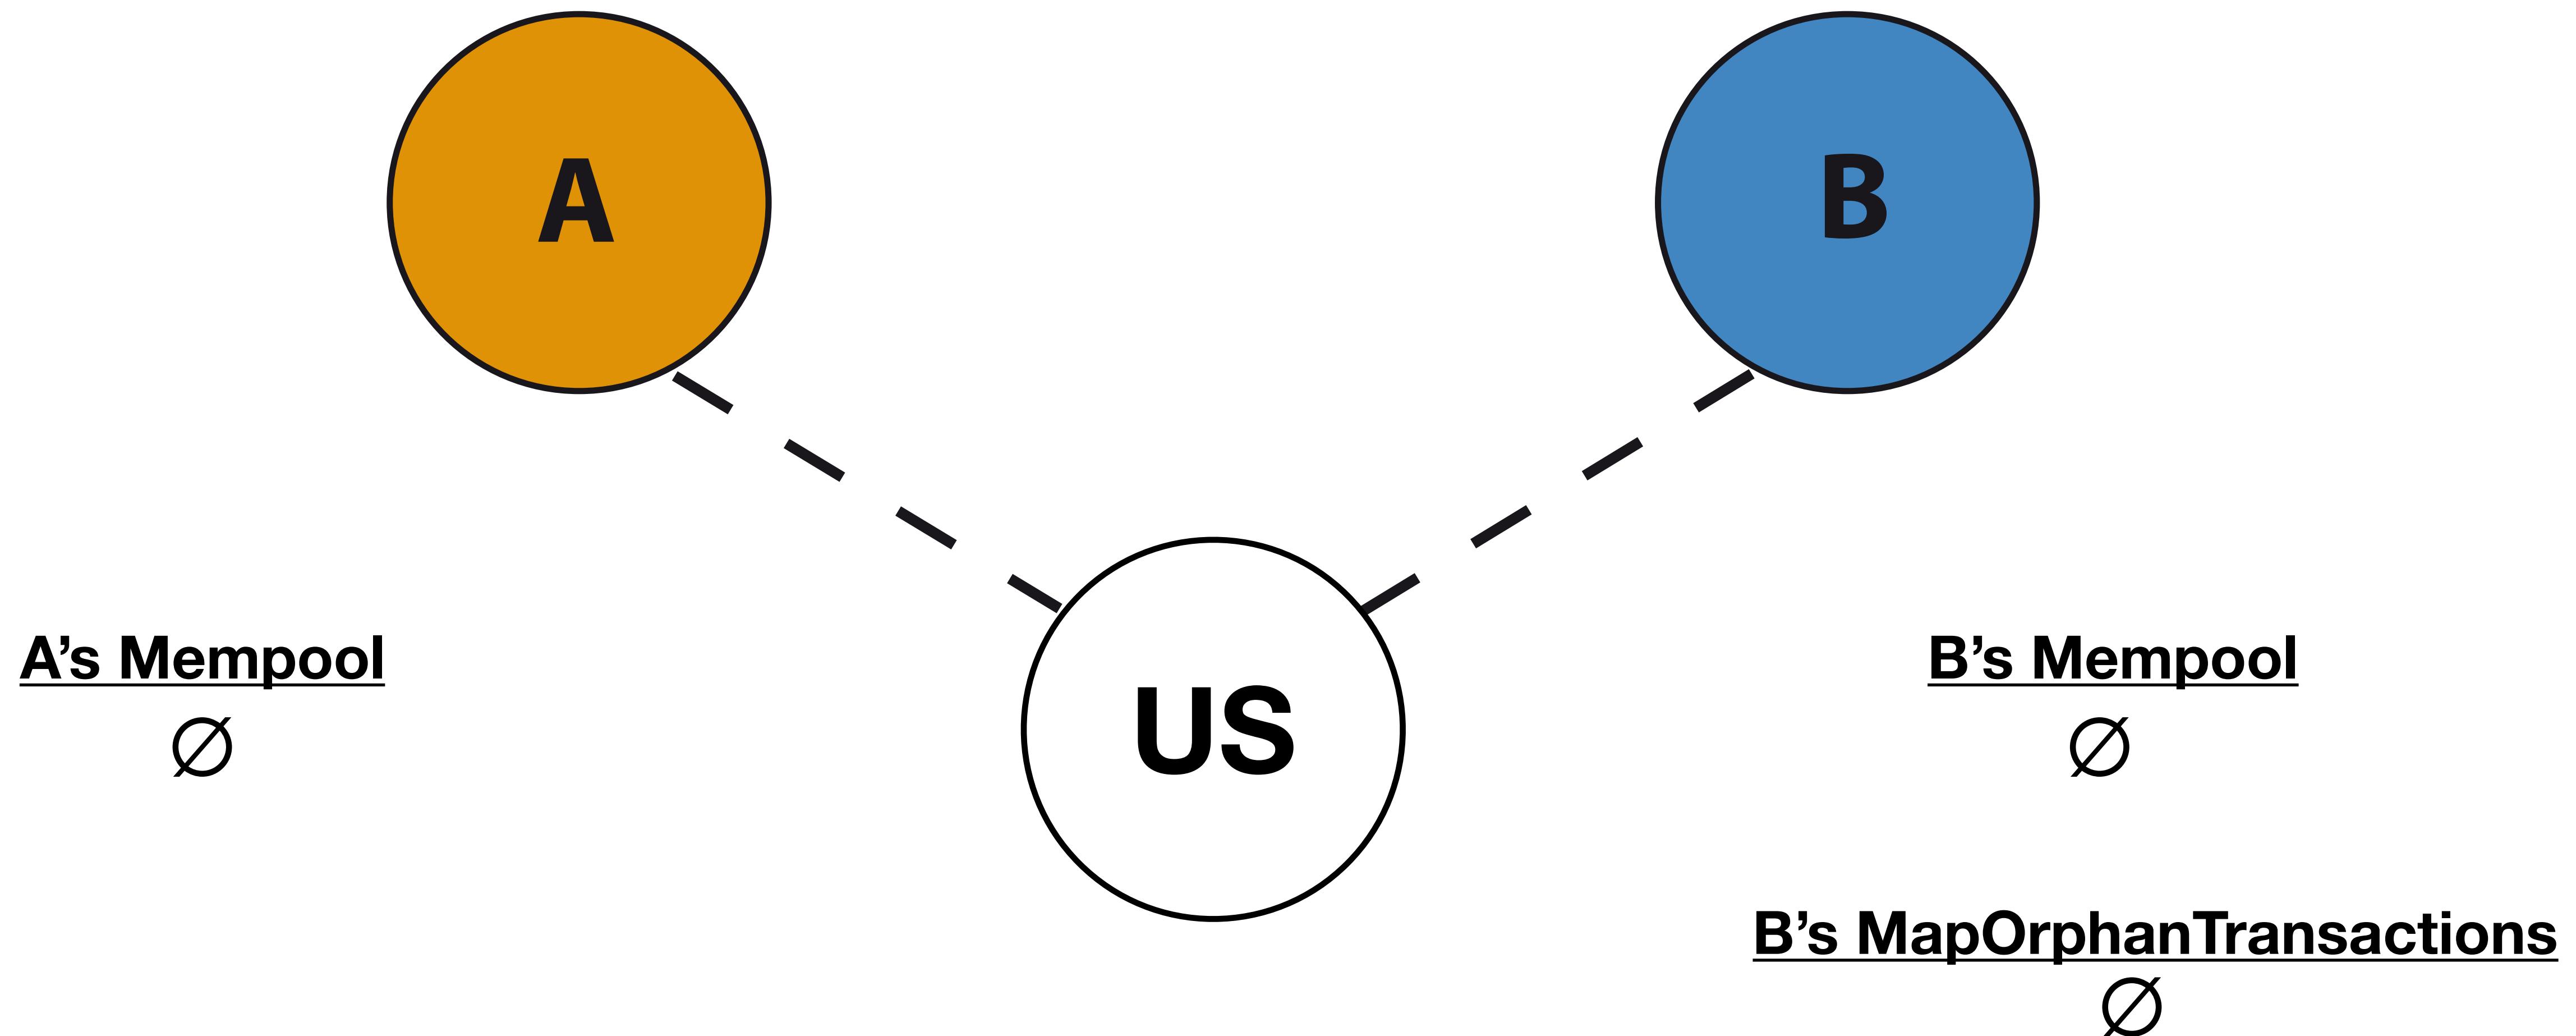
POSITIVE INFERRING TECHNIQUE



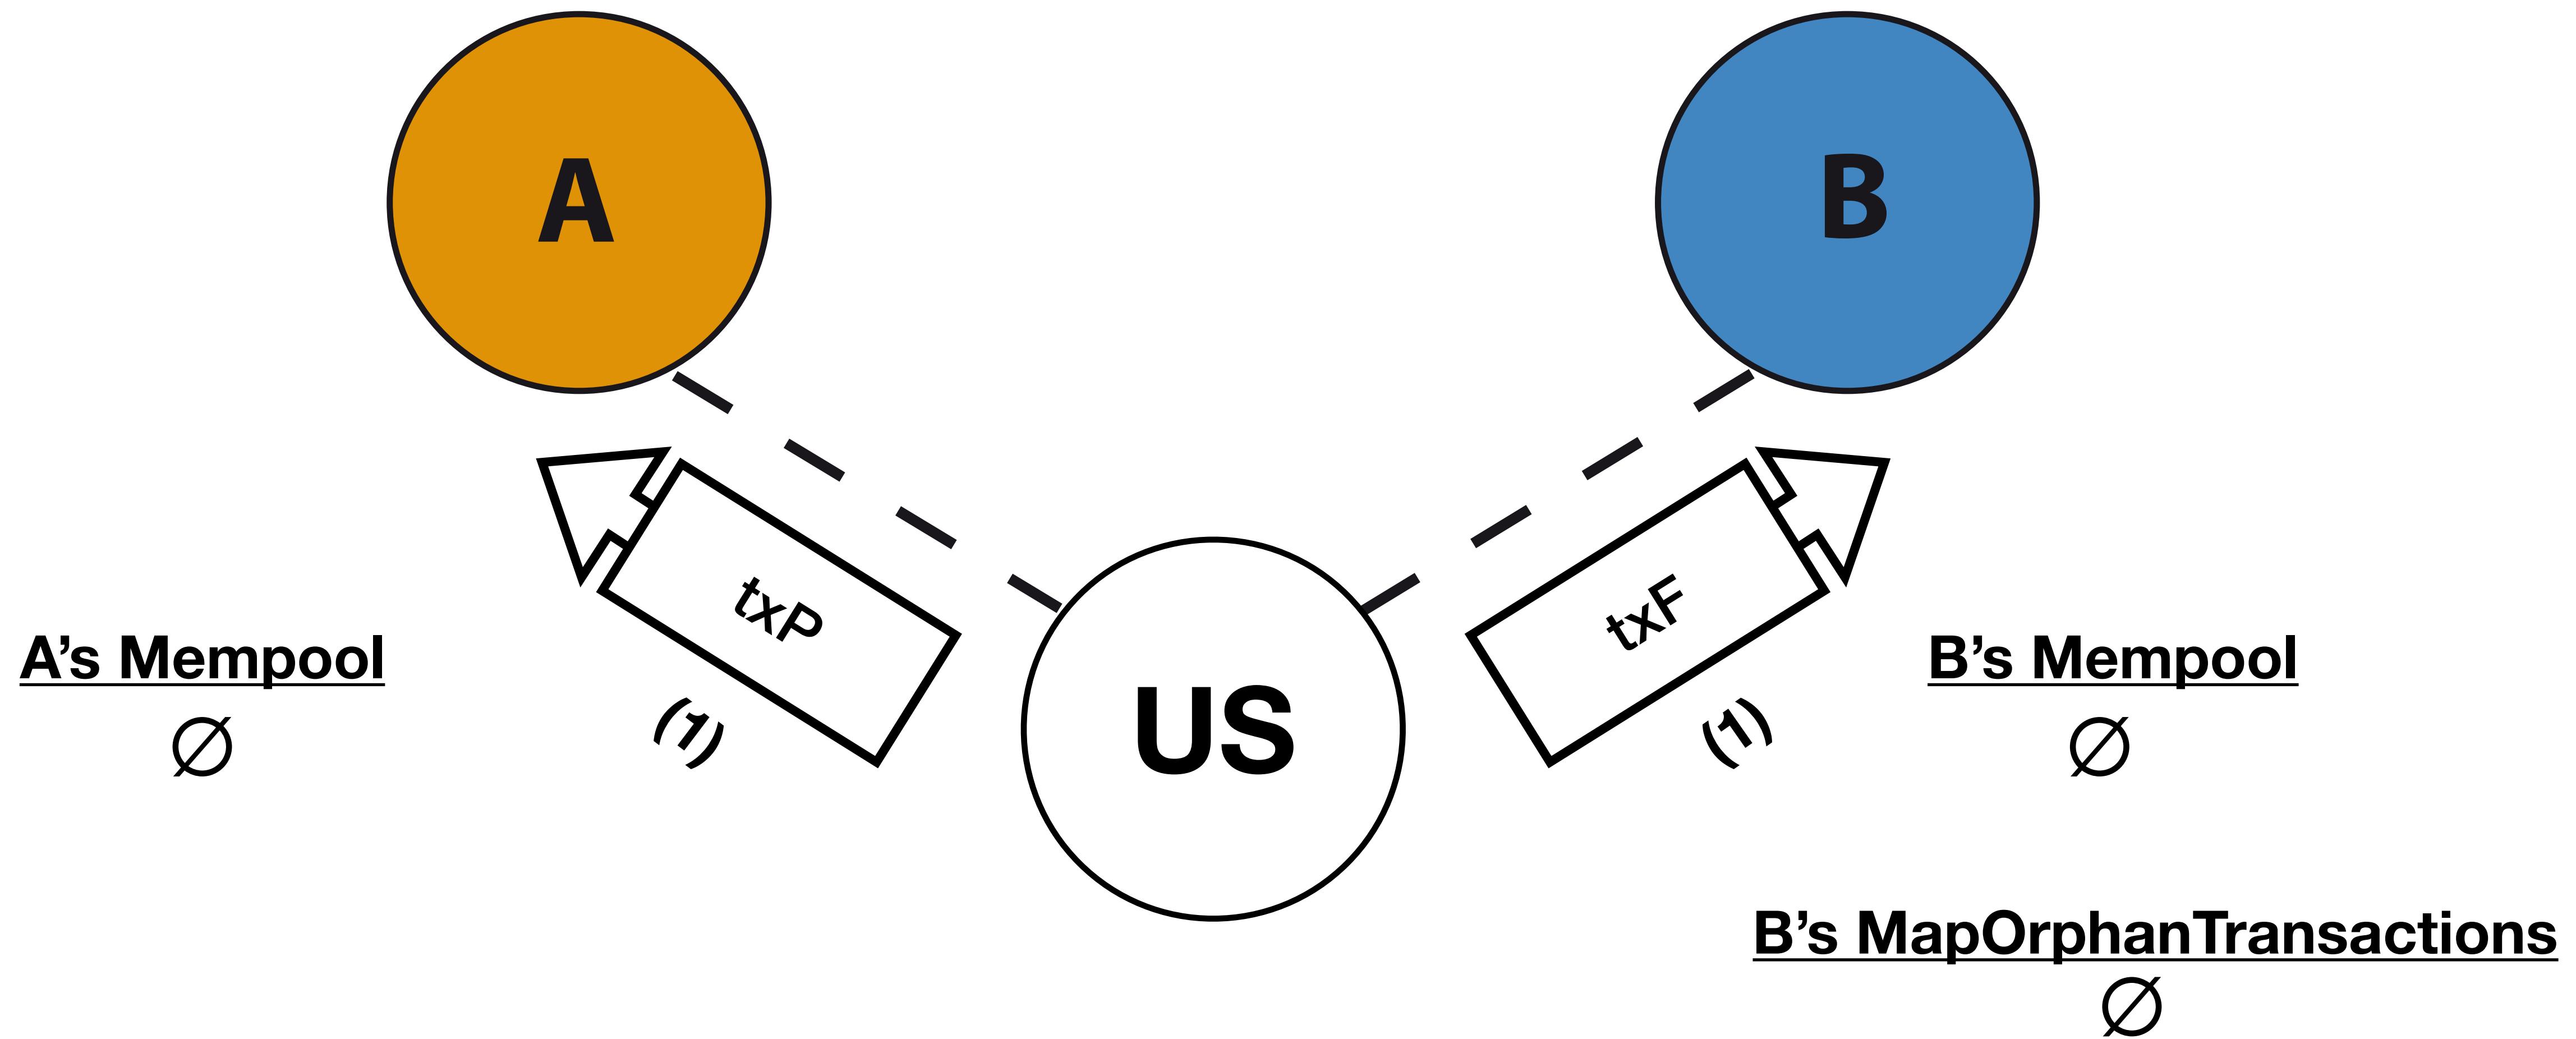
POSITIVE INFERRING TECHNIQUE



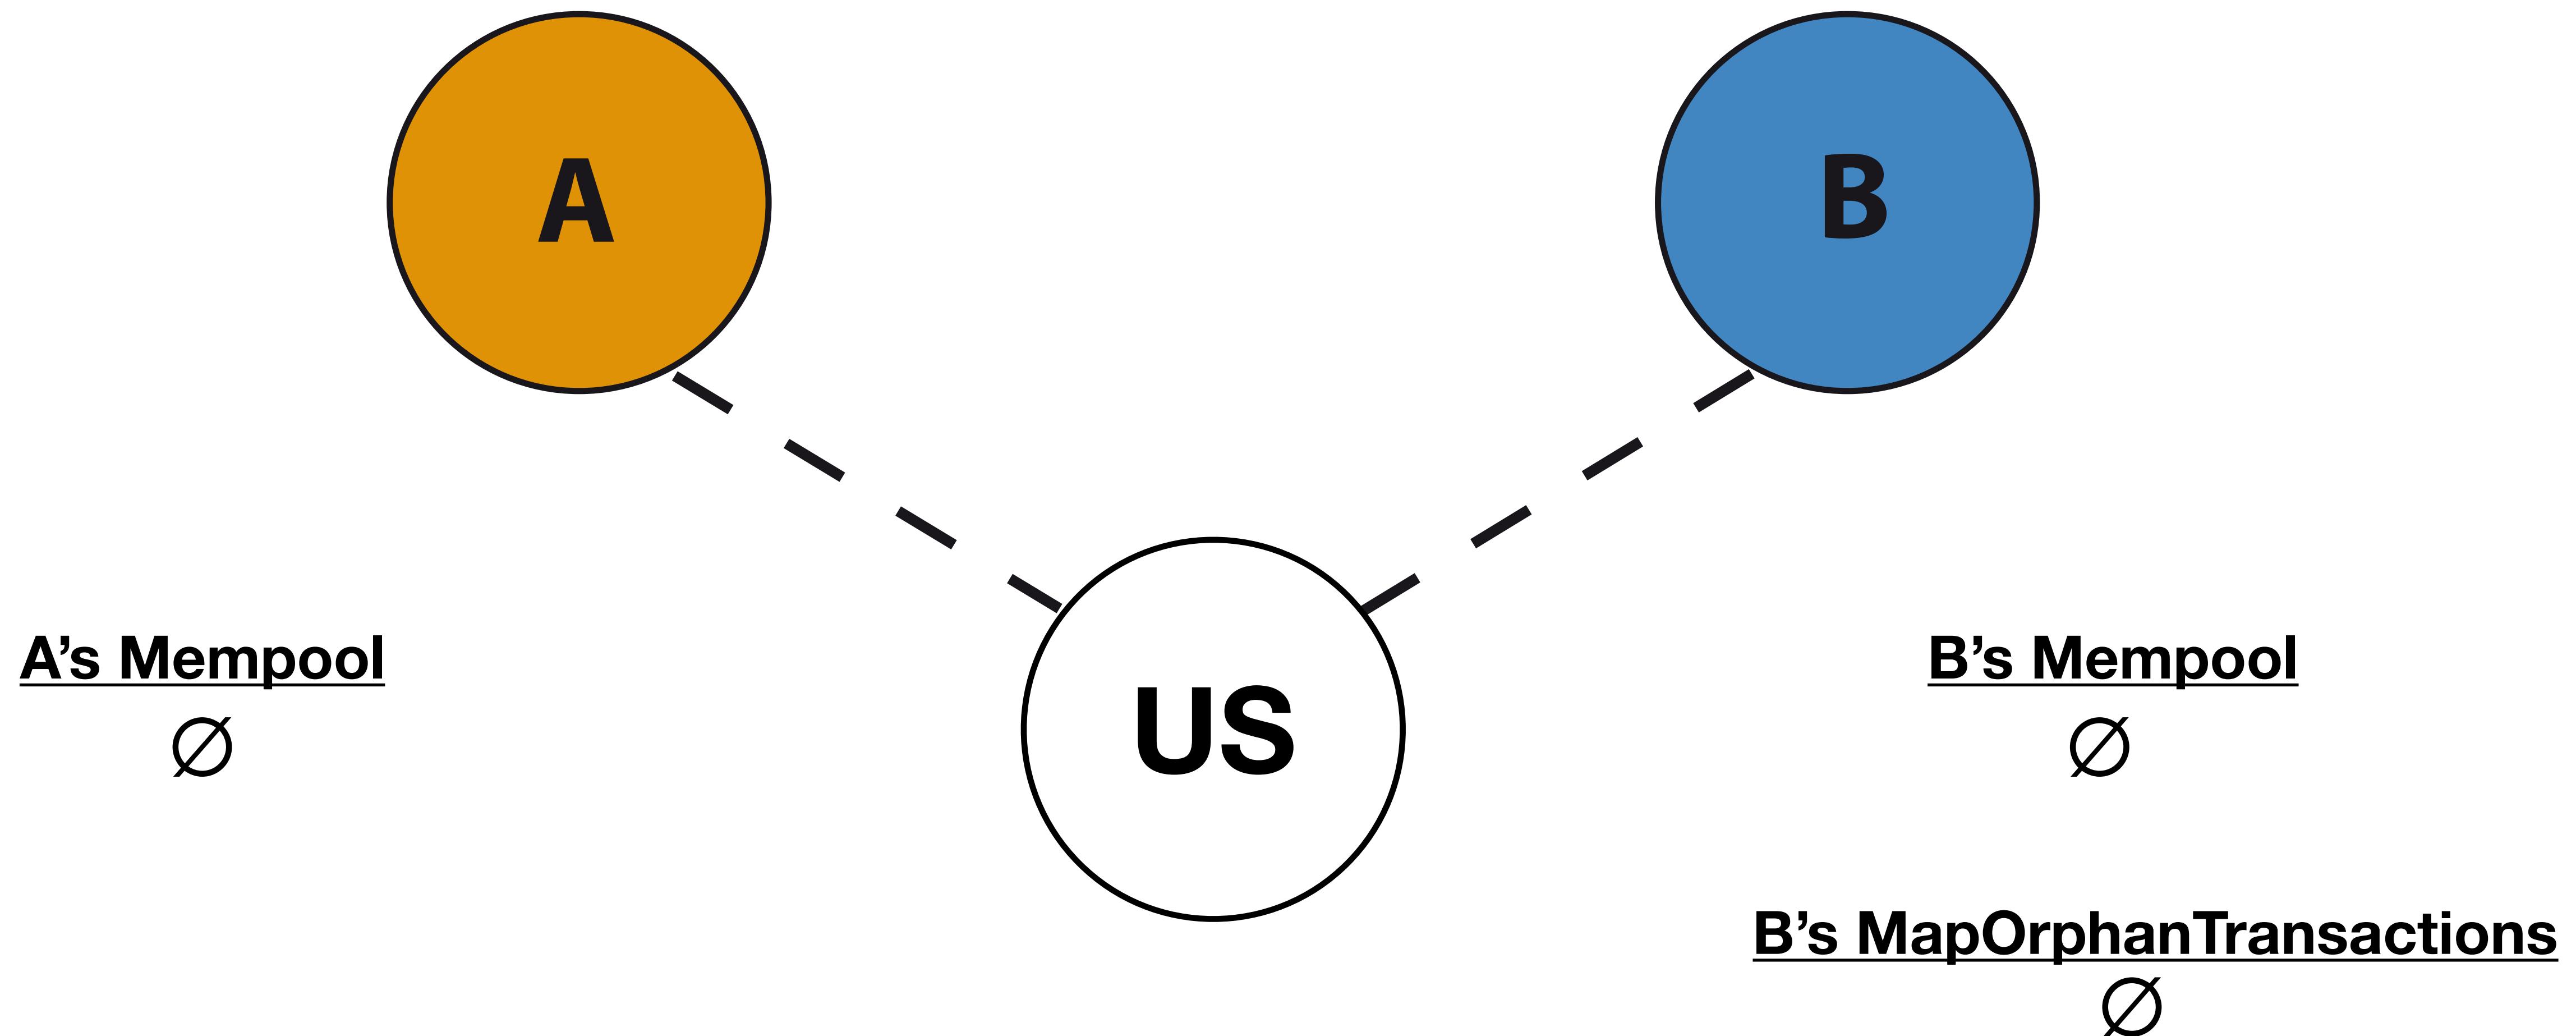
NEGATIVE INFERRING TECHNIQUE



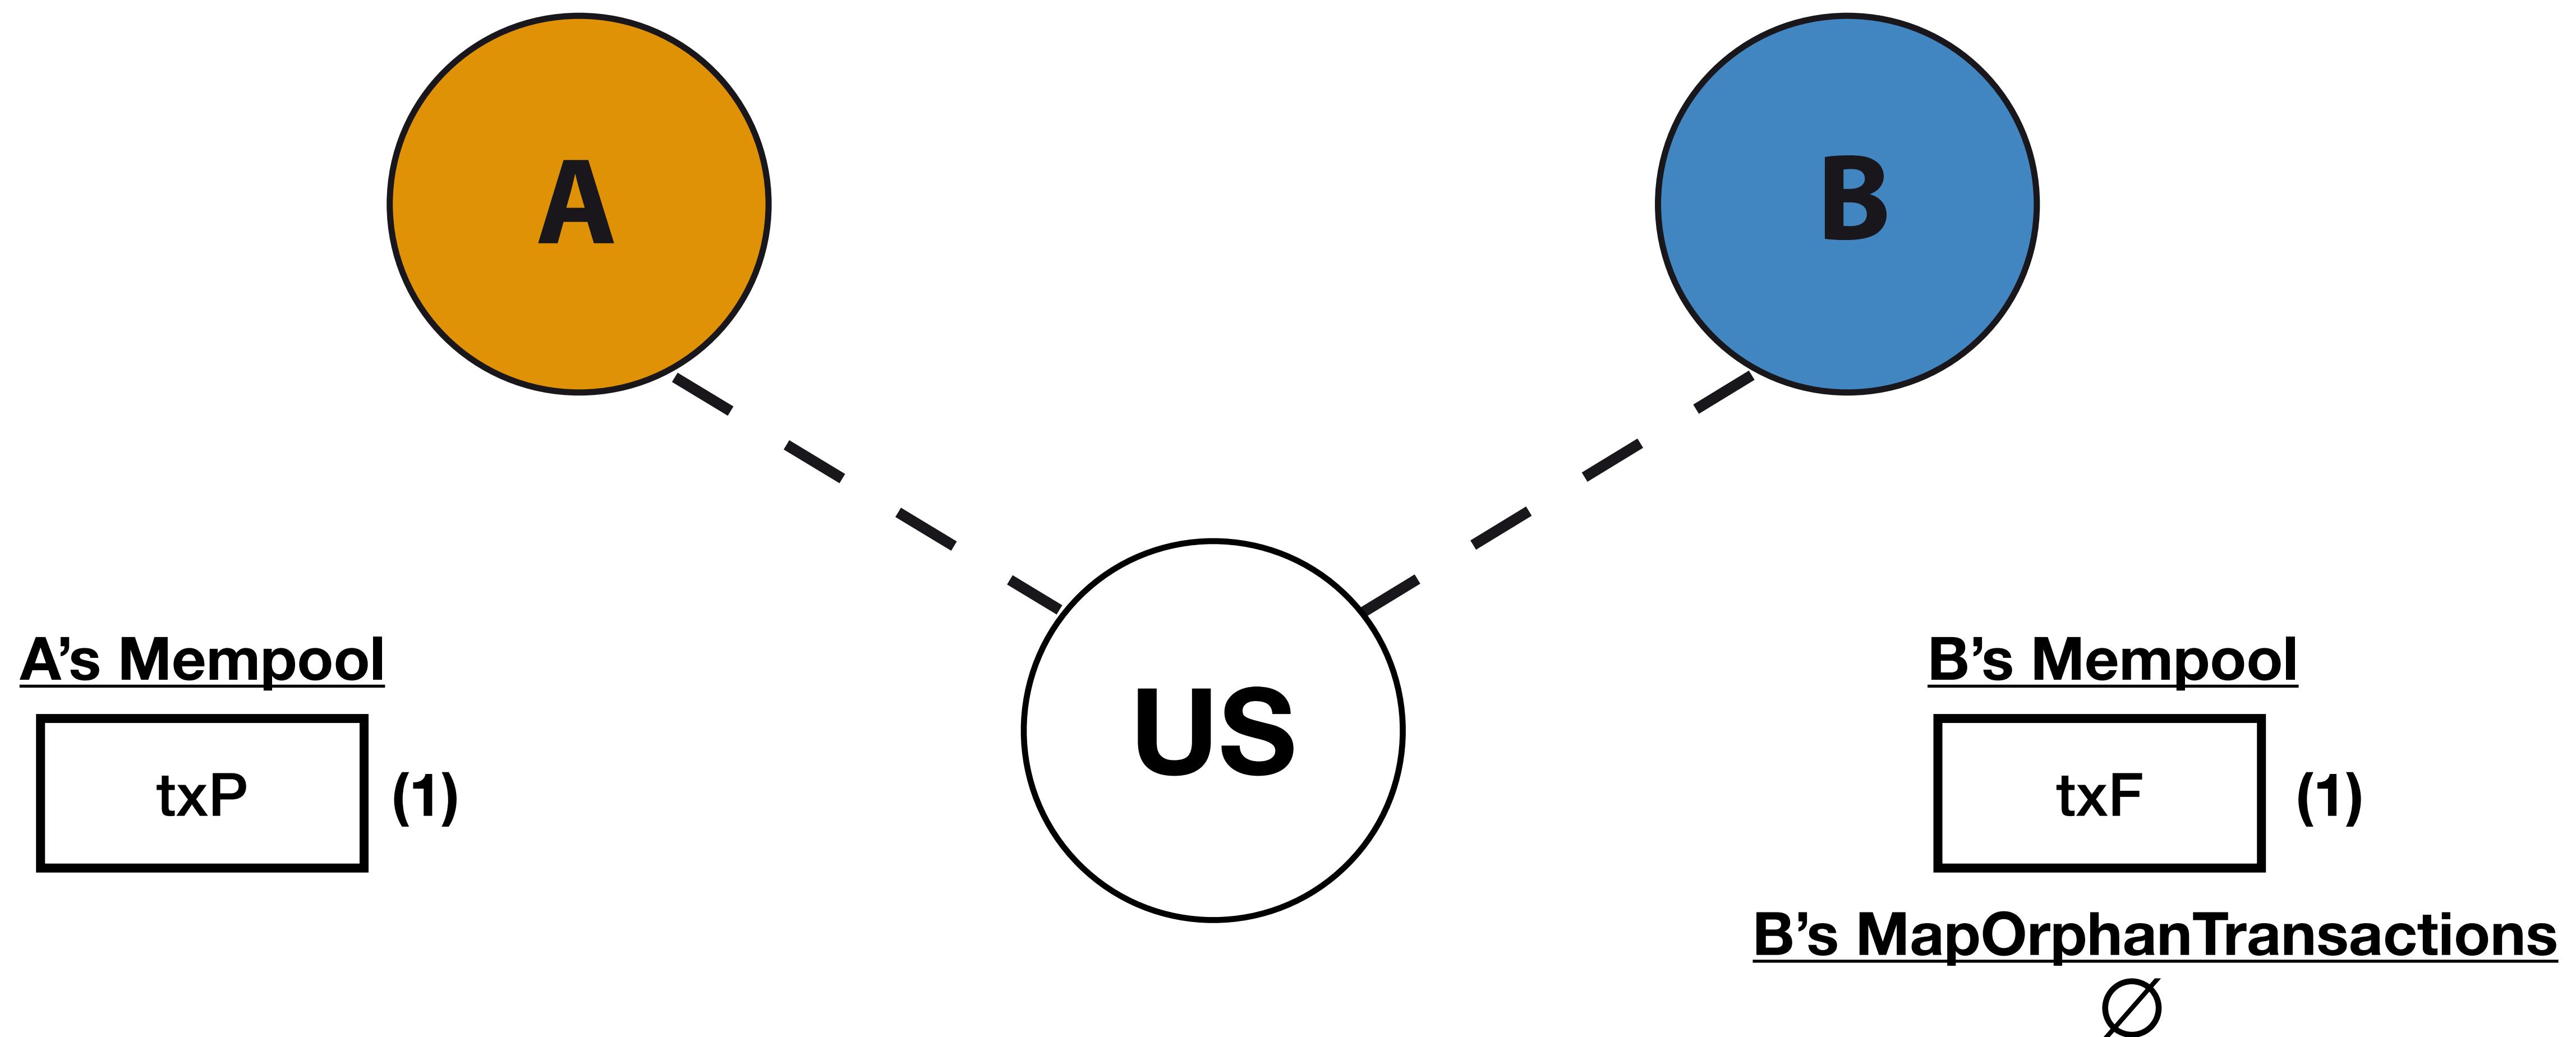
NEGATIVE INFERRING TECHNIQUE



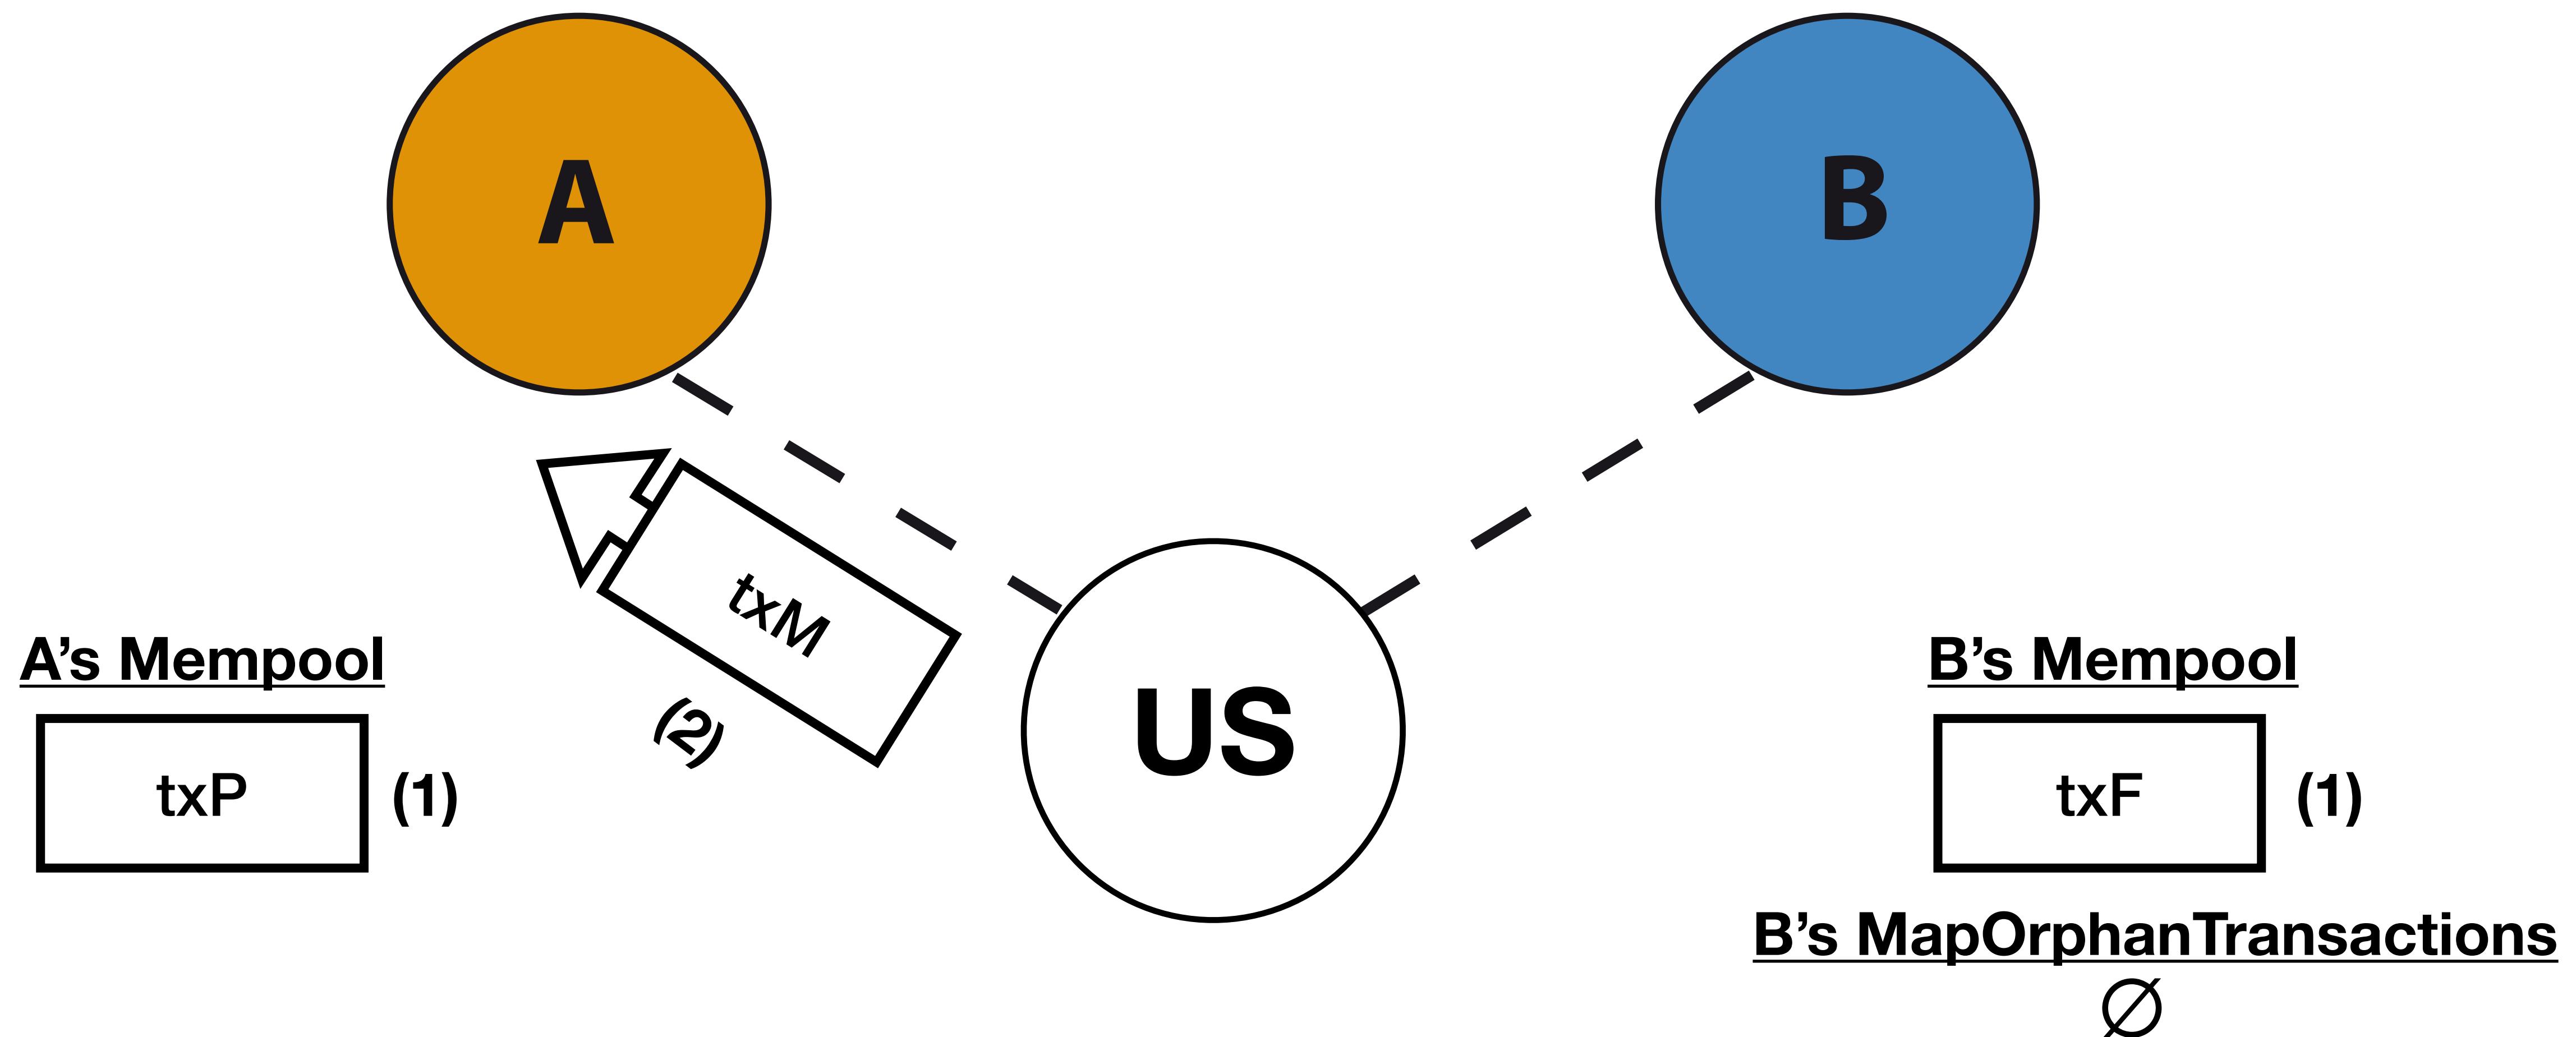
NEGATIVE INFERRING TECHNIQUE



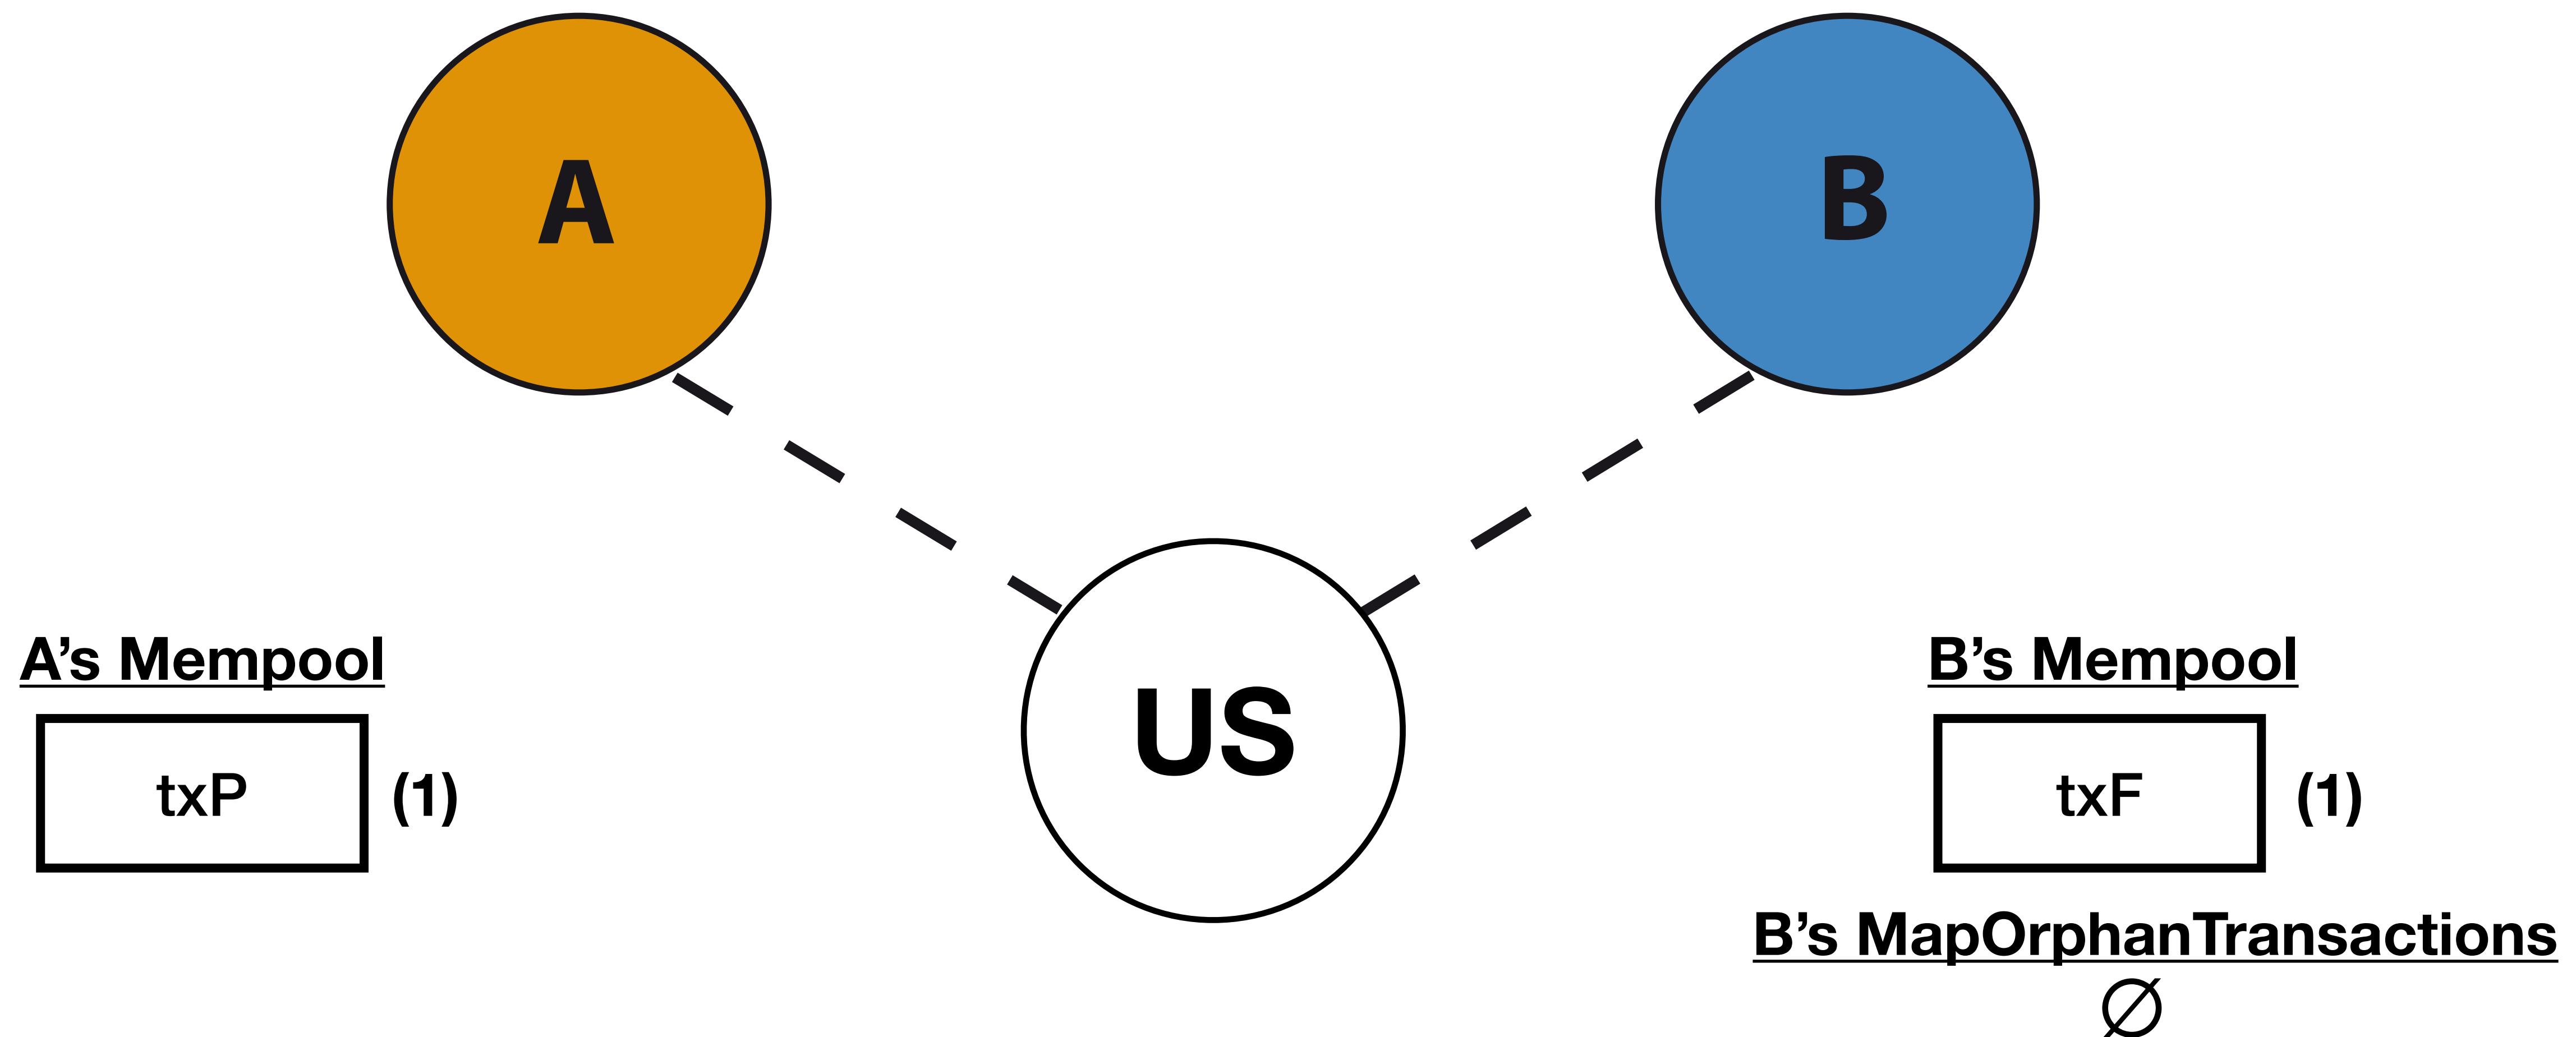
NEGATIVE INFERRING TECHNIQUE



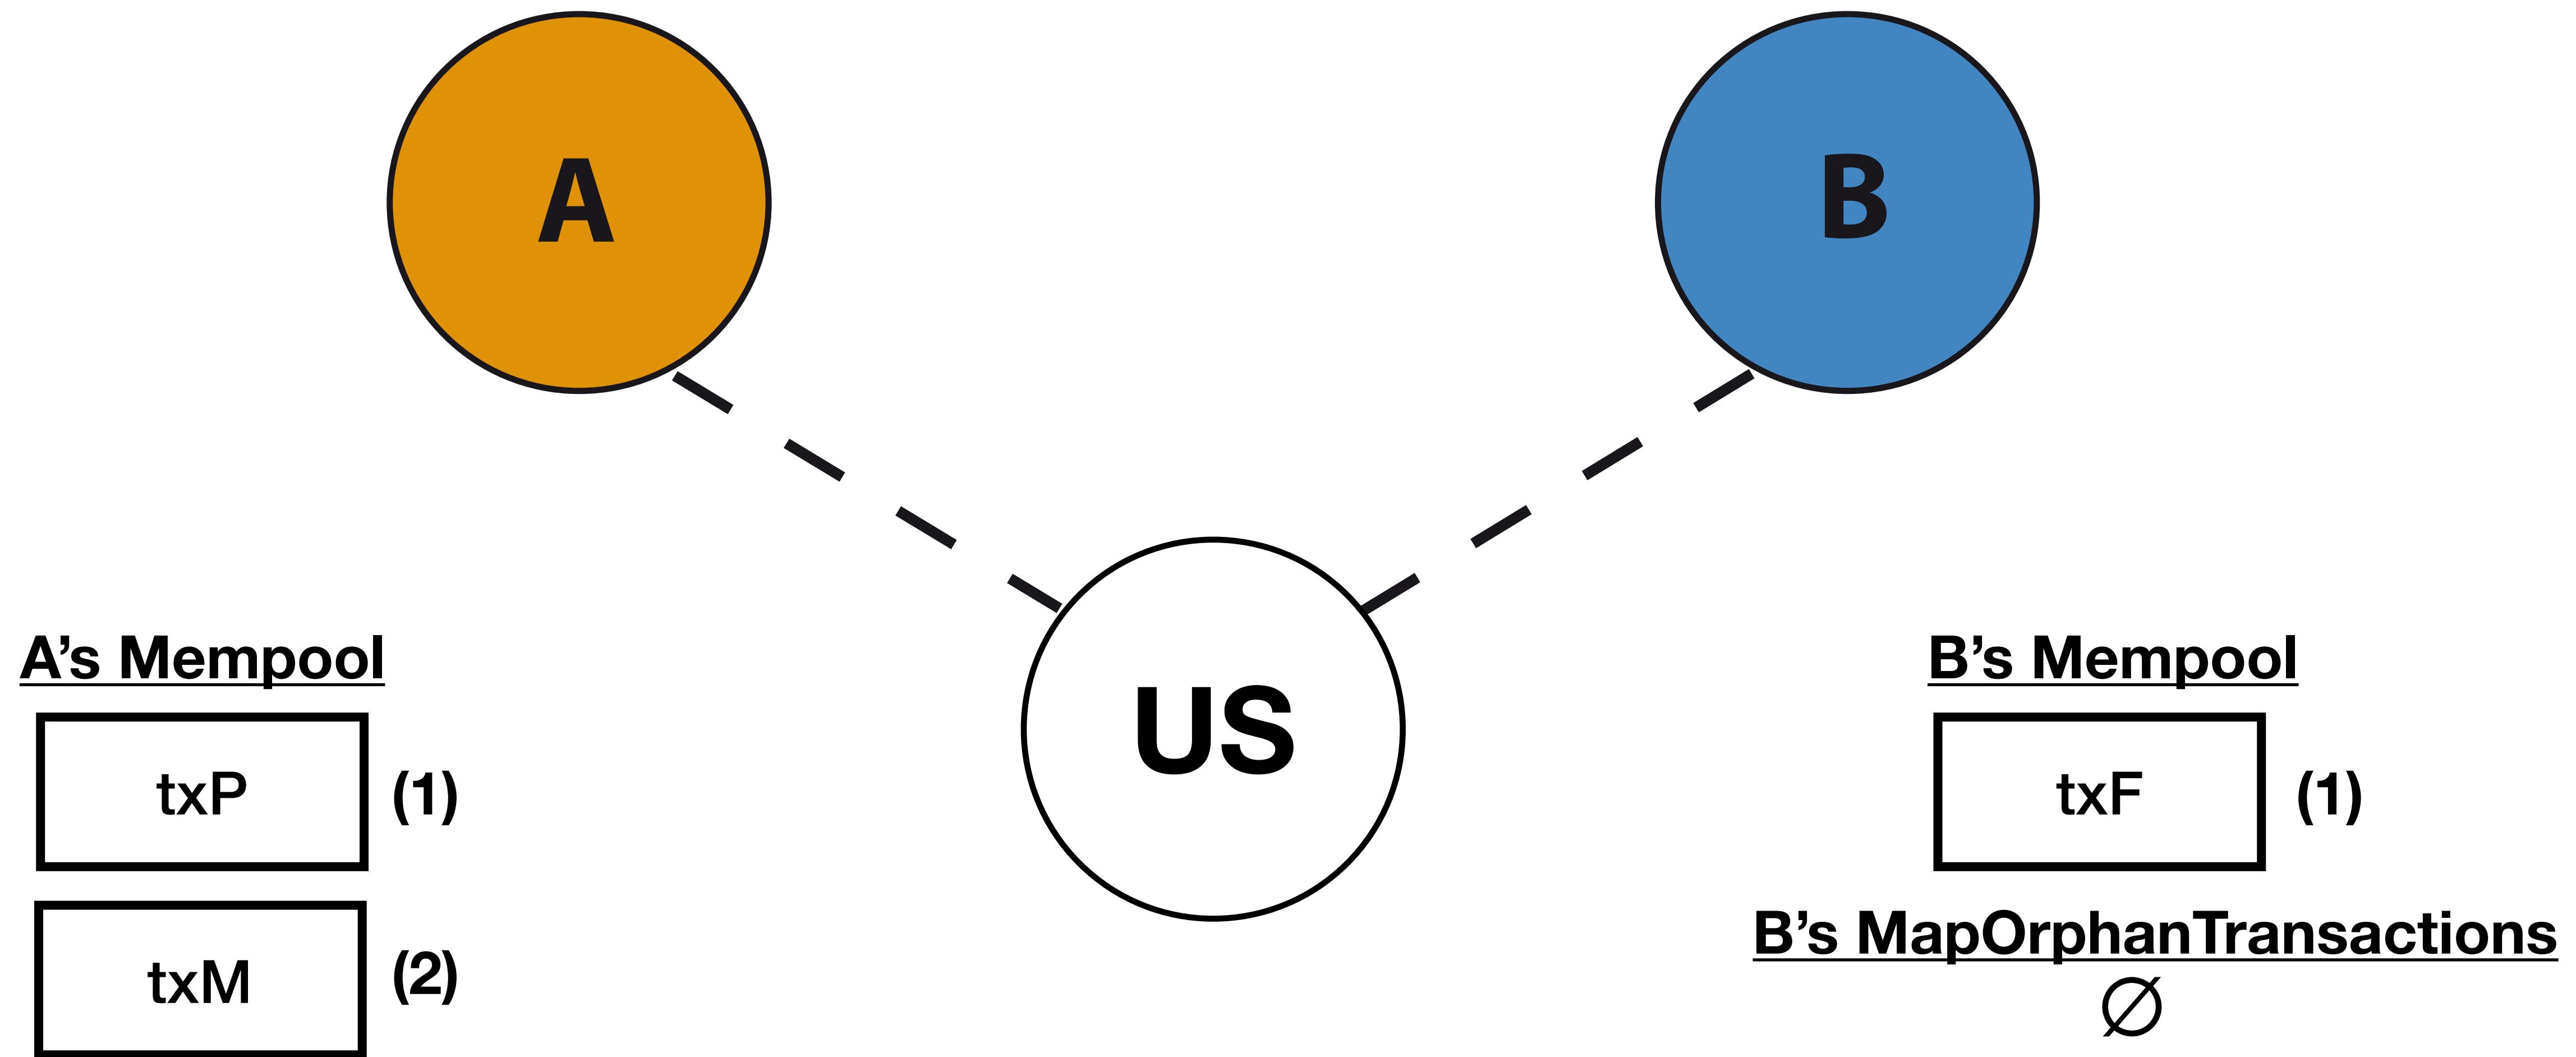
NEGATIVE INFERRING TECHNIQUE



NEGATIVE INFERRING TECHNIQUE

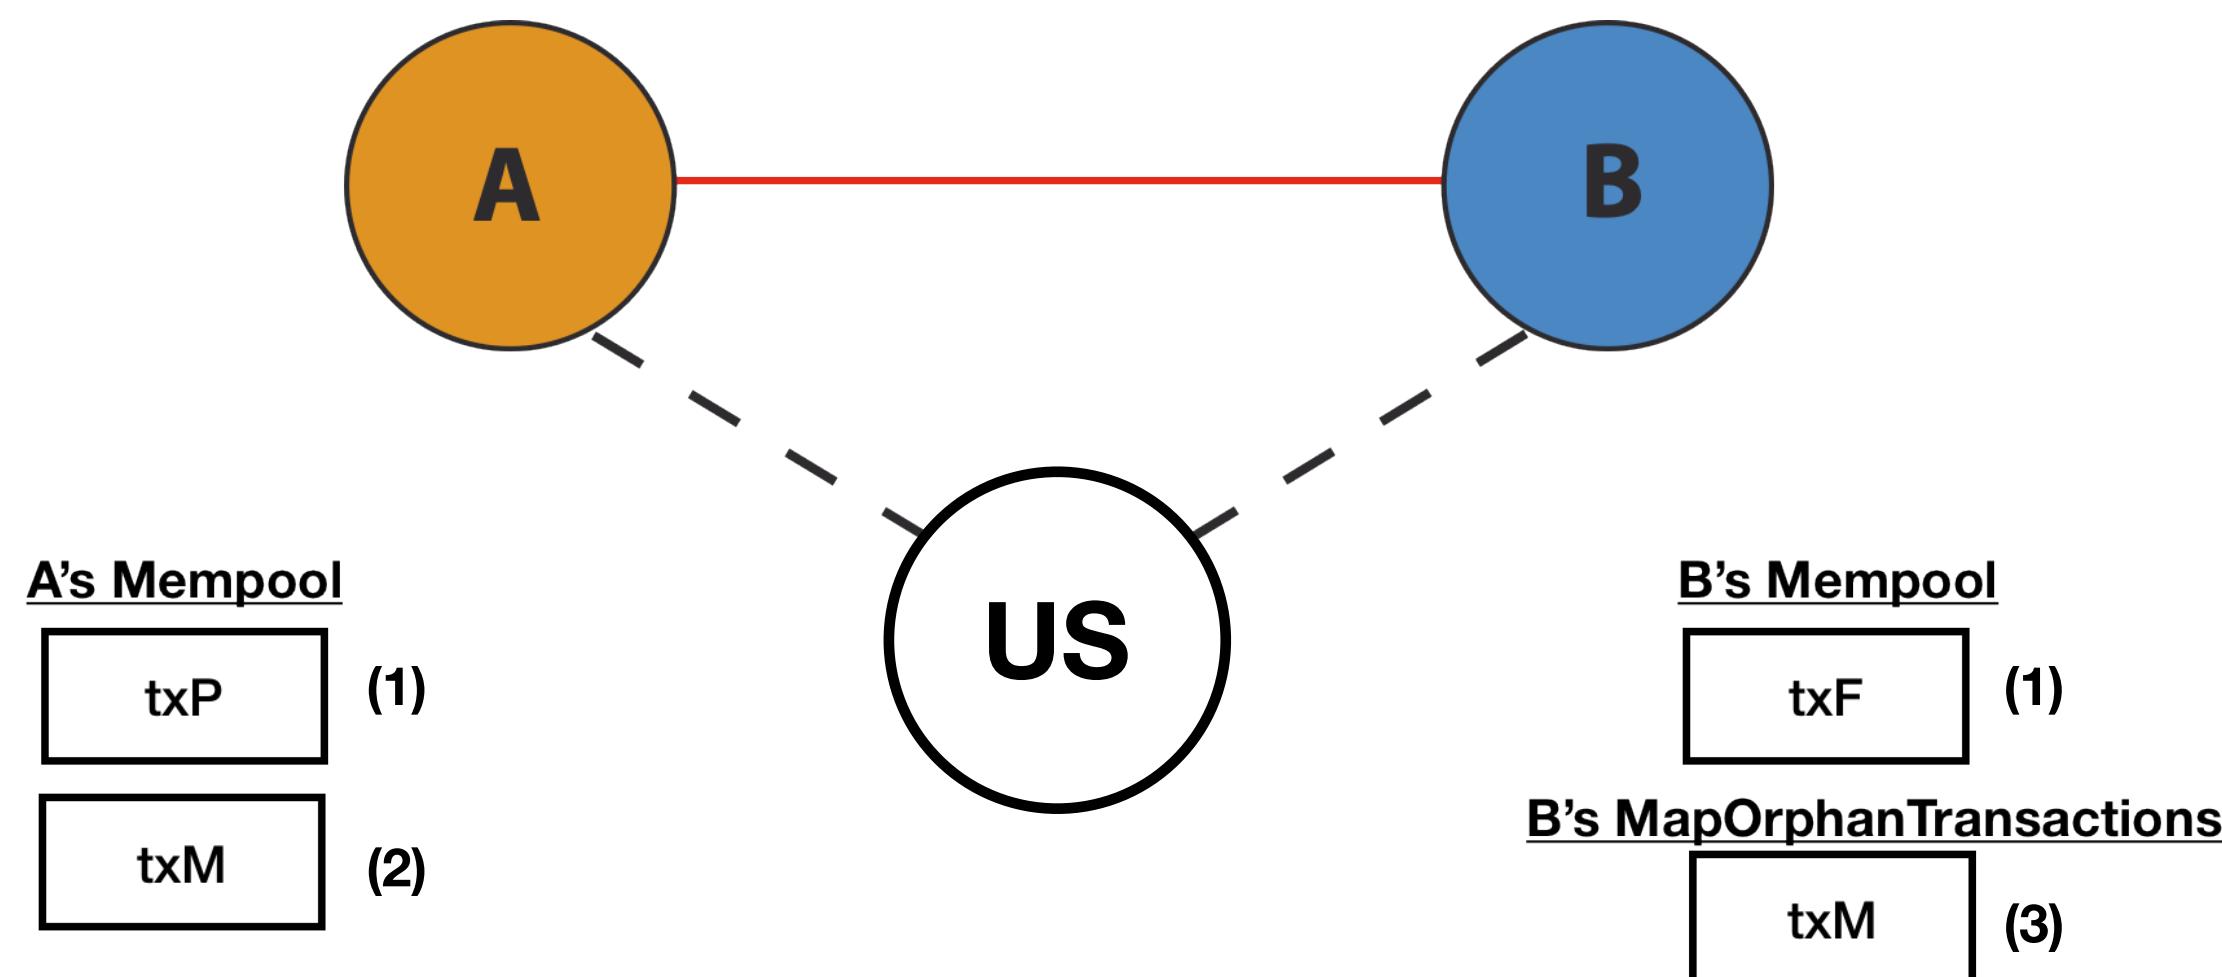


NEGATIVE INFERRING TECHNIQUE

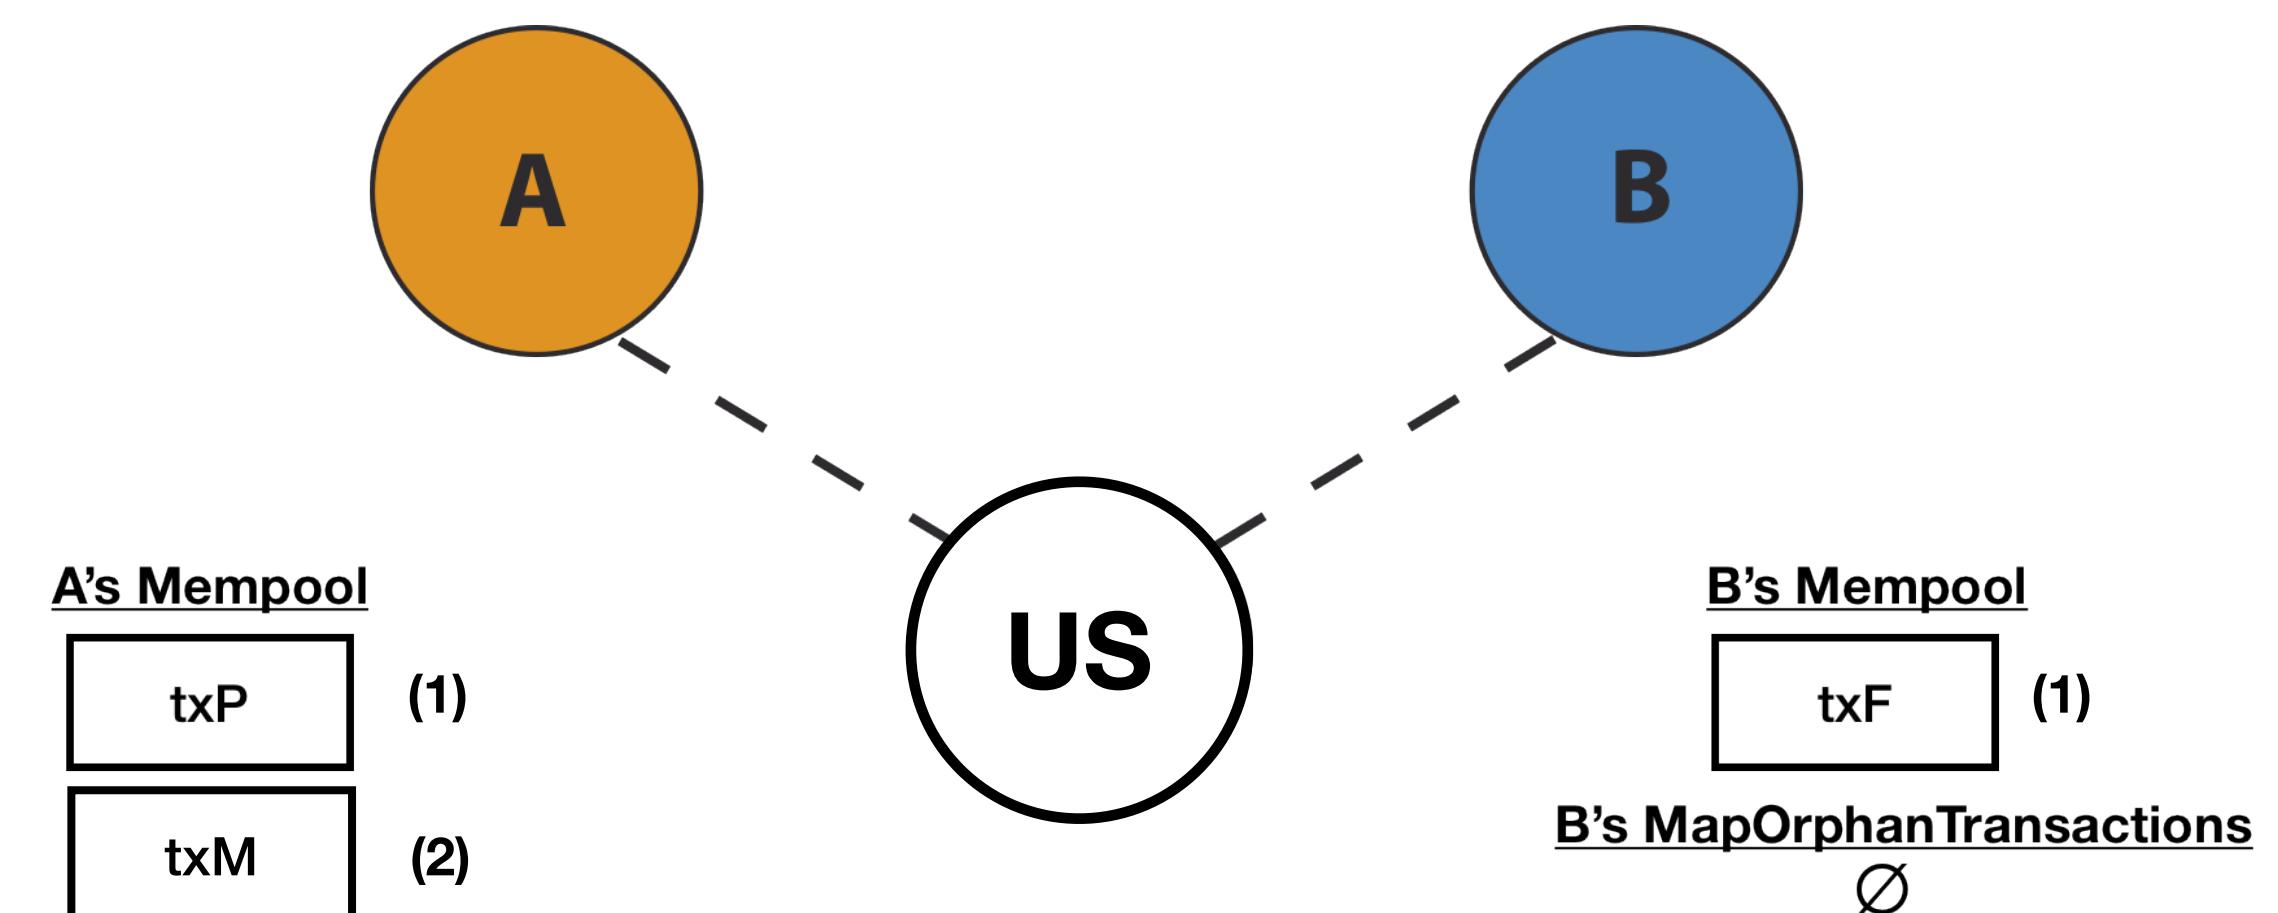


A BASIC TOPOLOGY INFERRING TECHNIQUE

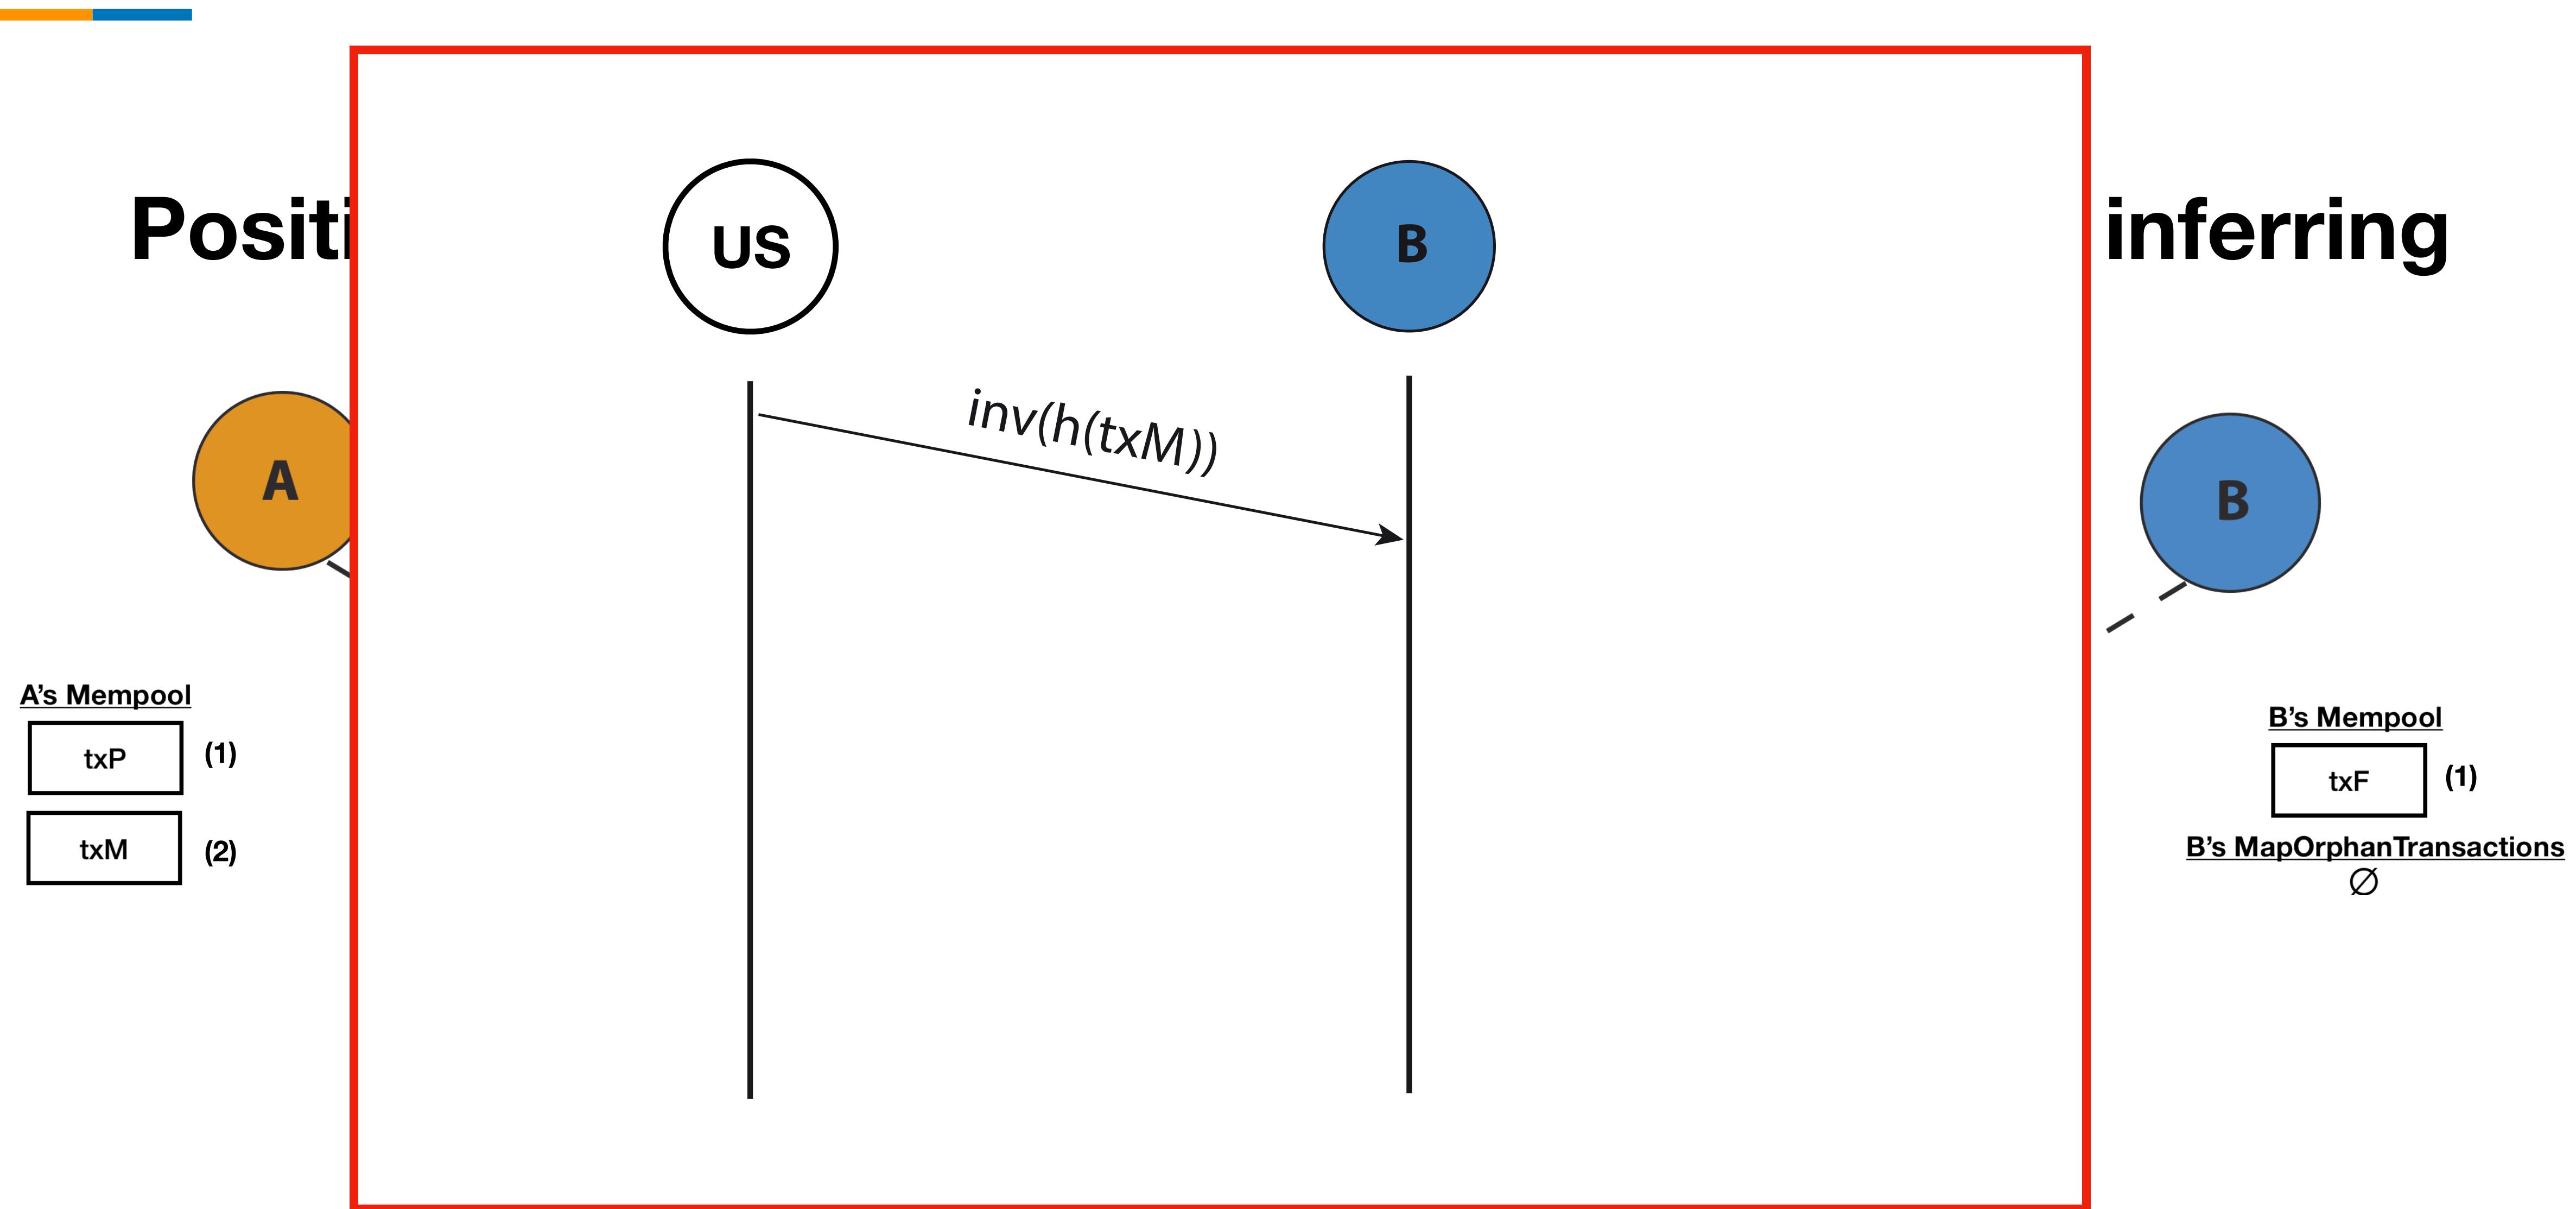
Positive edge inferring



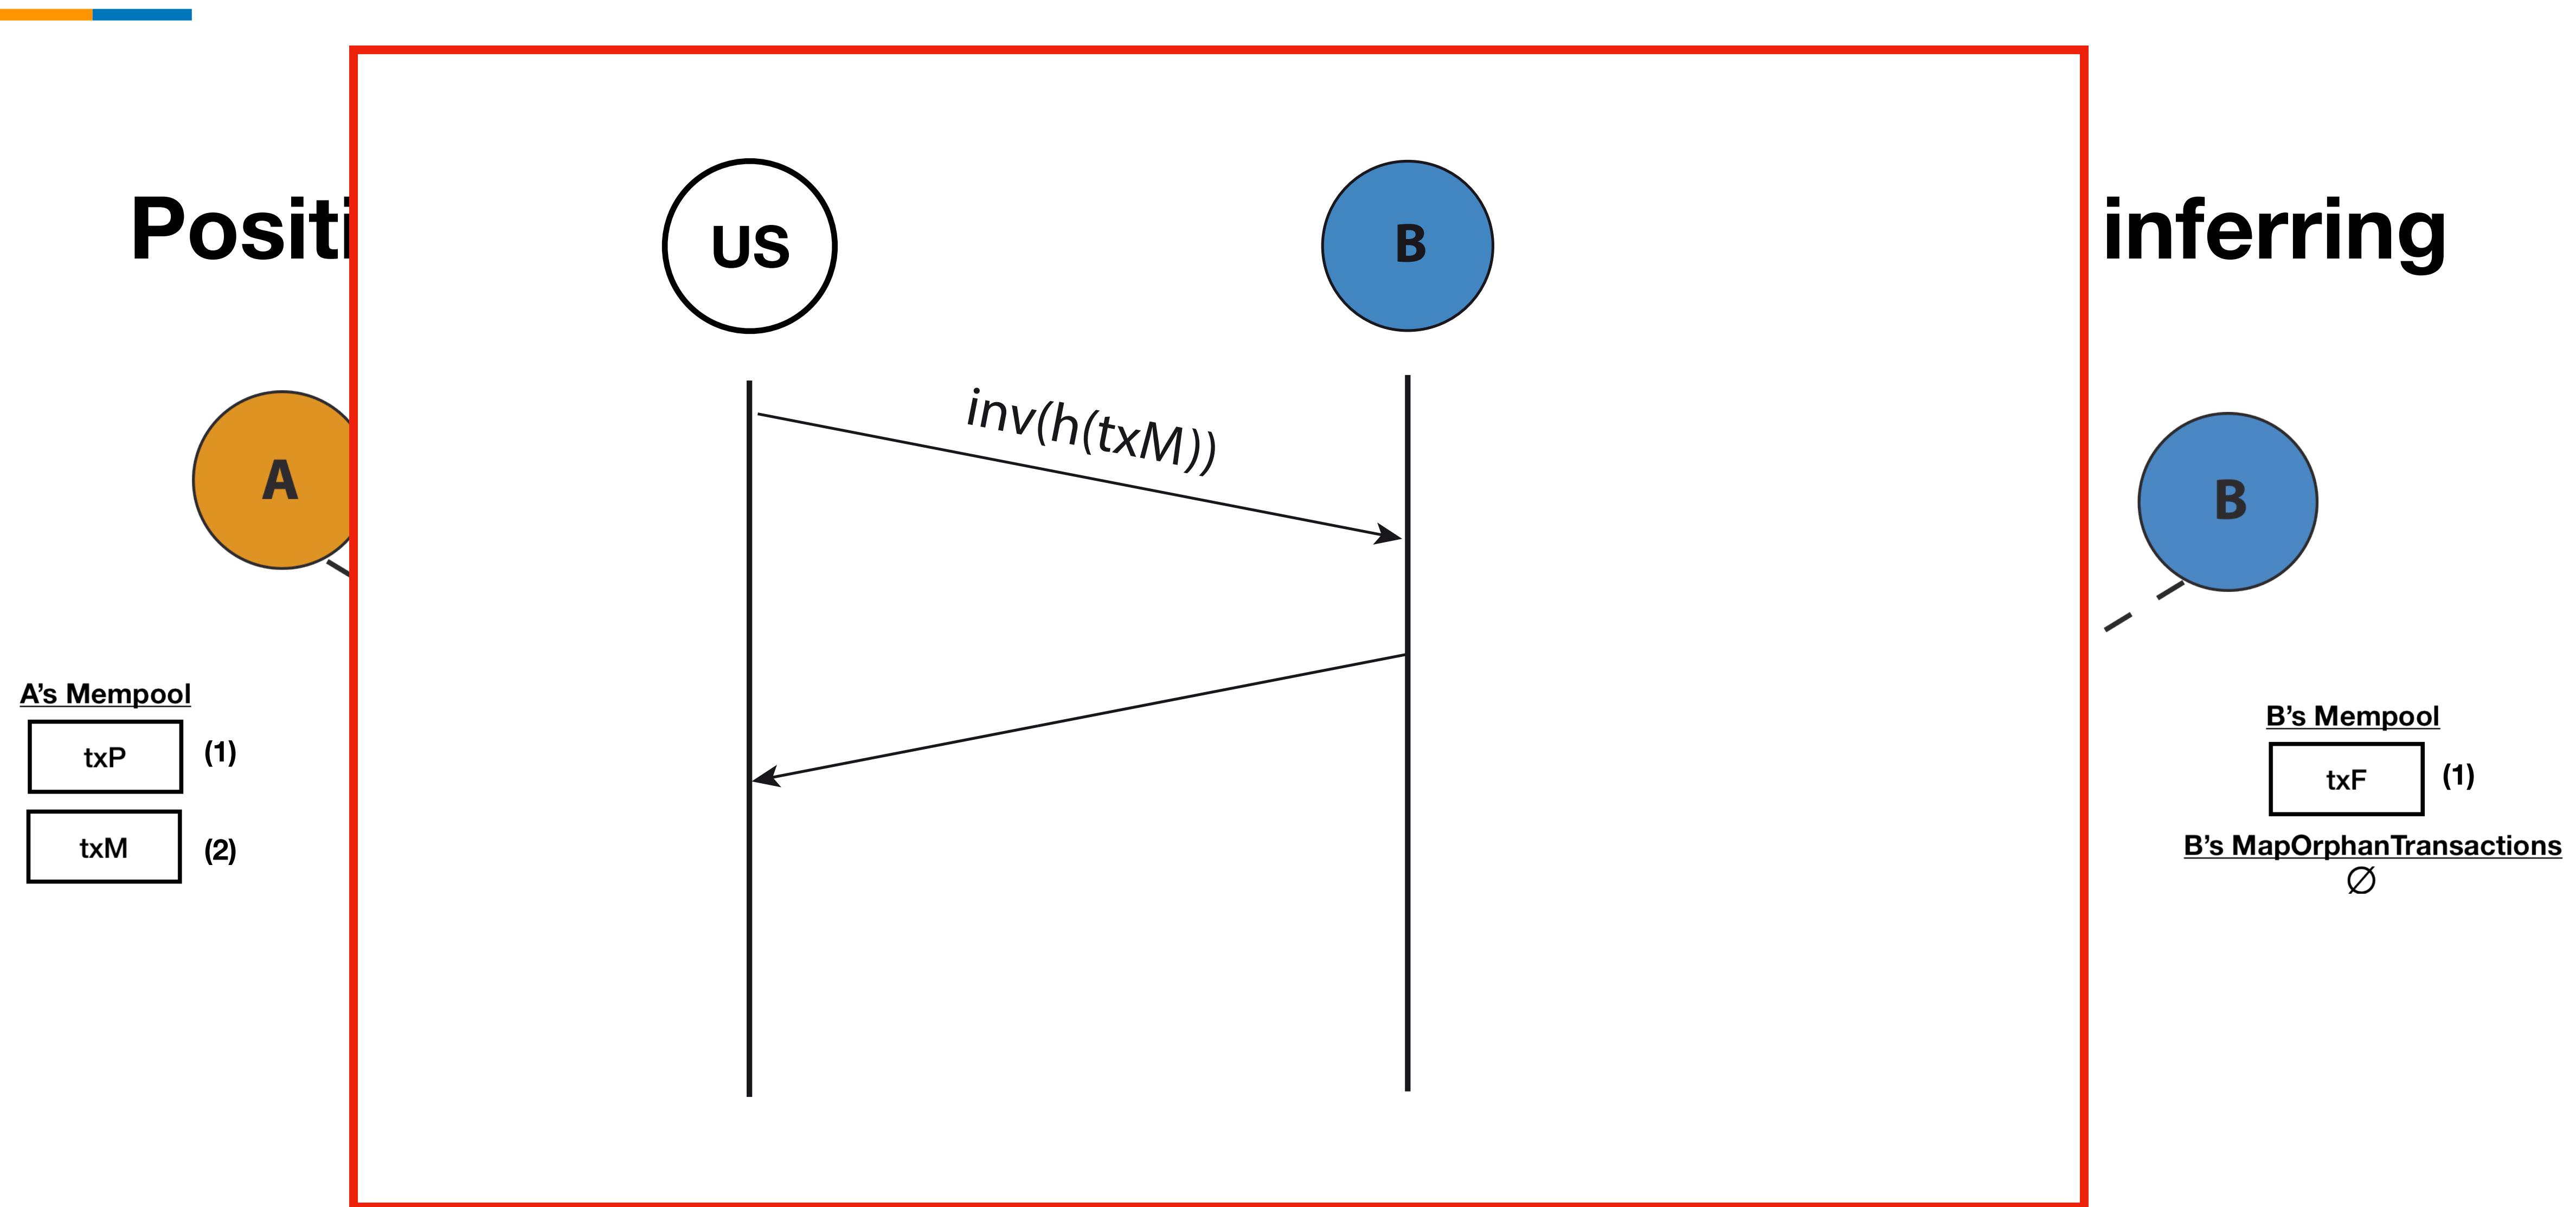
Negative edge inferring



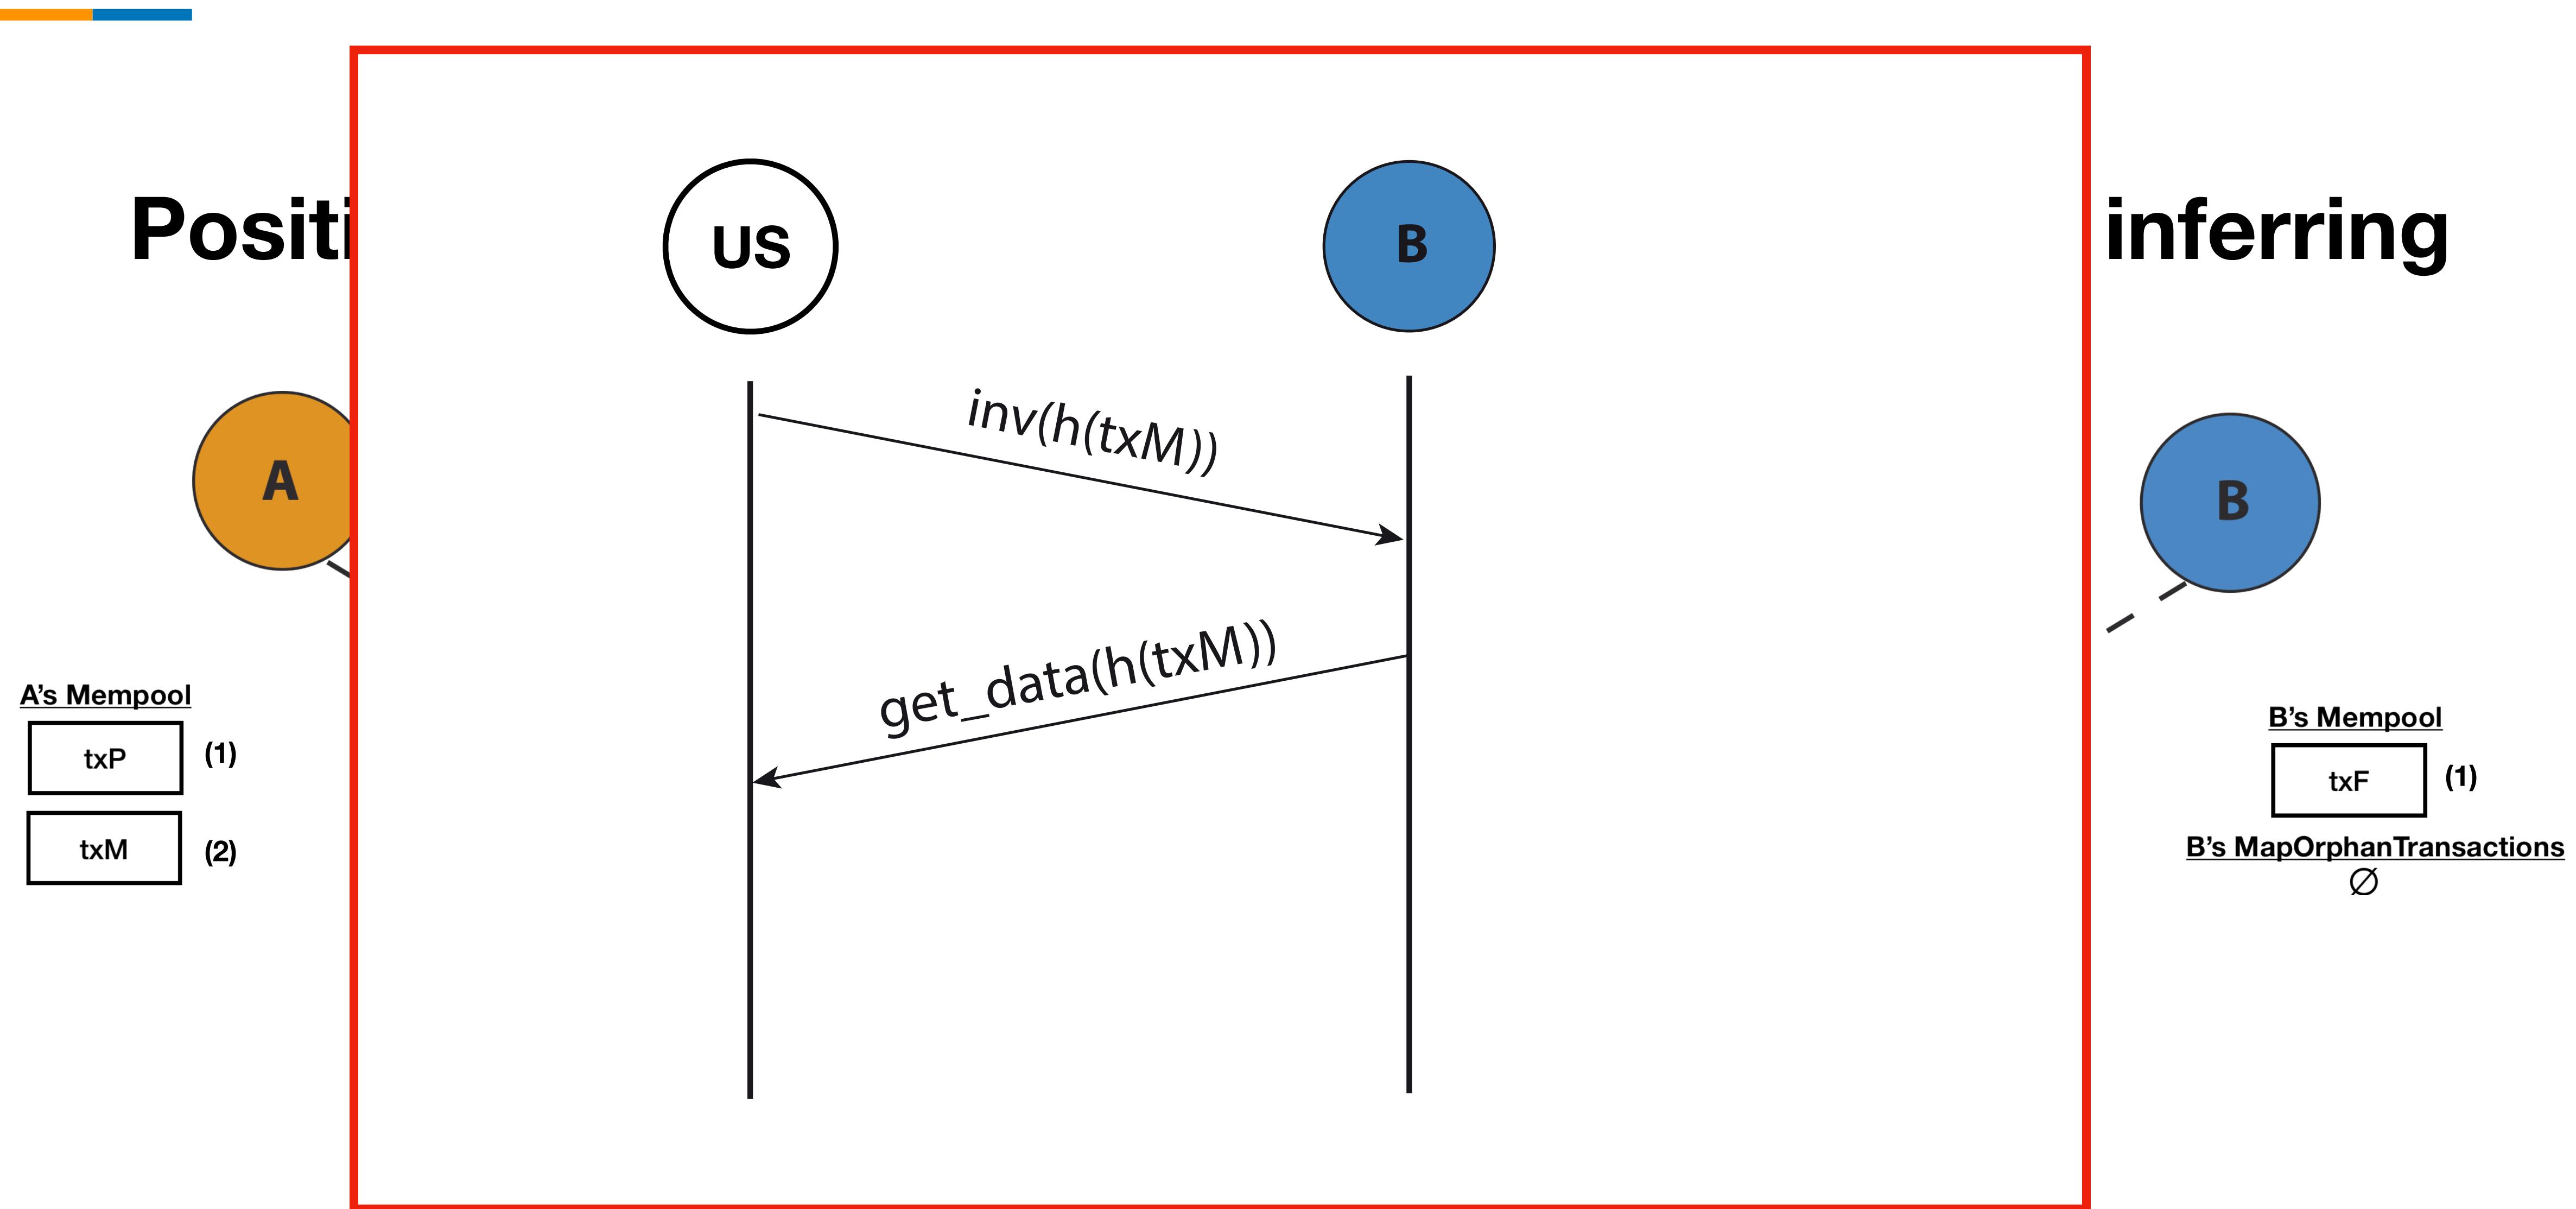
A BASIC TOPOLOGY INFERRING TECHNIQUE



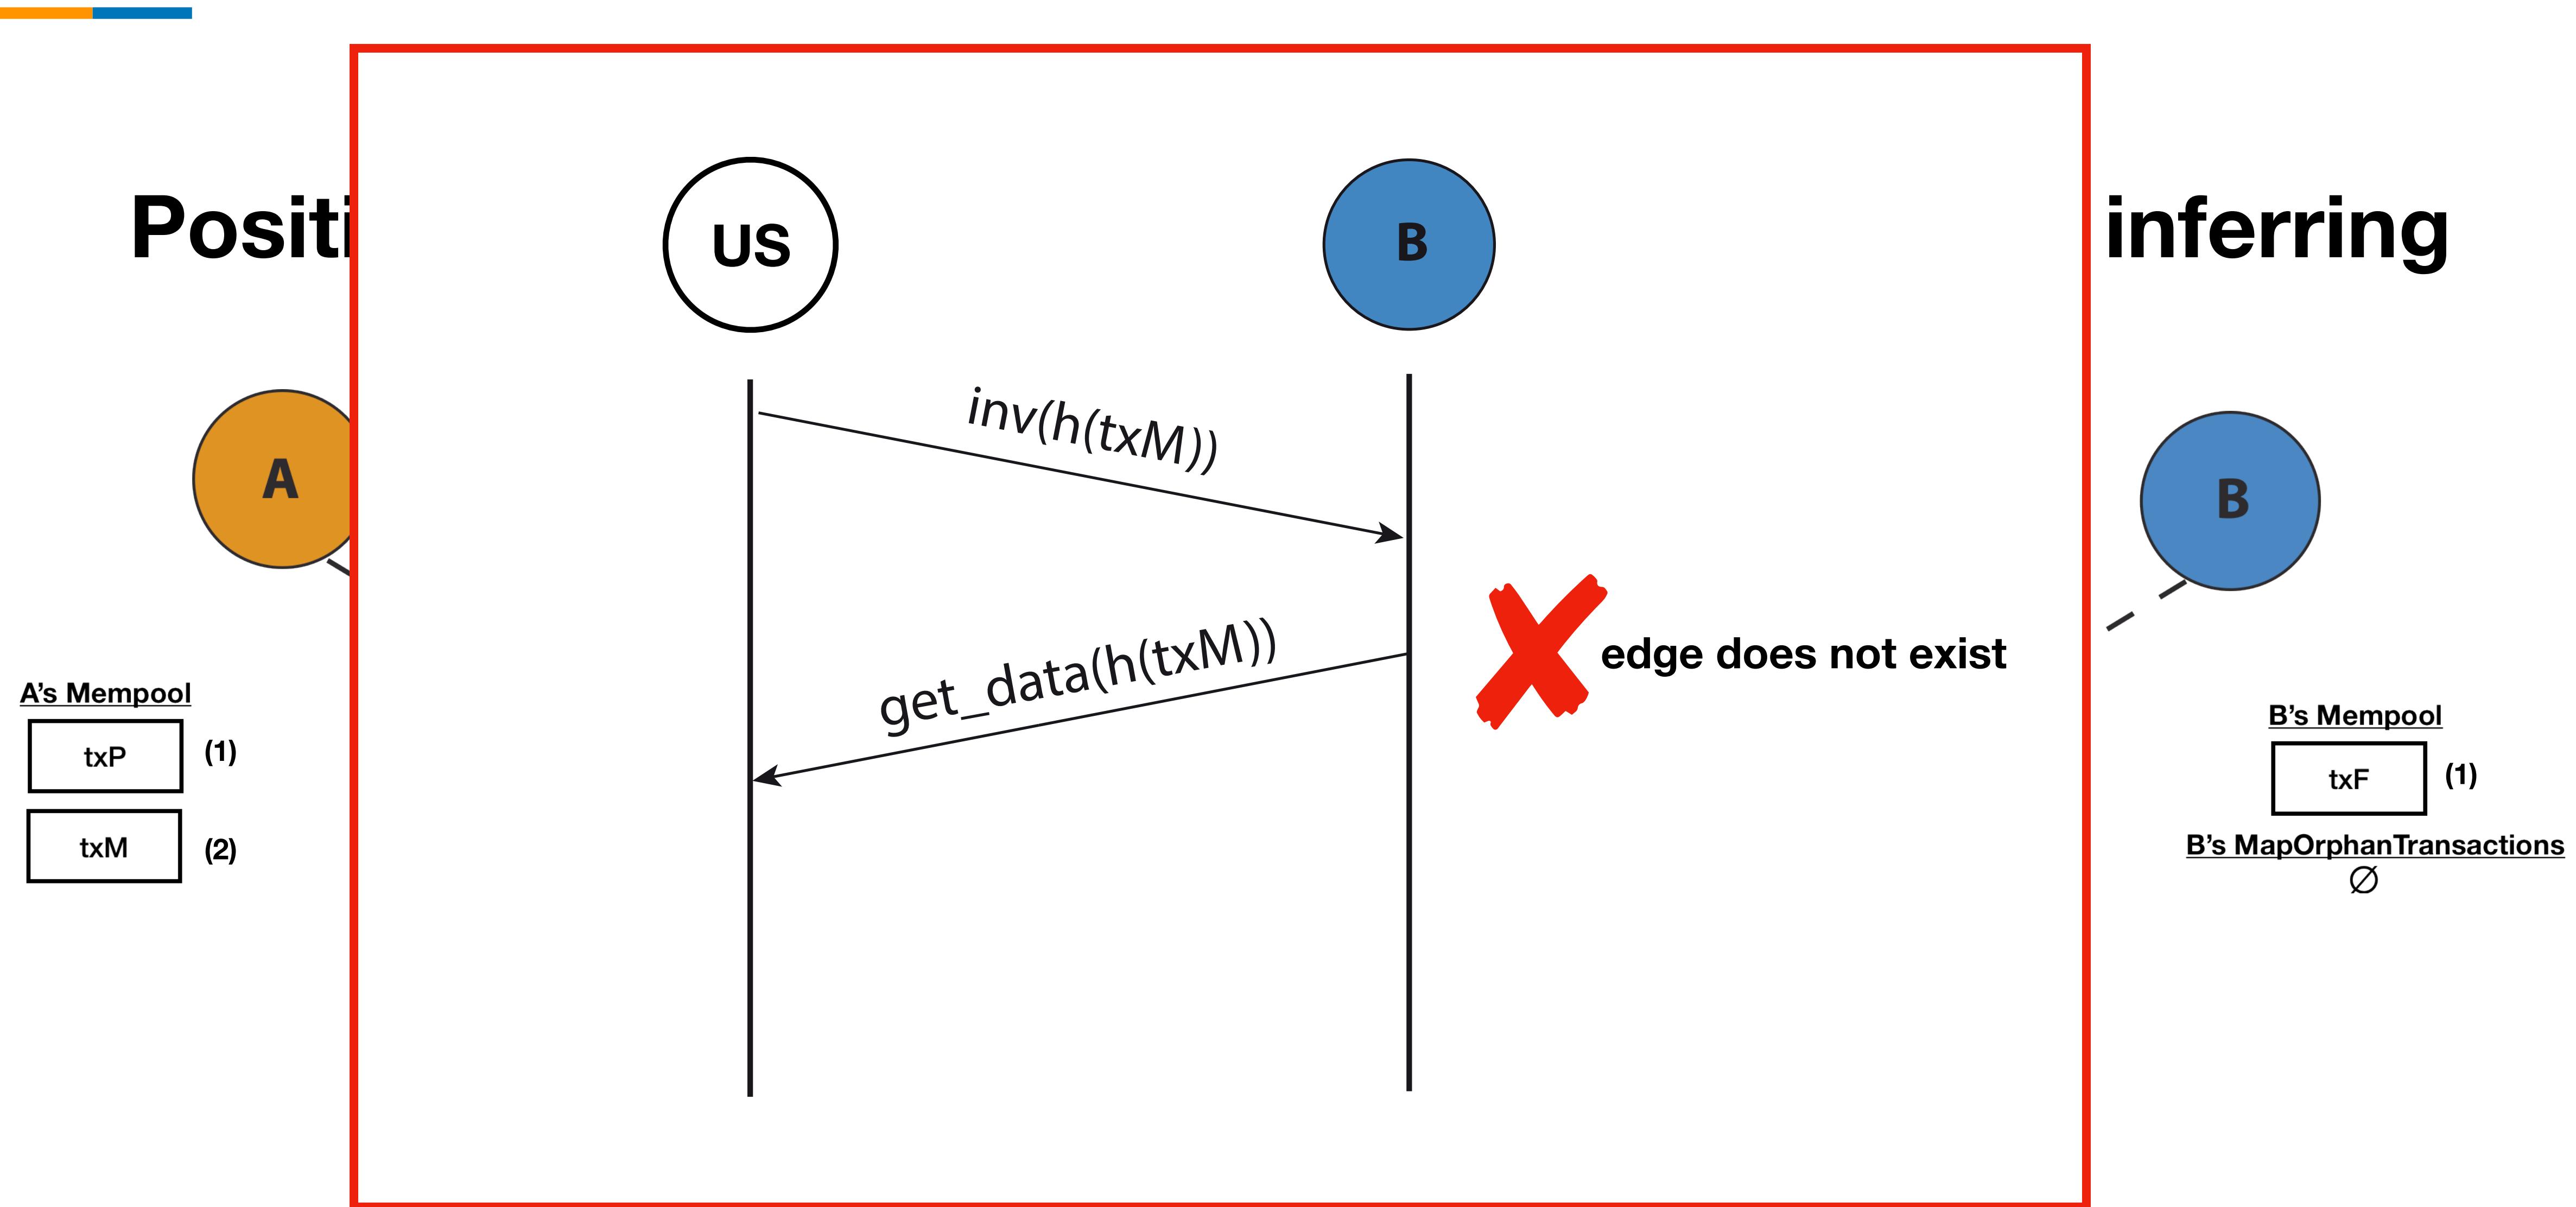
A BASIC TOPOLOGY INFERRING TECHNIQUE



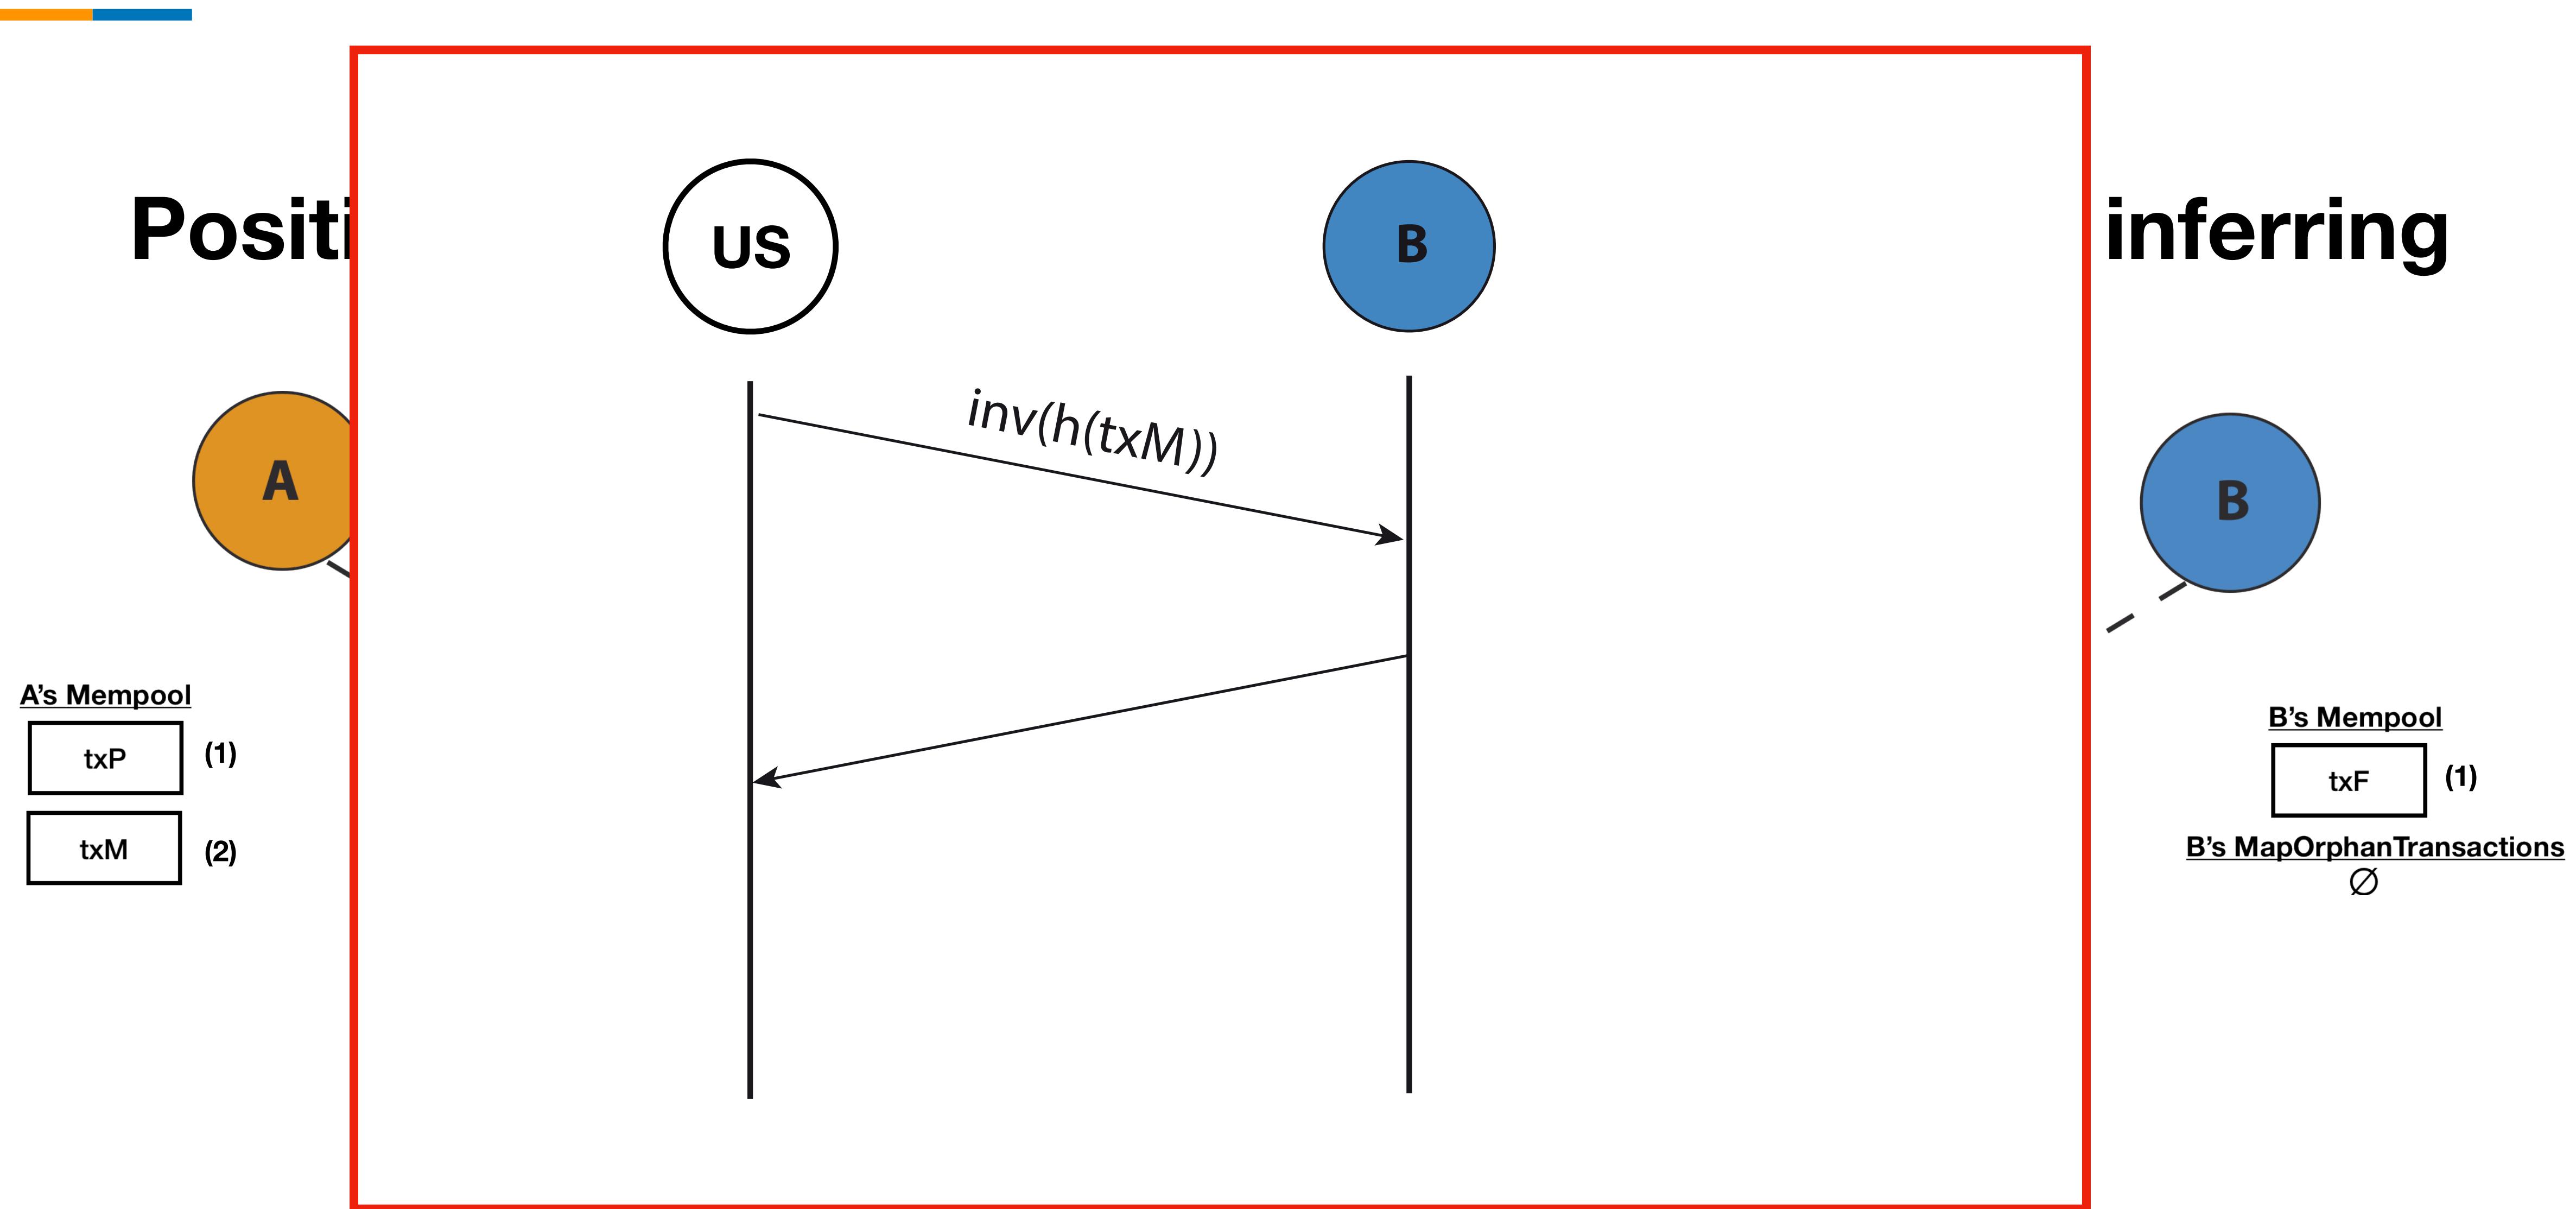
A BASIC TOPOLOGY INFERRING TECHNIQUE



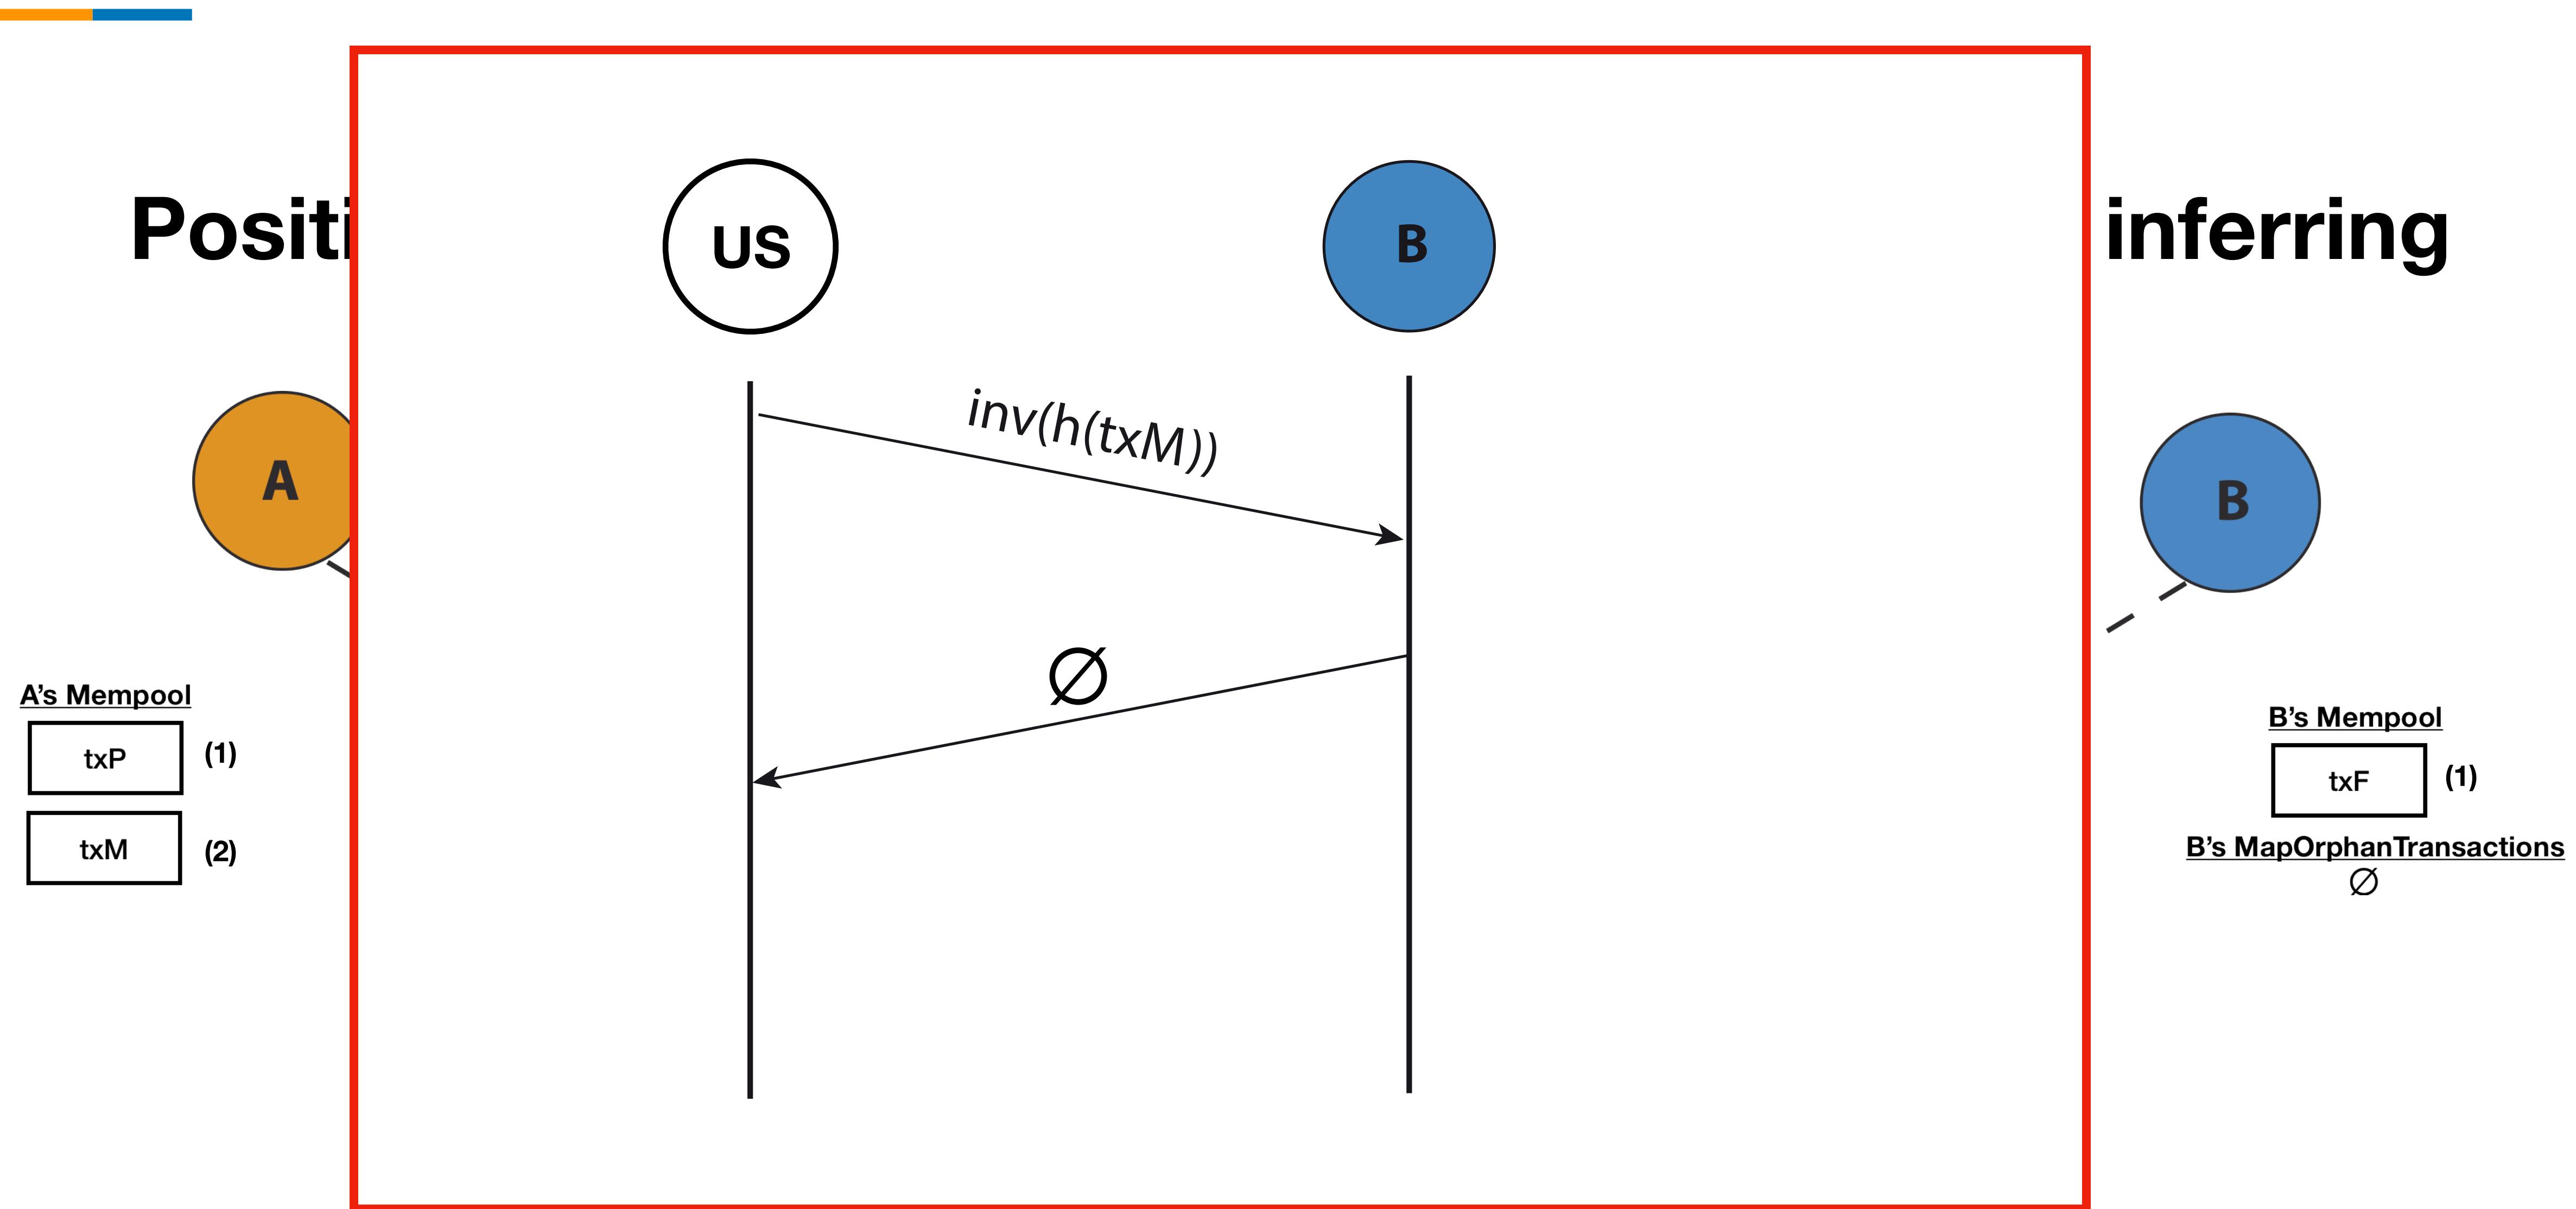
A BASIC TOPOLOGY INFERRING TECHNIQUE



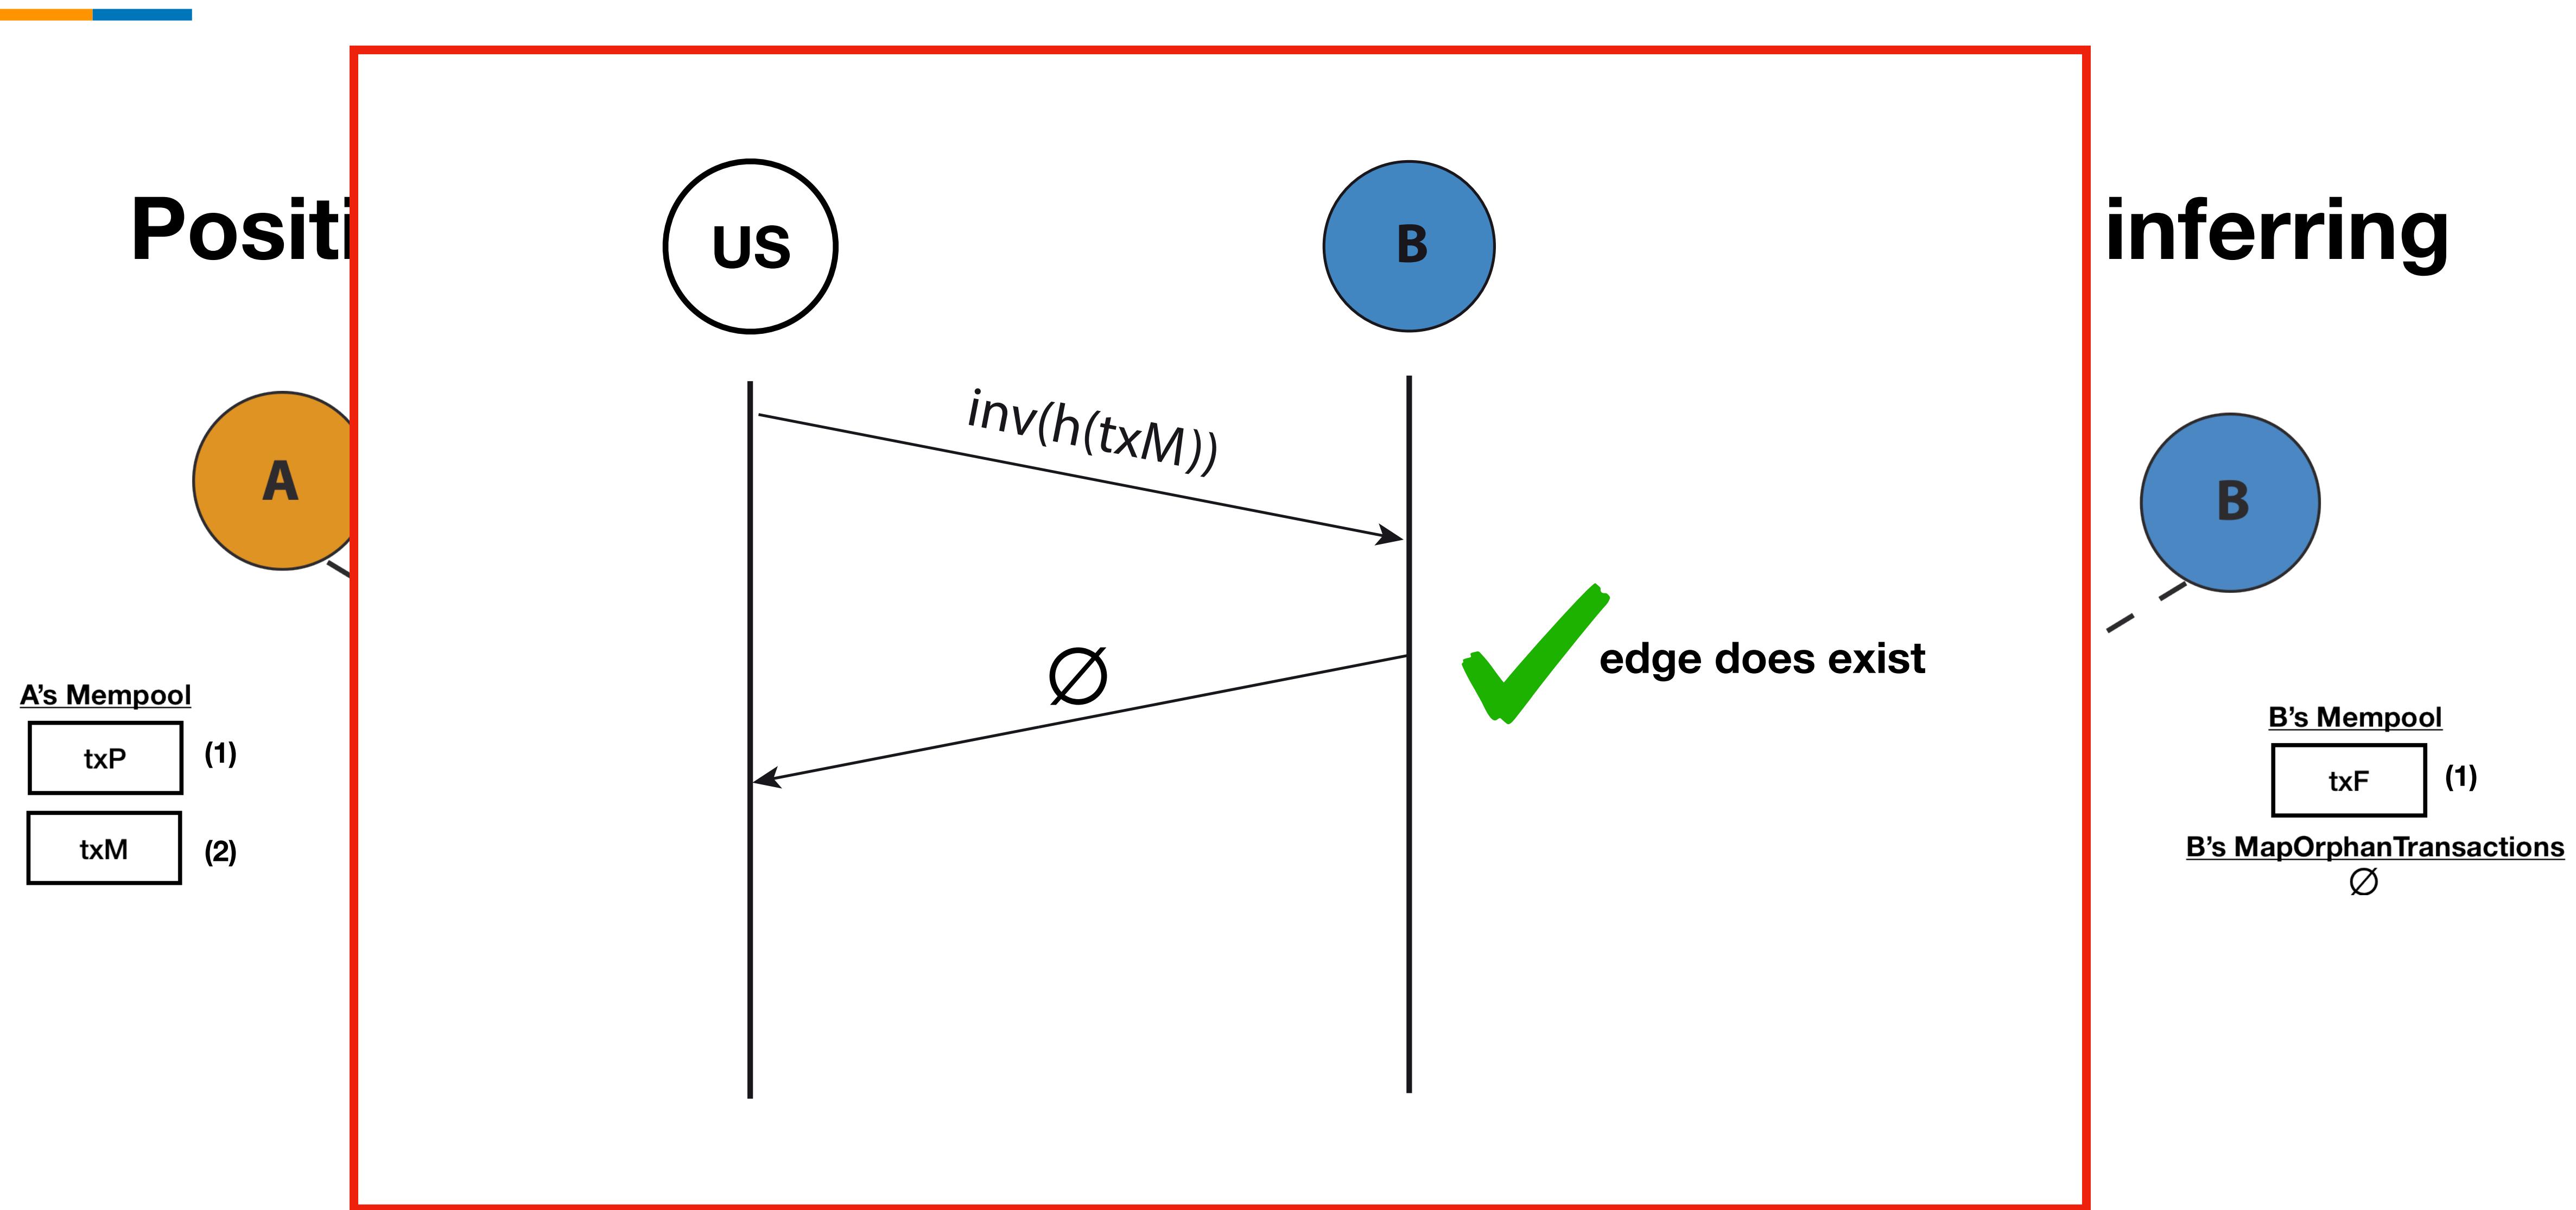
A BASIC TOPOLOGY INFERRING TECHNIQUE



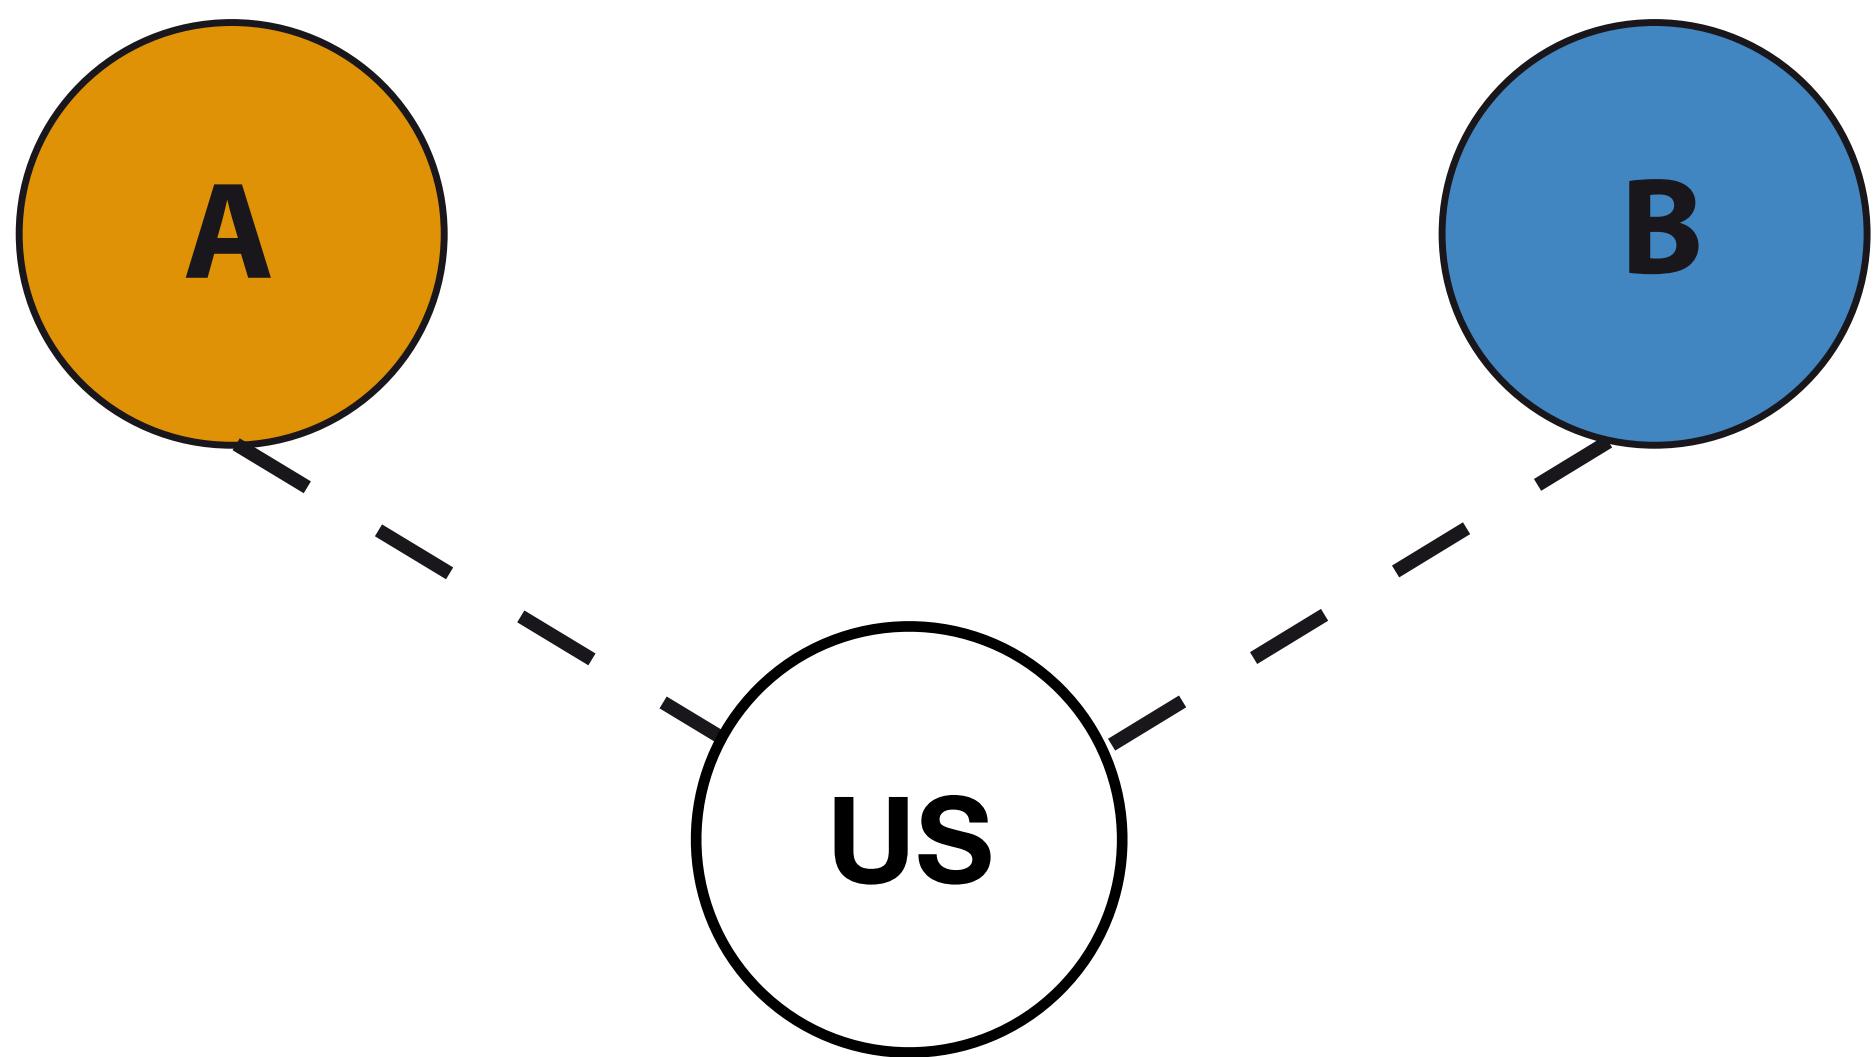
A BASIC TOPOLOGY INFERRING TECHNIQUE



A BASIC TOPOLOGY INFERRING TECHNIQUE

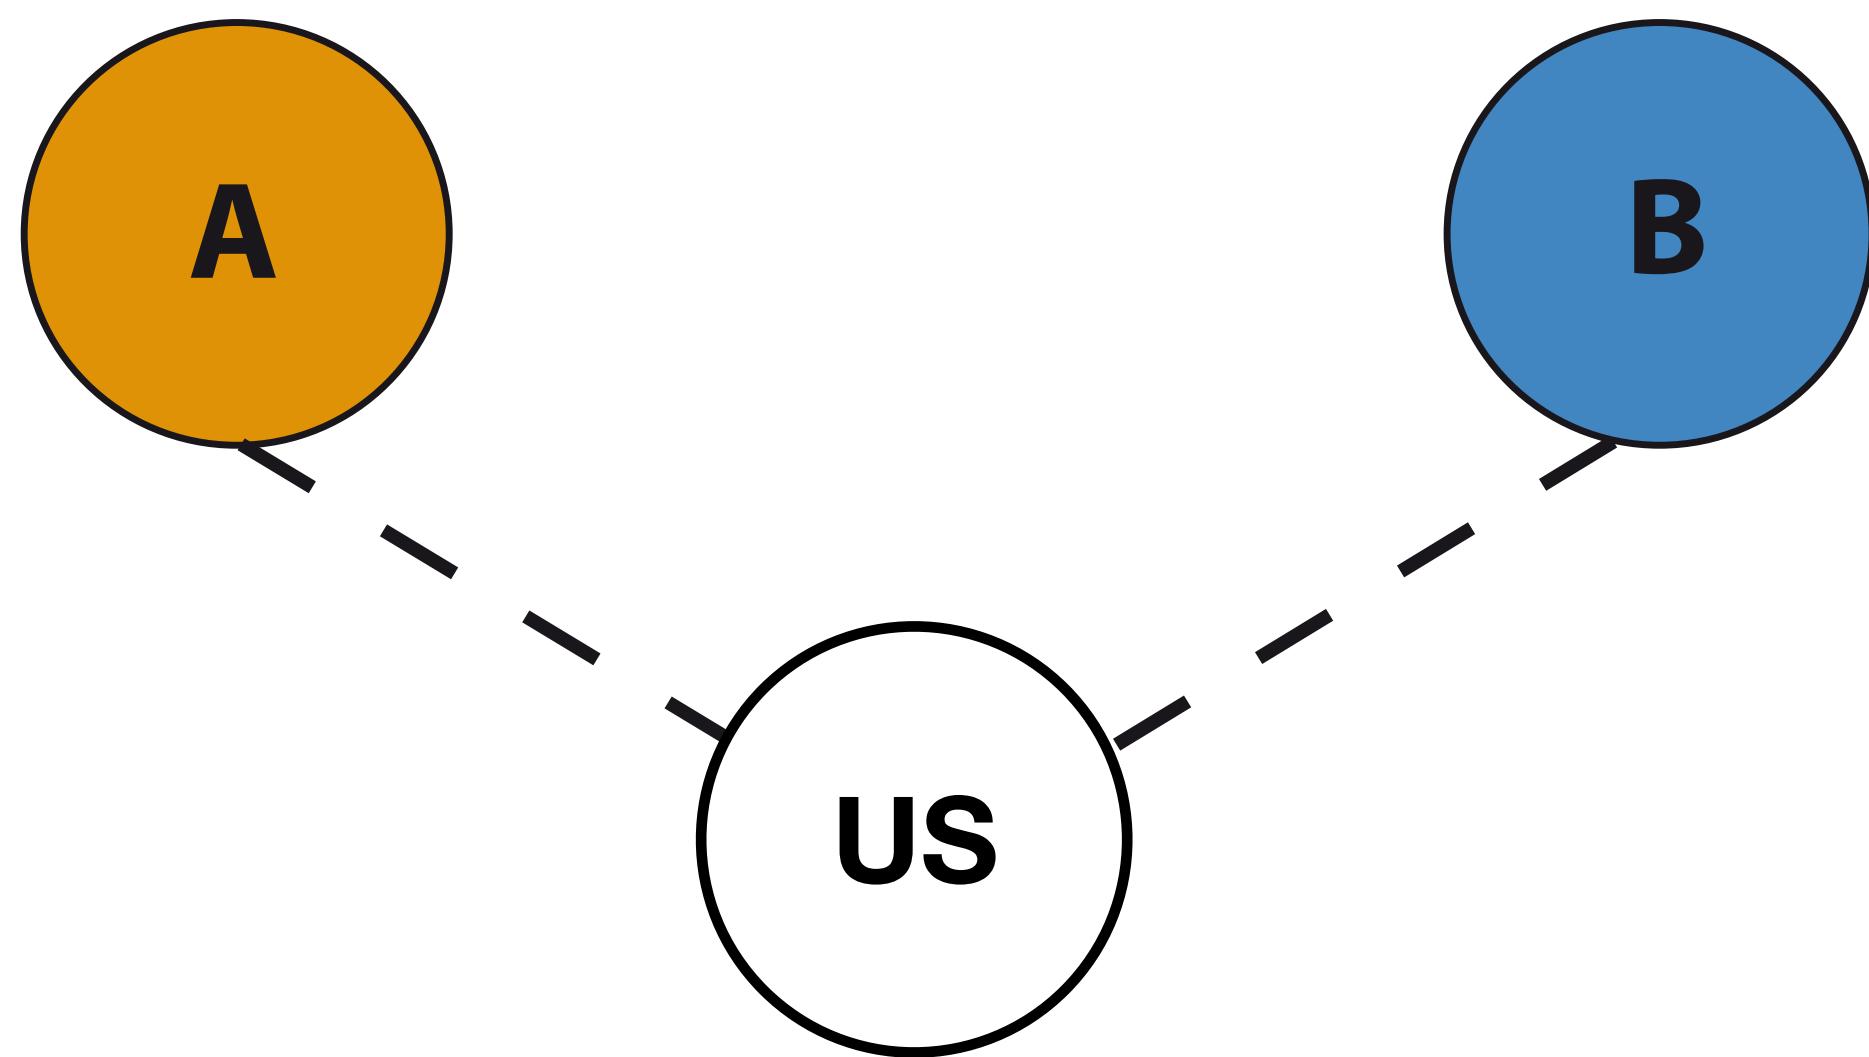


ITS NOT THAT EASY



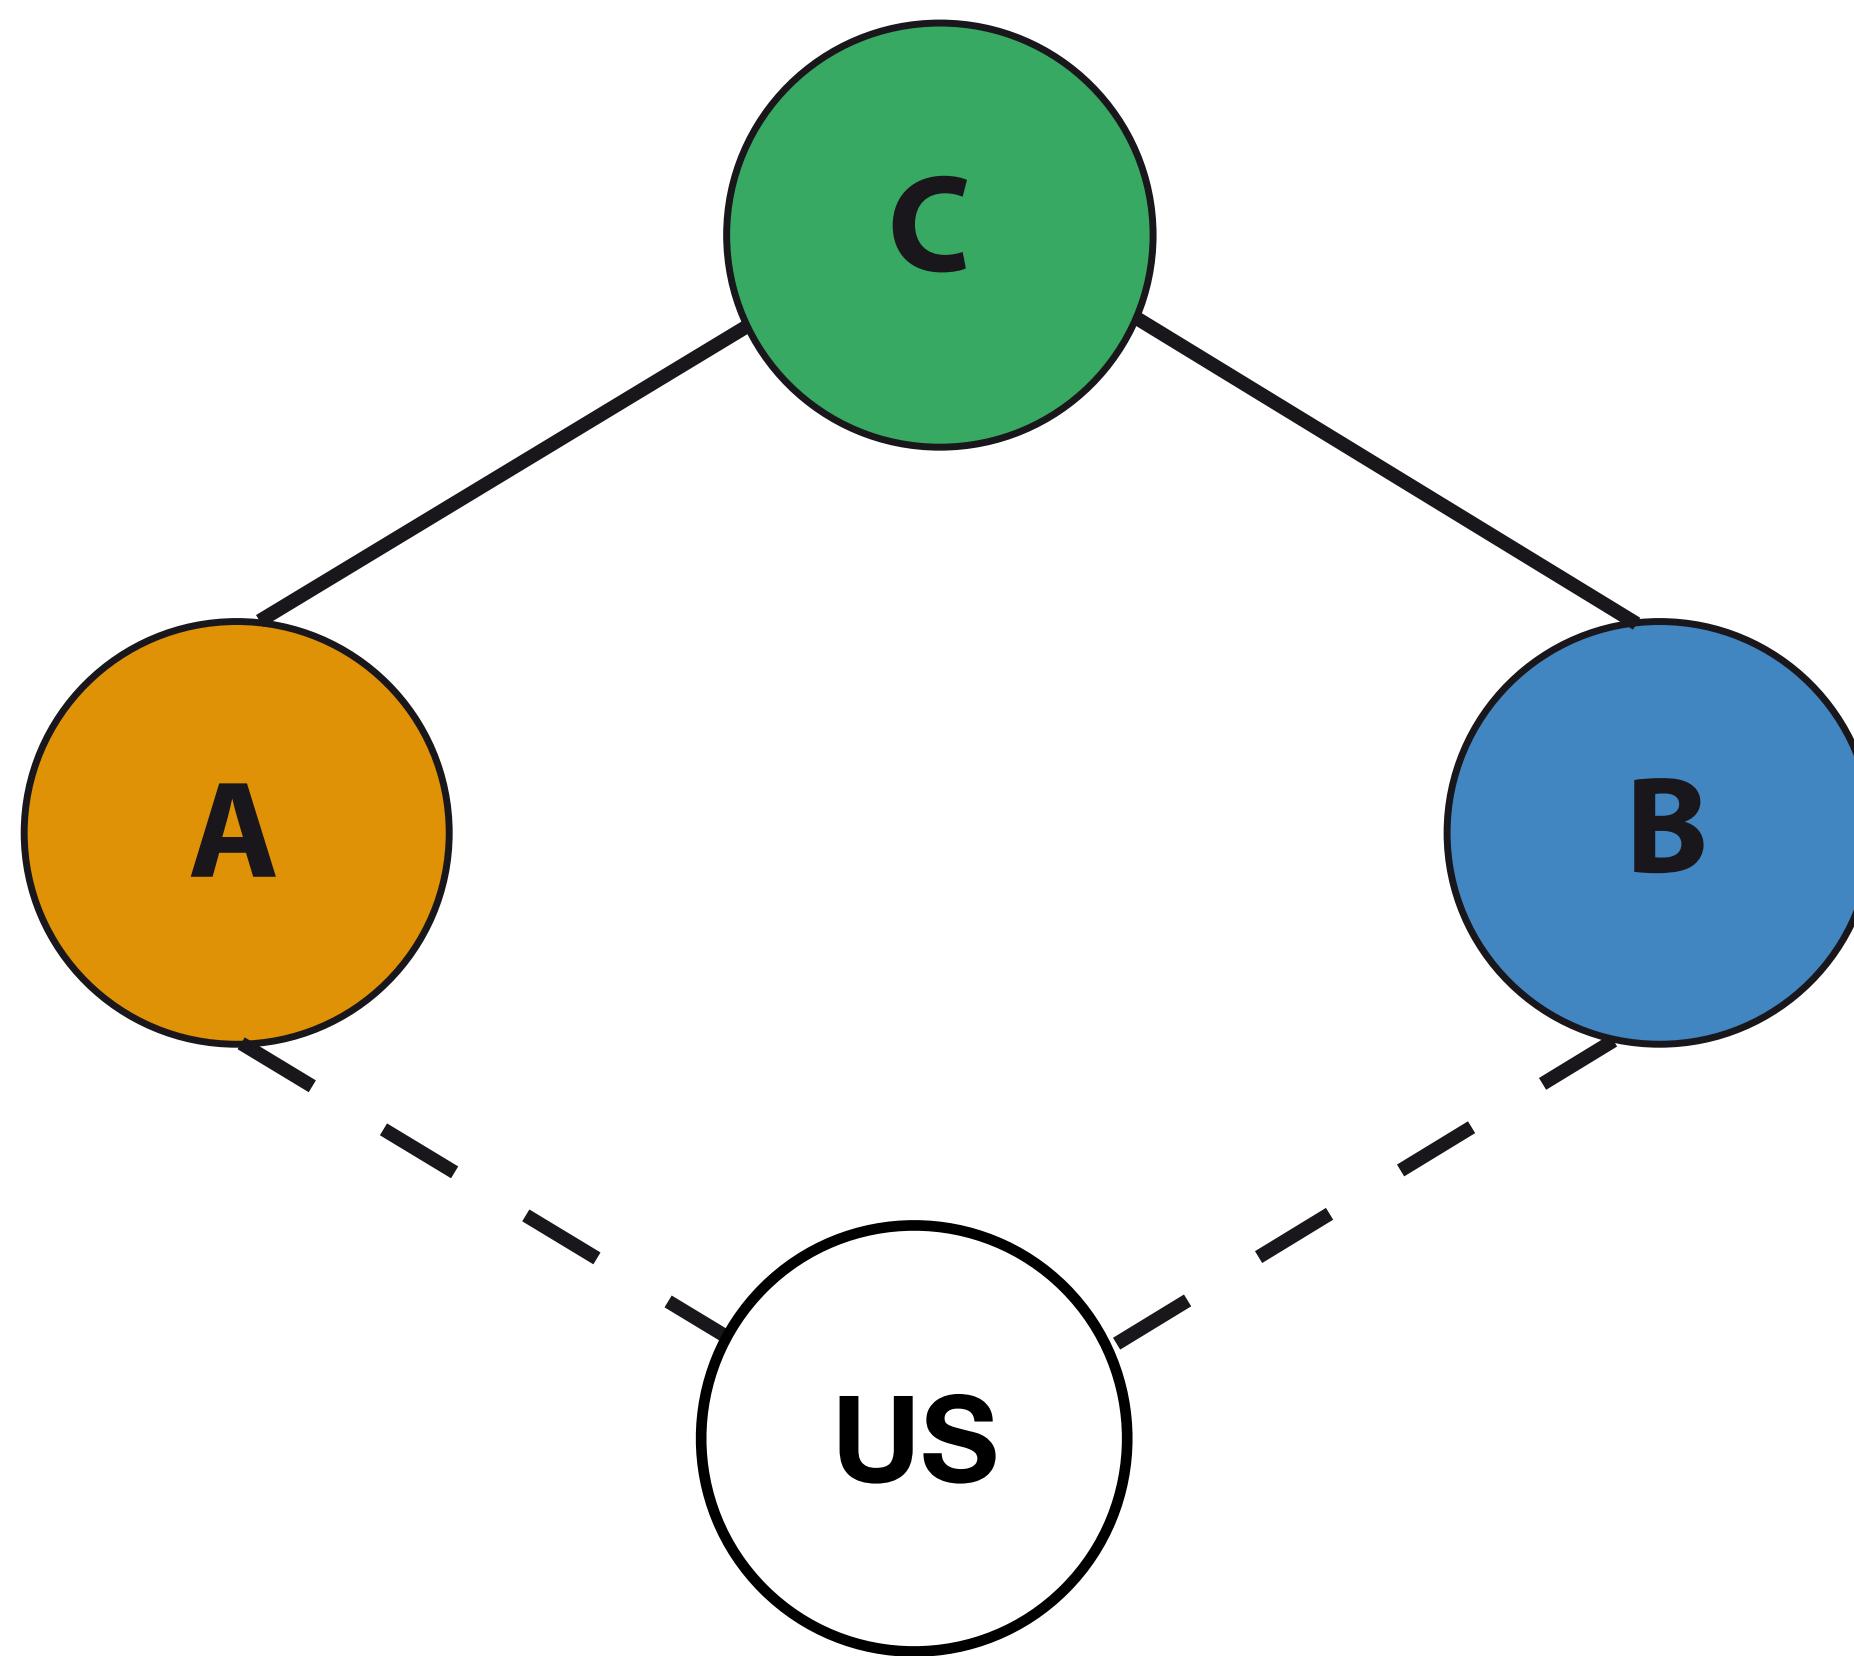
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



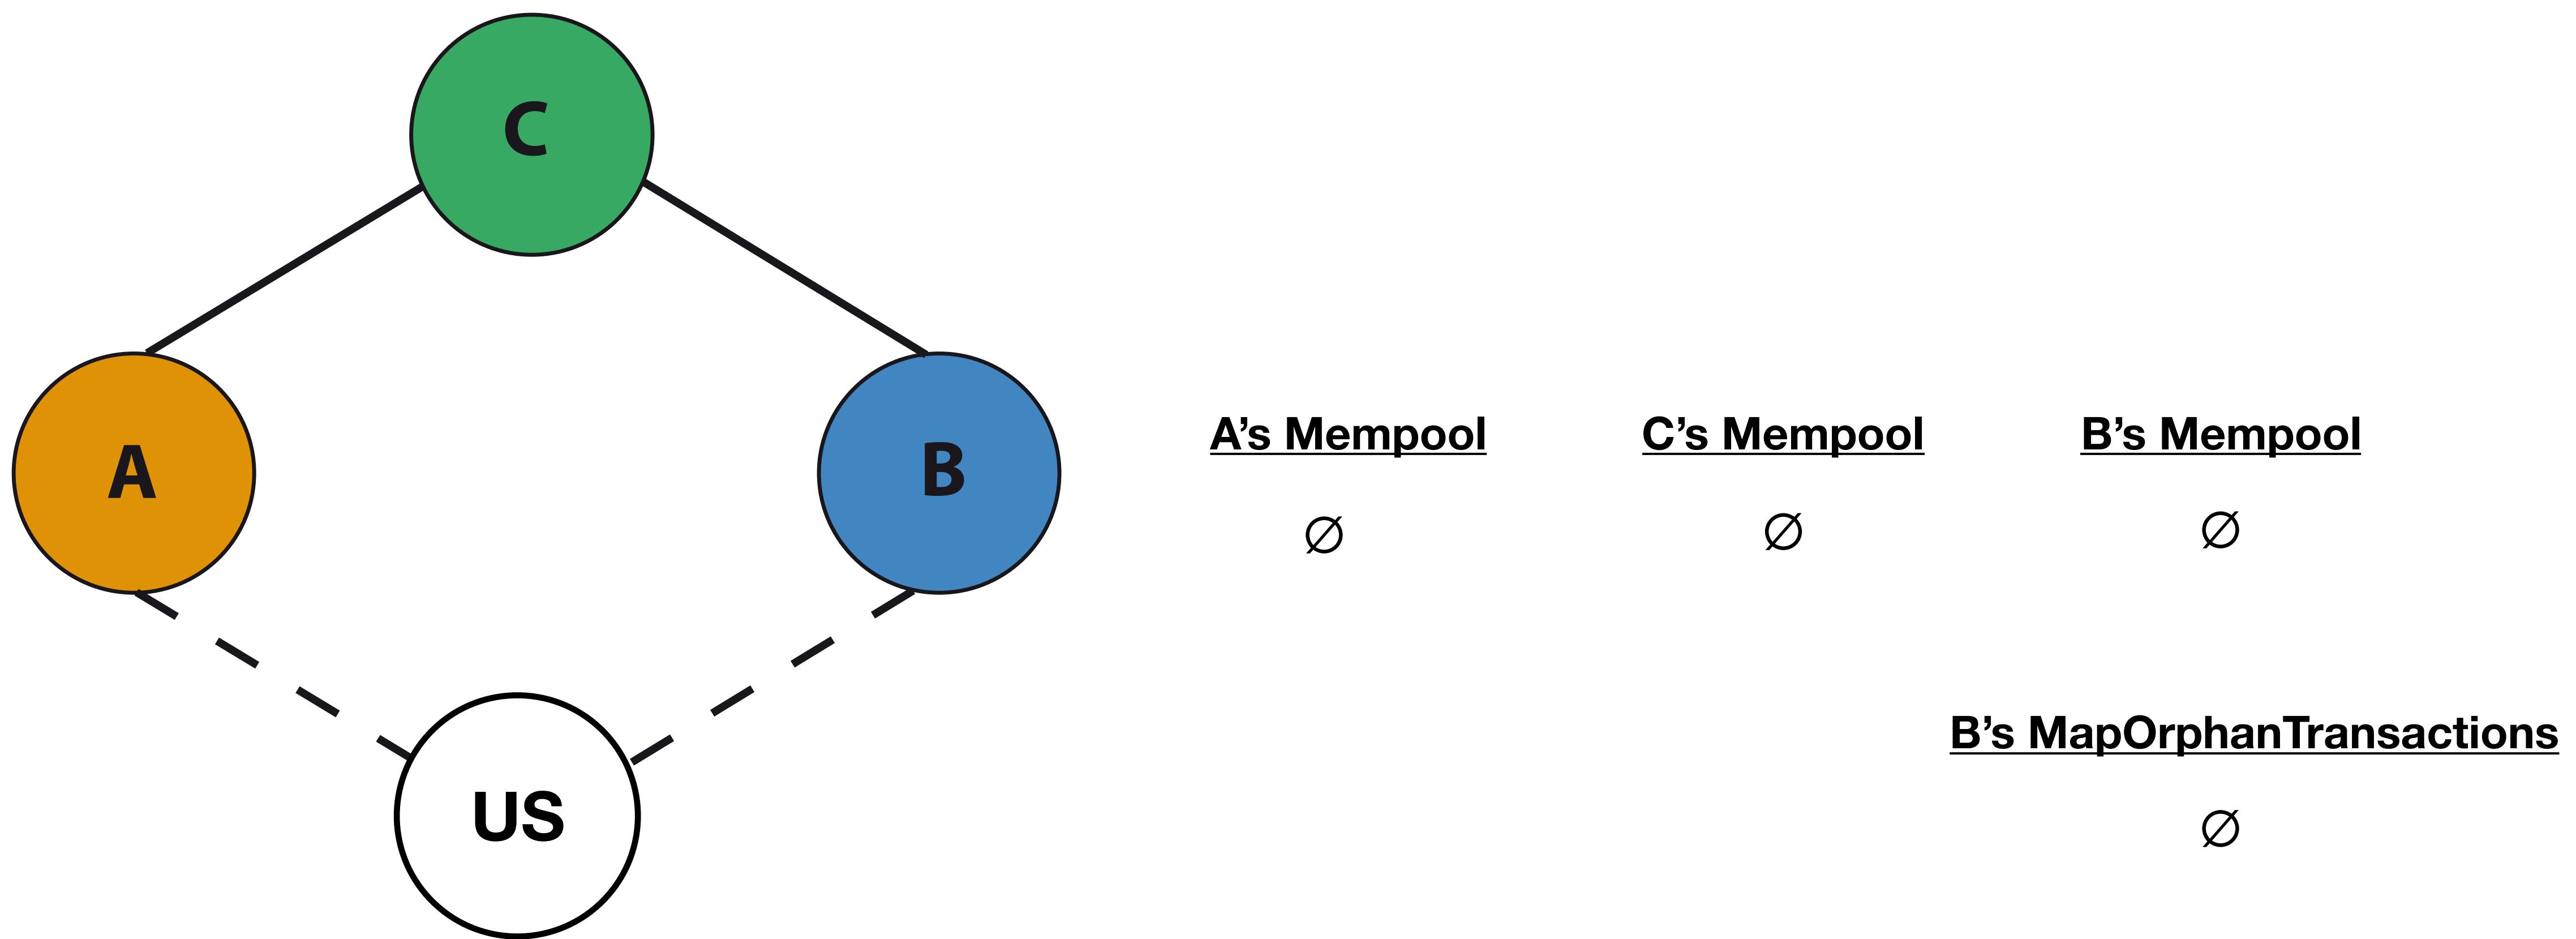
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



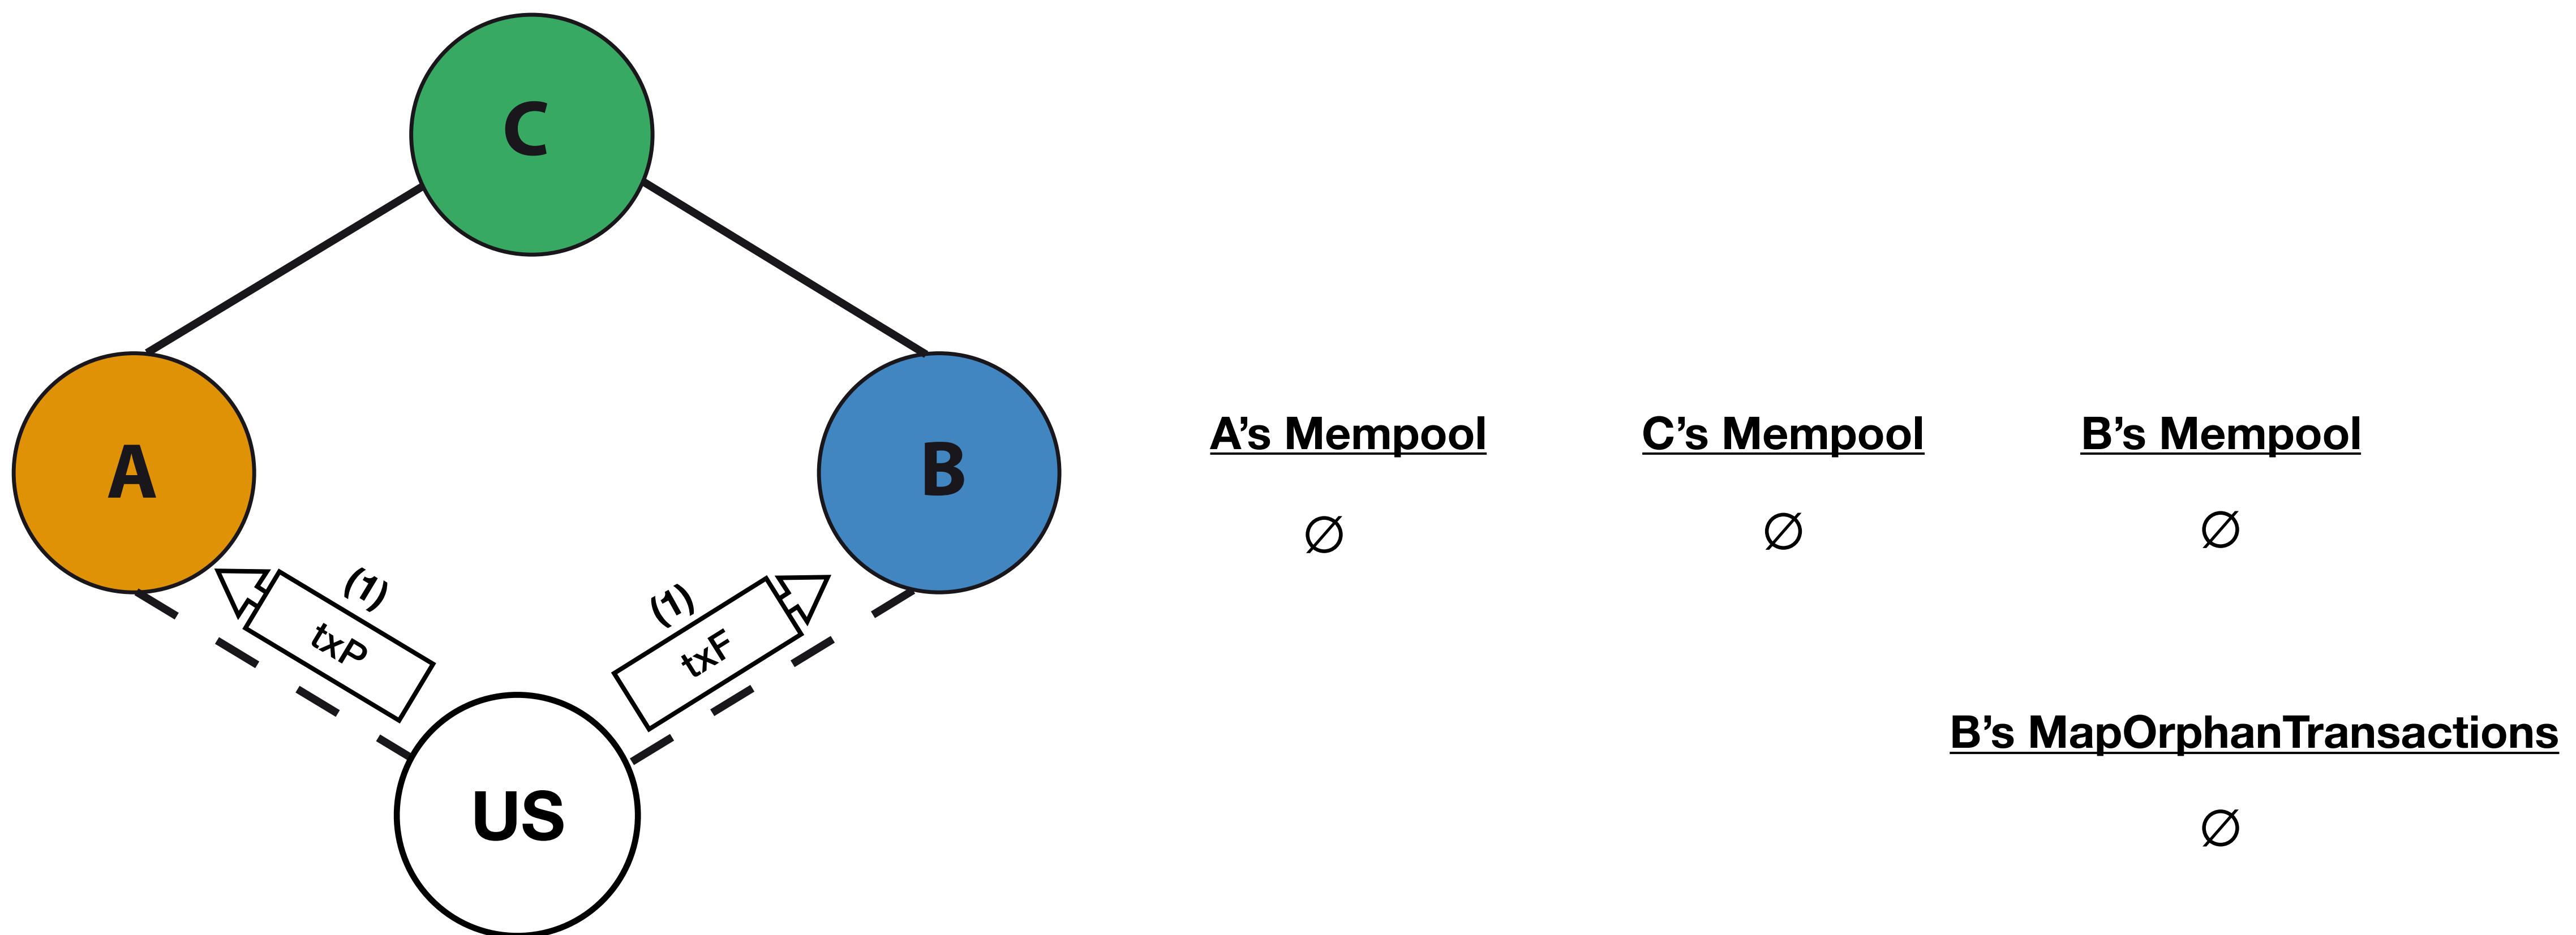
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



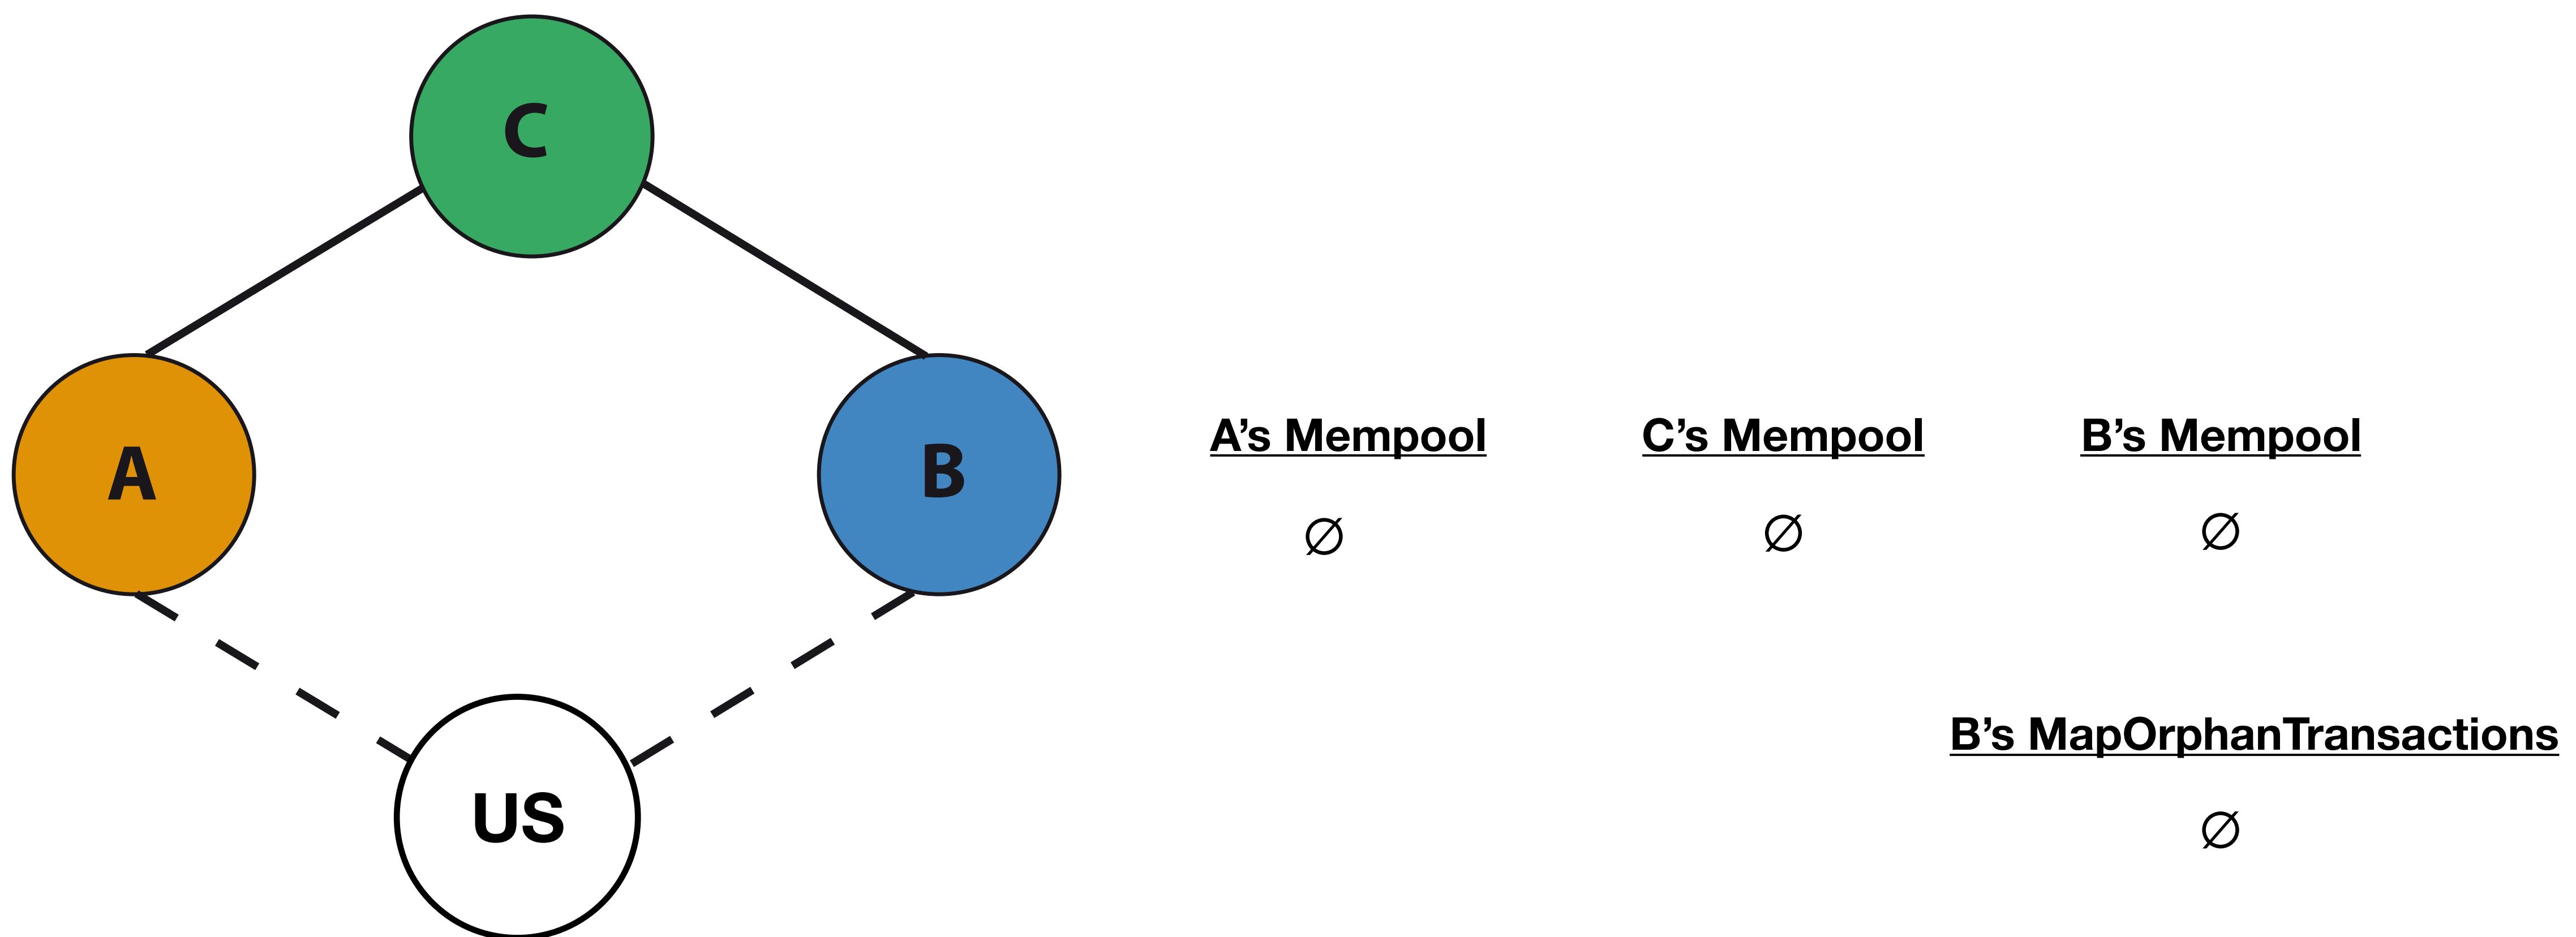
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



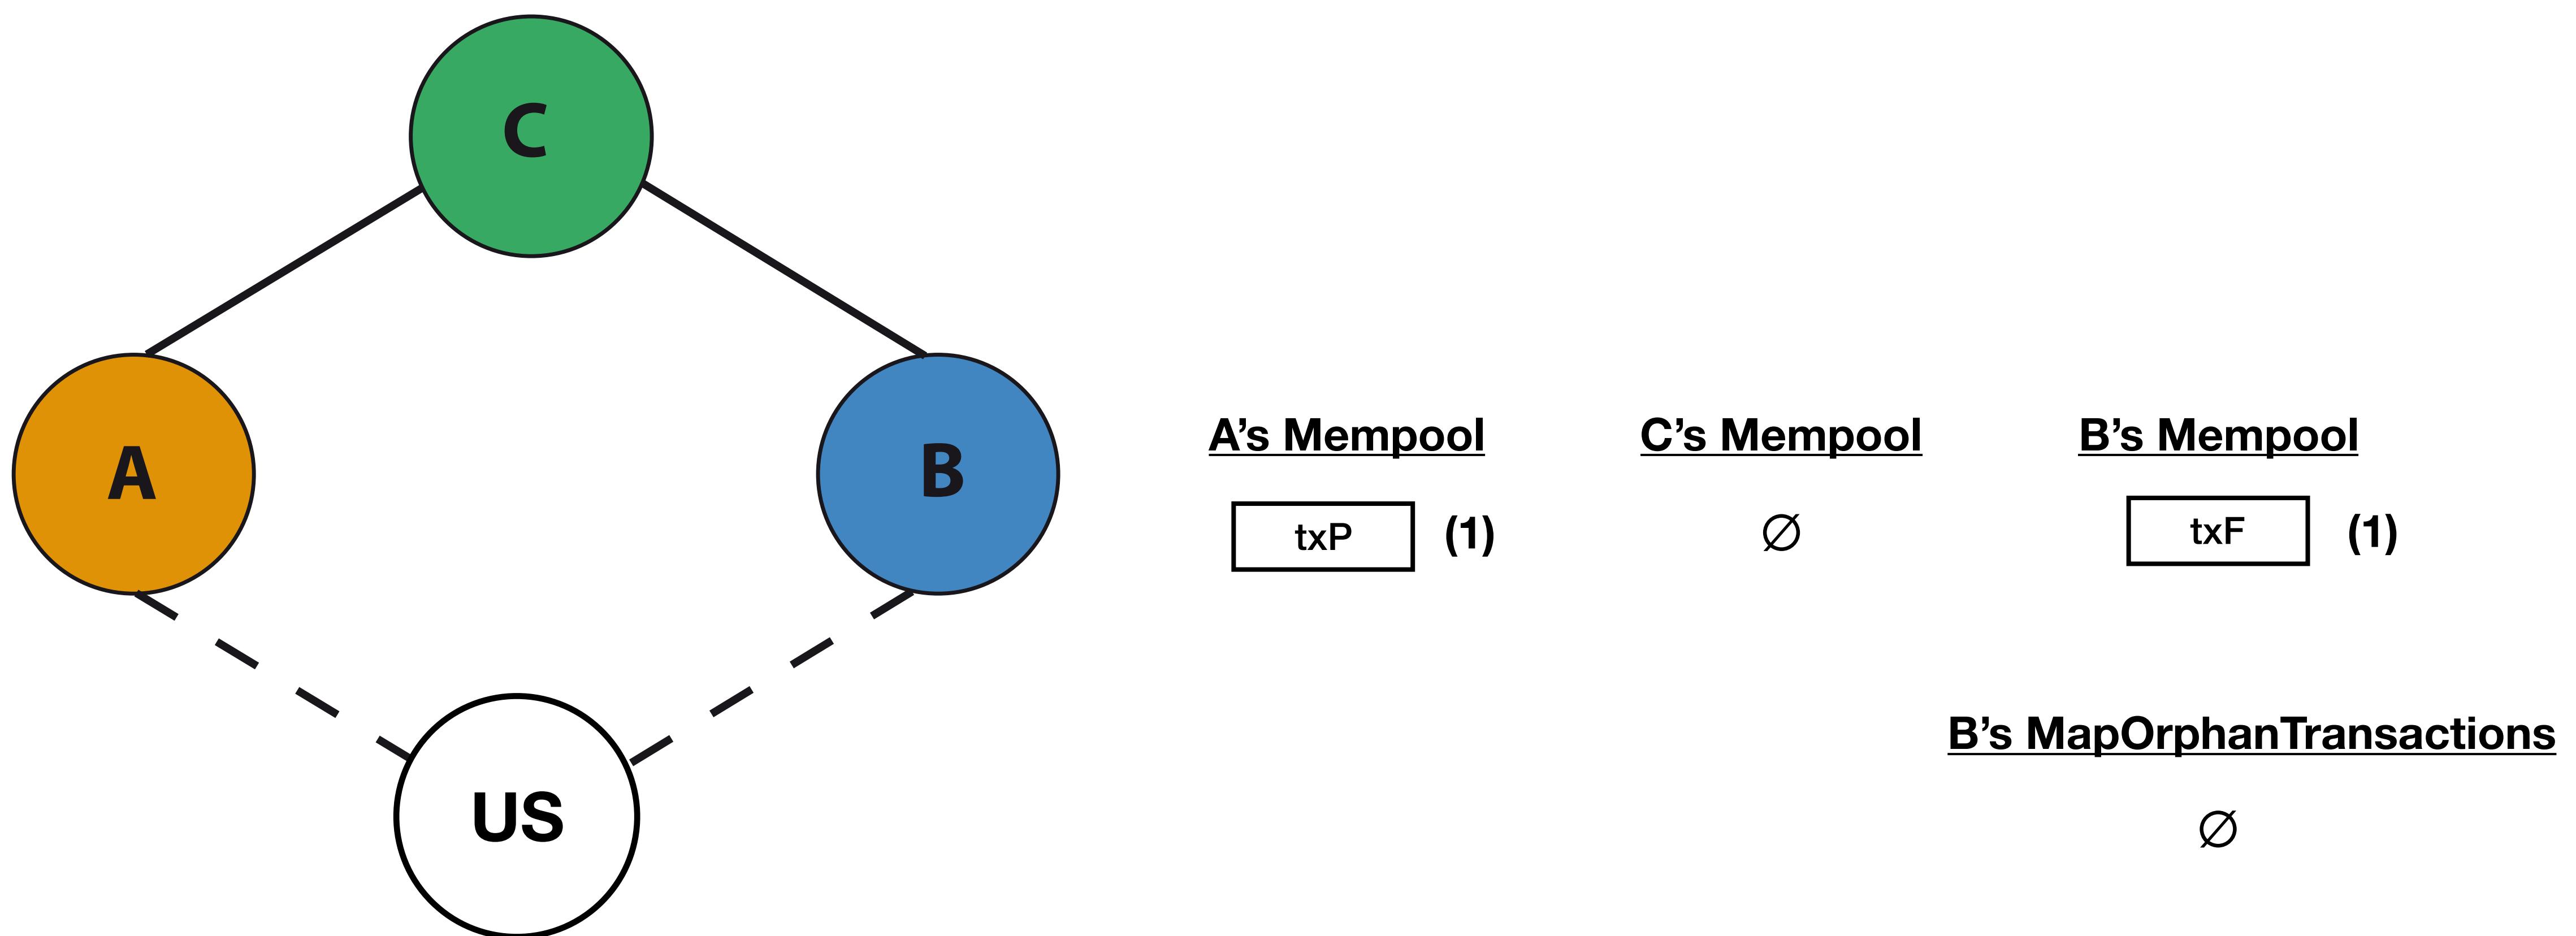
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



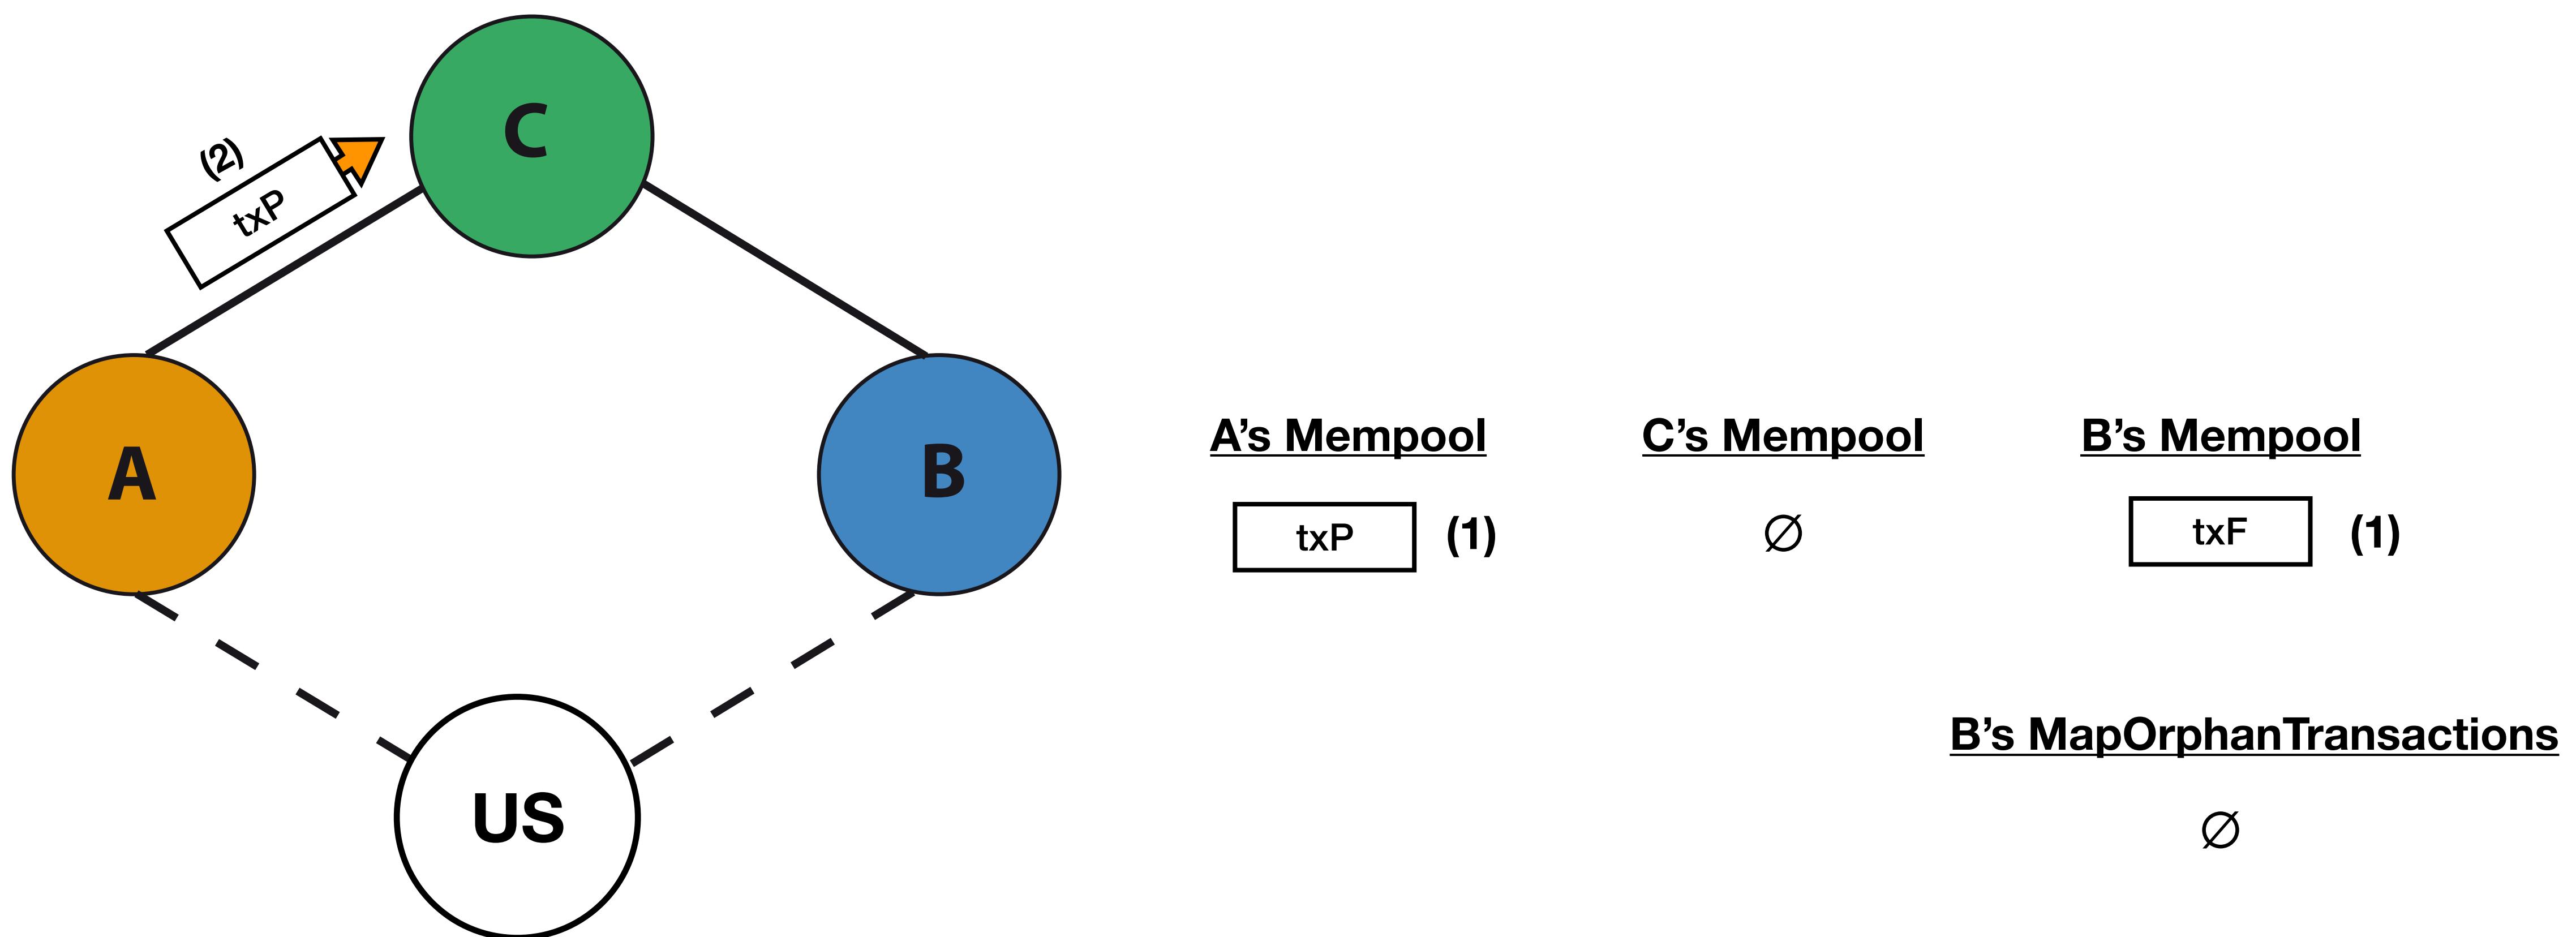
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



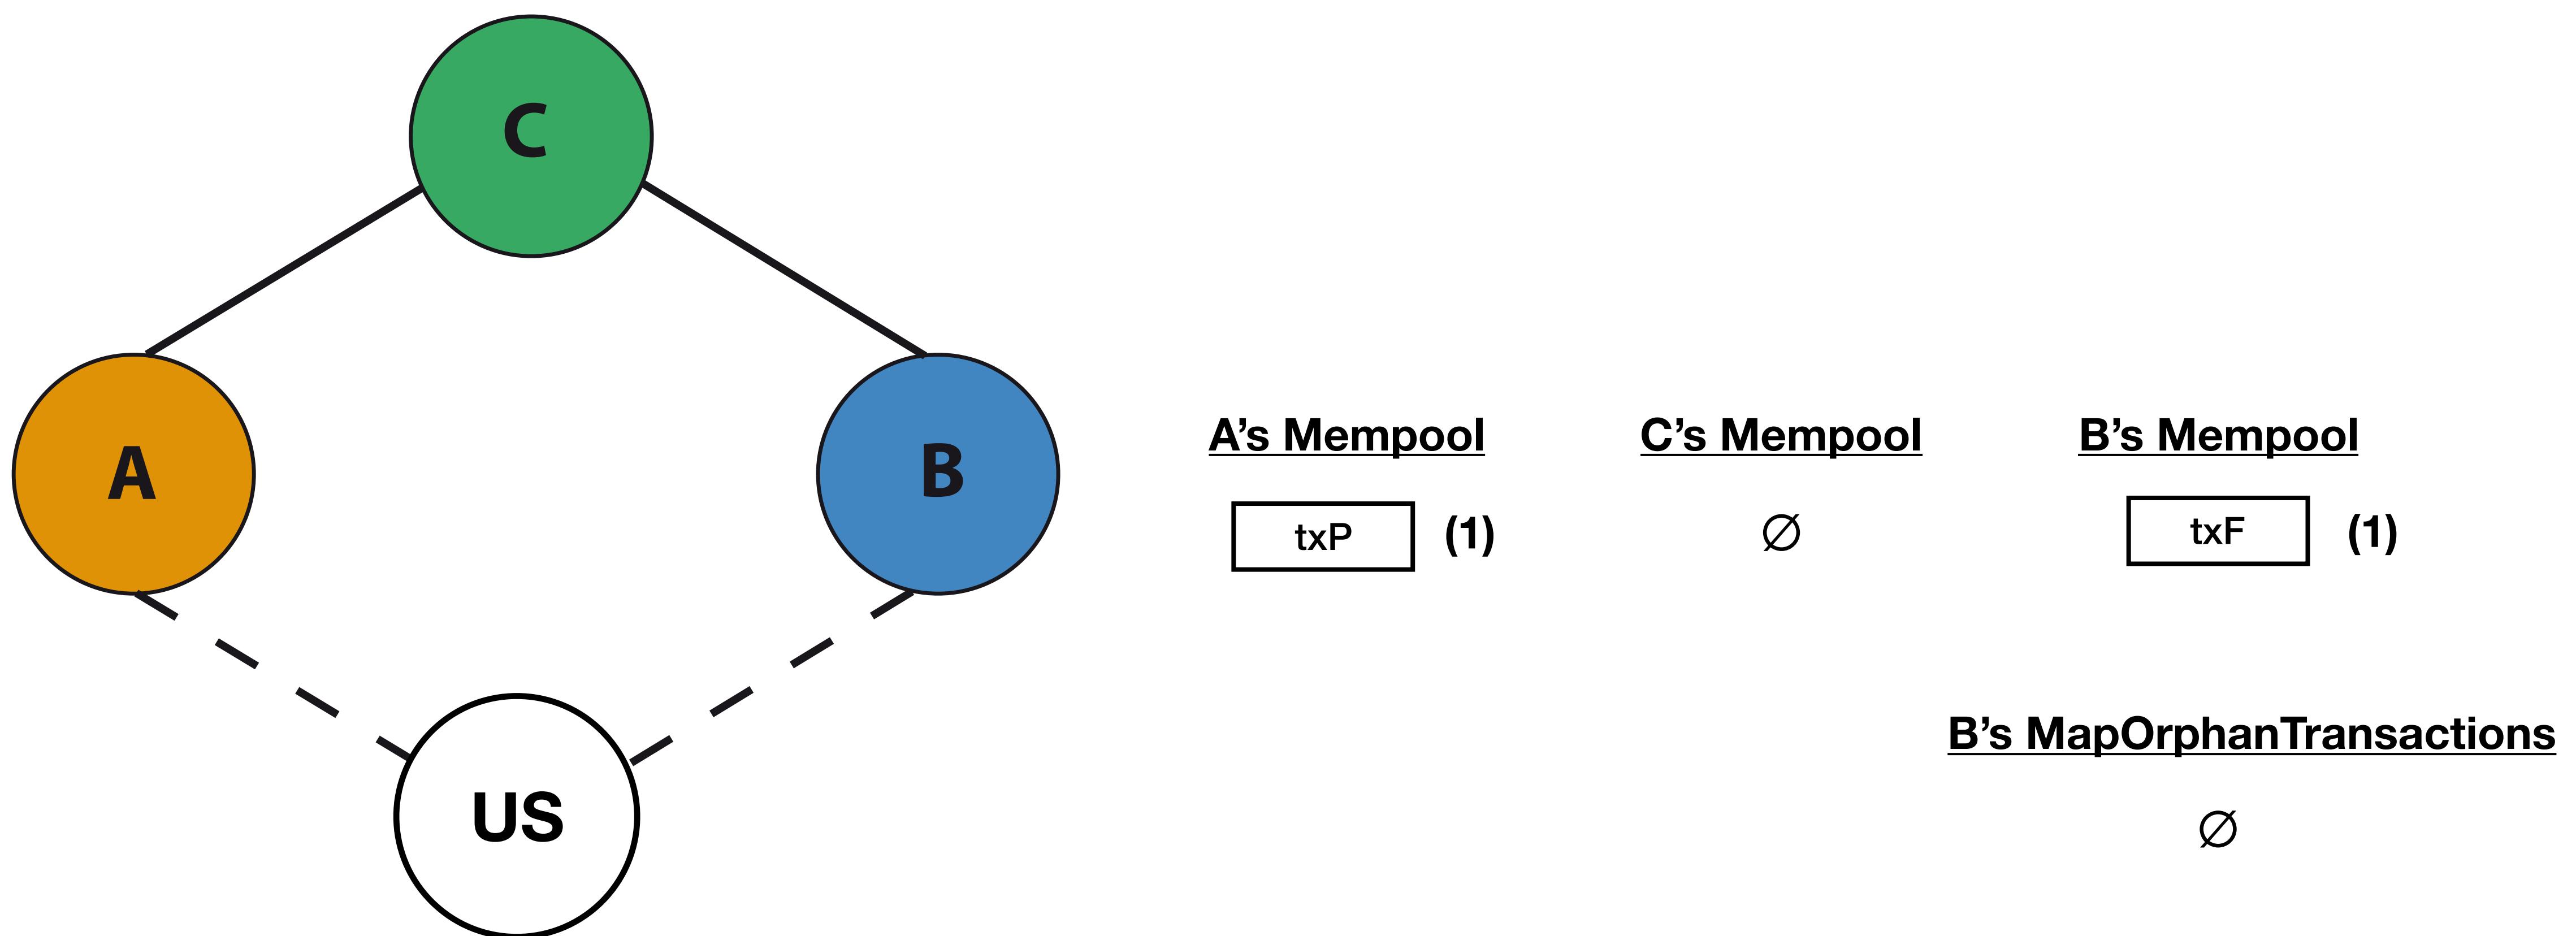
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



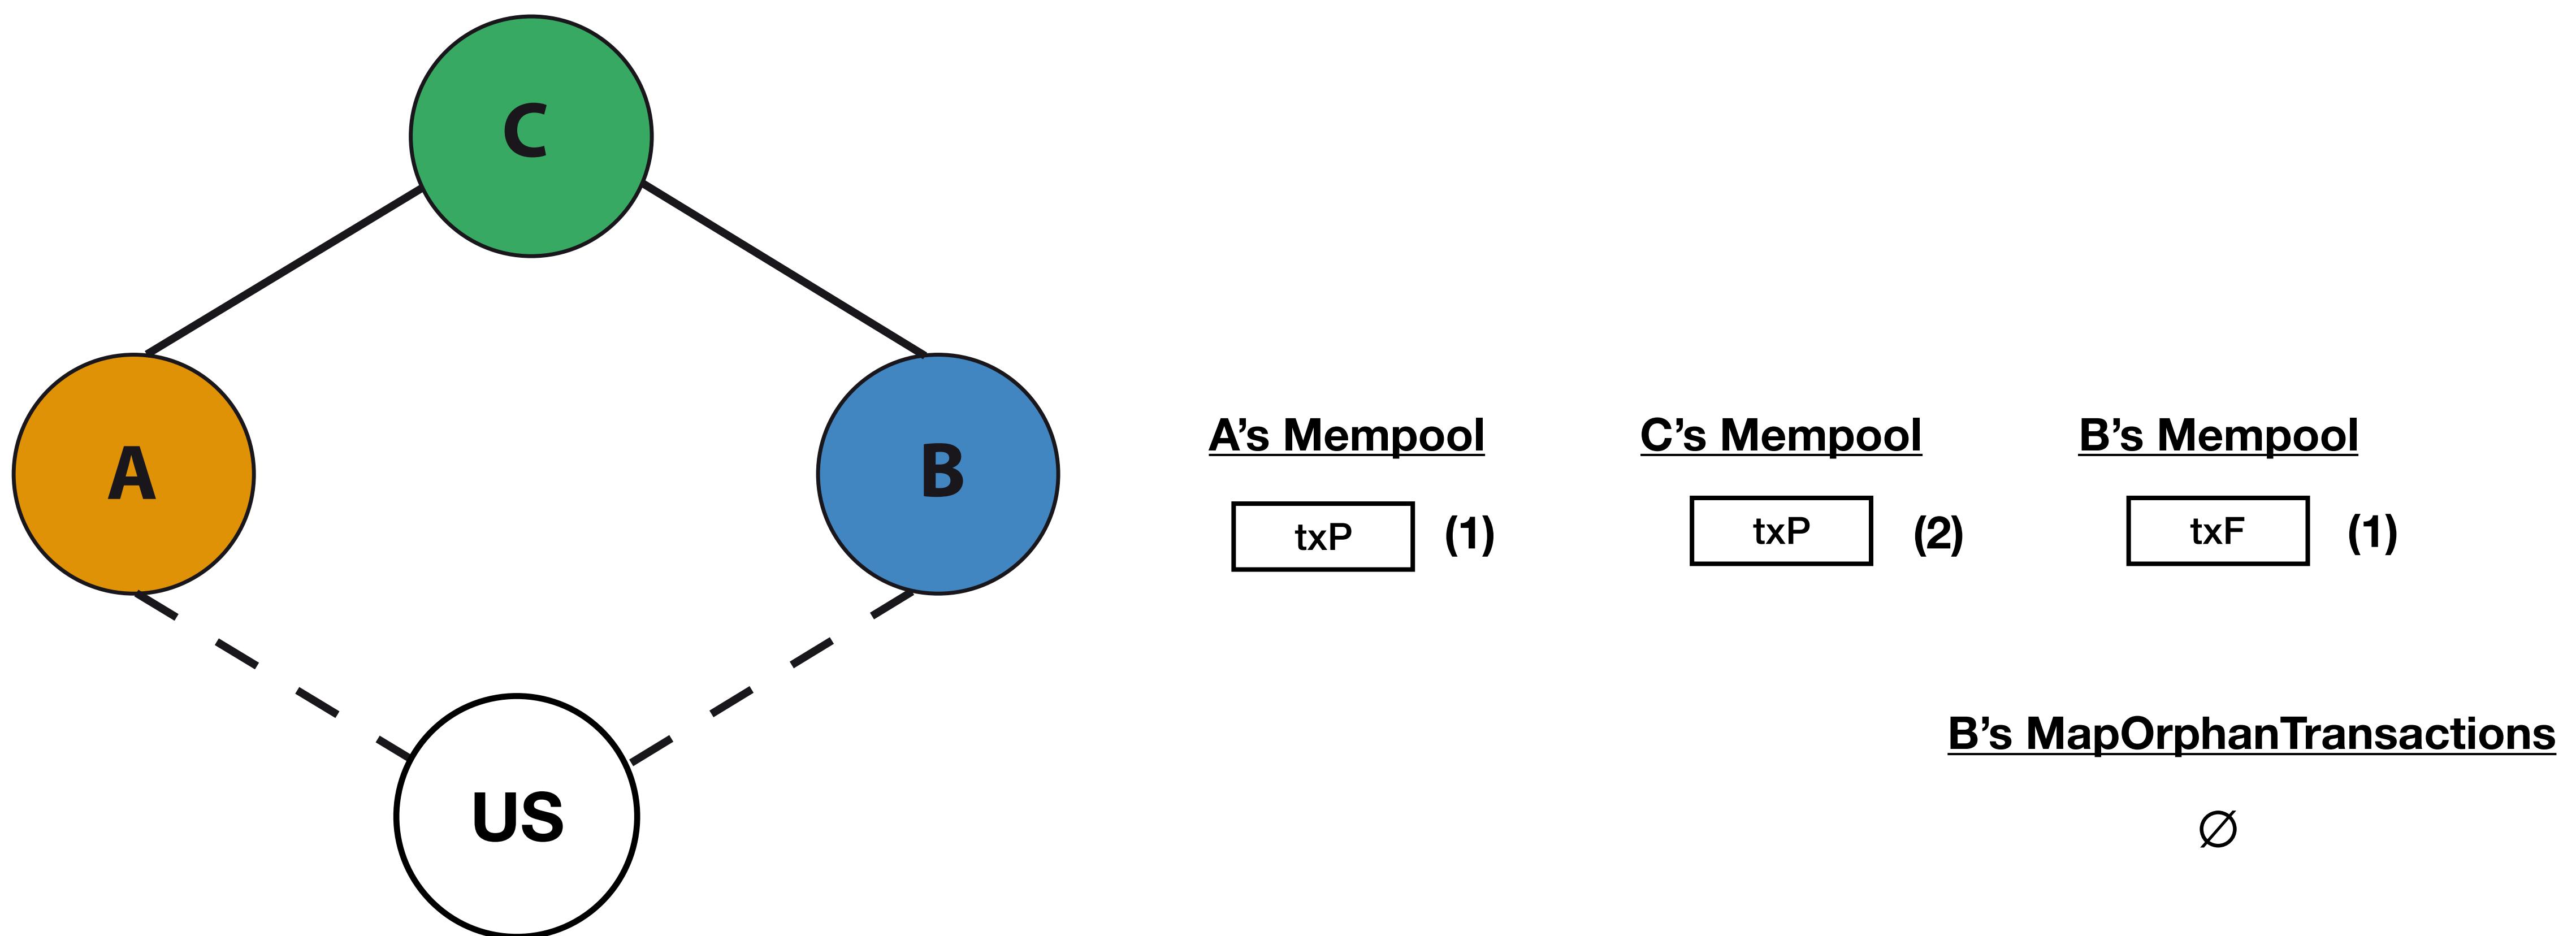
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



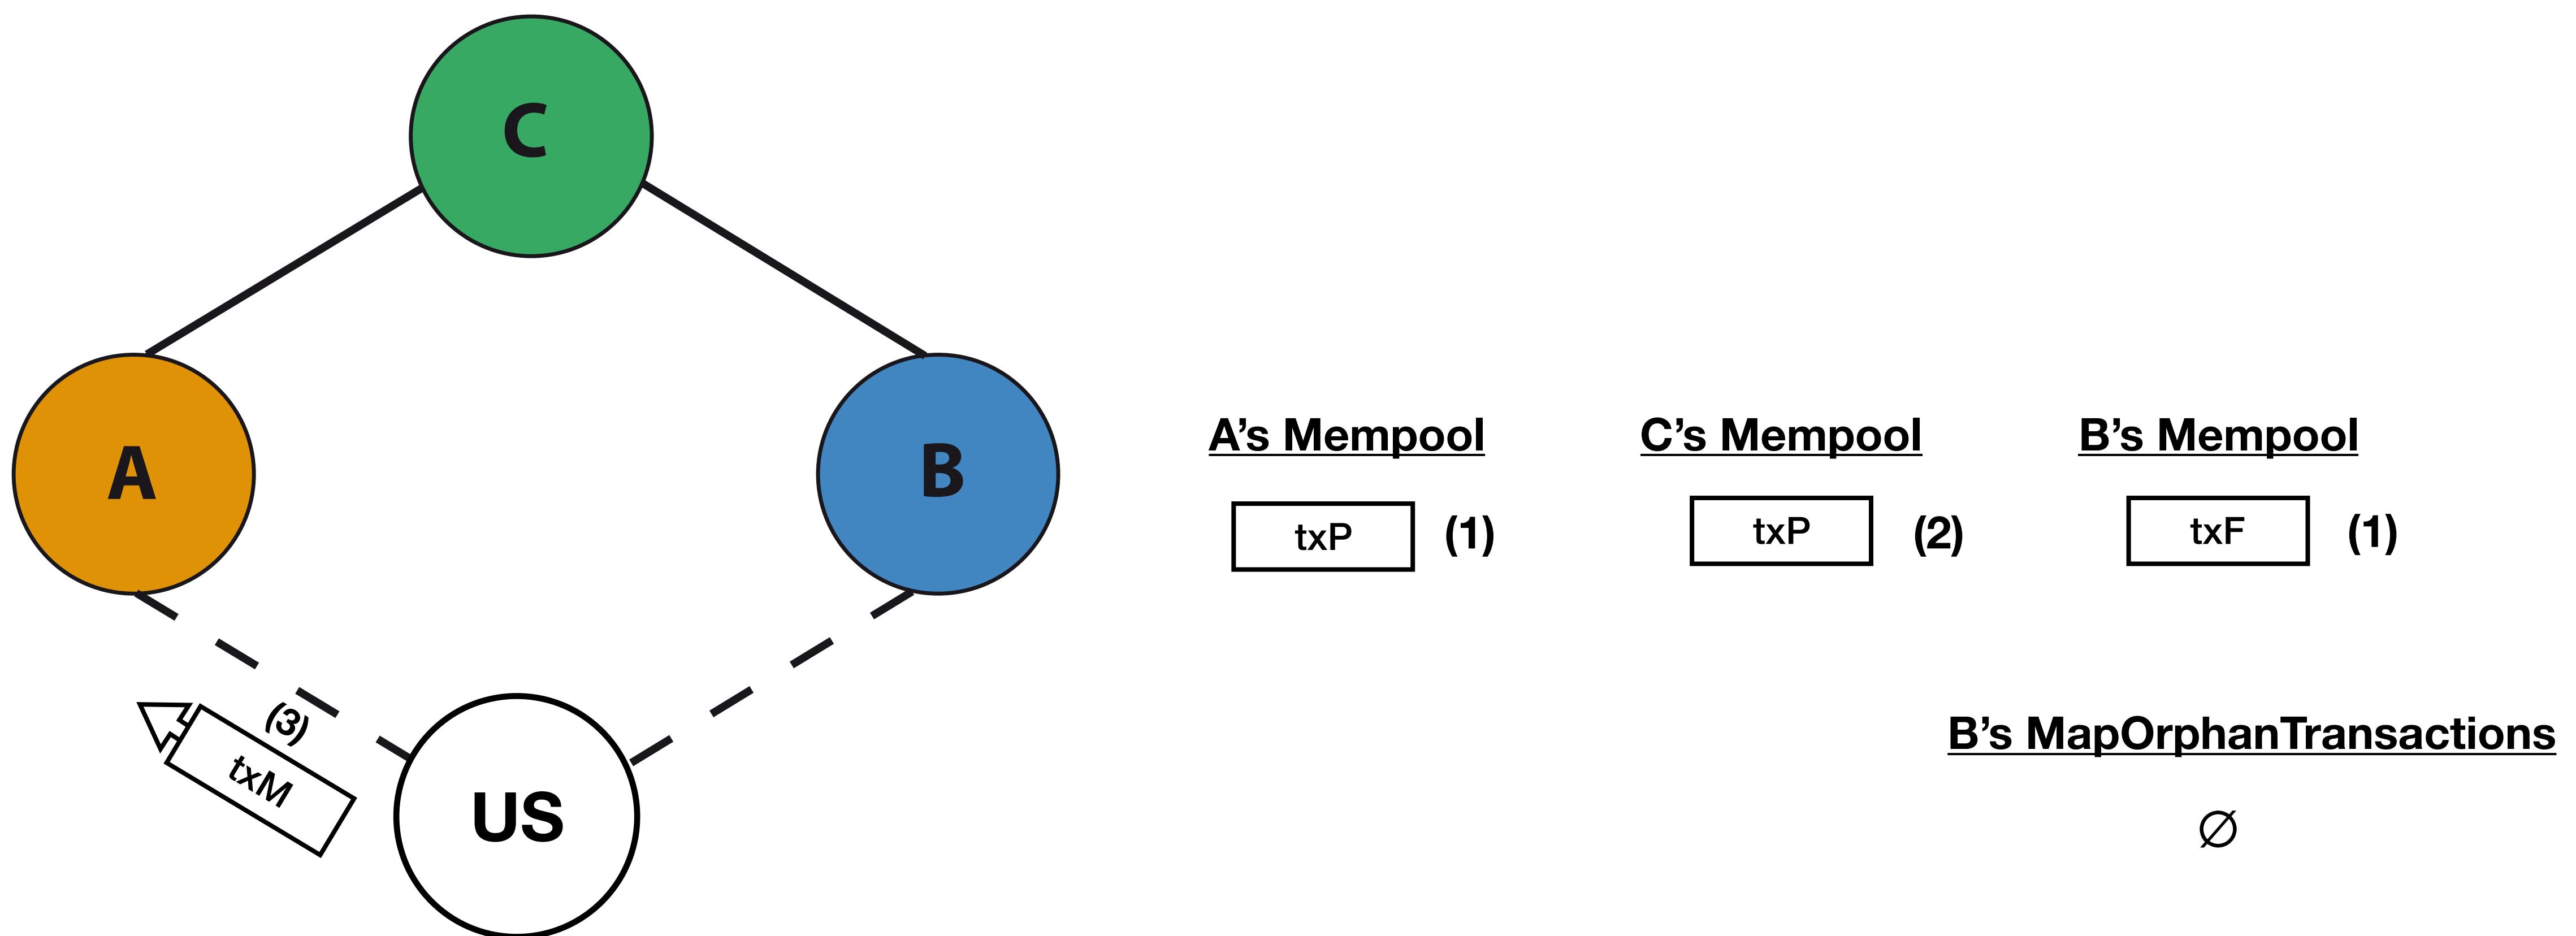
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



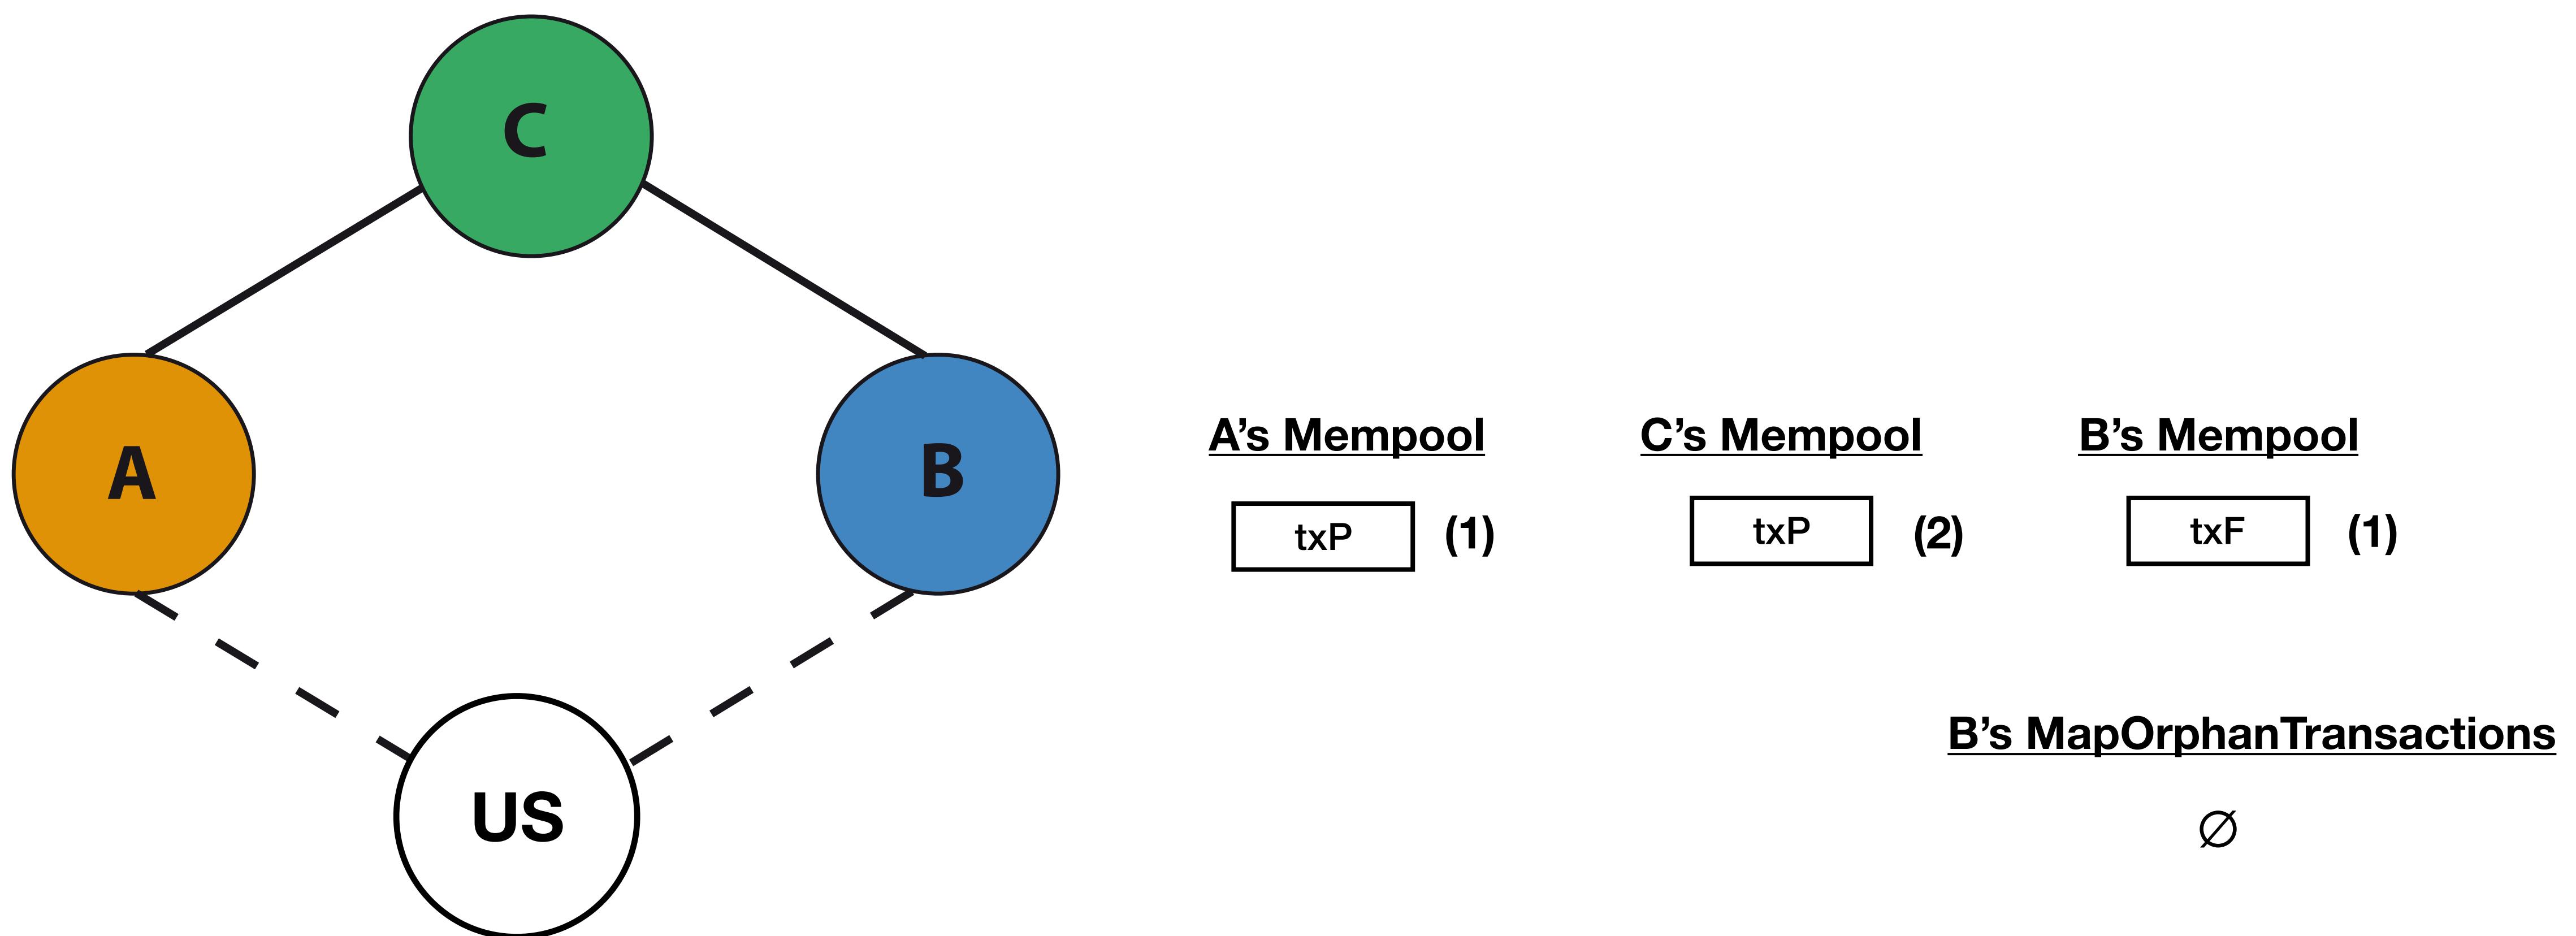
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



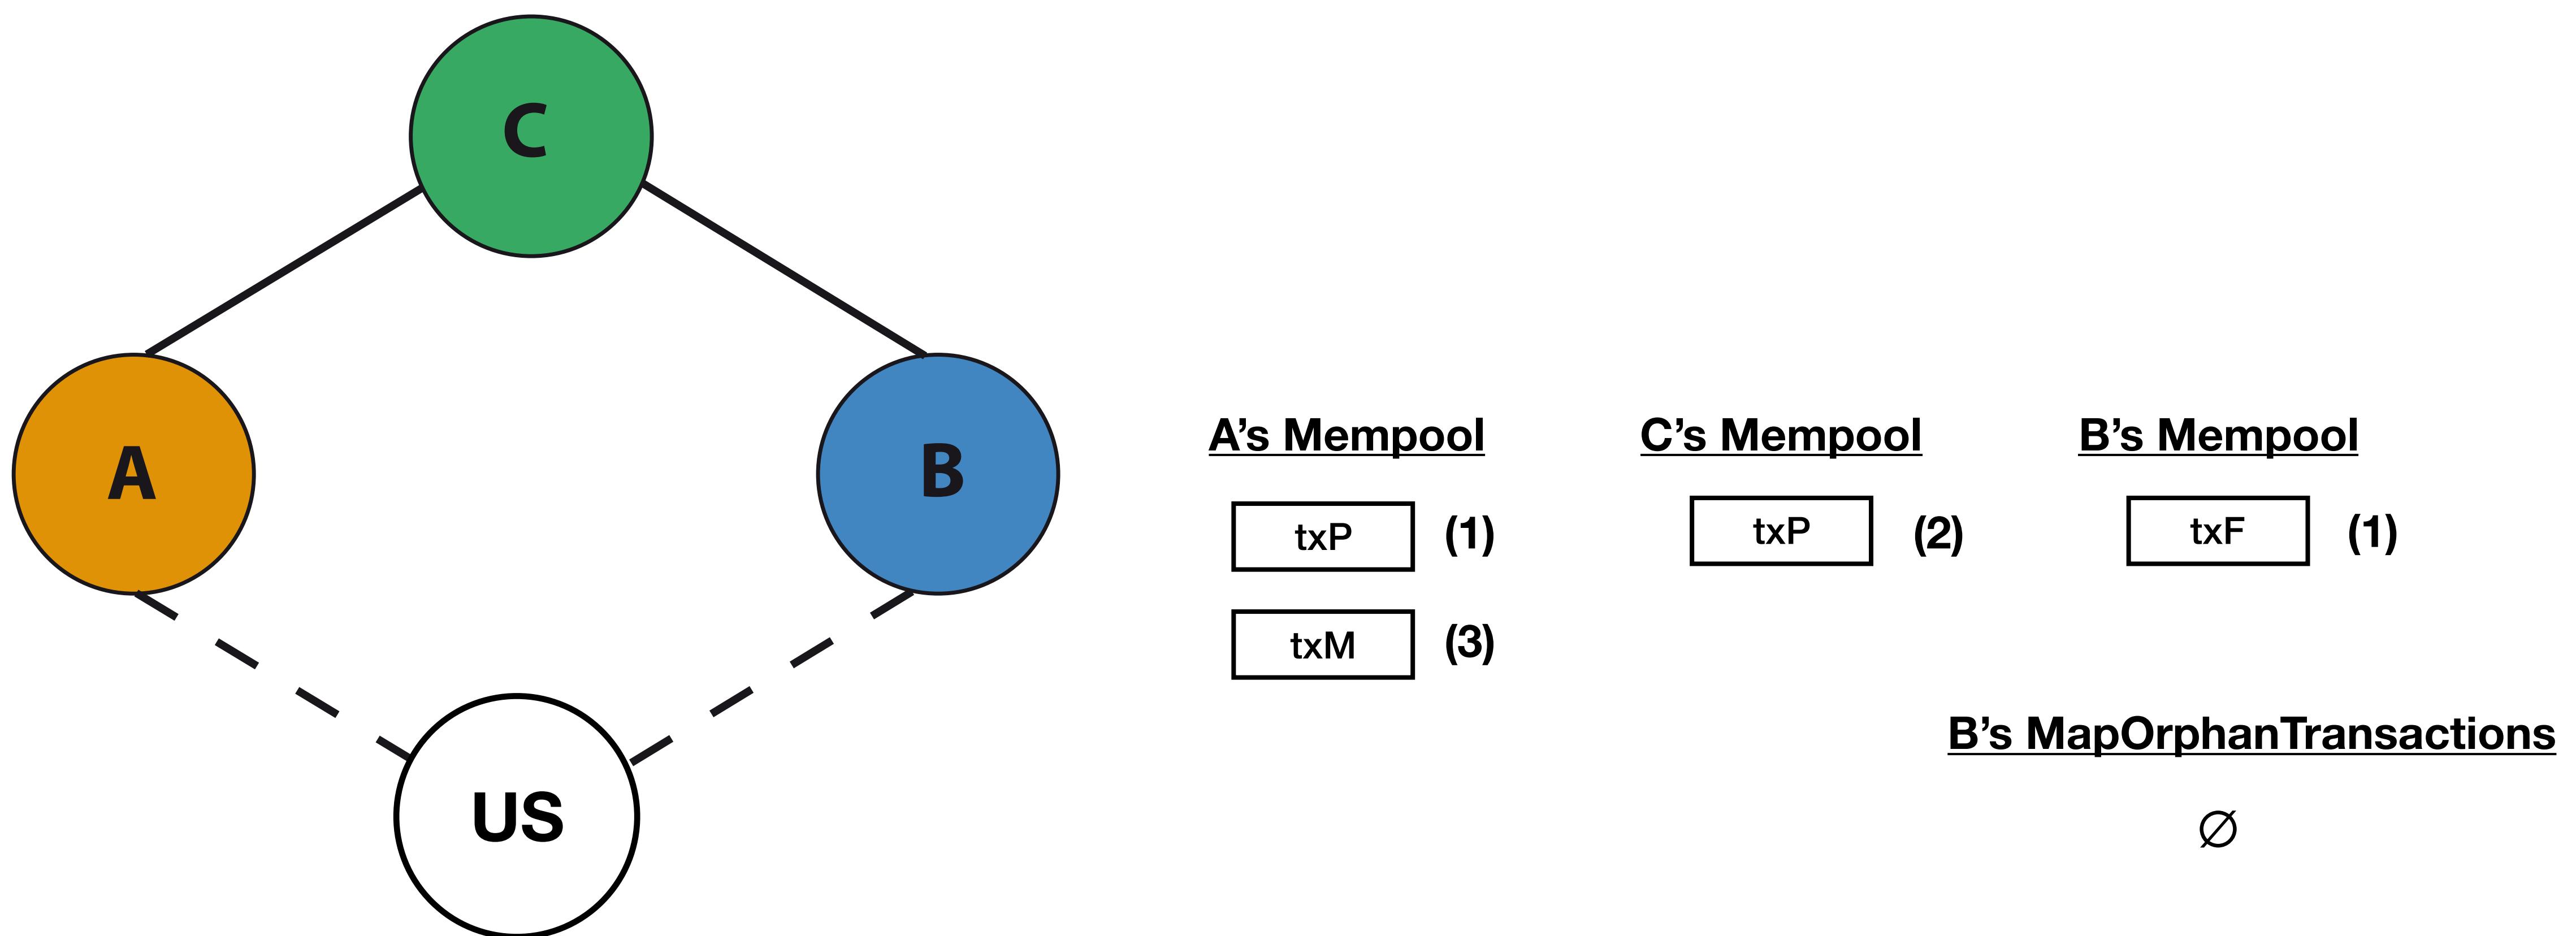
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



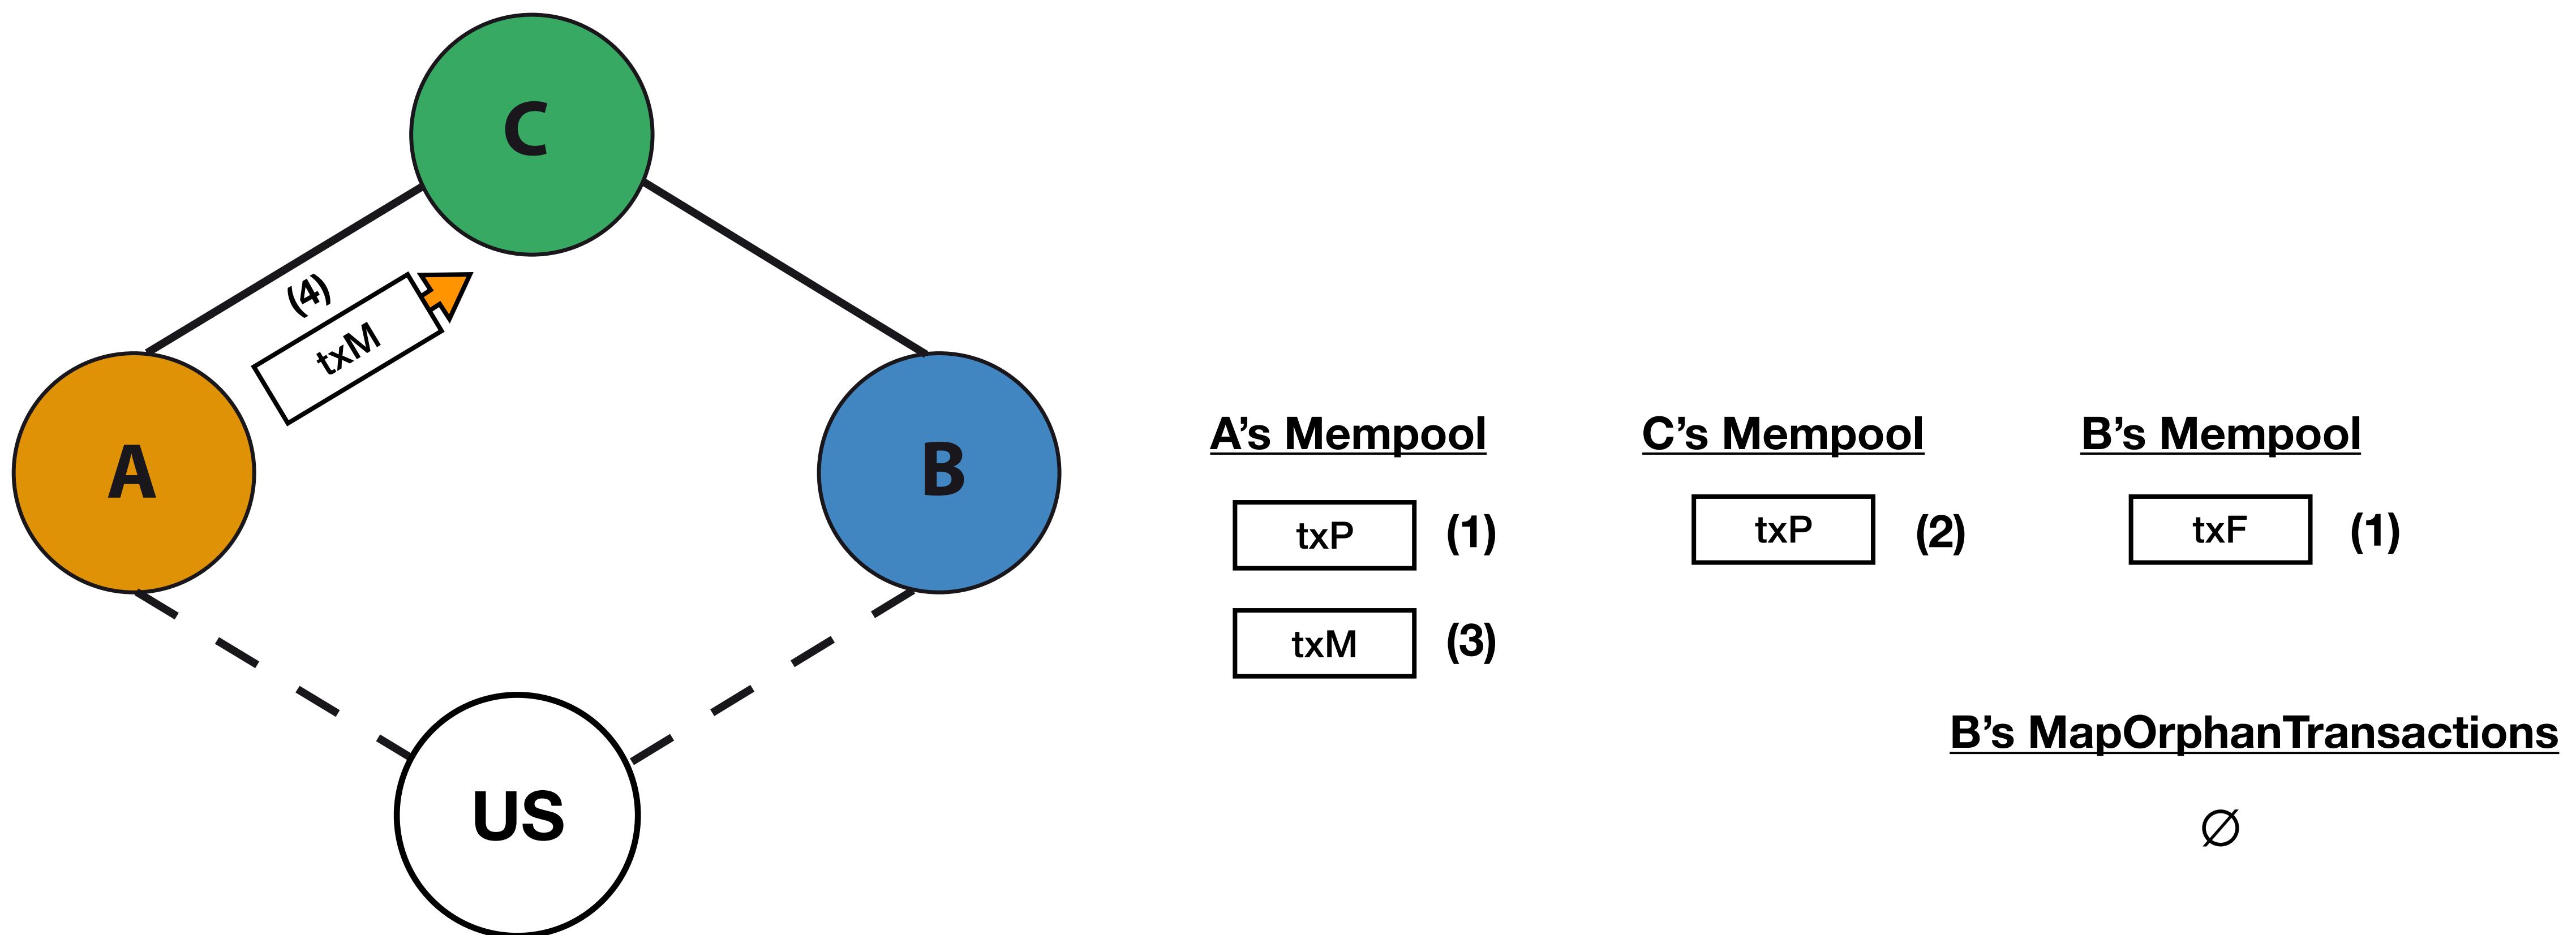
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



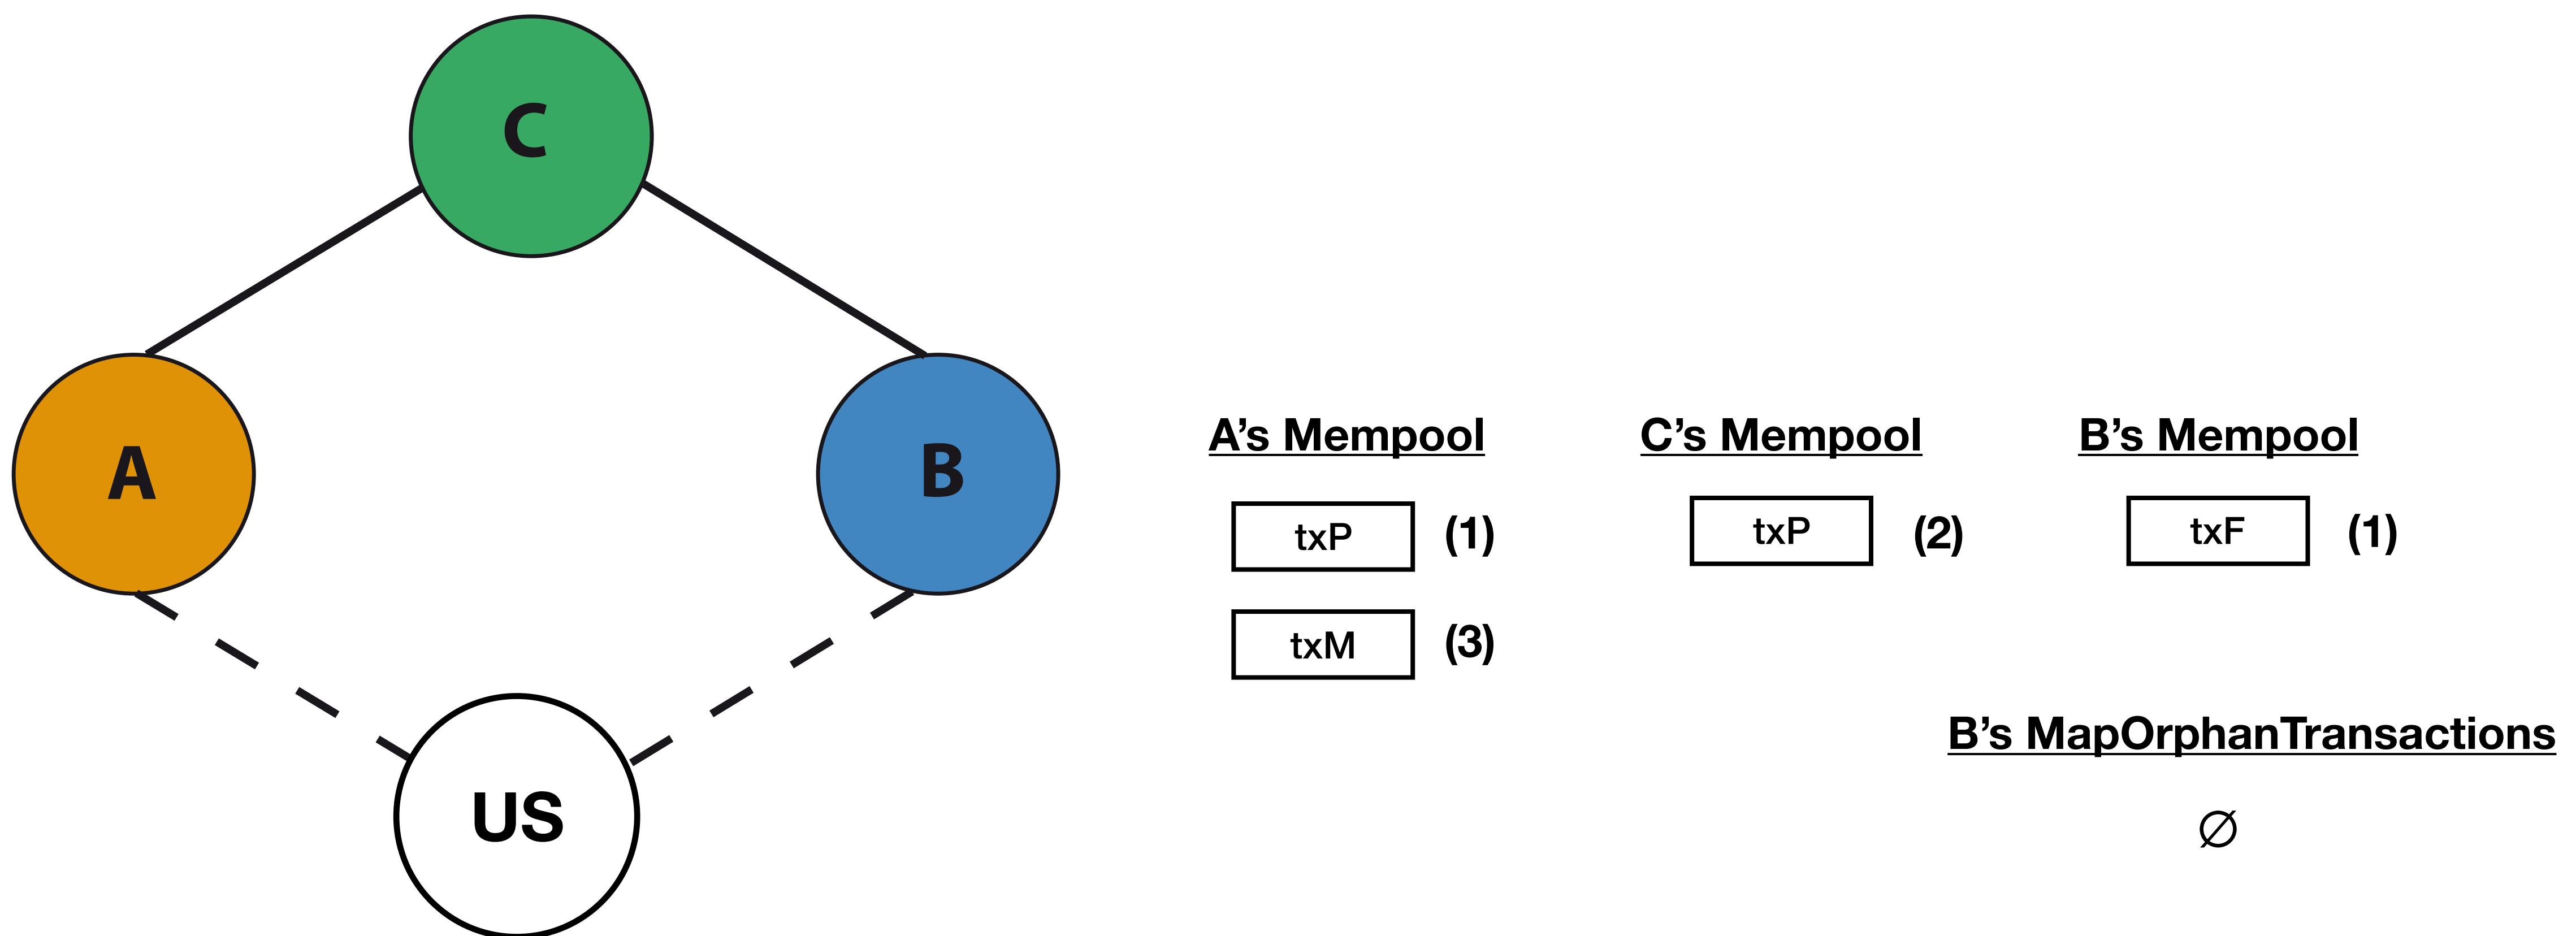
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



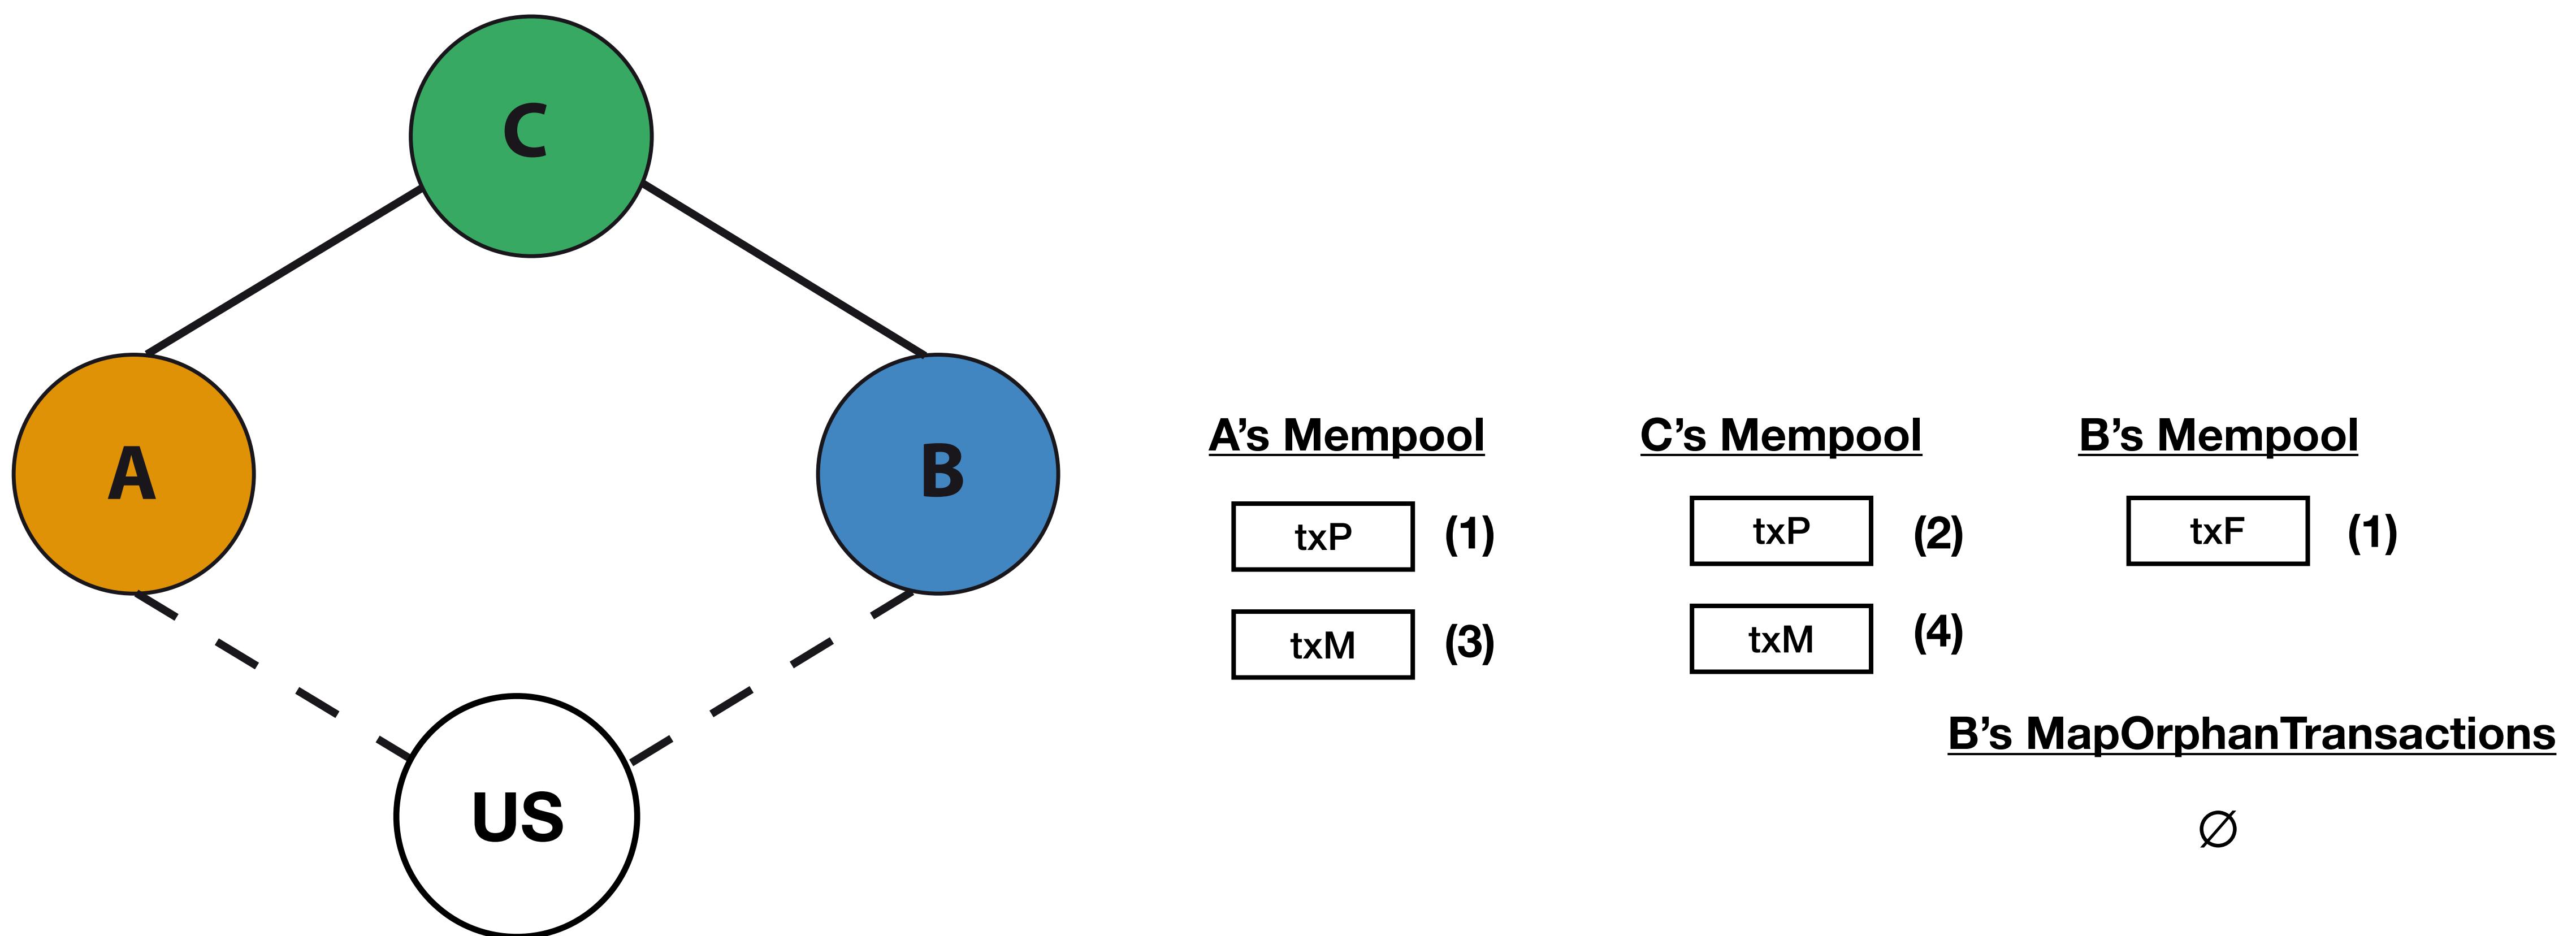
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



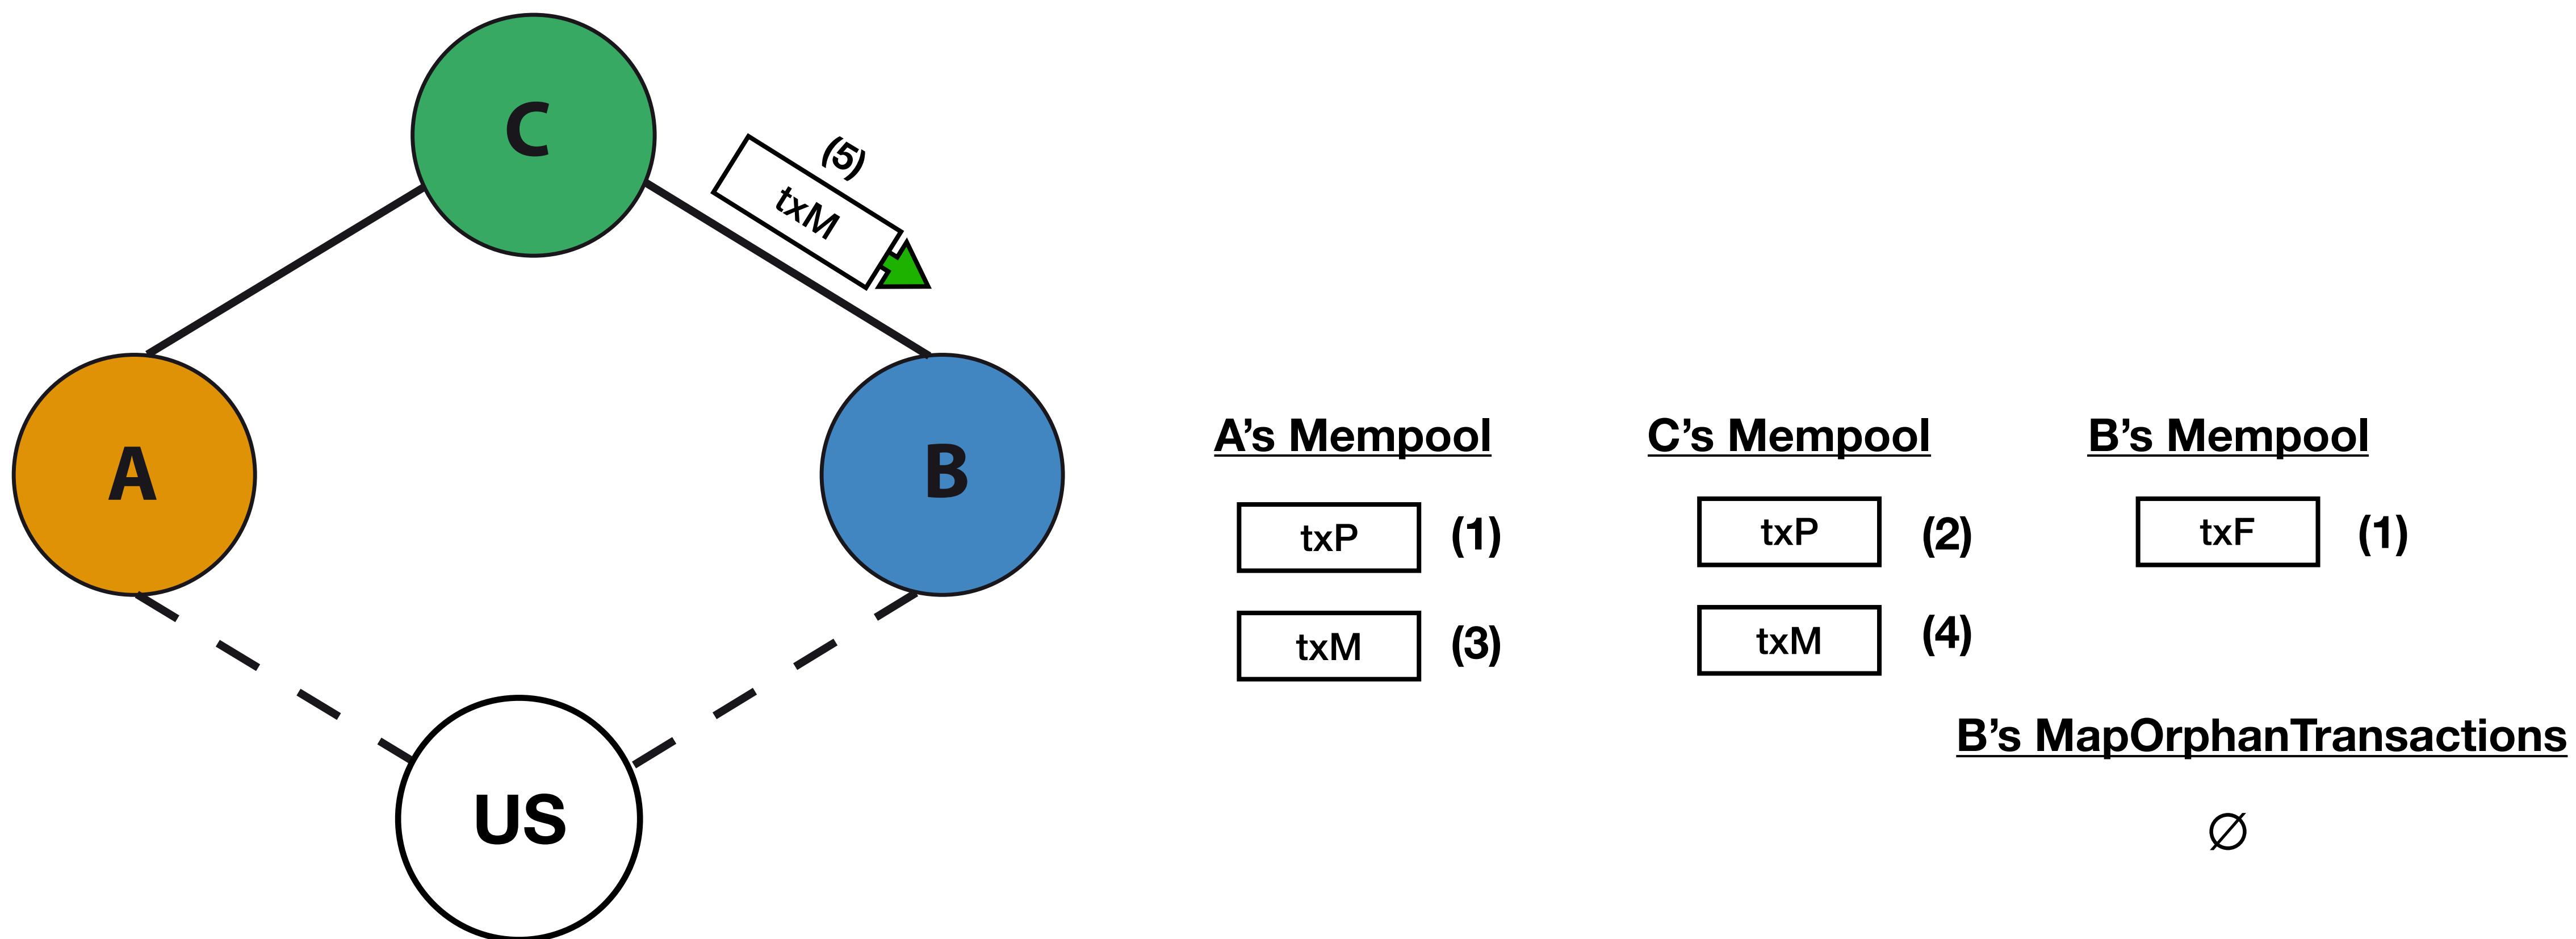
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



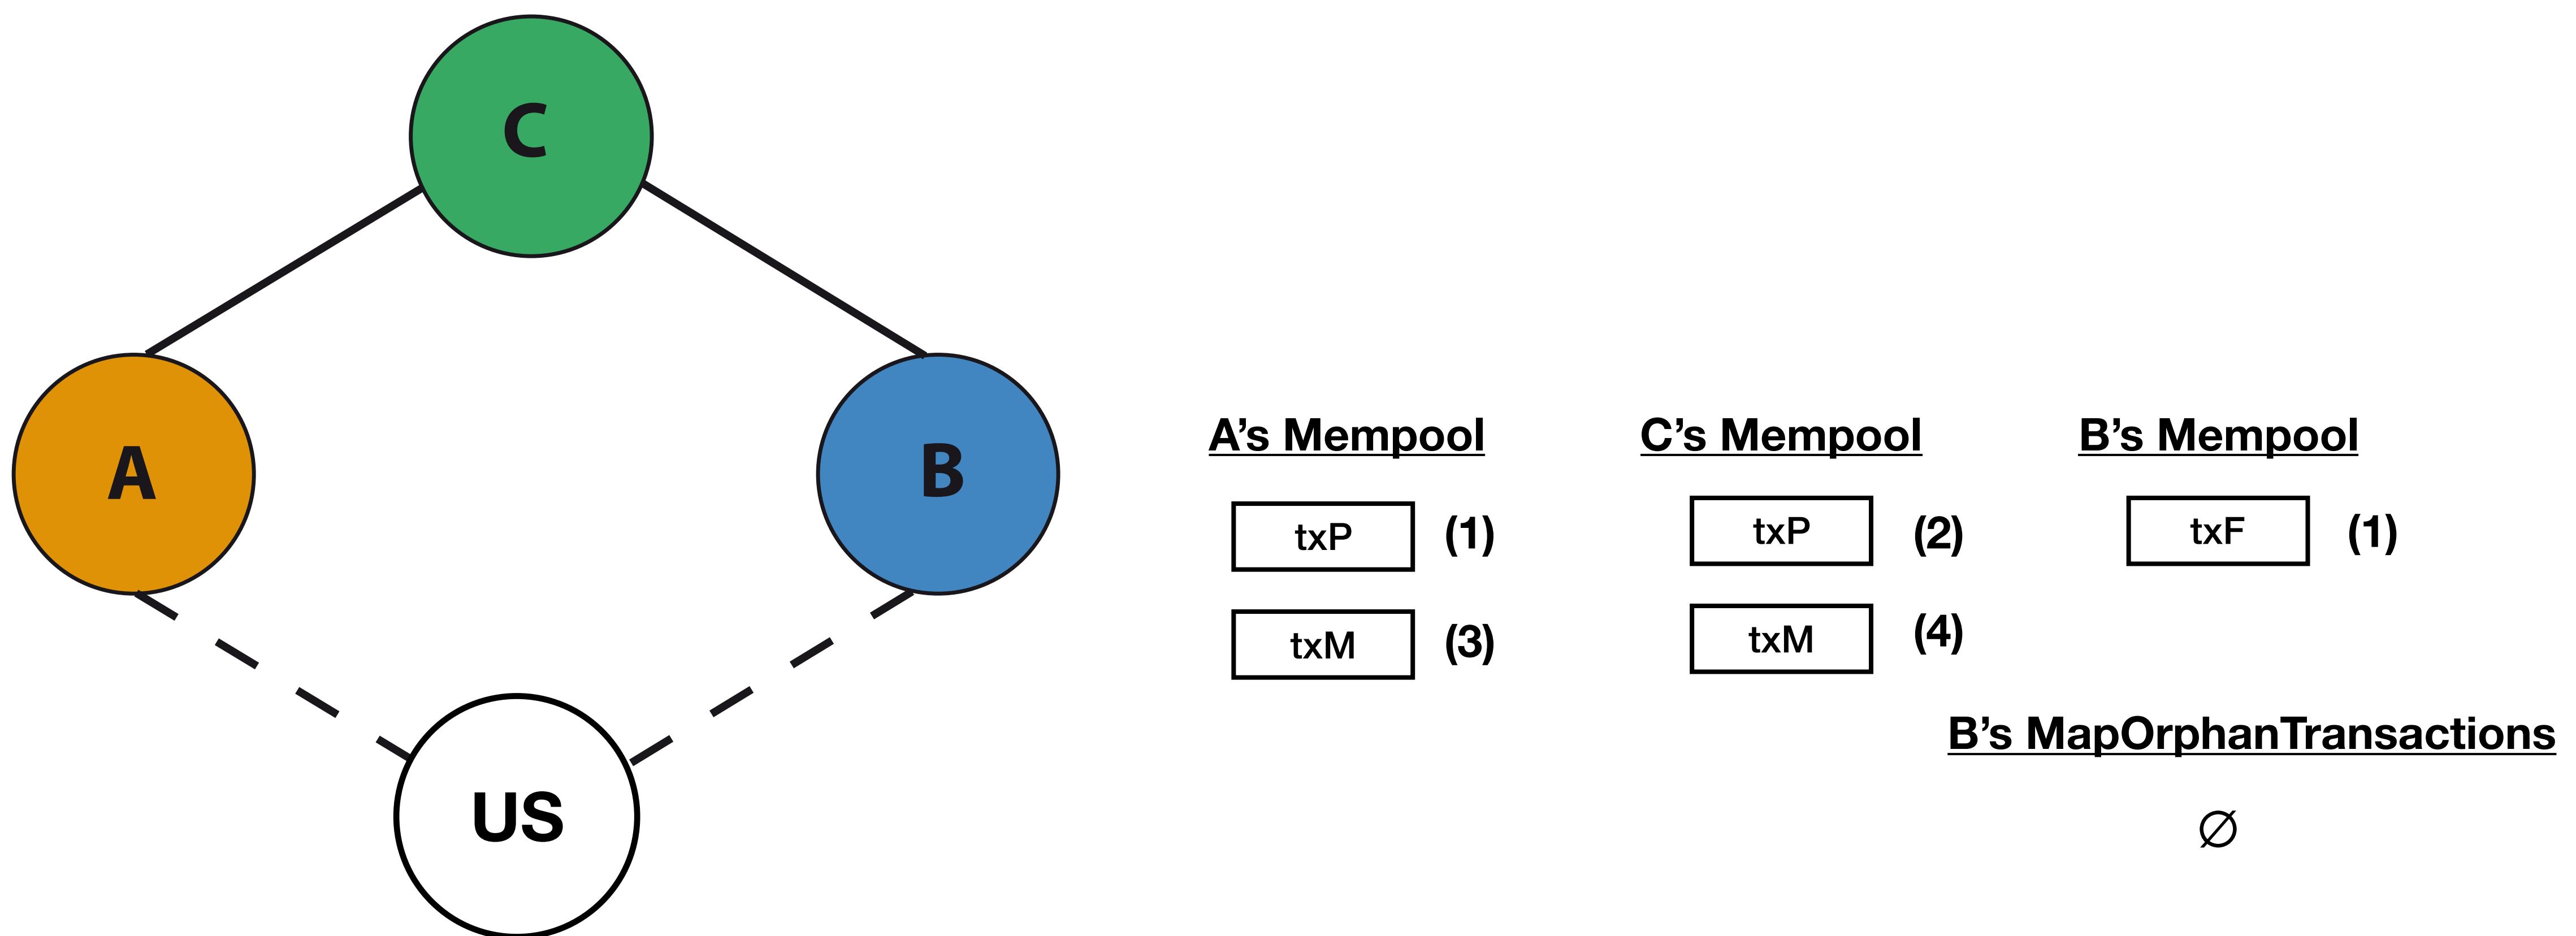
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



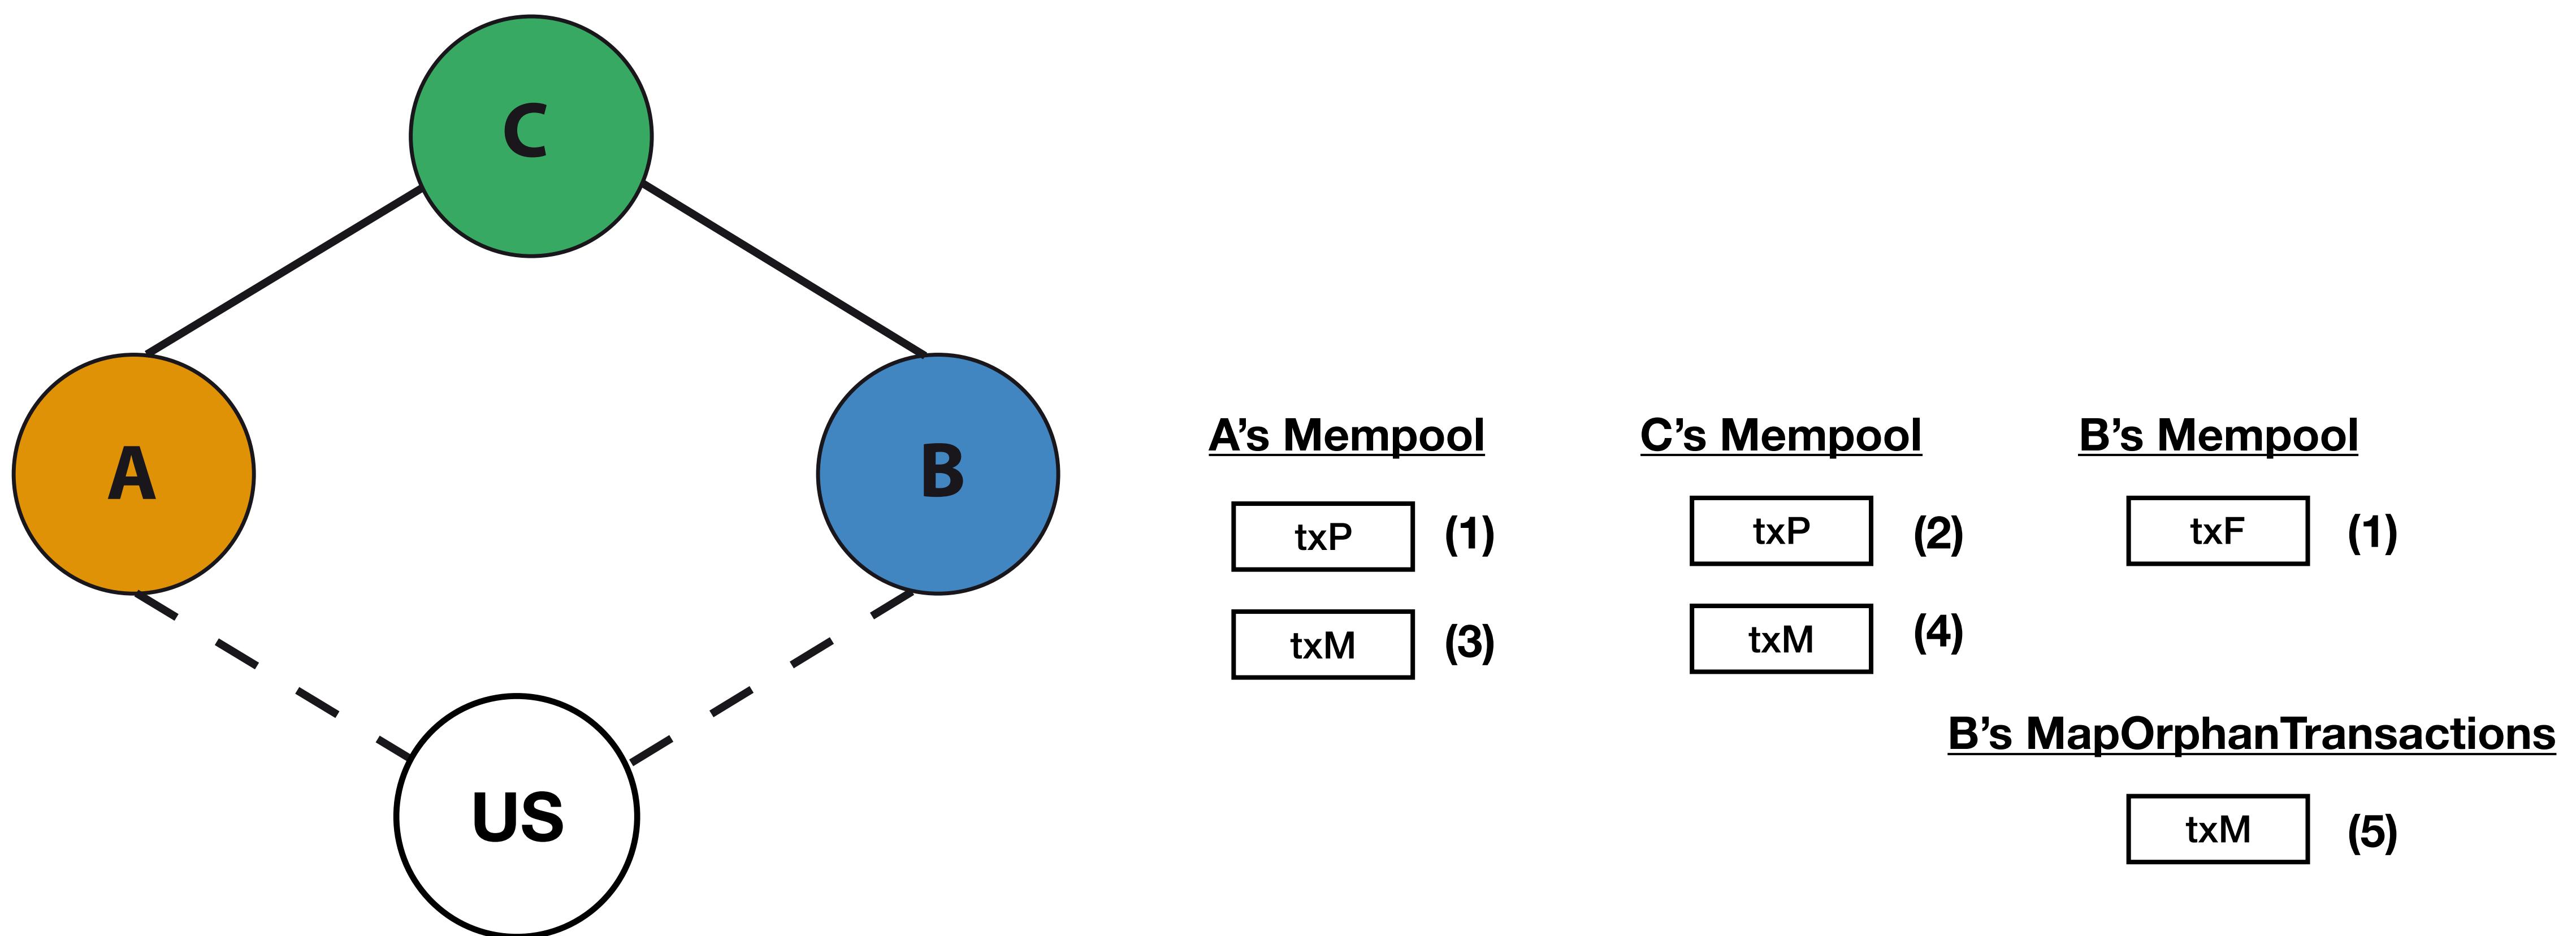
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture



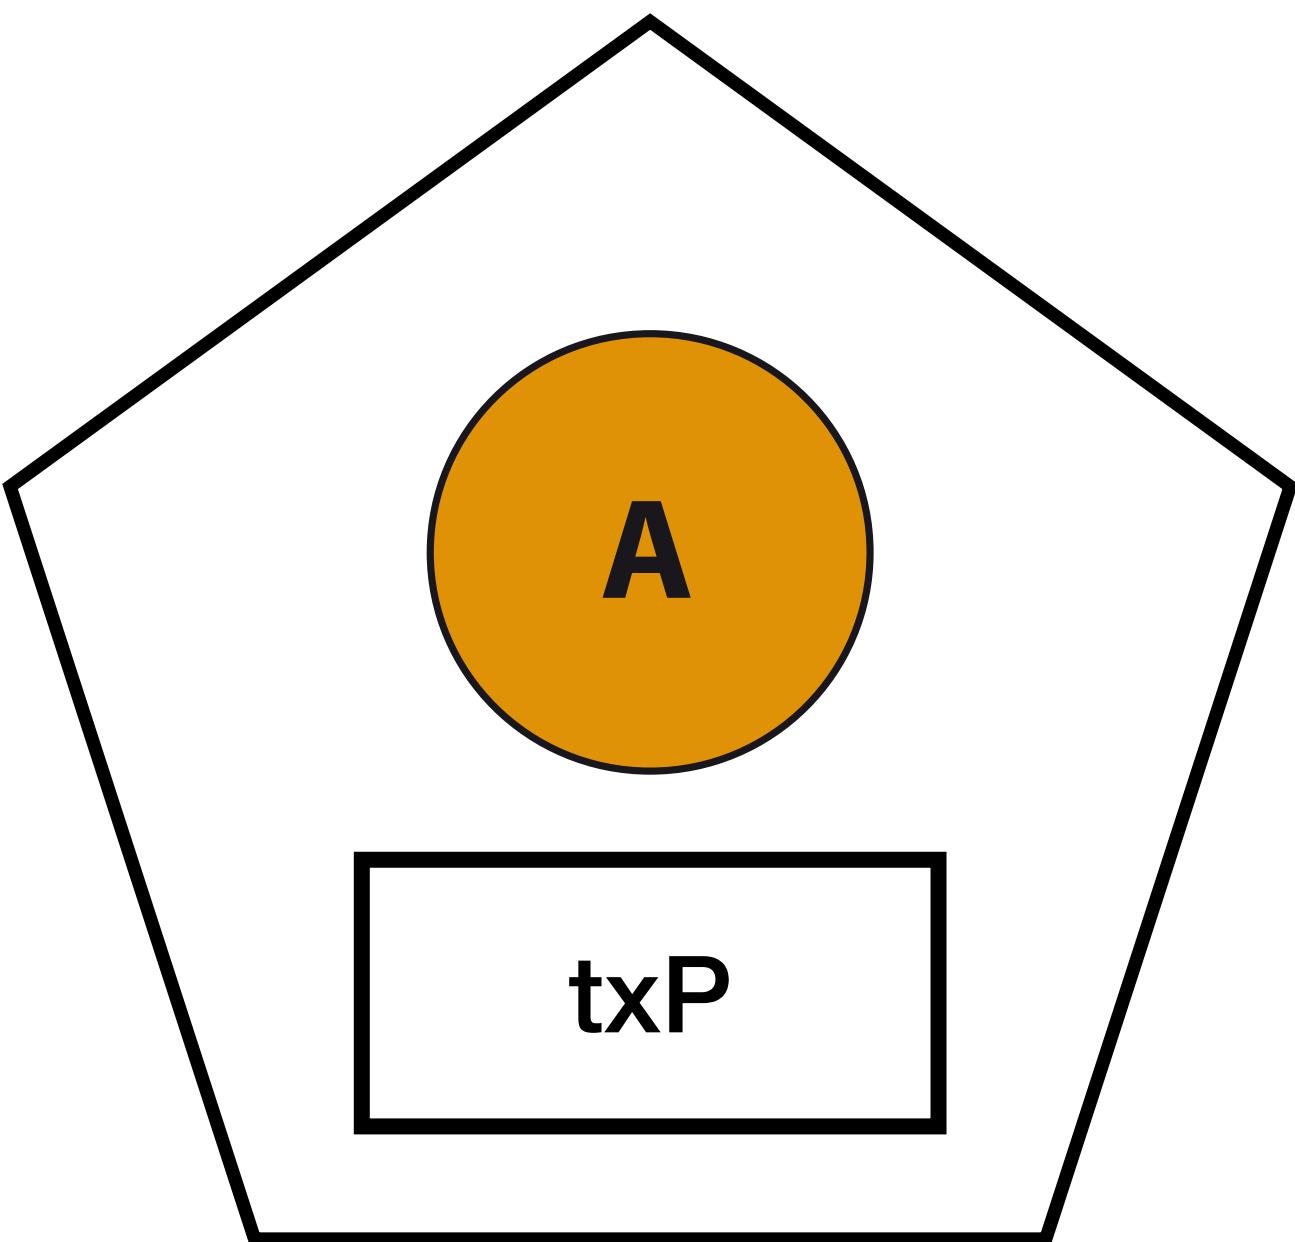
ITS NOT THAT EASY

Long story short, the technique will fail if we add an additional node to the picture

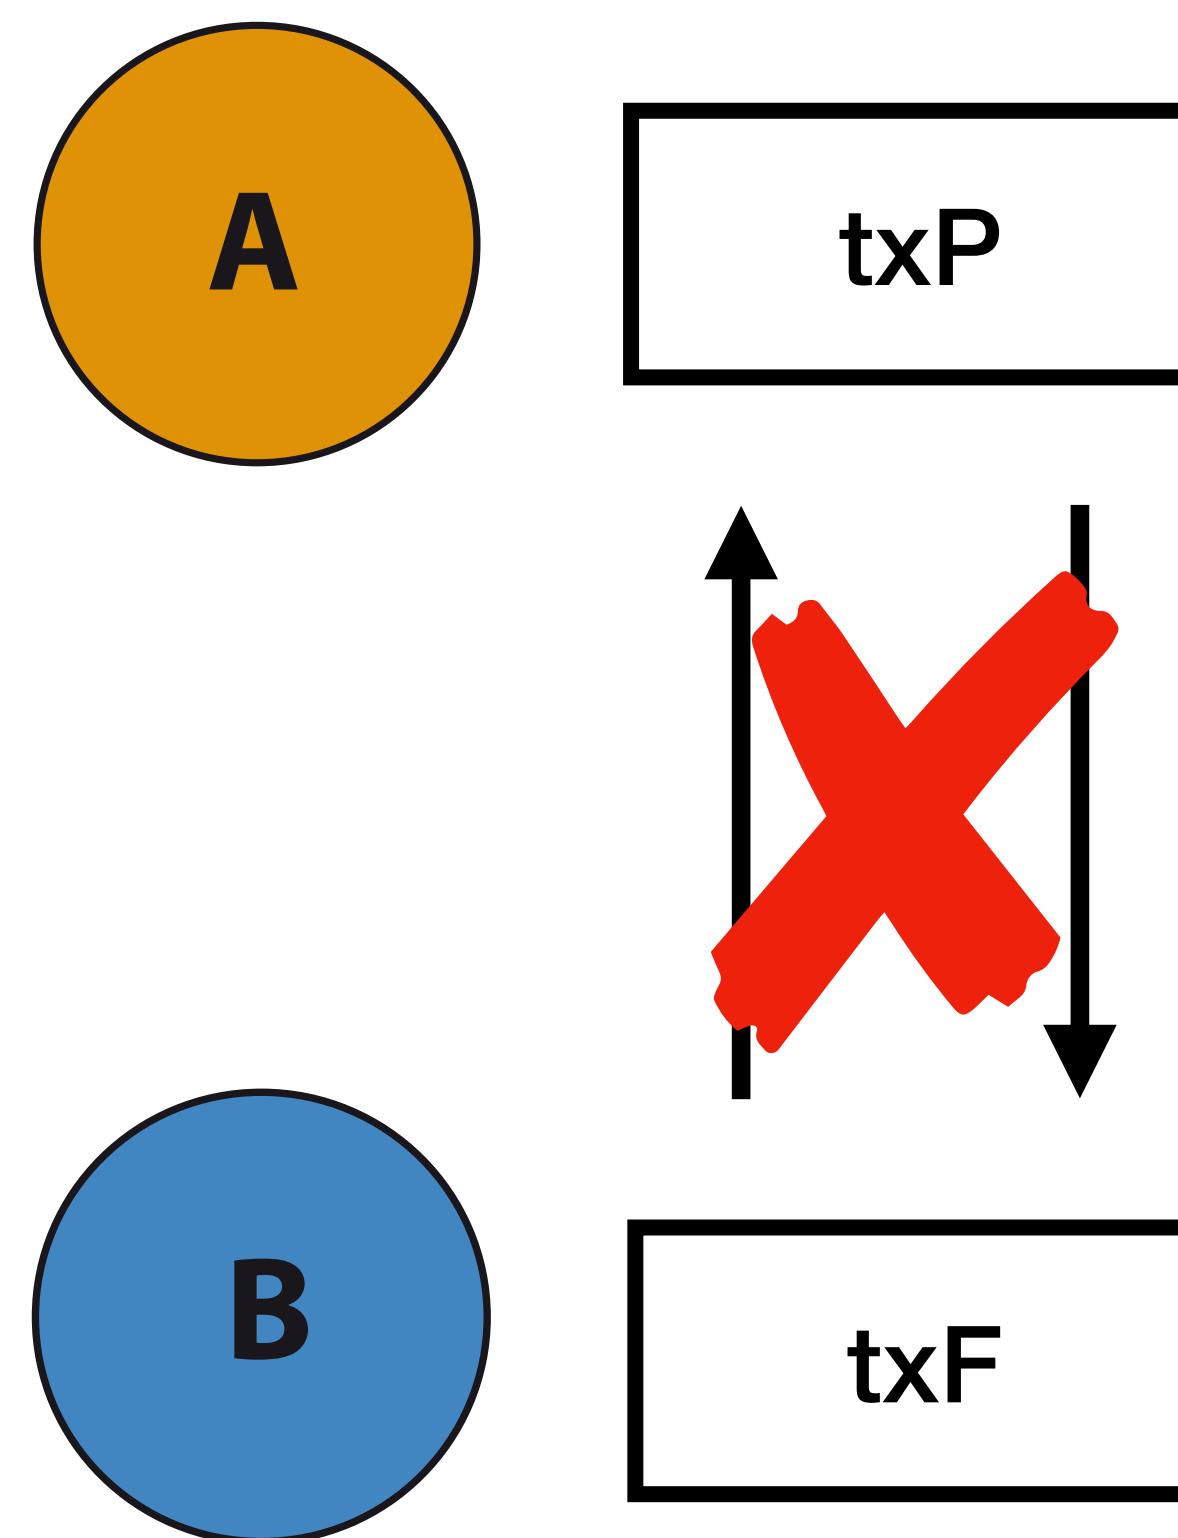


MAKE THIS WORK IN A REAL NETWORK

Isolation



Synchrony



Efficiency

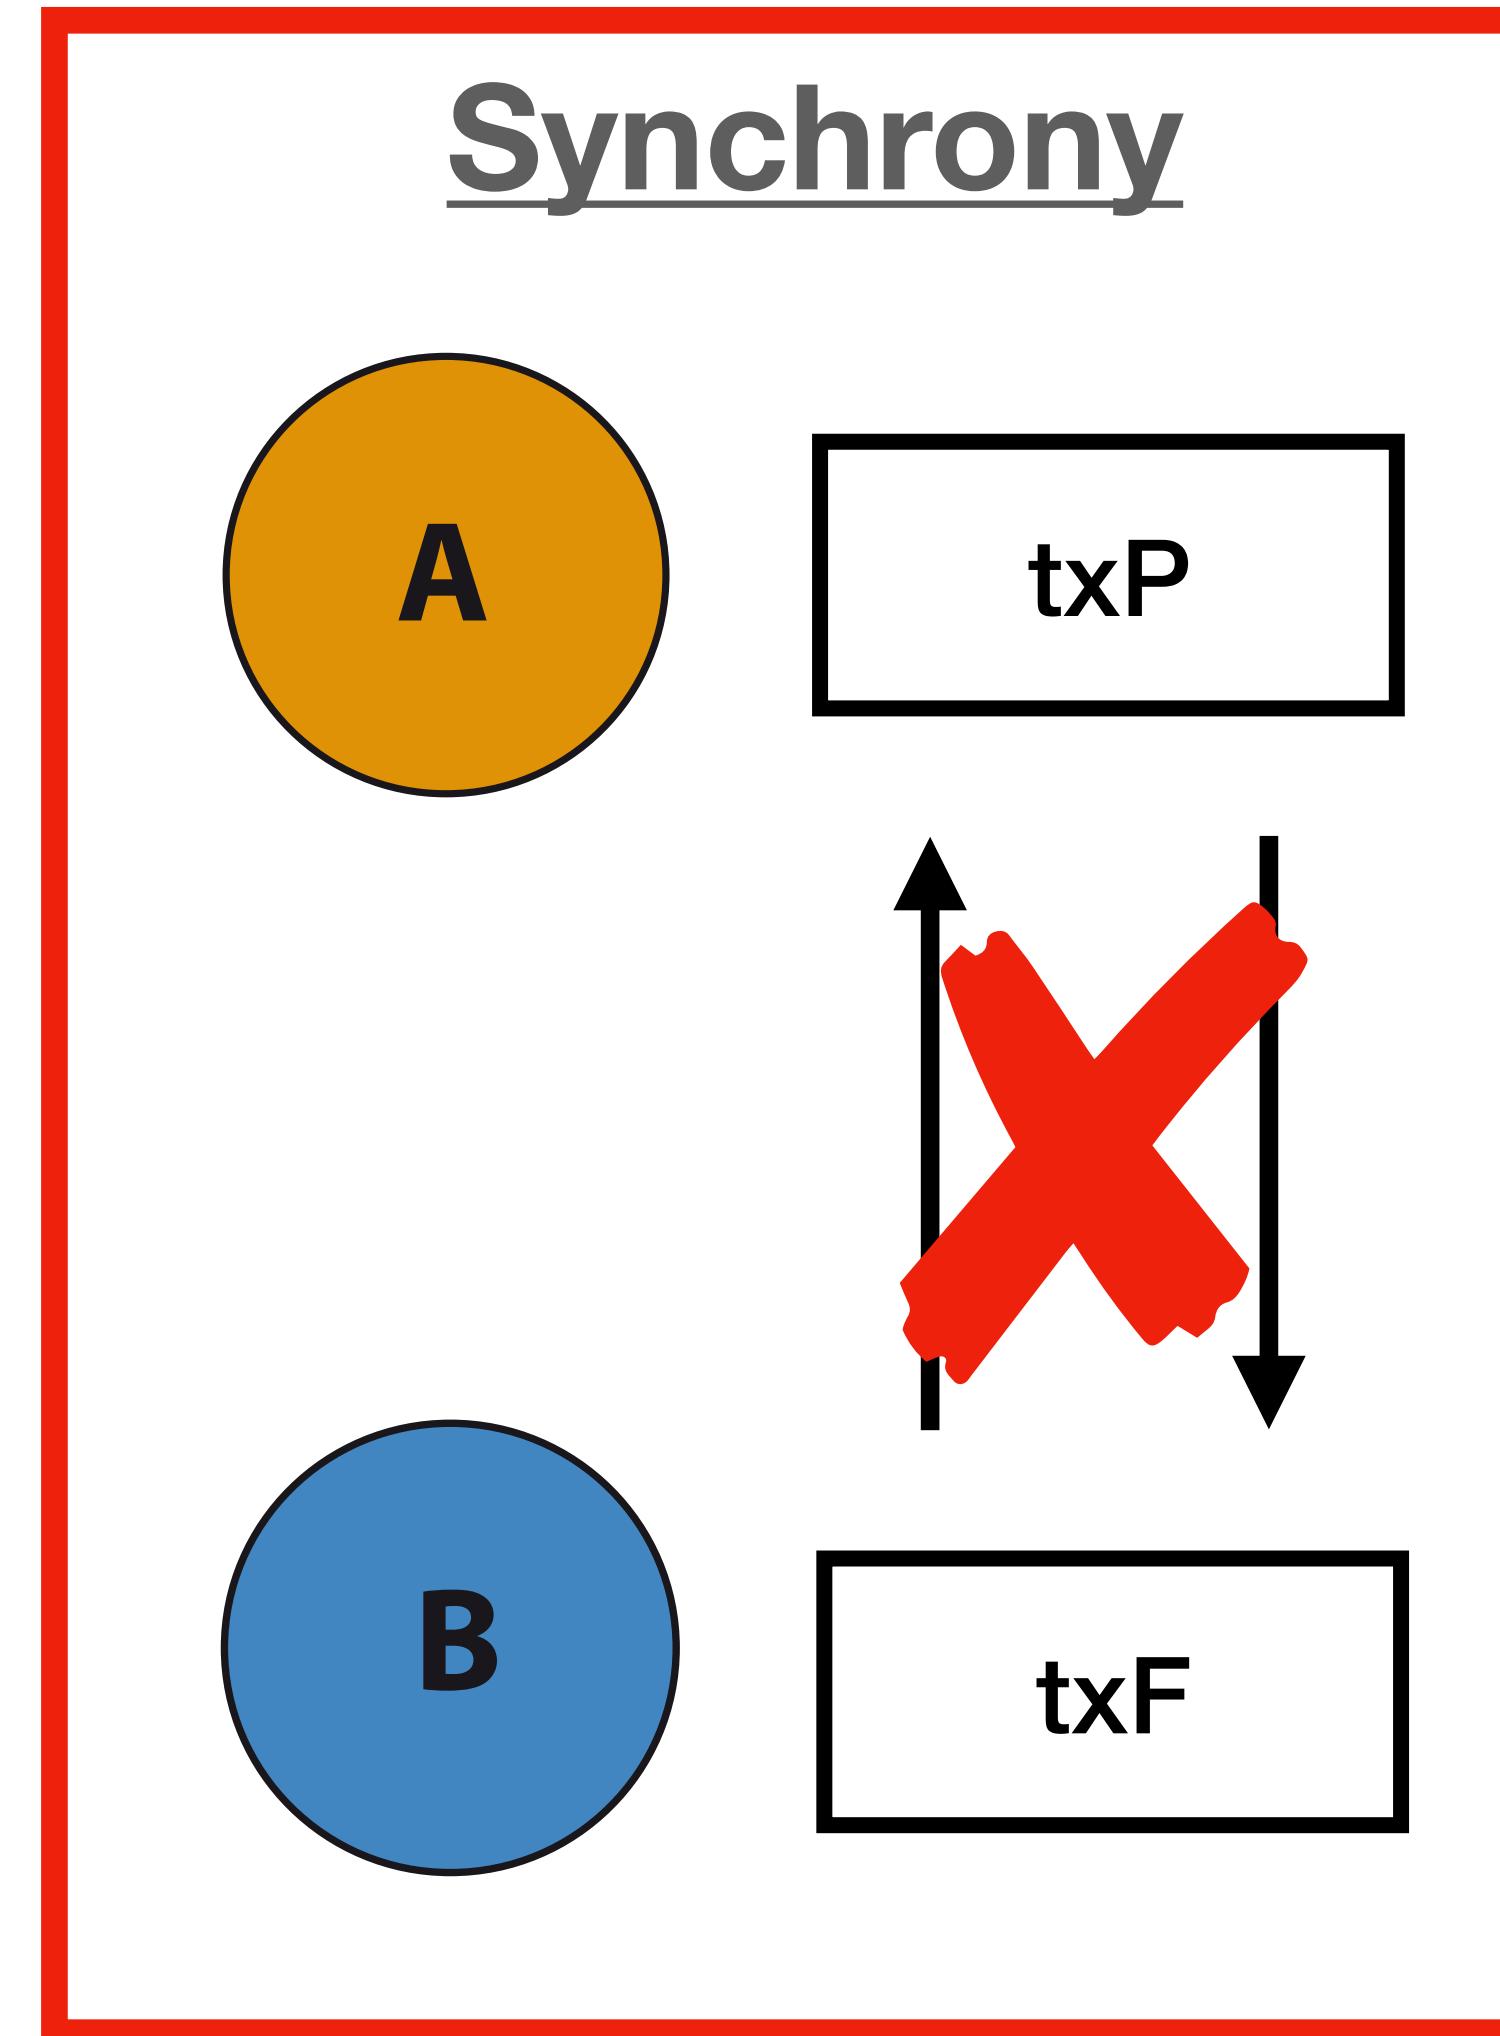
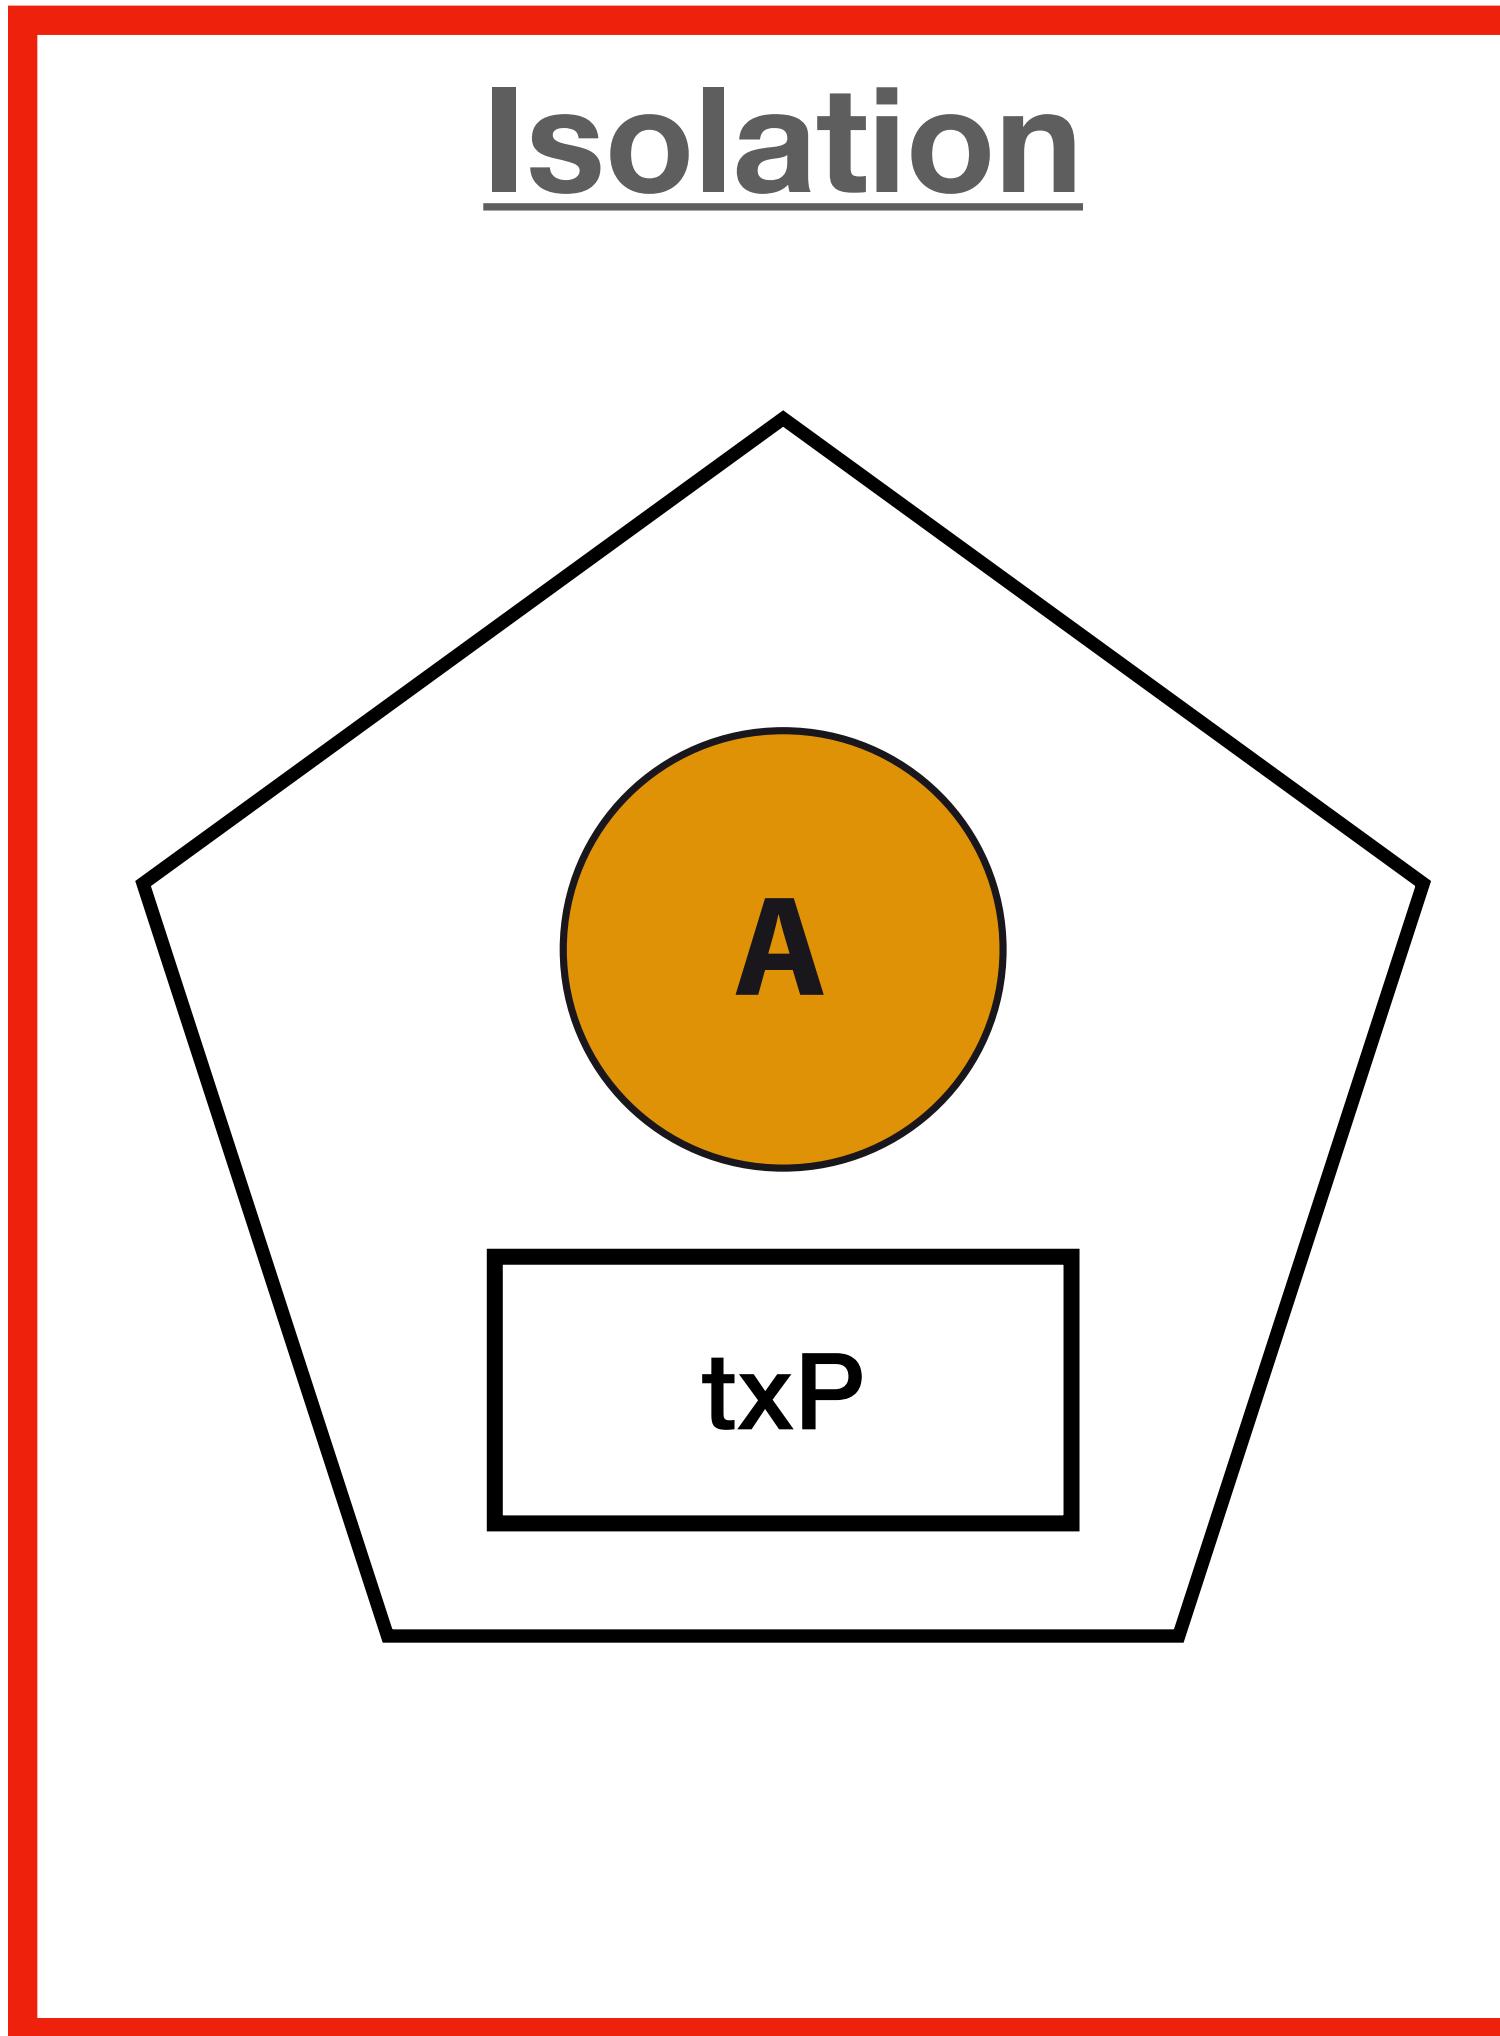
$\approx O(n)$

↓

$\approx O(\sqrt{n})$

$n = \#nodes$

MAKE THIS WORK IN A REAL NETWORK



Efficiency

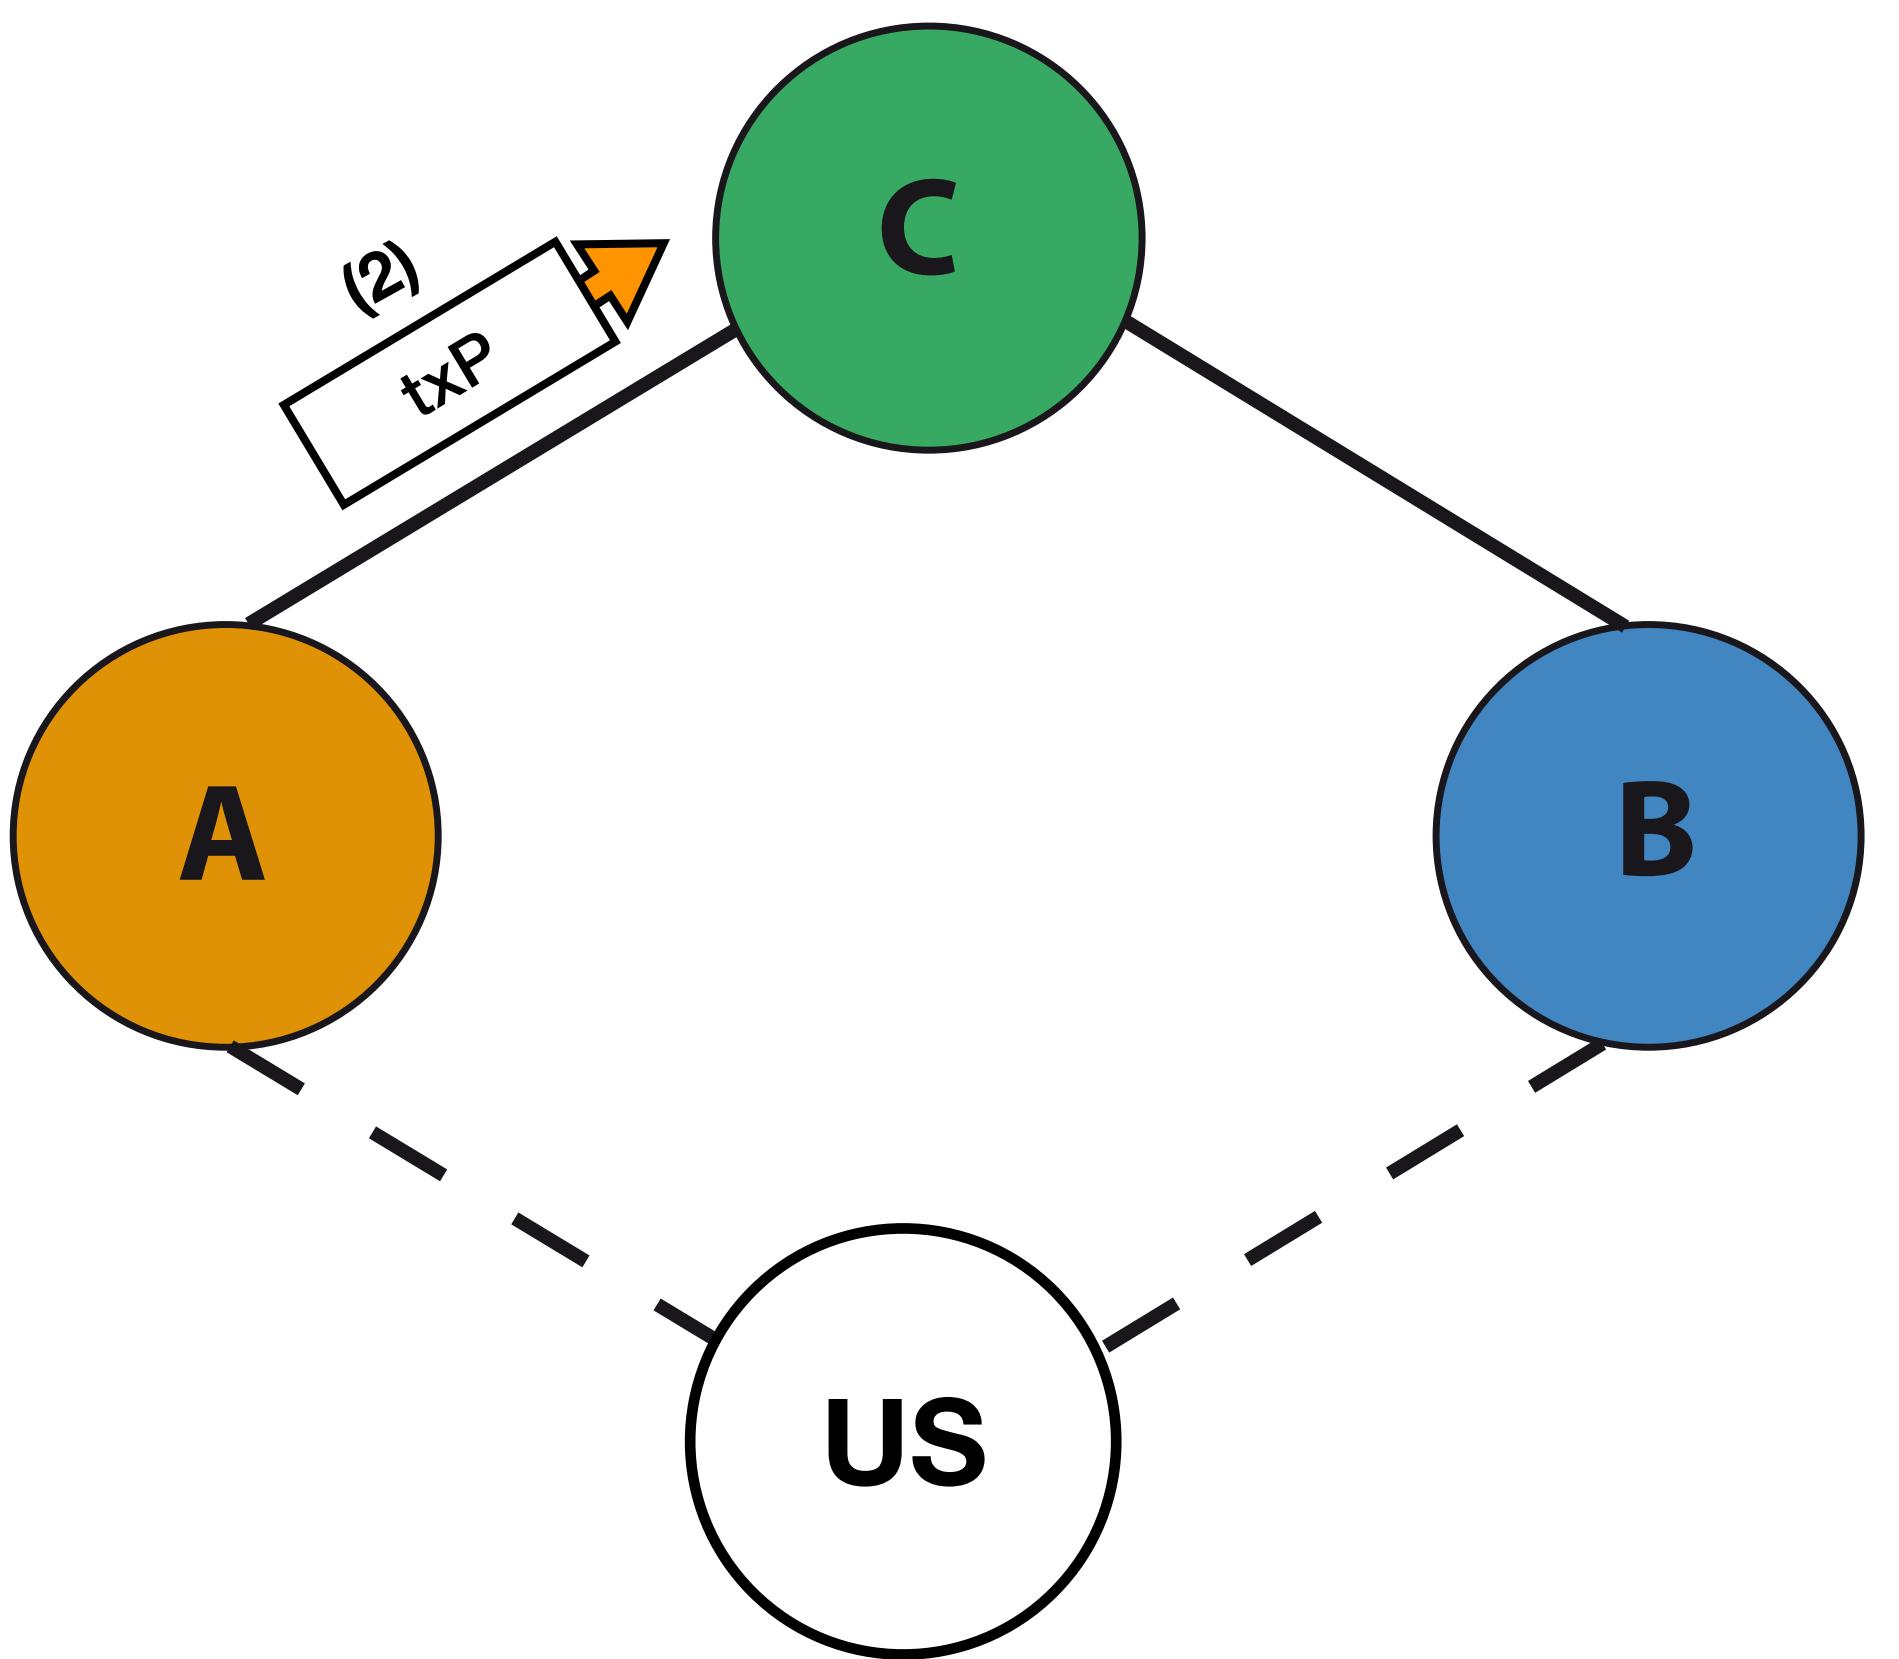
$\approx O(n)$

$\approx O(\sqrt{n})$

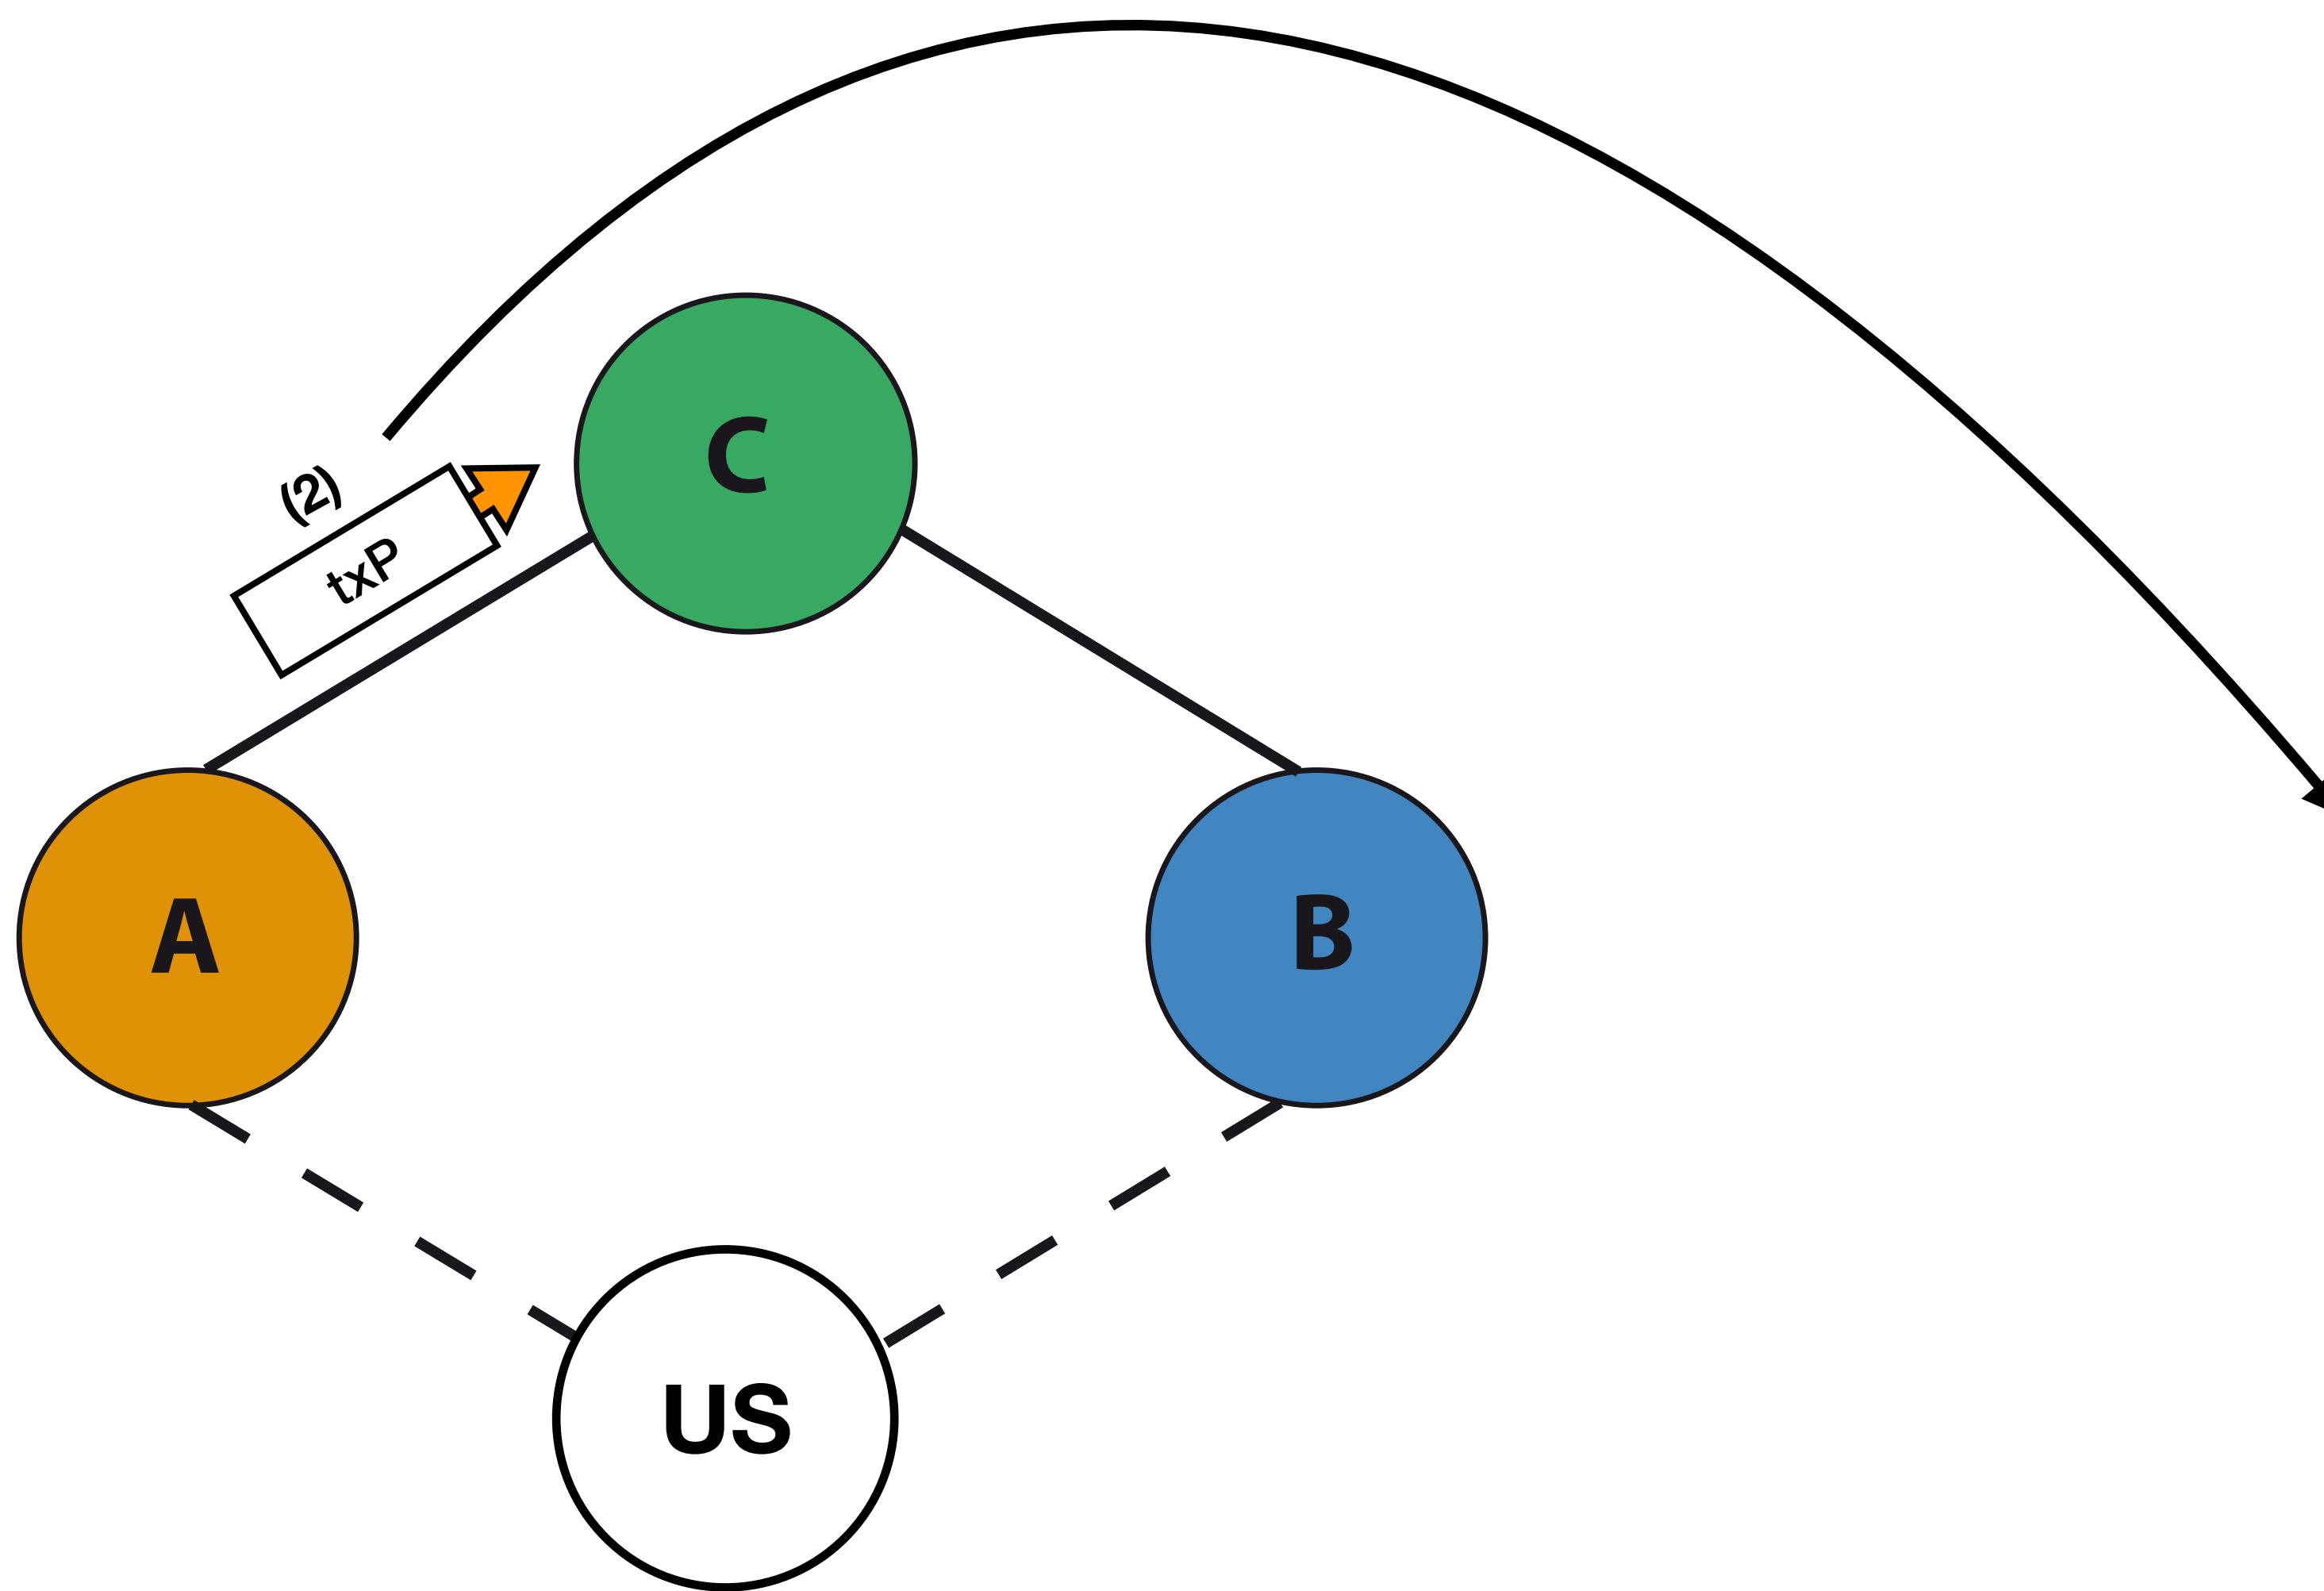
$n = \#nodes$

The text discusses efficiency levels. It starts with $\approx O(n)$, followed by a downward-pointing pink arrow, and then $\approx O(\sqrt{n})$. Below this, the text $n = \#nodes$ is written.

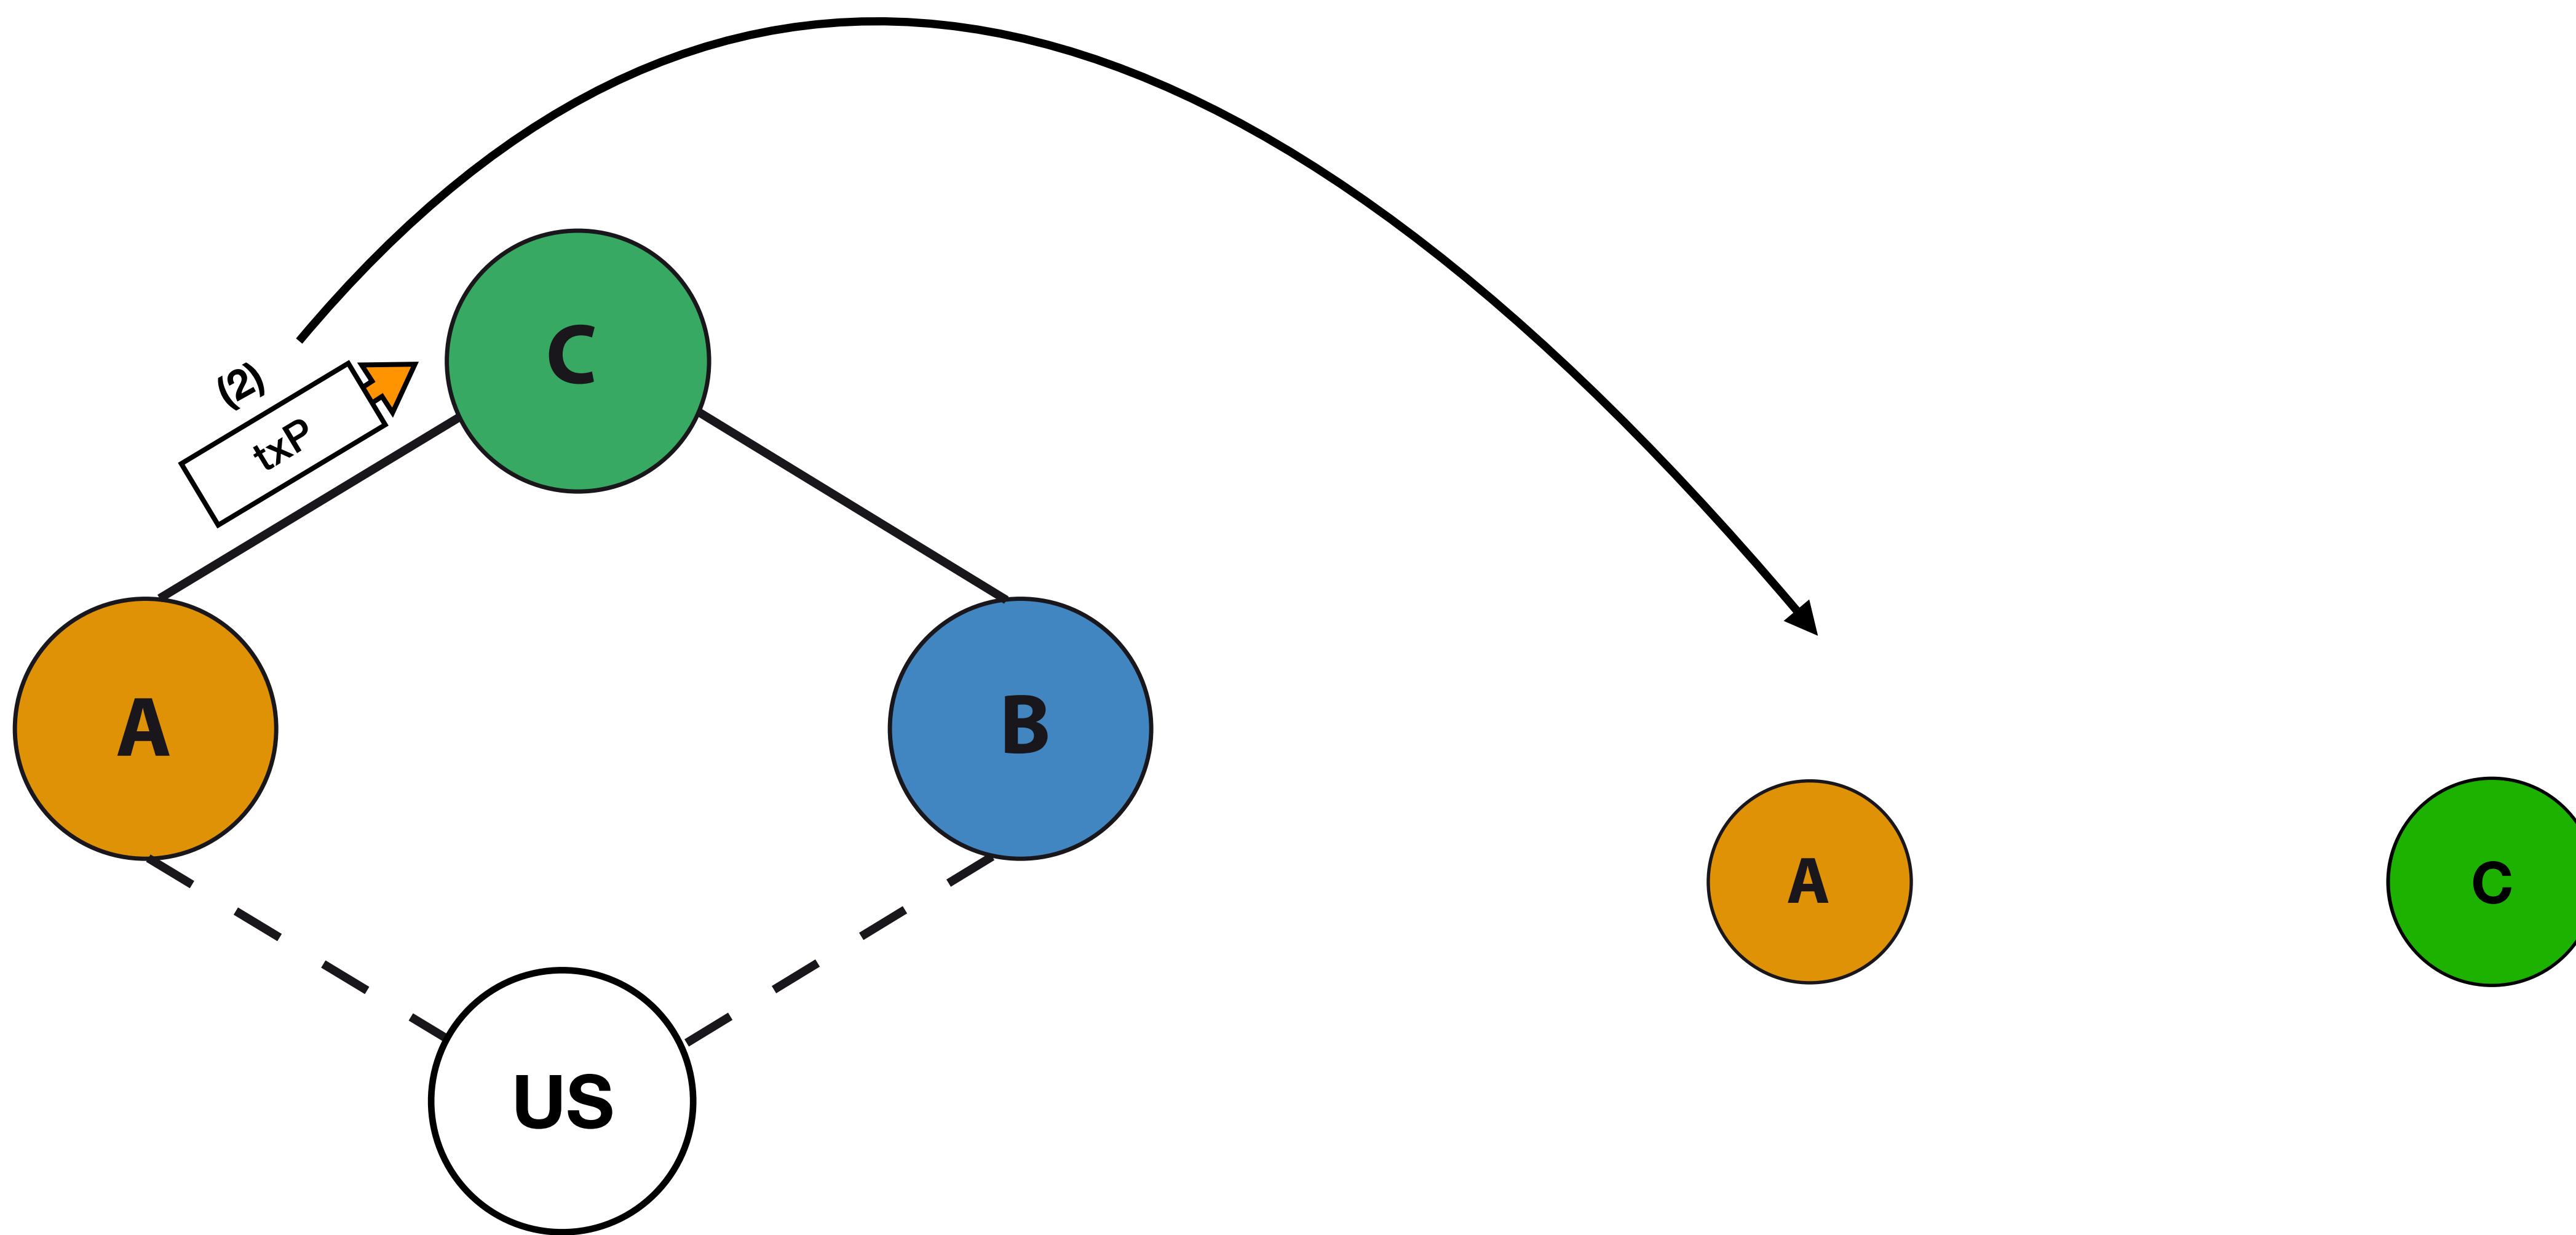
ISOLATION



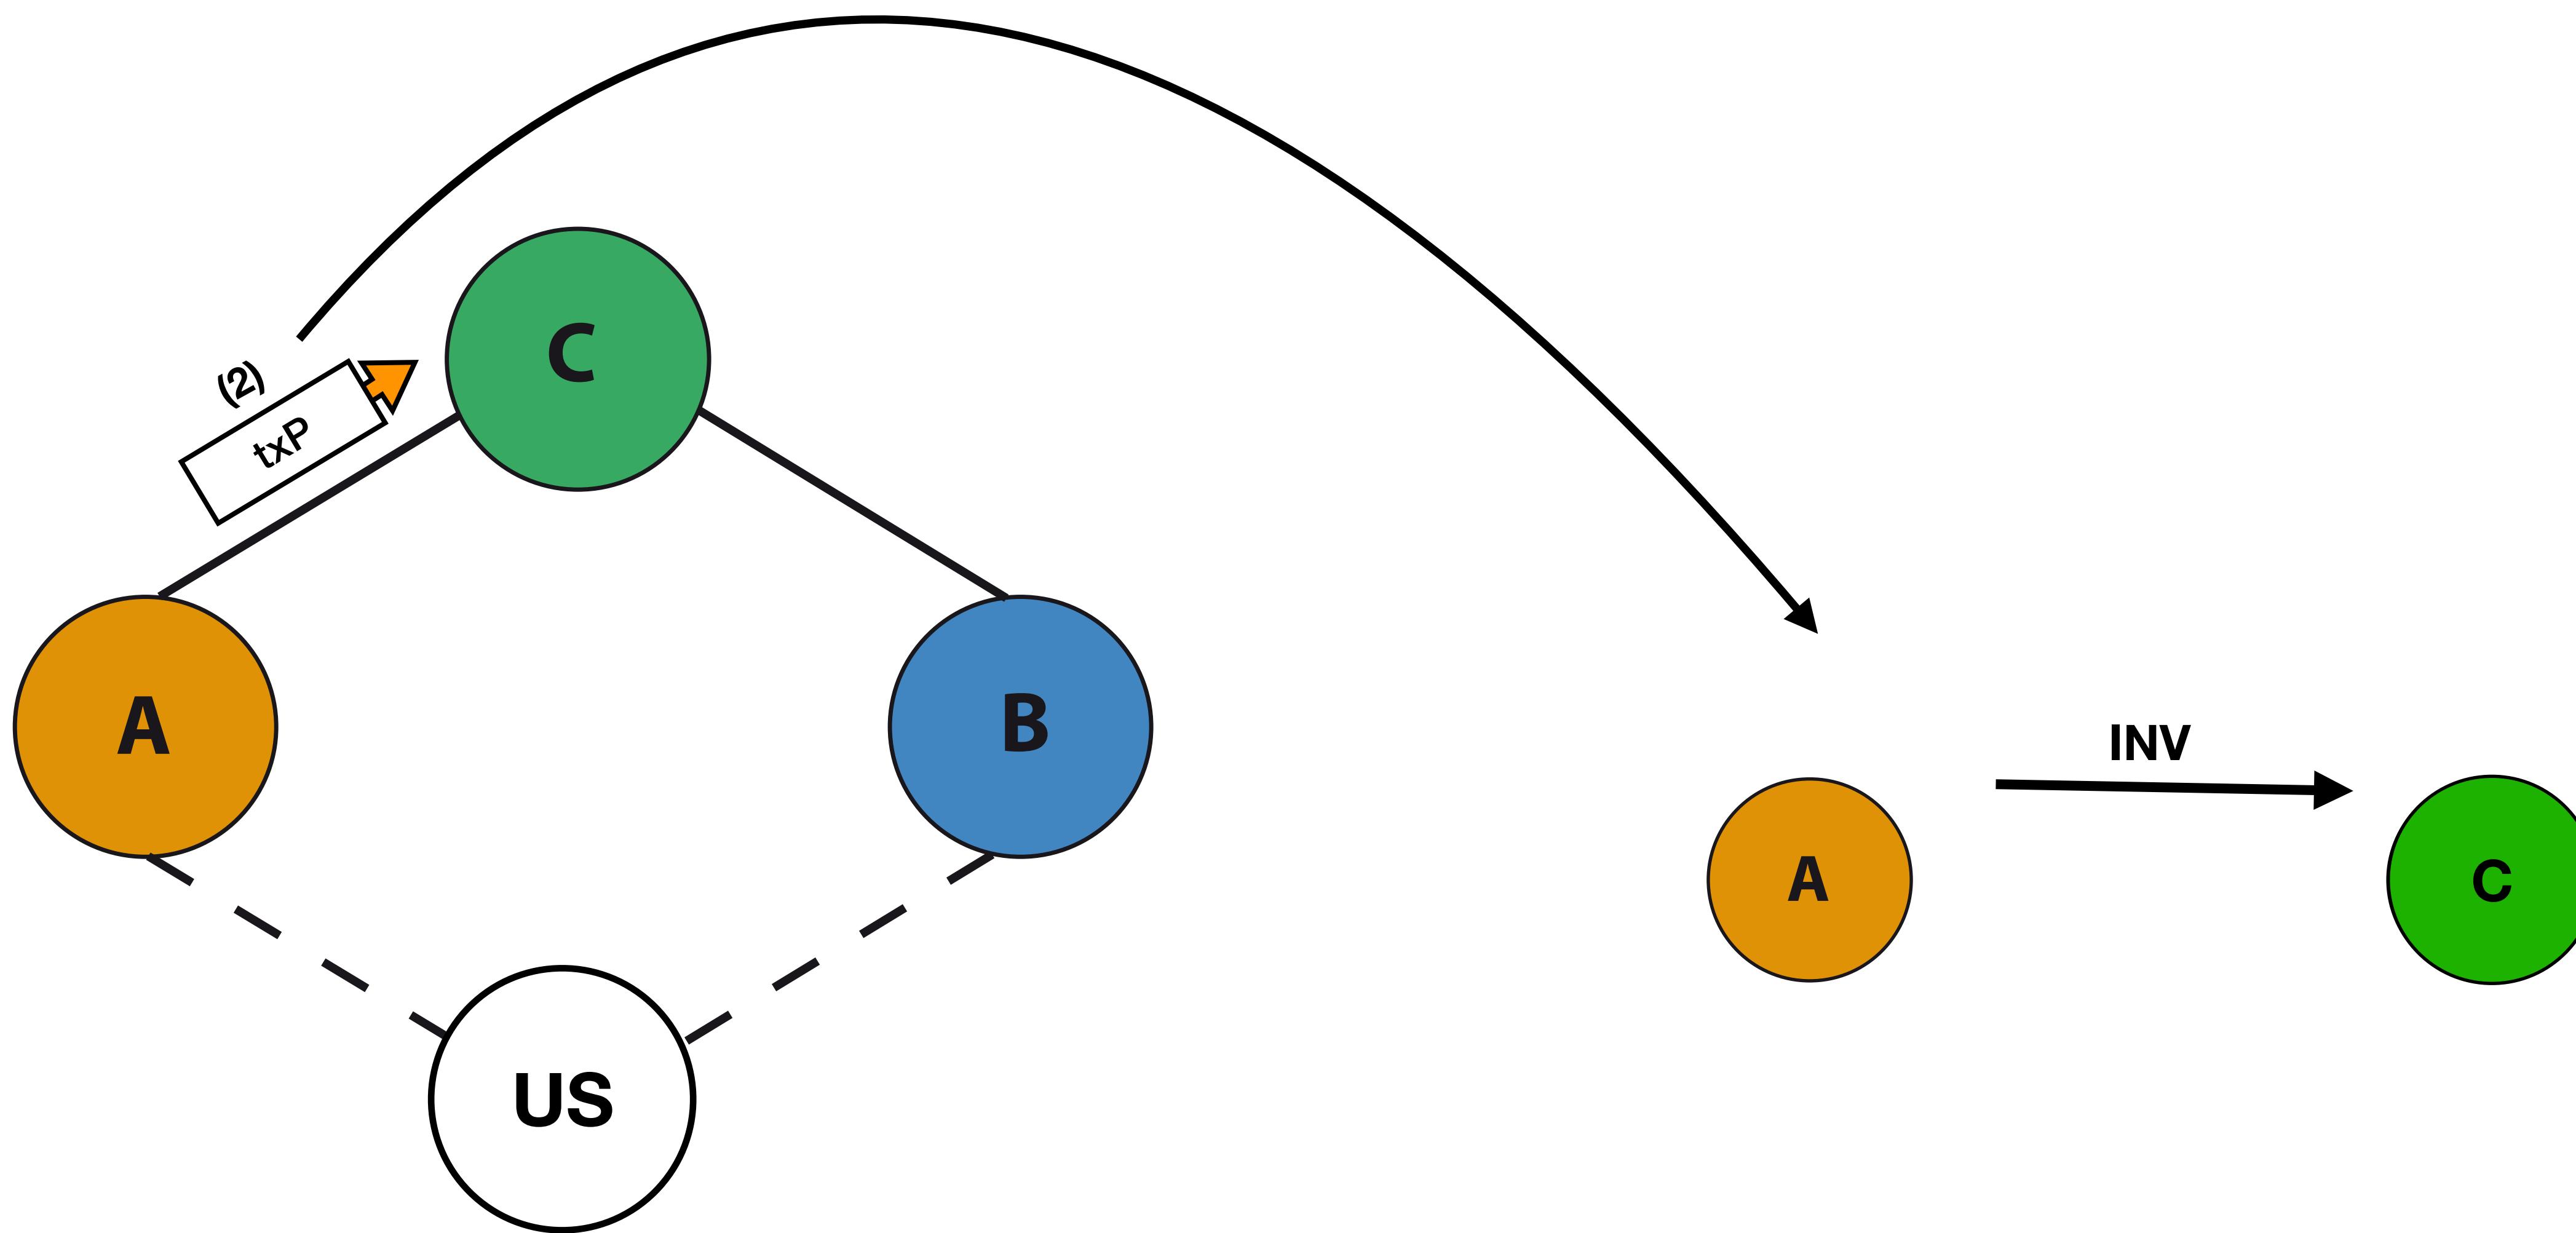
ISOLATION



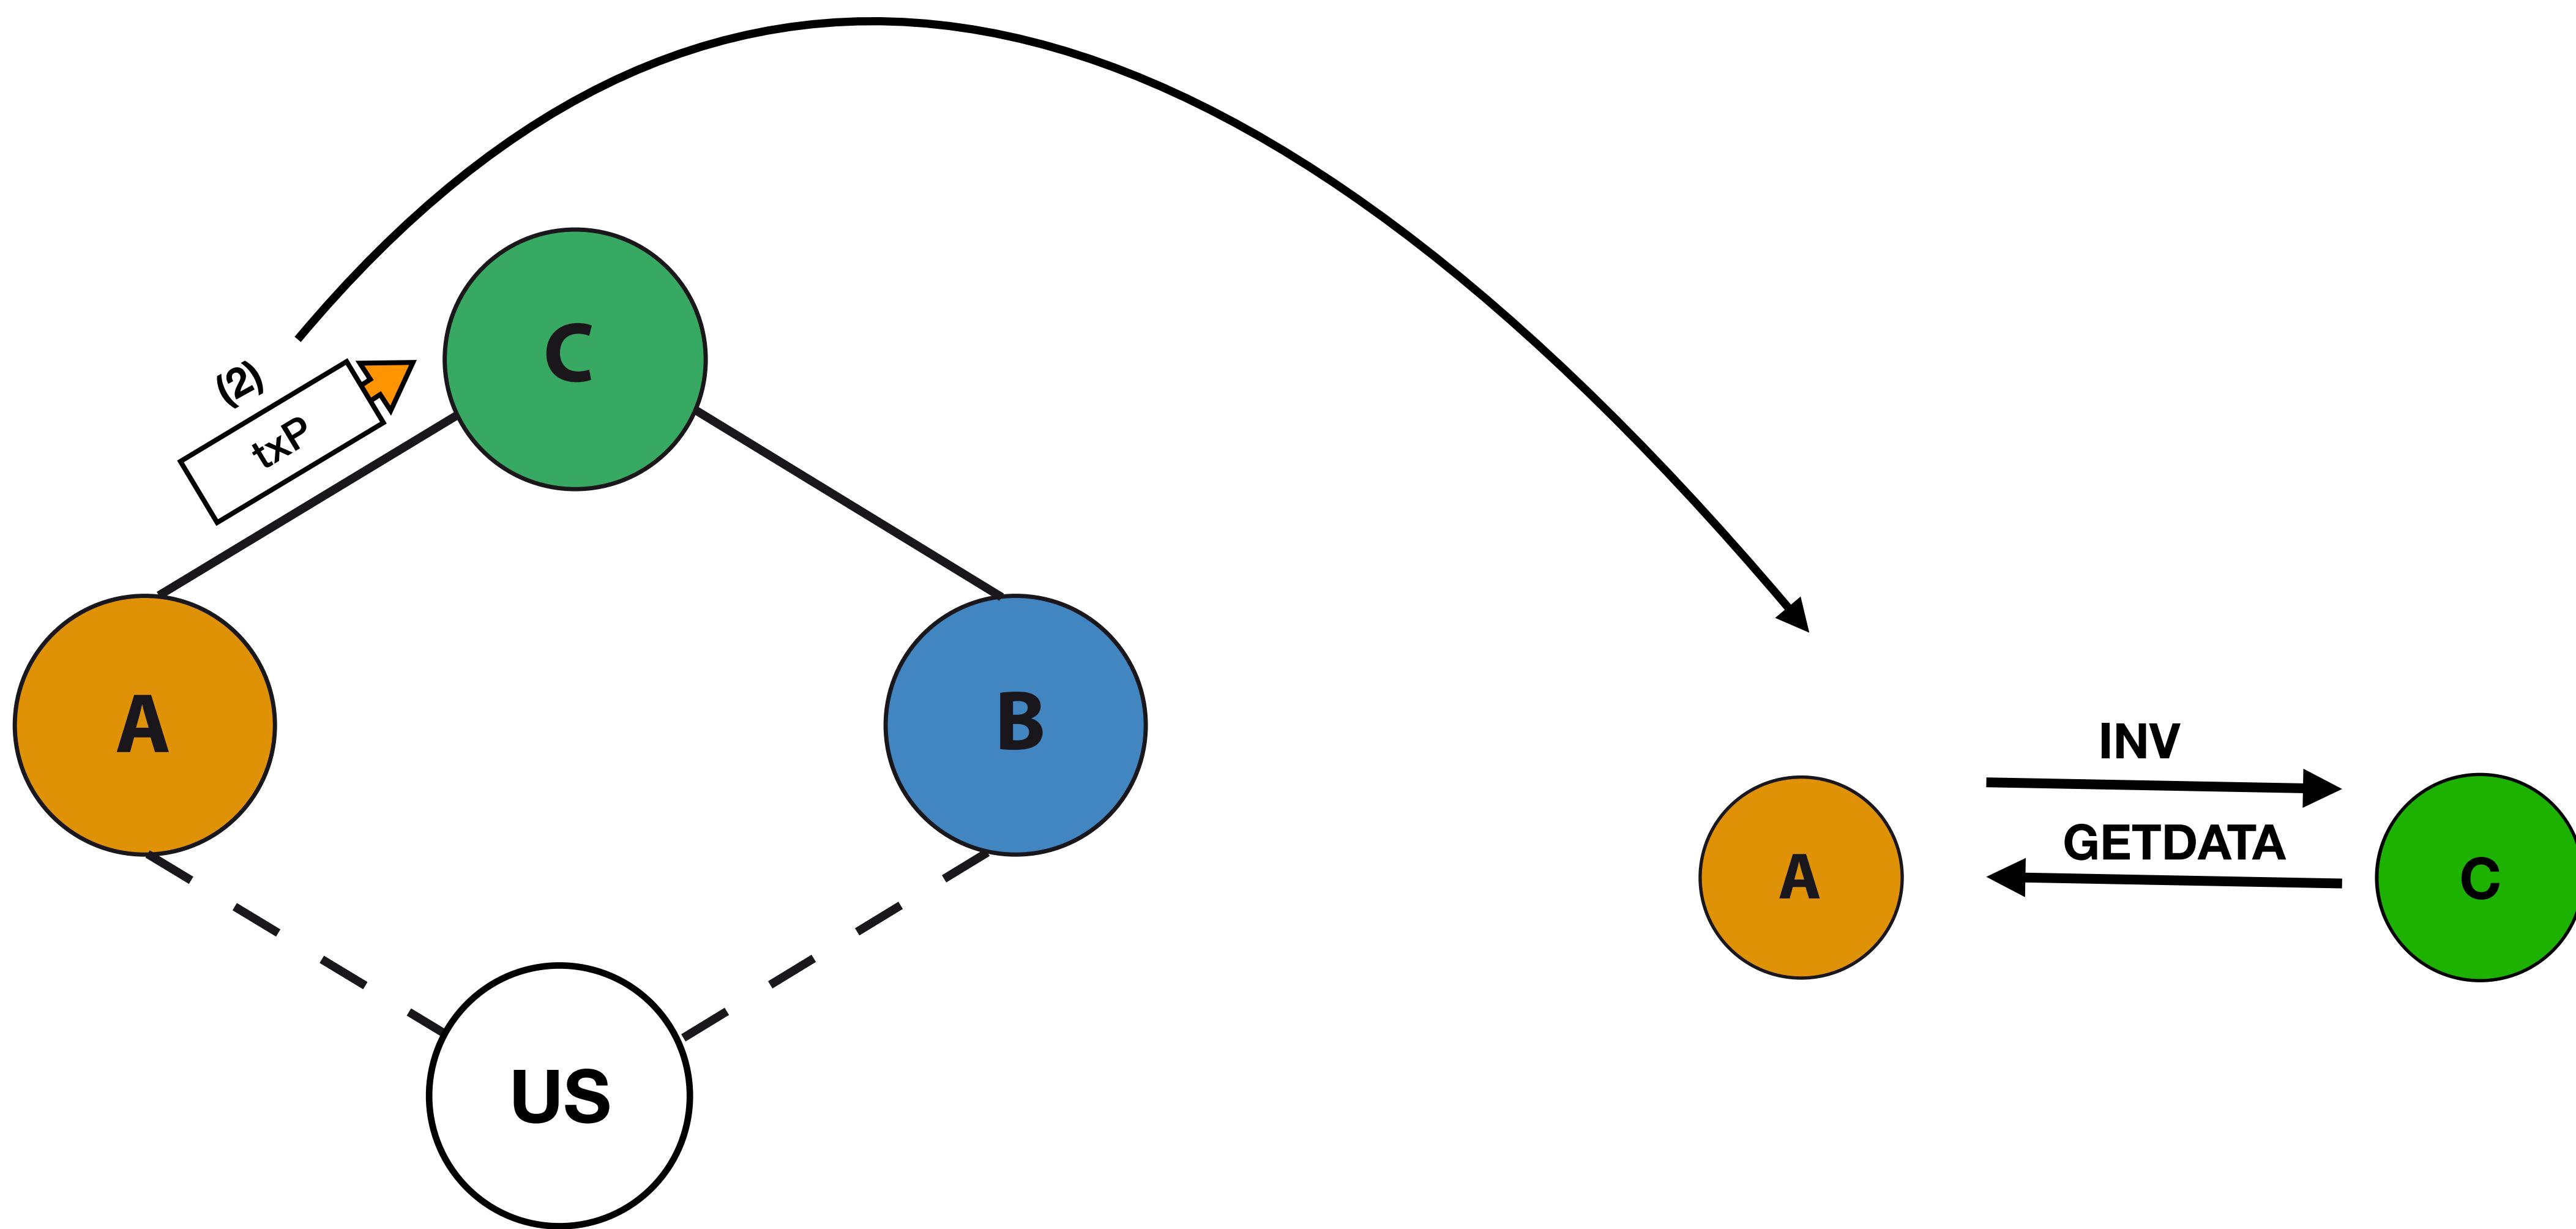
ISOLATION



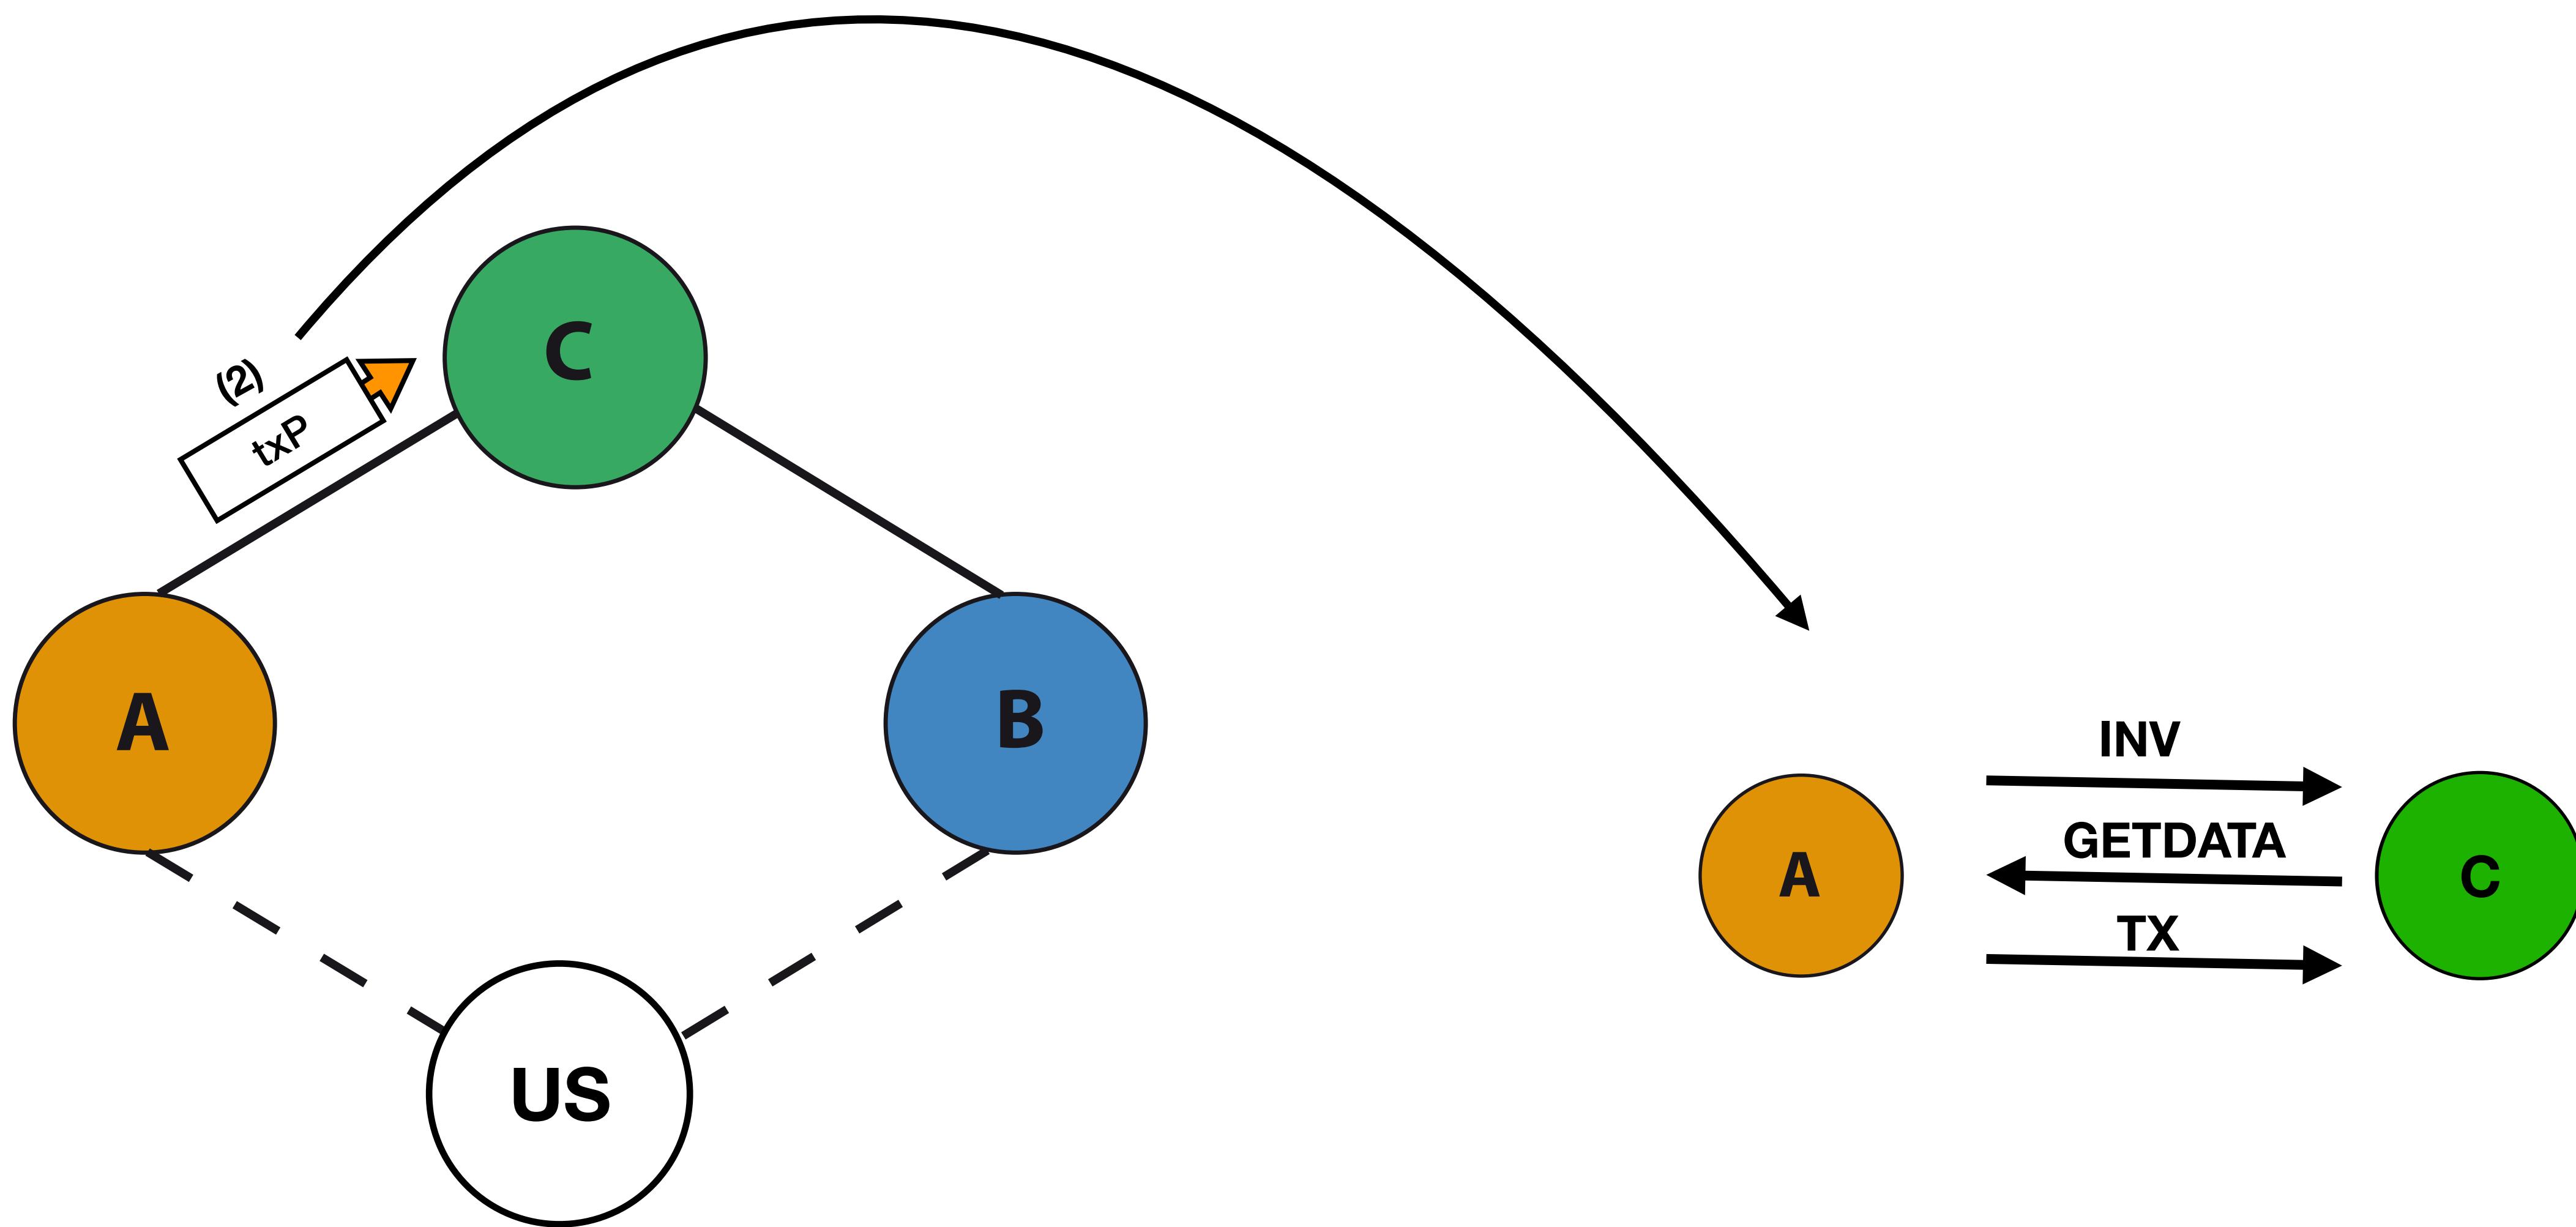
ISOLATION



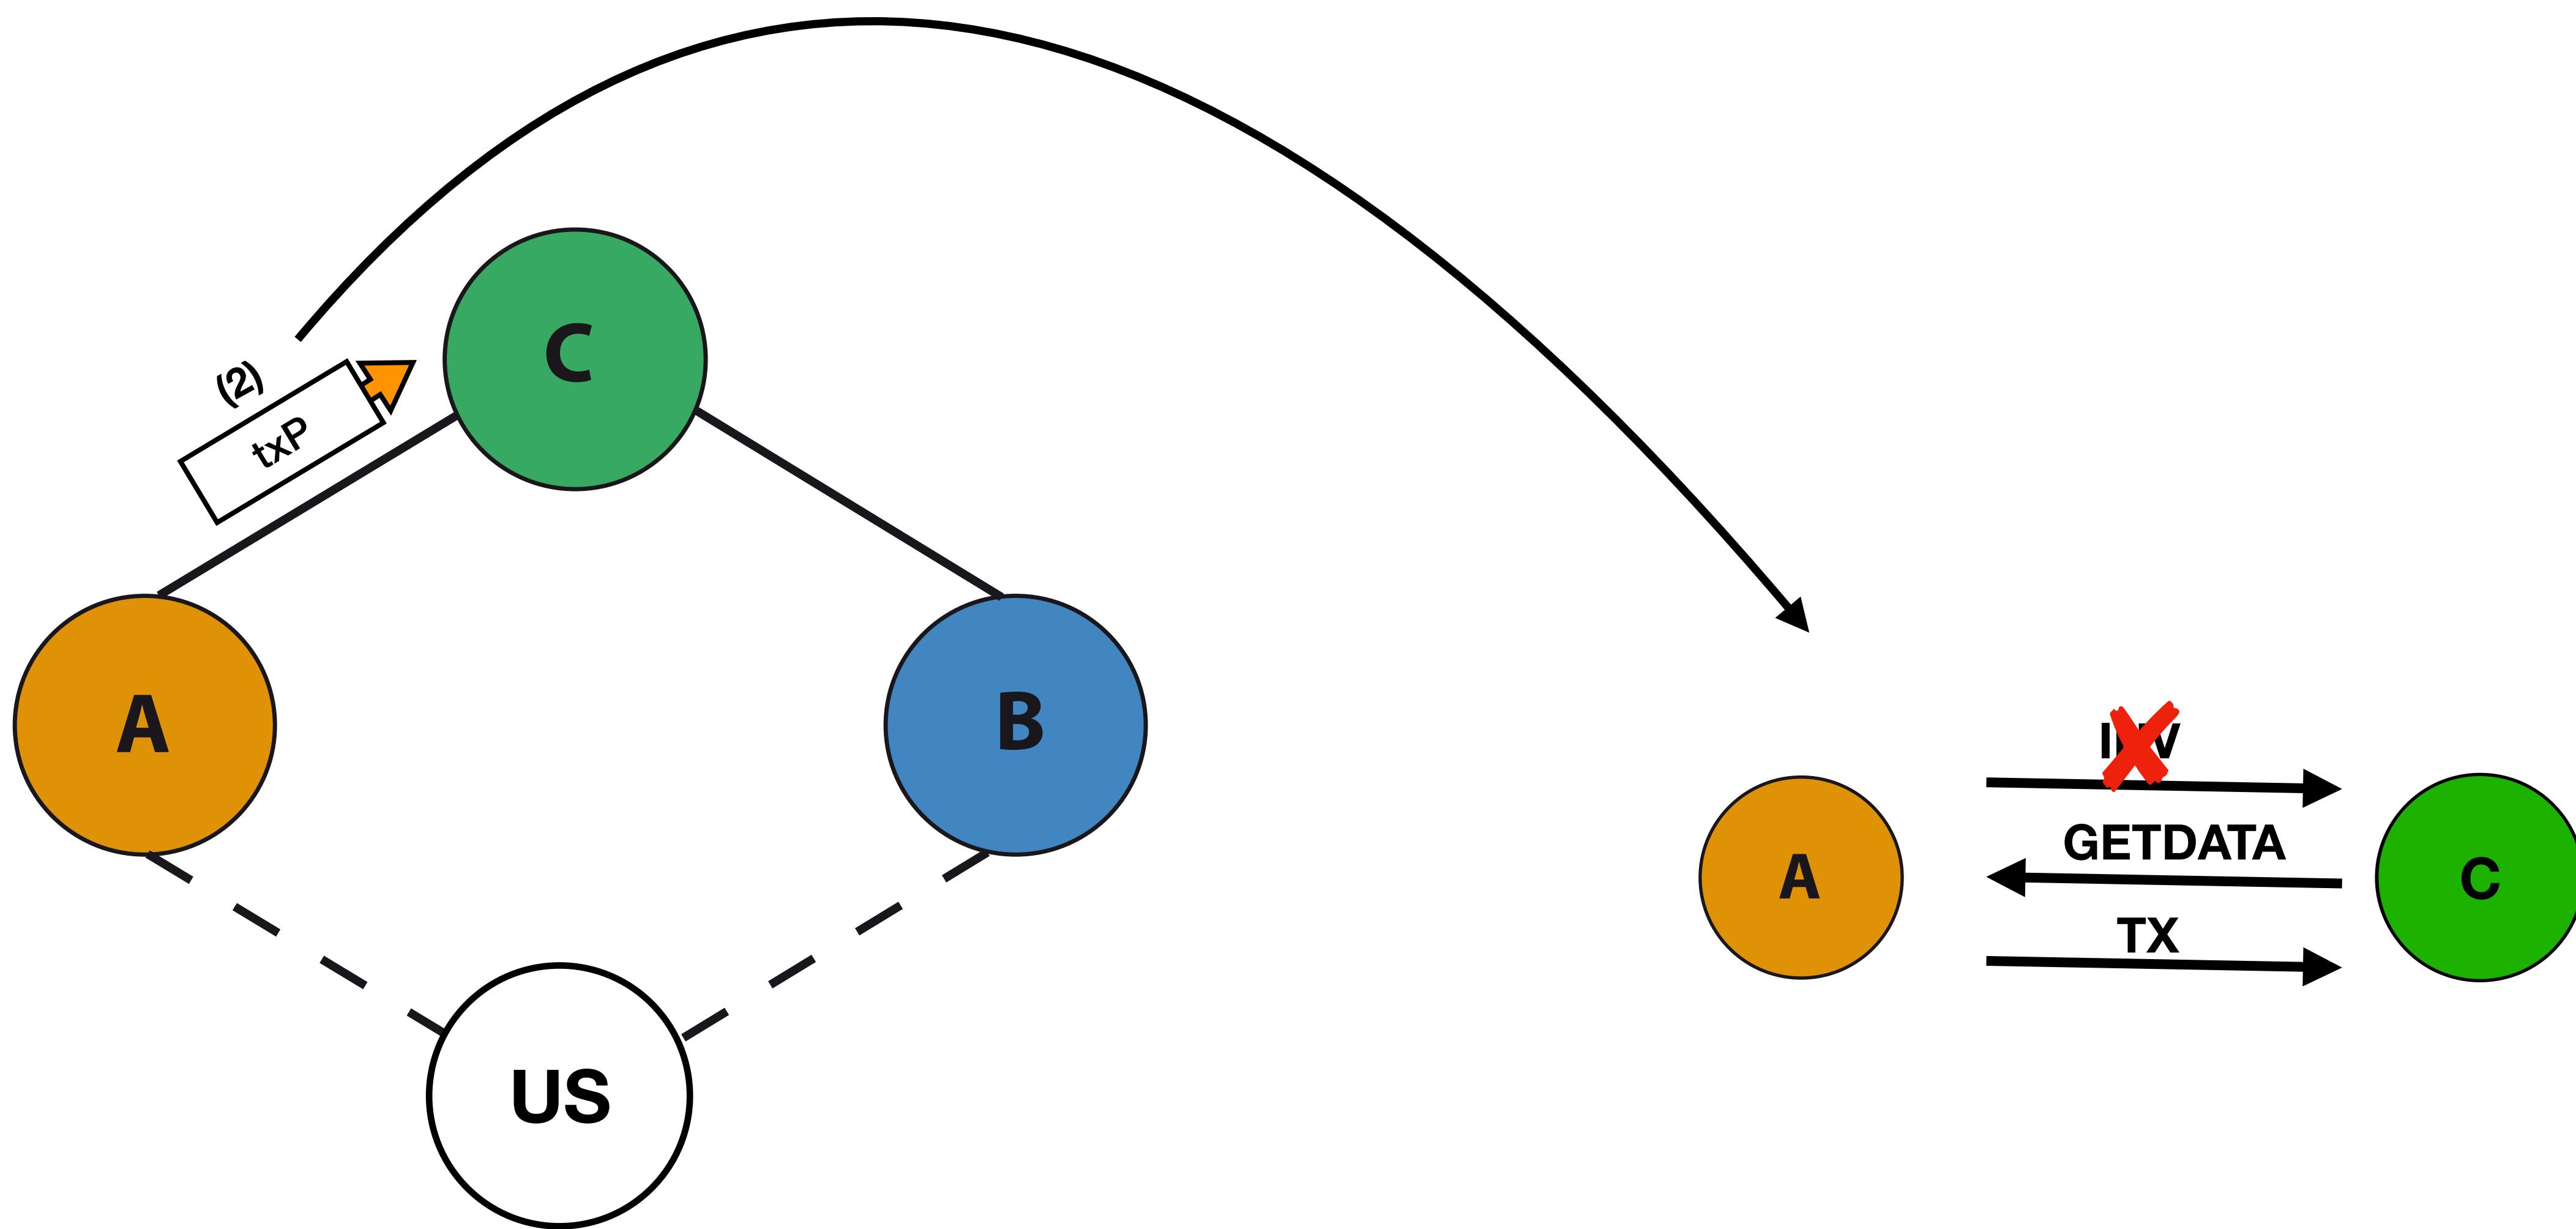
ISOLATION



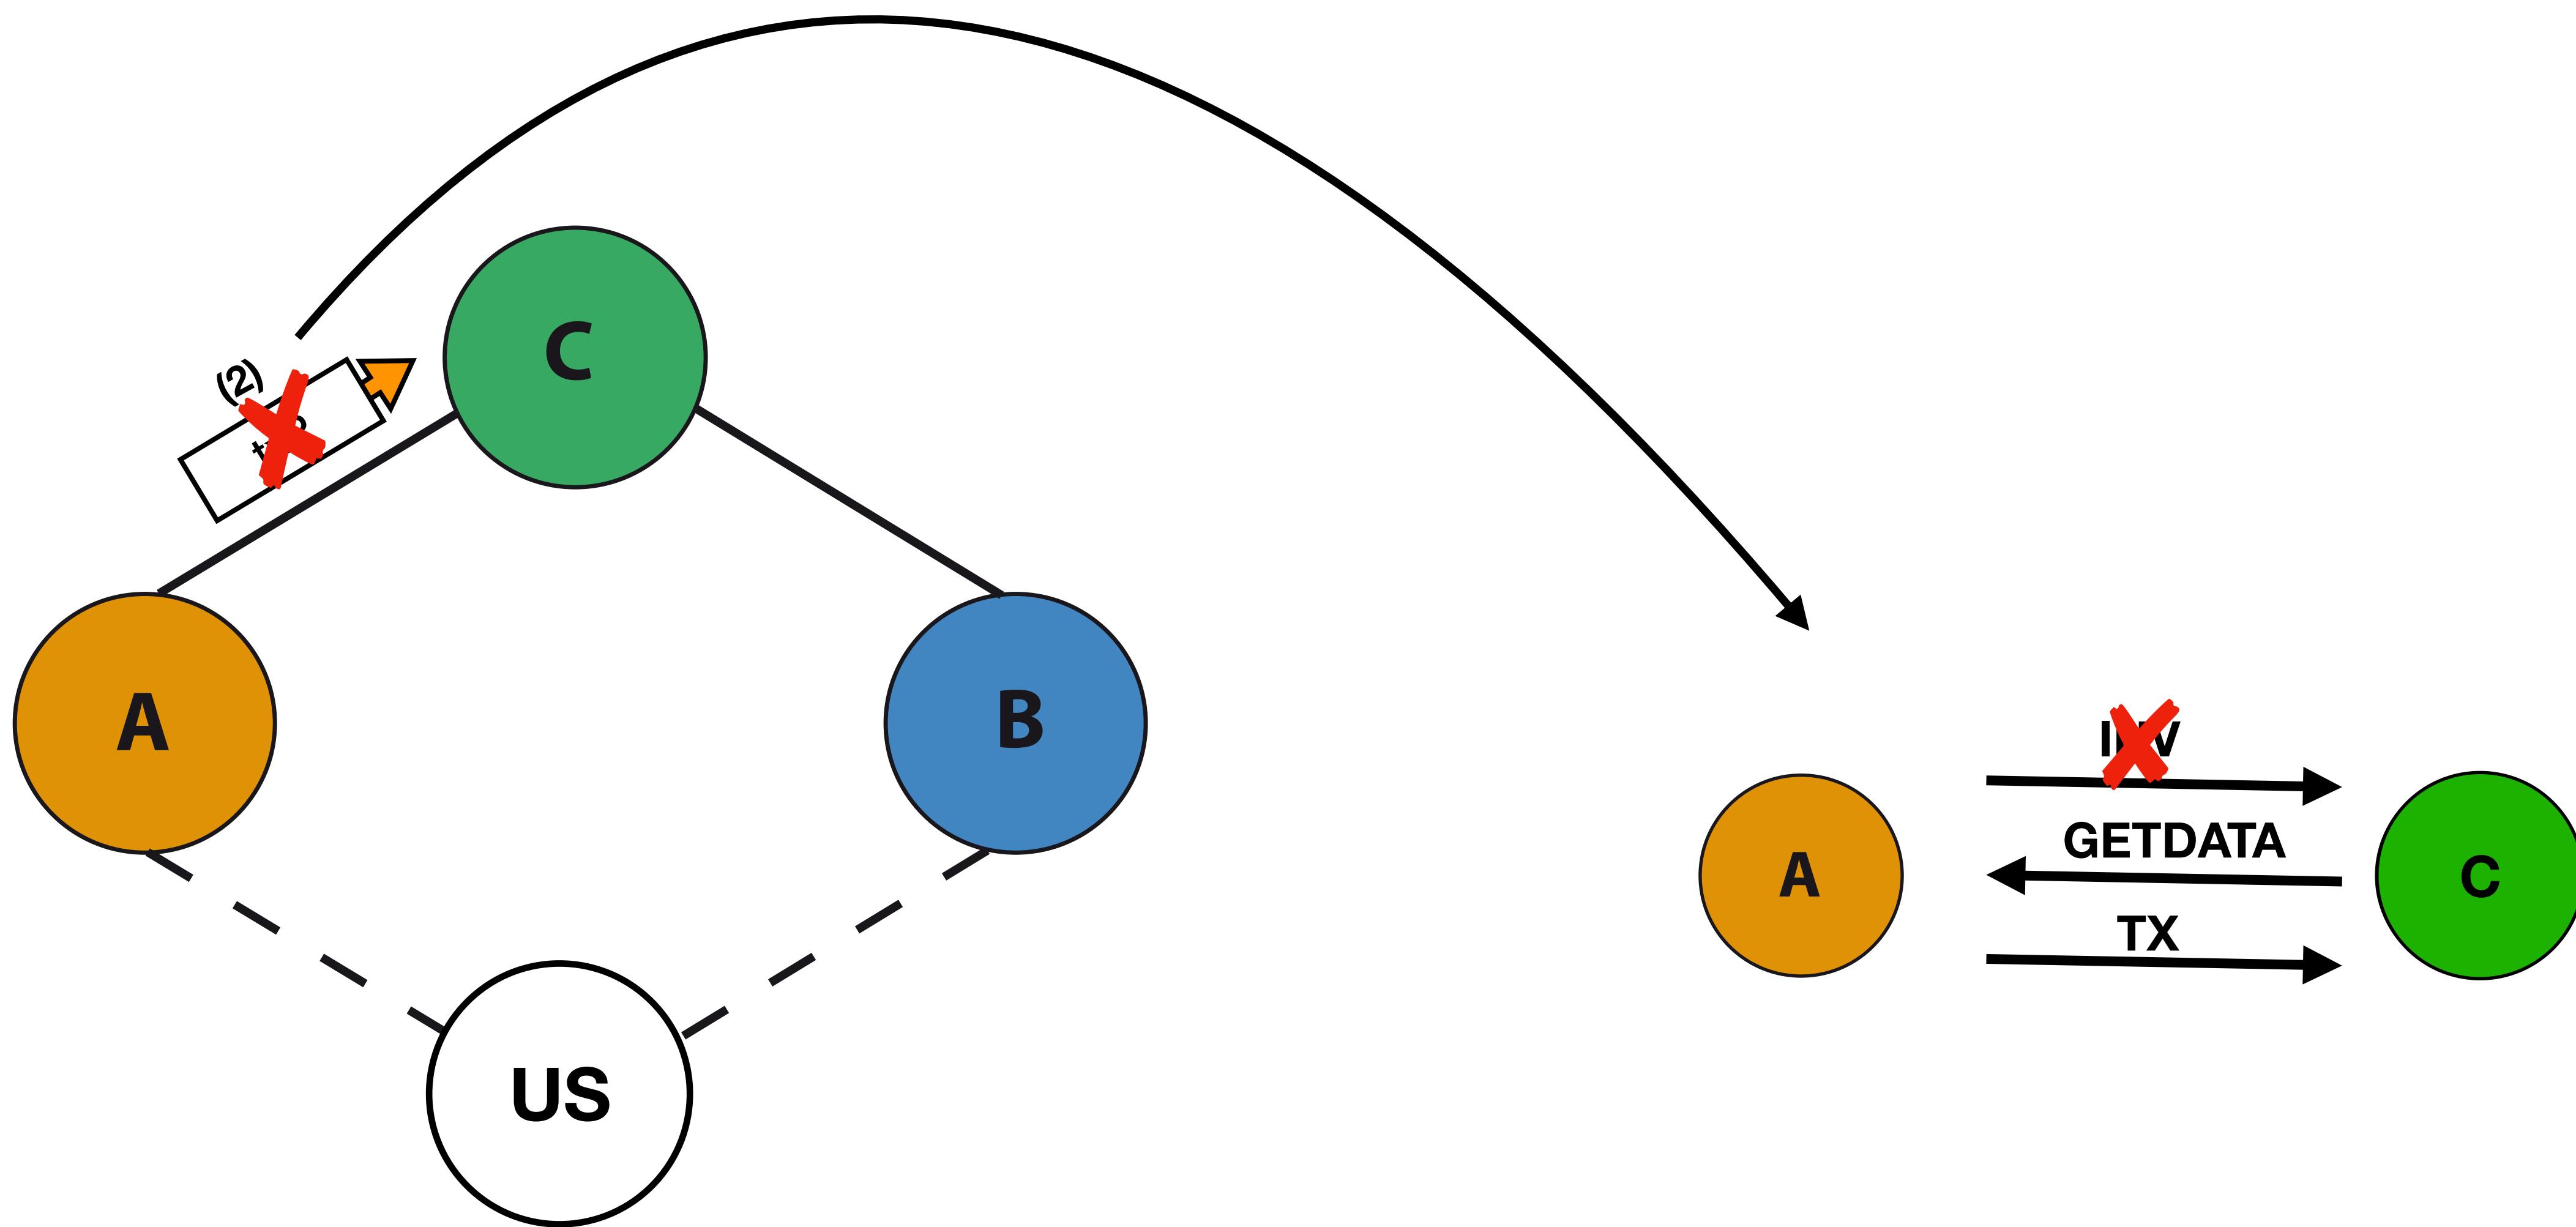
ISOLATION



ISOLATION

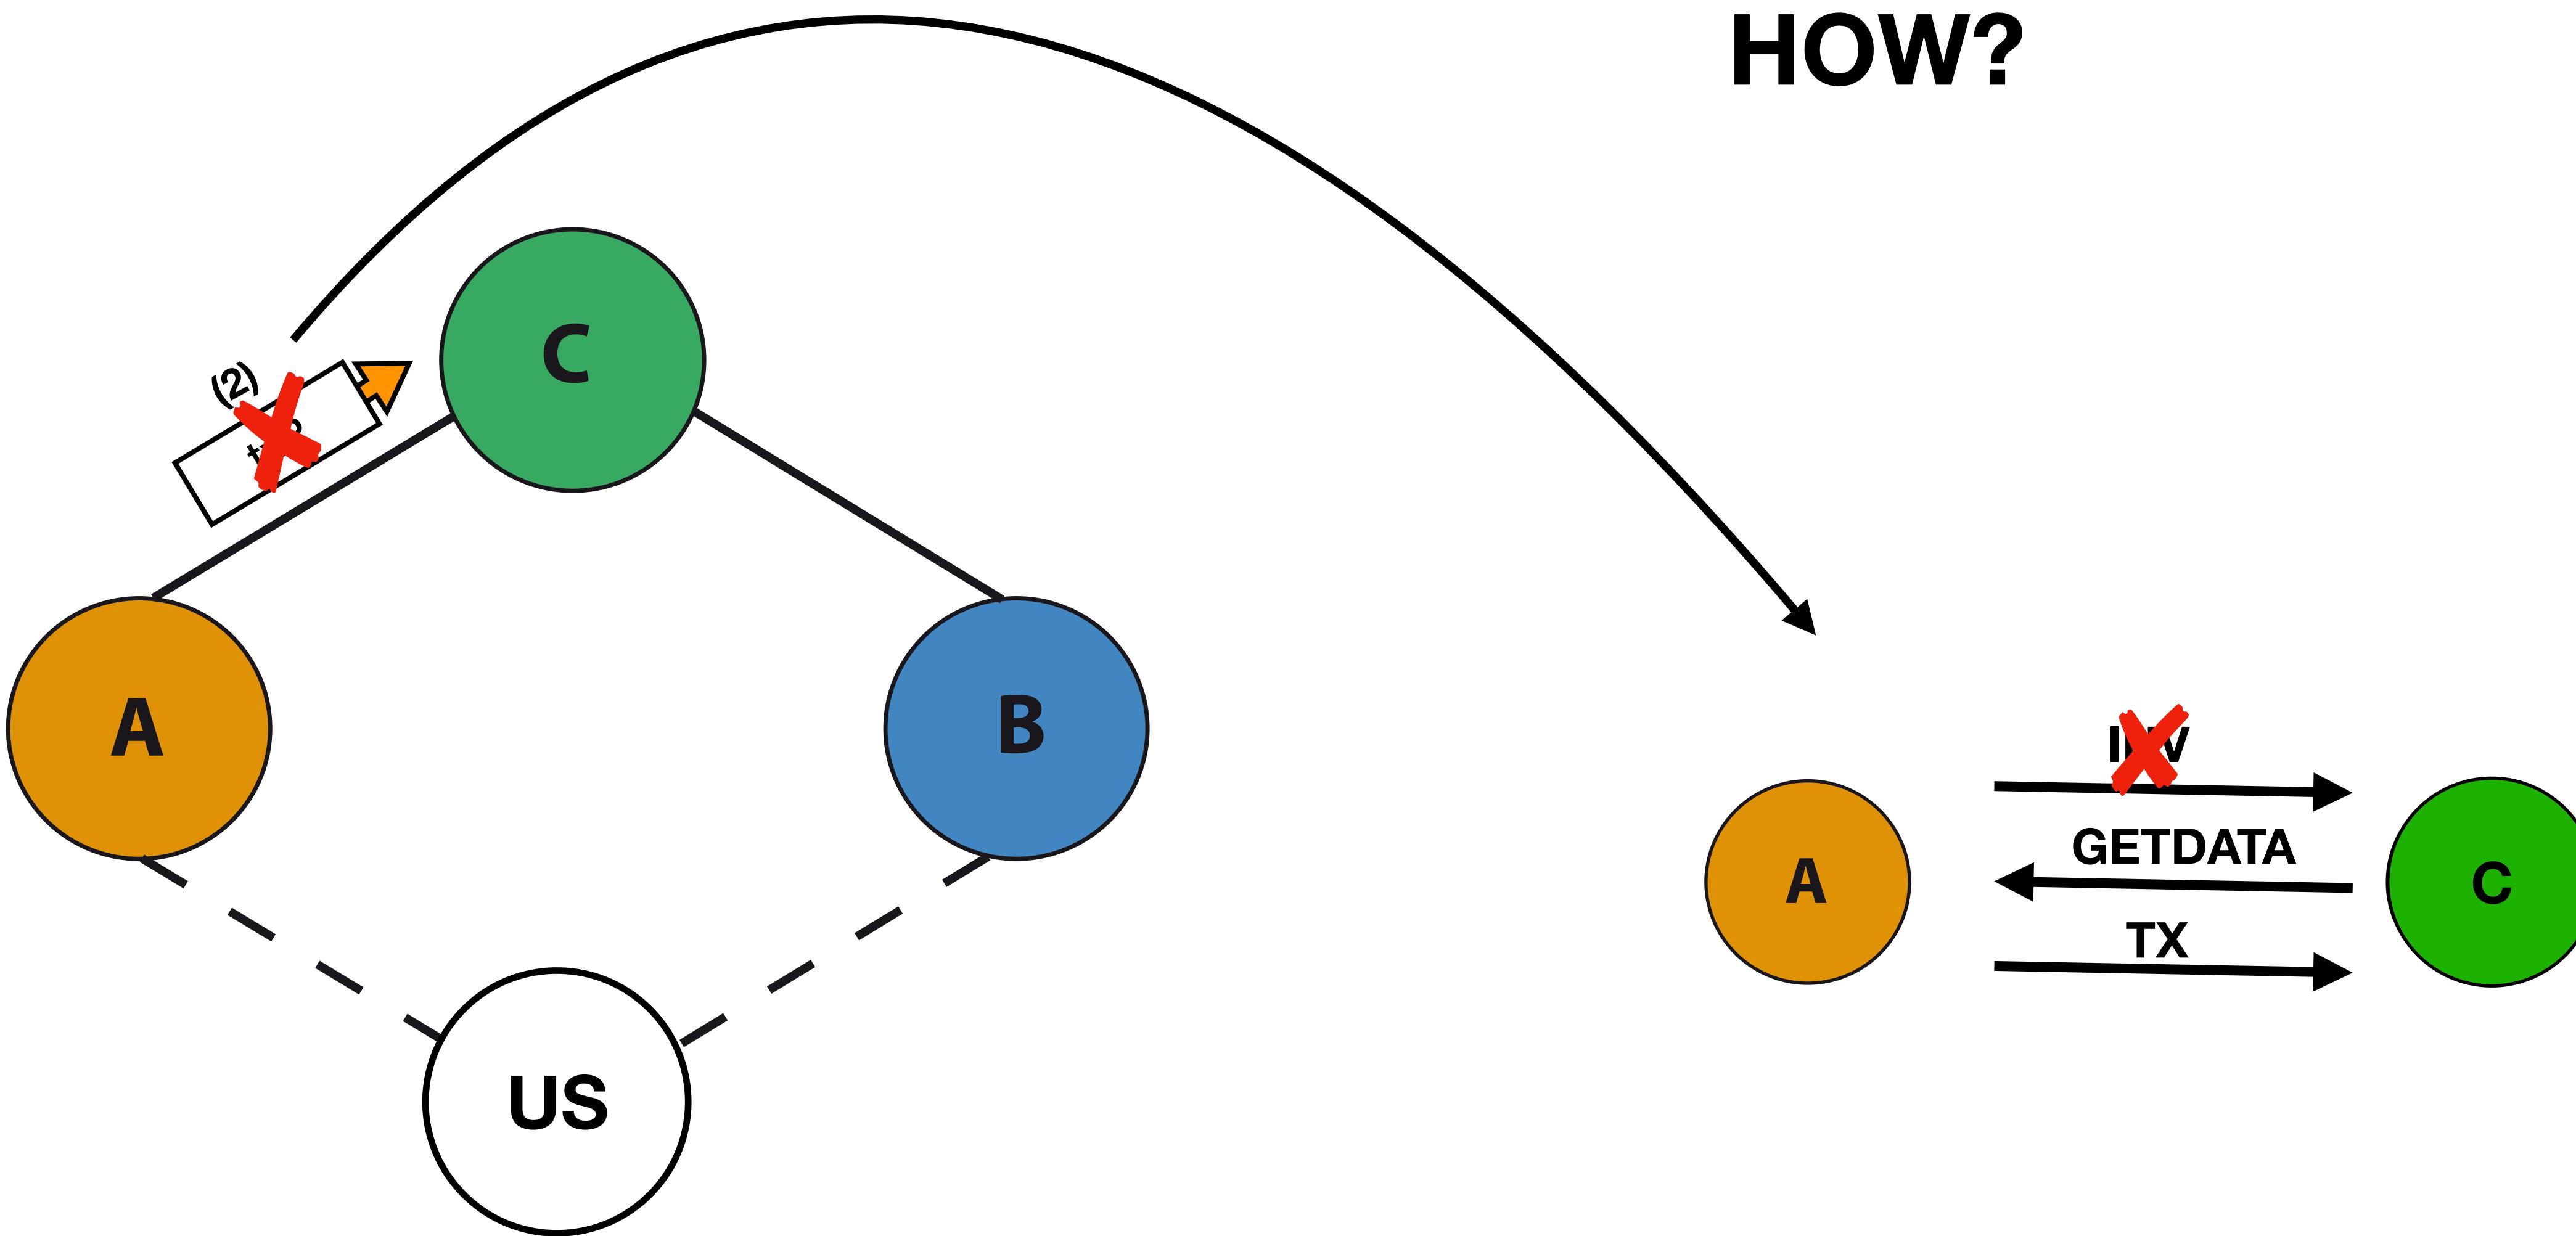


ISOLATION

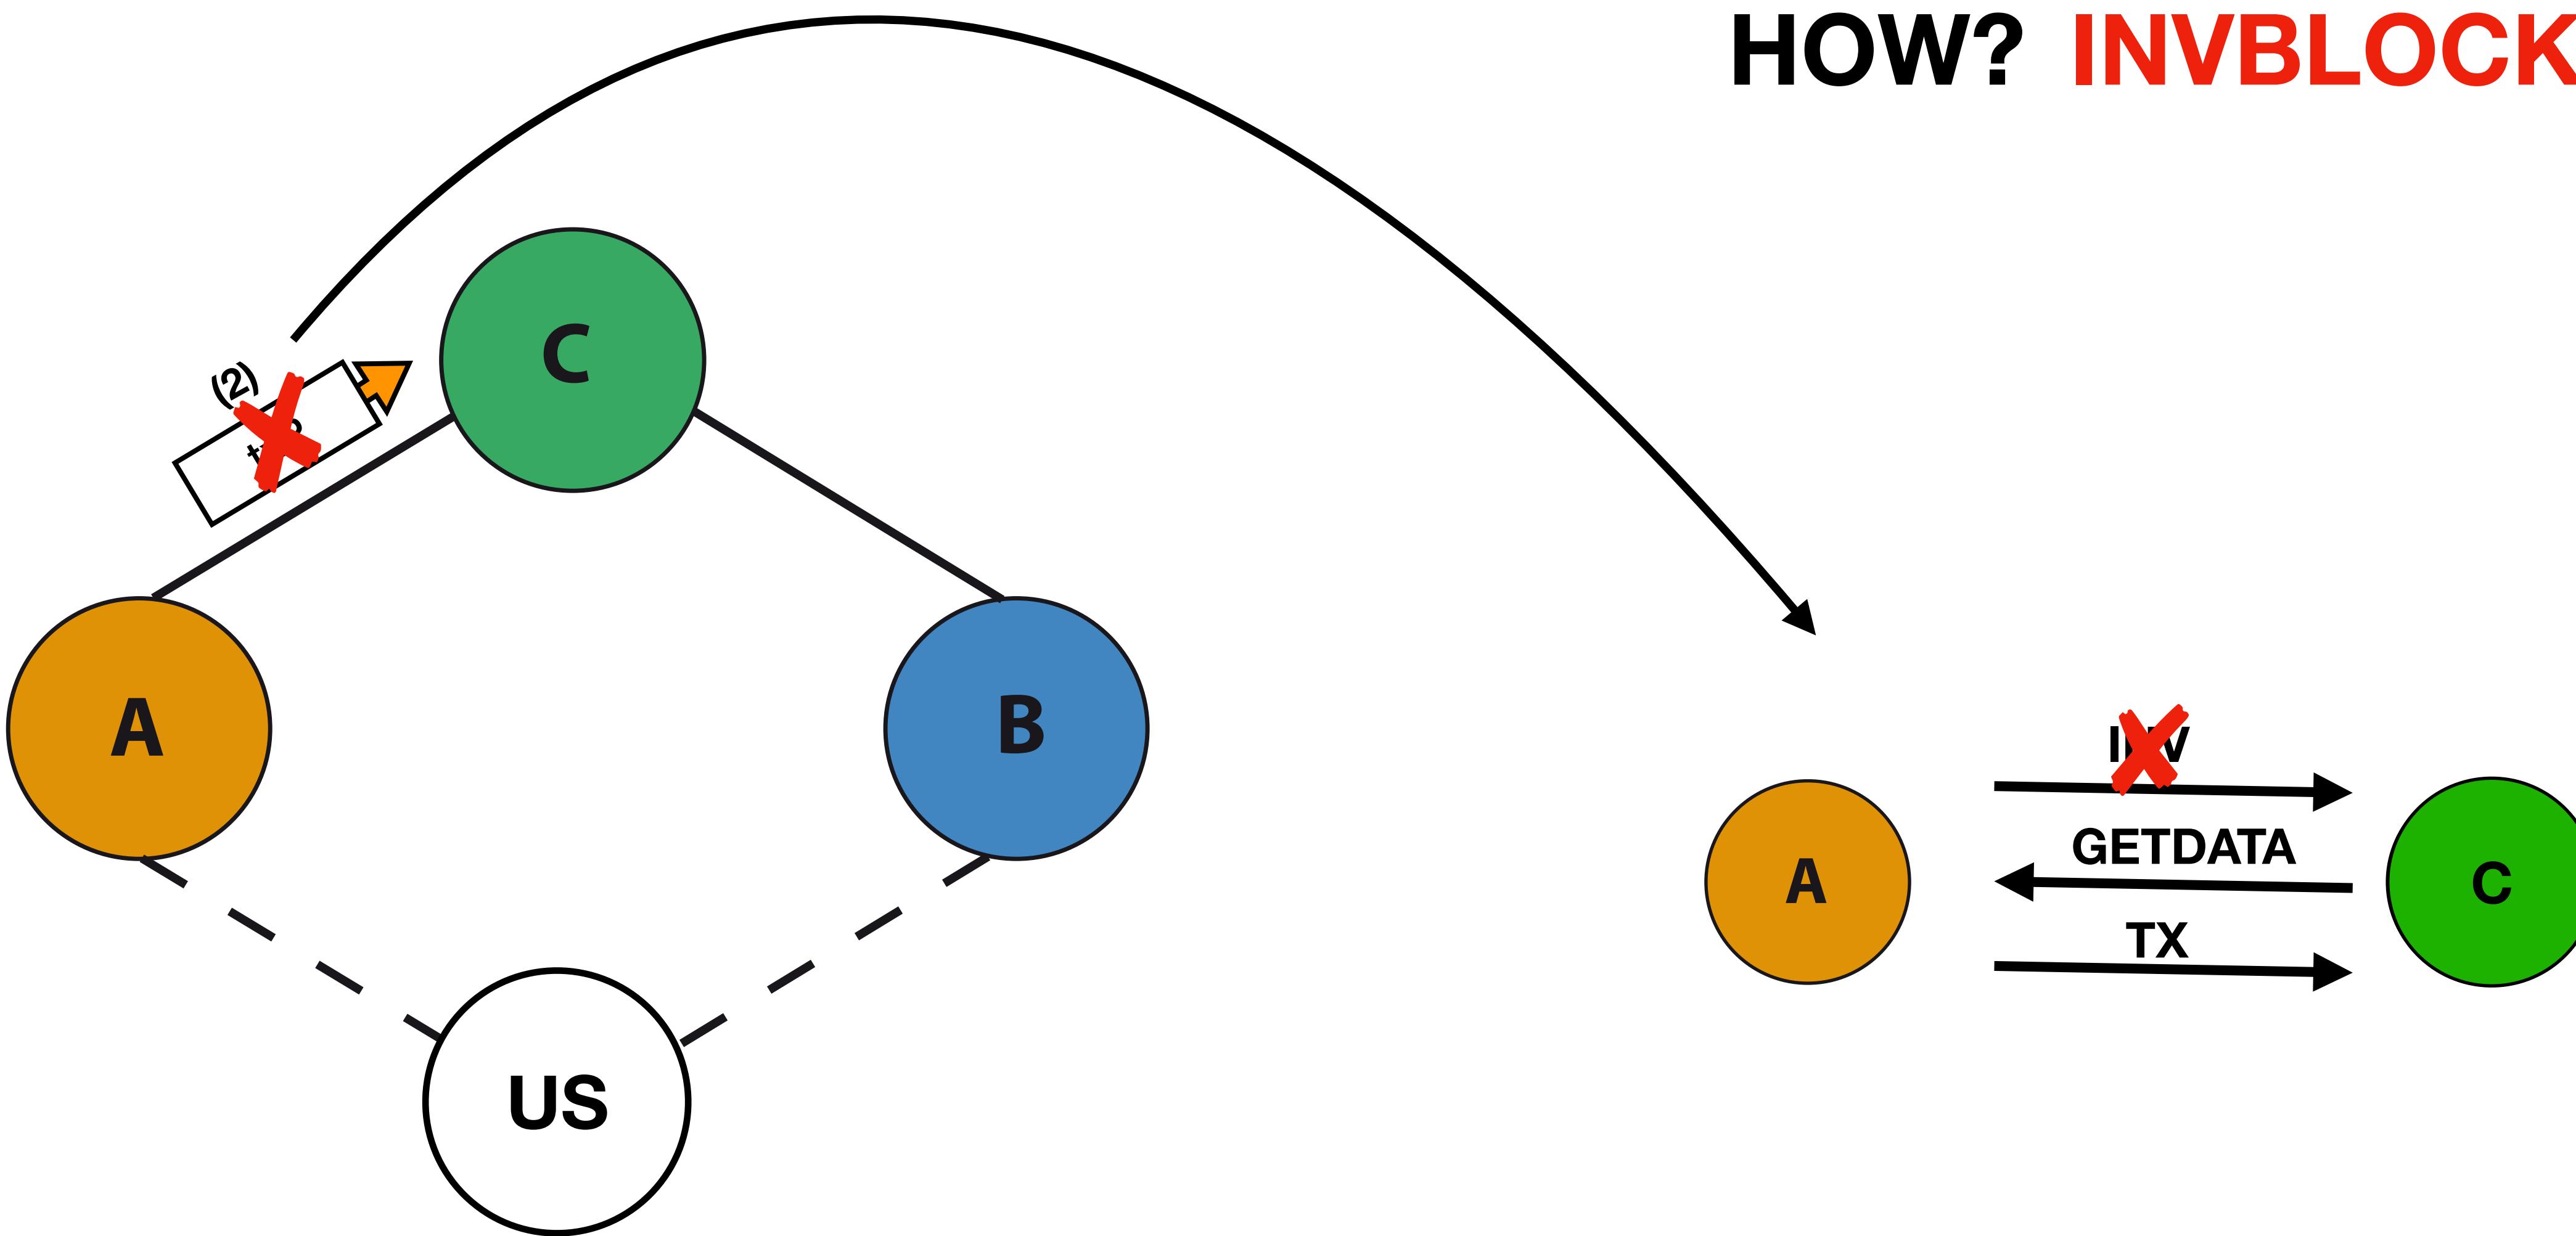


ISOLATION

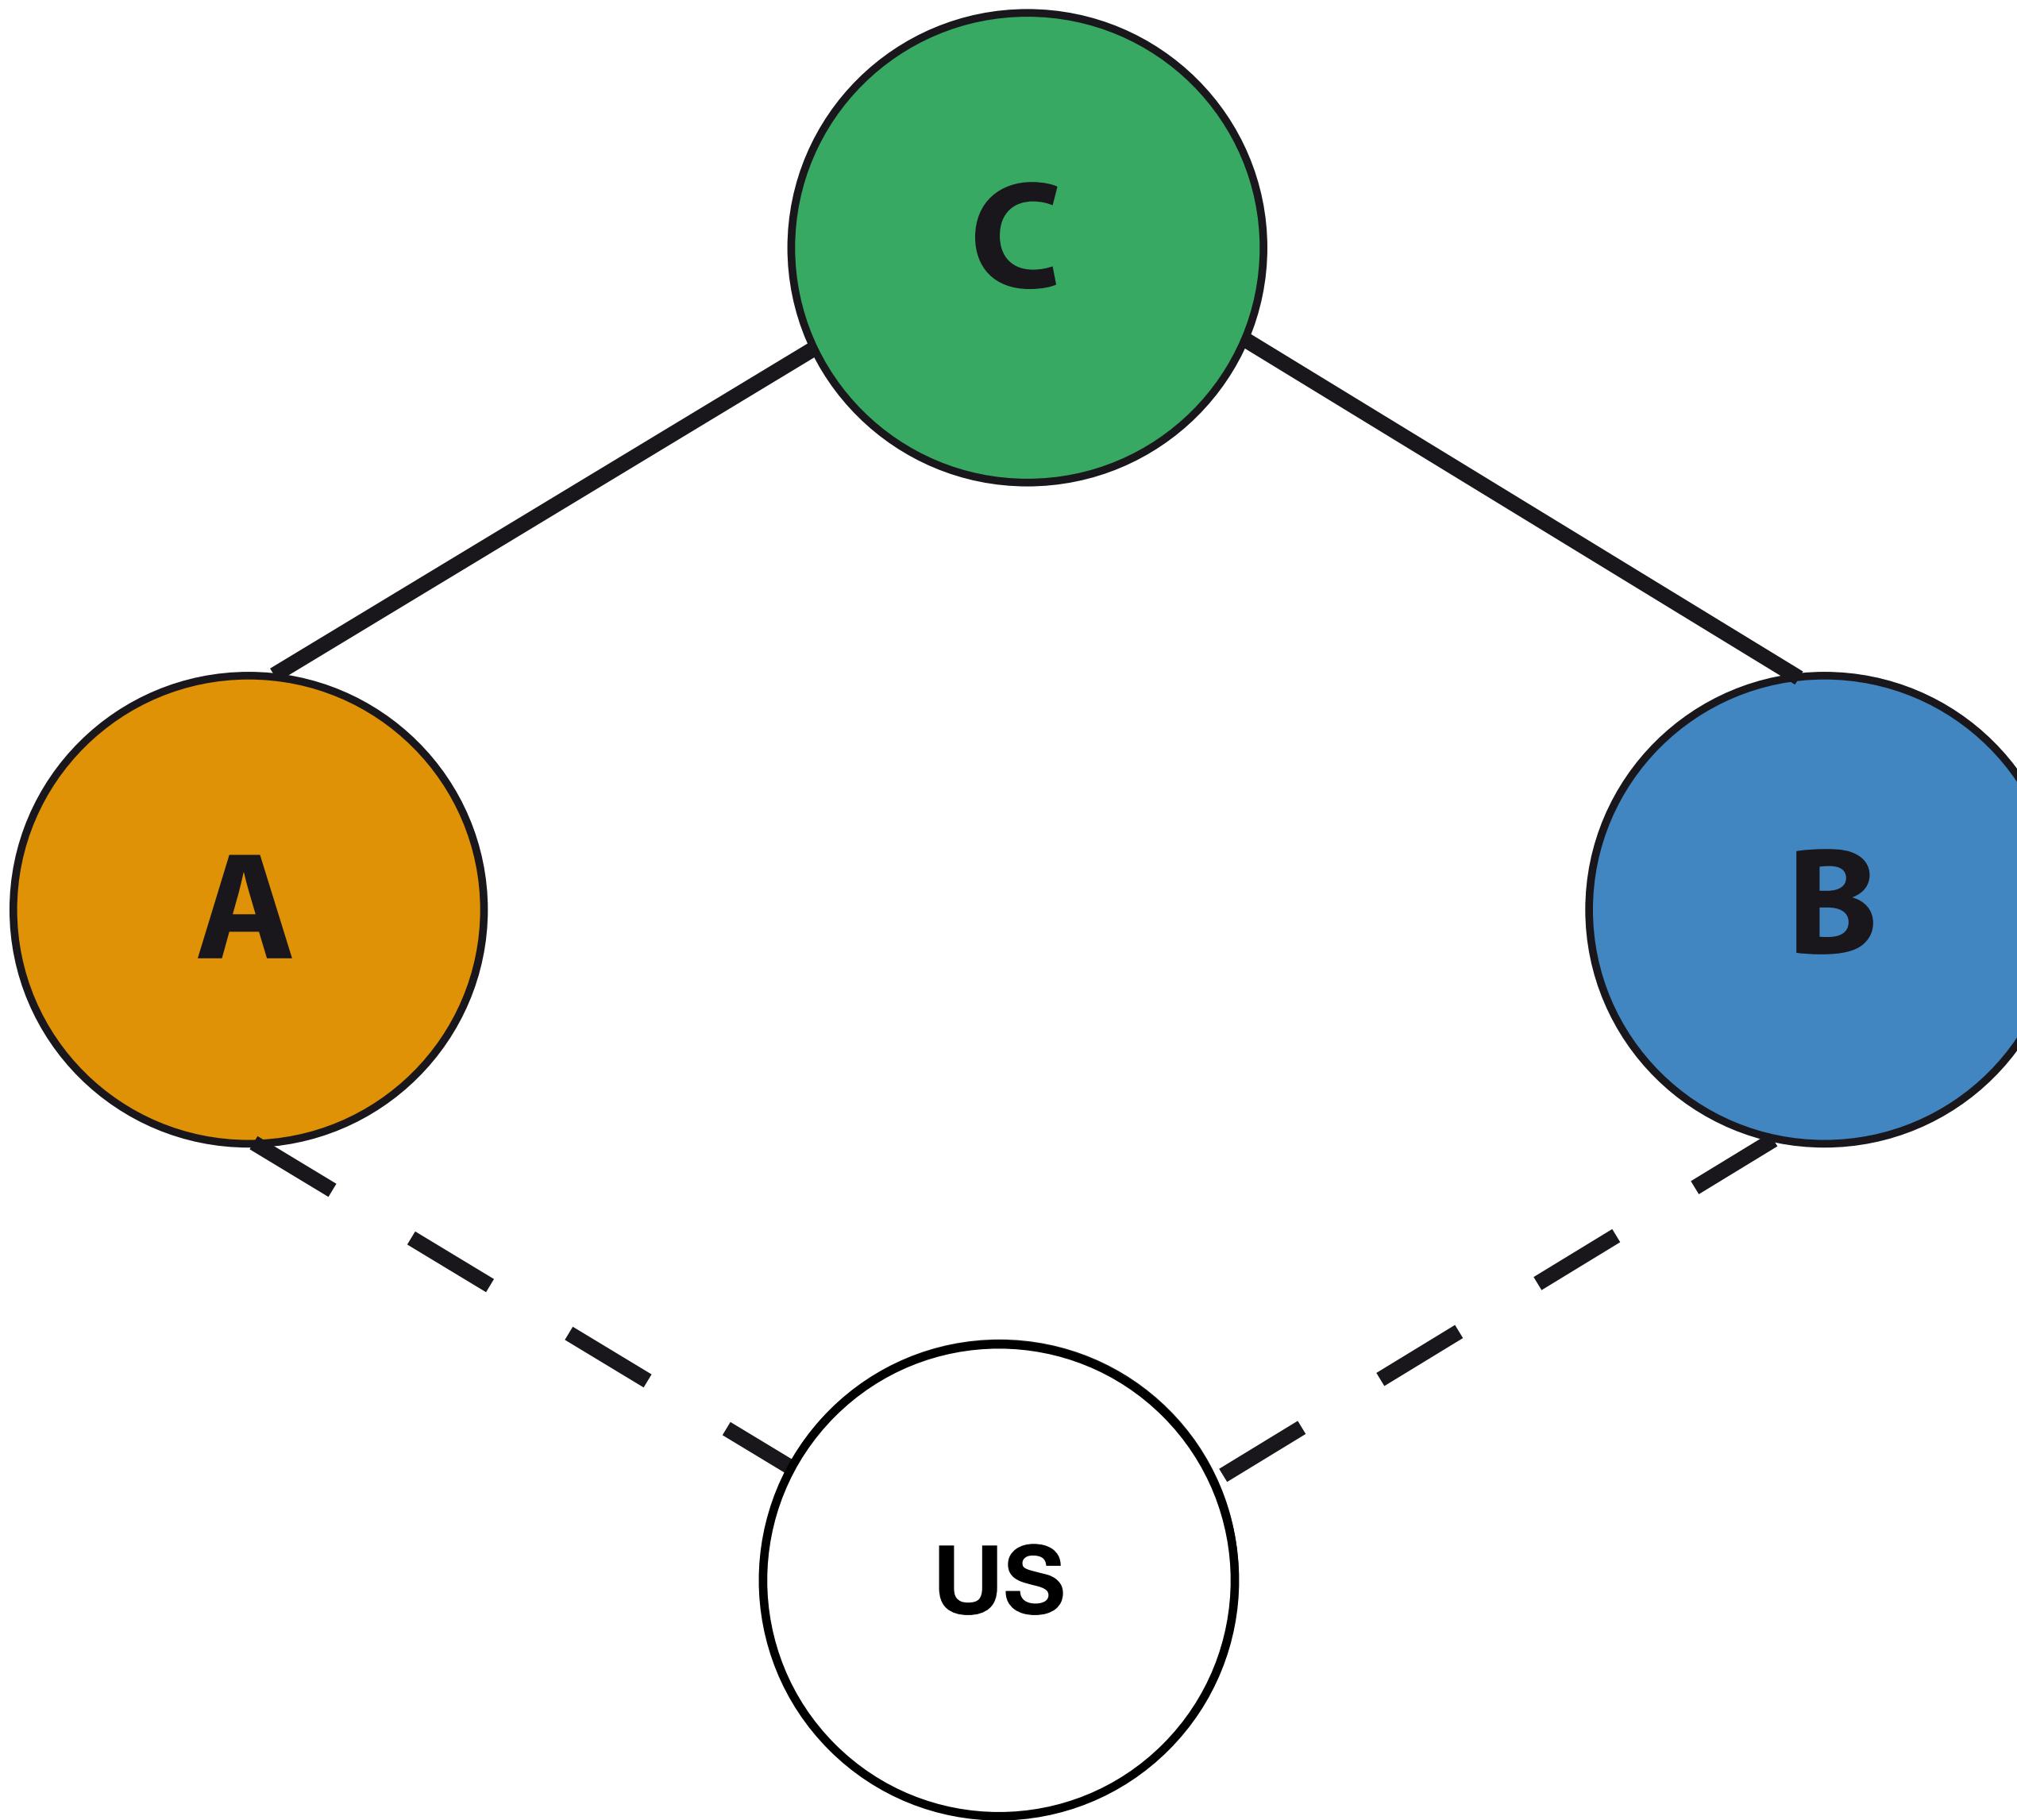
HOW?



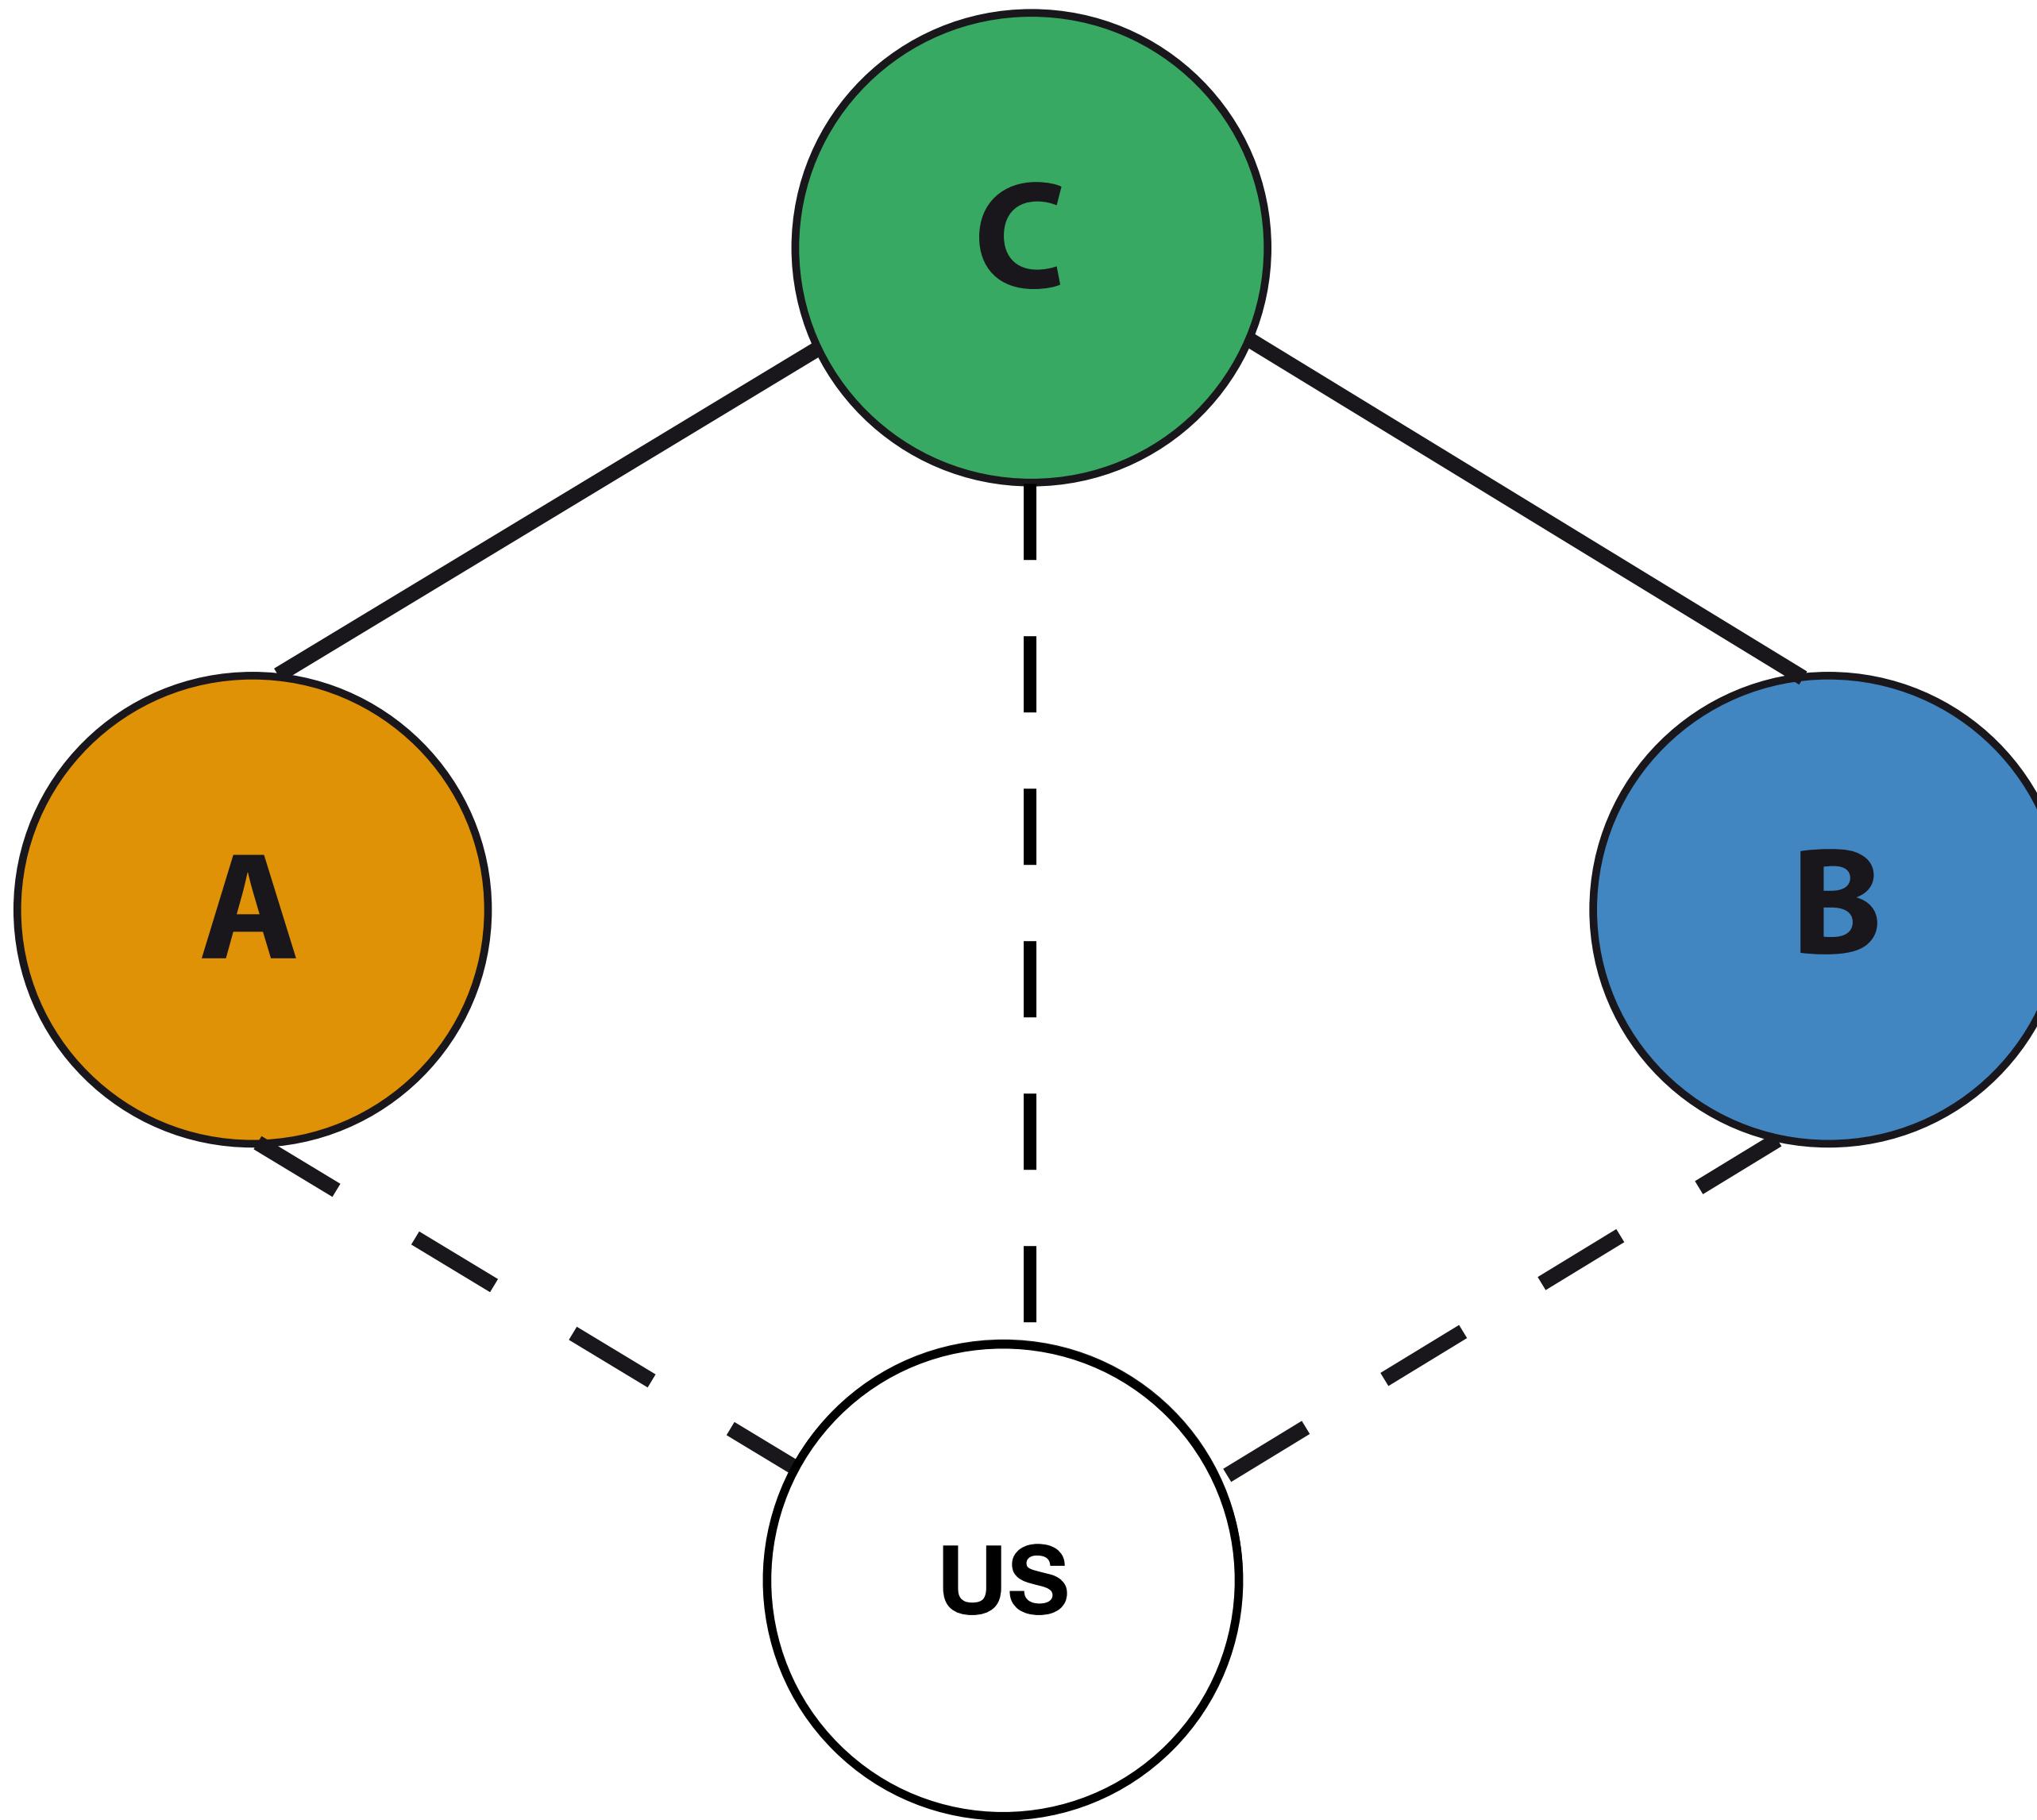
ISOLATION



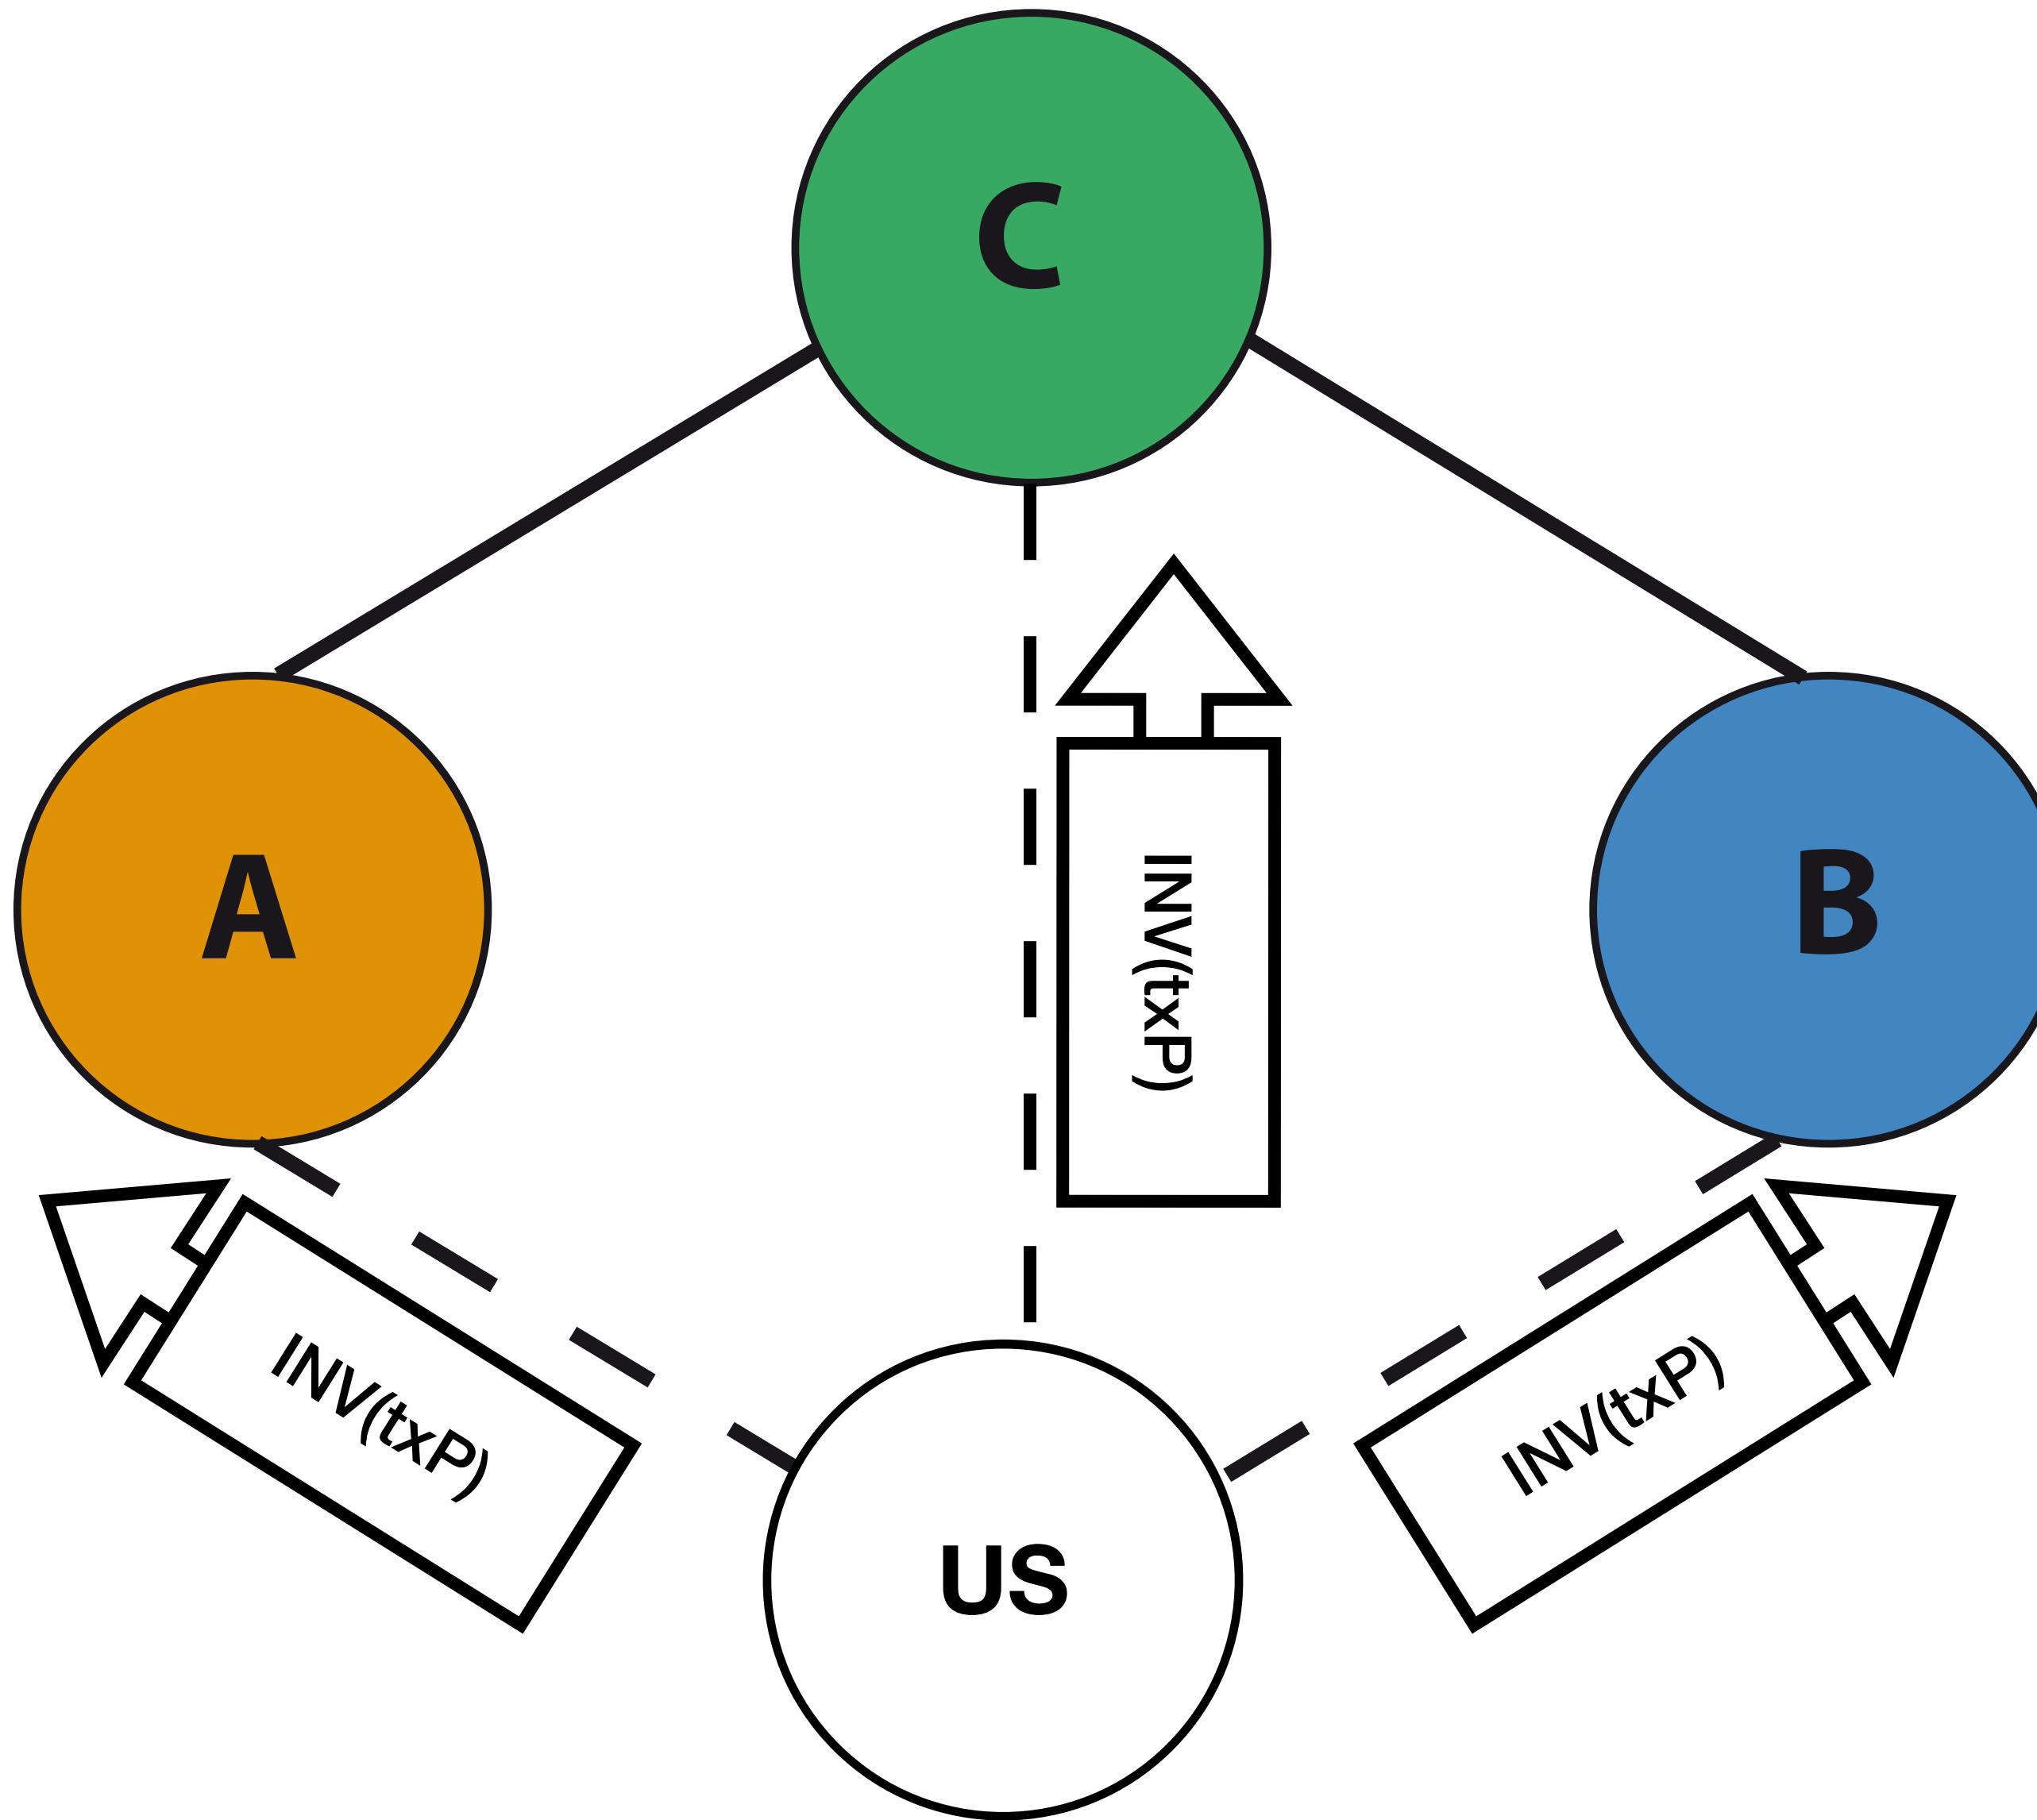
INVBLOCKING



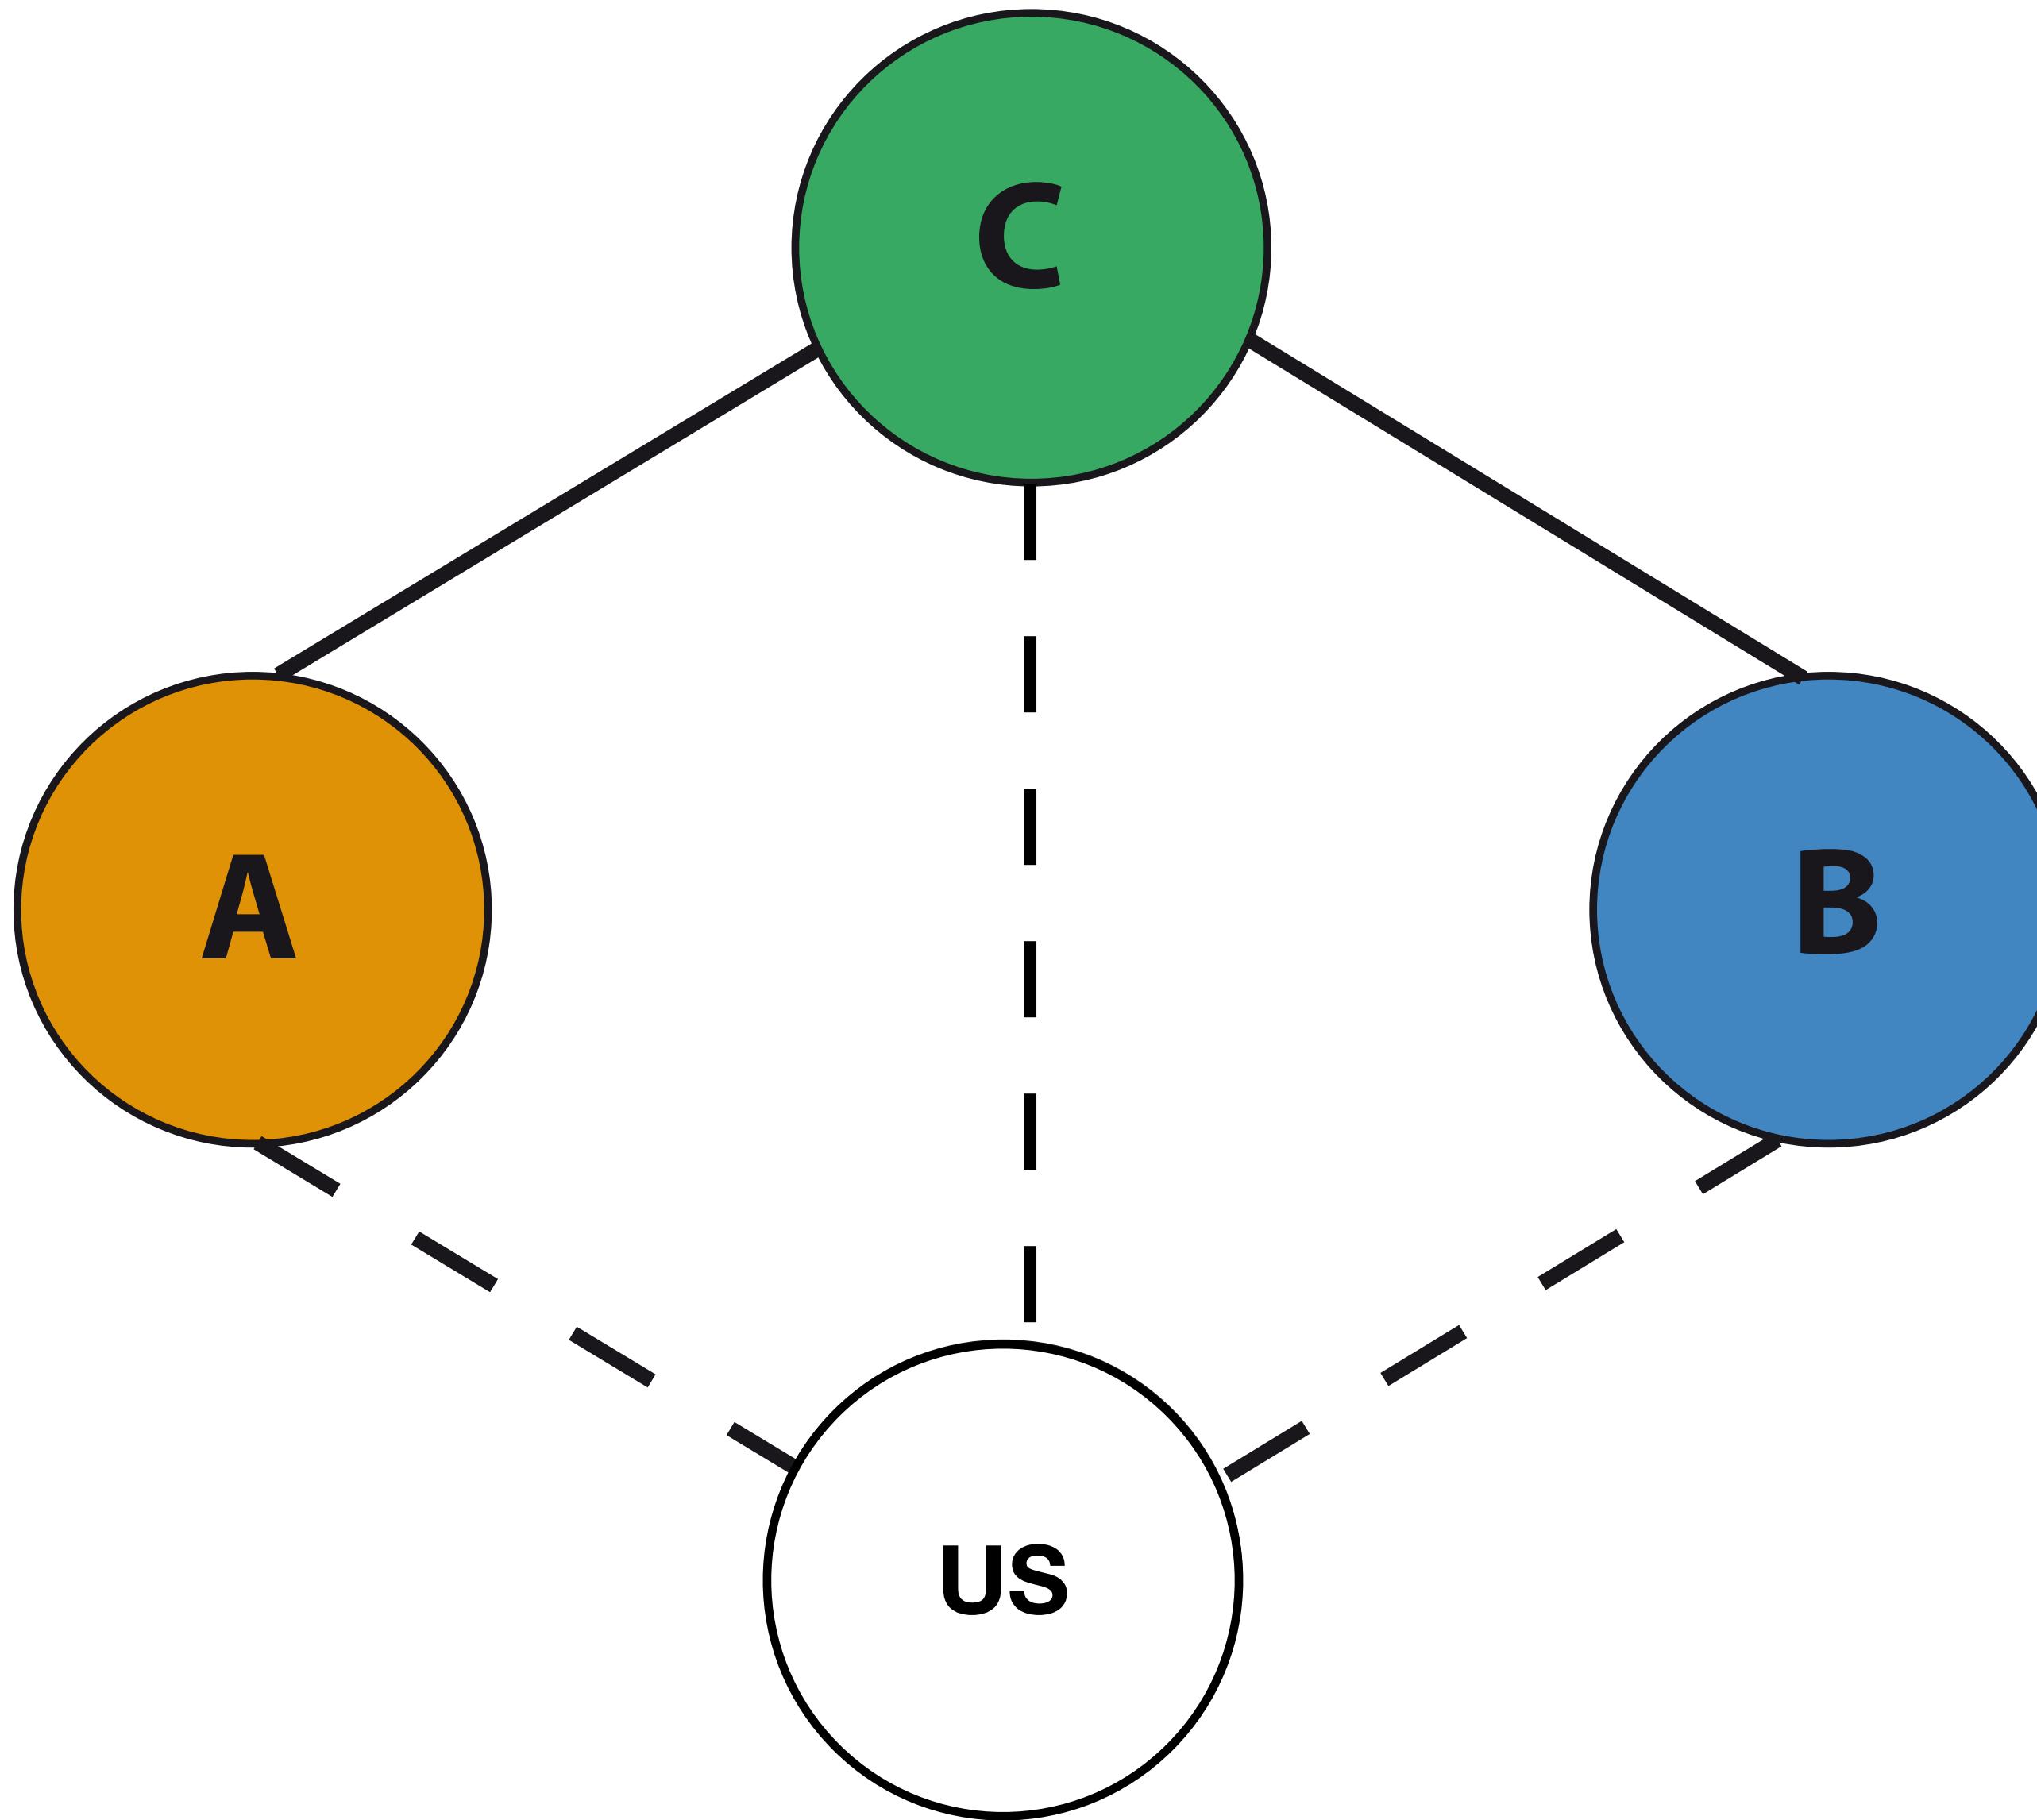
INVBLOCKING



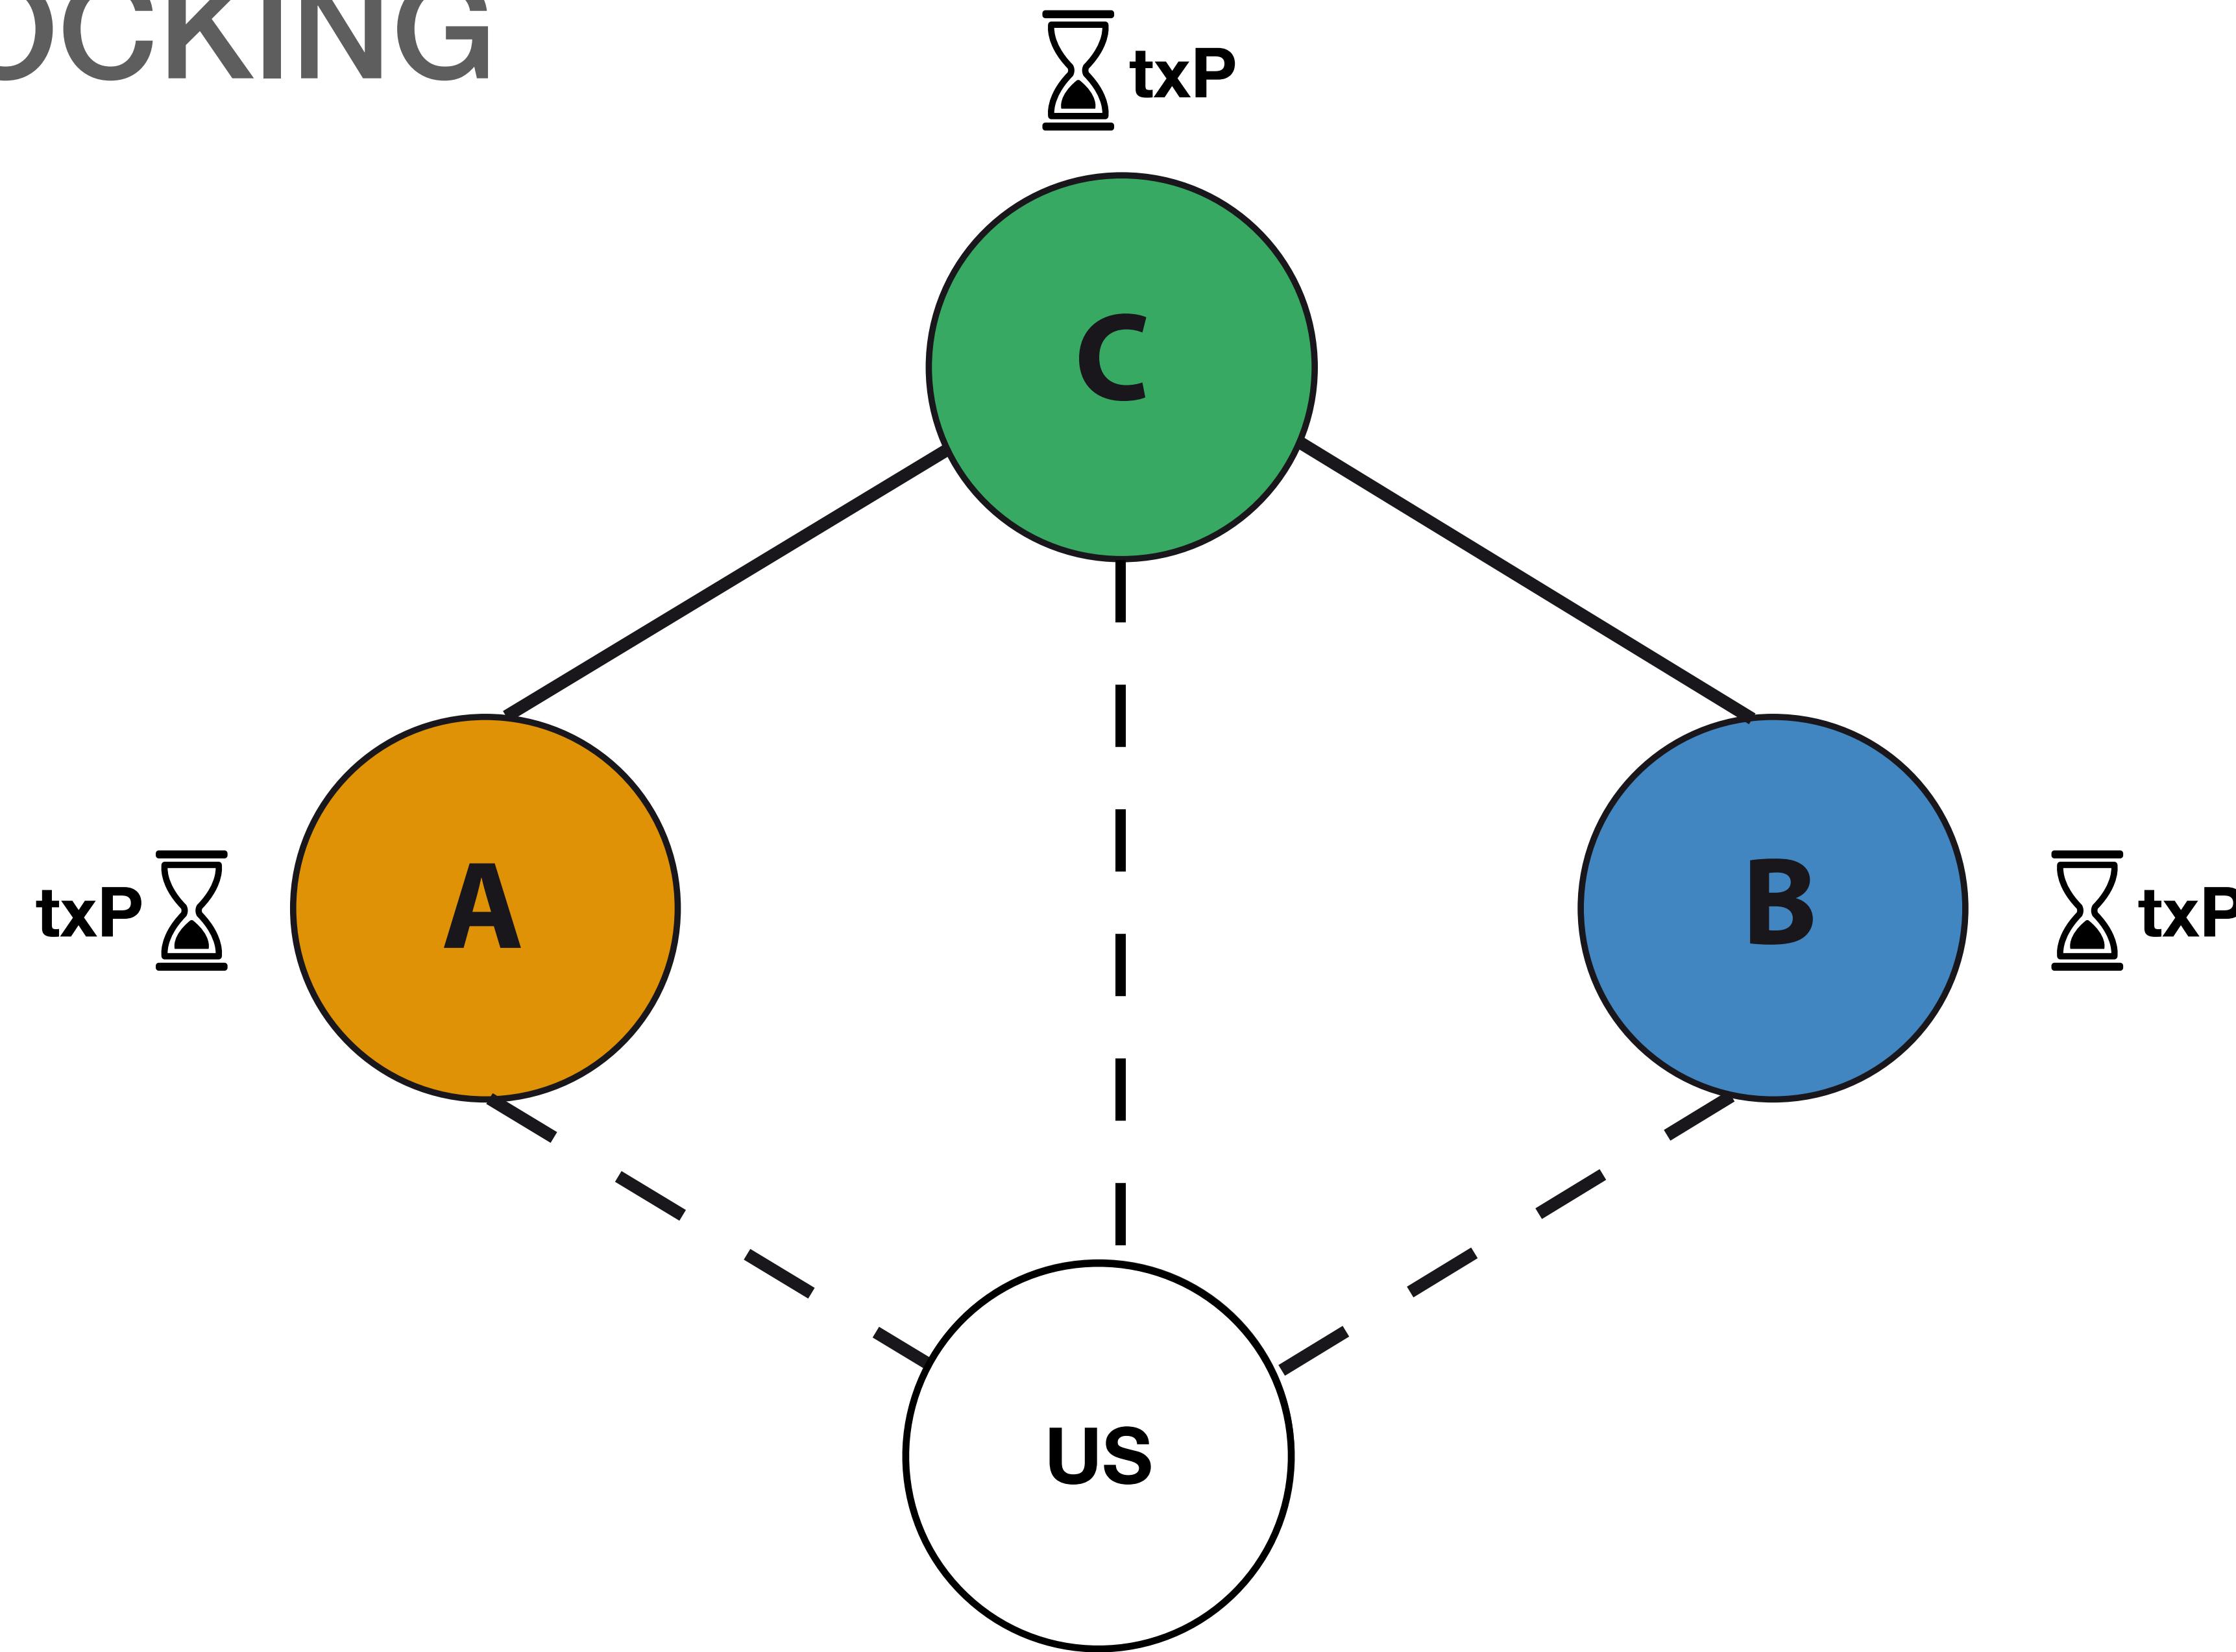
INVBLOCKING



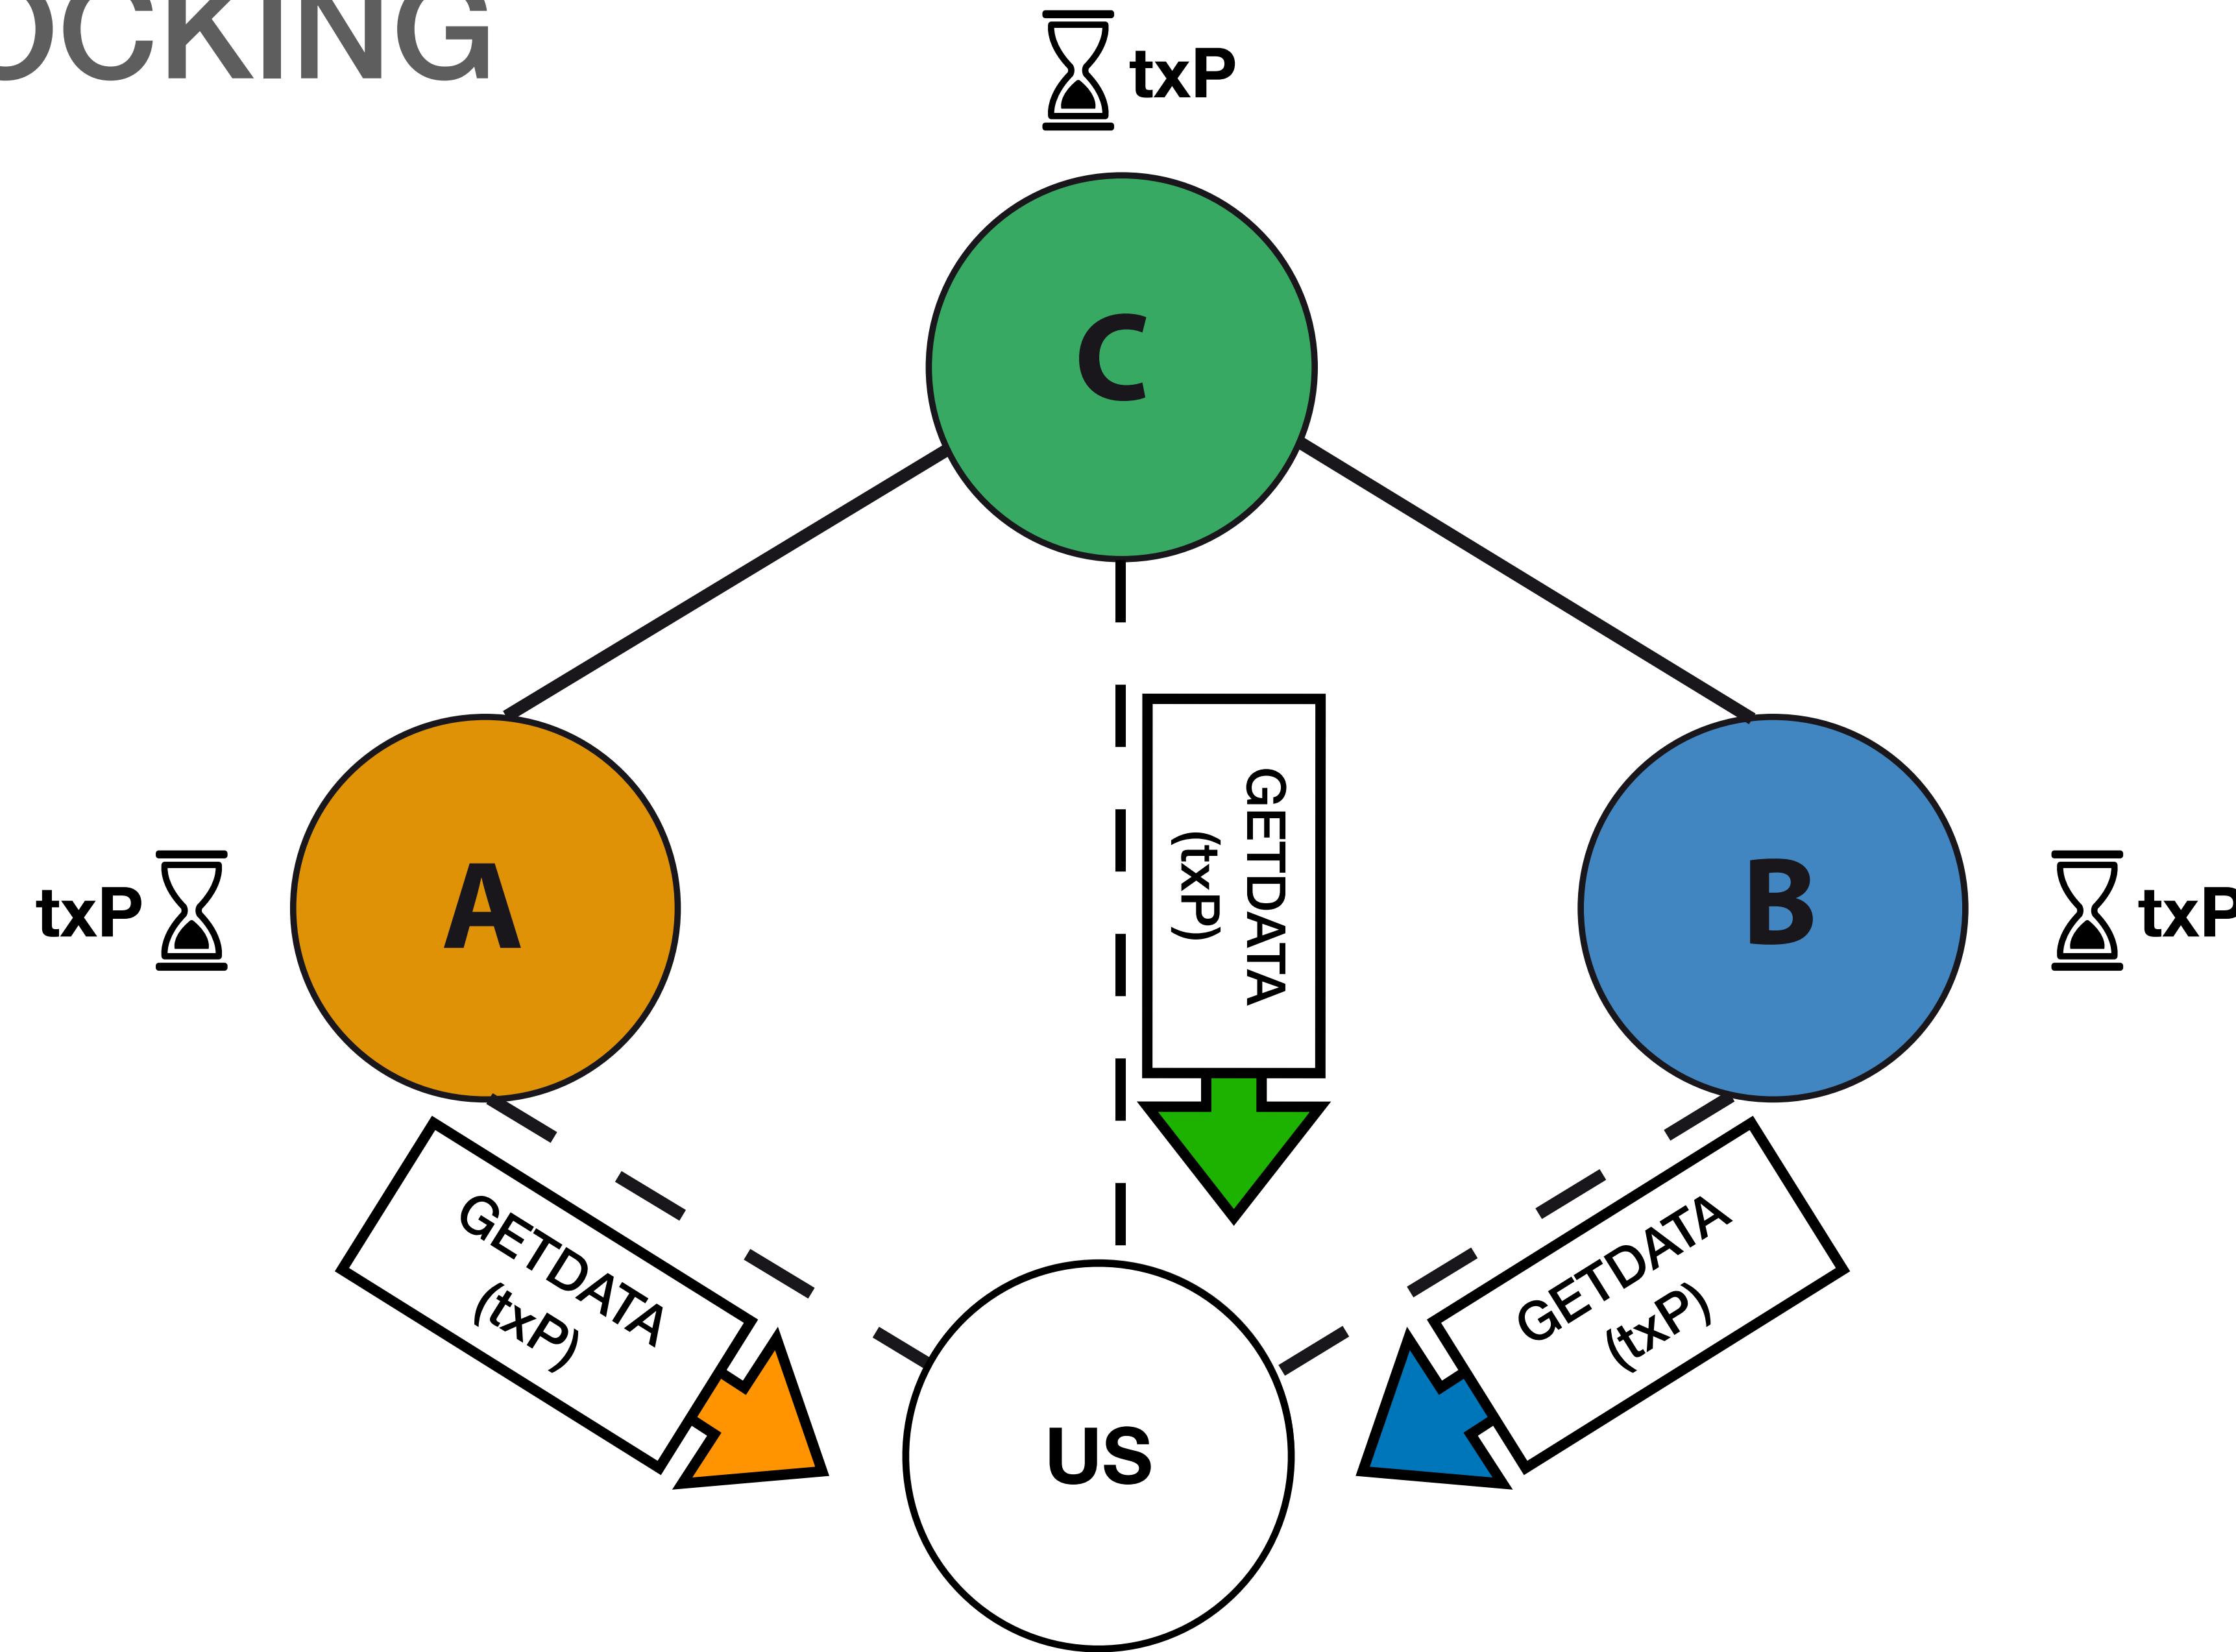
INVBLOCKING



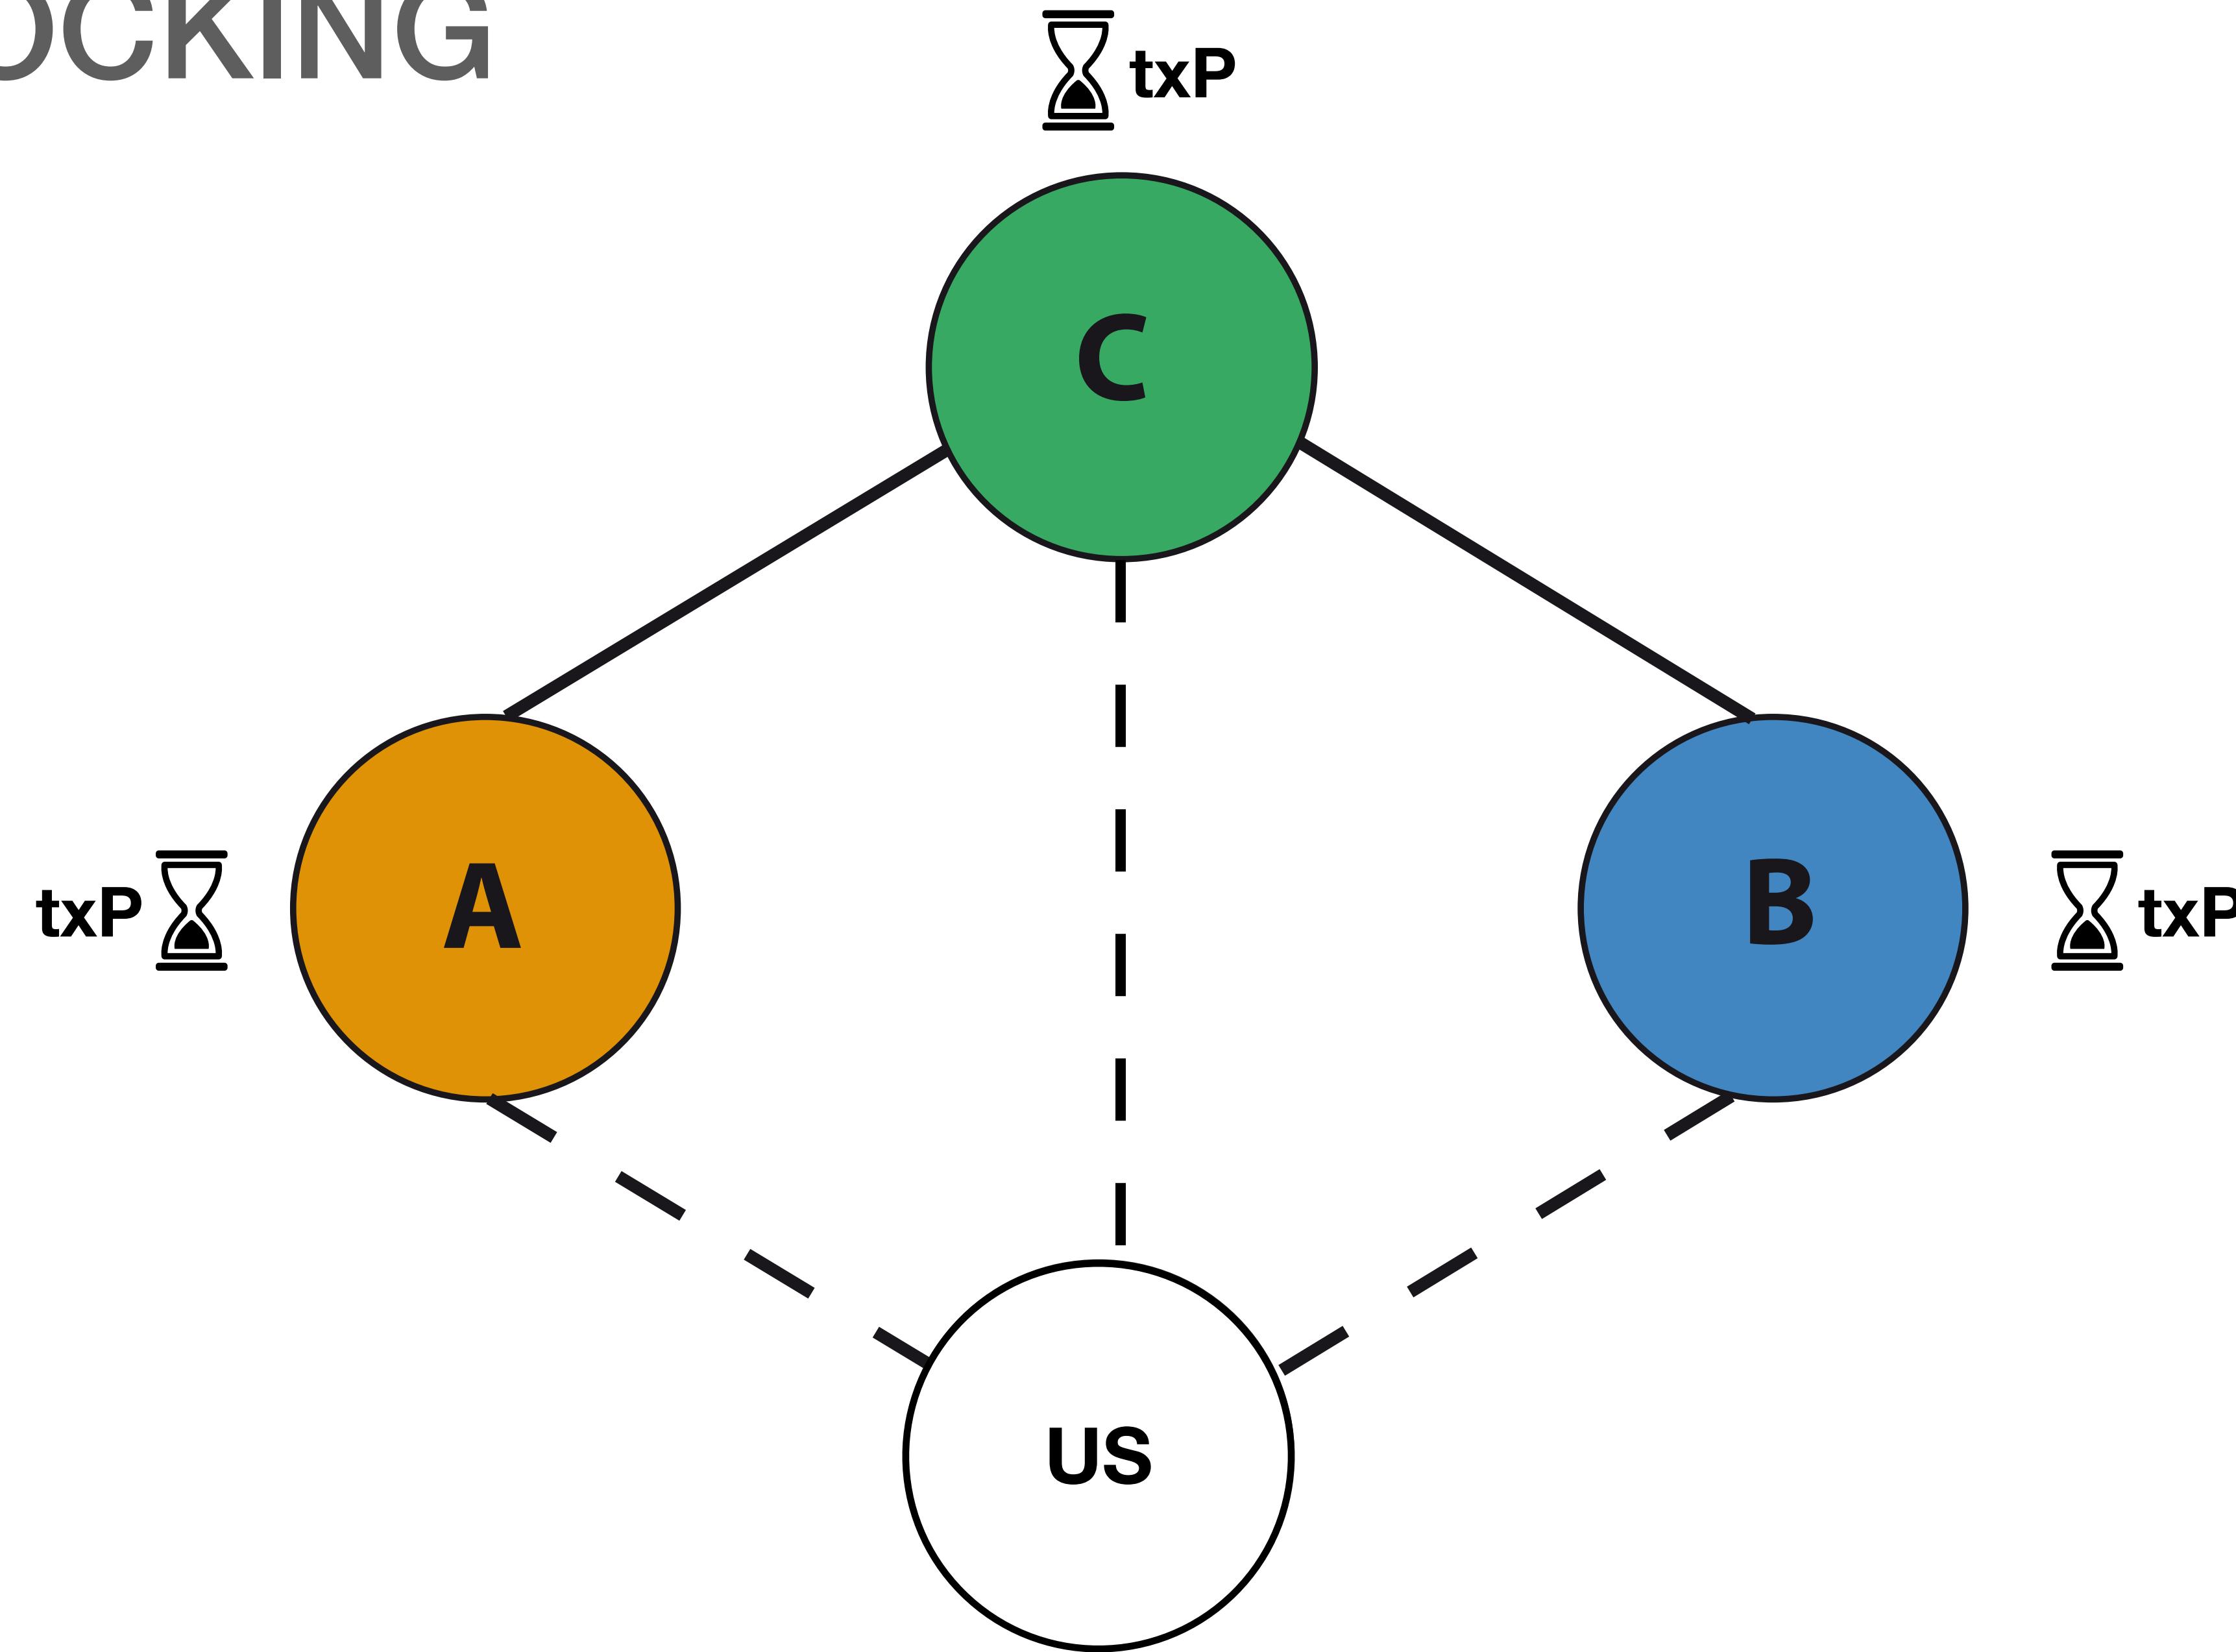
INVBLOCKING



INVBLOCKING



INVBLOCKING



INVBLOCKING

The diagram shows two nodes: a green circle labeled 'C' at the top and a white circle labeled 'US' at the bottom. A red horizontal bar spans the width of the slide, centered between the two nodes. At each end of this bar, there is a small circular marker: an orange one on the left and a blue one on the right. Solid black lines connect node 'C' to both markers. Dashed black lines connect node 'US' to both markers. Above the red bar, there is a small icon of an hourglass followed by the text 'txP'. A red rectangular box surrounds the text in the center of the slide.

We have a 2-min window where isolation and synchrony are
not a problem!

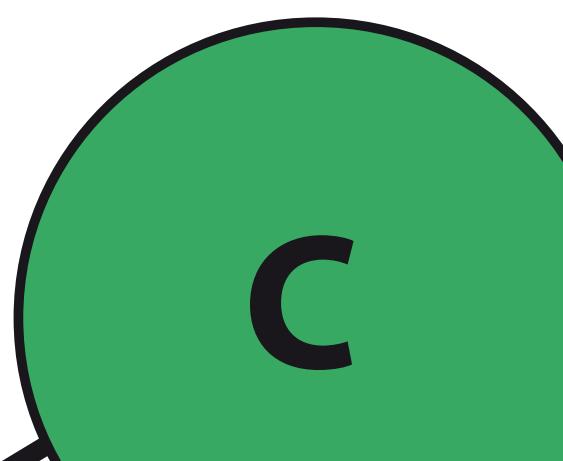
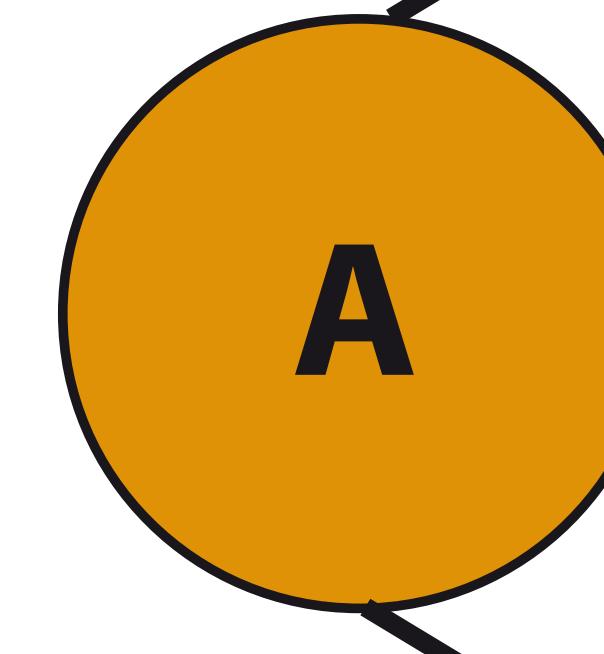
SIMPLIFIED TXPROBE



A's Mempool



txP



txP

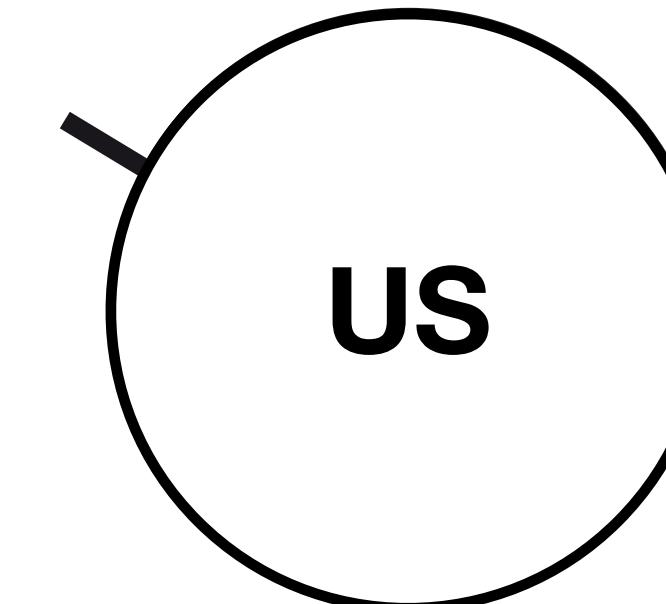
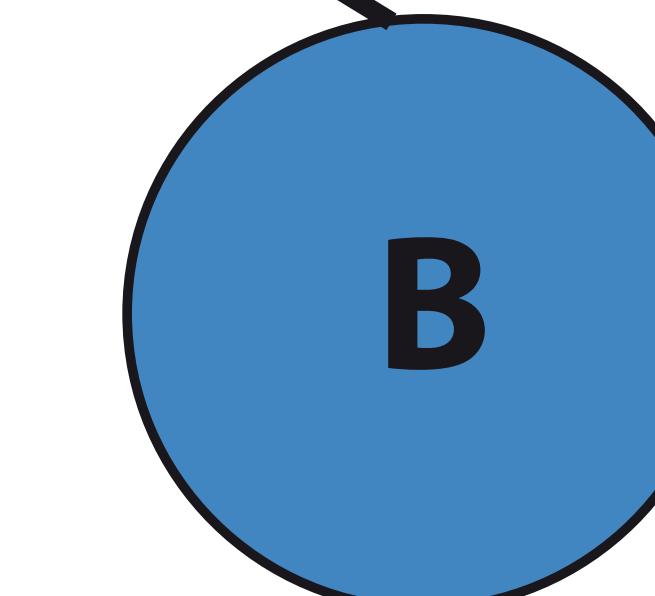
C's Mempool
∅

C's Orphanpool
∅

B's Mempool



txP



B's Orphanpool



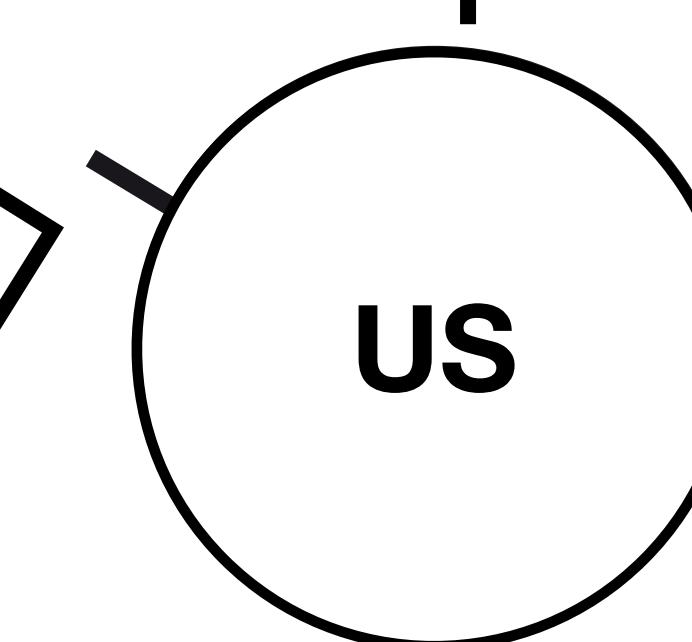
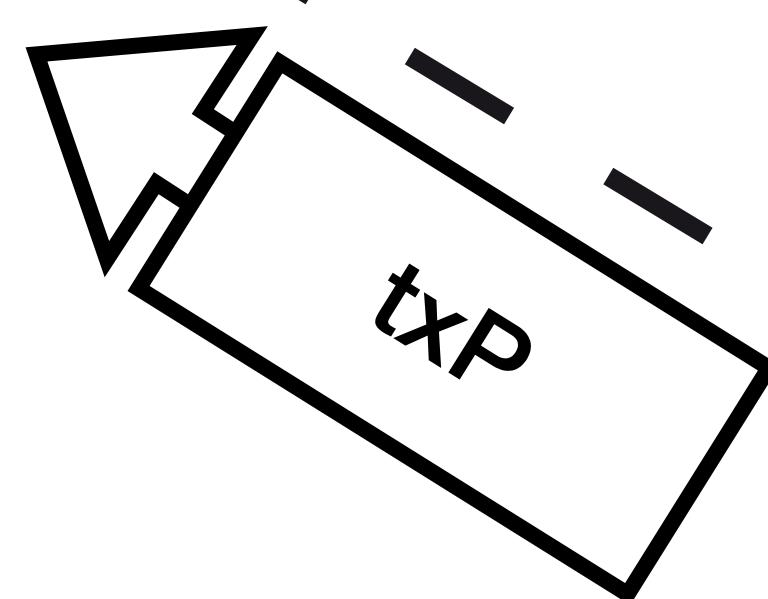
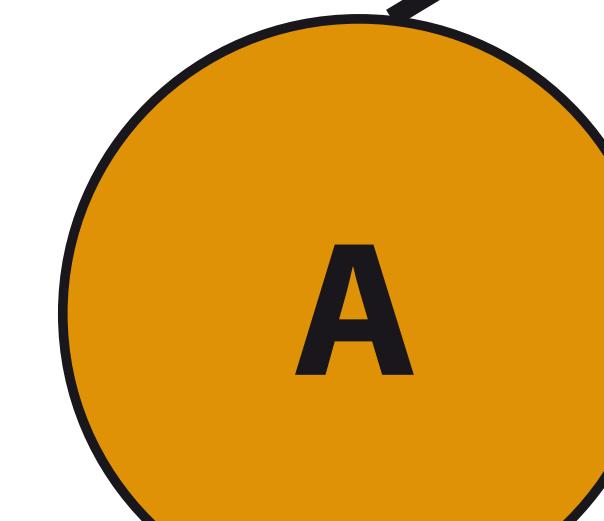
SIMPLIFIED TXPROBE



A's Mempool



txP



txP

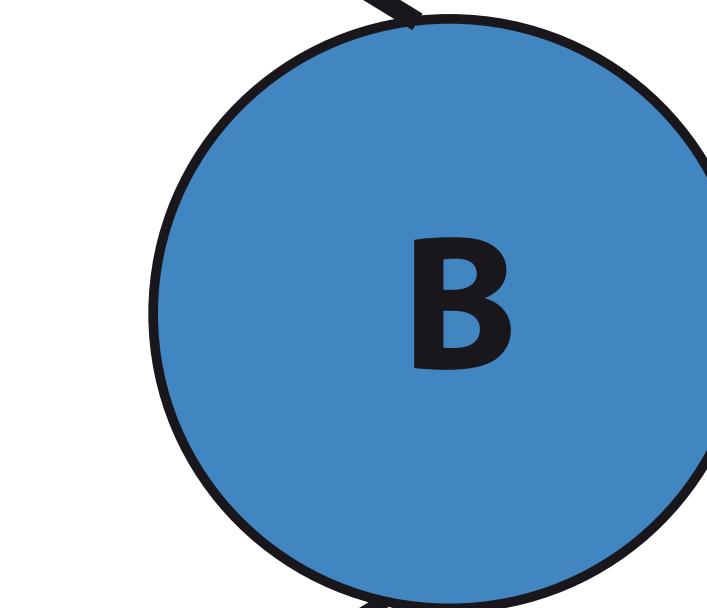
C's Mempool
∅

C's Orphanpool
∅

B's Mempool

∅

txP



B's Orphanpool

∅

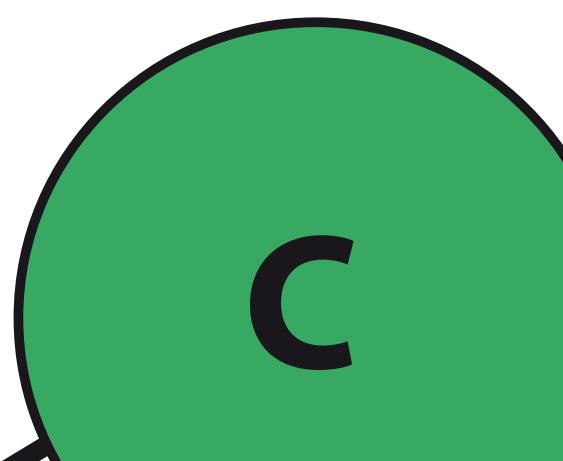
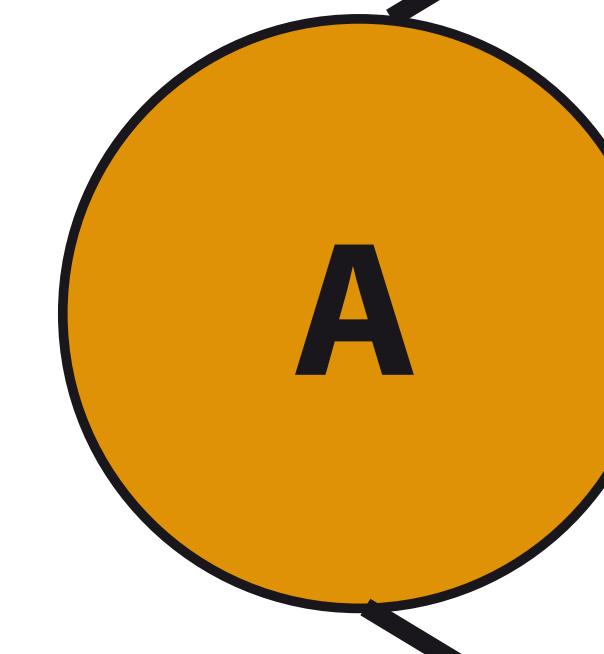
SIMPLIFIED TXPROBE



A's Mempool



txP



txP

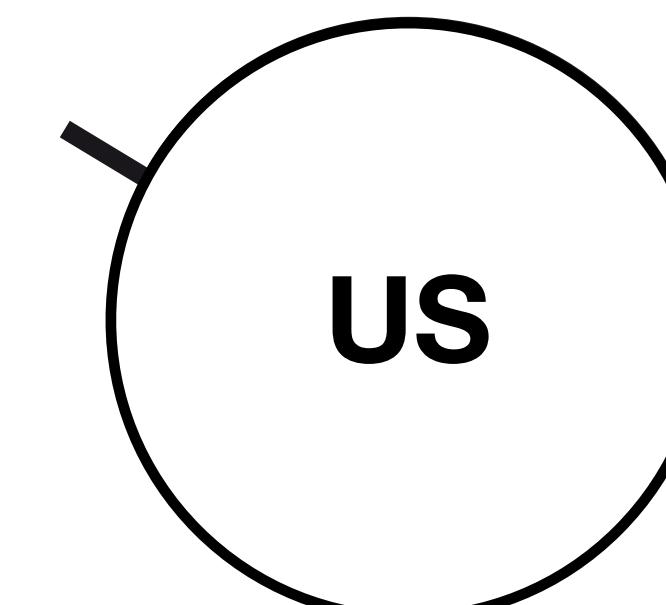
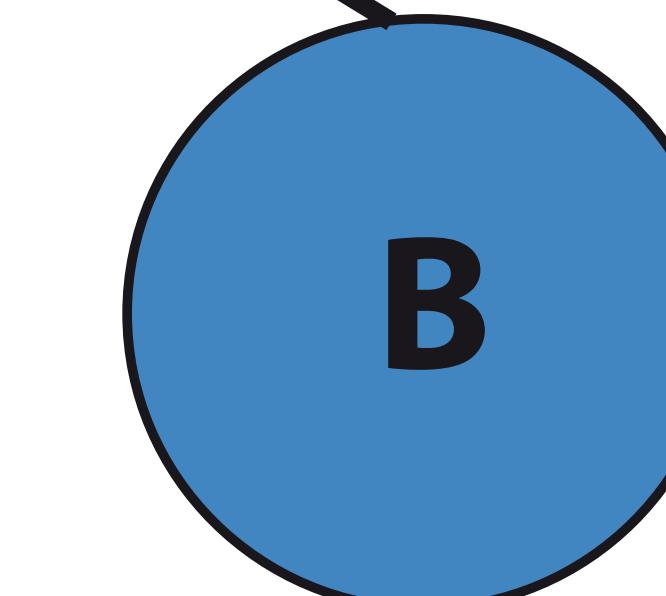
C's Mempool
∅

C's Orphanpool
∅

B's Mempool



txP



B's Orphanpool

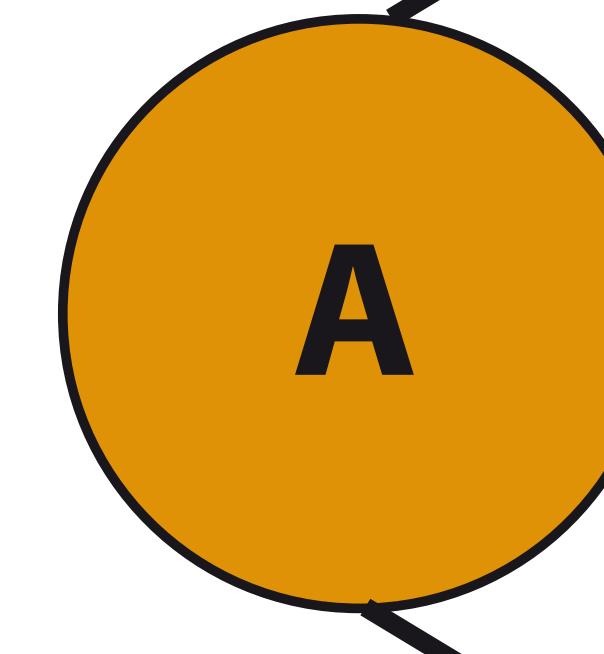


SIMPLIFIED TXPROBE

A's Mempool



txP



txP

C's Mempool
∅

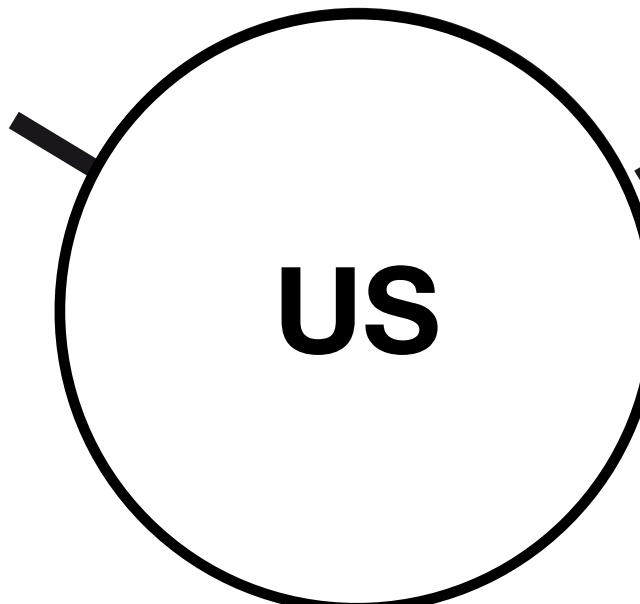
C's Orphanpool
∅

B's Mempool

txP



B's Orphanpool

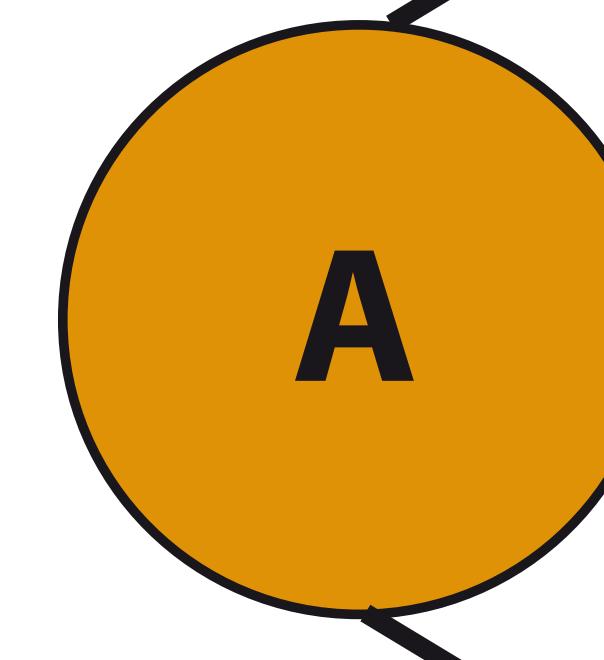


SIMPLIFIED TXPROBE

A's Mempool



txP



INV(txP)

txP

C's Mempool



C's Orphanpool

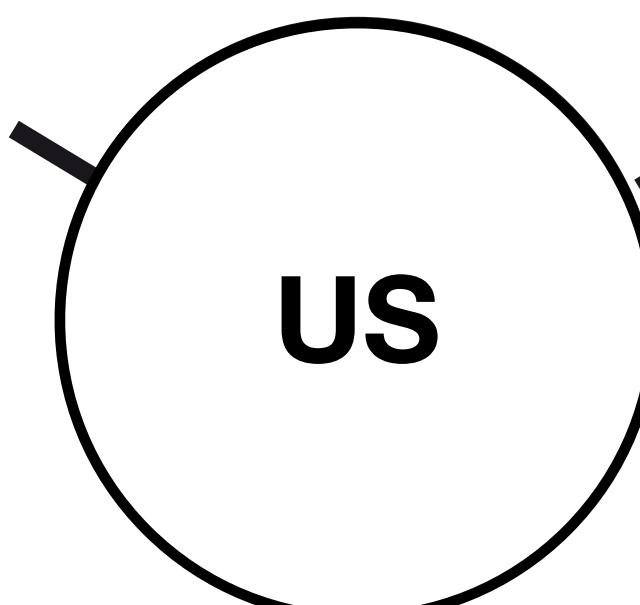


B's Mempool

txP



B's Orphanpool

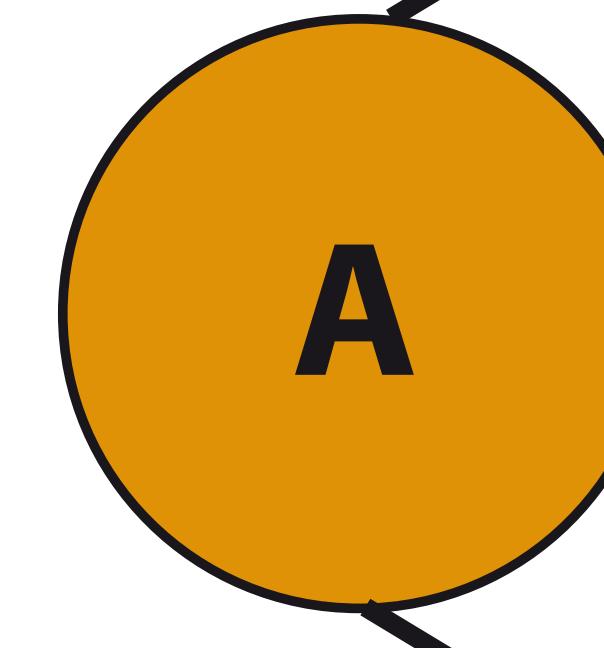


SIMPLIFIED TXPROBE

A's Mempool



txP



txP

C's Mempool
∅

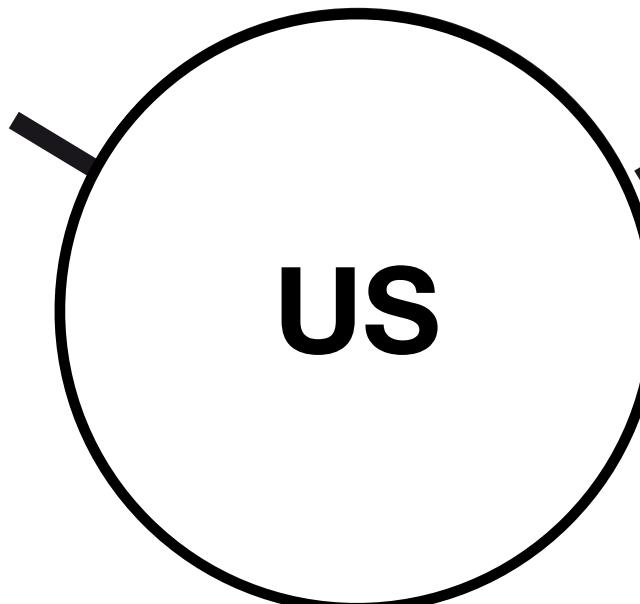
C's Orphanpool
∅

B's Mempool

txP

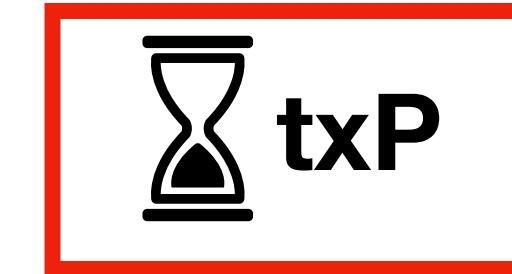
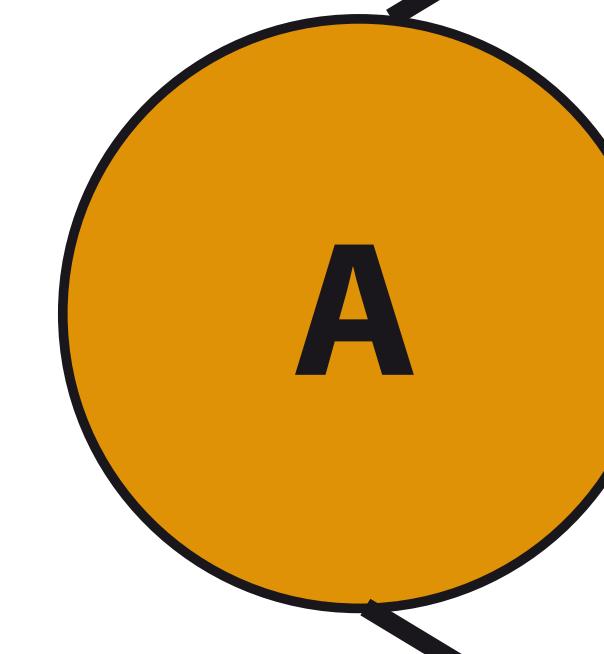


B's Orphanpool



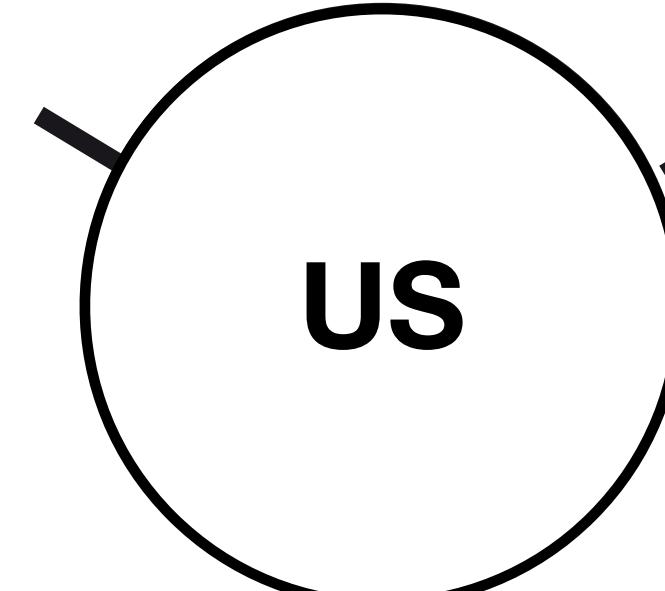
SIMPLIFIED TXPROBE

A's Mempool

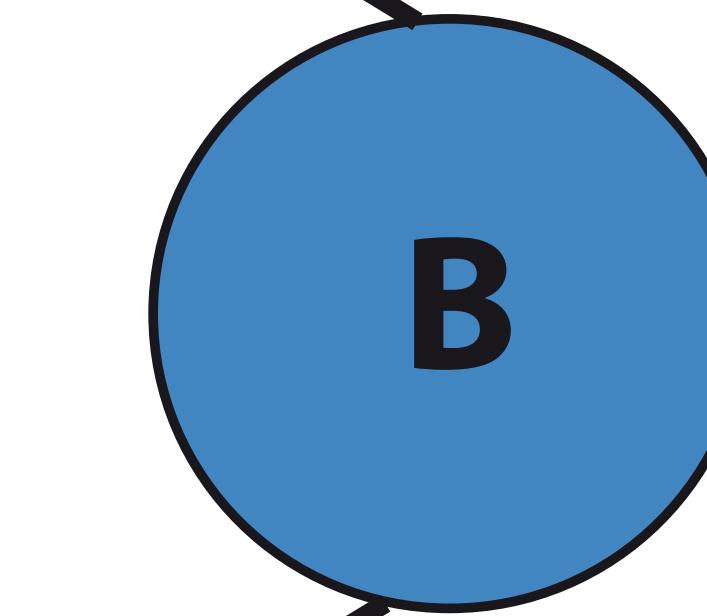


C's Mempool
∅

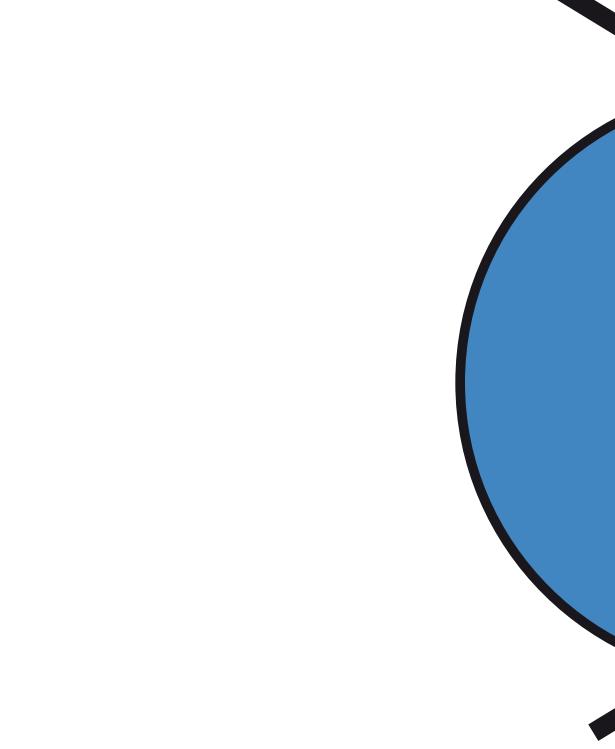
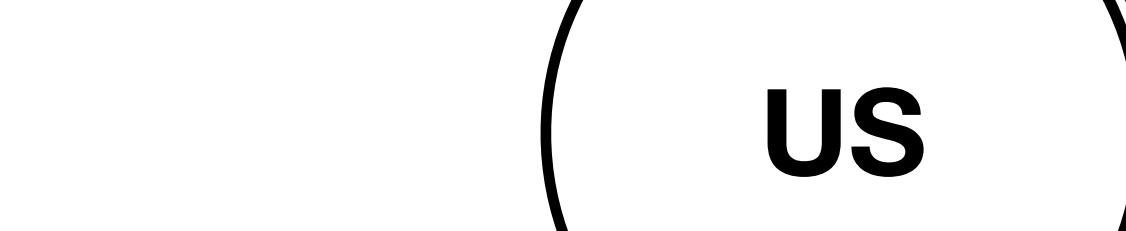
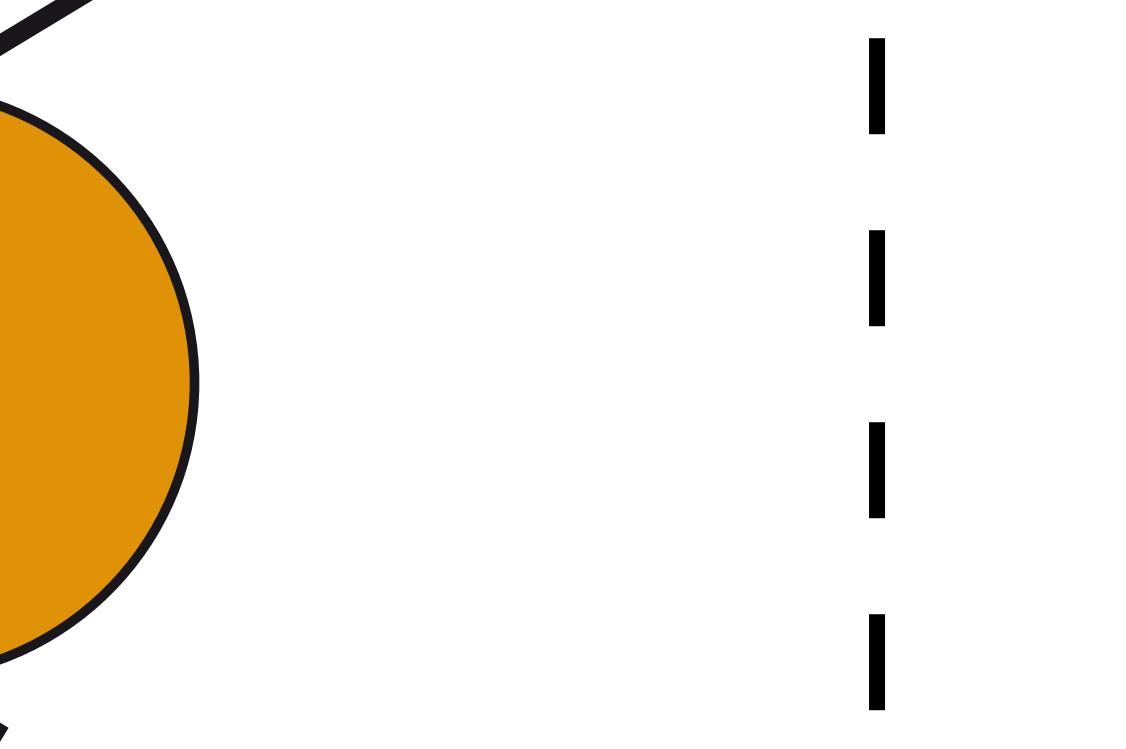
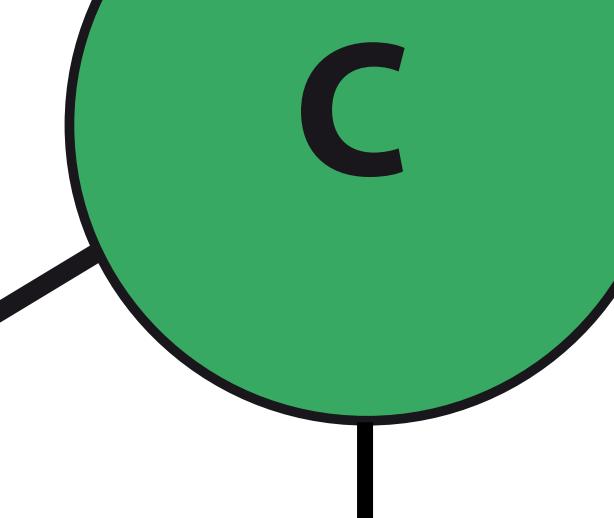
C's Orphanpool
∅



B's Mempool



B's Orphanpool

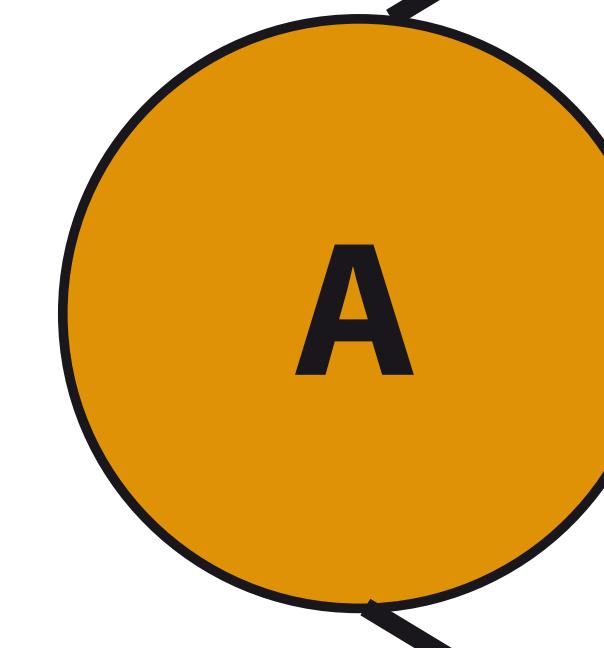


SIMPLIFIED TXPROBE

A's Mempool



txP



txP

C's Mempool
∅

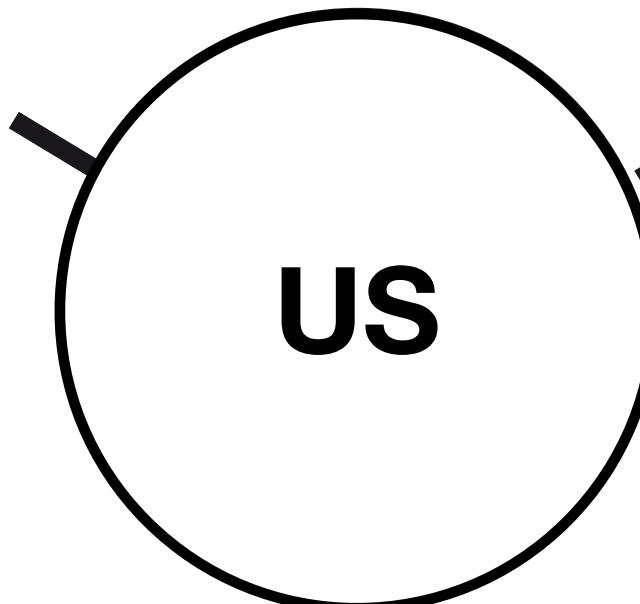
C's Orphanpool
∅

B's Mempool

txP



B's Orphanpool

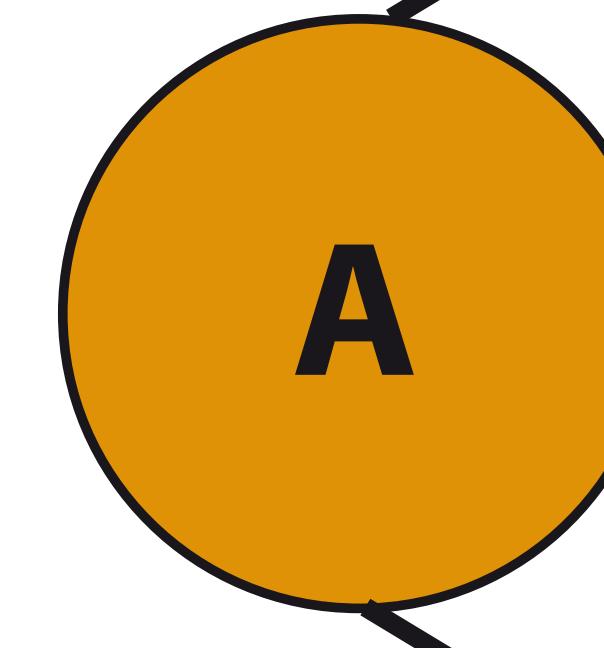


SIMPLIFIED TXPROBE

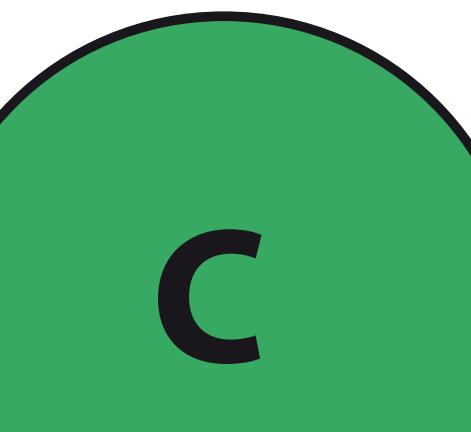
A's Mempool



txP



txP



C's Mempool



C's Orphanpool

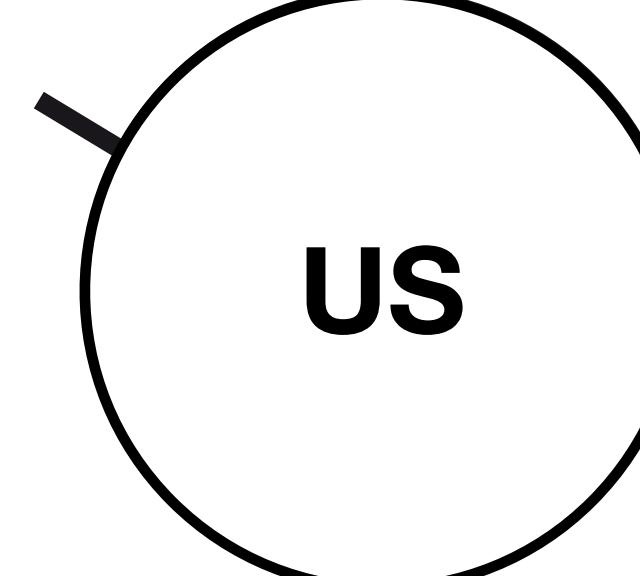


B's Mempool

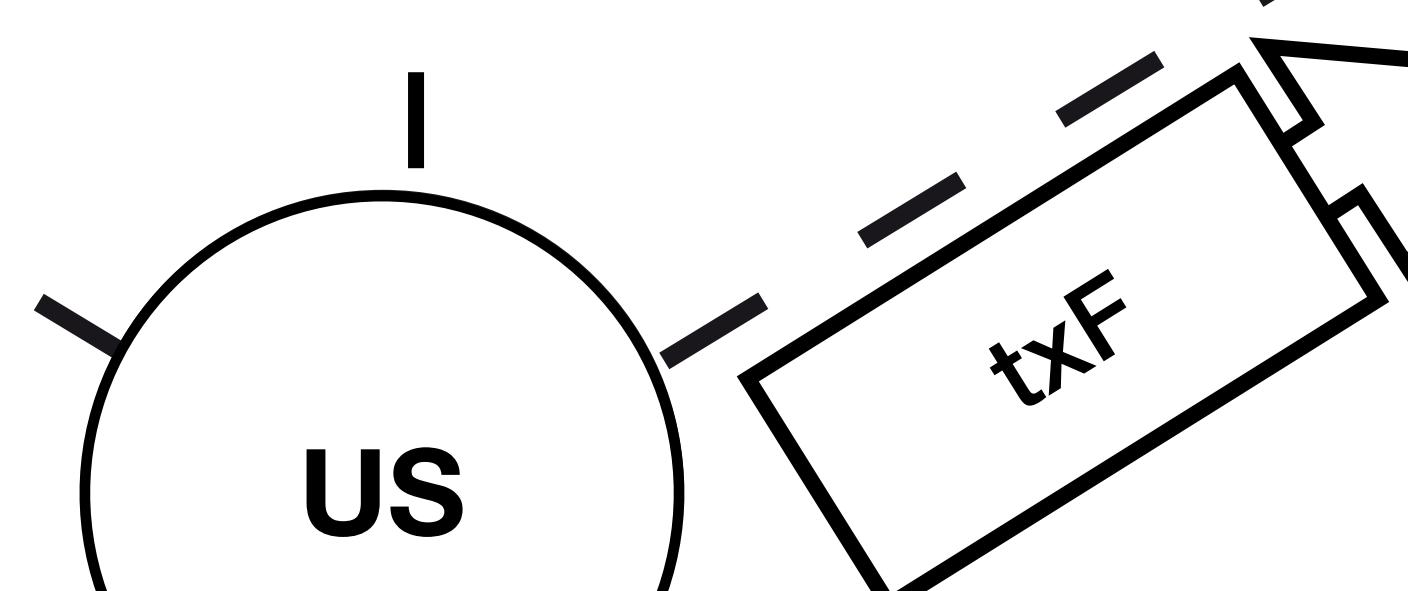
txP



B's Orphanpool



txF

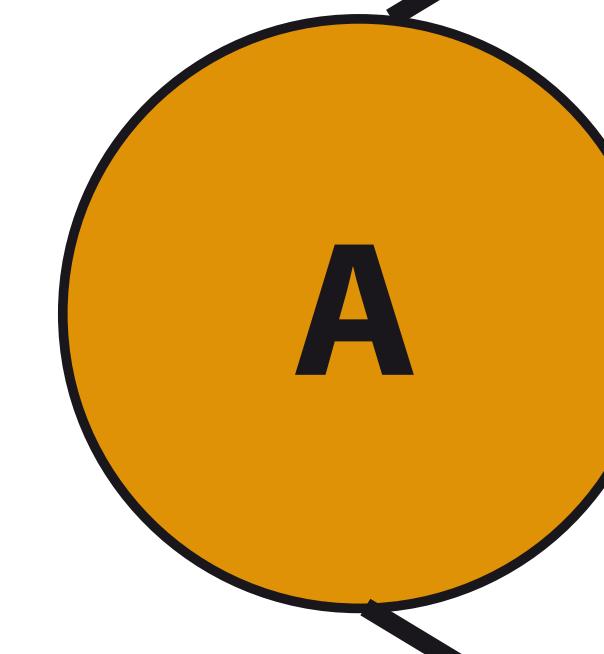


SIMPLIFIED TXPROBE

A's Mempool



txP



txP

C's Mempool
∅

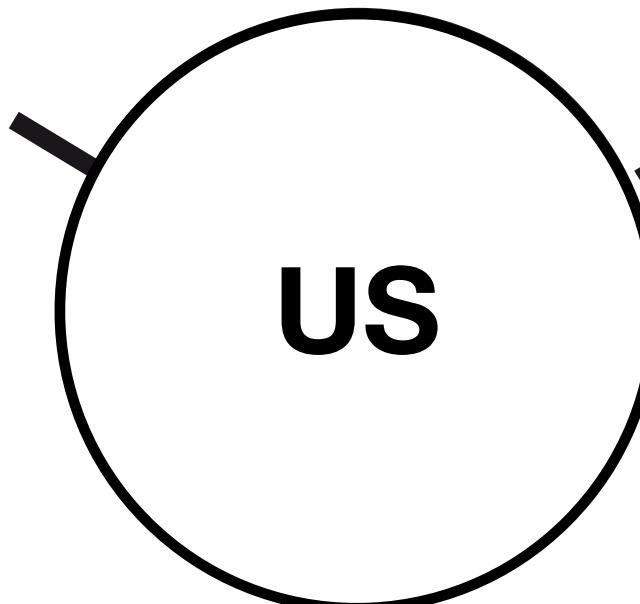
C's Orphanpool
∅

B's Mempool

txP

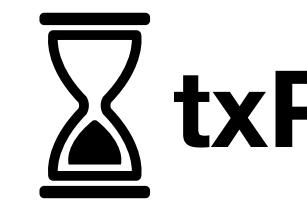
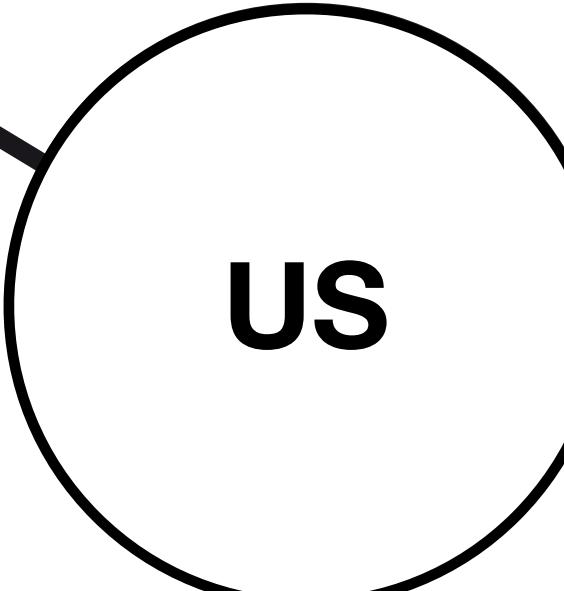
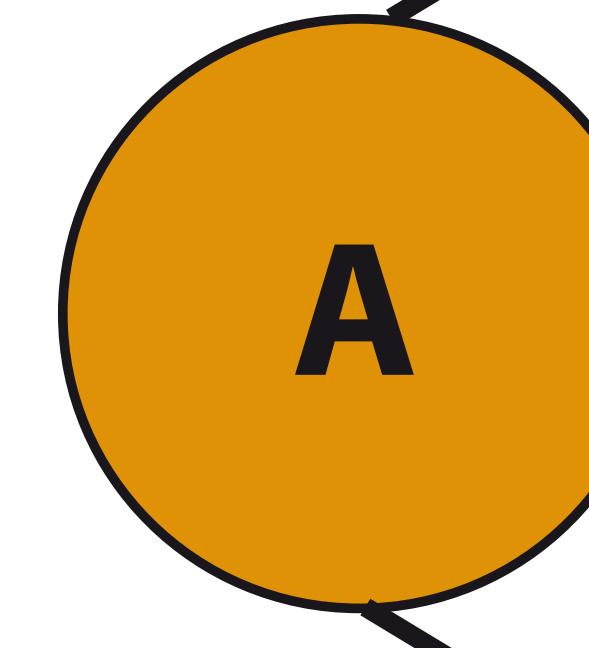


B's Orphanpool



SIMPLIFIED TXPROBE

A's Mempool



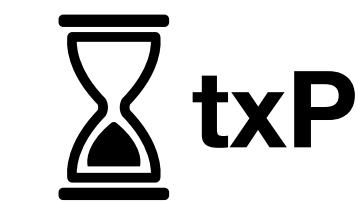
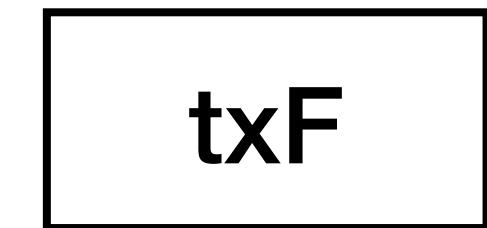
C's Mempool



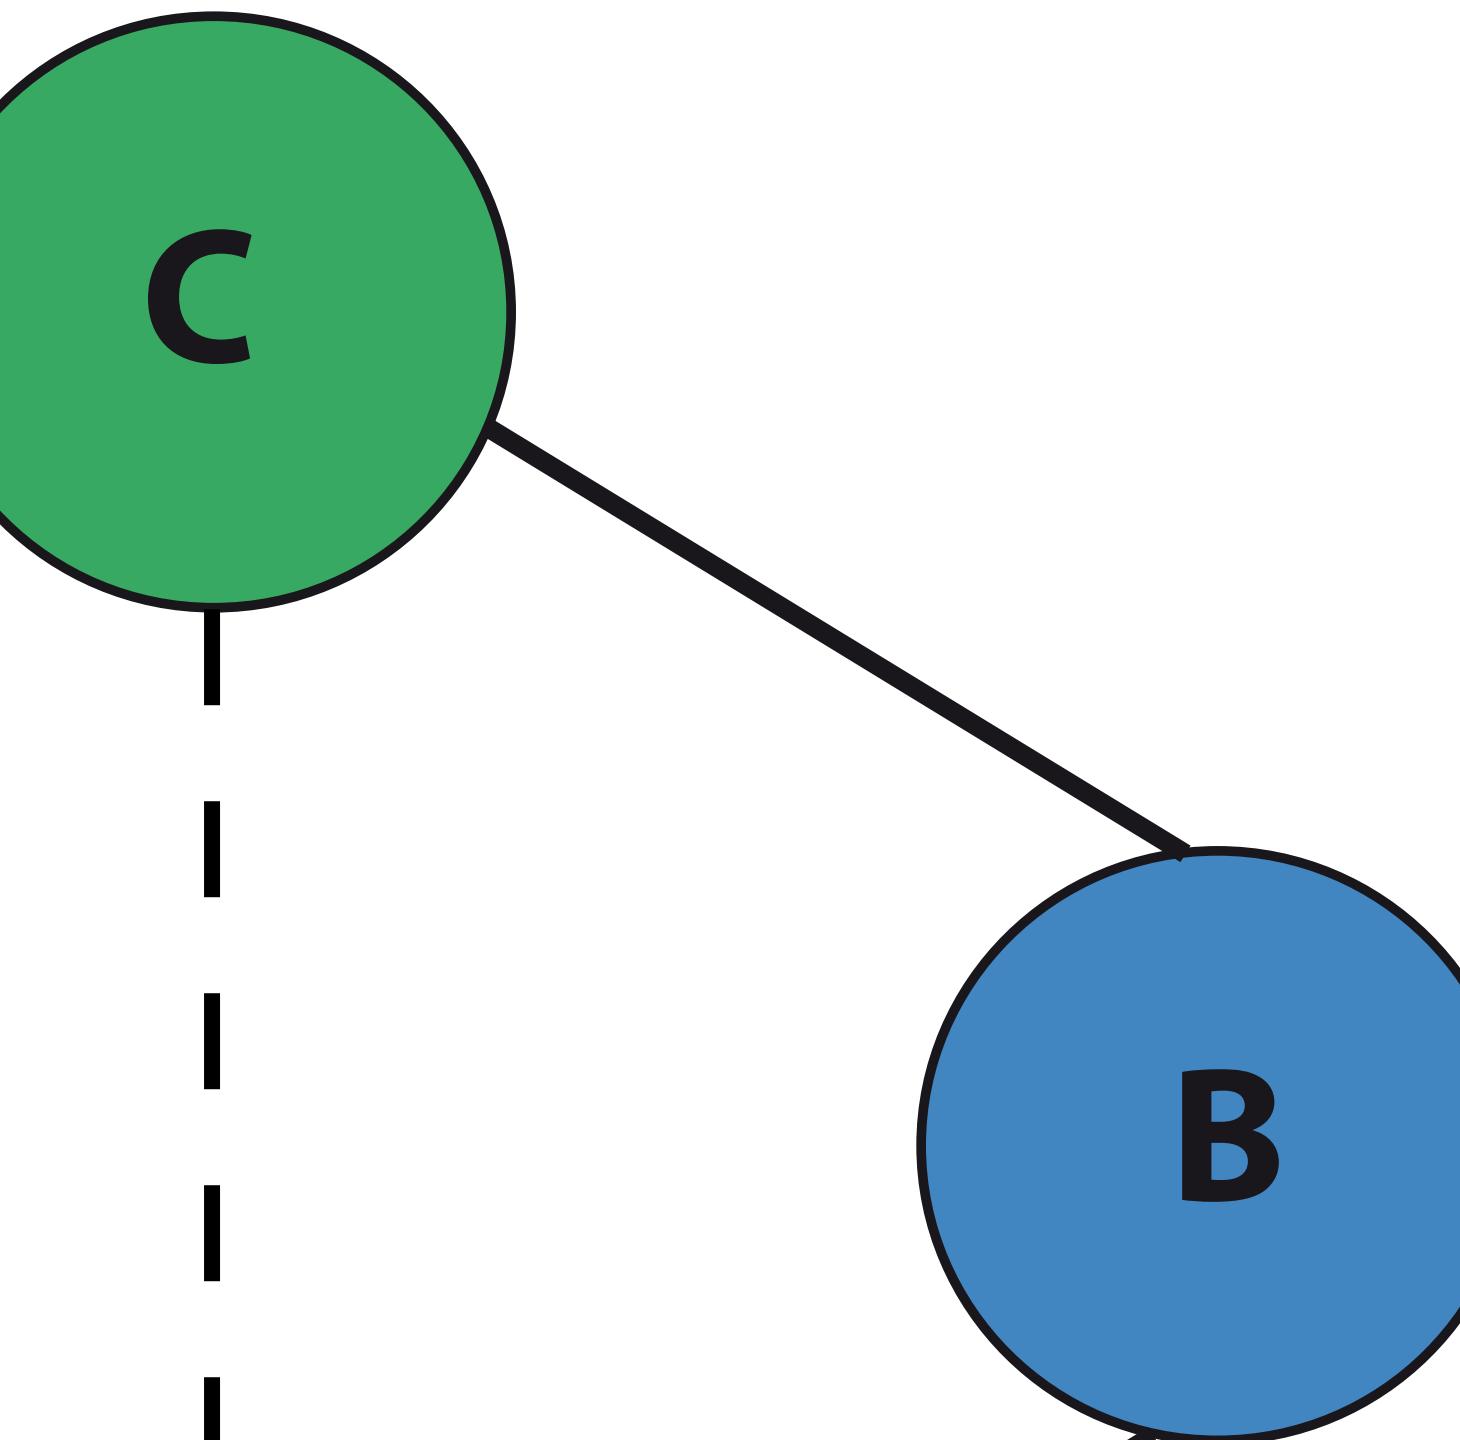
C's Orphanpool



B's Mempool

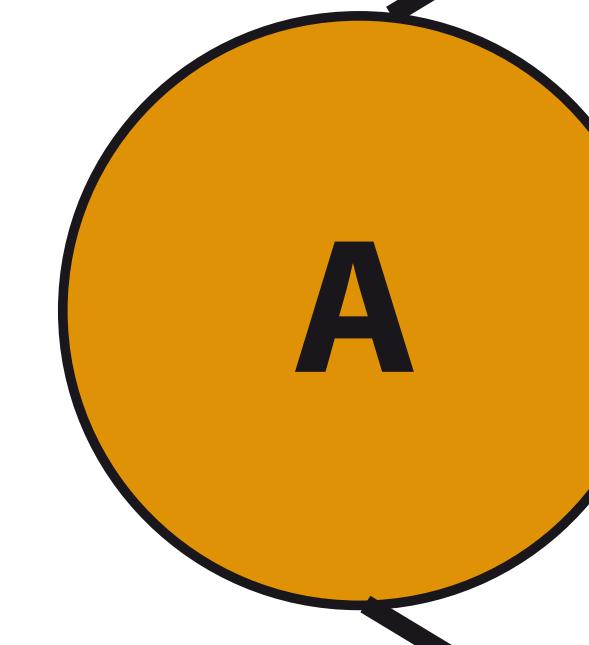


B's Orphanpool



SIMPLIFIED TXPROBE

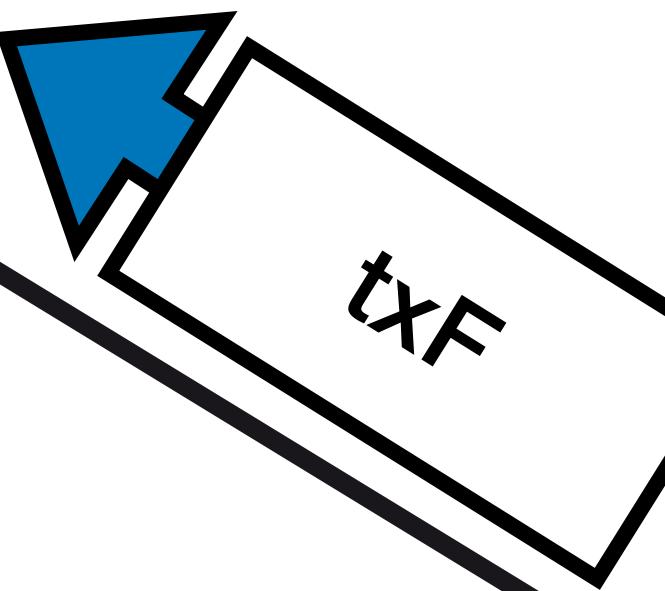
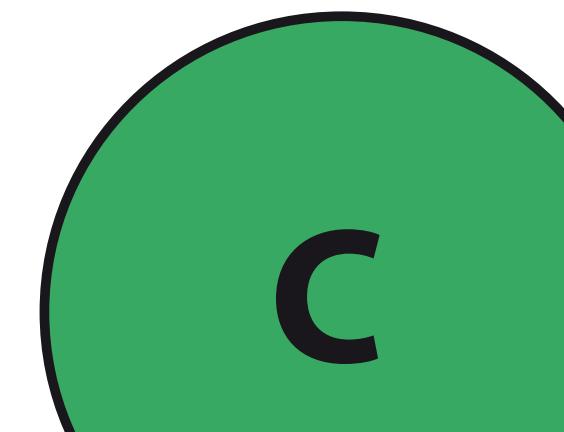
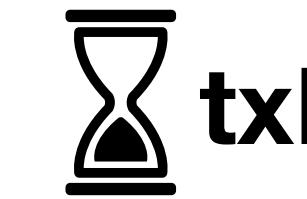
A's Mempool

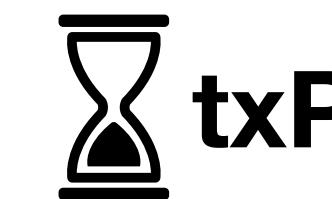
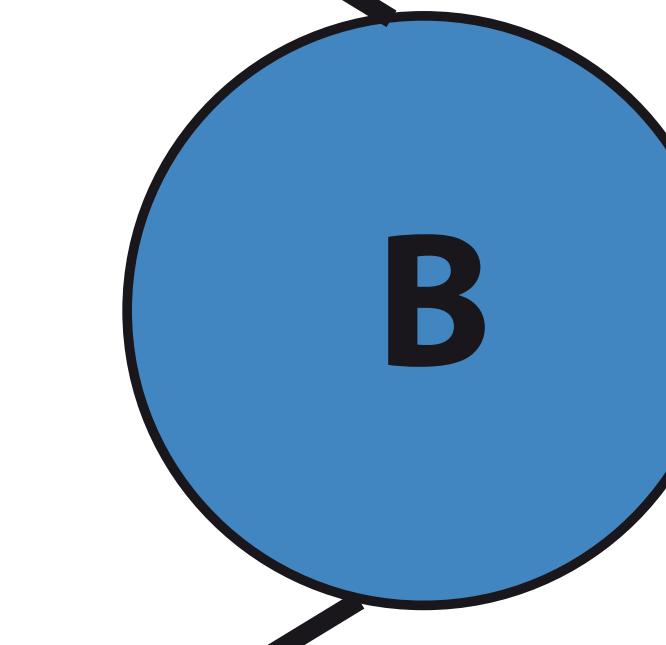


C's Mempool



C's Orphanpool

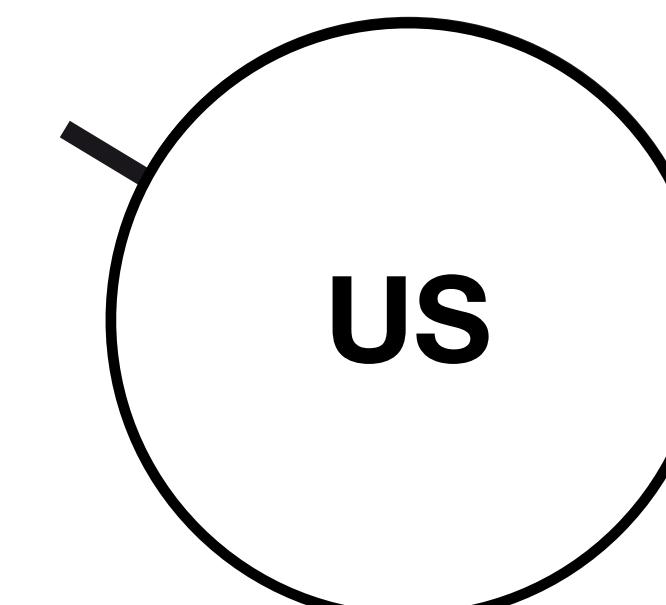




B's Mempool



B's Orphanpool

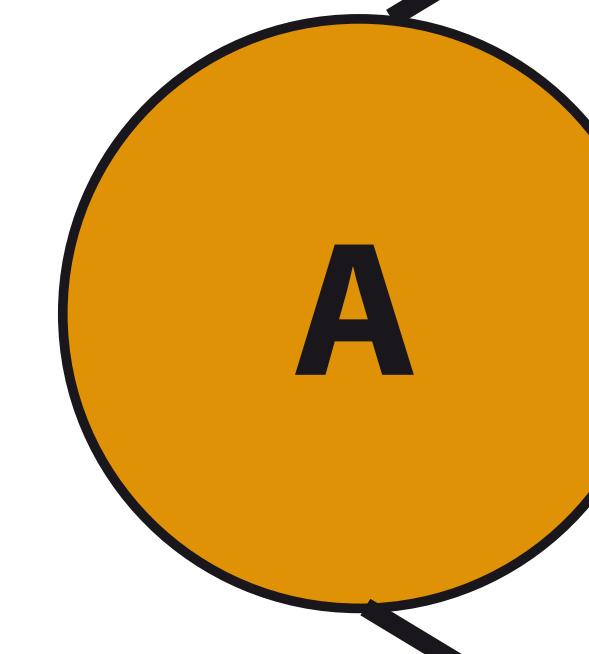


SIMPLIFIED TXPROBE

A's Mempool



txP



A

US

txP

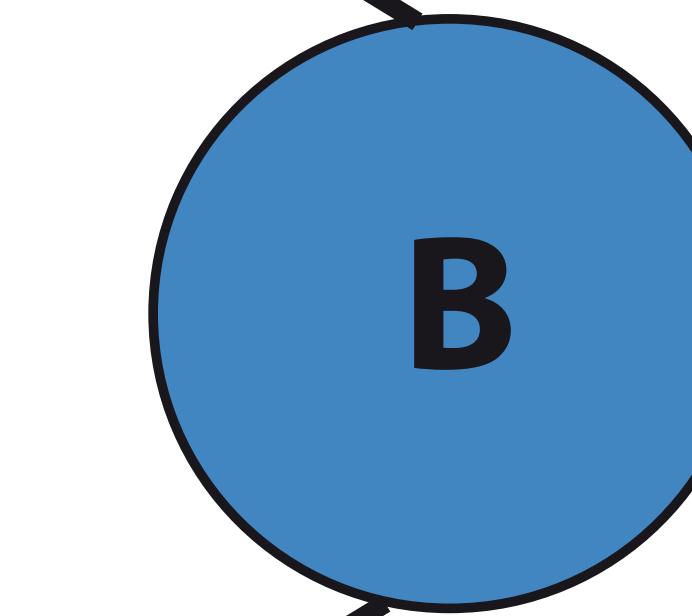
C's Mempool
∅

C's Orphanpool
∅

B's Mempool



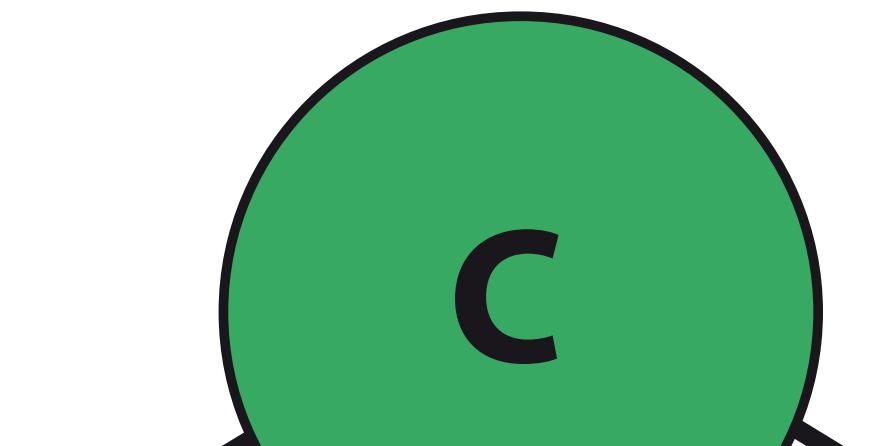
txP



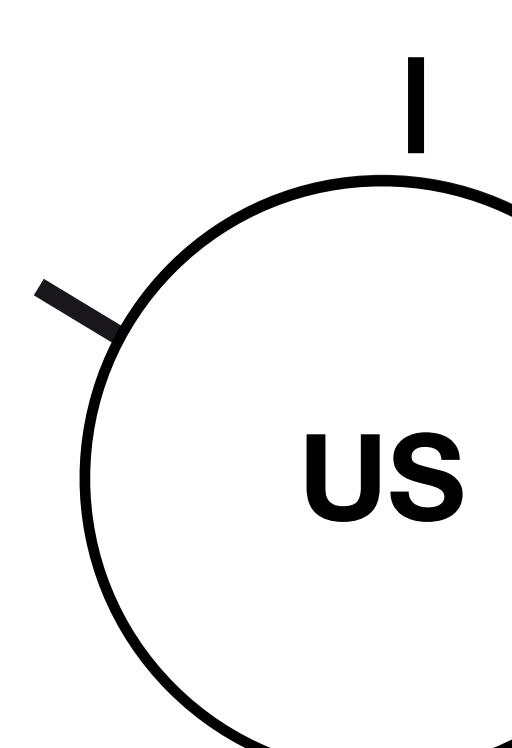
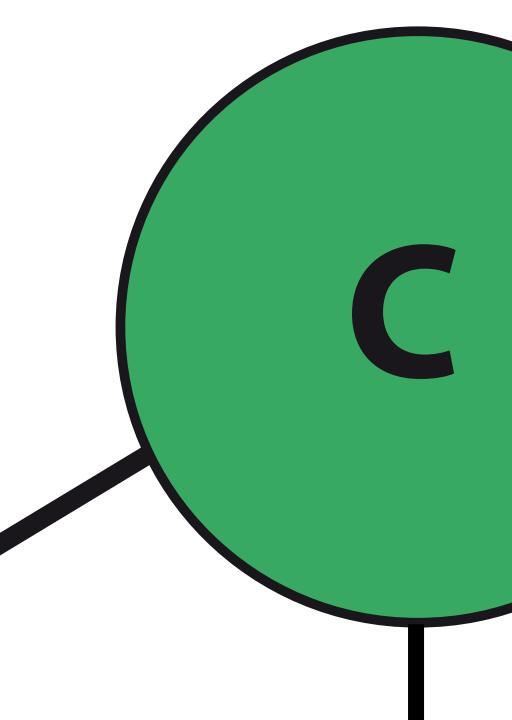
B

B's Orphanpool

∅



C



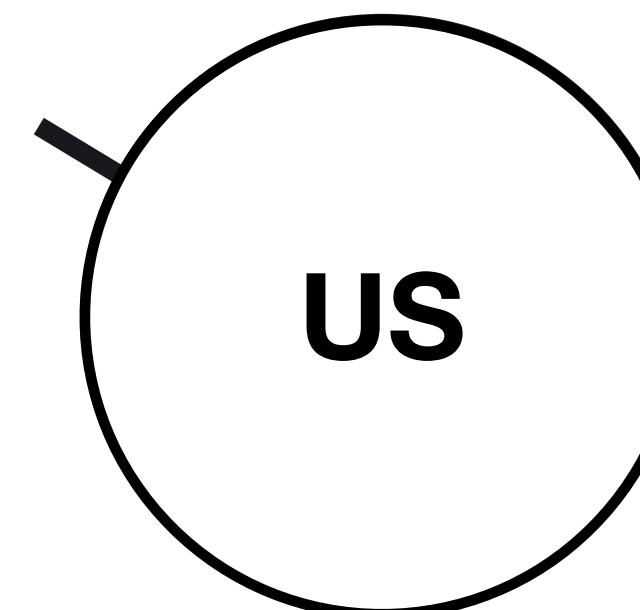
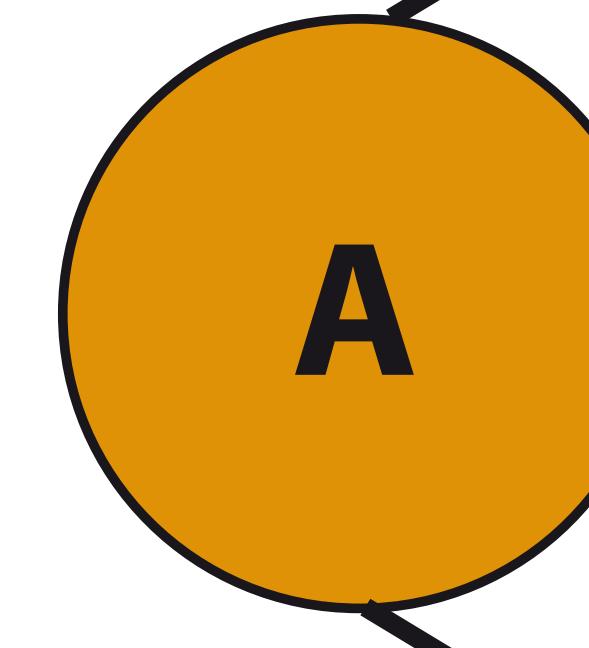
SIMPLIFIED TXPROBE



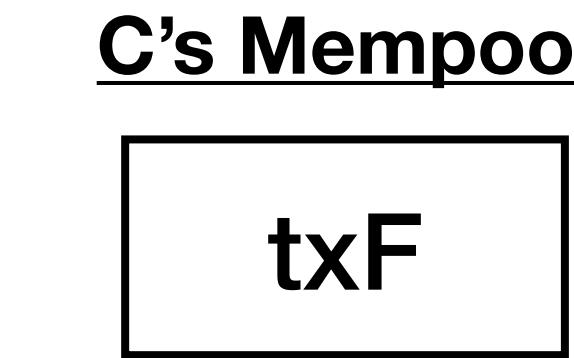
A's Mempool



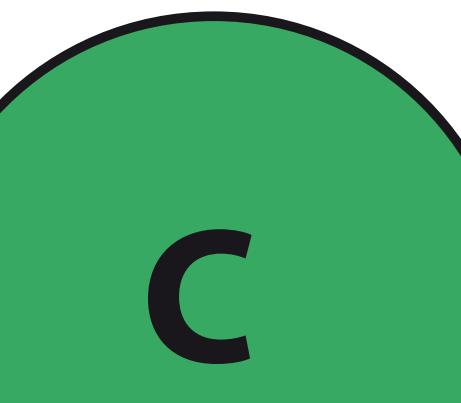
txP



txP



C's Orphanpool
∅



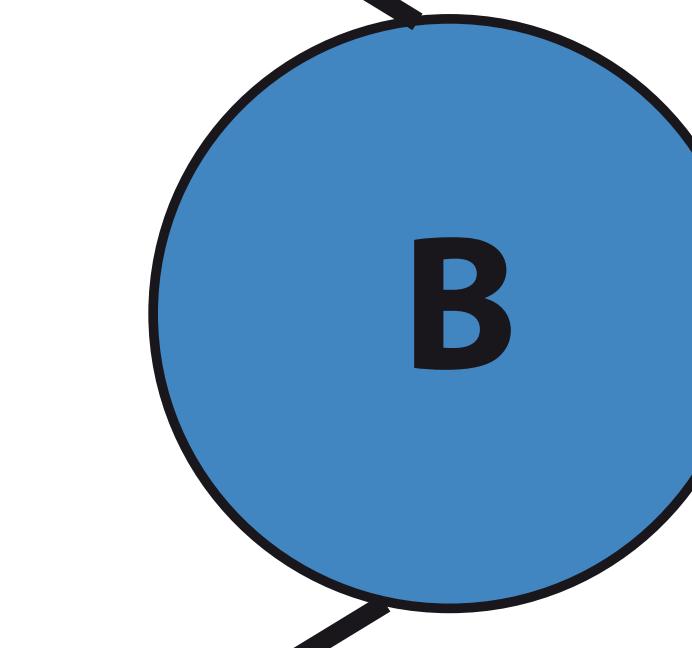
txF



B's Mempool



txP



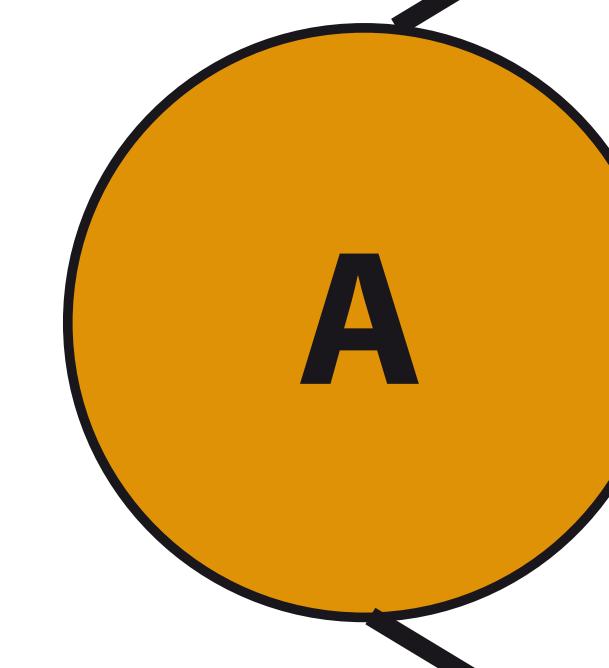
B's Orphanpool
∅

SIMPLIFIED TXPROBE

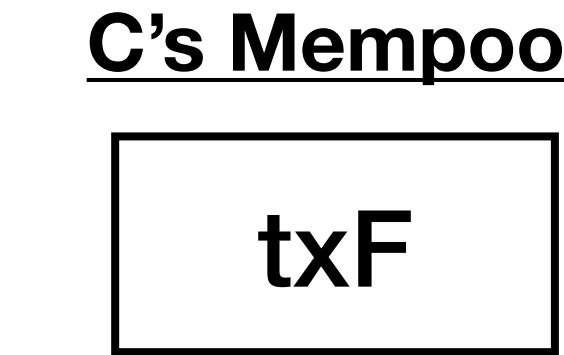
A's Mempool



txP



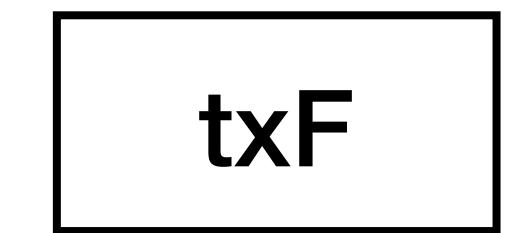
txP



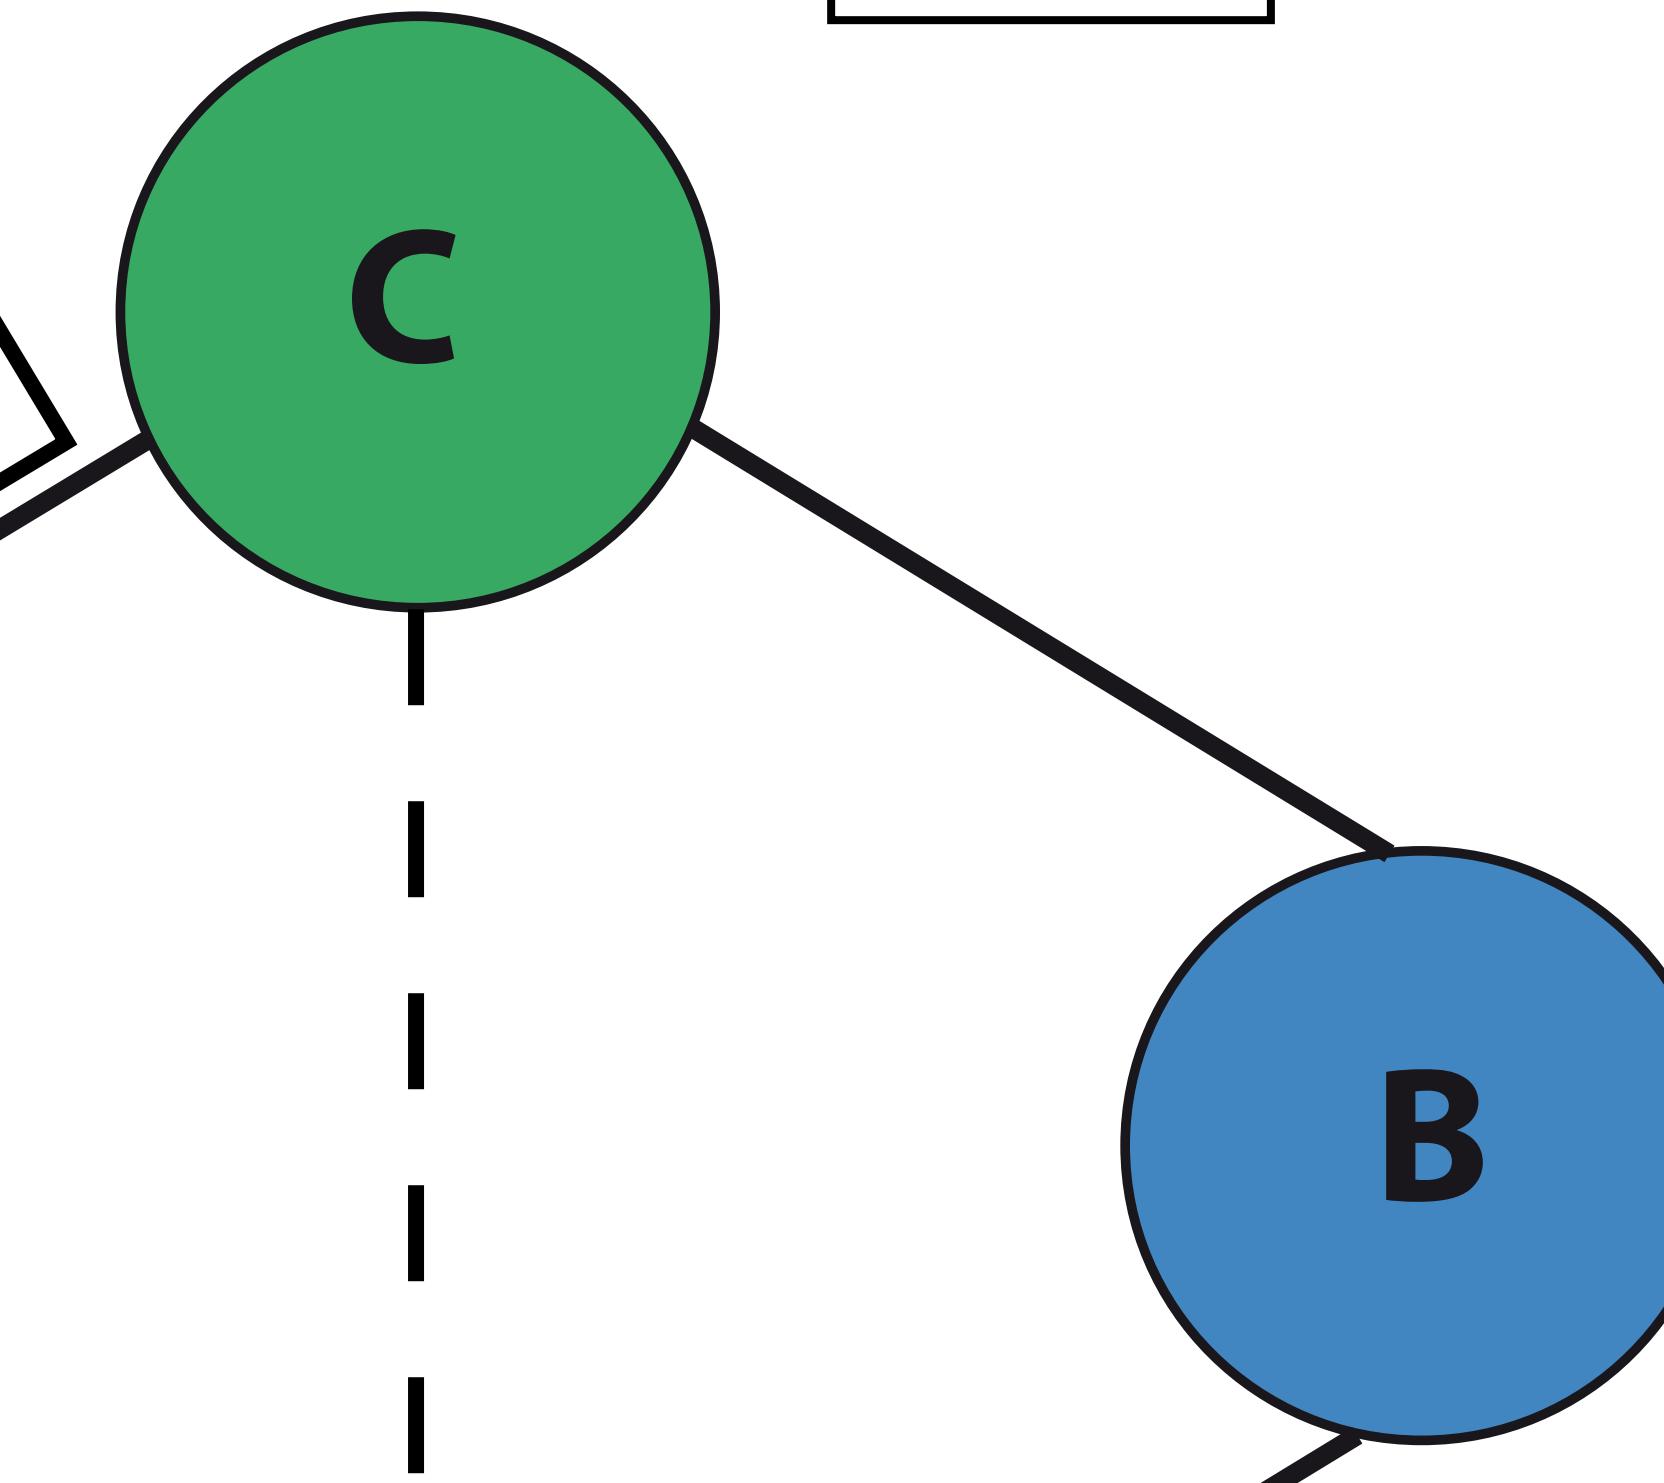
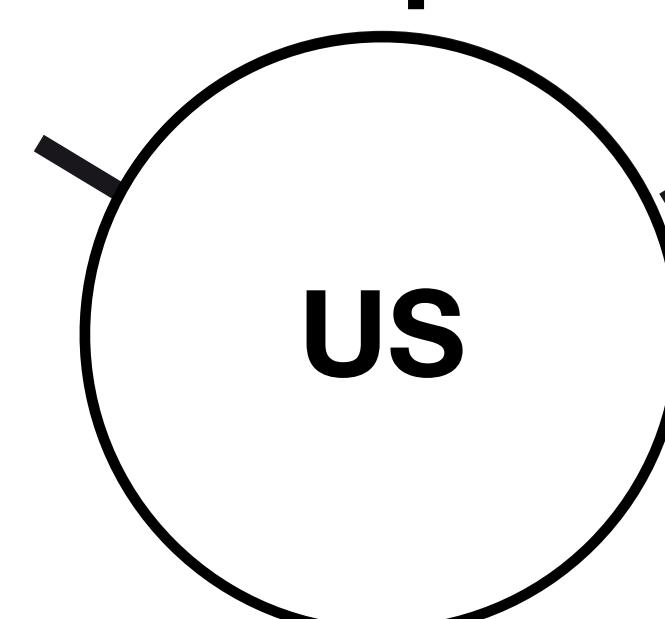
C's Orphanpool



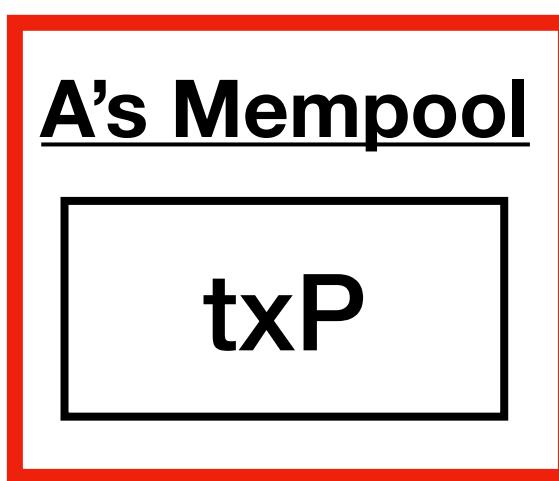
B's Mempool



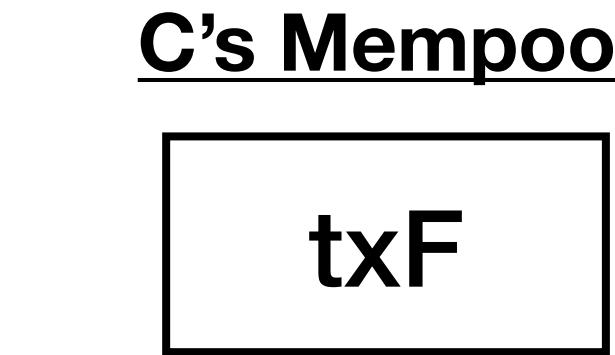
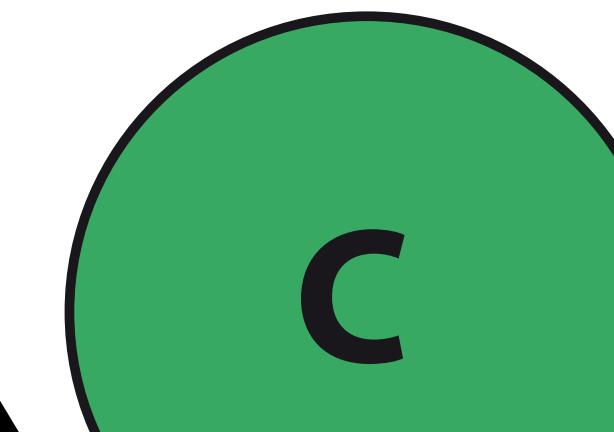
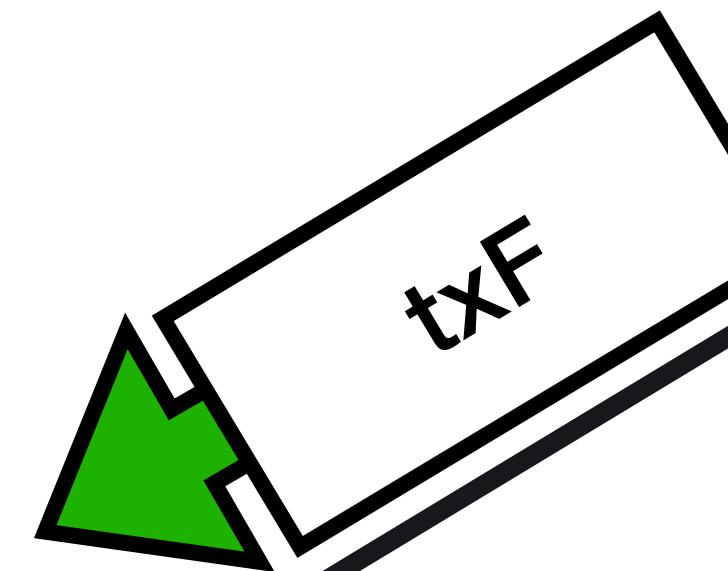
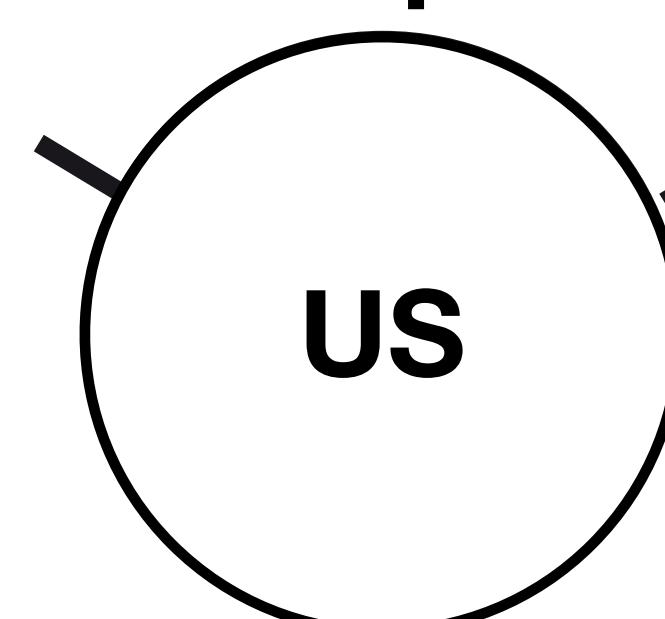
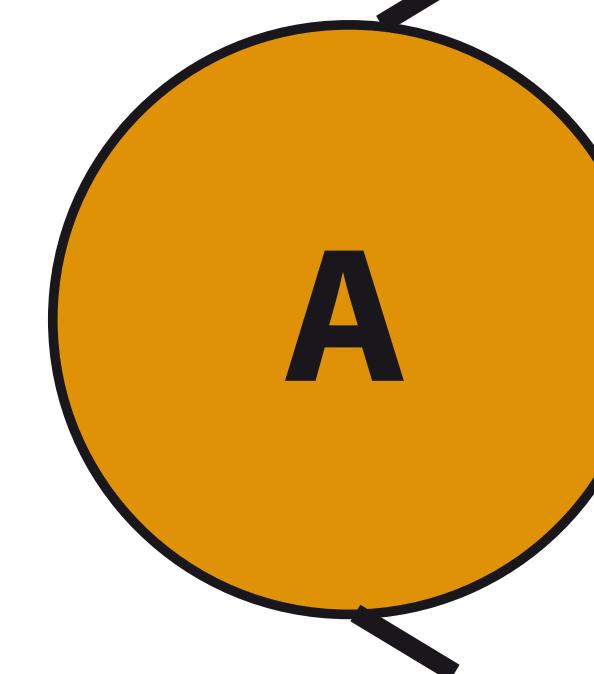
txP



SIMPLIFIED TXPROBE



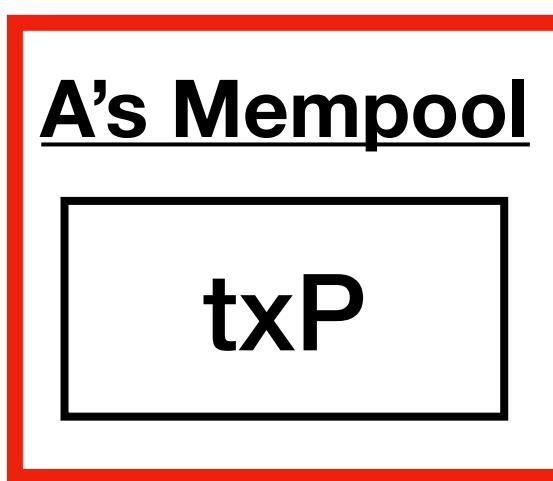
txP



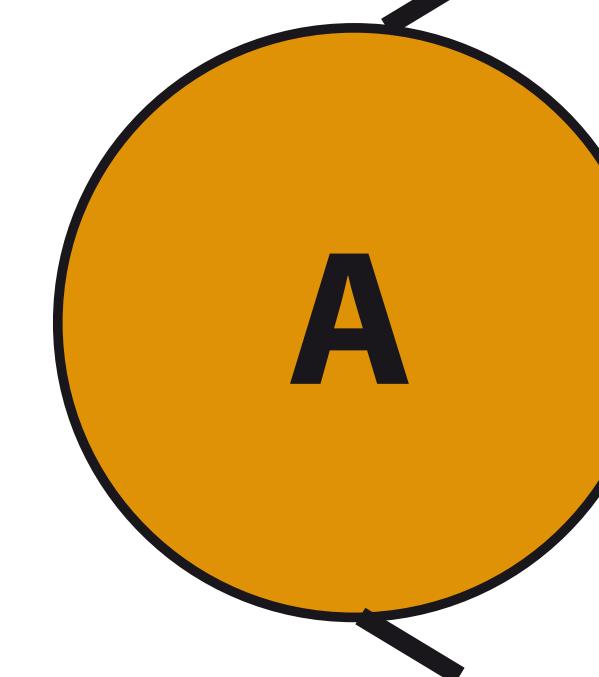
txP

txF

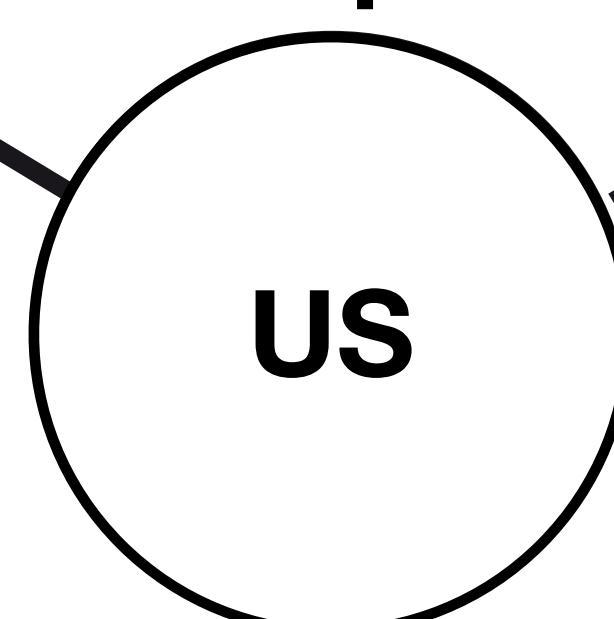
SIMPLIFIED TXPROBE



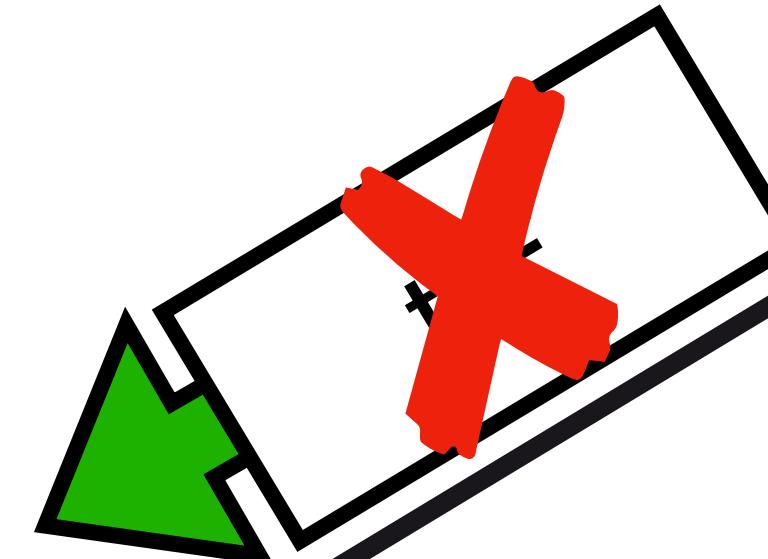
txP



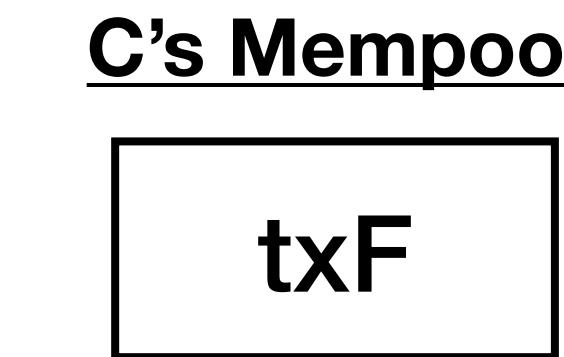
A



US



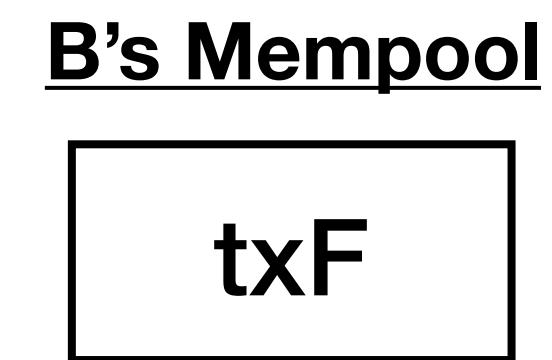
txP



txF



txP



txF



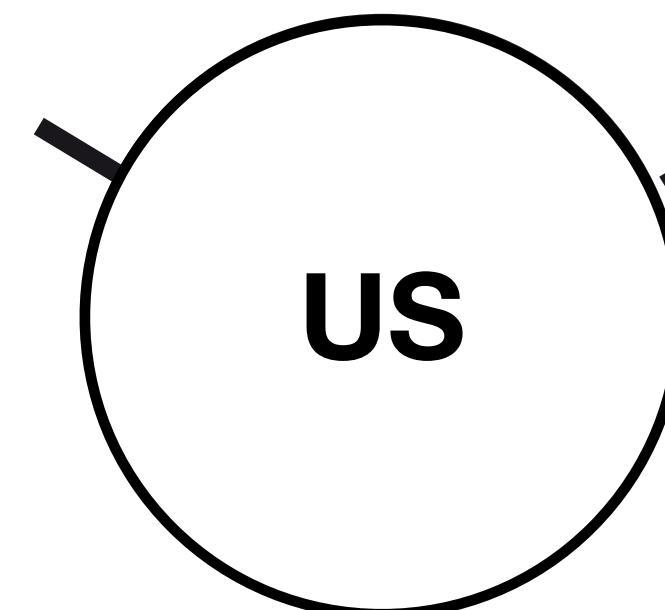
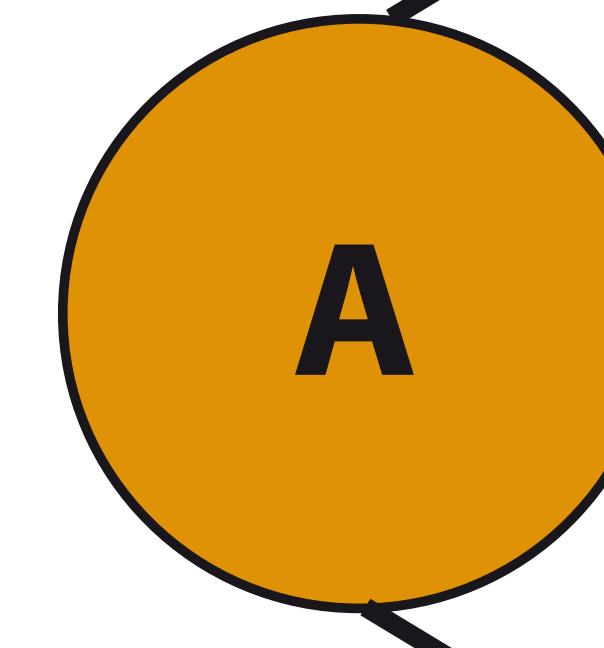
SIMPLIFIED TXPROBE



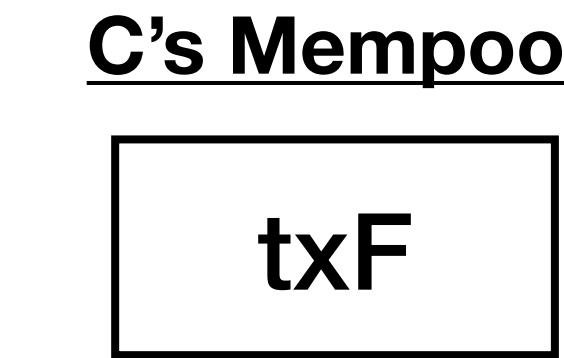
A's Mempool



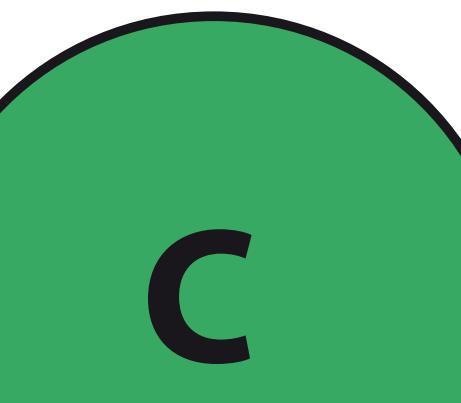
txP



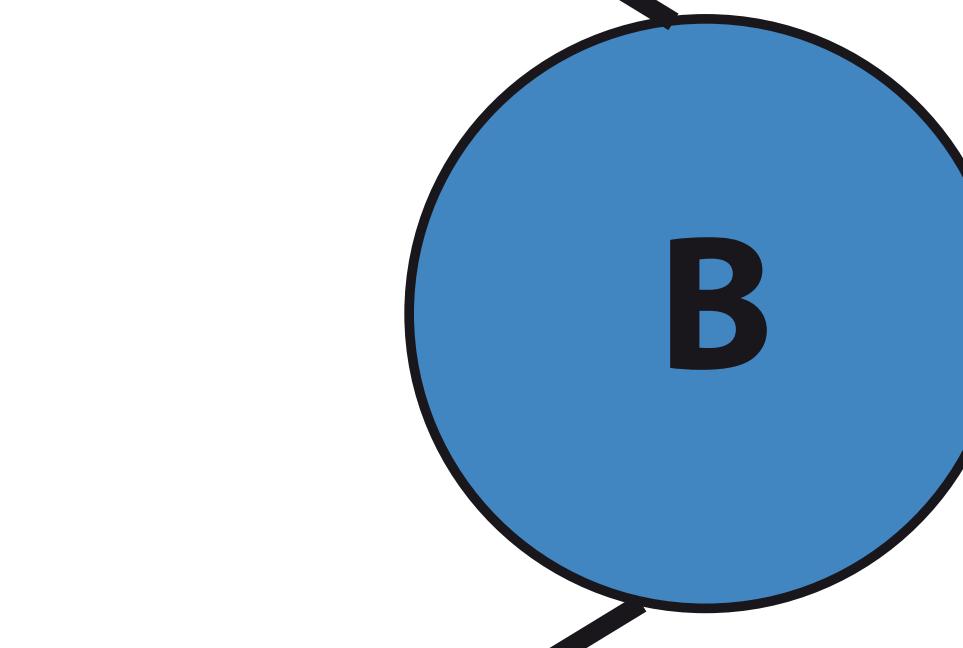
txP



C's Orphanpool
∅



txF



txP



B's Mempool

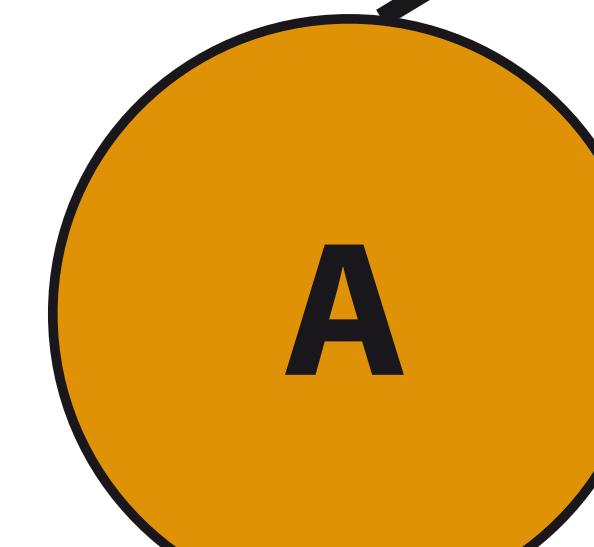
B's Orphanpool
∅

SIMPLIFIED TXPROBE

A's Mempool



txP



txP

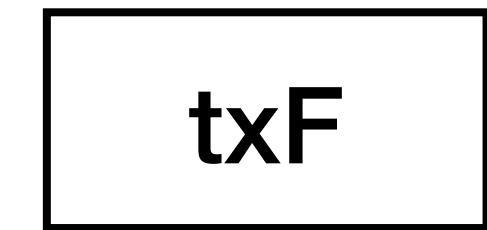
C's Mempool



C's Orphanpool



B's Mempool

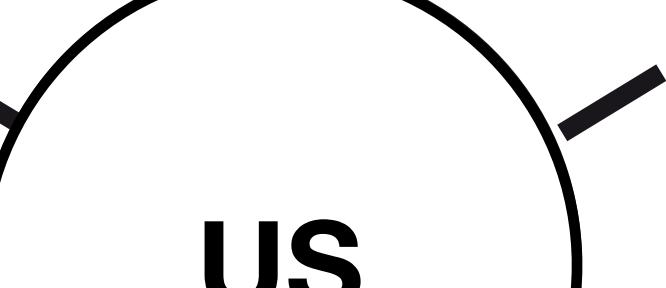
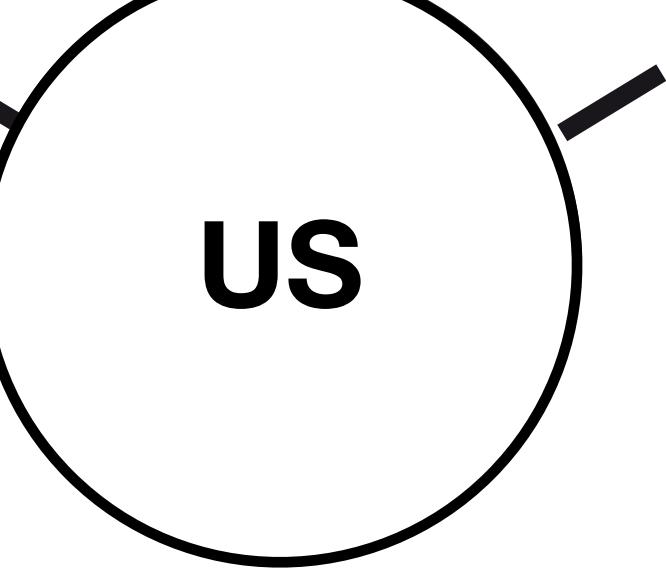
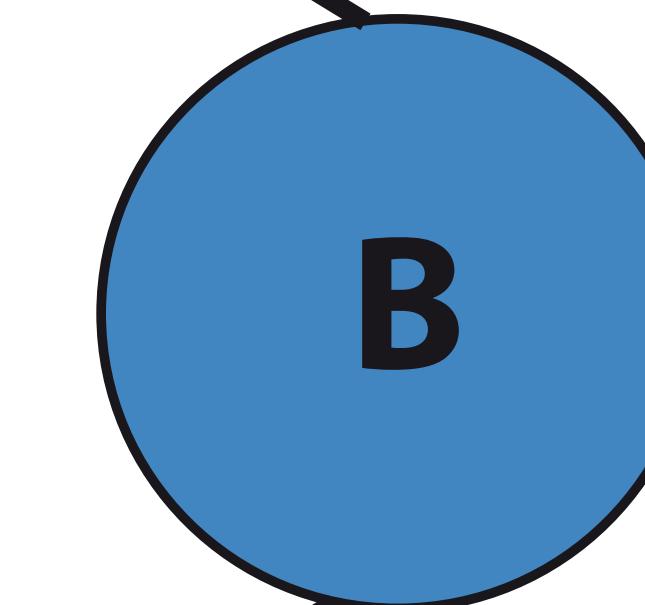
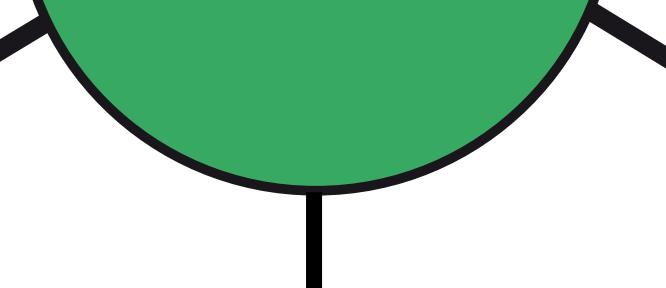
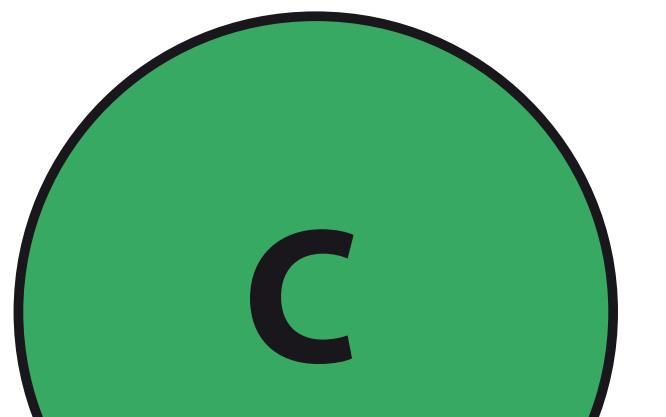
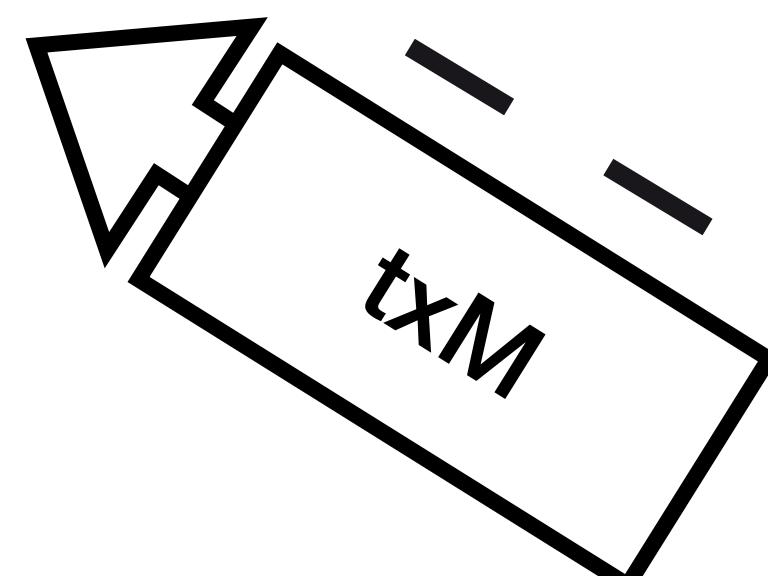


txP

B's Orphanpool



US



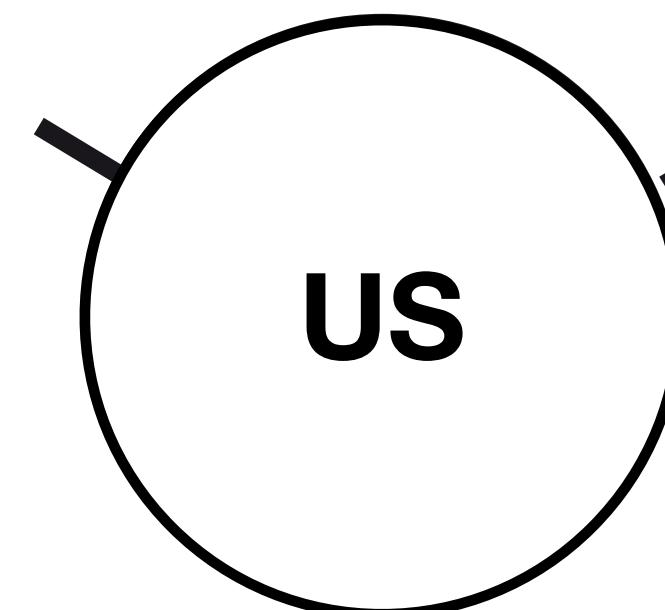
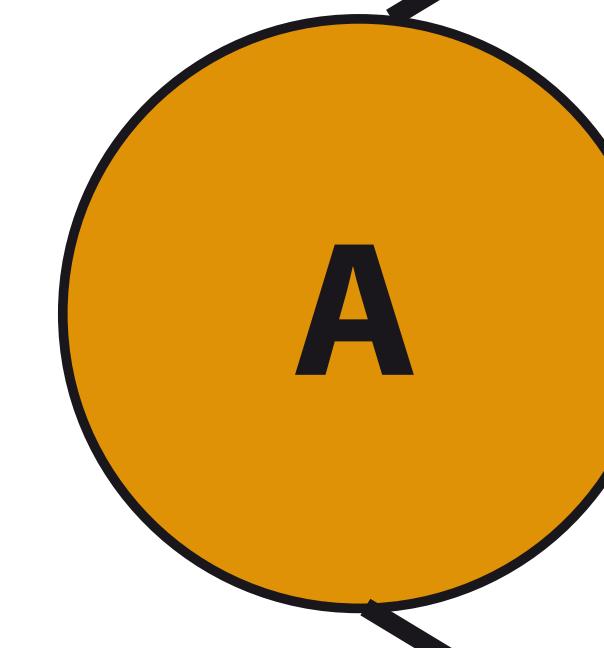
SIMPLIFIED TXPROBE



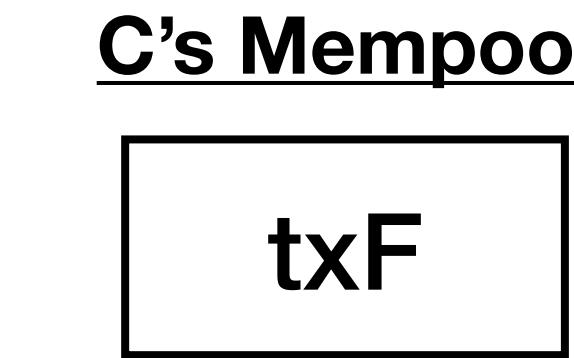
A's Mempool



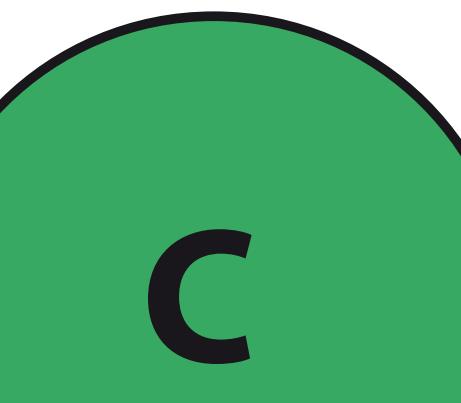
txP



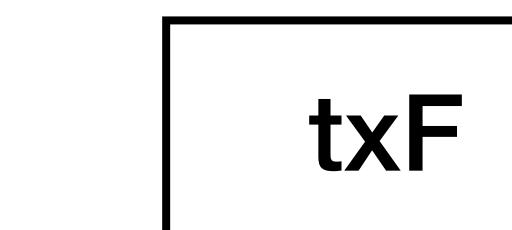
txP



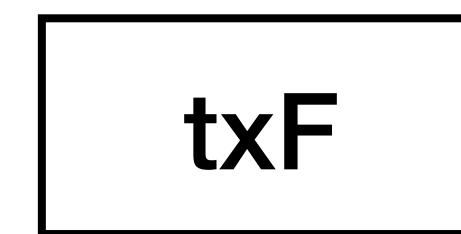
C's Orphanpool
∅



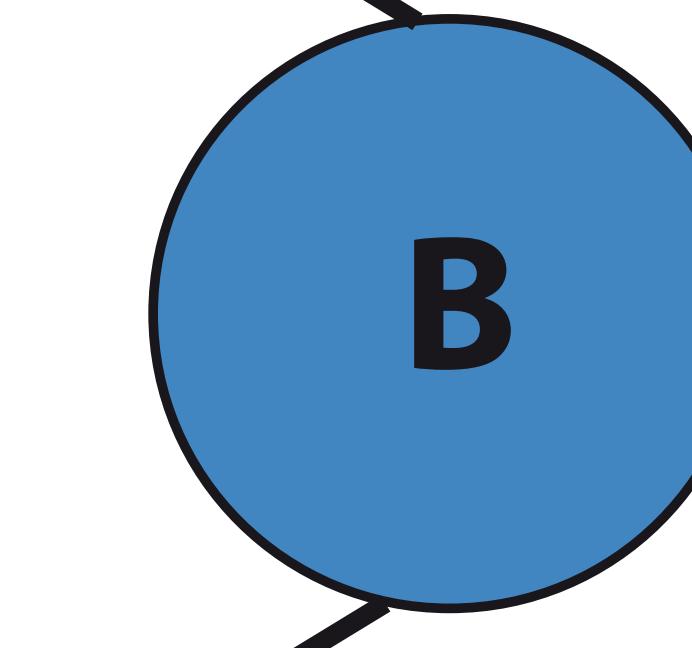
txF



B's Mempool



txP



B's Orphanpool
∅

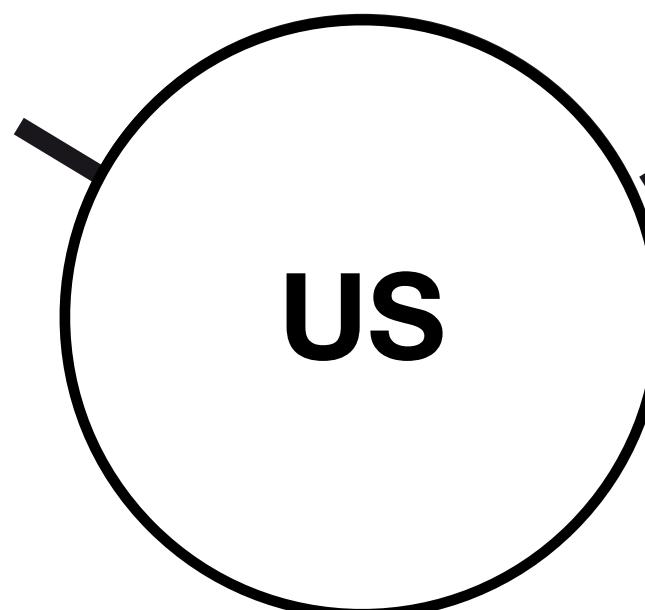
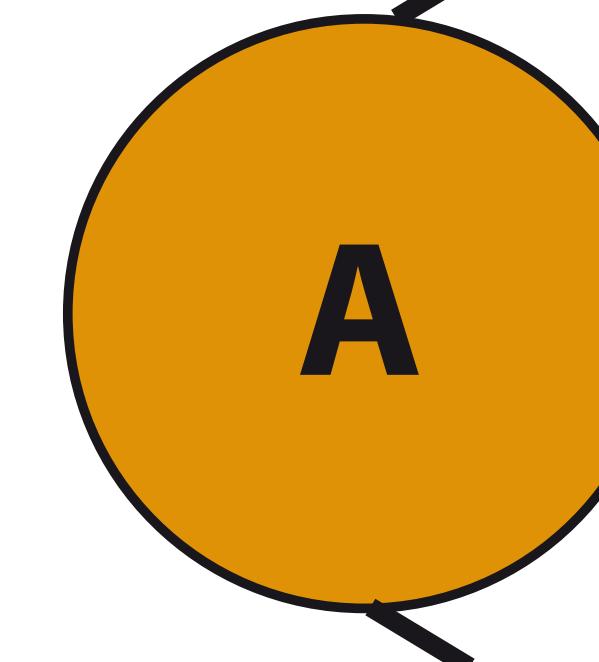
SIMPLIFIED TXPROBE



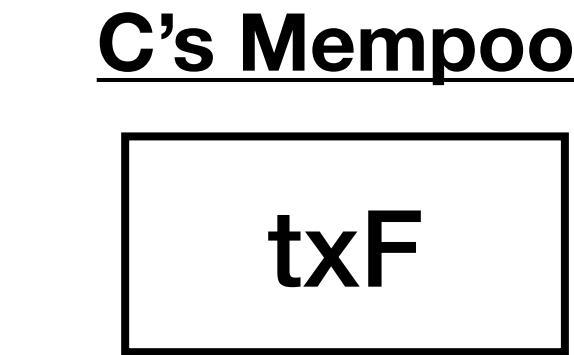
A's Mempool

txP
txM

txP



txP



C's Orphanpool
∅



txF

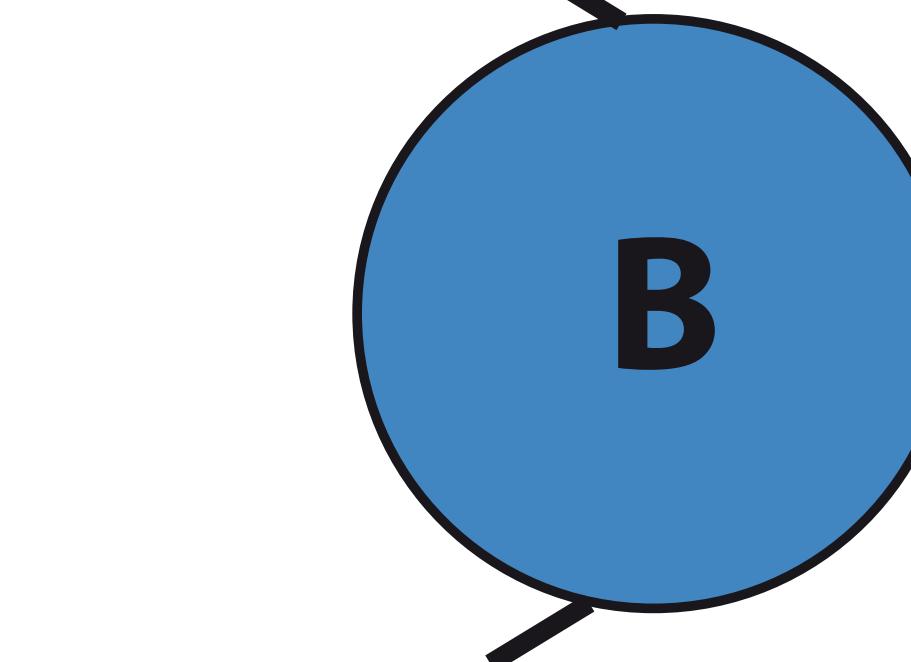
txP

B's Mempool

txF

B's Orphanpool

∅

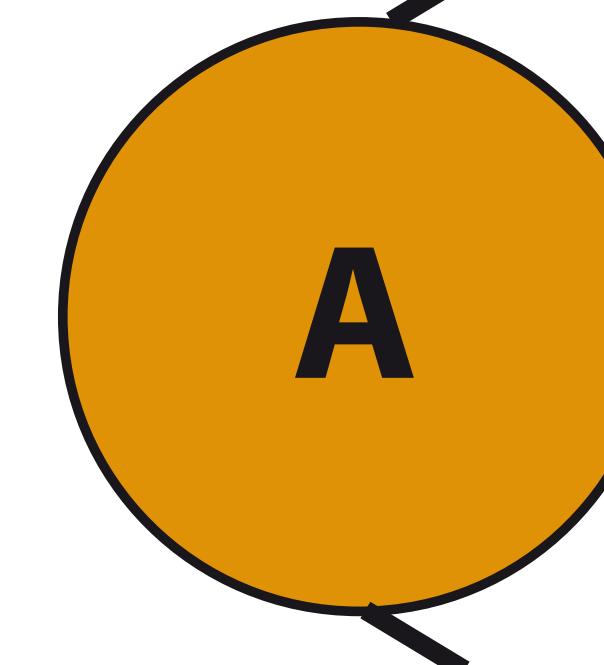


SIMPLIFIED TXPROBE

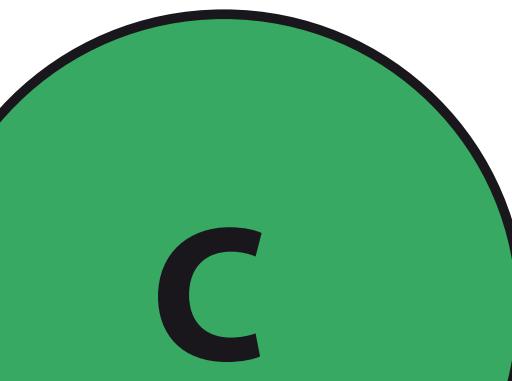
A's Mempool

txP
txM

txP



txP



C's Mempool

txF

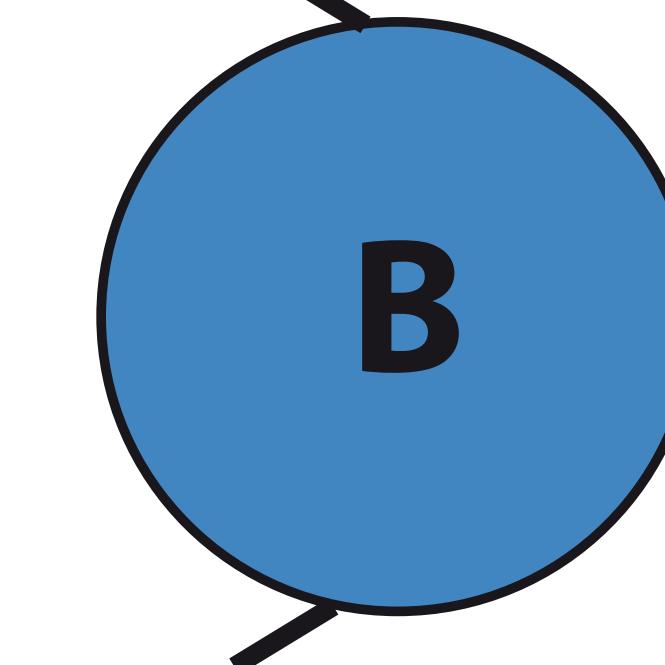
C's Orphanpool

∅

B's Mempool

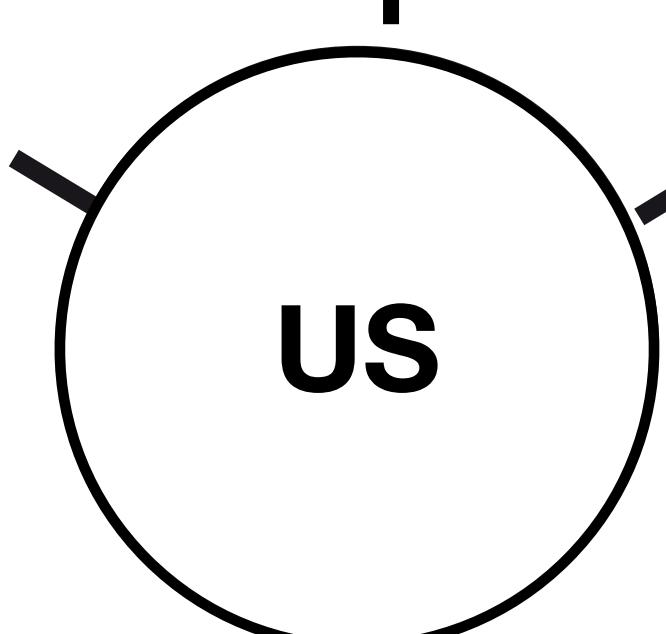
txP
txF

txP



B's Orphanpool

∅



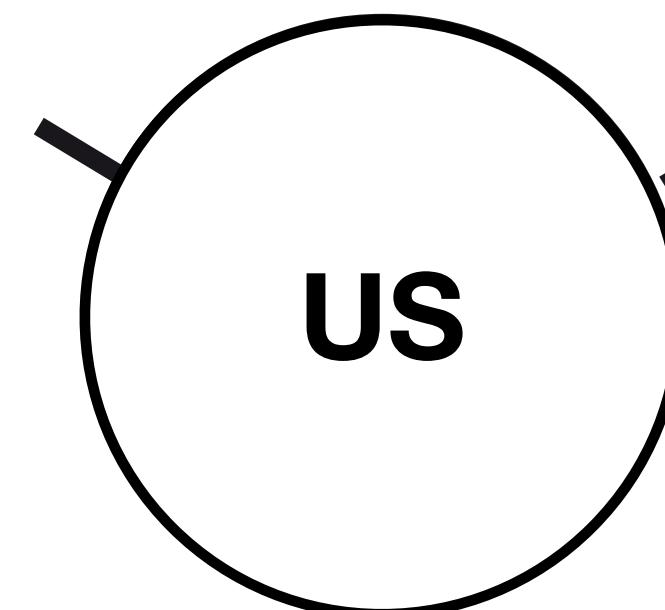
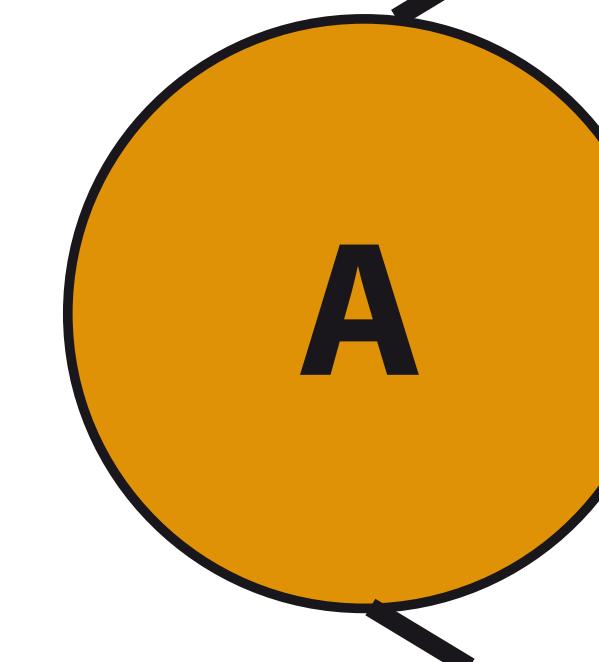
SIMPLIFIED TXPROBE



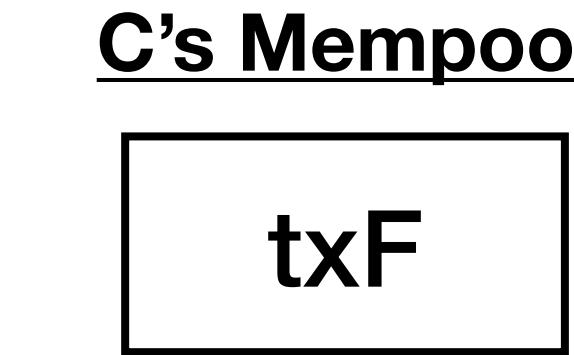
A's Mempool

txP
txM

txP



txP



C's Orphanpool
∅

txP

B's Mempool

txF

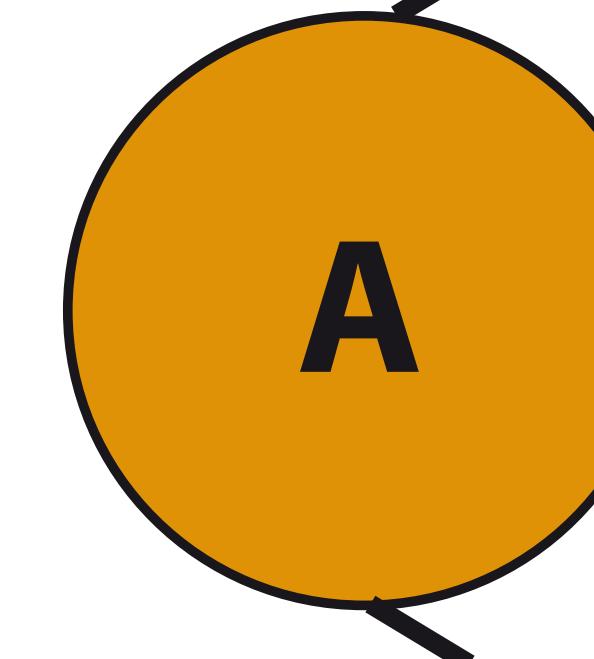
B's Orphanpool
∅

SIMPLIFIED TXPROBE

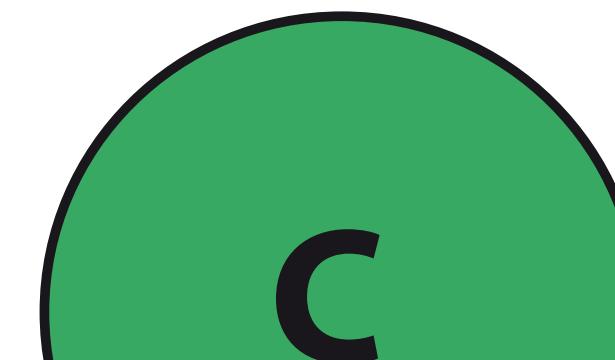
A's Mempool

txP
txM

txP



txP



C's Mempool

txF

C's Orphanpool

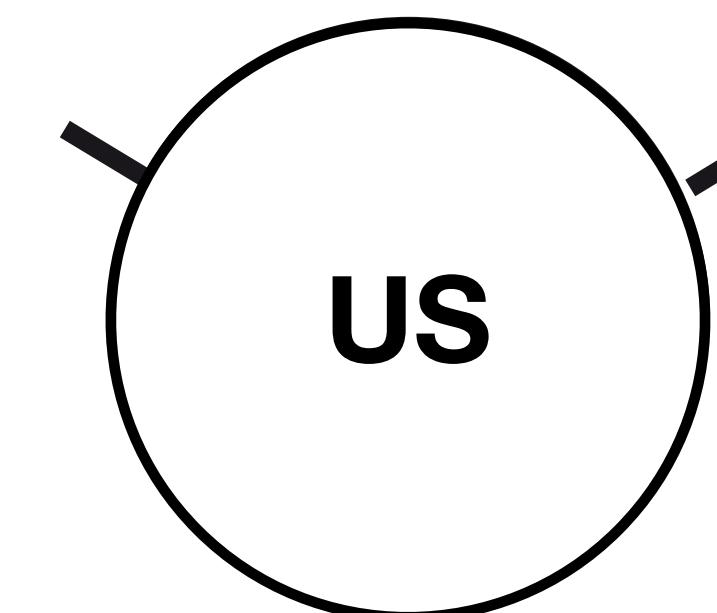
txM

B's Mempool

txP

∅

B's Orphanpool

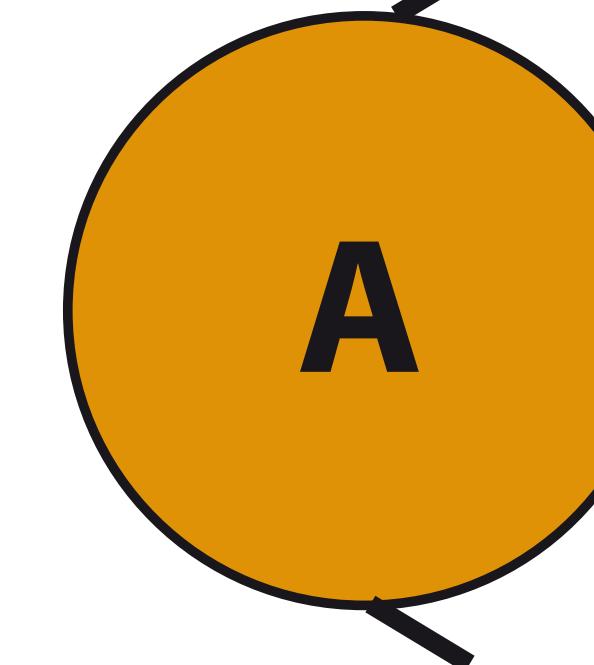


SIMPLIFIED TXPROBE

A's Mempool

txP
txM

txP



txP



C's Mempool

txF

C's Orphanpool

txM

B's Mempool

txP
txF

B's Orphanpool

∅

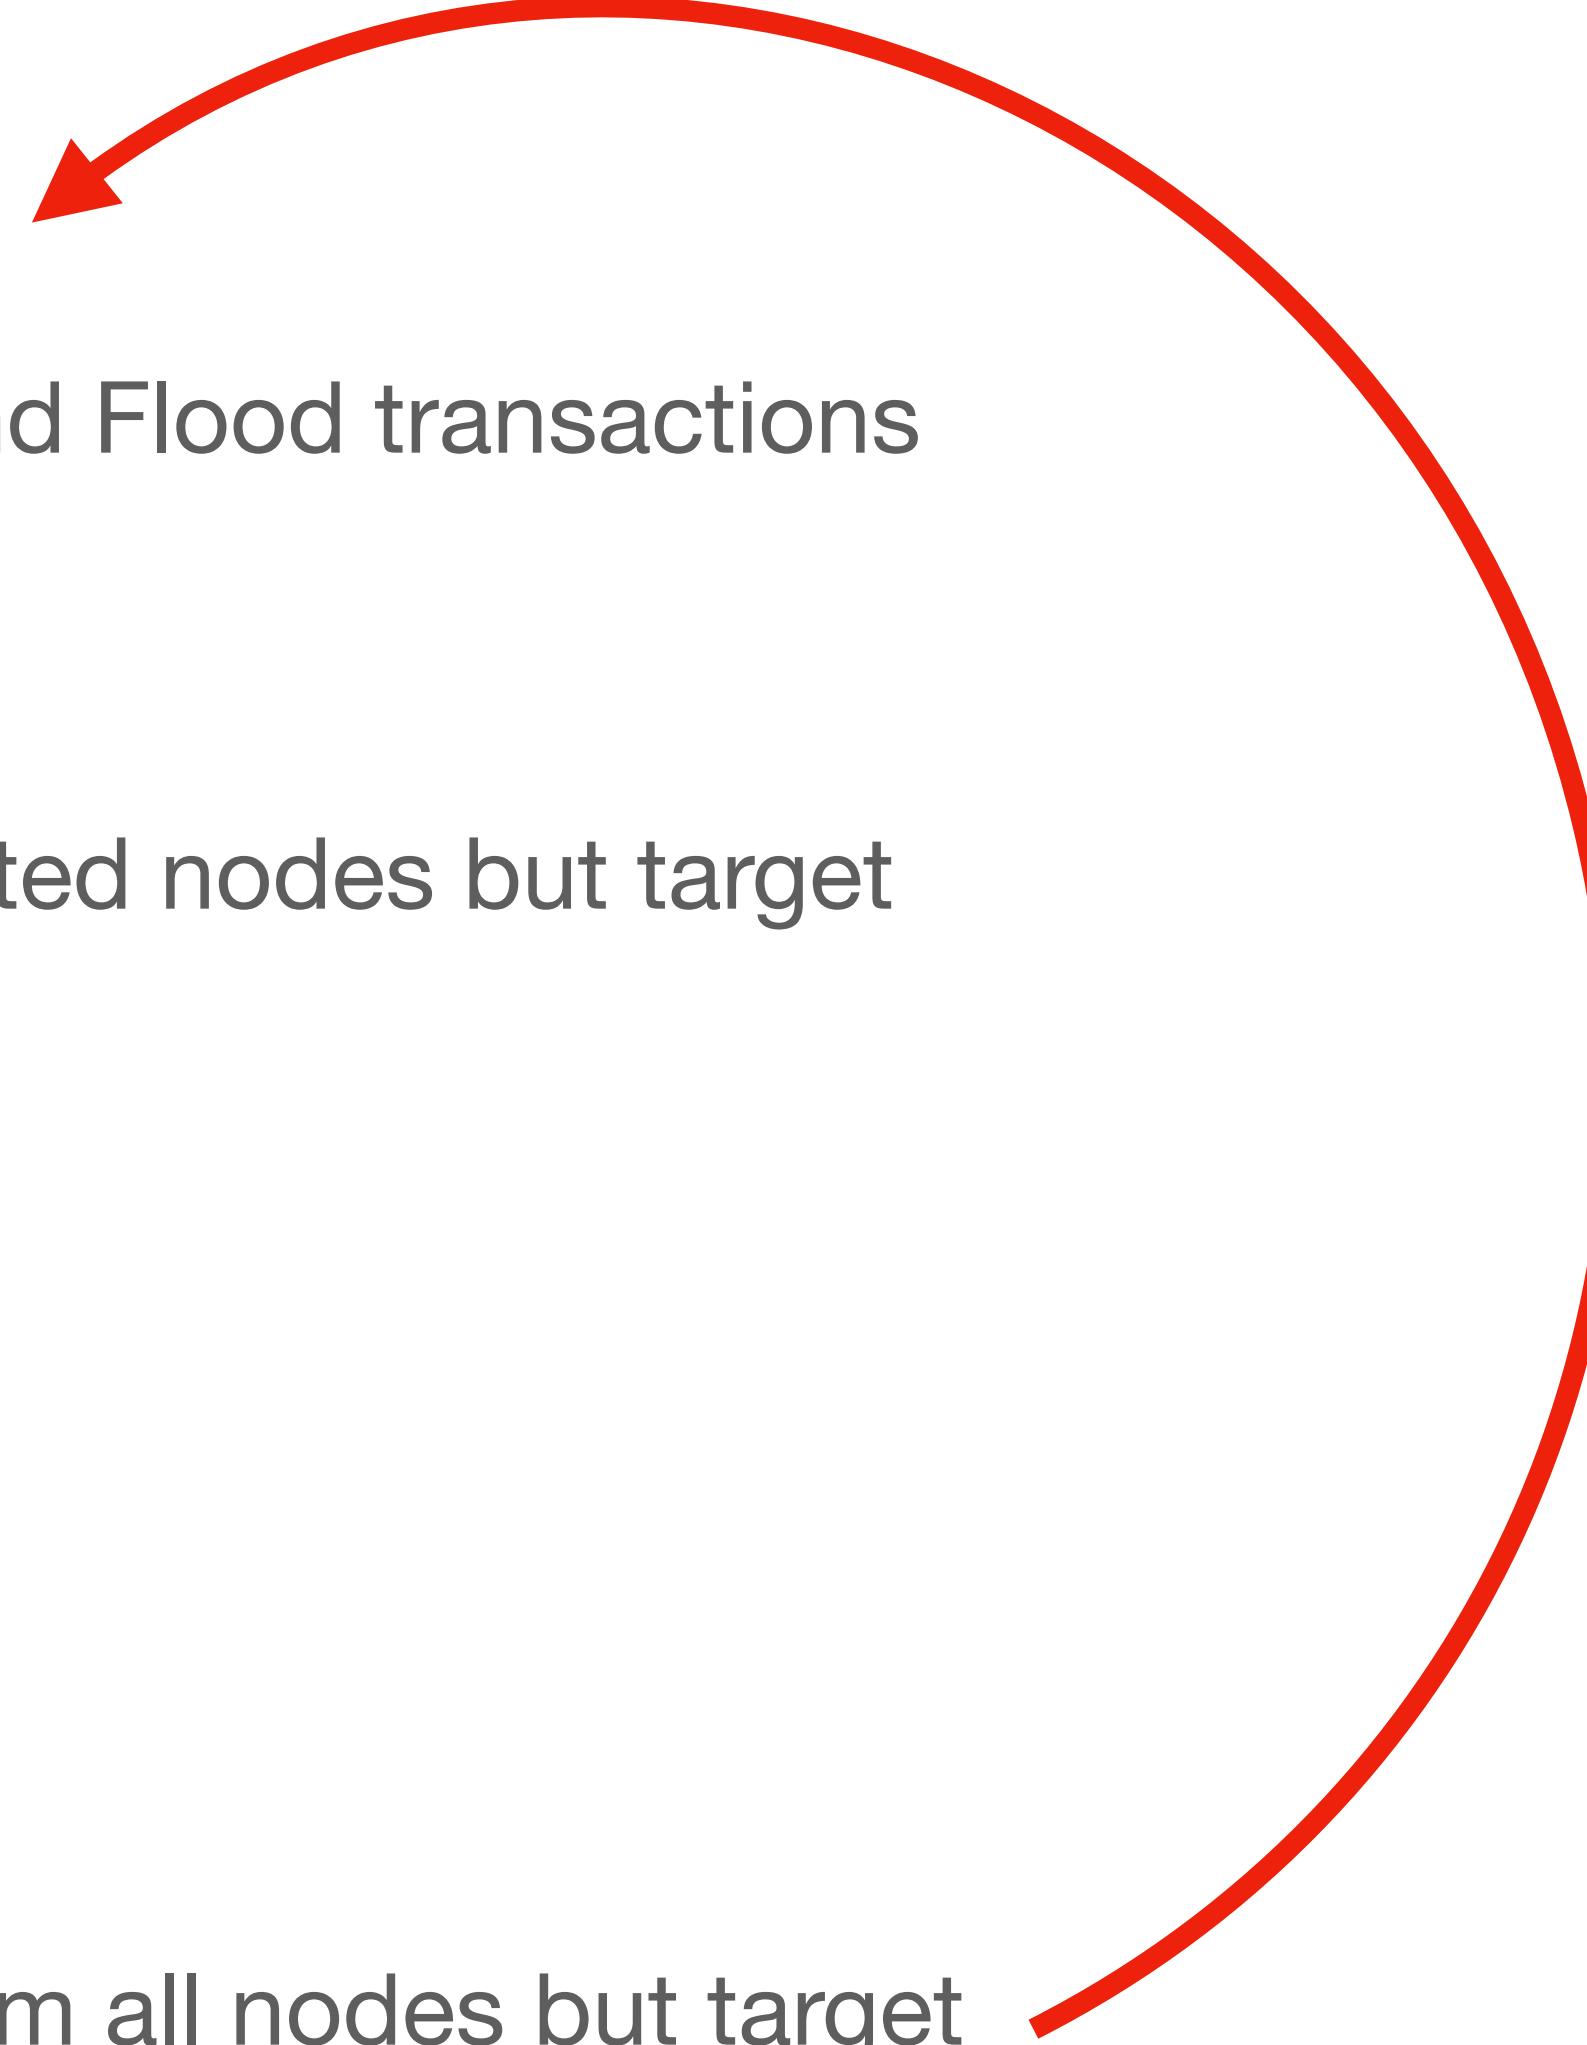
US

TXPROBE - PROTOCOL OVERVIEW

- **Choose** a target node
- **Create** Parent, Marker and Flood transactions
- **INVBLOCK** the network
- **Send** Flood to all connected nodes but target
- **Let** Flood propagate
- **Send** Parent to target
- **Send** Marker to target
- **Let** Marker propagate
- **Request** marker back from all nodes but target

TXPROBE - PROTOCOL OVERVIEW

- **Choose** a target node
- **Create** Parent, Marker and Flood transactions
- **INVBLOCK** the network
- **Send** Flood to all connected nodes but target
- **Let** Flood propagate
- **Send** Parent to target
- **Send** Marker to target
- **Let** Marker propagate
- **Request** marker back from all nodes but target



For every node in the
network

TXPROBE - COSTS ESTIMATION

For a network like **Bitcoin mainnet**:

nodes \approx 10000

time \approx 8.25 hours

cost = 573210-764280 satoshi (5 sat/byte) \approx **\$ $(20-30)$**

TXPROBE - DATA VALIDATION (TESTNET)

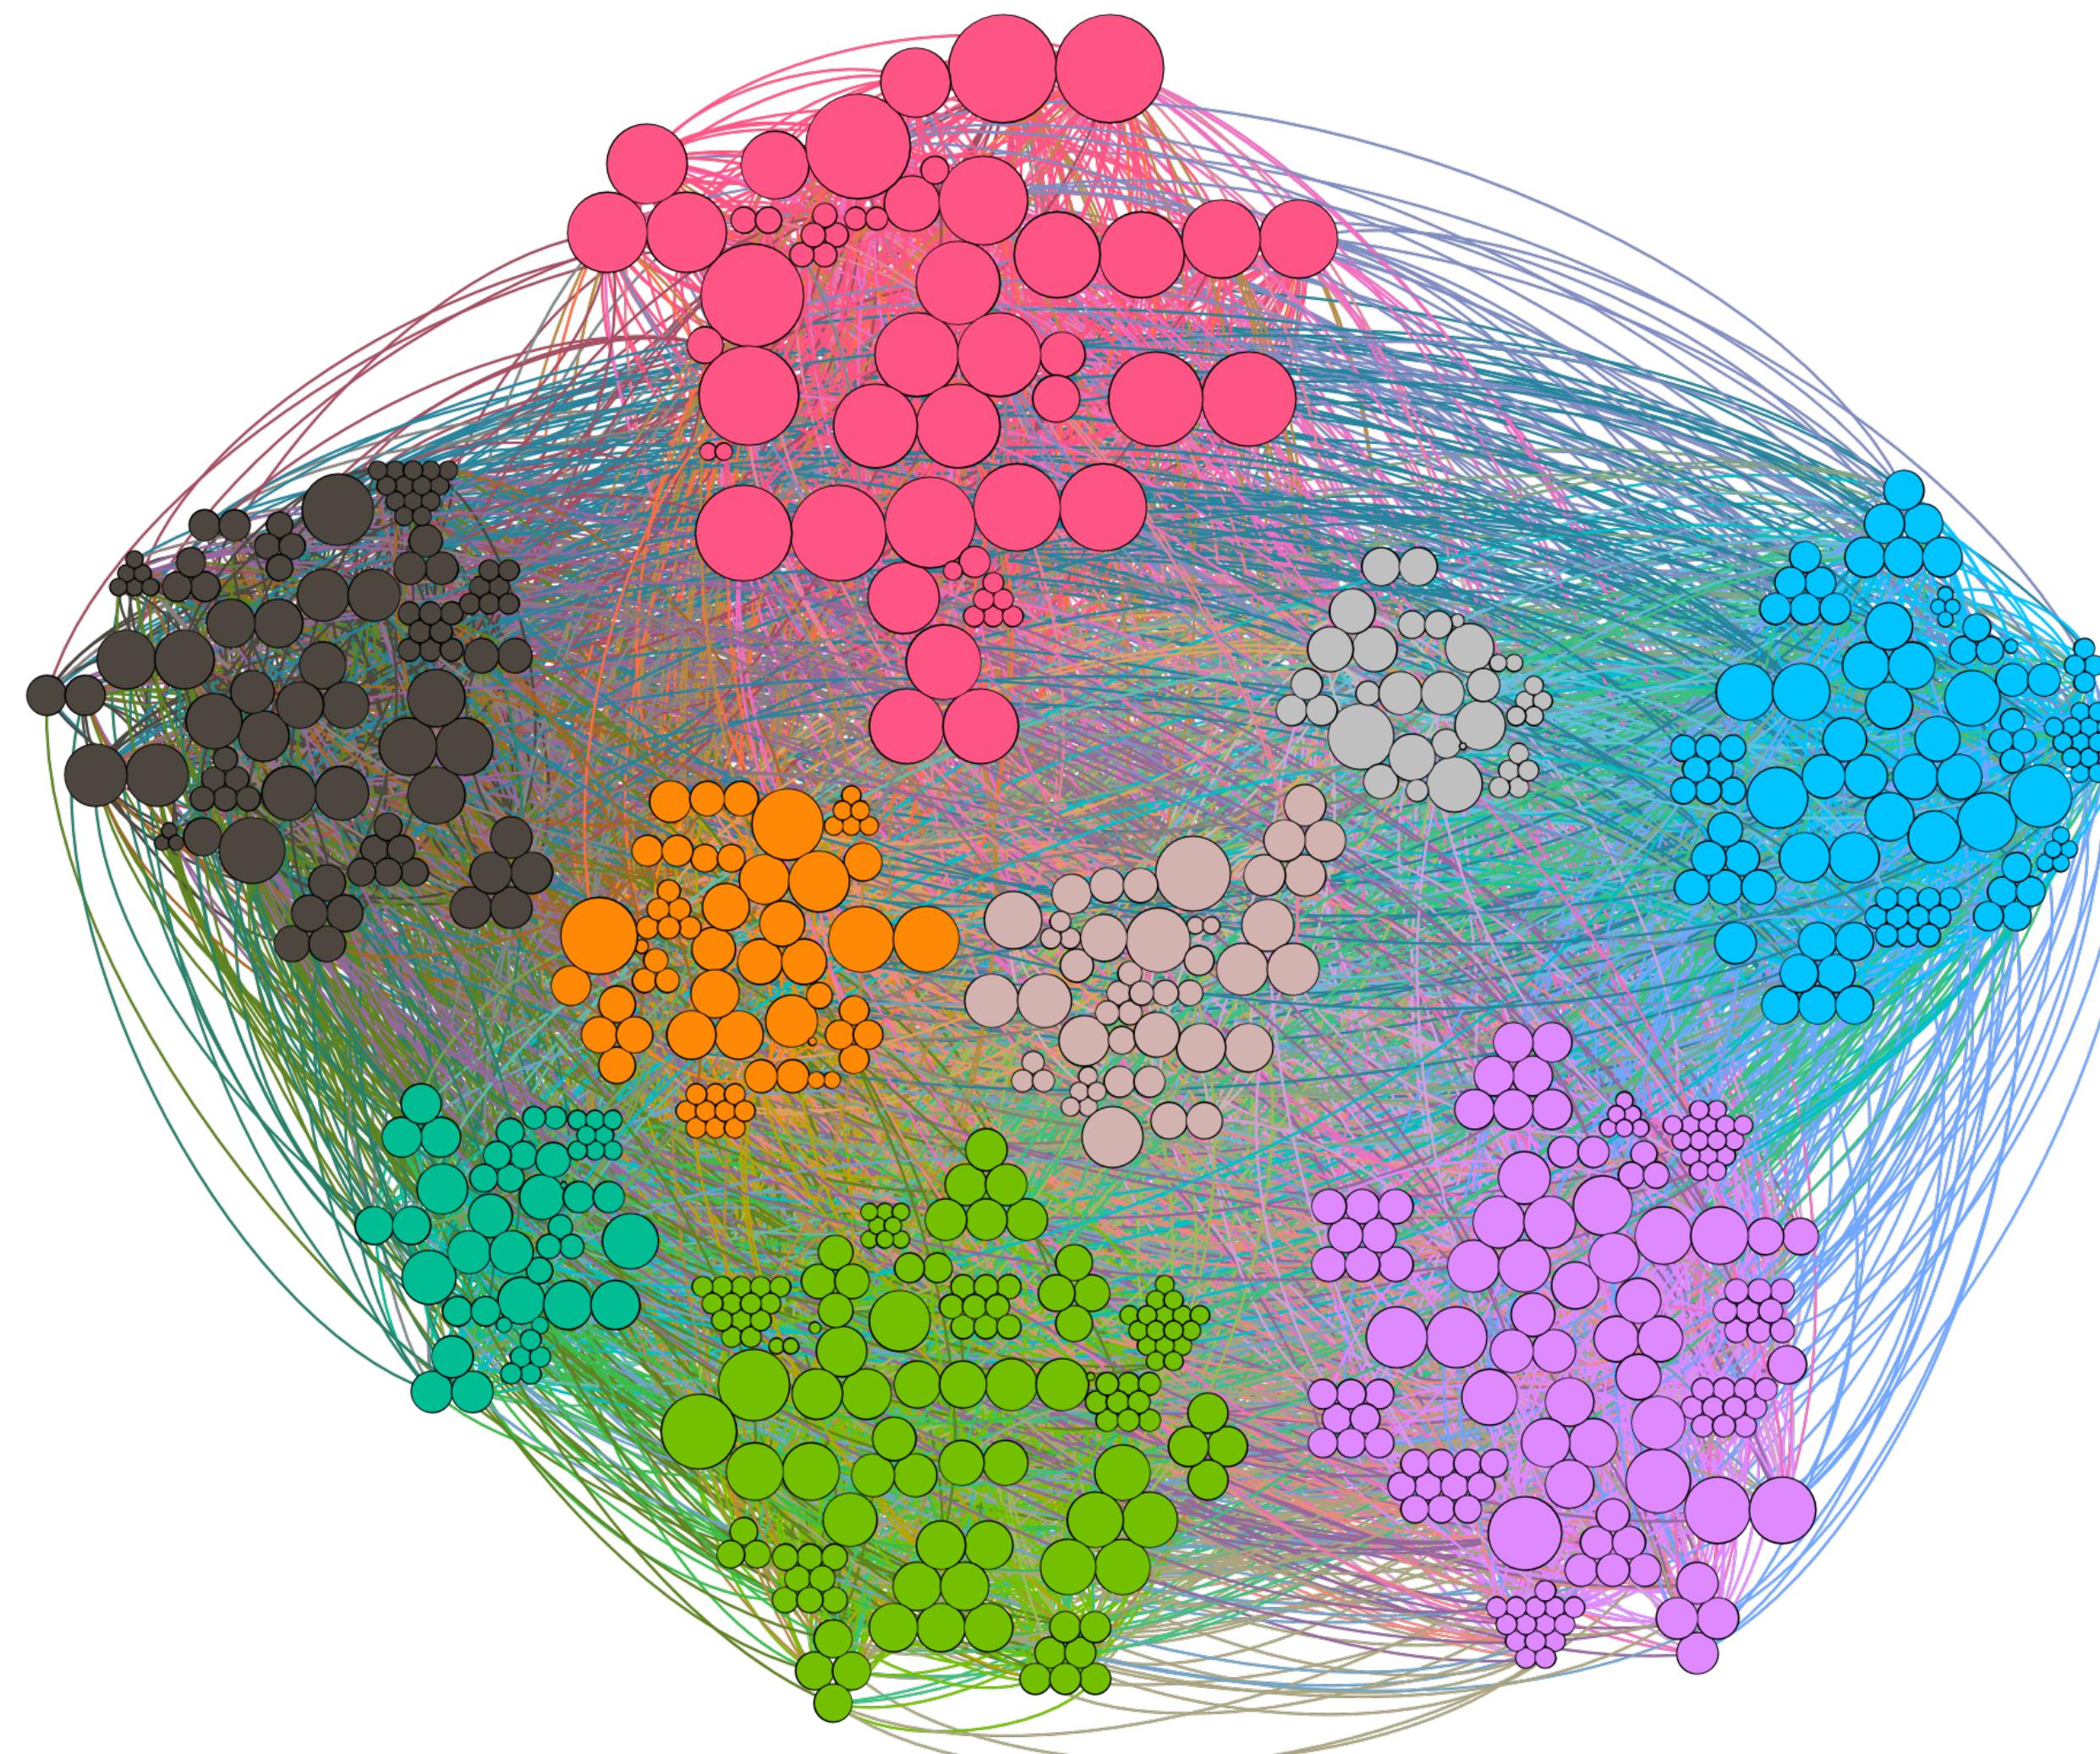
We run 5 Bitcoin Core nodes as **ground truth**

We define our **precision / recall** by checking how well can we infer the ground truth nodes connections

Over 40 trials and with 95% confidence:

- **Precision = 100%**
- **Recall = 93.86% - 95.45%**

TXPROBE - TESTNET TOPOLOGY



precision = 100%

recall = 97.40%

size → degree

color → Community
unfolding

Higher **community structure**
and **modularity** than random
graph

CONCLUSIONS

CONCLUSIONS

Select orphan transaction uniformly for eviction #14626

 Merged MarcoFalke merged 1 commit into [bitcoin:master](#) from [sipa:201810_uniform_orphan_eviction](#) 3 days ago

 Conversation 20  Commits 1  Checks 0  Files changed 1

 **sipa** commented on 31 Oct 2018 Member + ...

The previous code was biased towards evicting transactions whose txid has a larger gap (lexicographically) with the previous txid in the orphan pool.

CONCLUSIONS

Select orphan transaction uniformly for eviction #14626

 Merged MarcoFalke merged 1 commit into [bitcoin:master](#) from [sipa:201810_uniform_orphan_eviction](#) 3 days ago

 Conversation 20  Commits 1  Checks 0  Files changed 1

 **sipa** commented on 31 Oct 2018 Member + ...

The previous code was biased towards evicting transactions whose txid has a larger gap (lexicographically) with the previous txid in the orphan pool.

randomize GETDATA(tx) request order and introduce bias toward outbound #14897

 Merged **sipa** merged 1 commit into [bitcoin:master](#) from [naumenkogs:master](#) 10 days ago

 Conversation 115  Commits 1  Checks 0  Files changed 6

 **naumenkogs** commented on 8 Dec 2018 • edited by MarcoFalke Contributor + ...

This code makes executing two particular (and potentially other) attacks harder.

InvBlock

CONCLUSIONS

Select orphan transaction uniformly for eviction #14626

Merged MarcoFalke merged 1 commit into bitcoin:master from sipa:201810_uniform_orphan_eviction 3 days ago

Conversation 20 Commits 1 Checks 0 Files changed 1

 sipa commented on 31 Oct 2018 Member + ...

The previous code was biased towards evicting transactions whose txid has a larger gap (lexicographically) with the previous txid in the orphan pool.

randomize GETDATA(tx) request order and introduce bias toward outbound #14897

Merged sipa merged 1 commit into bitcoin:master from naumenkogs:master 10 days ago

Conversation 115 Commits 1 Checks 0 Files changed 6

 naumenkogs commented on 8 Dec 2018 • edited by MarcoFalke Contributor + ...

This code makes executing two particular (and potentially other) attacks harder.

InvBlock



CONCLUSIONS

Select orphan transaction uniformly for eviction #14626

Merged MarcoFalke merged 1 commit into bitcoin:master from sipa:201810_uniform_orphan_eviction 3 days ago

Conversation 20 Commits 1 Checks 0 Files changed 1

sipa commented on 31 Oct 2018 Member + ...

The previous commit... (lexicographical order)

randomize G toward outbo...

Merged sipa merged

Conversation 115 Commits 1 Checks 0 Files changed 6

Is topology hiding a design goal, or is it a mean to achieve other goals (e.g: Transaction privacy)?



naumenkogs commented on 8 Dec 2018 • edited by MarcoFalke Contributor + ...

This code makes executing two particular (and potentially other) attacks harder.

InvBlock

QUESTIONS



BONUS TRACK

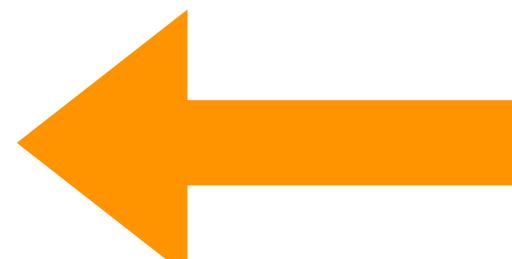
Testnet vs Mainnet

INVBLOCKING (no-link)

Efficiency / Orphan pool eviction

WHY TESTNET AND NO MAINNET?

- TxProbe is rather invasive: it empties the **MapOrphanTransactions pool** of all nodes in the network every round
- We could not measure the implication that such behavior may have had on the **propagation of regular transactions**

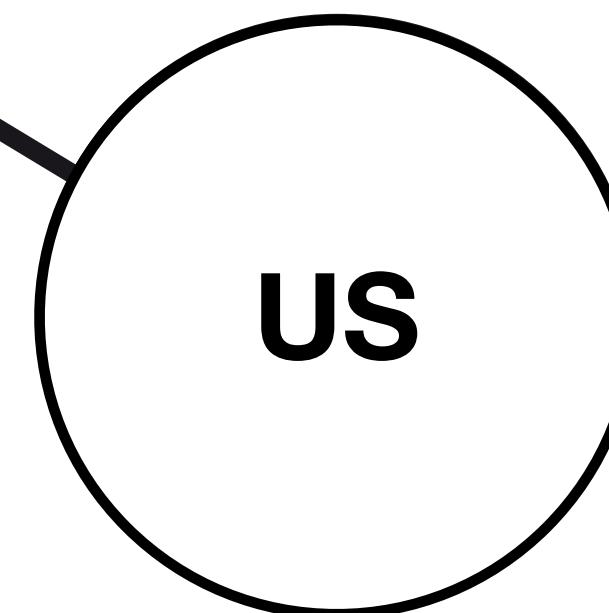
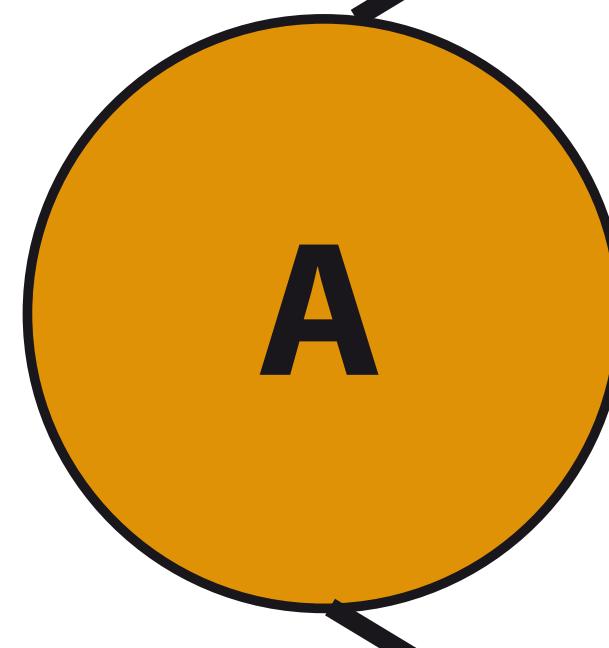


INVBLOCKING V2



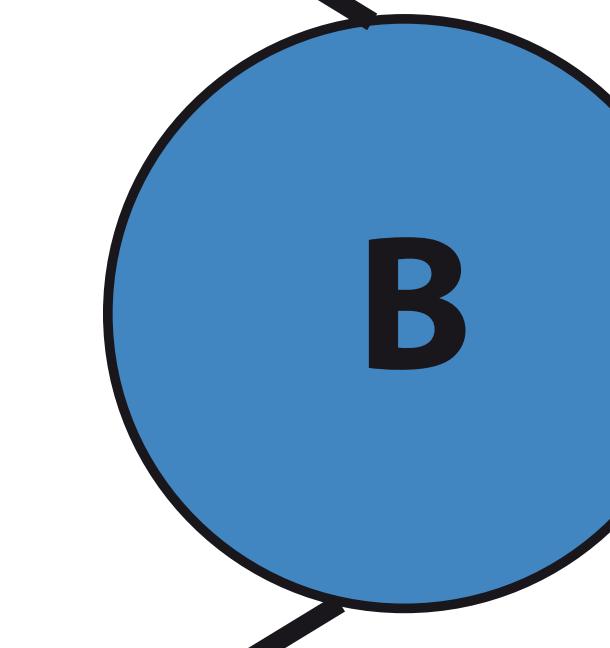
A's Mempool

\emptyset



C's Mempool

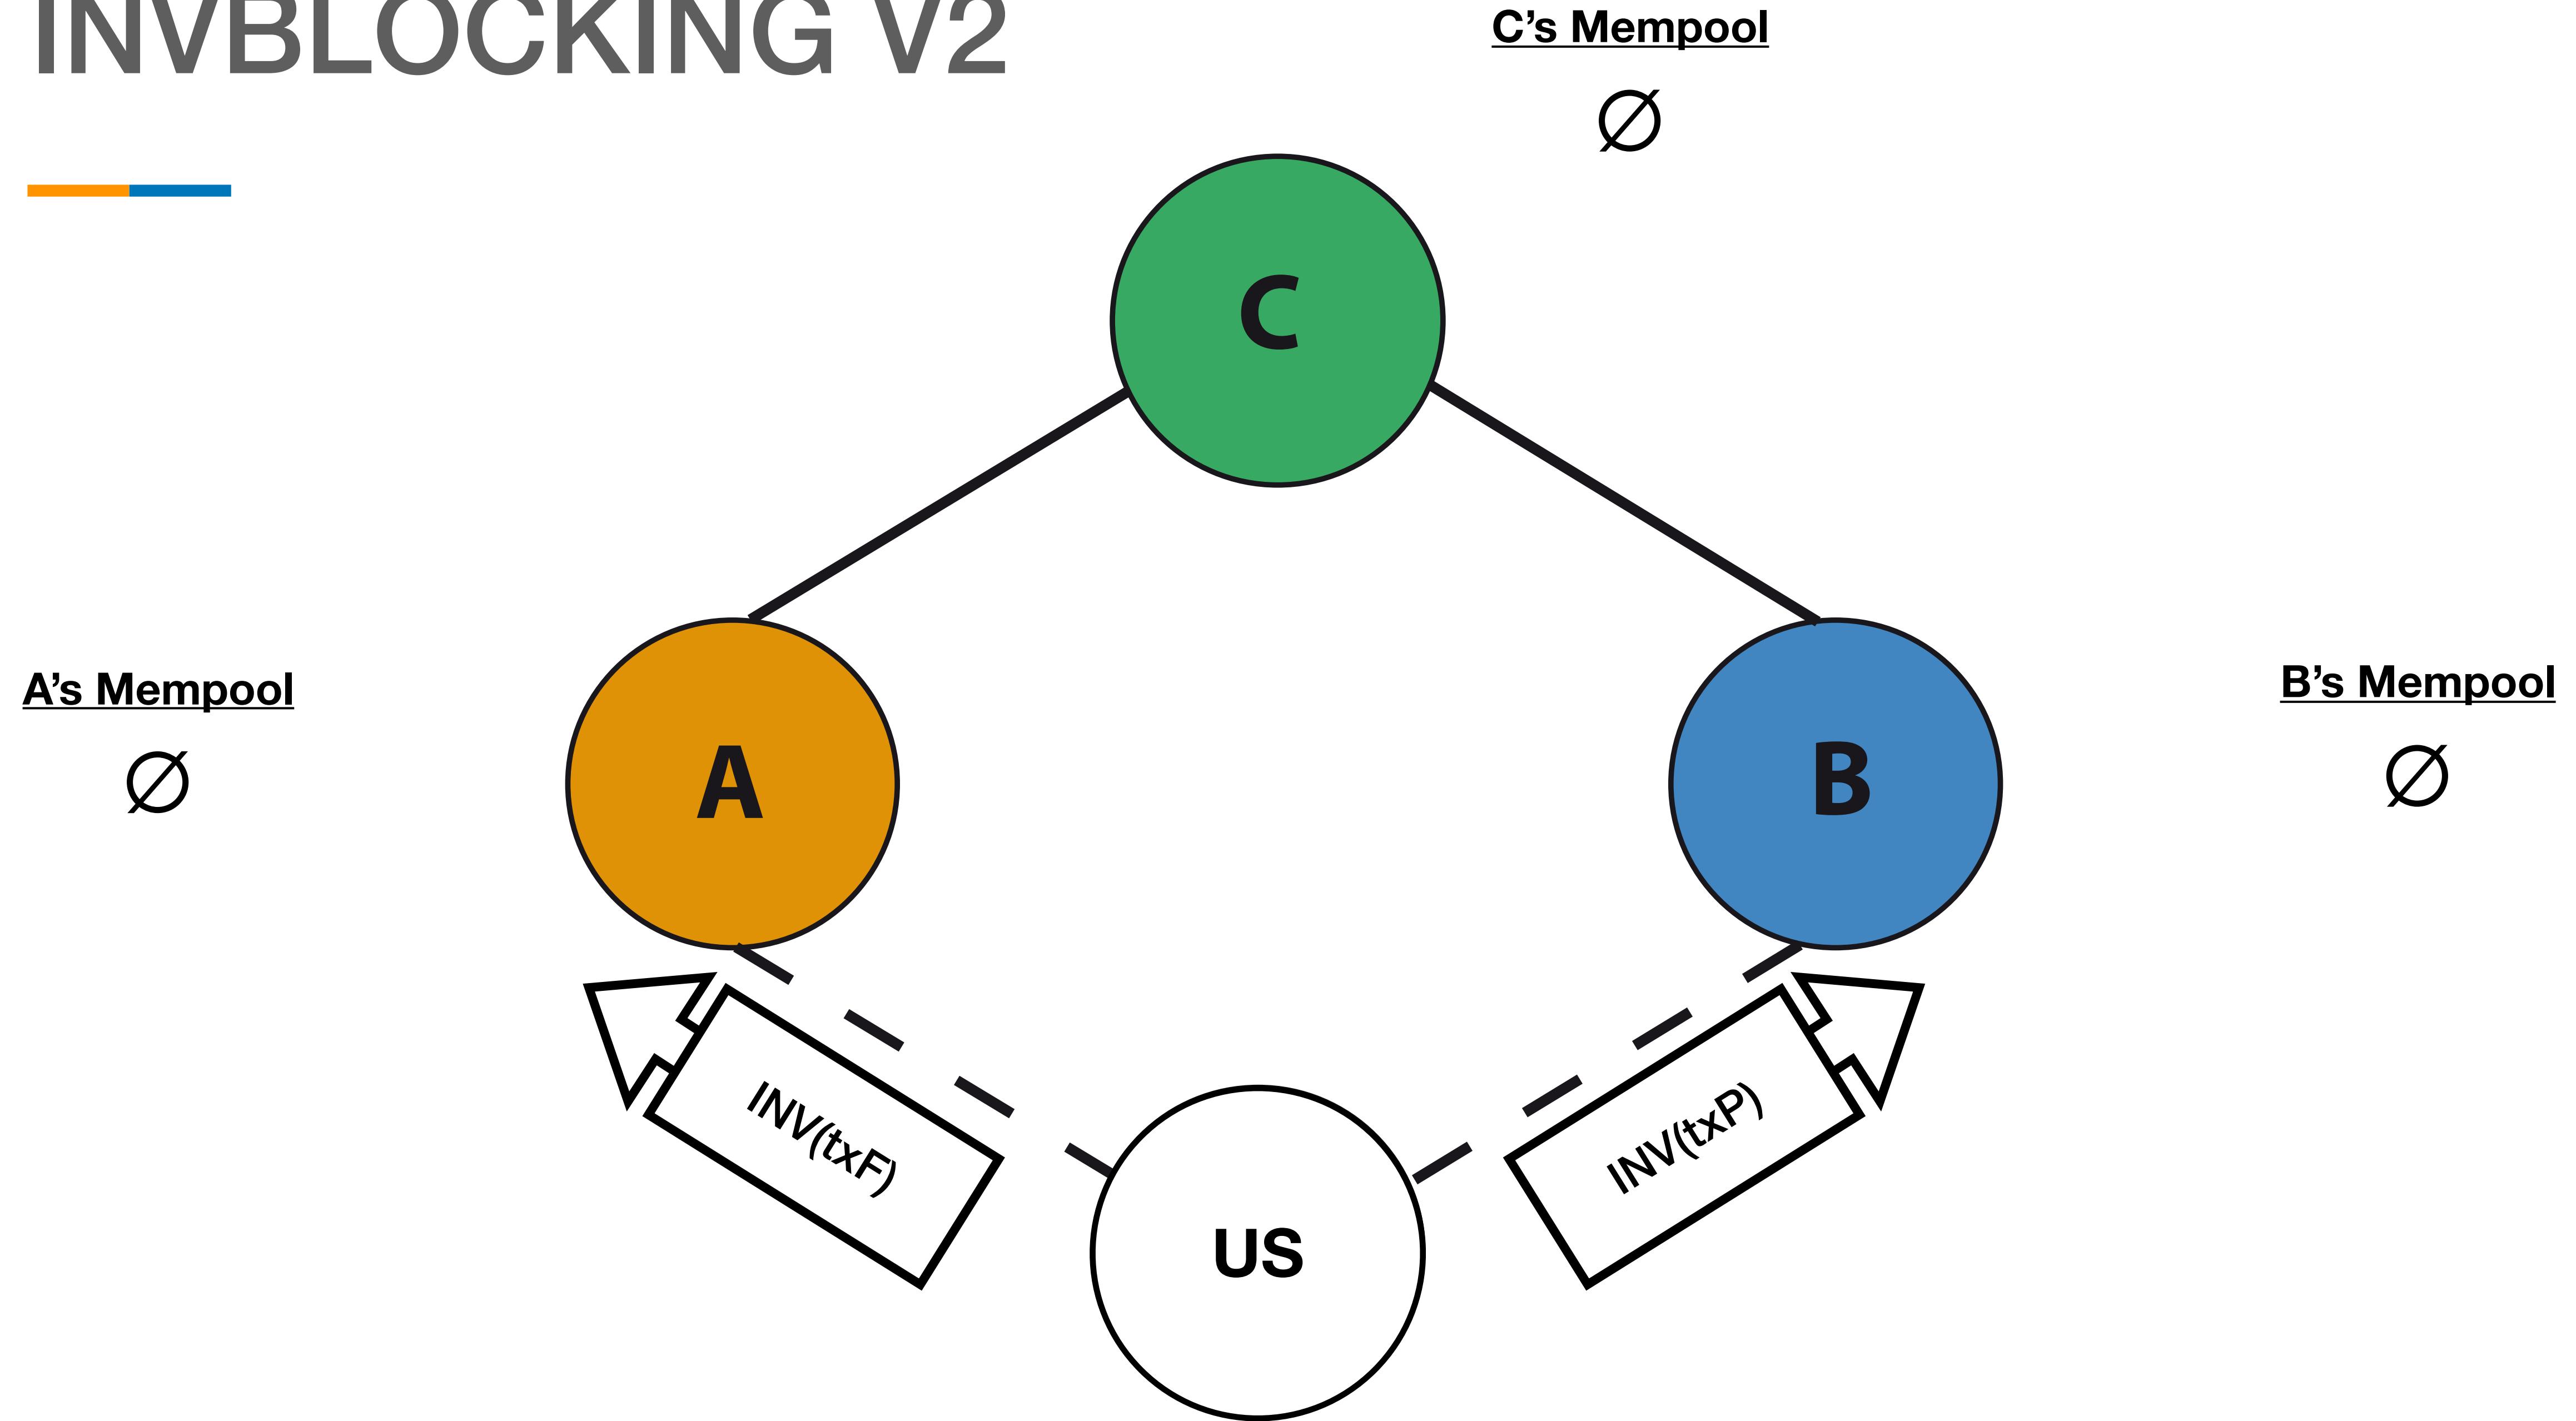
\emptyset



B's Mempool

\emptyset

INVBLOCKING V2

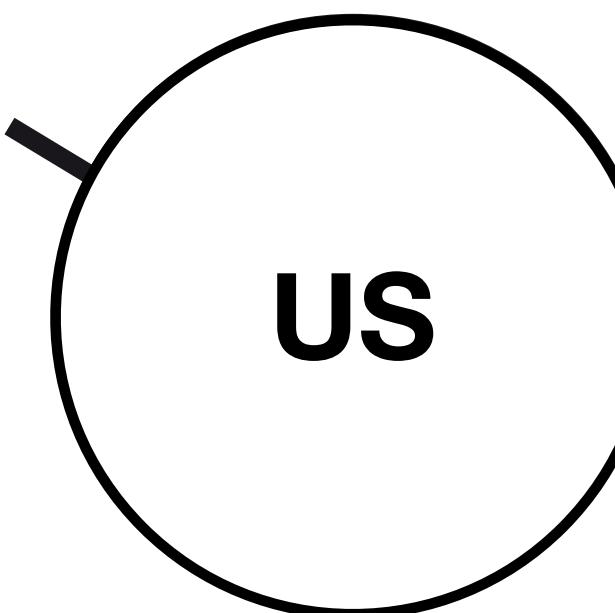
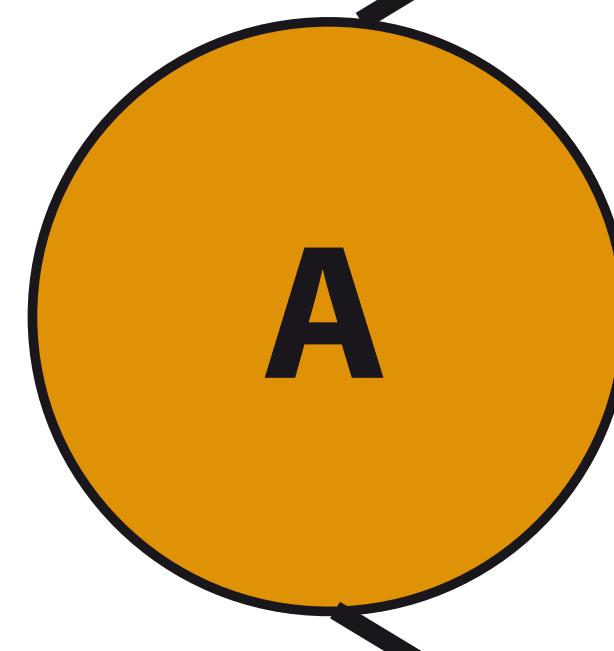


INVBLOCKING V2



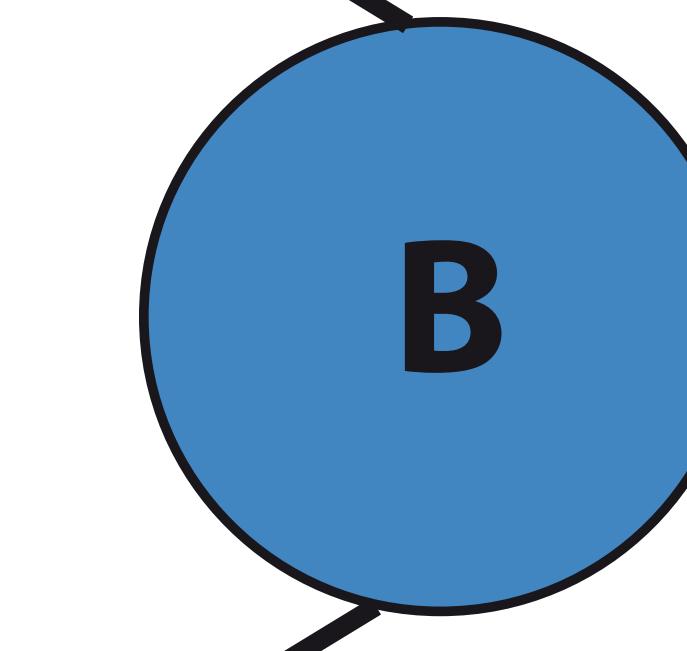
A's Mempool

\emptyset



C's Mempool

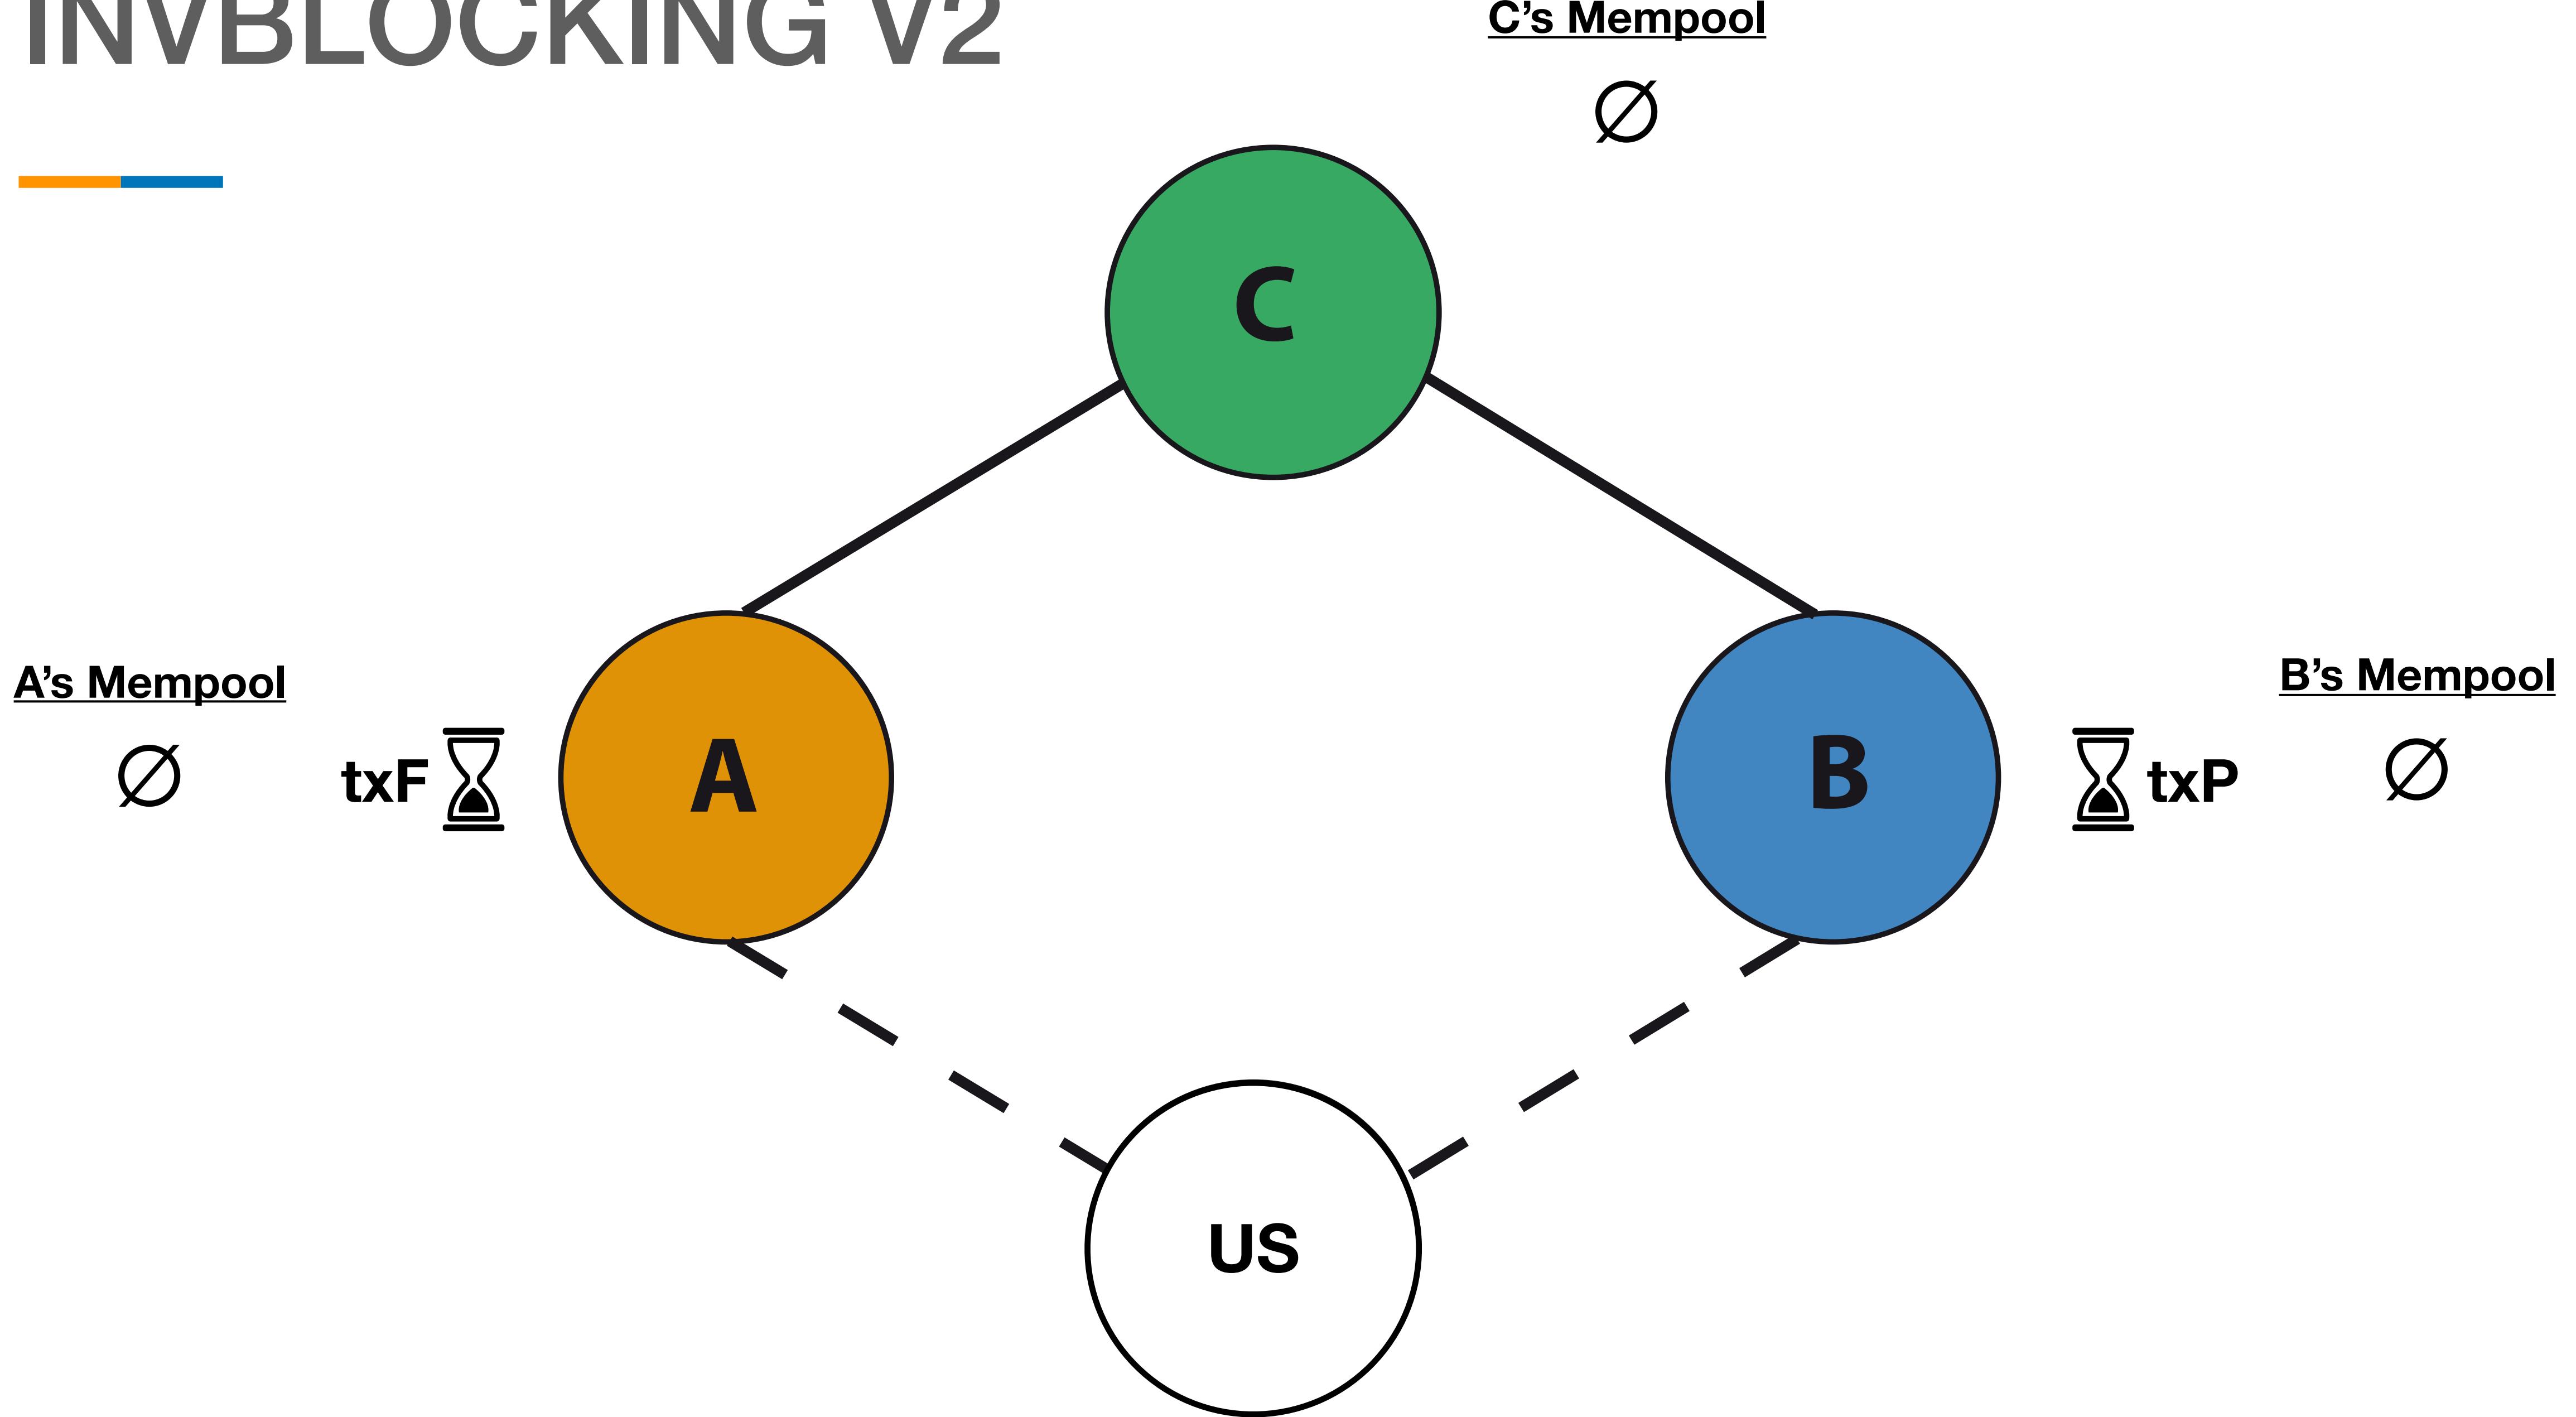
\emptyset



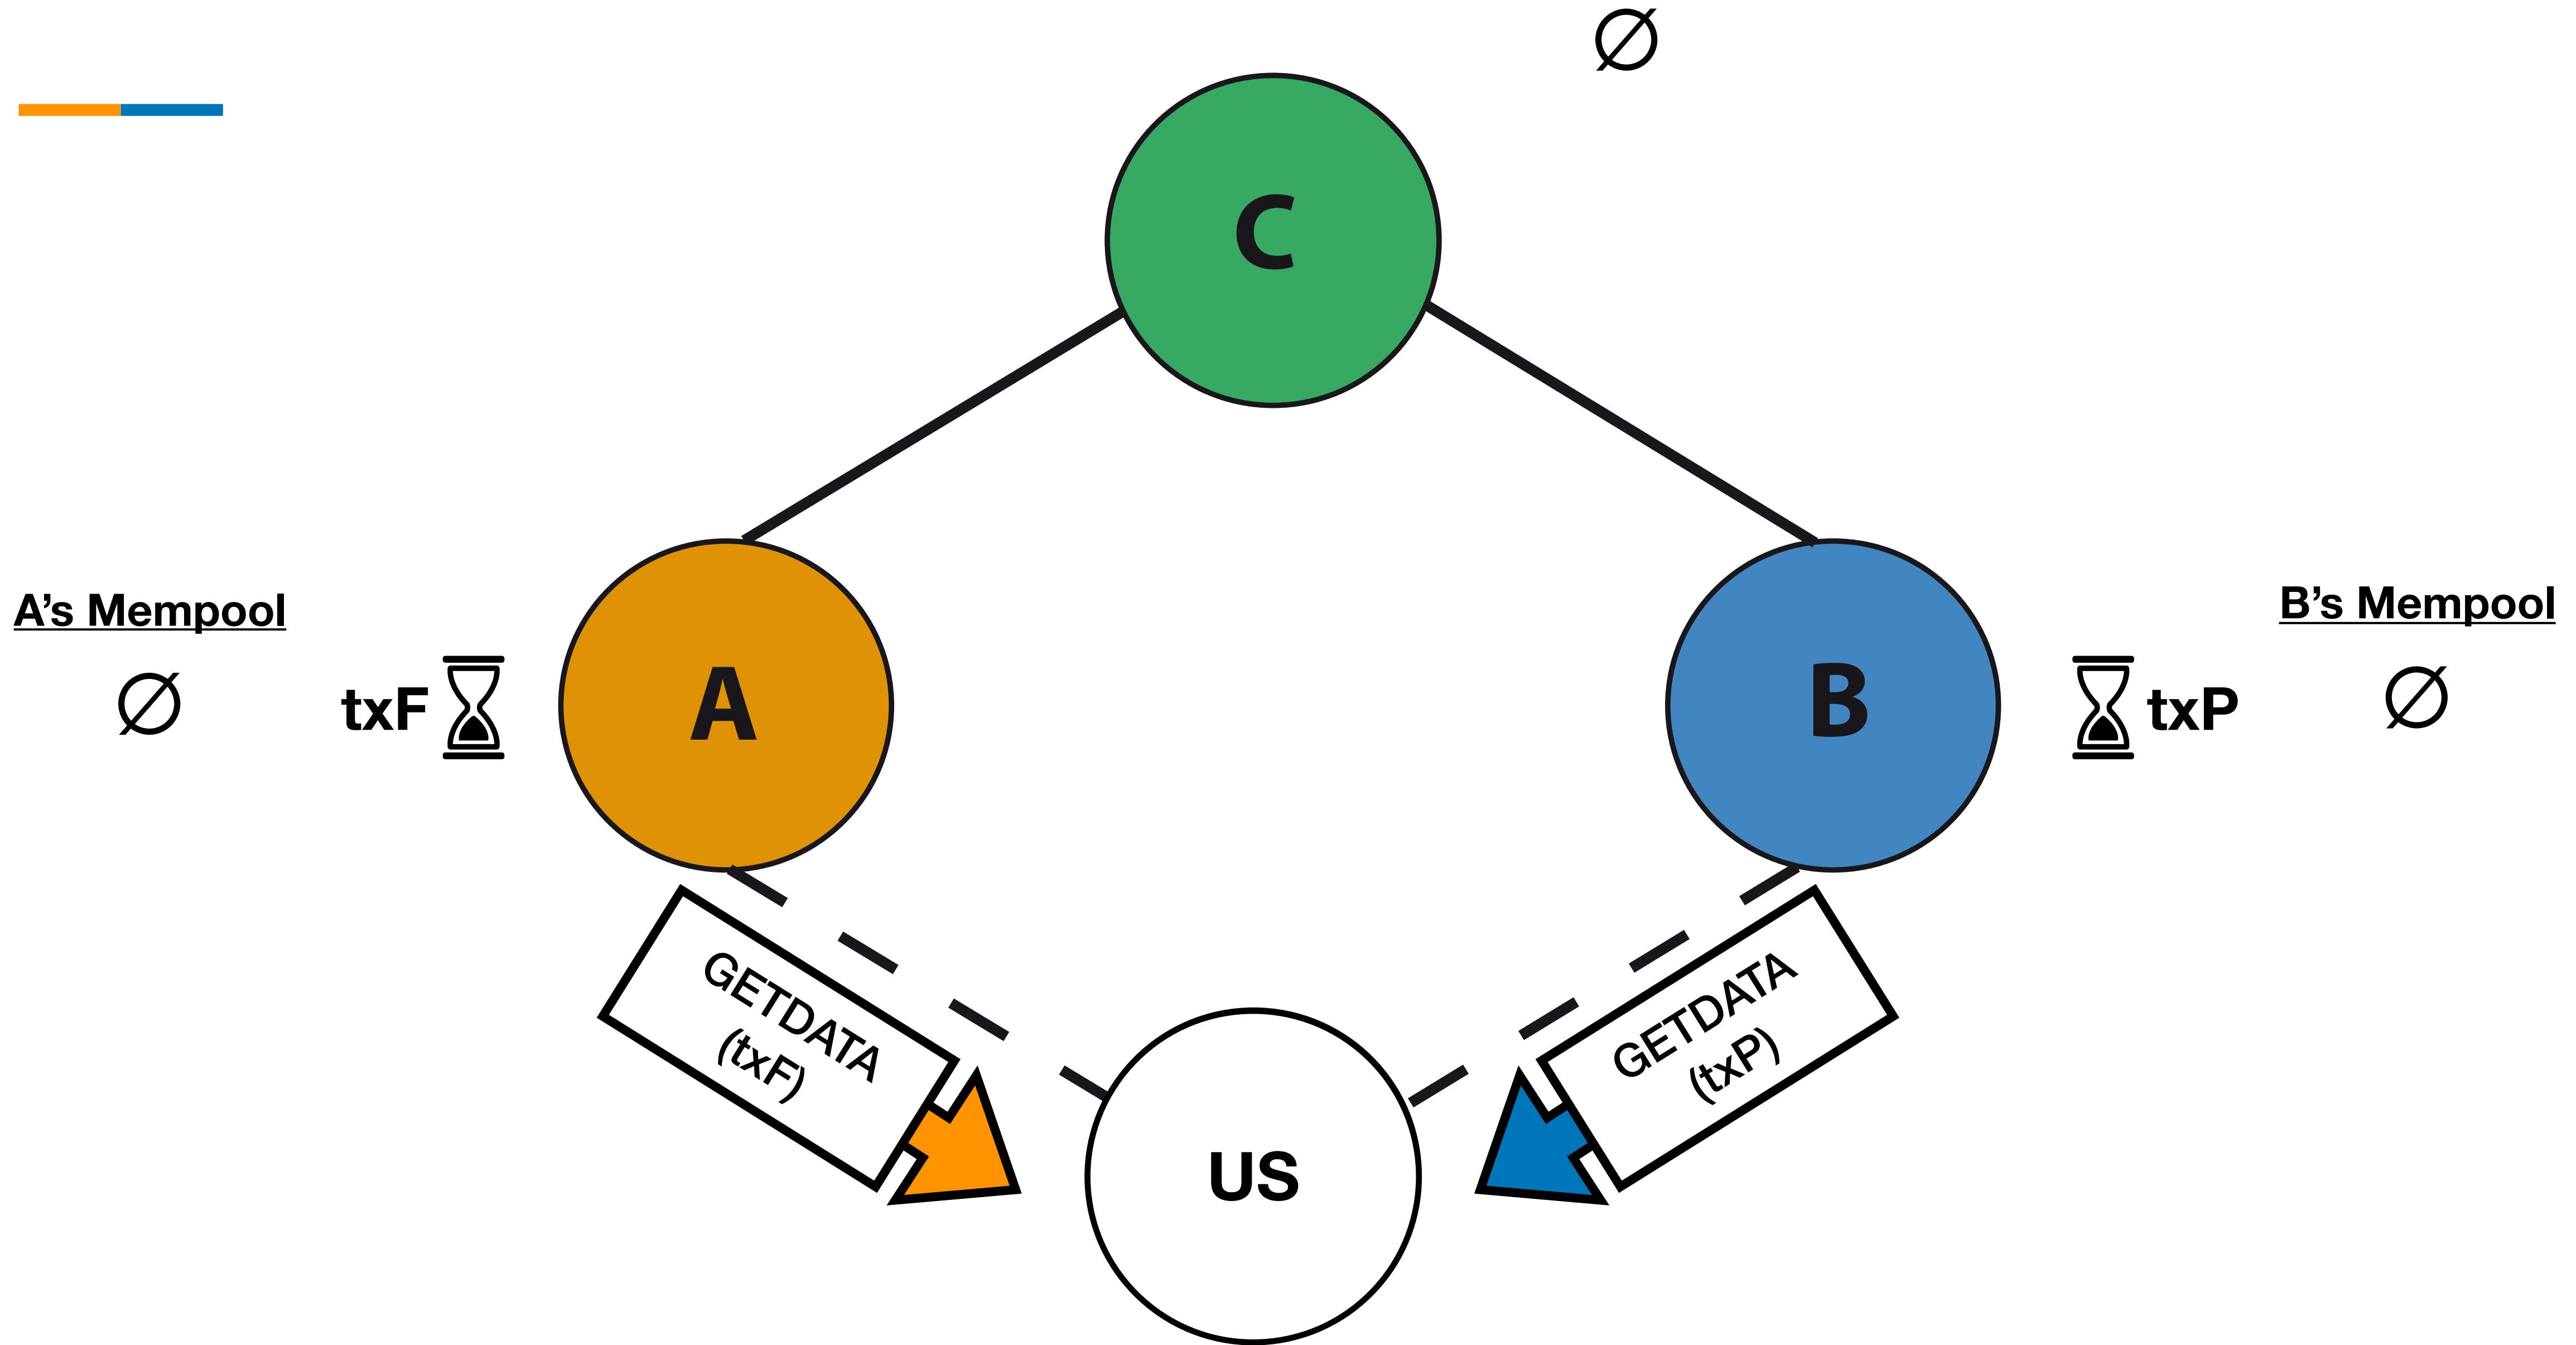
B's Mempool

\emptyset

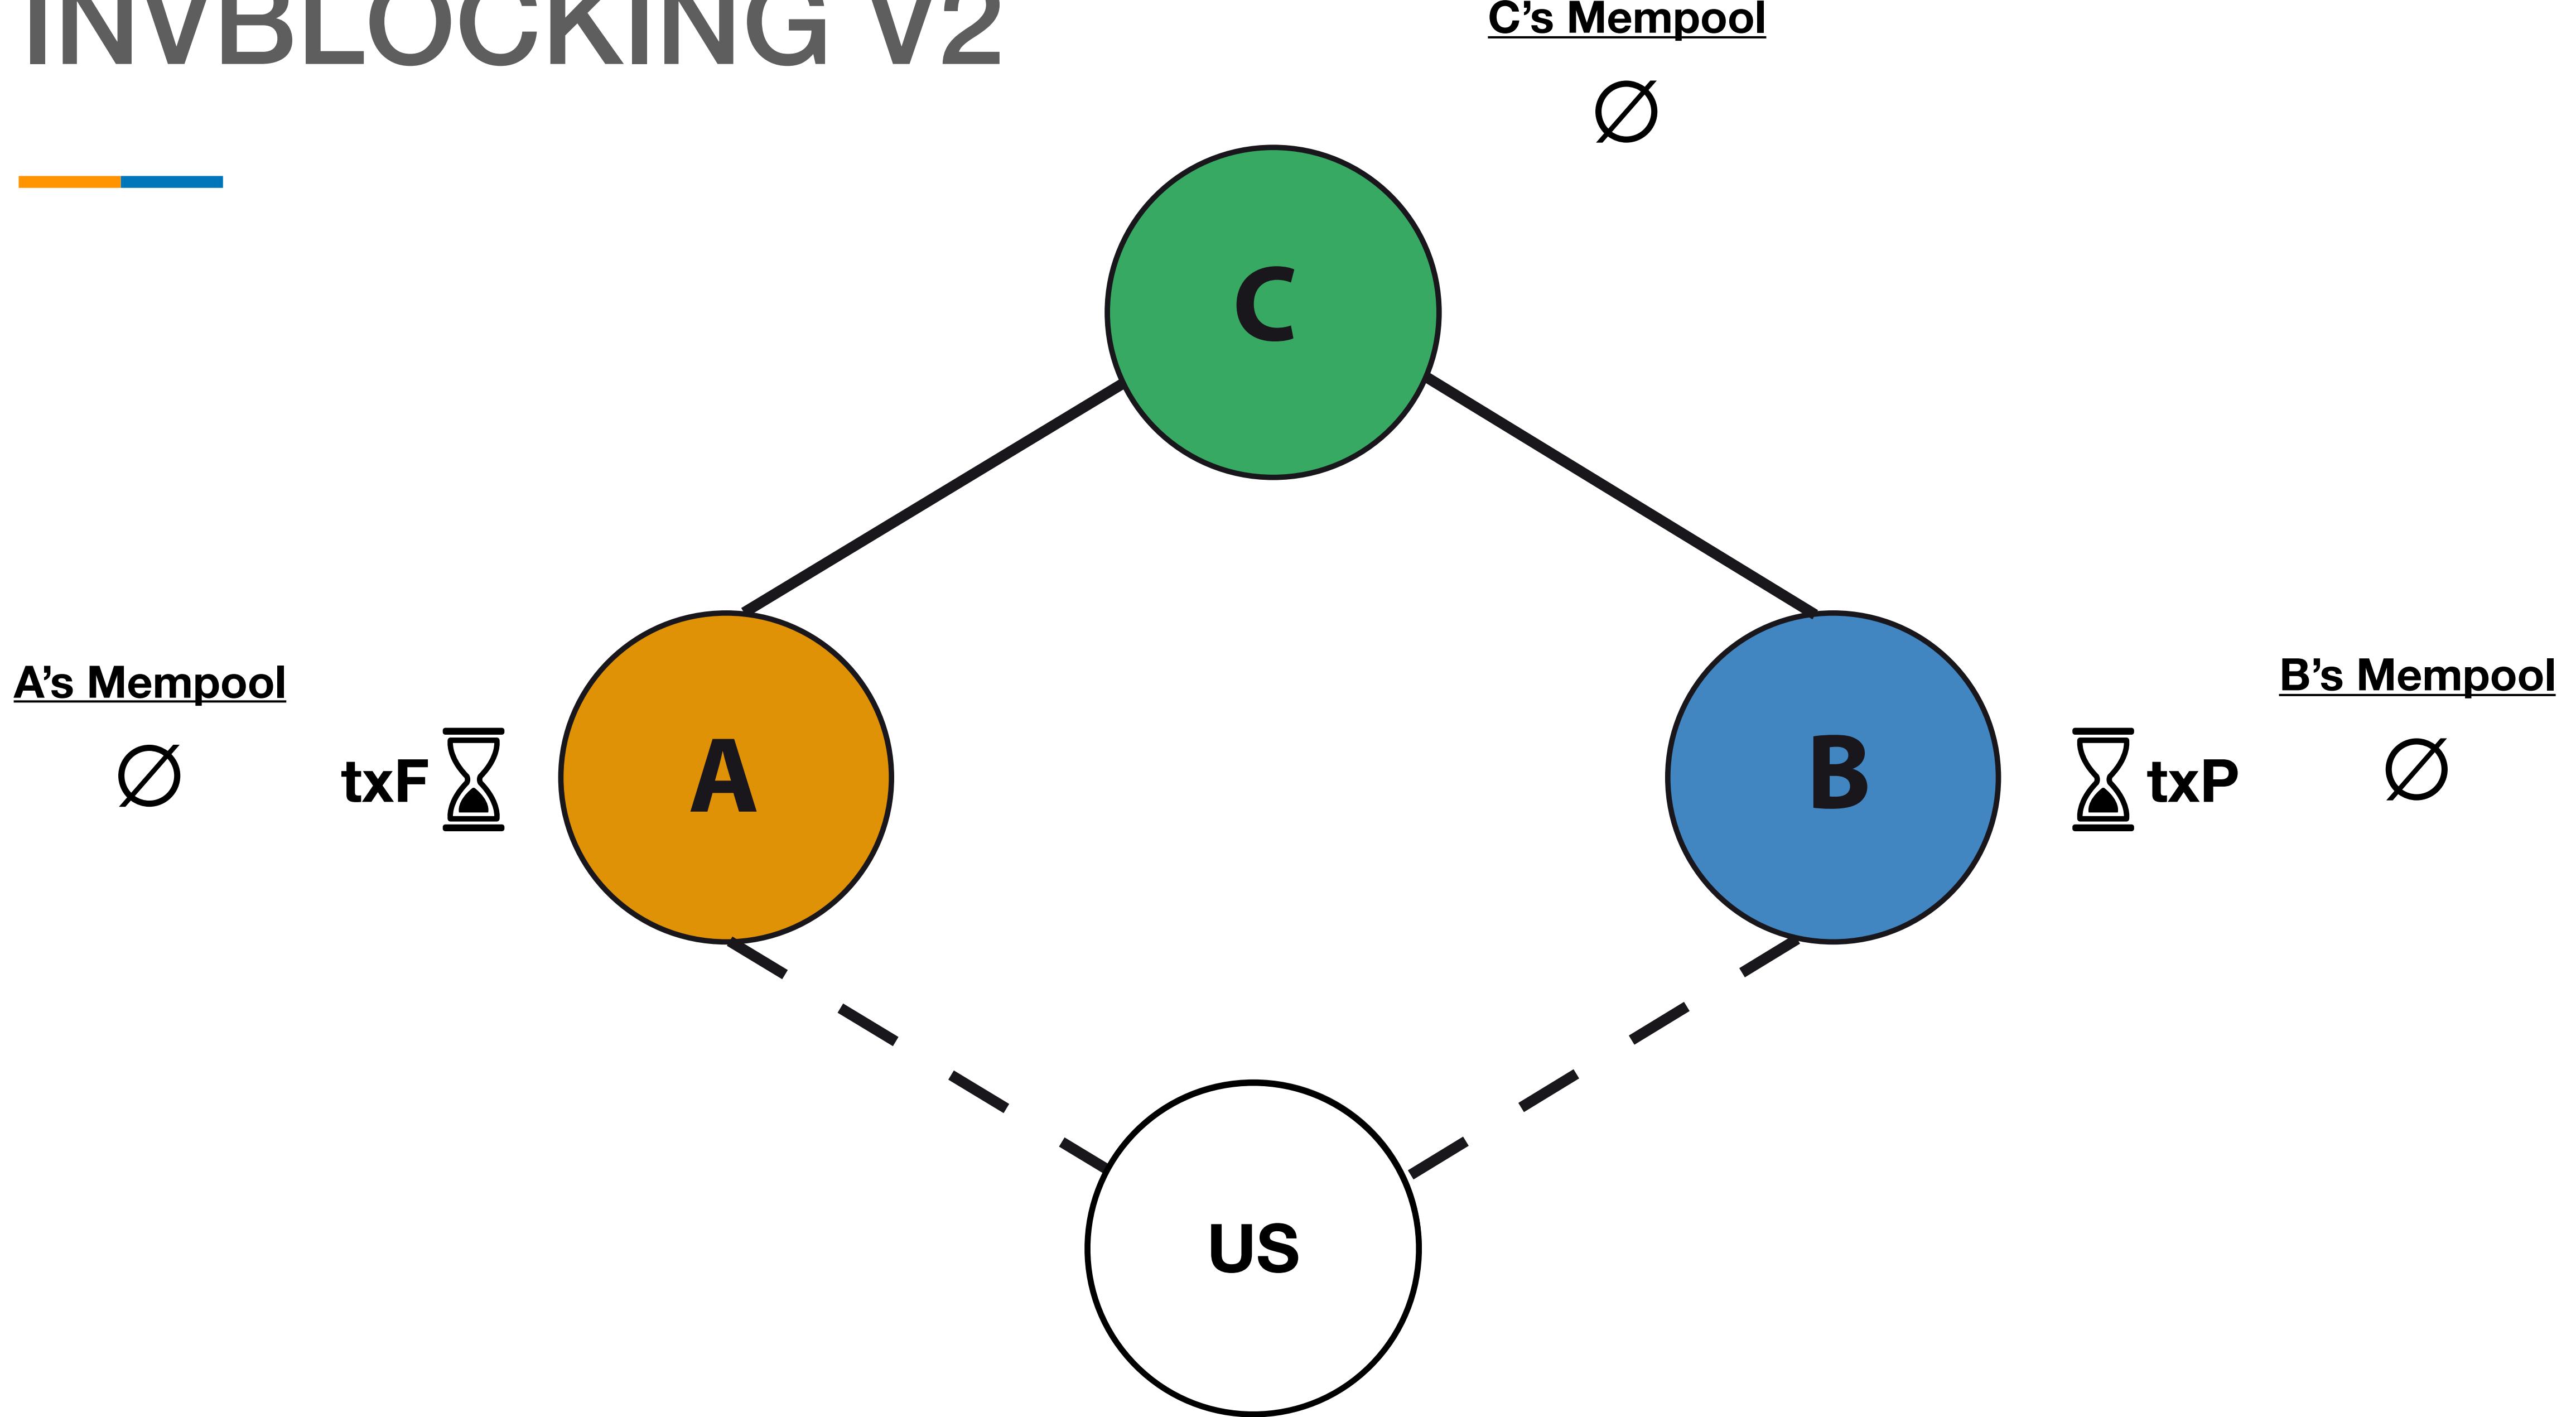
INVBLOCKING V2



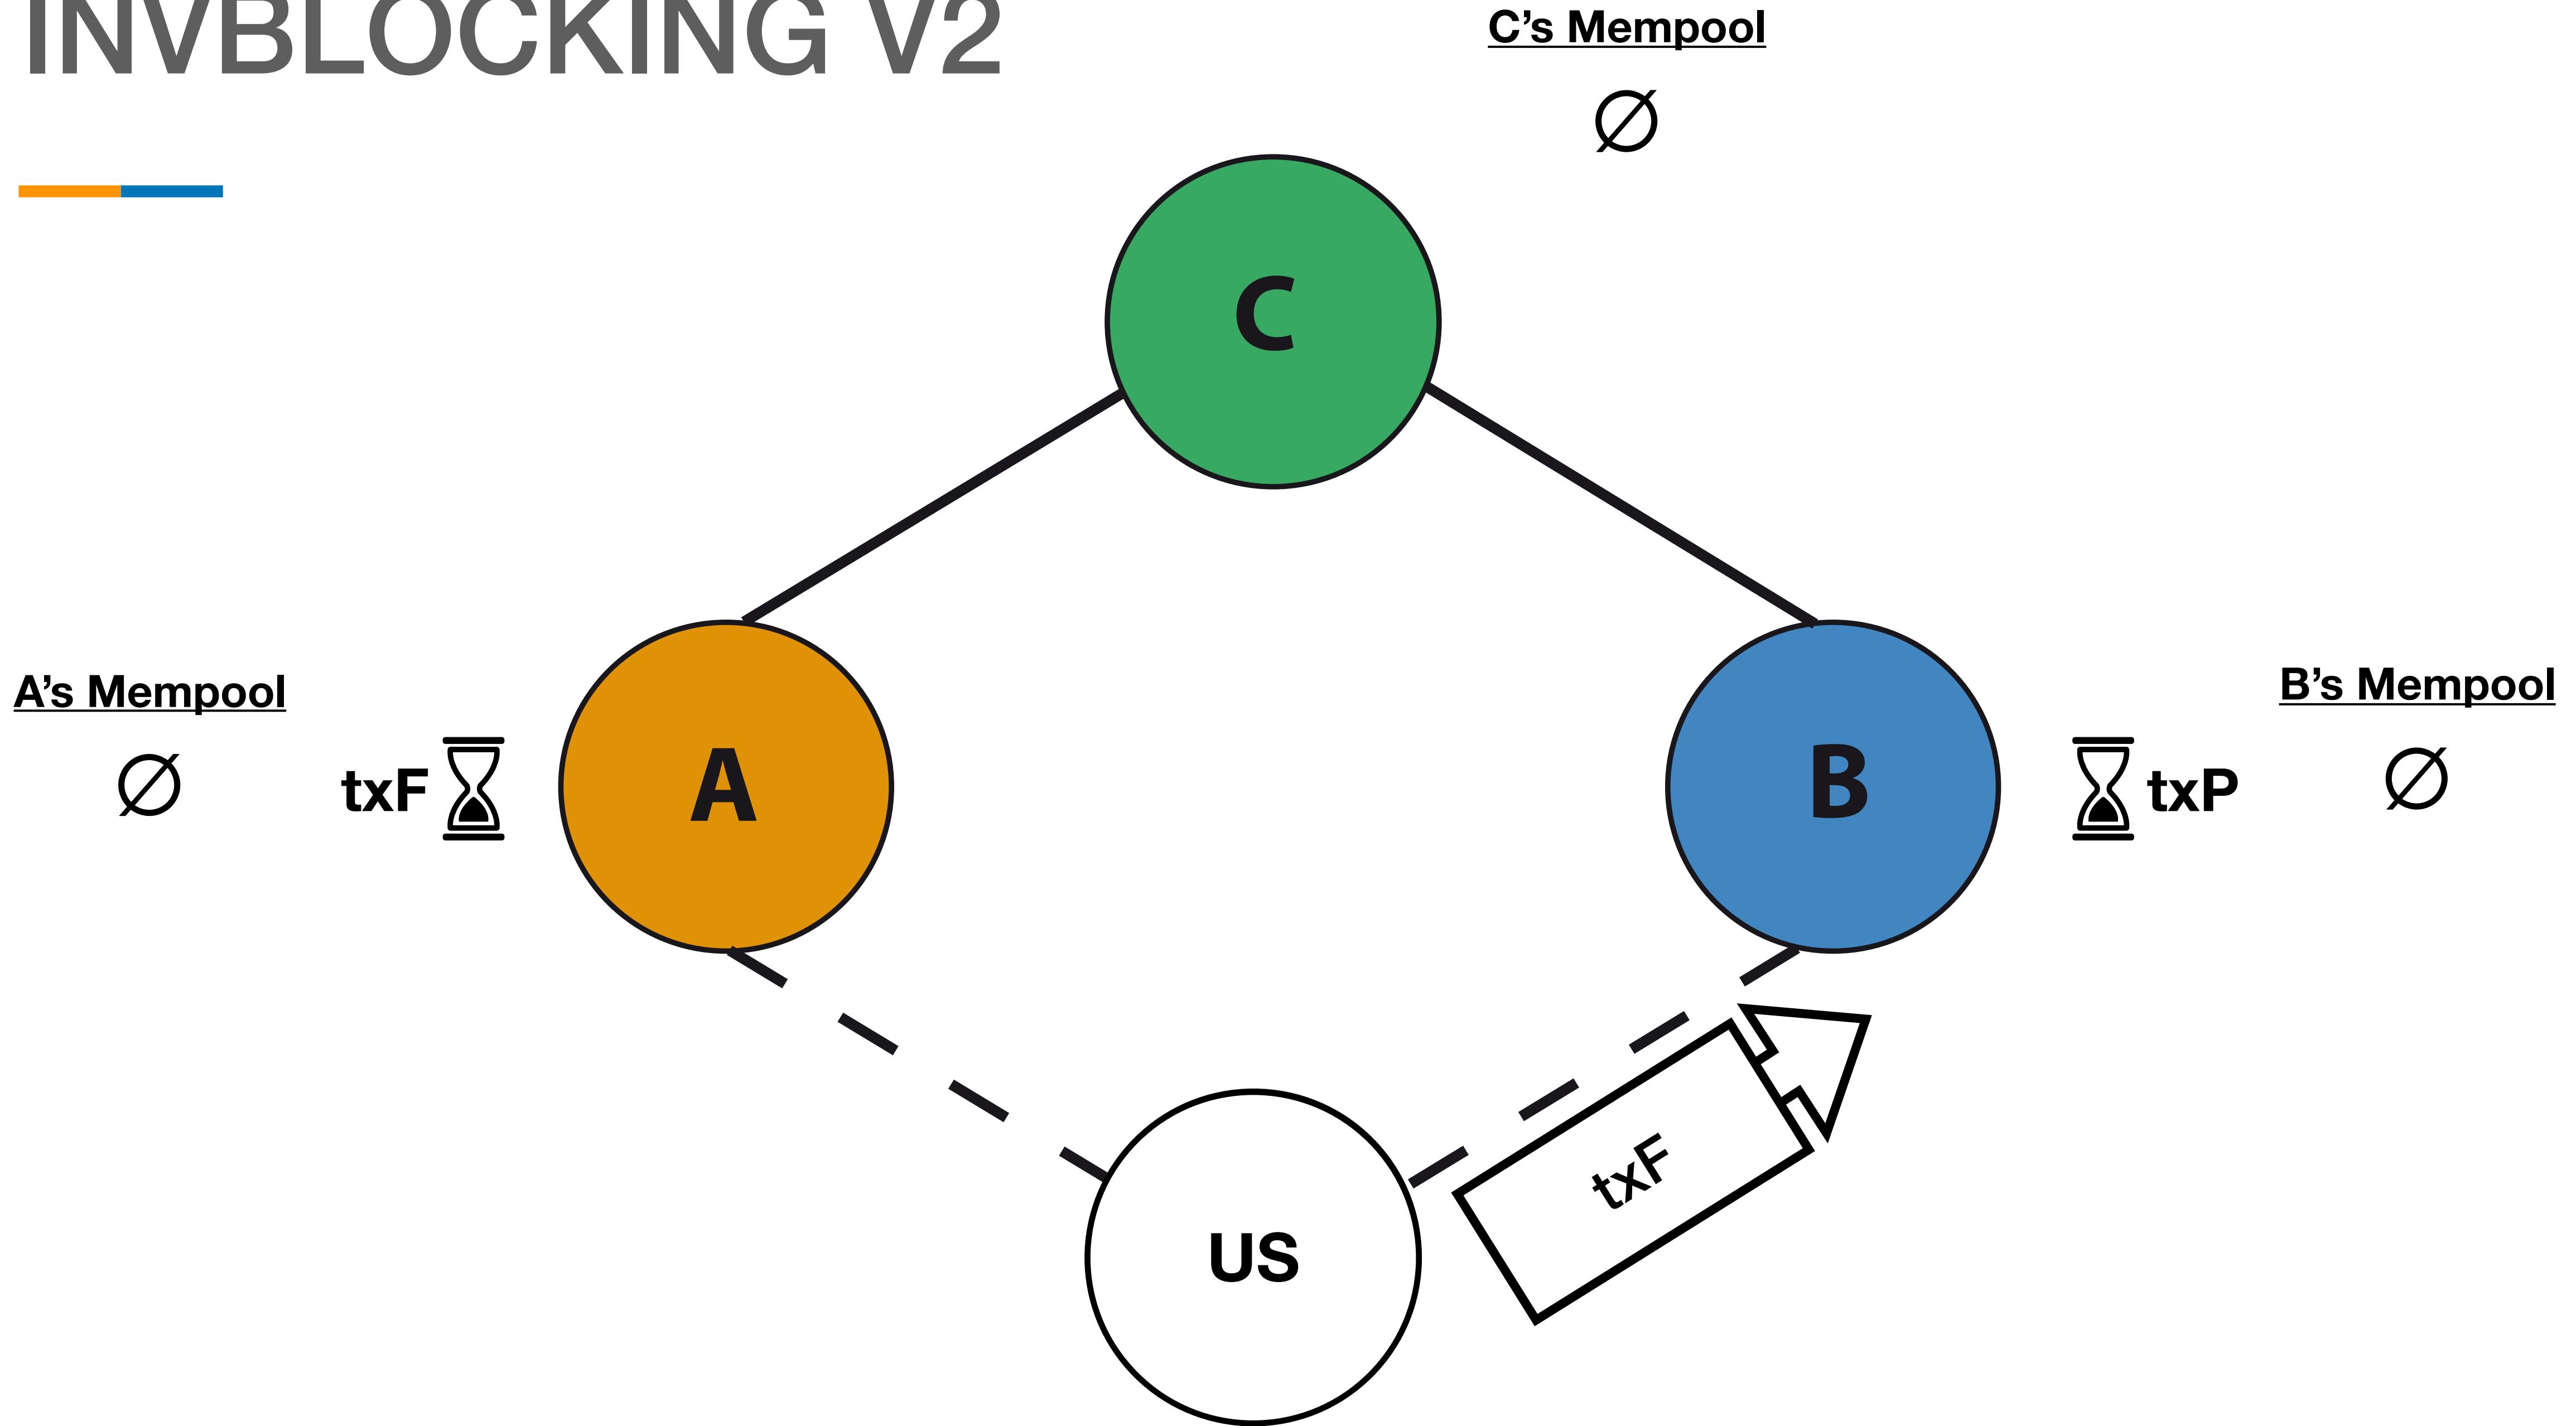
INVBLOCKING V2



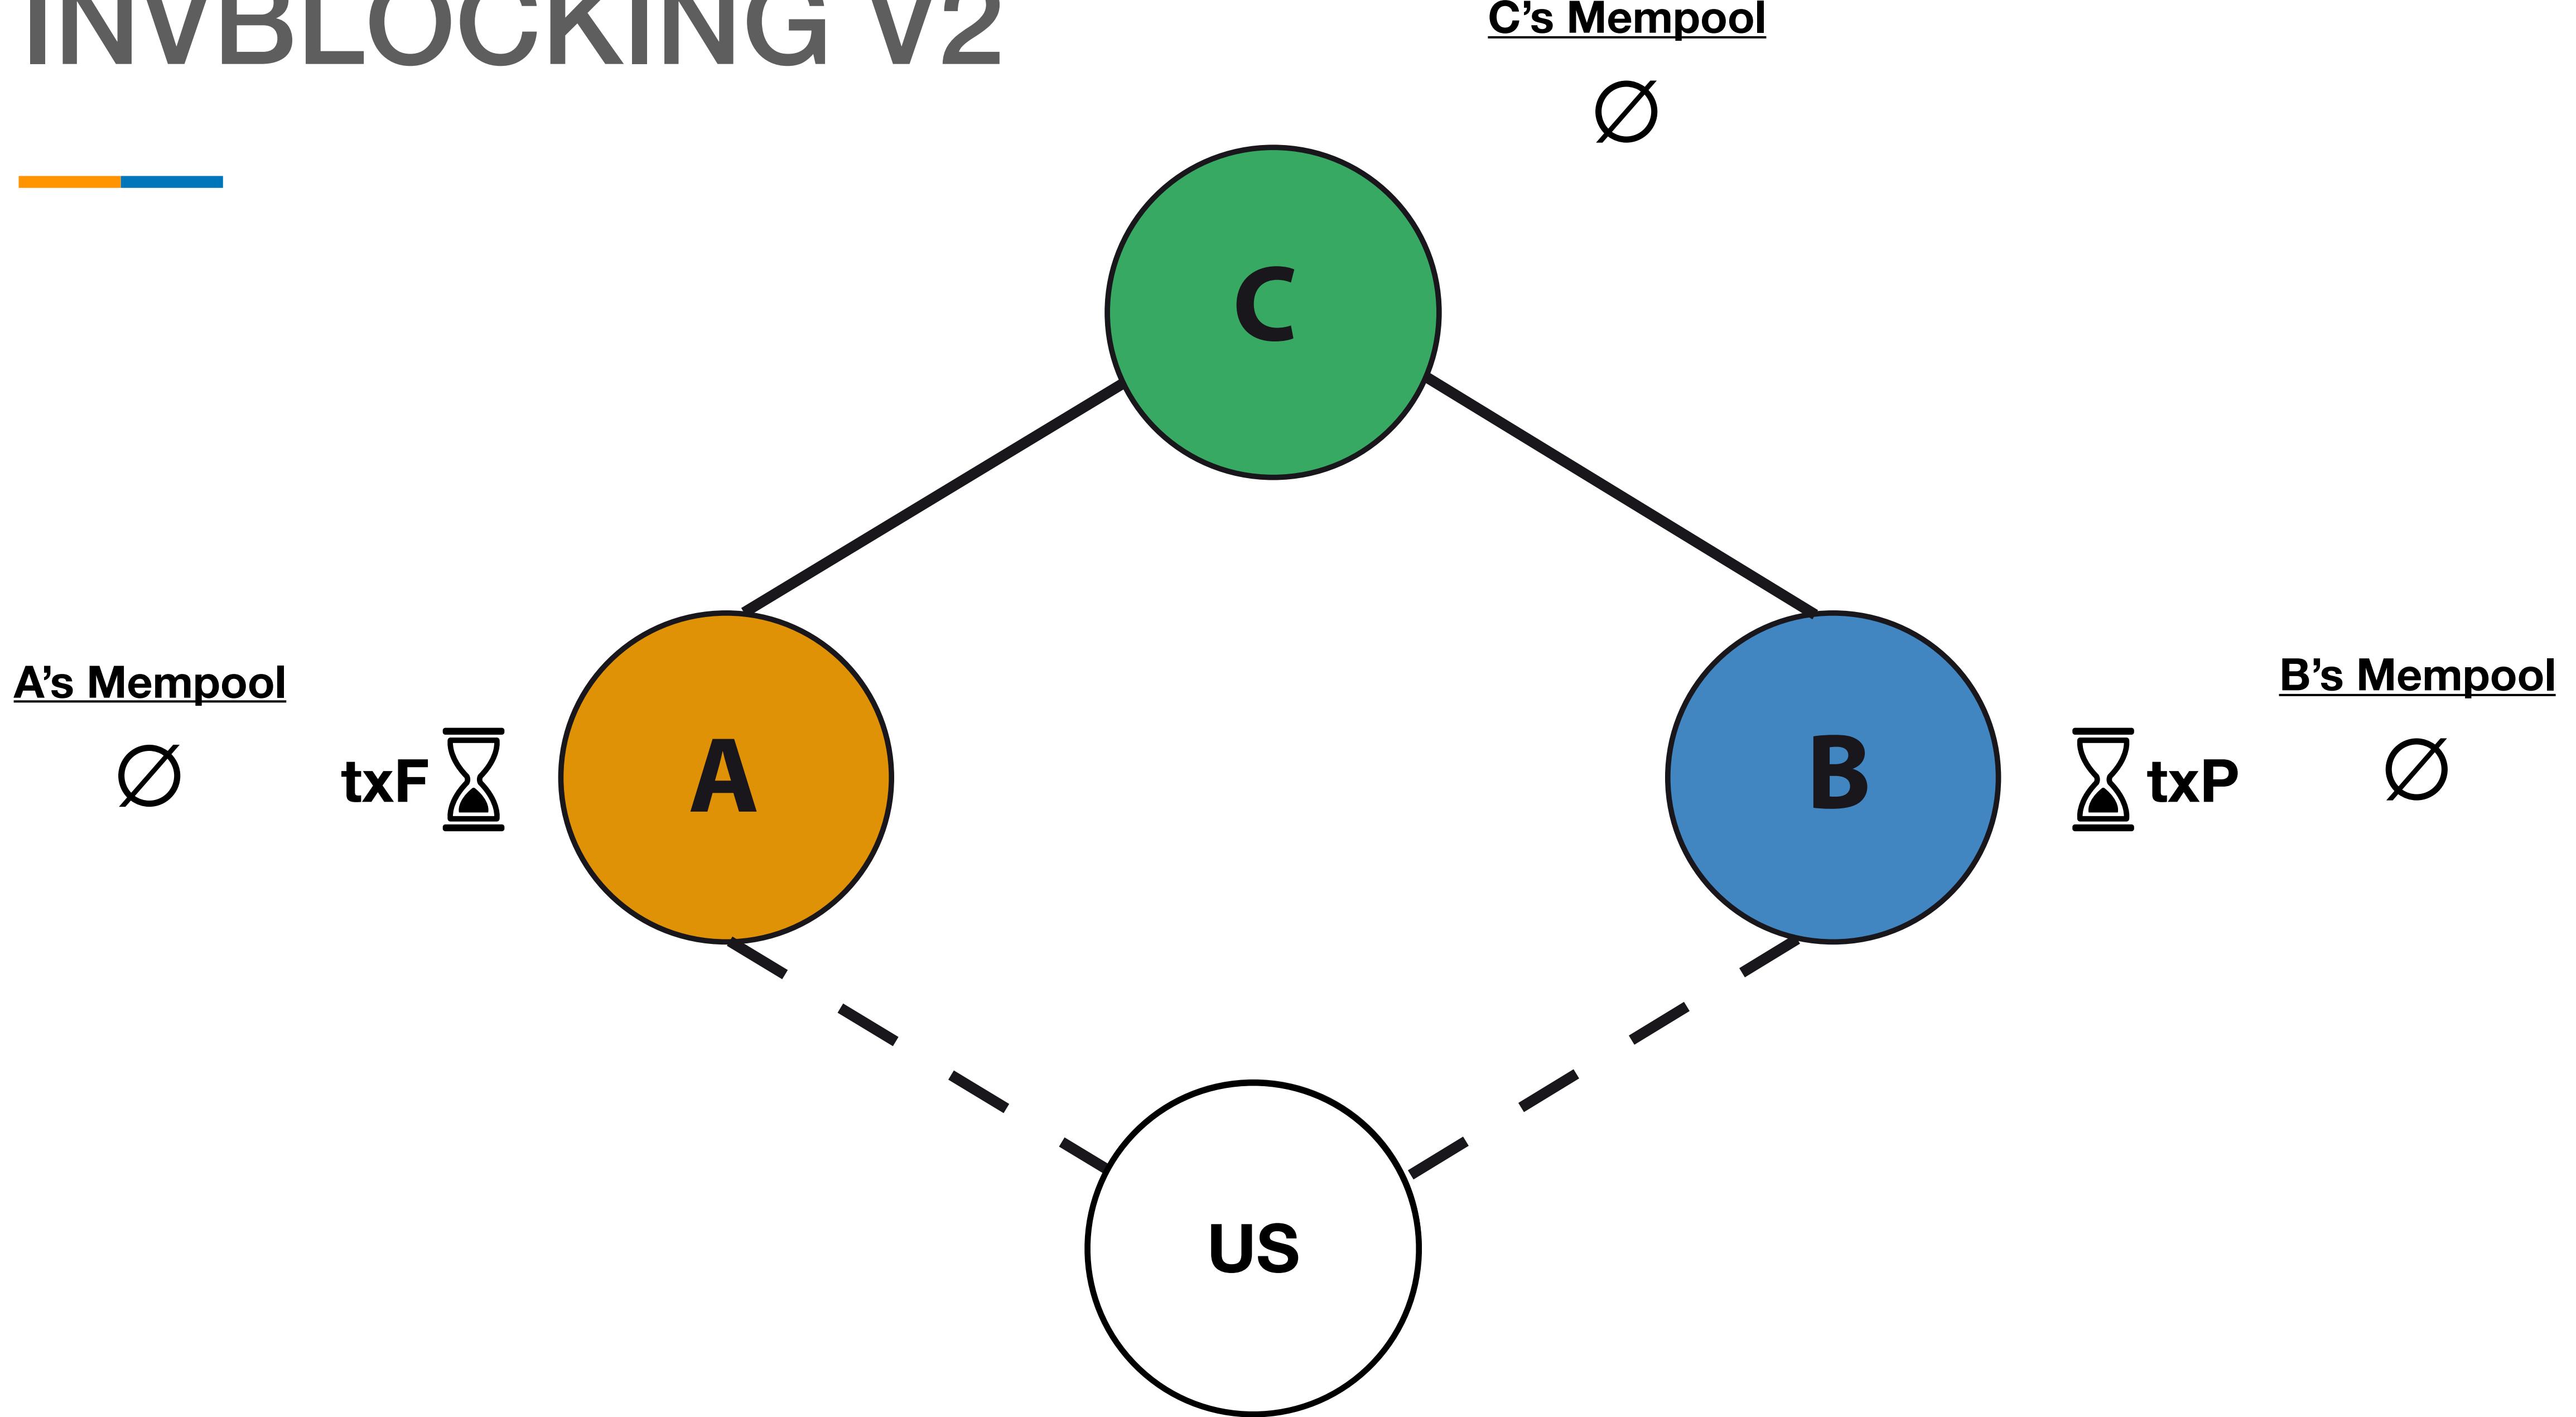
INVBLOCKING V2



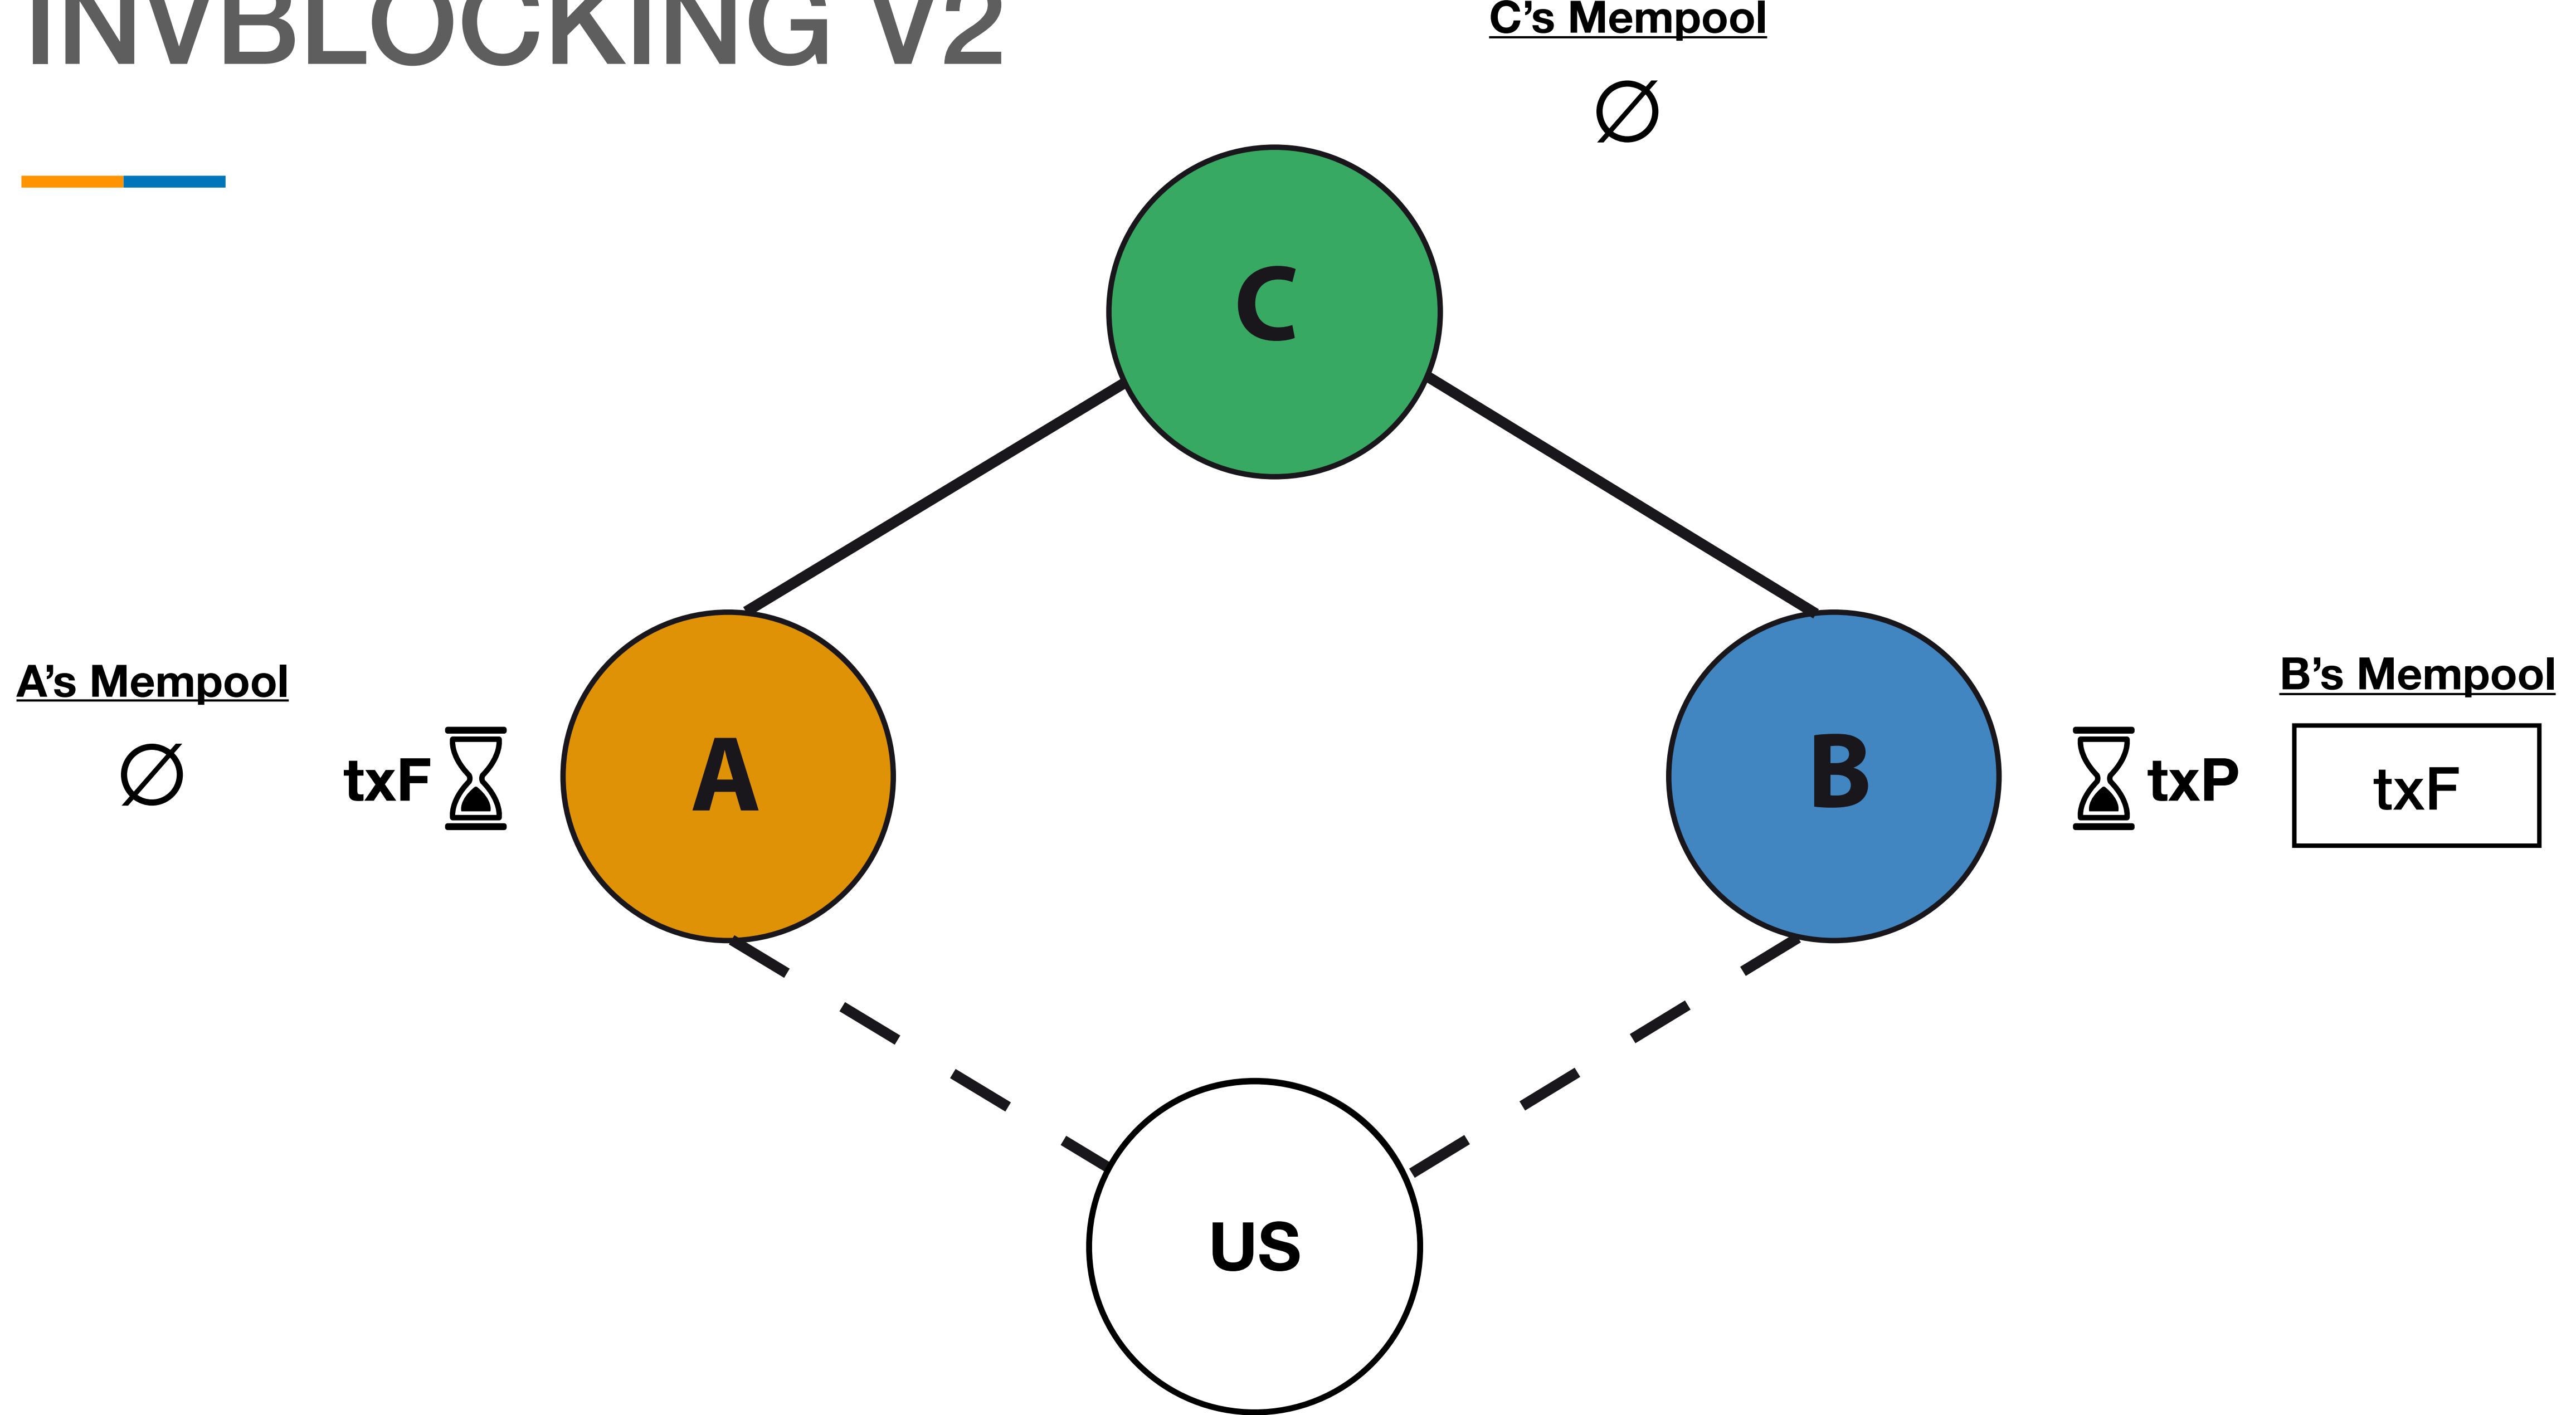
INVBLOCKING V2



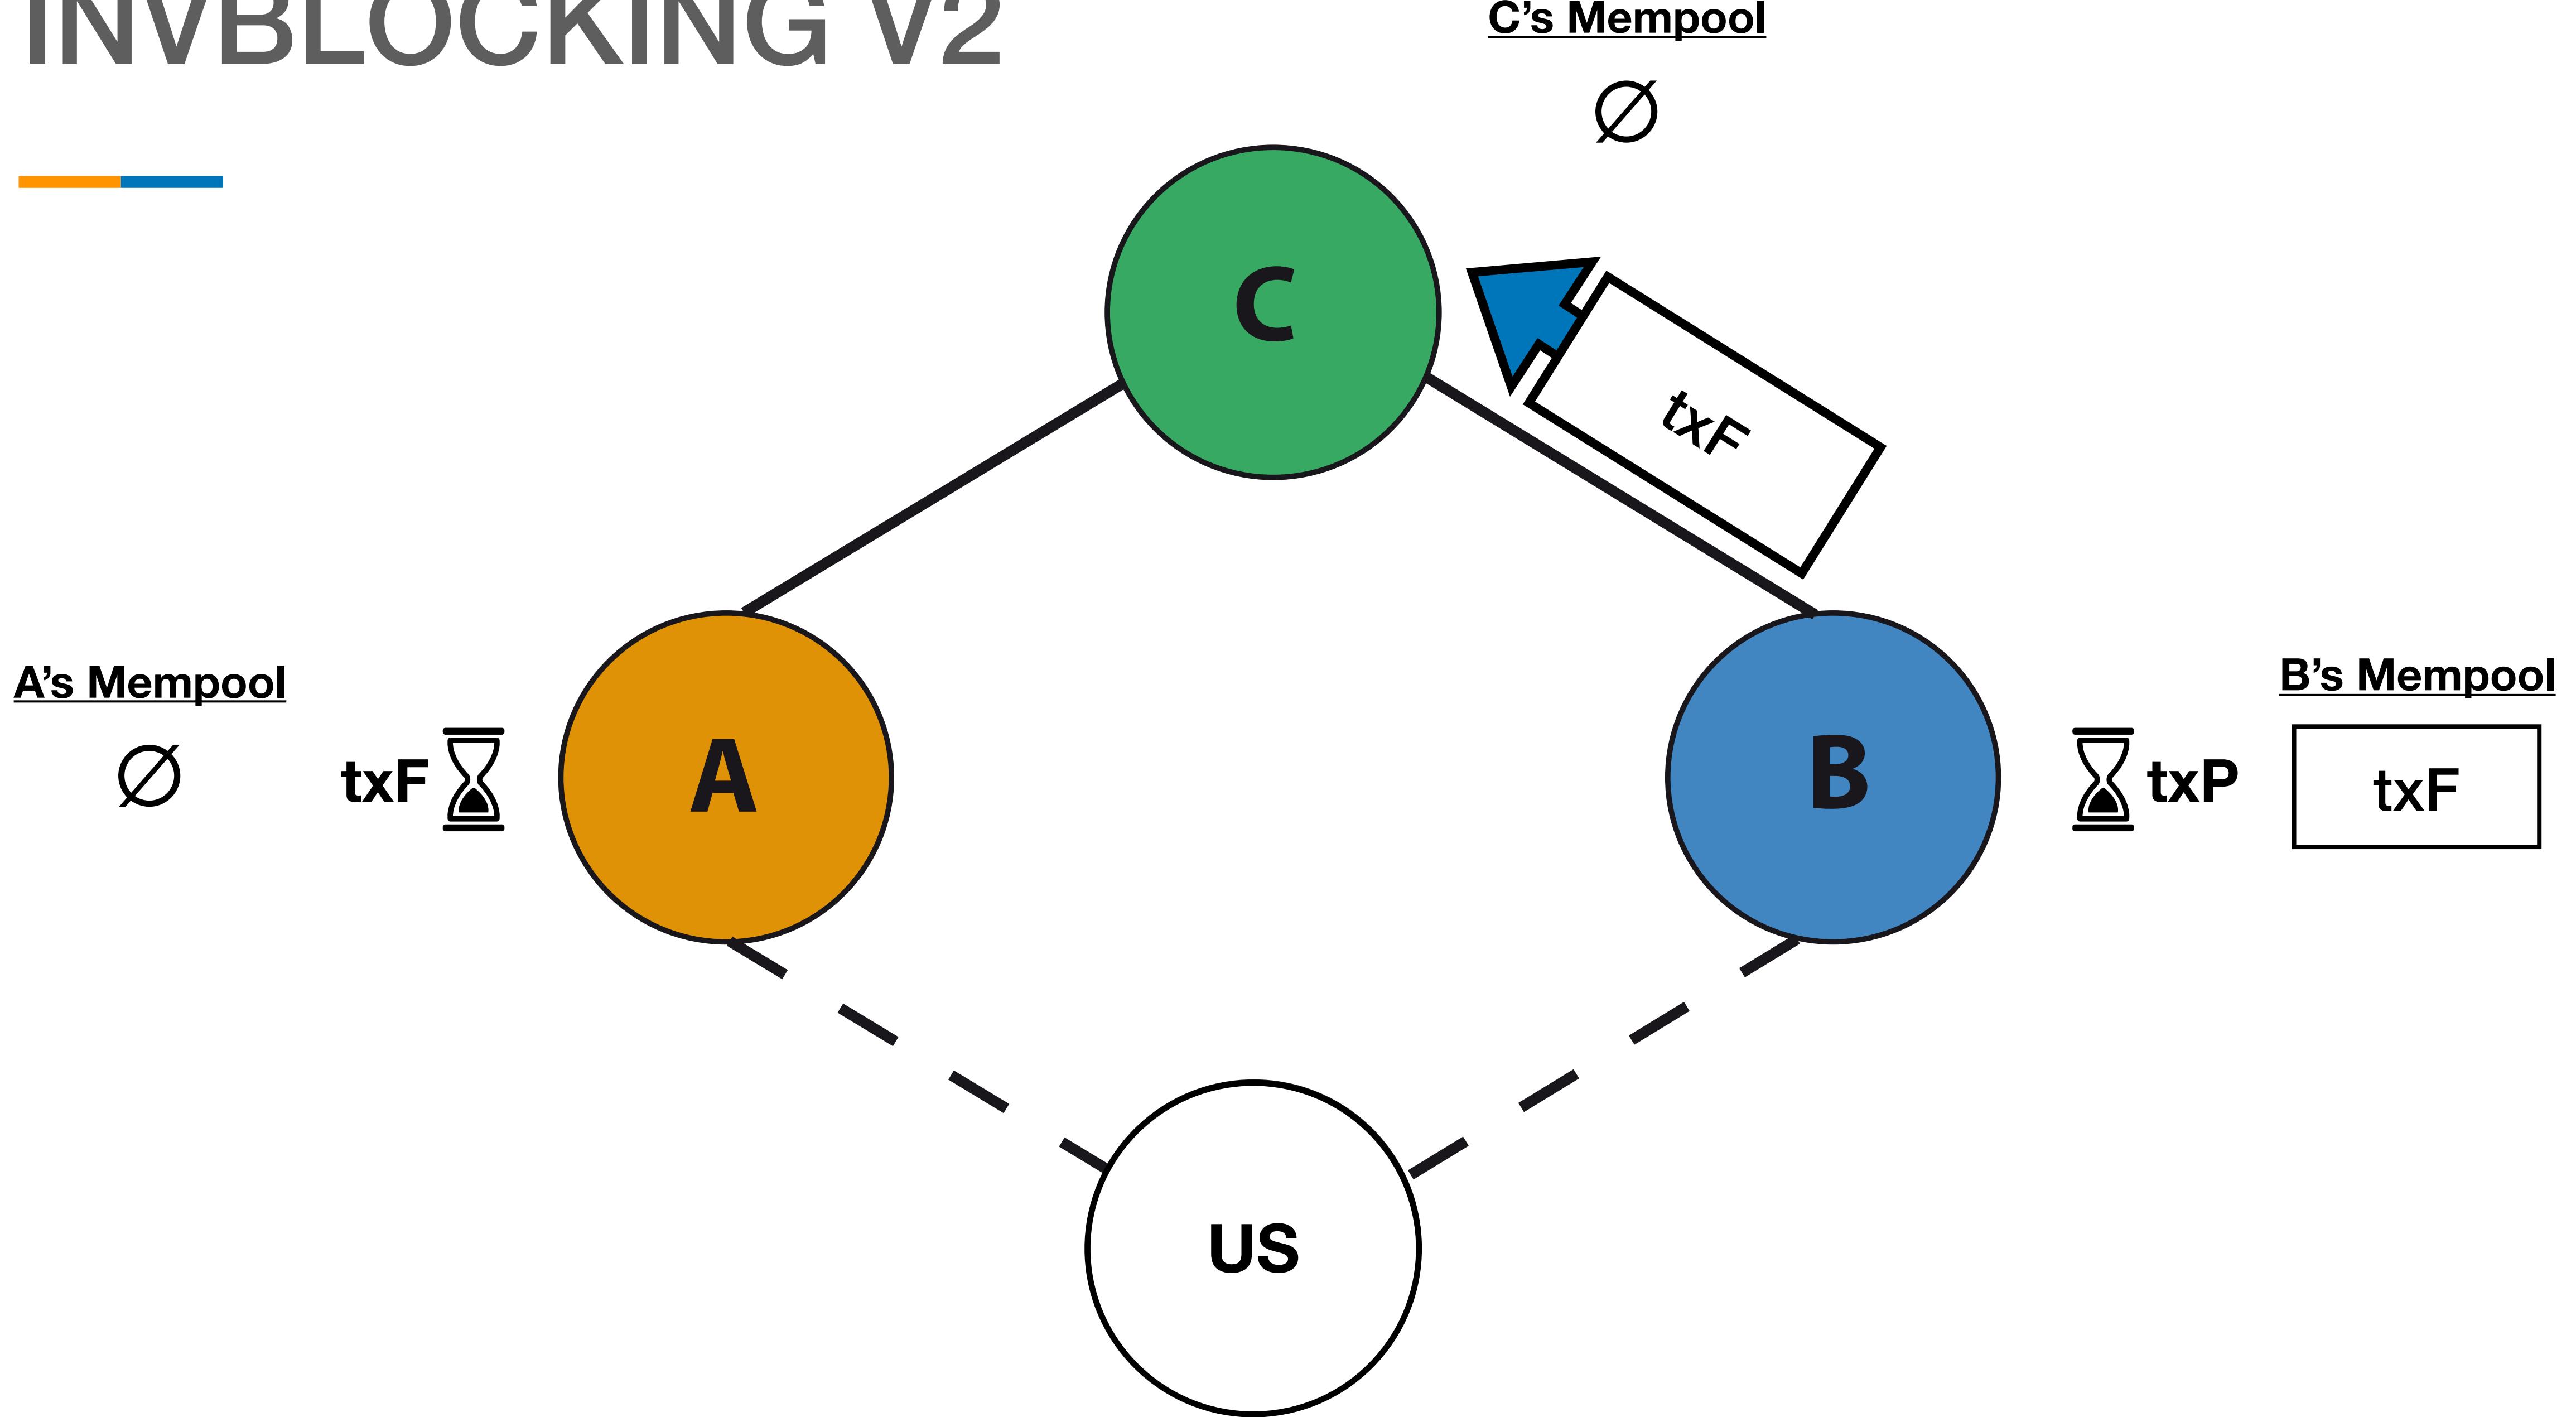
INVBLOCKING V2



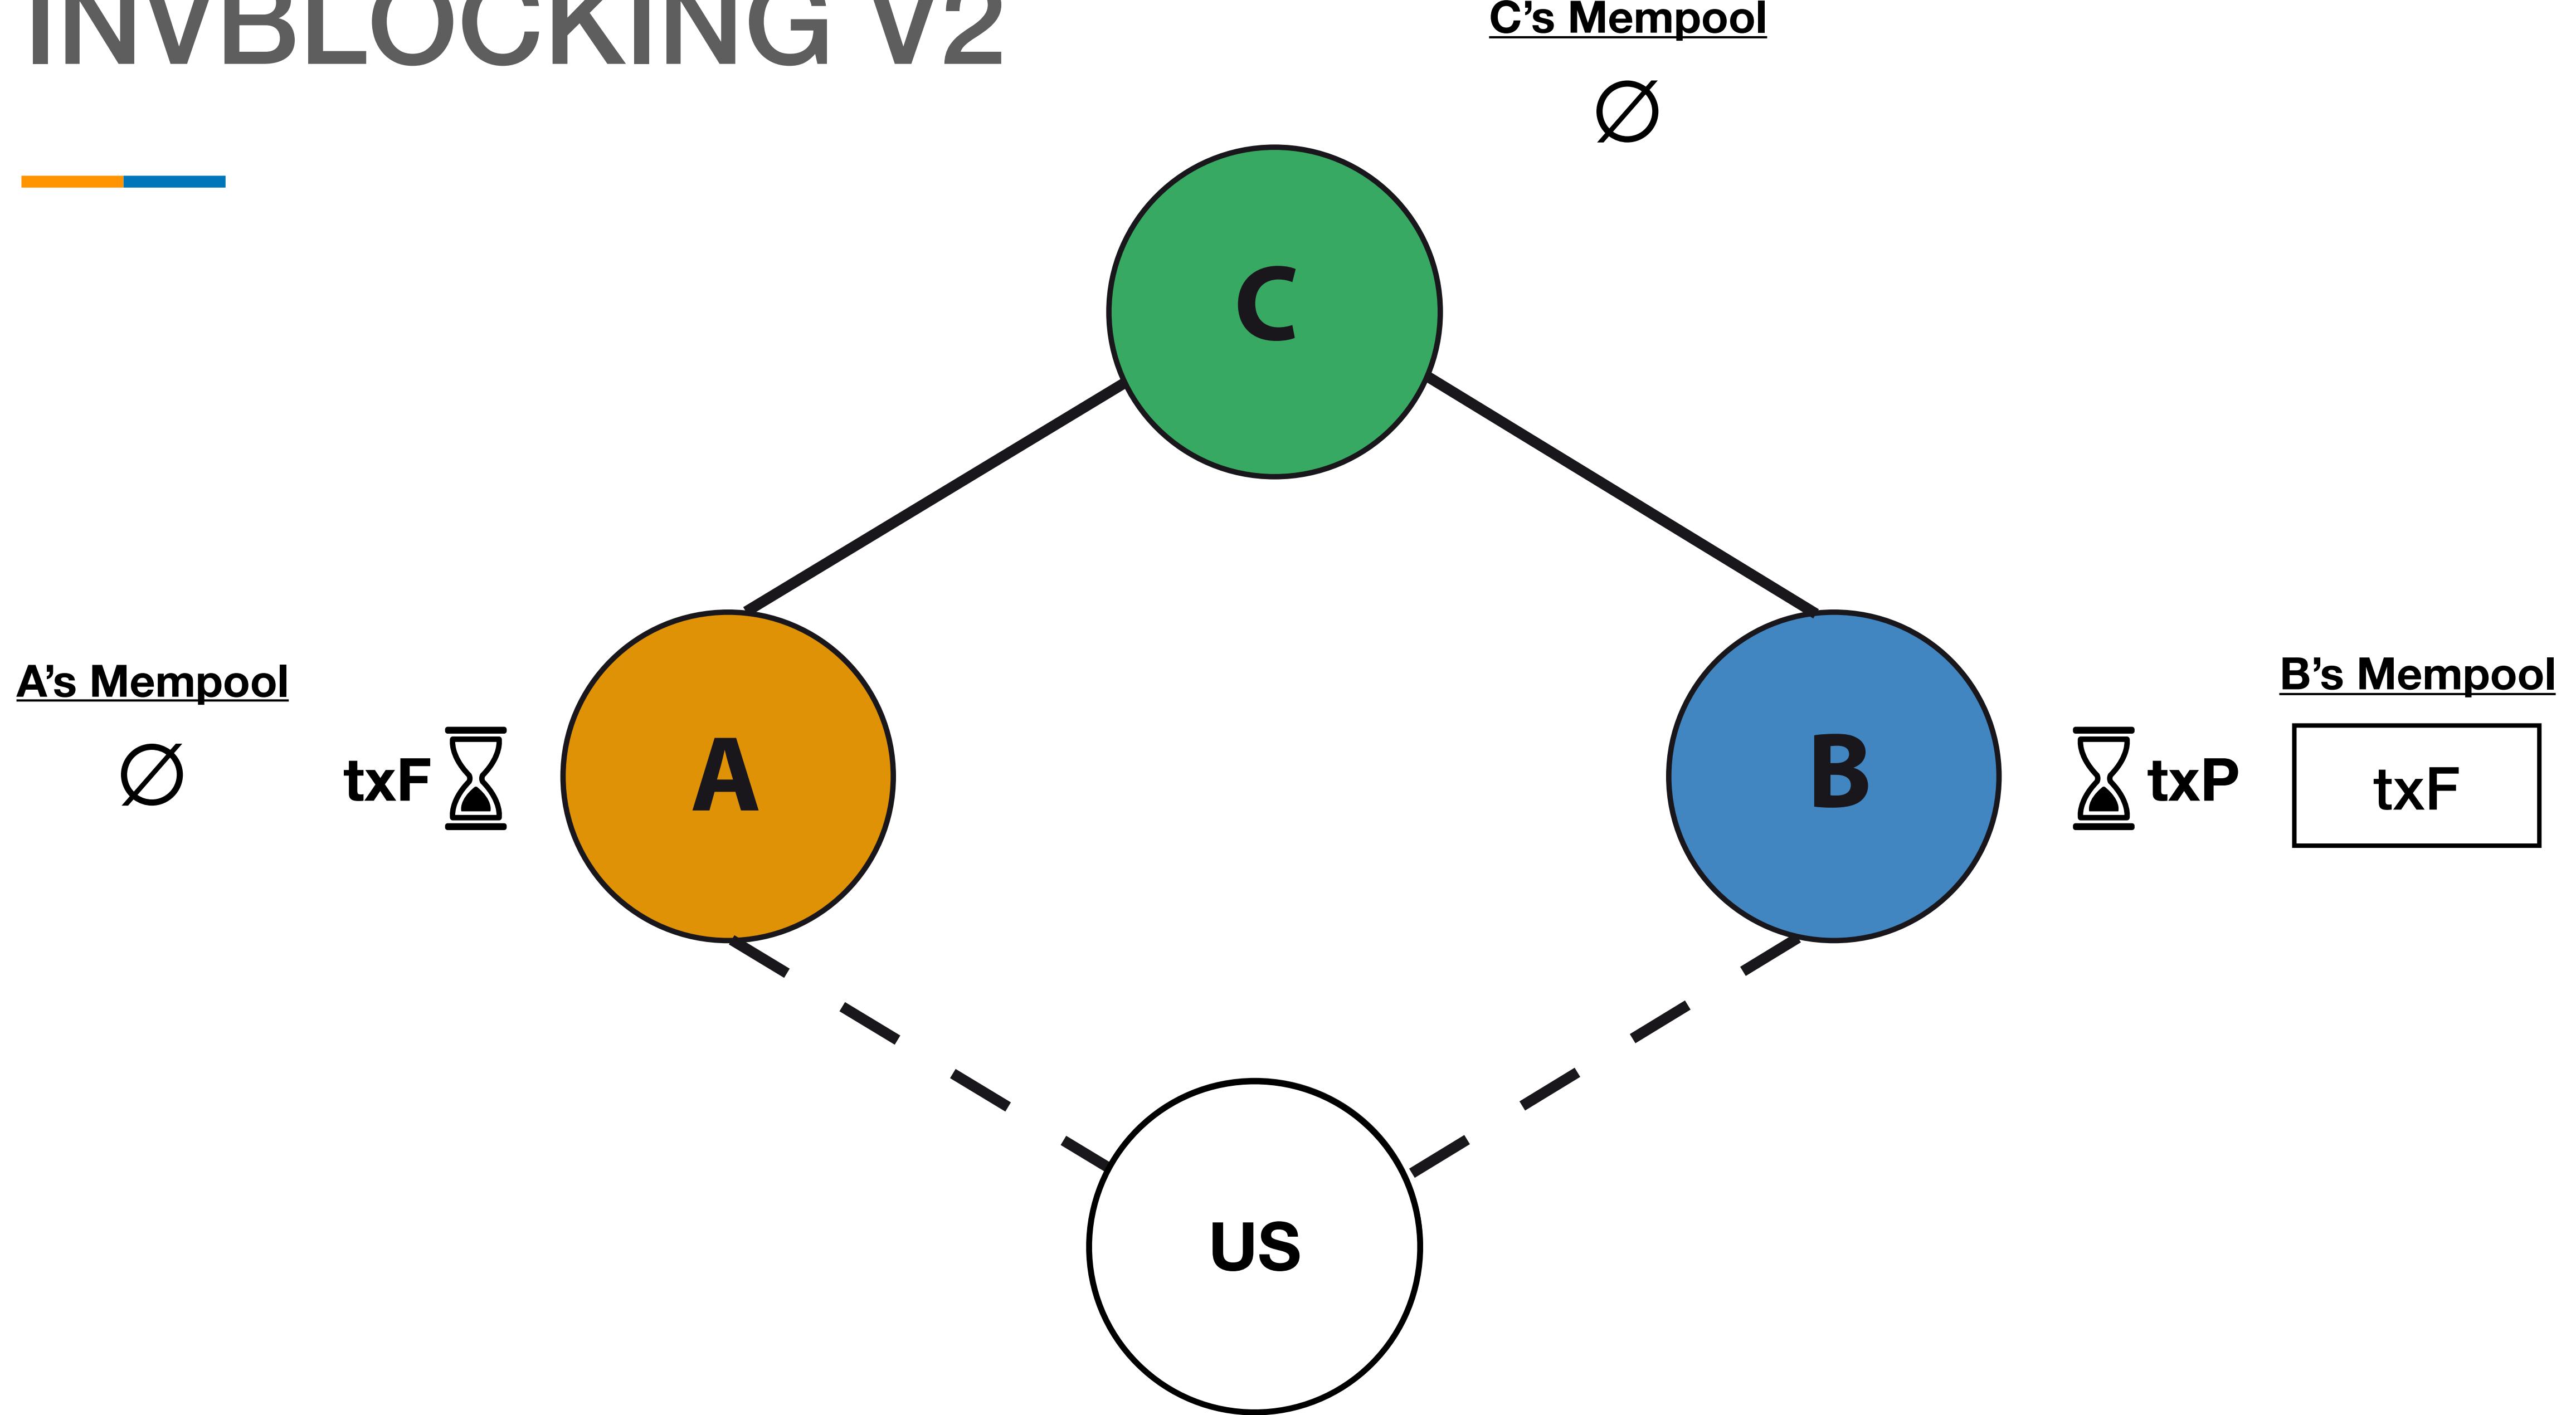
INVBLOCKING V2



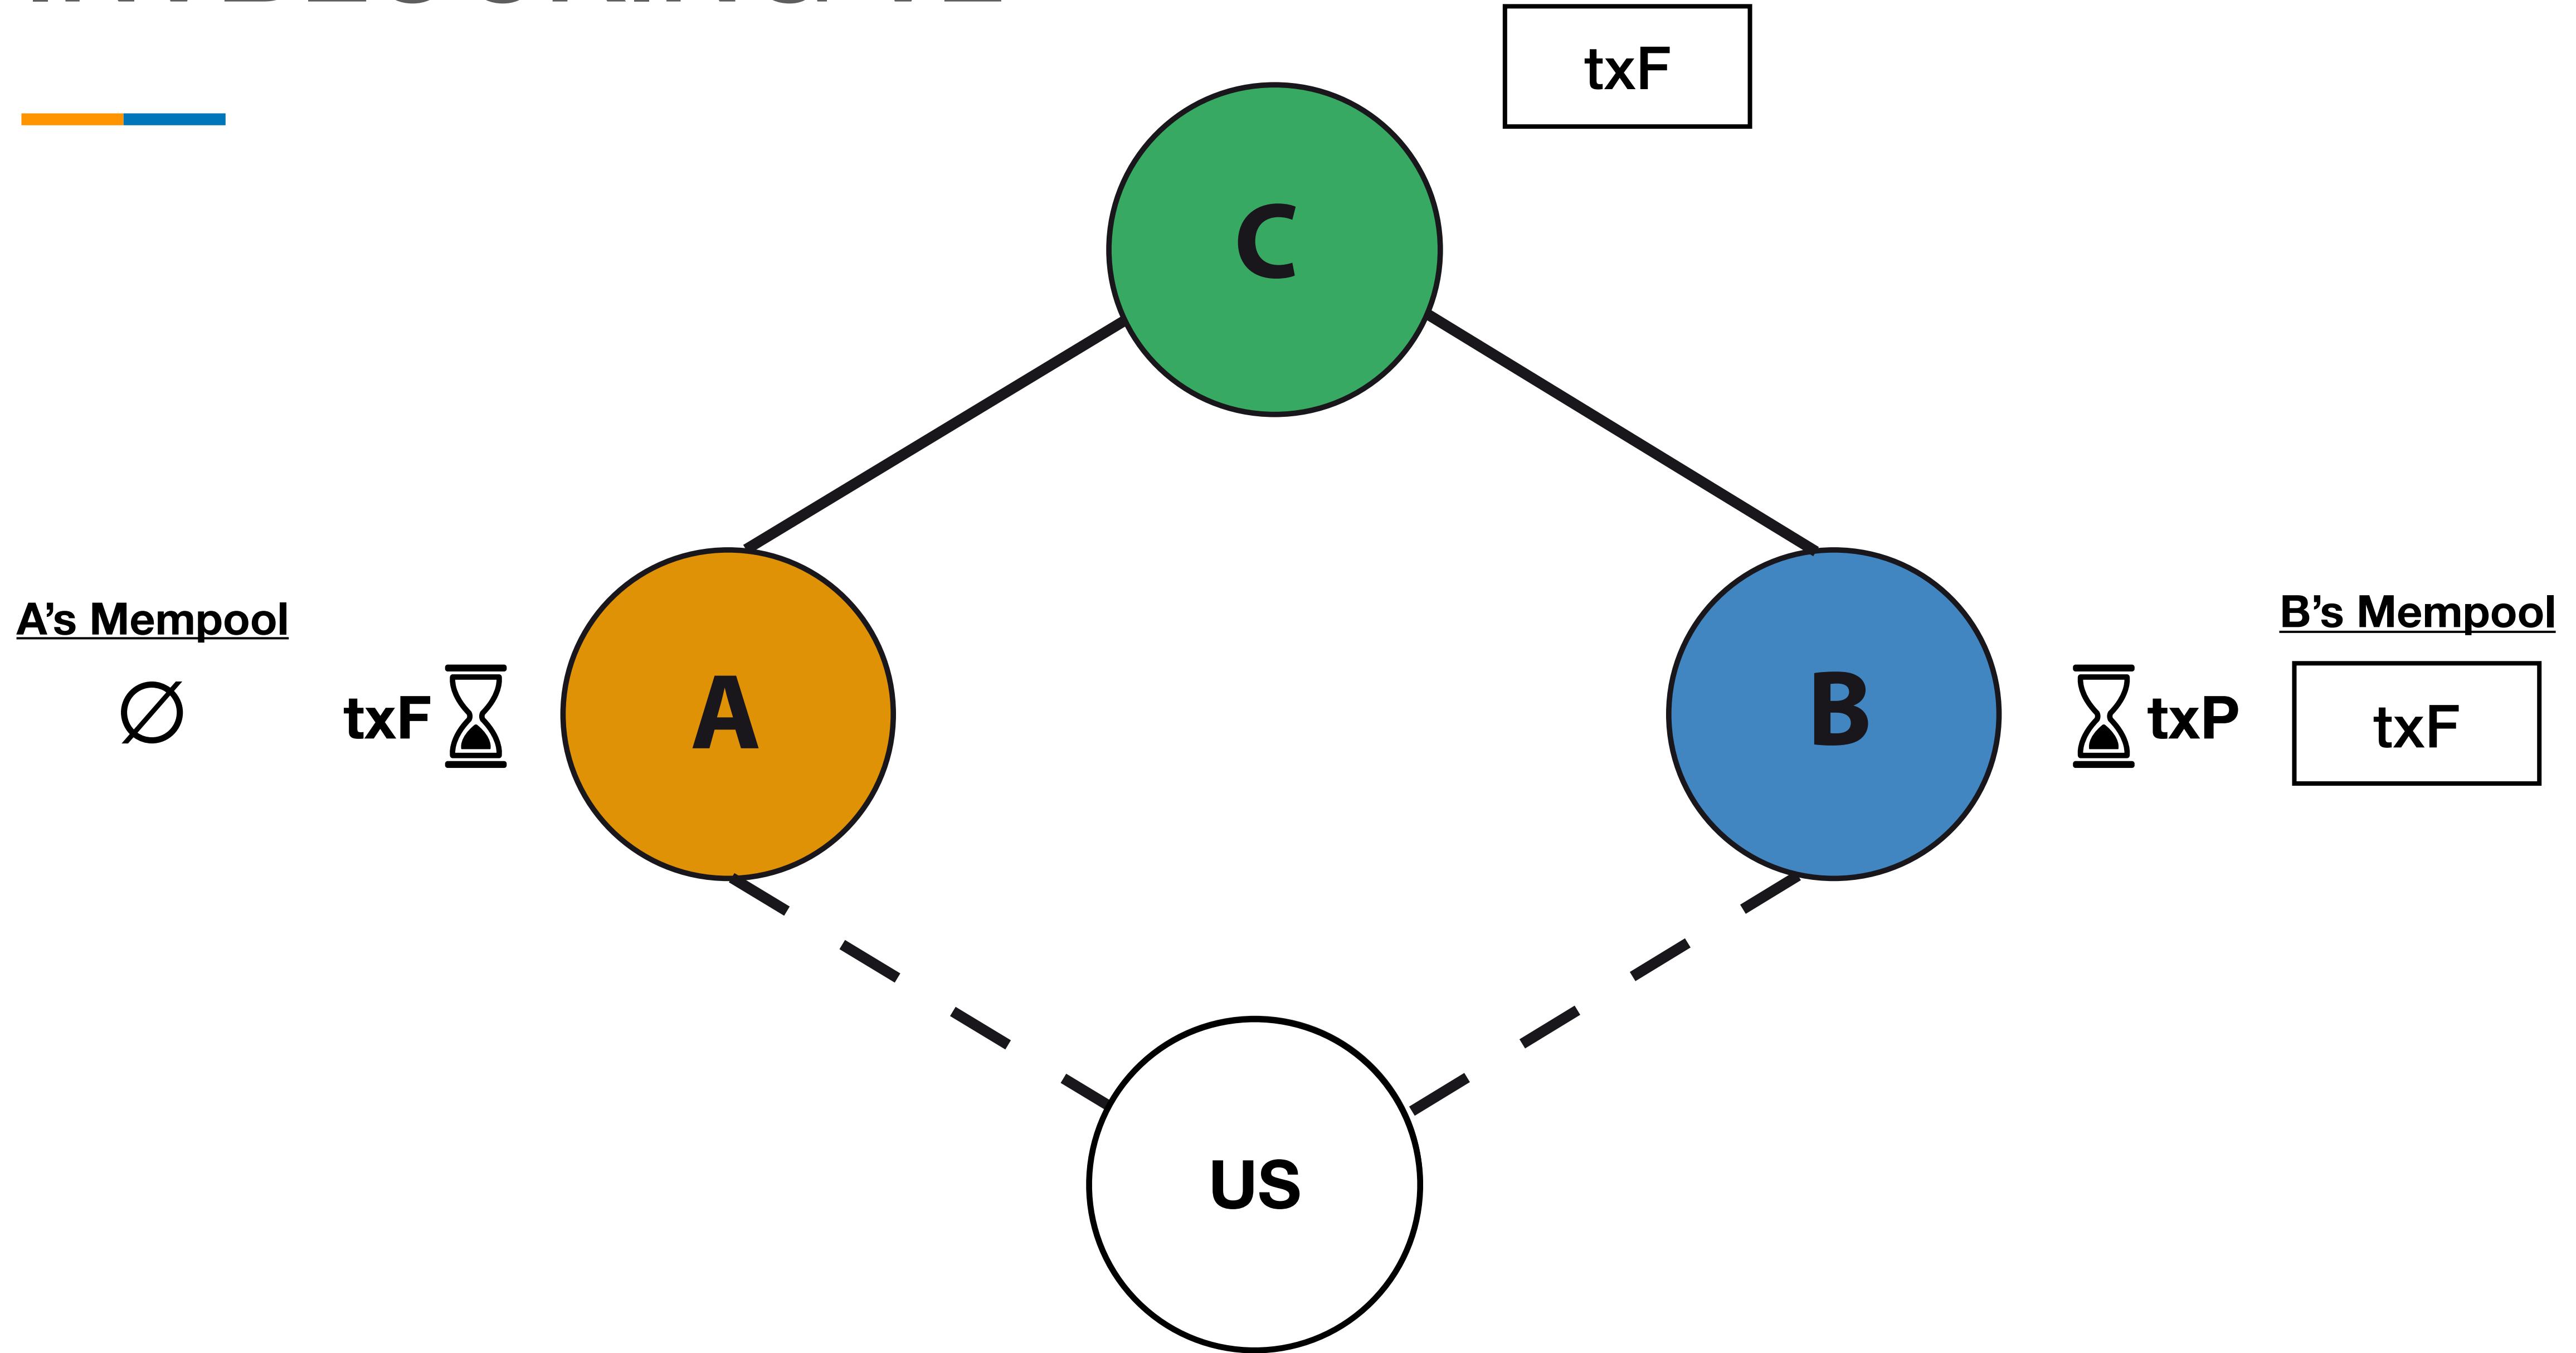
INVBLOCKING V2



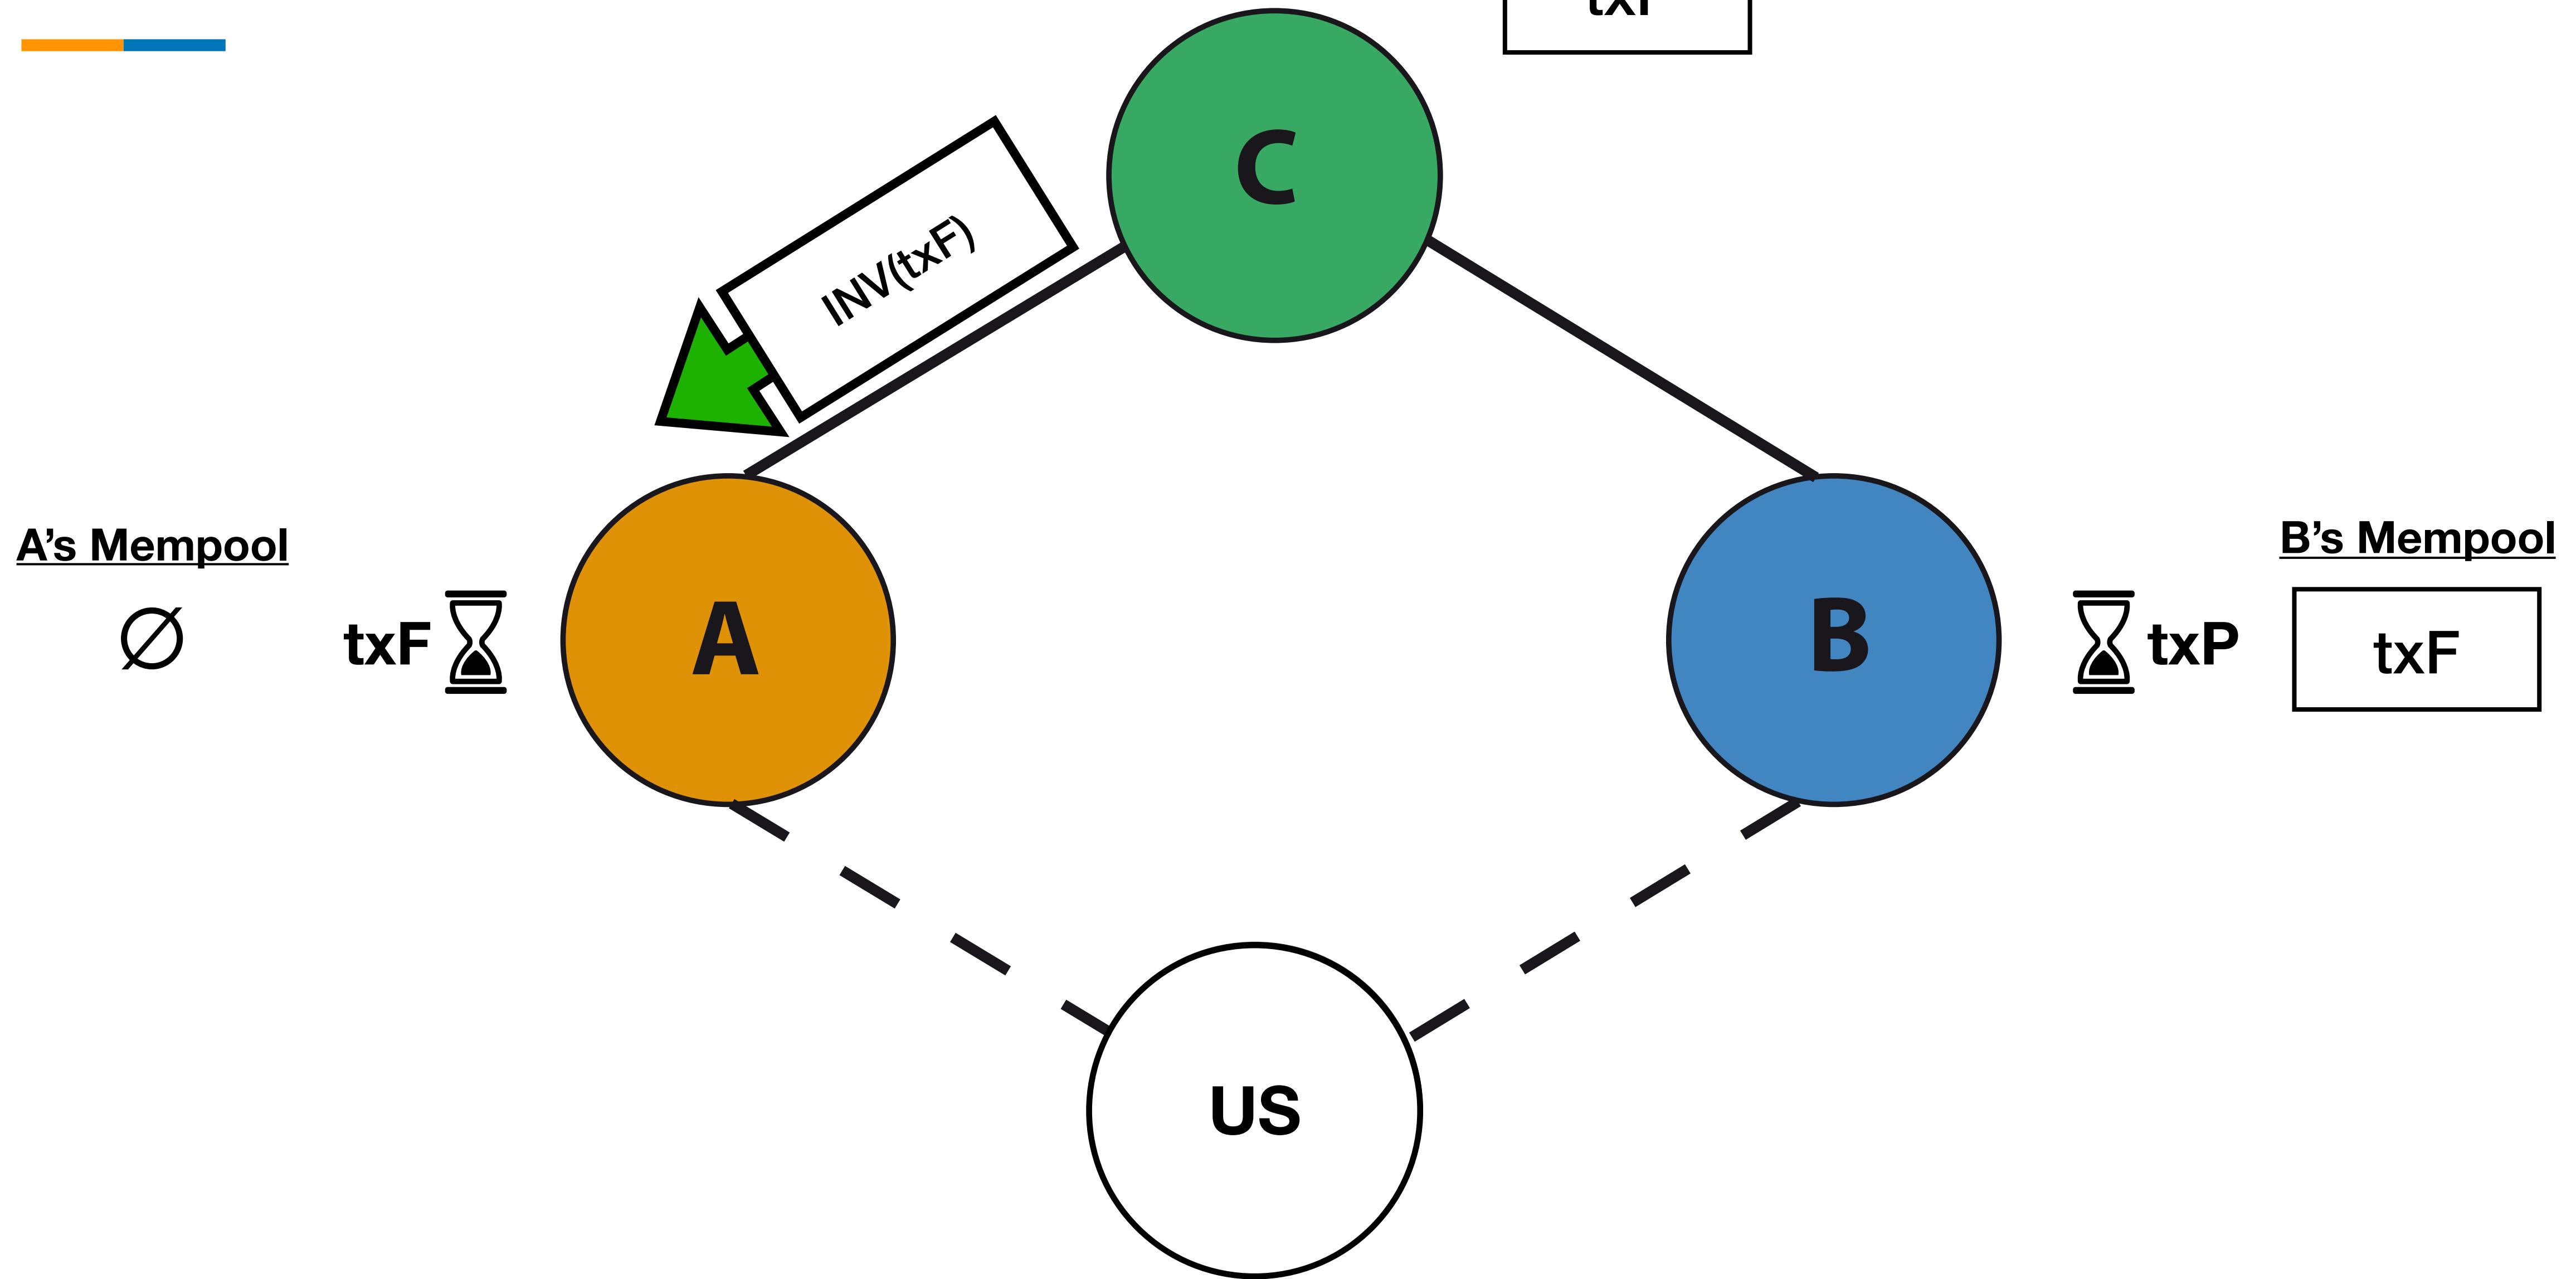
INVBLOCKING V2



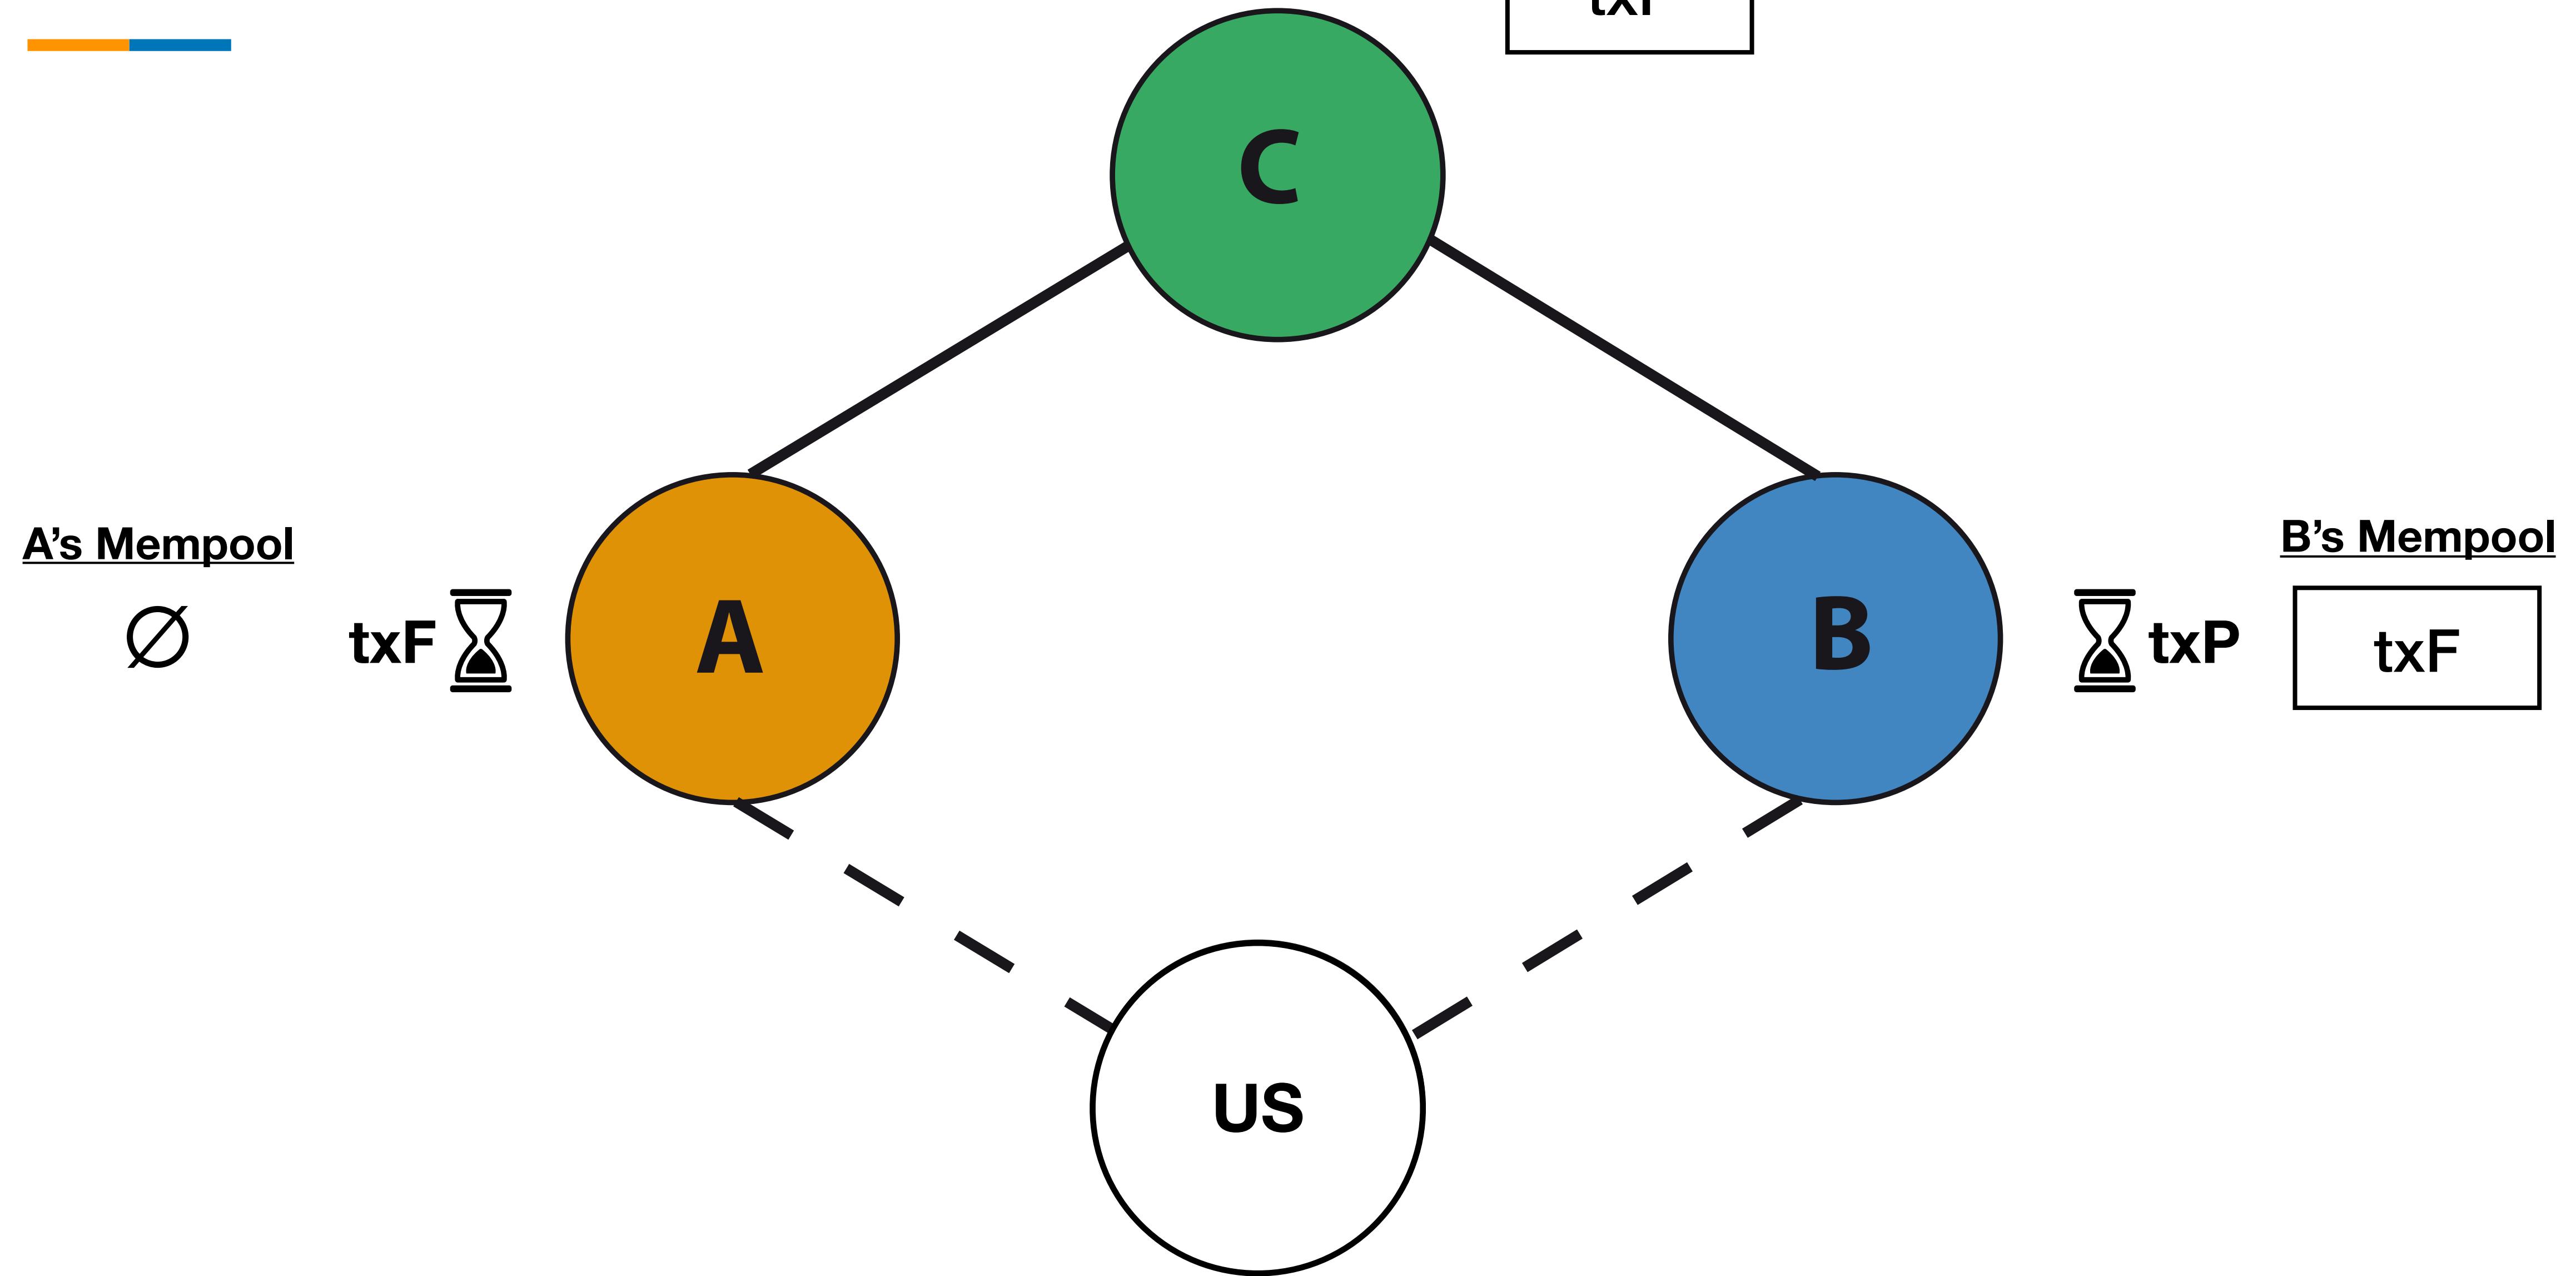
INVBLOCKING V2



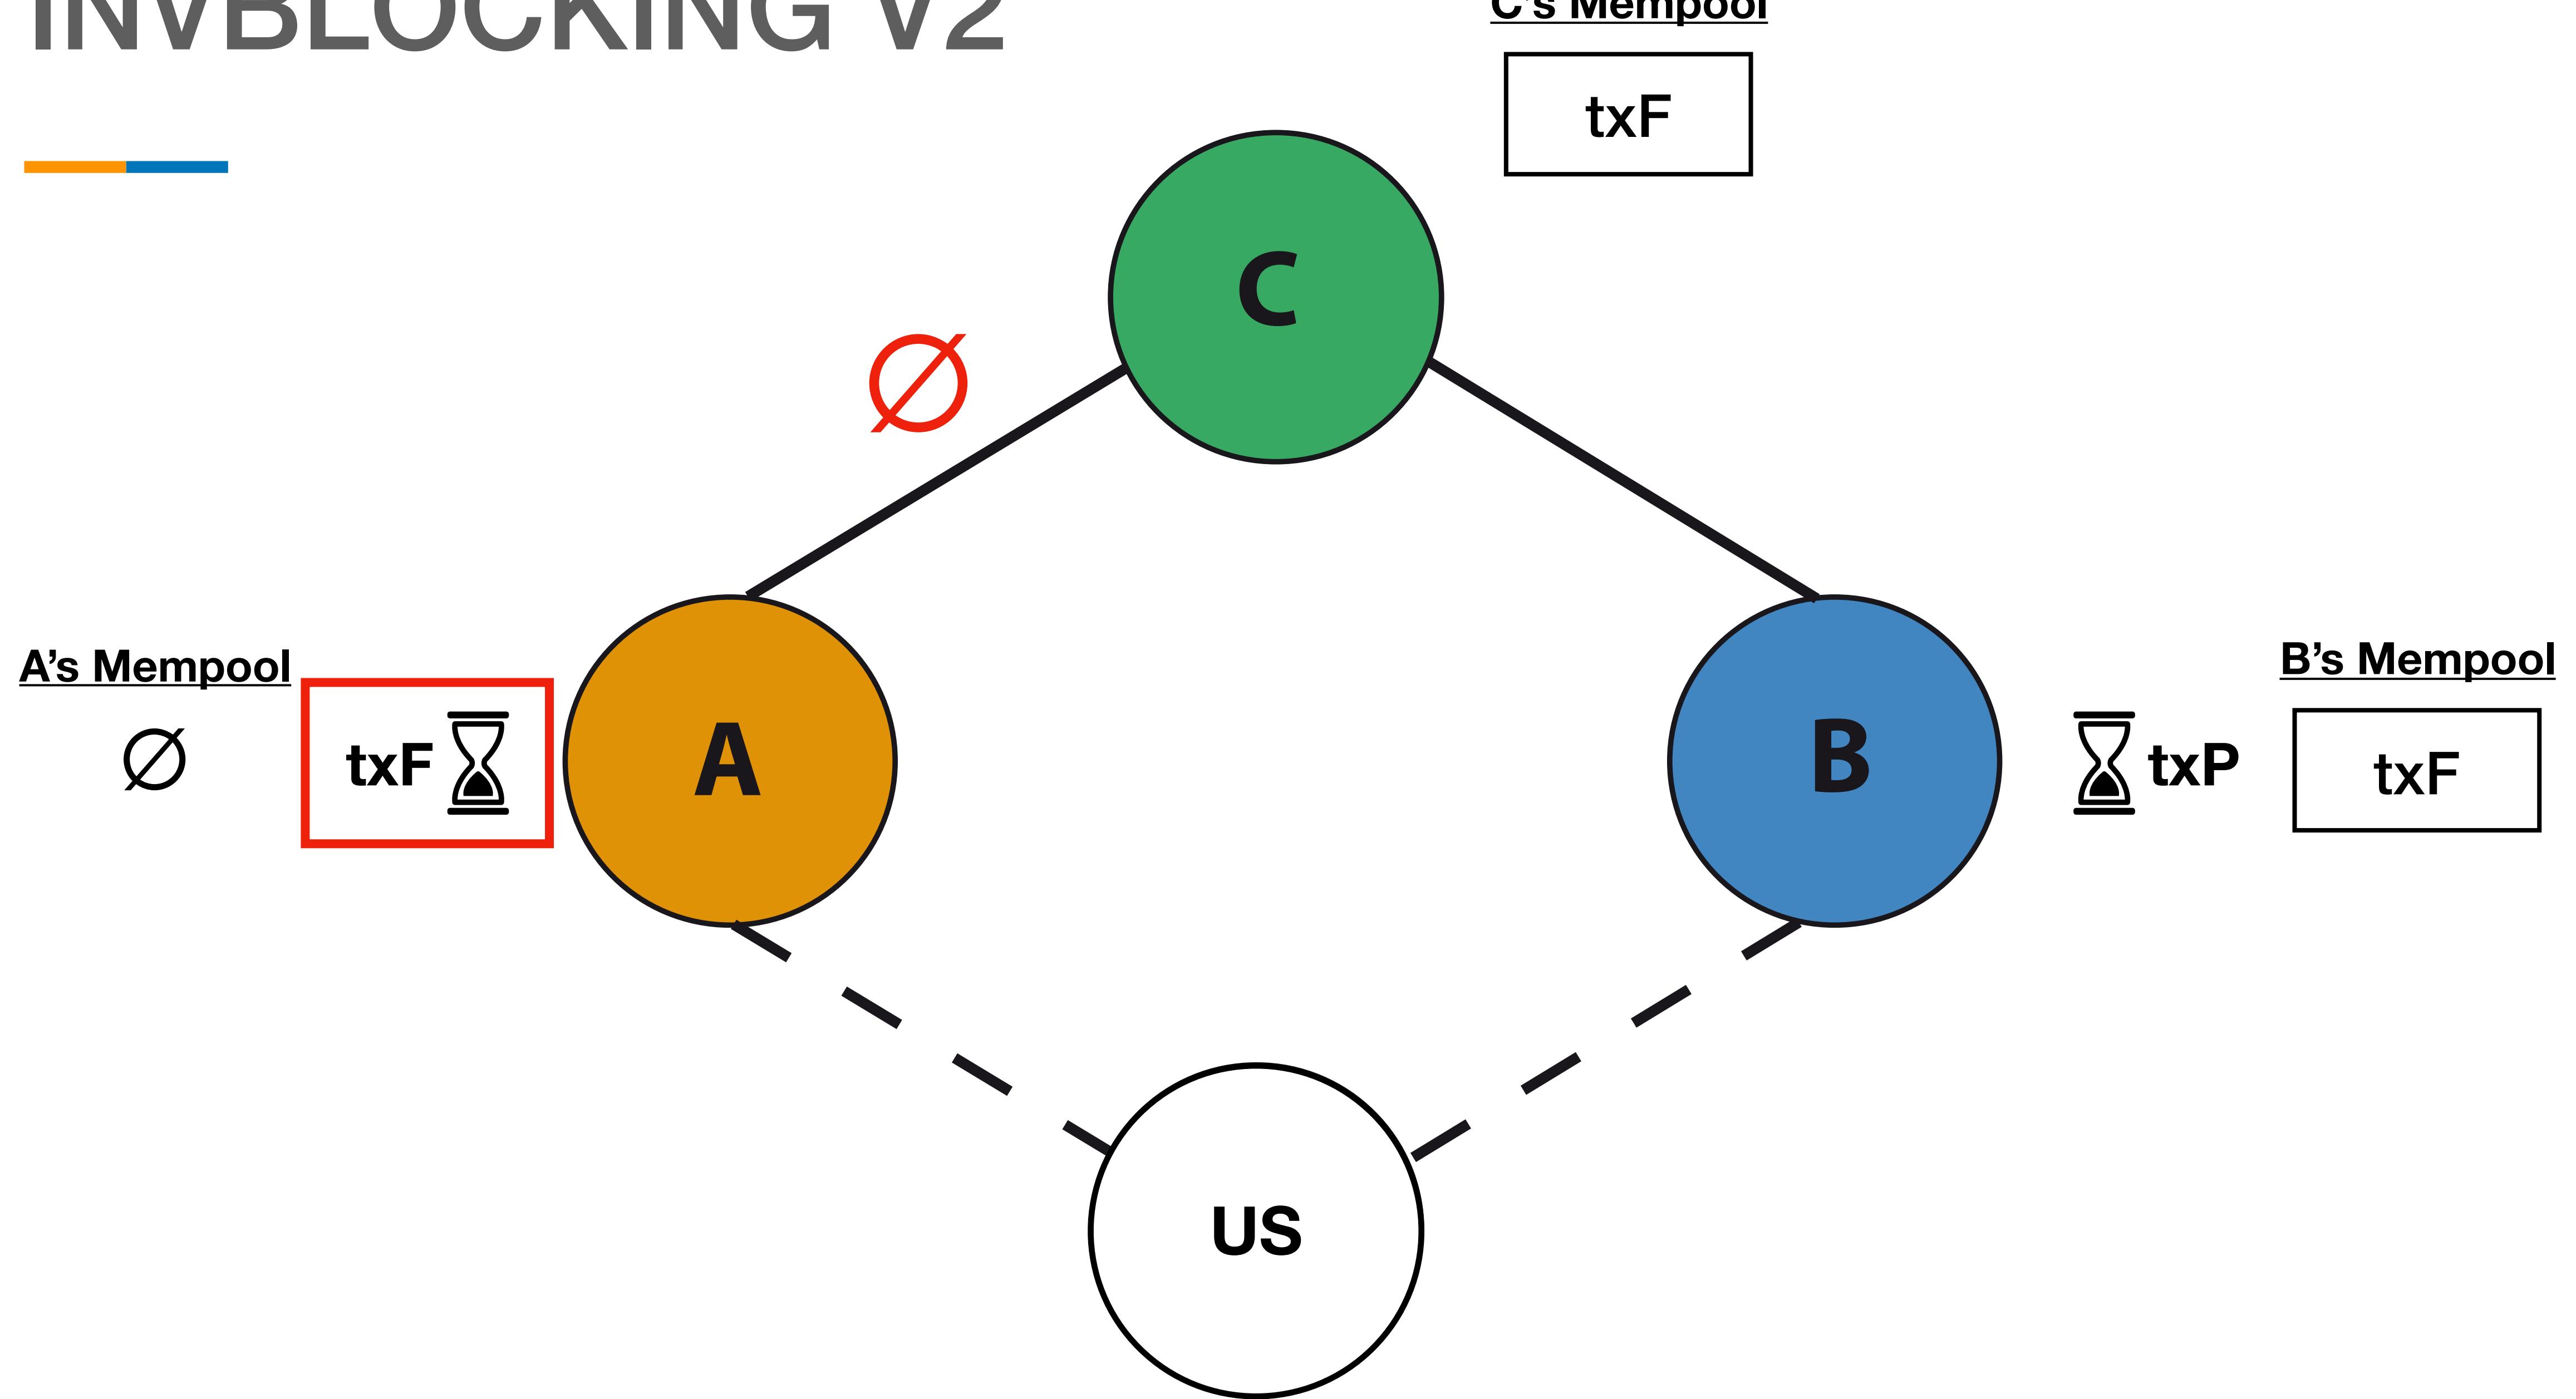
INVBLOCKING V2



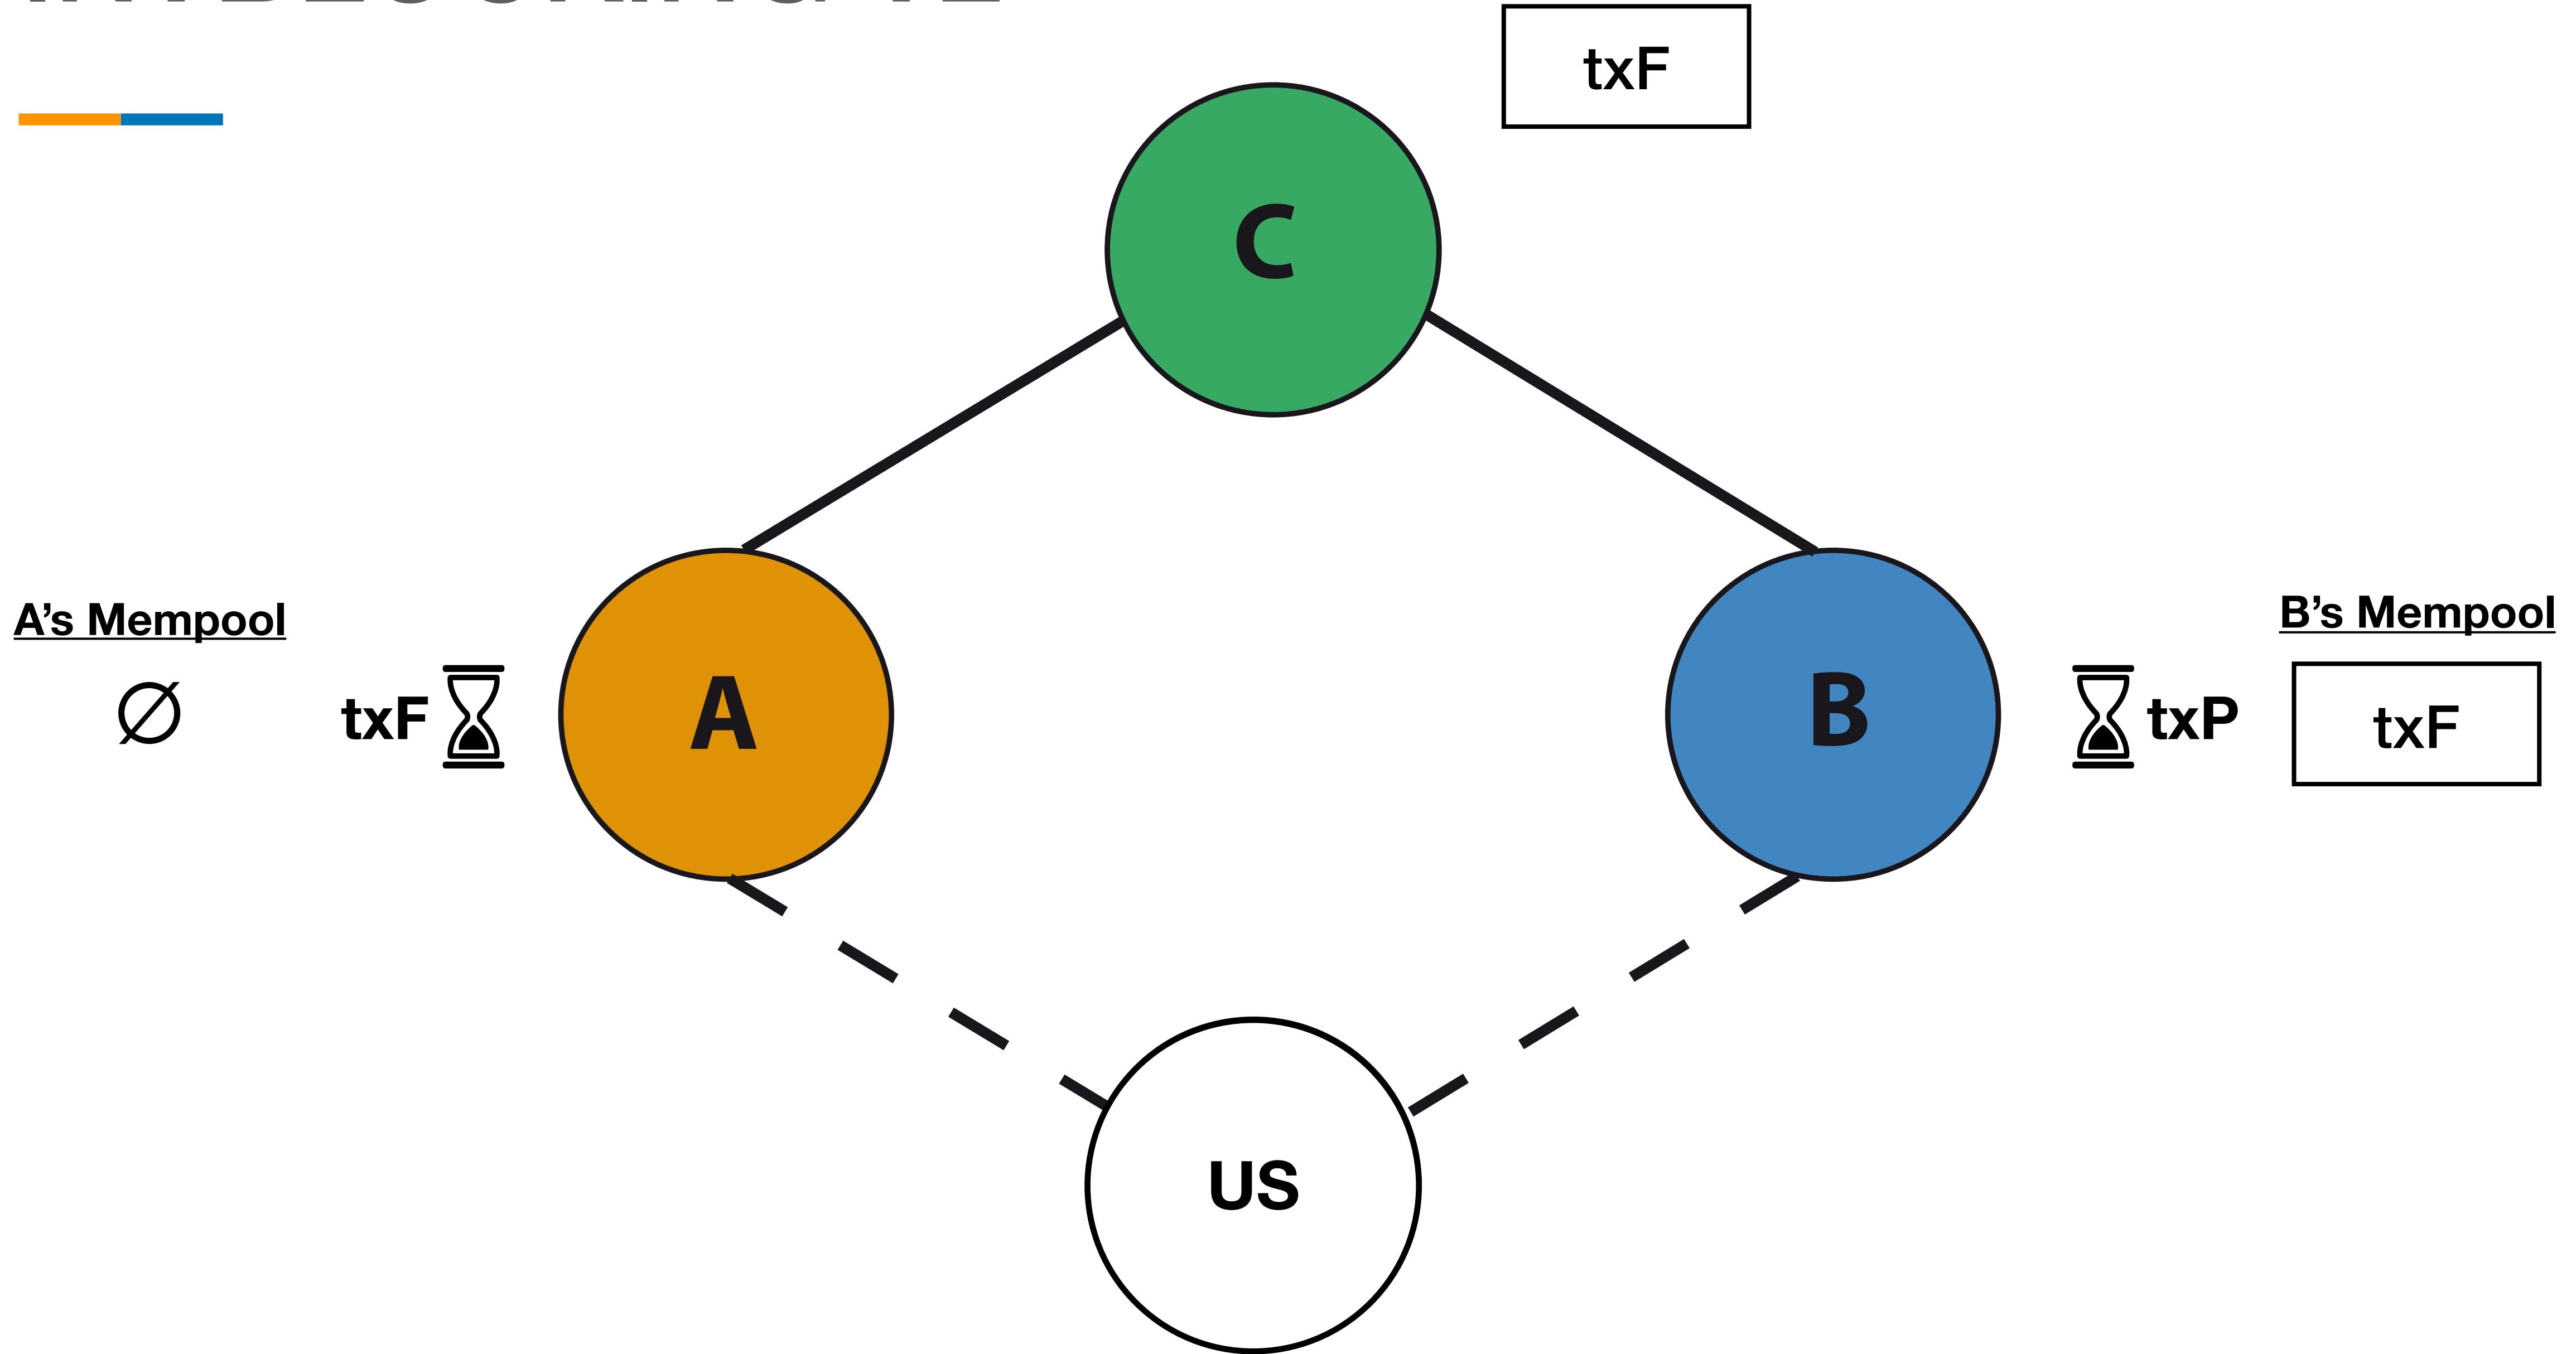
INVBLOCKING V2



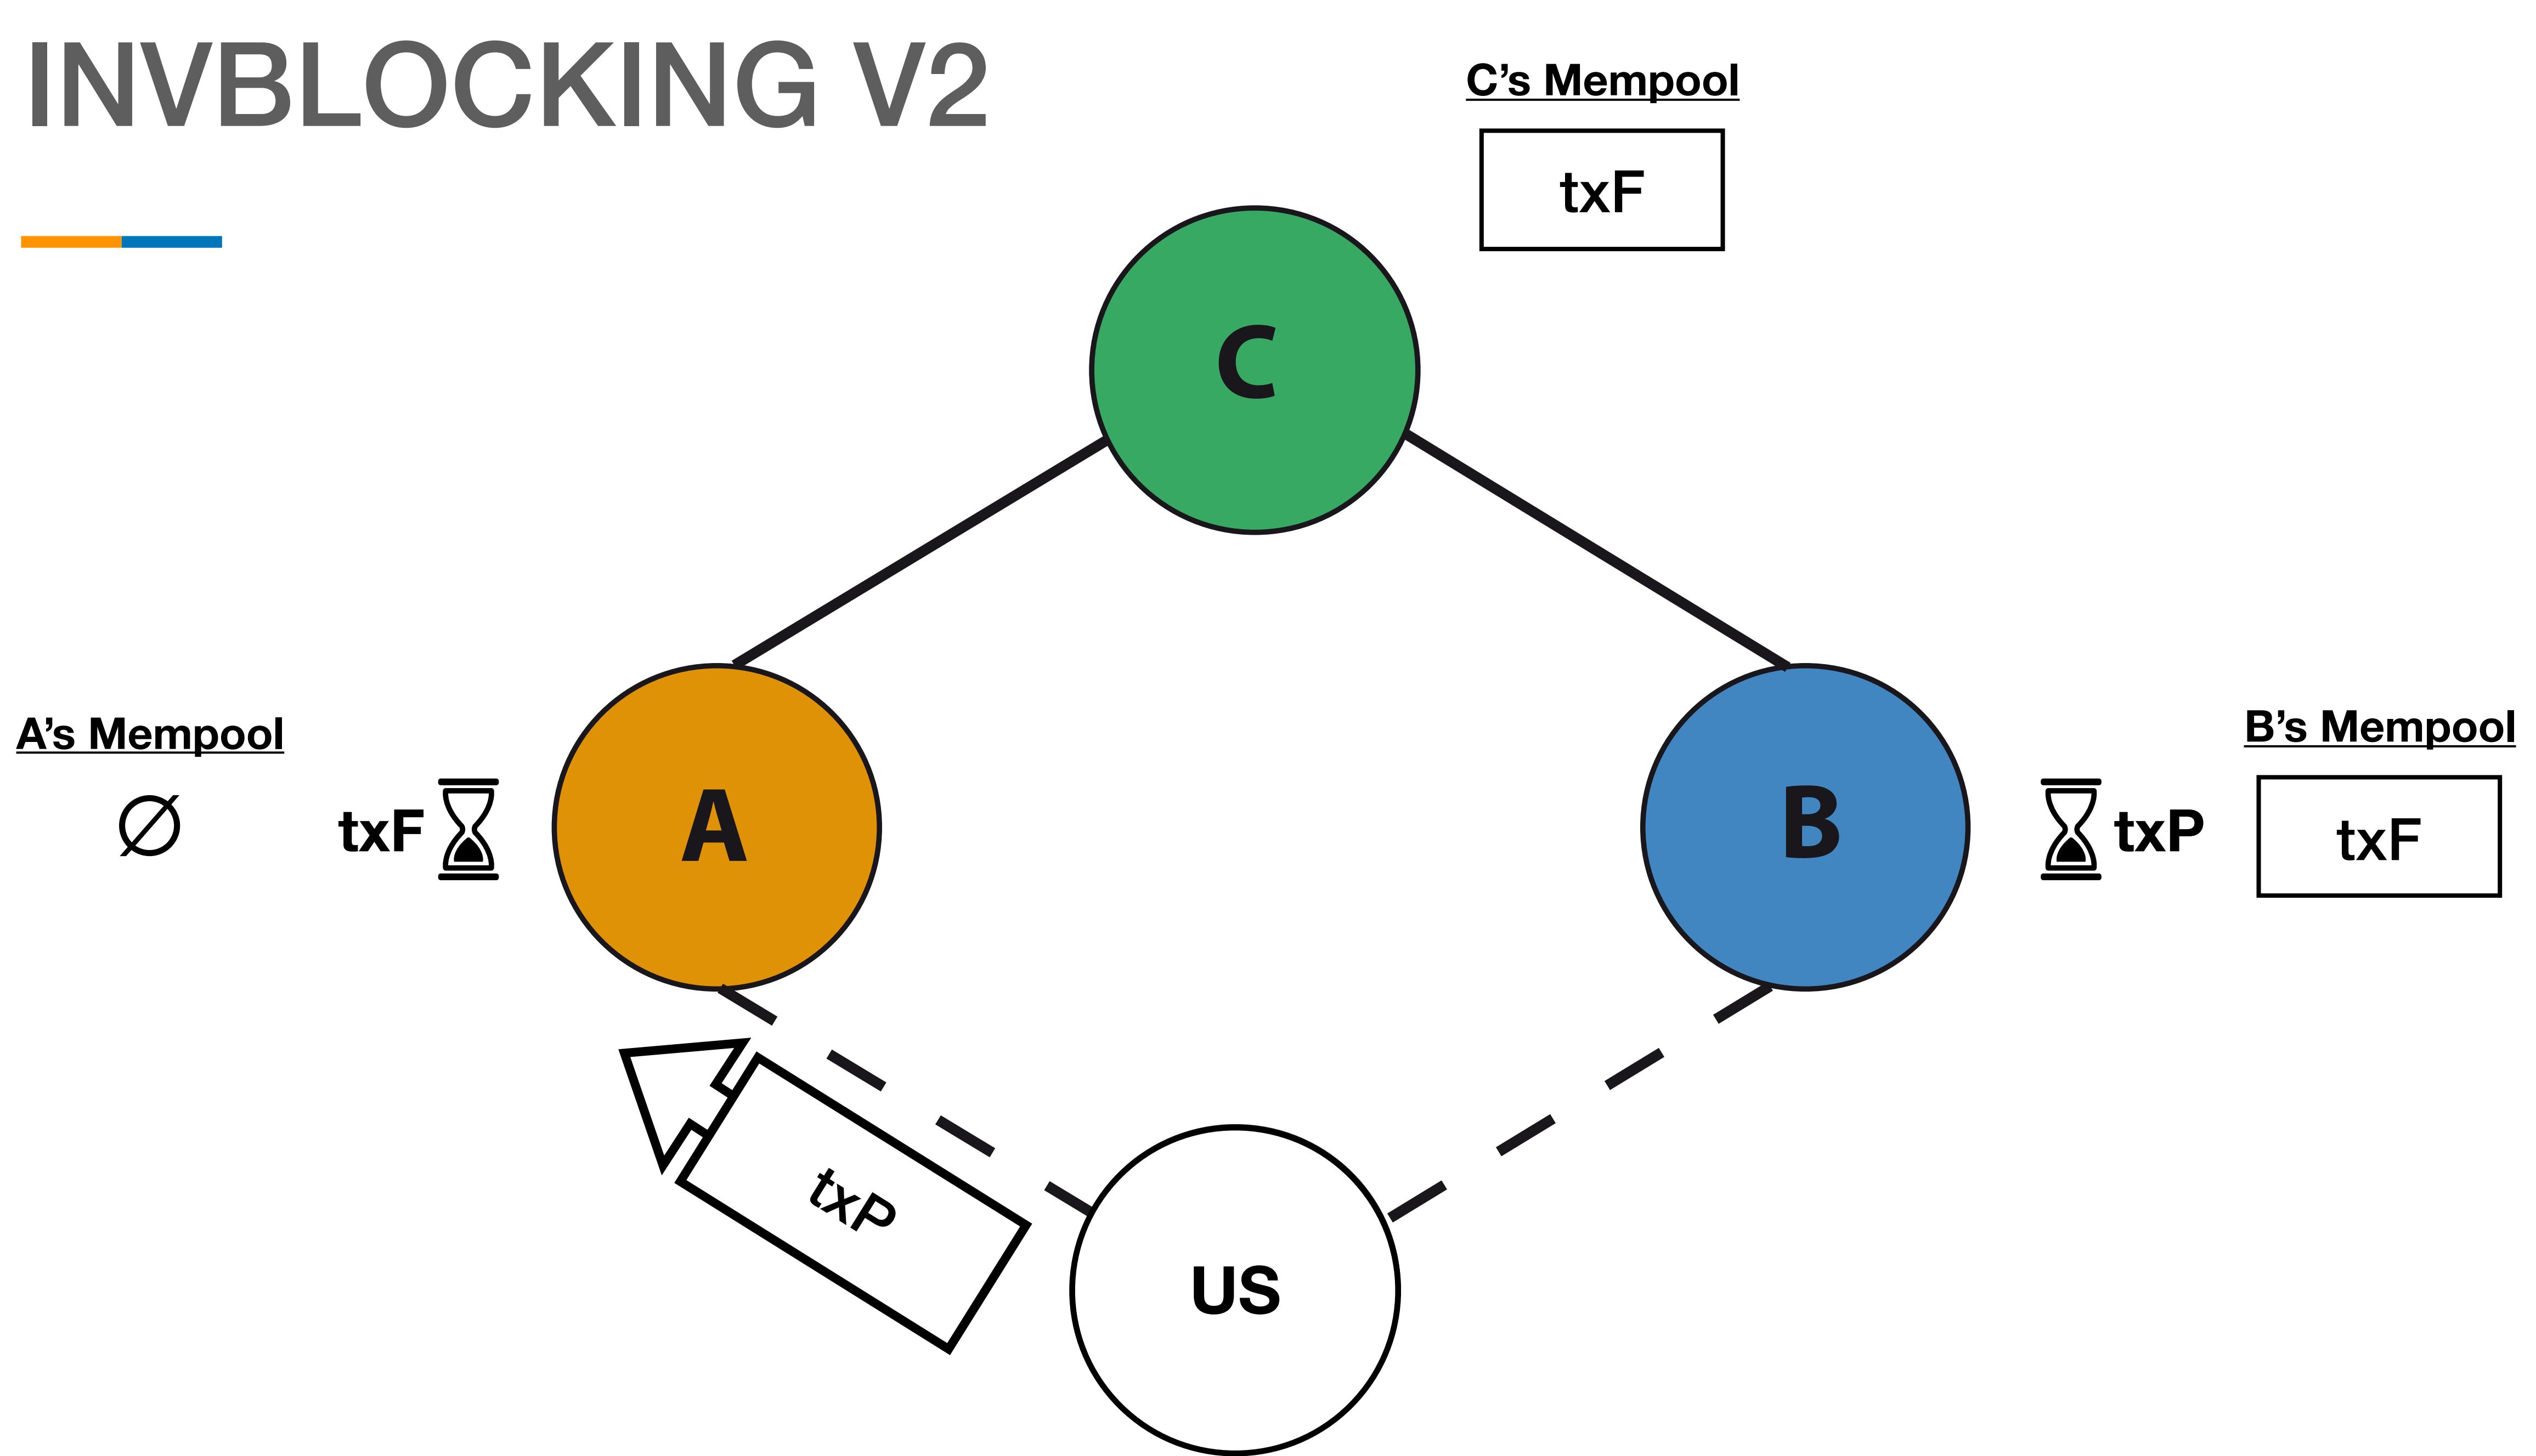
INVBLOCKING V2



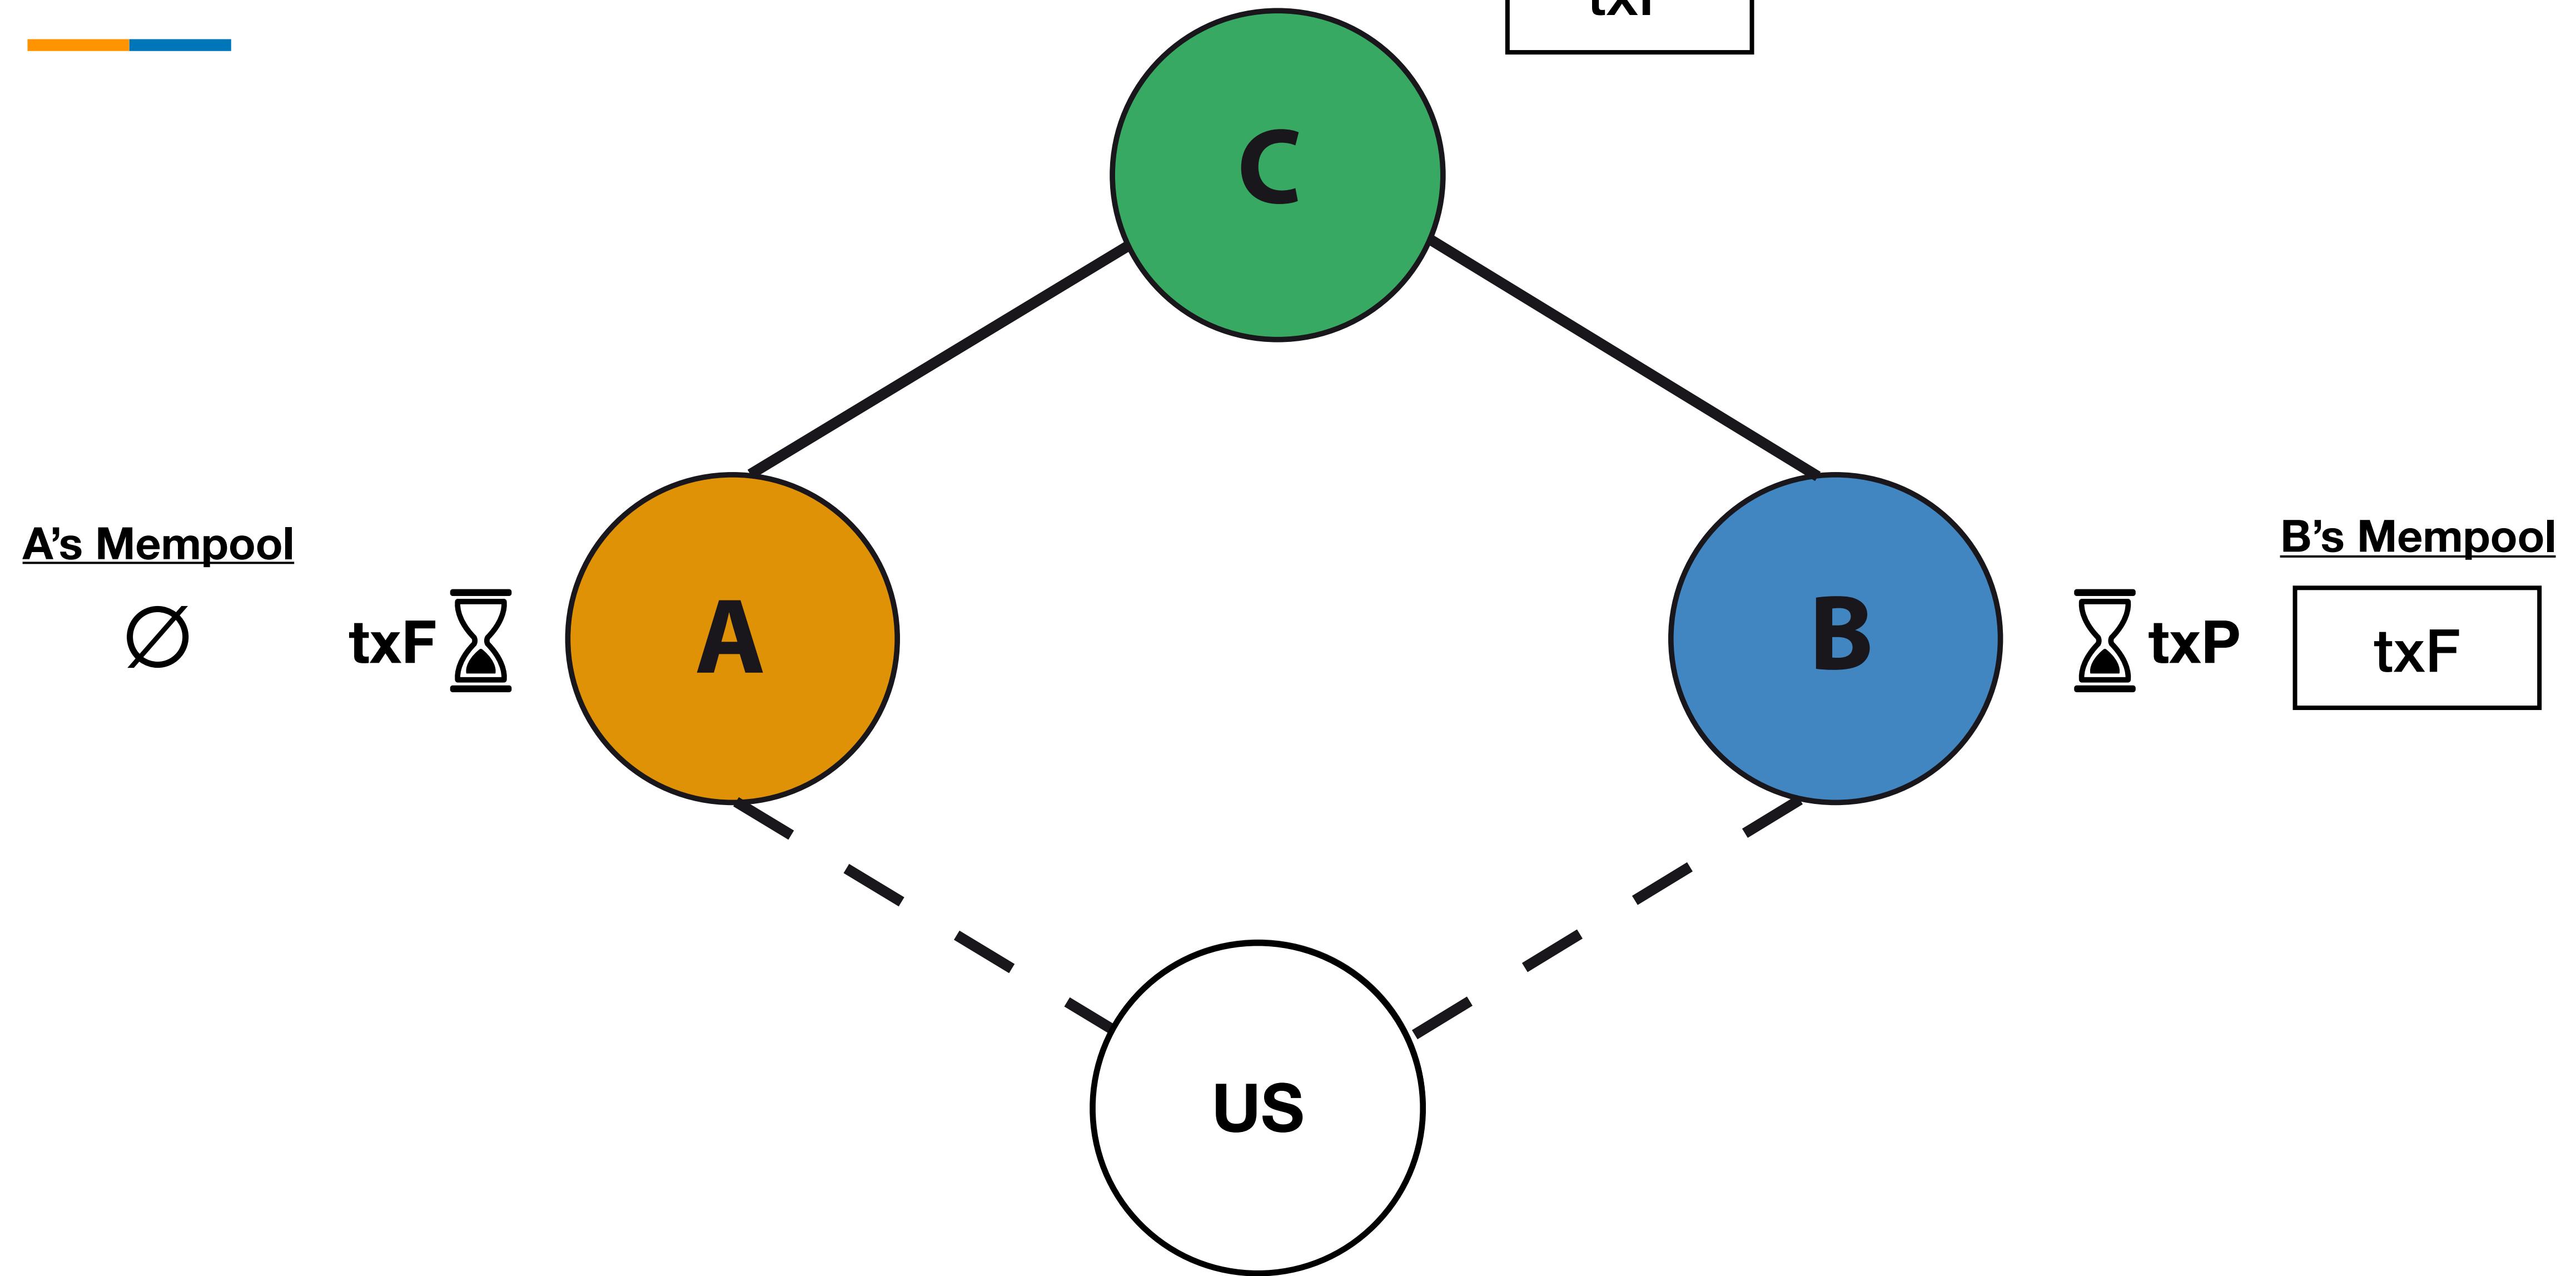
INVBLOCKING V2



INVBLOCKING V2



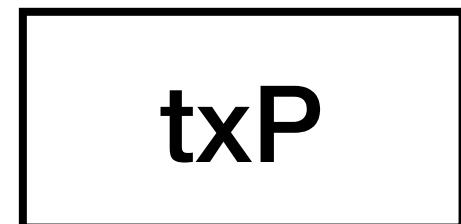
INVBLOCKING V2



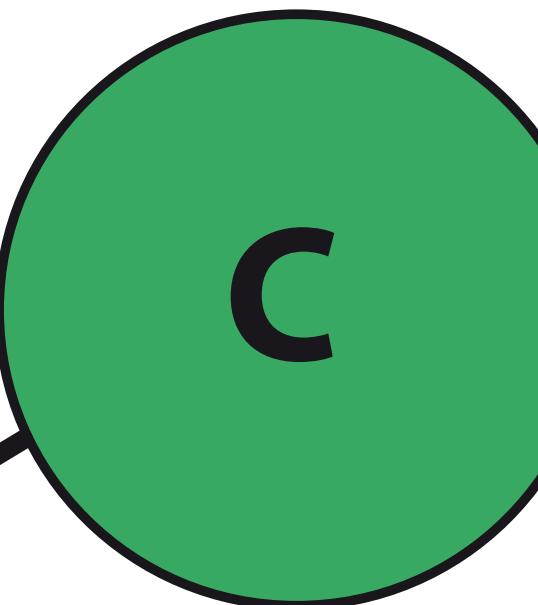
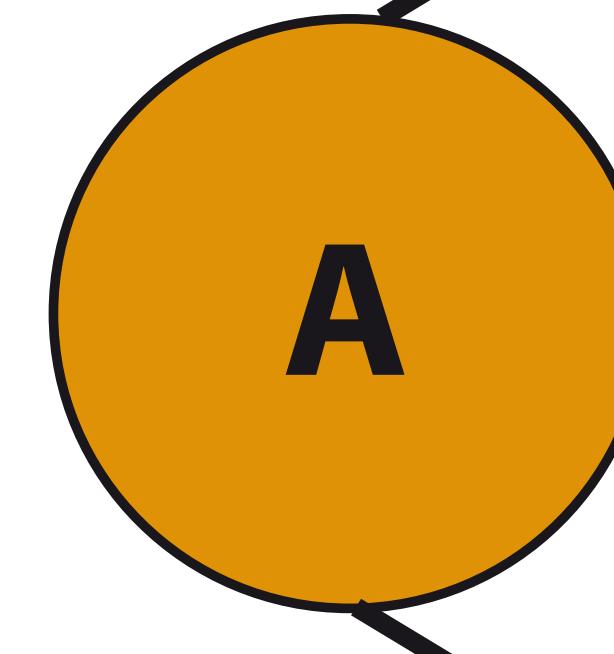
INVBLOCKING V2



A's Mempool



txF



C's Mempool

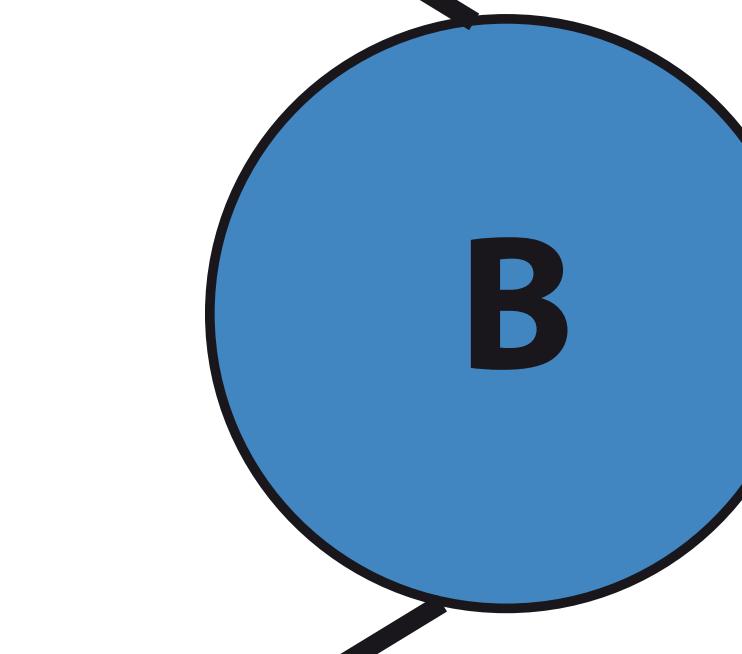
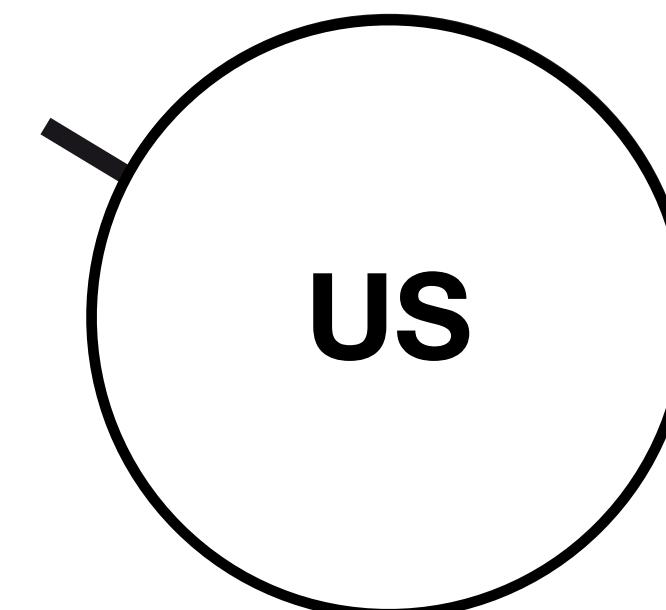


txF

B's Mempool



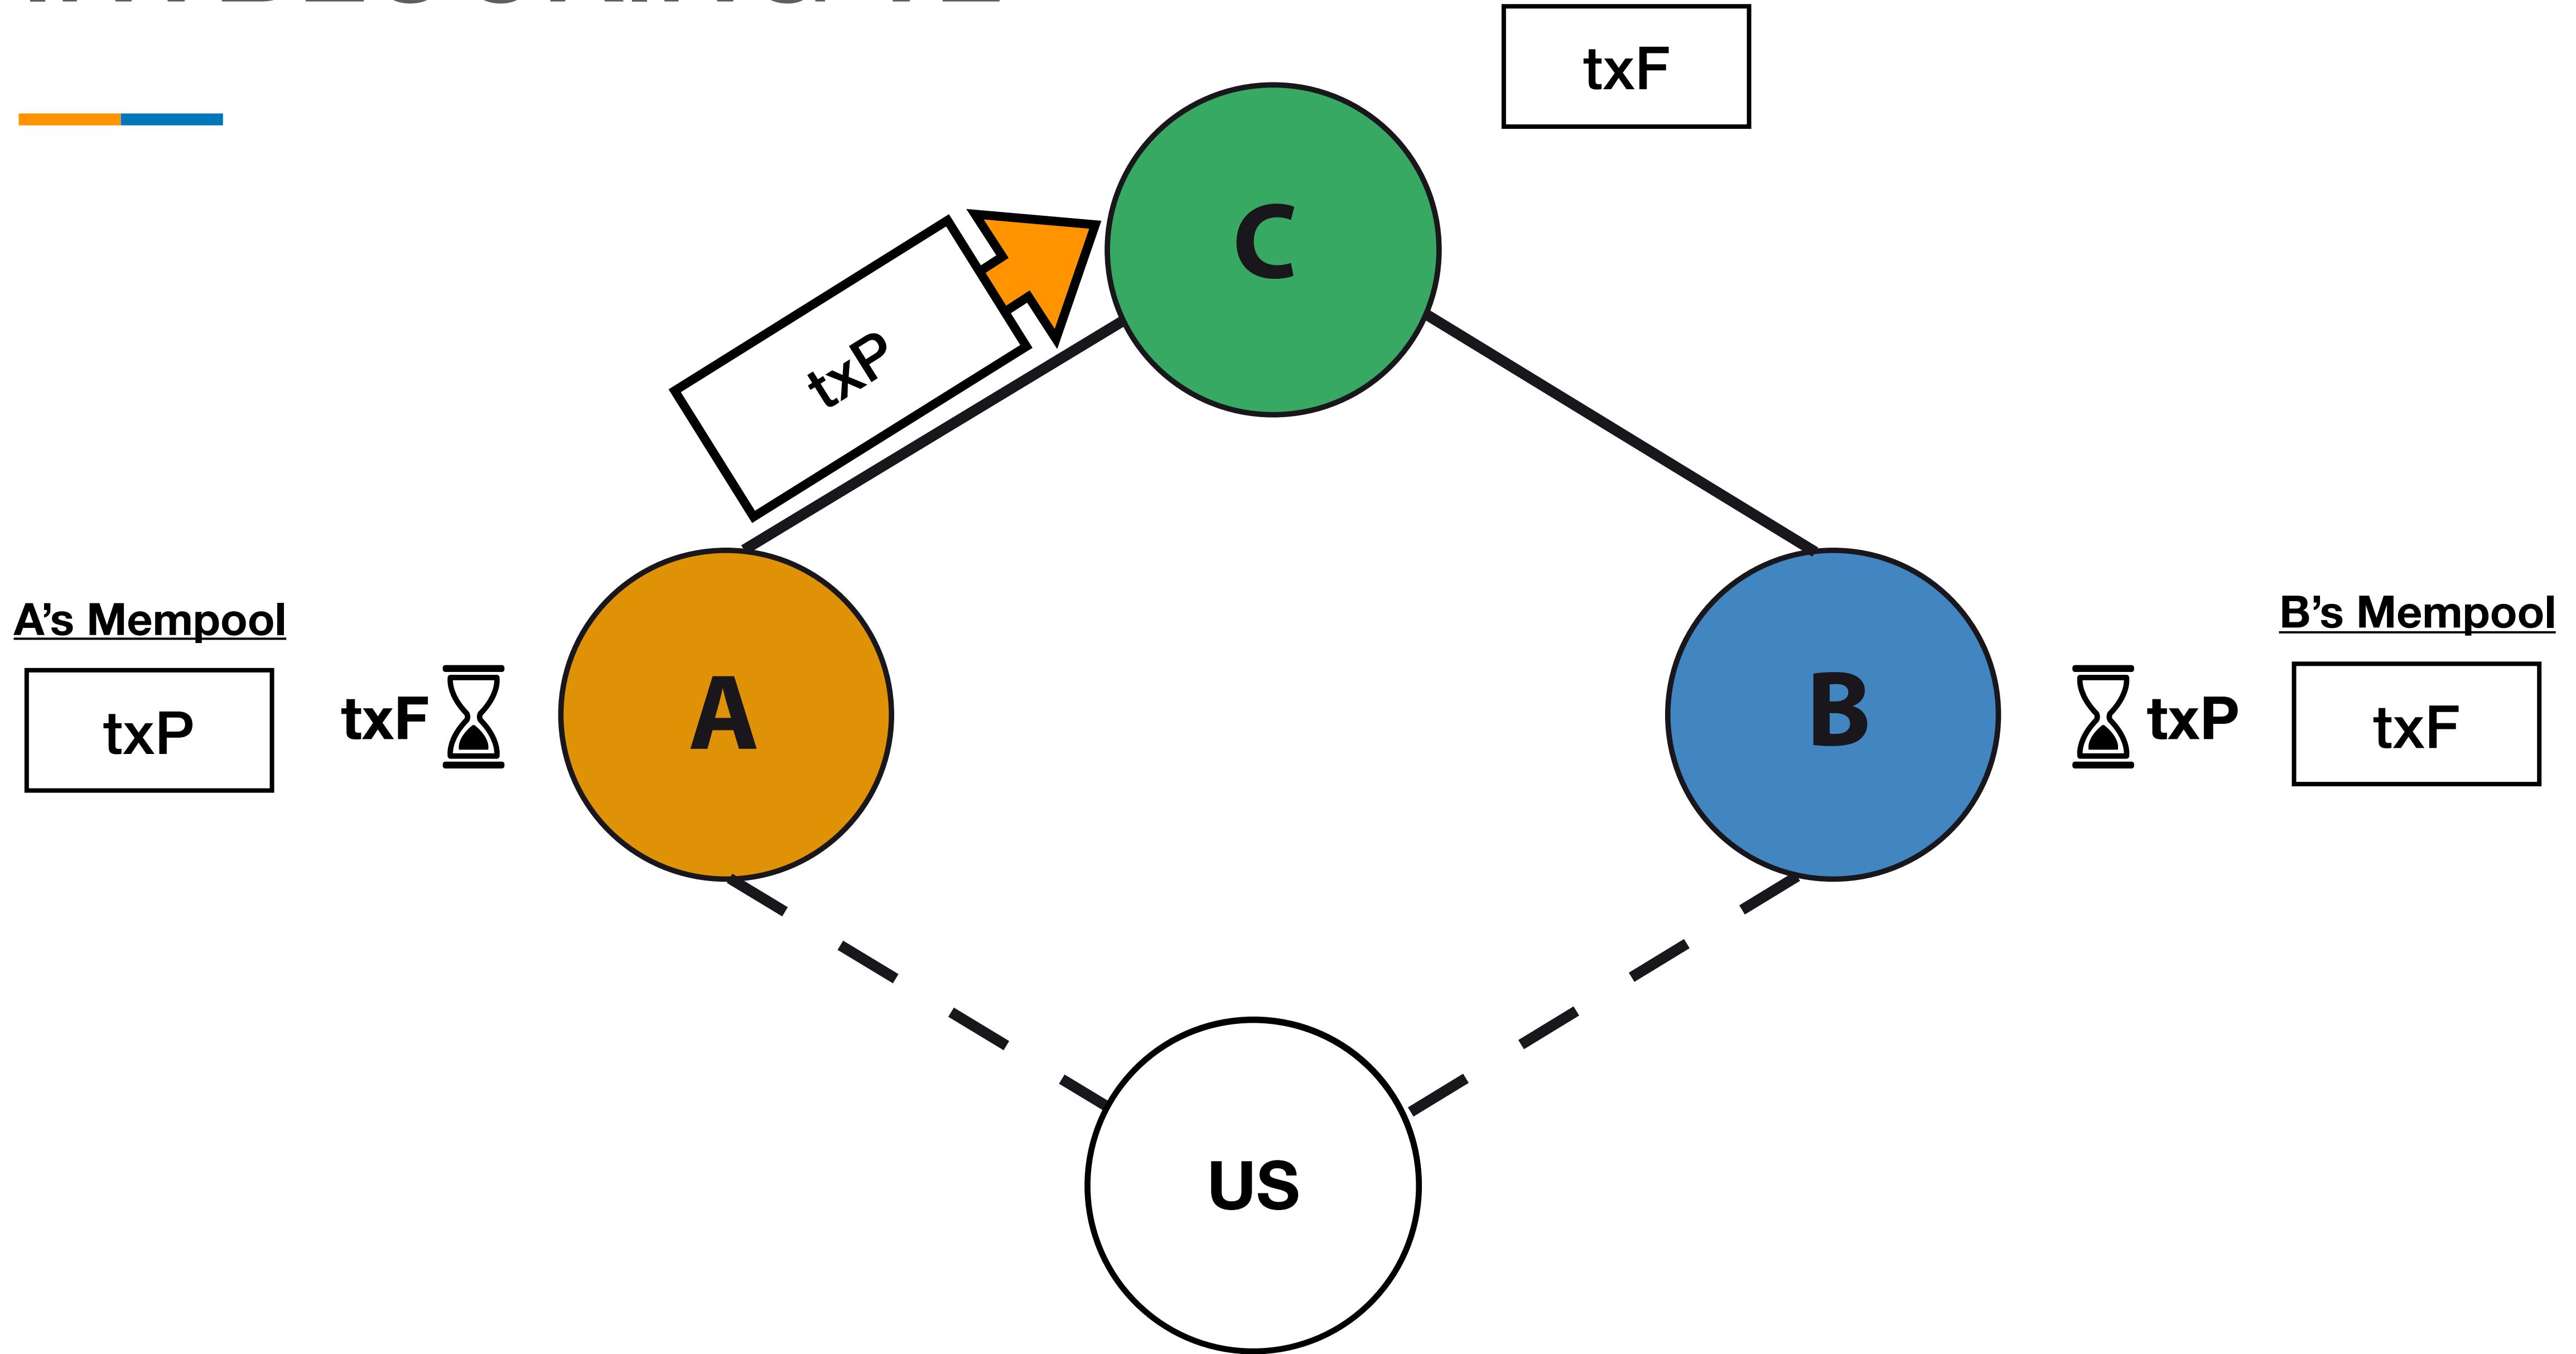
txP



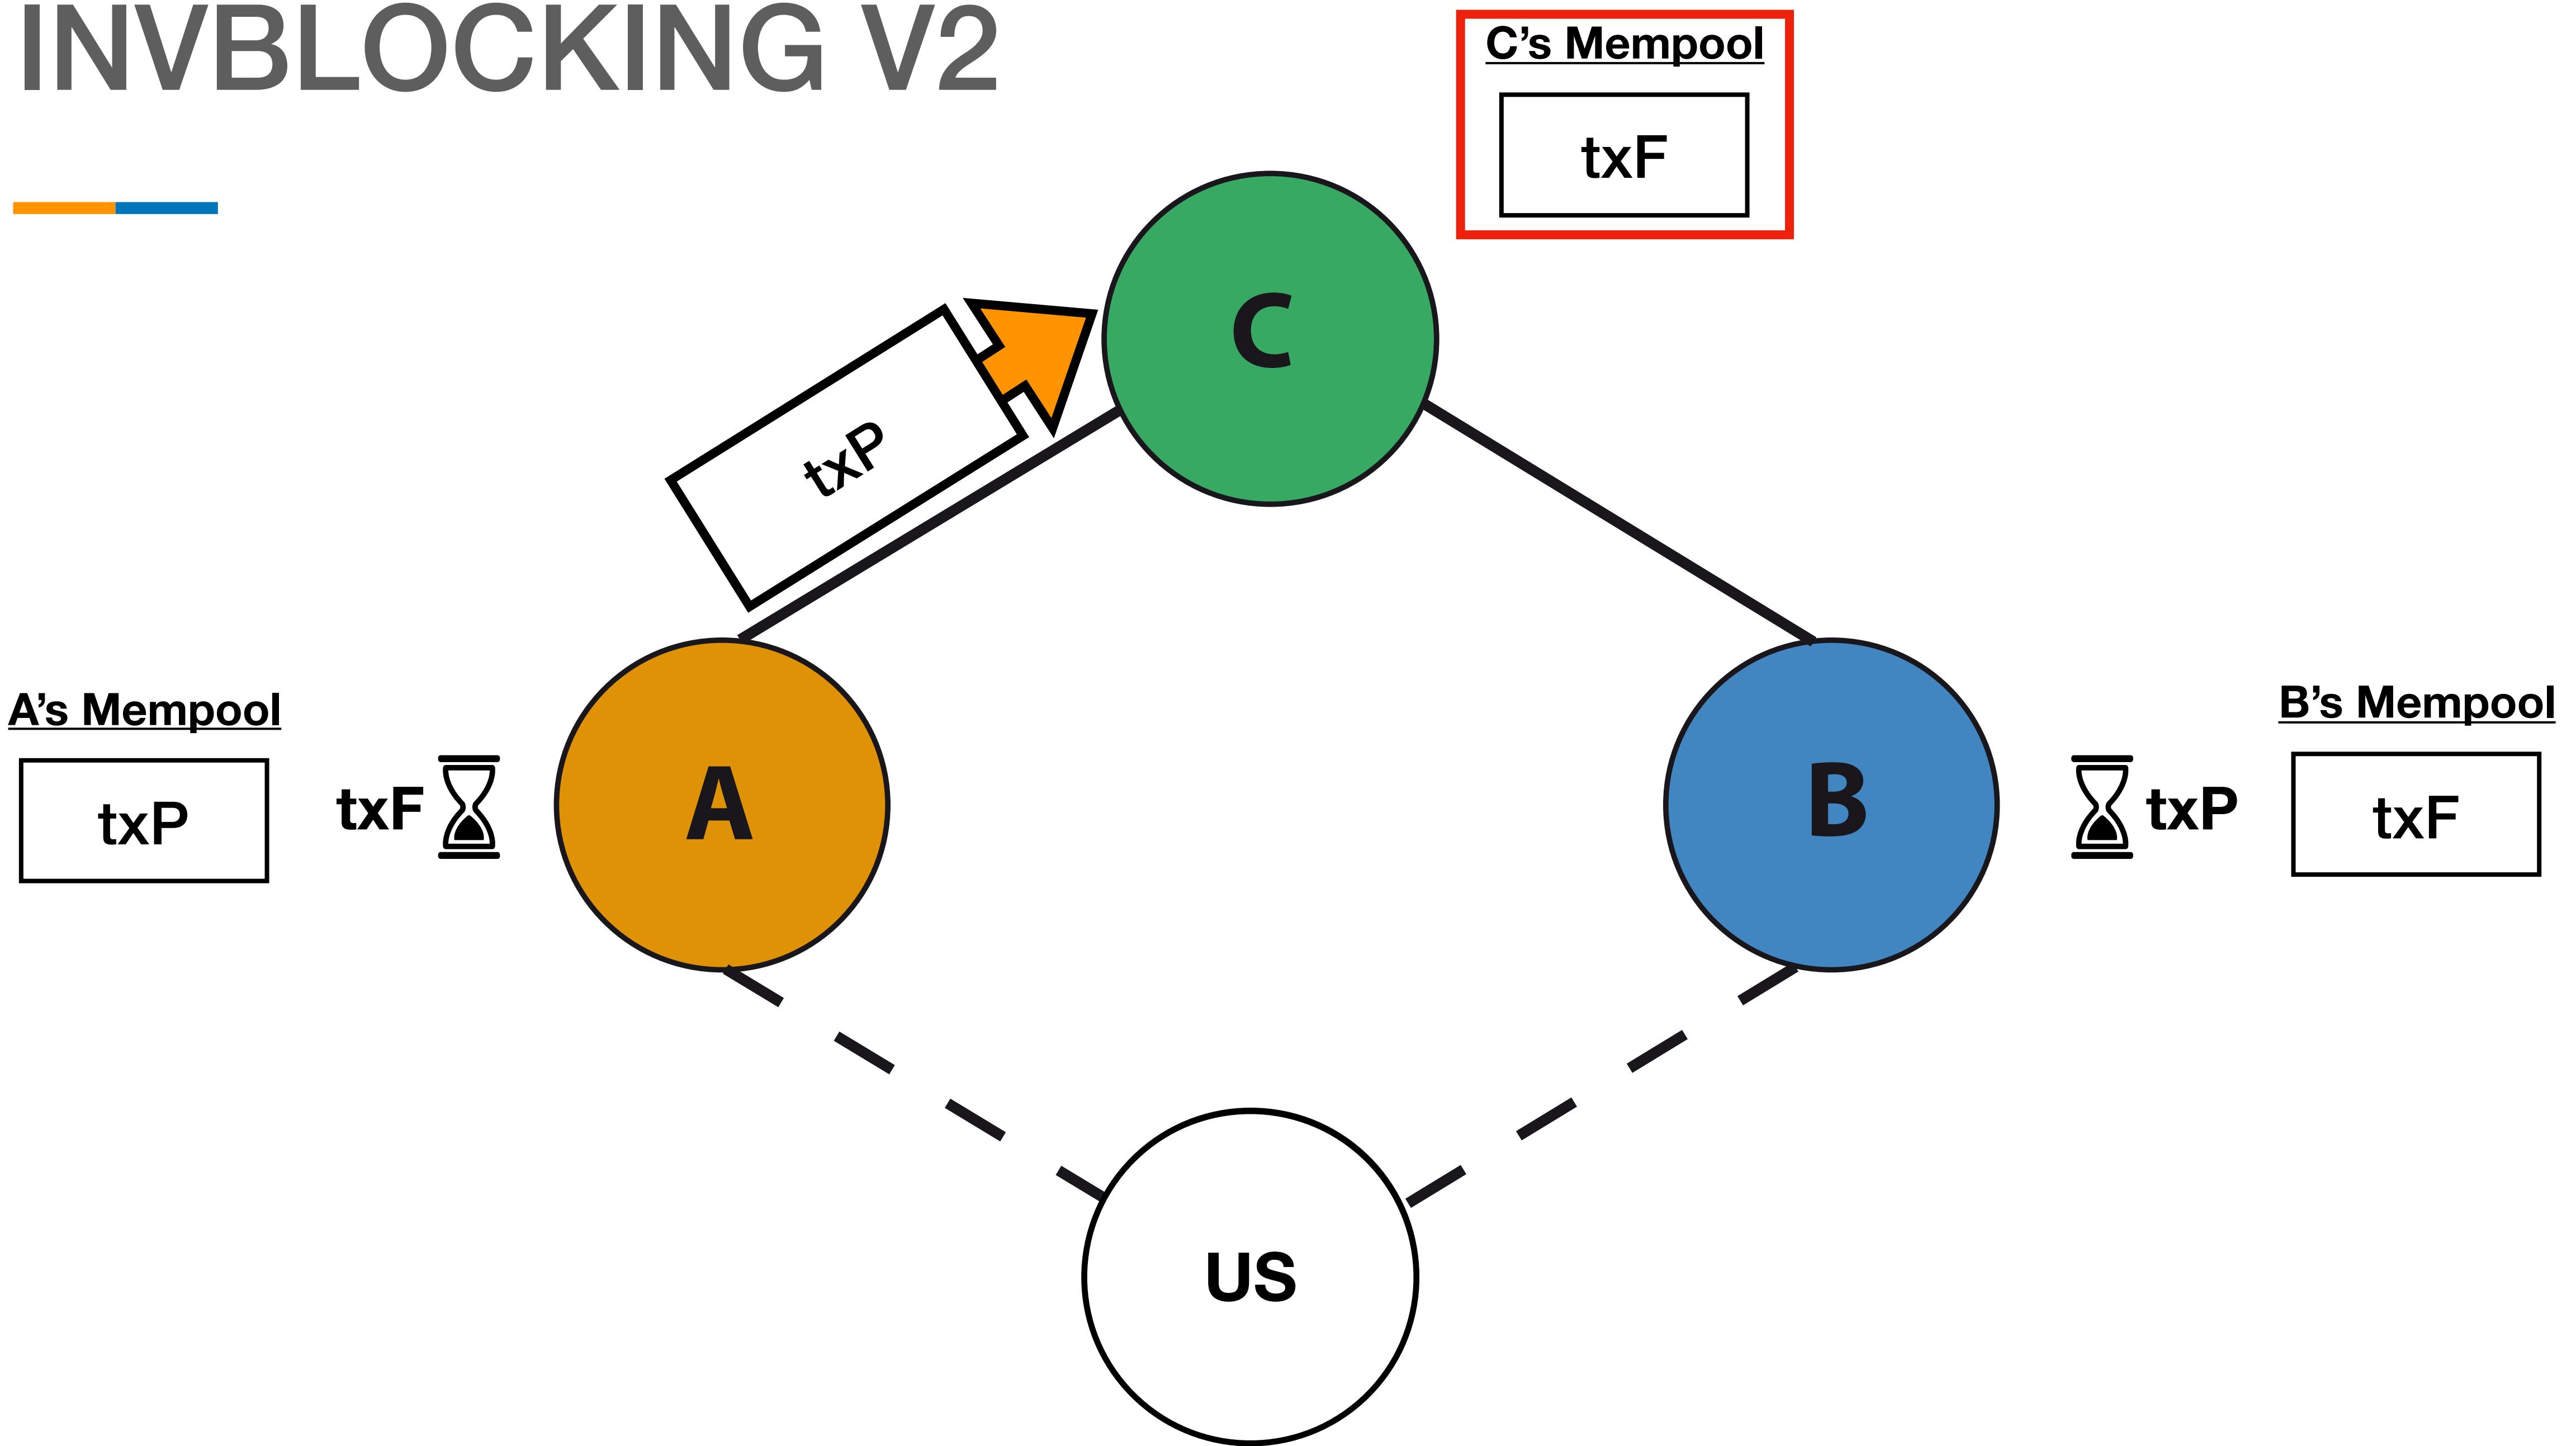
US



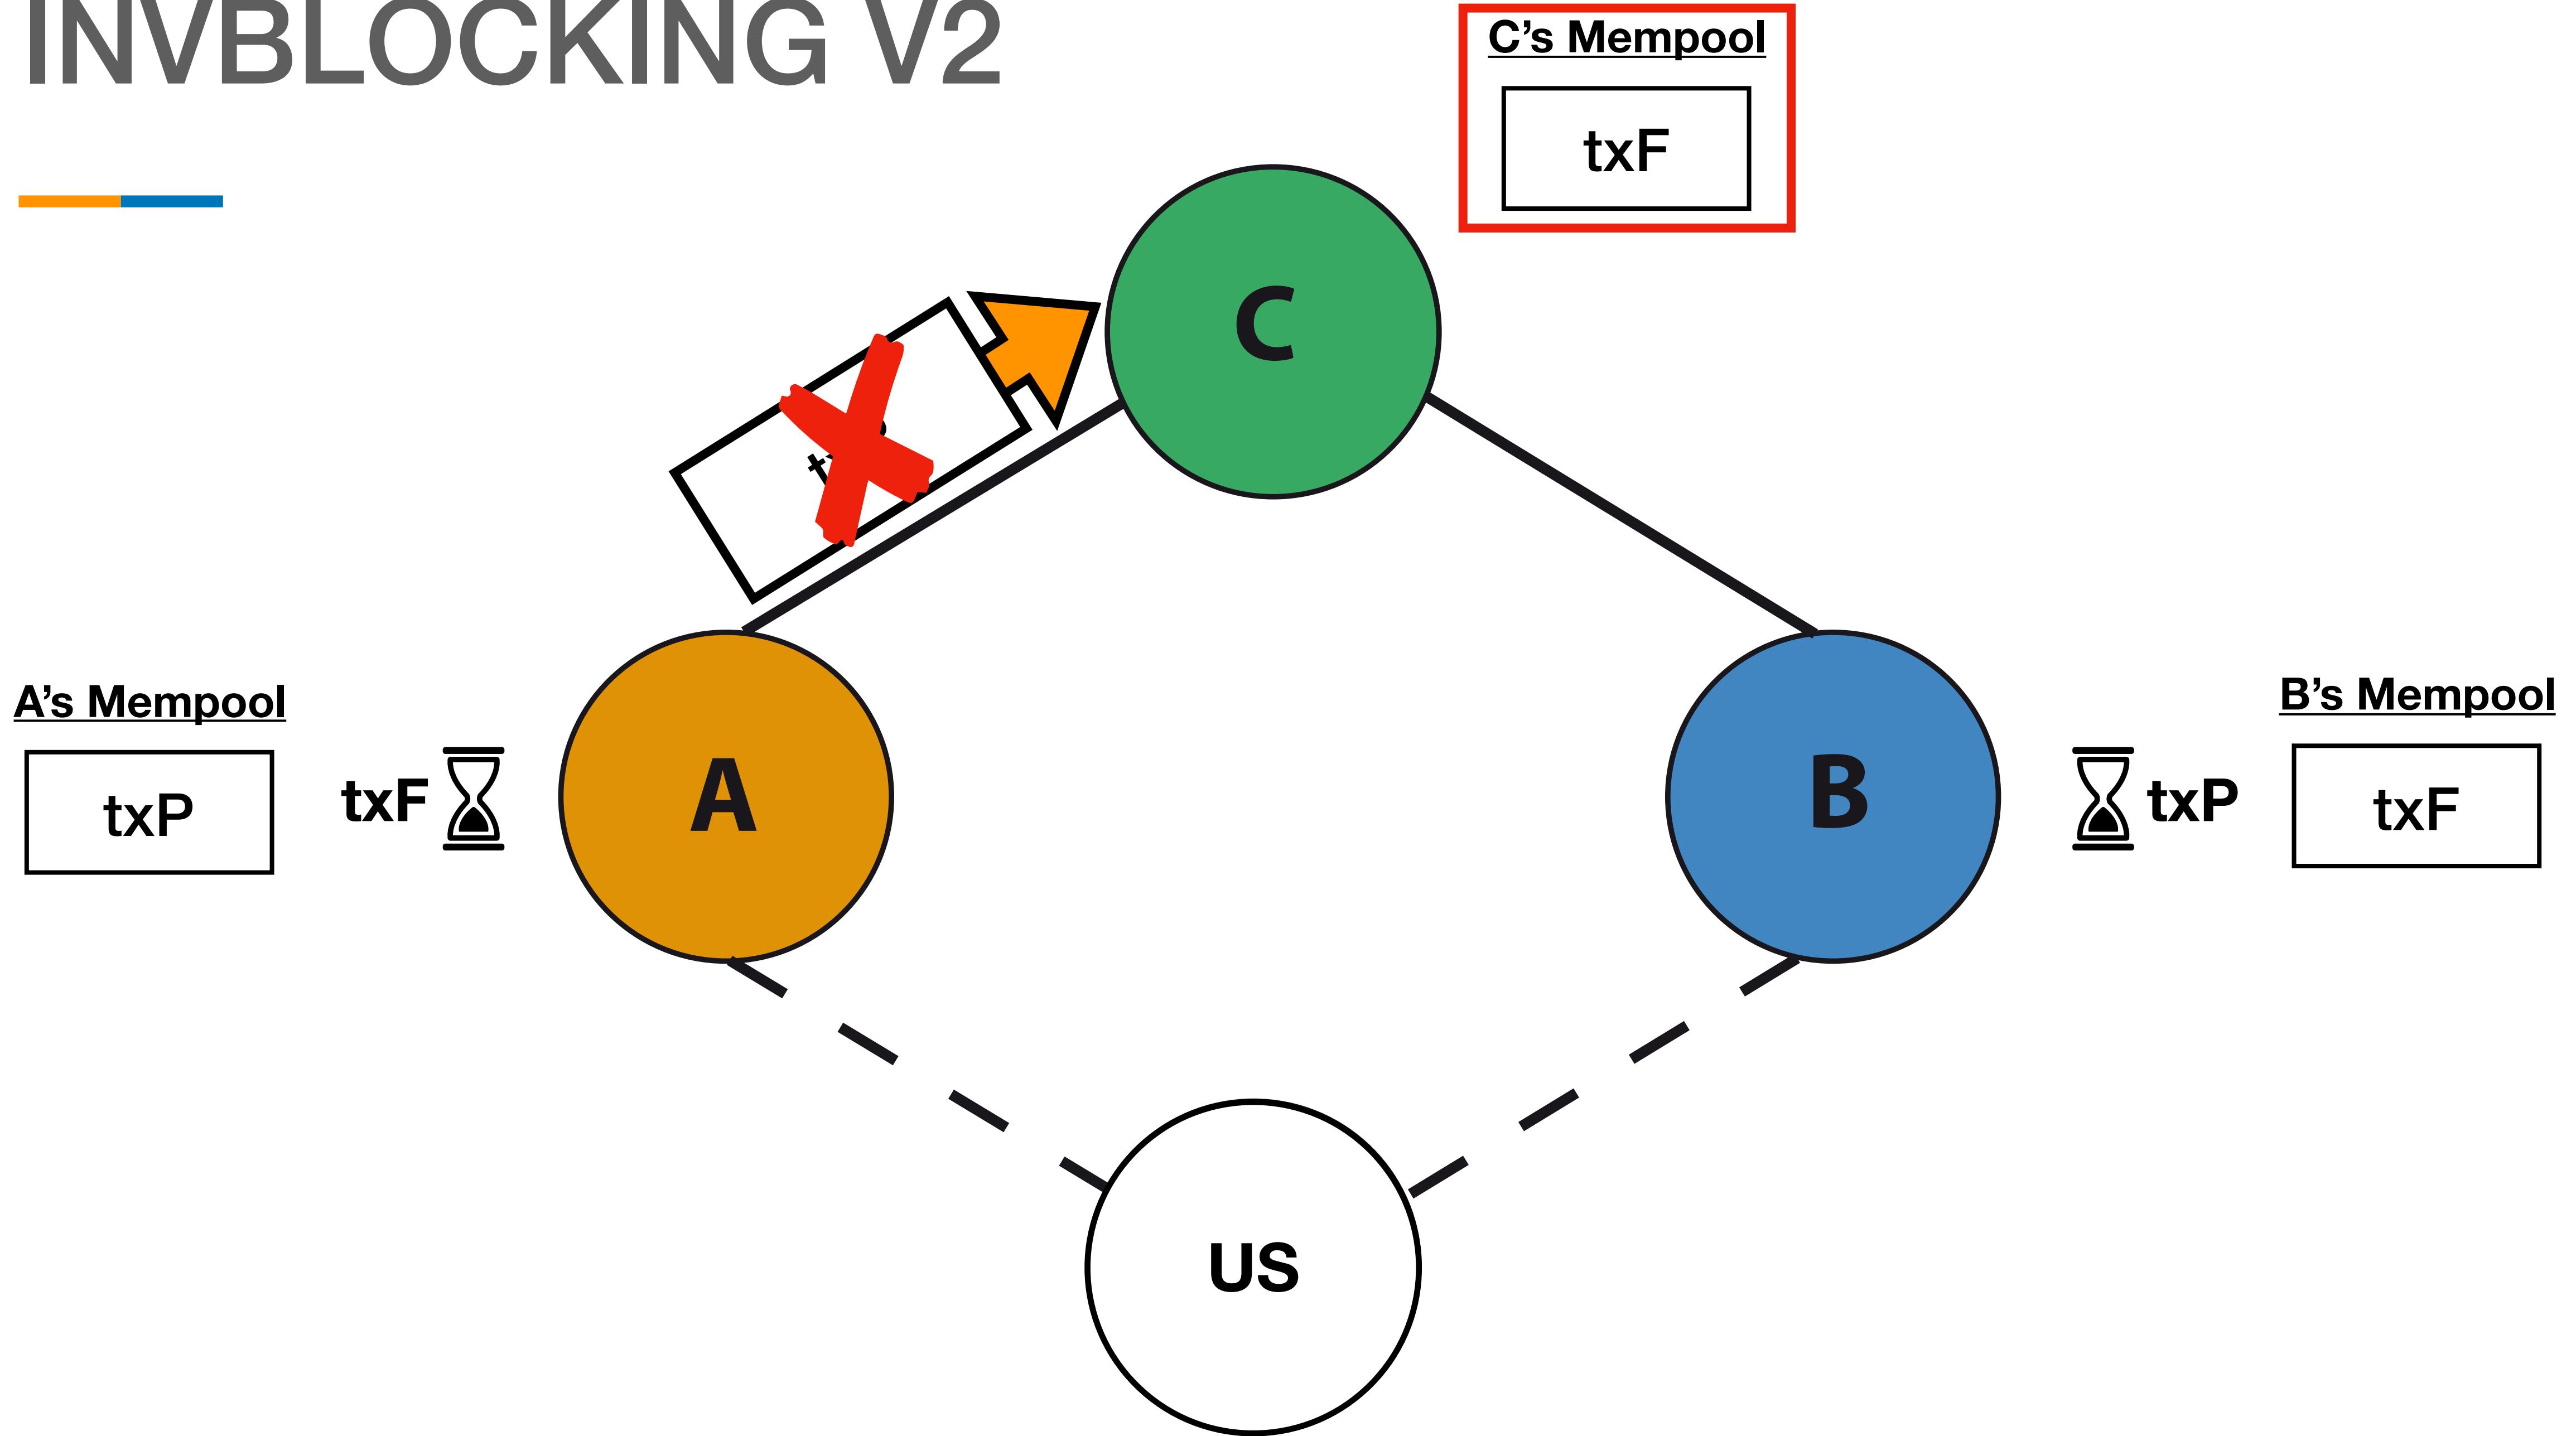
INVBLOCKING V2



INVBLOCKING V2



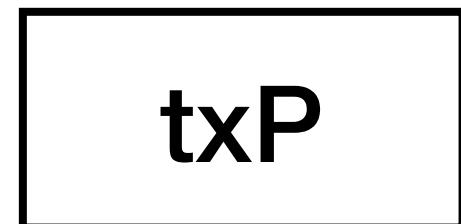
INVBLOCKING V2



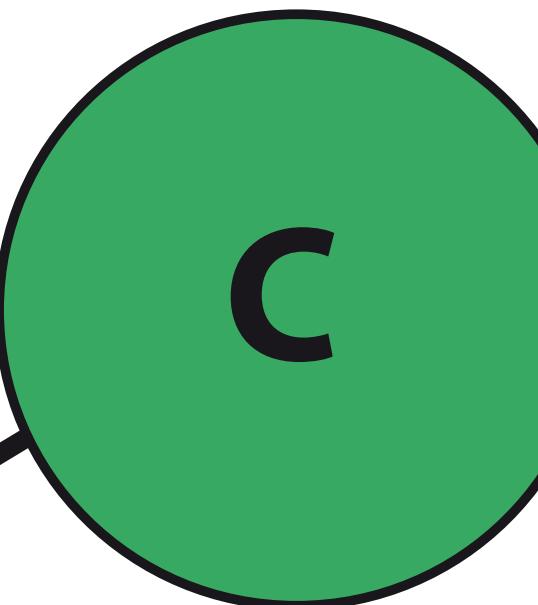
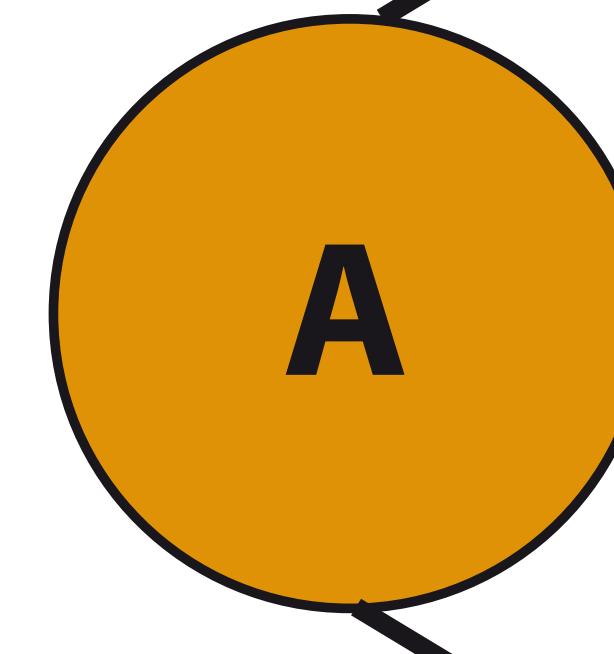
INVBLOCKING V2



A's Mempool



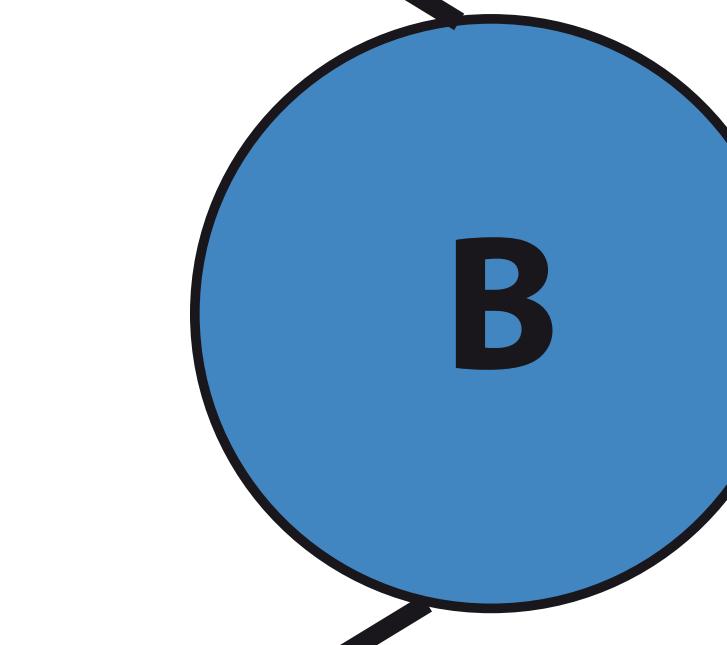
txF



C's Mempool



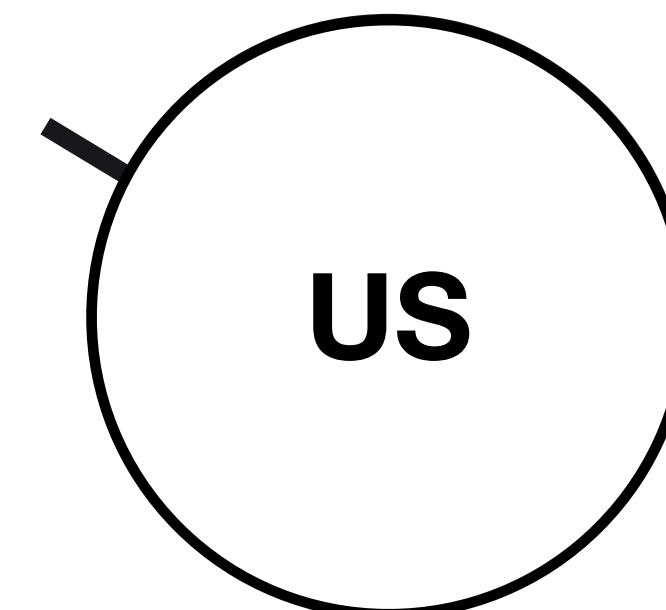
txF



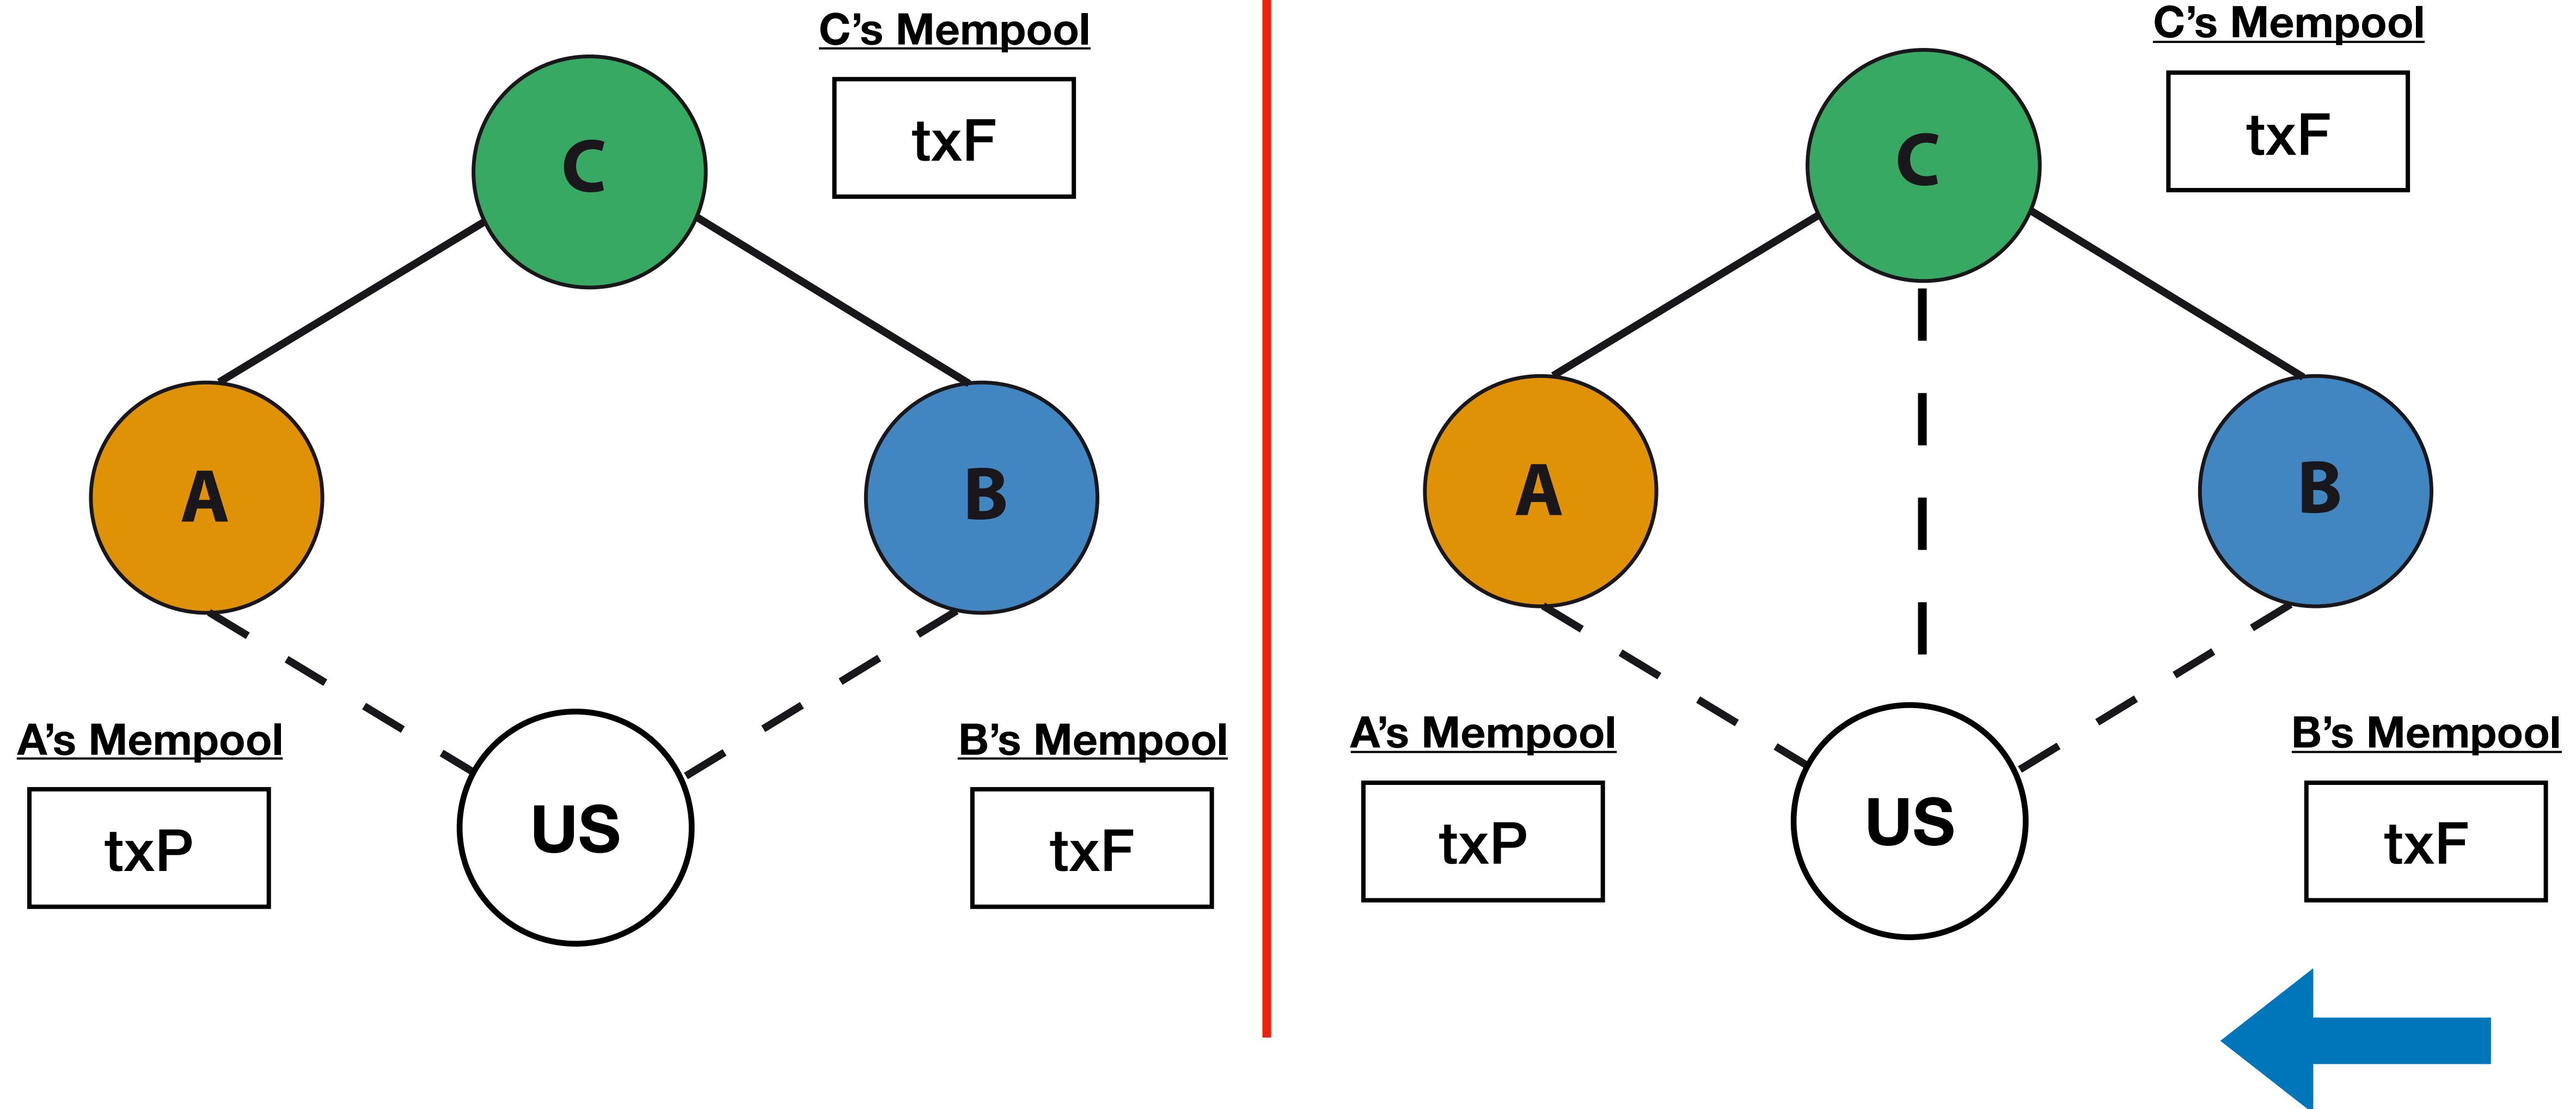
B's Mempool



txF

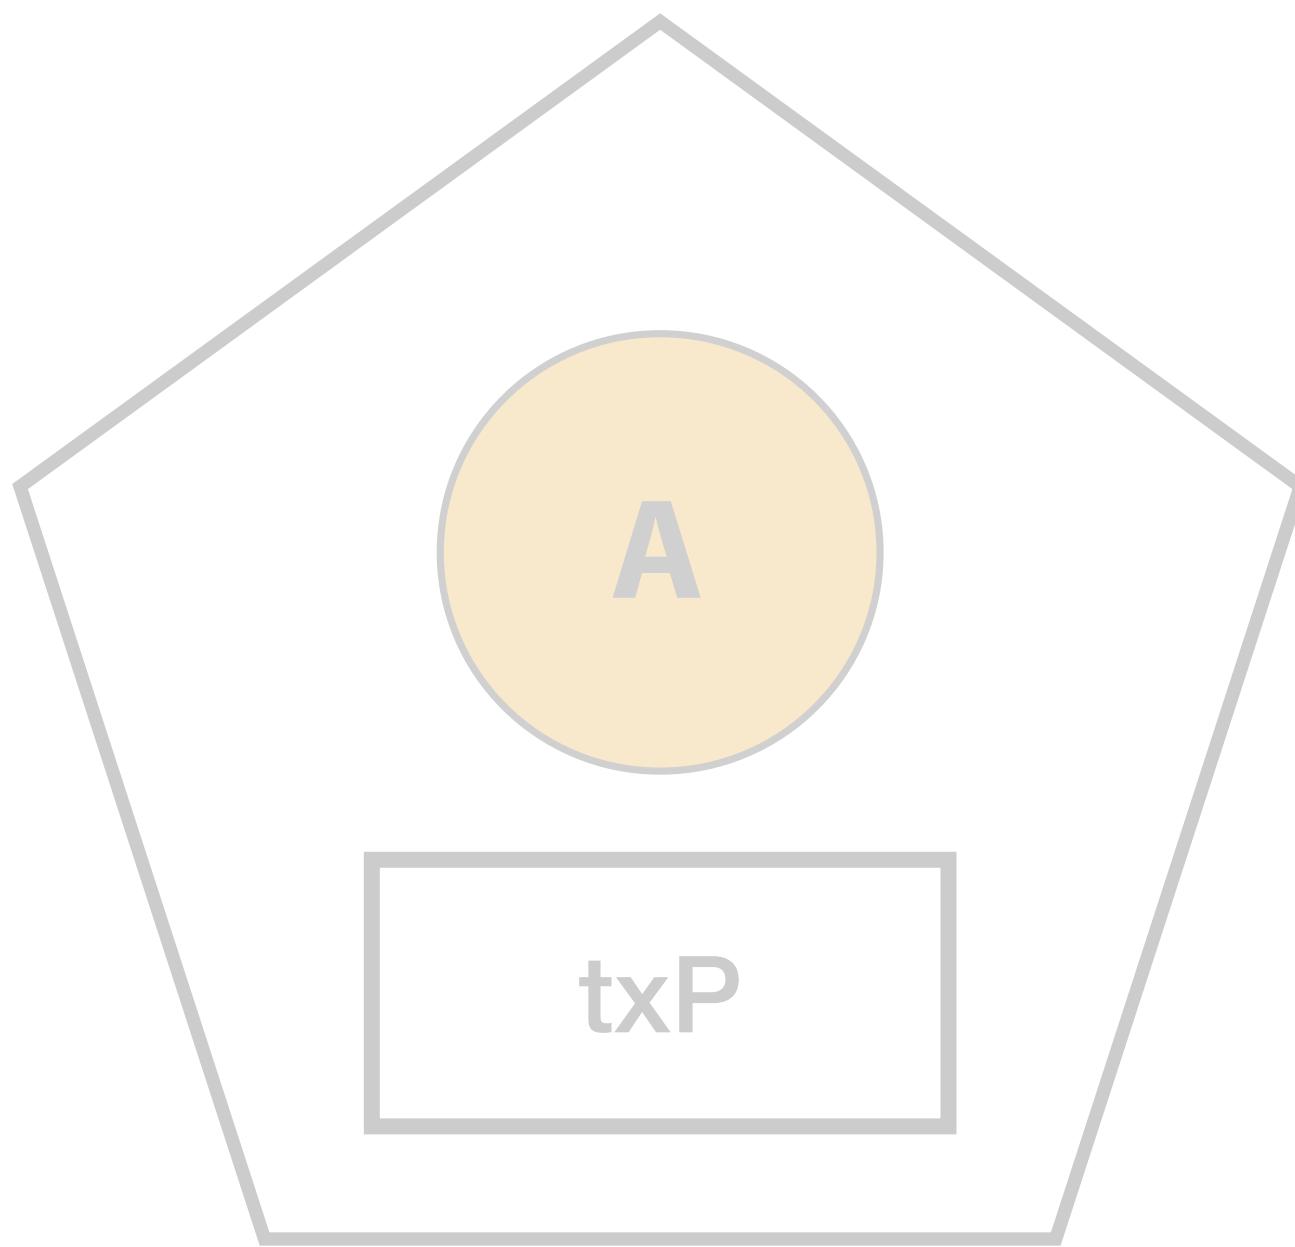


INVBLOCKING COMPARISON

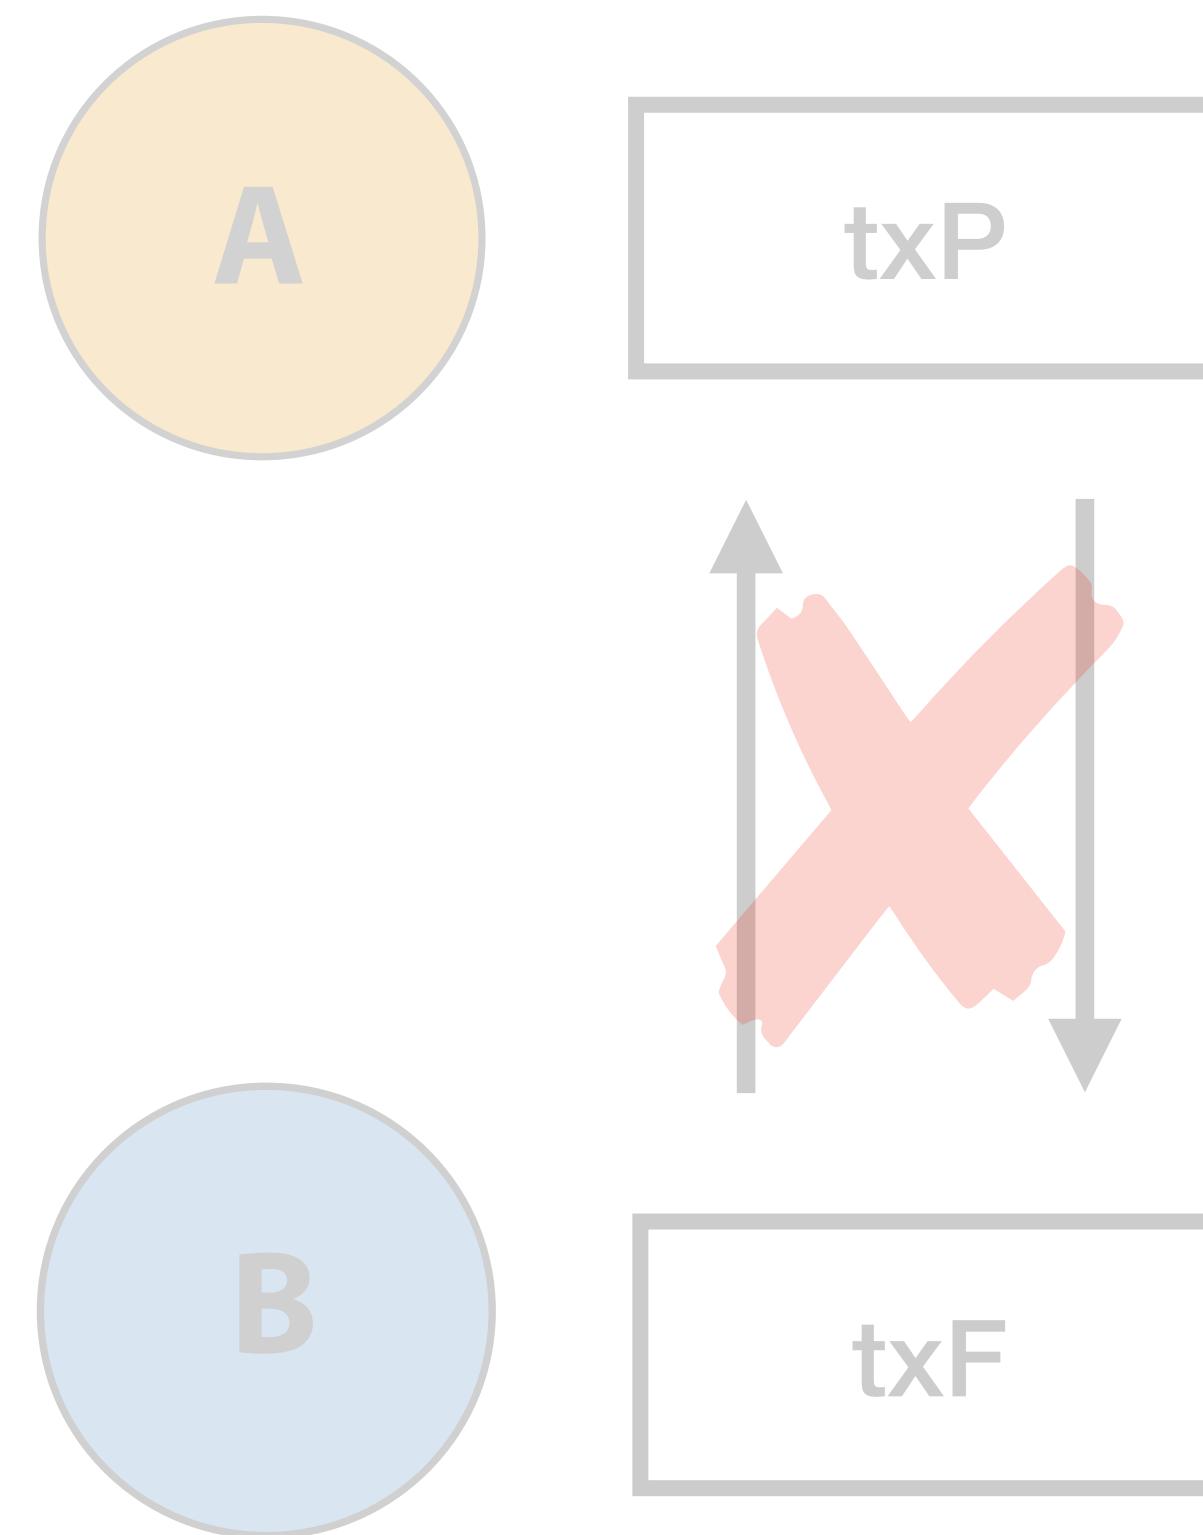


MAKE THIS WORK IN A REAL NETWORK

Isolation

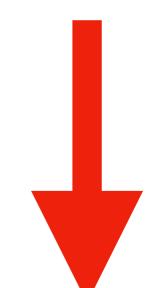


Synchrony



Efficiency

$\approx O(n)$



$\approx O(\sqrt{n})$

$n = \#nodes$

ORPHANPOOL EVICTION (BEFORE v0.18)

```
while (mapOrphanTransactions.size() > nMaxOrphans)

{
    // Evict a random orphan:

    uint256 randomhash = rng.rand256();

    std::map<uint256, COrphanTx>::iterator it = mapOrphanTransactions.lower_bound(randomhash);

    if (it == mapOrphanTransactions.end())
        it = mapOrphanTransactions.begin();

    EraseOrphanTx(it->first);

    ++nEvicted;
}
```

source: https://github.com/bitcoin/bitcoin/blob/273d025/src/net_processing.cpp#L783-L791

ORPHANPOOL EVICTION (BEFORE v0.18)

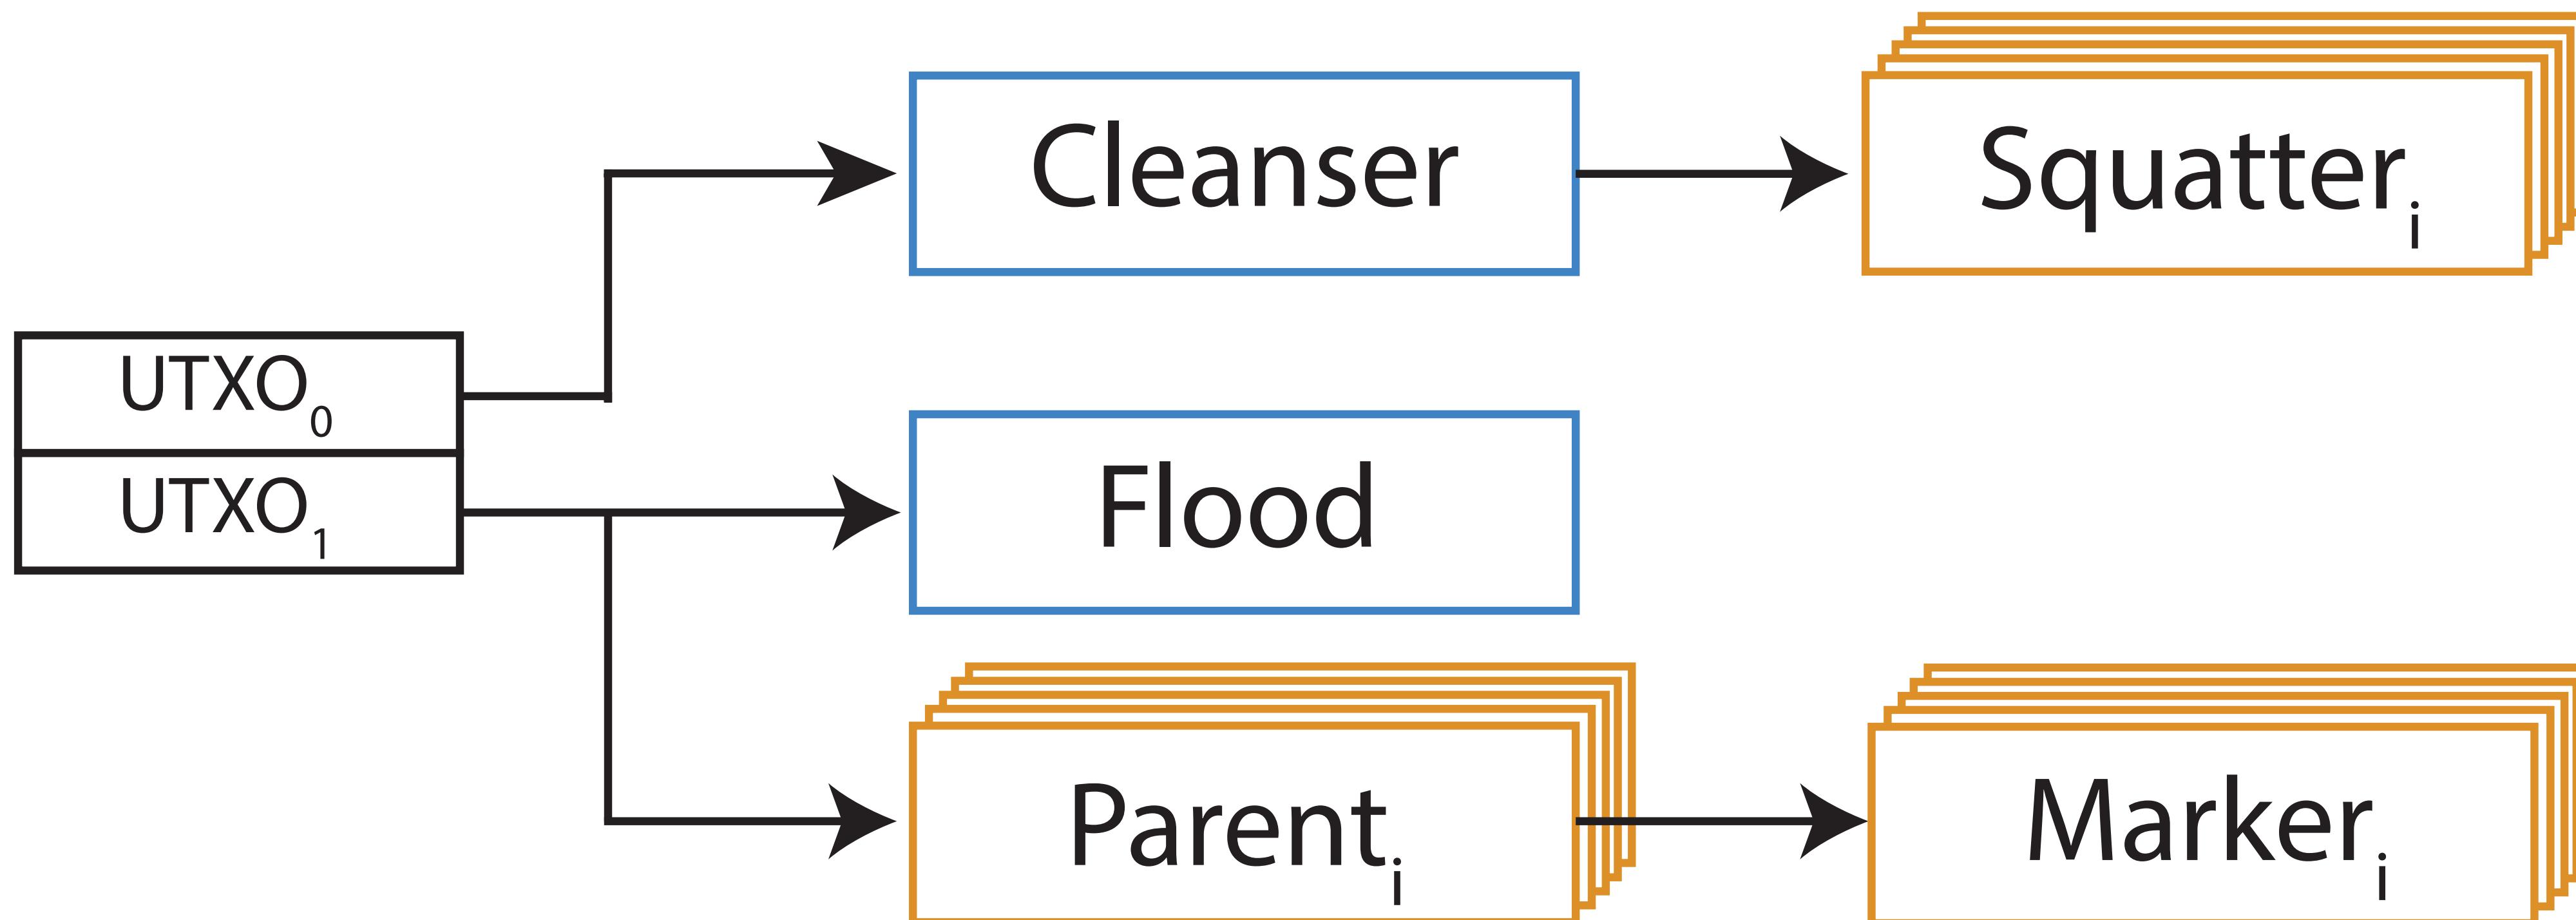
```
while (mapOrphanTransactions.size() > nMaxOrphans)
```

```
{
```

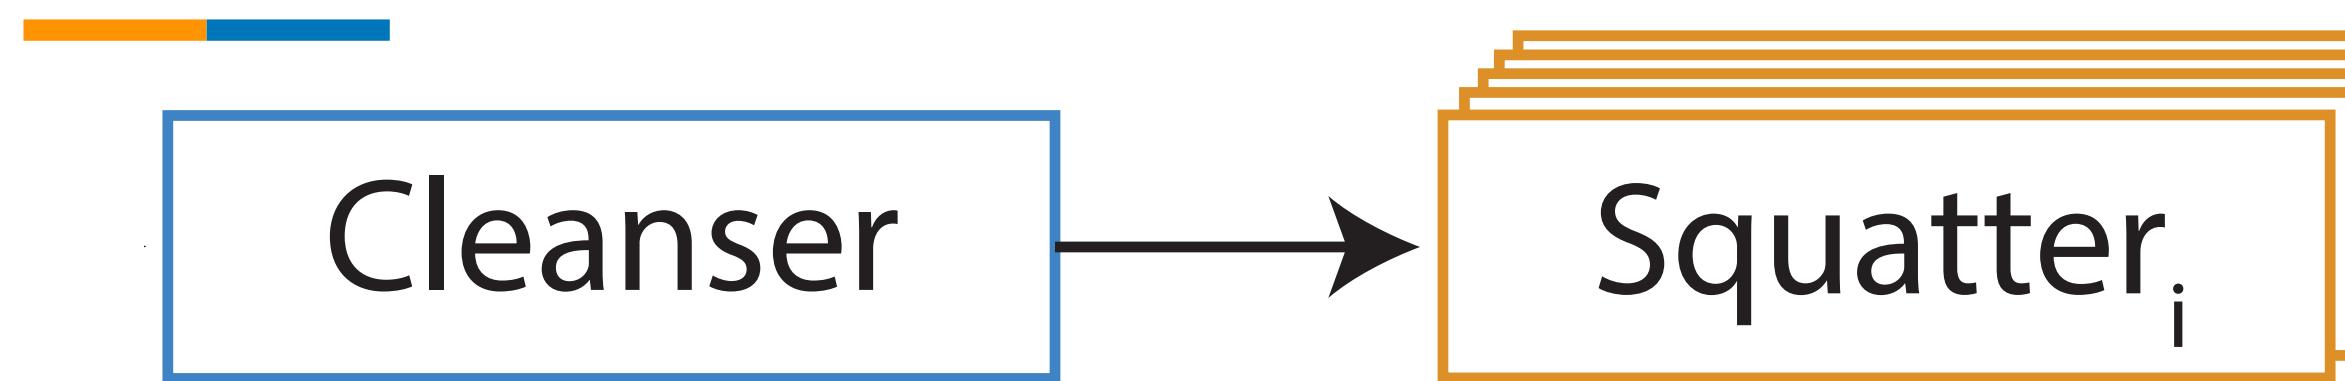
- Pick a random 256-bit value **R**
- Get the orphan transaction (**O**) with hash **closer to, but greater than, R**
- **Evict O**
- Repeat until mapOrphanTransaction **is not full (default: 100)**
- **Double-spends are not checked for orphans**

```
}
```

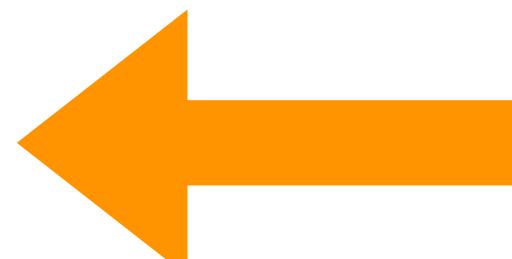
TXPROBE TRANSACTIONS OVERVIEW



MAKE ROOM IN THE ORPHANPOOL



- Create the cleanser (regular transaction) and **100 squatters (double-spends between each other)**
- Every squatter is created in a POW-ish way (e.g. re-sign until its hash falls below a certain threshold)
- All squatters are sent to the **flood set nodes** to replace any existing orphan.
- Finally, the cleanser is sent to empty the orphanpool



QUESTIONS

